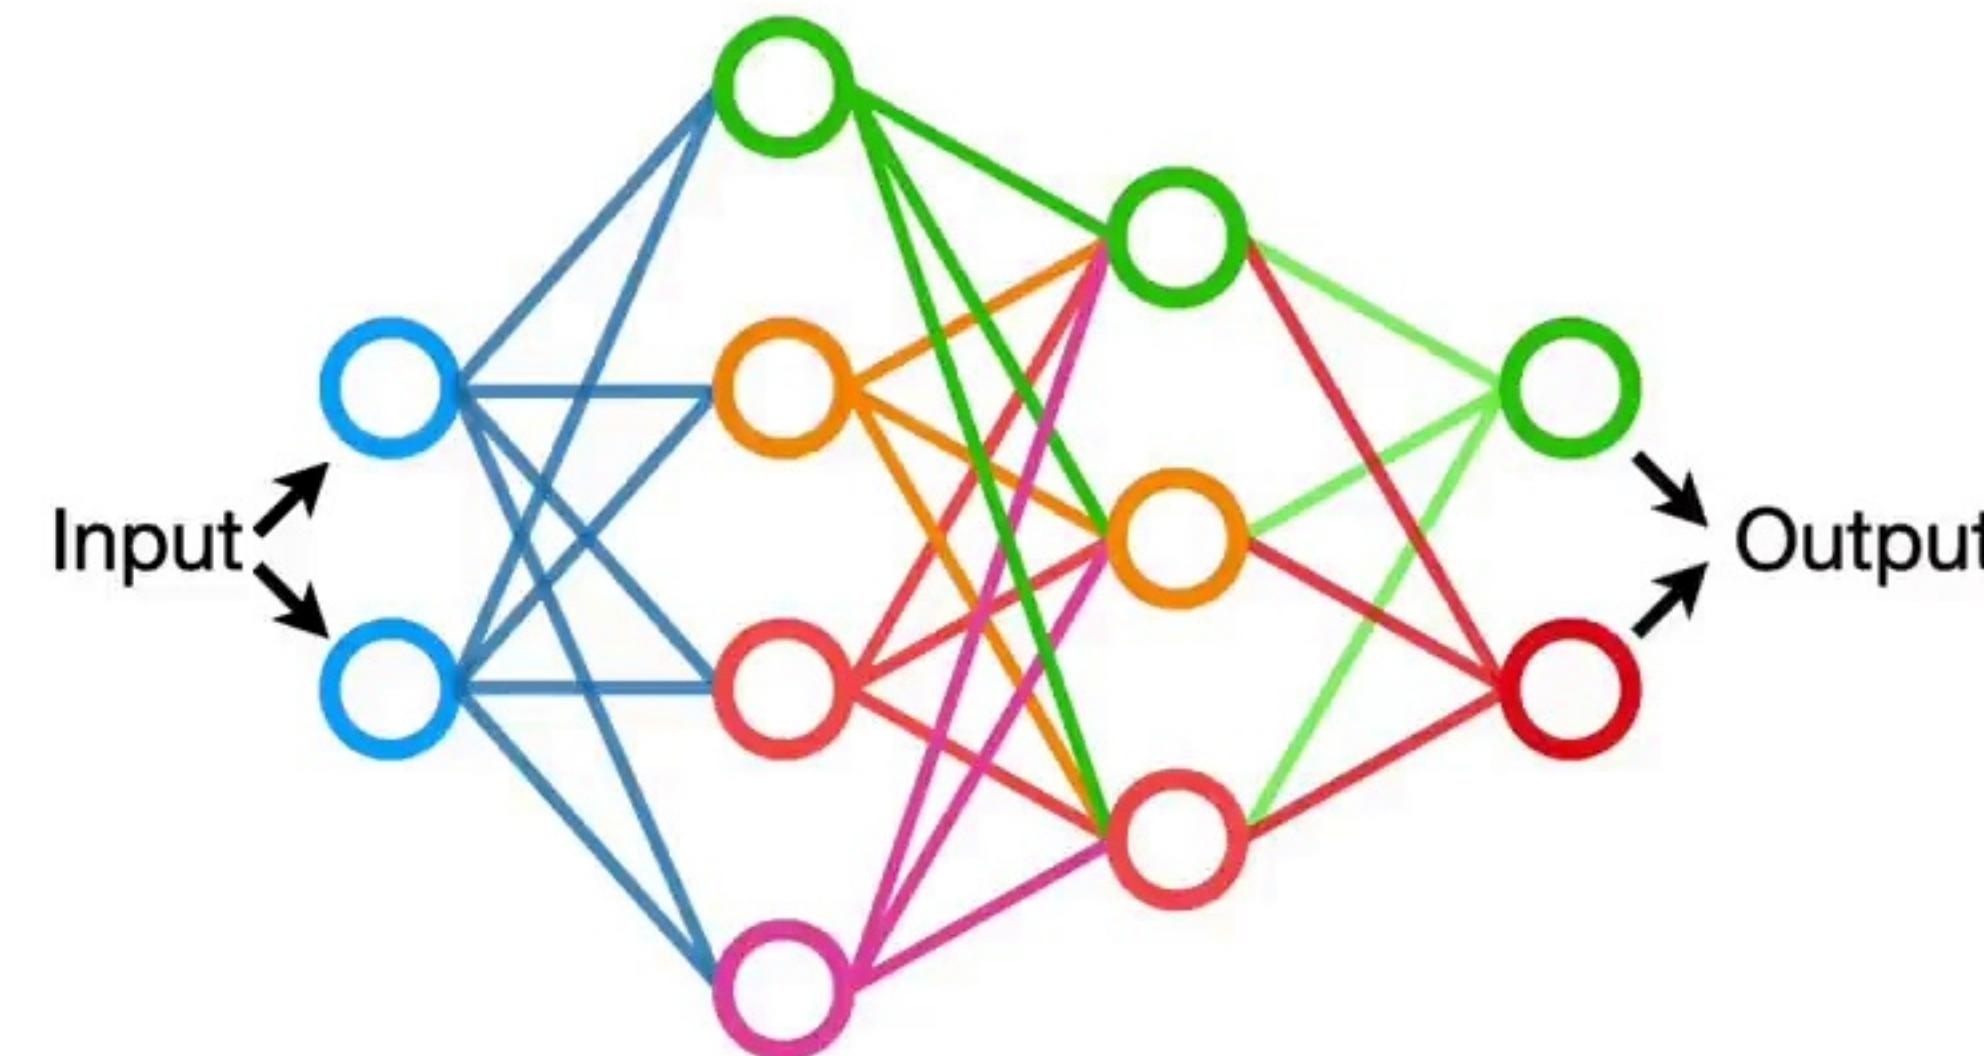




**Neural Networks**, one of the most popular algorithms in **Machine Learning**, cover a broad range of concepts and techniques.



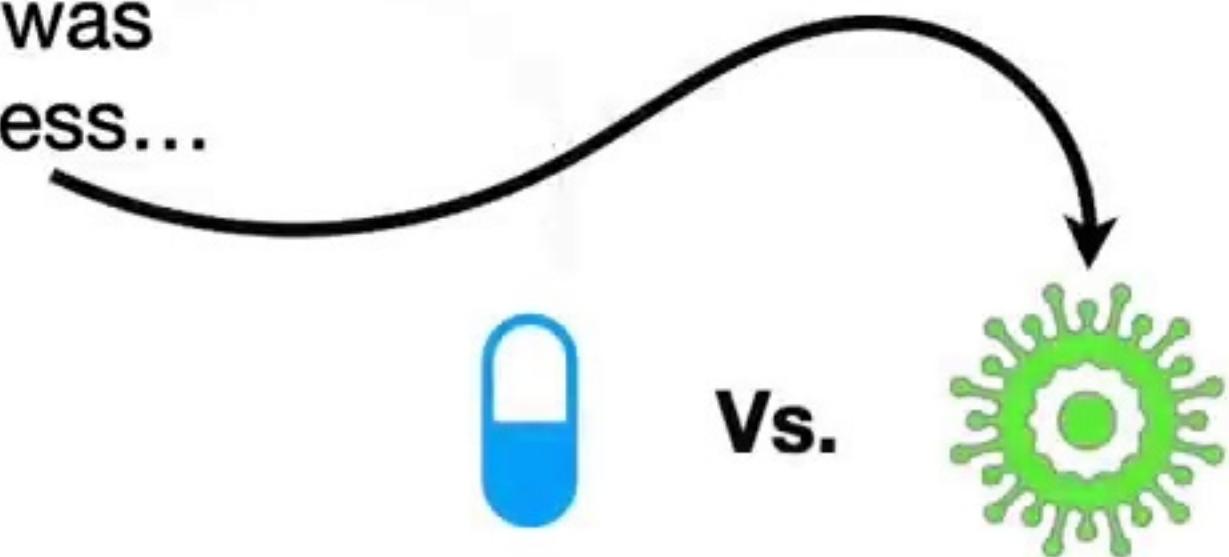


So, with that said, let's imagine  
we tested a drug that was  
designed to treat an illness...



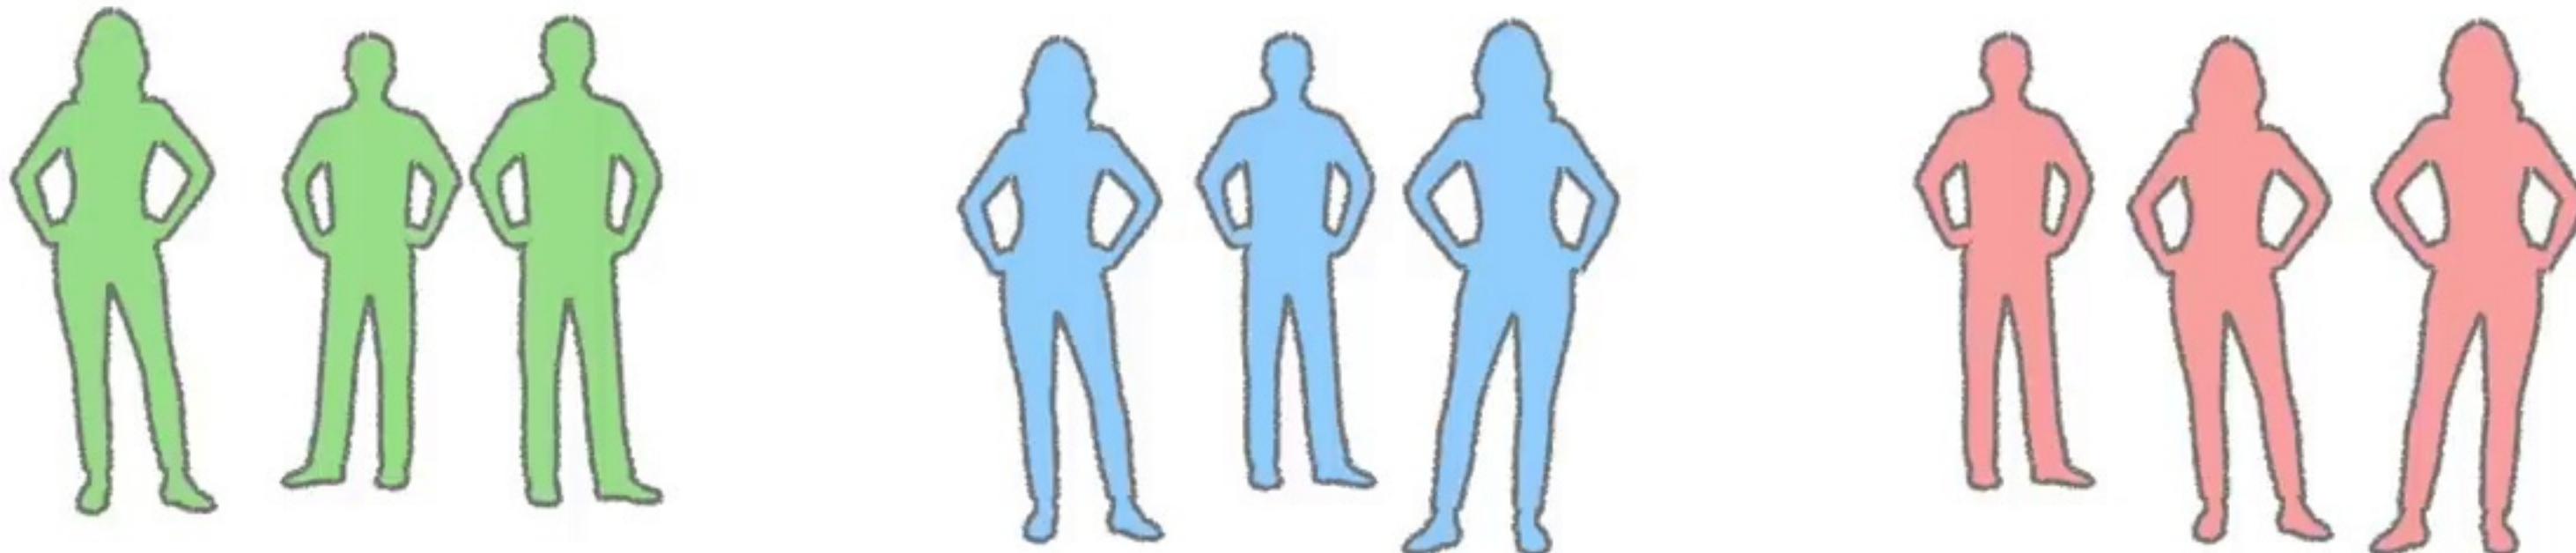
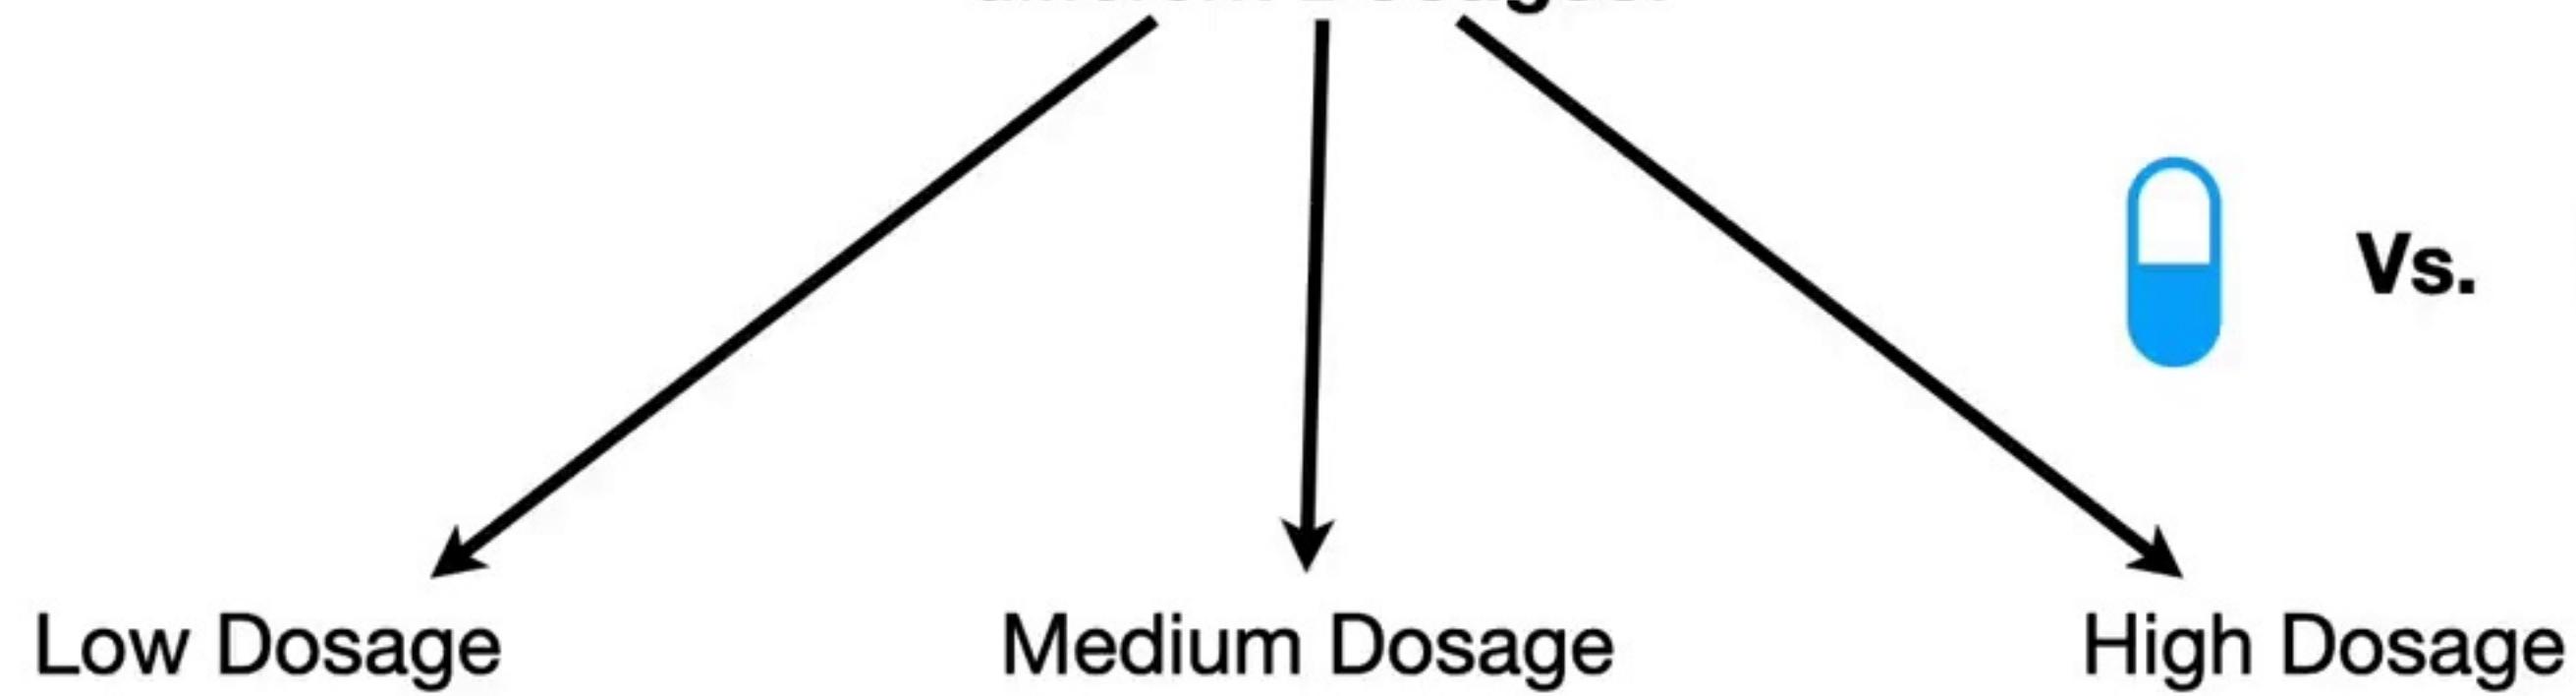


So, with that said, let's imagine  
we tested a drug that was  
designed to treat an illness...



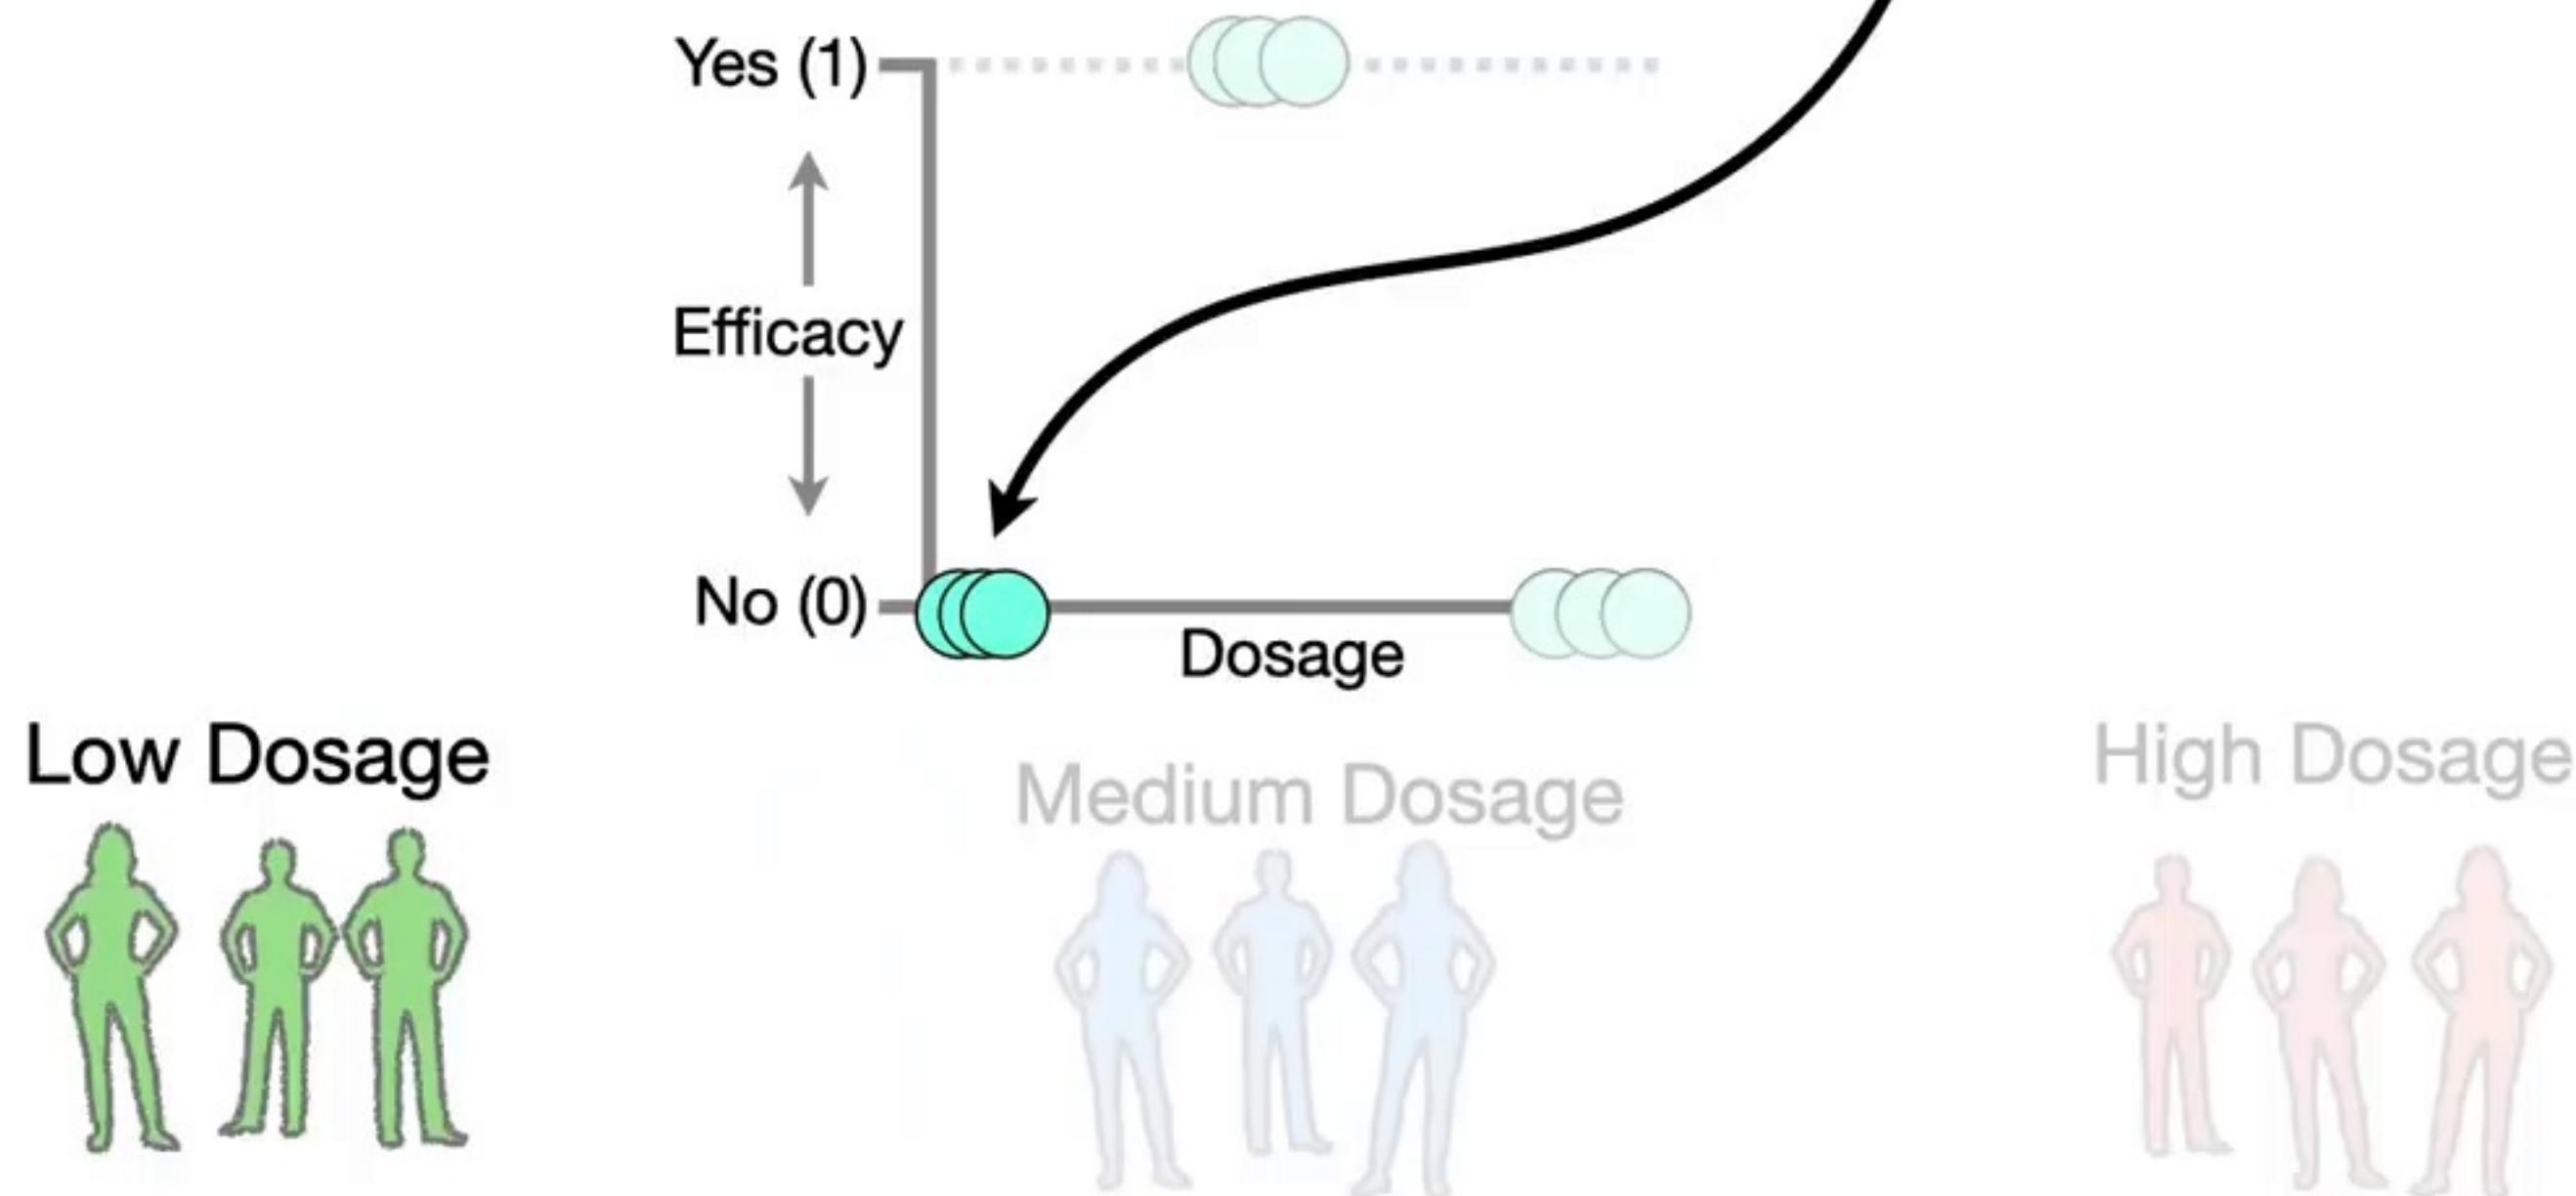


...and we gave the drug to **3**  
different groups of people with **3**  
different **Dosages**.



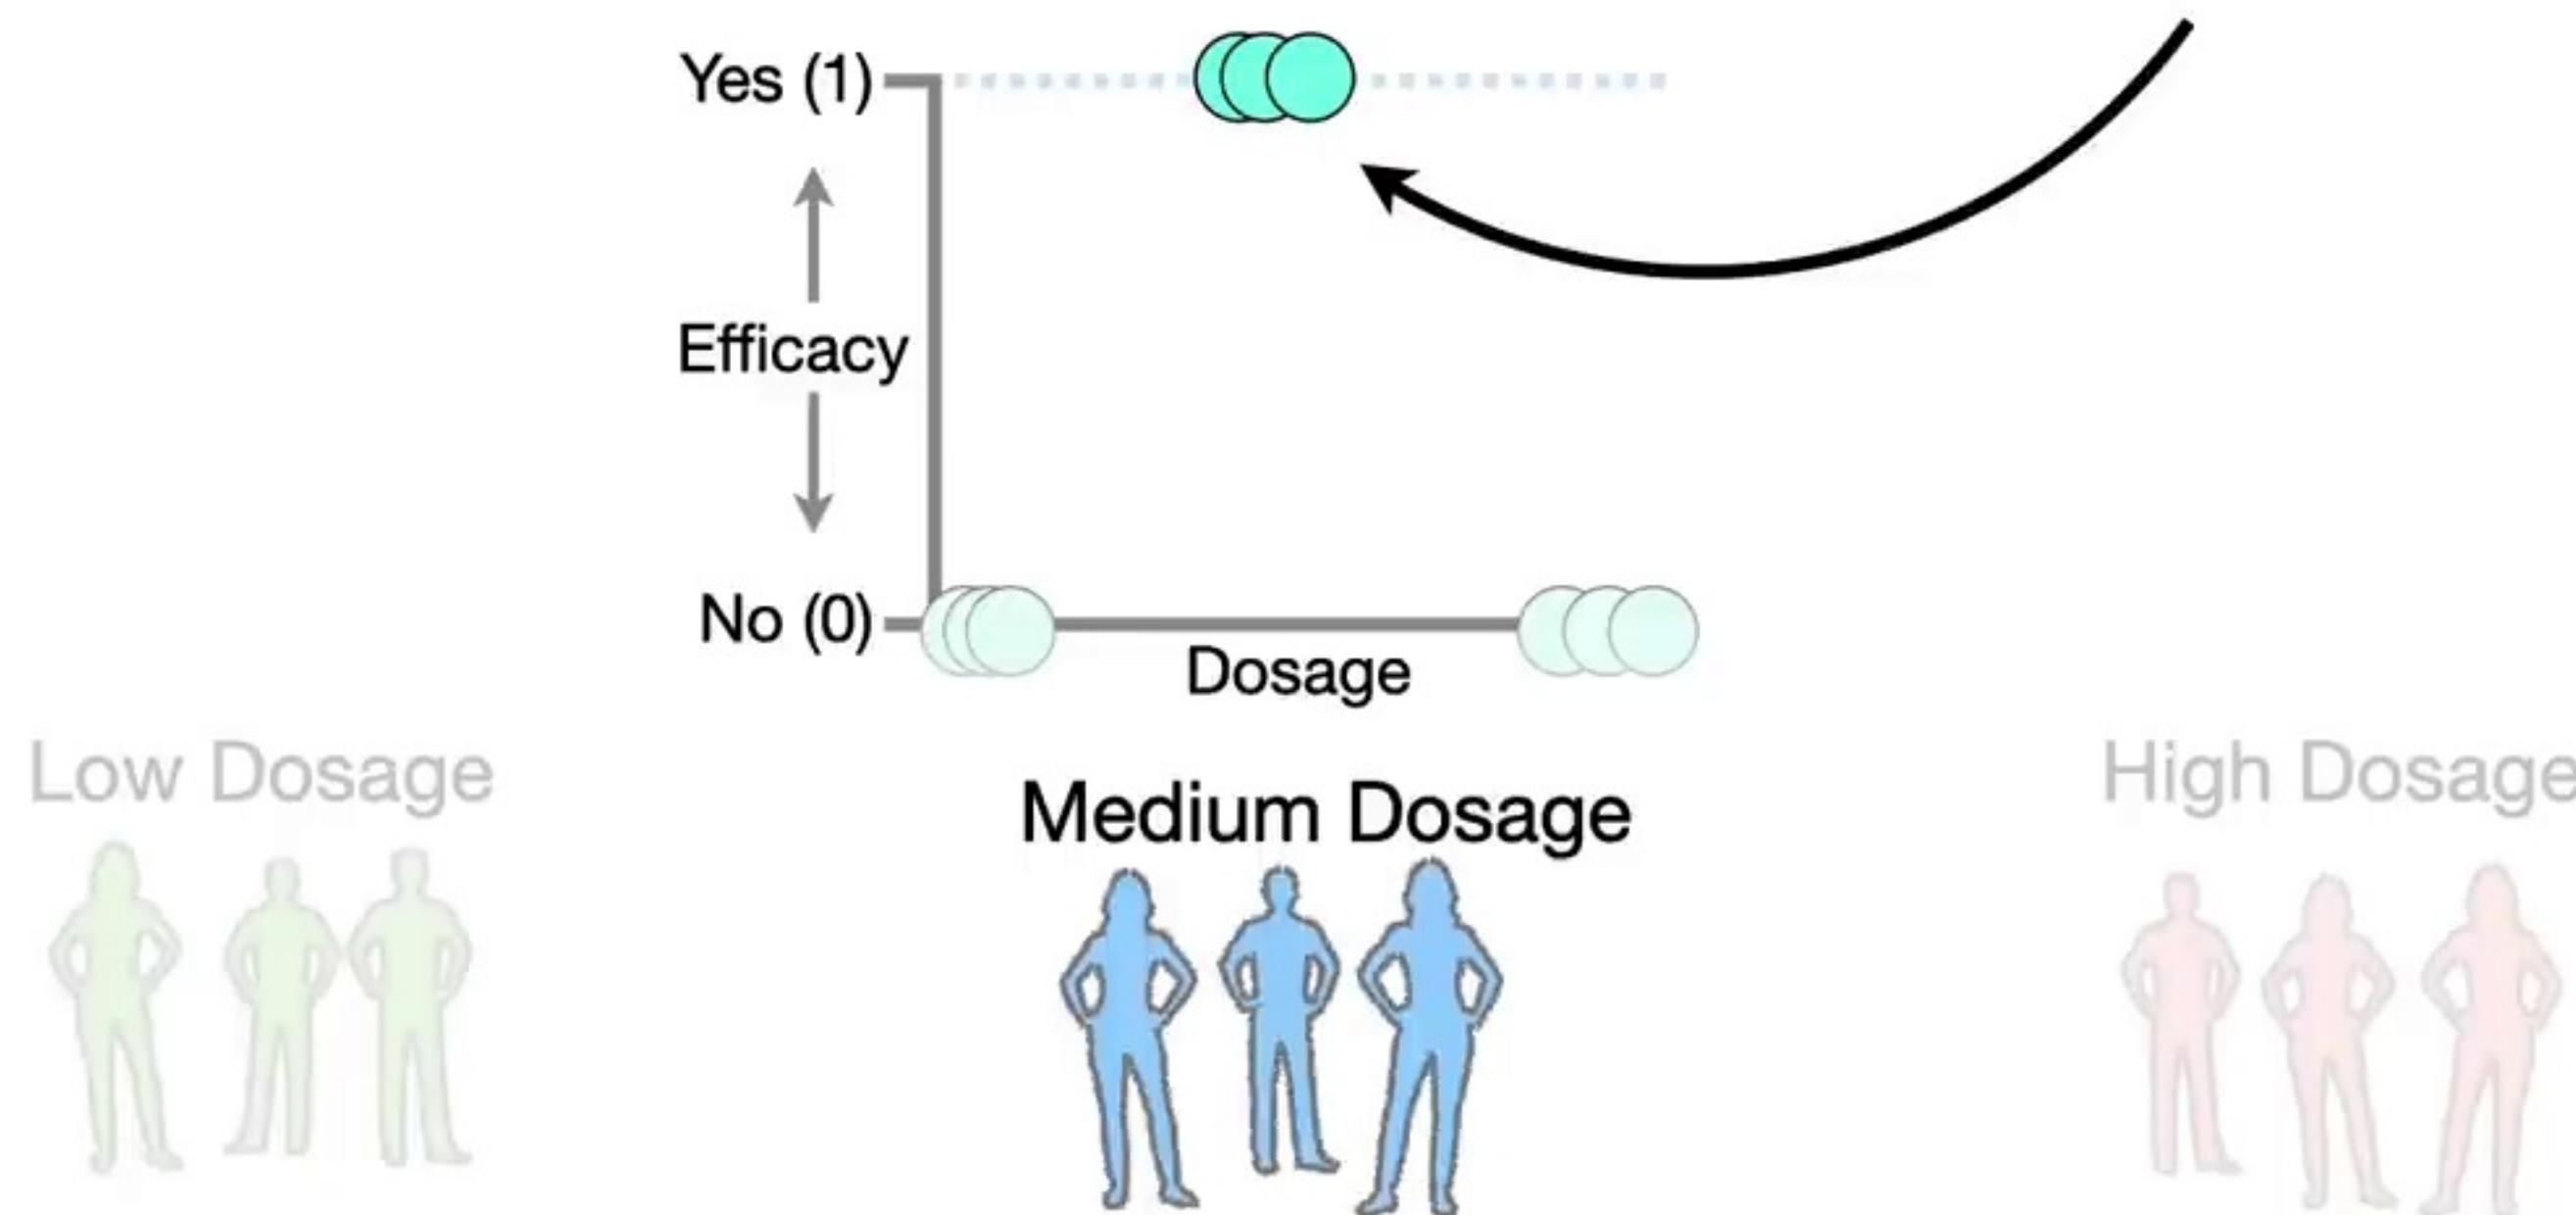


The low **Dosages** were ***not Effective***, so we set them to **0** on this graph.



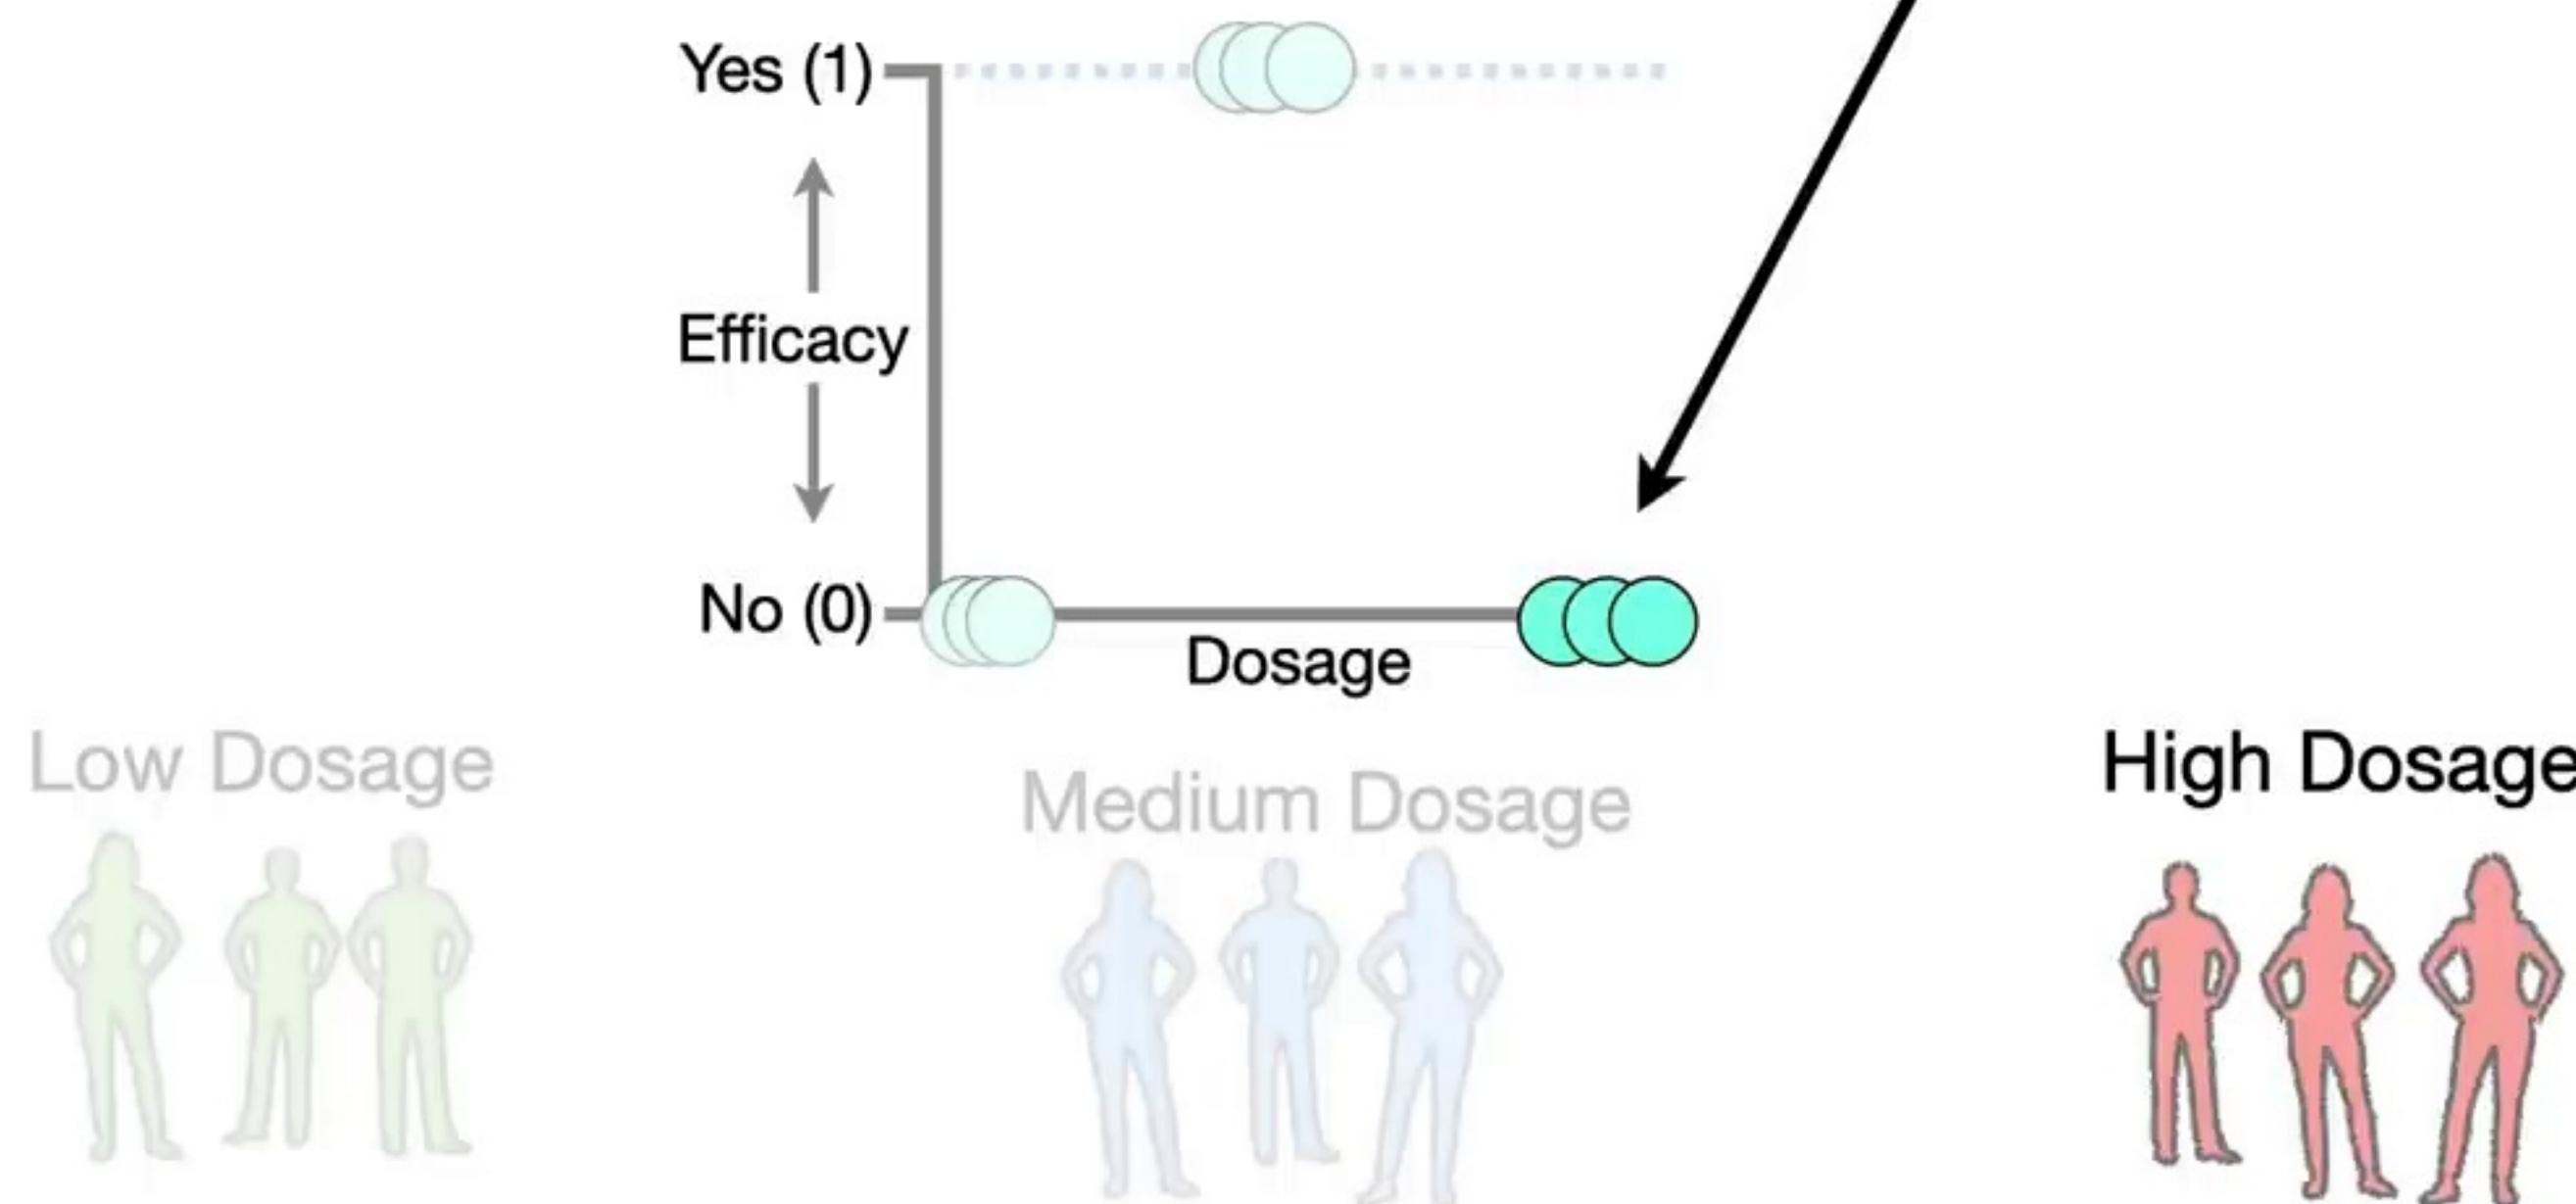


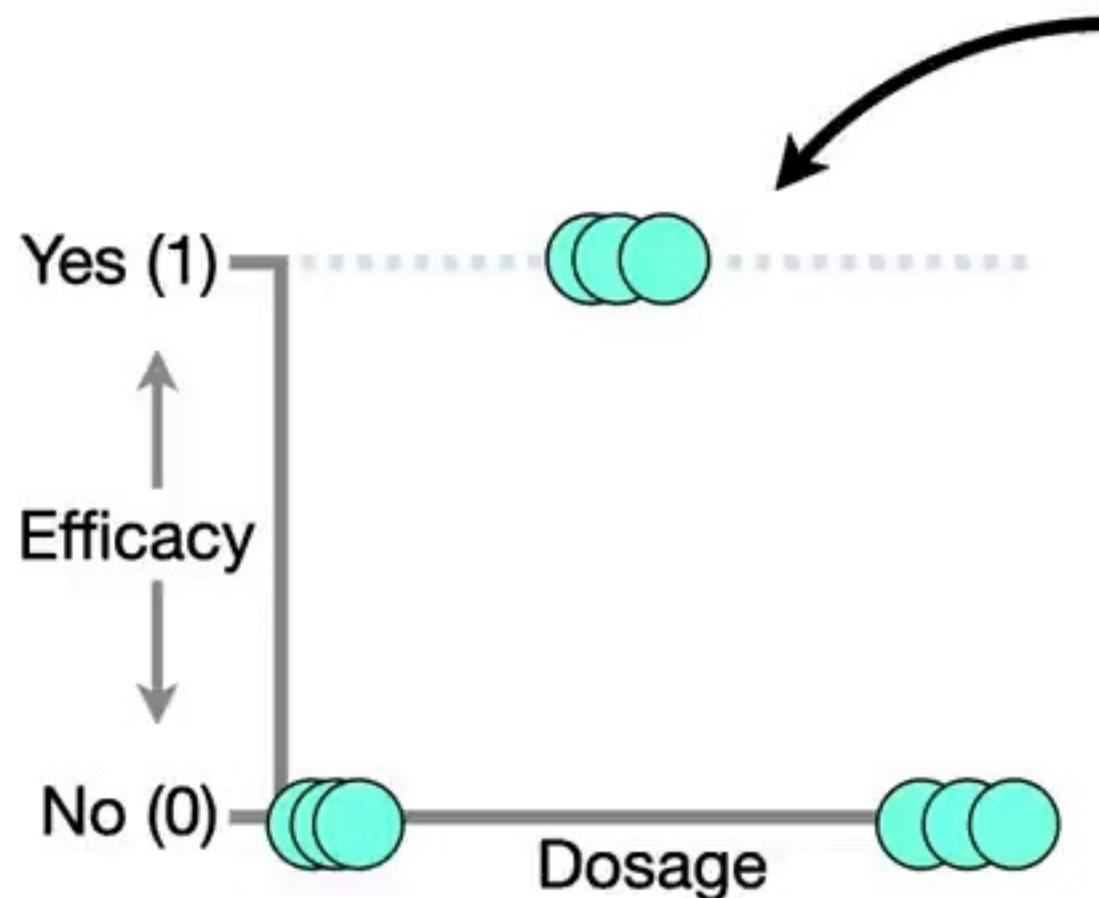
In contrast, the medium **Dosages** were **Effective**, so we set them to 1.





And the high  
**Dosages** were **not**  
**Effective**, so those  
are set to **0**.

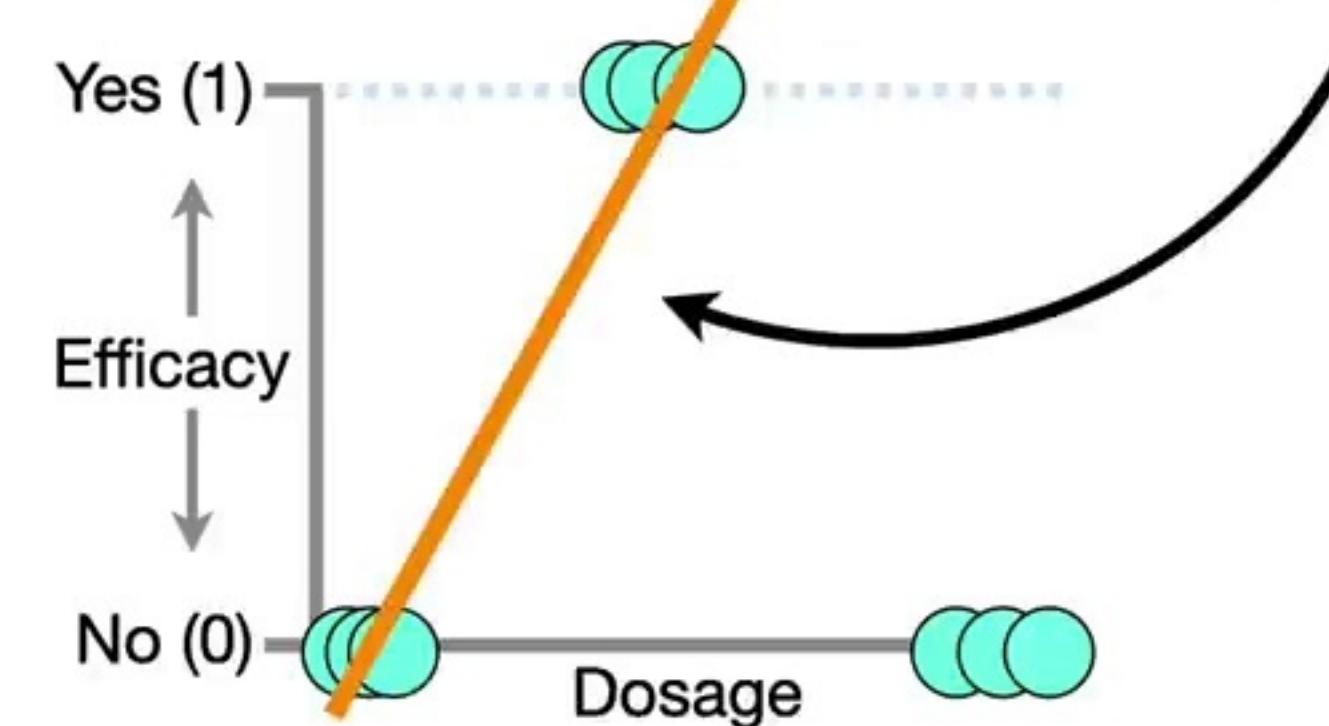


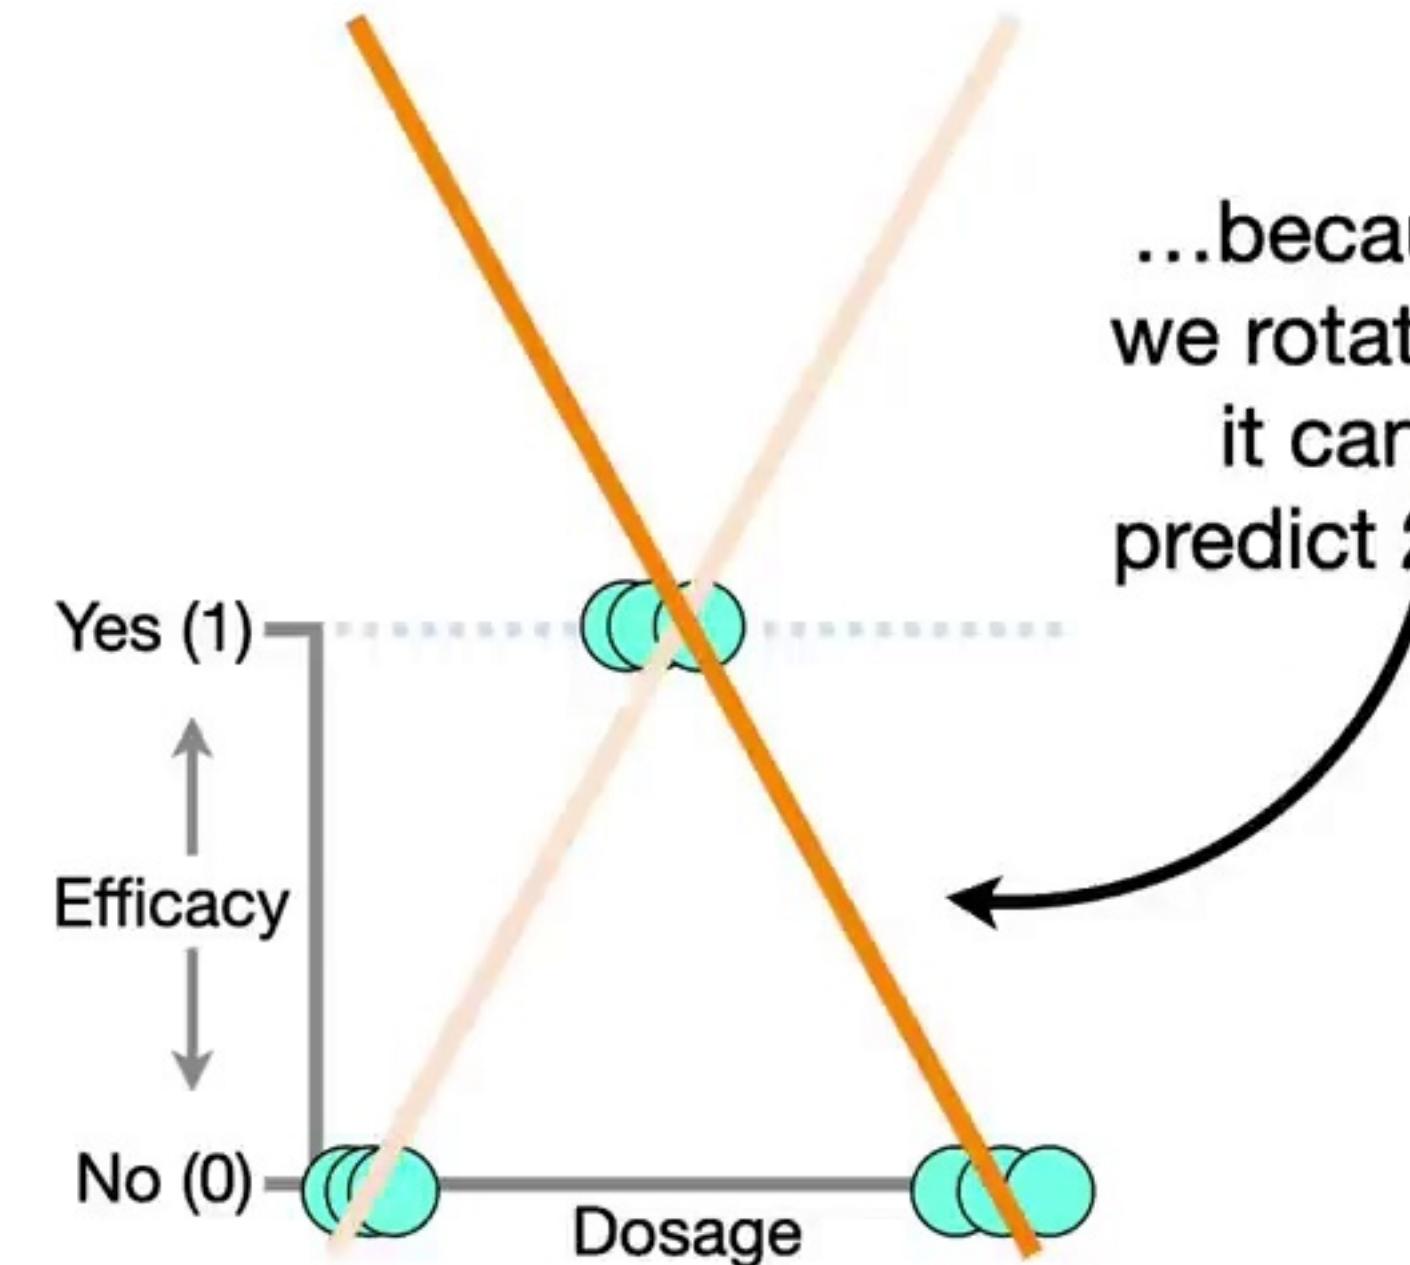


Now that we have this data, we would like to use it to predict whether or not a future **Dosage** will be Effective.

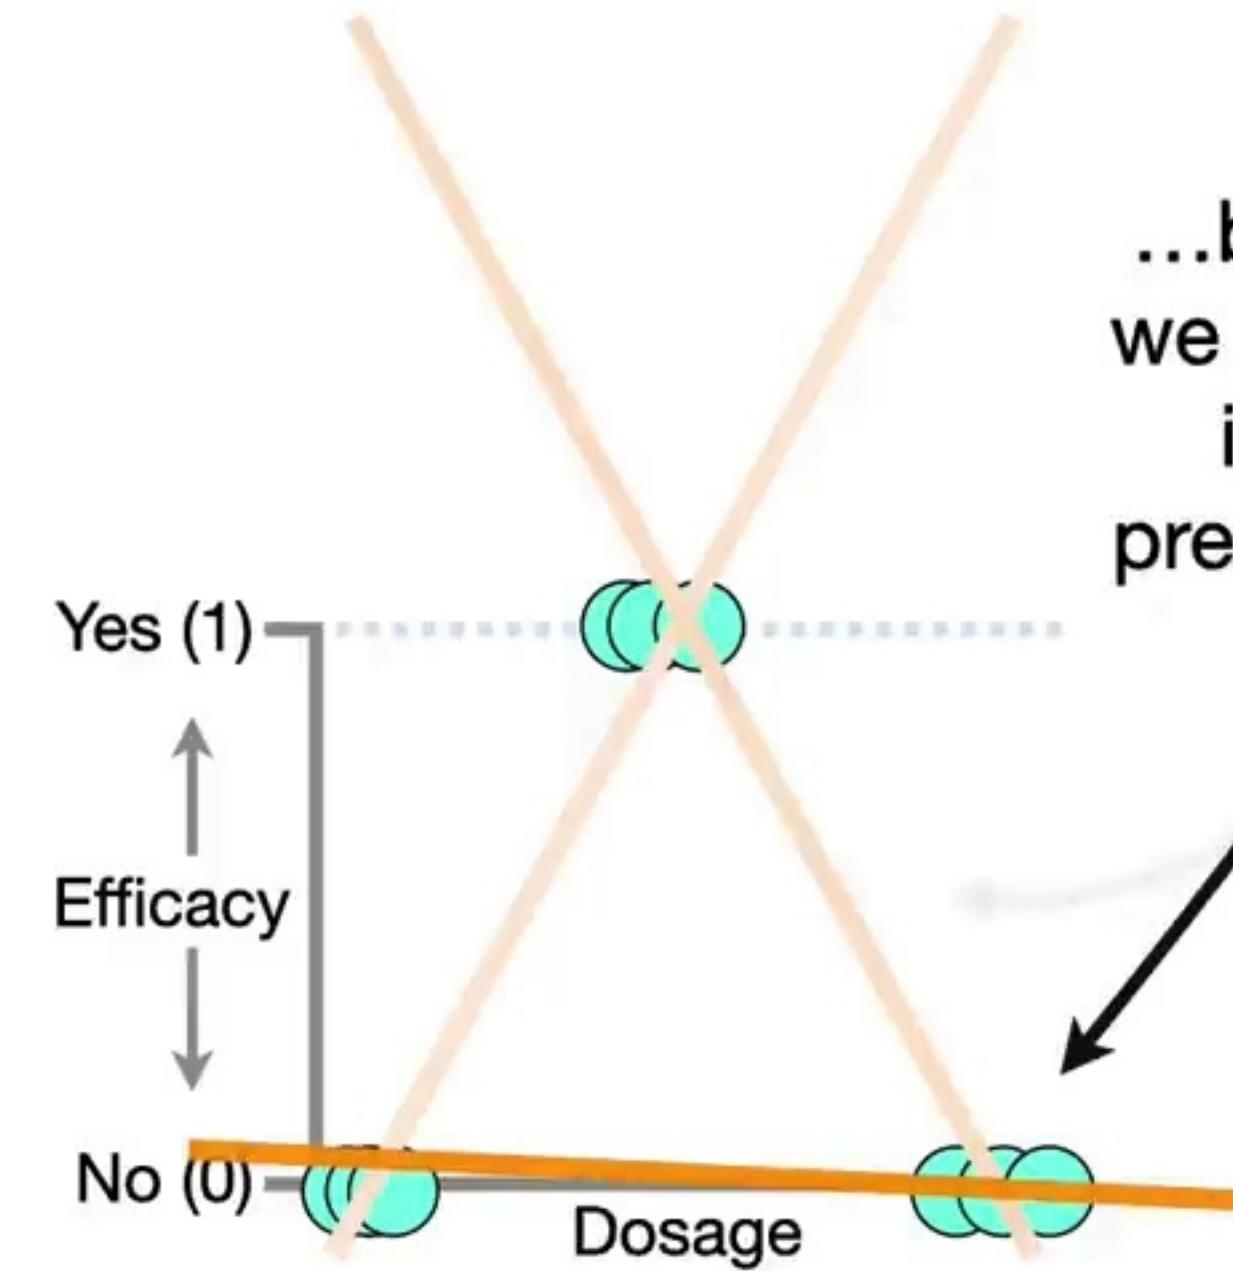


However, we can't just fit a **straight line** to the data to make predictions...

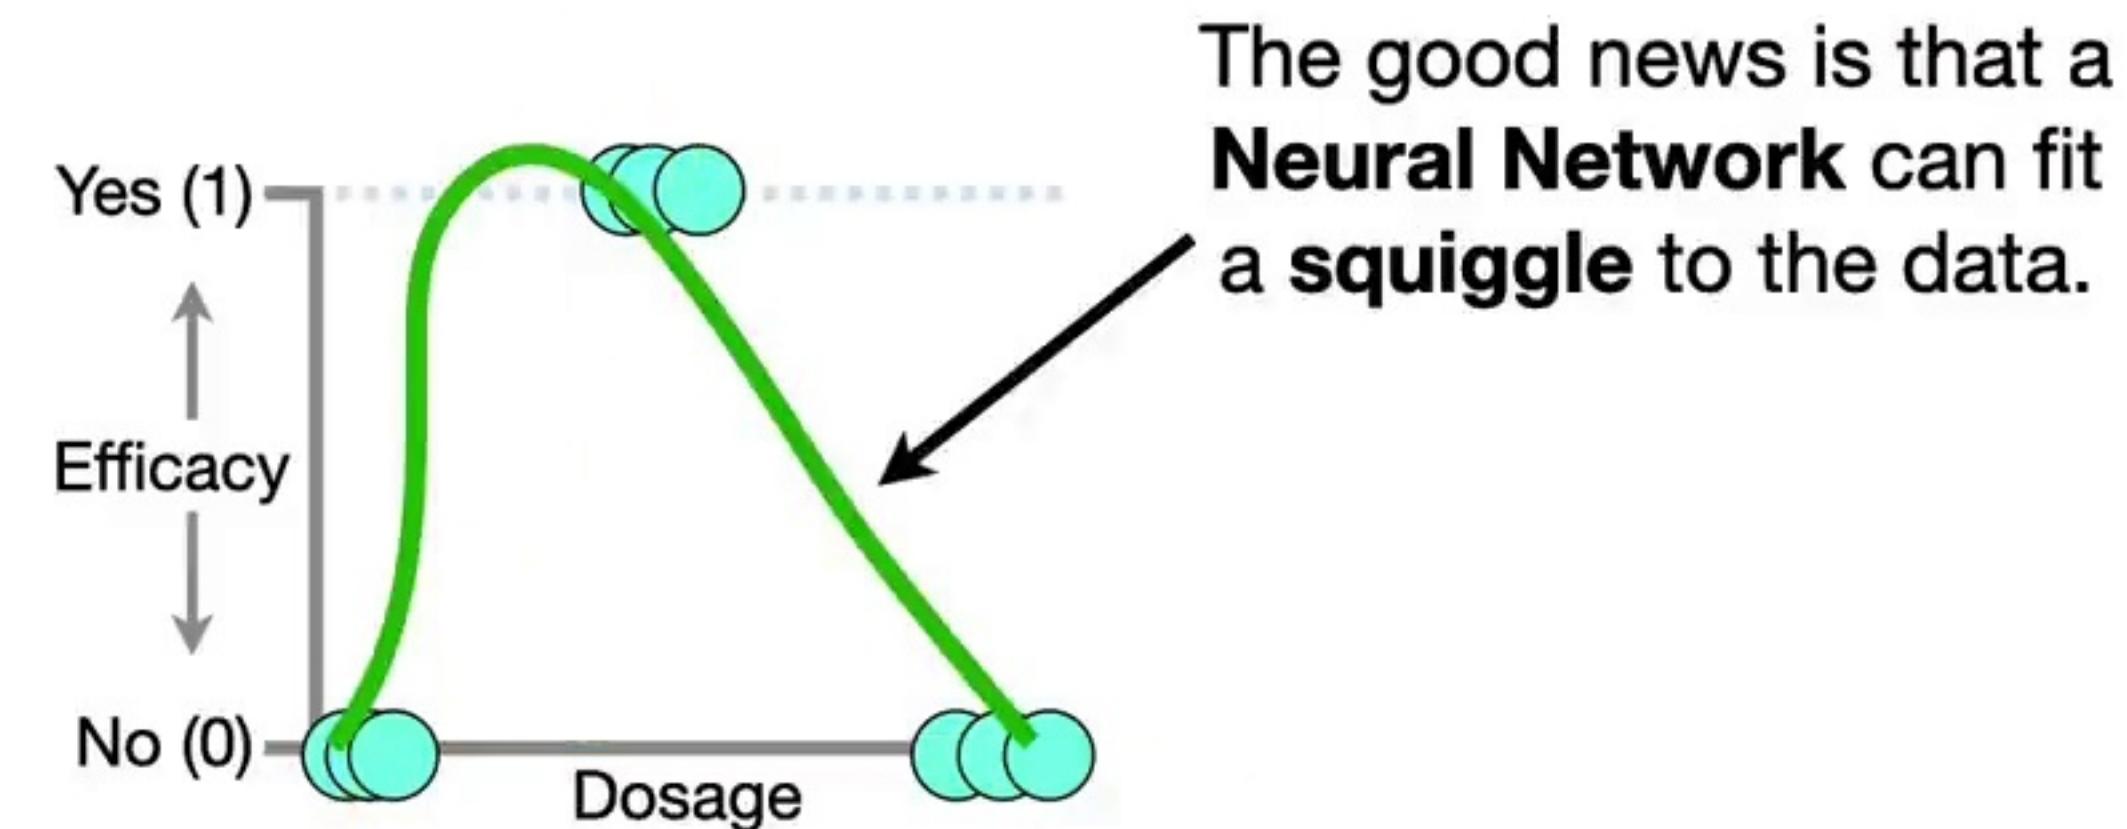


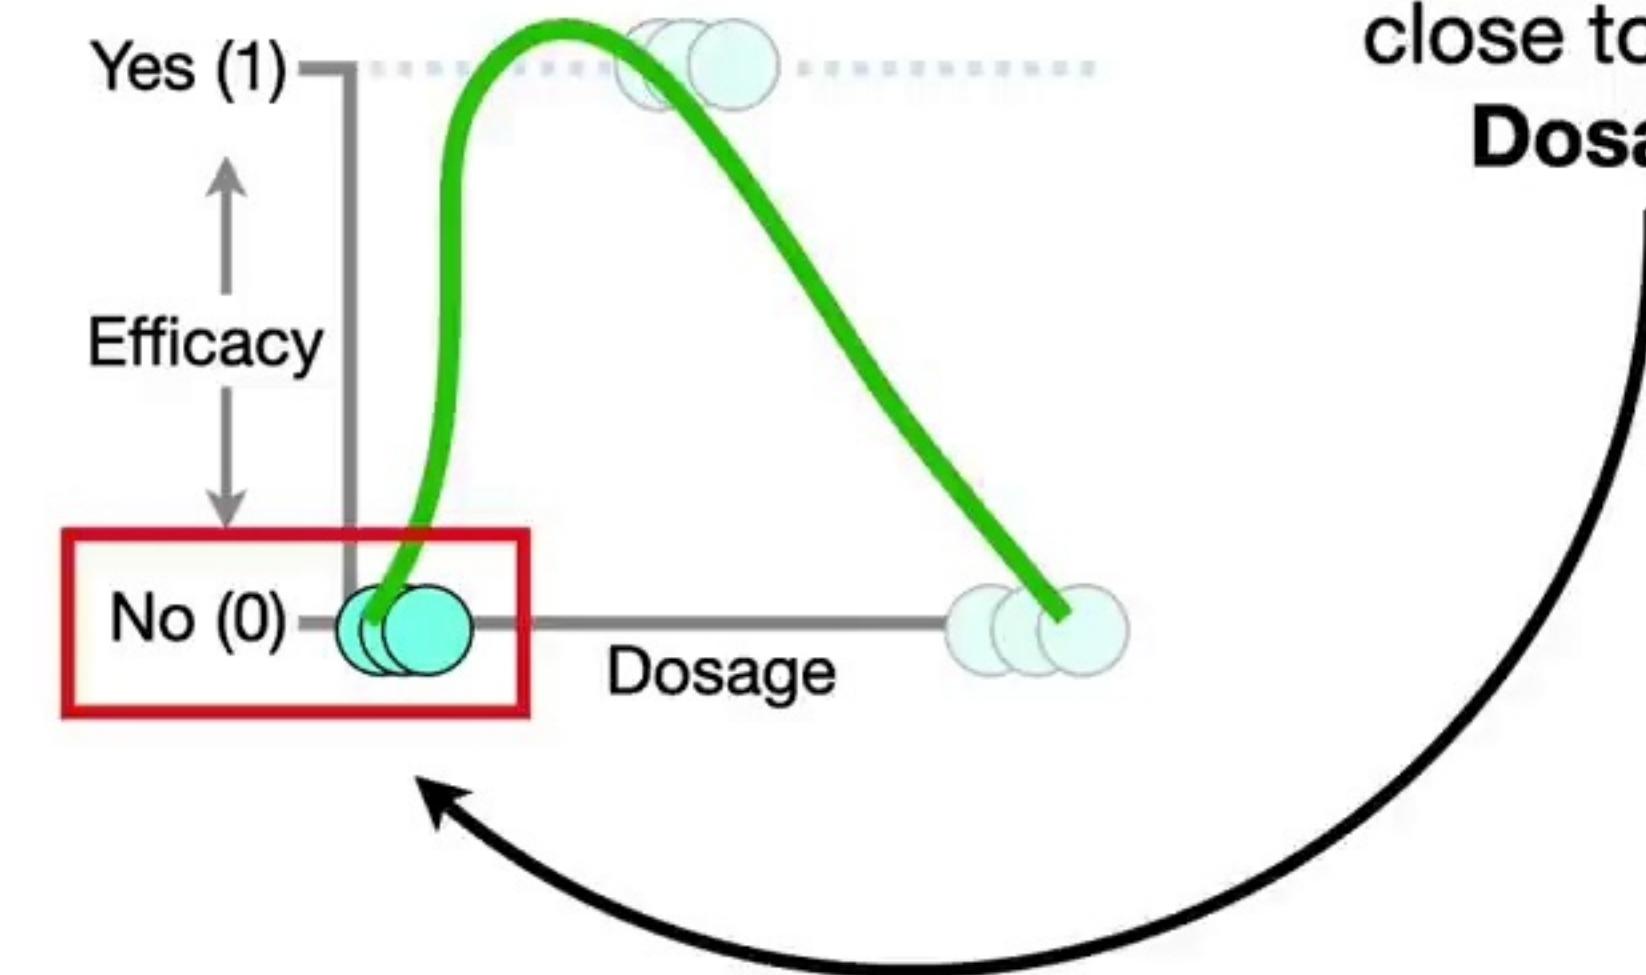


...because no matter how we rotate the **straight line**, it can only accurately predict **2** of the **3** dosages.

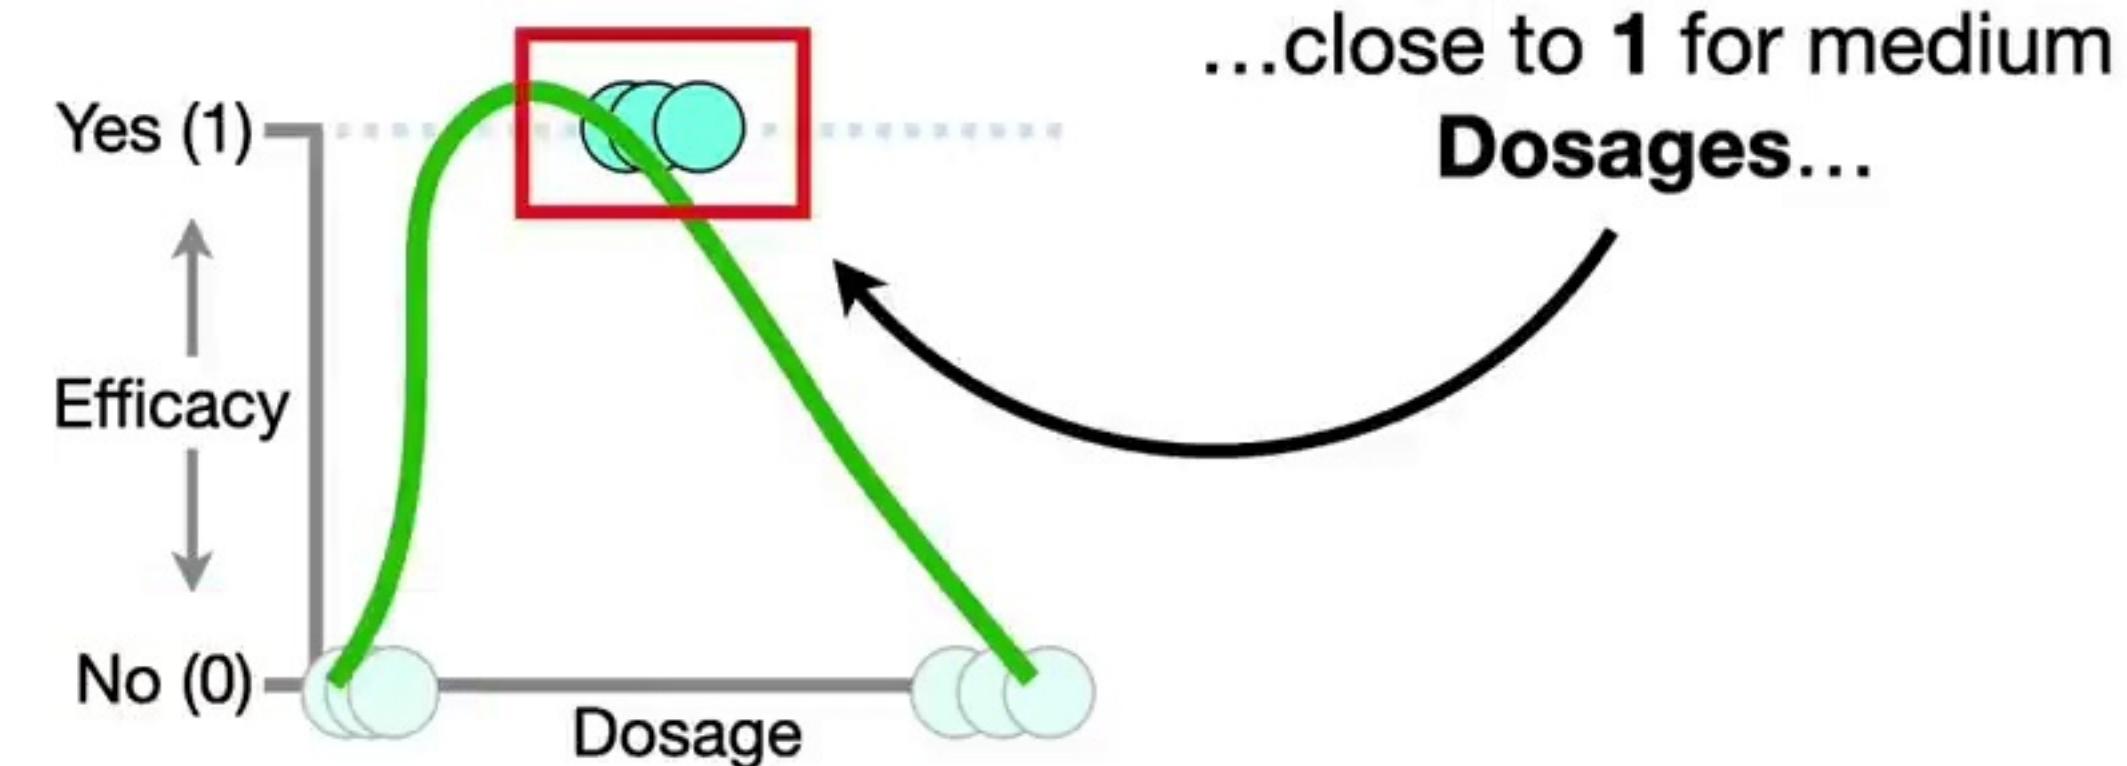


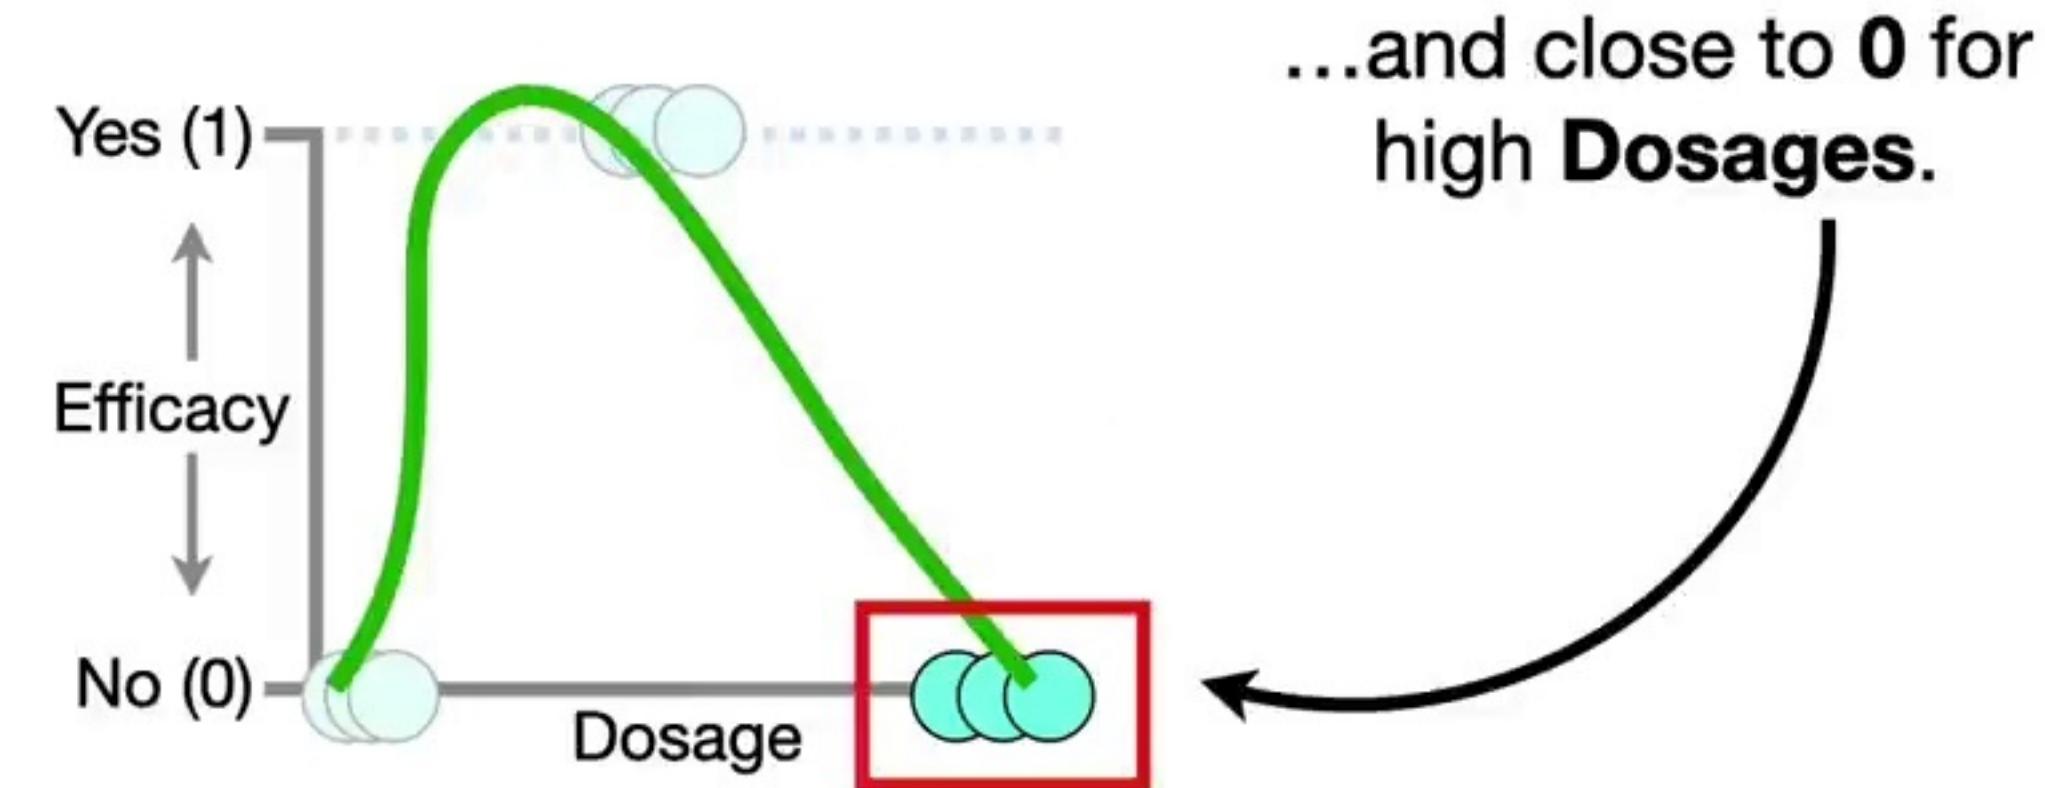
...because no matter how we rotate the **straight line**, it can only accurately predict **2** of the **3** dosages.

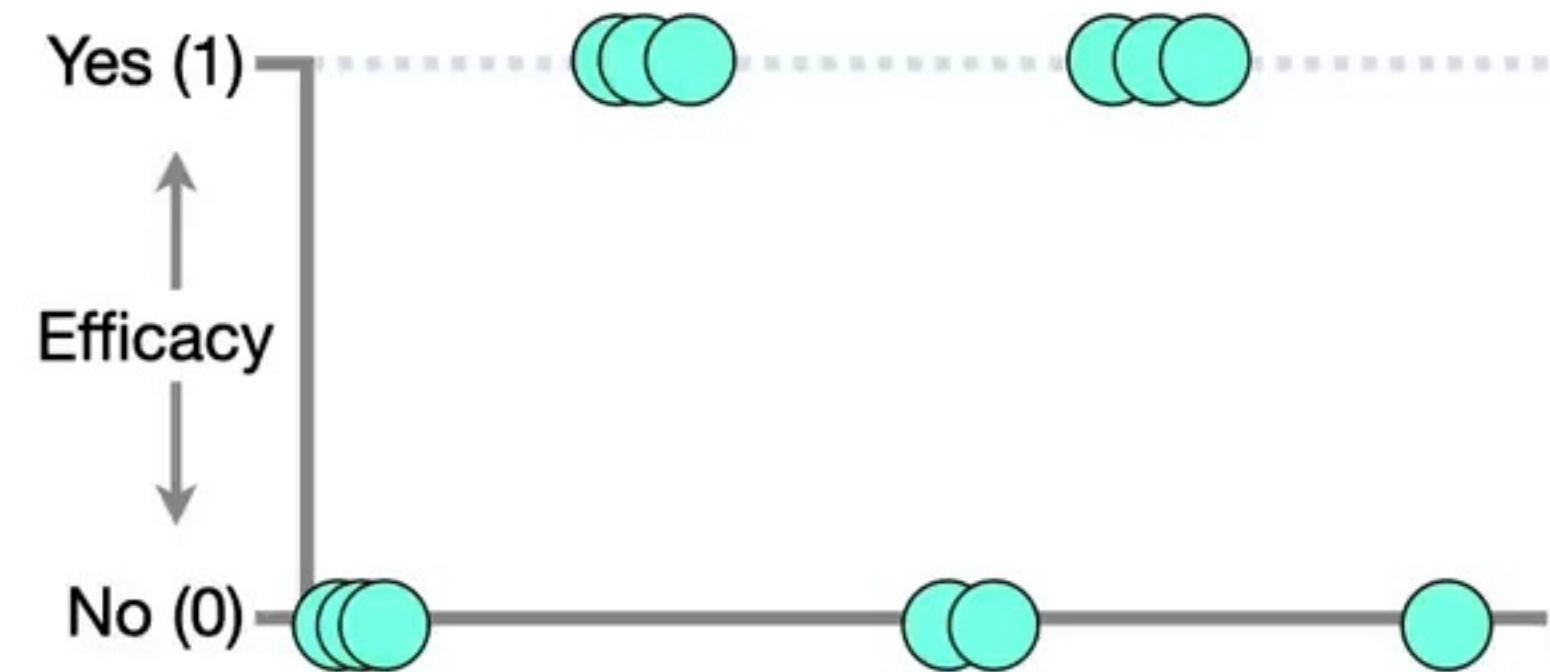




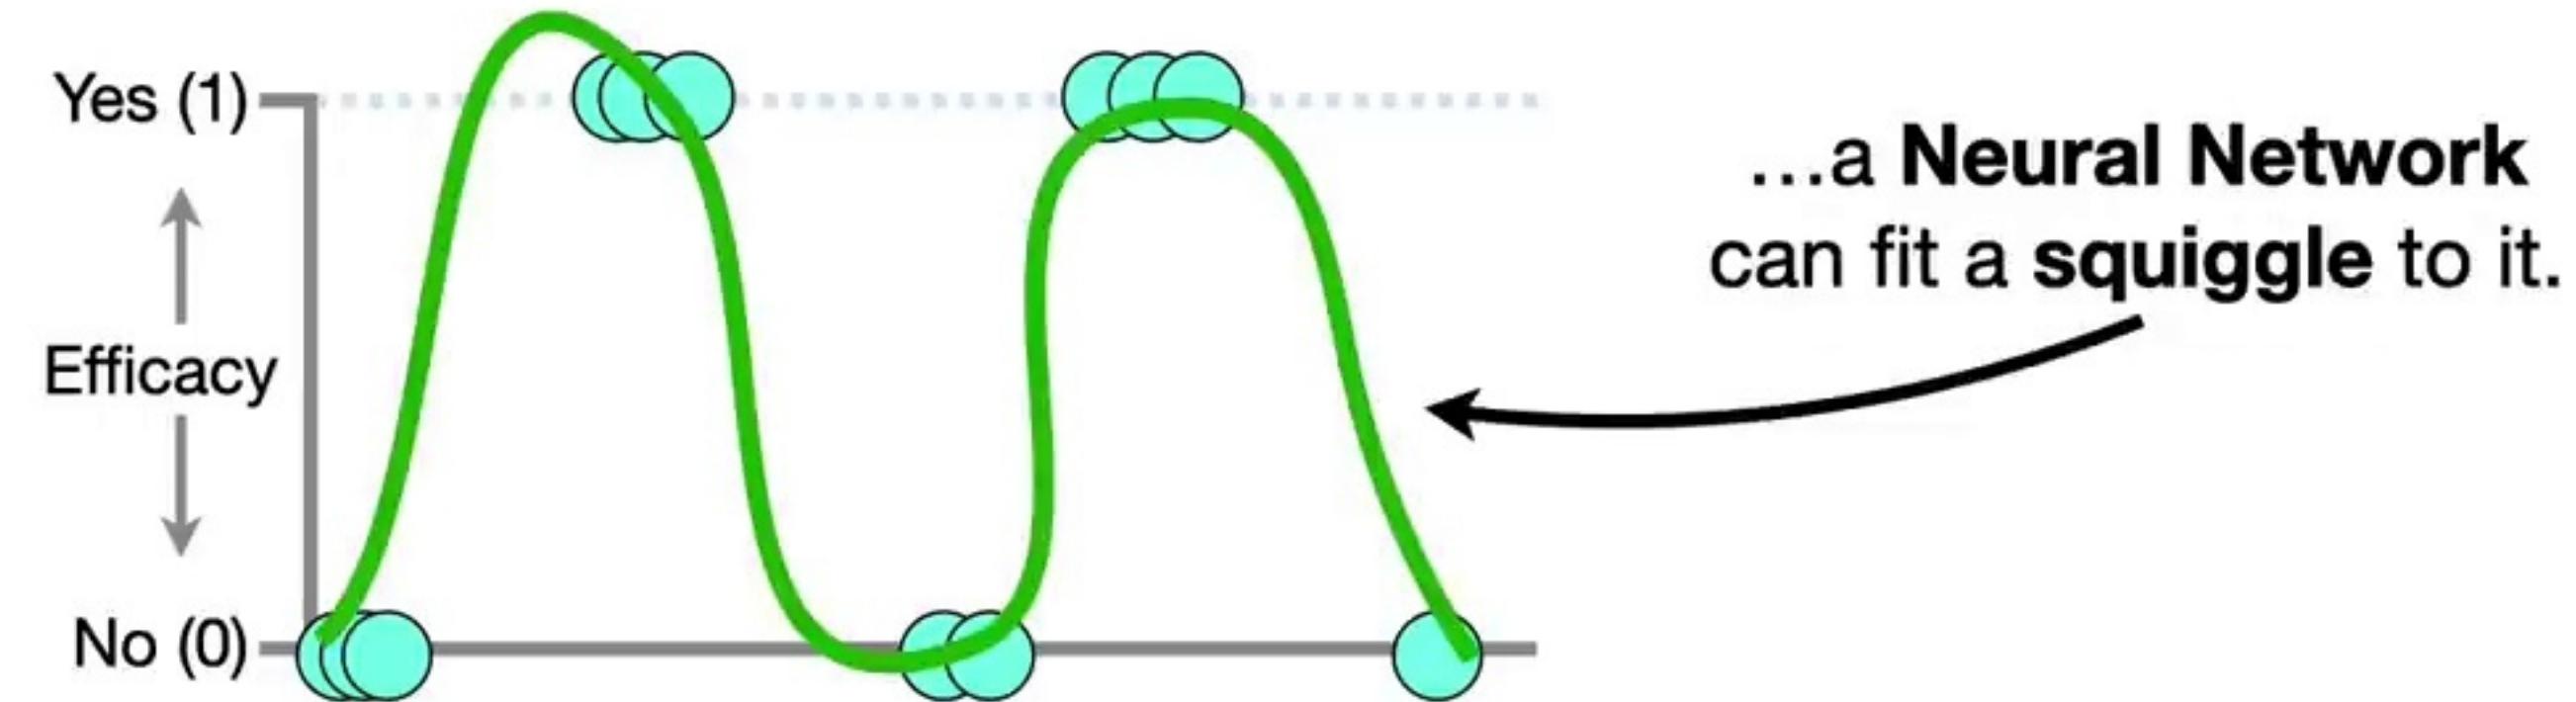
The **green squiggle** is  
close to **0** for low  
**Dosages...**





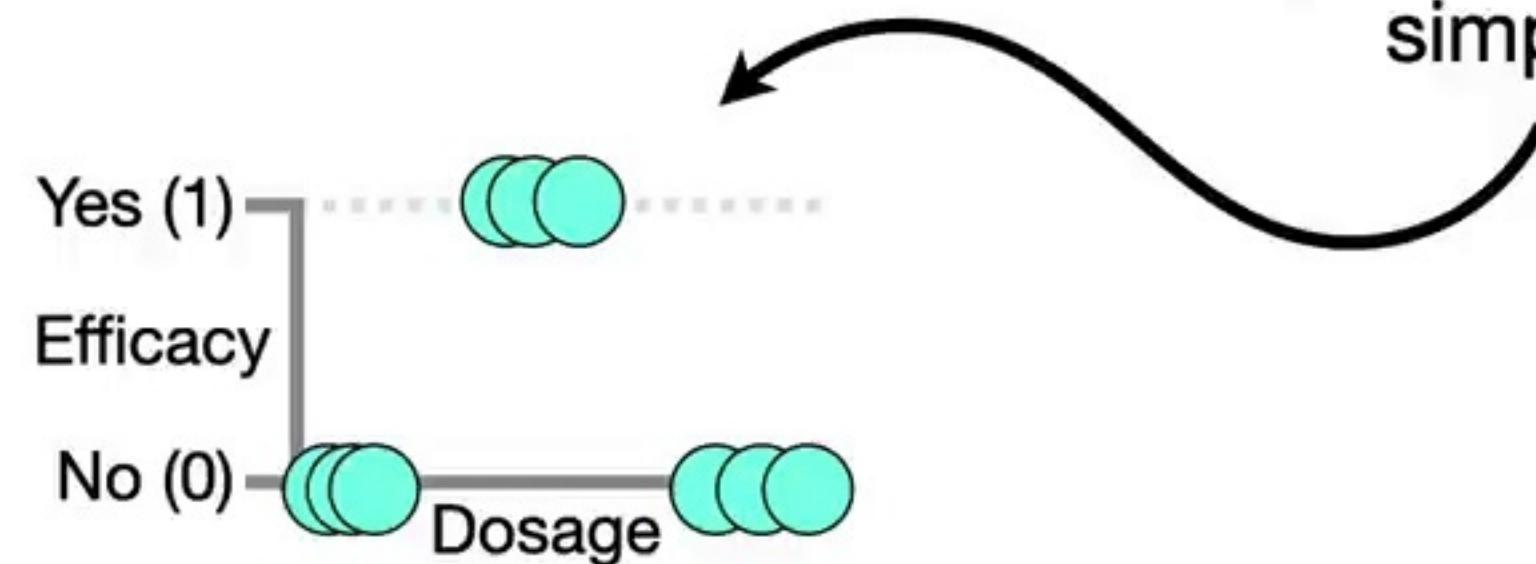


And even if we have  
really complicated  
dataset like this...



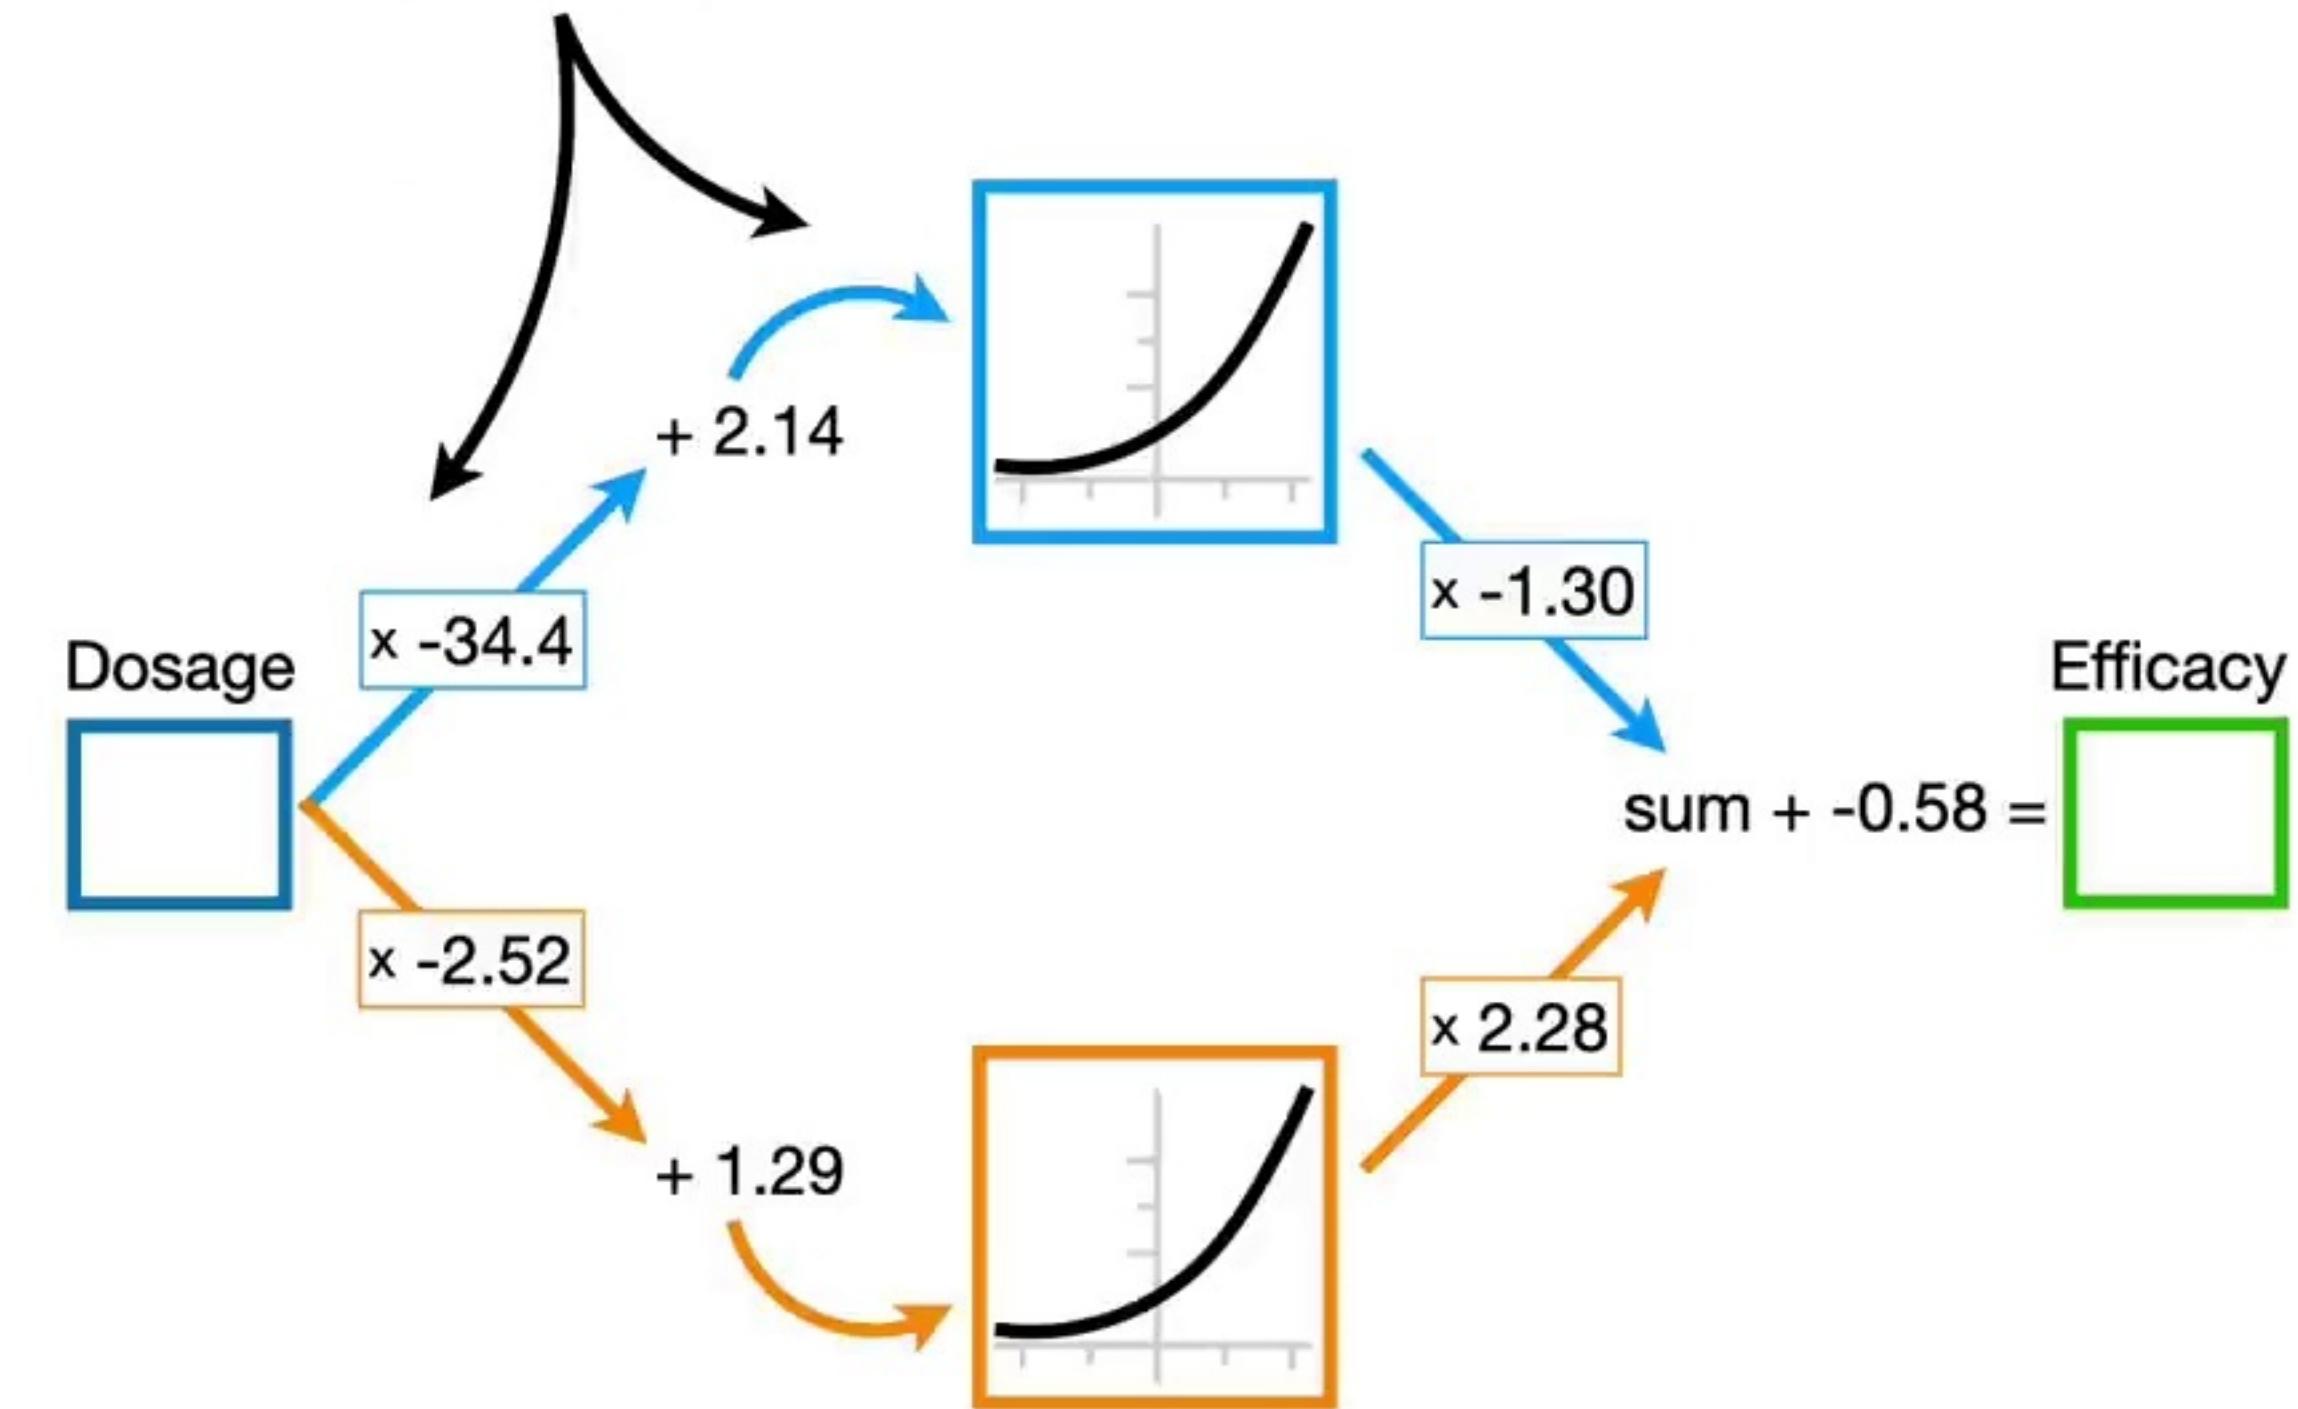
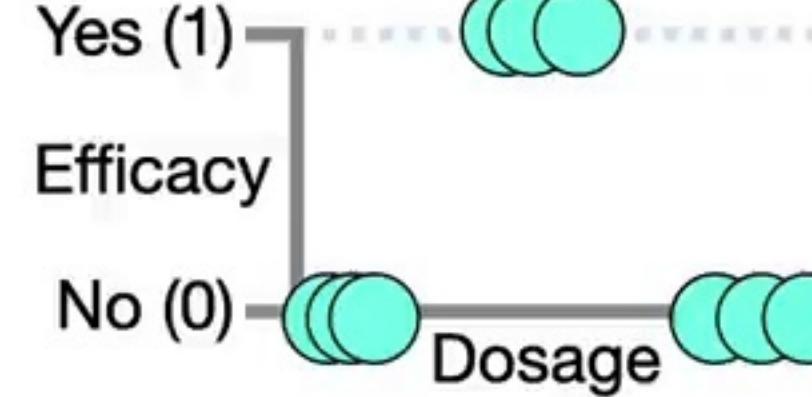


In this **StatQuest**, we're going to use this super simple dataset...



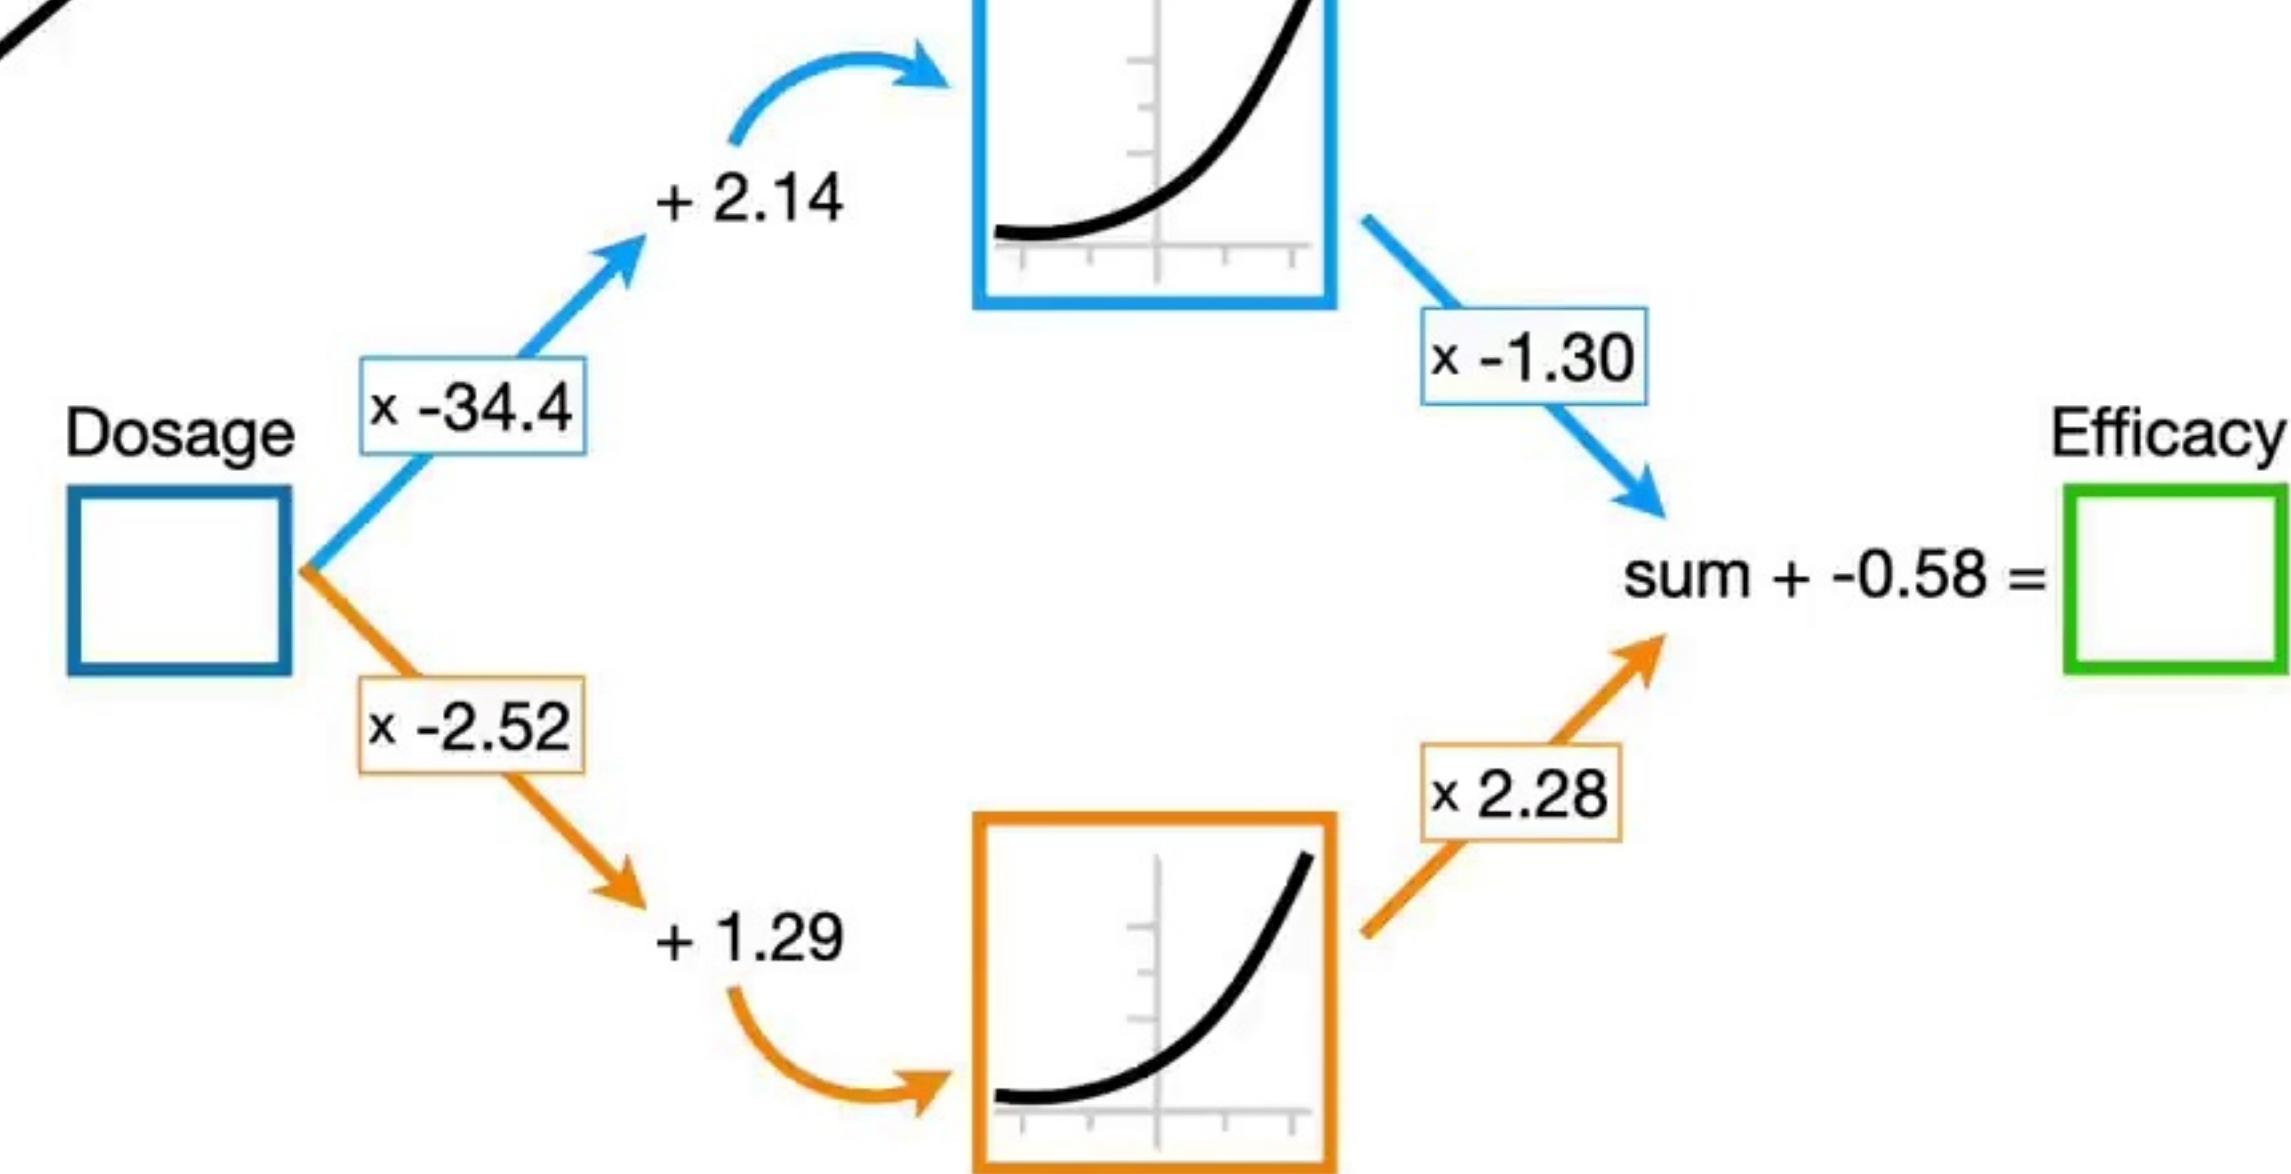
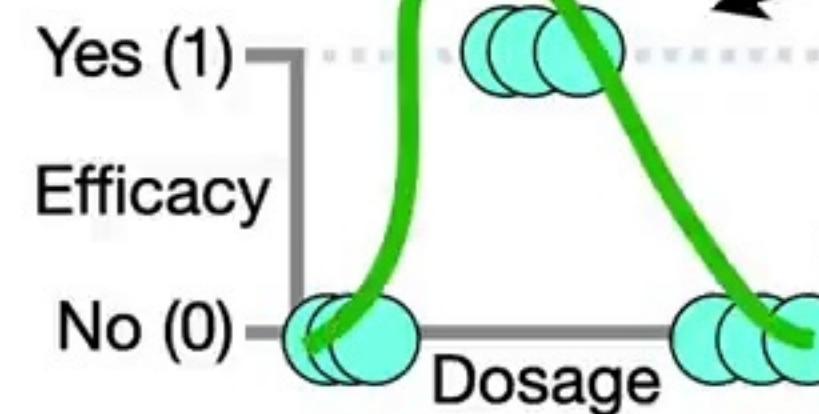


...and show how this  
**Neural Network...**



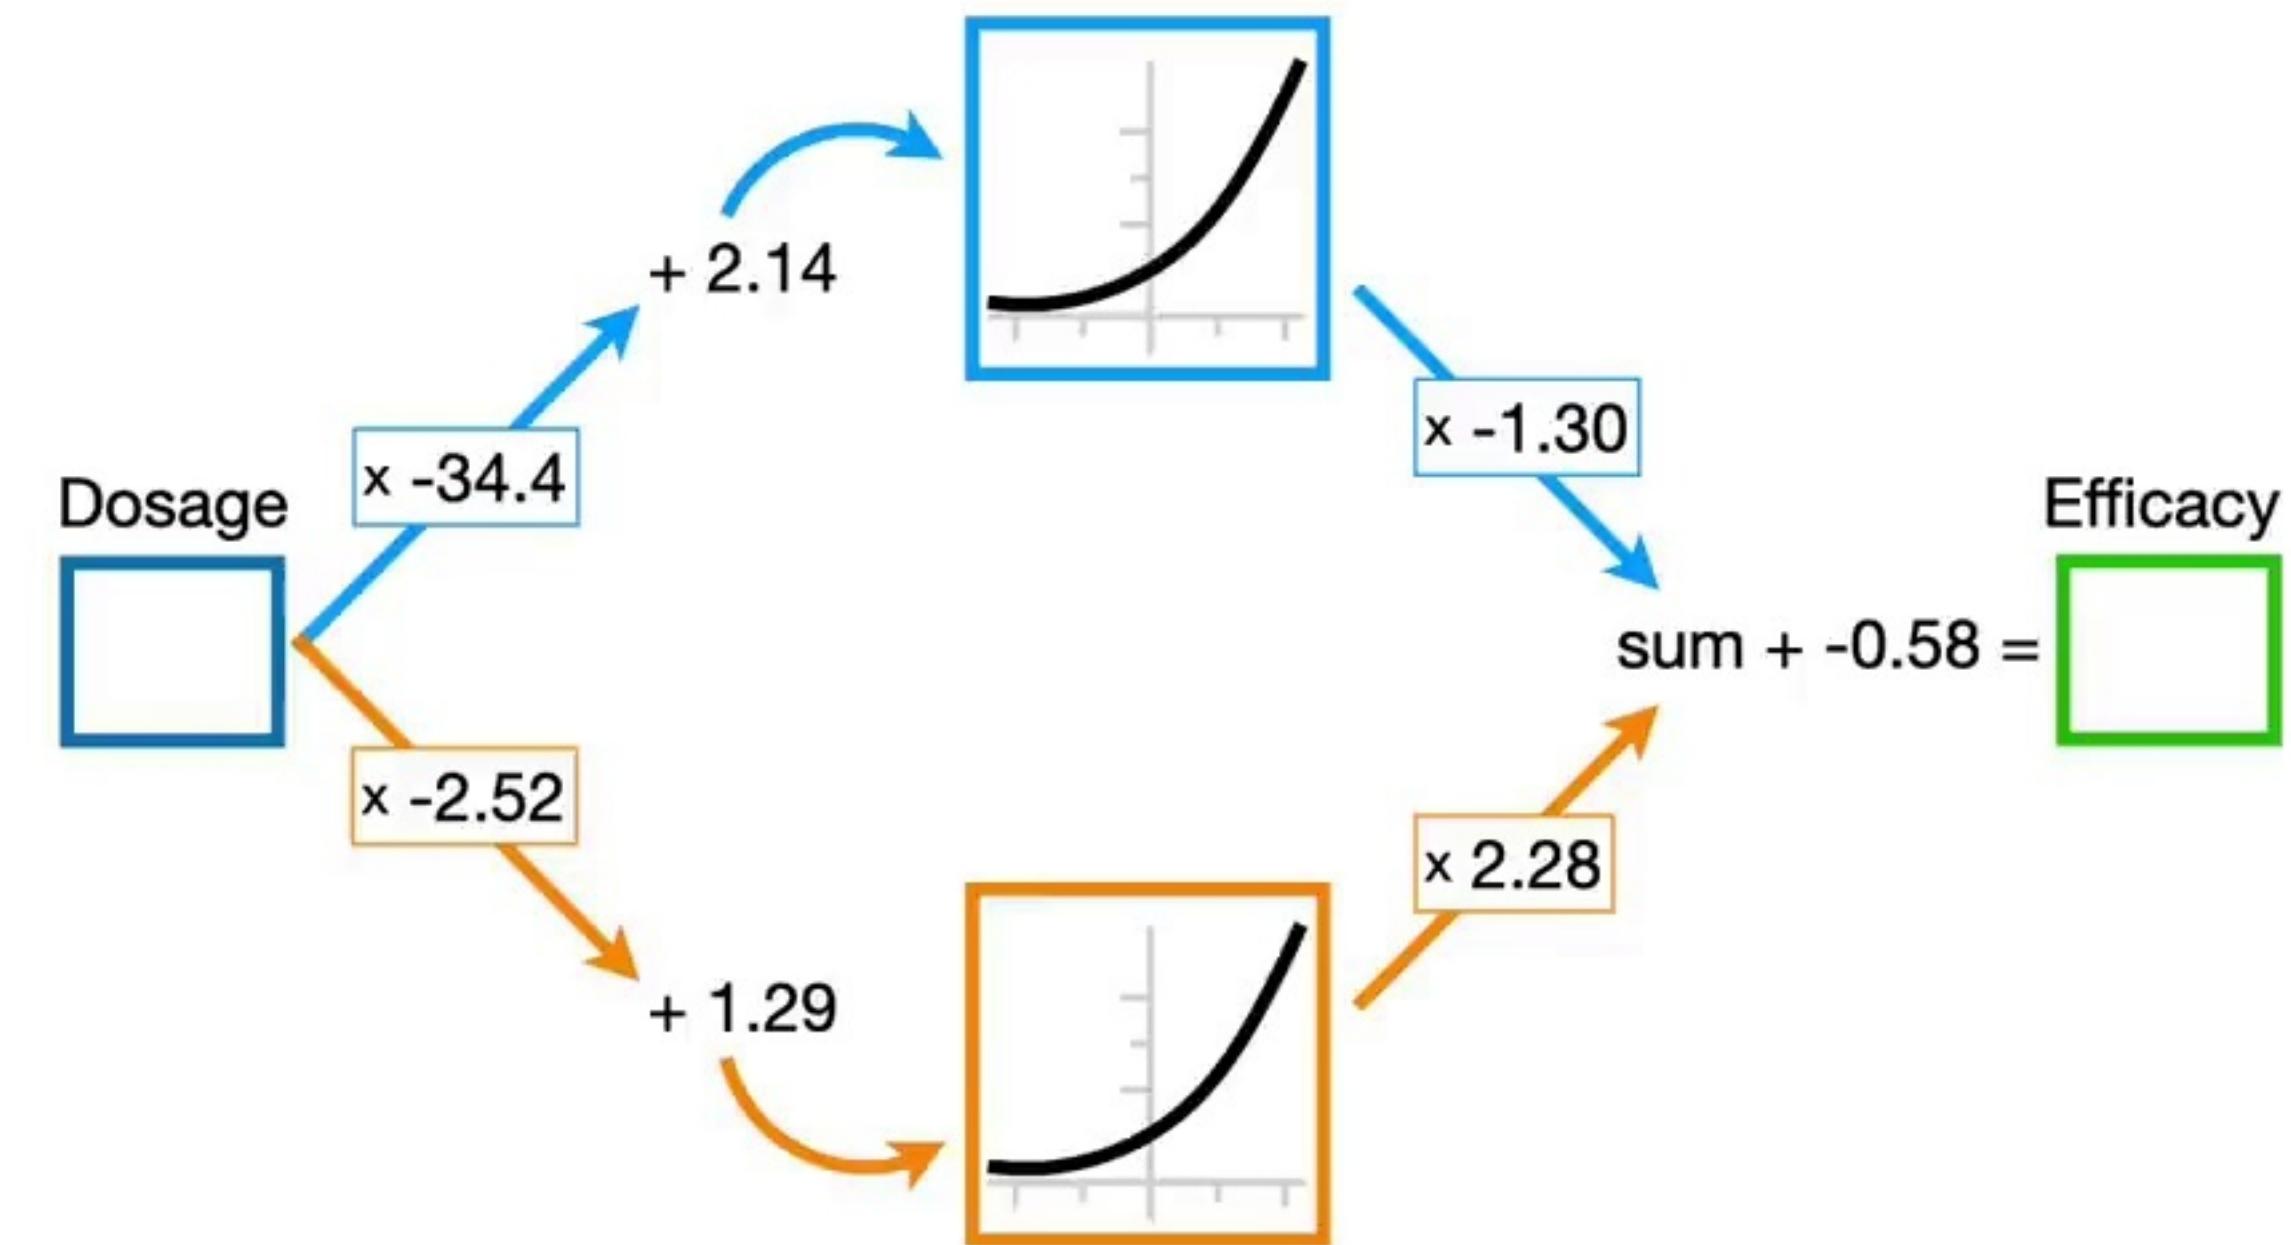


...creates this **green** squiggle.



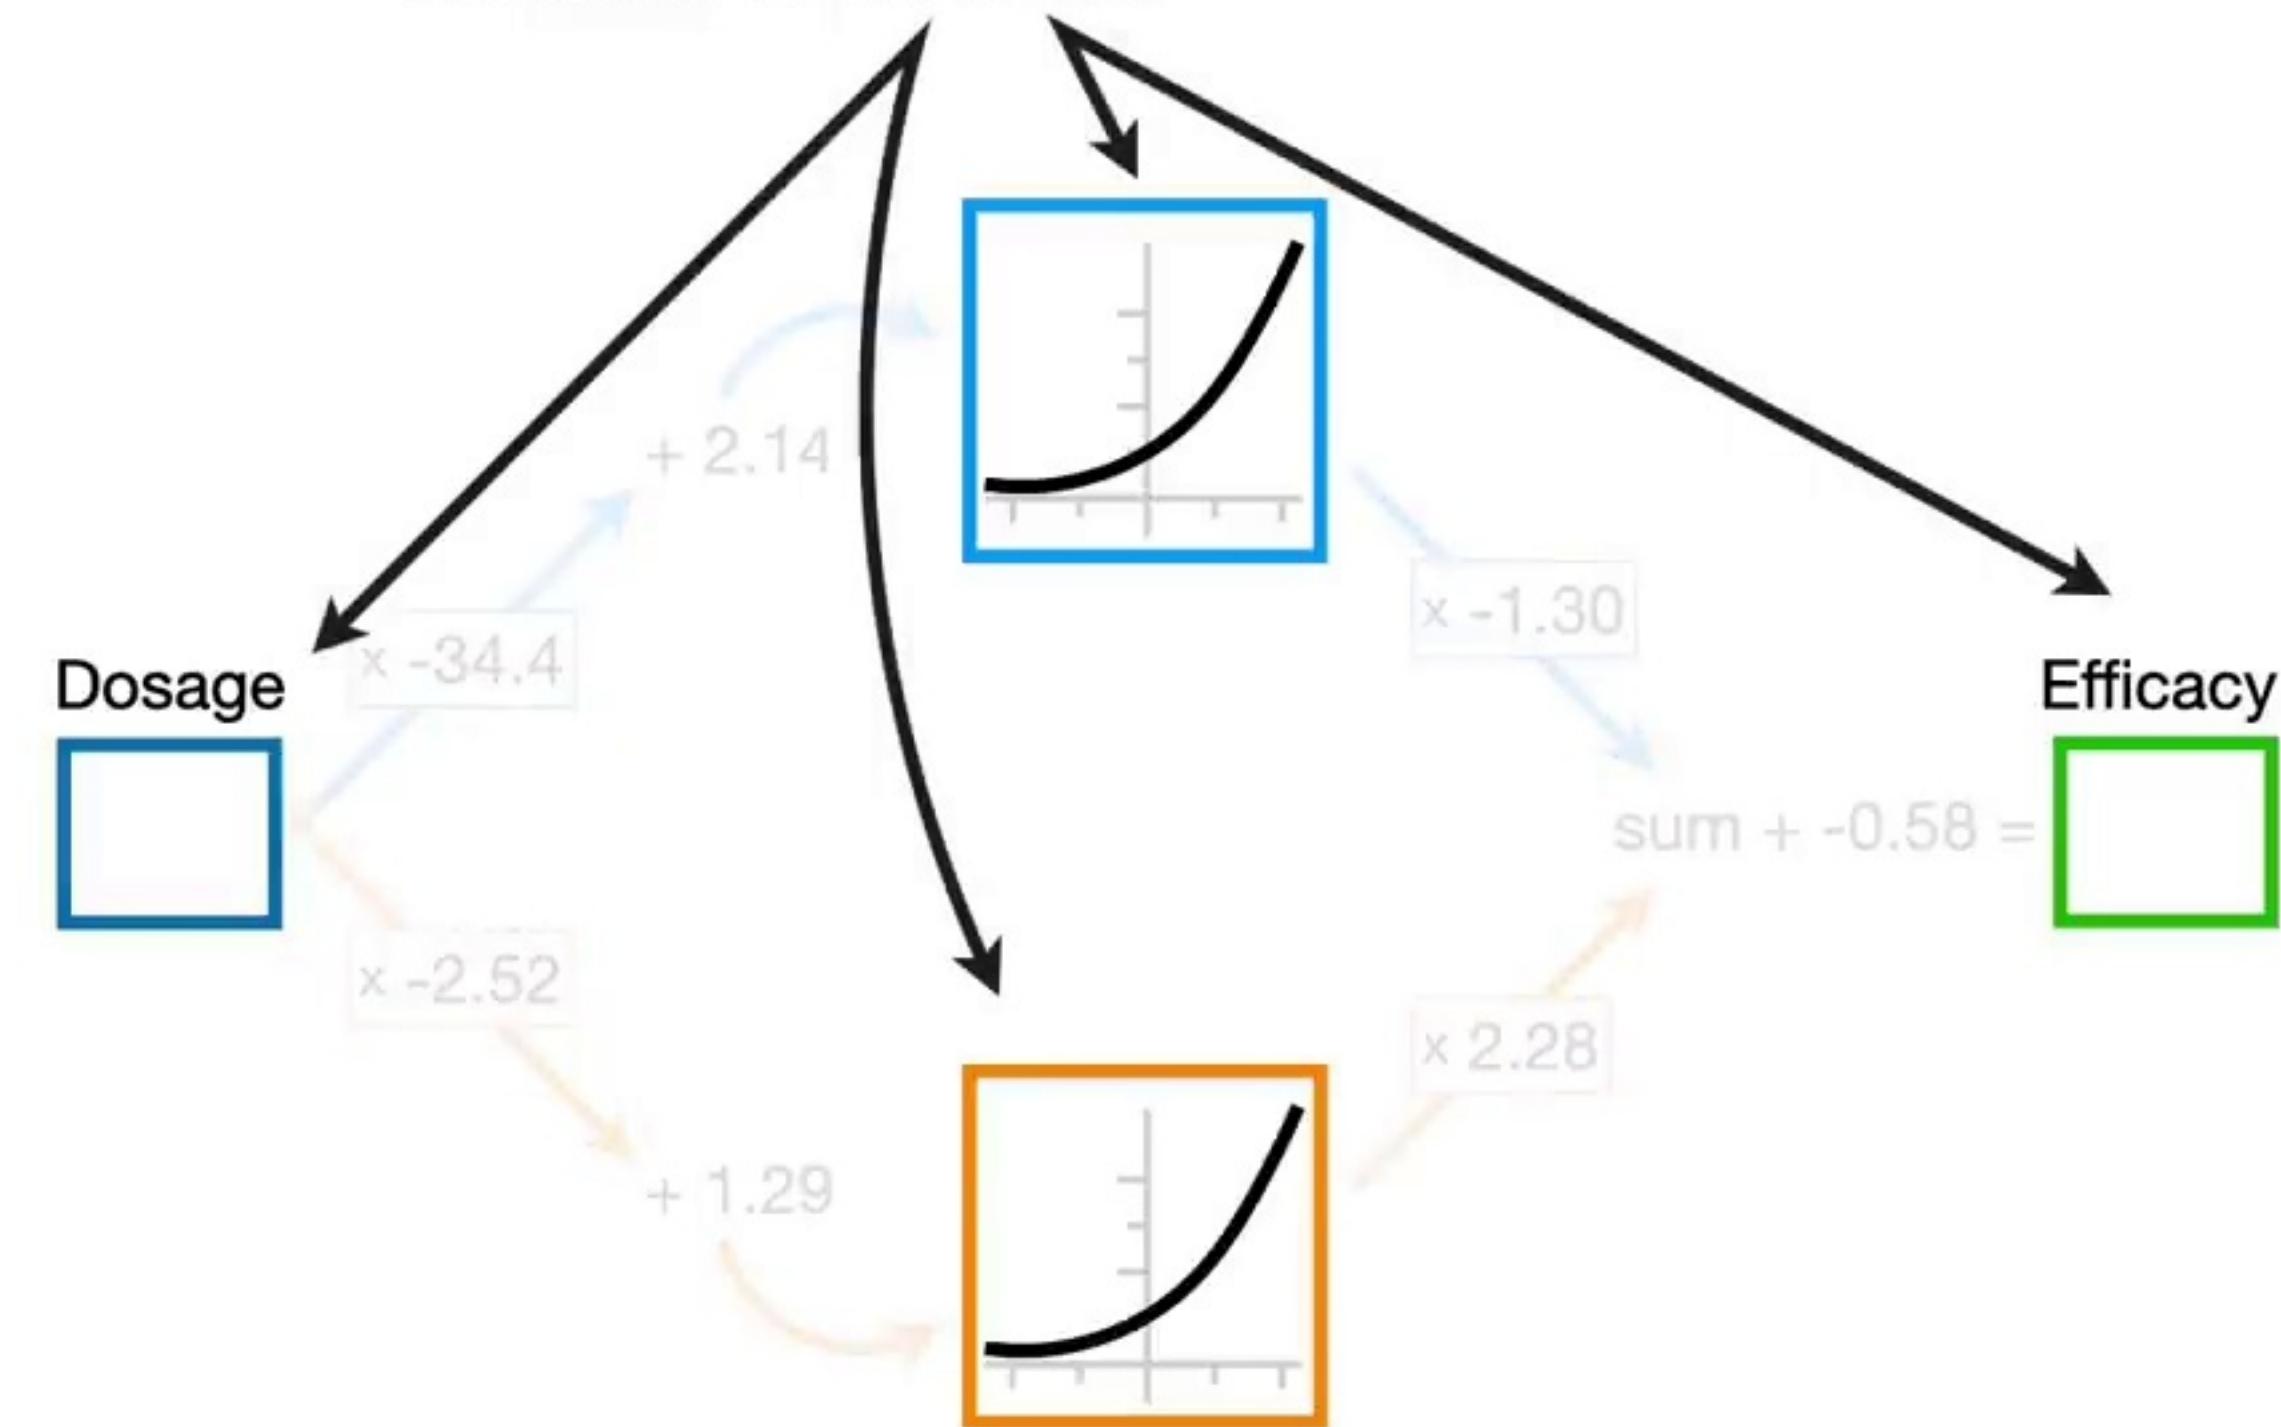


But first let's just talk about what a **Neural Network** is.



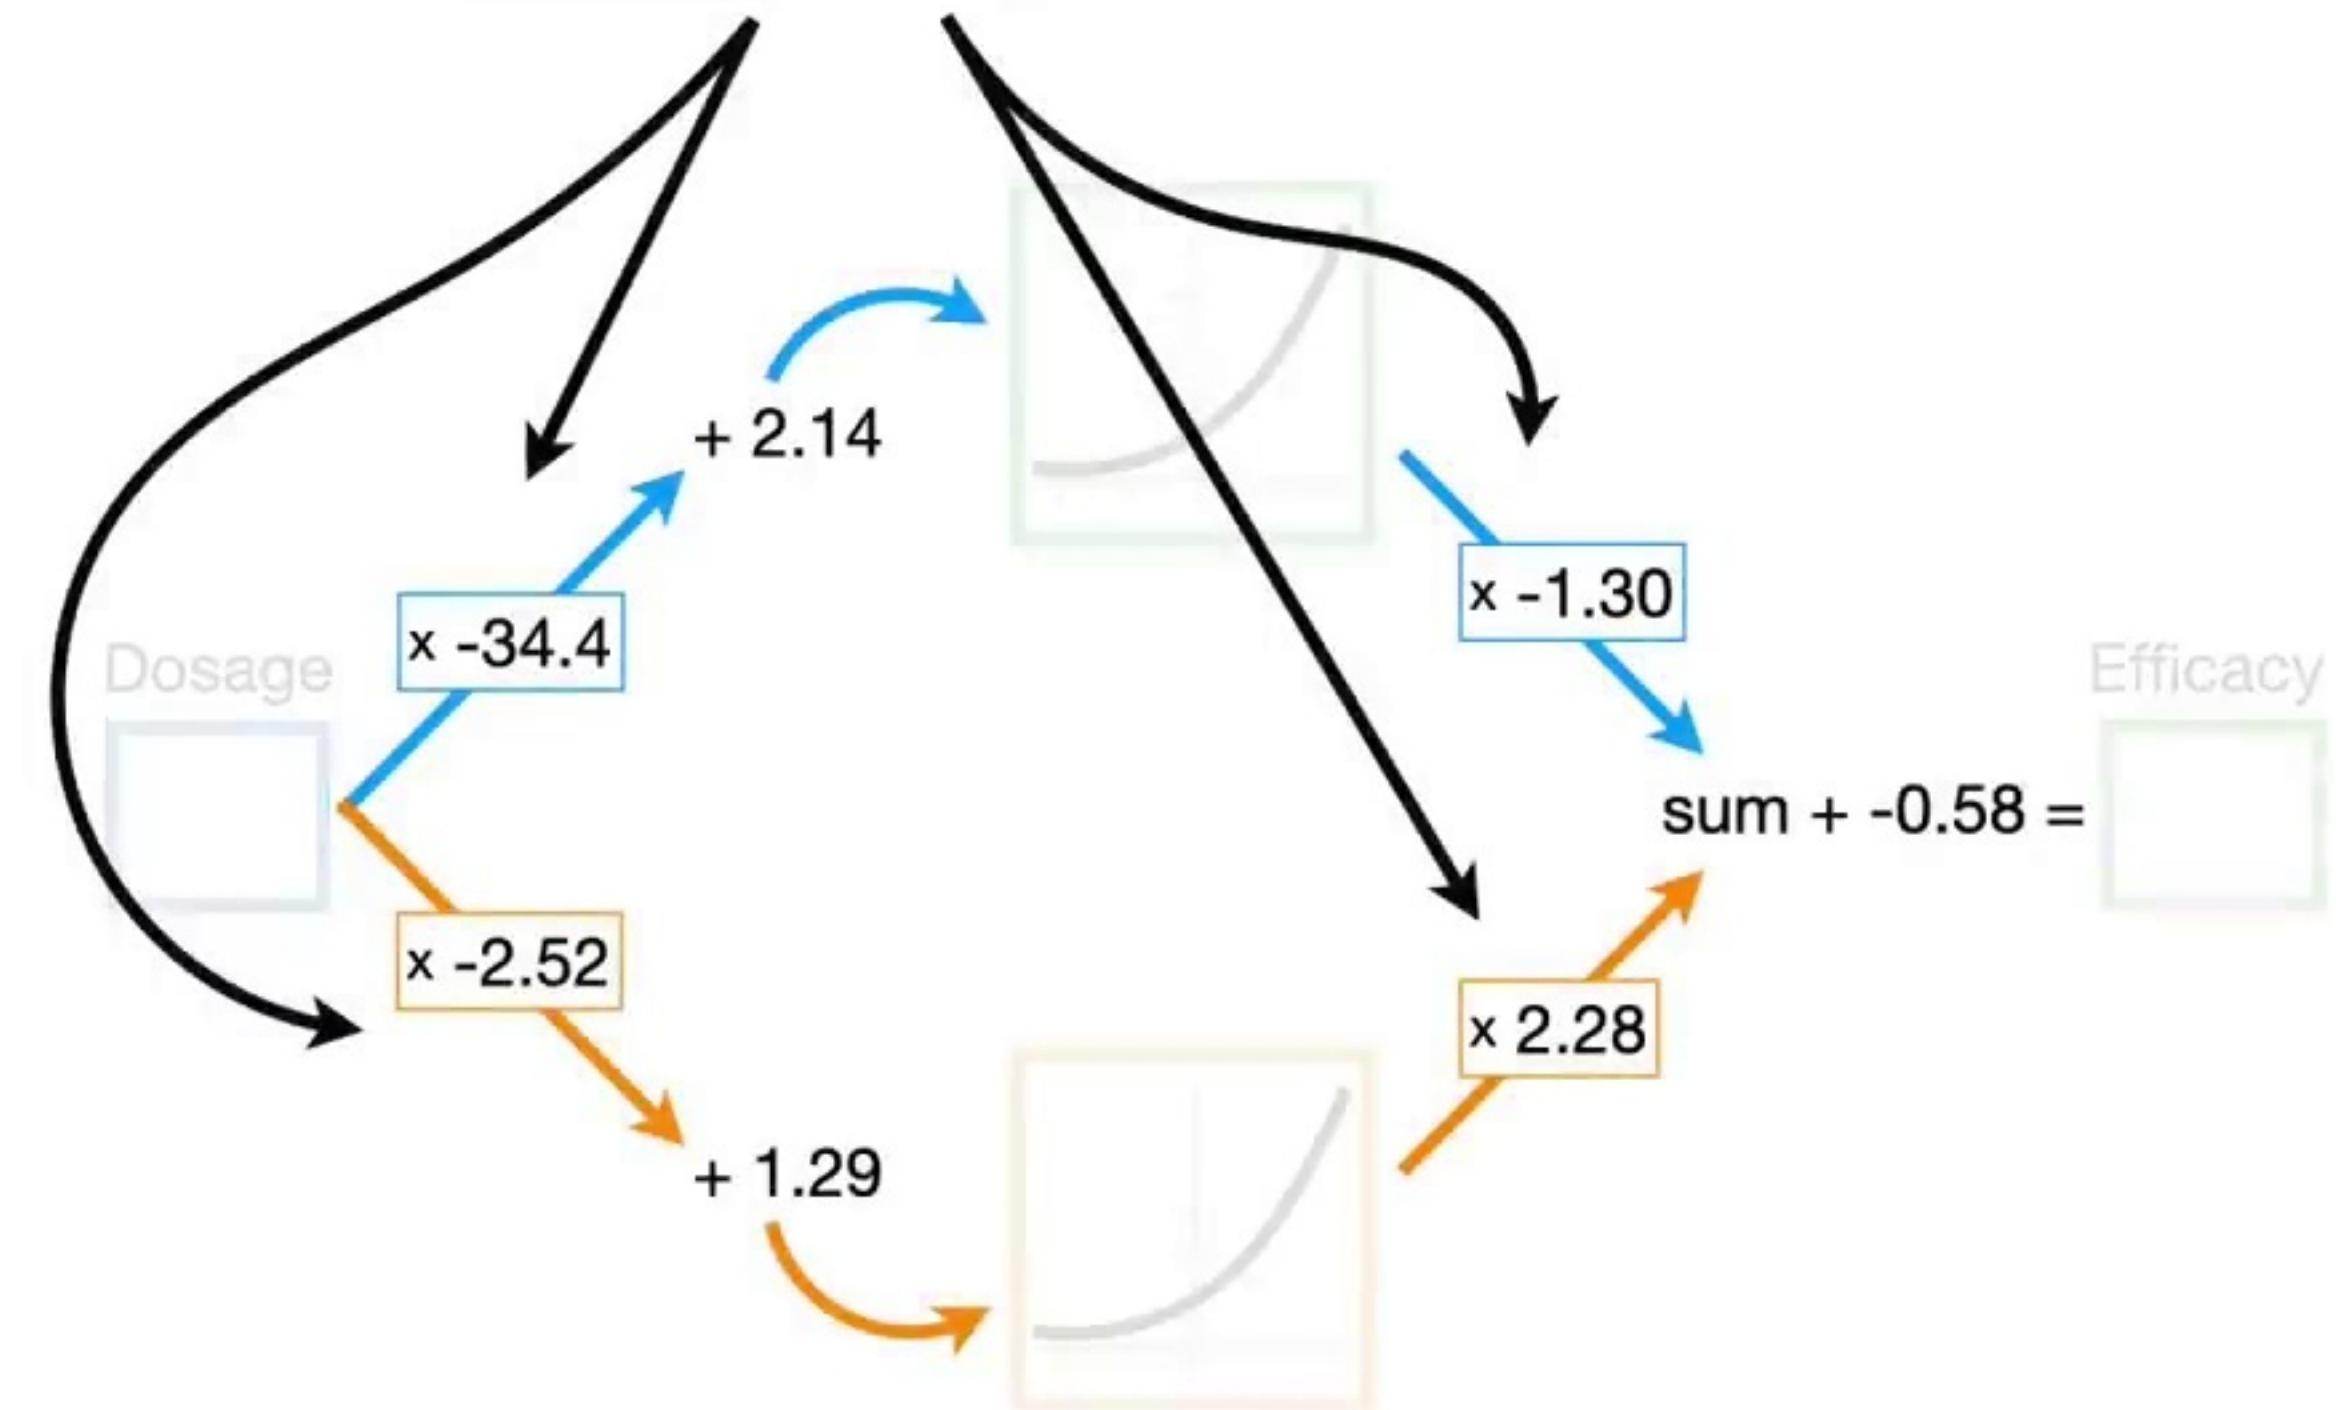


## A Neural Network consists of Nodes...



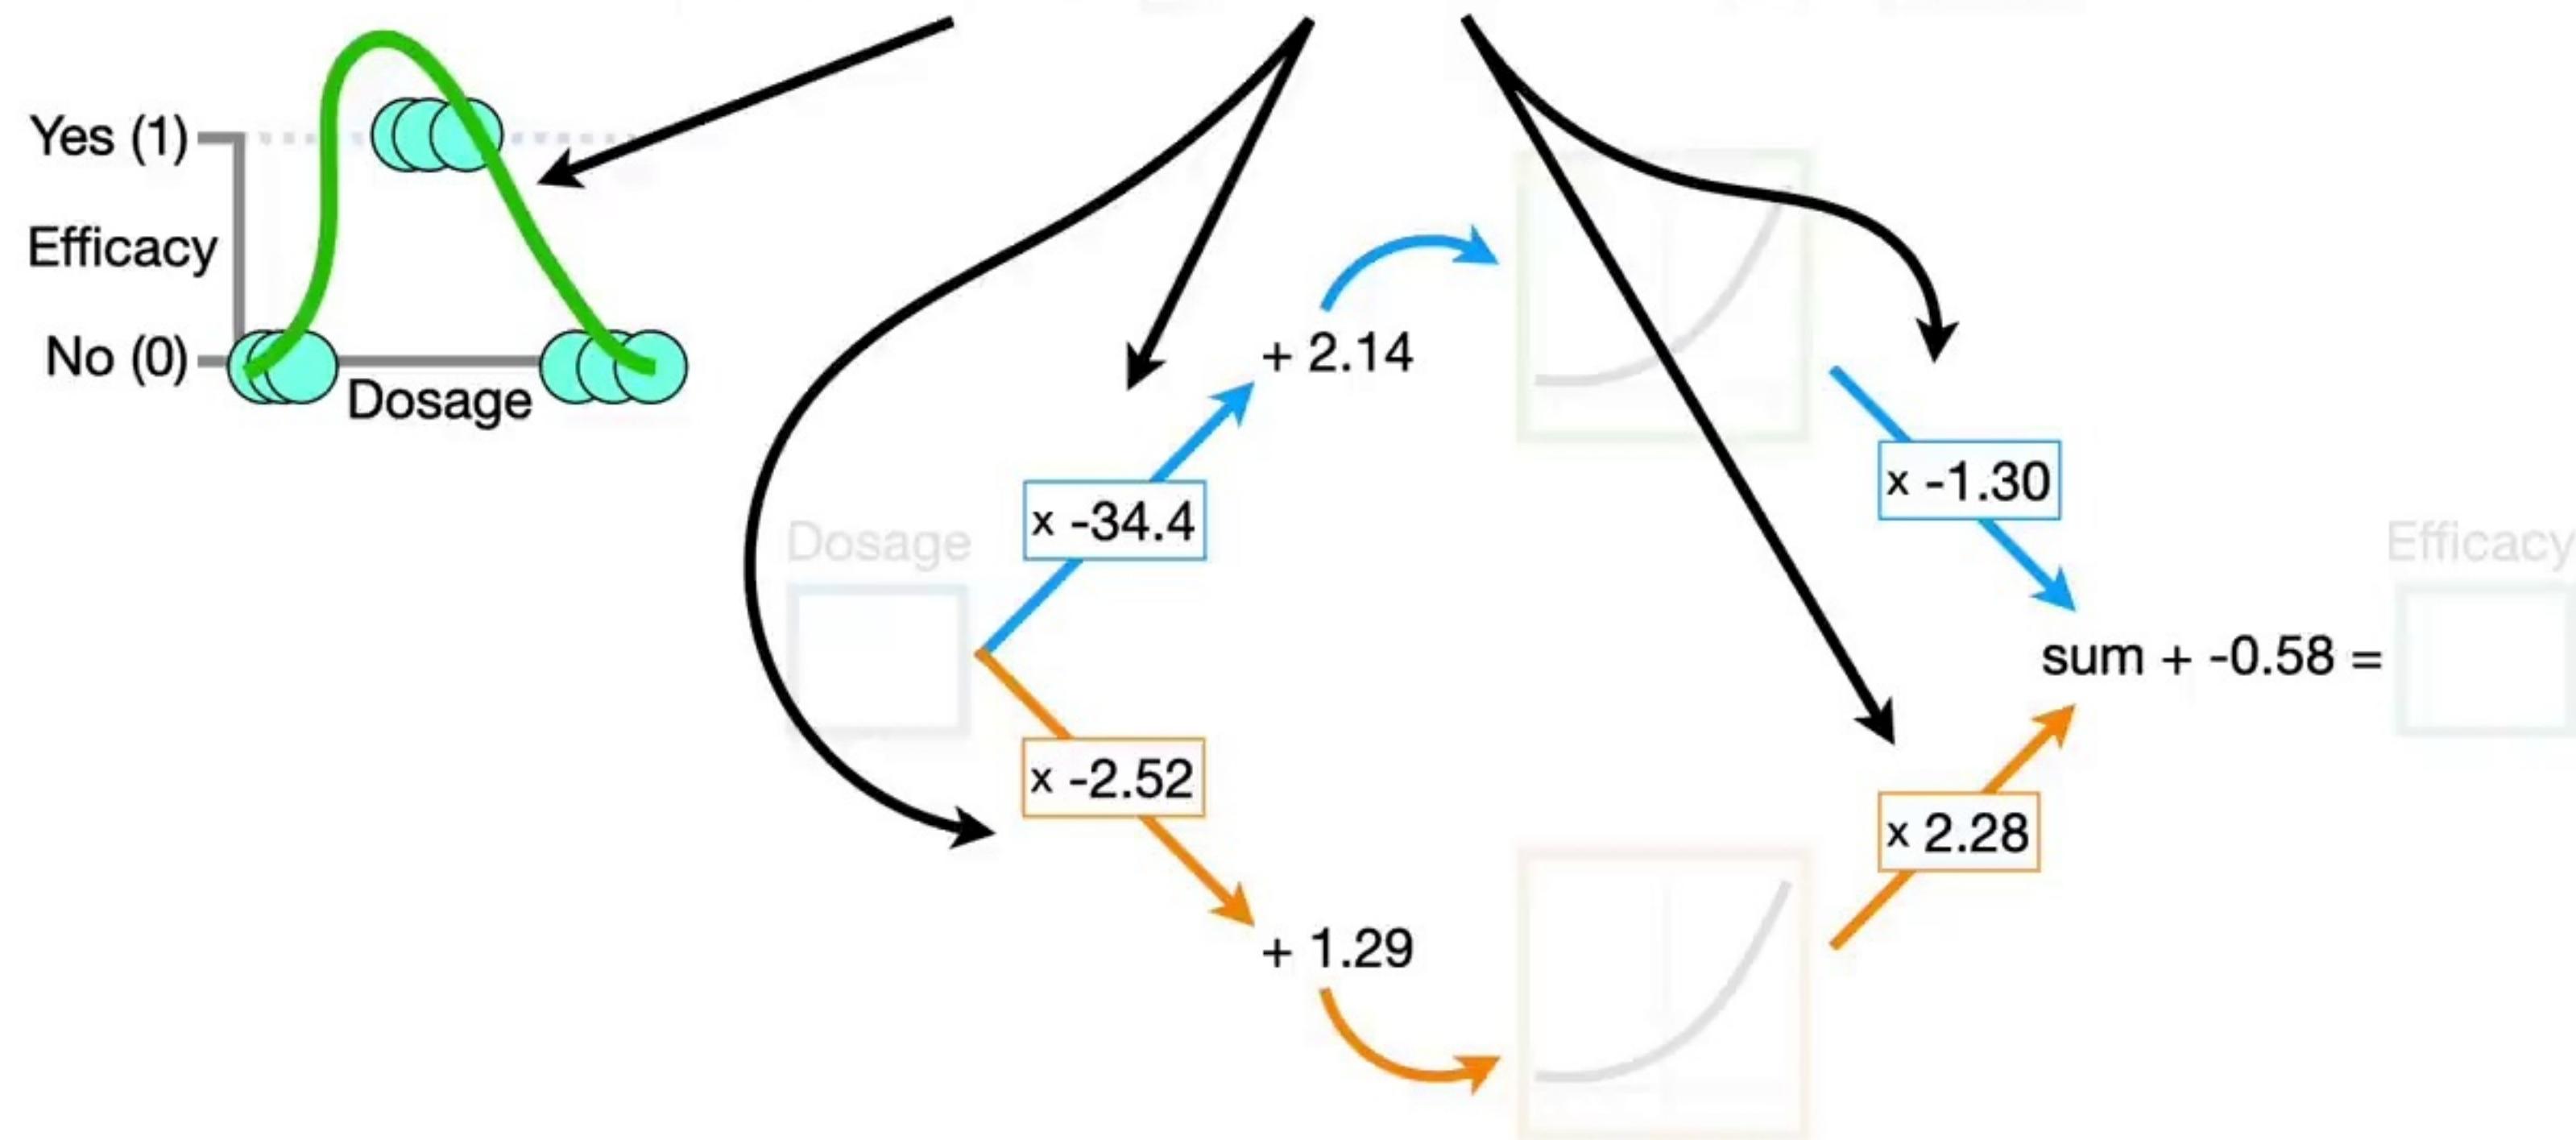


...and **connections**  
between the nodes.



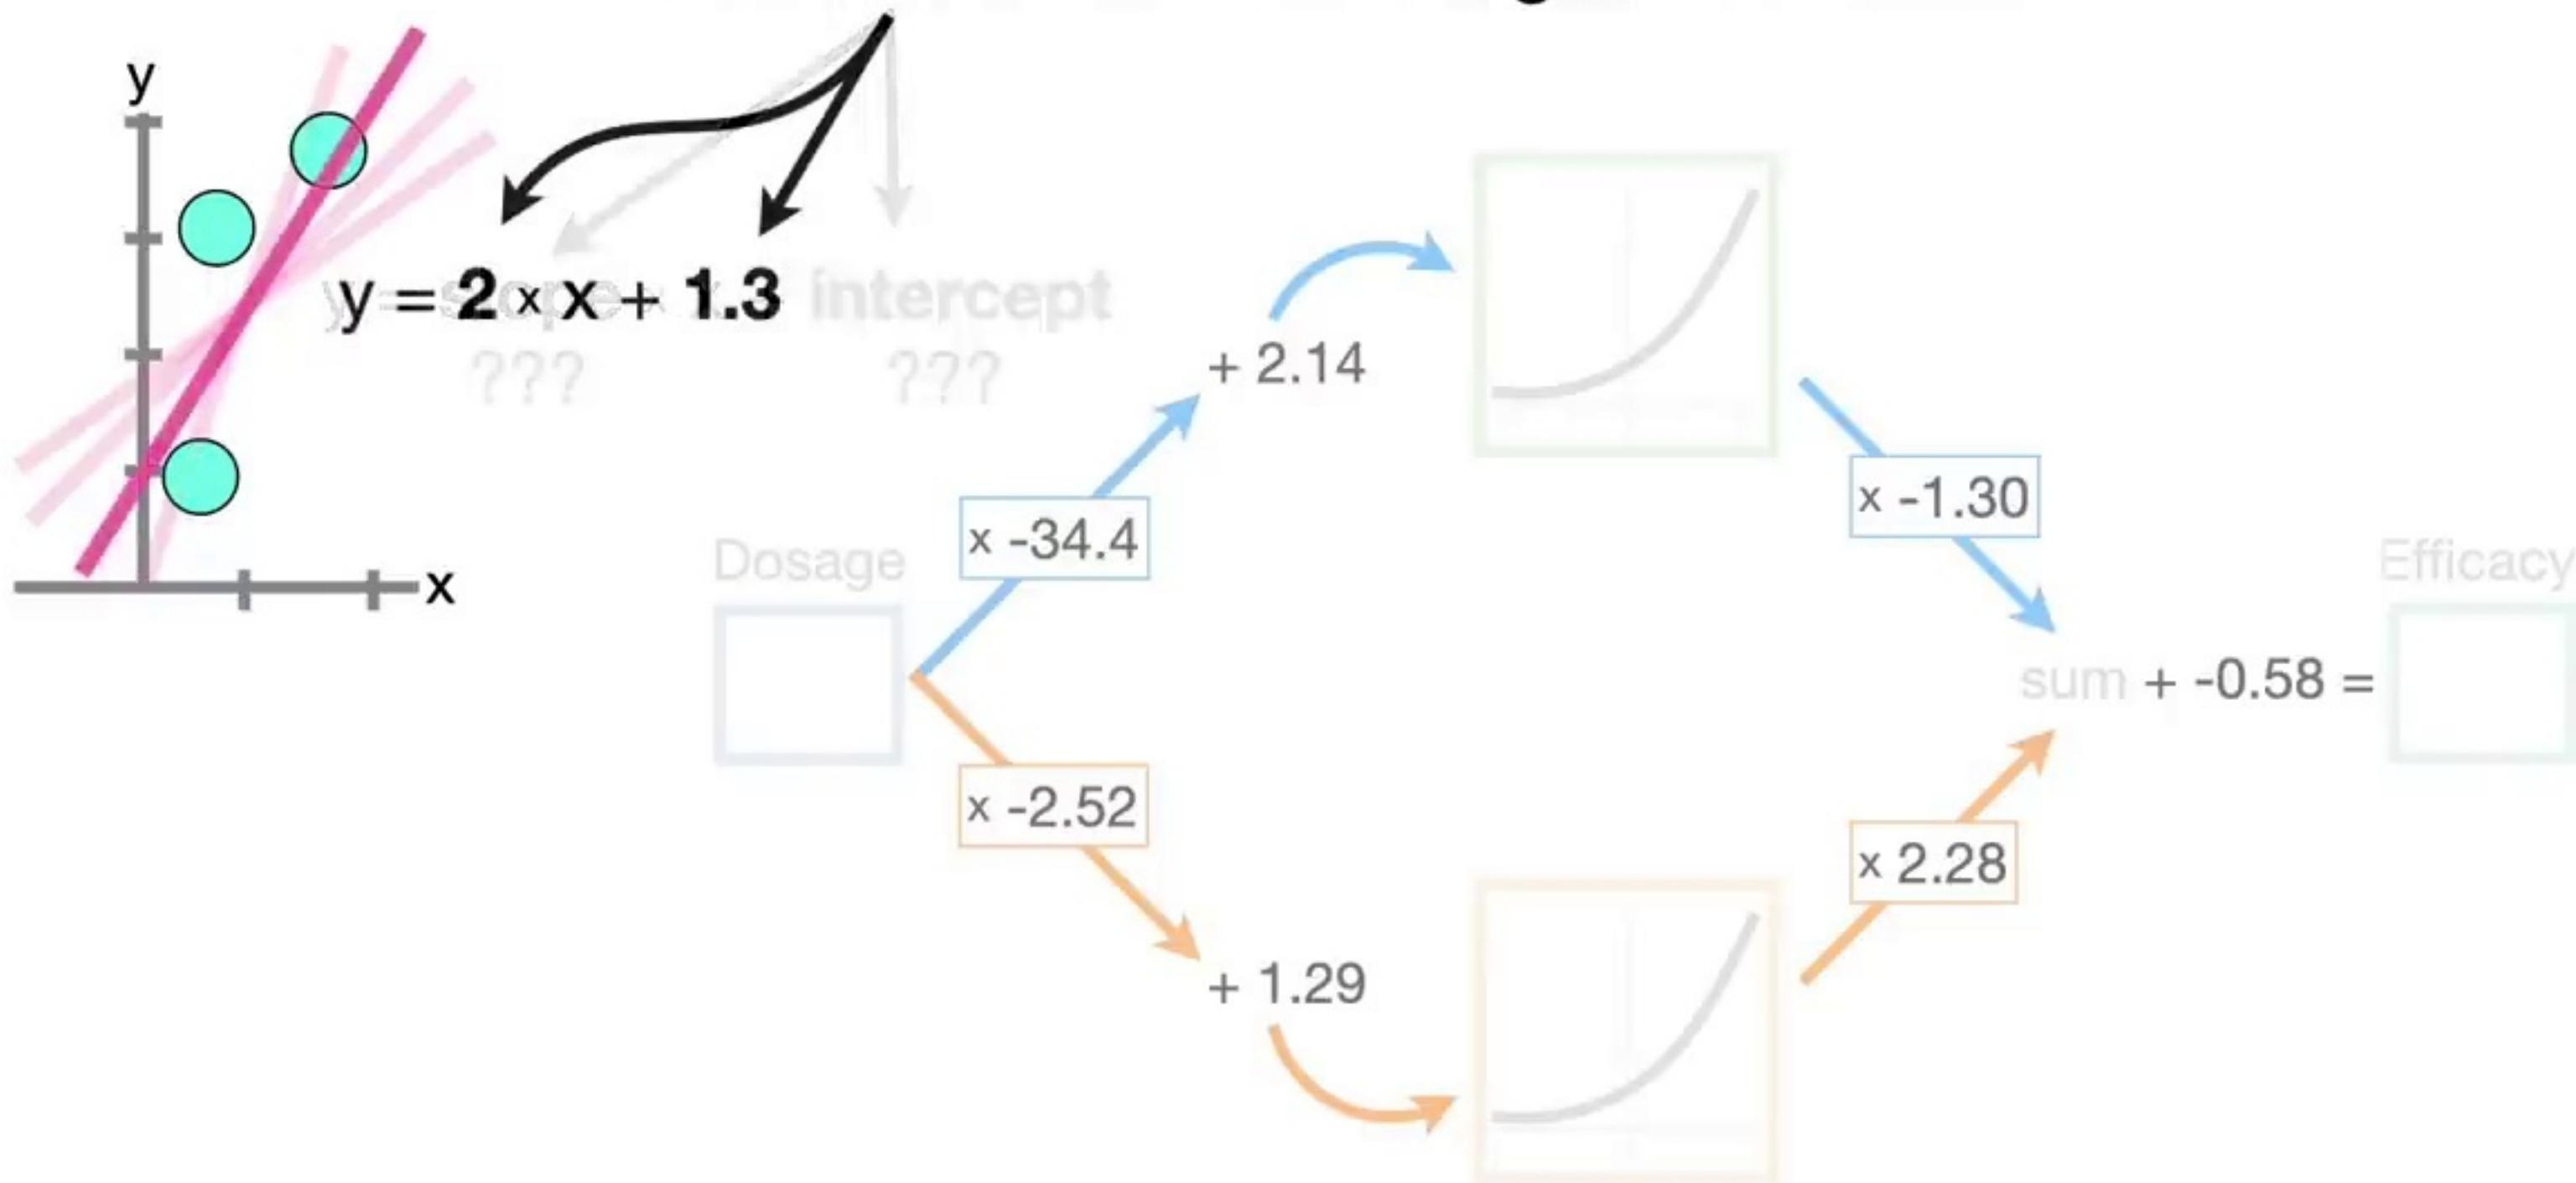


**NOTE:** The numbers along each connection represent parameter values that were estimated when this **Neural Network** was fit to the data.



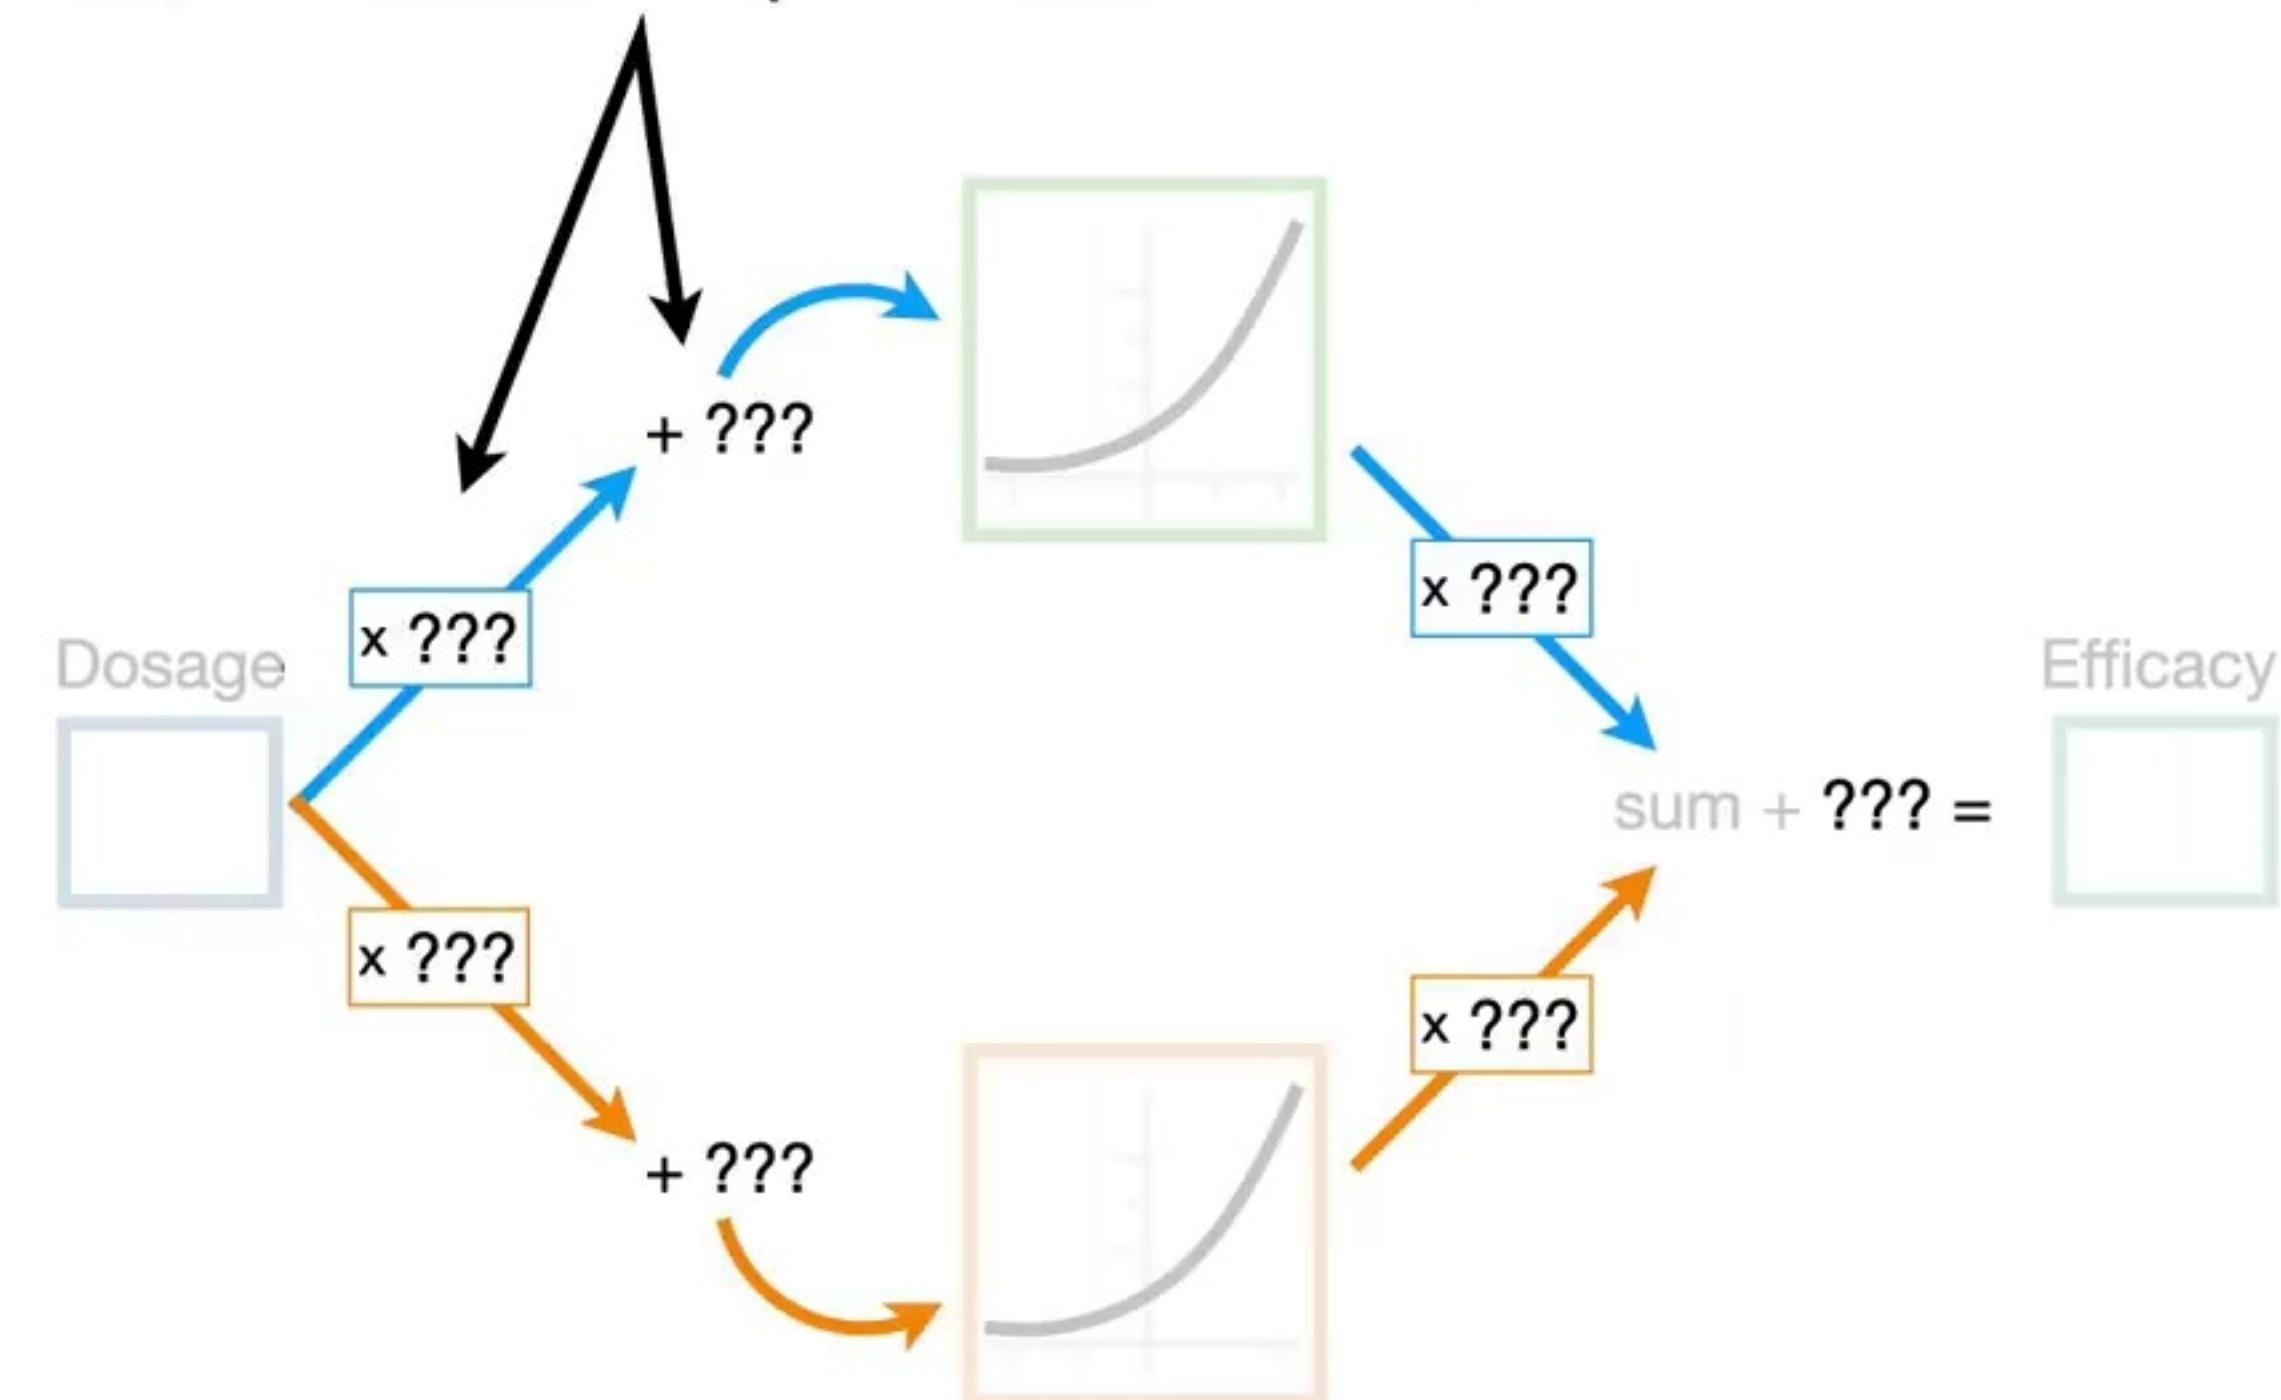


For now, just know that these parameter estimates are analogous to the **slope** and **intercept** values that we solve for when we fit a **straight line** to data.



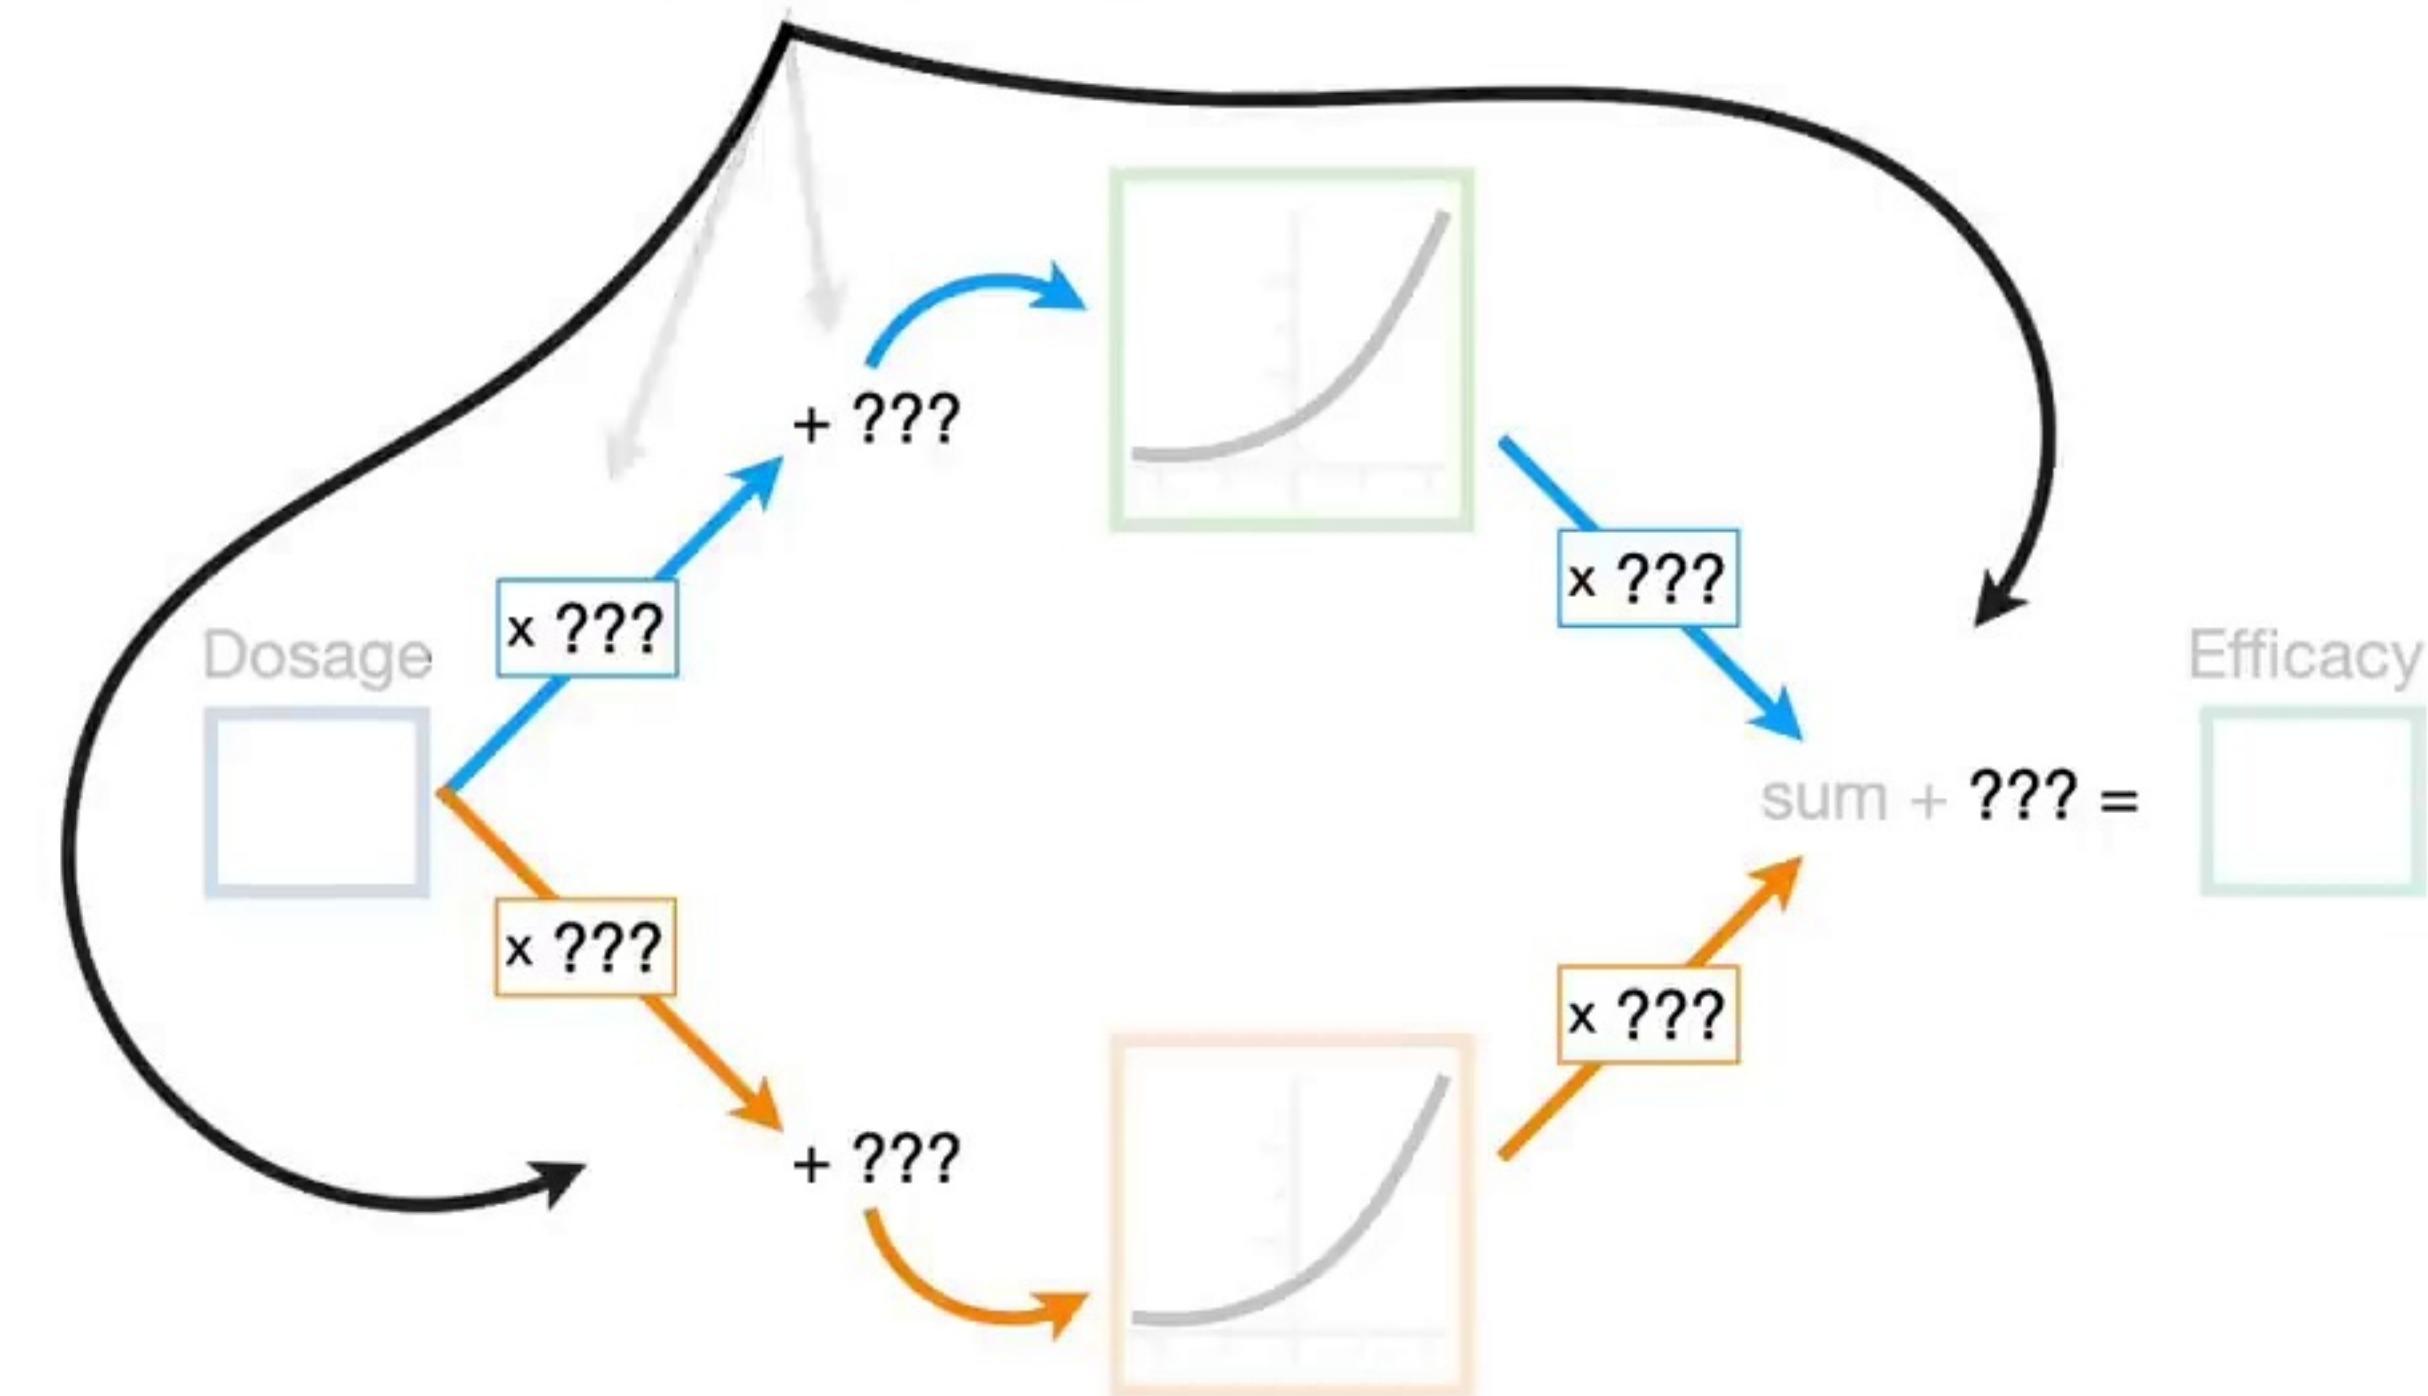


Likewise, a **Neural Network** starts out with unknown parameter values...



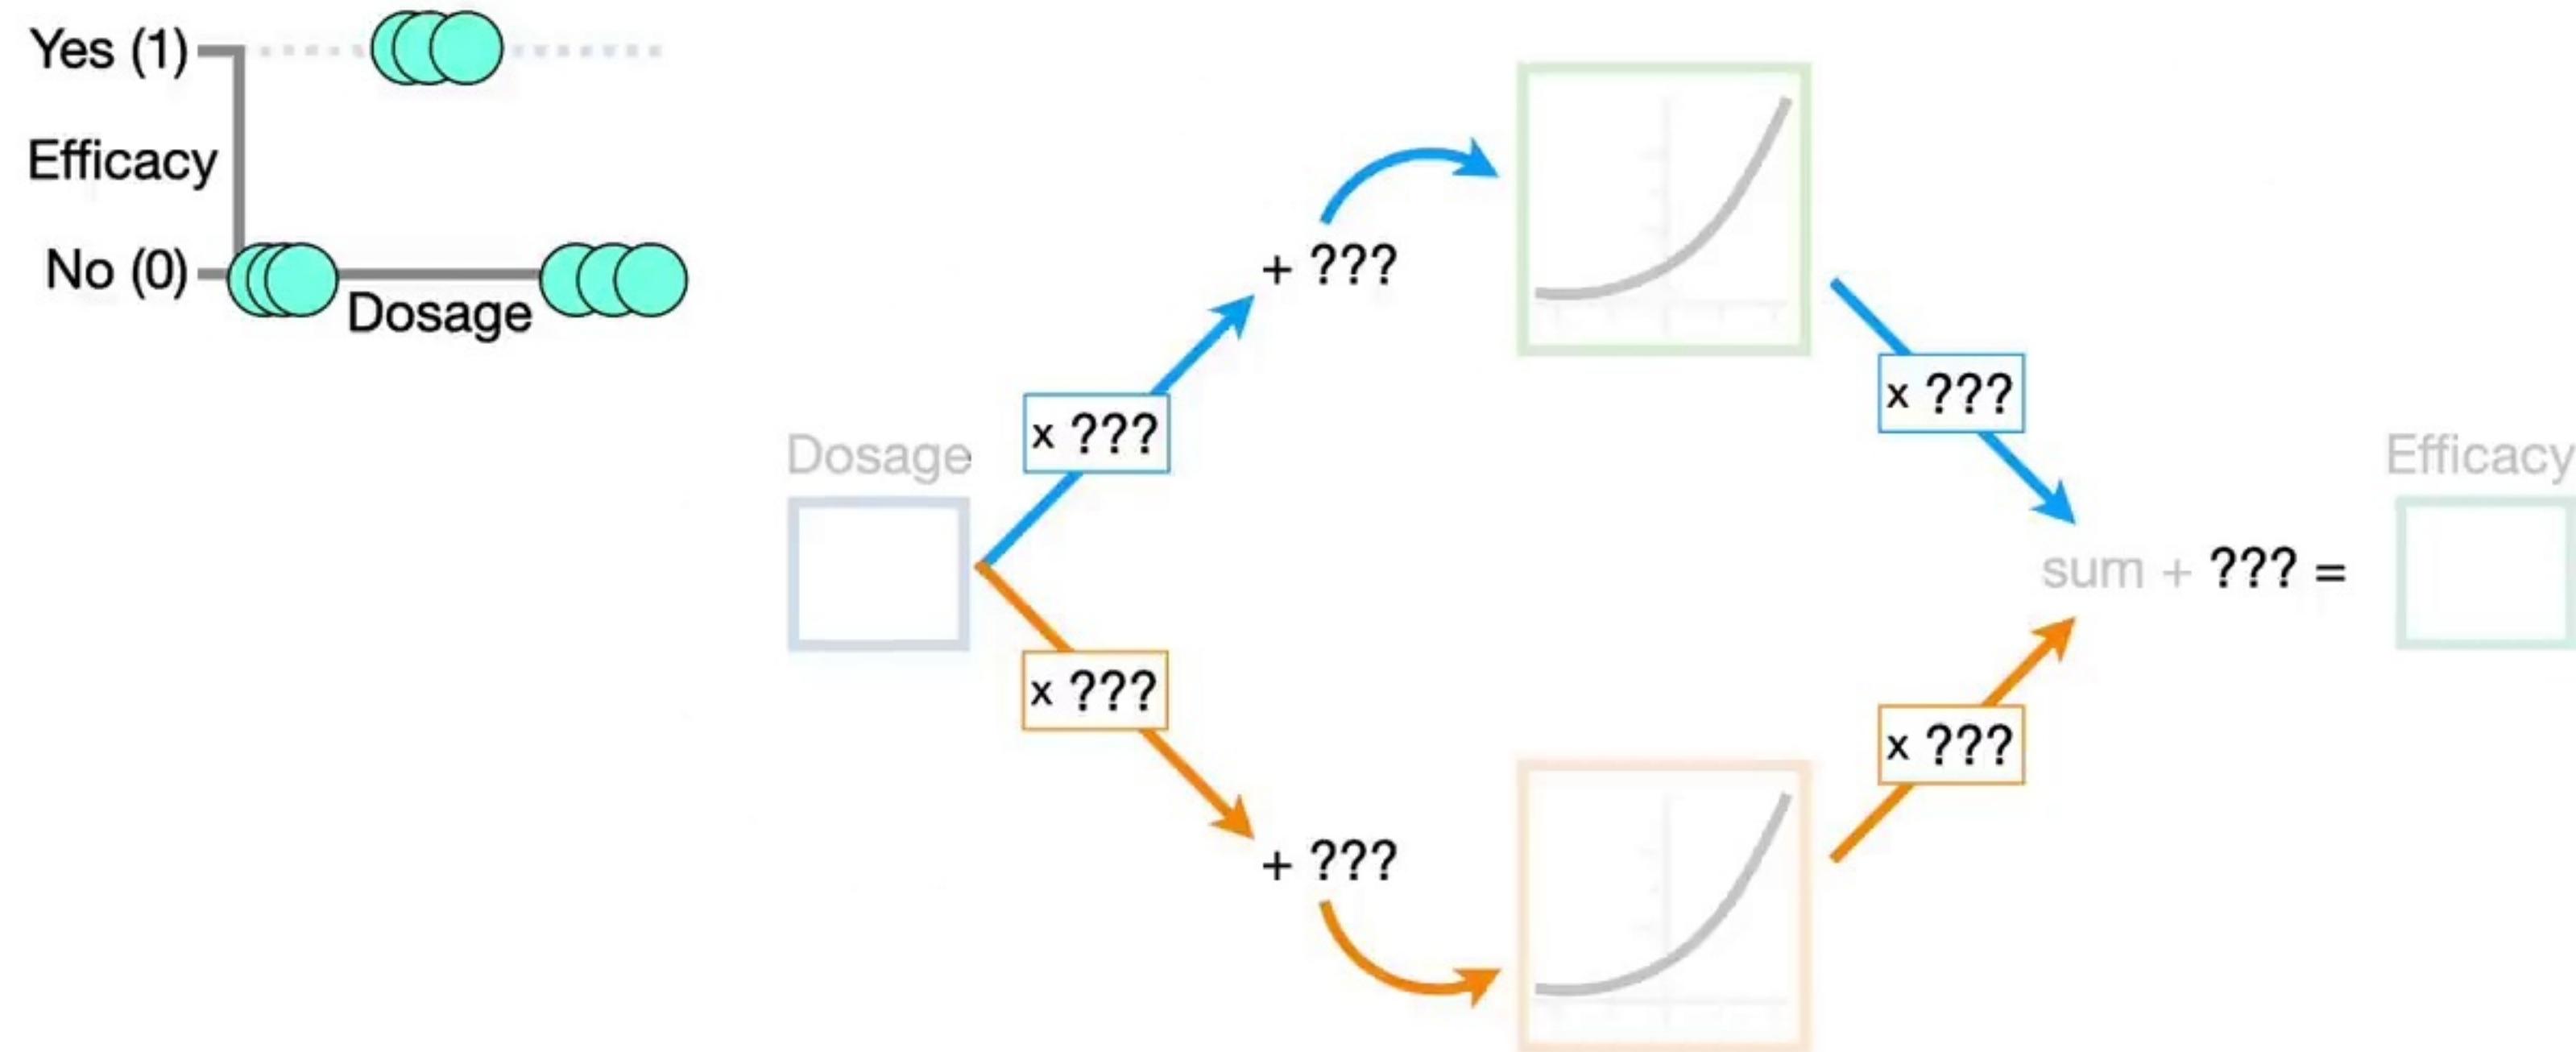


Likewise, a **Neural Network** starts out with unknown parameter values...



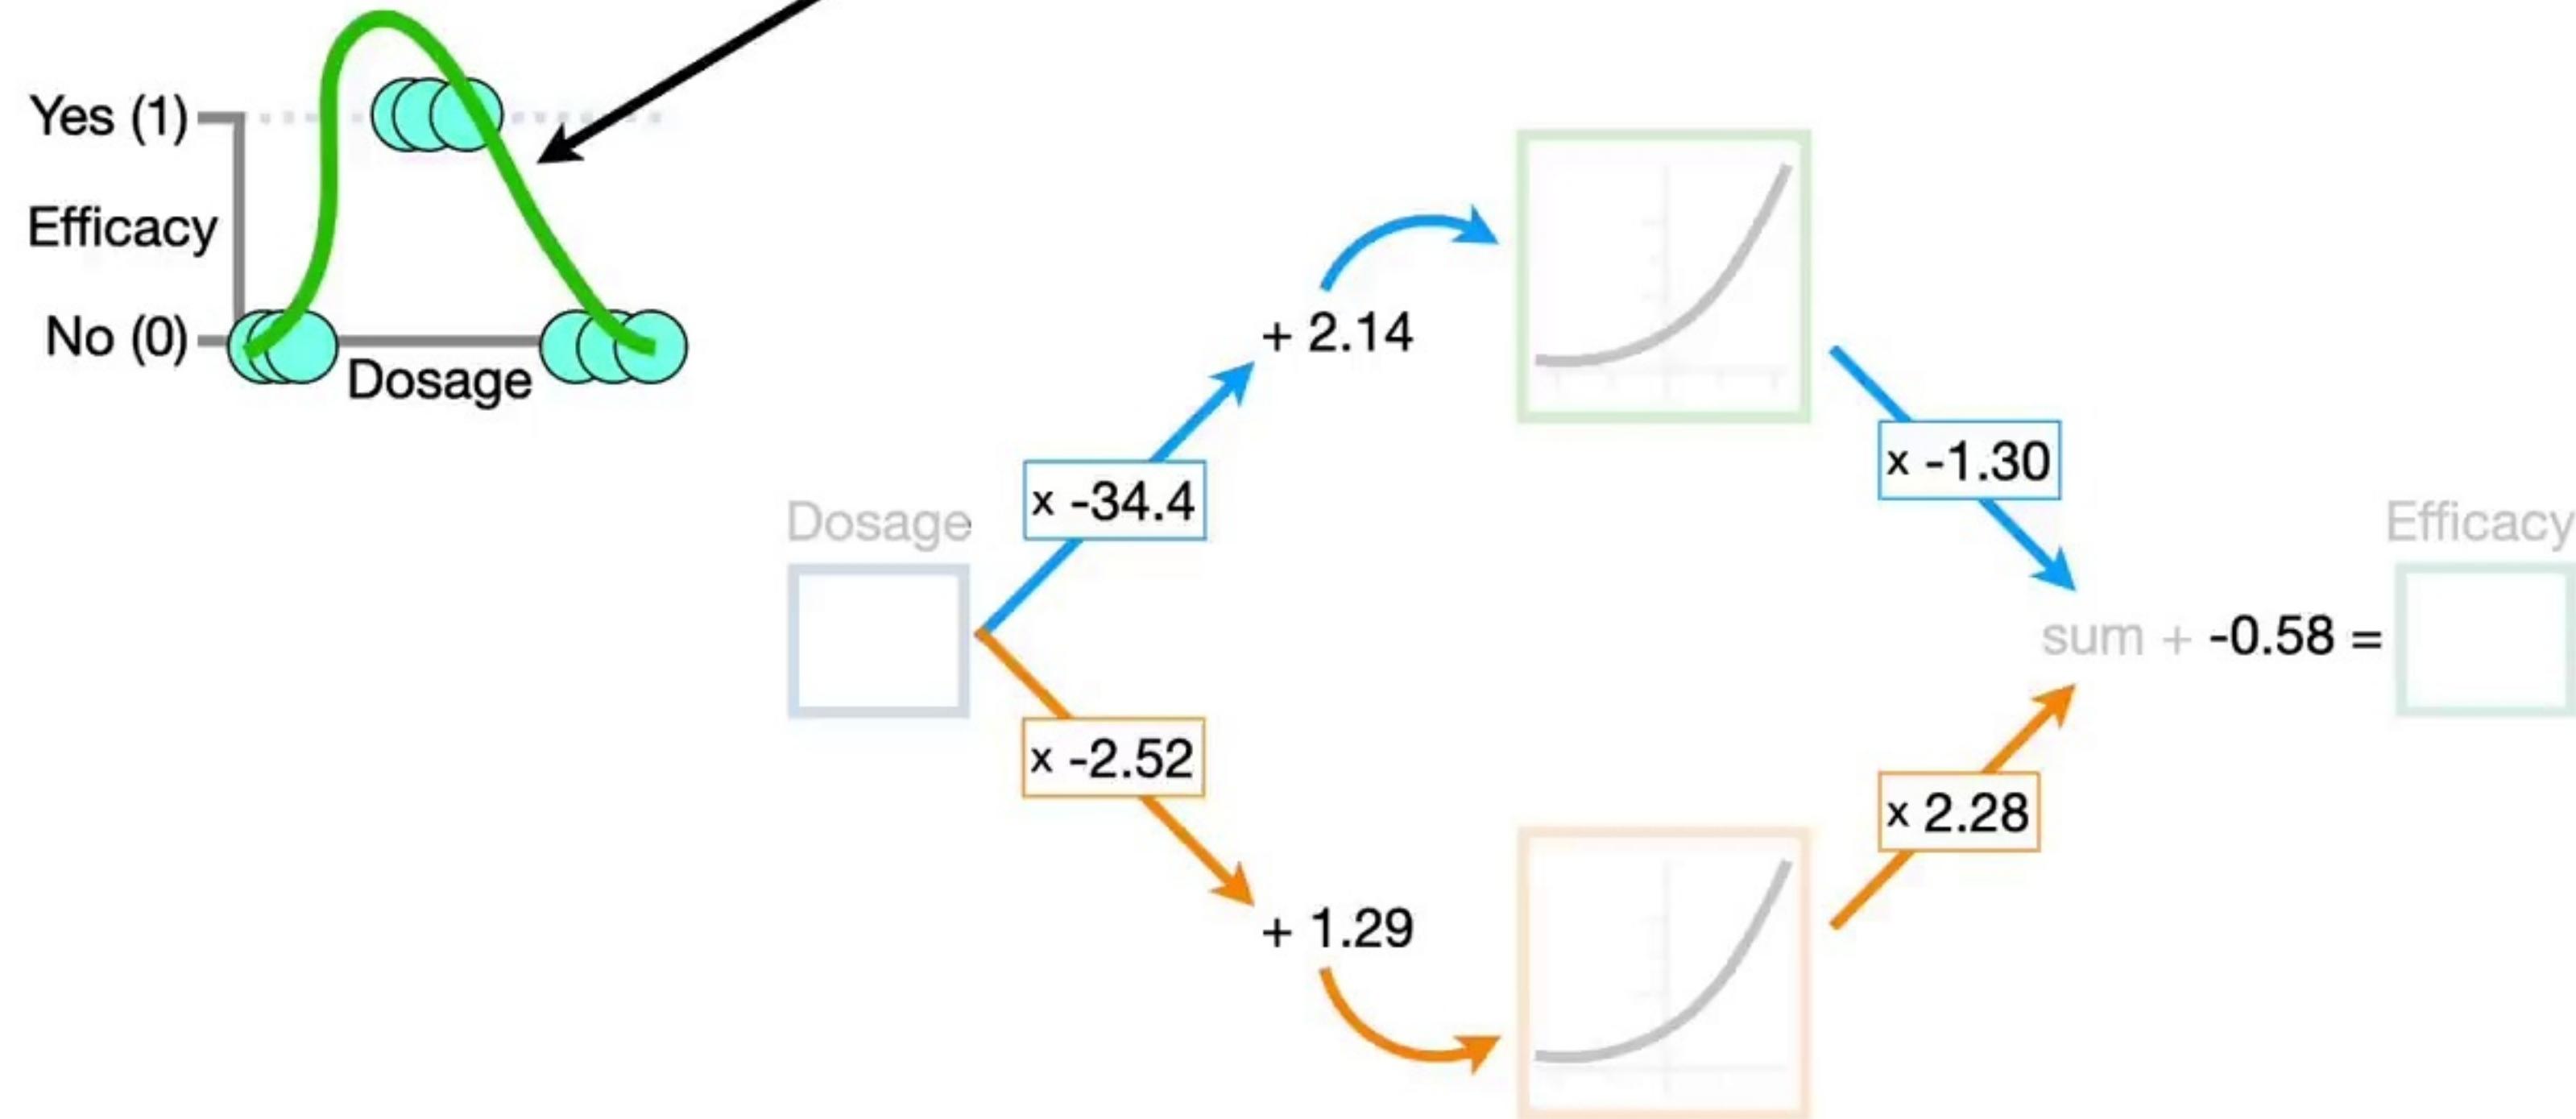


...that are estimated when we fit the **Neural Network** to a dataset using a method called **Backpropagation**.



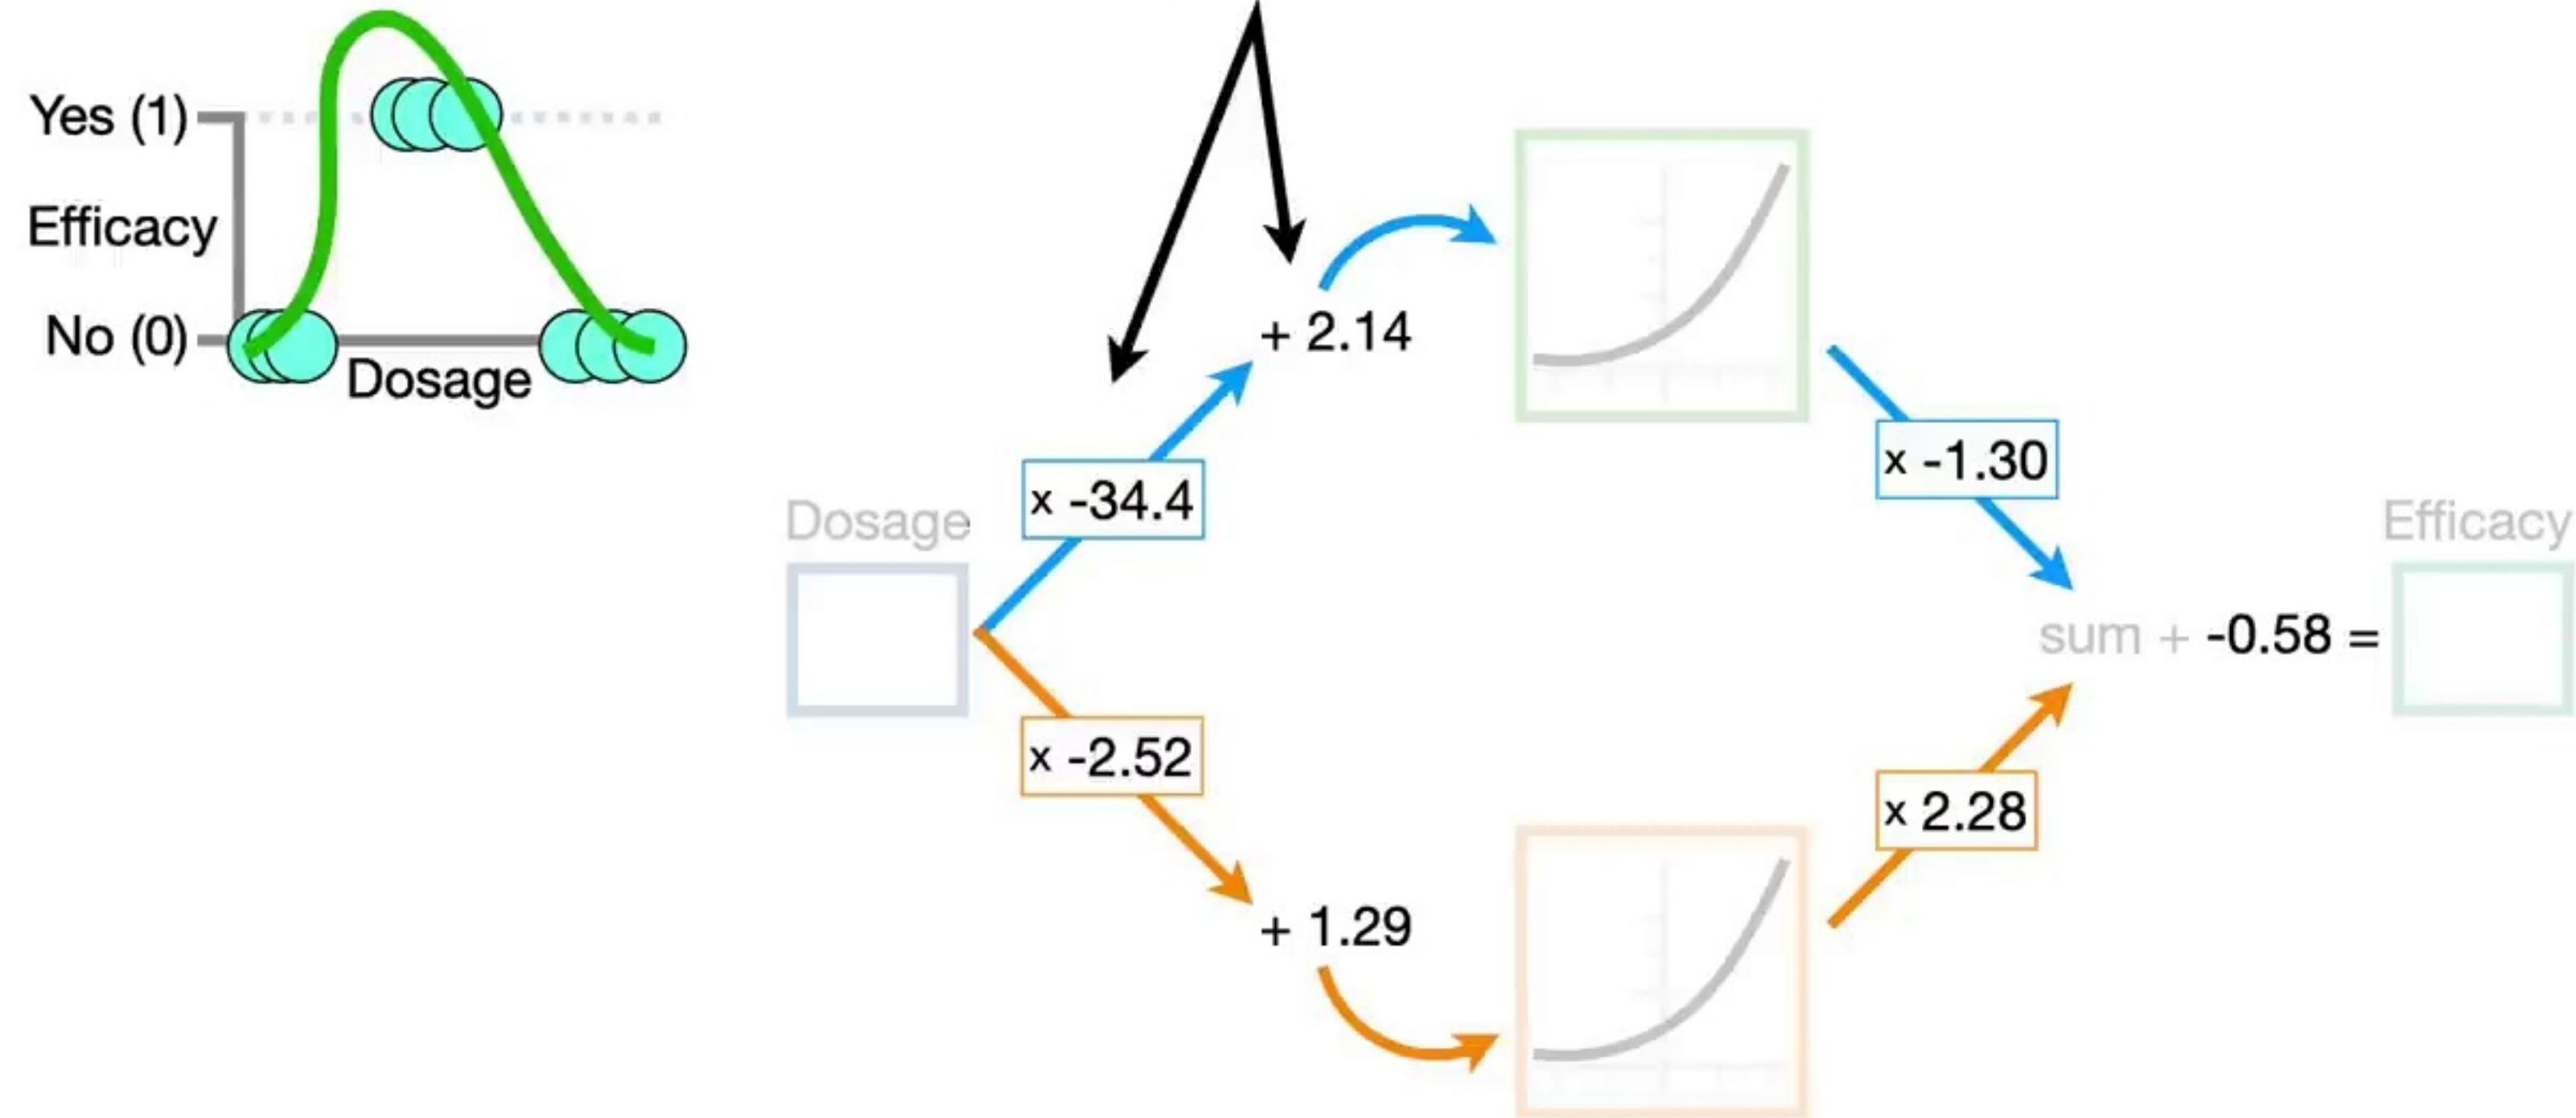


...but for now, just assume that  
we've already fit this **Neural**  
**Network** to this specific dataset...



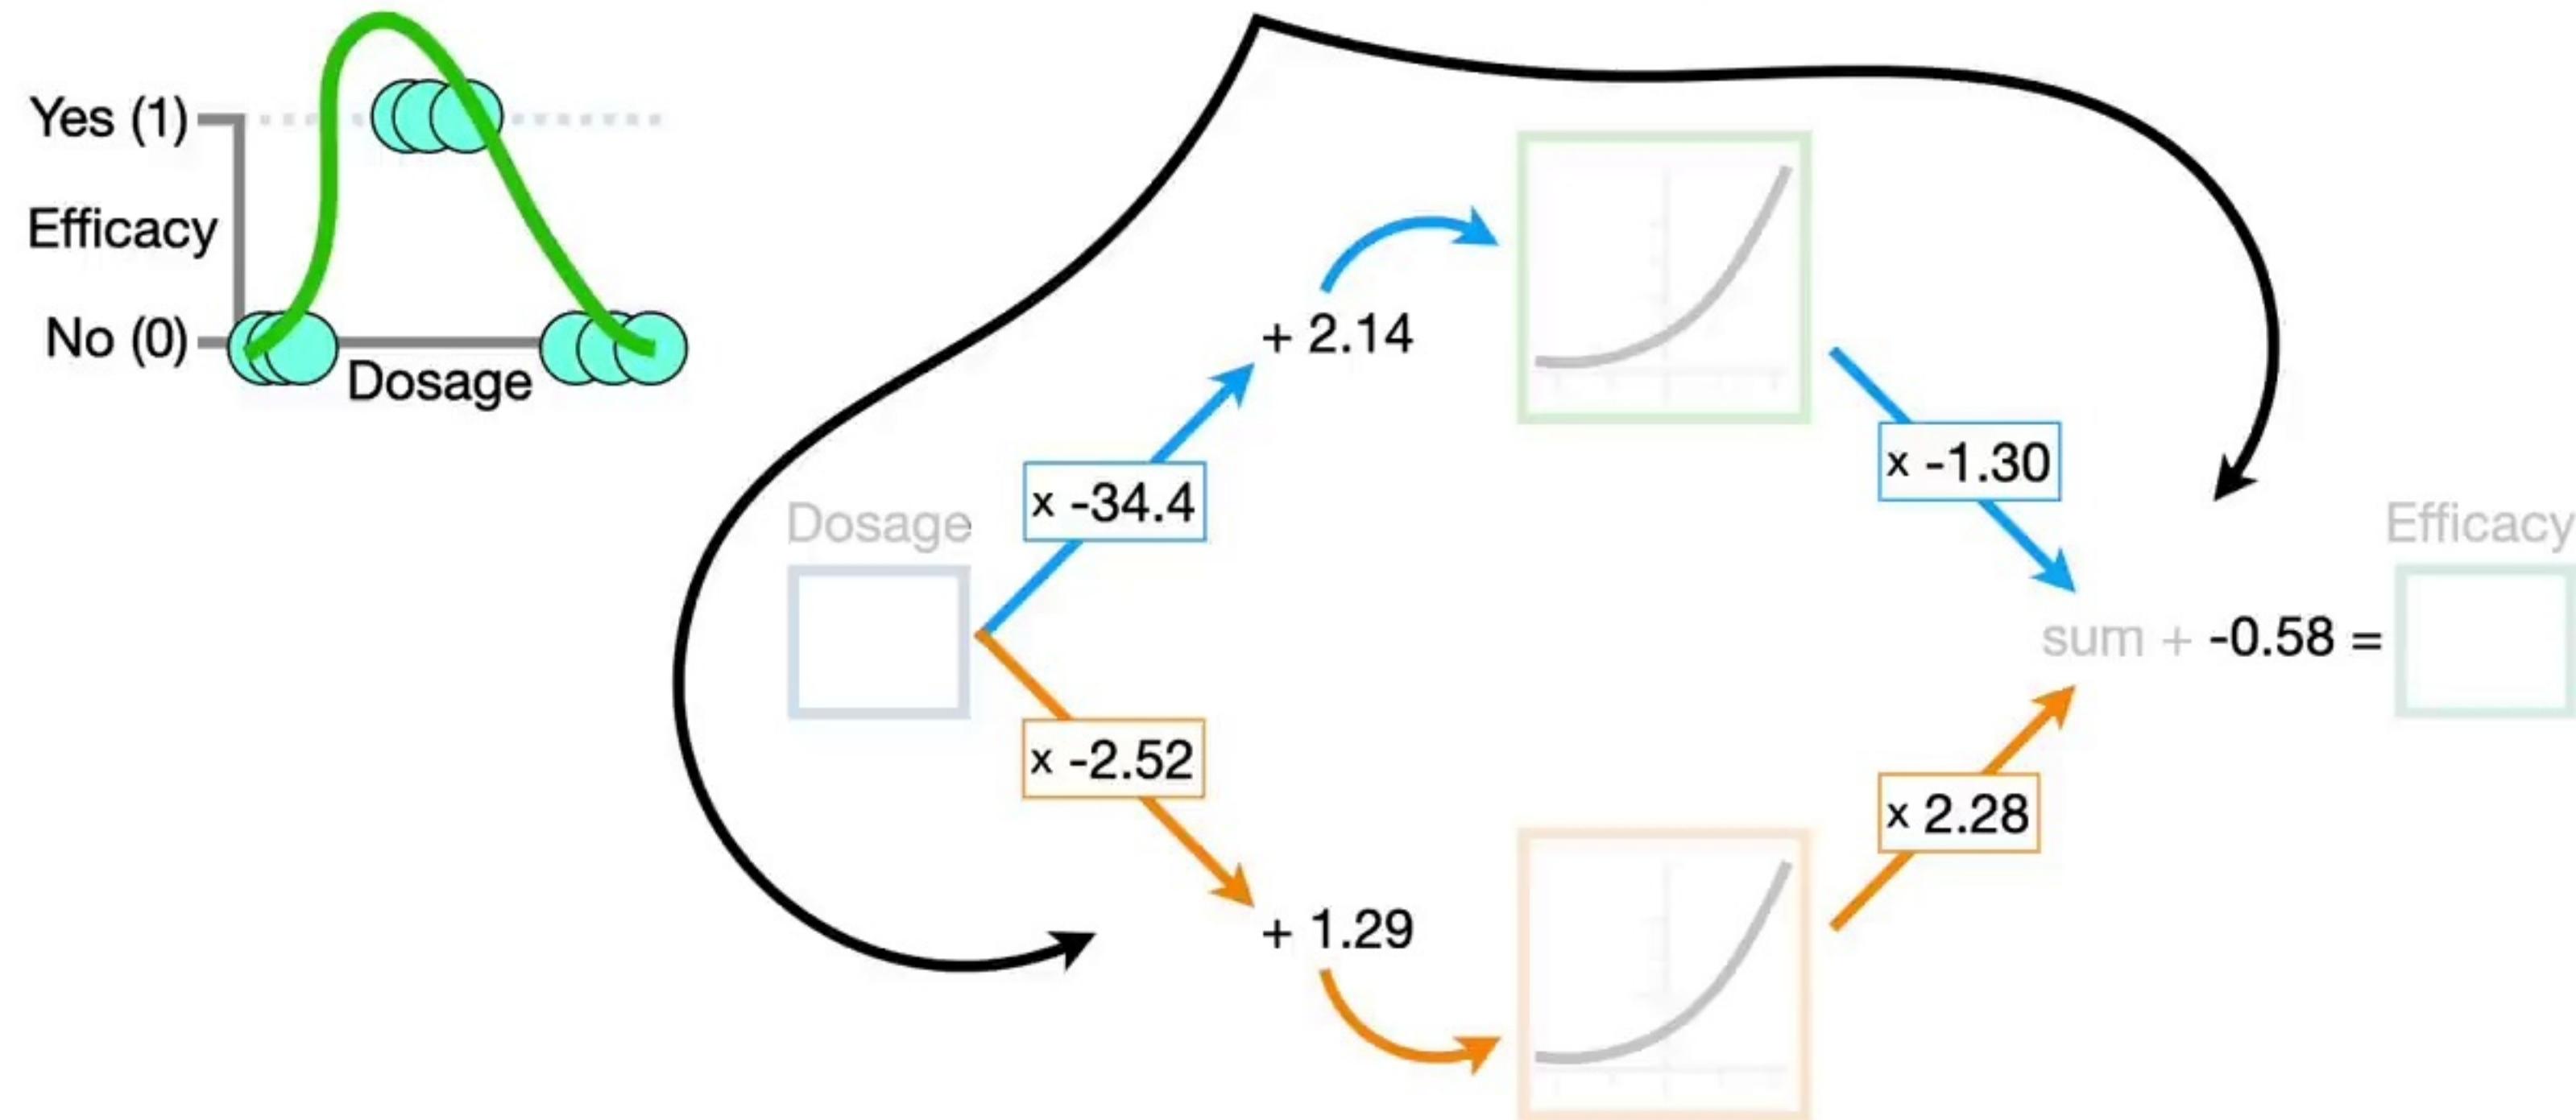


...and that means we have already estimated these parameters.



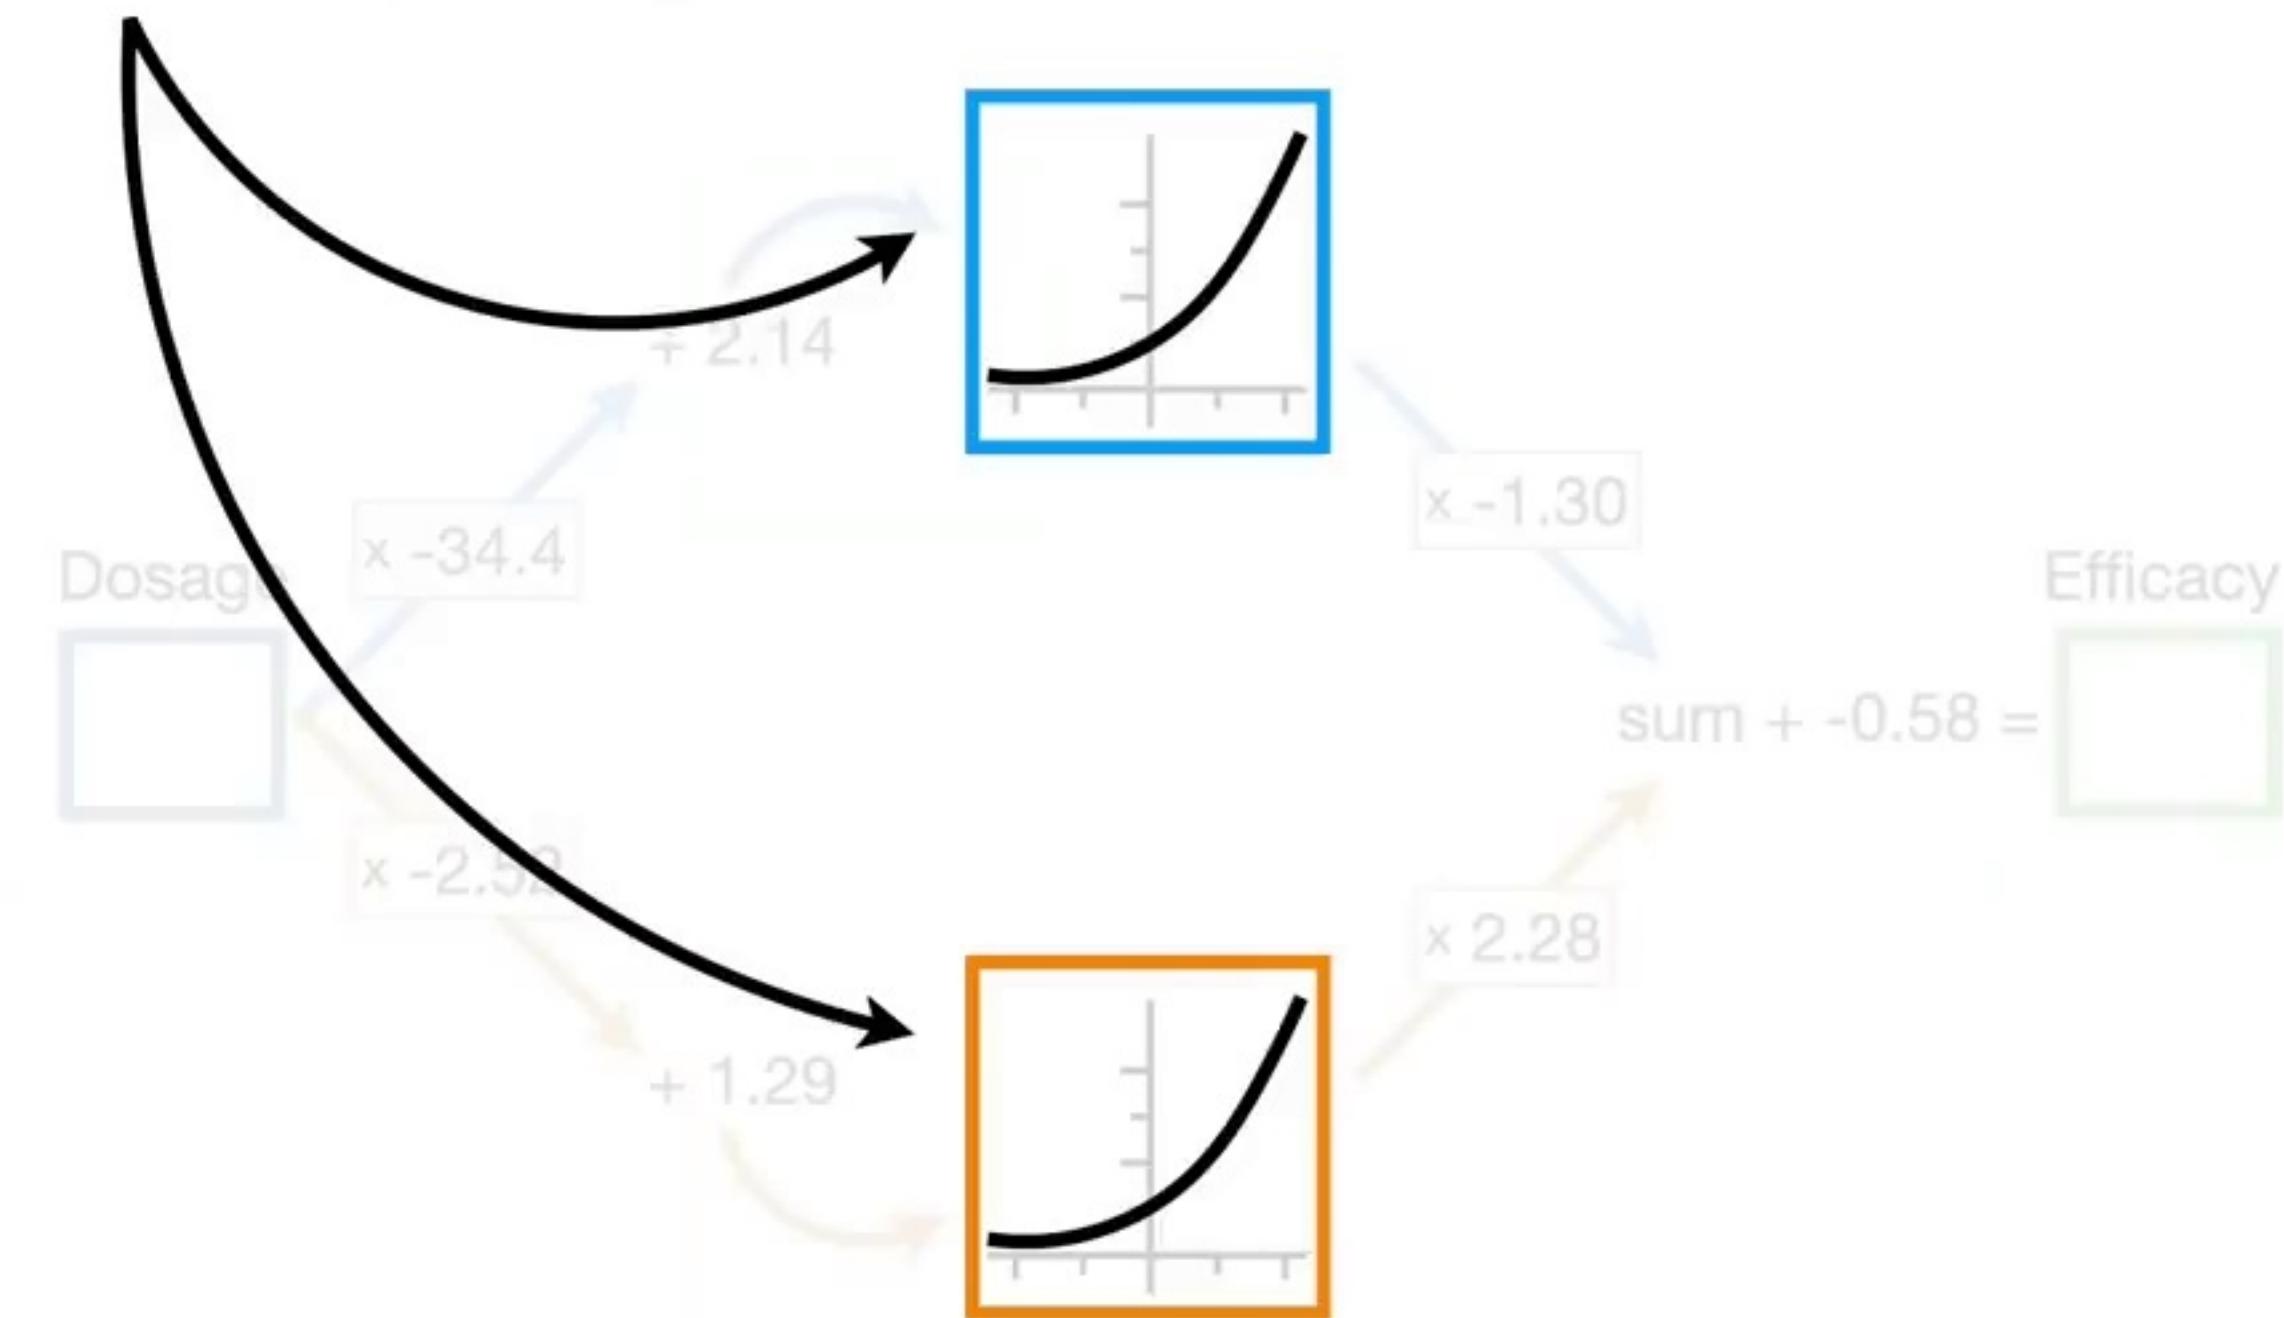


...and that means we have already estimated these parameters.



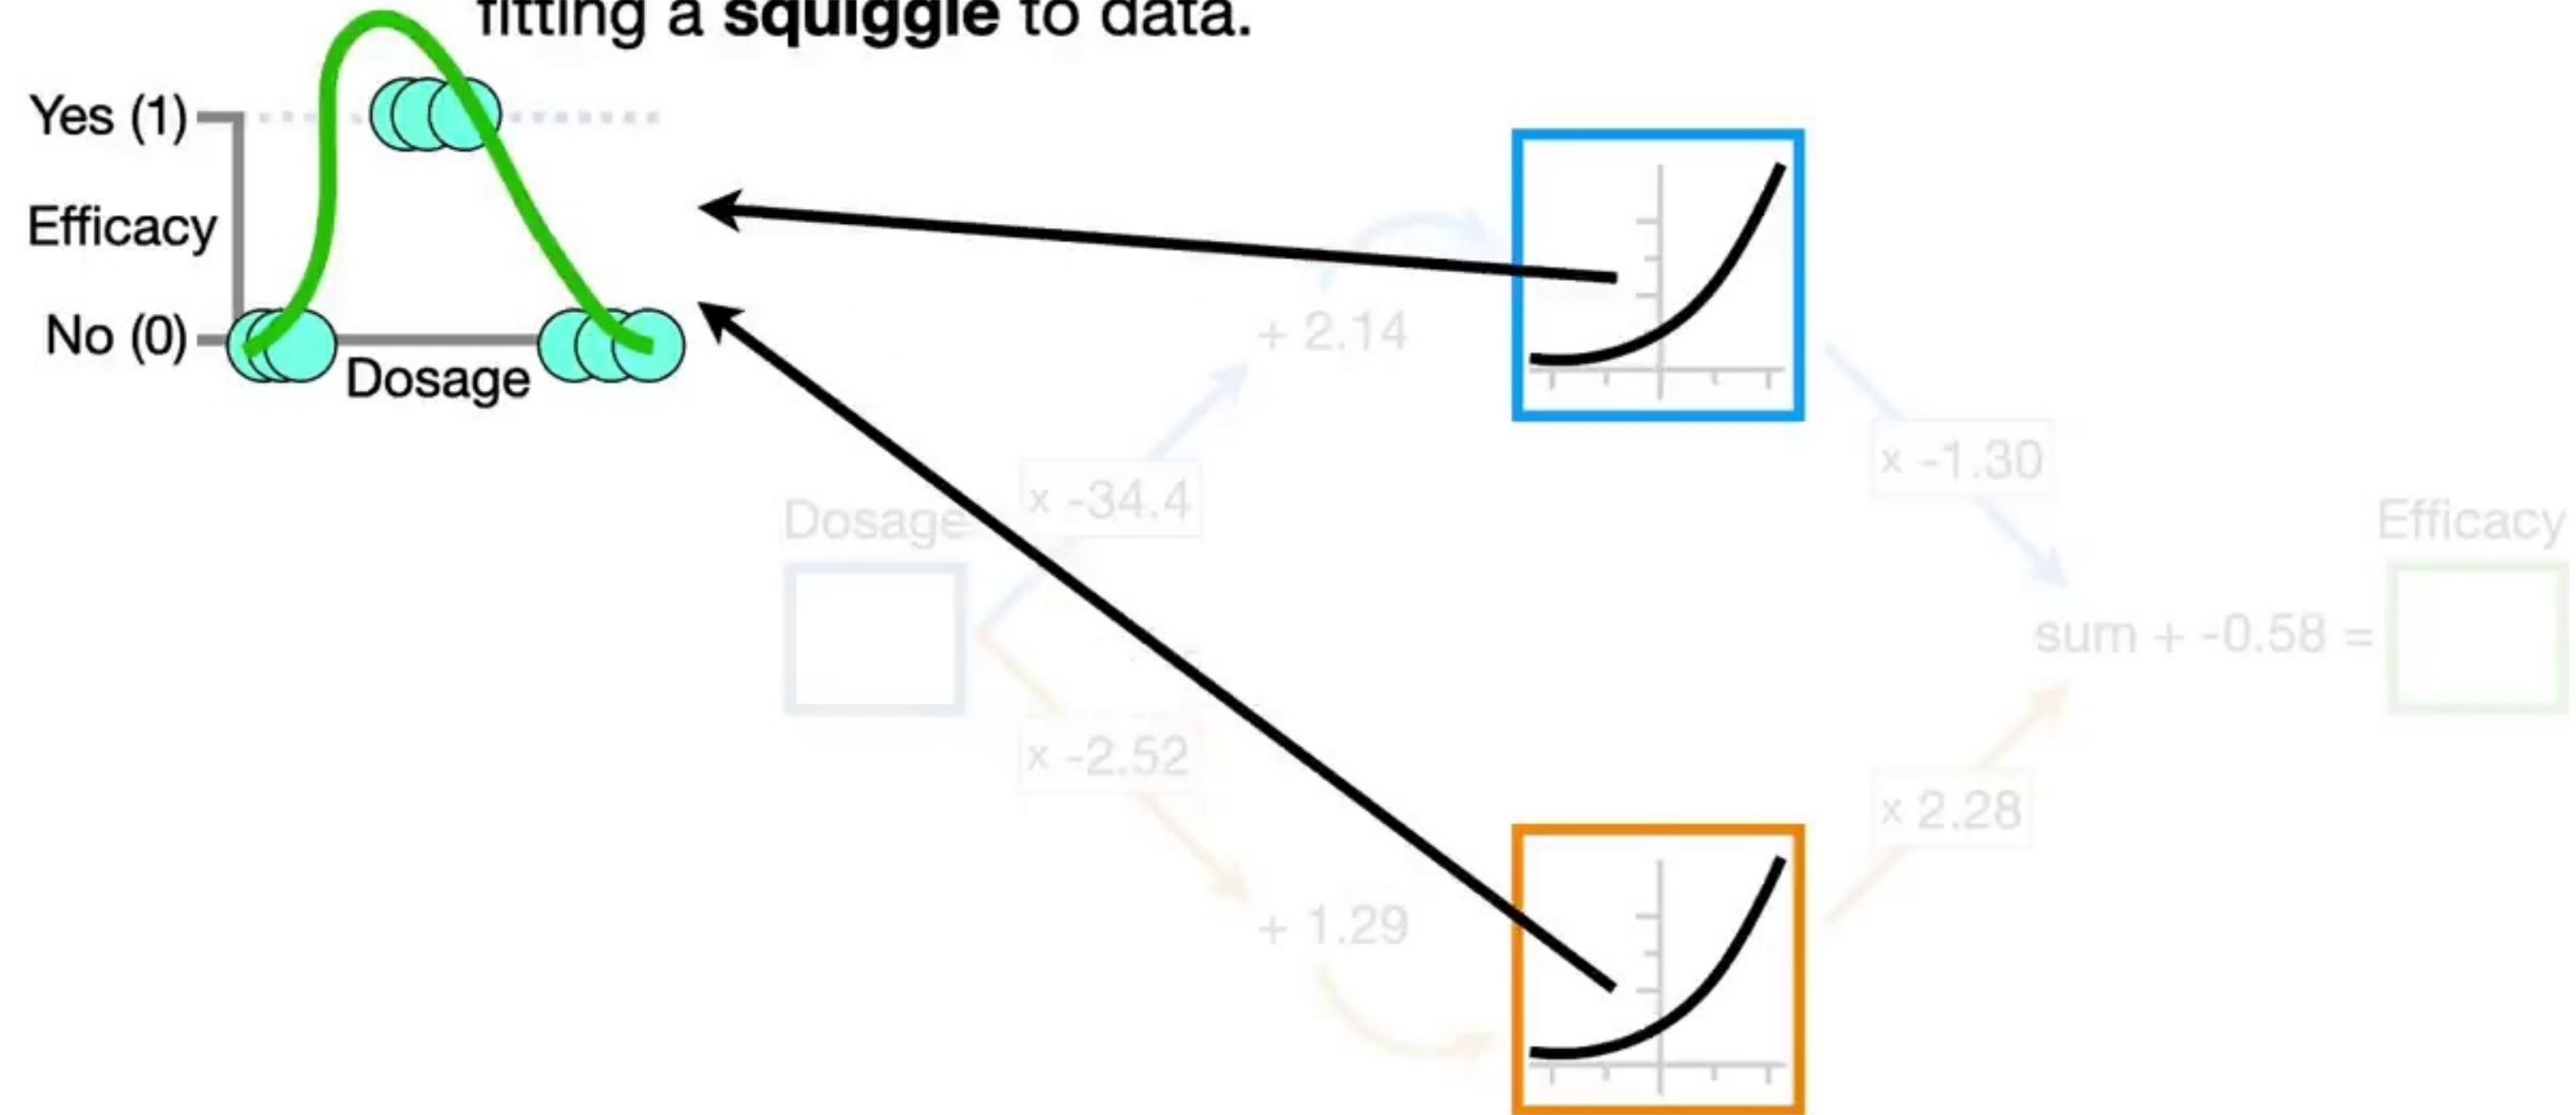


Also, you may have noticed that some of the **Nodes** have **curved lines** inside of them.



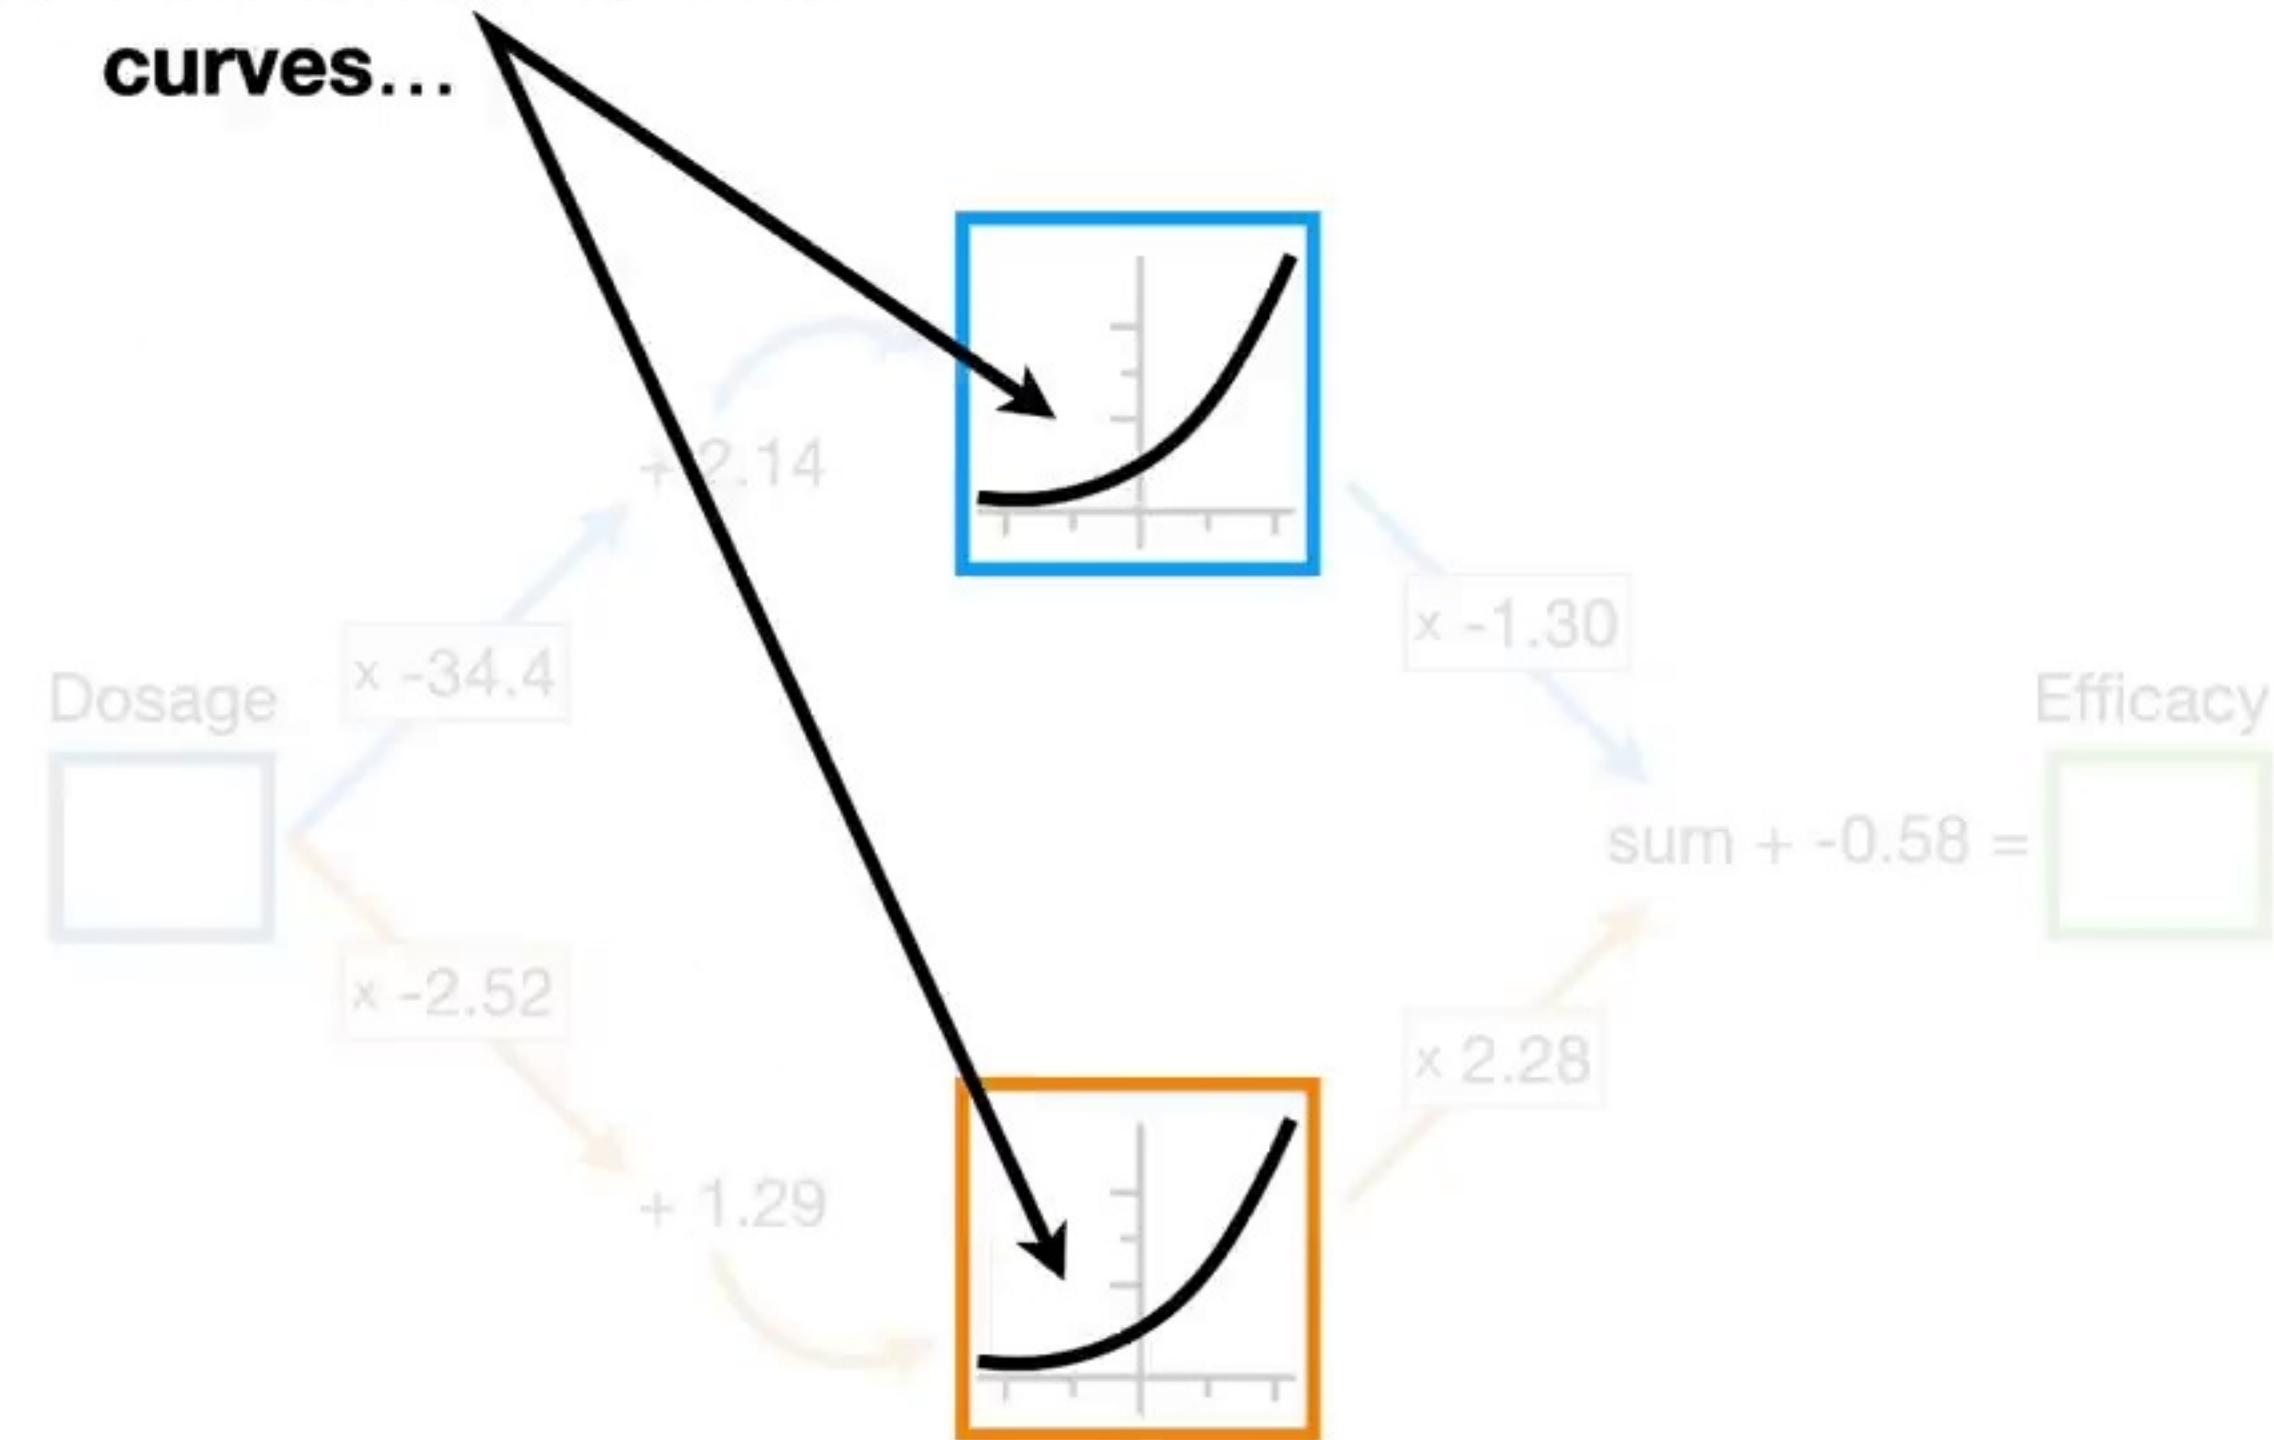
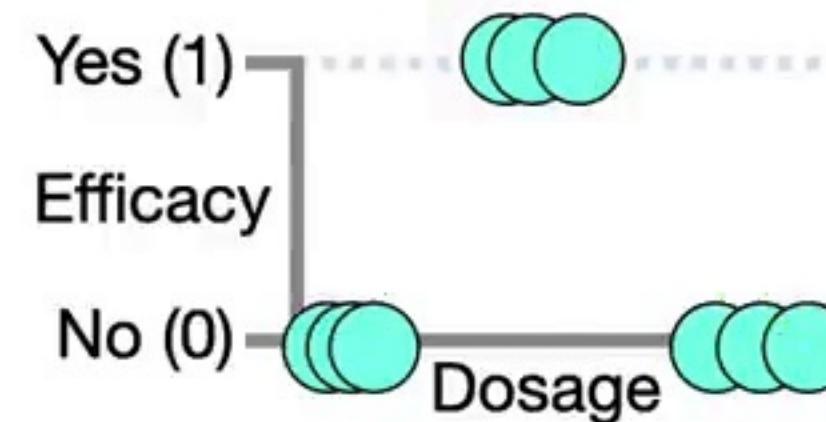


These **bent** or **curved** lines  
are the building blocks for  
fitting a **squiggle** to data.



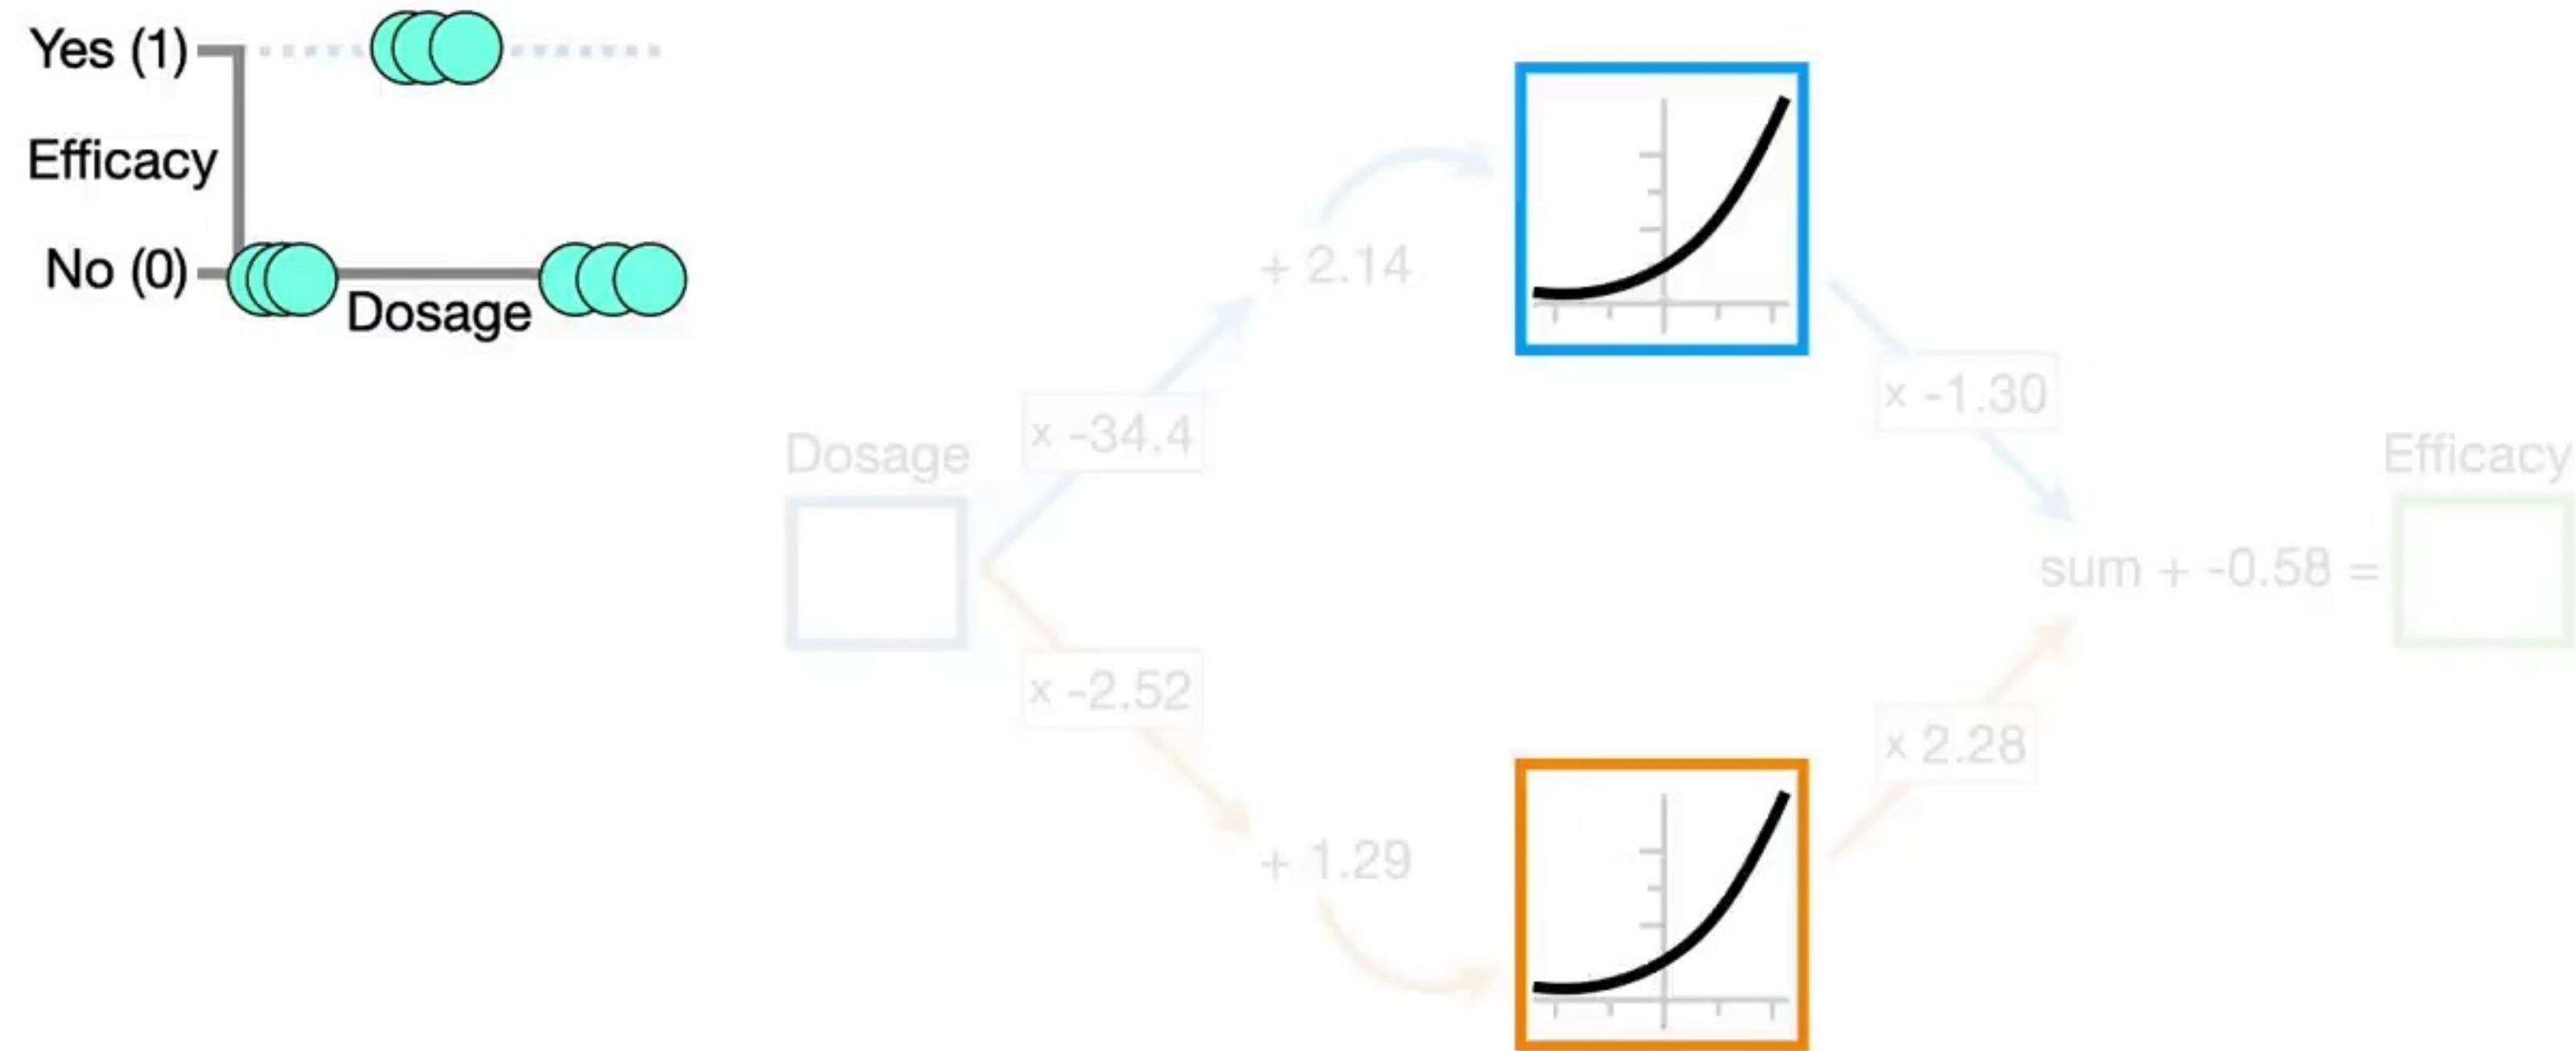


The goal of this **StatQuest** is to  
show you how these identical  
**curves...**



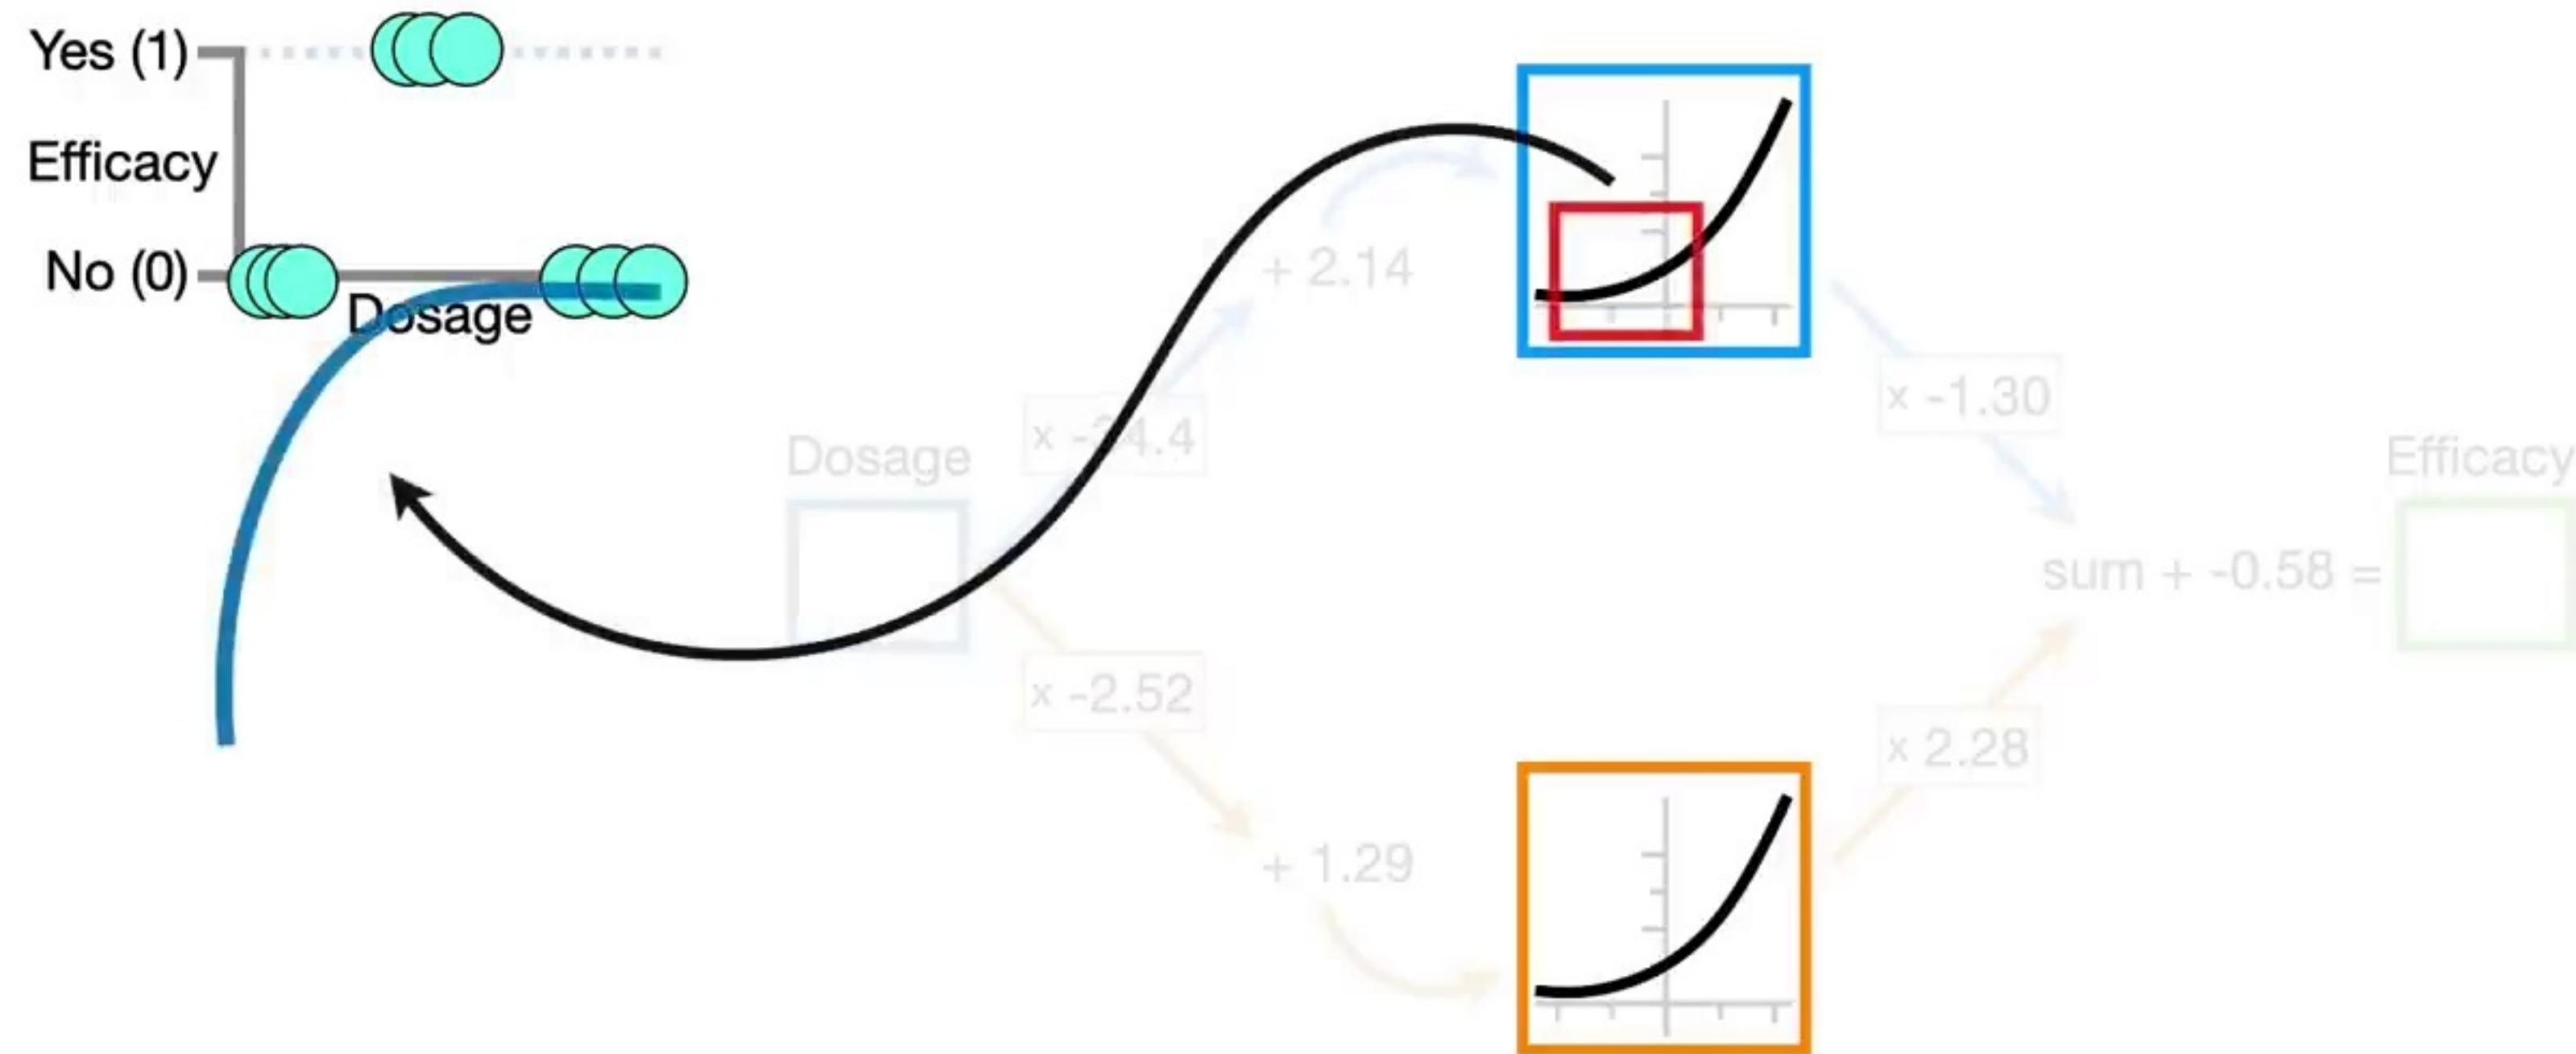


...can be reshaped by the parameter values...



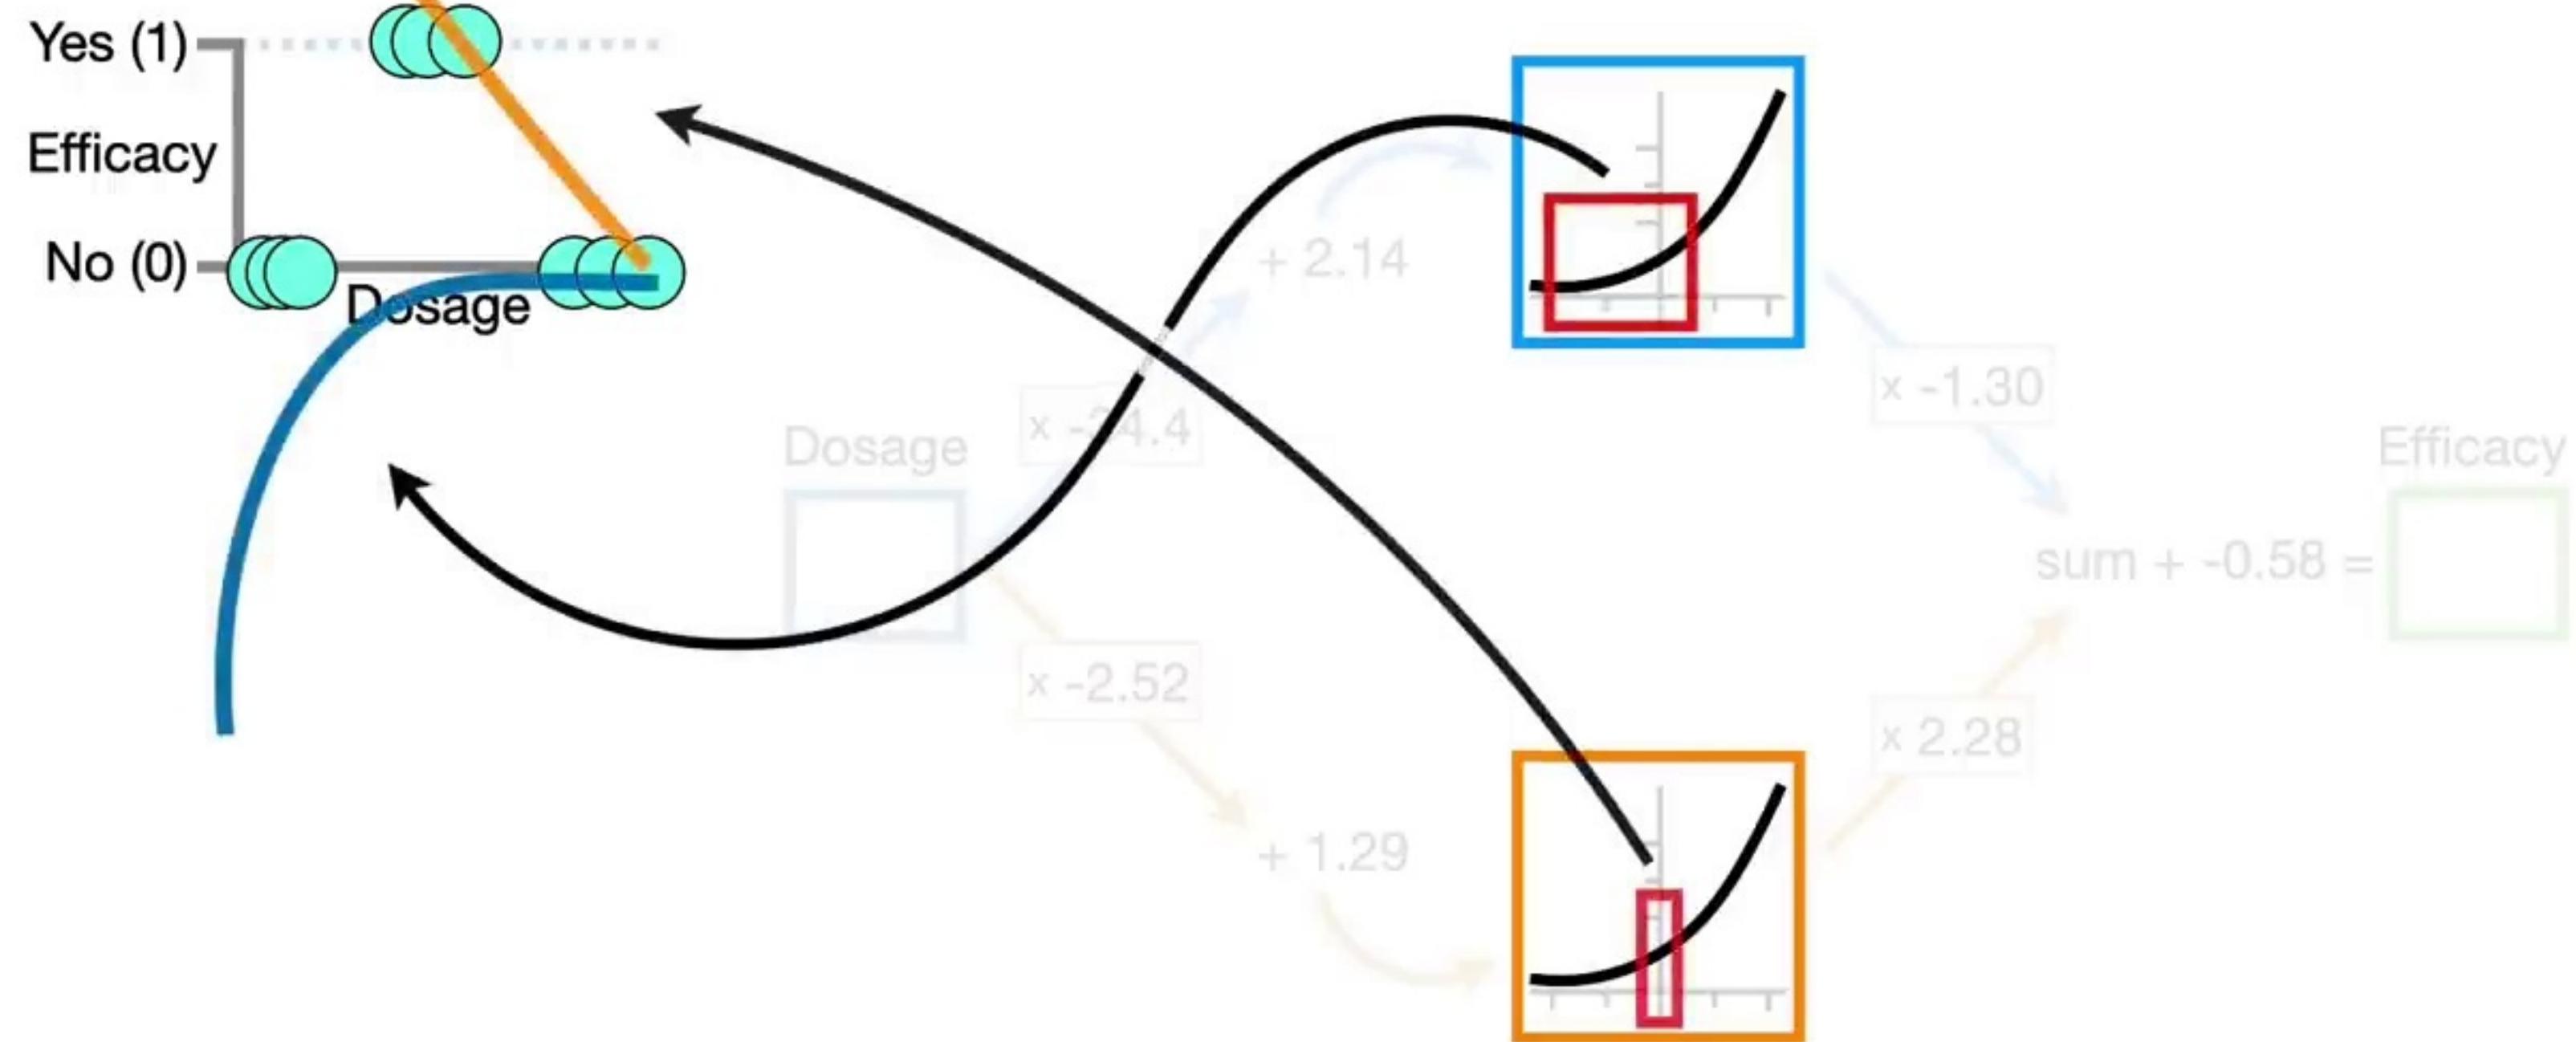


...can be reshaped by the parameter values...



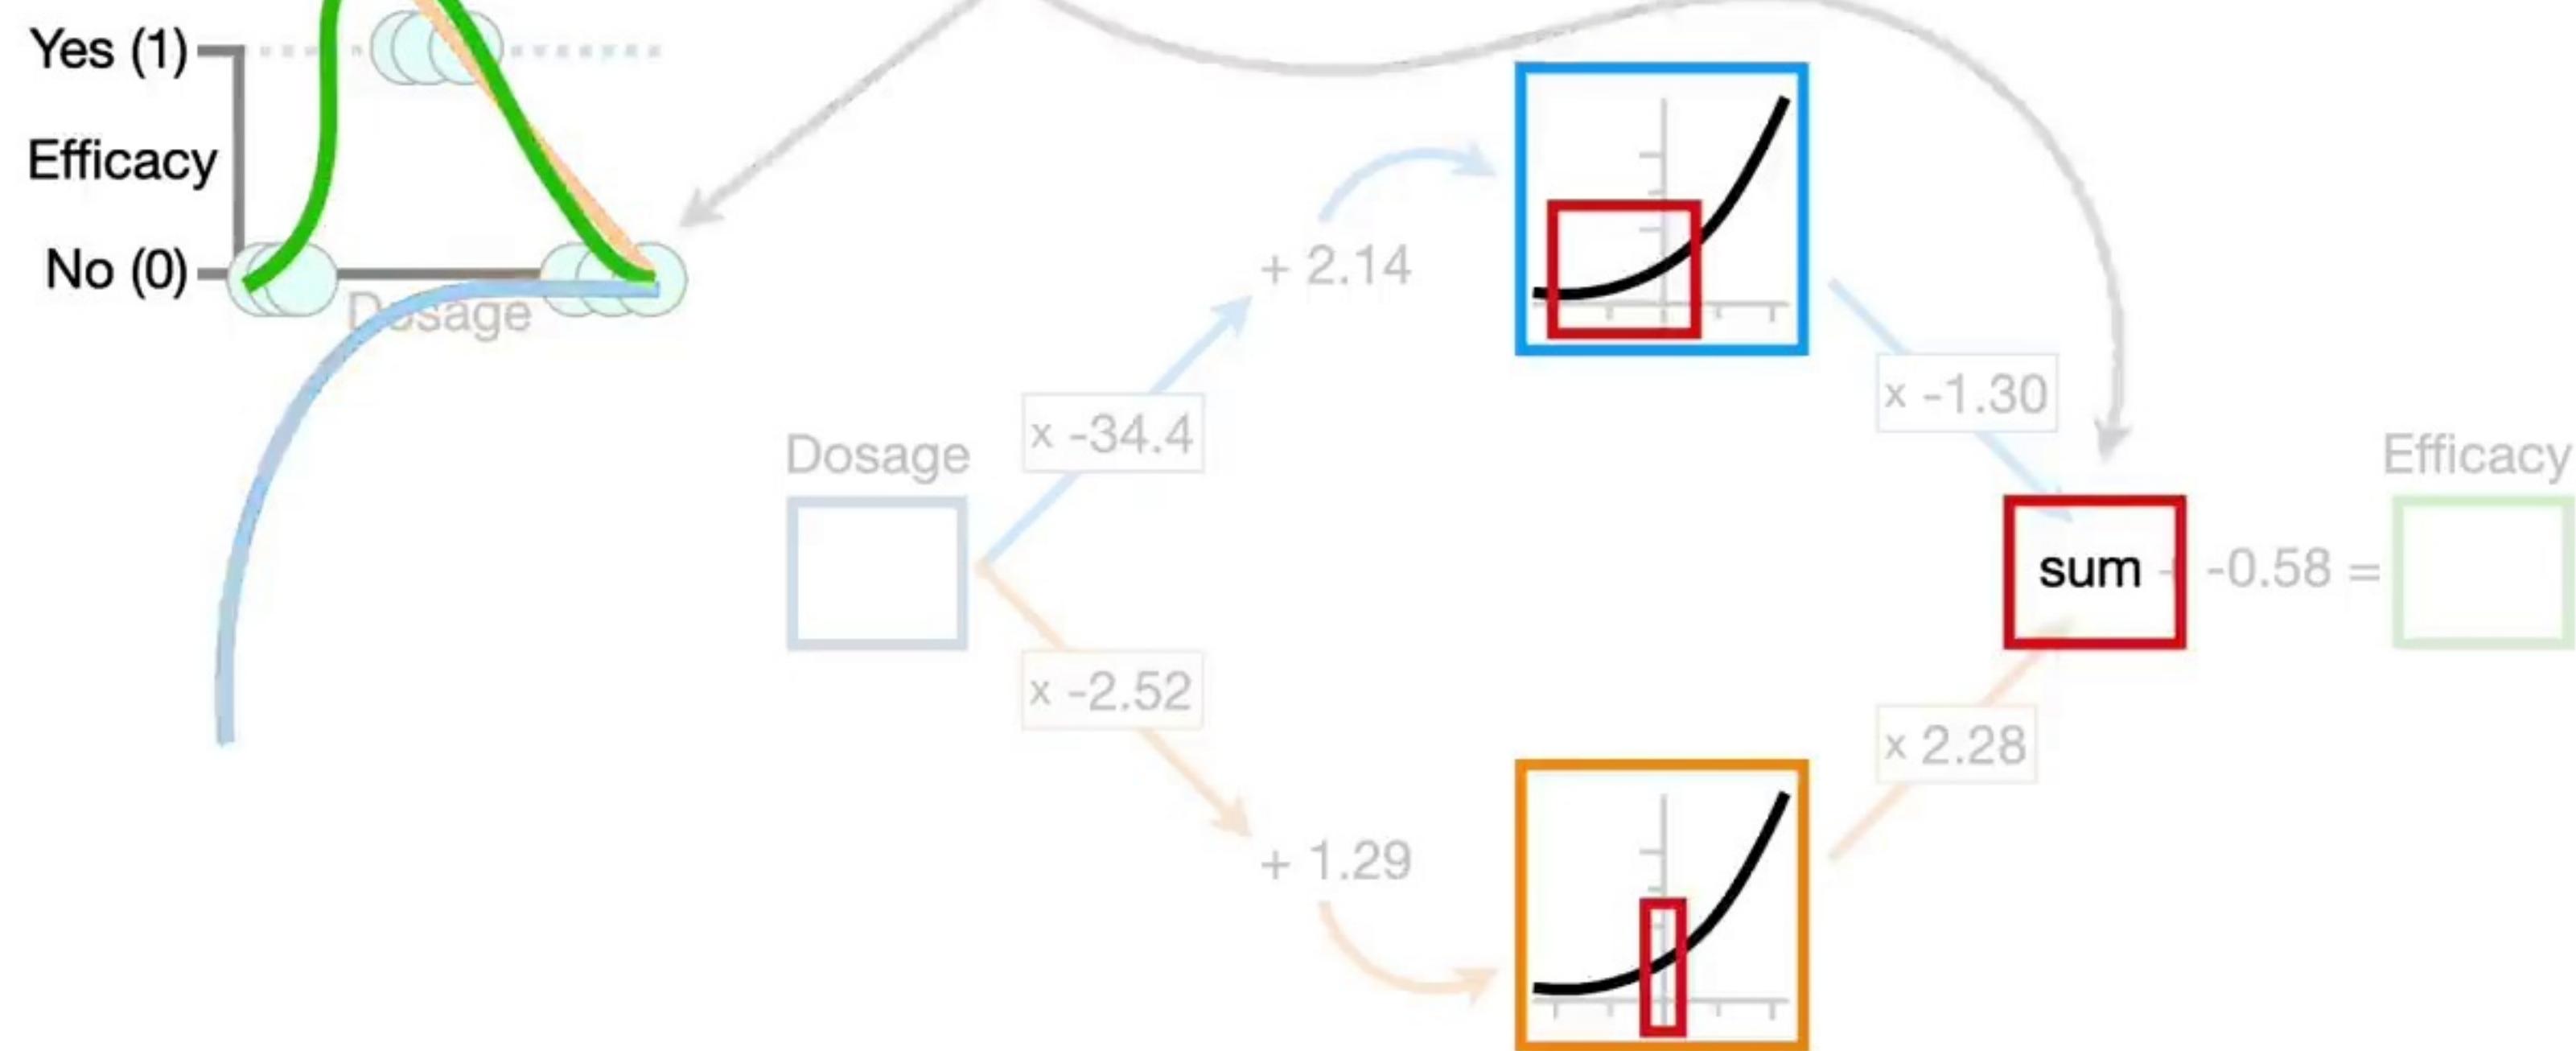


...can be reshaped by the parameter values...



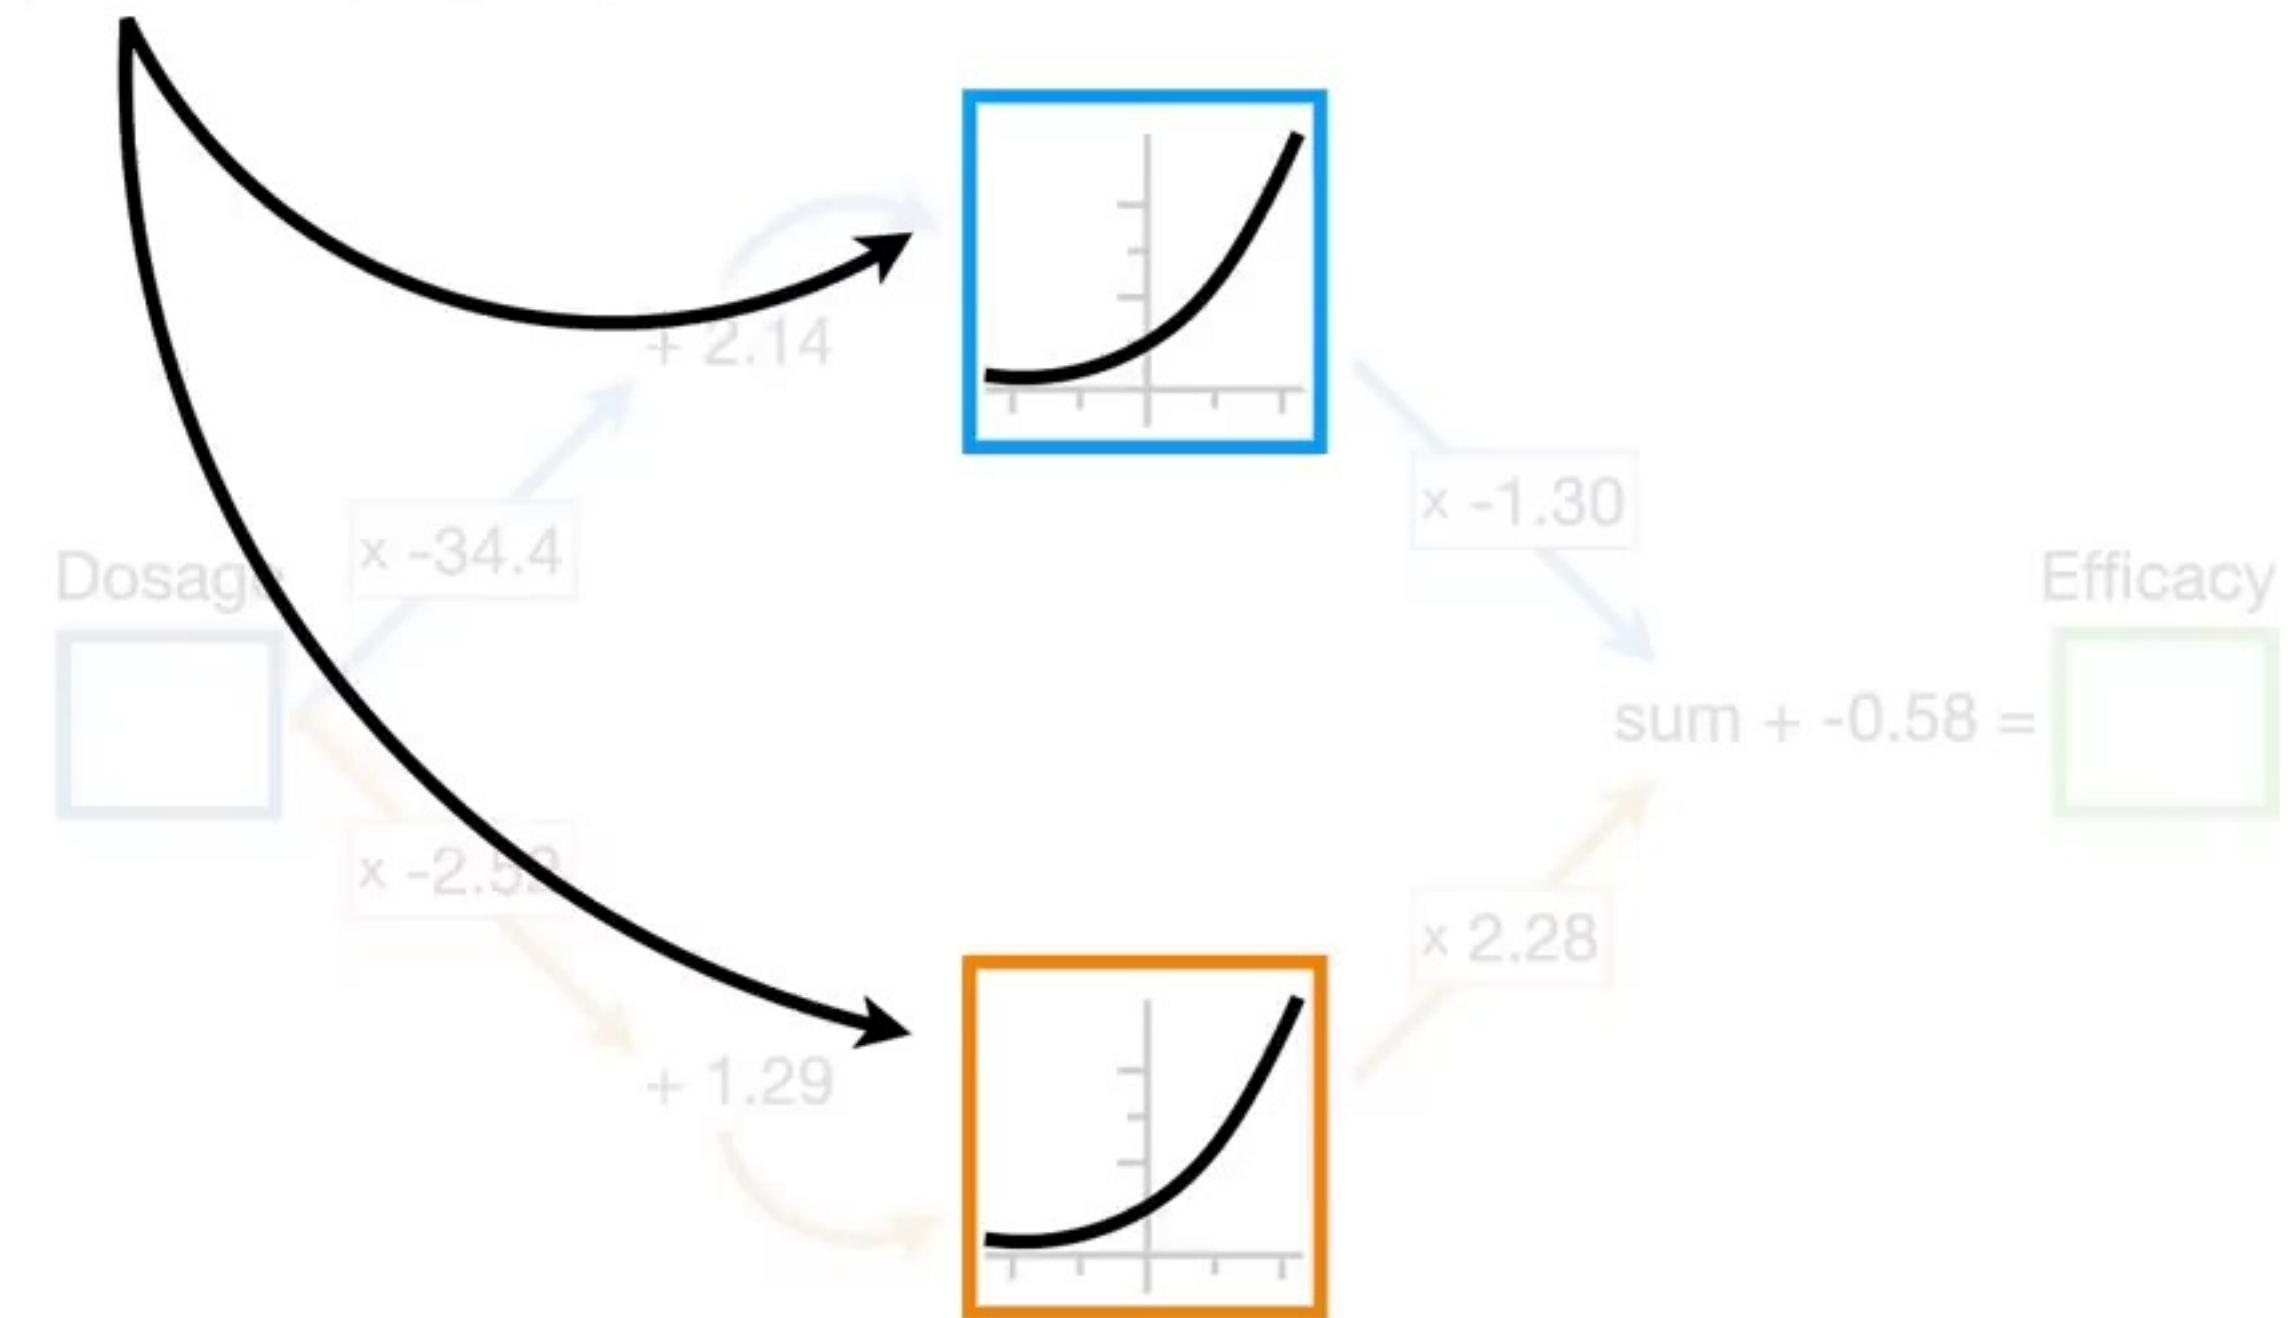
Double  
BAM!!  
**SQ!**

...and then added together  
to get a **green squiggle** that  
fits the data.



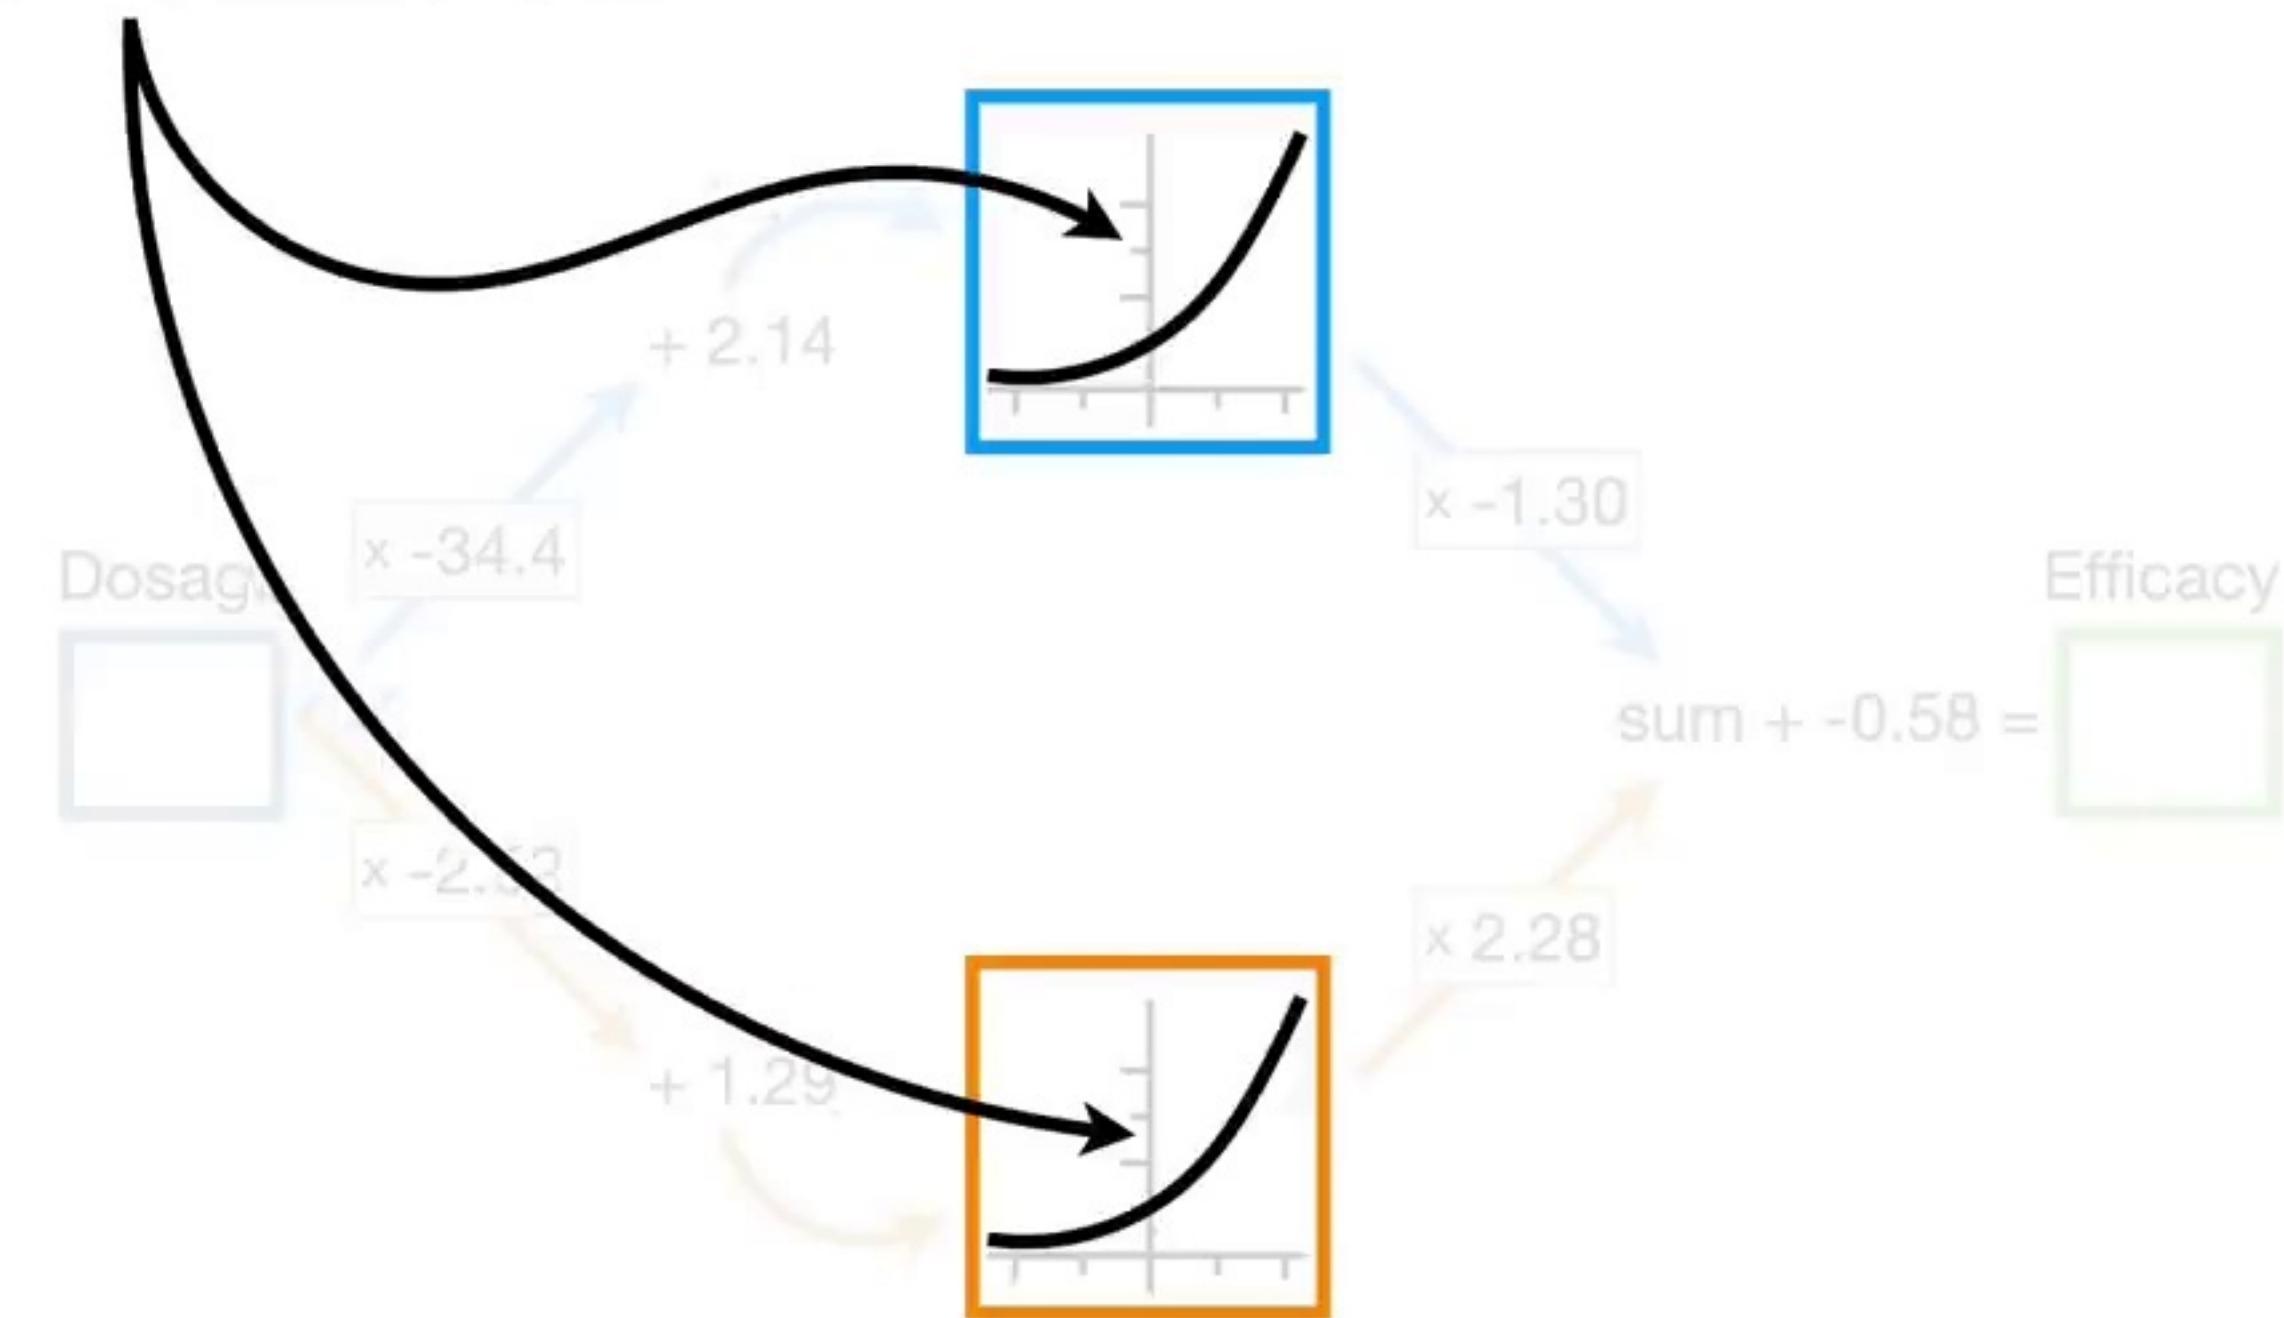


**NOTE:** There are many common bent or curved lines that we can choose for a Neural Network.



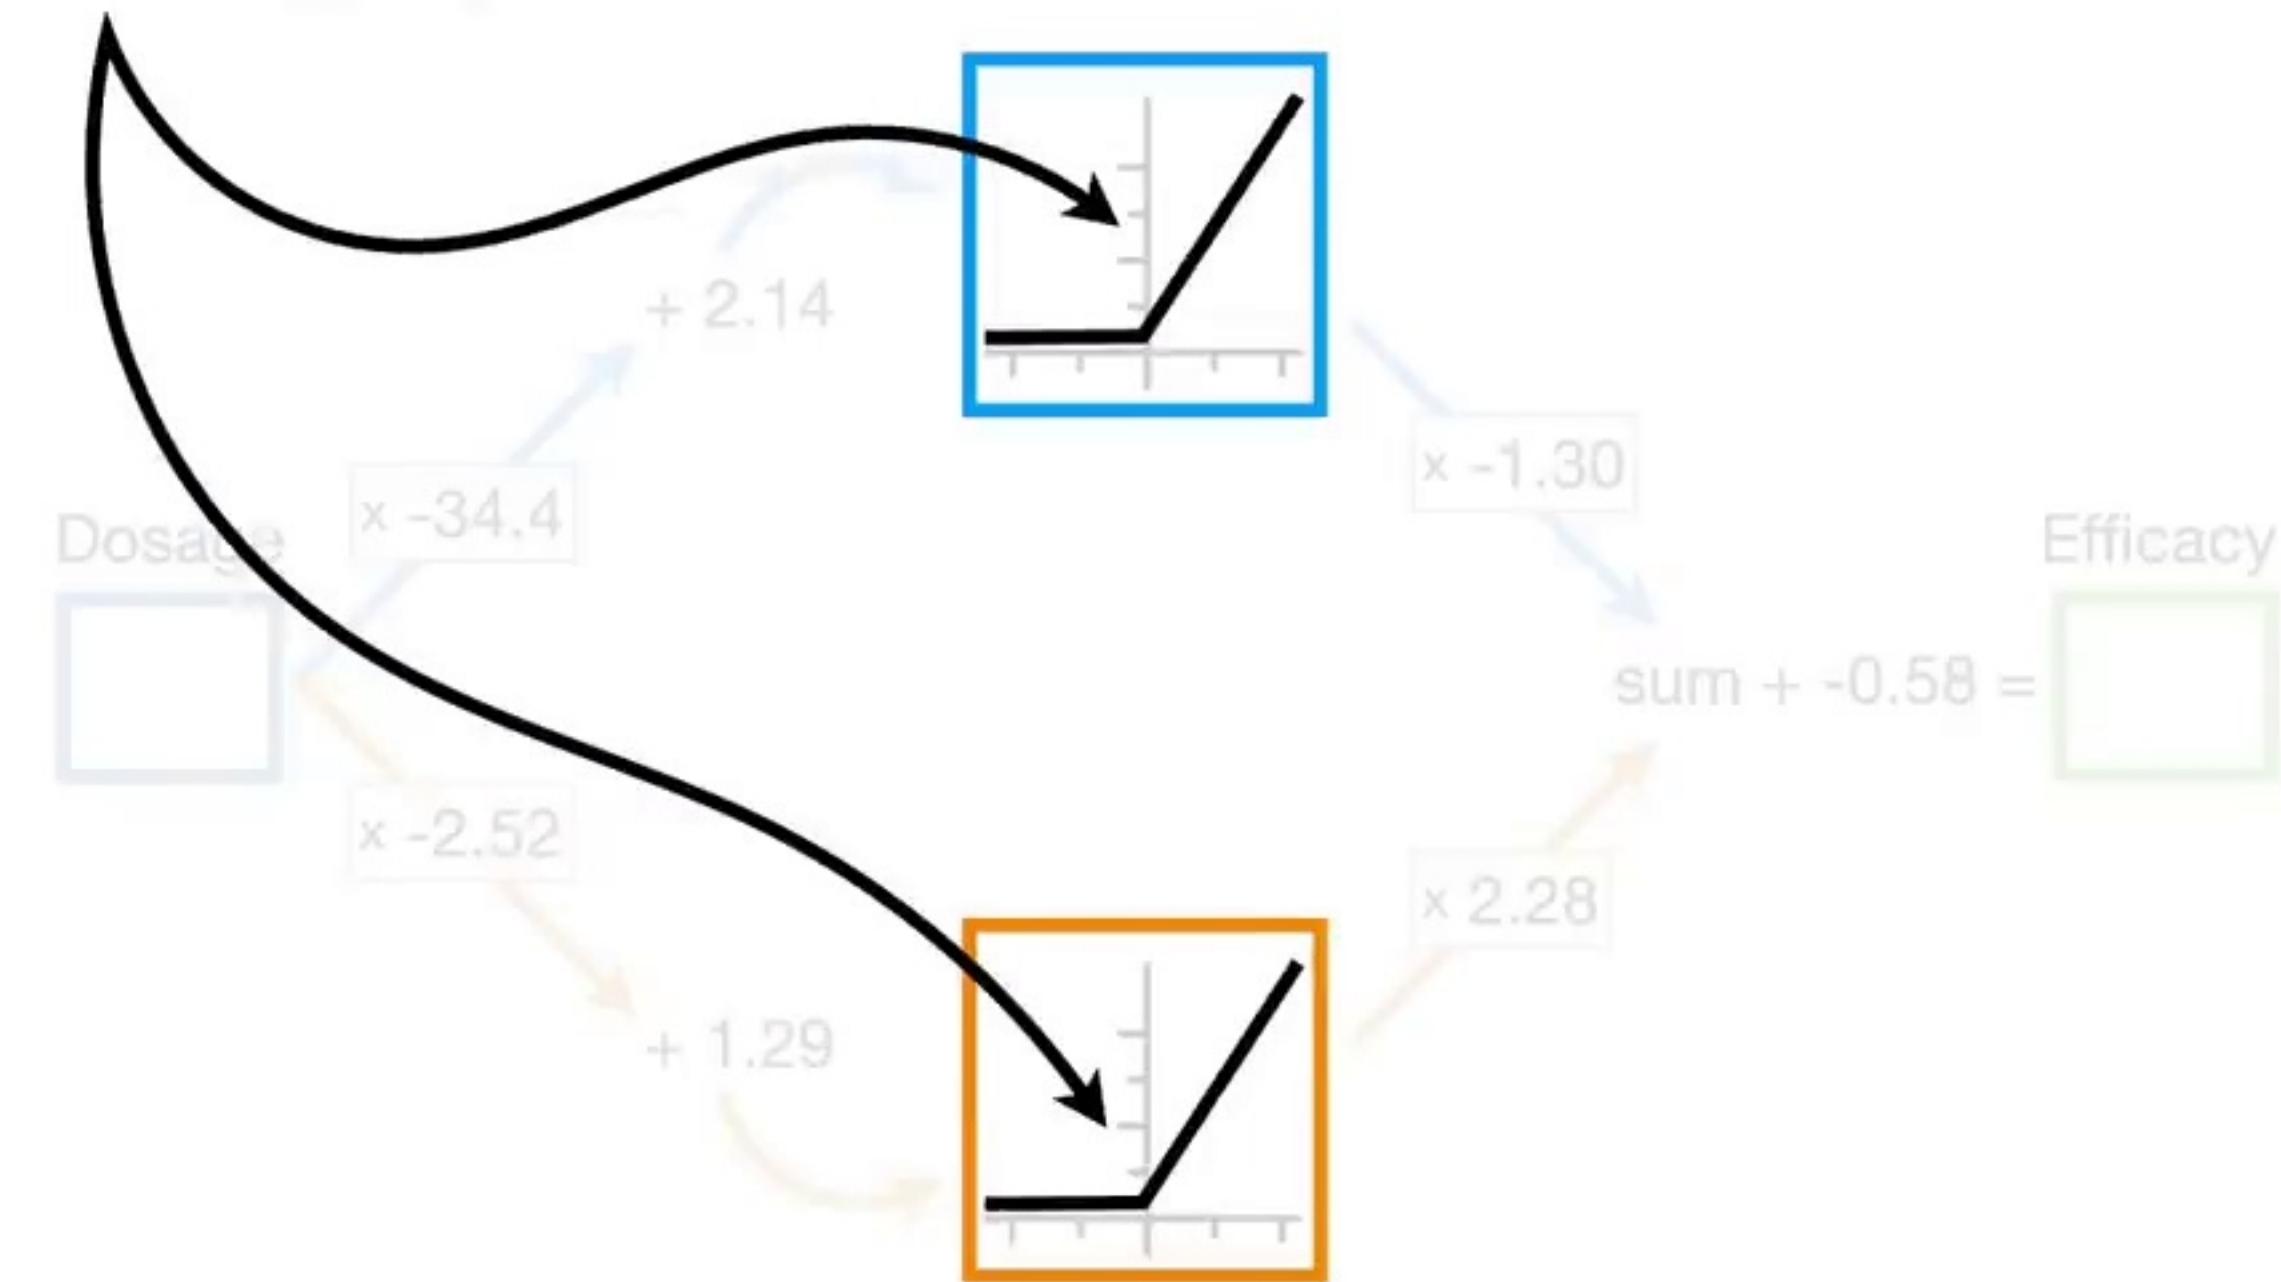


This specific **curved line** is called **softplus**, which sounds like a brand of toilet paper.



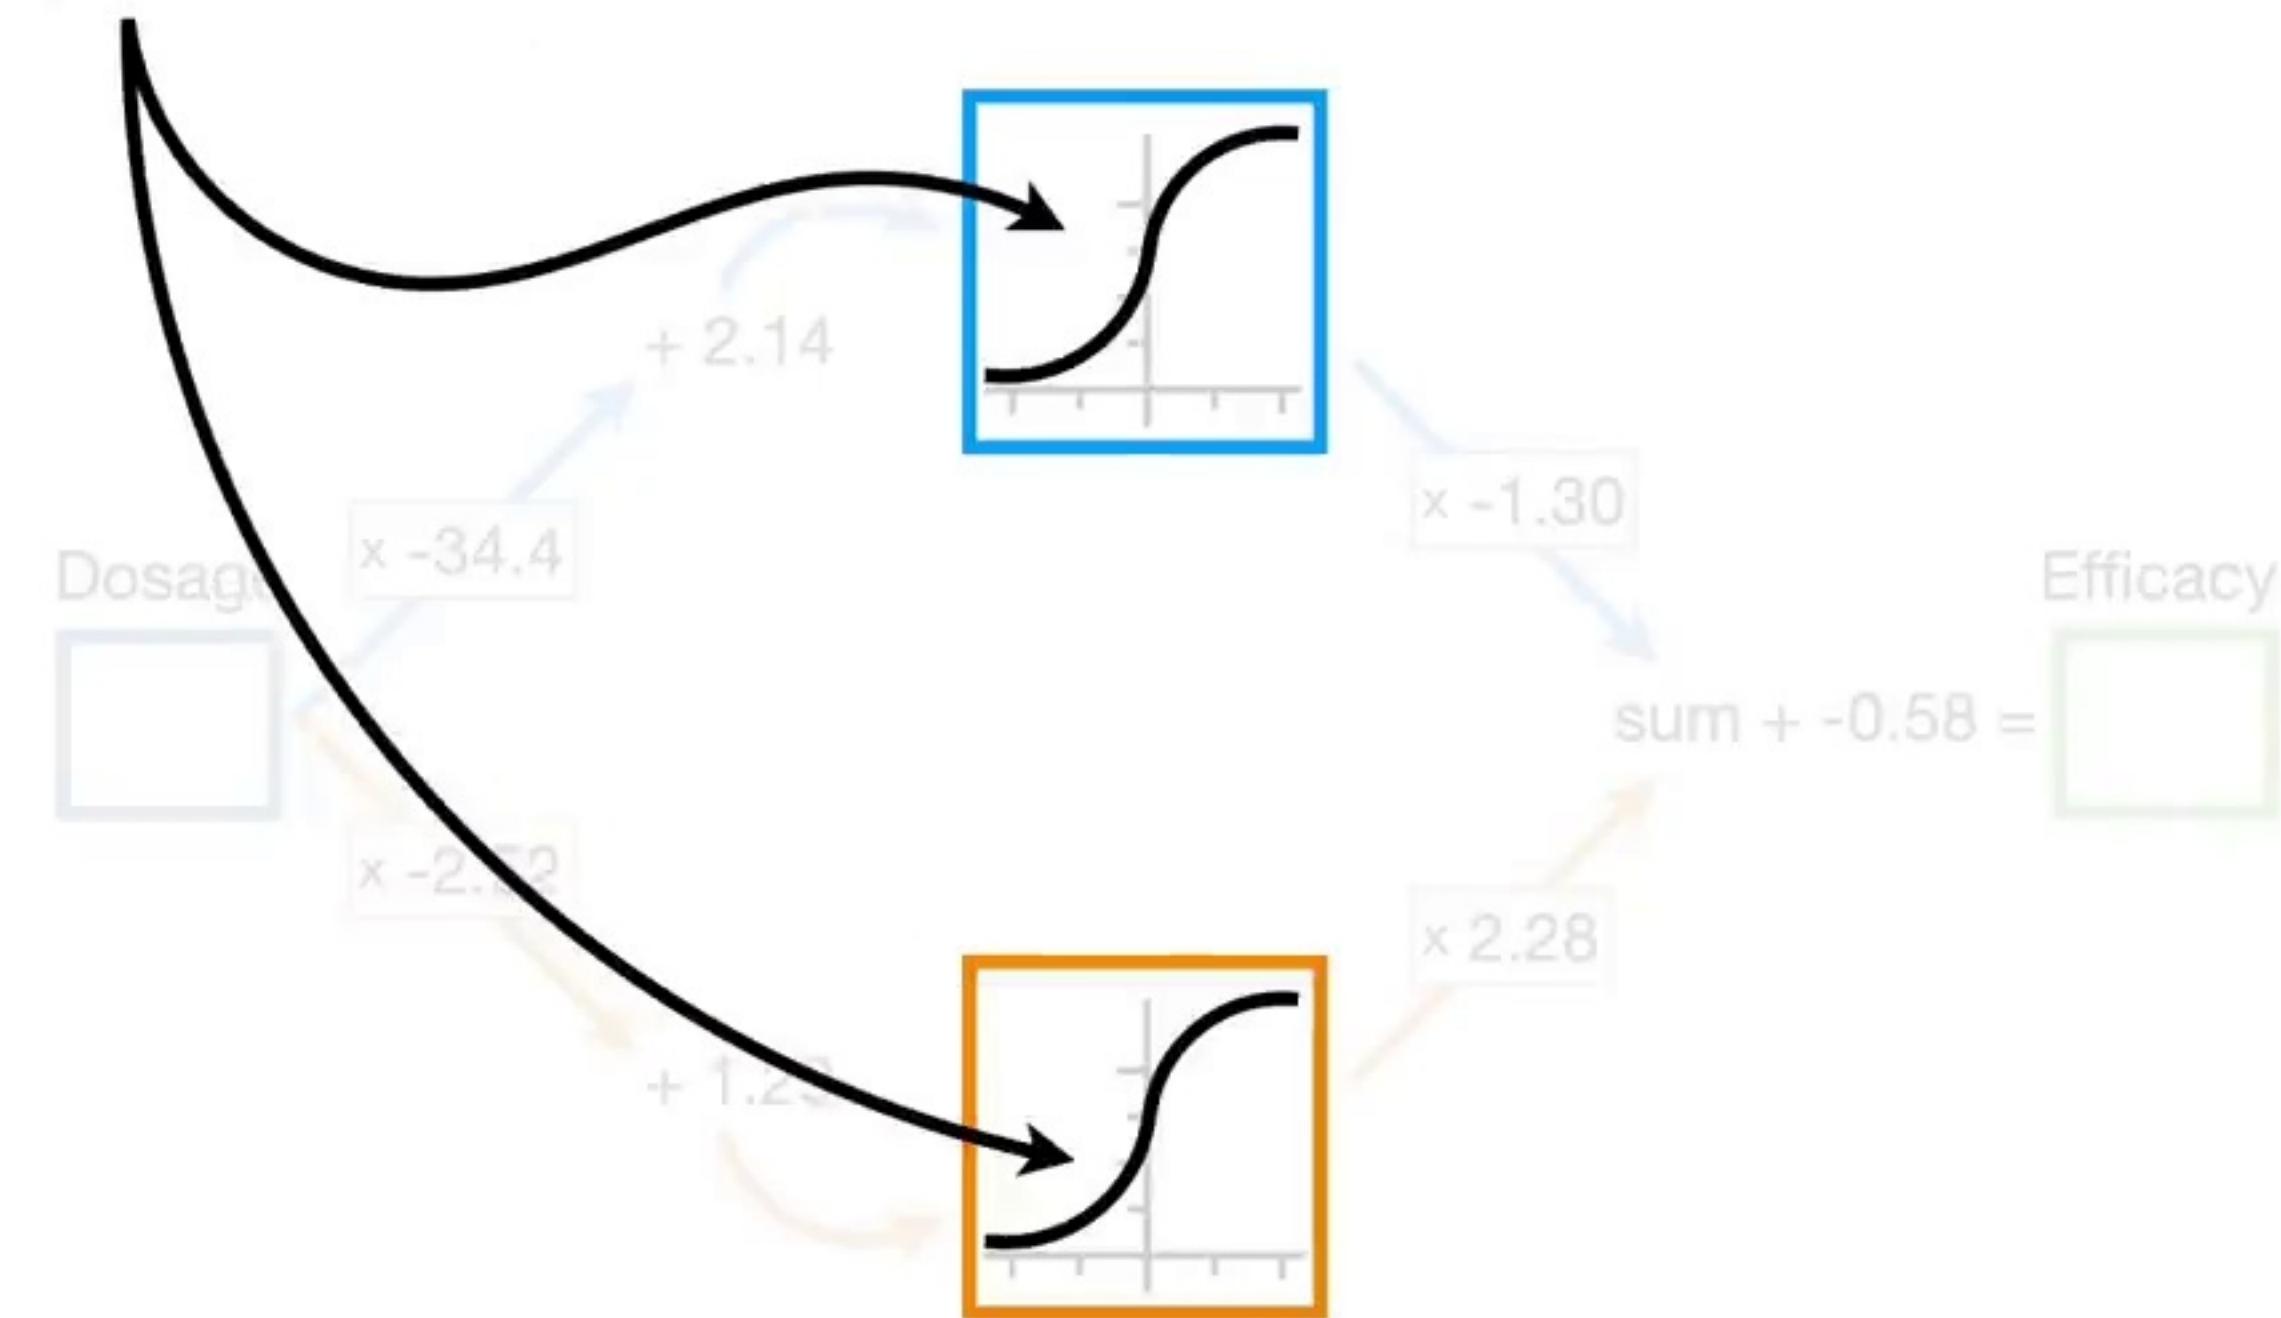


Alternatively, we could use this **bent line**, called, **ReLU**, which is short for **Rectified Linear Unit** and sounds like a robot.



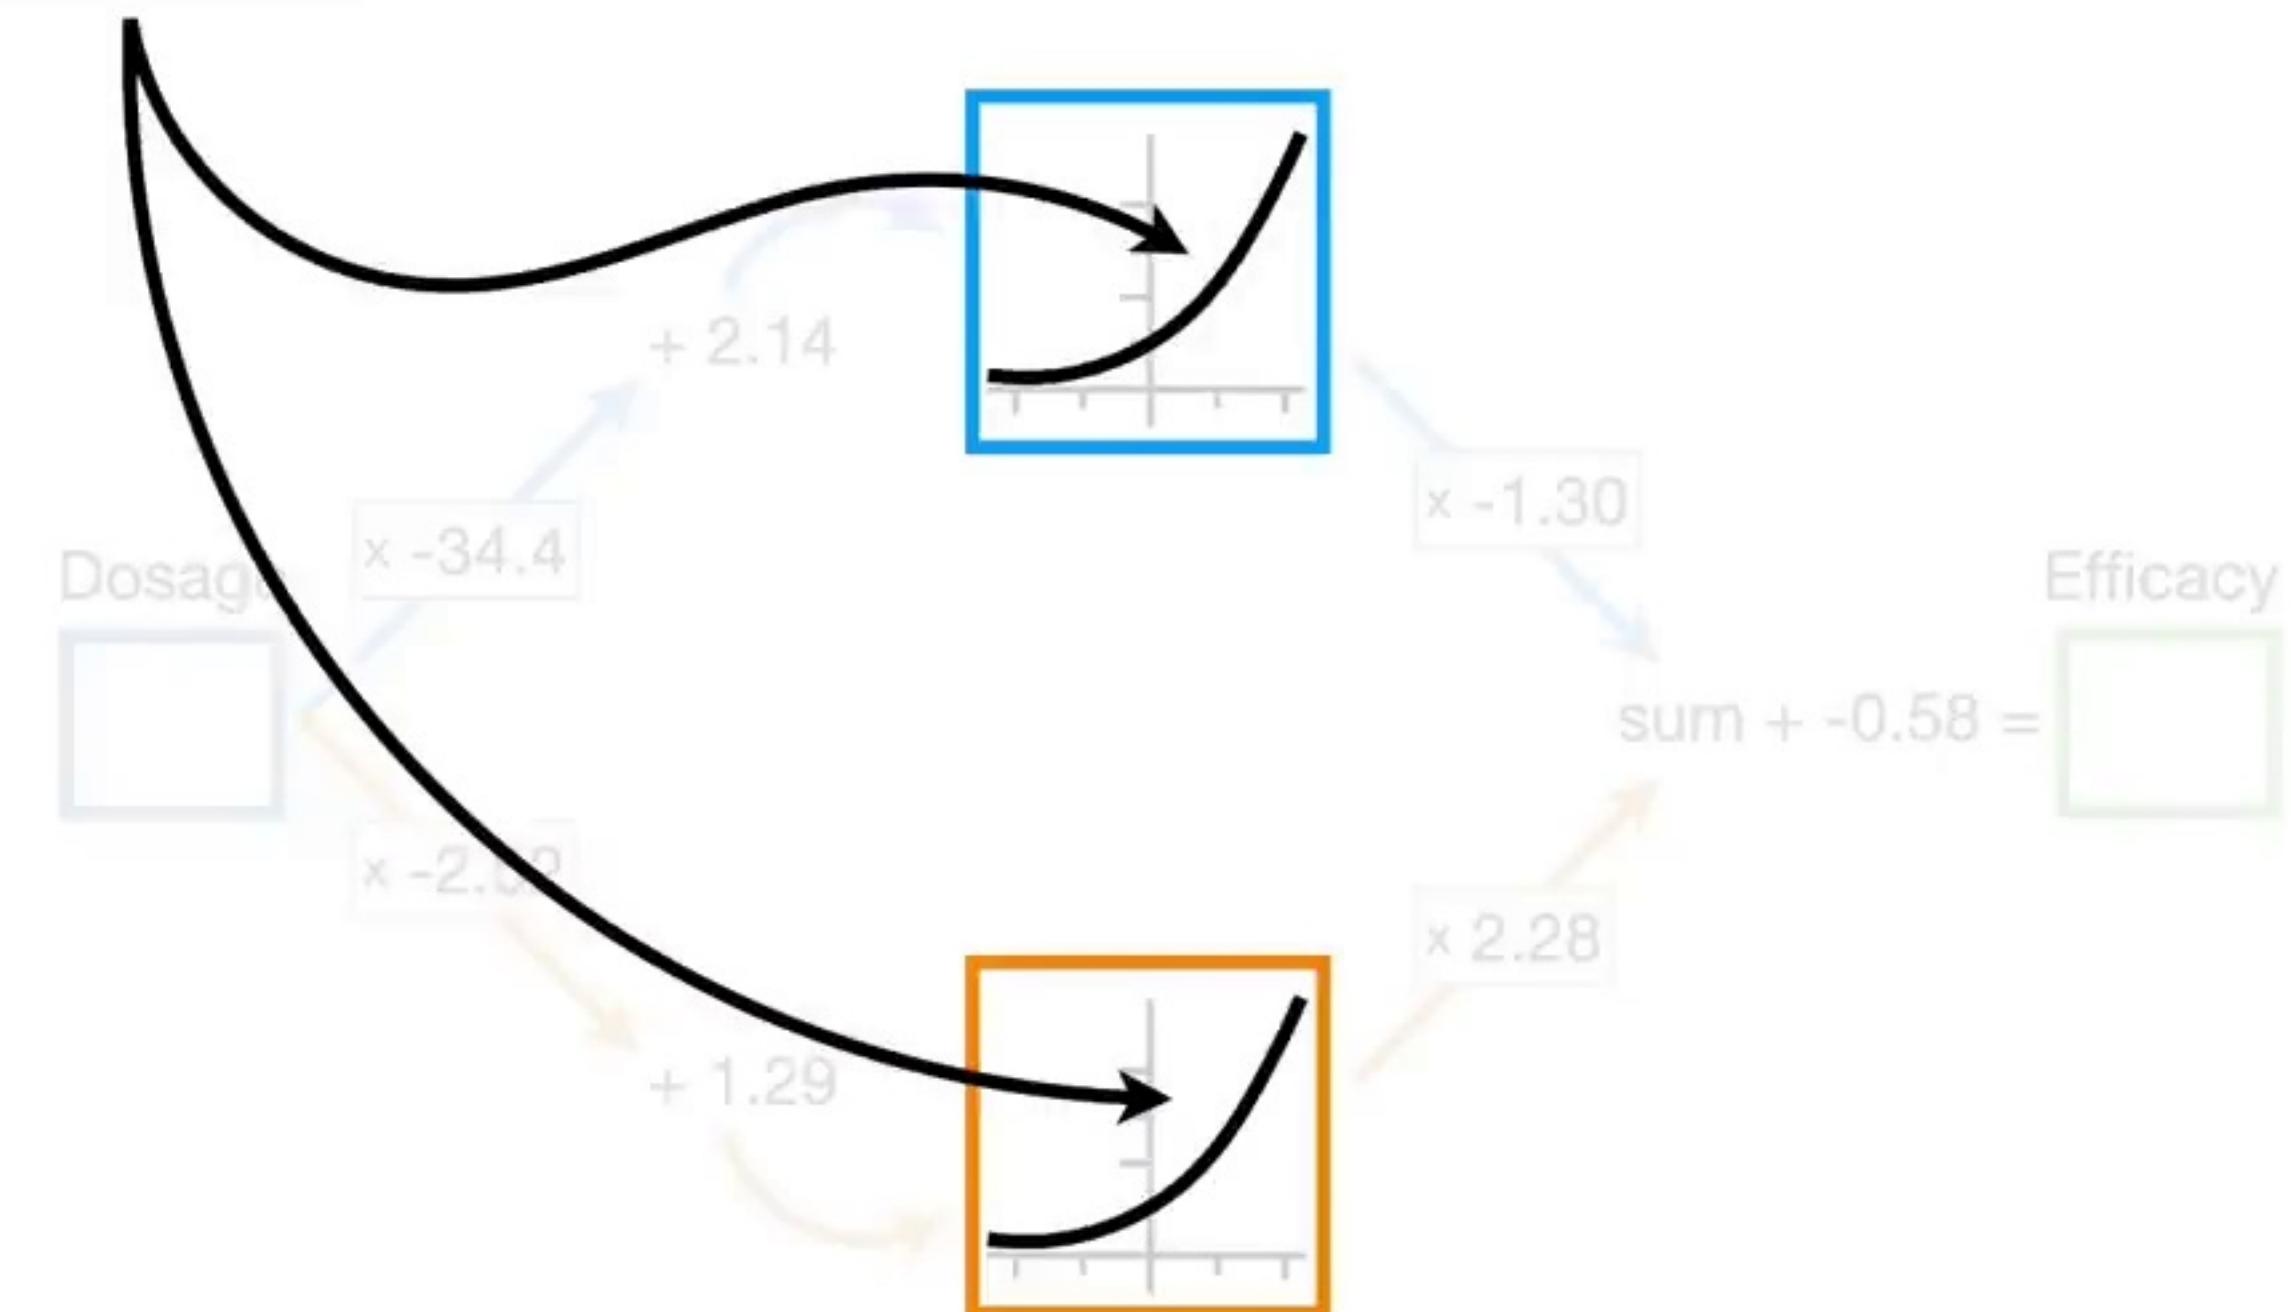


Or we could use a **sigmoid** shape or any other **bent** or **curved line**.



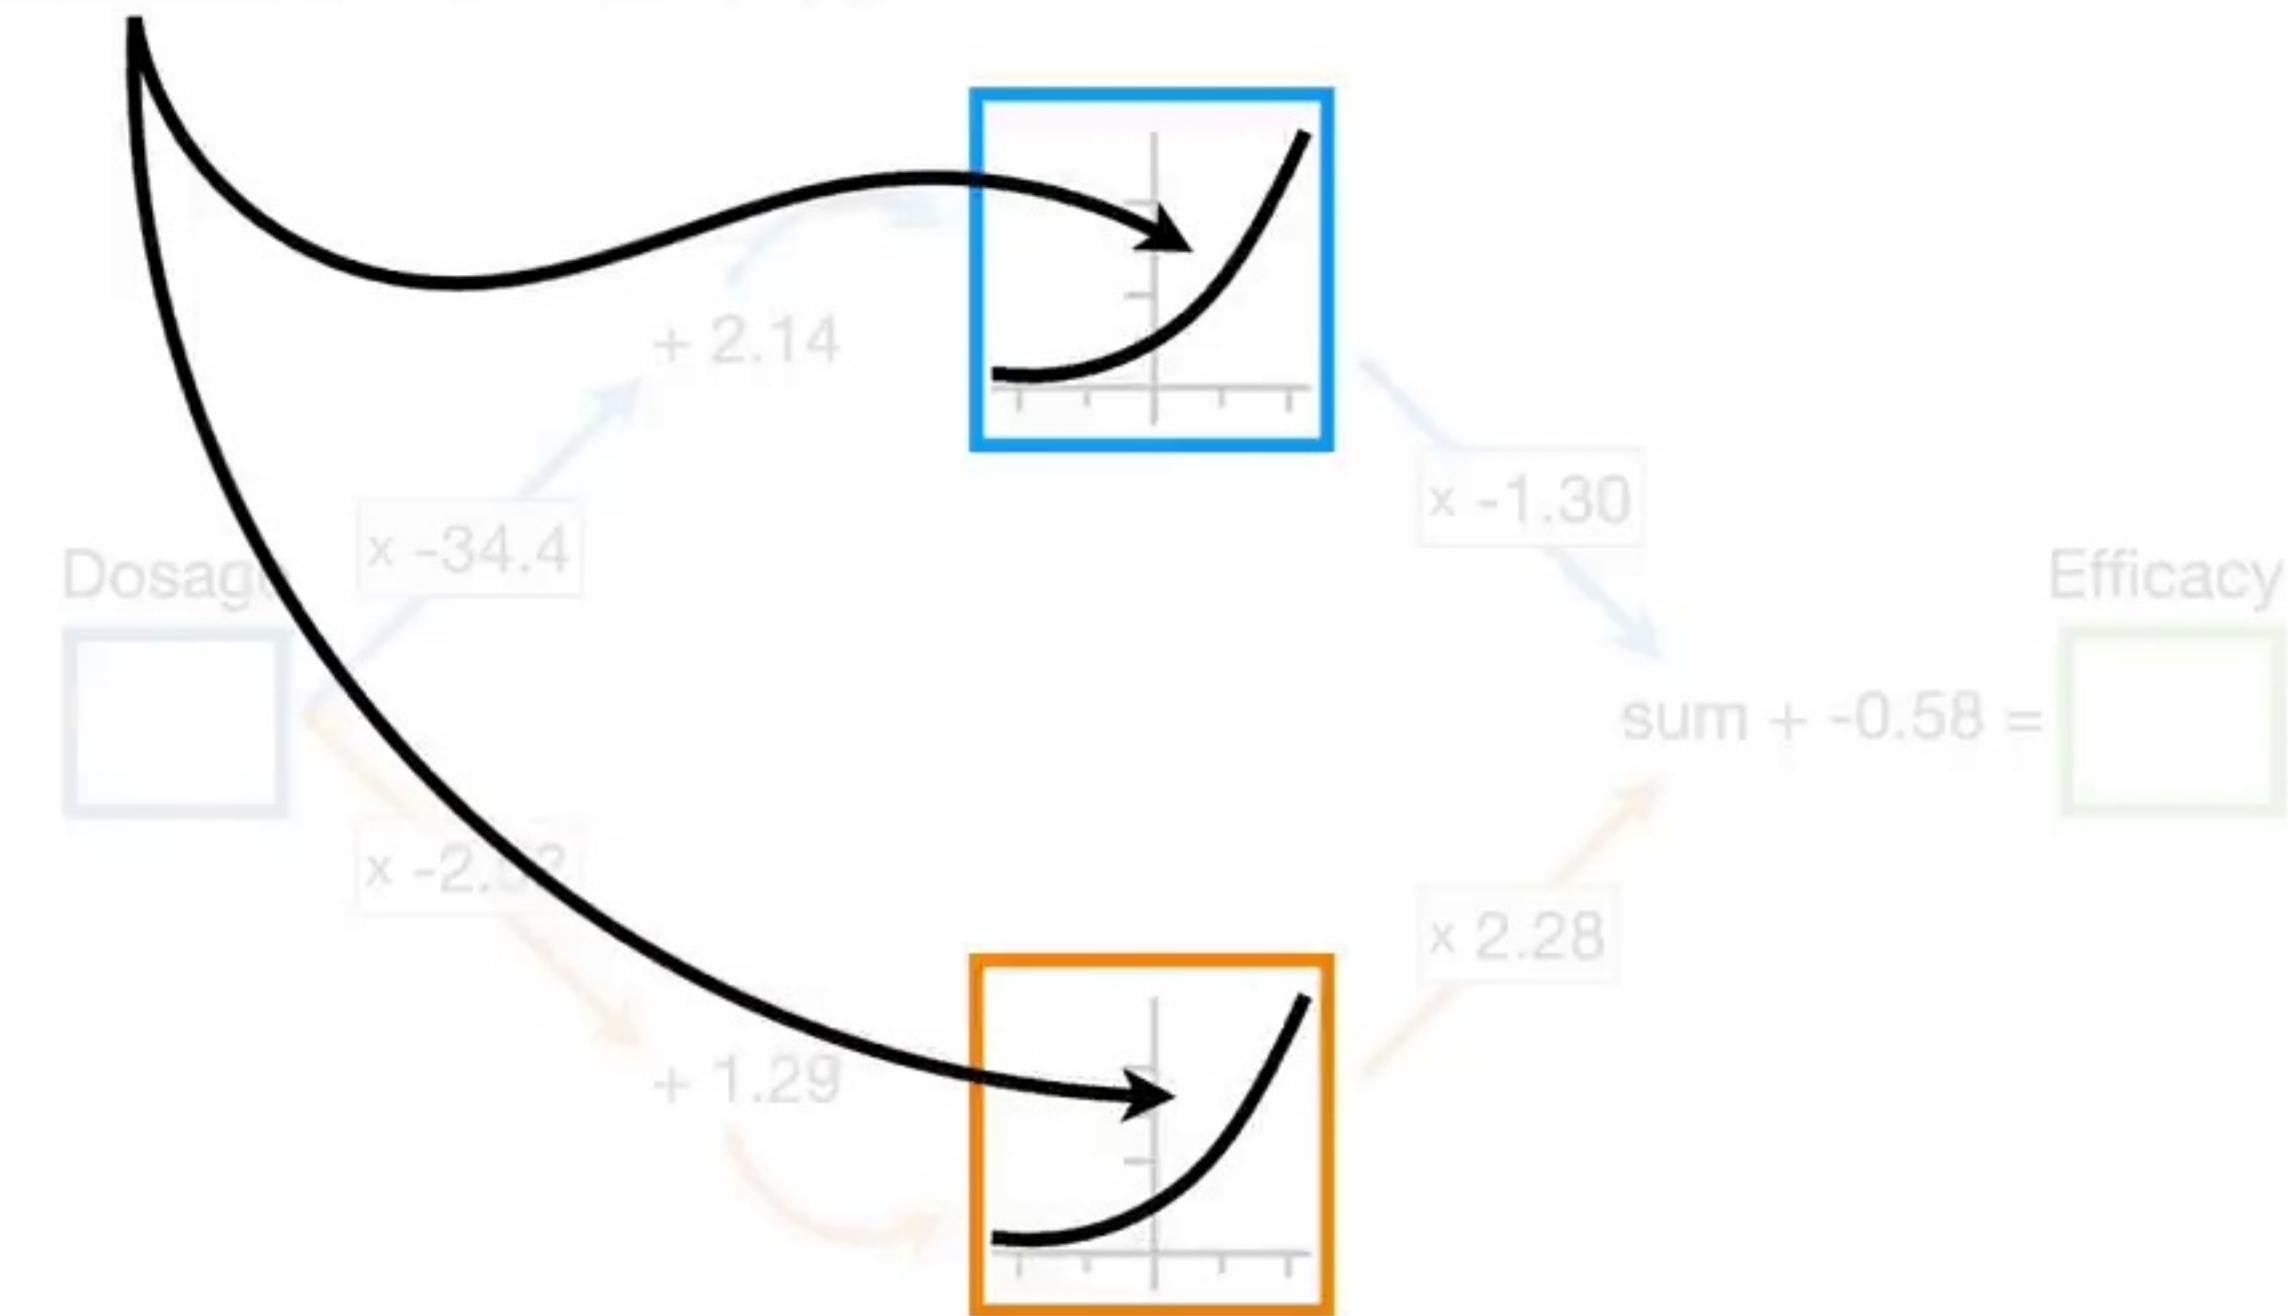


The **curved or bent lines**  
are called **Activation  
Functions.**



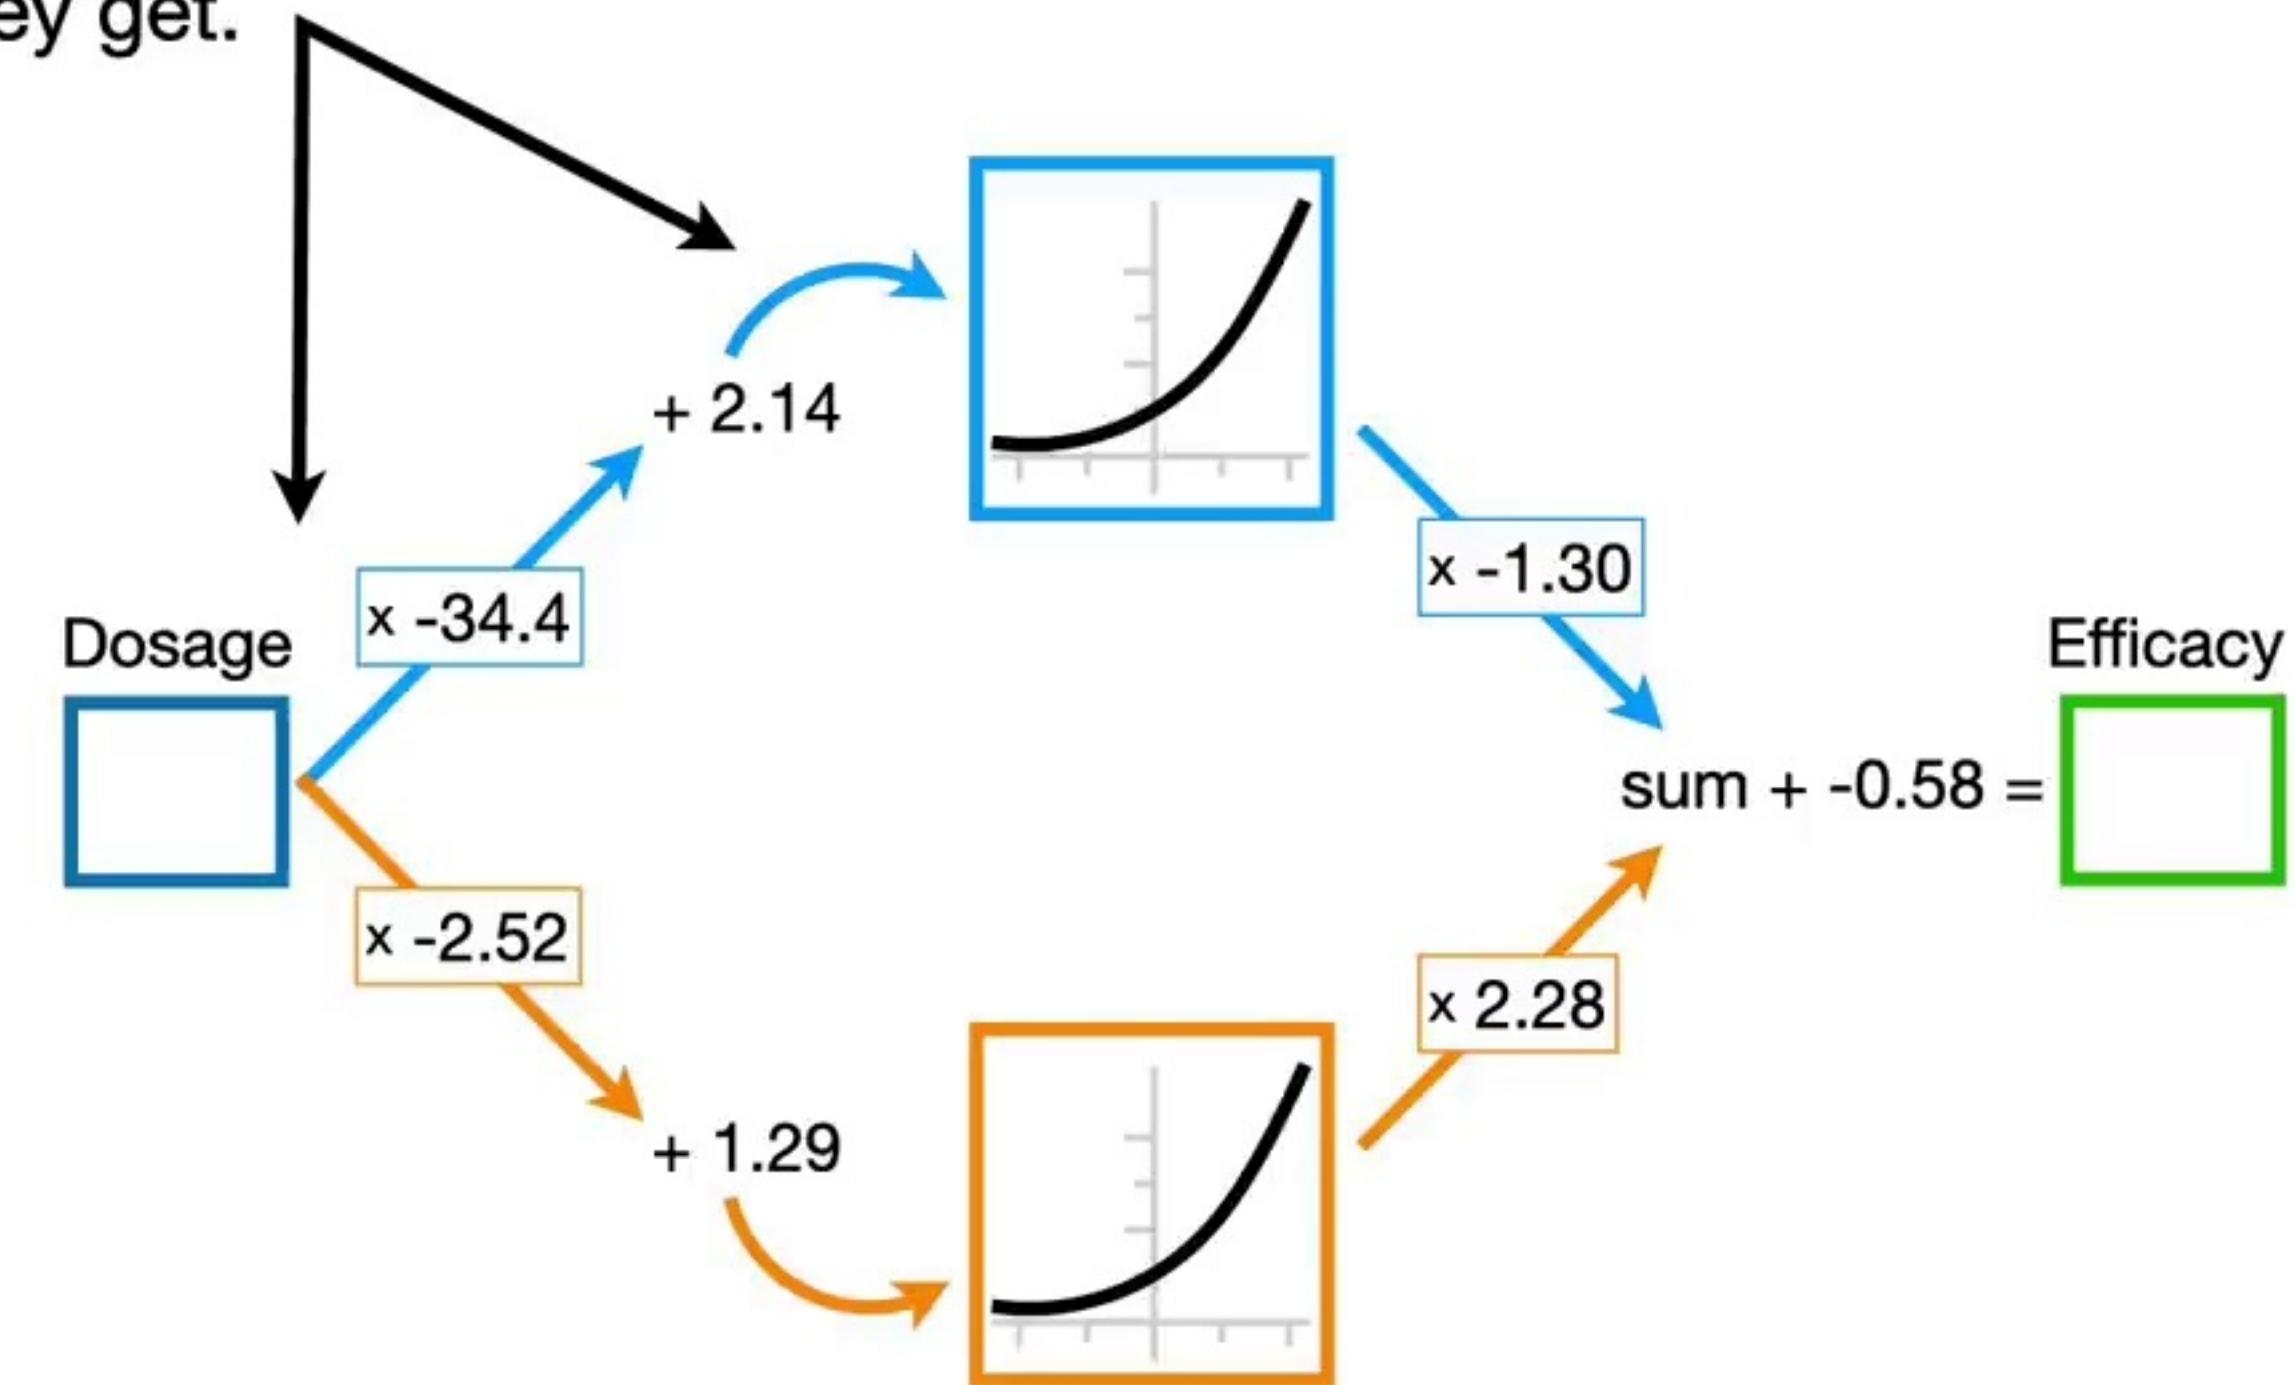


When you build a **Neural Network**, you have to decide which **Activation Function or Functions** you want to use.



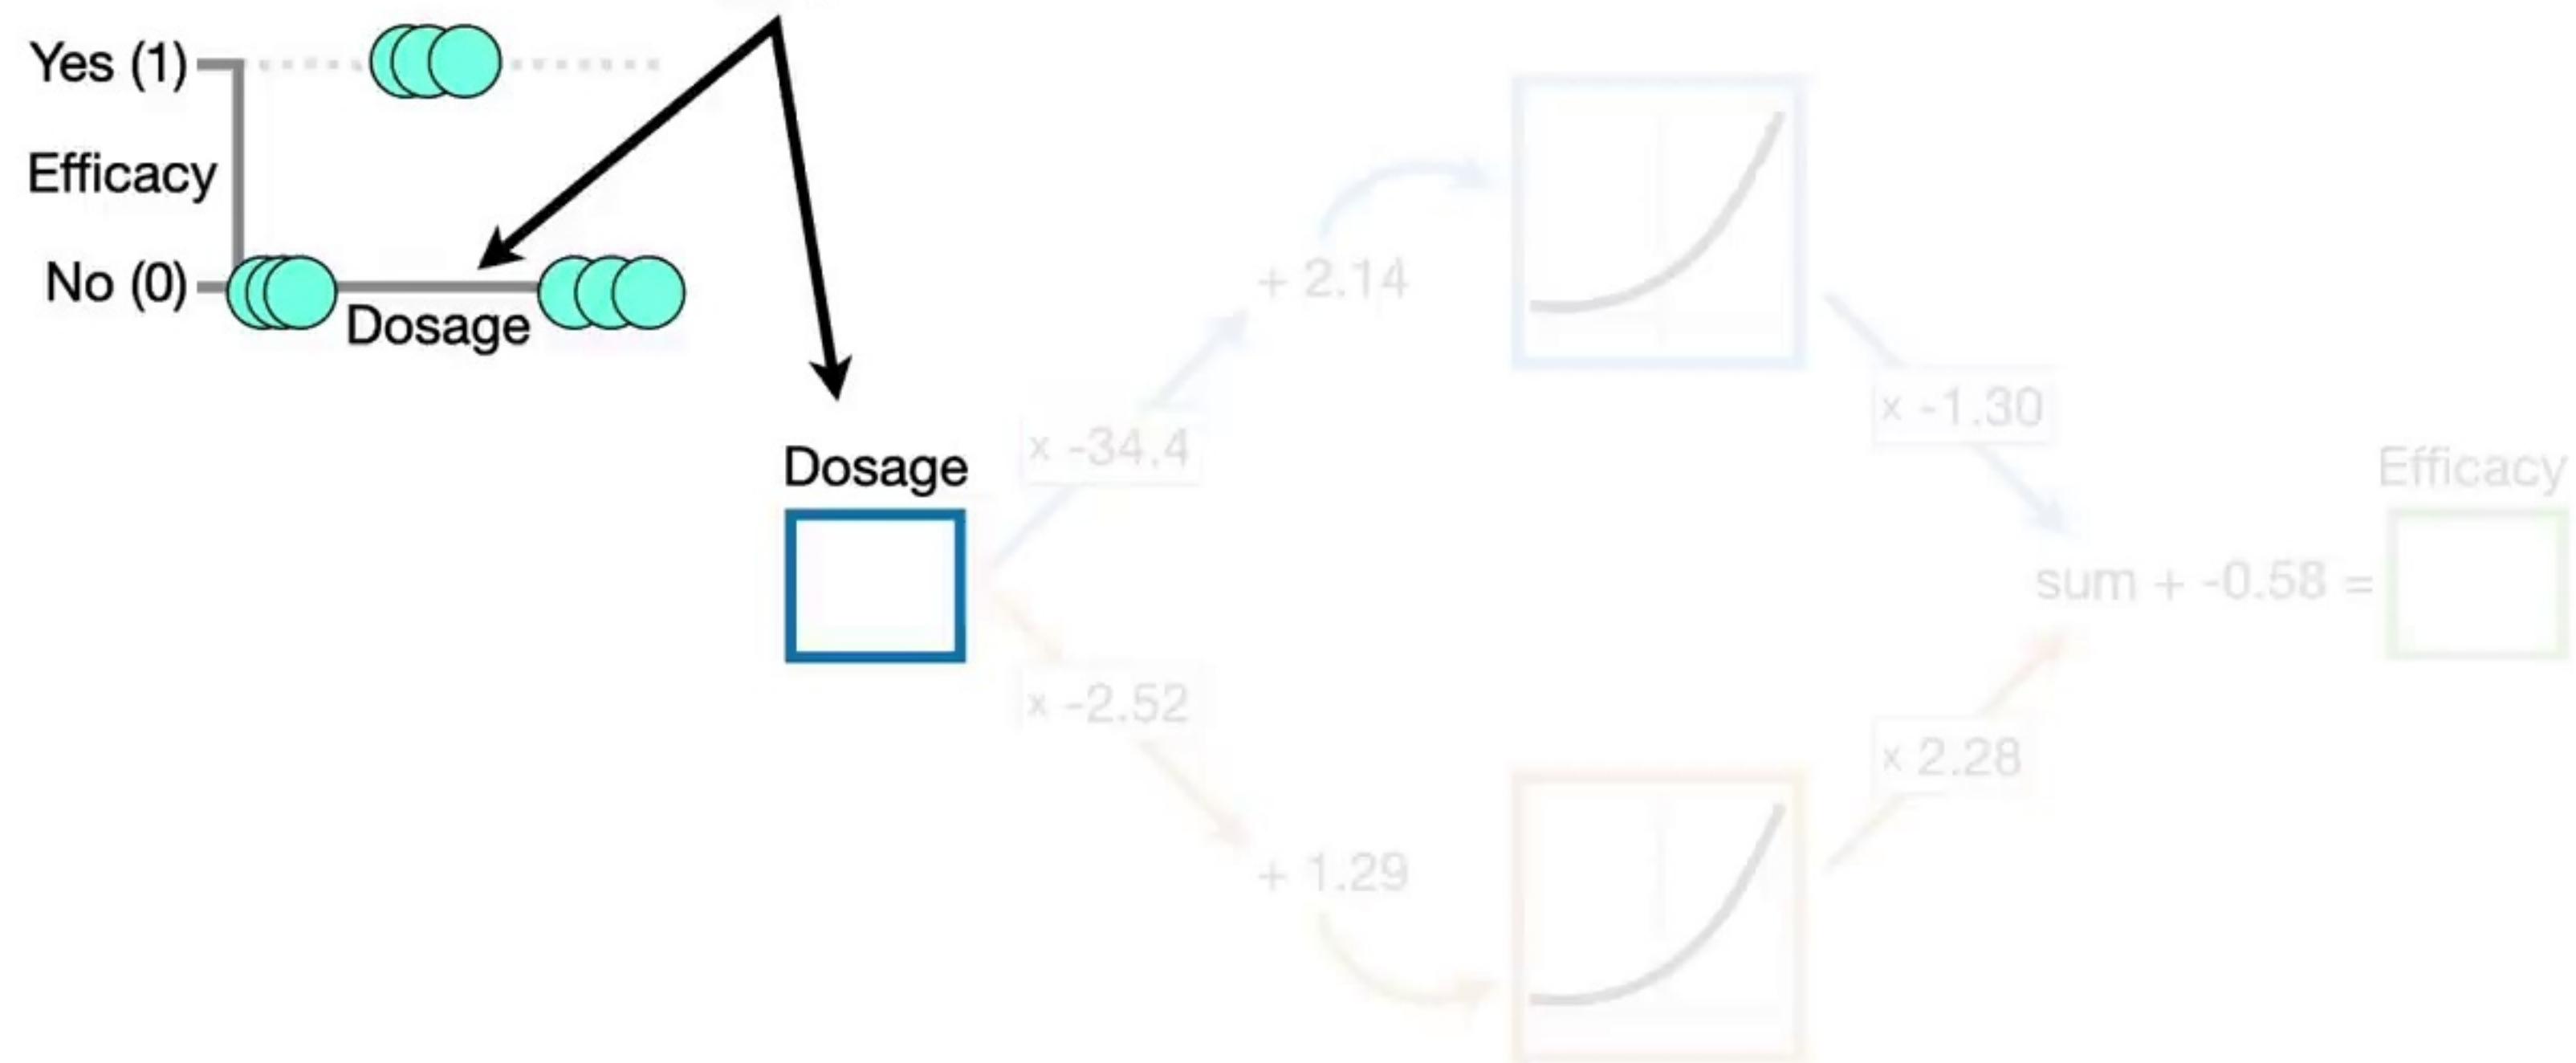


**NOTE:** This specific **Neural Network** is about as simple as they get.



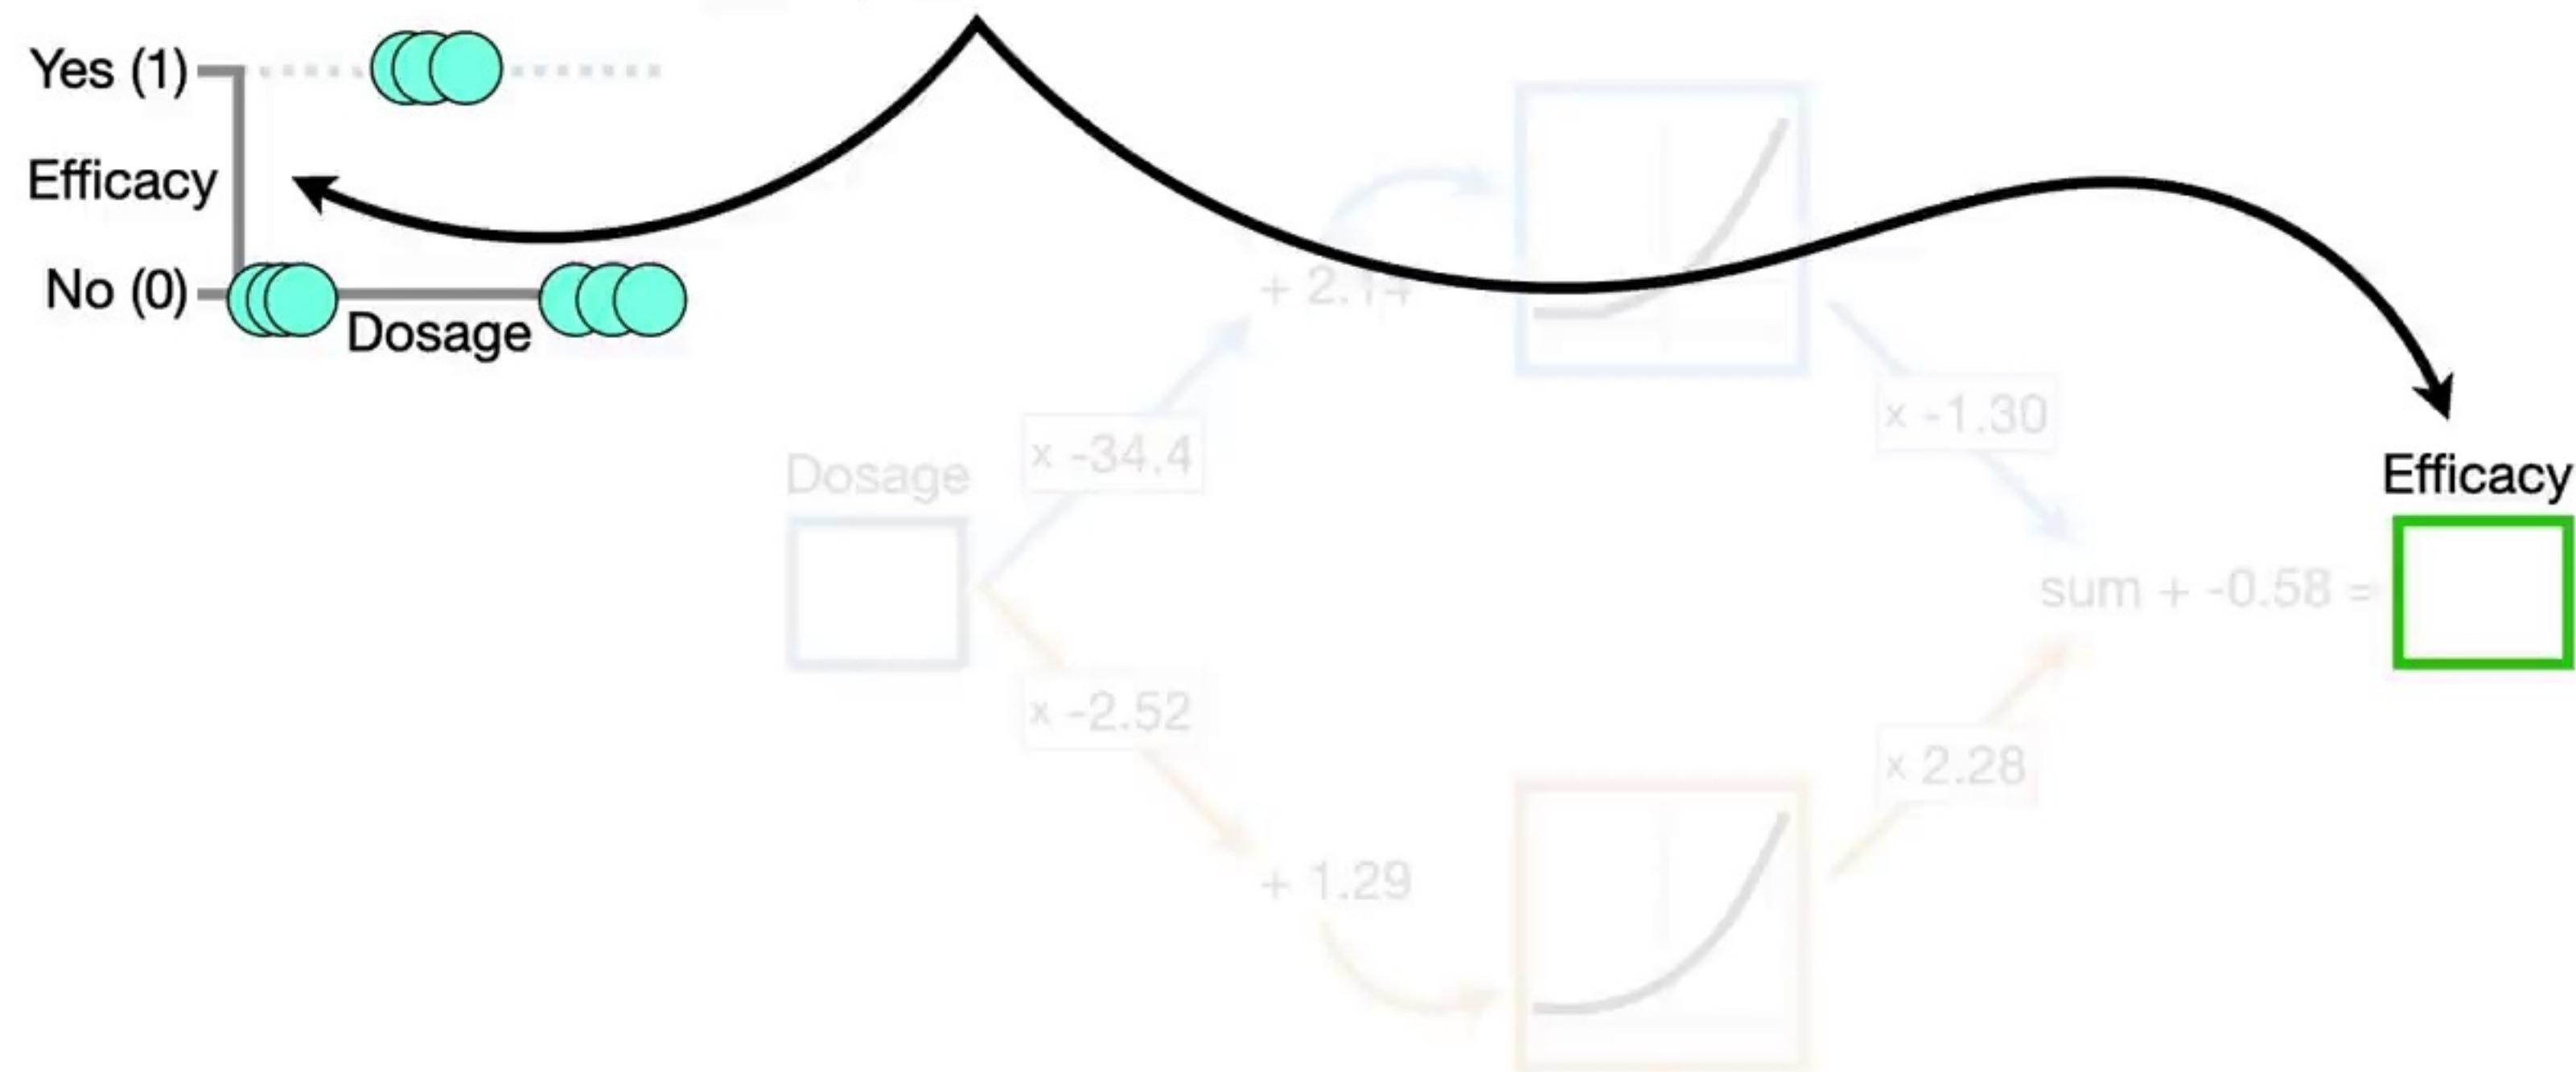


It only has **1 Input Node**  
where we plug in the  
**Dosage...**



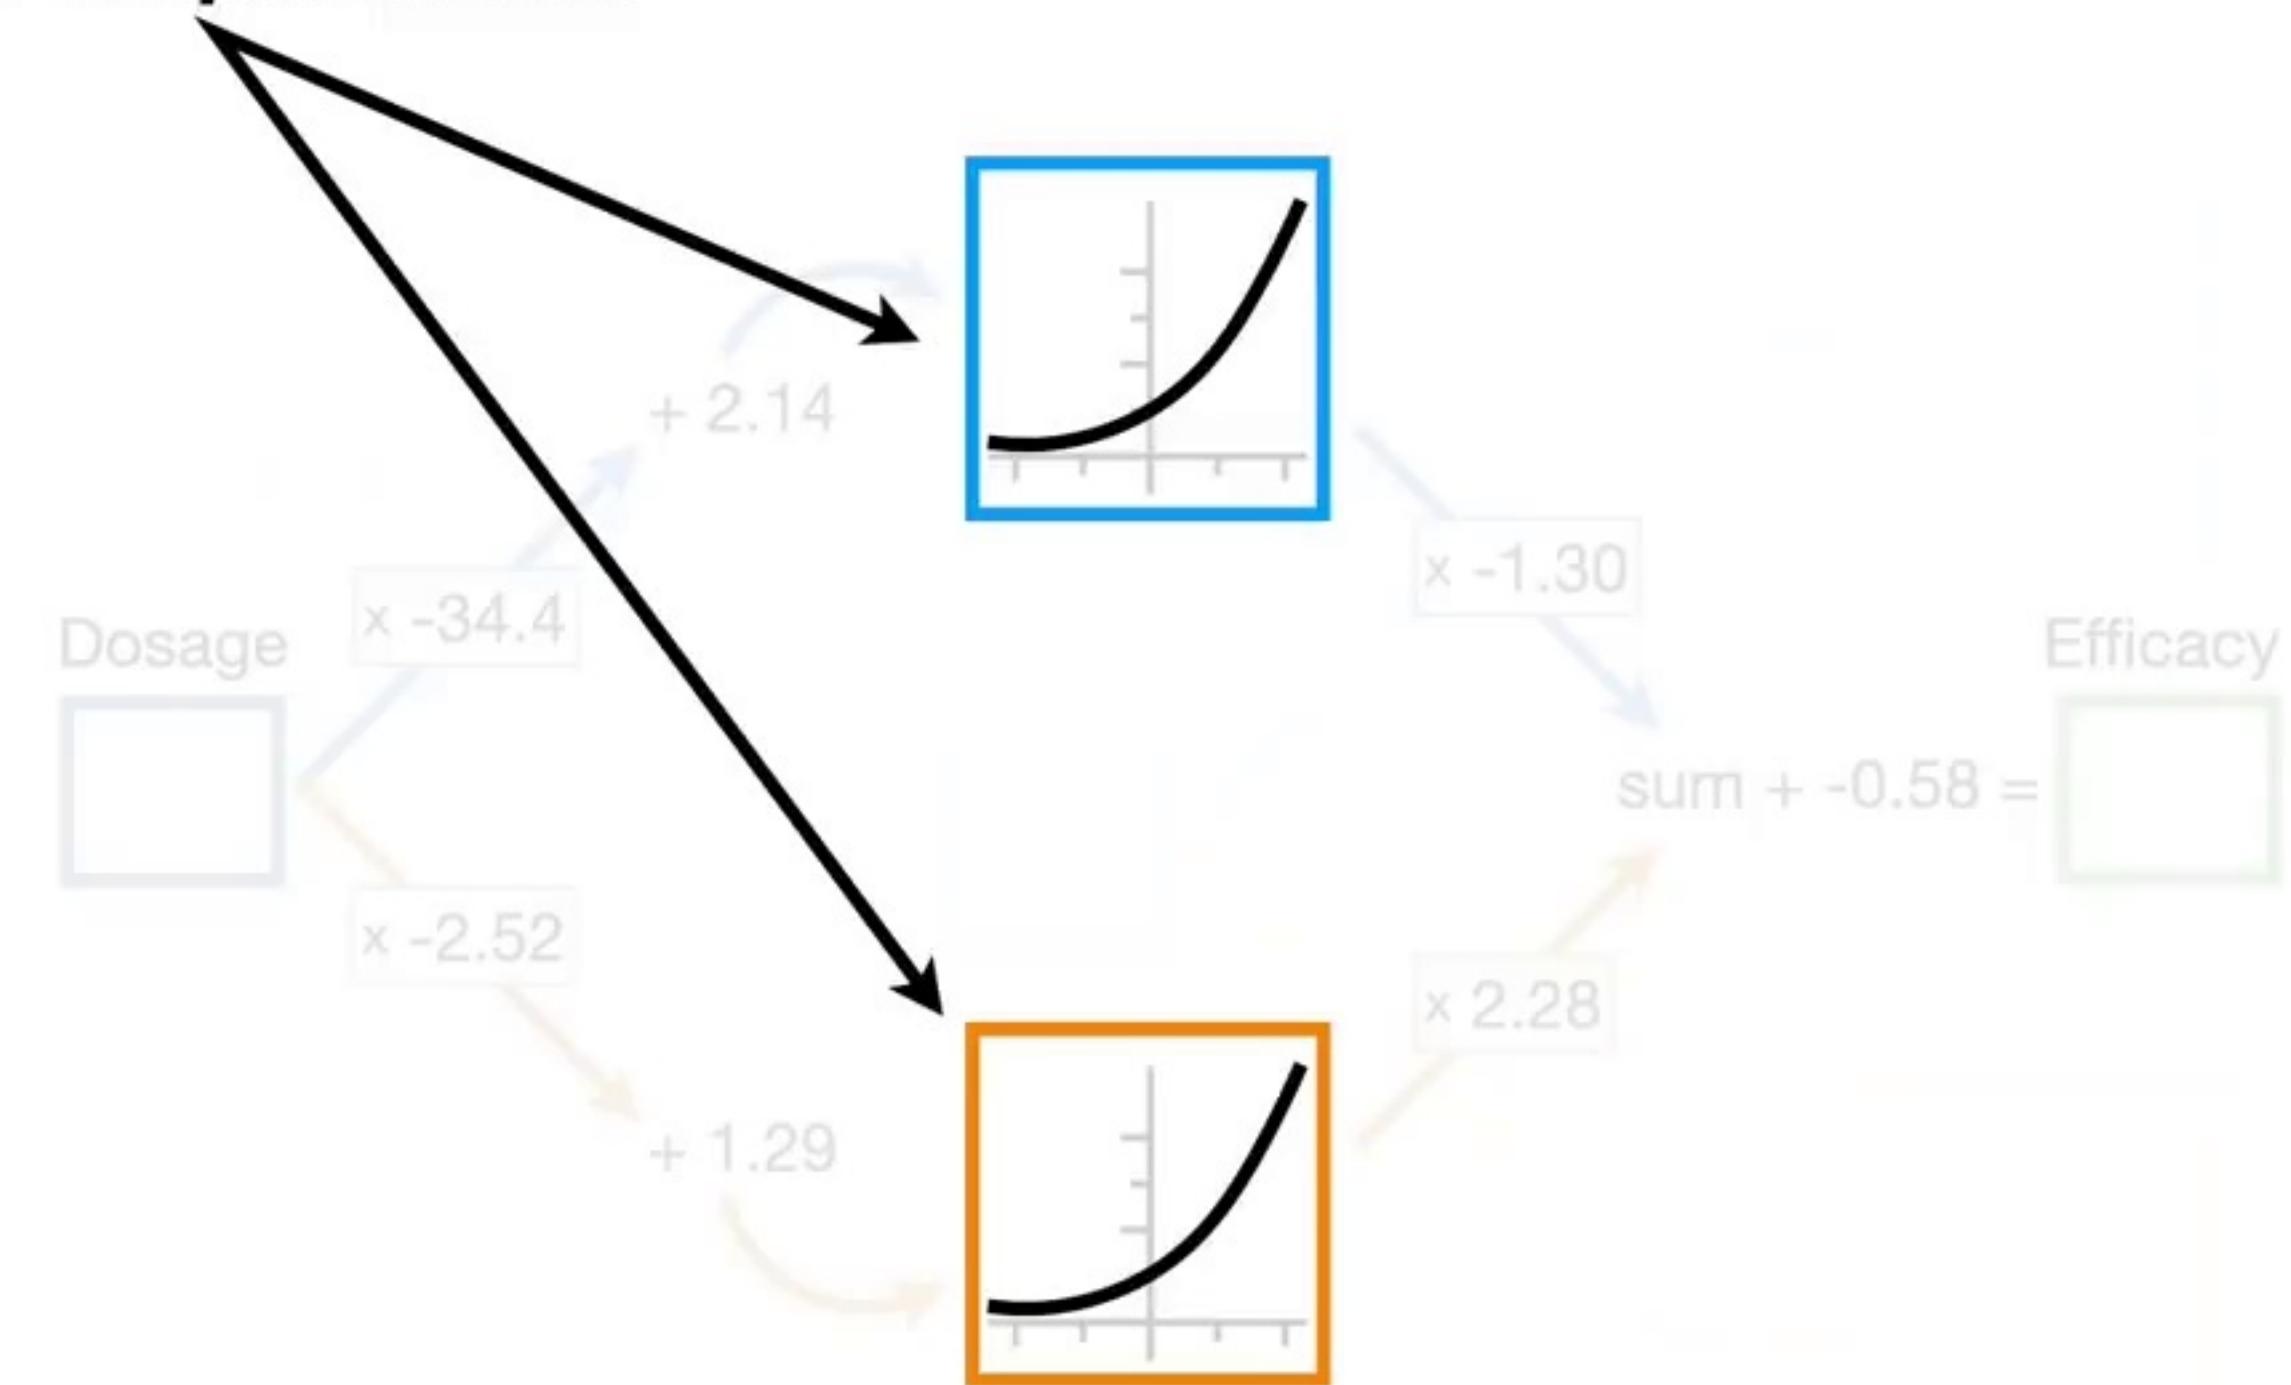


...only 1 Output Node to  
tell us the predicted  
**Effectiveness**...



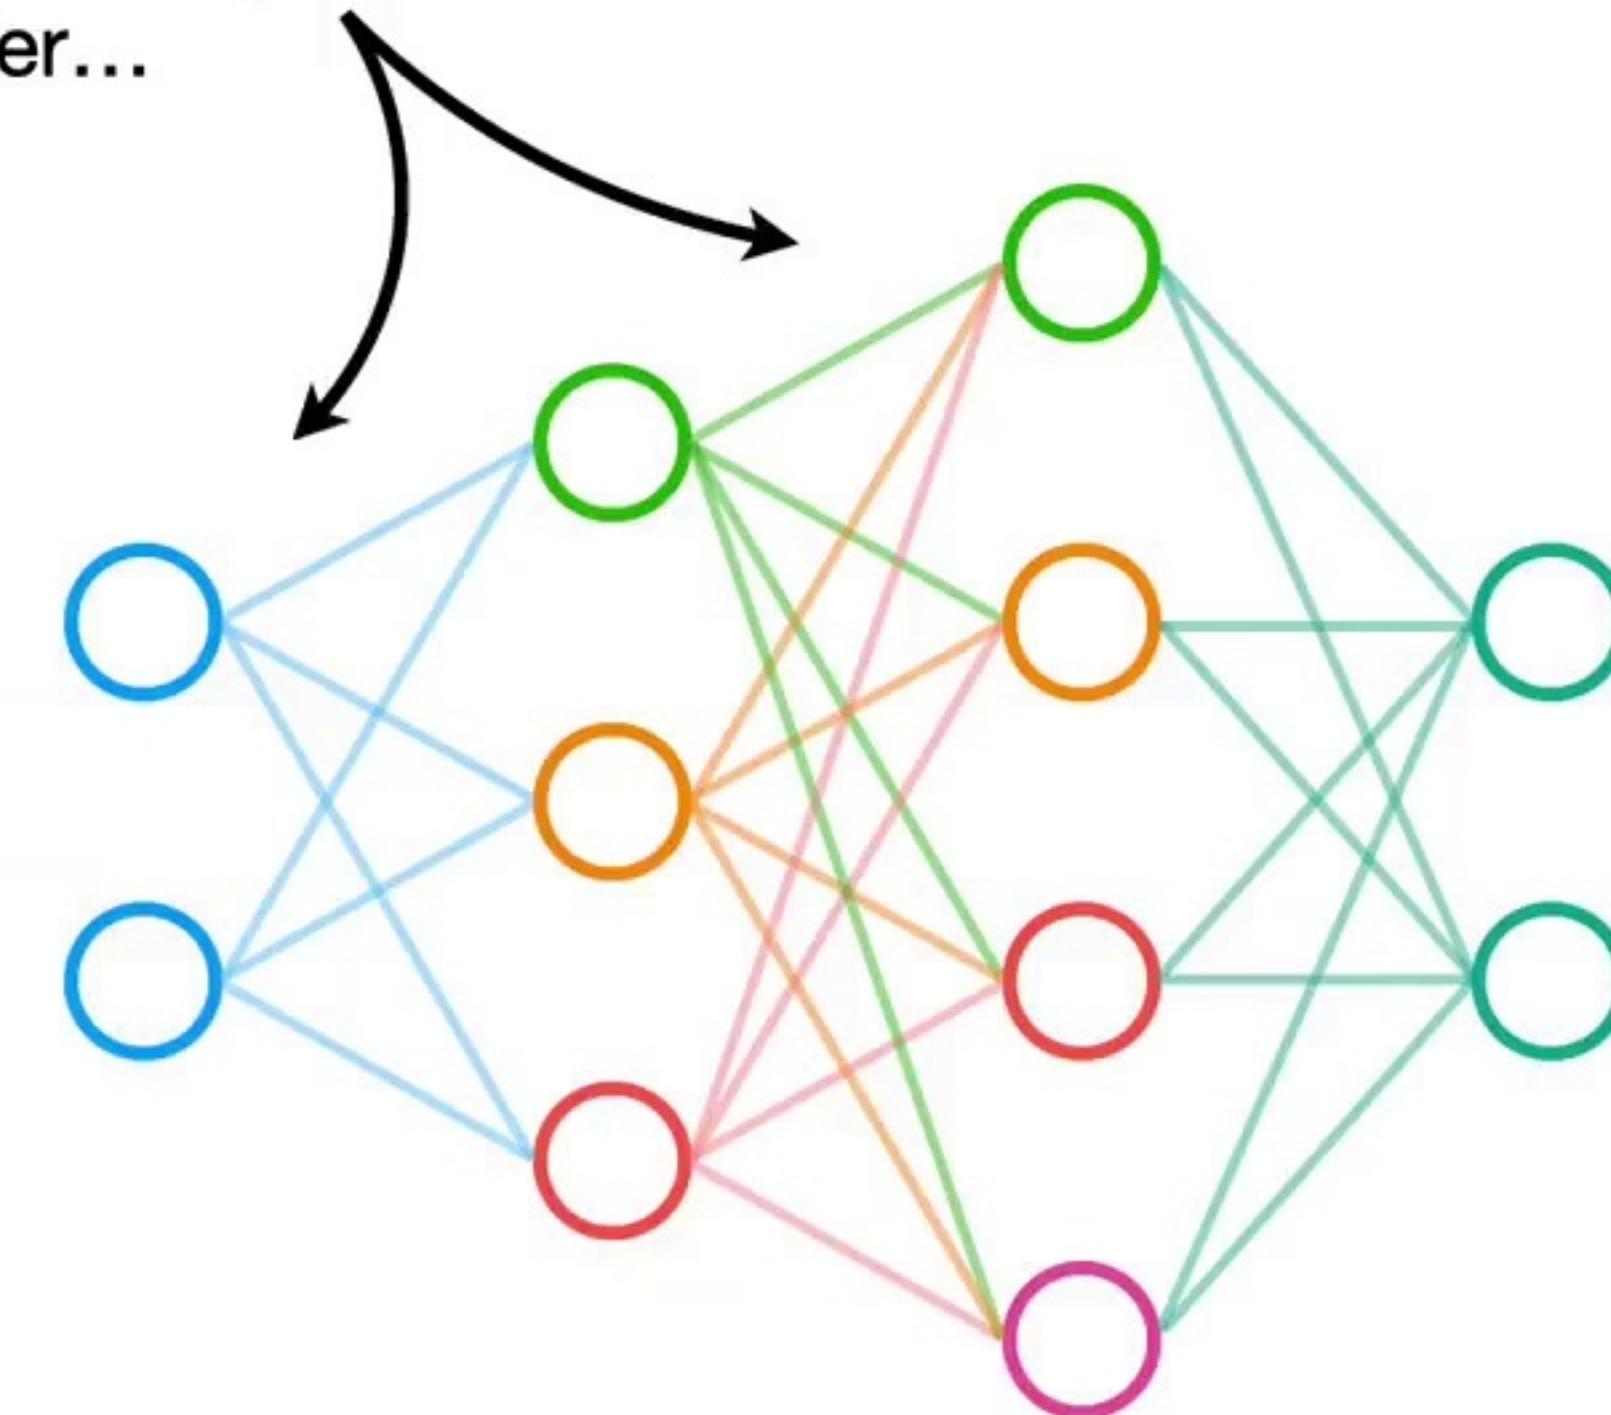


...and only **2 Nodes** between  
the **Input and Output Nodes**.



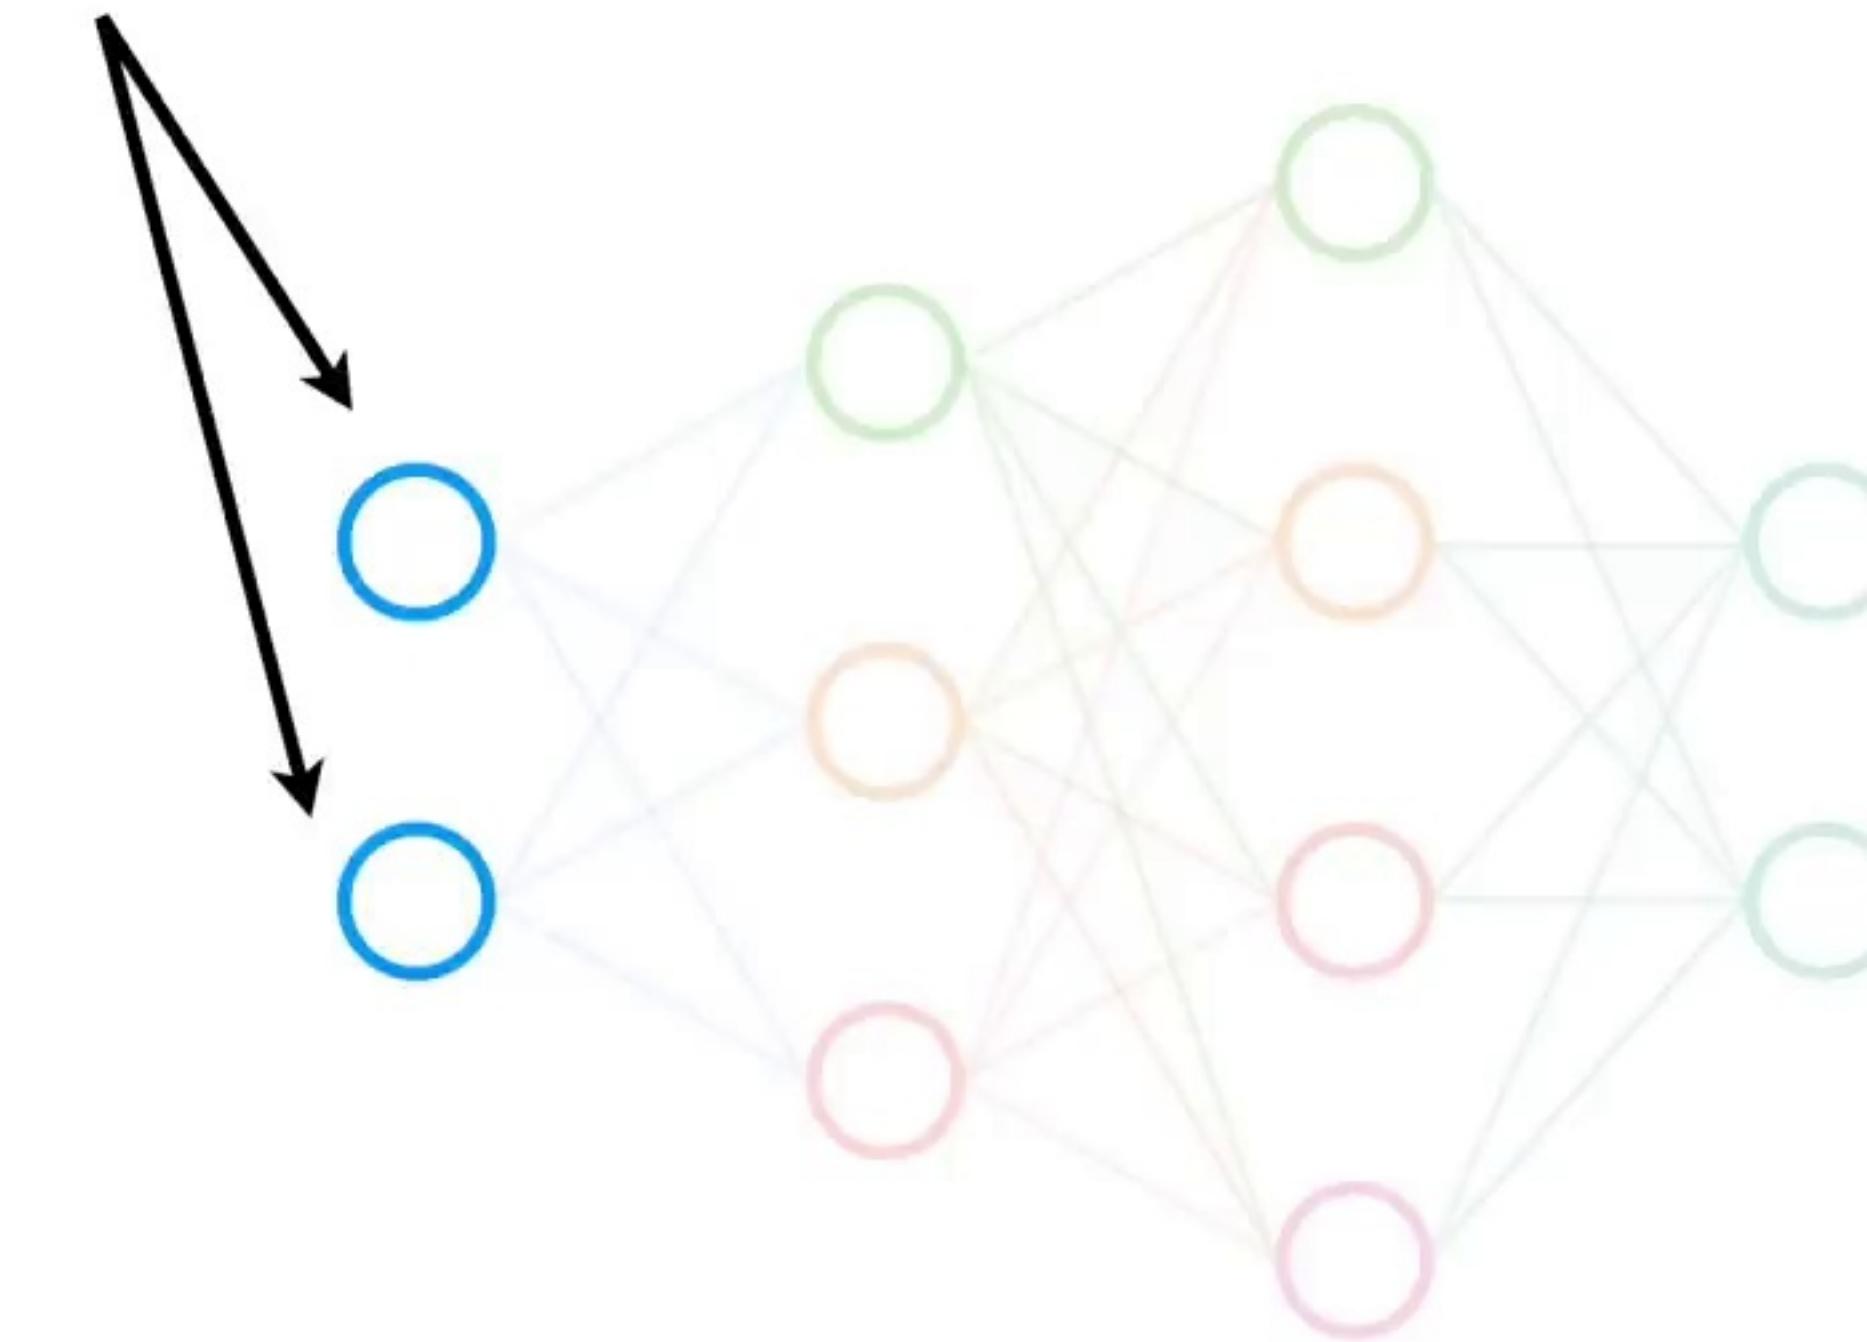


However, in practice, **Neural Networks** are usually much fancier...



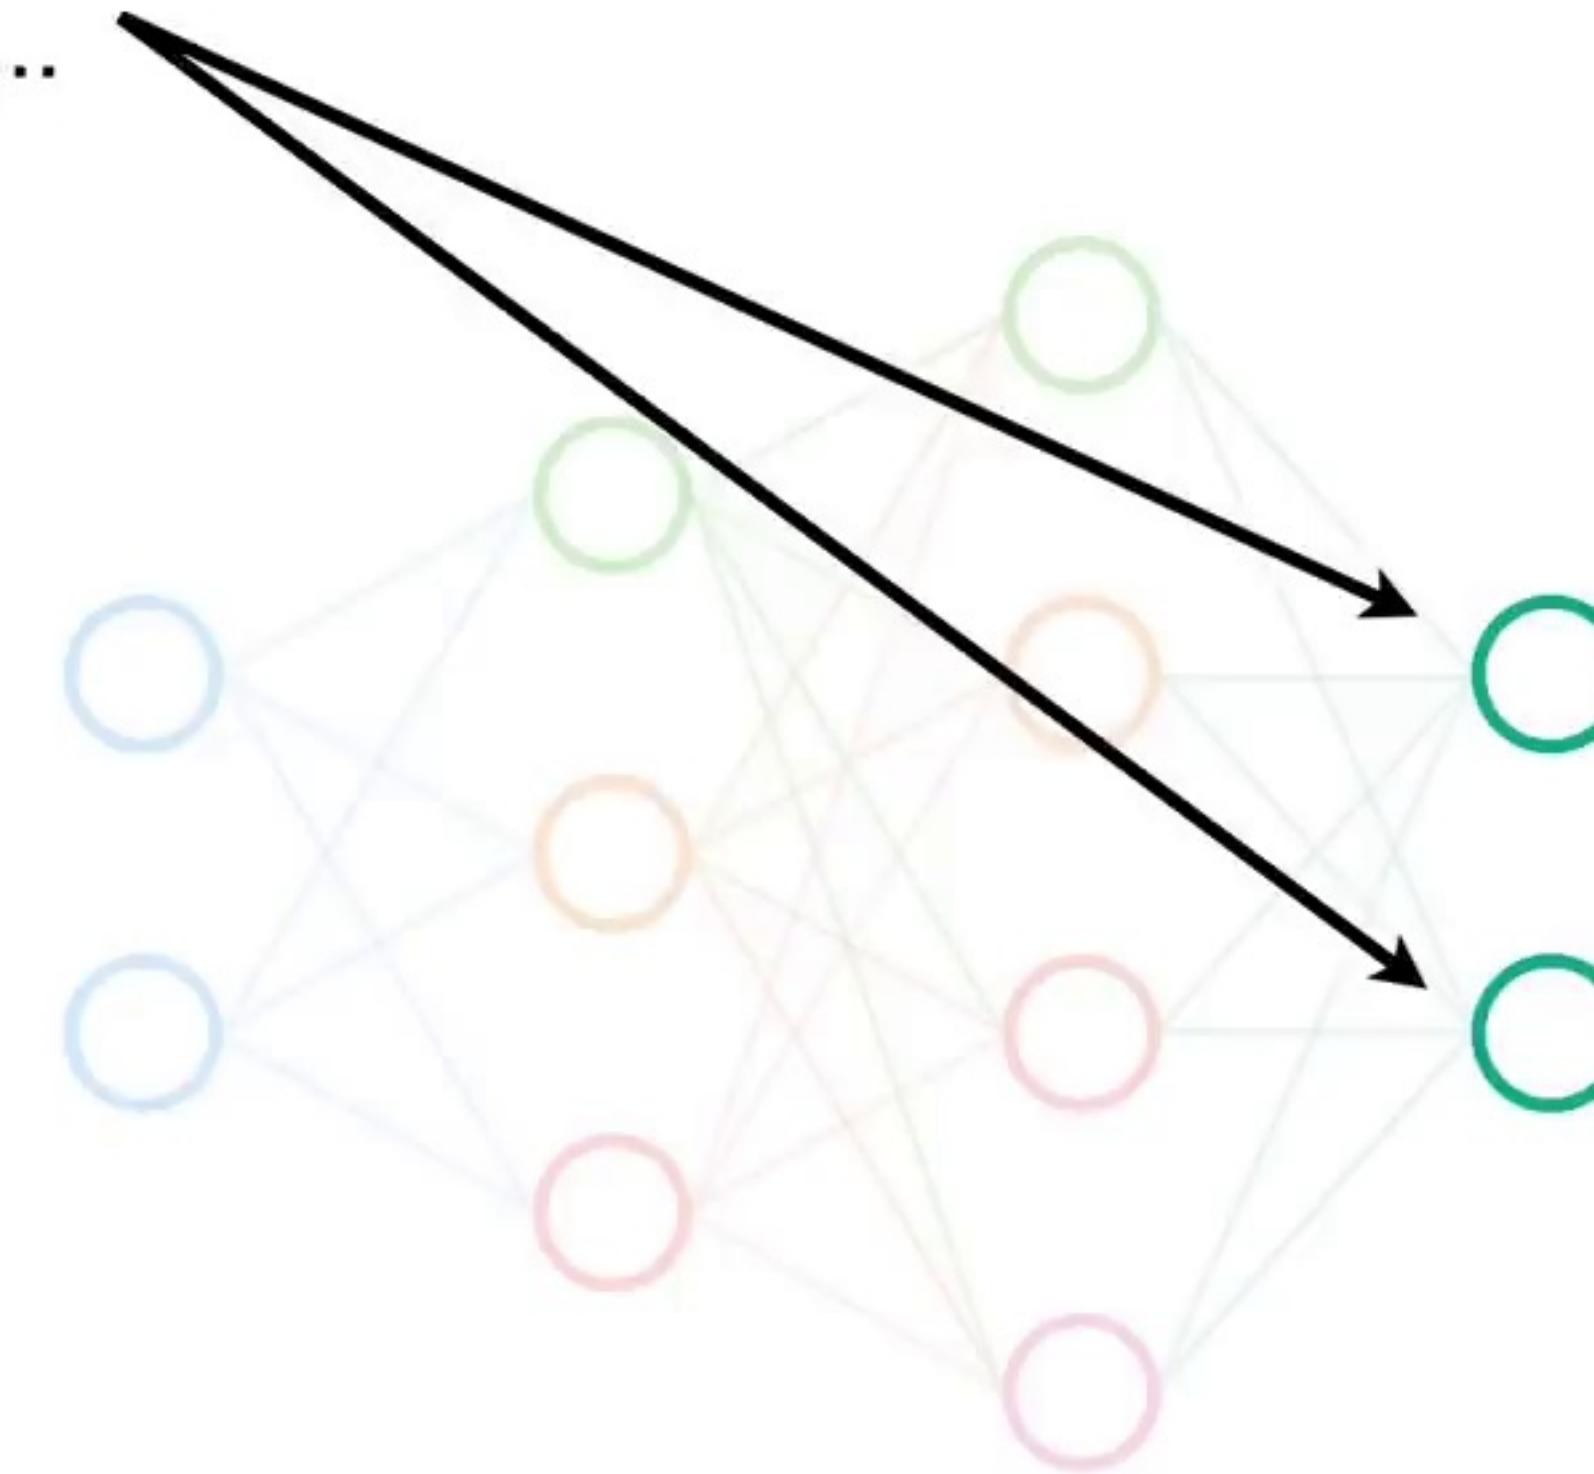


...and have more  
than 1 Input Node...



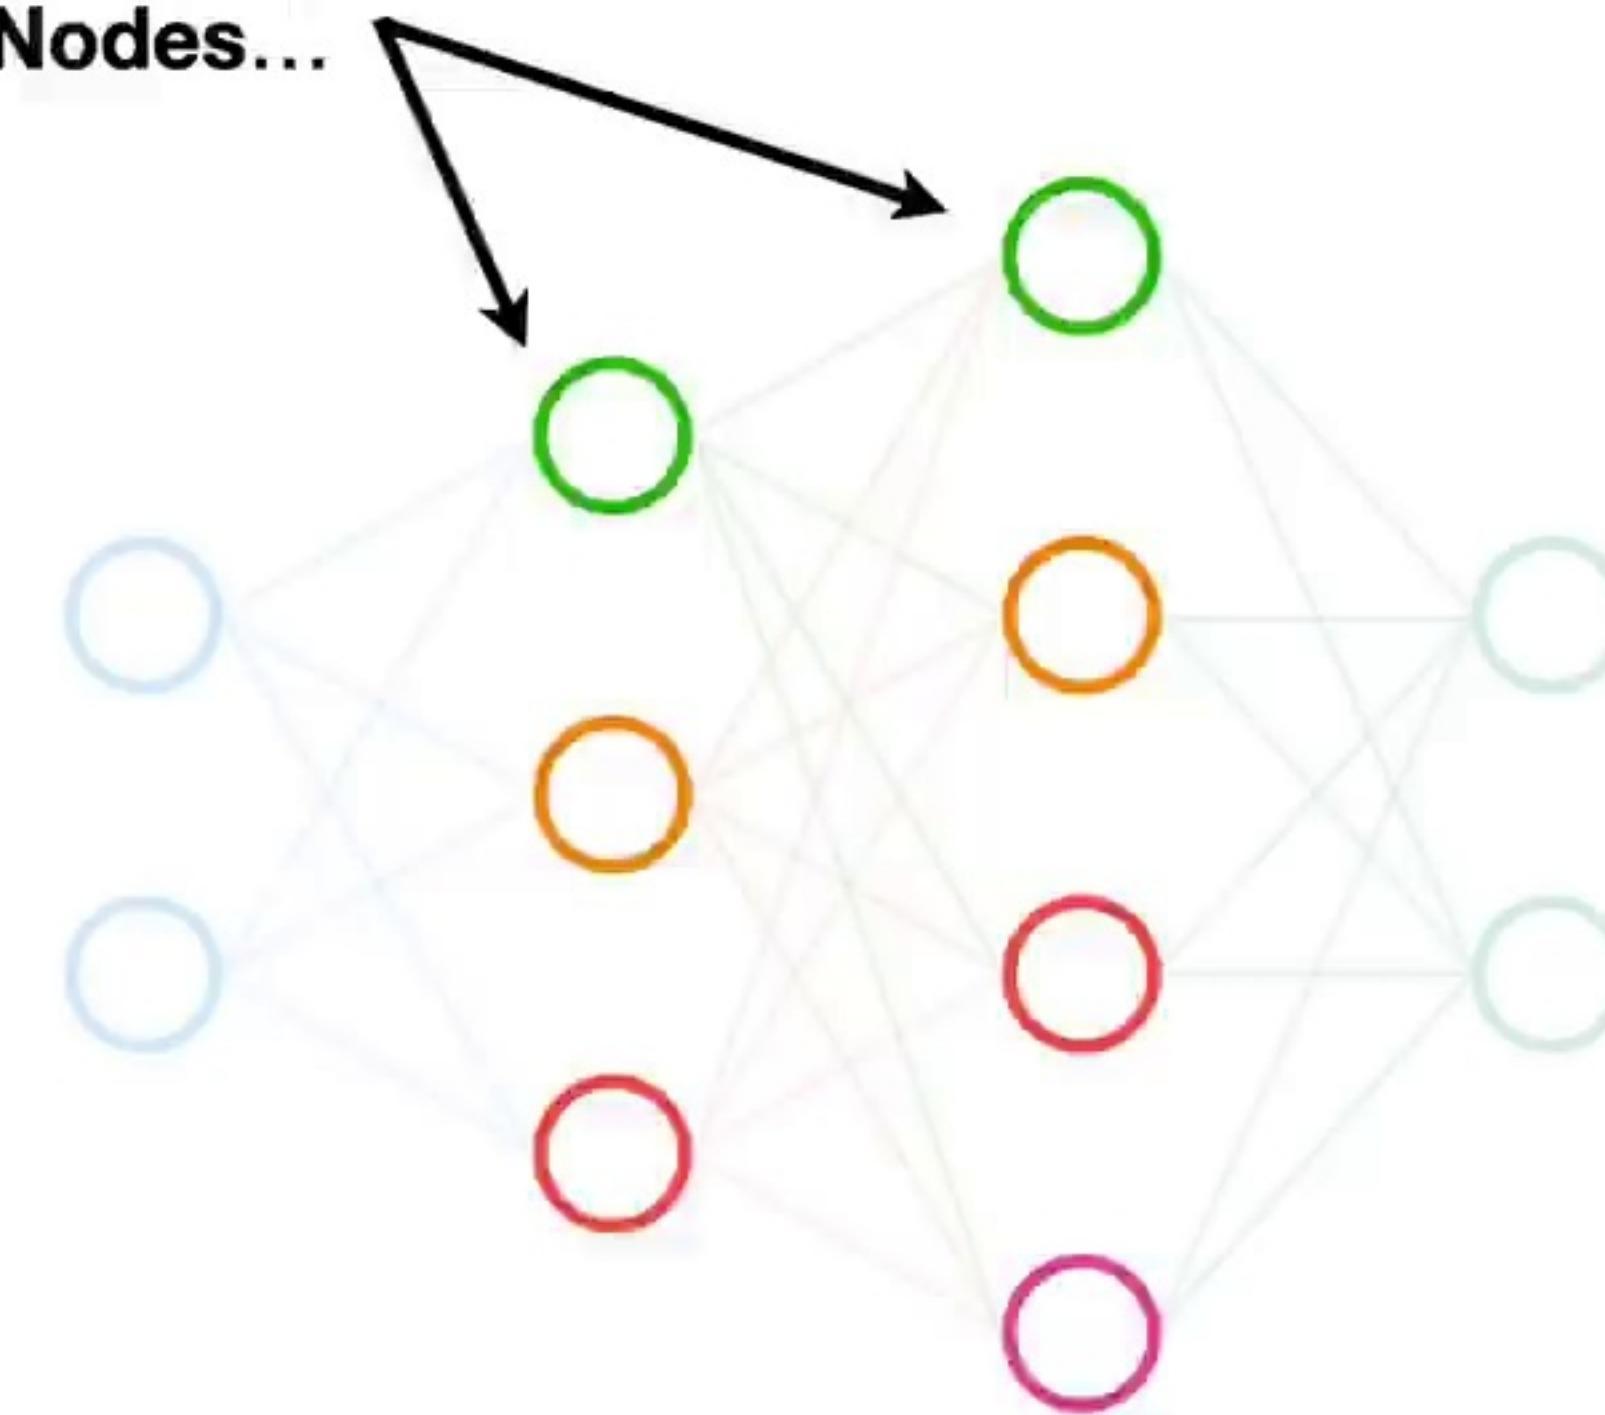


**...more than 1 Output  
Node...**



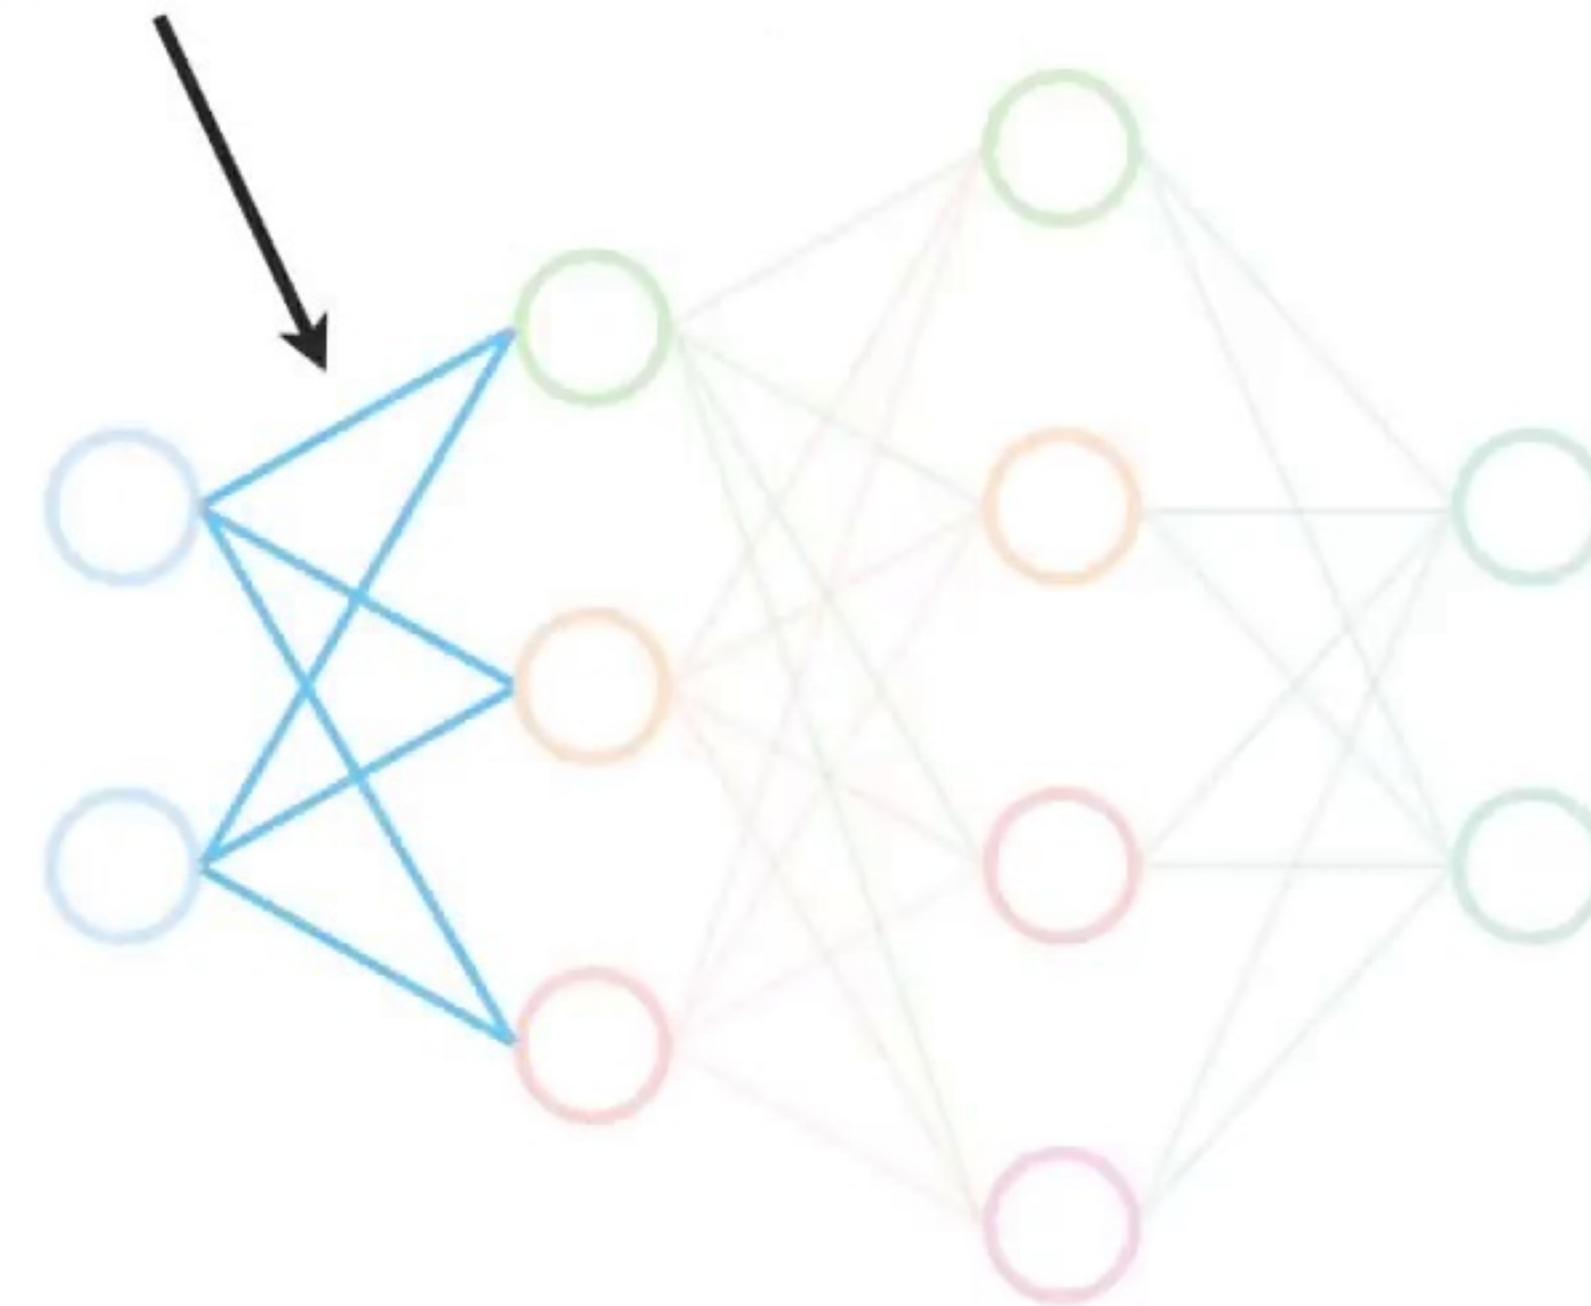


...different layers of  
**Nodes** between the **Input**  
and **Output Nodes**...



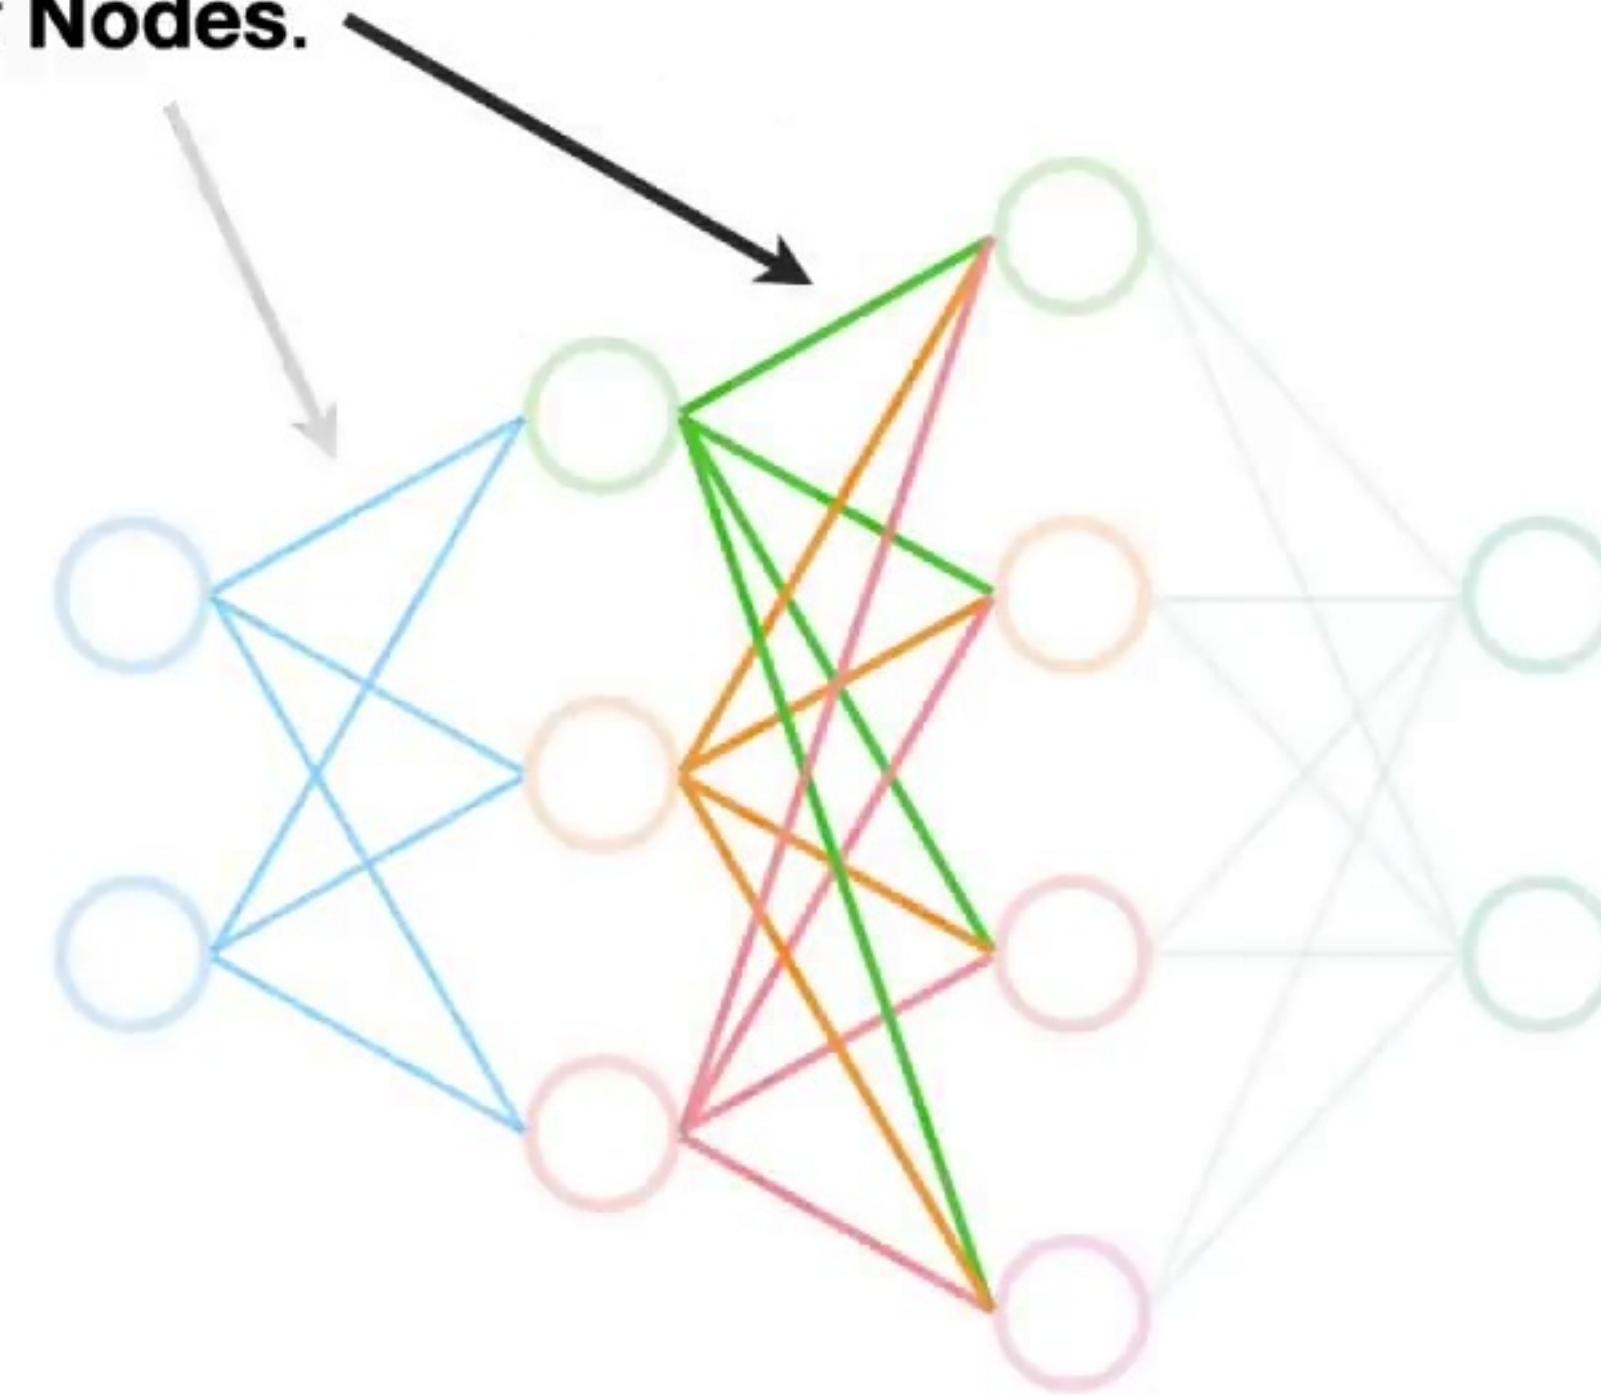


...and a spider web of connections between each layer of **Nodes**.



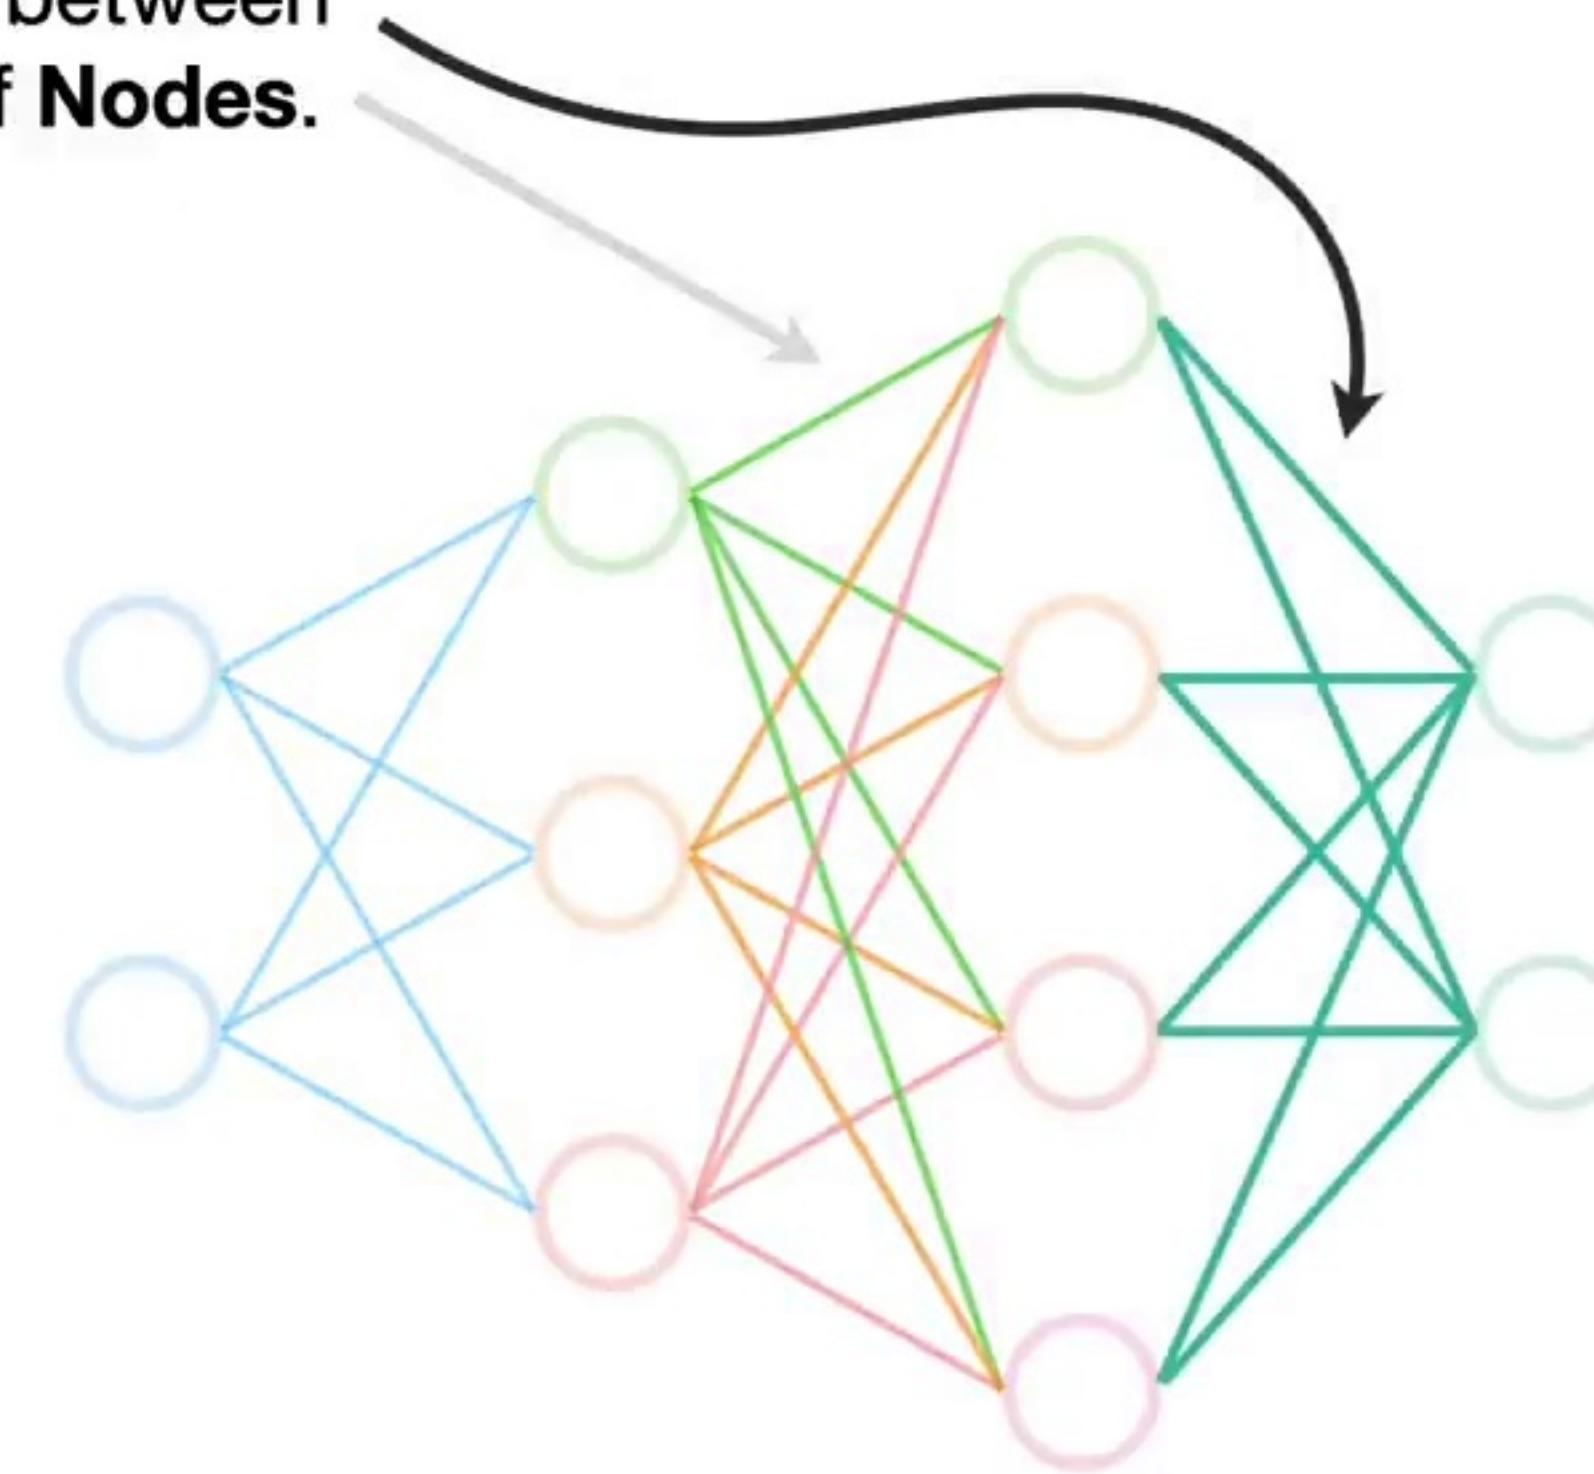


...and a spider web of connections between each layer of **Nodes**.



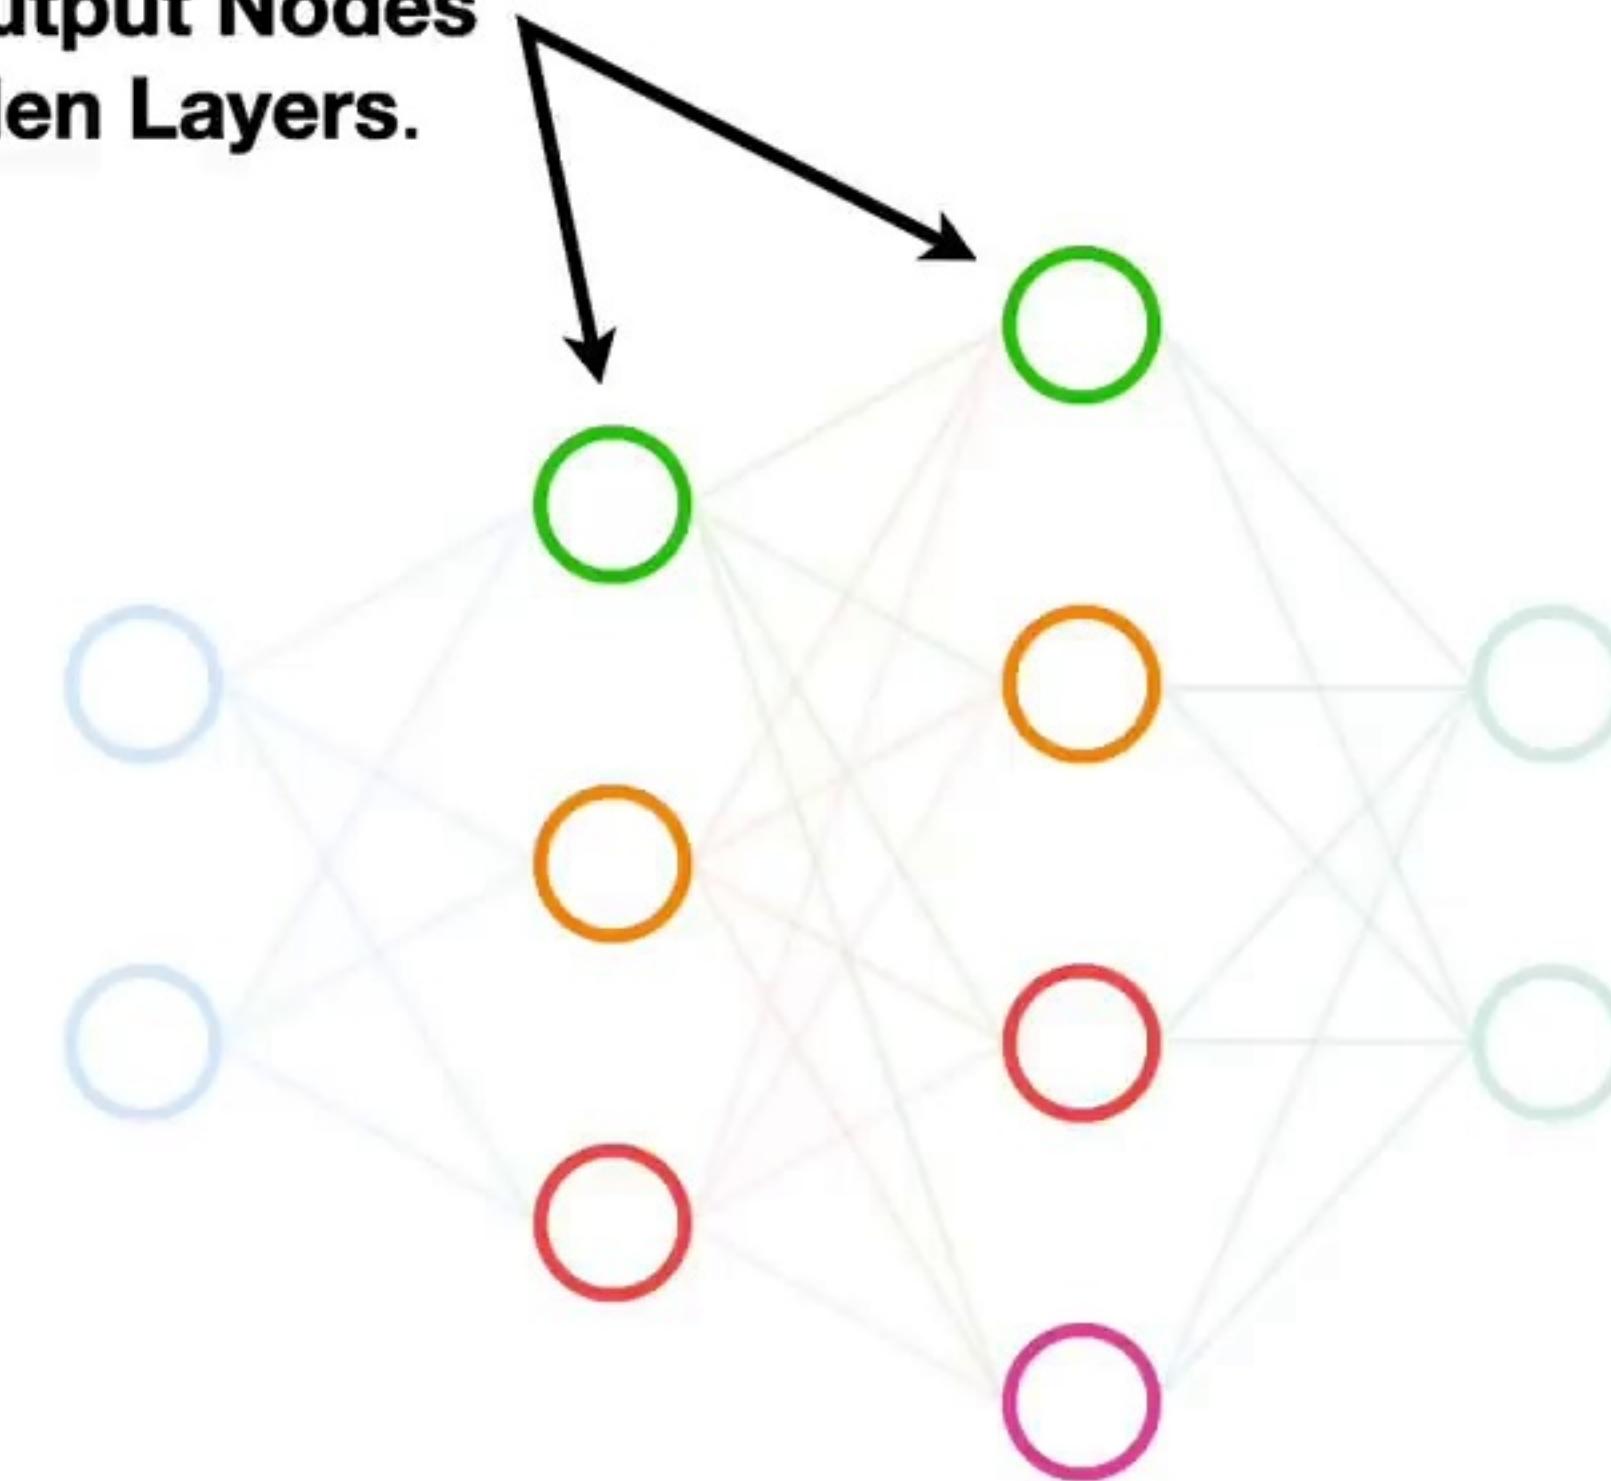


...and a spider web of connections between each layer of **Nodes**.



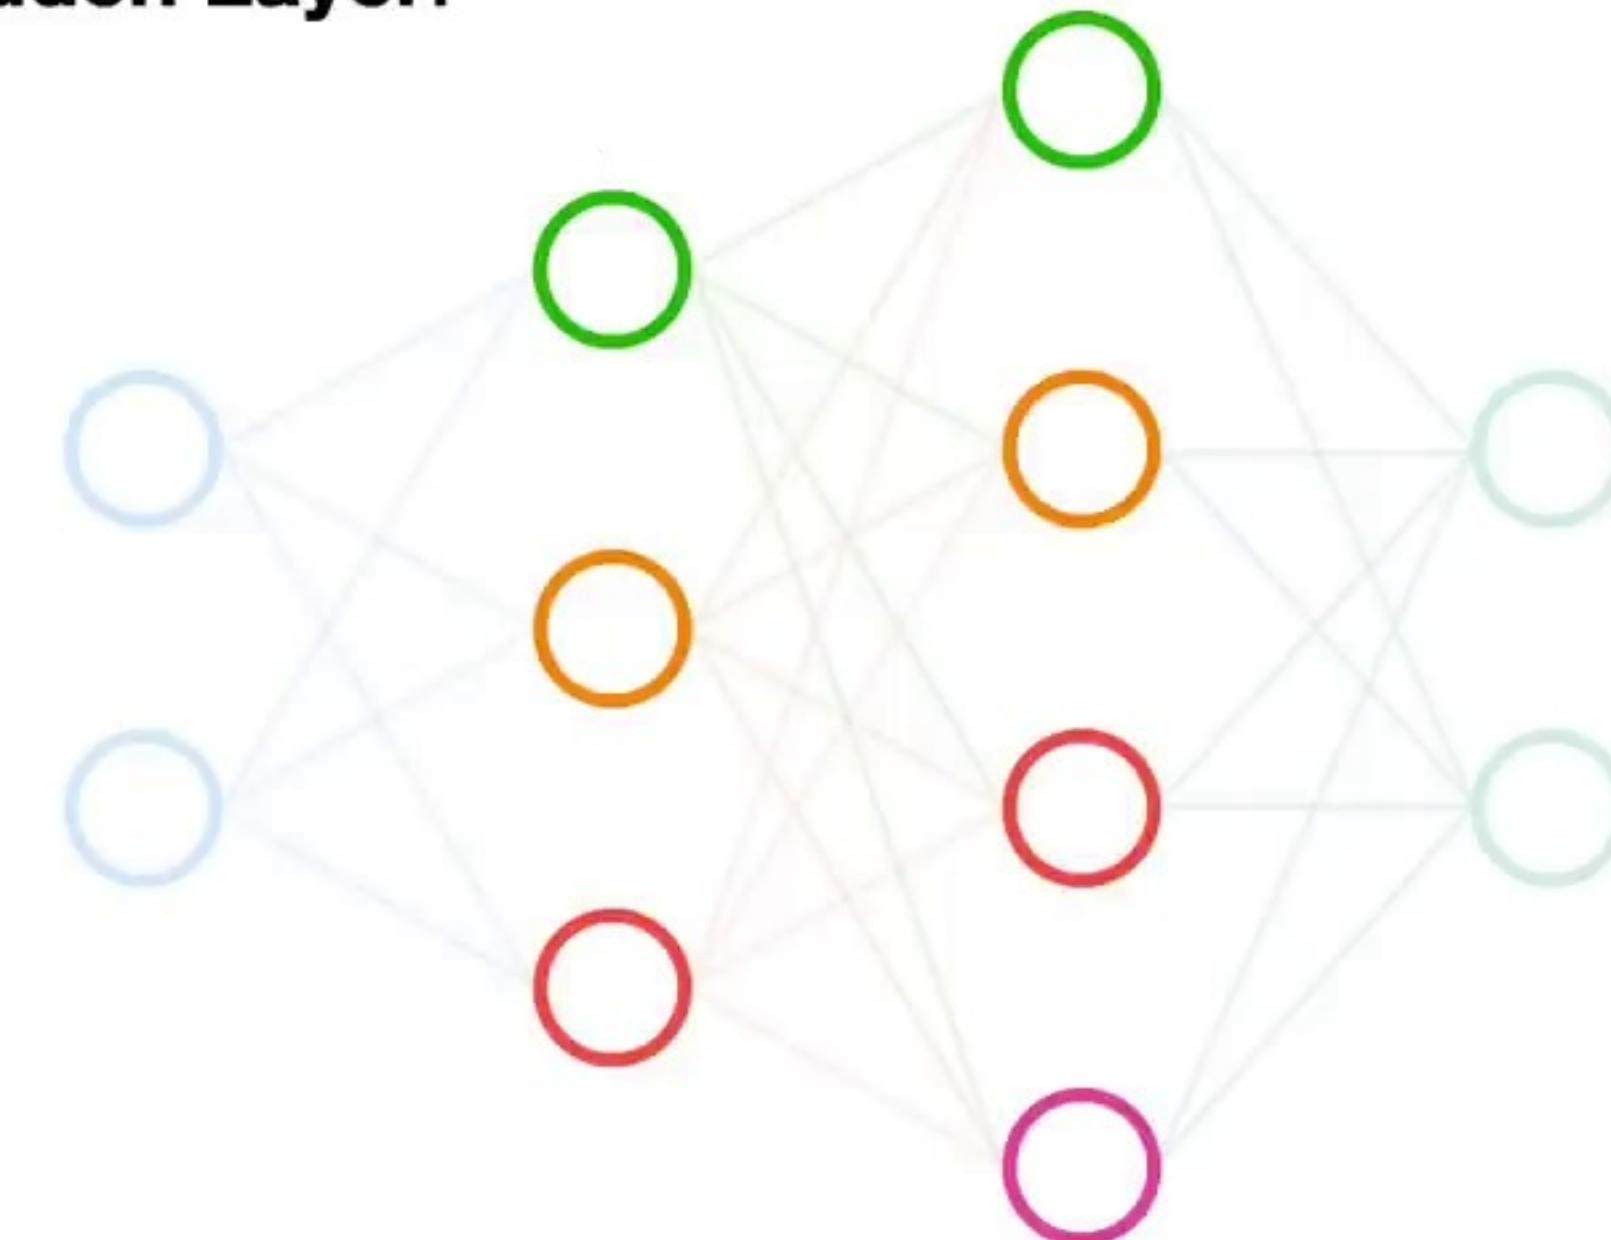


The layers of **Nodes** between  
the **Input** and **Output Nodes**  
are called **Hidden Layers**.



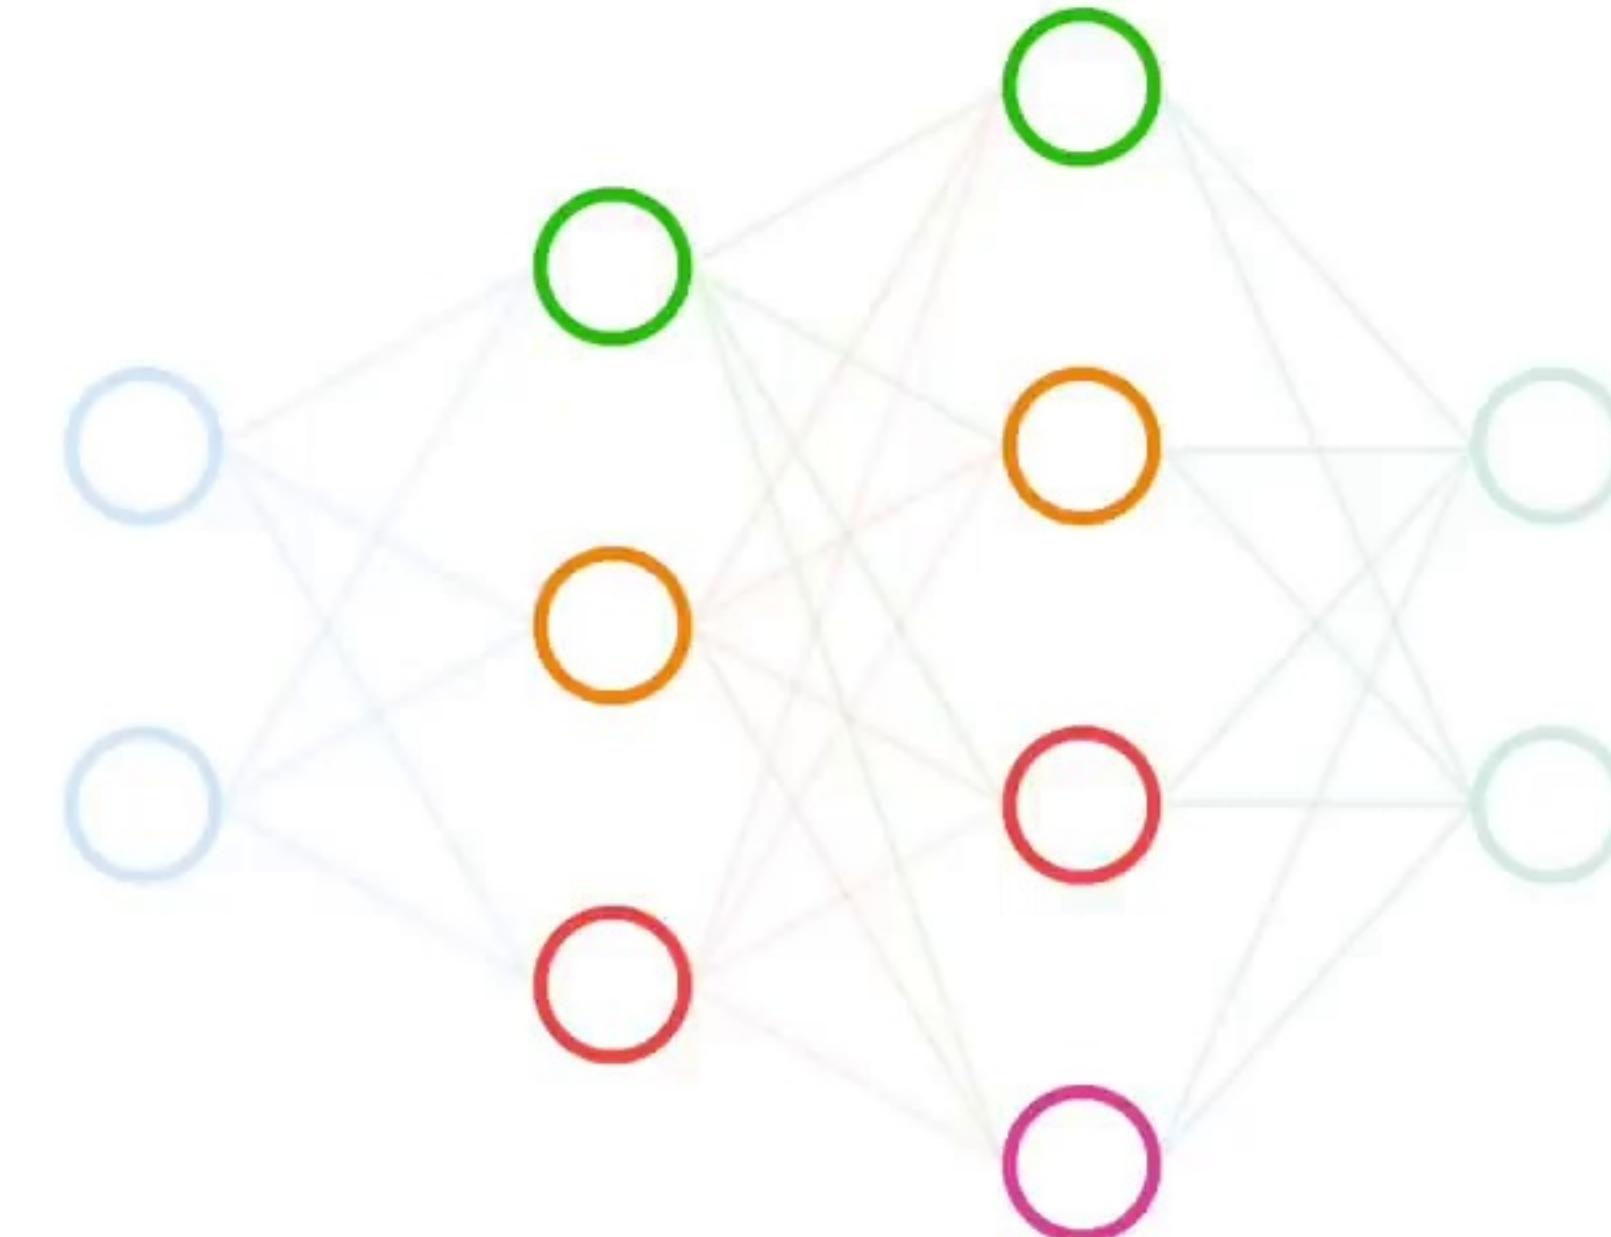


When you build a **Neural Network**,  
one of the first things you do is  
decide how many **Hidden Layers**  
you want, and how many **Nodes**  
go into each **Hidden Layer**.



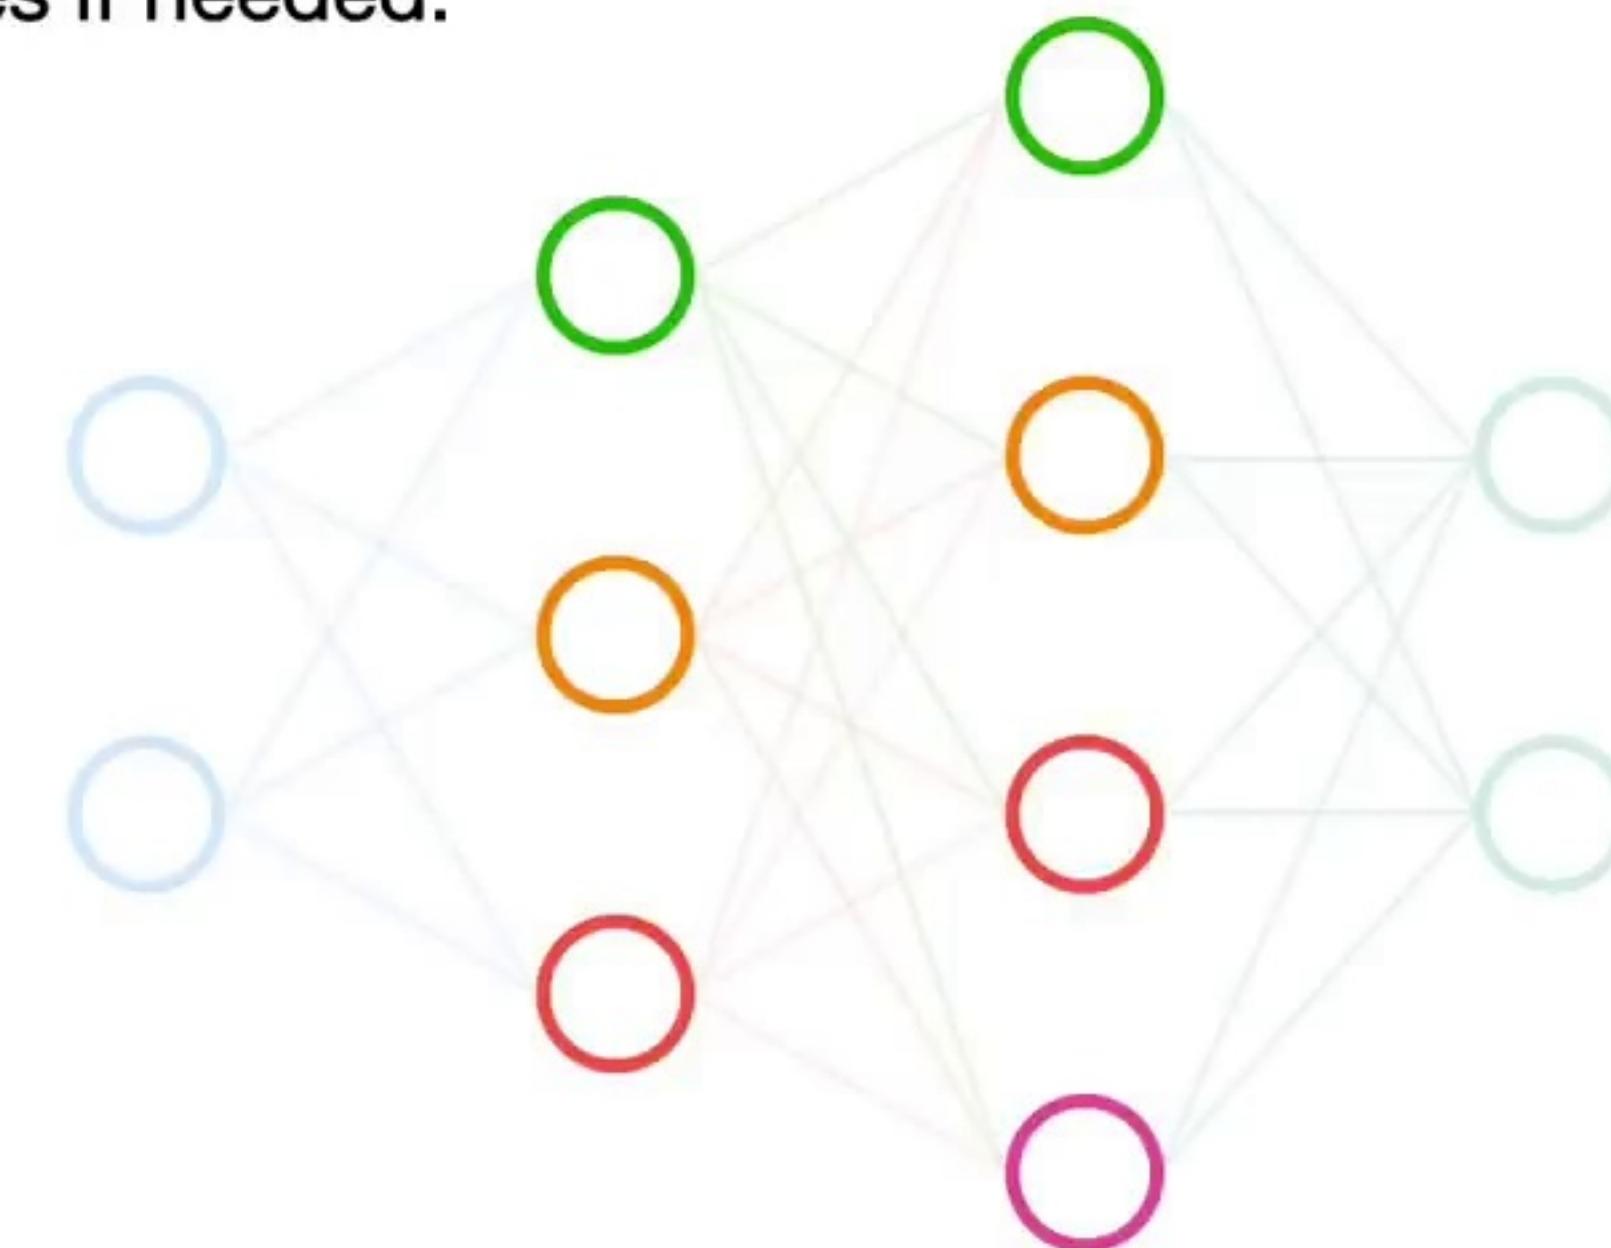


Although there are rules of thumb  
for making decisions about the  
**Hidden Layers...**



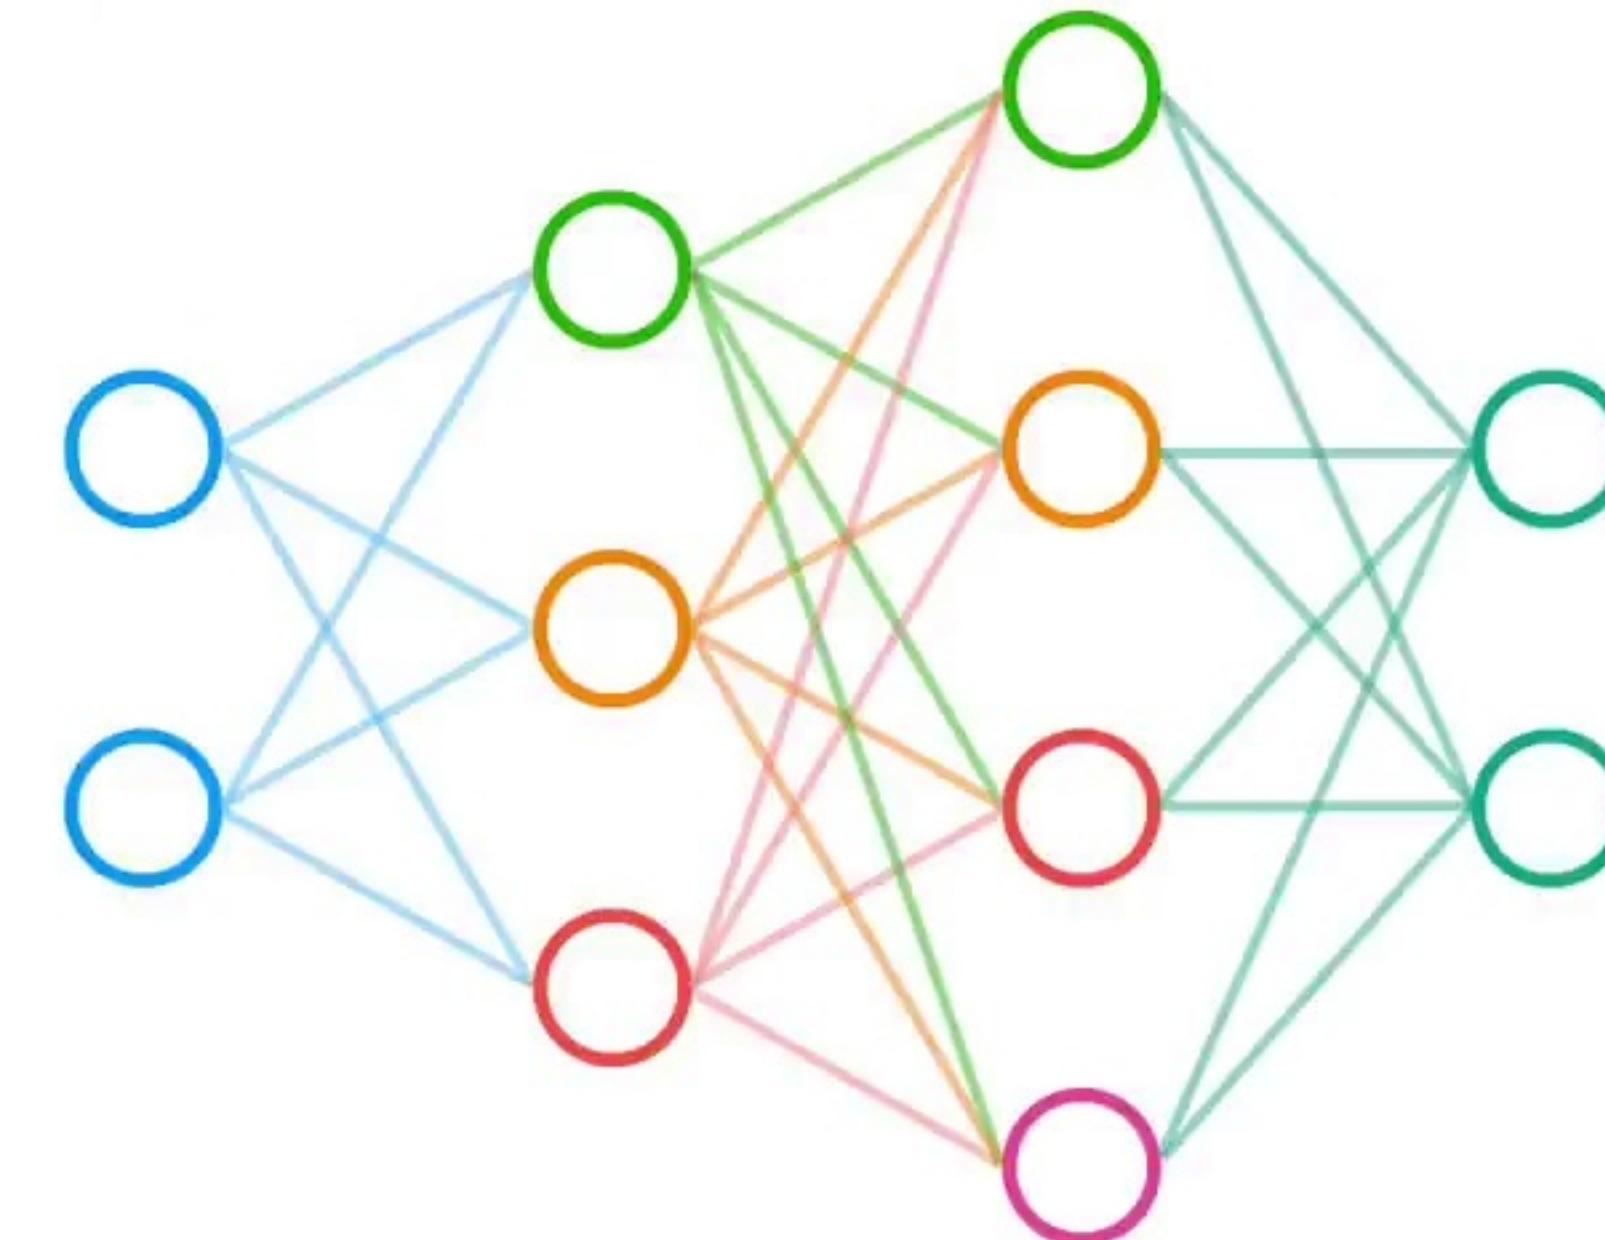


...you essentially make a guess  
and see how well the **Neural**  
**Network** performs, adding more  
layers and nodes if needed.



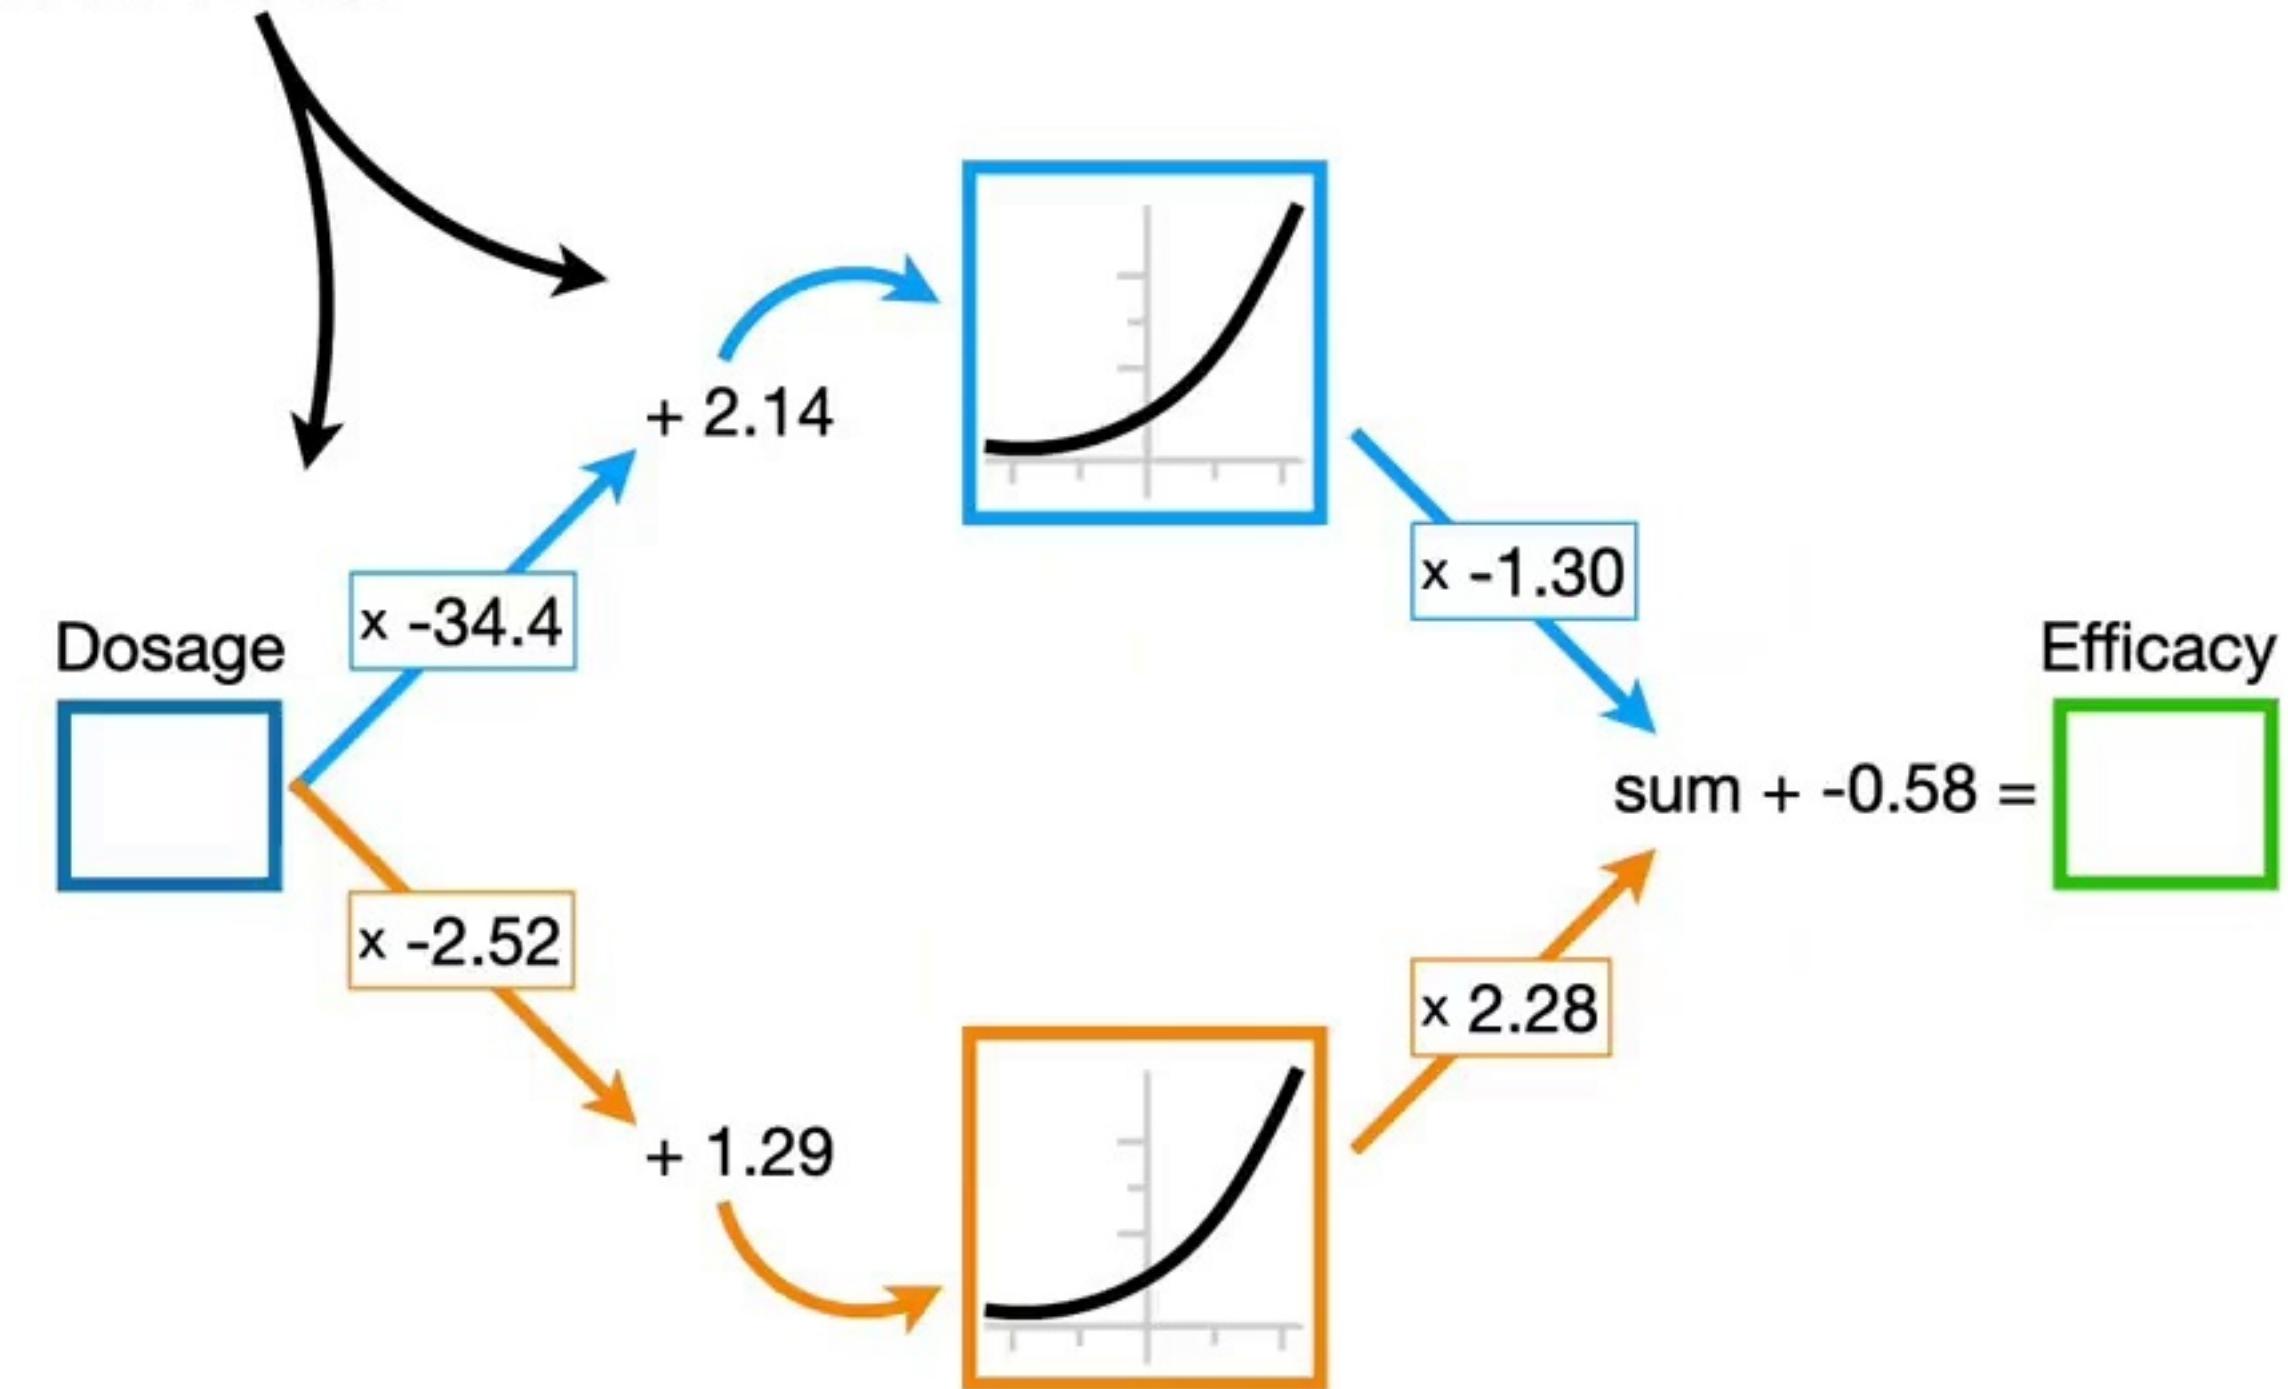


Now, even though this **Neural Network** looks fancy, it is still made from the same parts...



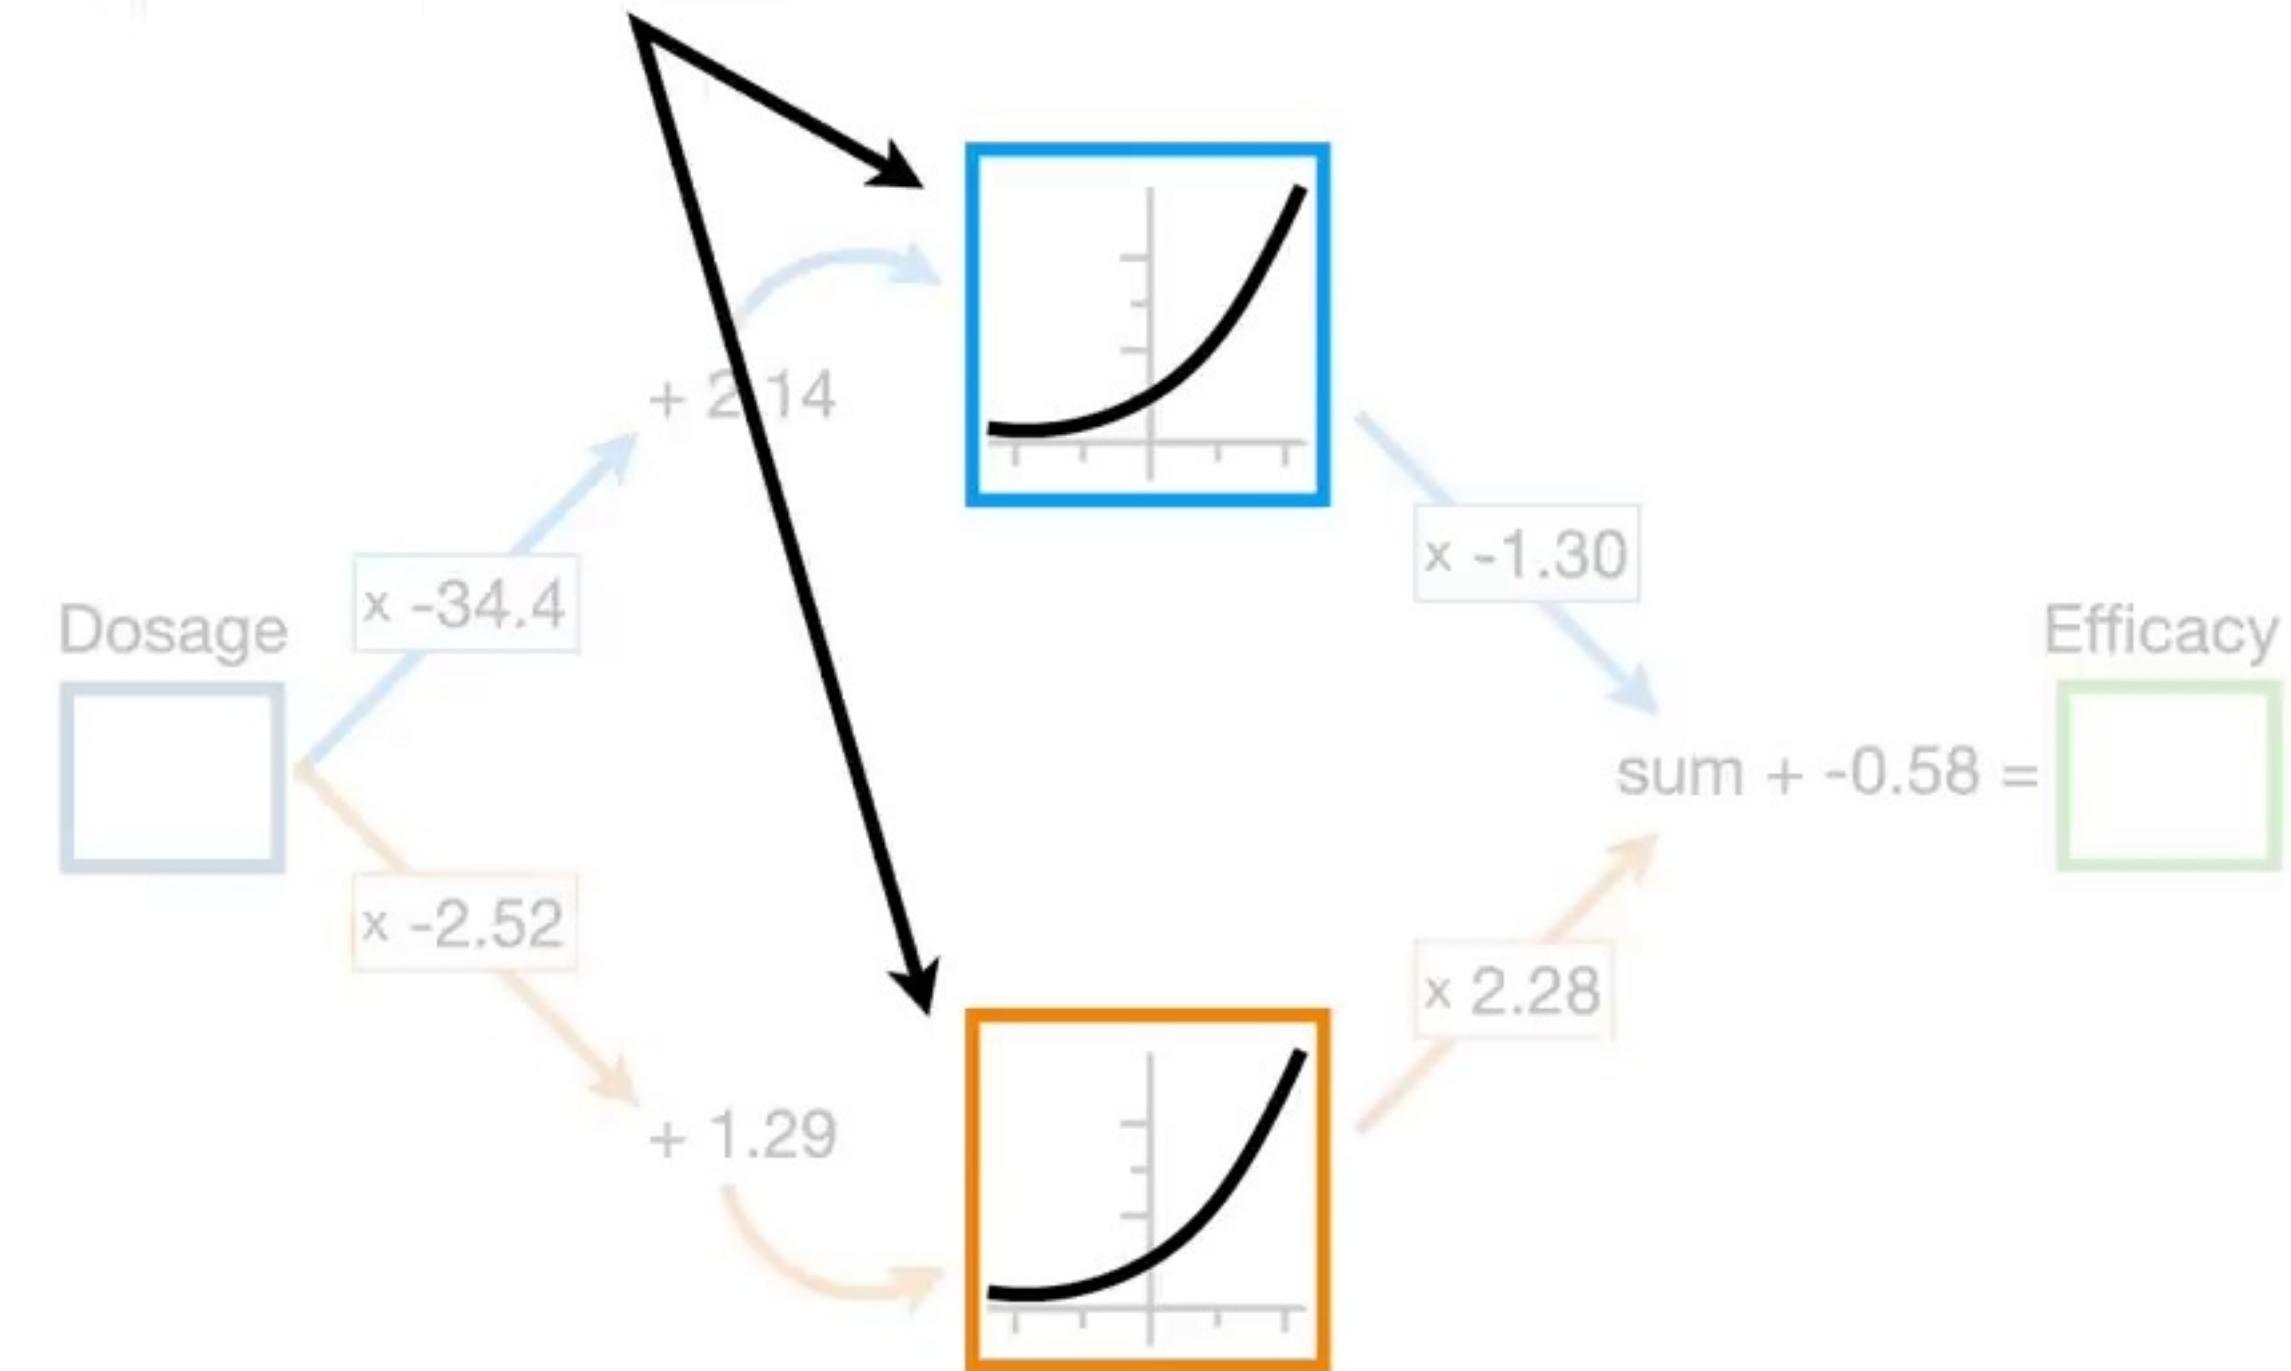


...used in this simple  
**Neural Network**...



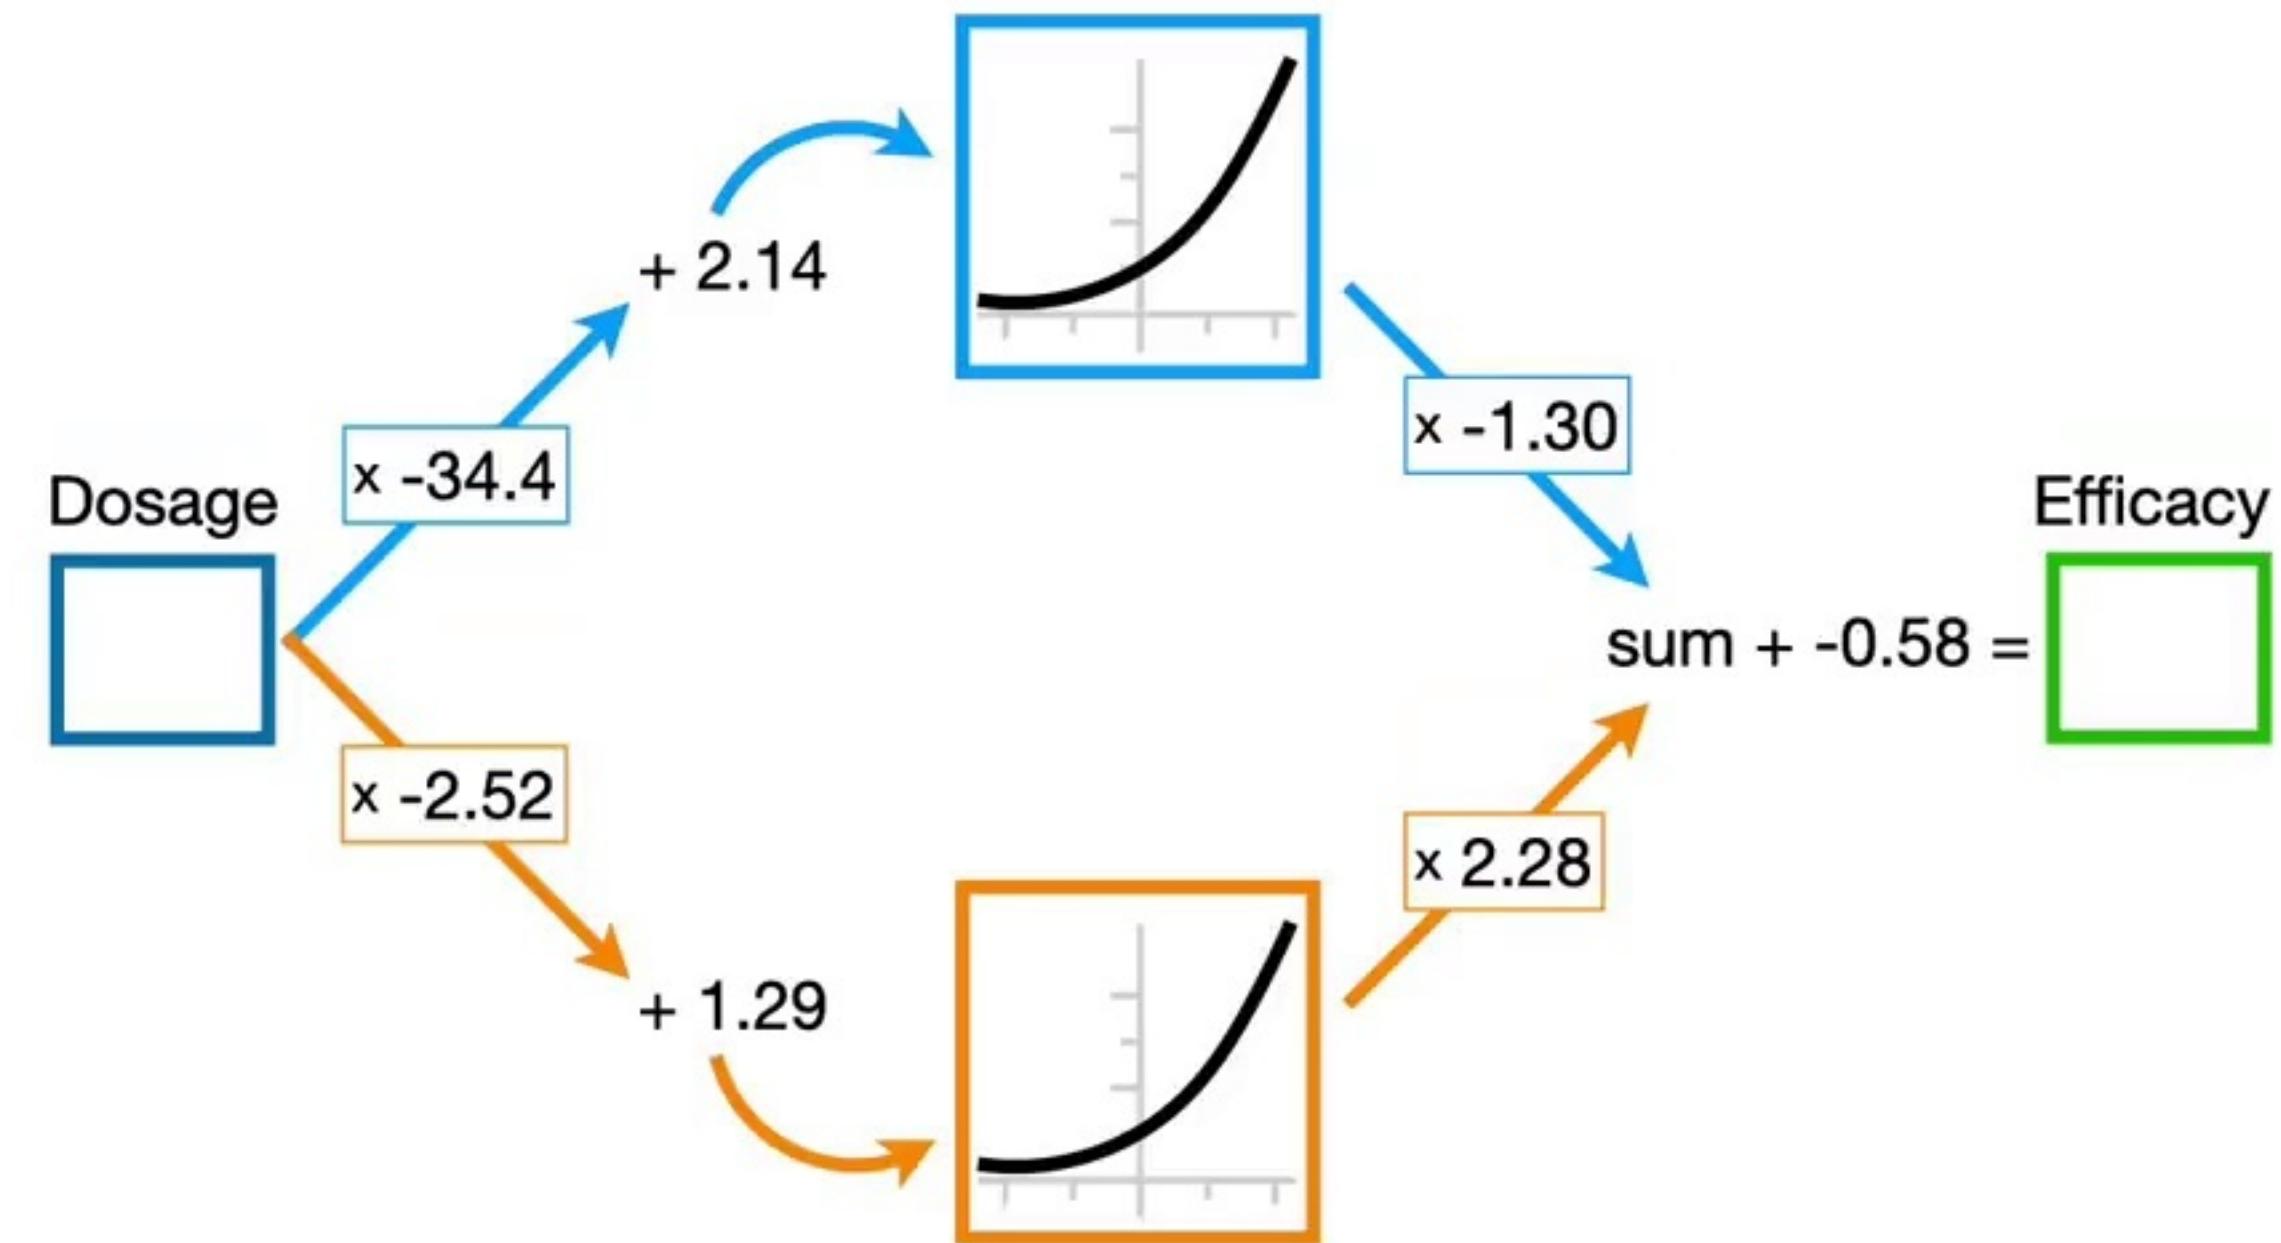


...which only has 1 Hidden  
Layer with 2 Nodes.



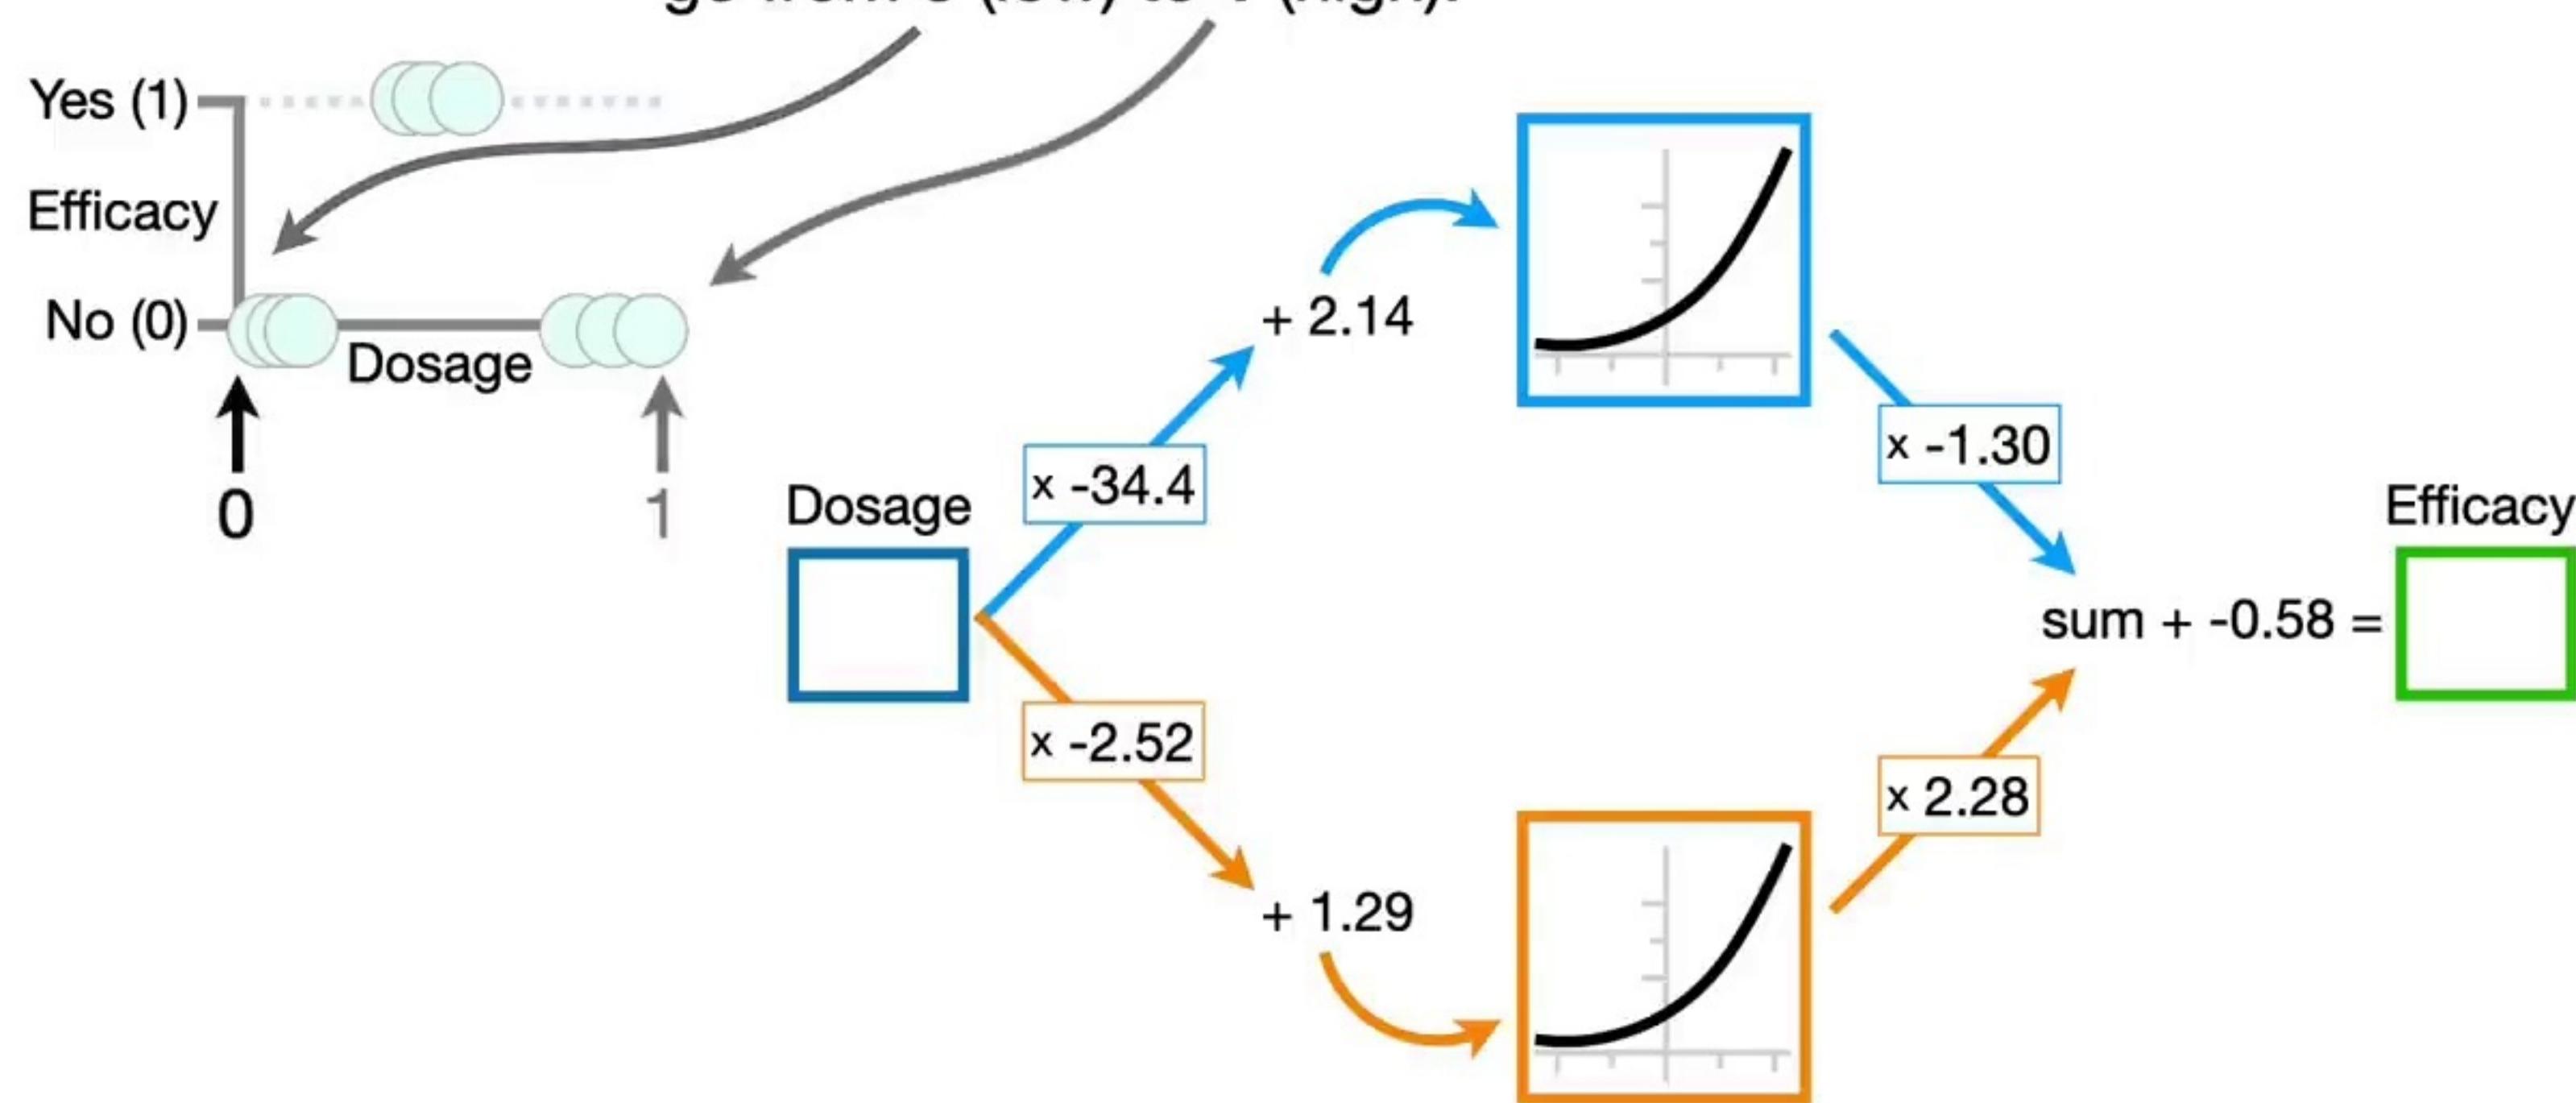


So let's learn how this **Neural Network** creates new shapes from the **curved or bent lines** in the **Hidden Layer**...



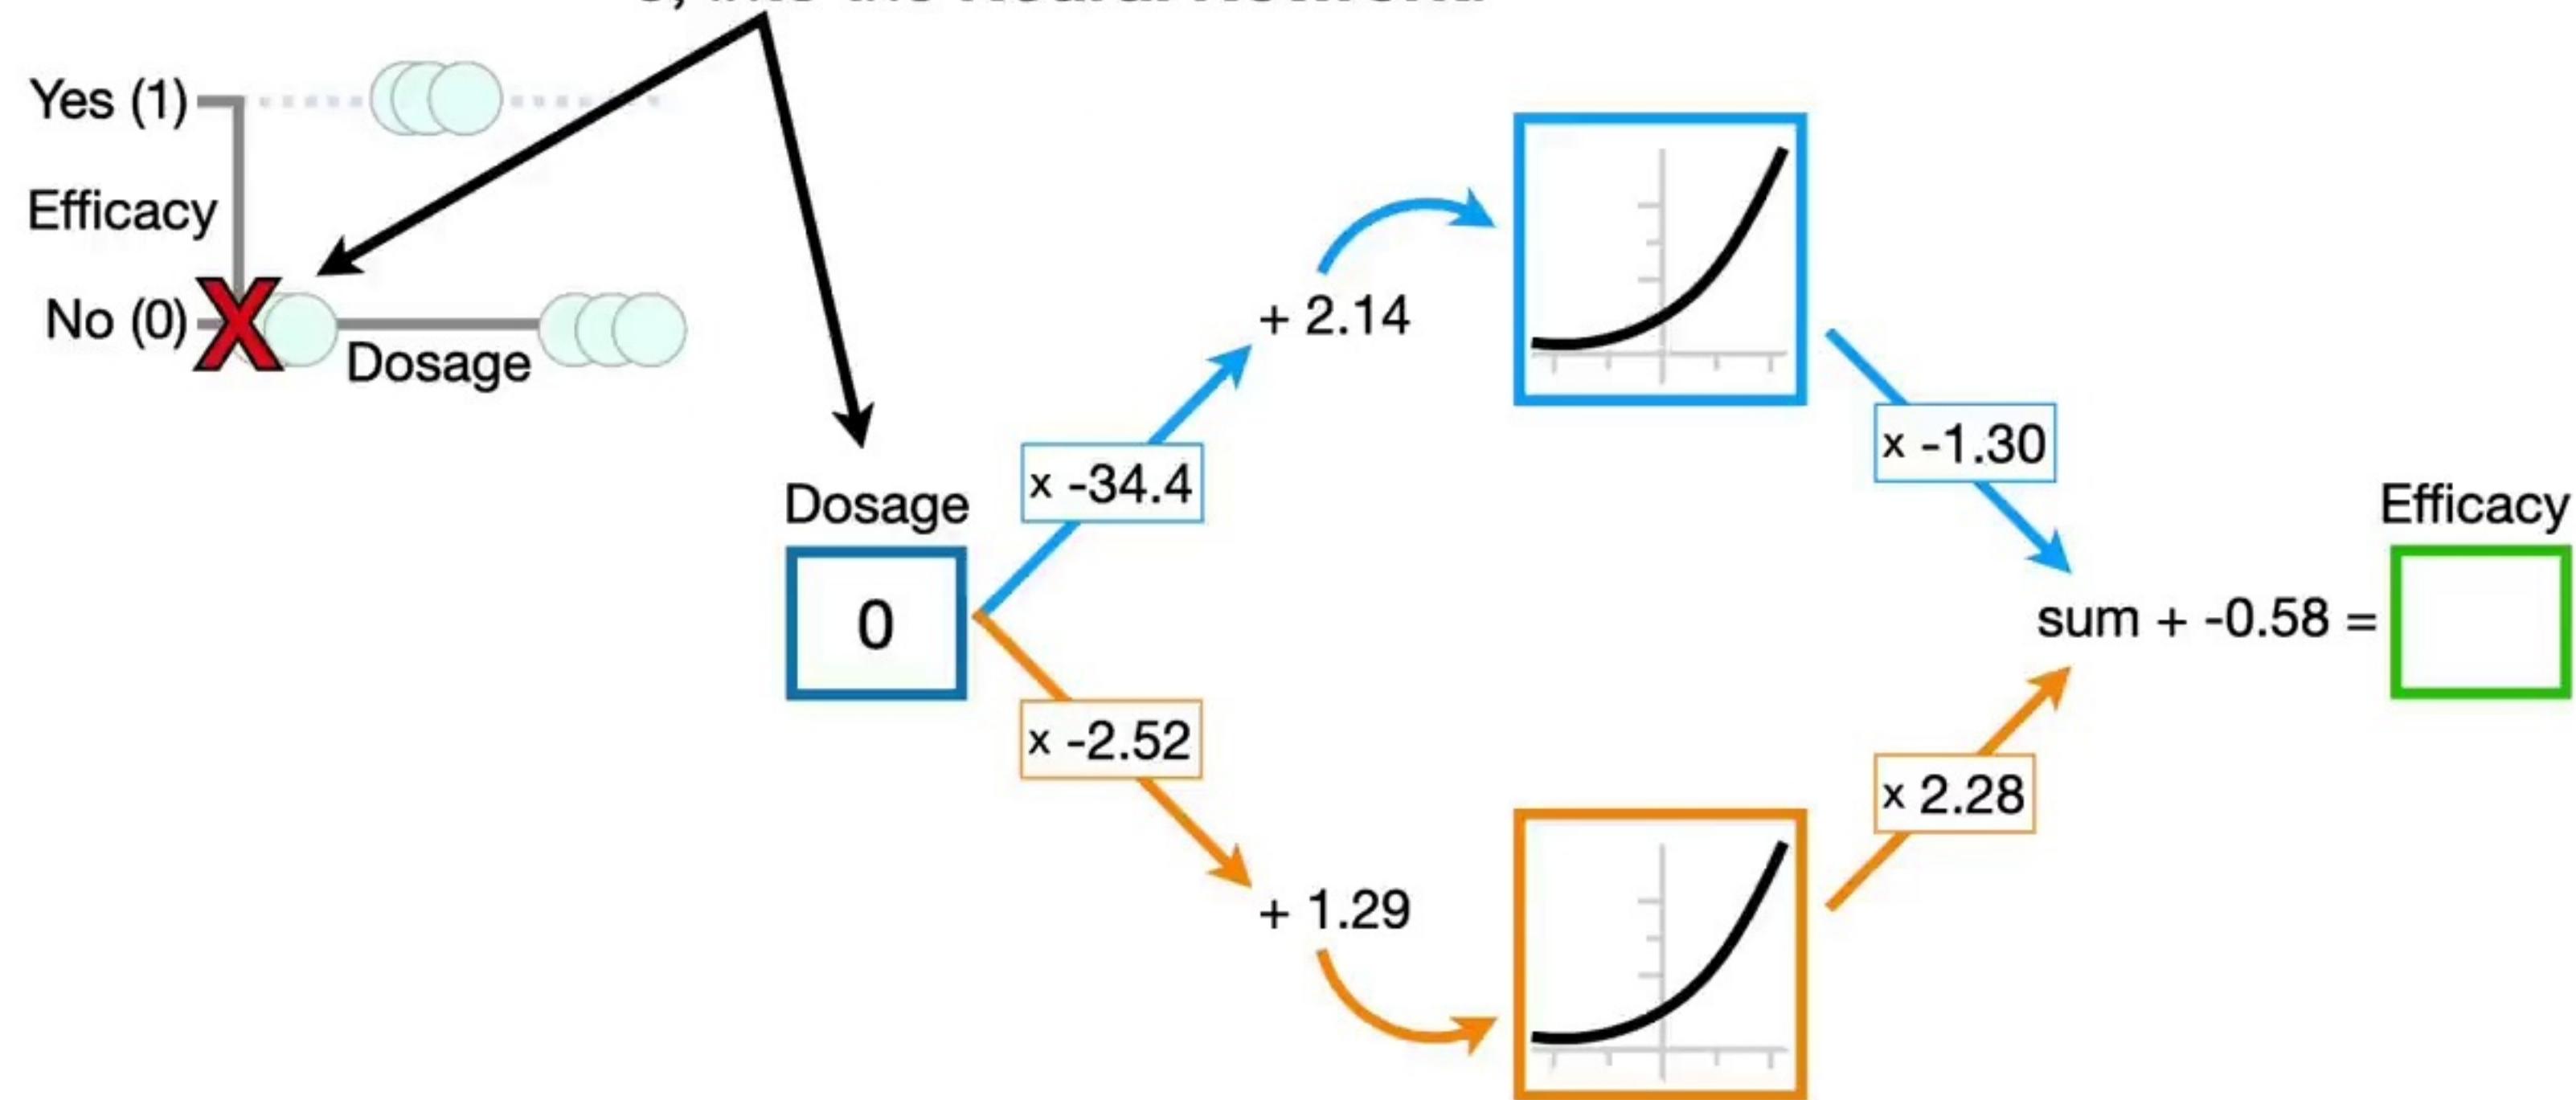


**NOTE:** To keep the math simple, let's assume **Dosages** go from **0** (low) to **1** (high).



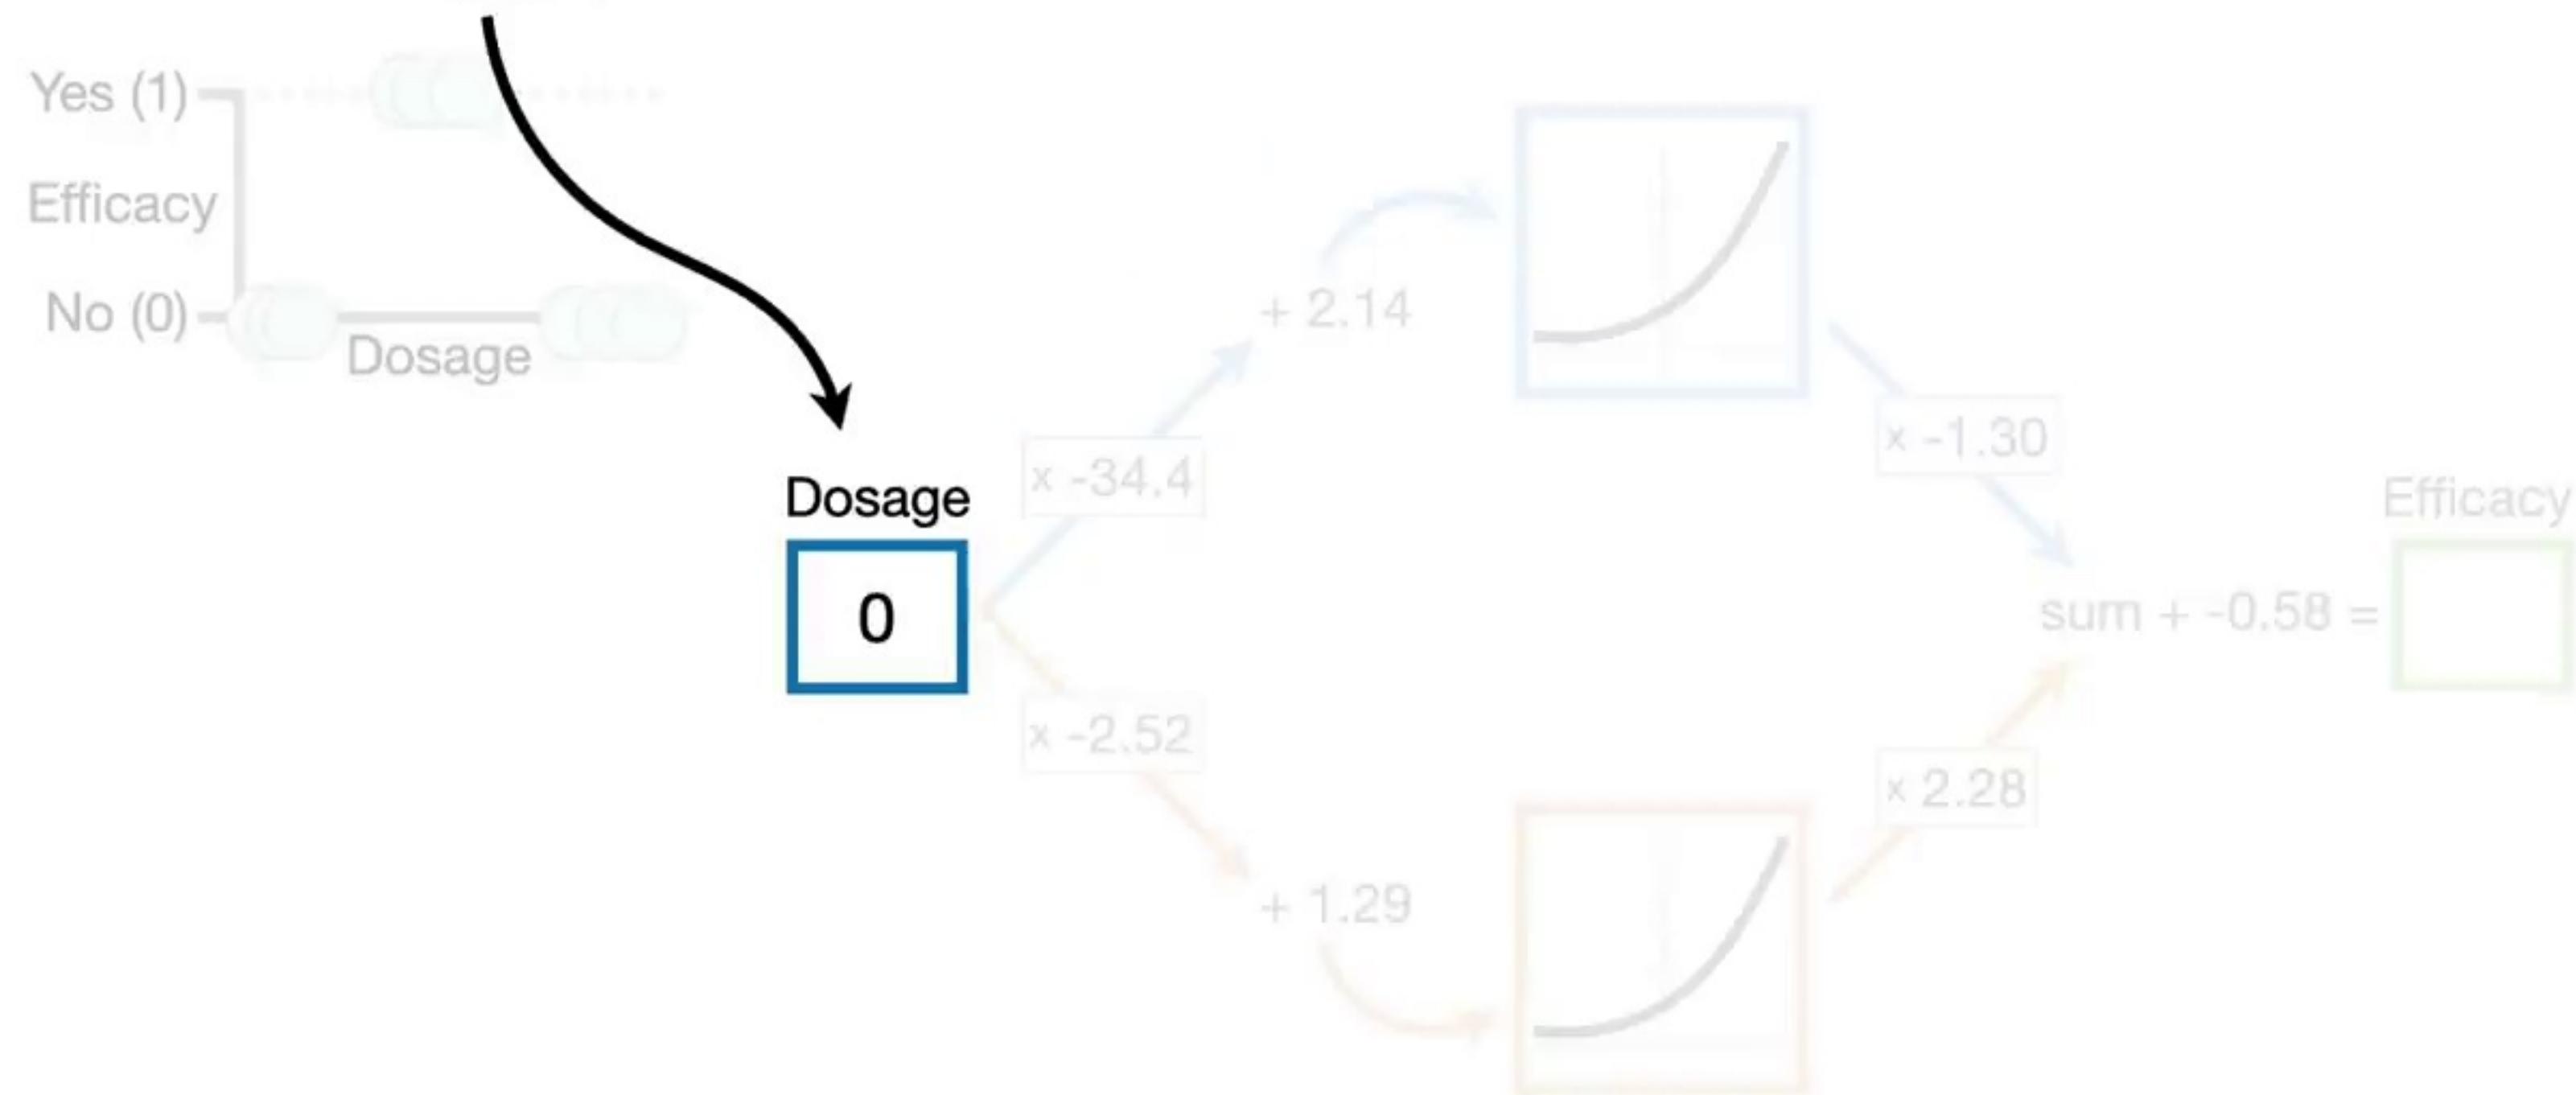


The first thing we are going to do is plug the lowest **Dosage**, **0**, into the **Neural Network**.



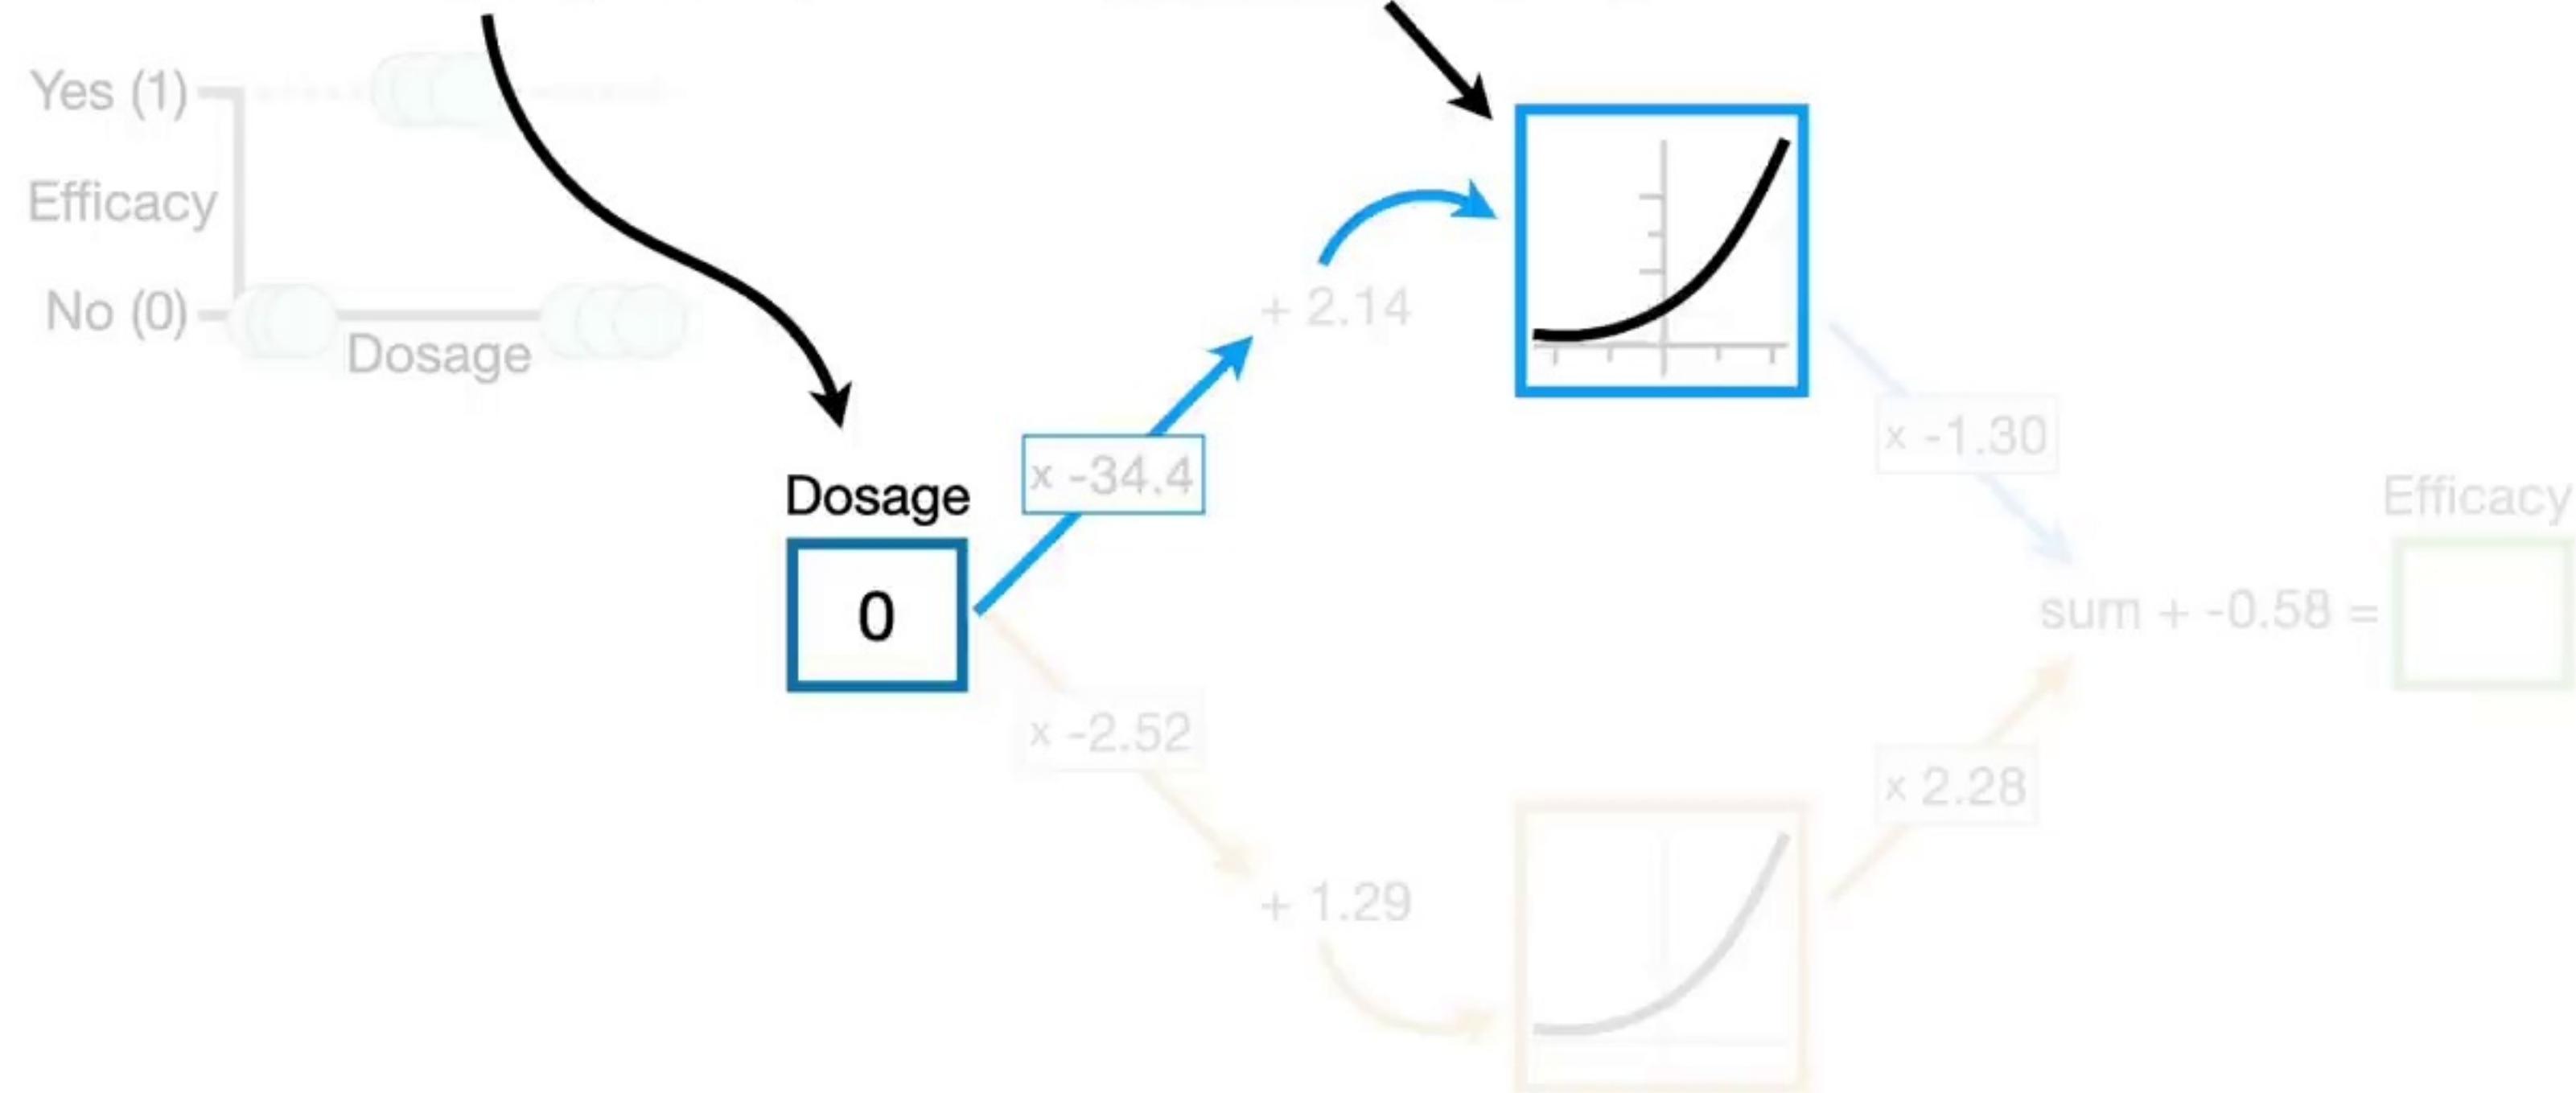


Now, to get from the  
**Input Node...**





Now, to get from the  
**Input Node**...





(Dosage  $\times$  -34.4)

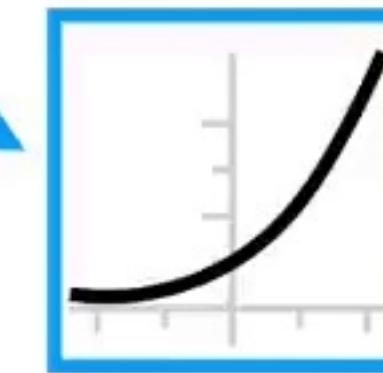
...this connection multiplies  
the **Dosage** by **-34.4**...



Dosage

$\times$  -34.4

+ 2.14



+ 1.29



$\times$  -1.30

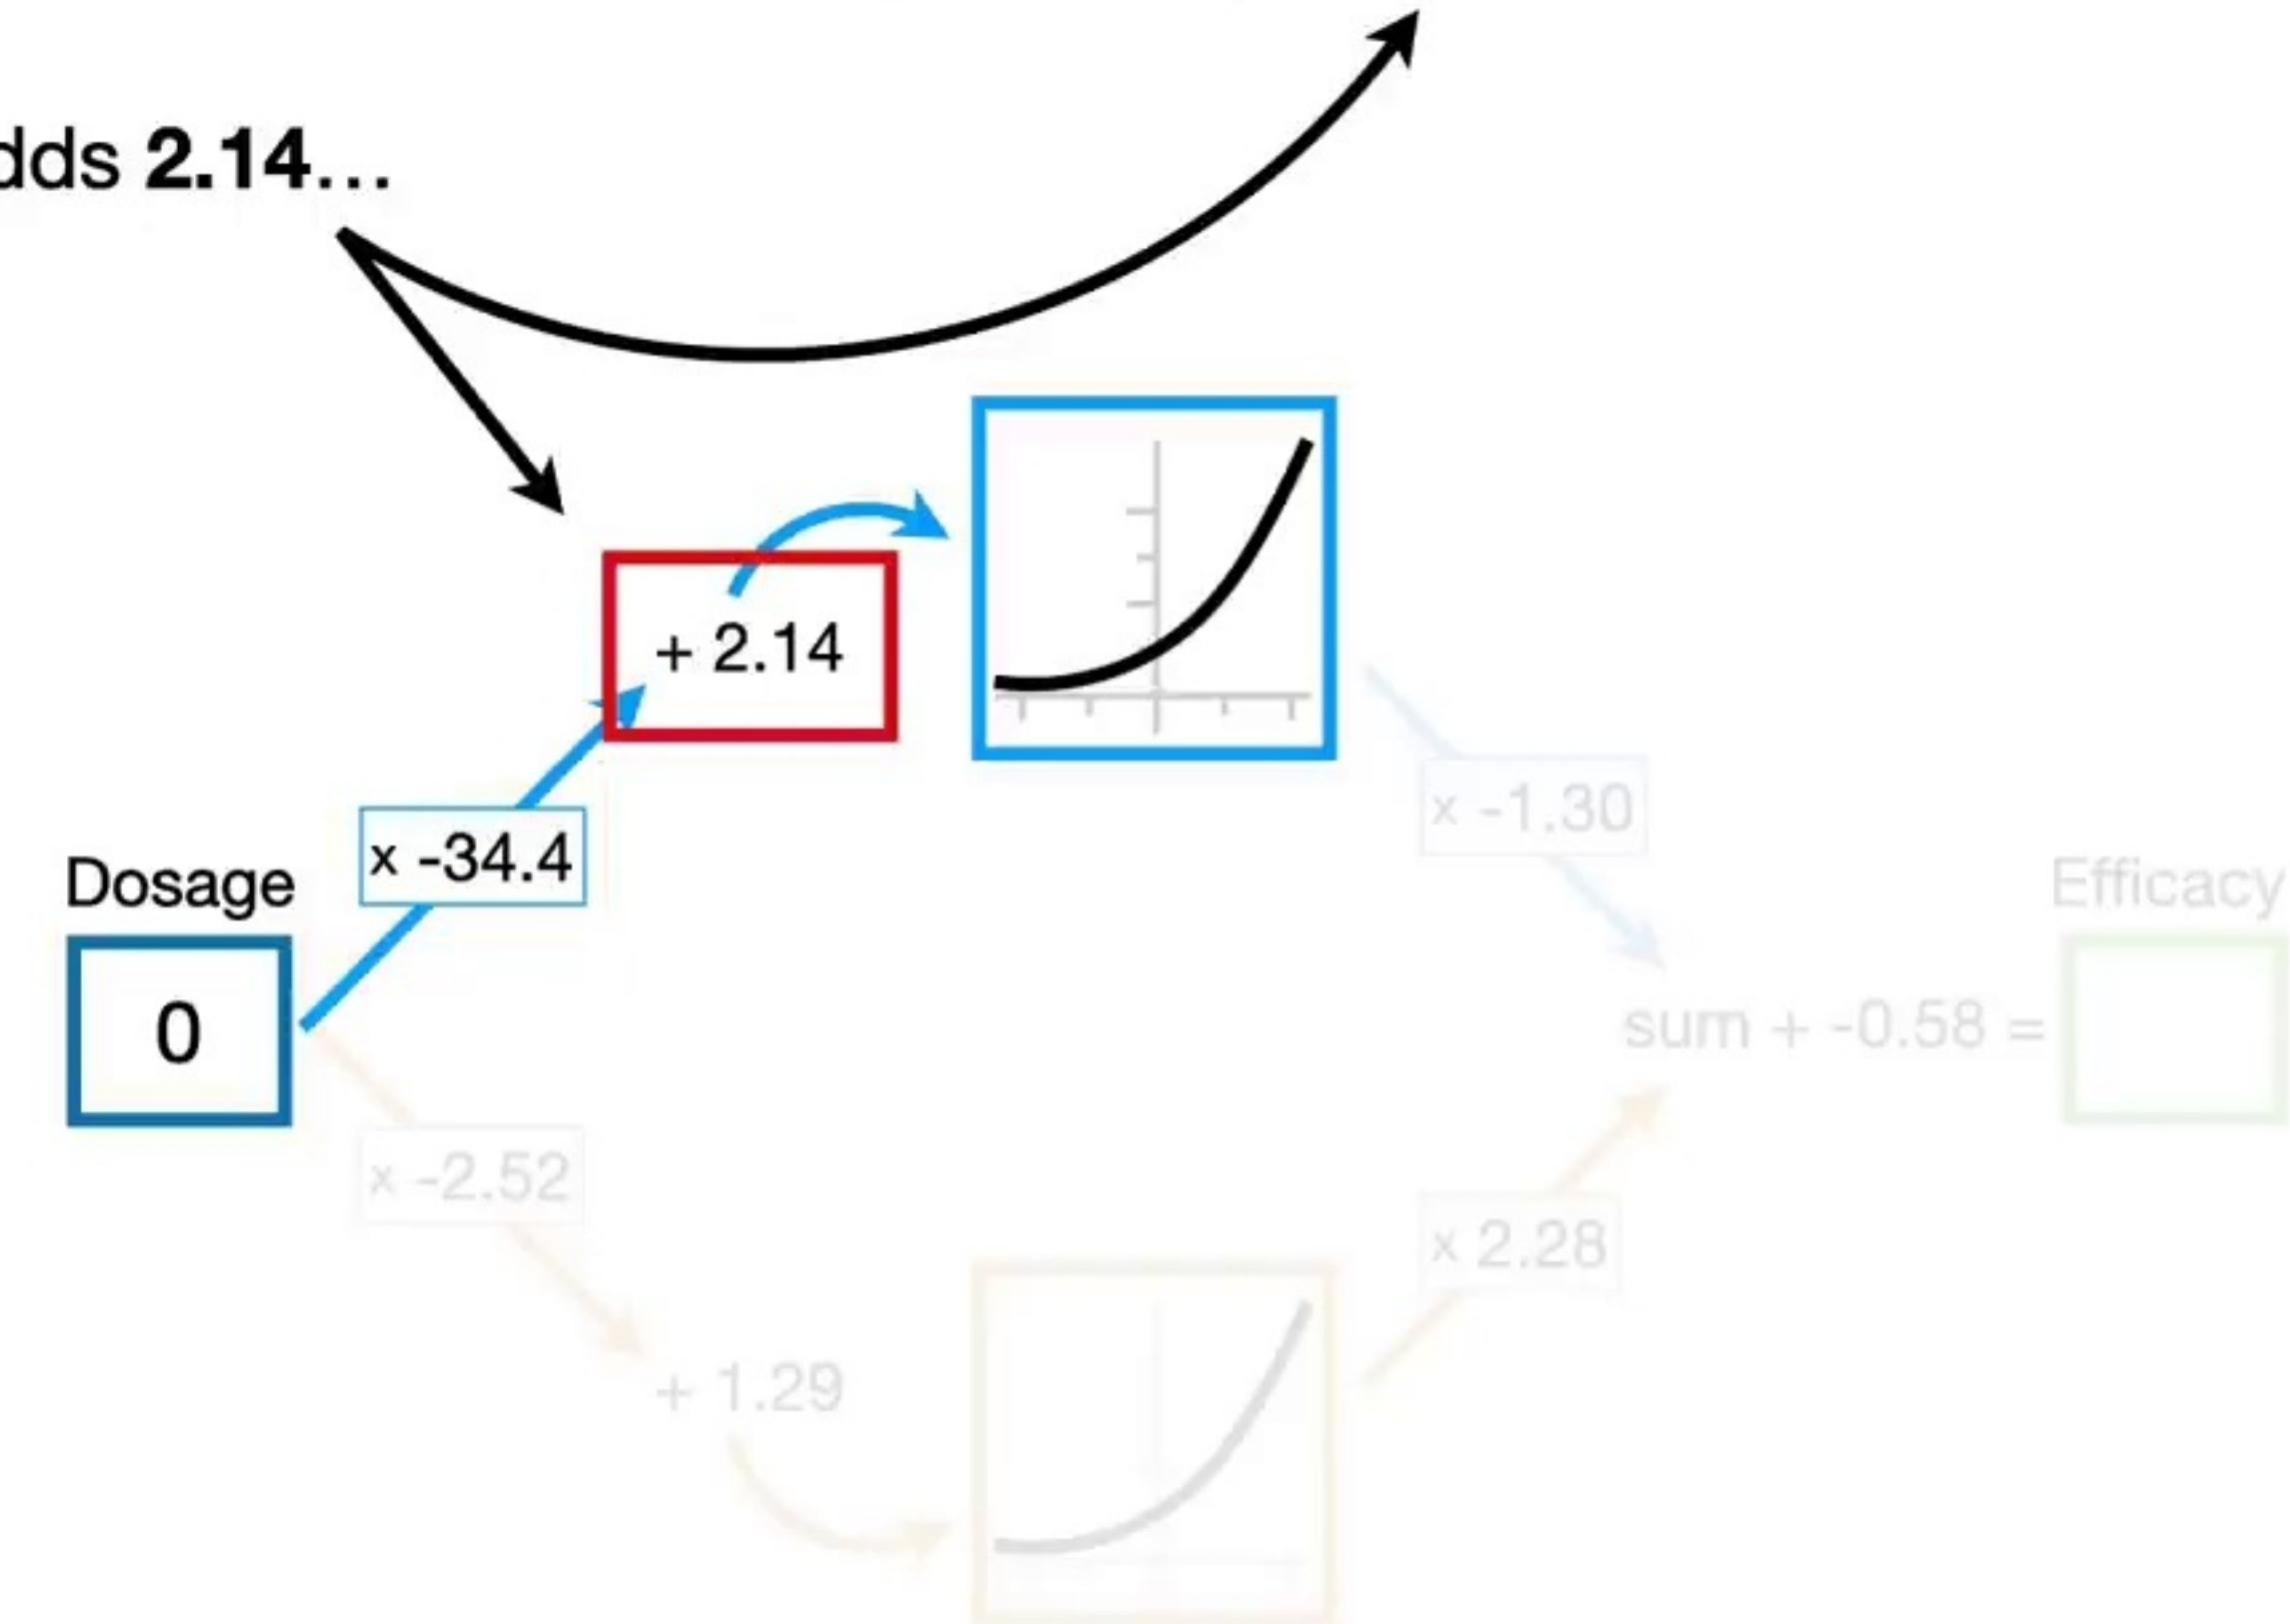
sum + -0.58 =

$\times$  2.28



$$(\text{Dosage} \times -34.4) + 2.14$$

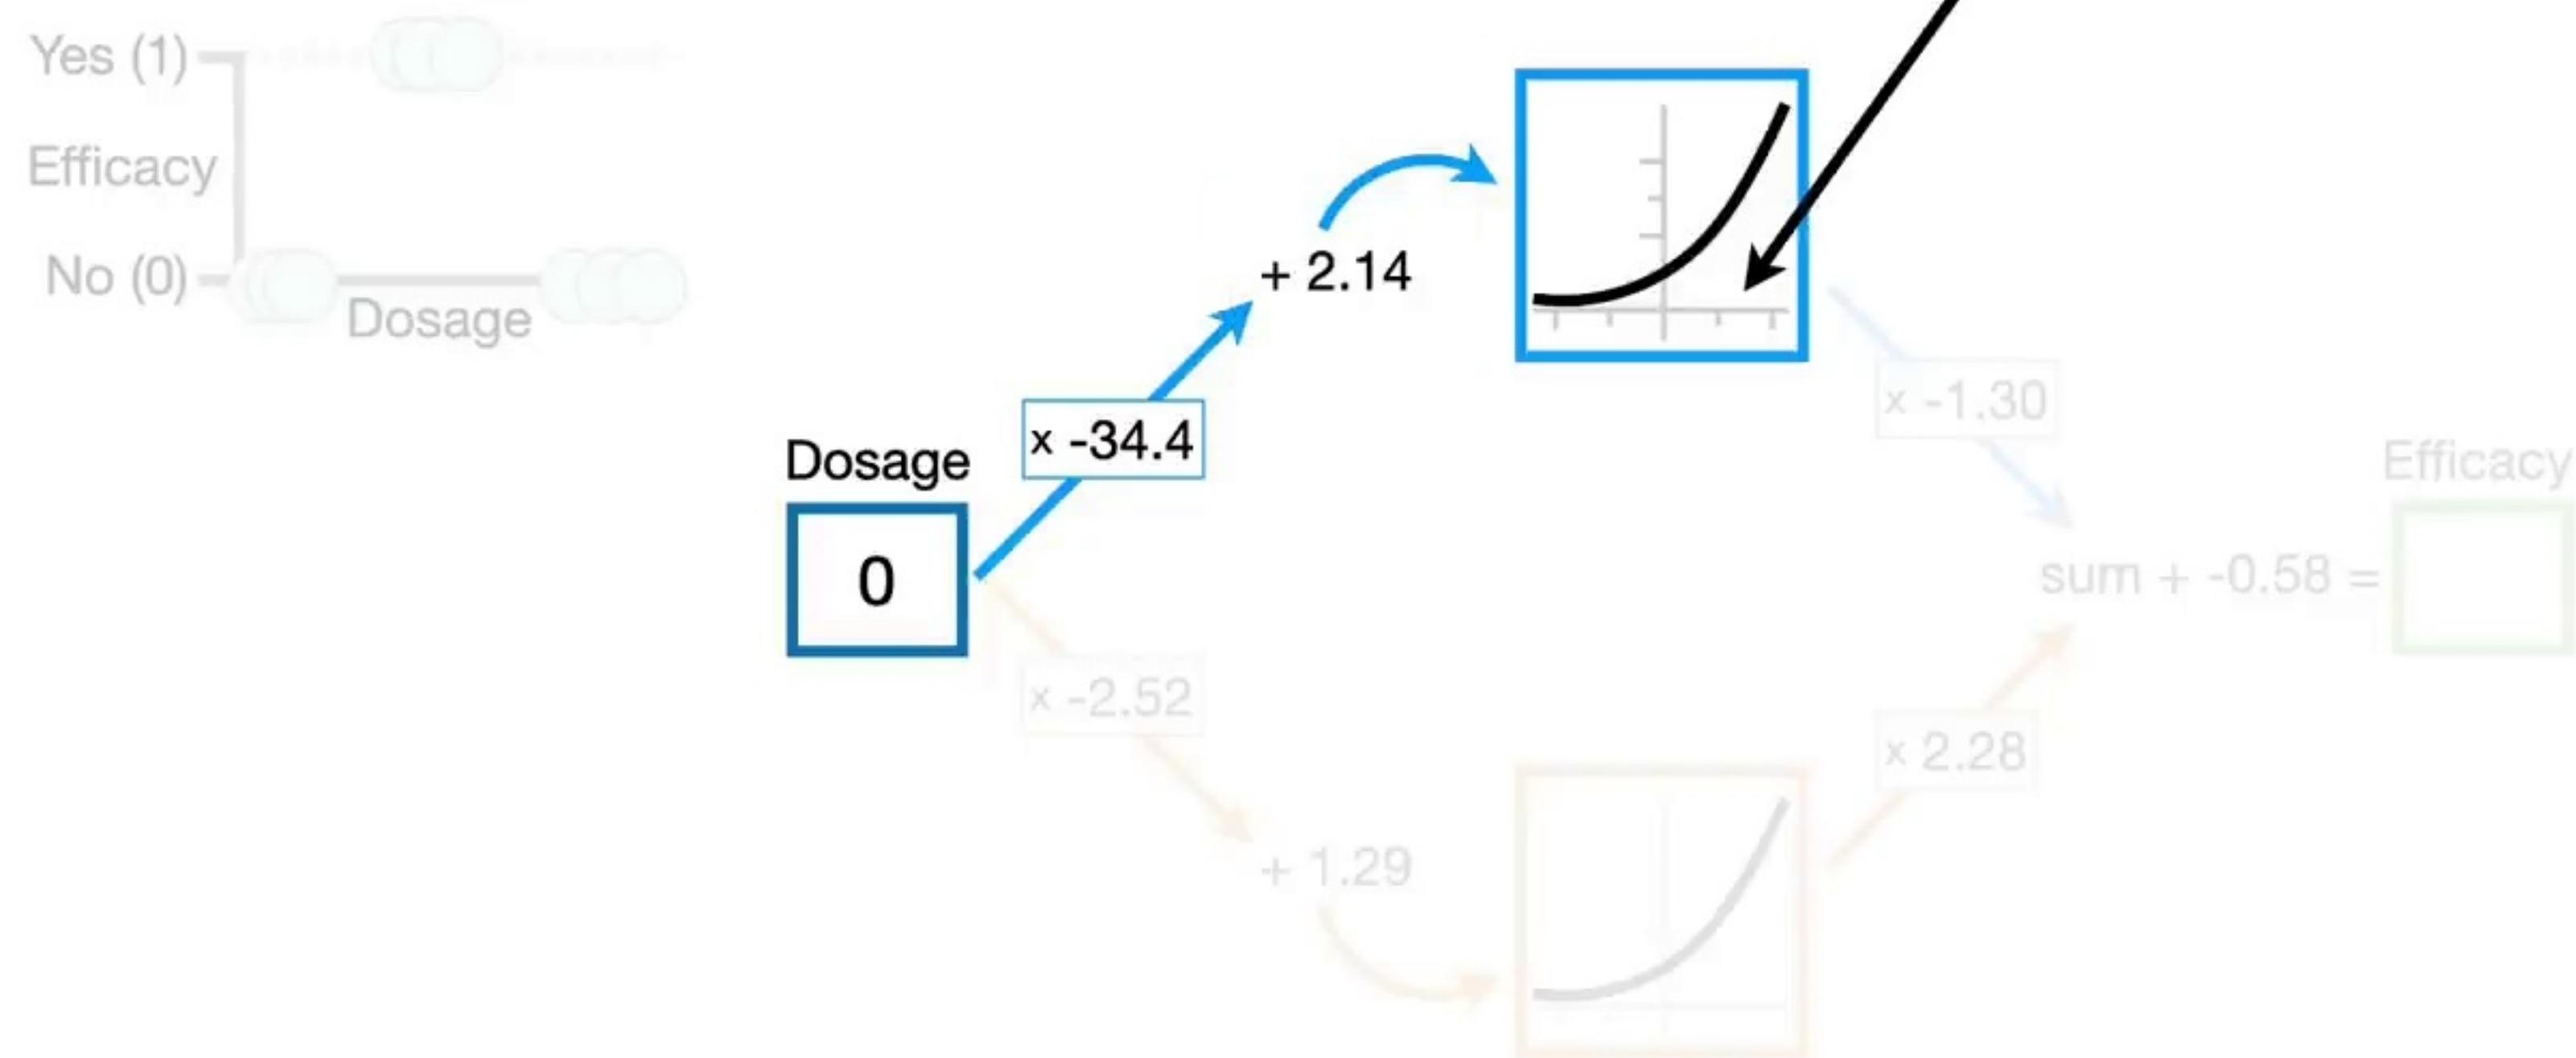
...and then adds 2.14...





...and the result is an x-axis coordinate for the **Activation Function**.

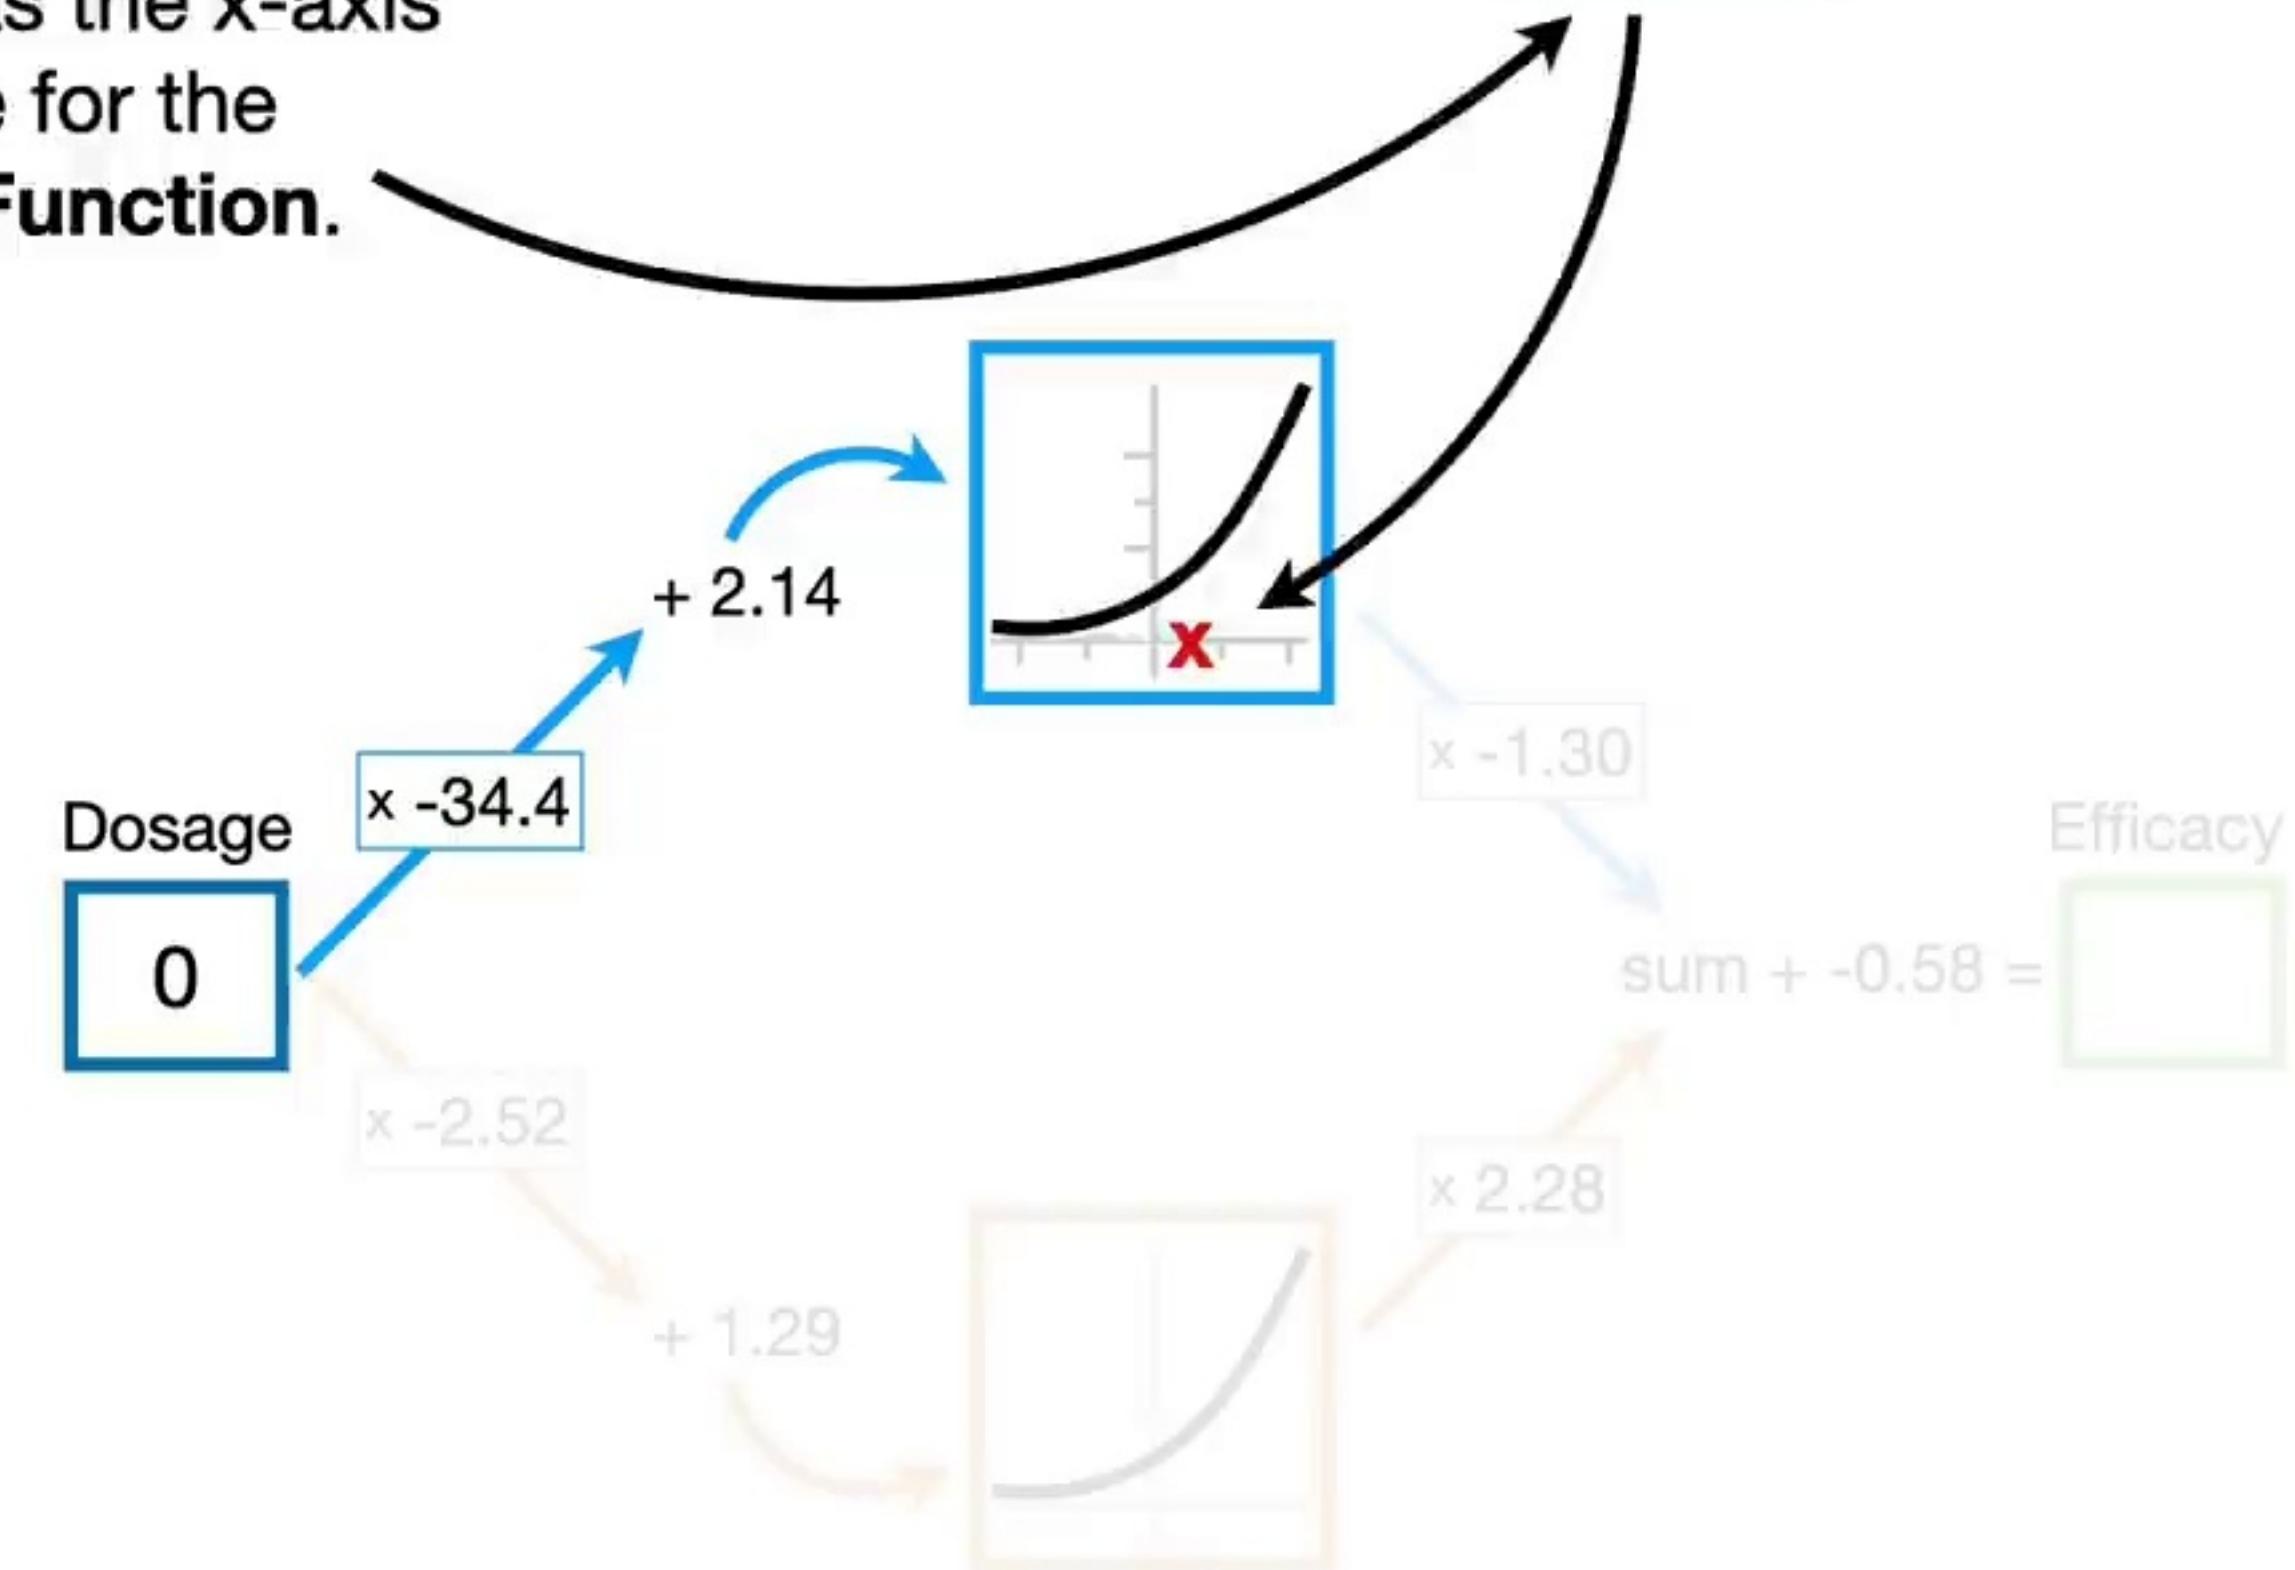
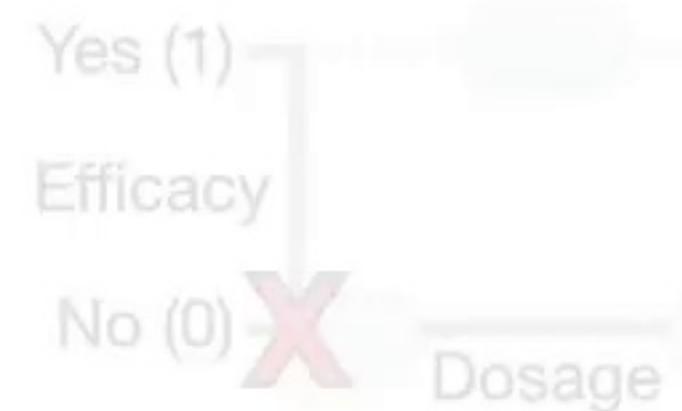
$$(\text{Dosage} \times -34.4) + 2.14 = \text{x-axis coordinate}$$





...to get **2.14** as the x-axis coordinate for the **Activation Function**.

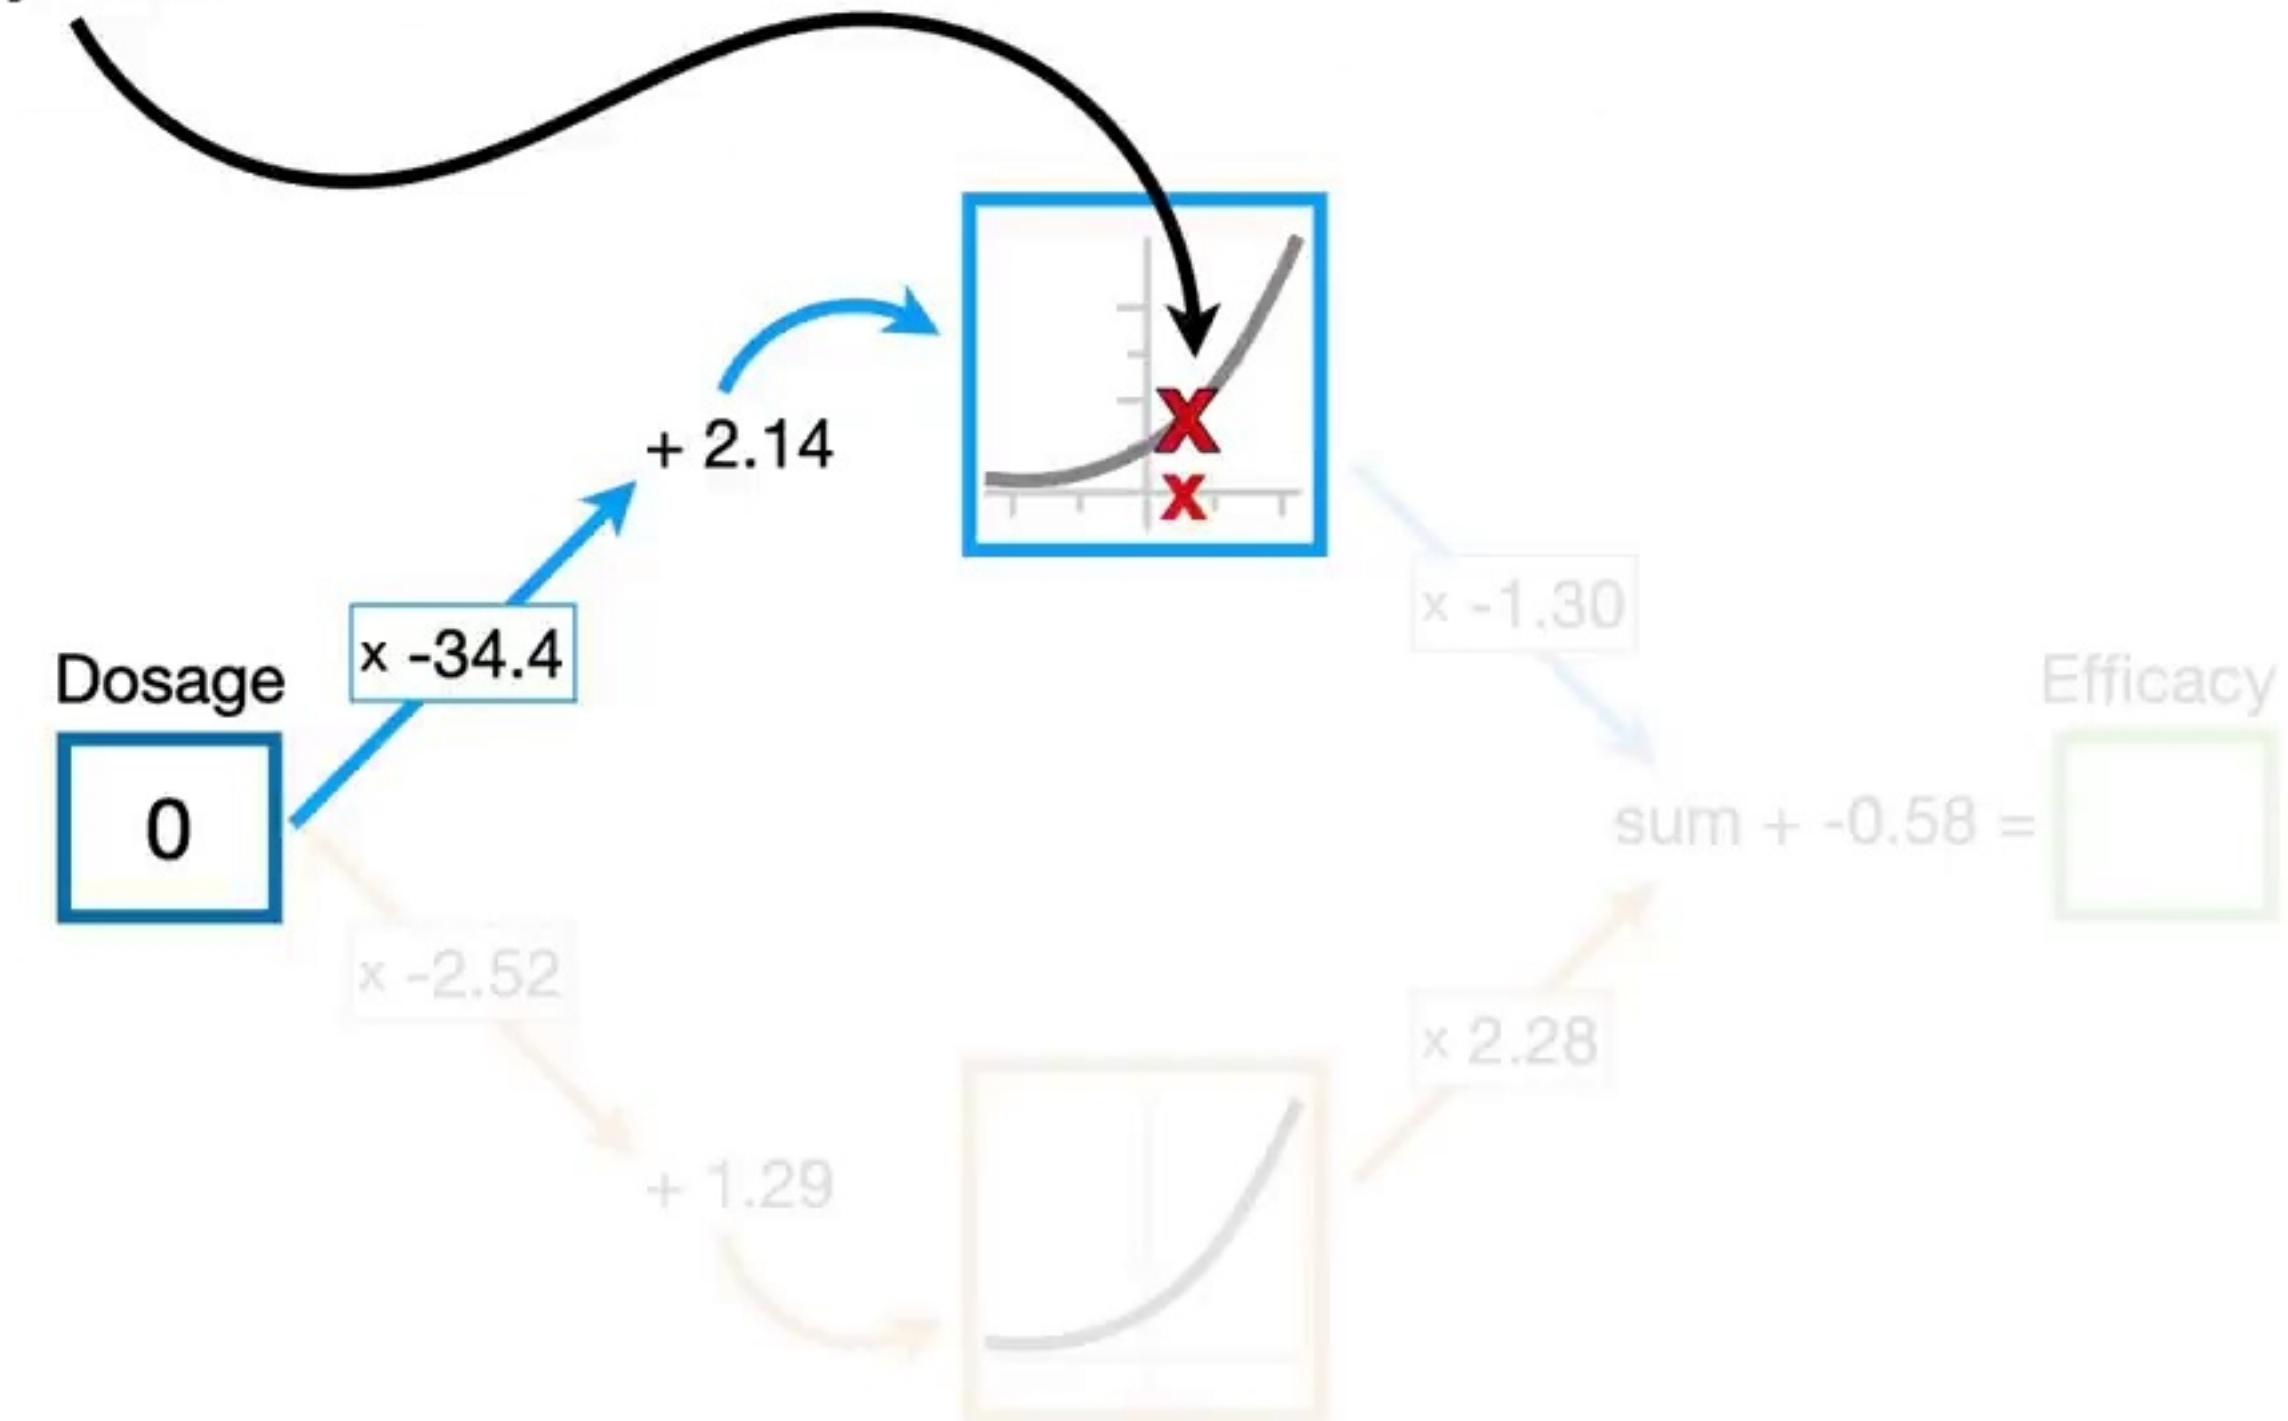
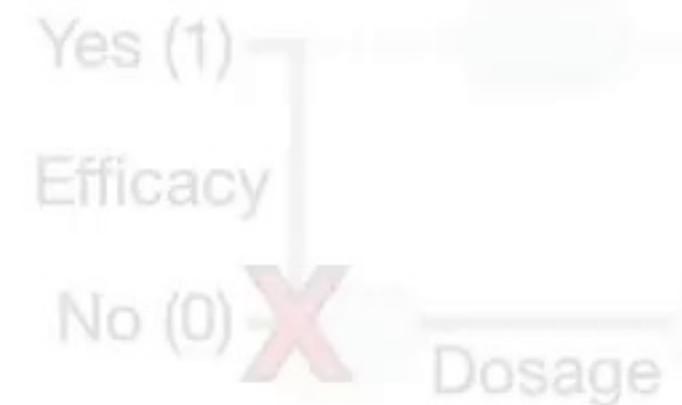
$$0 + 2.14 = 2.14$$





$$0 + 2.14 = 2.14$$

To get the corresponding y-axis value...



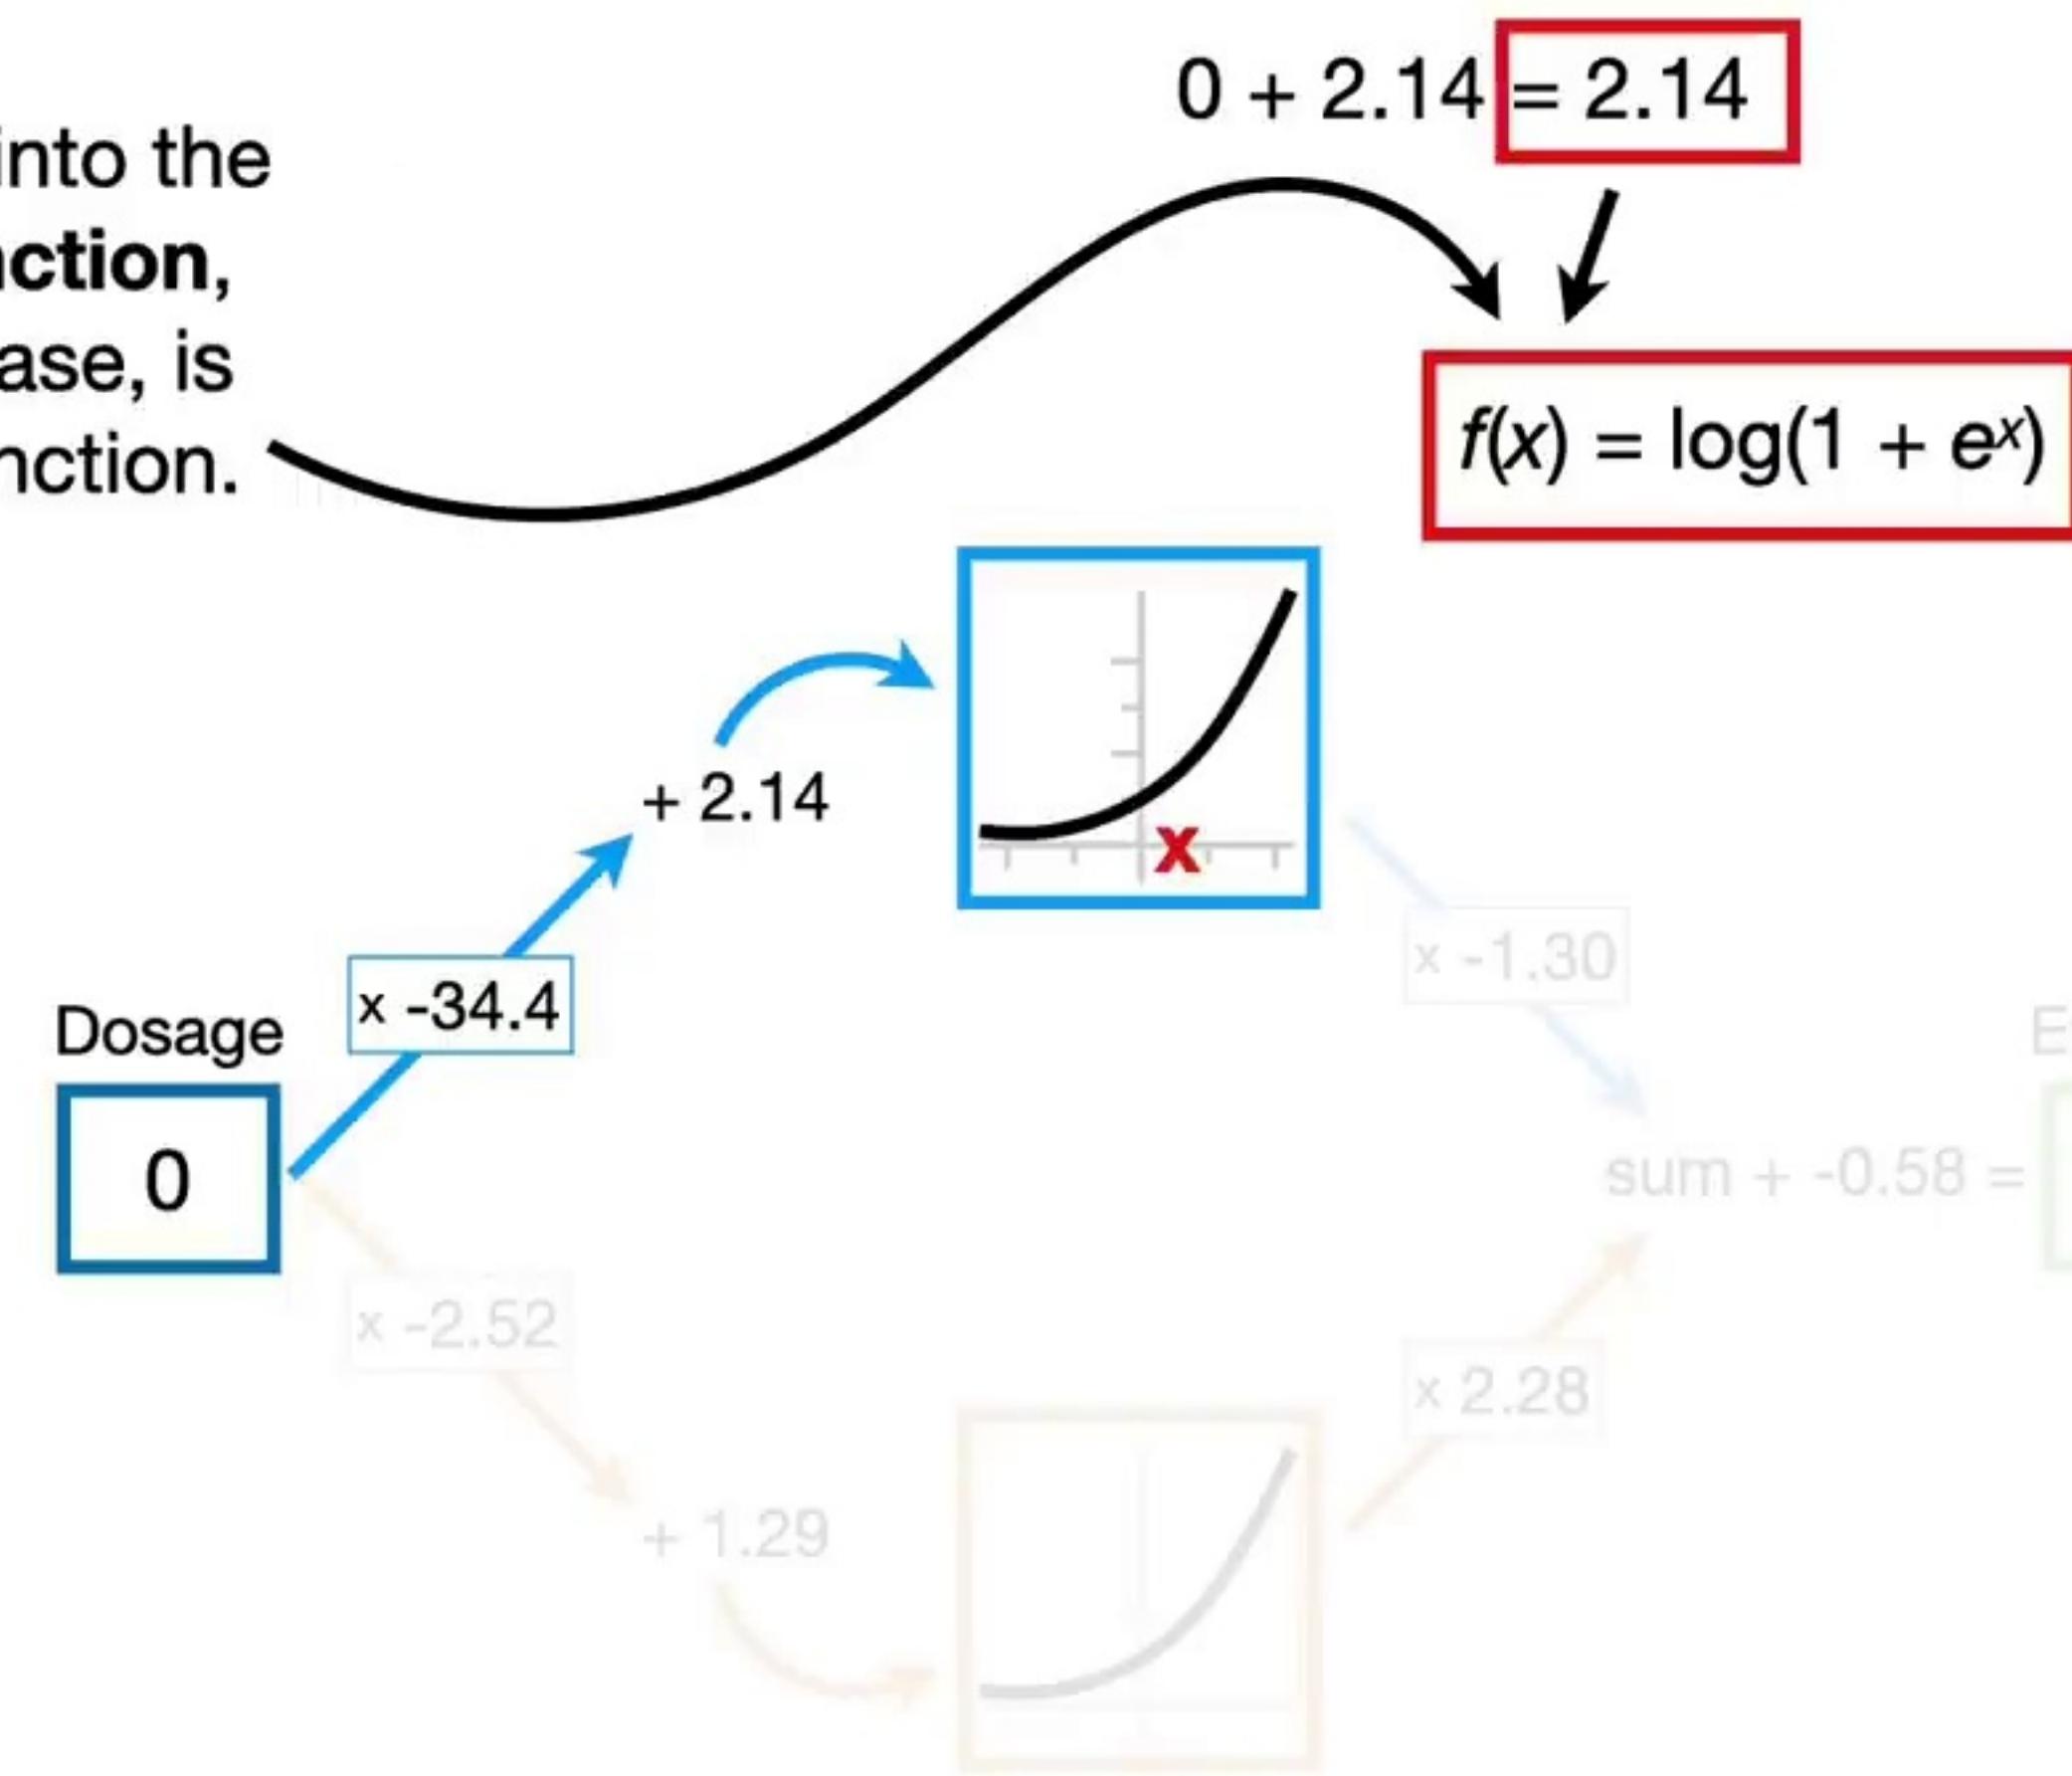


...we plug **2.14** into the **Activation Function**, which, in this case, is the **softplus** function.

$$0 + 2.14 = 2.14$$

$$f(x) = \log(1 + e^x)$$

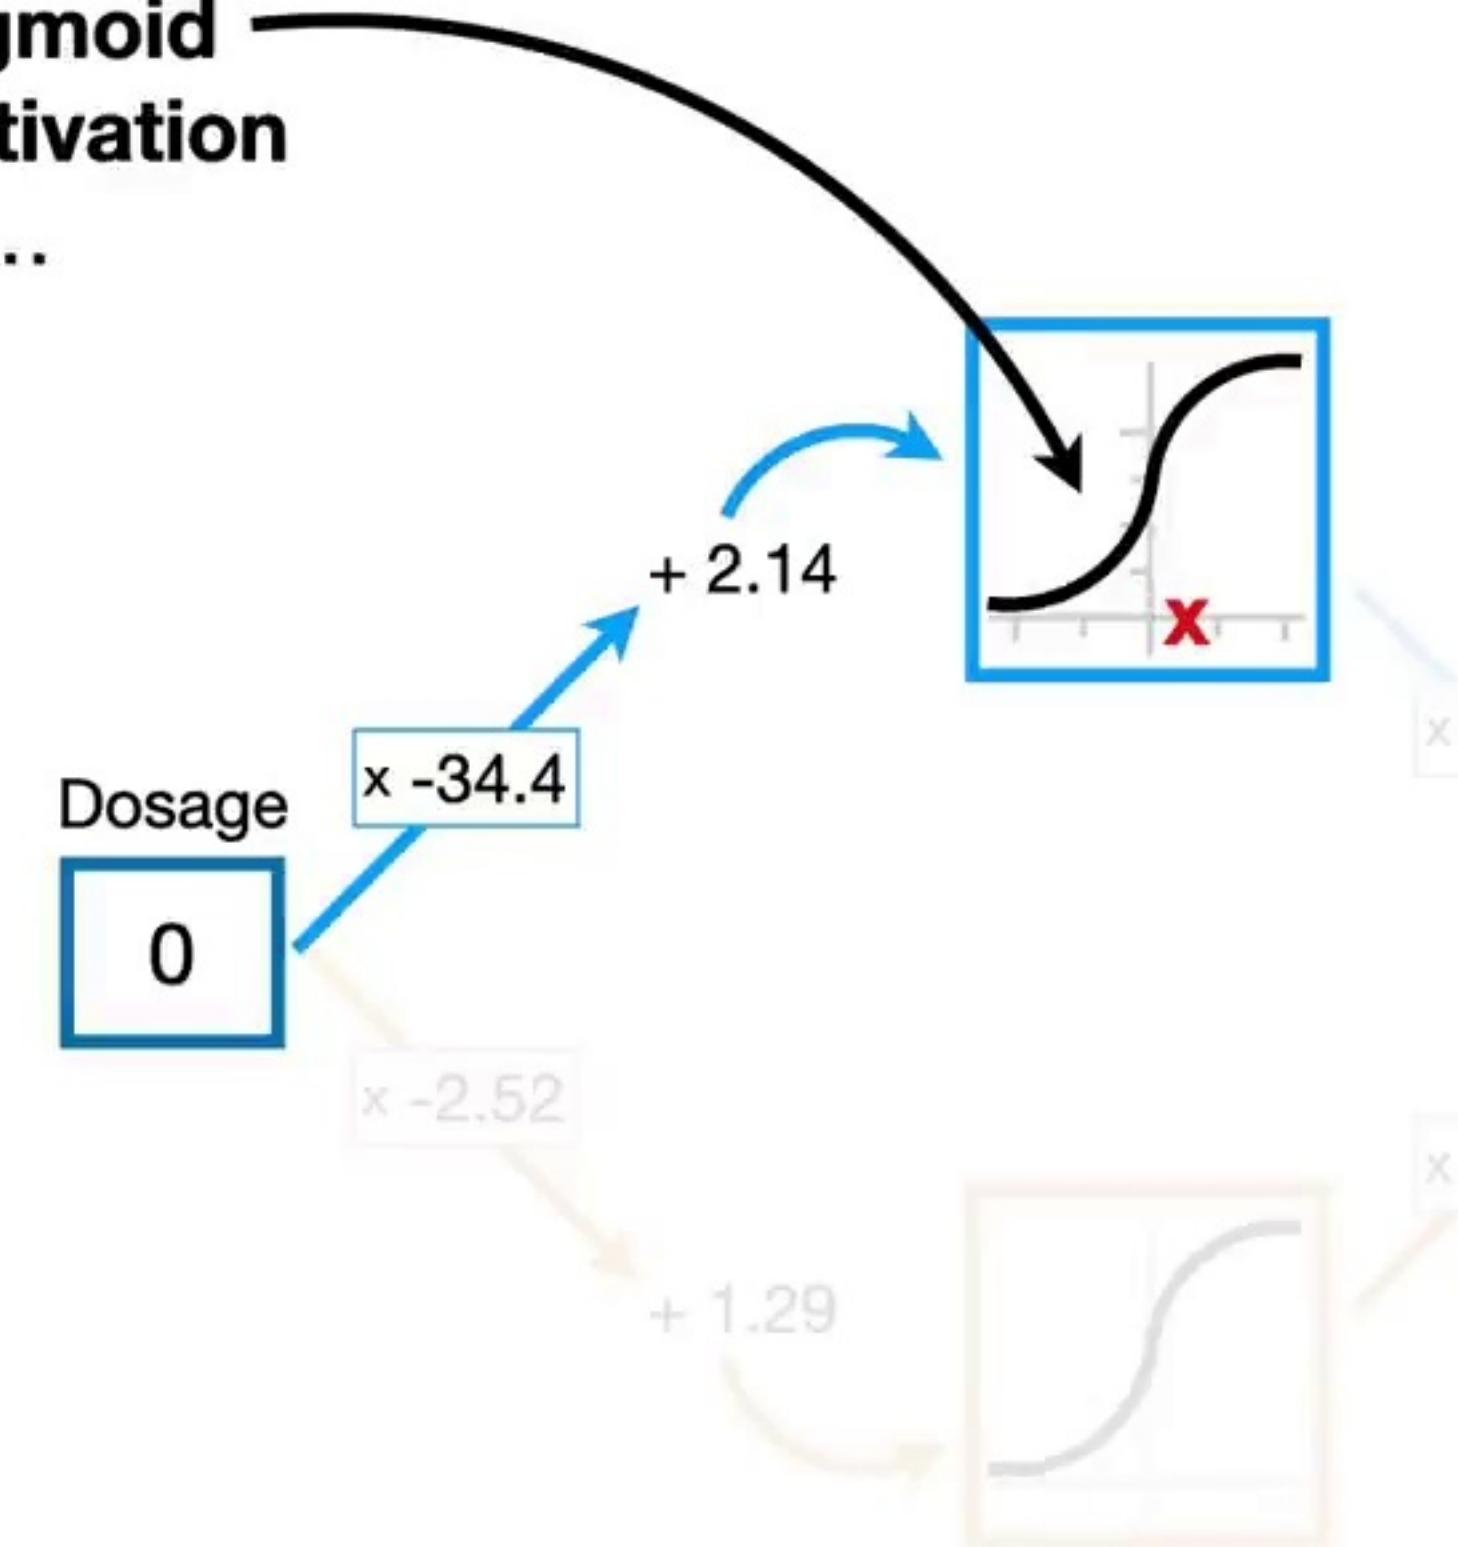
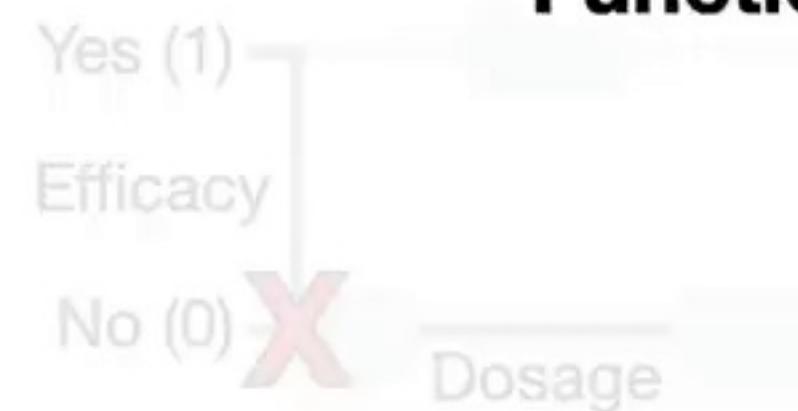
Yes (1)  
Efficacy  
No (0) **X**  
Dosage





**NOTE:** If we had chosen the **sigmoid** curve for the **Activation Function**...

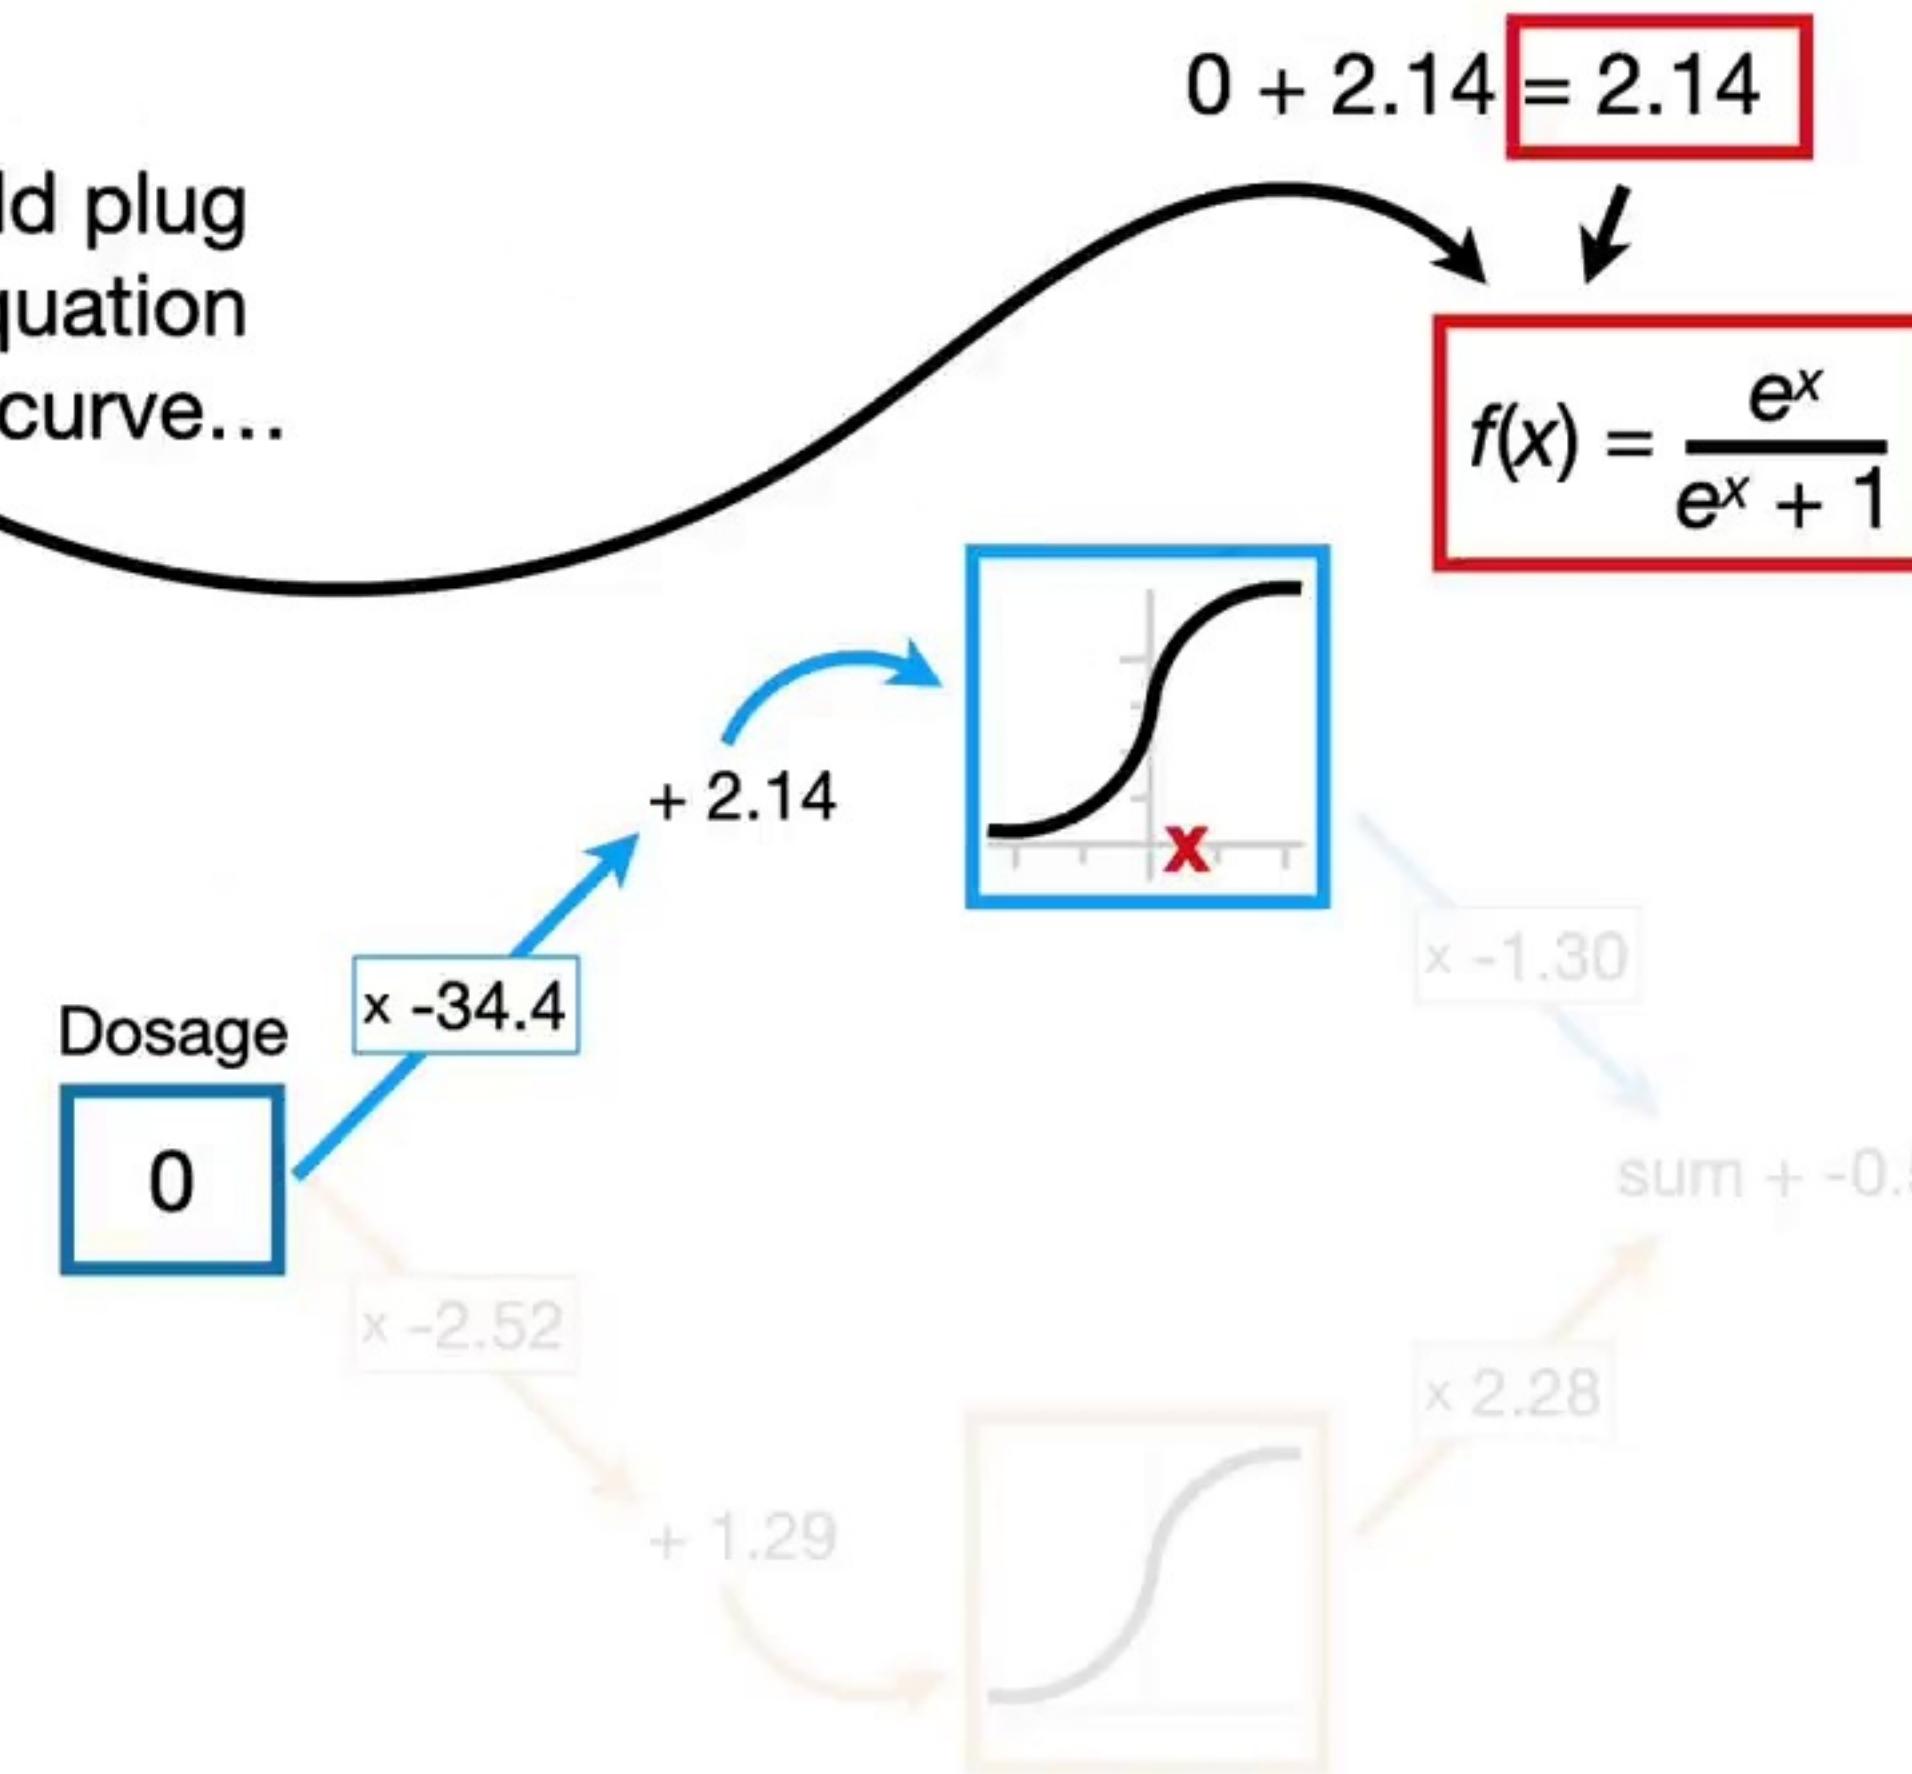
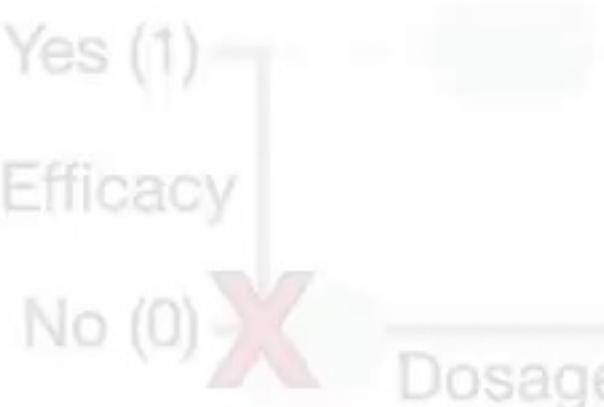
$$0 + 2.14 = 2.14$$



Efficacy



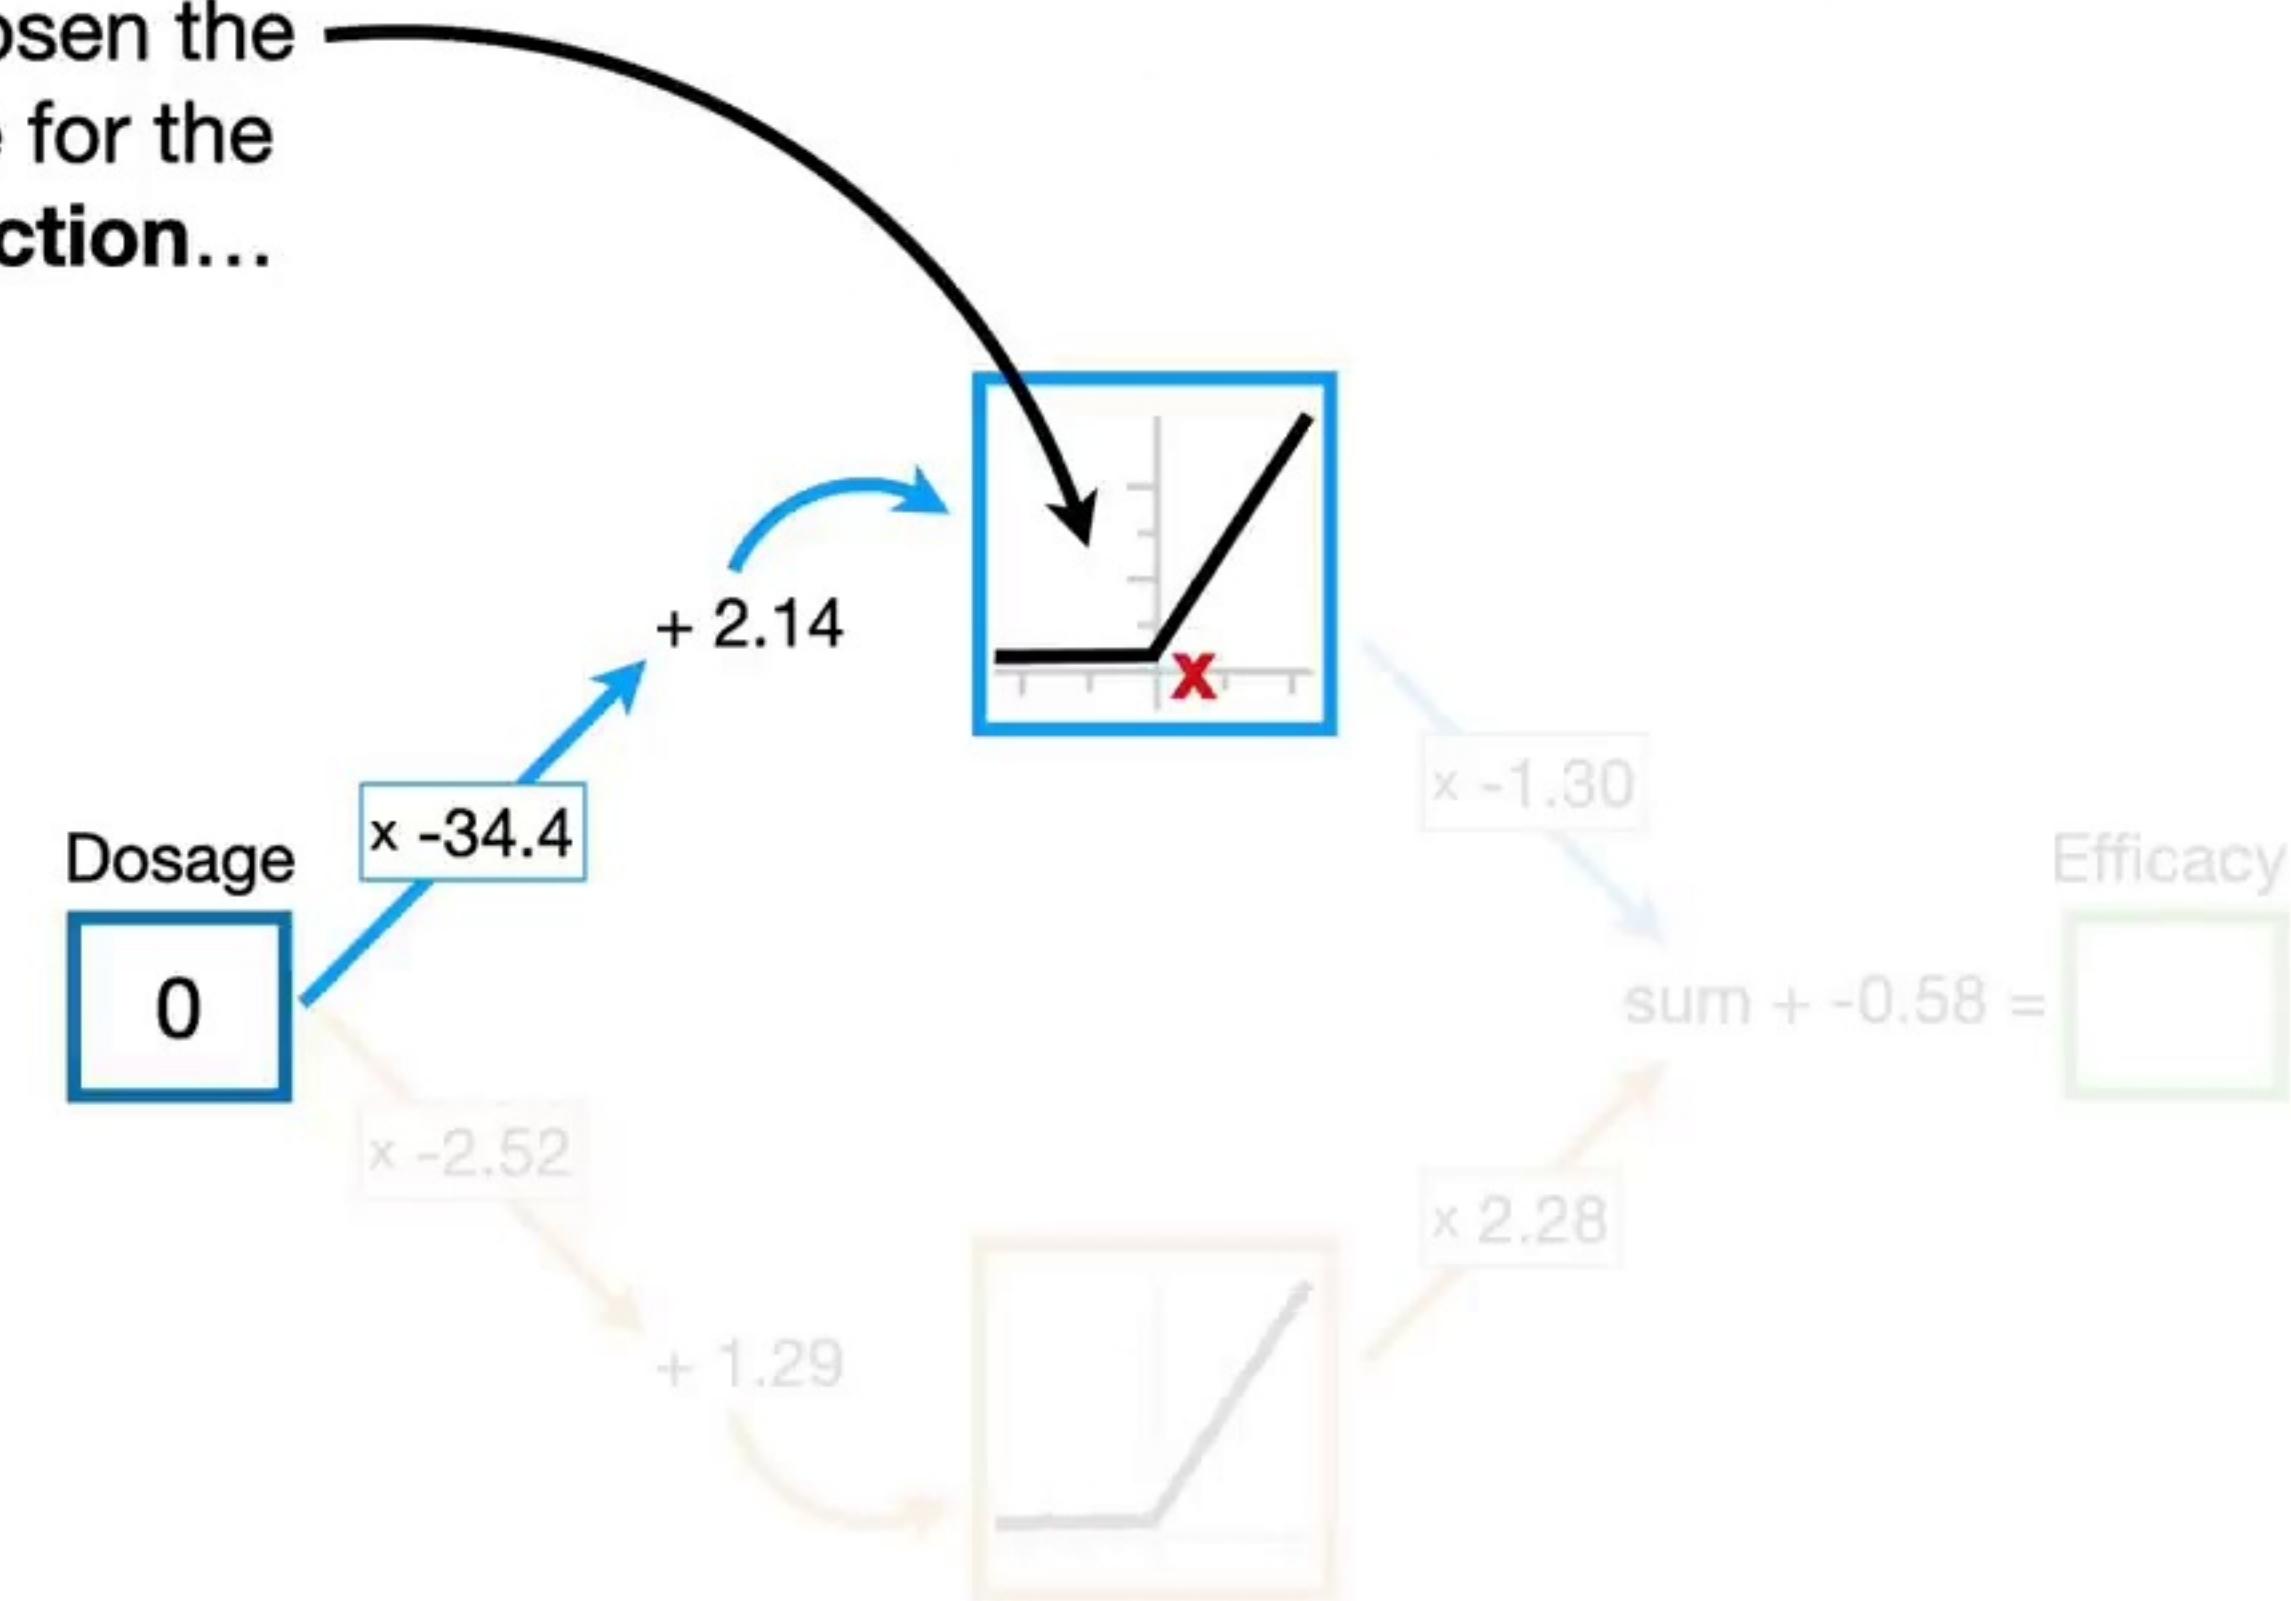
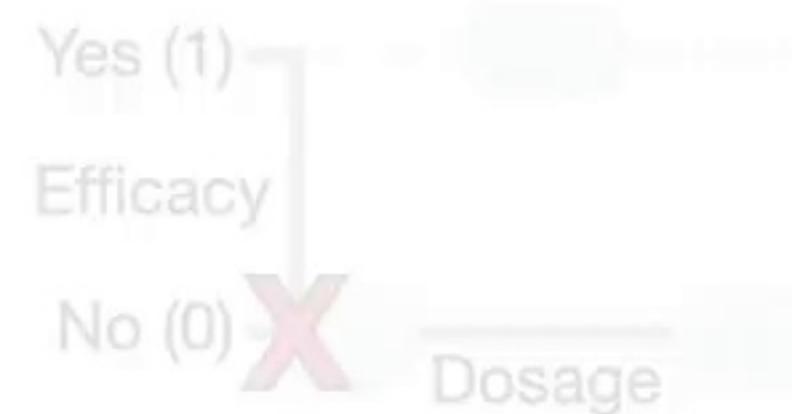
...then we would plug  
**2.14** into the equation  
for the **sigmoid curve**...





...and if had chosen the  
**ReLU bent line** for the  
**Activation Function...**

$$0 + 2.14 = 2.14$$

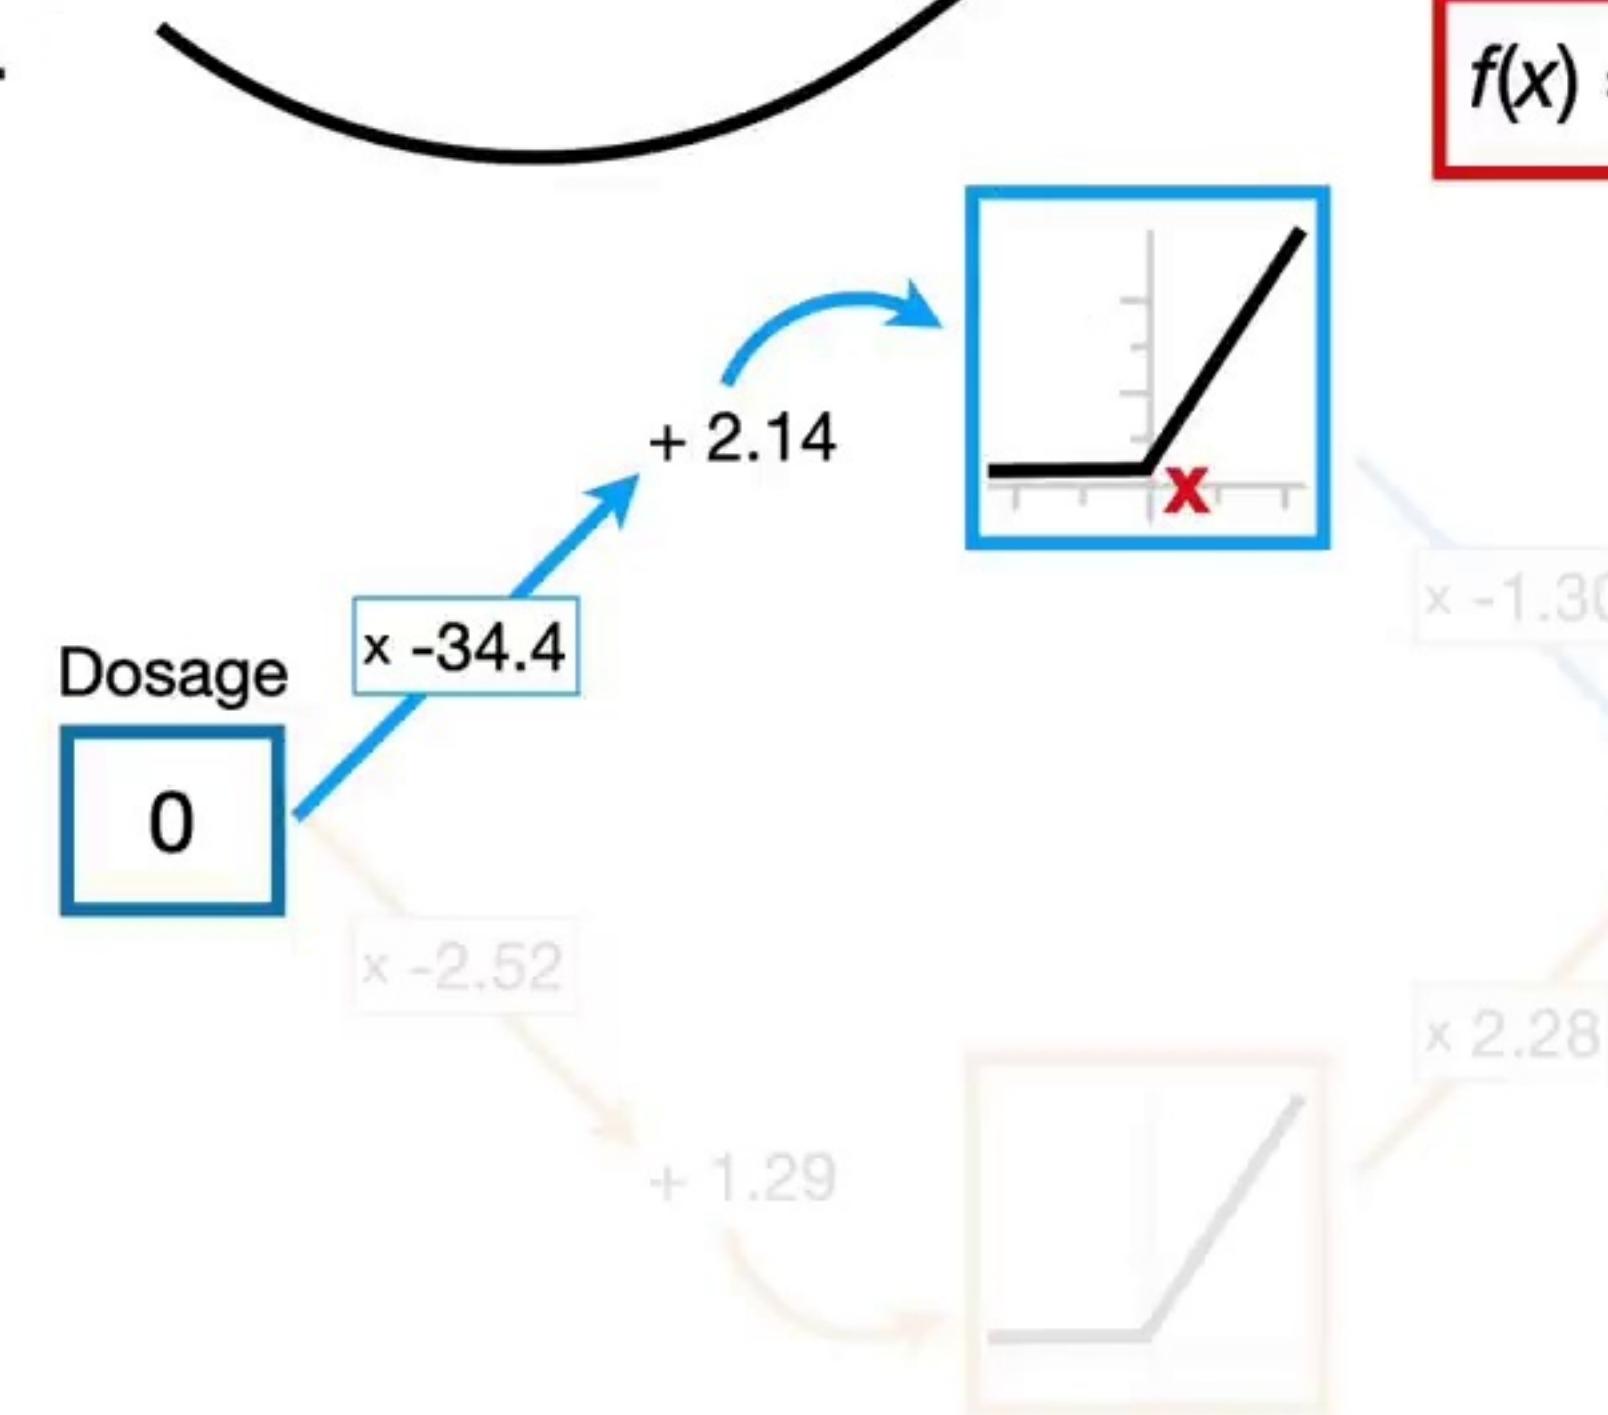
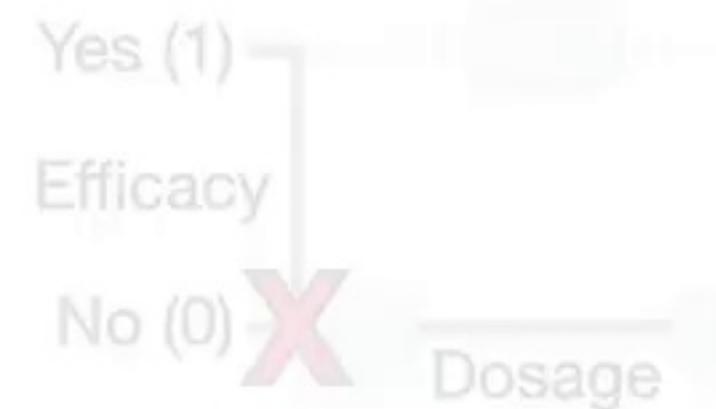




...then we would plug  
**2.14** into the ReLU  
equation.

$$0 + 2.14 = 2.14$$

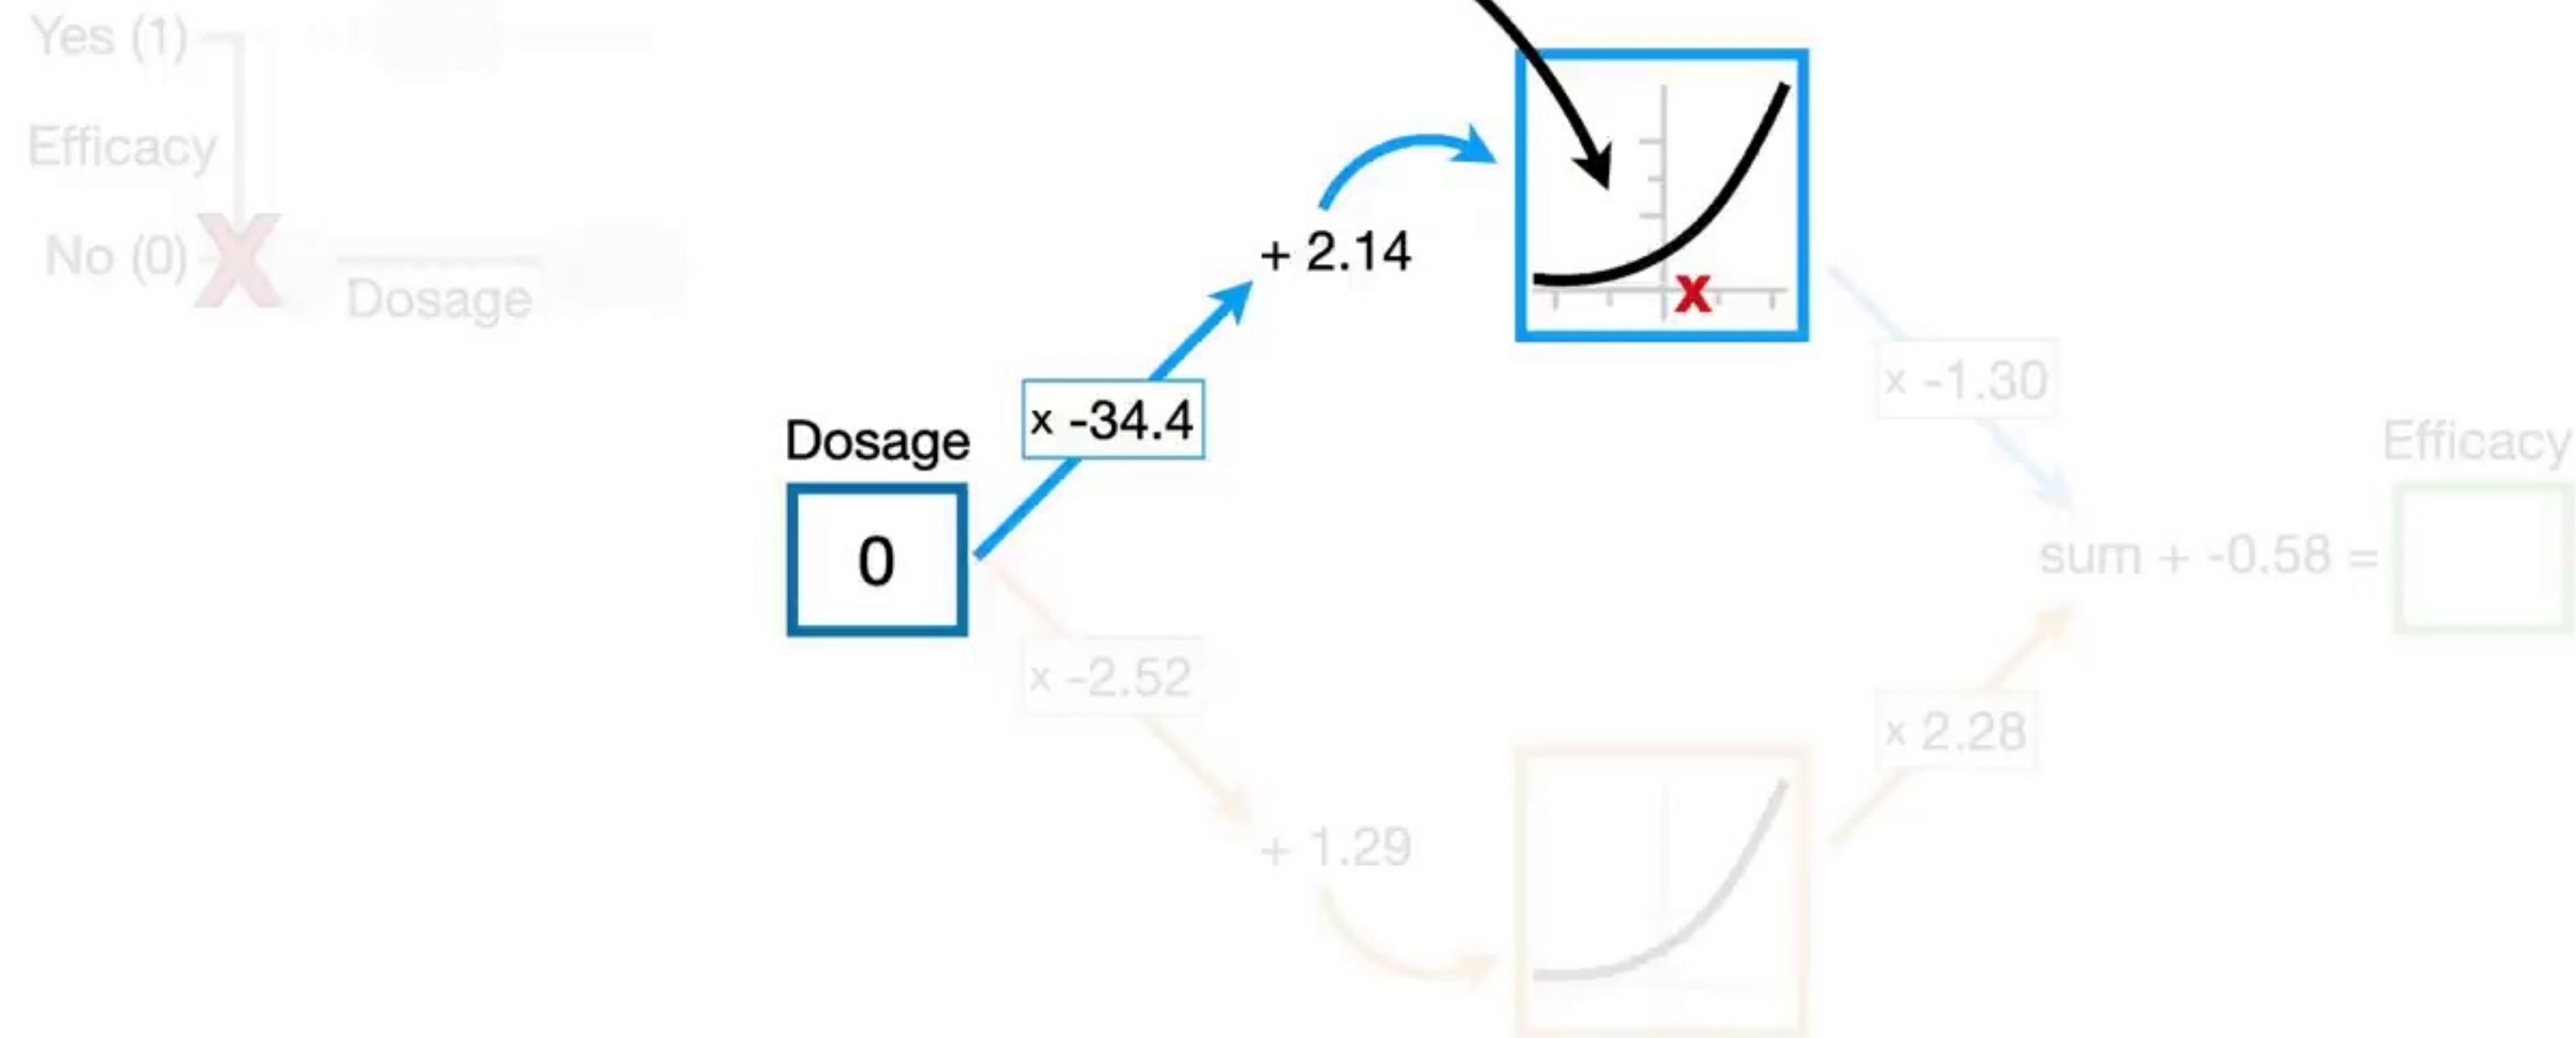
$$f(x) = \max(0, x)$$





$$0 + 2.14 = 2.14$$

But since we are using  
**softplus** for the  
**Activation Function...**

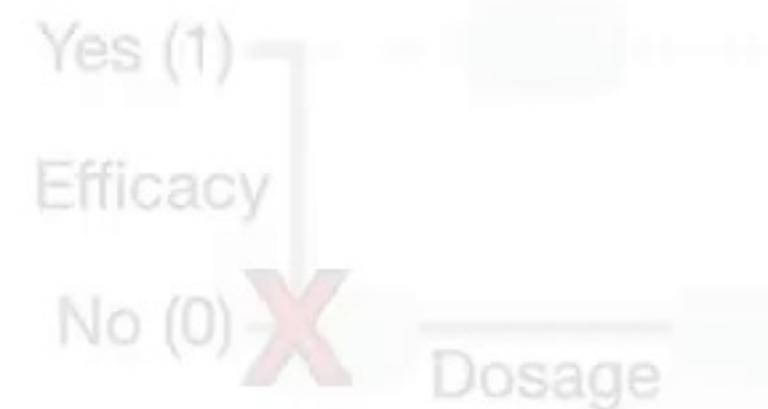




...we plug **2.14** into the  
**softplus** equation...

$$0 + 2.14 = 2.14$$

$$f(x) = \log(1 + e^x)$$



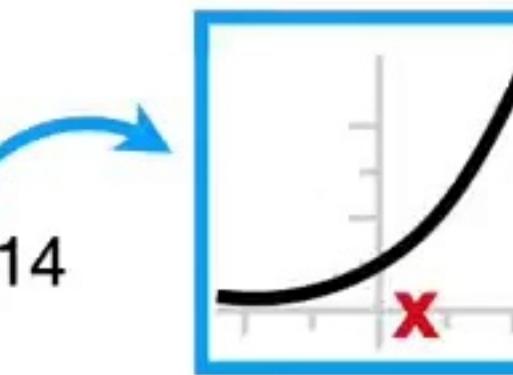
Dosage

0

x -34.4

x -2.52

+ 1.29



x -1.30

sum + -0.58 =

Efficacy



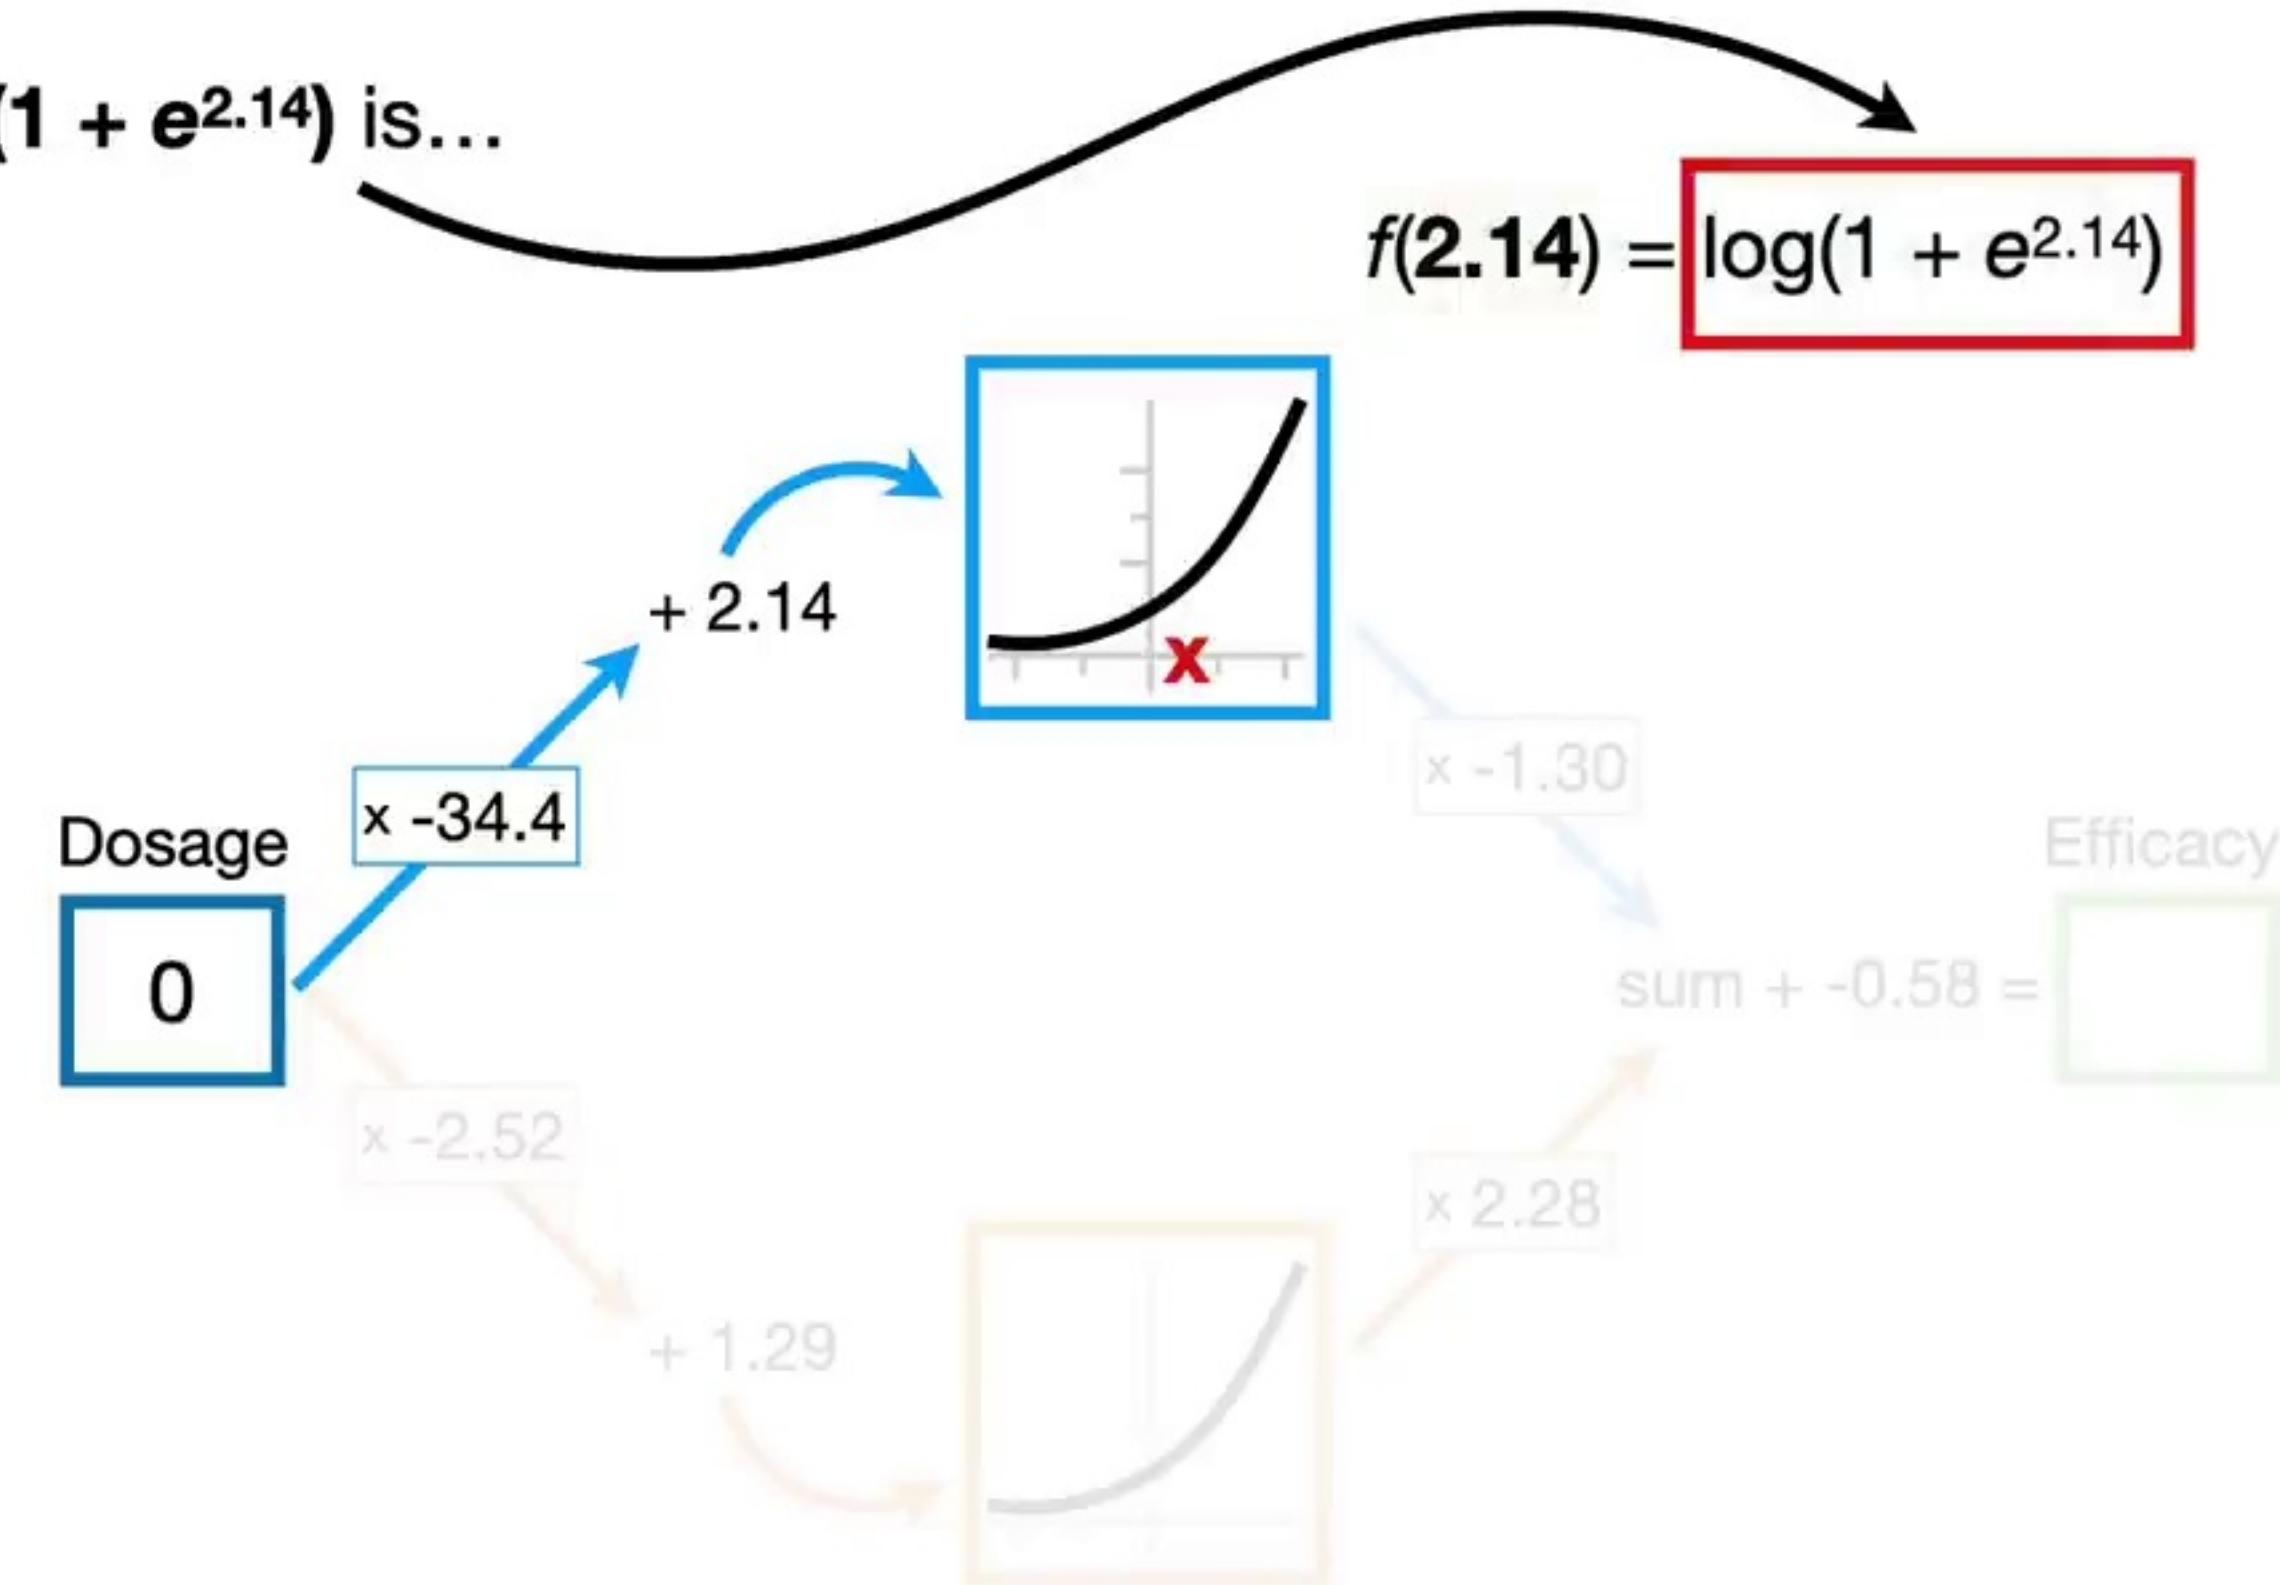
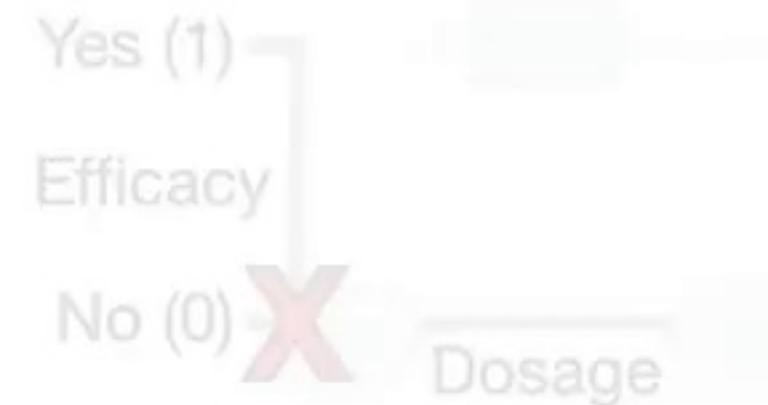
x 2.28



$$0 + 2.14 = 2.14$$

...and the  $\log(1 + e^{2.14})$  is...

$$f(2.14) = \boxed{\log(1 + e^{2.14})}$$





$$0 + 2.14 = 2.14$$

...2.25.

$$f(2.14) = \log(1 + e^{2.14}) = 2.25$$



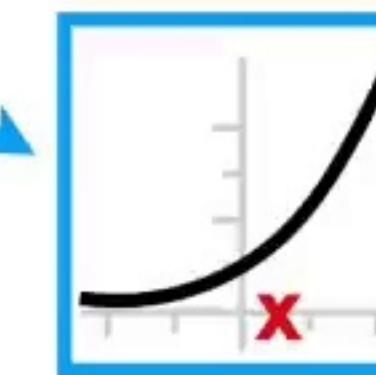
Dosage

$\times -34.4$

$\times -2.52$

$+ 1.29$

$+ 2.14$



$\times -1.30$

sum + -0.58 =

Efficacy



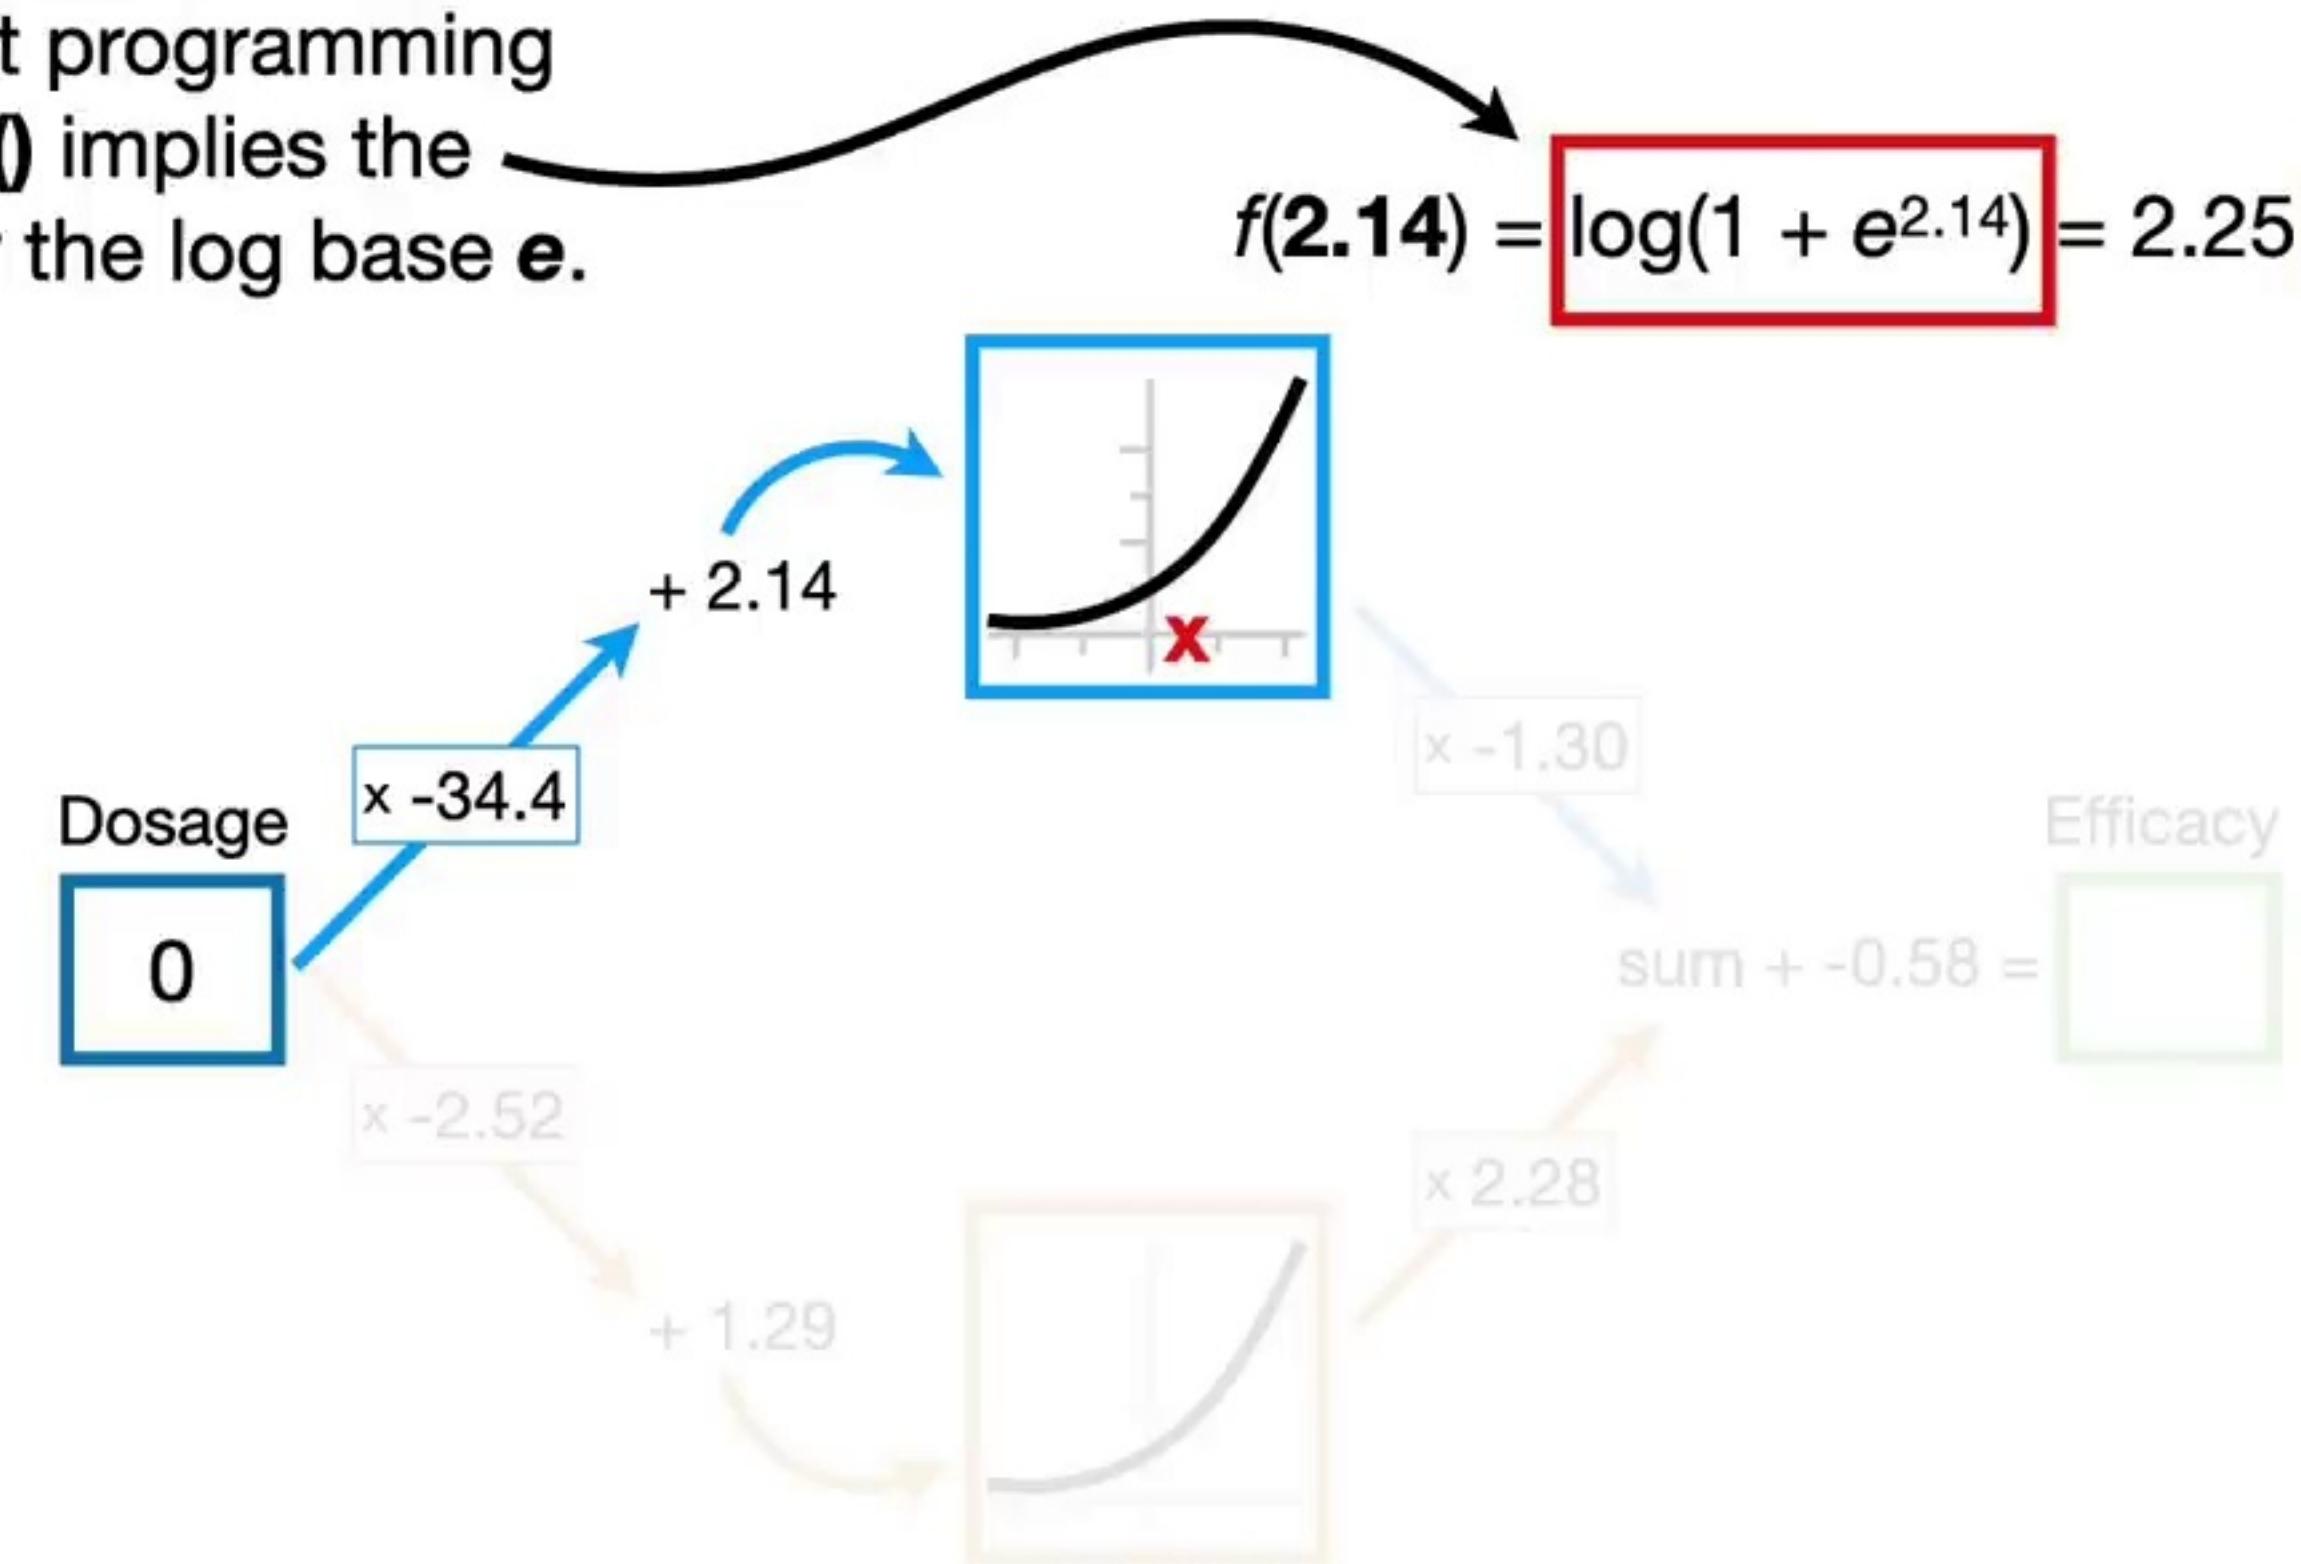
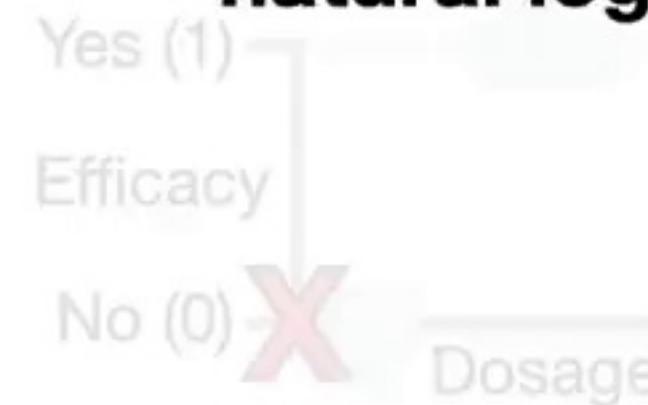
$\times 2.28$



**NOTE:** In statistics, machine learning and most programming languages, **log()** implies the **natural log (ln)**, or the log base **e**.

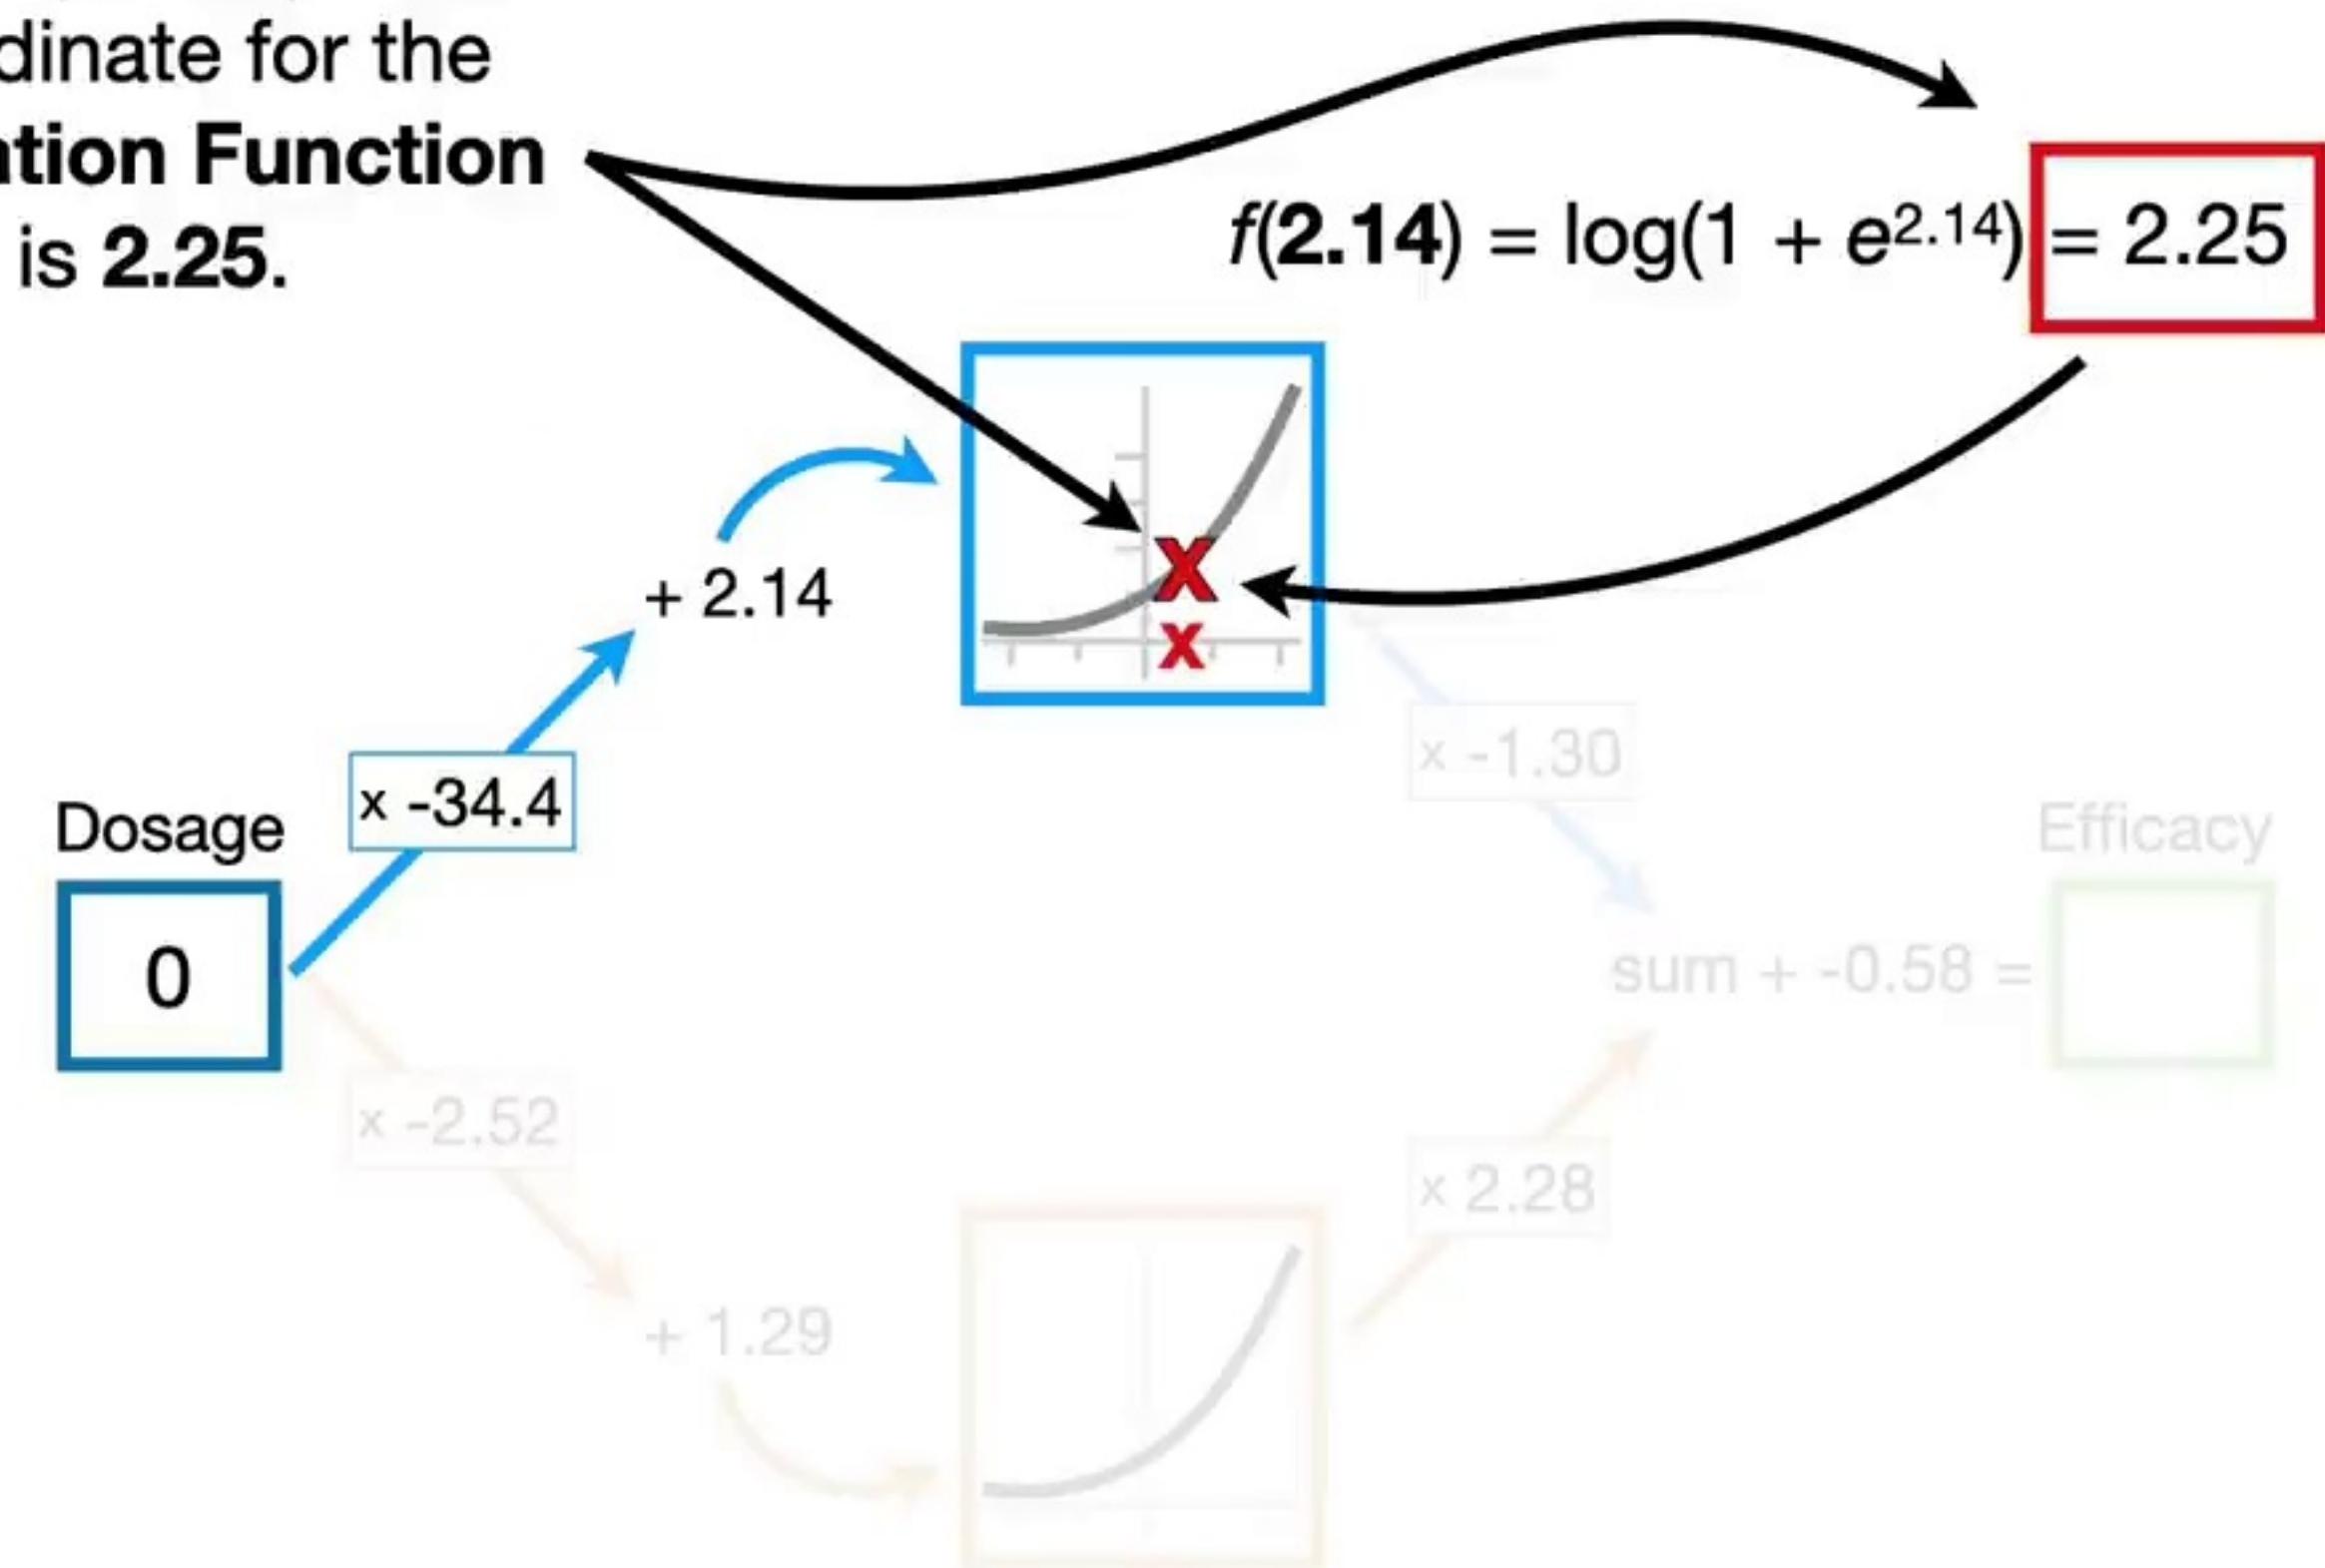
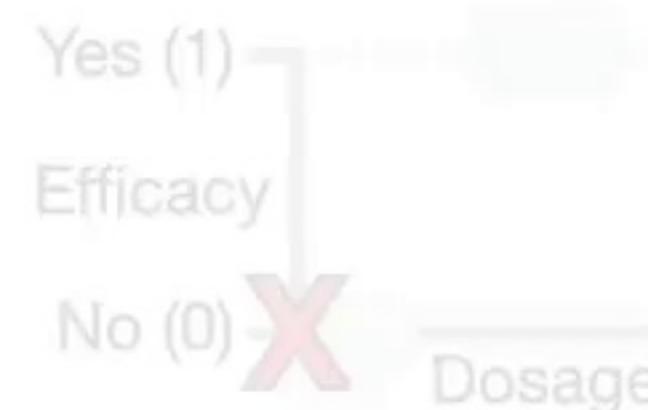
$$0 + 2.14 = 2.14$$

$$f(2.14) = \log(1 + e^{2.14}) = 2.25$$





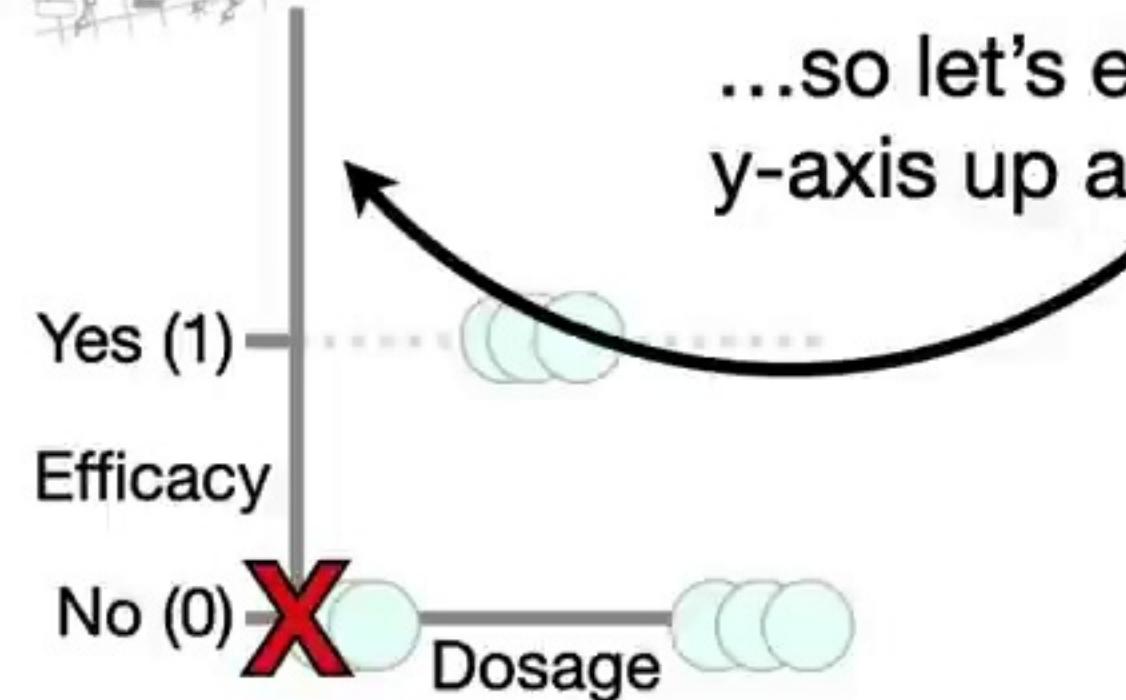
Anyway, the y-axis coordinate for the **Activation Function** is 2.25.



Double  
BAM!!  
**SO!**

$$(0 \times -34.4) + 2.14 = 2.14$$

...so let's extend this  
y-axis up a little bit...



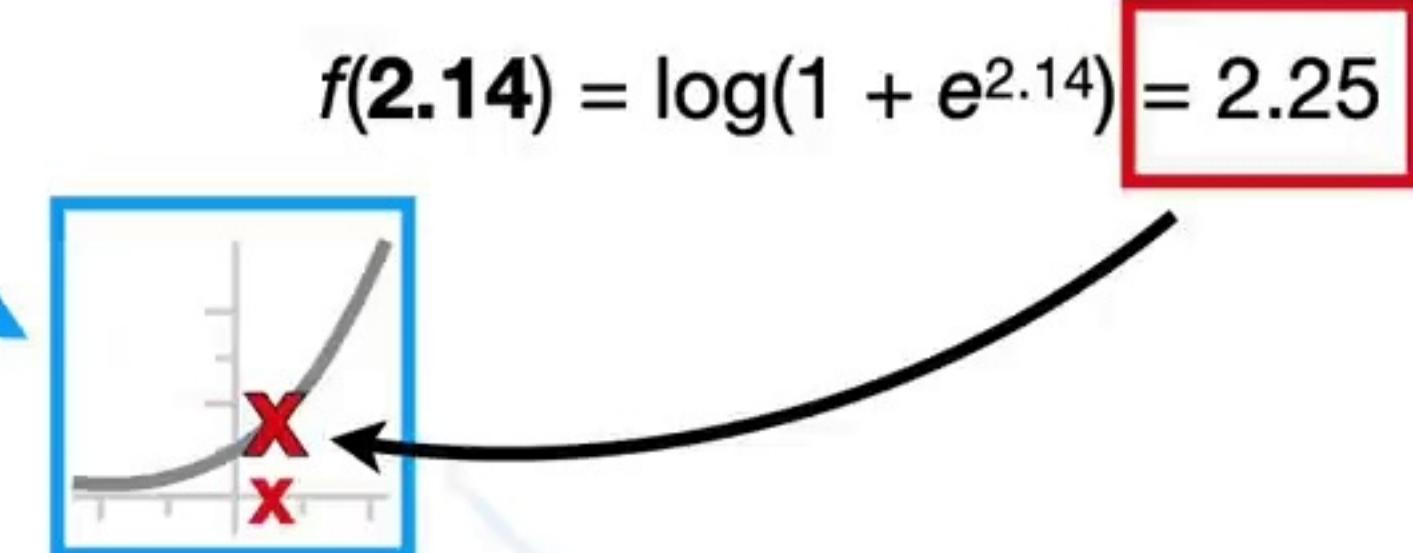
Dosage

$x -34.4$

$x -2.52$

+ 1.29

+ 2.14

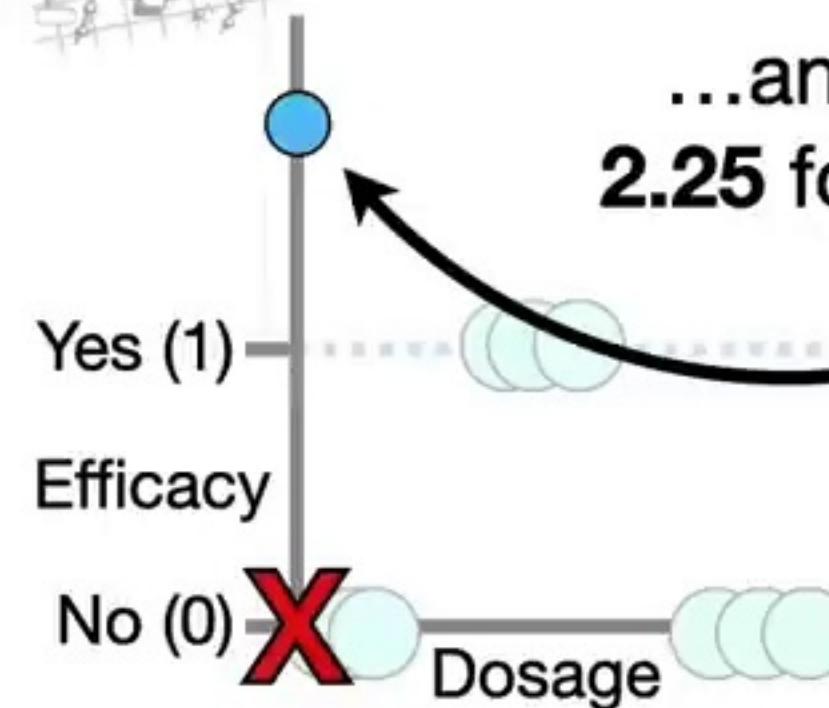


sum + -0.58 =

Double  
BAM!!  
**SO!**

$$(0 \times -34.4) + 2.14 = 2.14$$

...and put a **blue dot** at  
**2.25** for when **Dosage = 0**.



Dosage

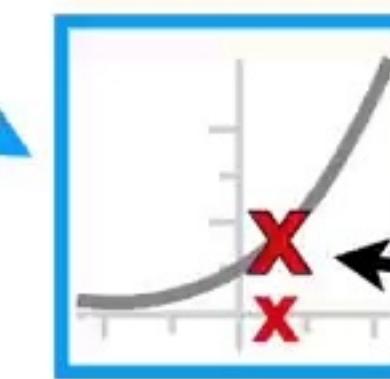
$\times -34.4$

$\times -2.52$

+ 1.29

$\times 2.28$

$$f(2.14) = \log(1 + e^{2.14}) = 2.25$$

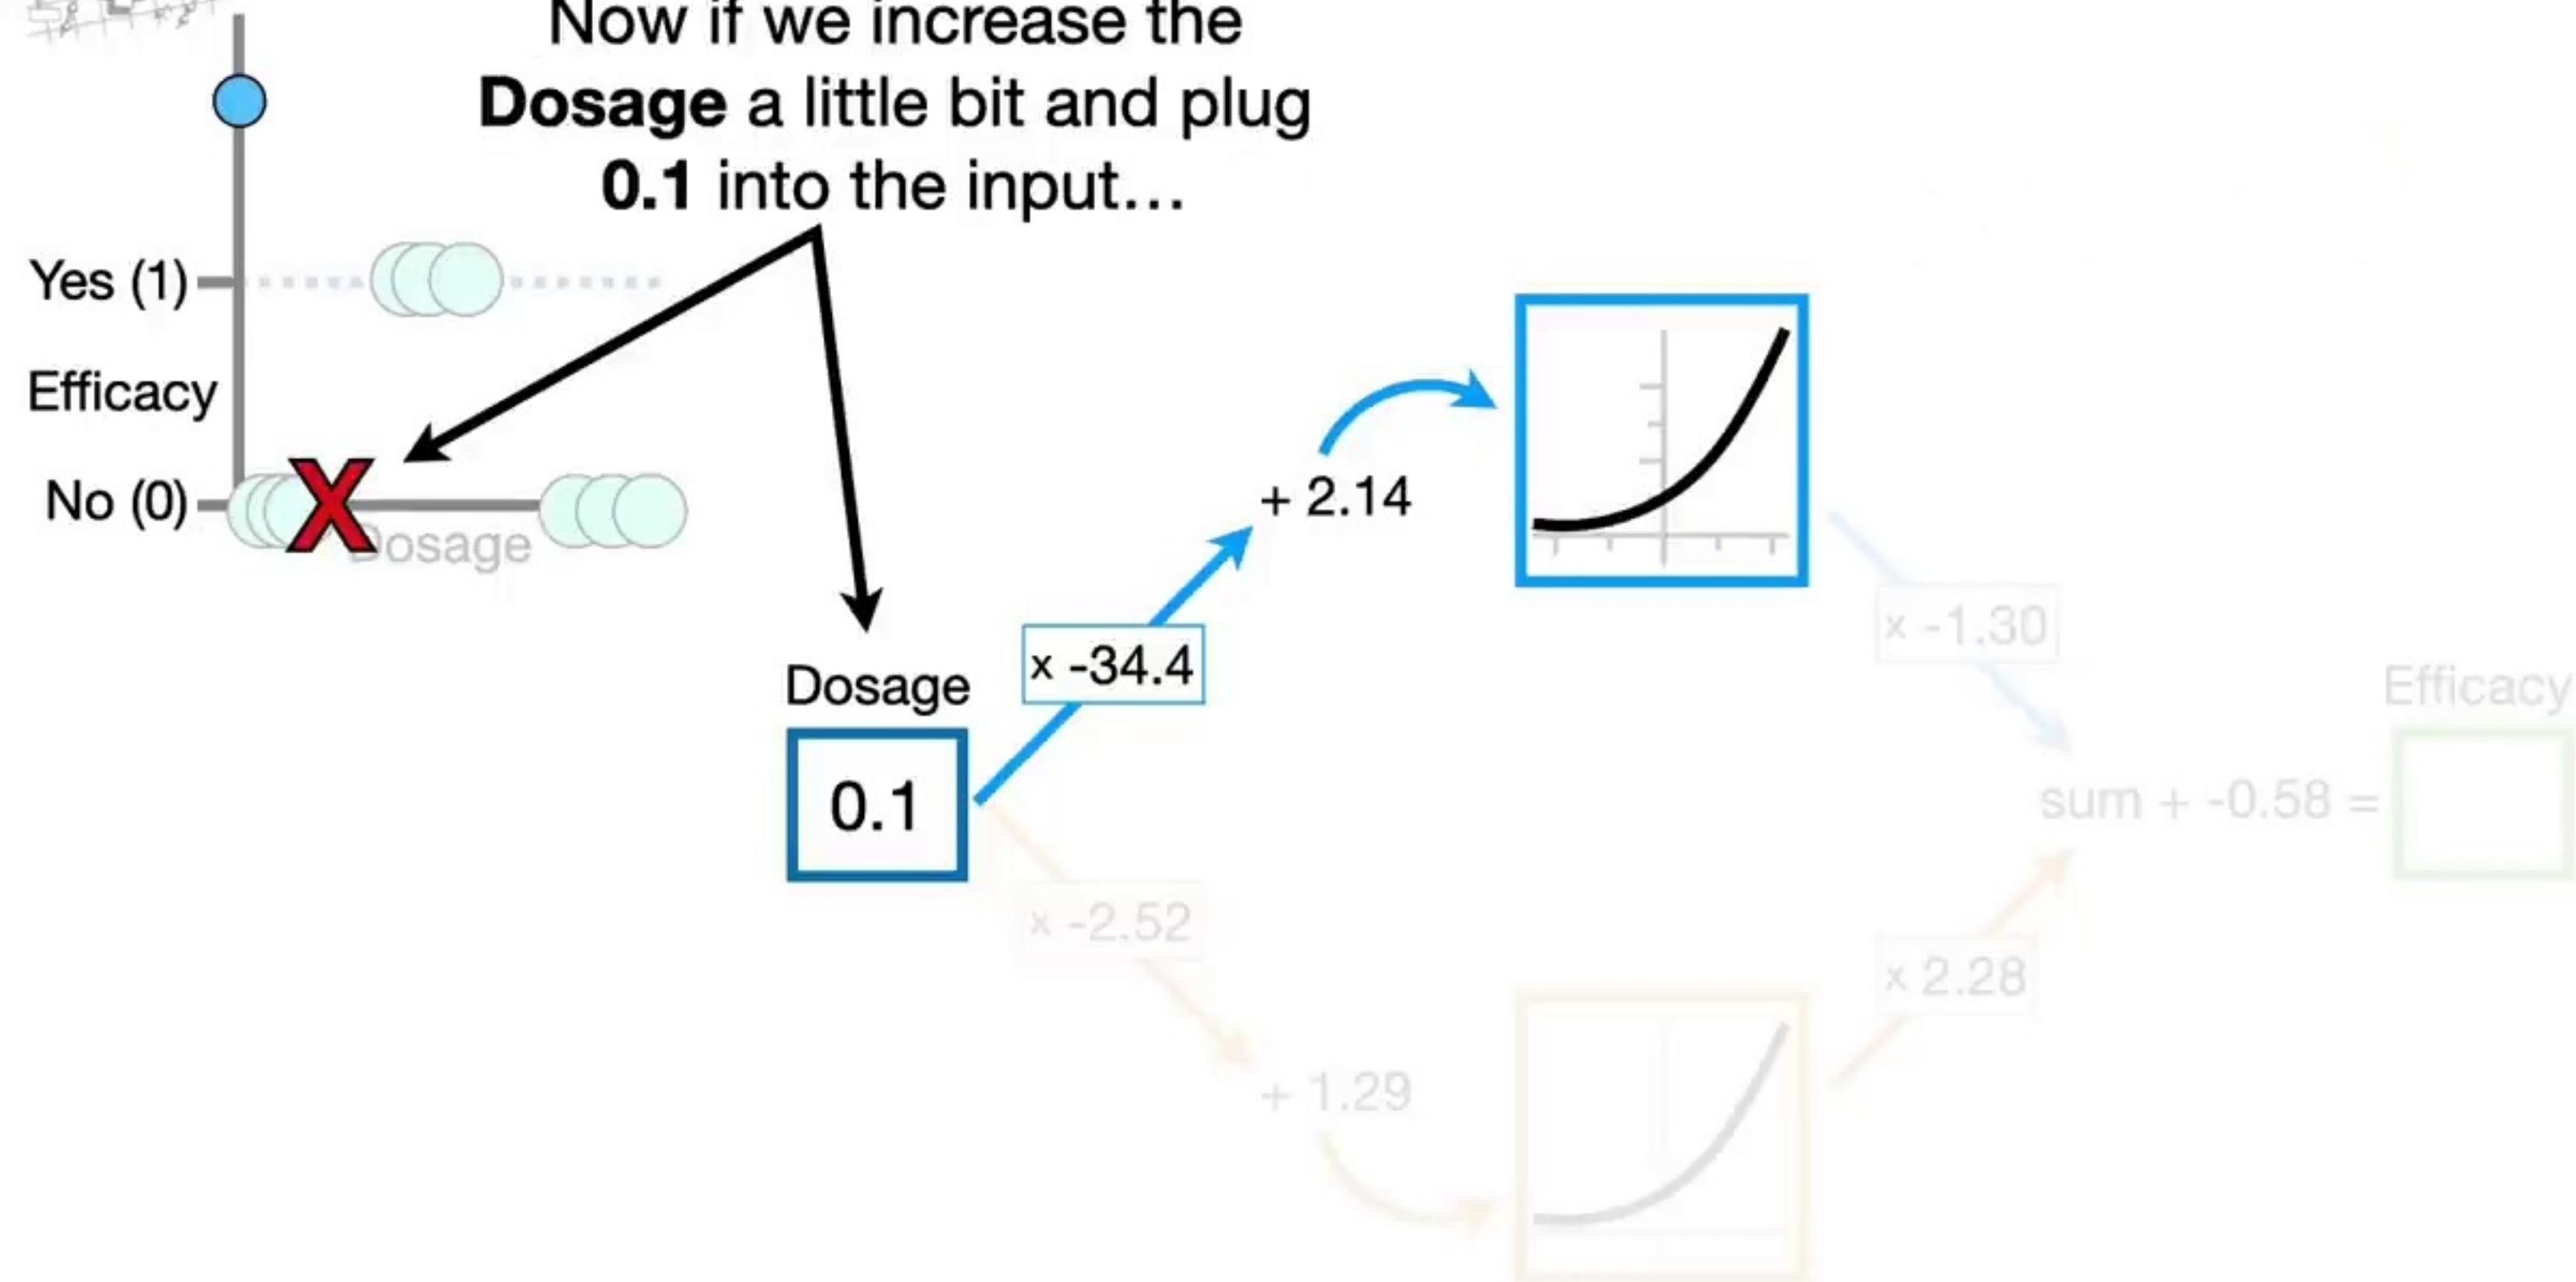


$\times -1.30$

sum + -0.58 =



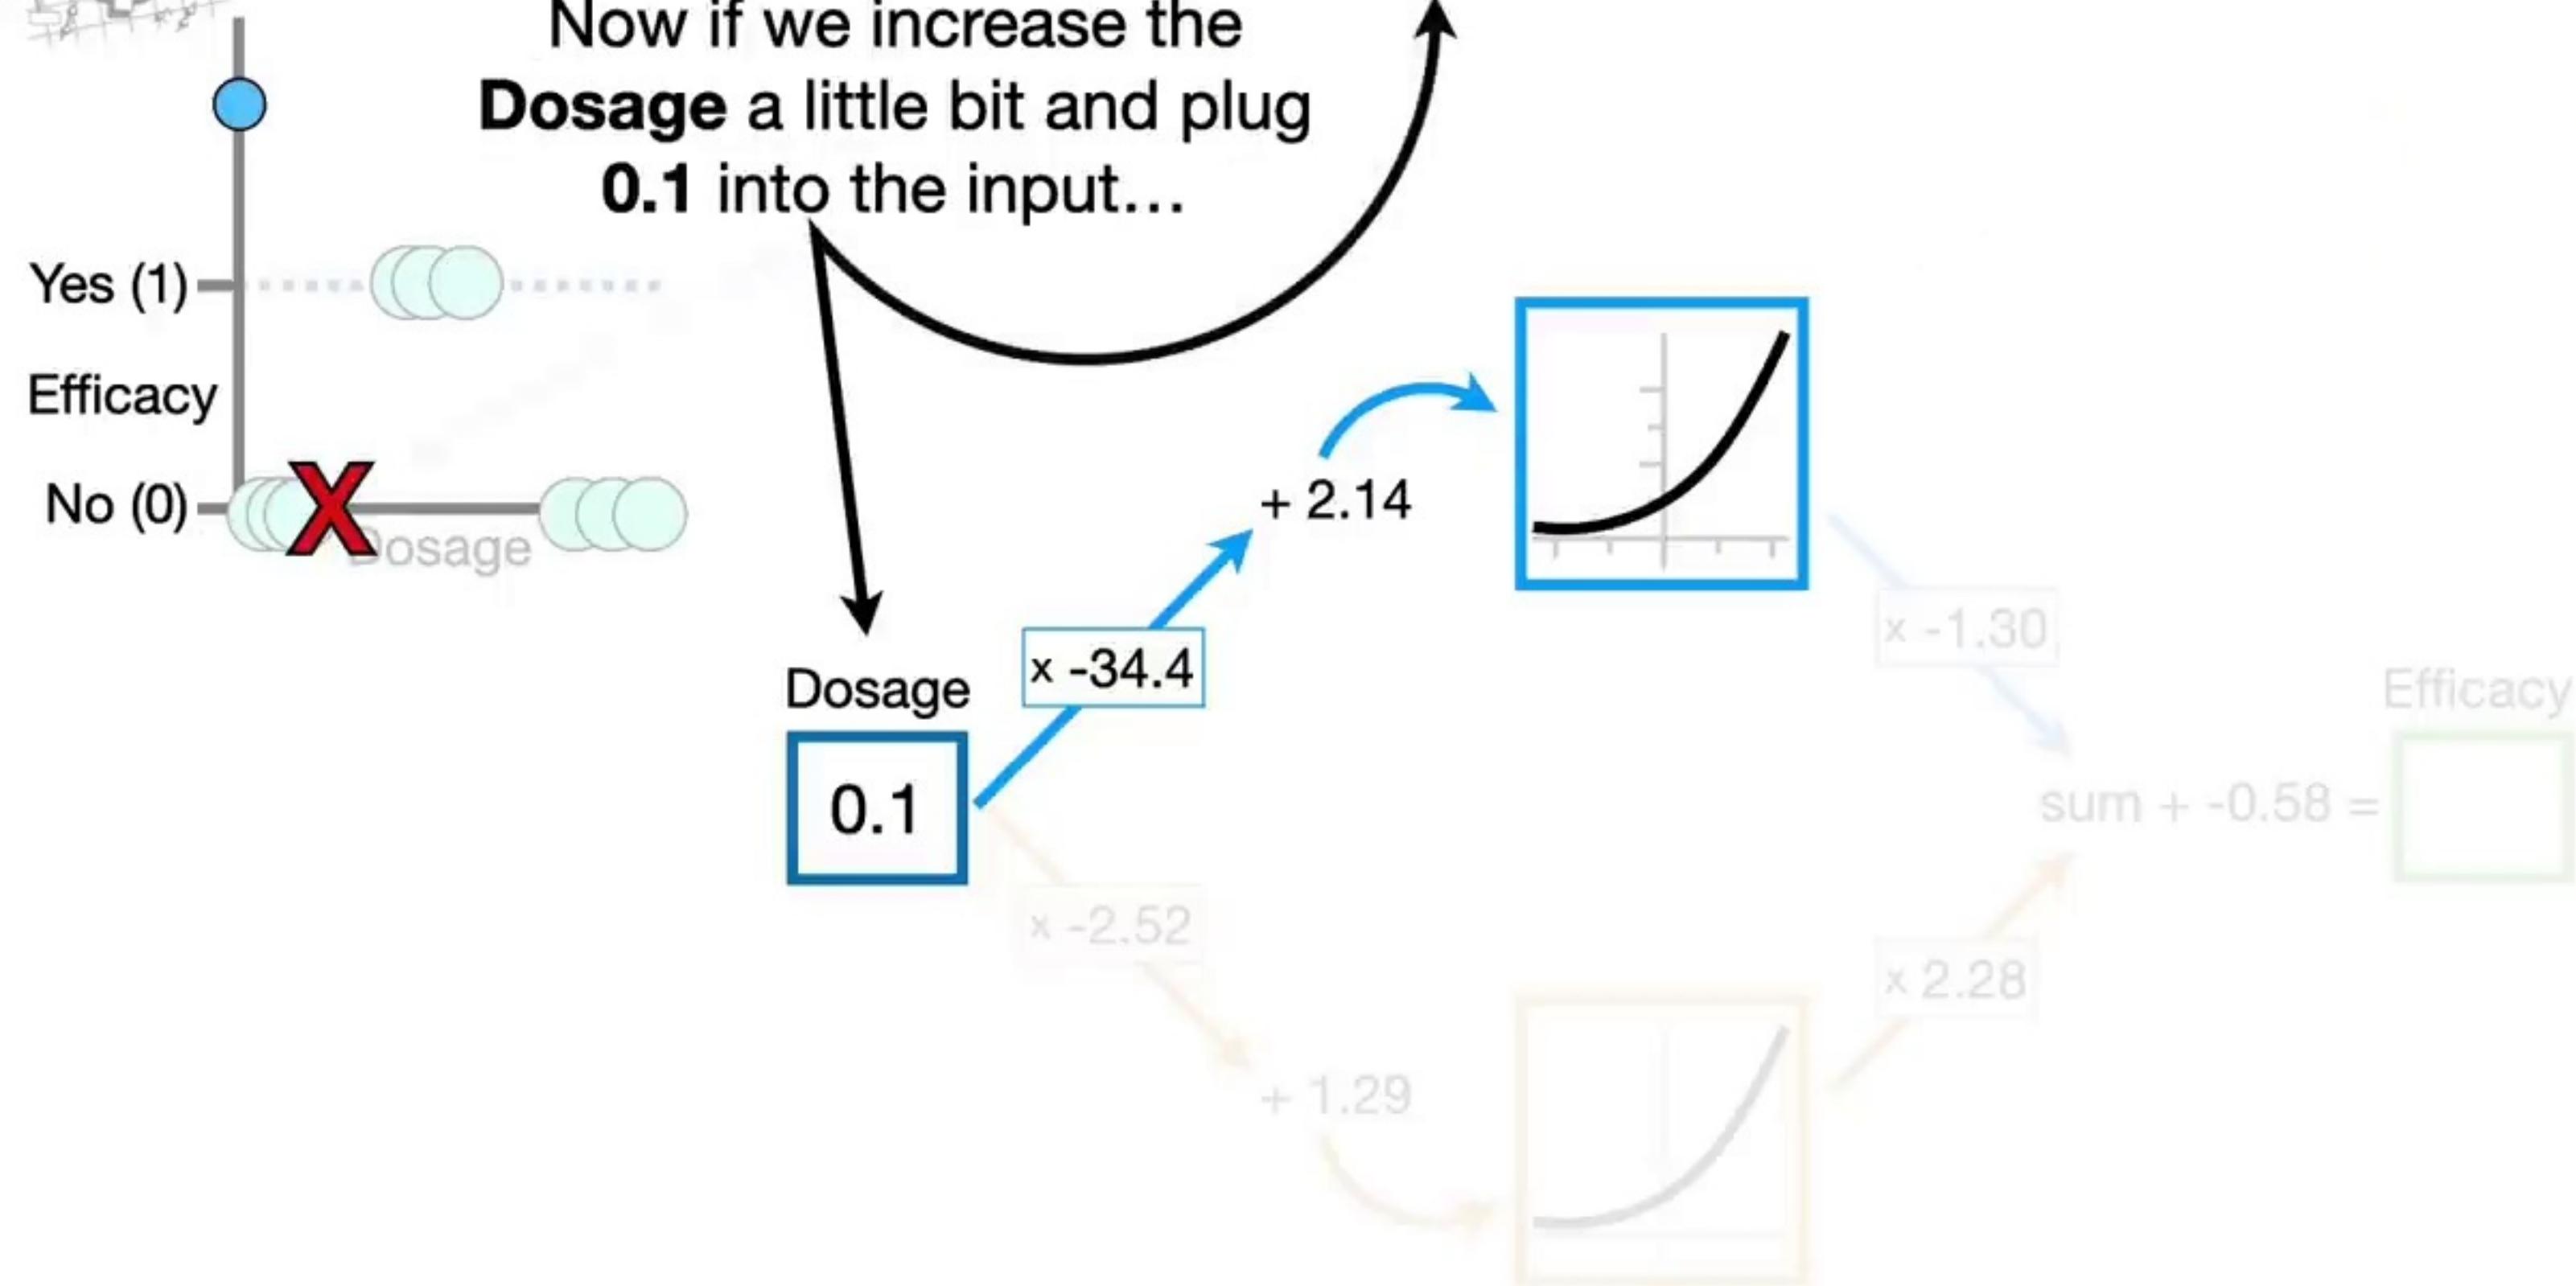
Now if we increase the  
**Dosage** a little bit and plug  
0.1 into the input...





**(Dosage  $\times$  -34.4) + 2.14 = x-axis coordinate**

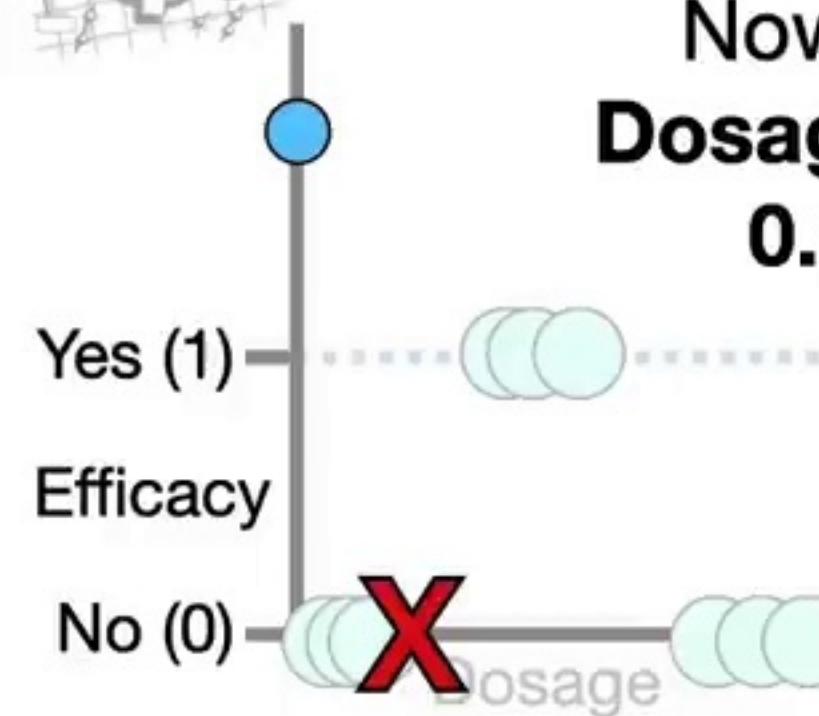
Now if we increase the  
**Dosage** a little bit and plug  
0.1 into the input...





$$(0.1 \times -34.4) + 2.14 = \text{x-axis coordinate}$$

Now if we increase the  
**Dosage** a little bit and plug  
0.1 into the input...



Dosage

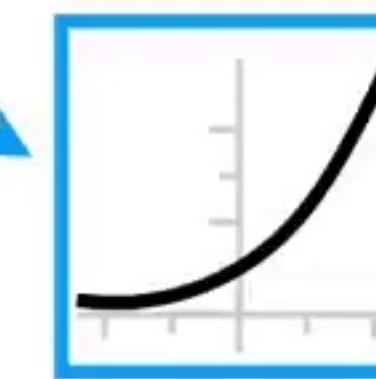
0.1

$\times -34.4$

$\times -2.52$

+ 1.29

+ 2.14



$\times -1.30$

sum + -0.58 =

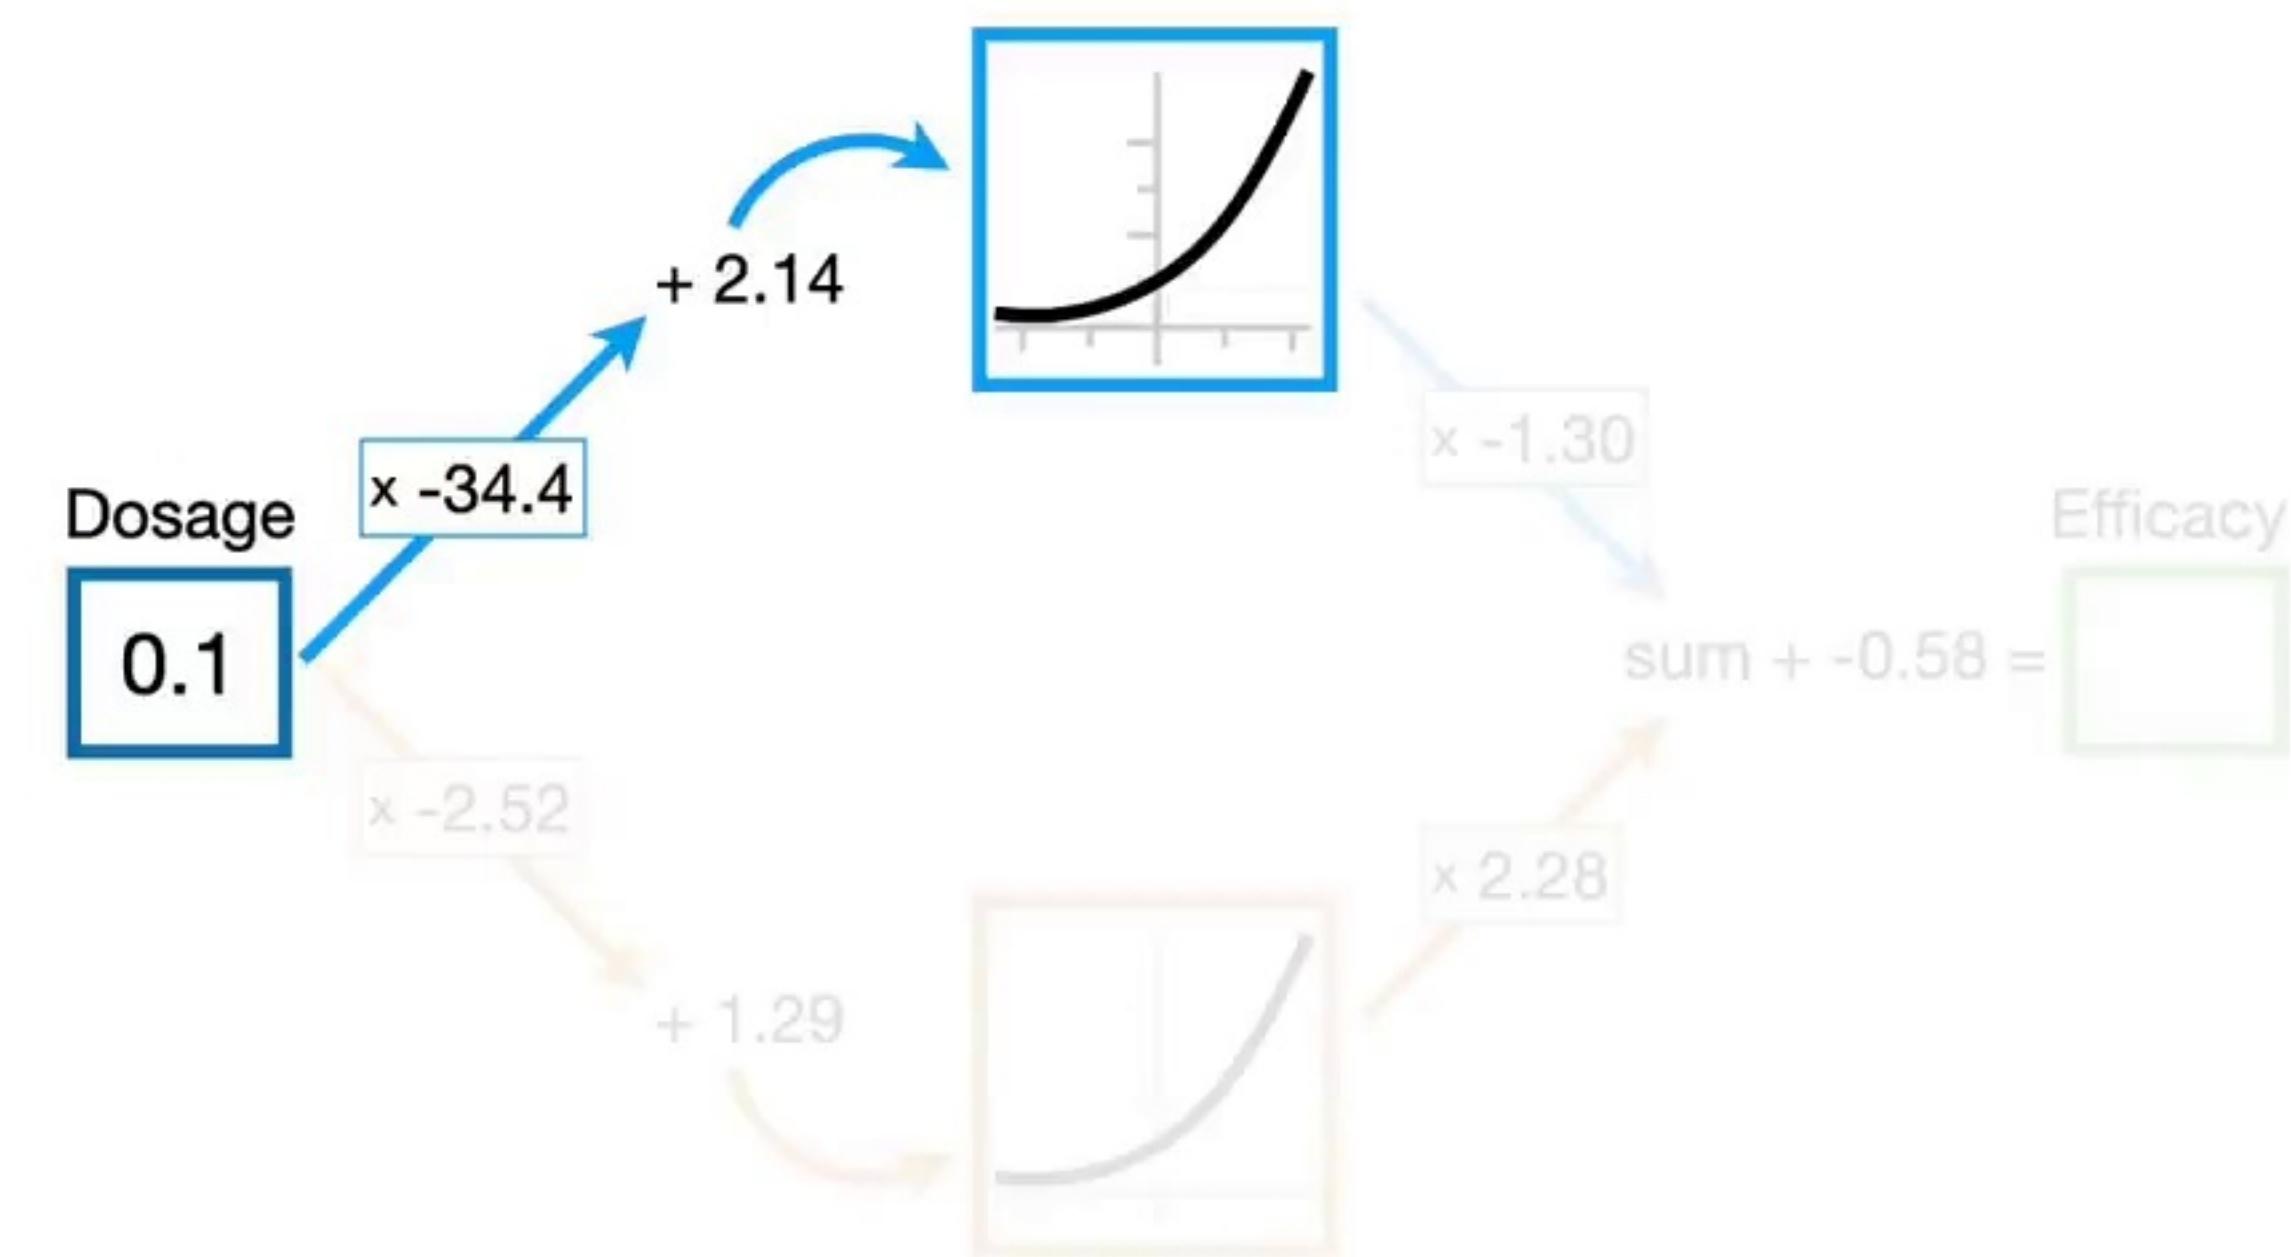
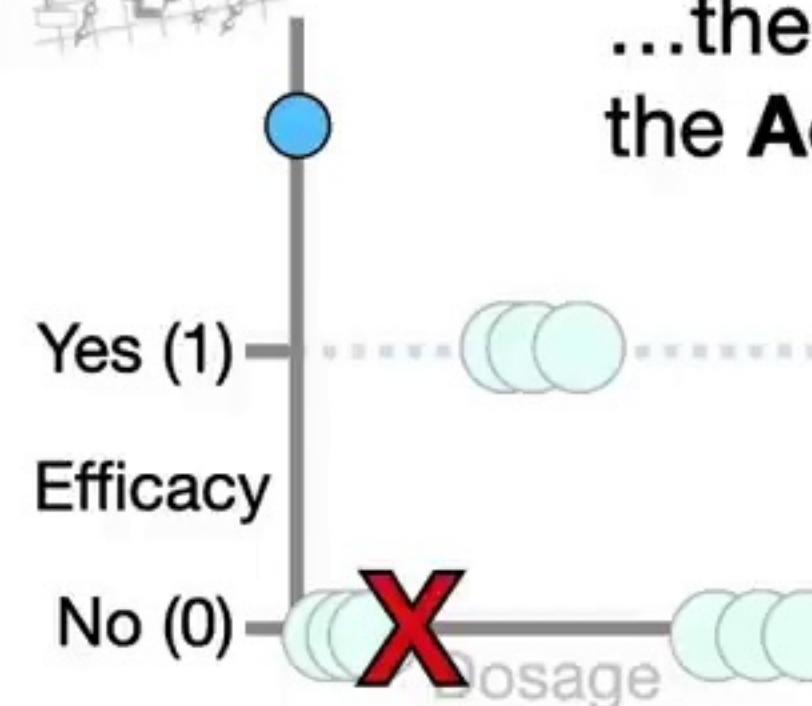


$\times 2.28$



$$(0.1 \times -34.4) + 2.14 = -1.3$$

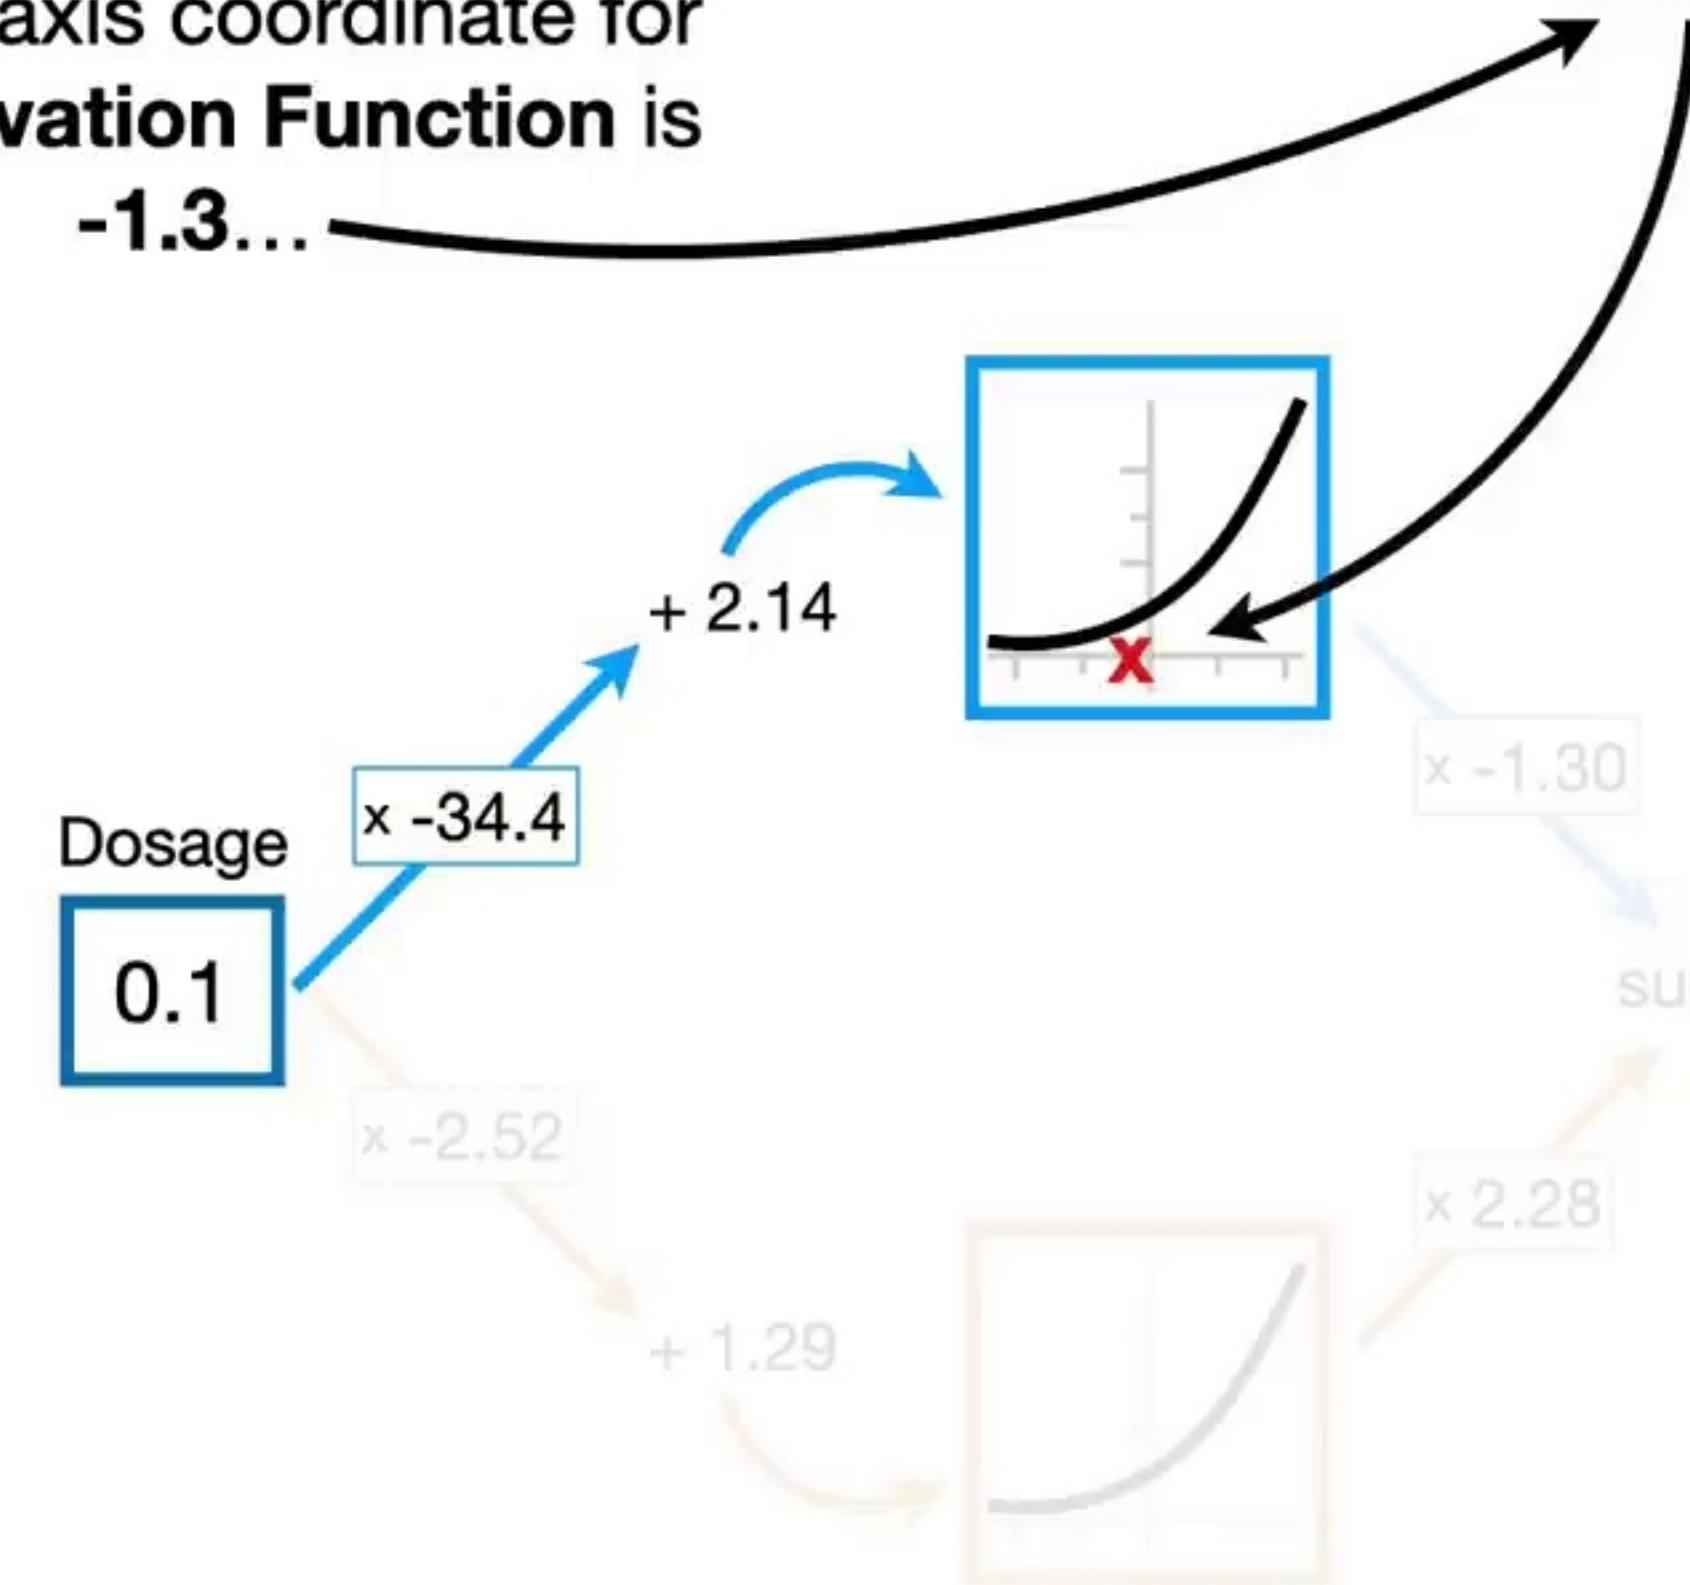
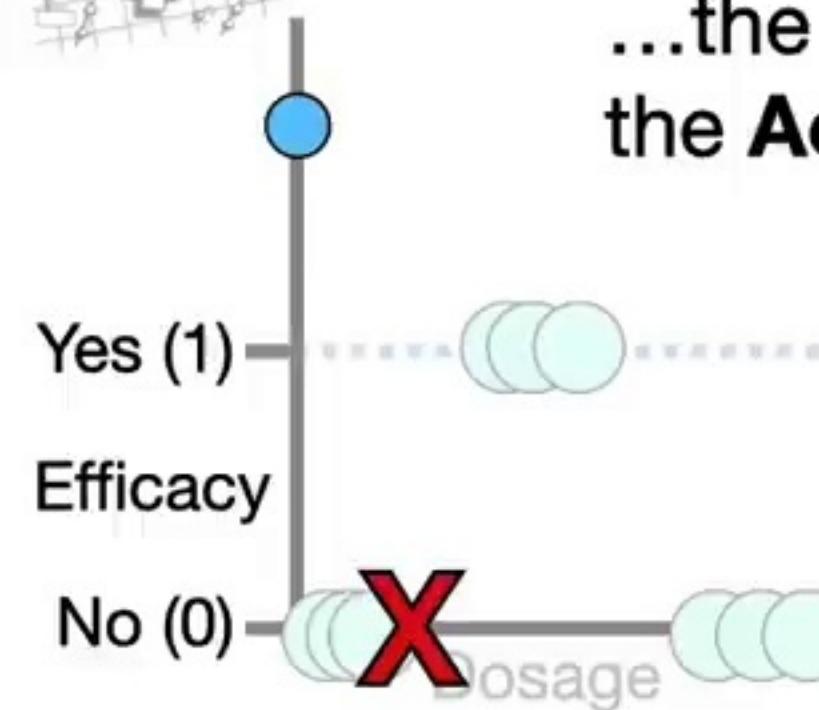
...the x-axis coordinate for  
the **Activation Function** is  
**-1.3...**





$$(0.1 \times -34.4) + 2.14 = -1.3$$

...the x-axis coordinate for  
the **Activation Function** is  
**-1.3...**

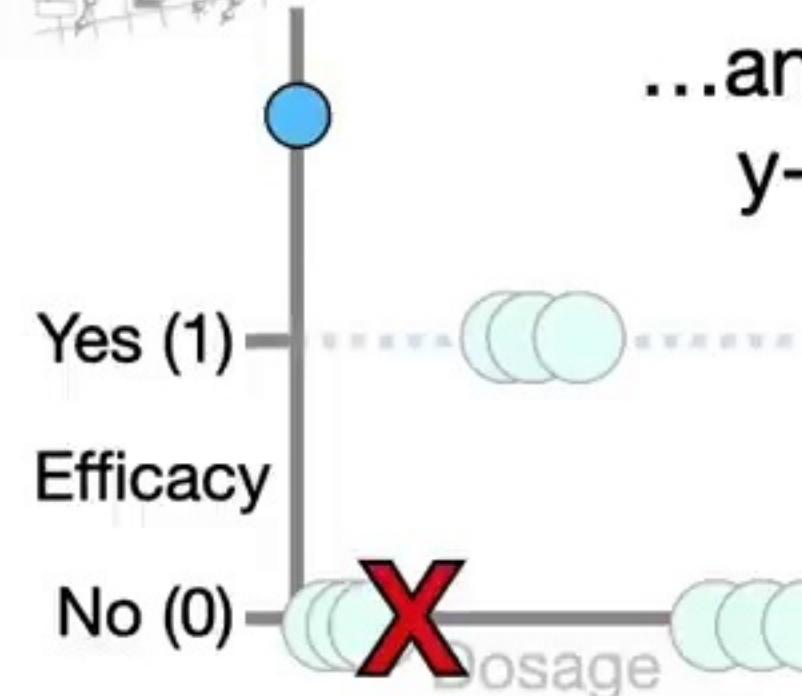




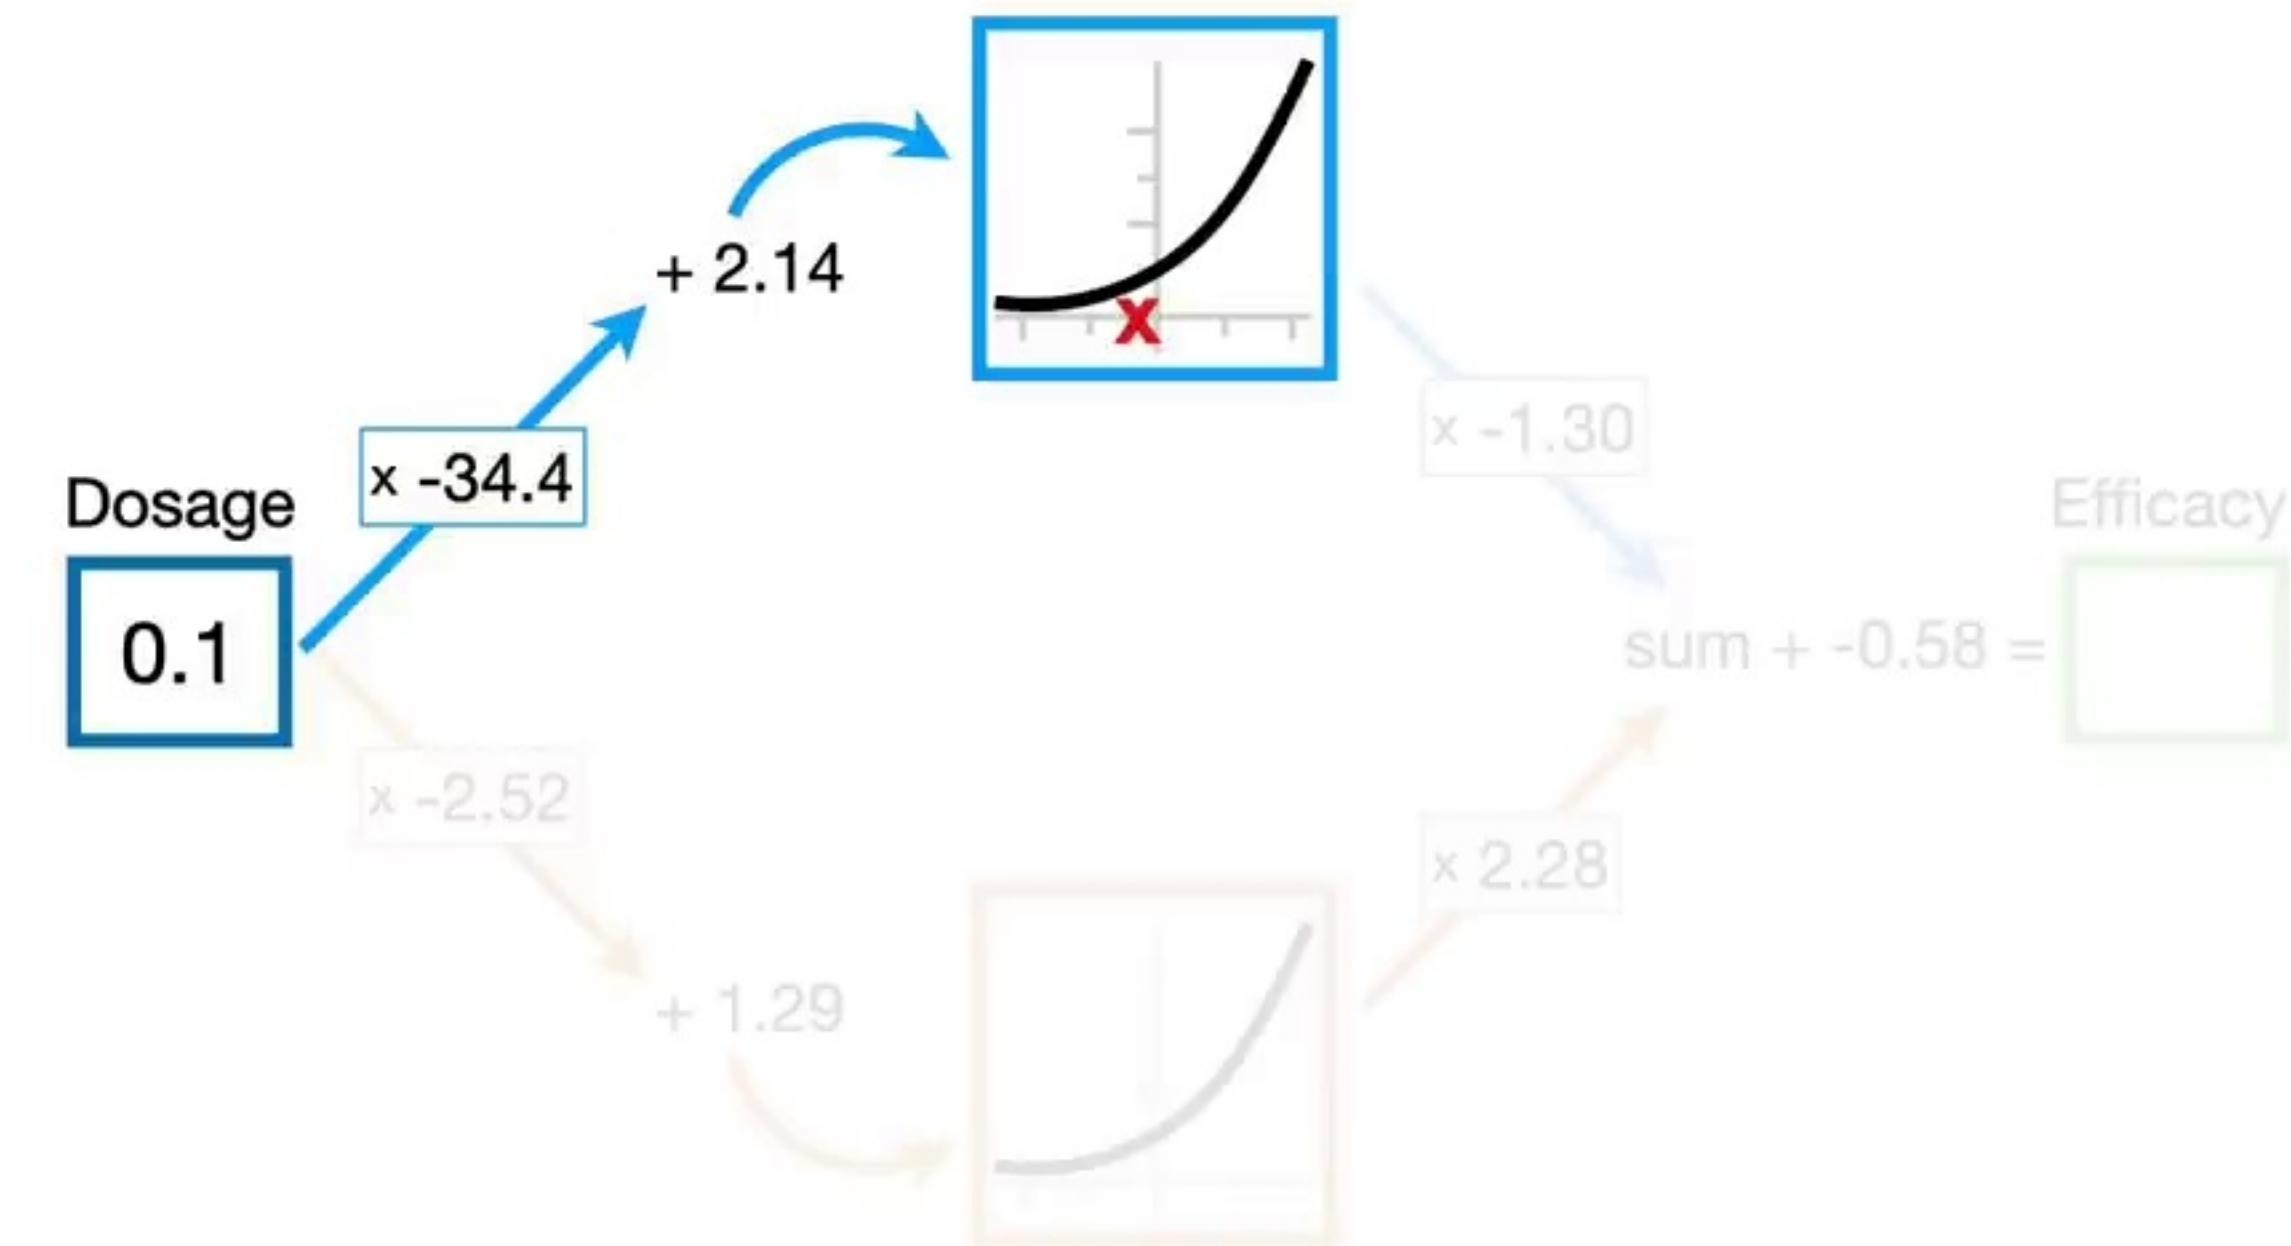
$$(0.1 \times -34.4) + 2.14 = -1.3$$



$$f(x) = \log(1 + e^x)$$



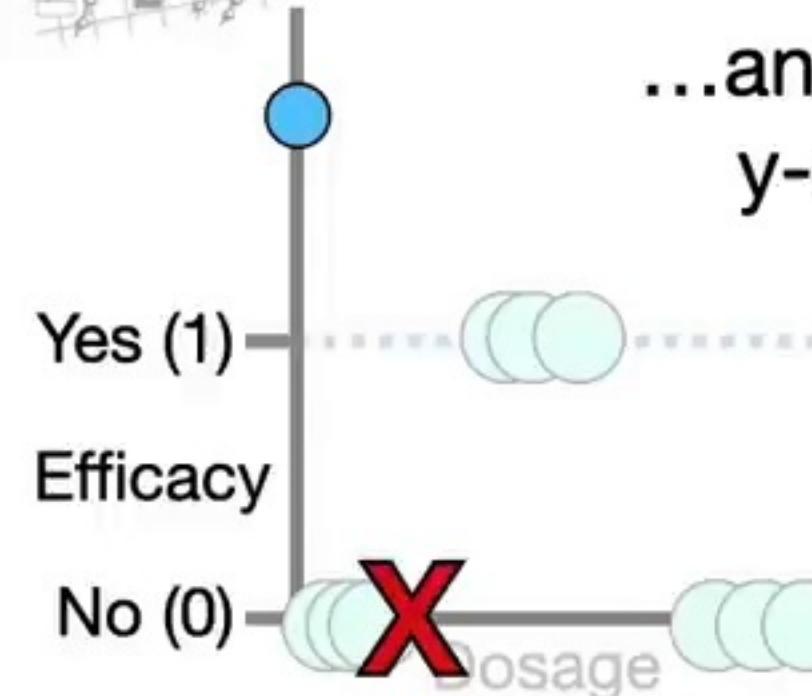
...and the corresponding  
y-axis value is **0.24**.





$$(0.1 \times -34.4) + 2.14 = -1.3$$

$$f(-1.3) = \log(1 + e^{-1.3})$$



Dosage  
0.1

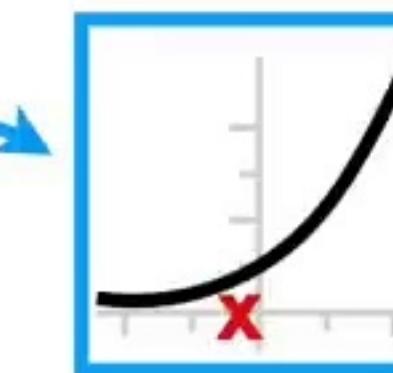
$\times -34.4$

$\times -2.52$

$+ 1.29$

$+ 2.14$

$\times 2.28$



$\times -1.30$

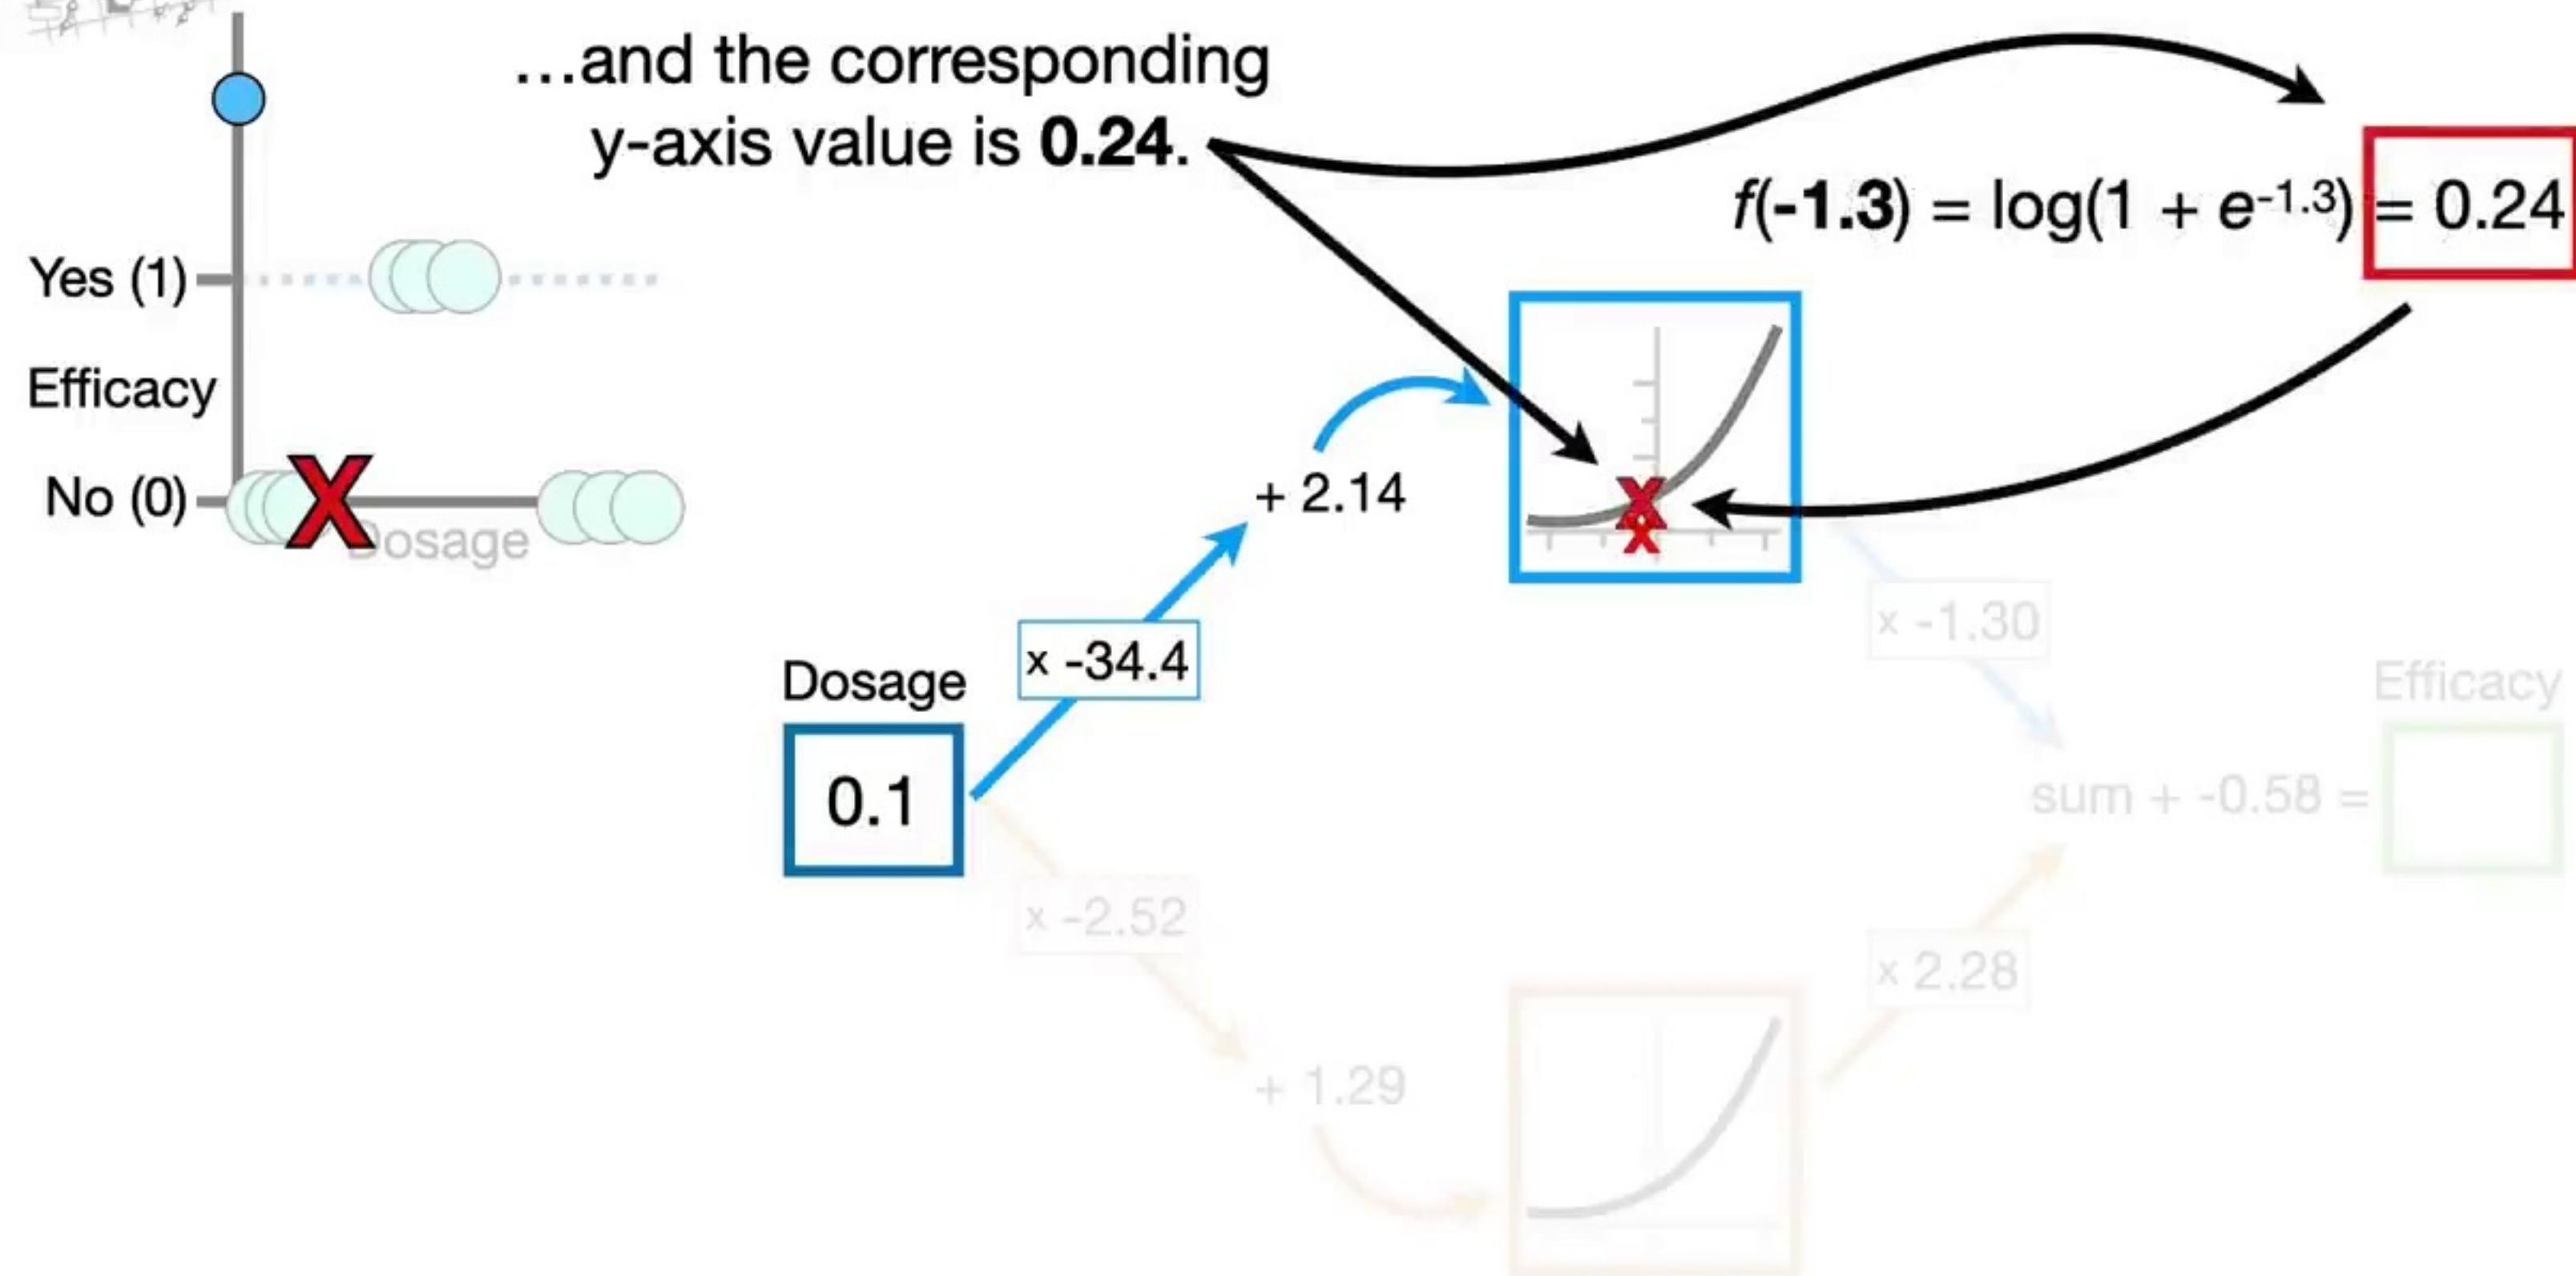
sum + -0.58 =

Efficacy



$$(0.1 \times -34.4) + 2.14 = -1.3$$

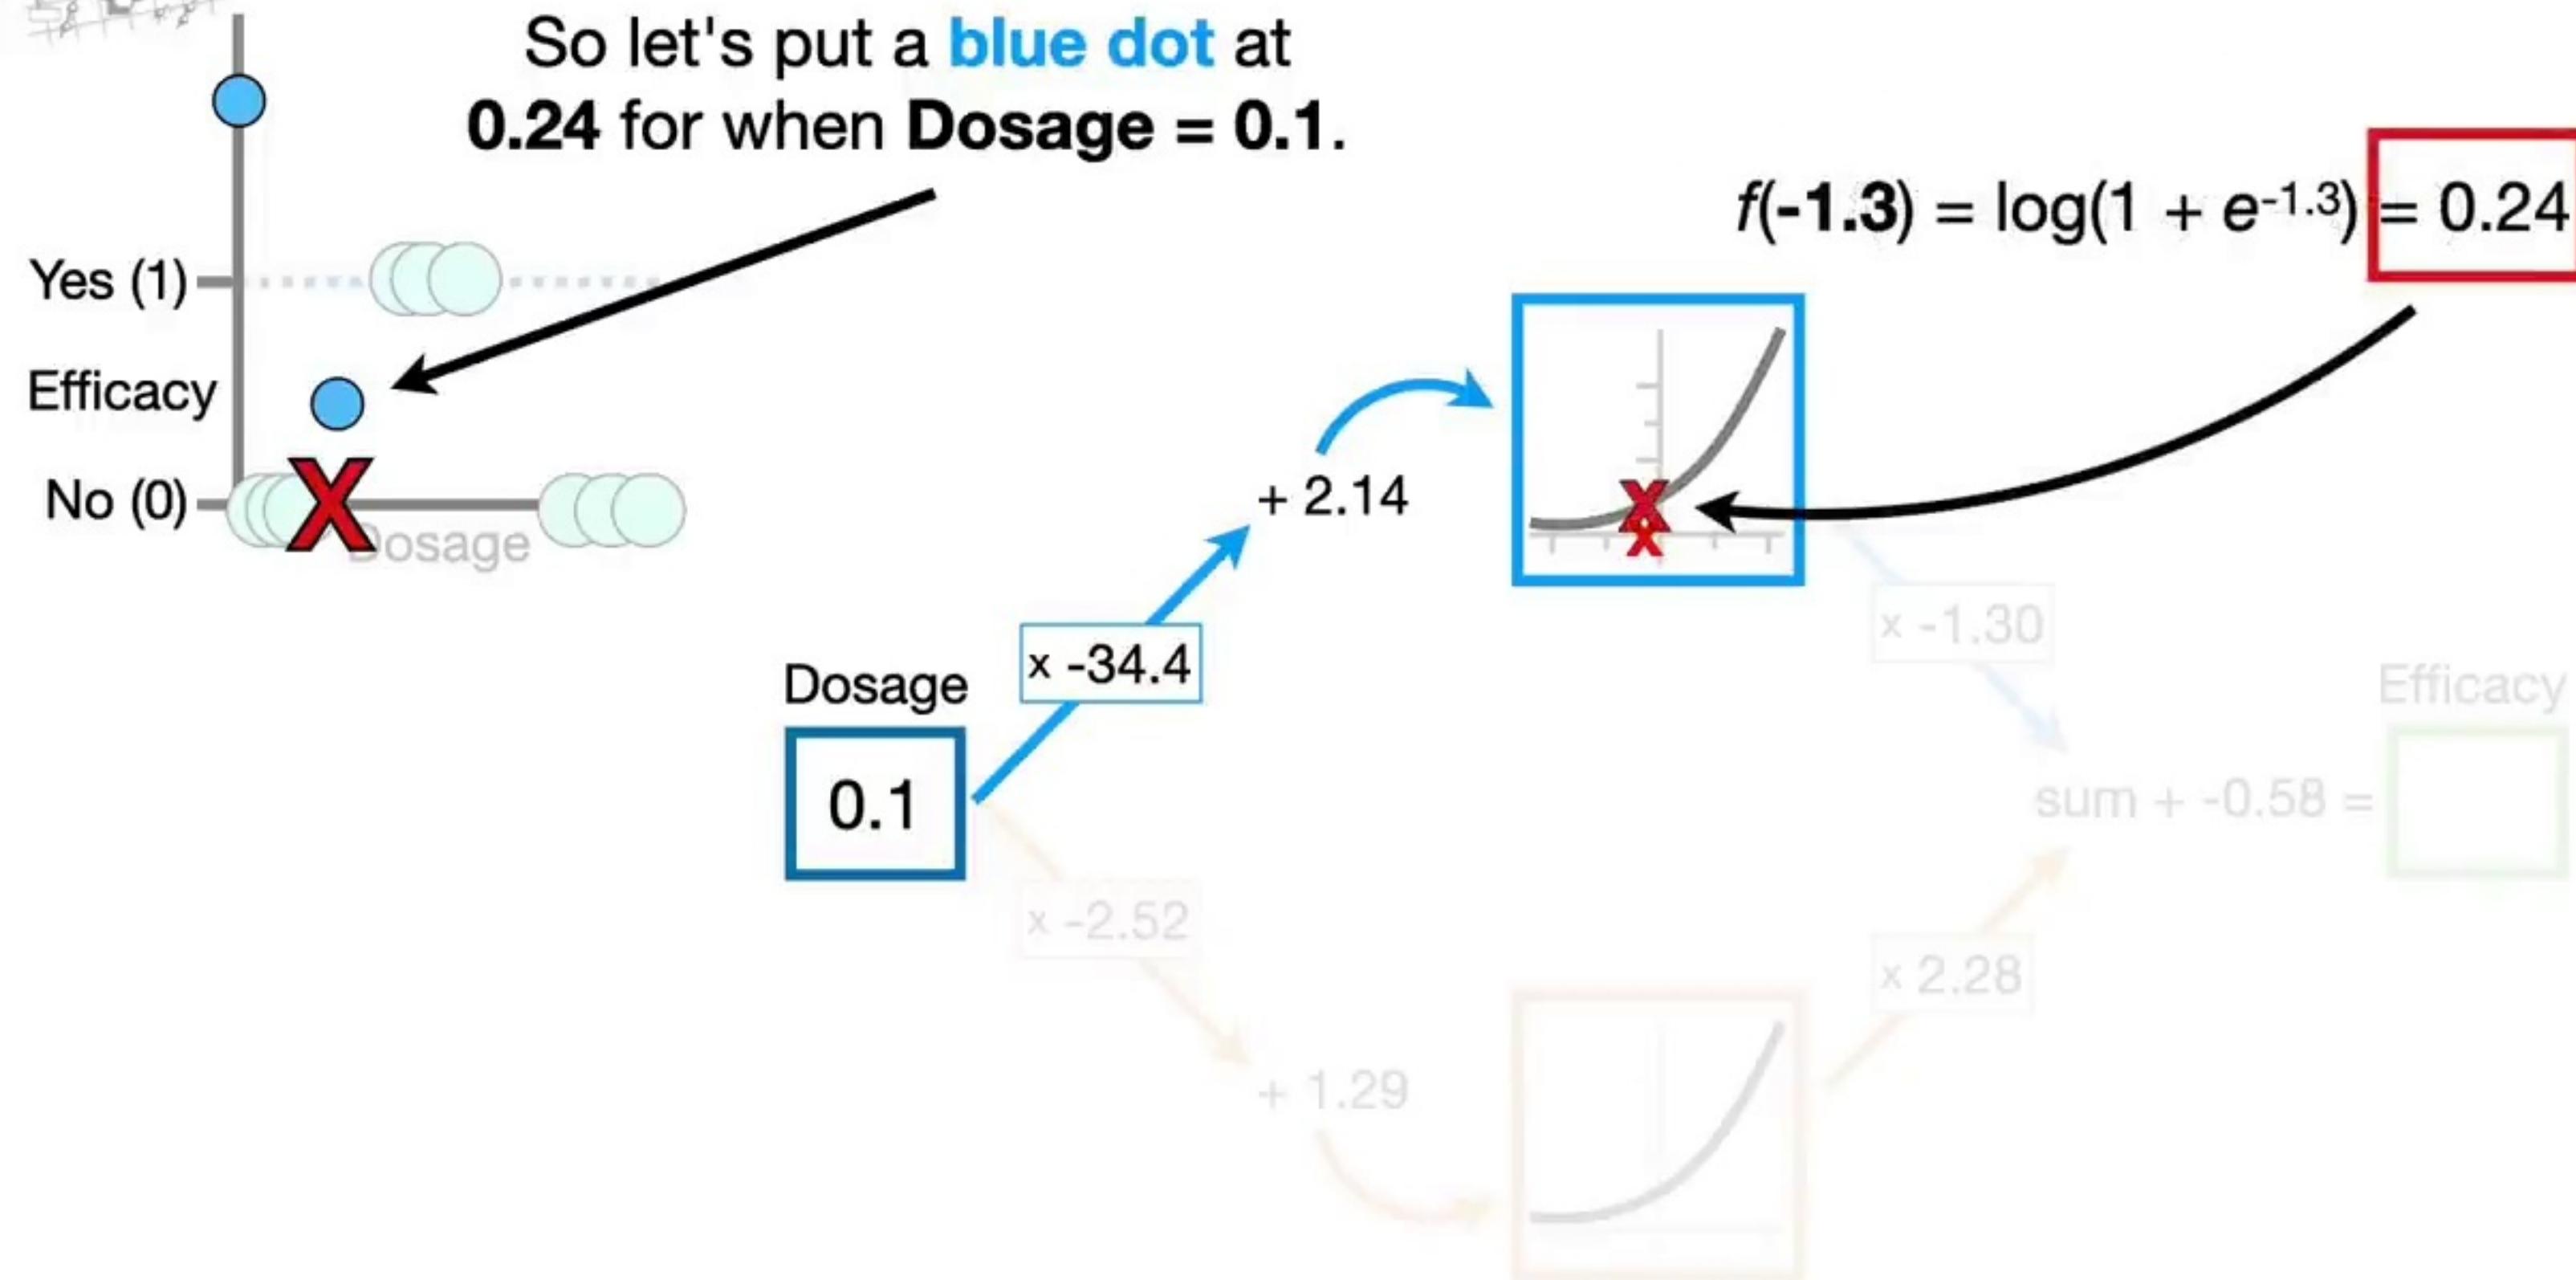
...and the corresponding  
y-axis value is **0.24**.





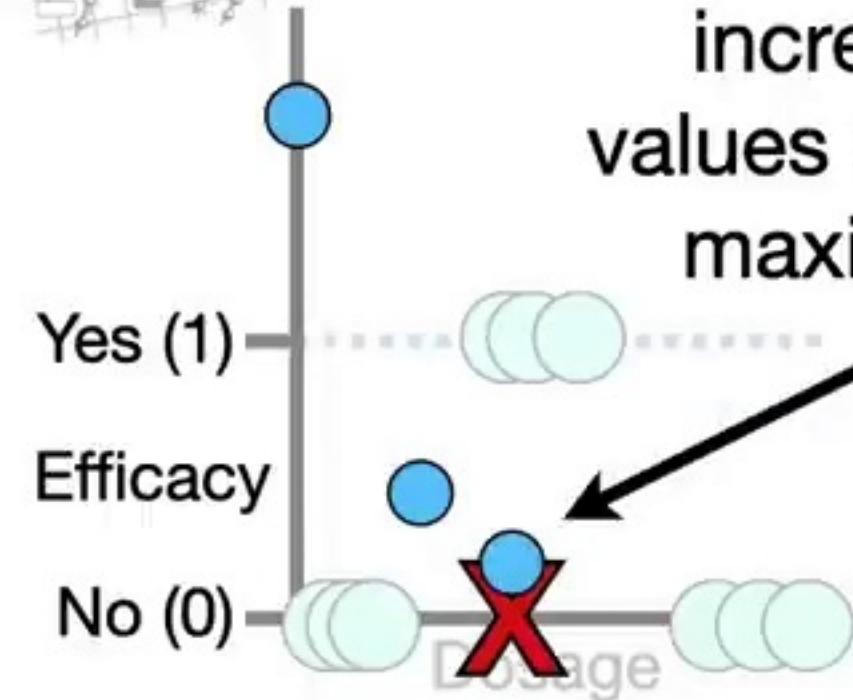
$$(0.1 \times -34.4) + 2.14 = -1.3$$

So let's put a **blue dot** at  
**0.24** for when **Dosage = 0.1**.





And if we continue to increase the **Dosage** values all the way to **1** (the maximum **Dosage**)...



Dosage

0.4

$x - 34.4$

$x - 2.52$

+ 1.29

+ 2.14



$$(\text{Dosage} \times -34.4) + 2.14 = \text{x-axis coordinate}$$

$$f(x) = \log(1 + e^x)$$

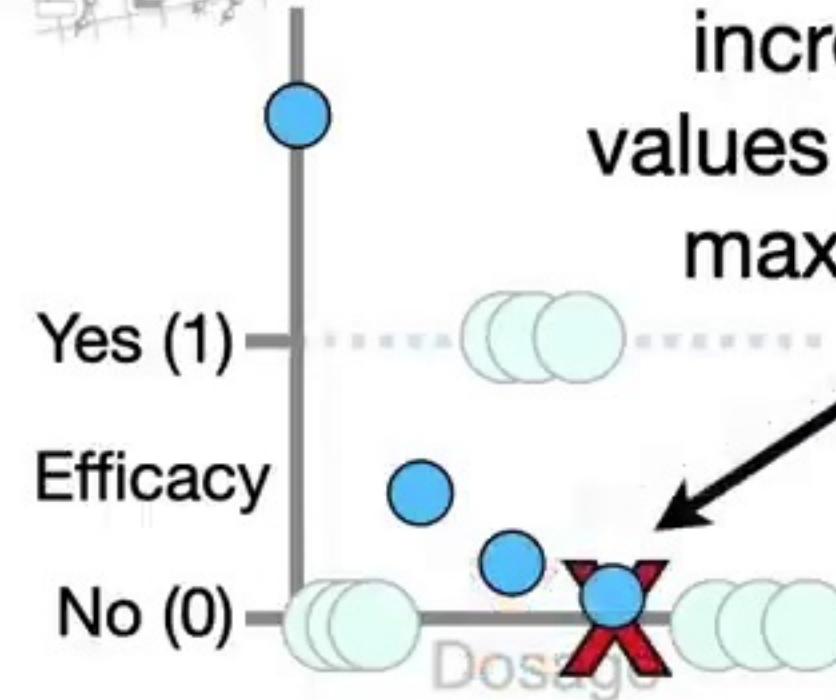
$x - 1.30$

sum + -0.58 =

Efficacy



And if we continue to increase the **Dosage** values all the way to **1** (the maximum **Dosage**)...



Dosage

0.6

$x - 34.4$

$x - 2.52$

+ 1.29

+ 2.14

$x 2.28$

$x - 1.30$

sum + -0.58 =



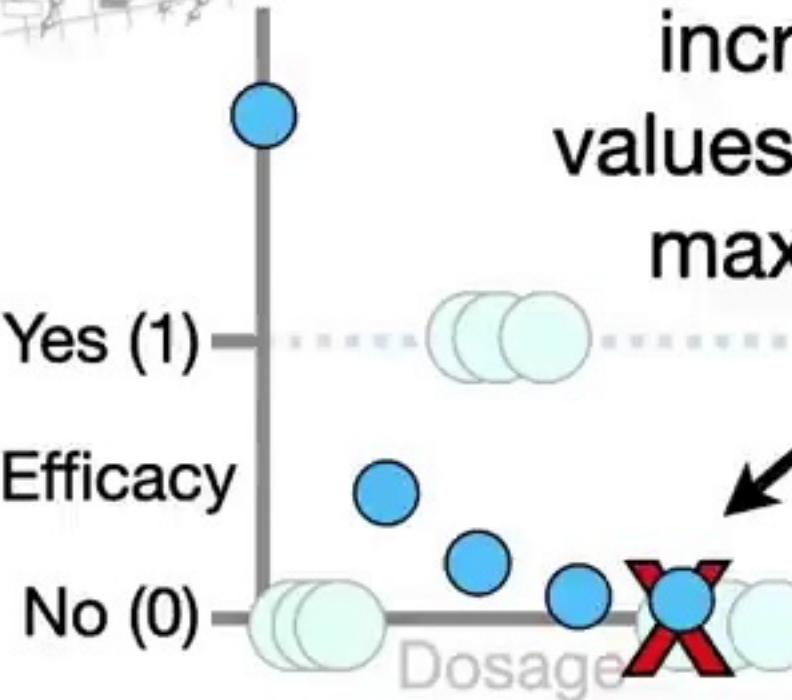
$$(\text{Dosage} \times -34.4) + 2.14 = x\text{-axis coordinate}$$

$$f(x) = \log(1 + e^x)$$

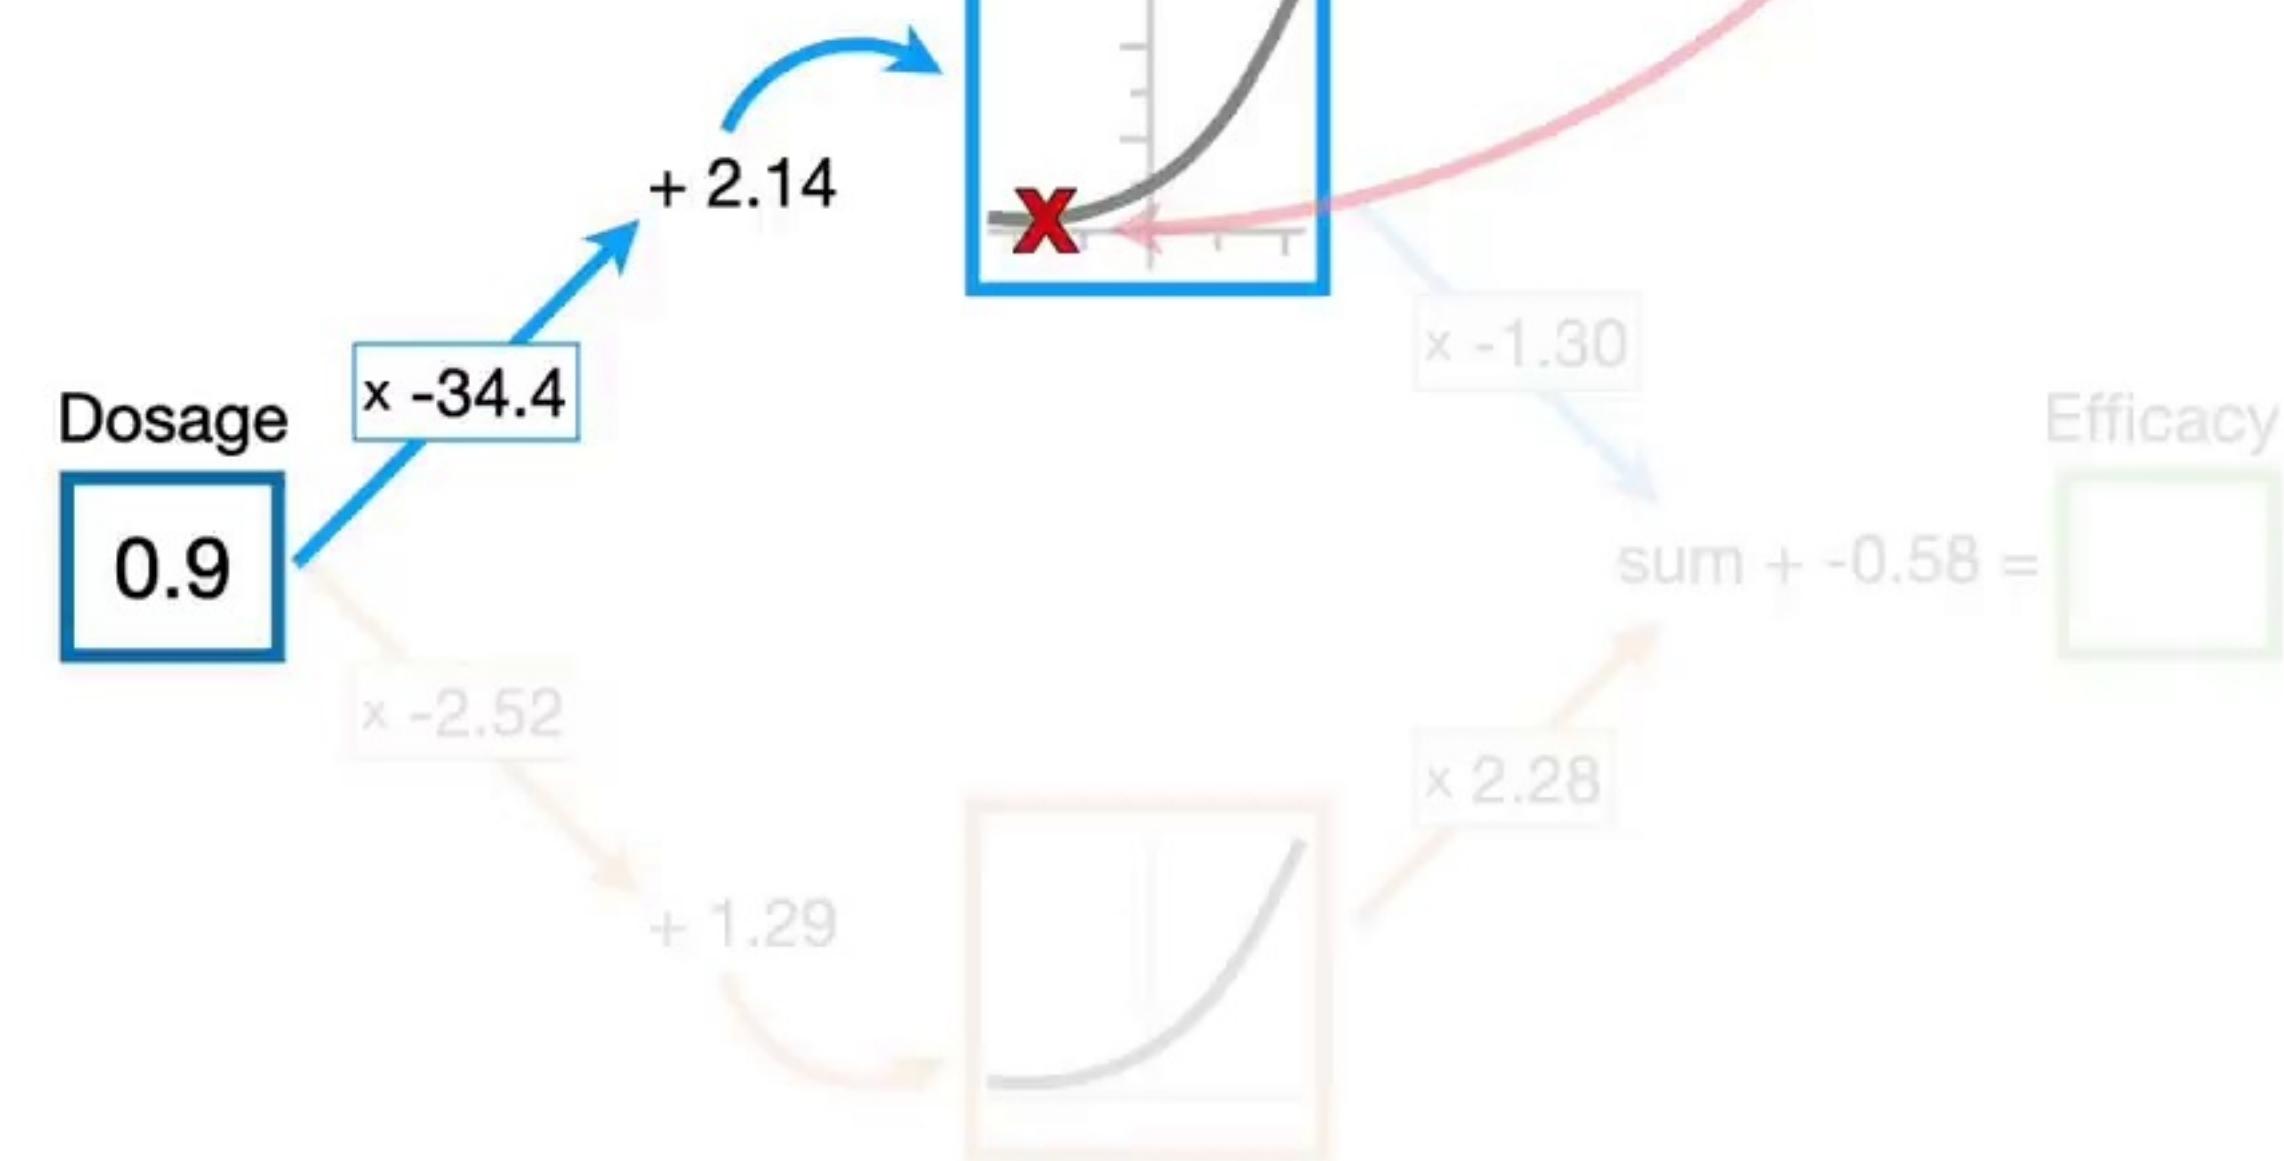


And if we continue to increase the **Dosage** values all the way to **1** (the maximum **Dosage**)...

$$(\text{Dosage} \times -34.4) + 2.14 = \text{x-axis coordinate}$$



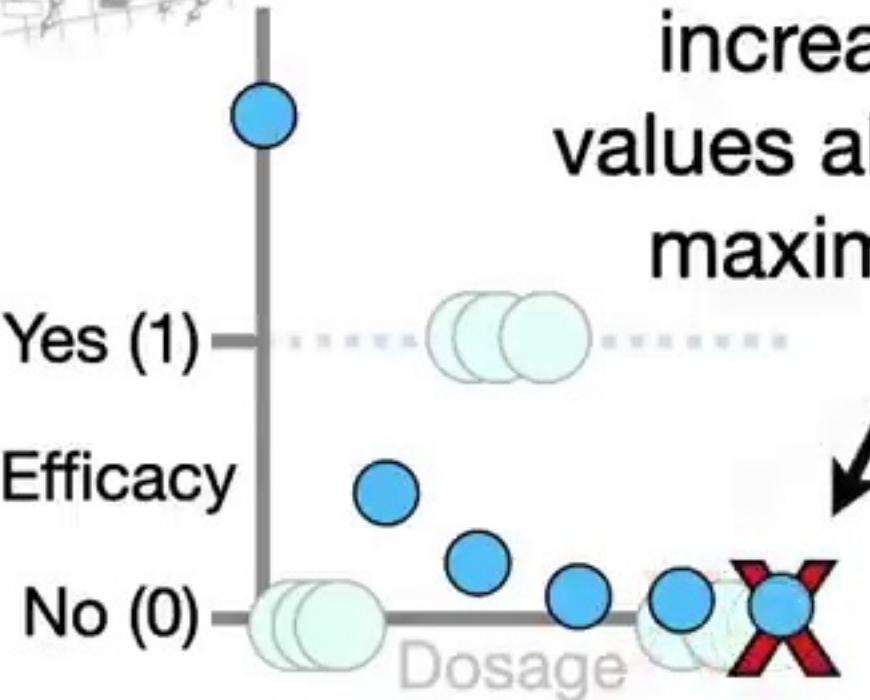
$$f(x) = \log(1 + e^x)$$





And if we continue to increase the **Dosage** values all the way to **1** (the maximum **Dosage**)...

$$(\text{Dosage} \times -34.4) + 2.14 = \text{x-axis coordinate}$$



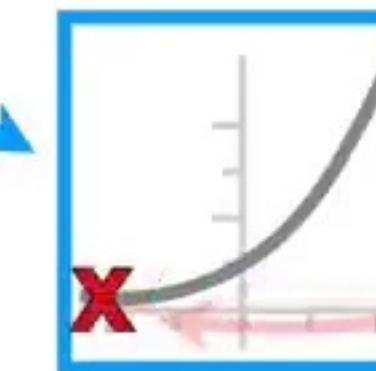
Dosage

$x - 34.4$

$x - 2.52$

$+ 1.29$

$+ 2.14$



$x - 1.30$

sum  $+ -0.58 =$

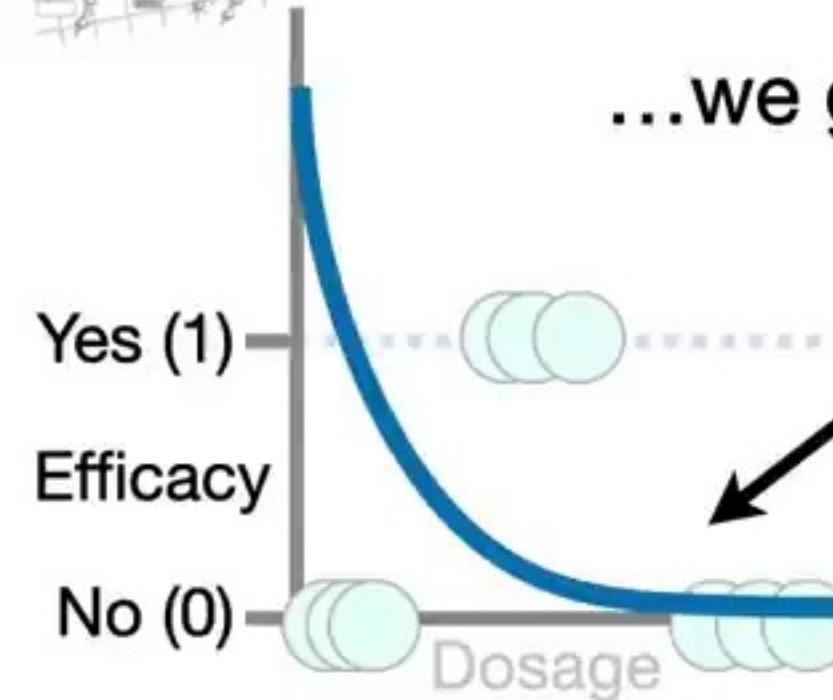
Efficacy



$x 2.28$

double  
BAM!!  
**SQ!**

...we get this **blue curve**.



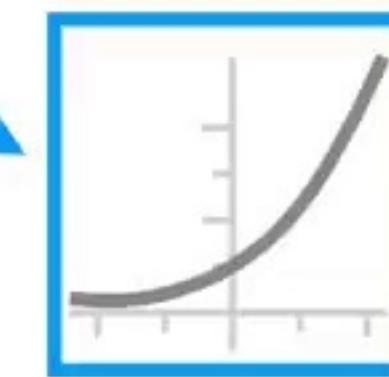
Dosage

$\times -34.4$

$\times -2.52$

$+ 1.29$

$+ 2.14$



$\times -1.30$

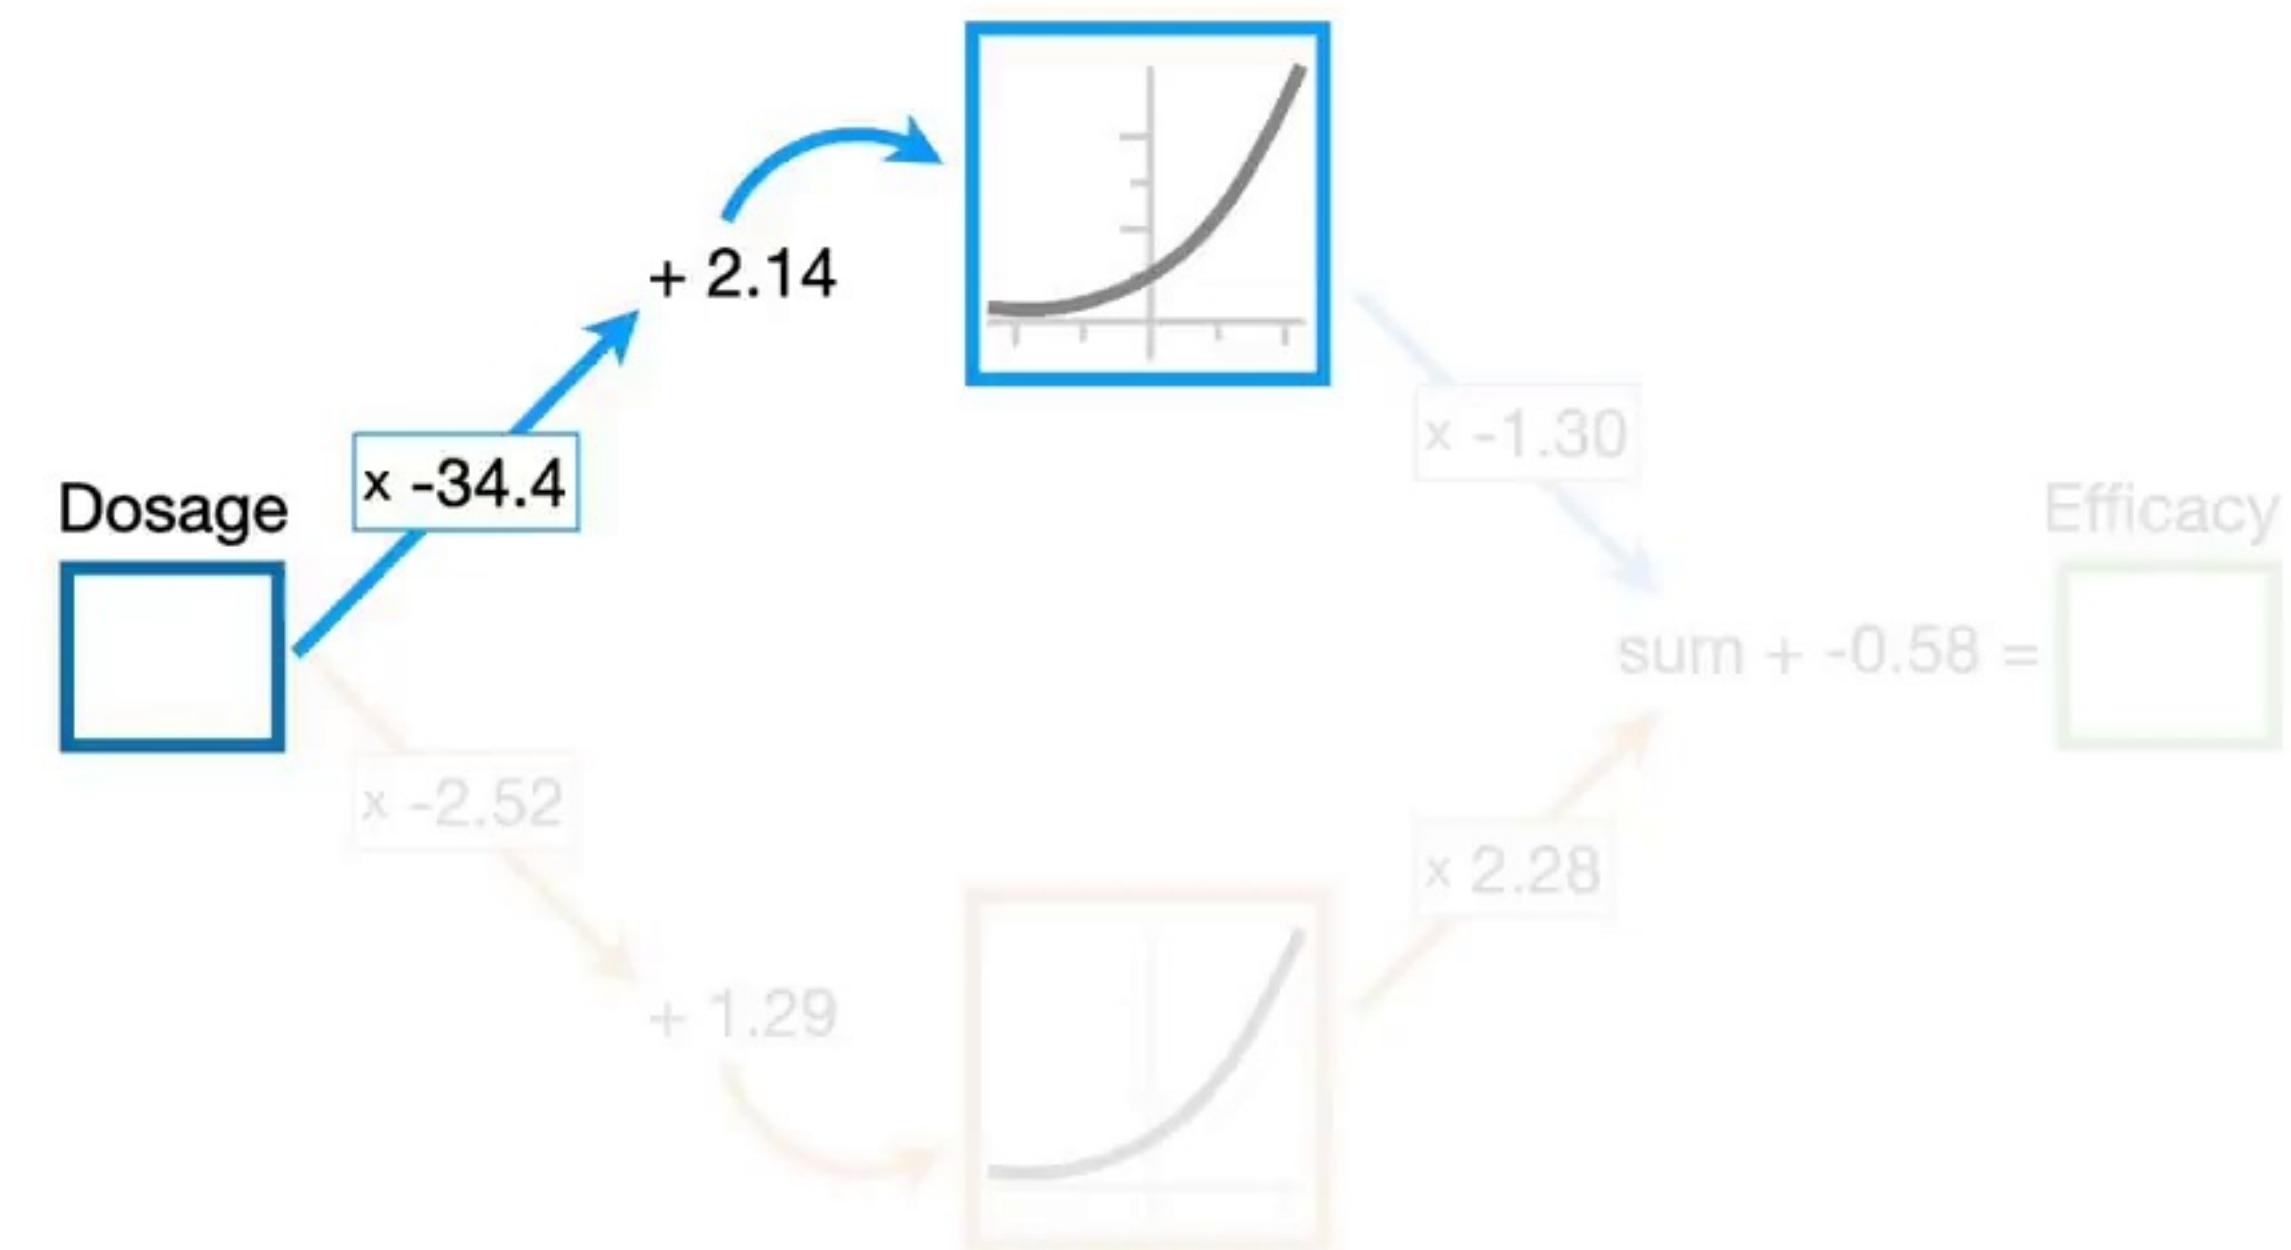
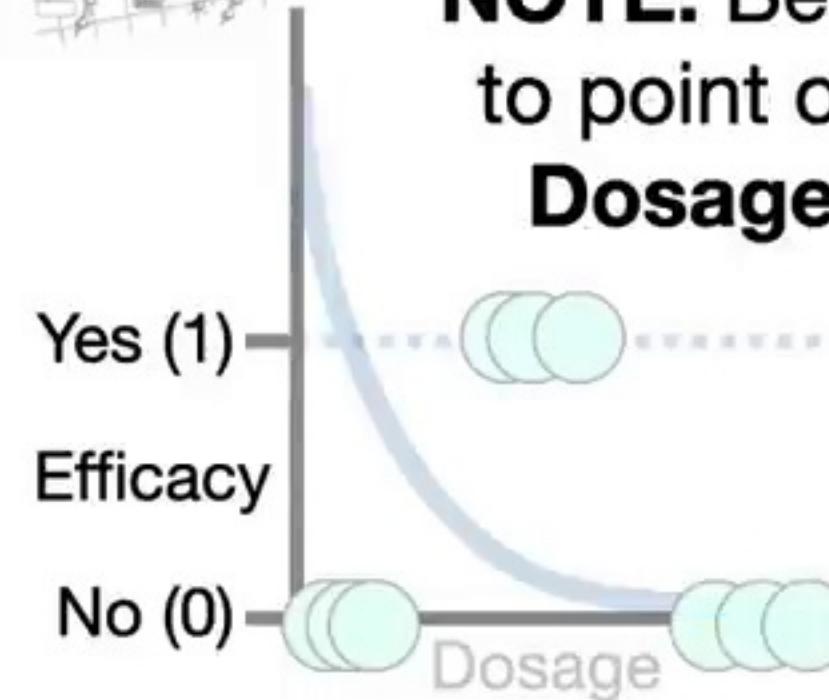
sum  $+ -0.58 =$



$\times 2.28$

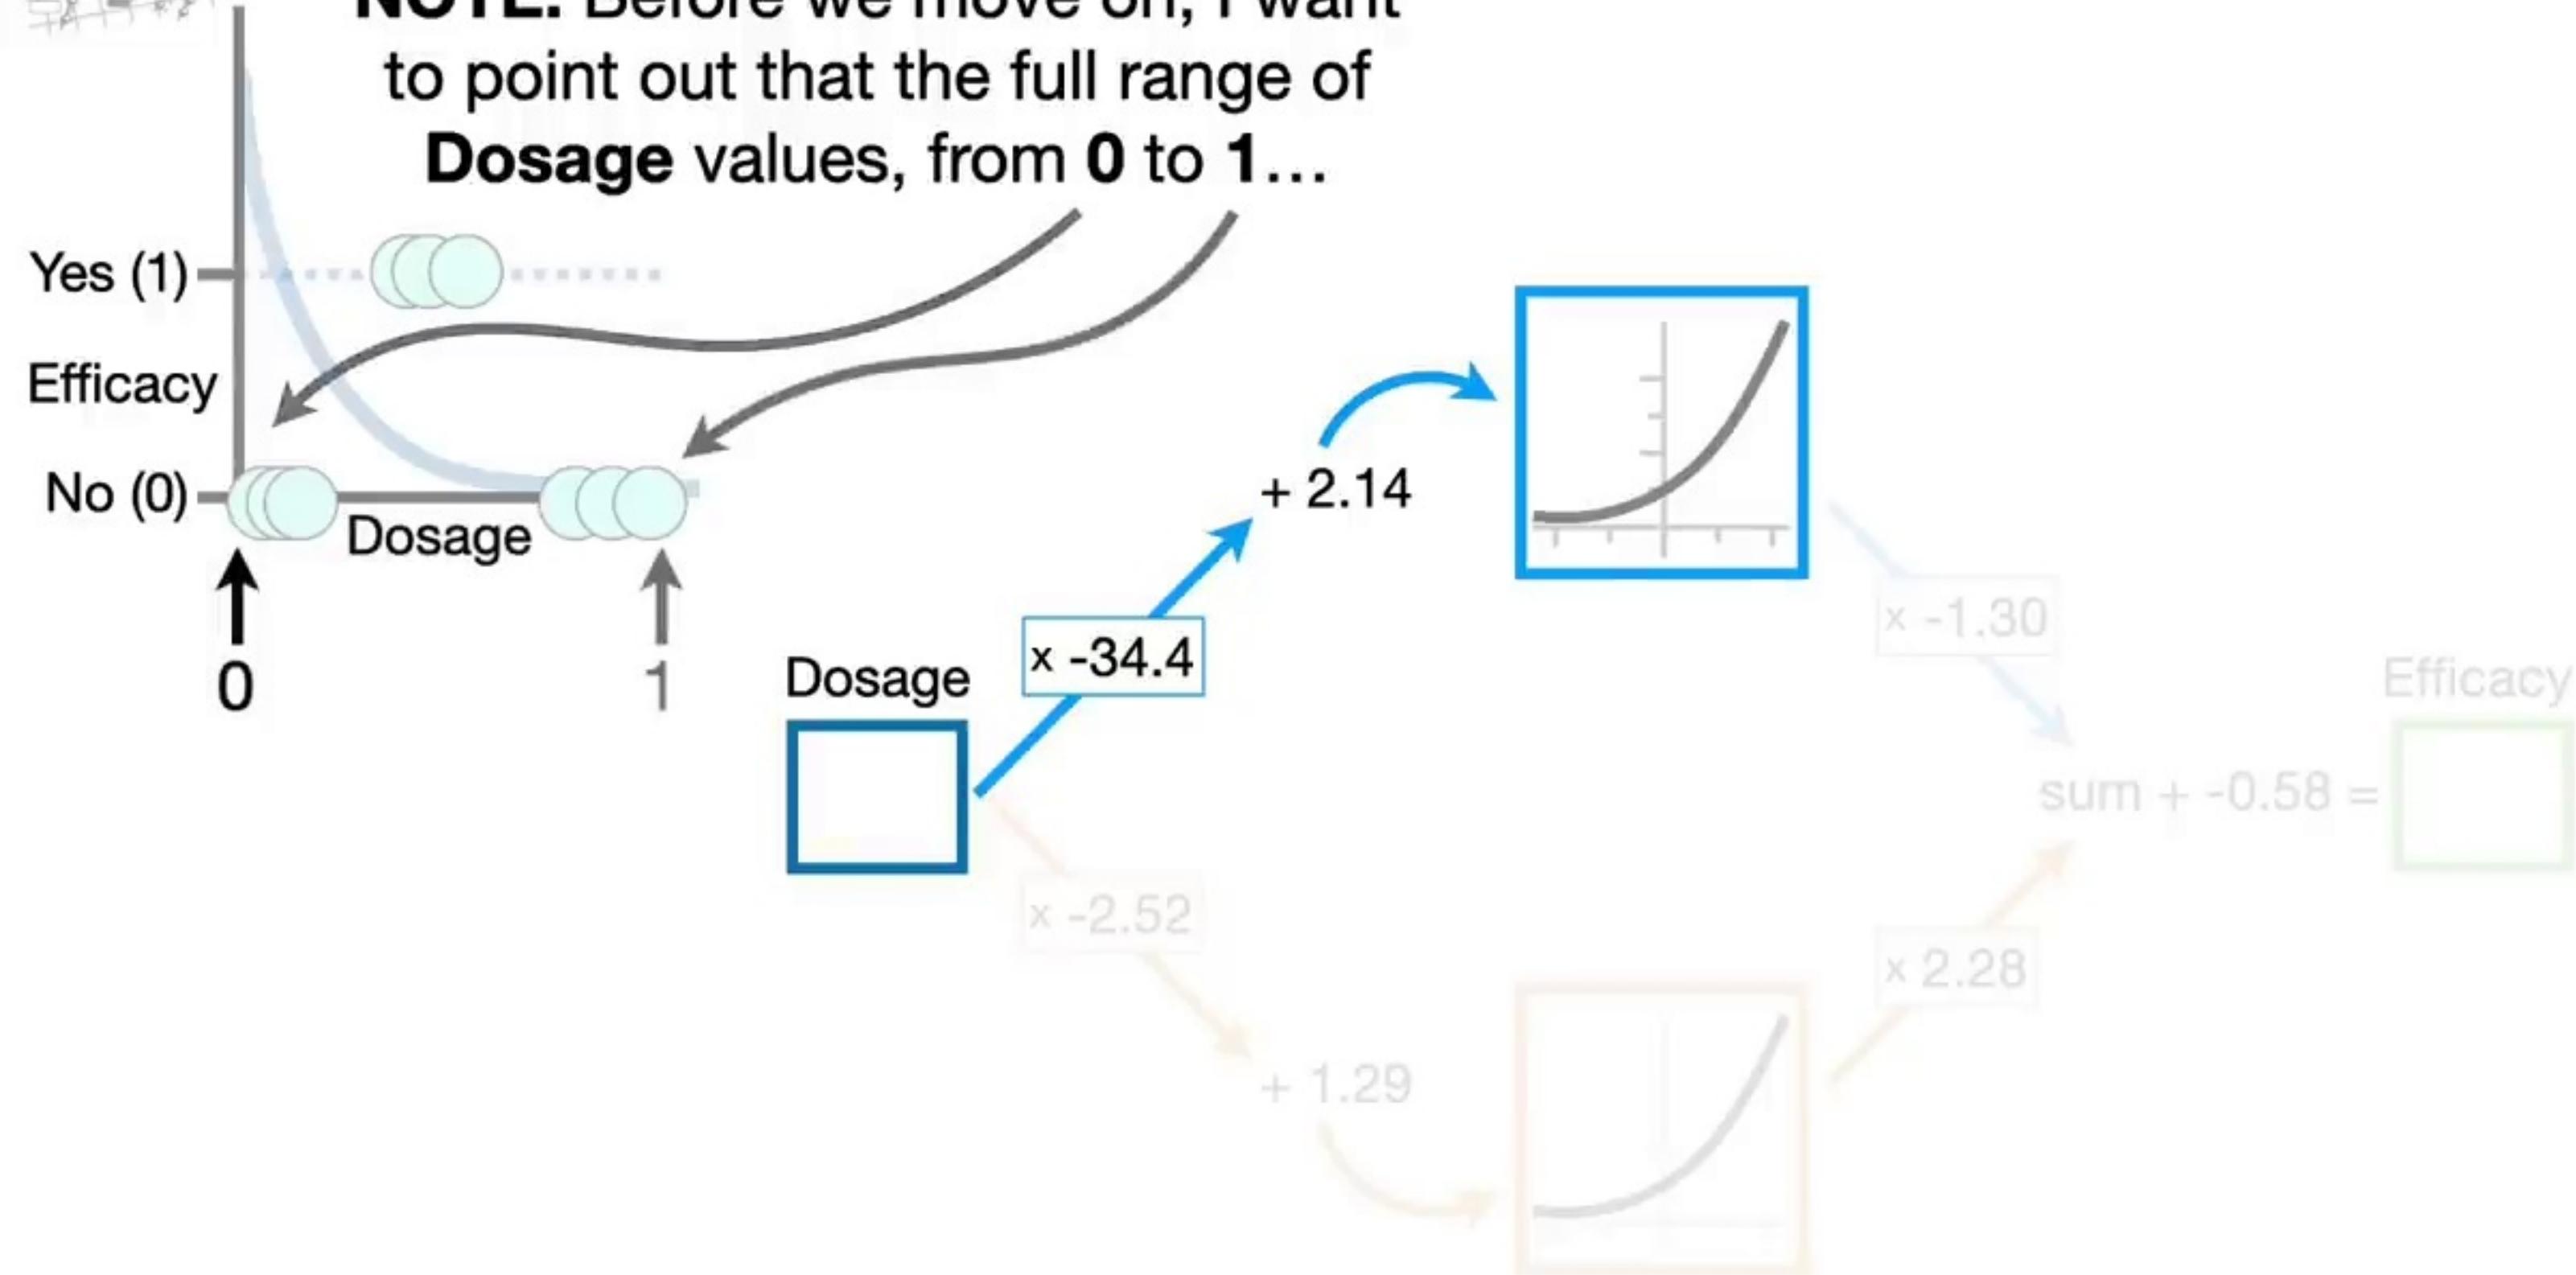


**NOTE:** Before we move on, I want to point out that the full range of **Dosage** values, from **0** to **1**...



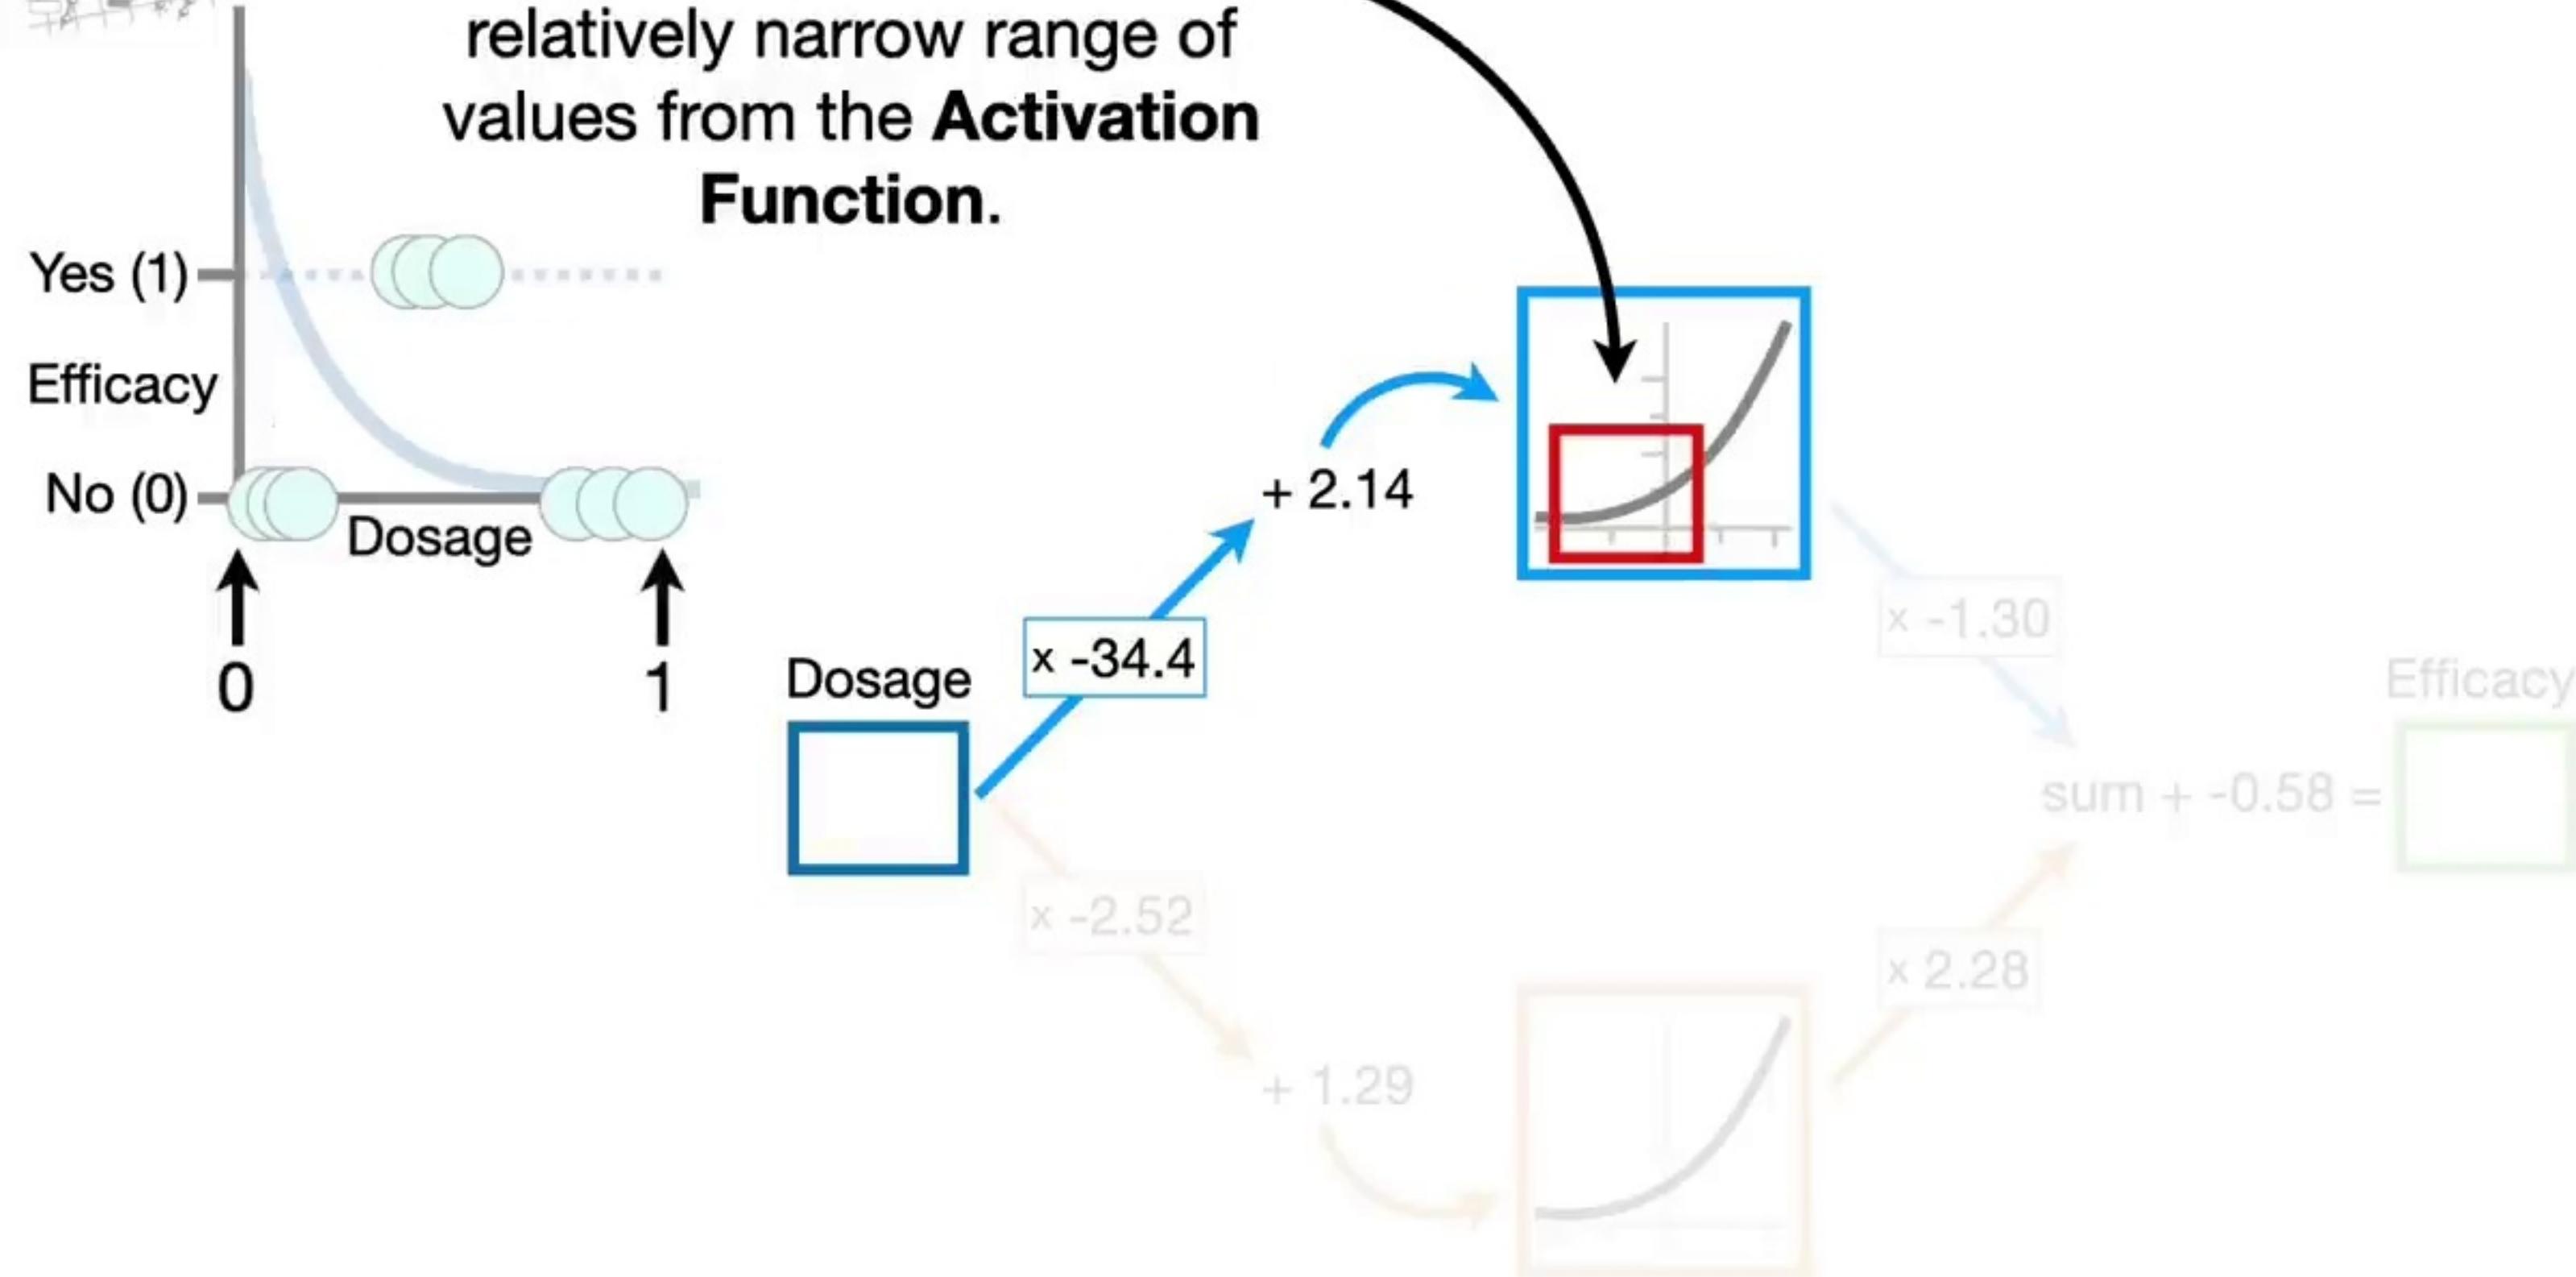


**NOTE:** Before we move on, I want to point out that the full range of **Dosage** values, from **0** to **1**...



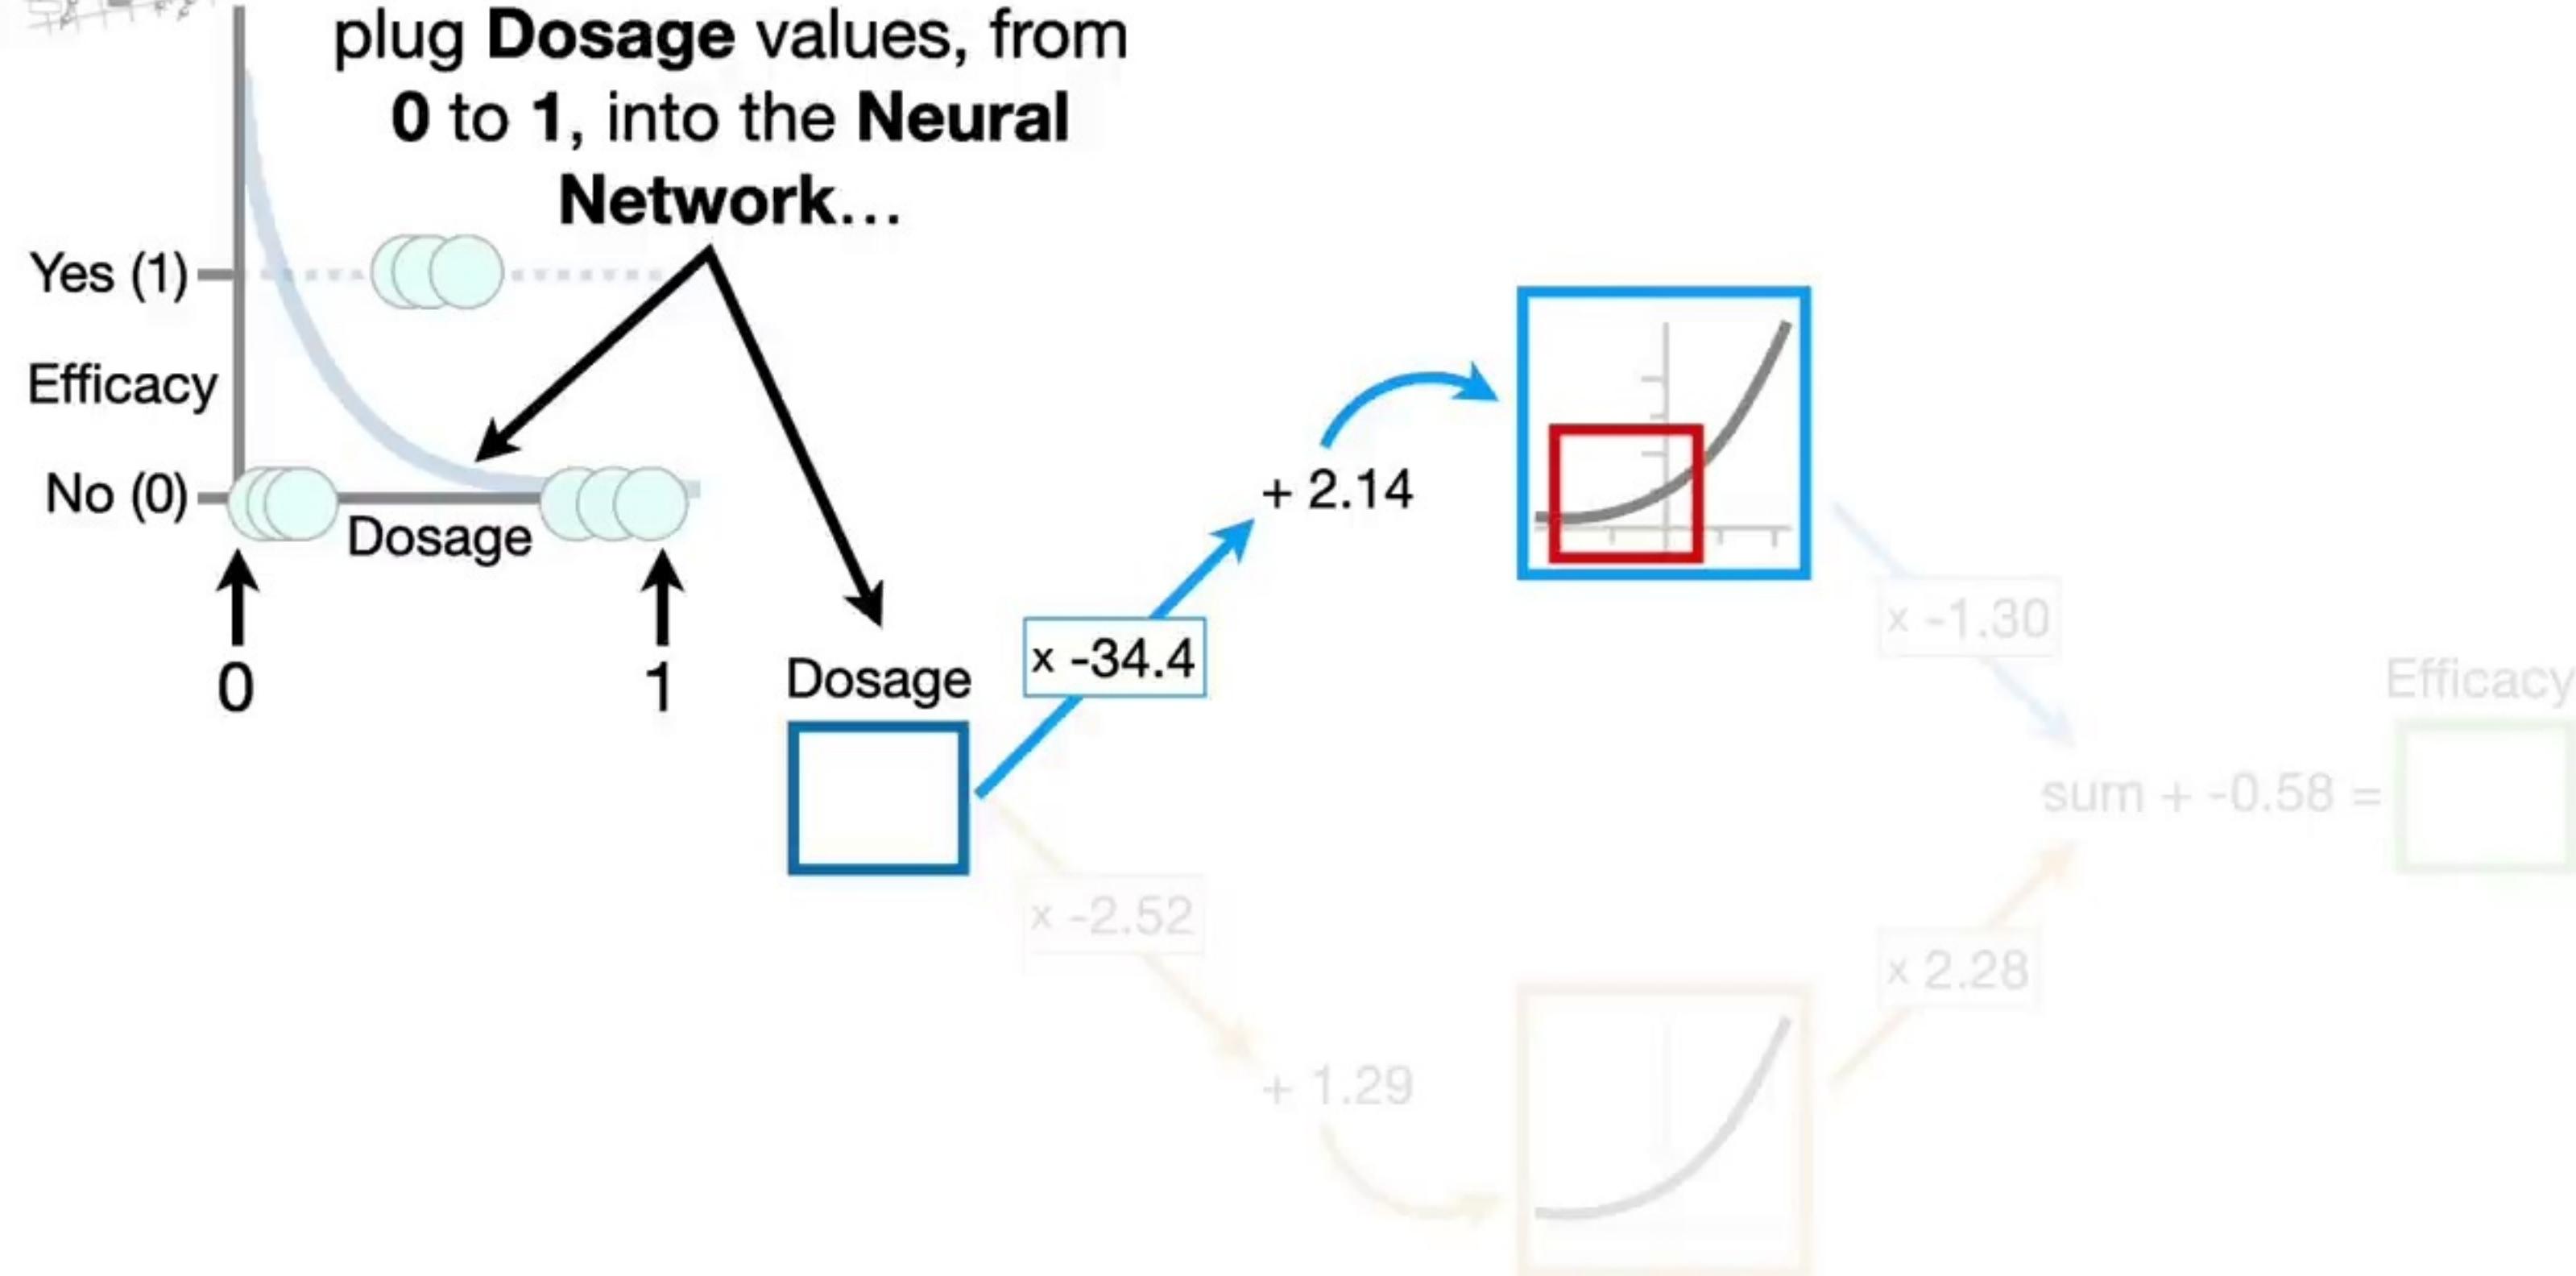


...corresponds to this relatively narrow range of values from the **Activation Function**.





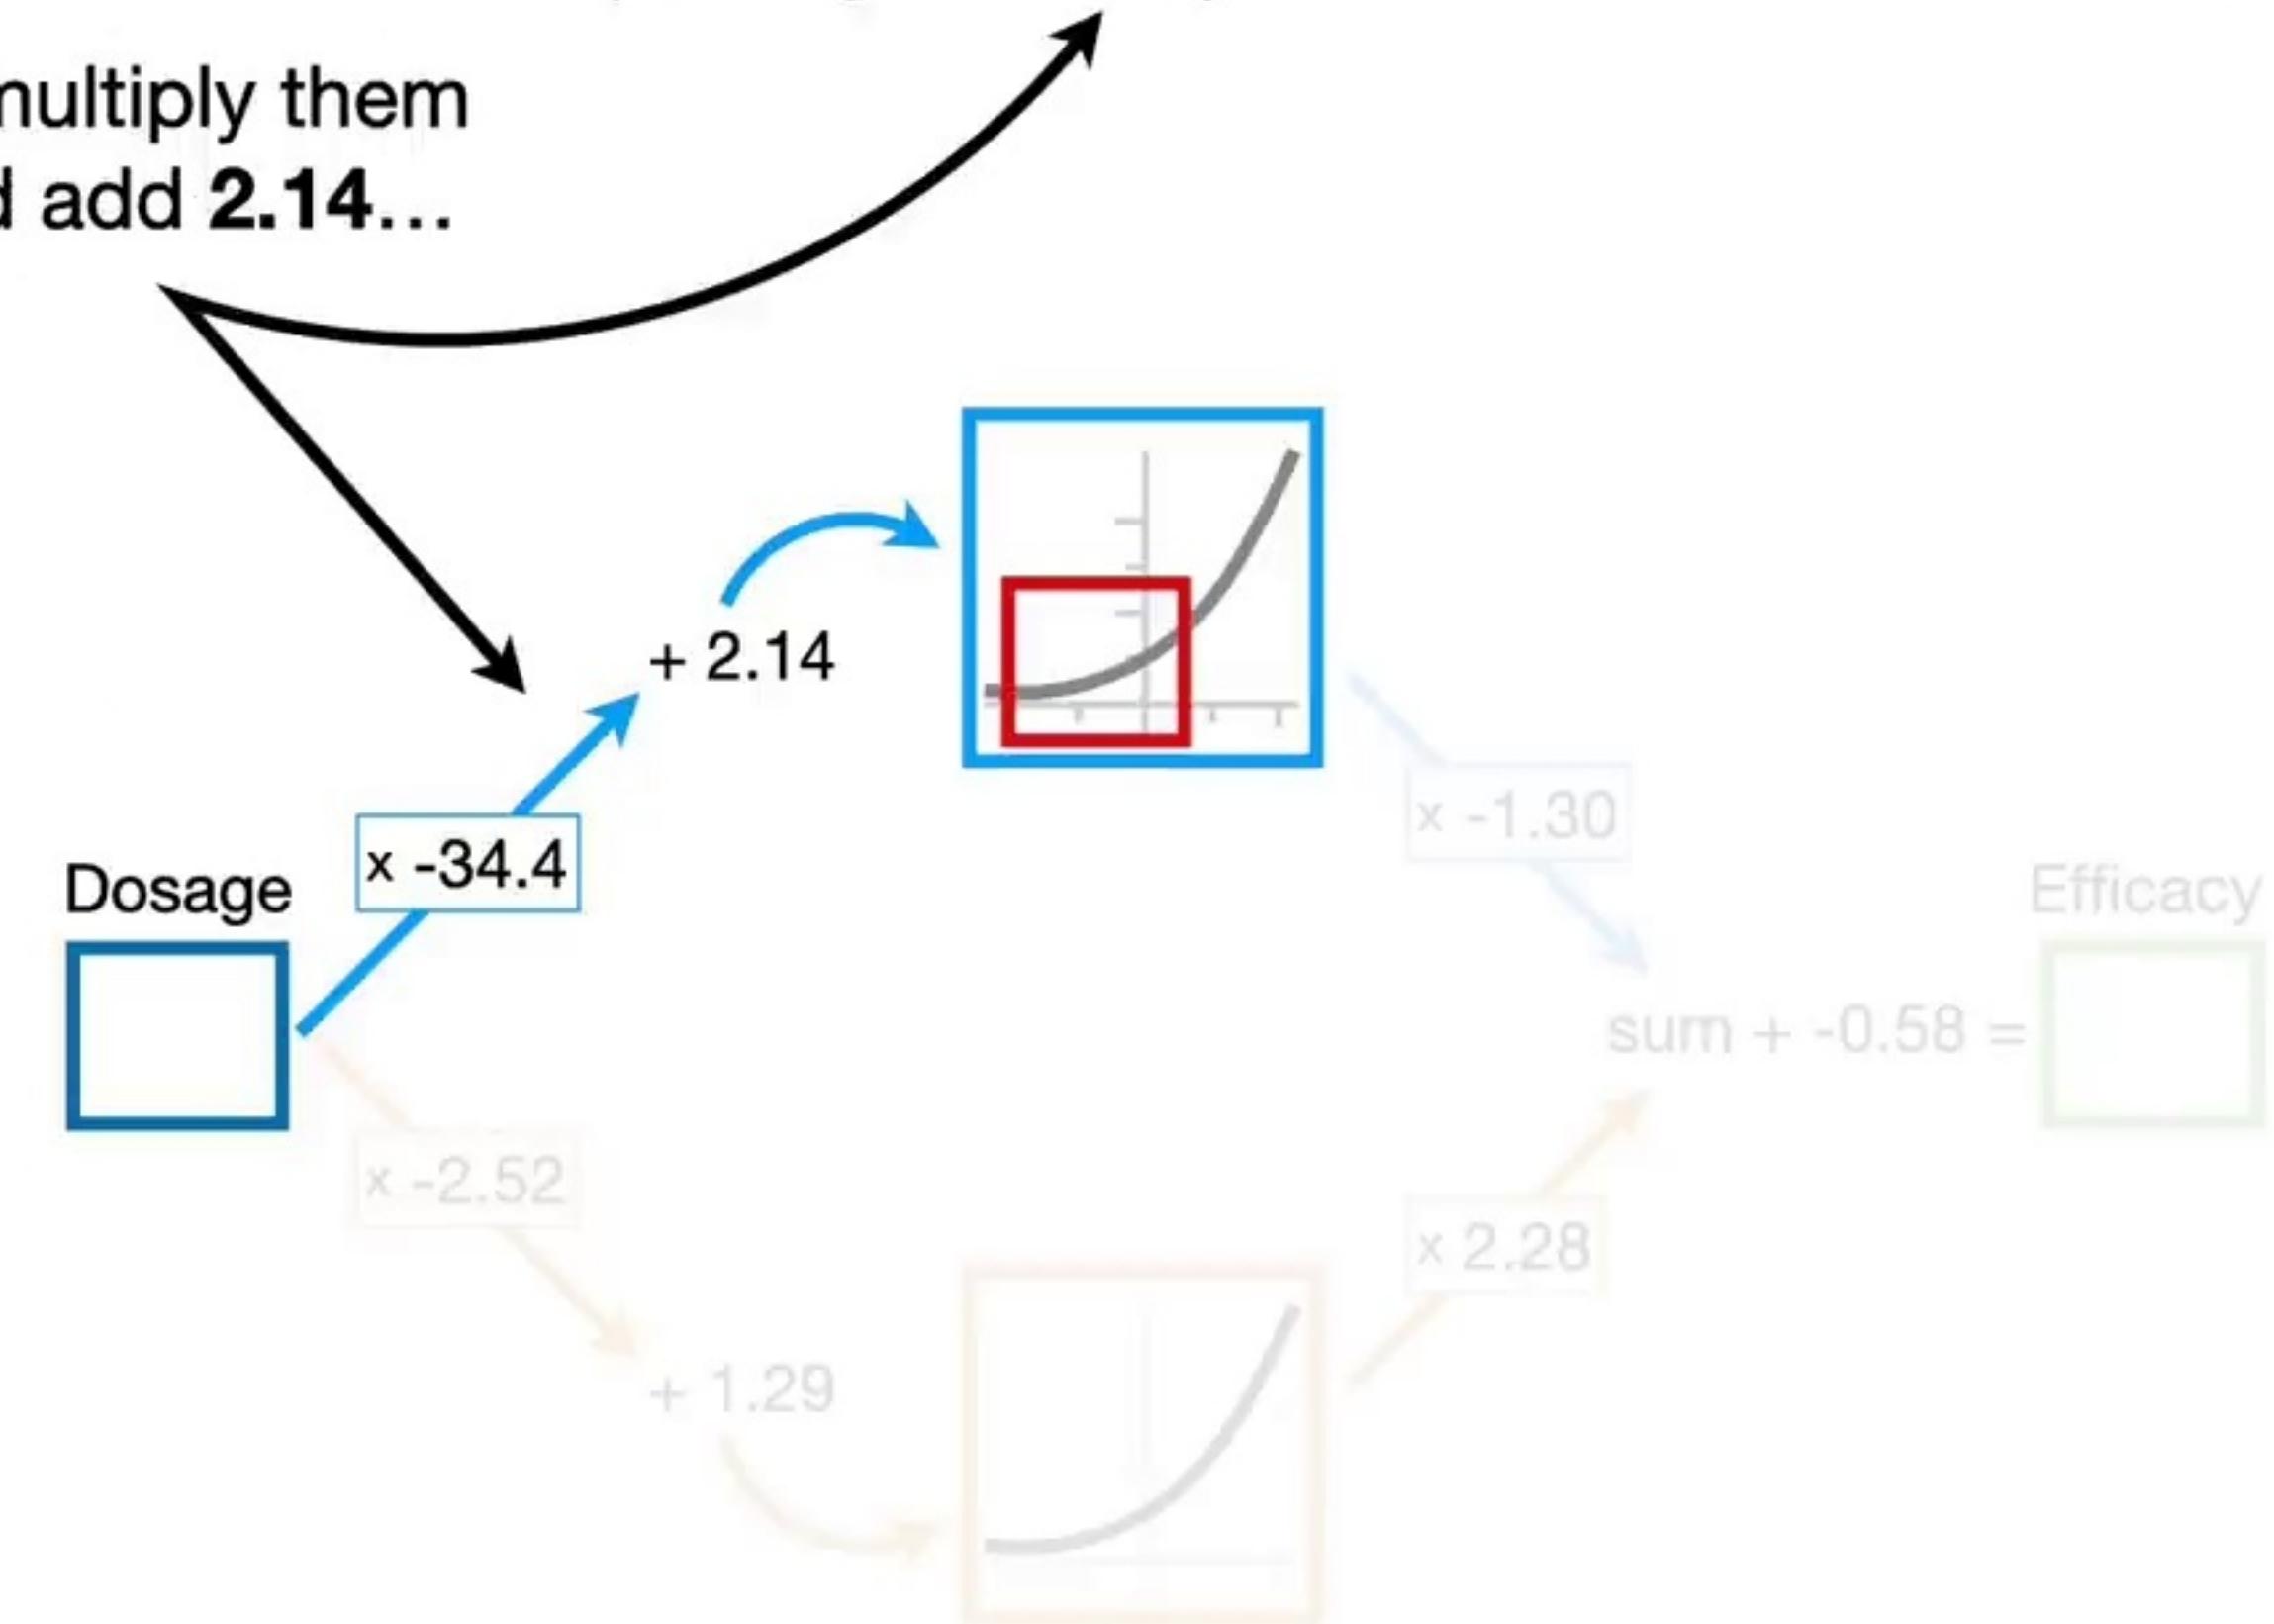
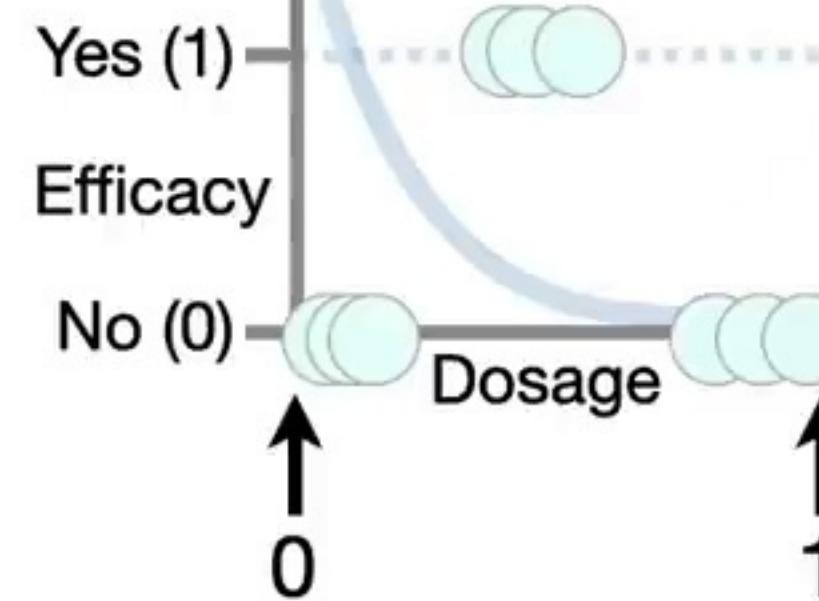
In other words, when we plug **Dosage** values, from 0 to 1, into the **Neural Network**...



Double  
BAM!!  
**SO!**

**(Dosage  $\times$  -34.4) + 2.14 = x-axis coordinate**

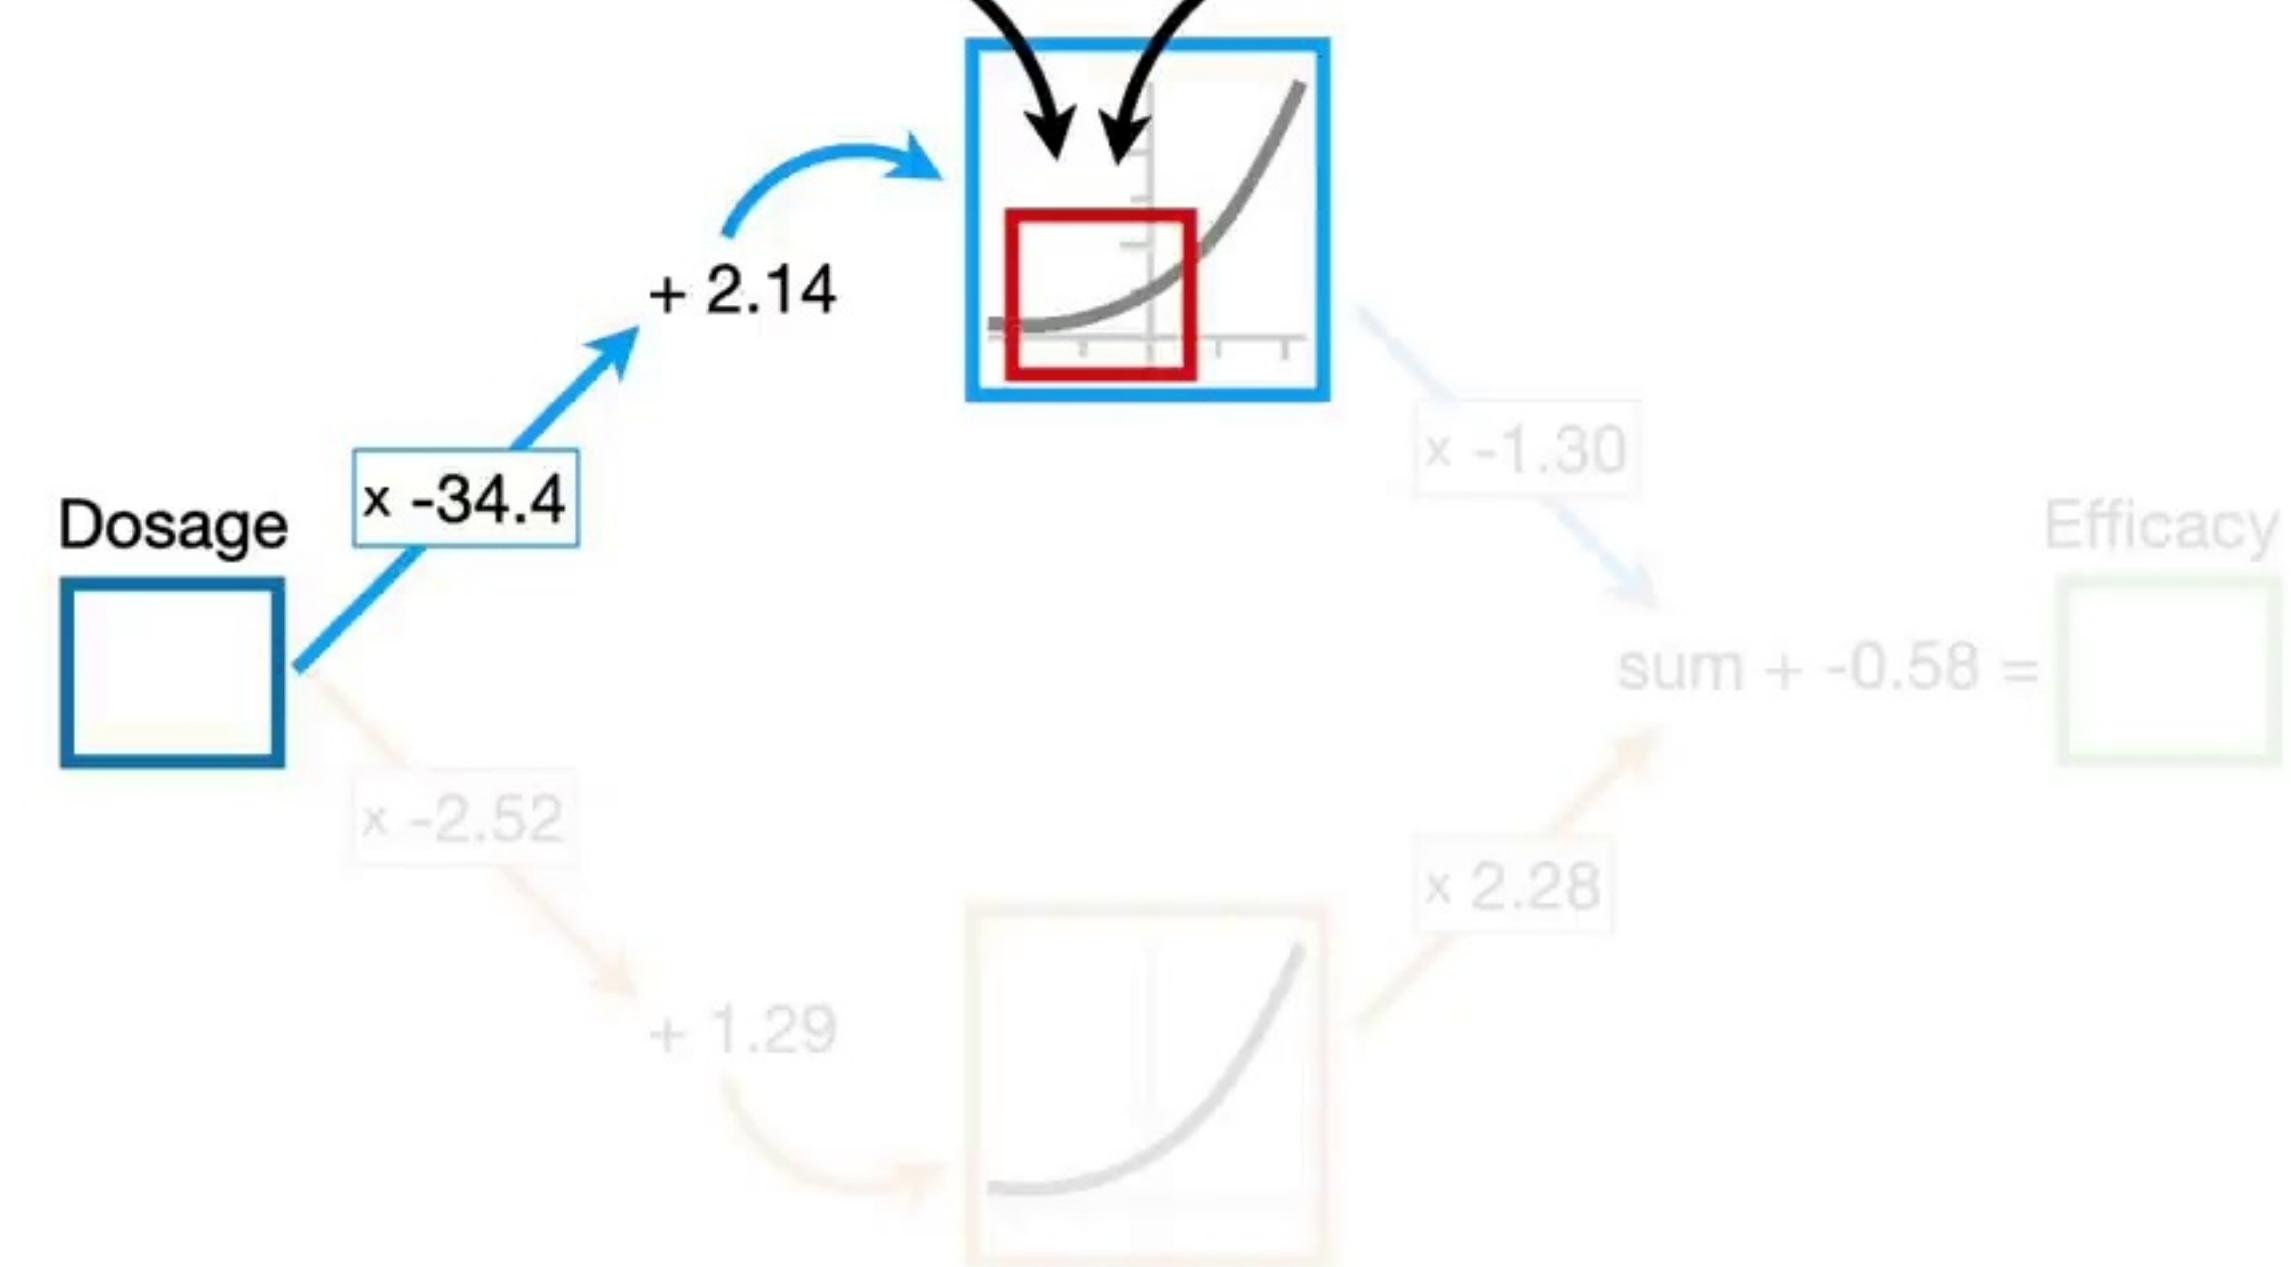
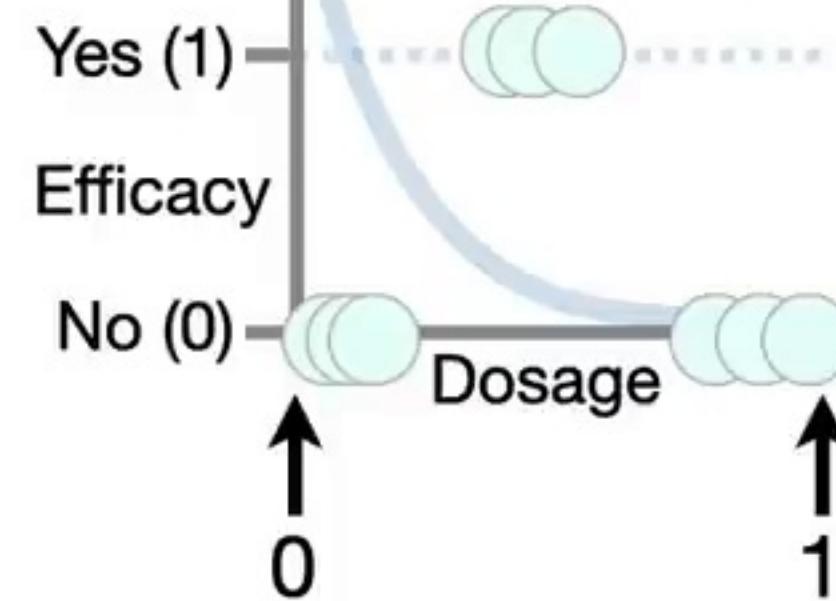
...and then multiply them  
by **-34.4** and add **2.14**...





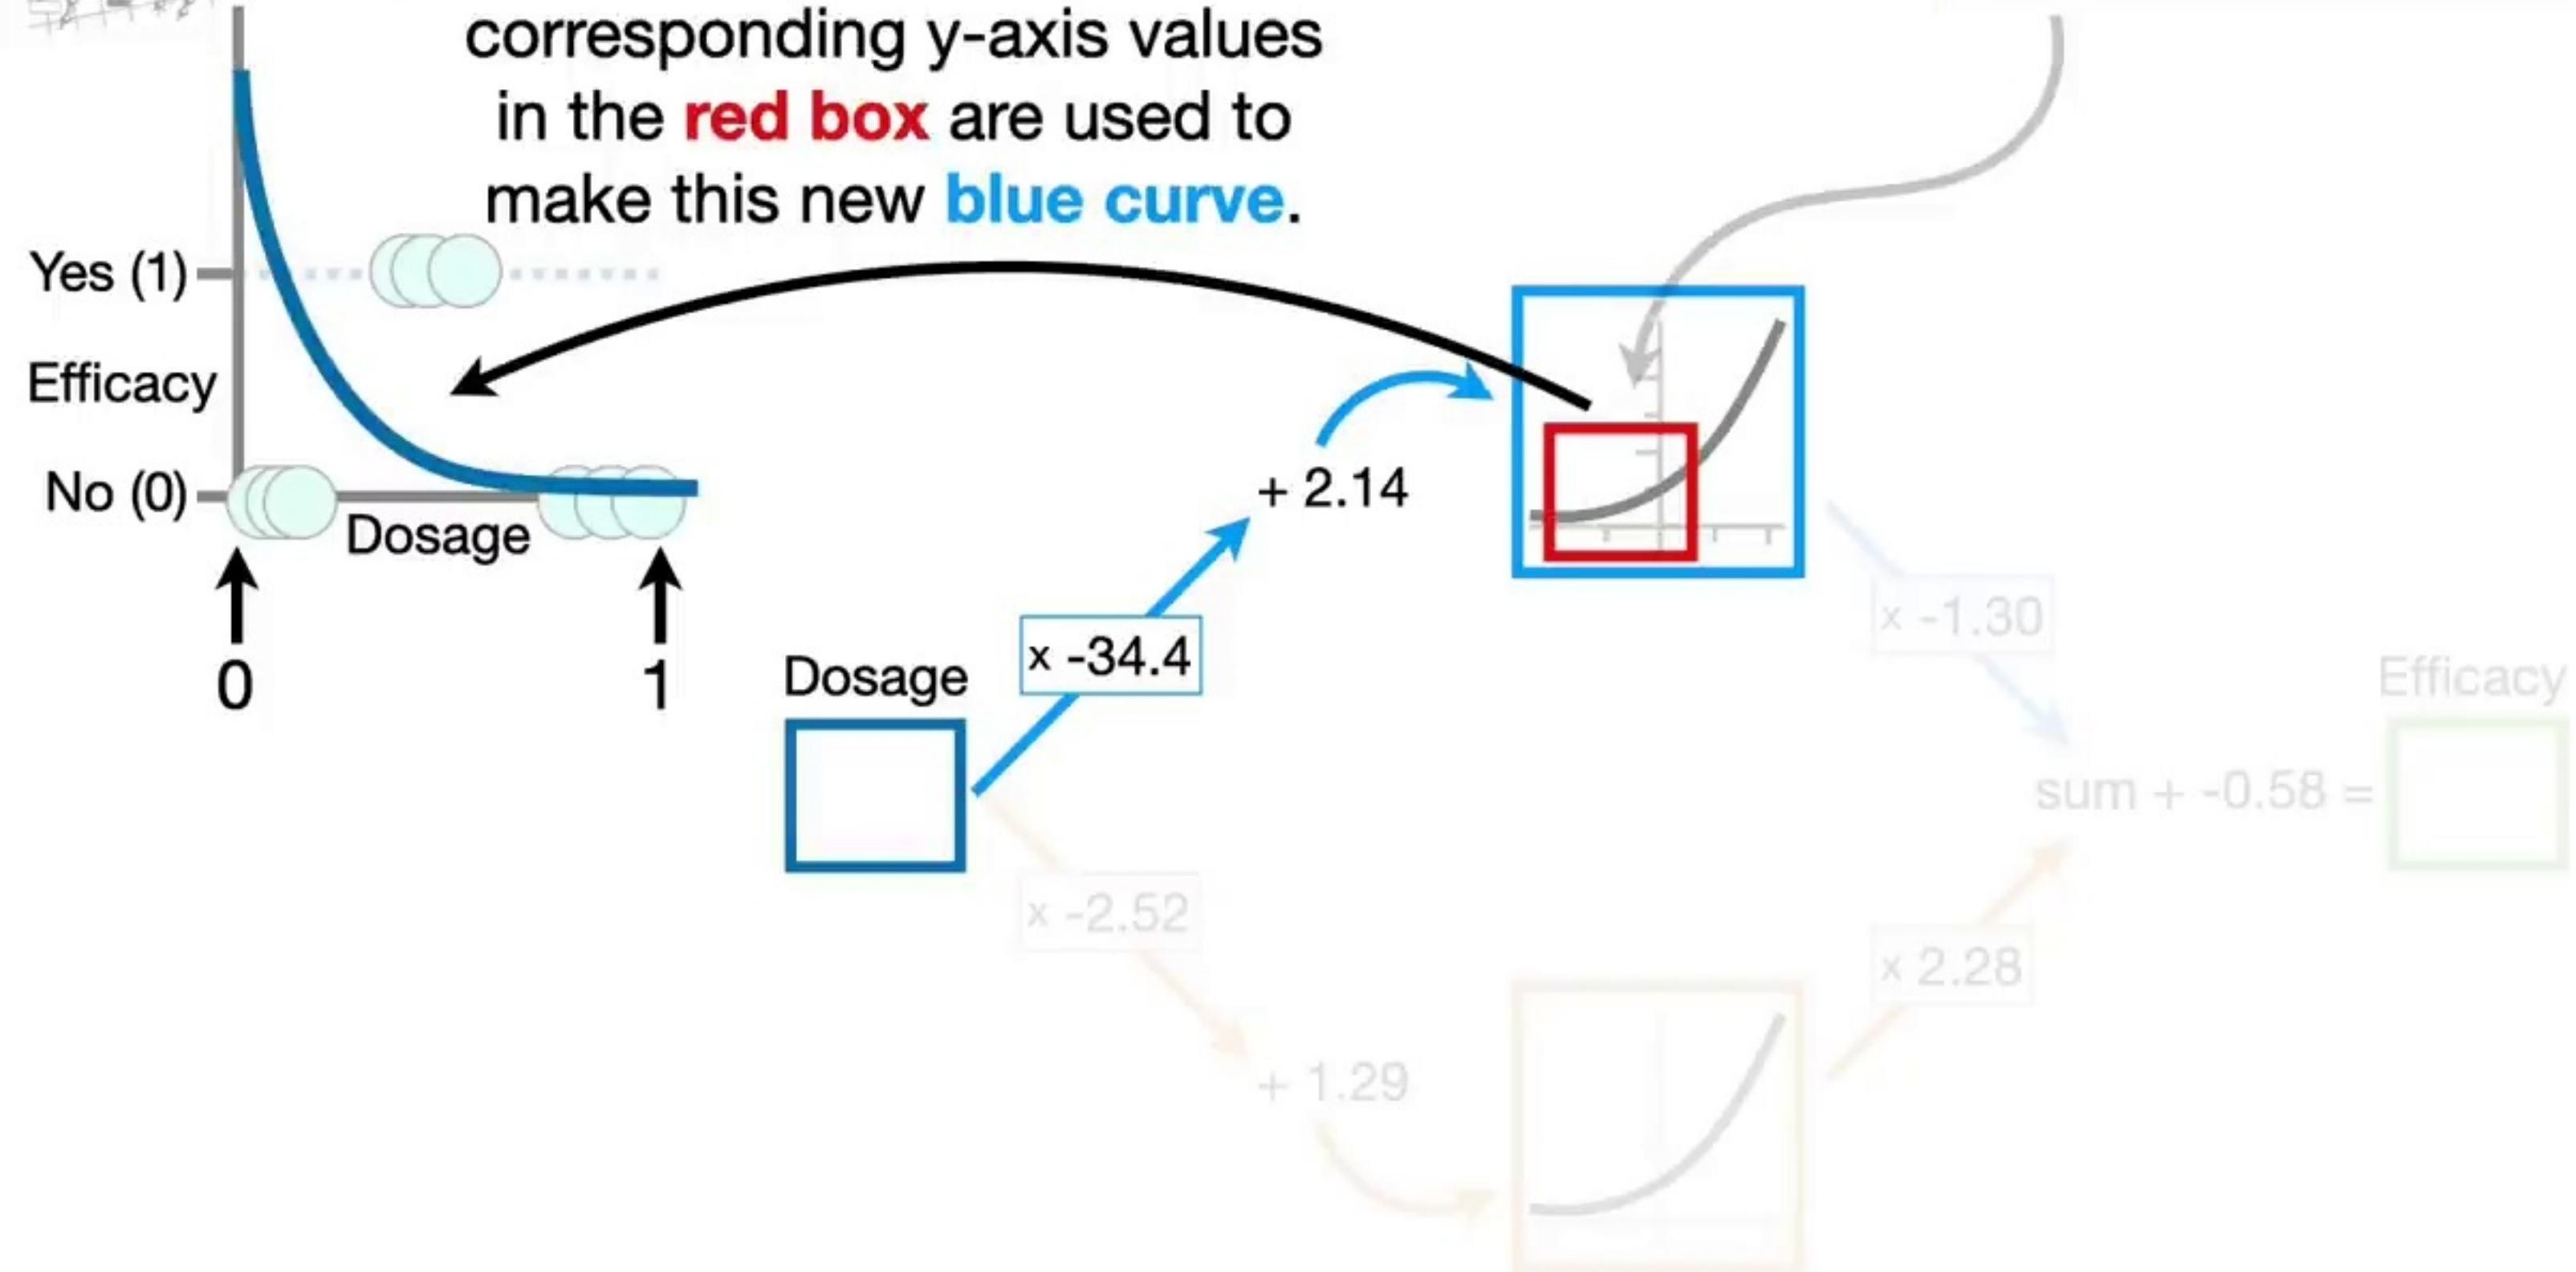
$$(\text{Dosage} \times -34.4) + 2.14 = \text{x-axis coordinate}$$

...then we only get x-axis coordinates that are within the **red box**.



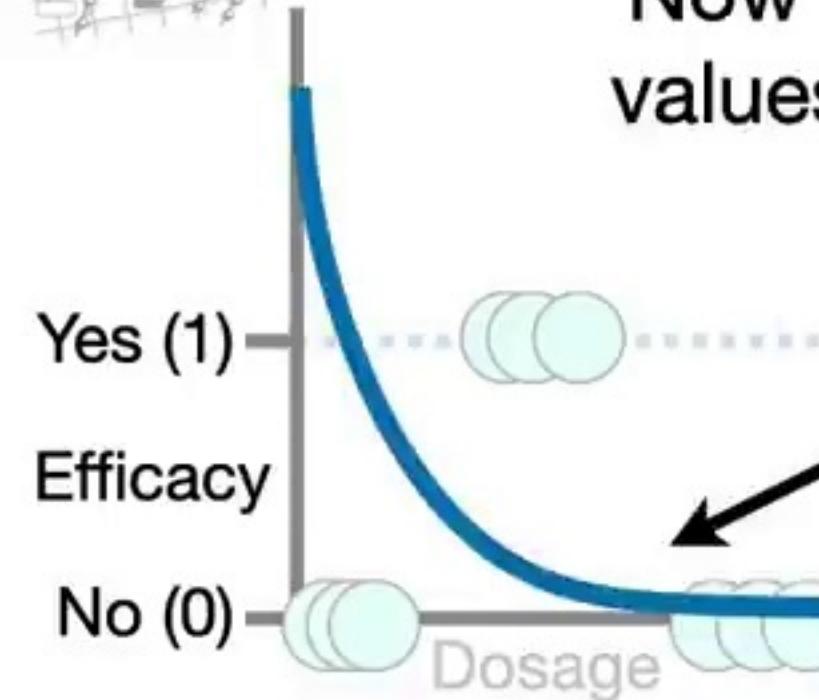
Double  
BAM!!  
**SQ!**

...and thus, only the corresponding y-axis values in the **red box** are used to make this new **blue curve**.





Now we scale the y-axis values for the **blue curve** by **-1.30**.



by -1.30.

Dosage

$\times -34.4$

$\times -2.52$

$+ 1.29$

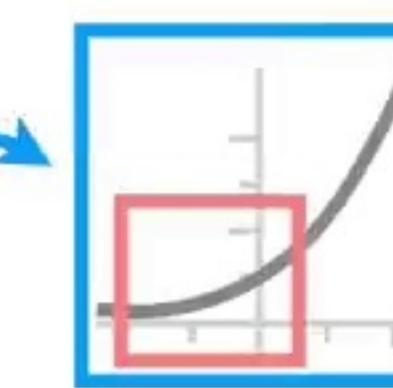
$+ 2.14$

$\times -1.30$

sum + -0.58 =

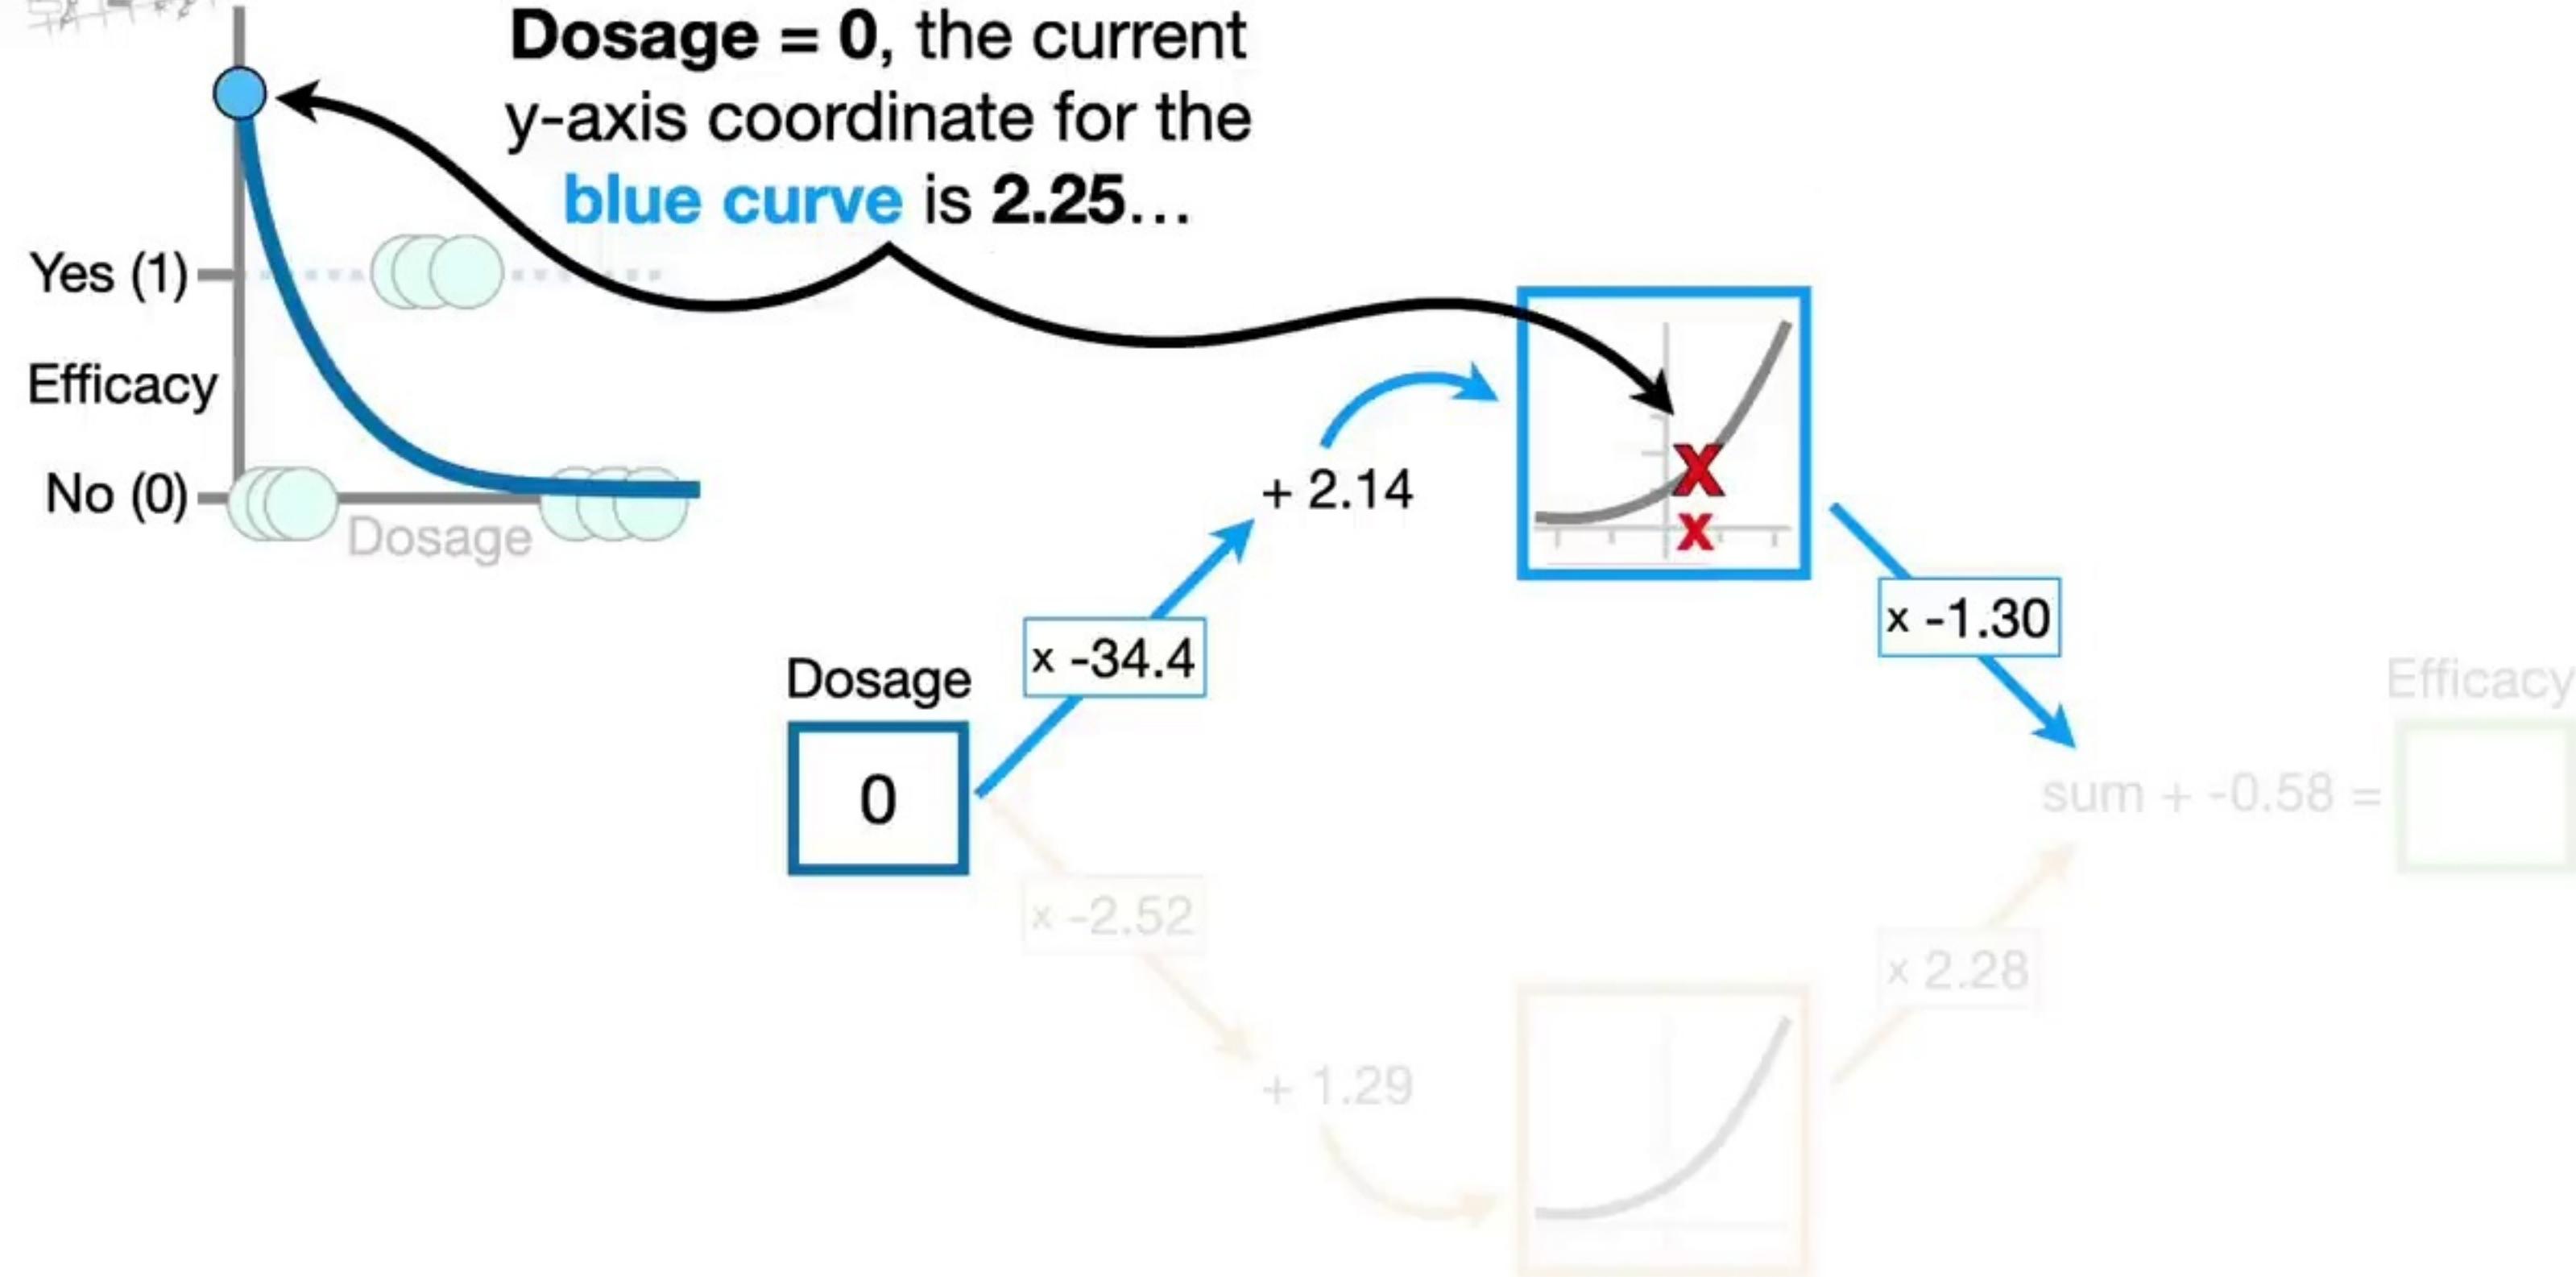
Efficacy

$\times 2.28$





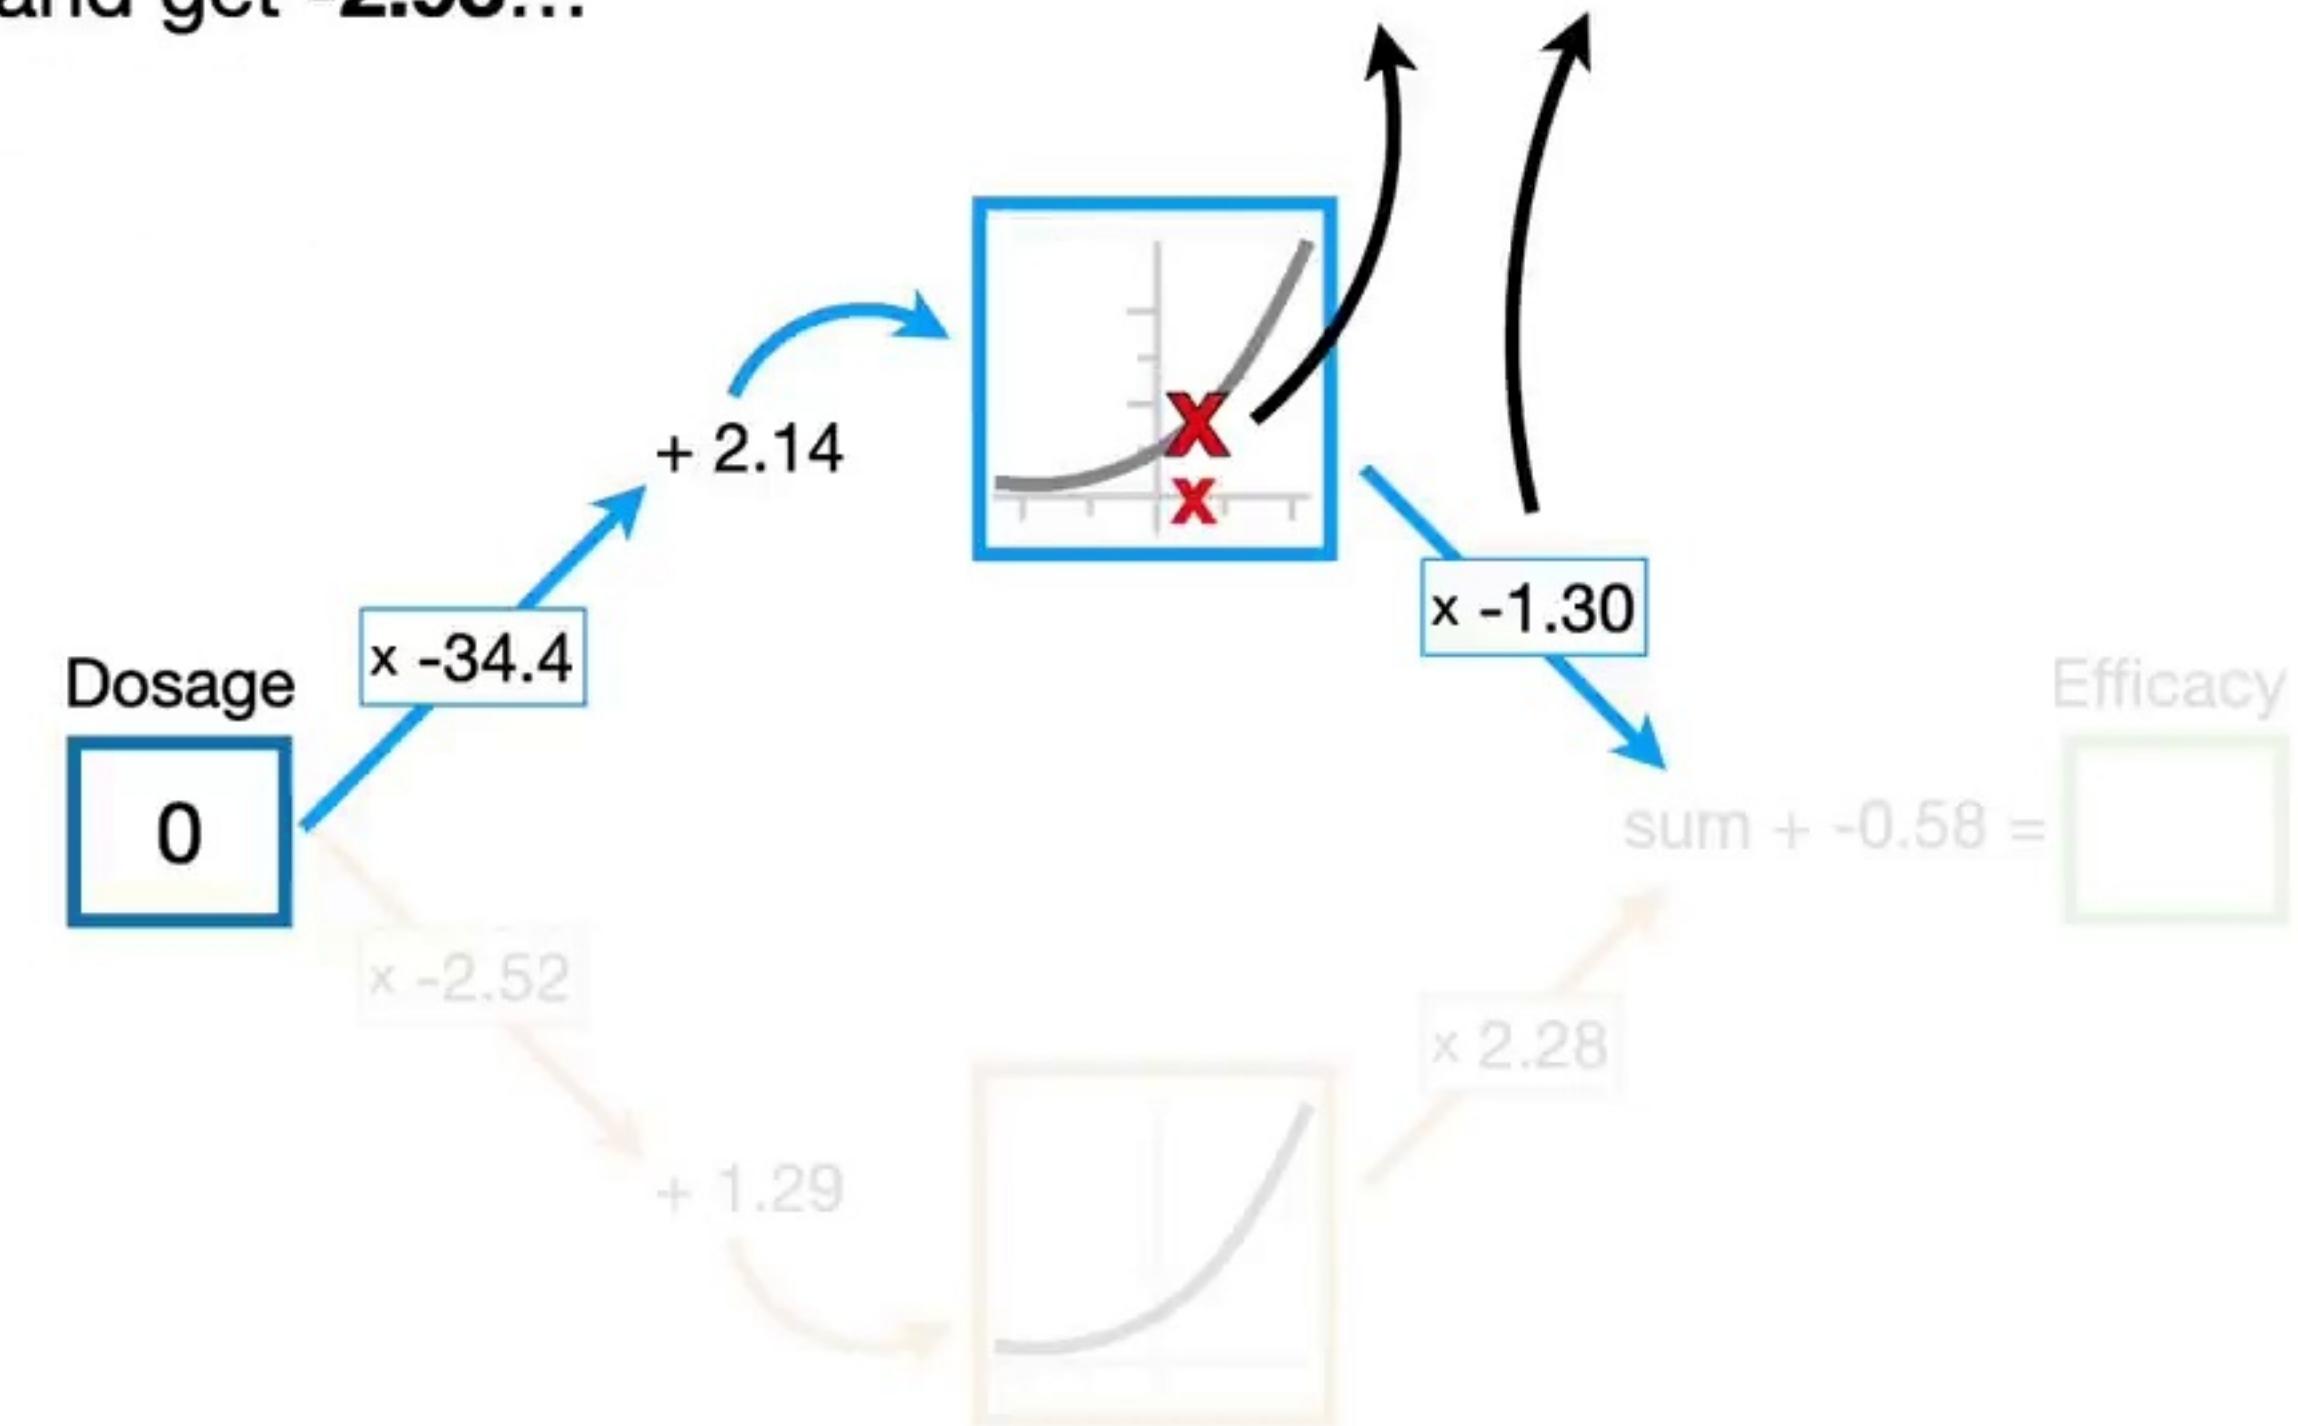
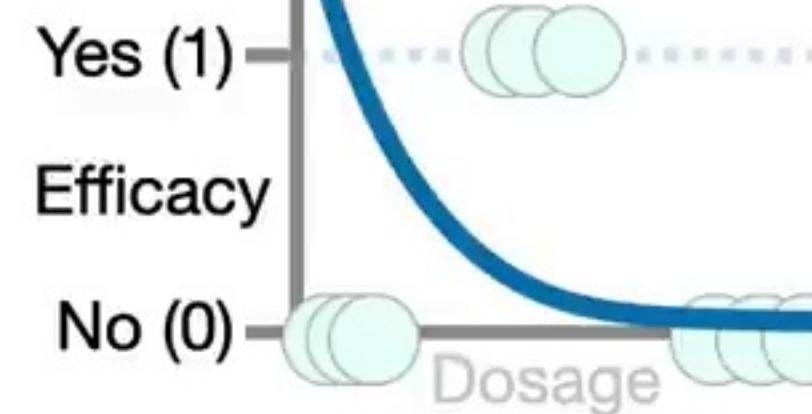
For example, when  
**Dosage = 0**, the current  
y-axis coordinate for the  
**blue curve** is **2.25...**



Double  
BAM!!!  
**SO!**

...so we multiply **2.25** by  
**-1.30** and get **-2.93**...

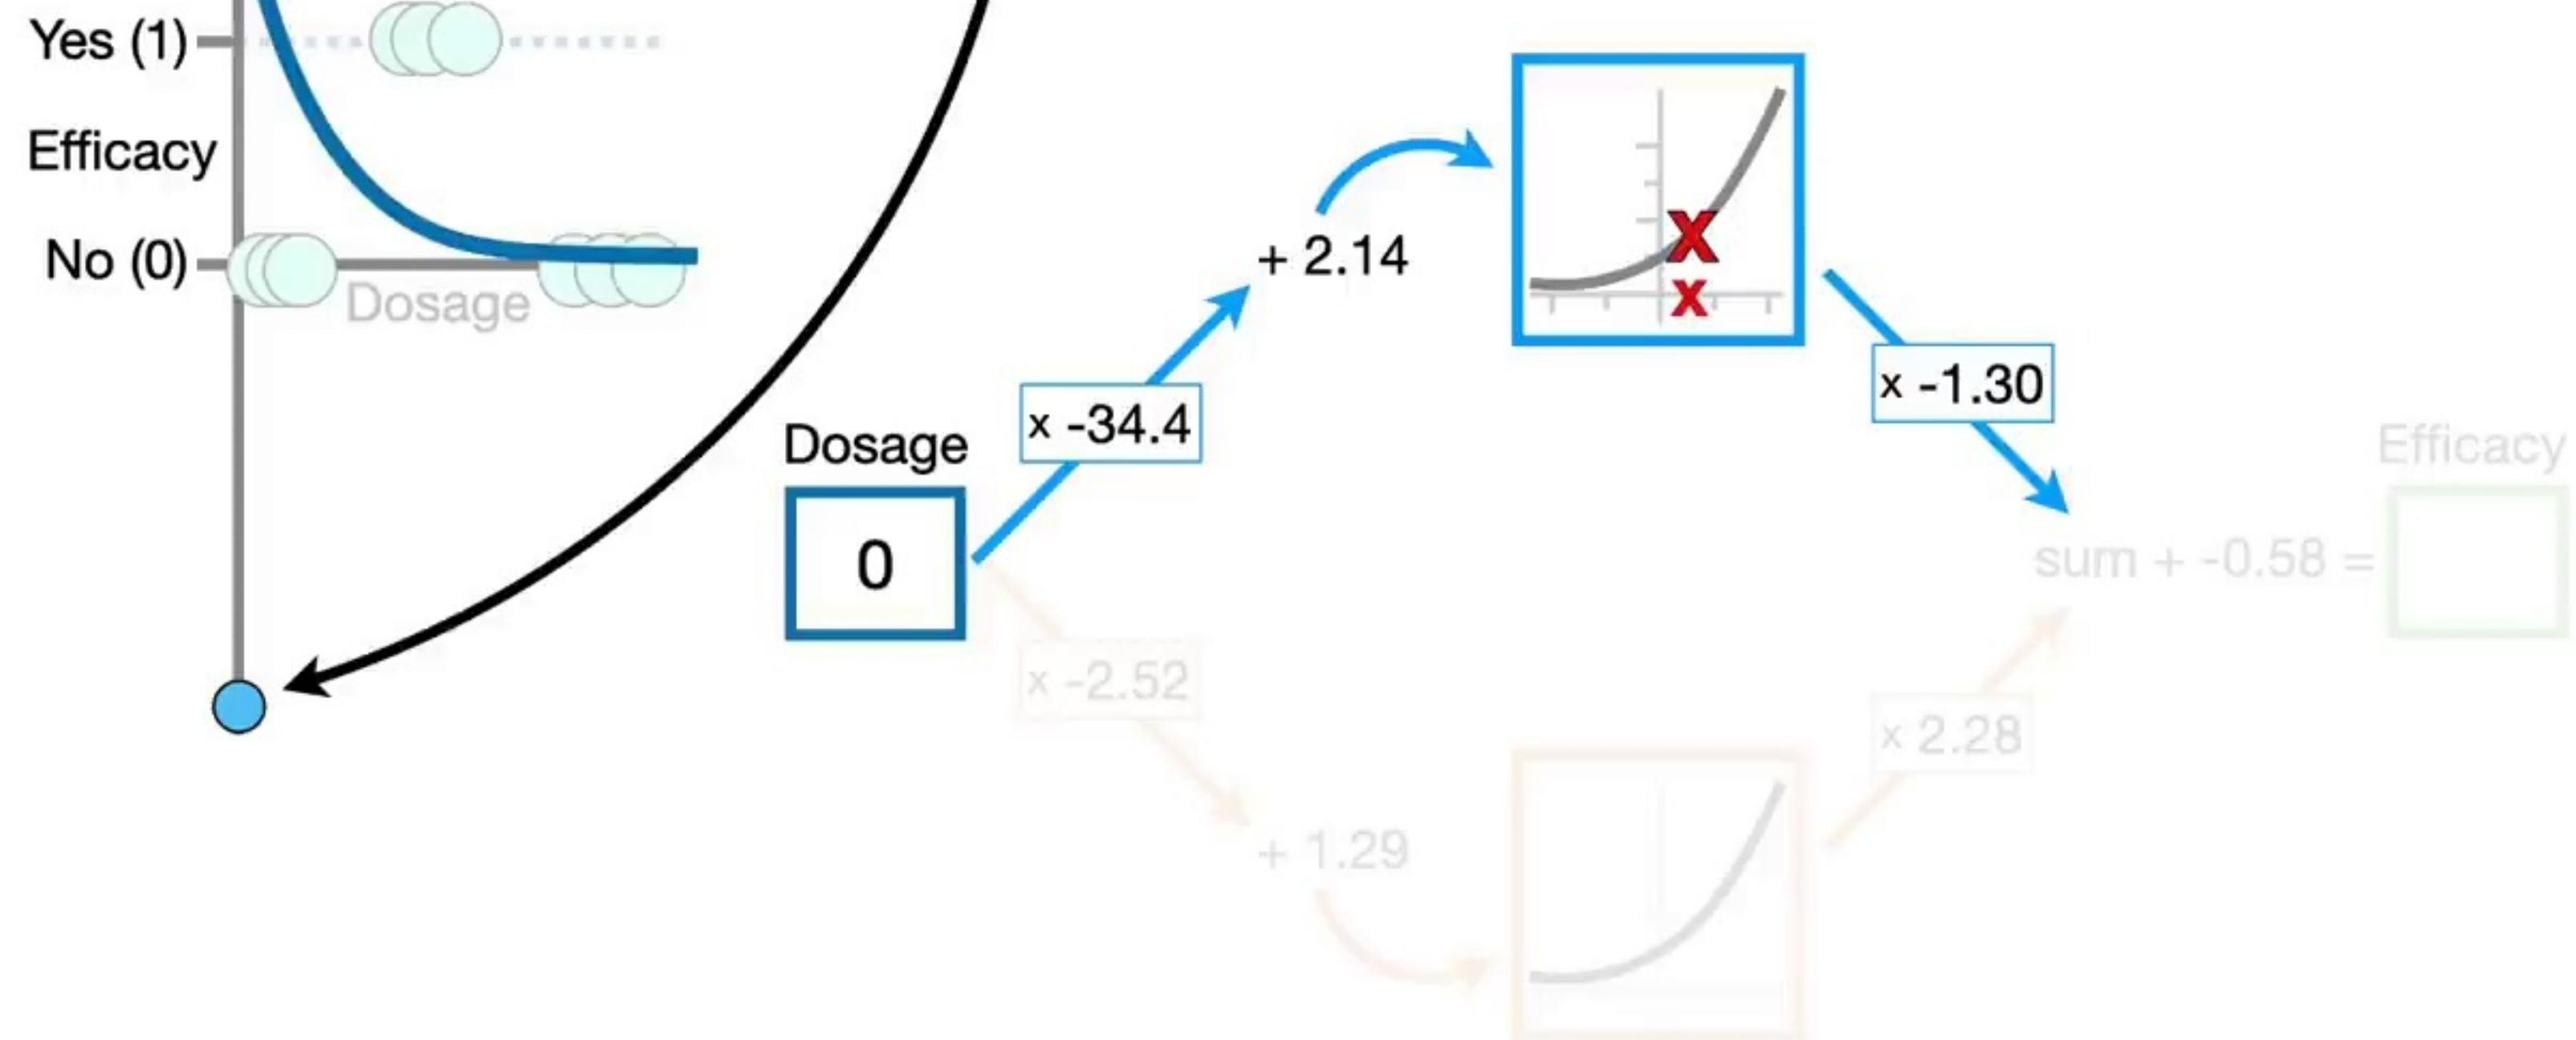
$$2.25 \times -1.30$$





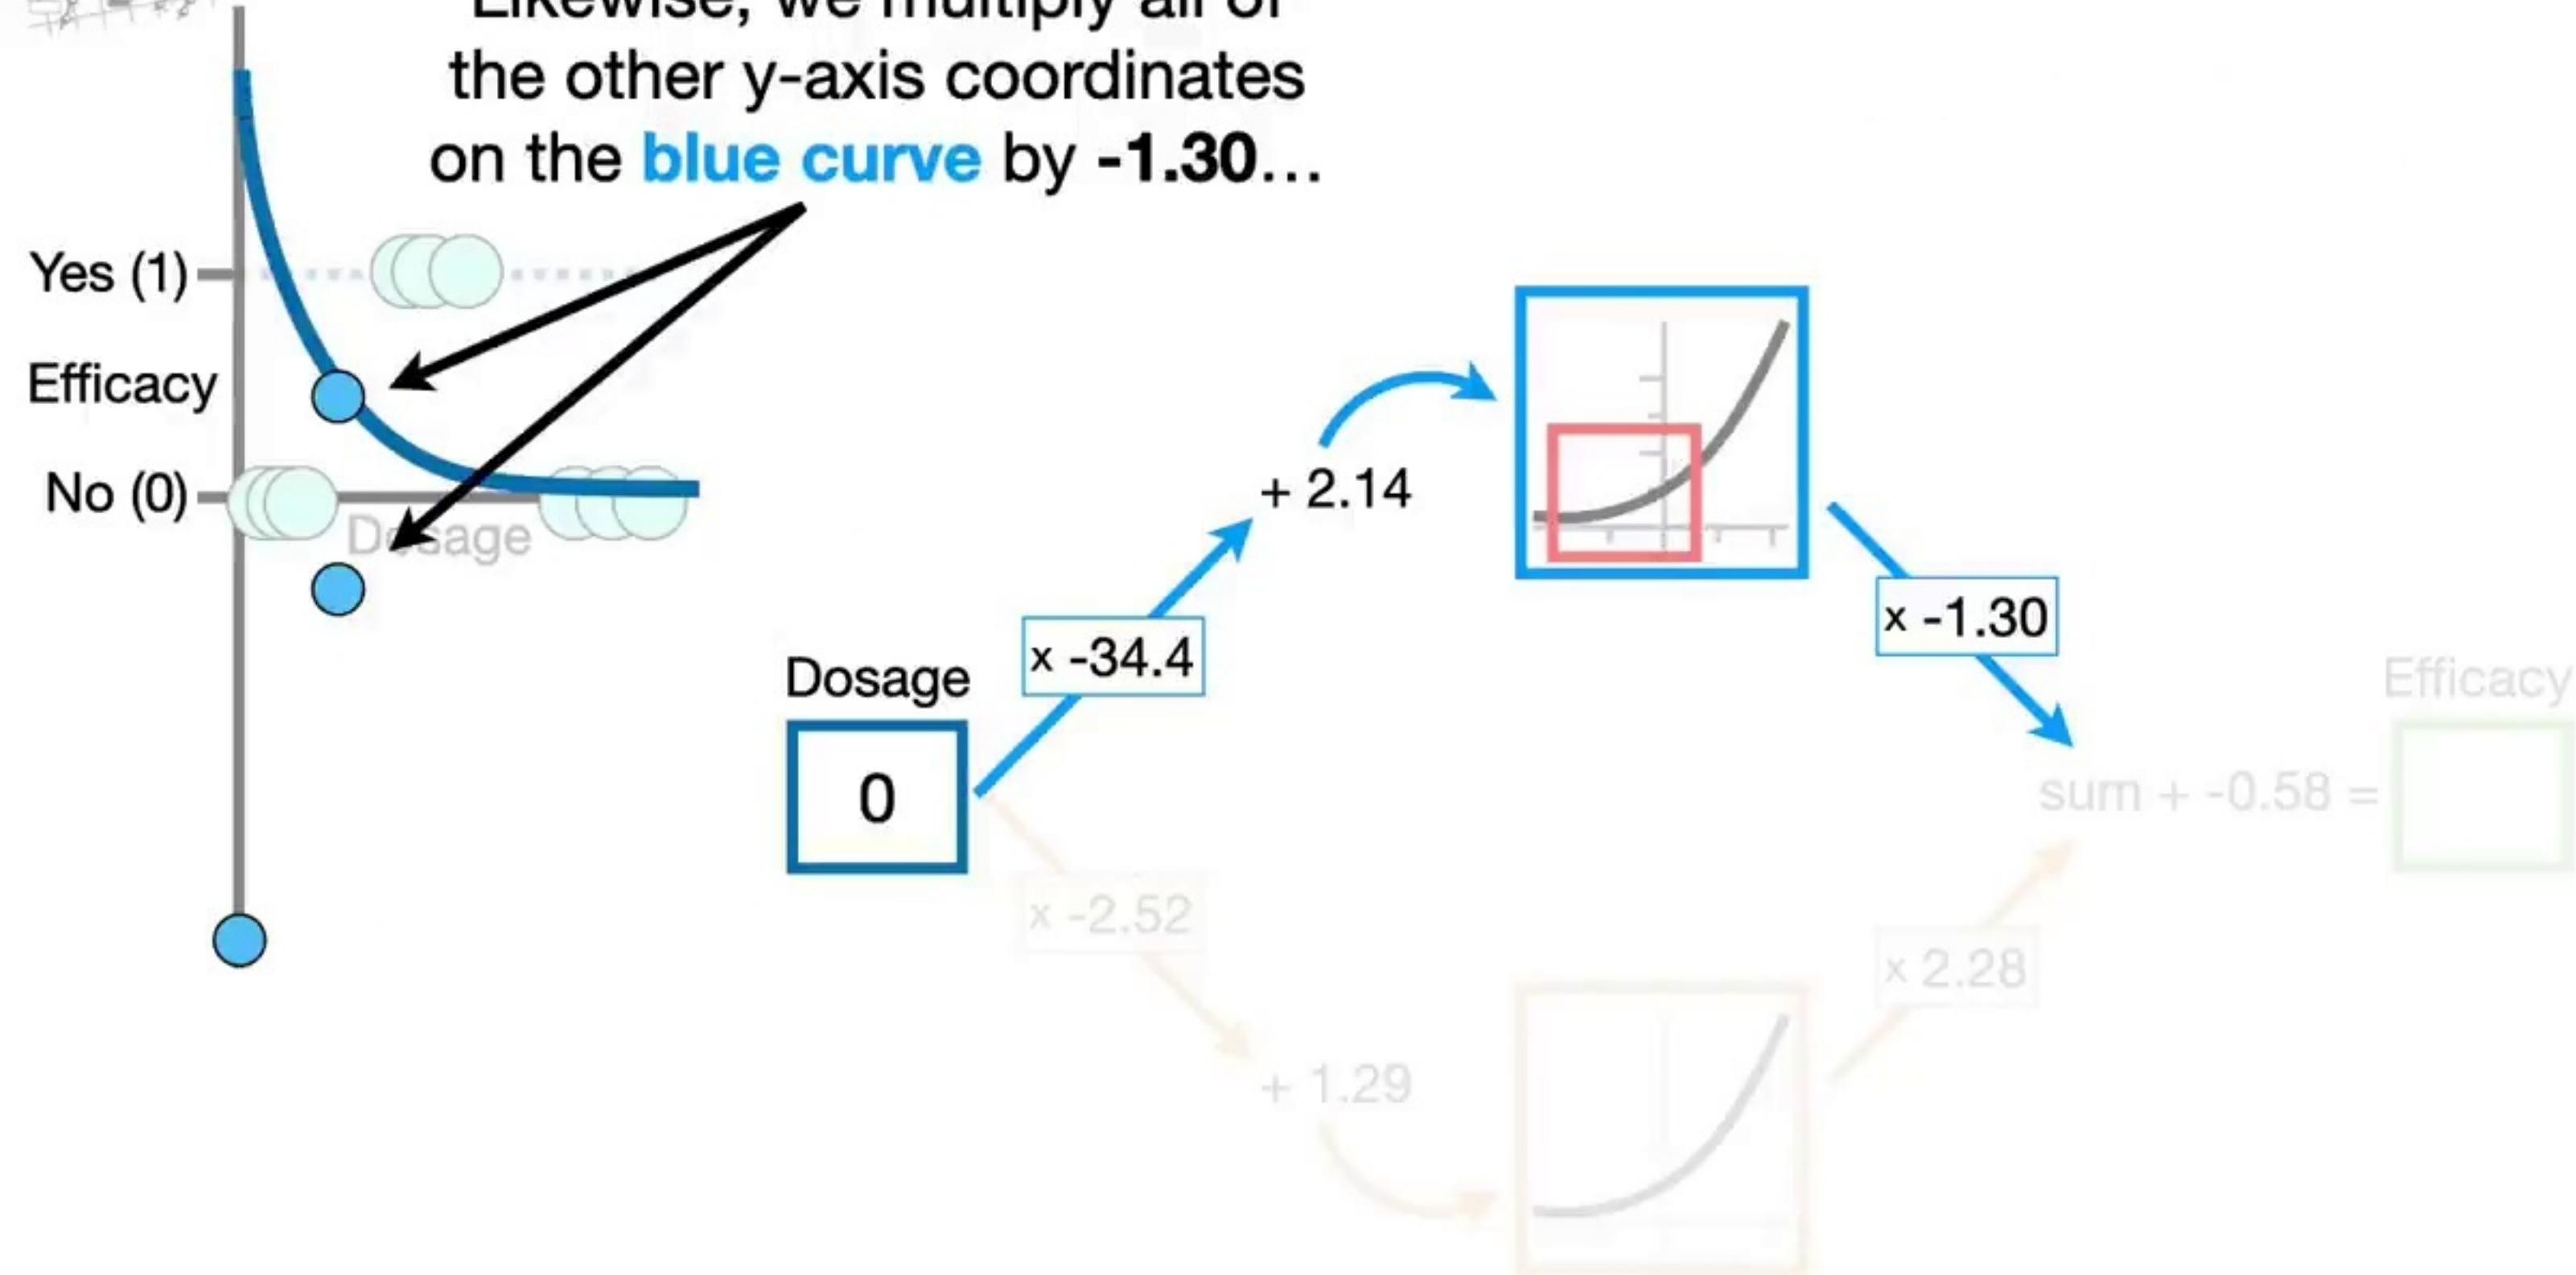
...and -2.93 corresponds to this position on the y-axis.

$$2.25 \times -1.30 = -2.93$$



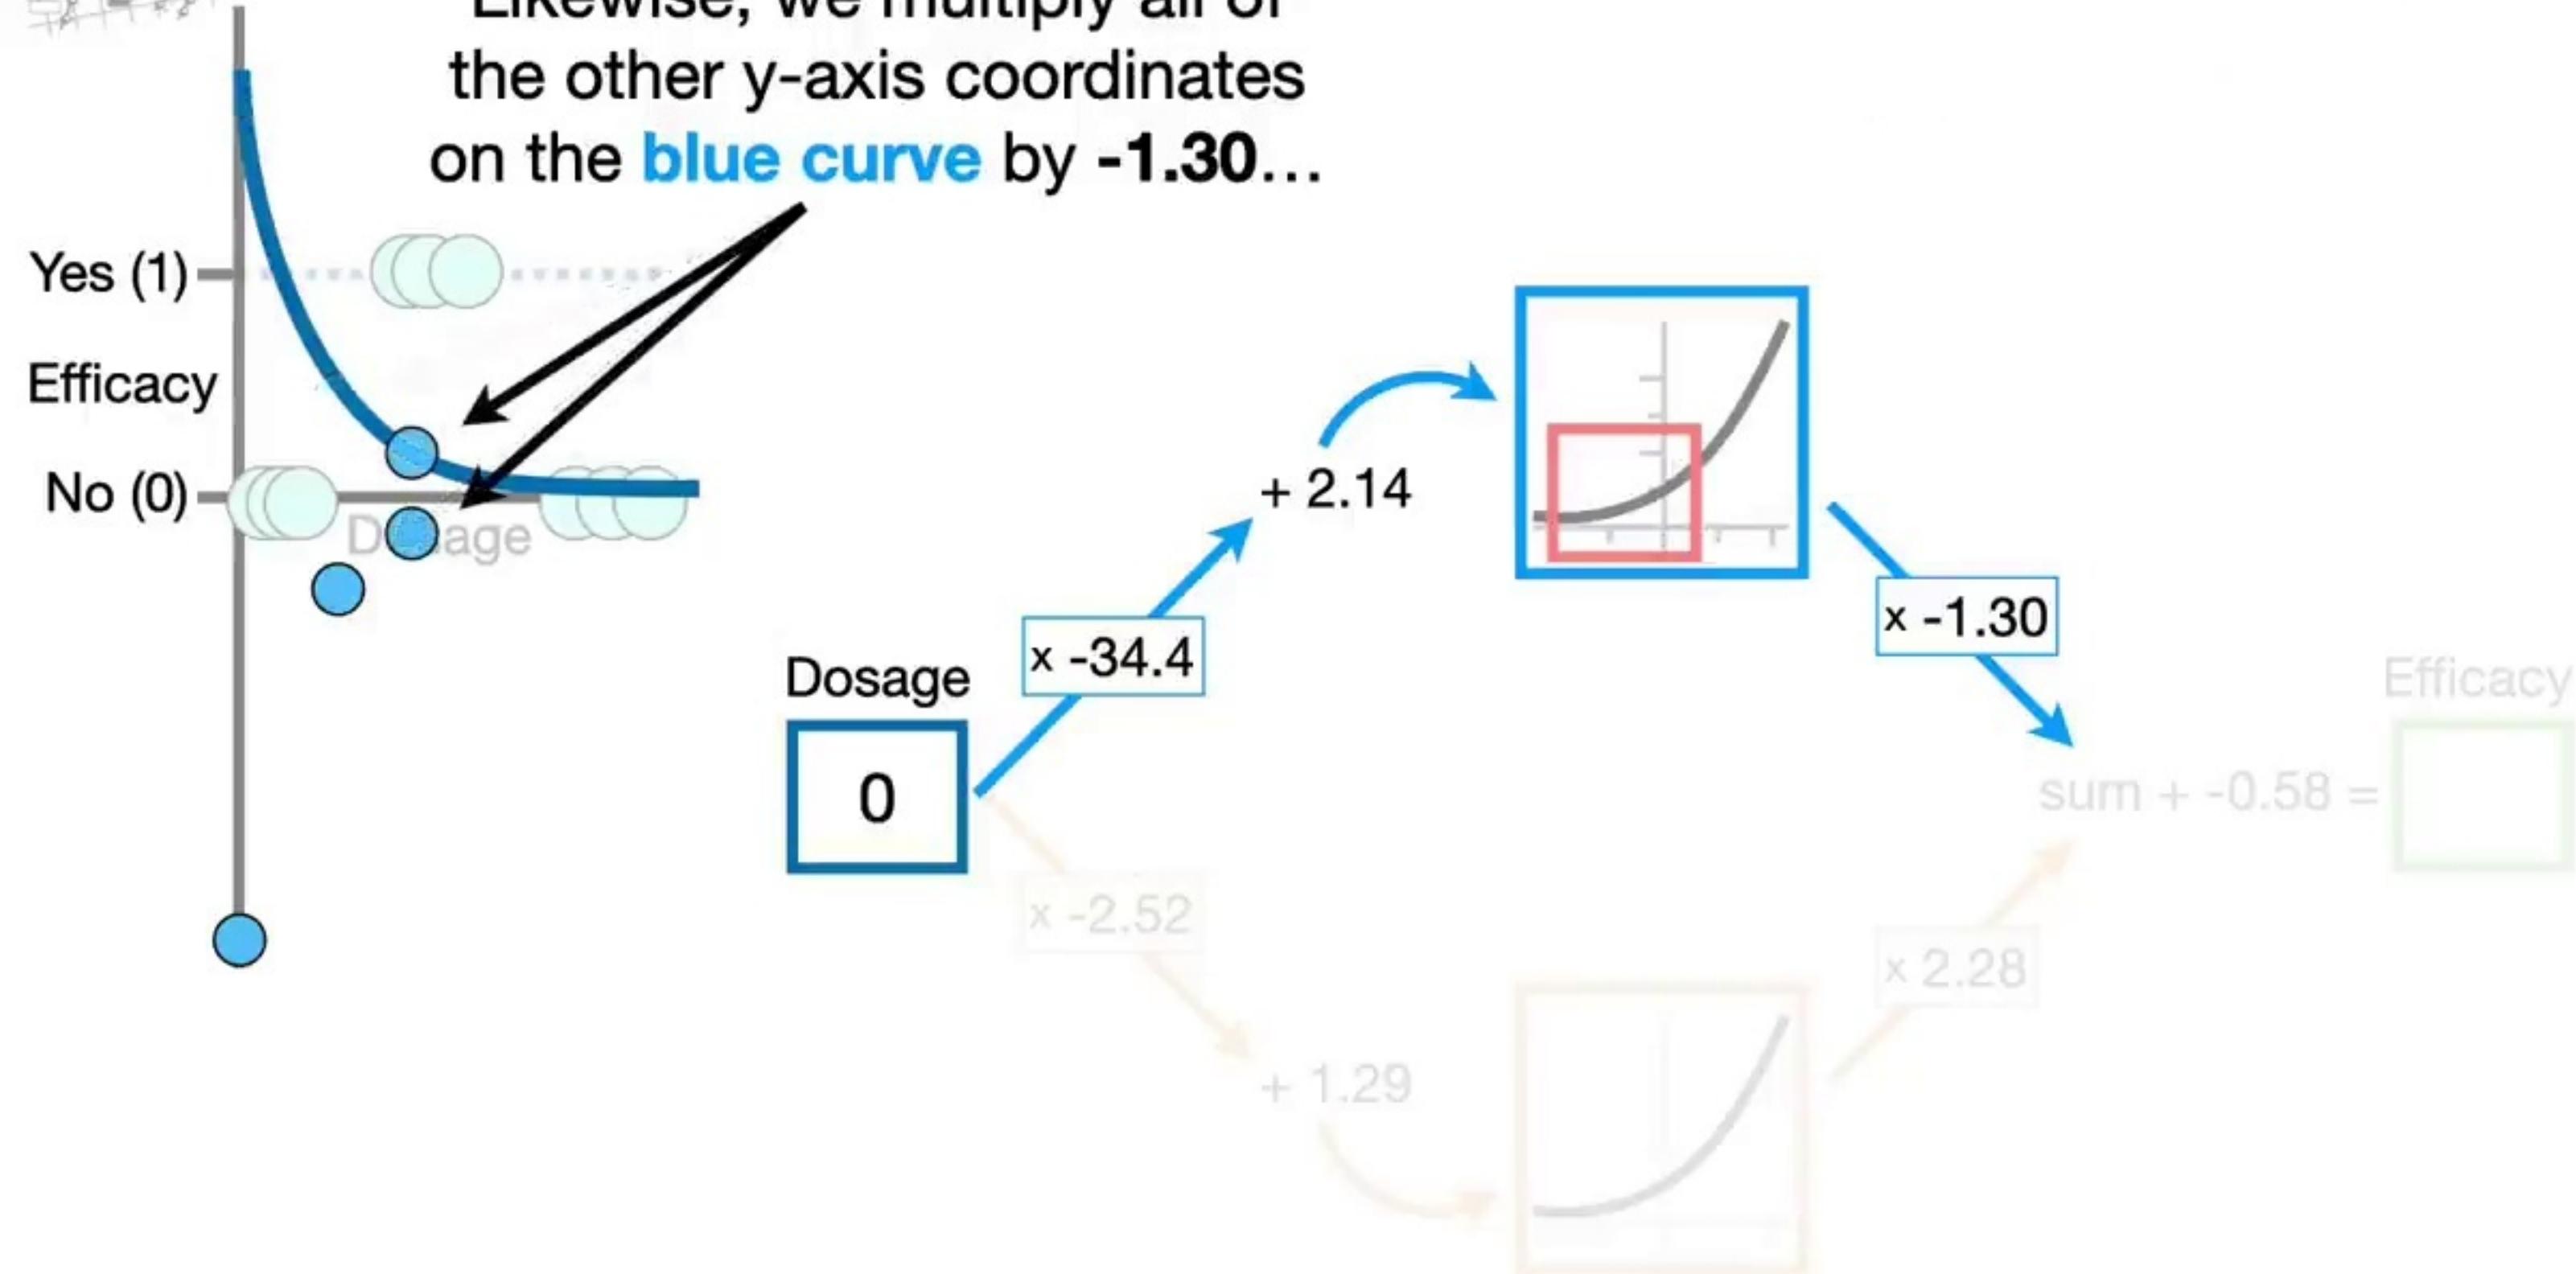


Likewise, we multiply all of the other y-axis coordinates on the **blue curve** by -1.30...



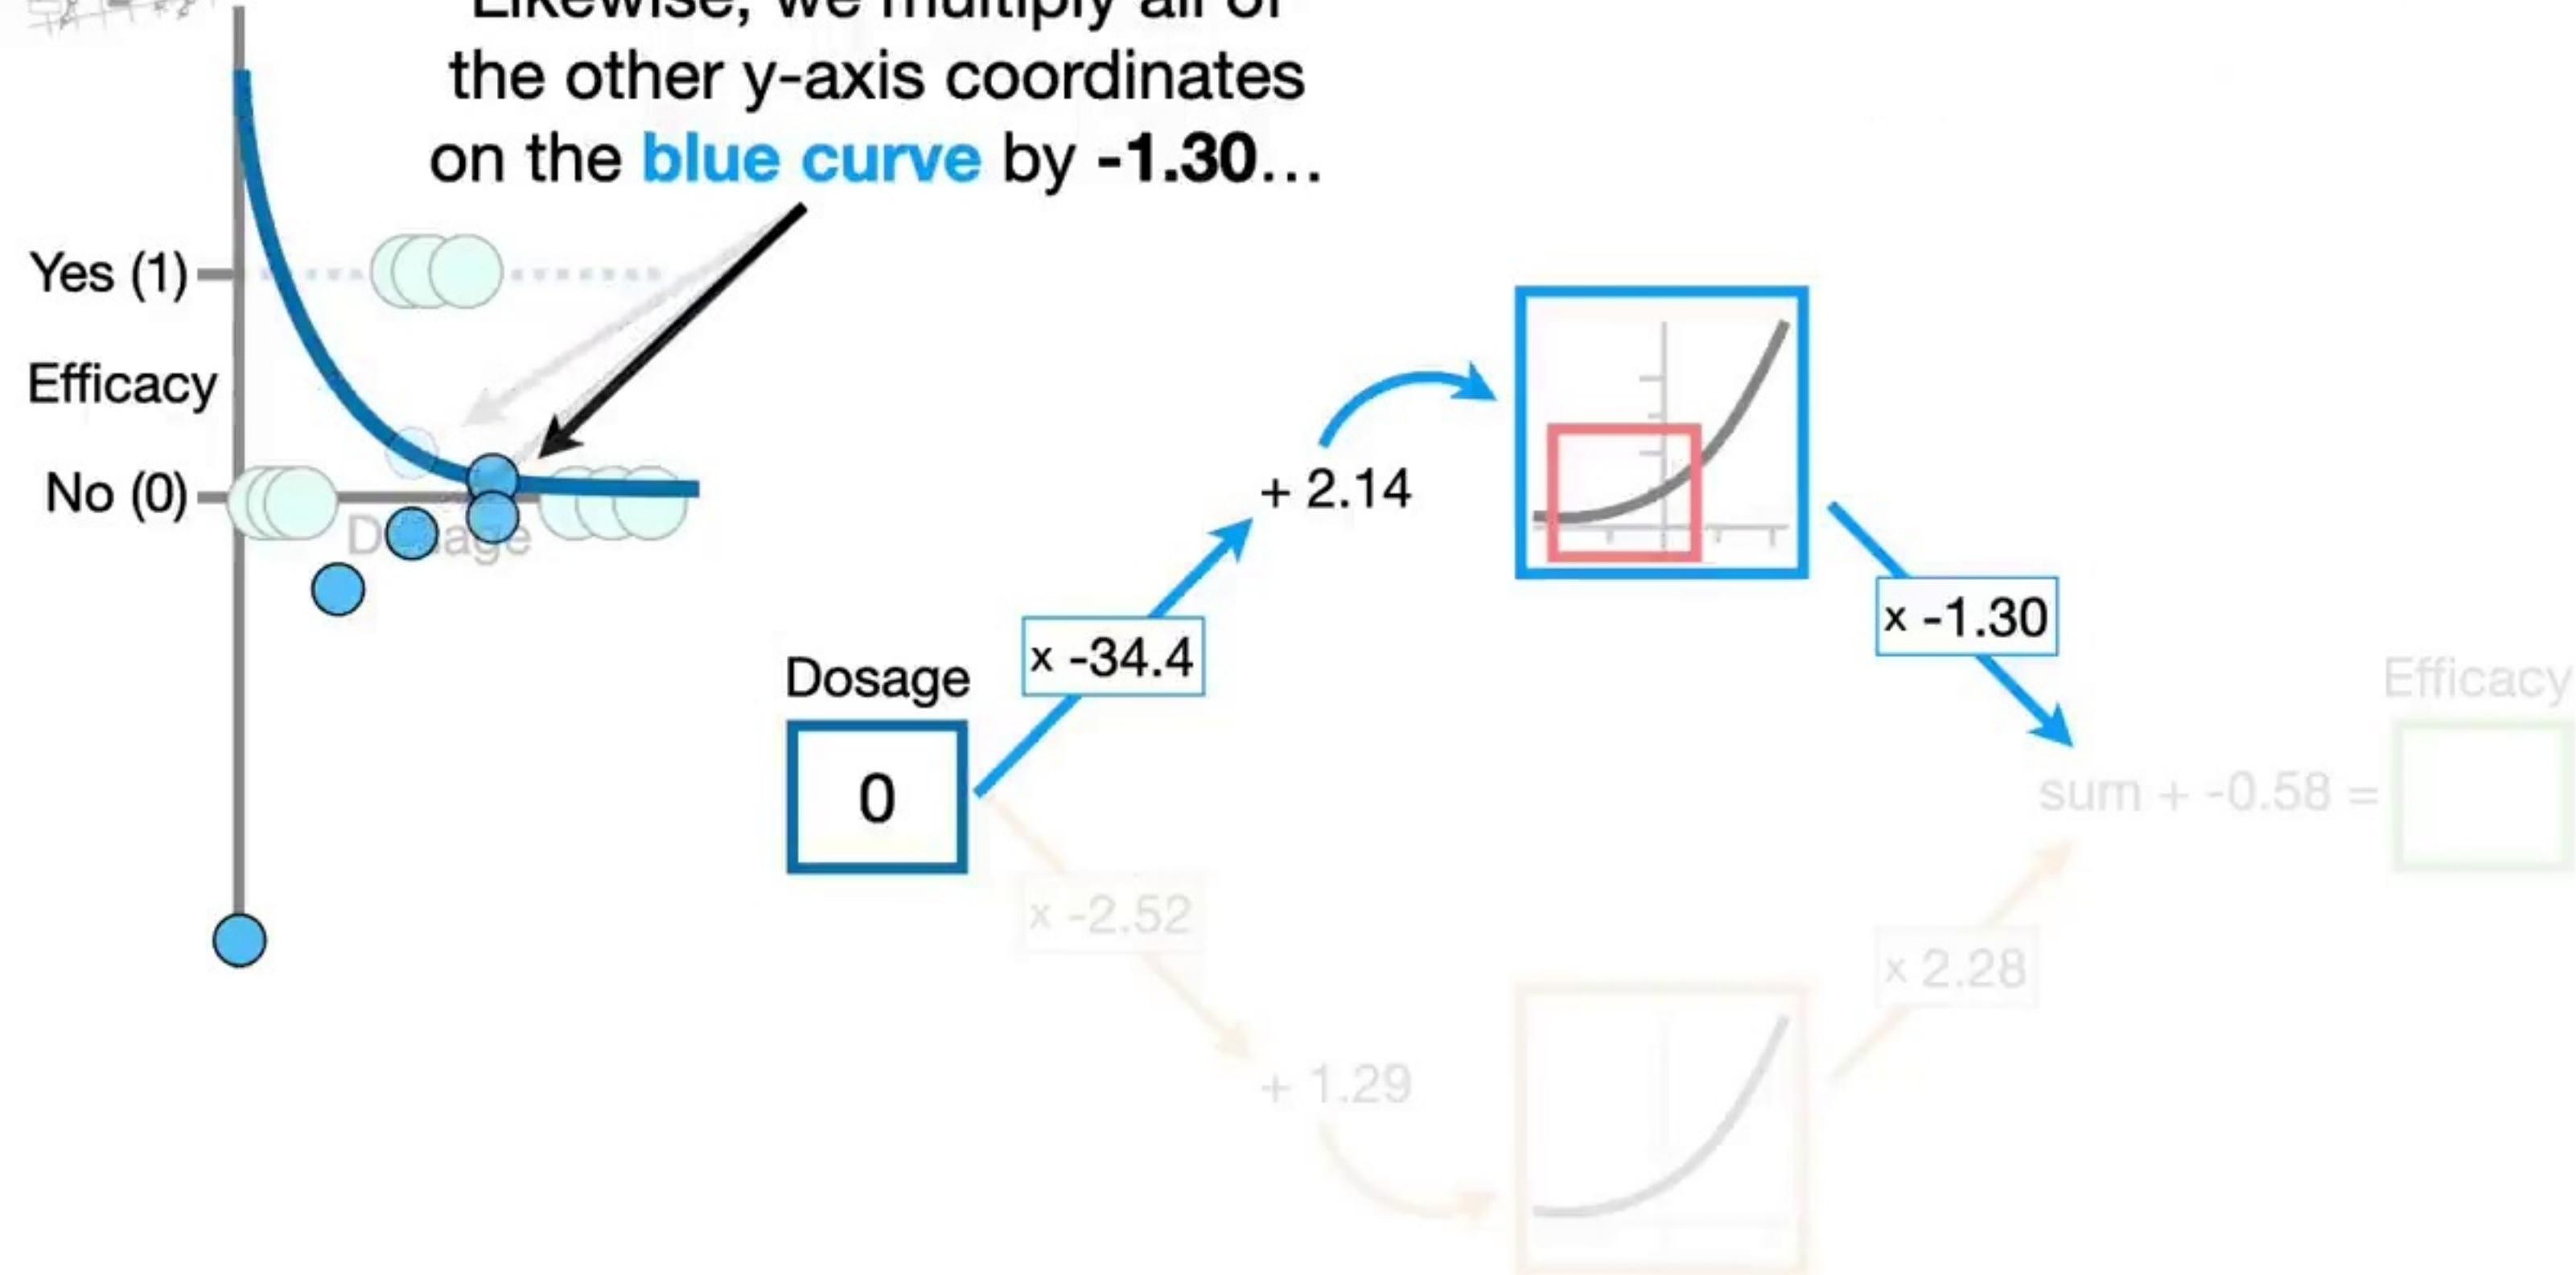


Likewise, we multiply all of the other y-axis coordinates on the **blue curve** by -1.30...



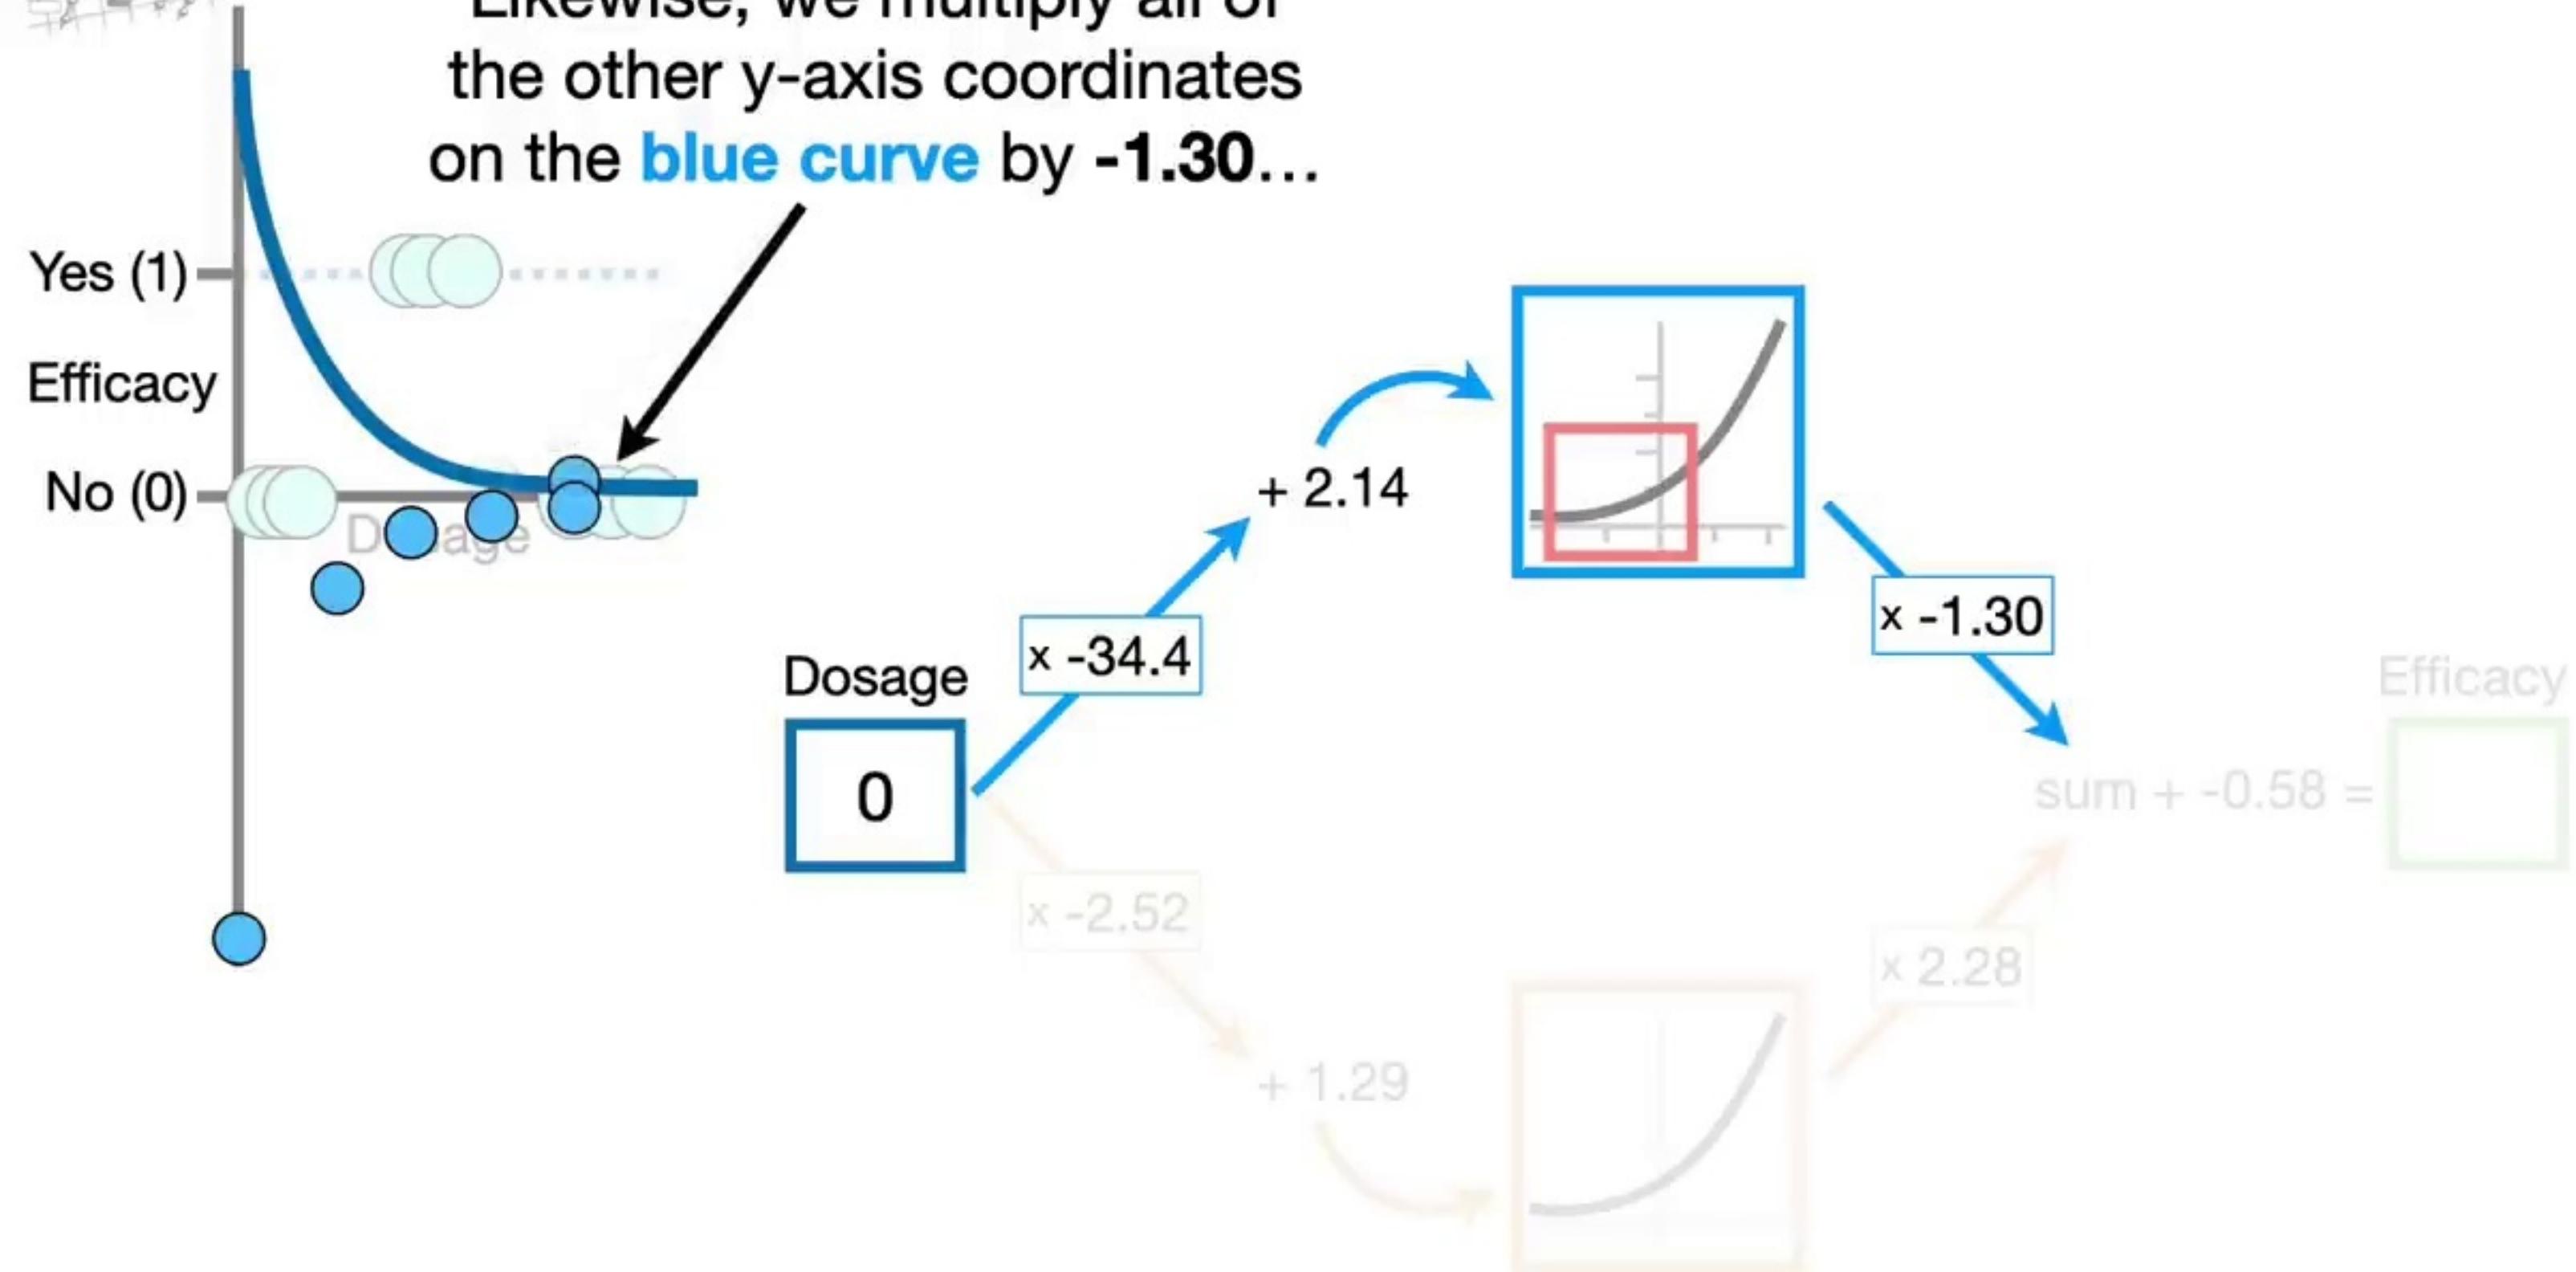


Likewise, we multiply all of the other y-axis coordinates on the **blue curve** by -1.30...



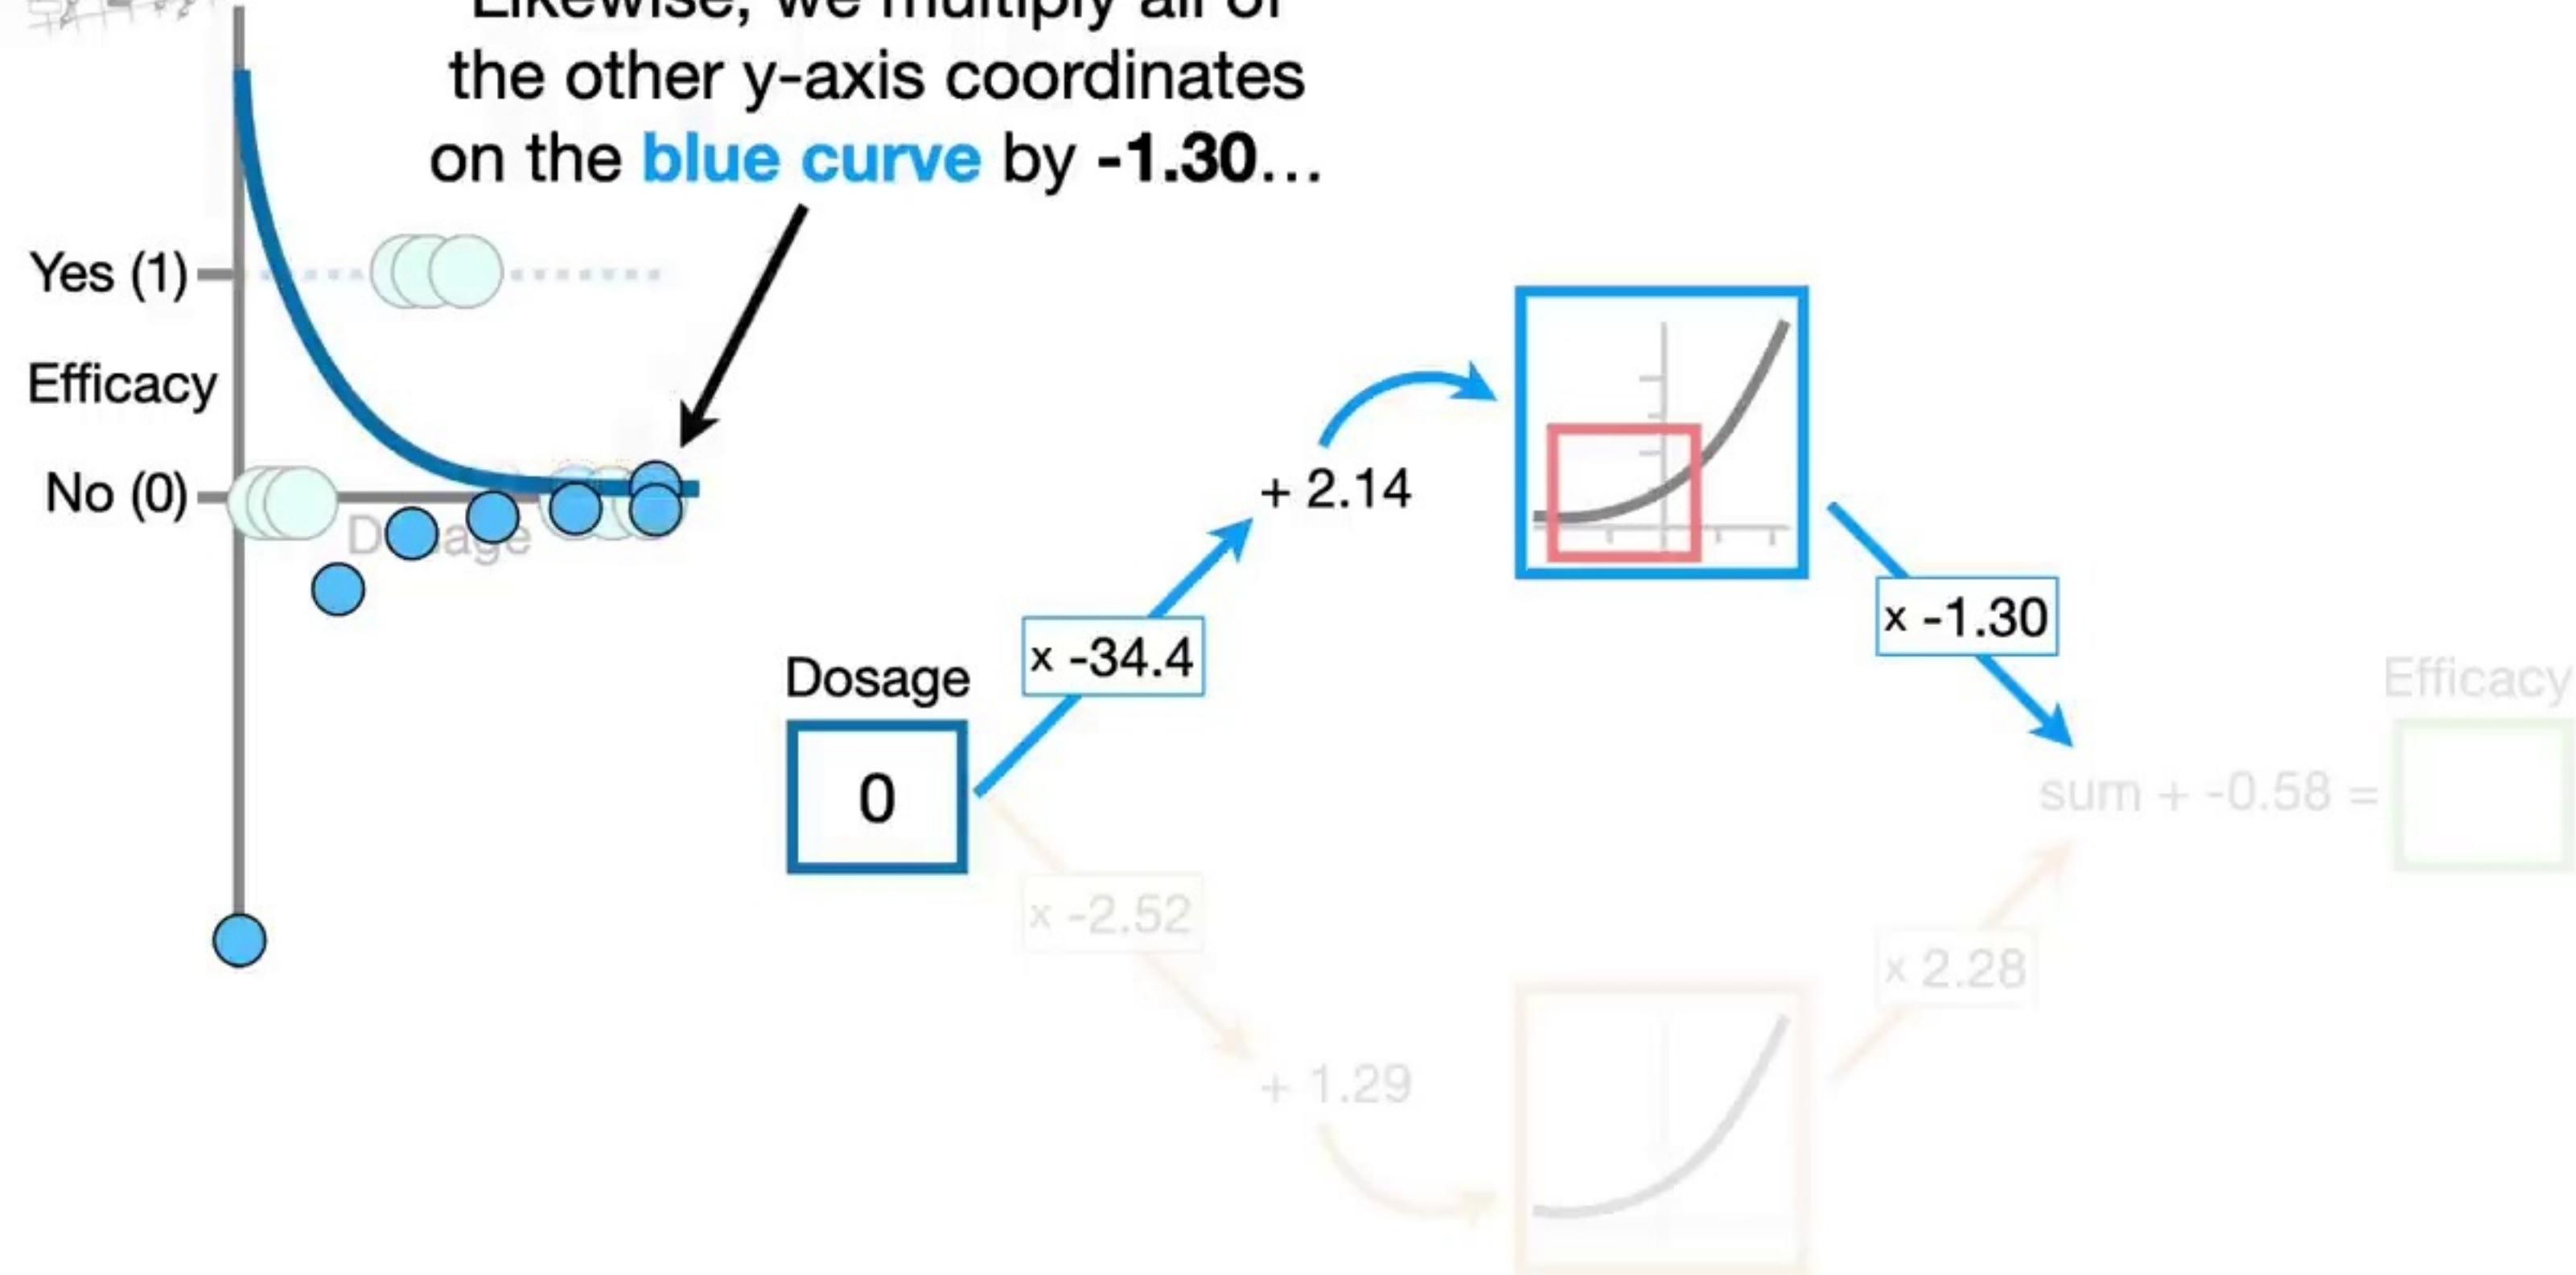


Likewise, we multiply all of the other y-axis coordinates on the **blue curve** by -1.30...



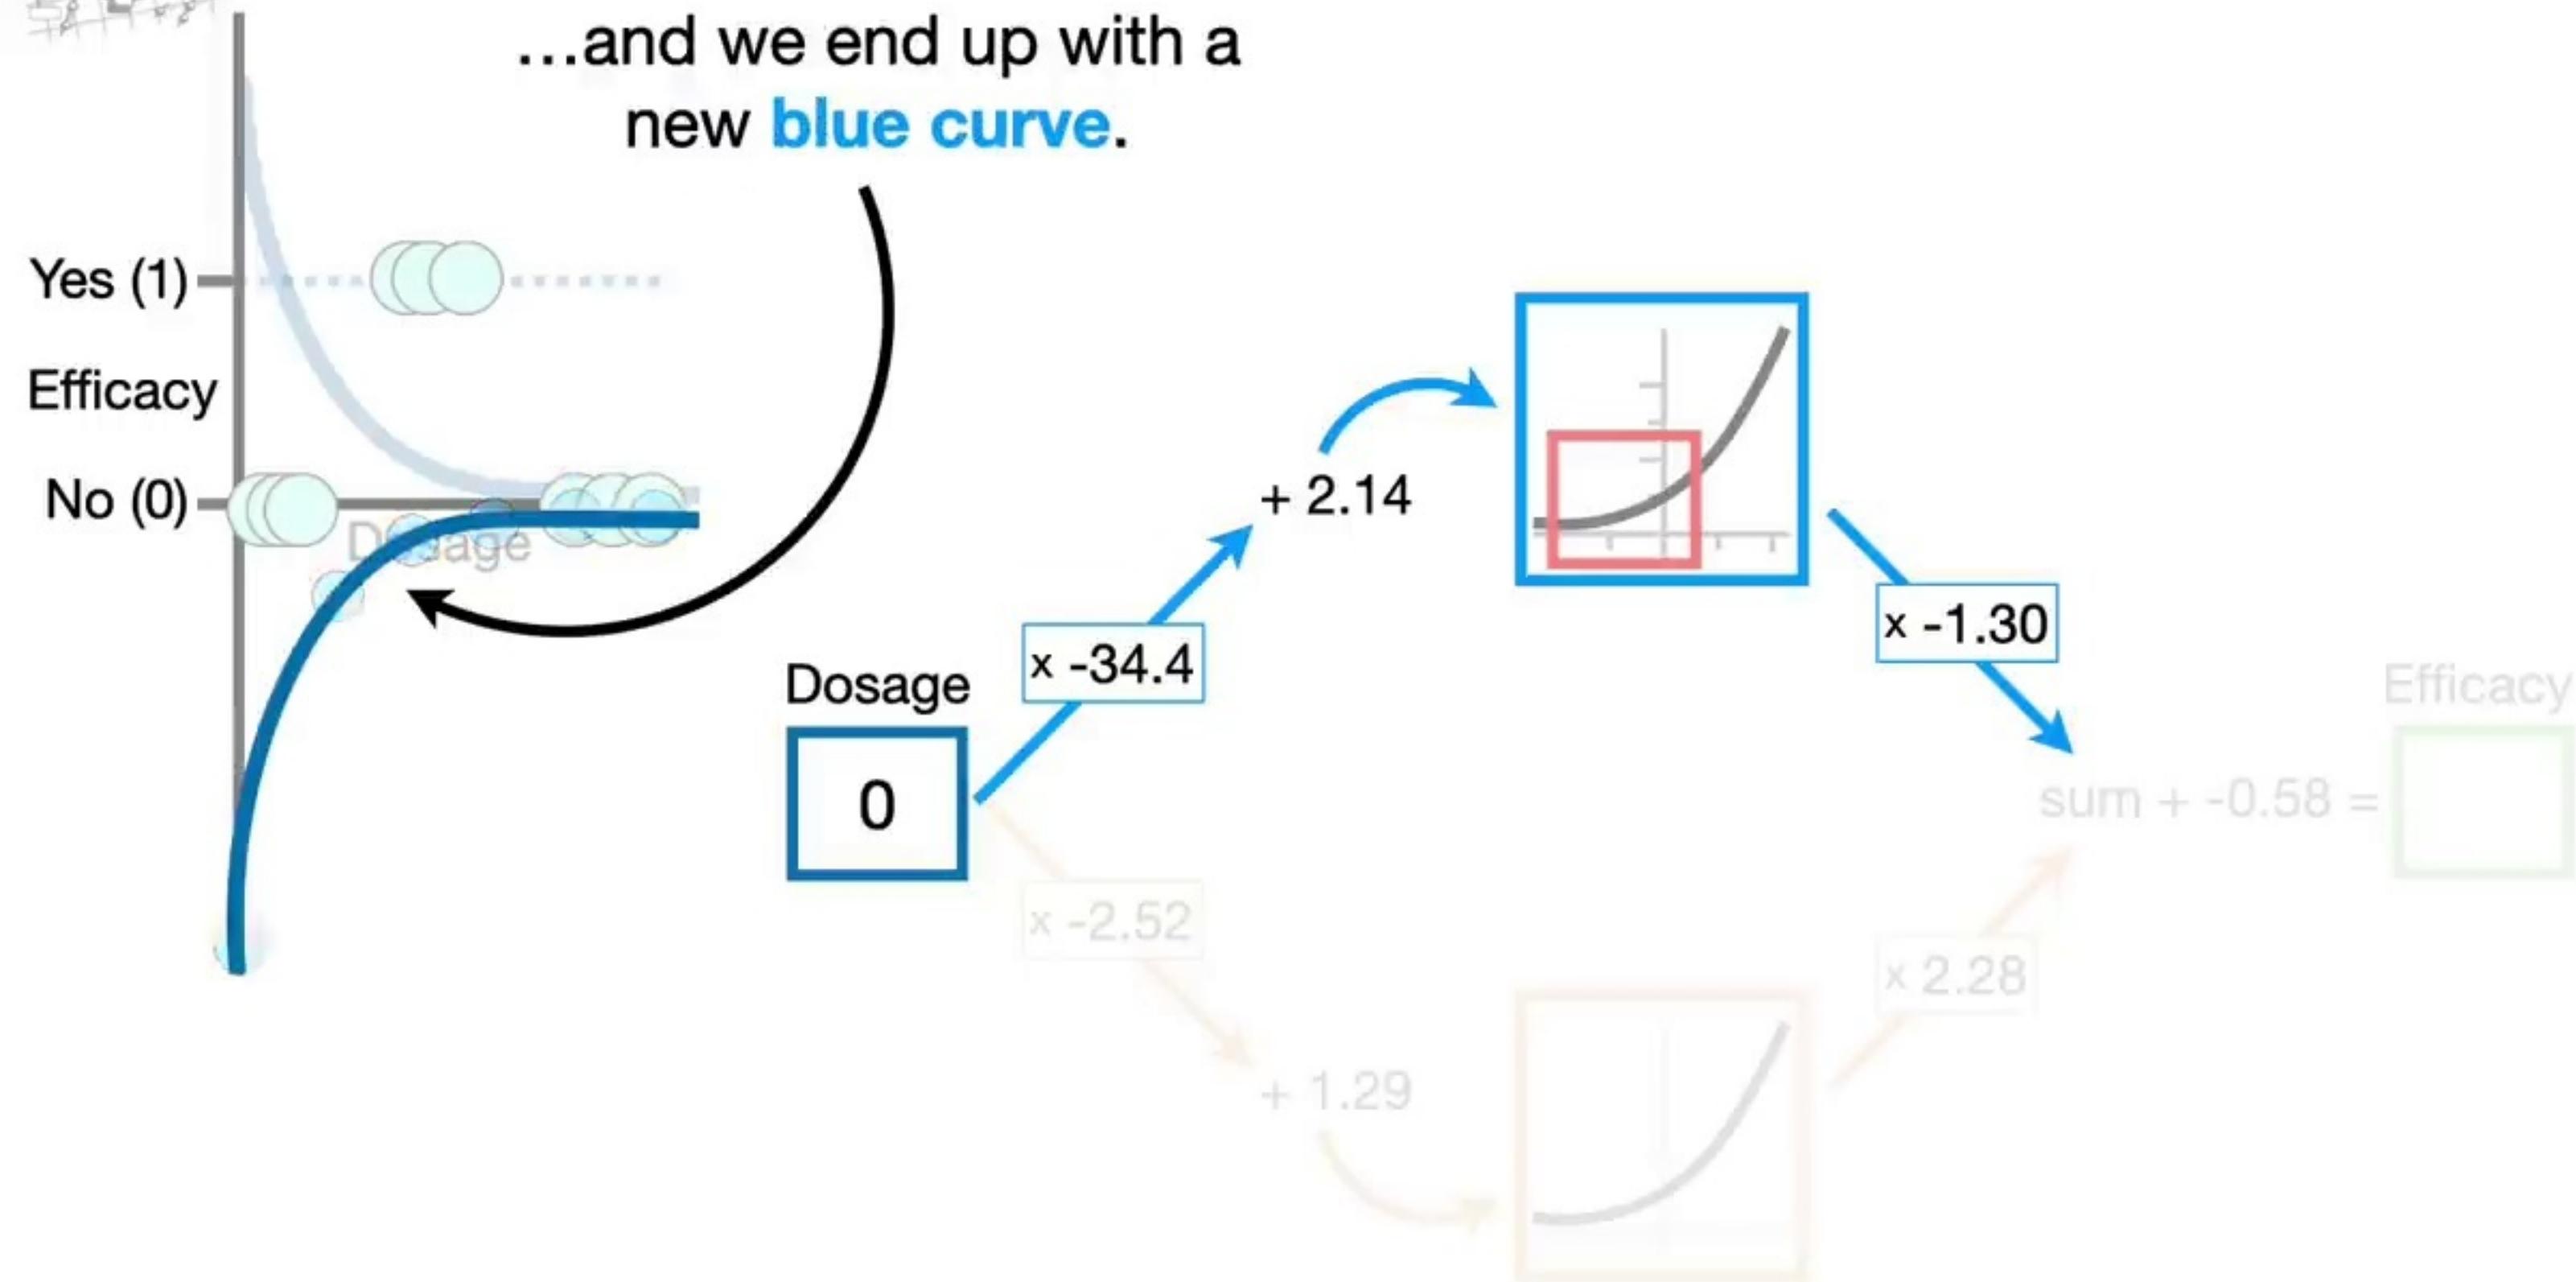


Likewise, we multiply all of the other y-axis coordinates on the **blue curve** by -1.30...



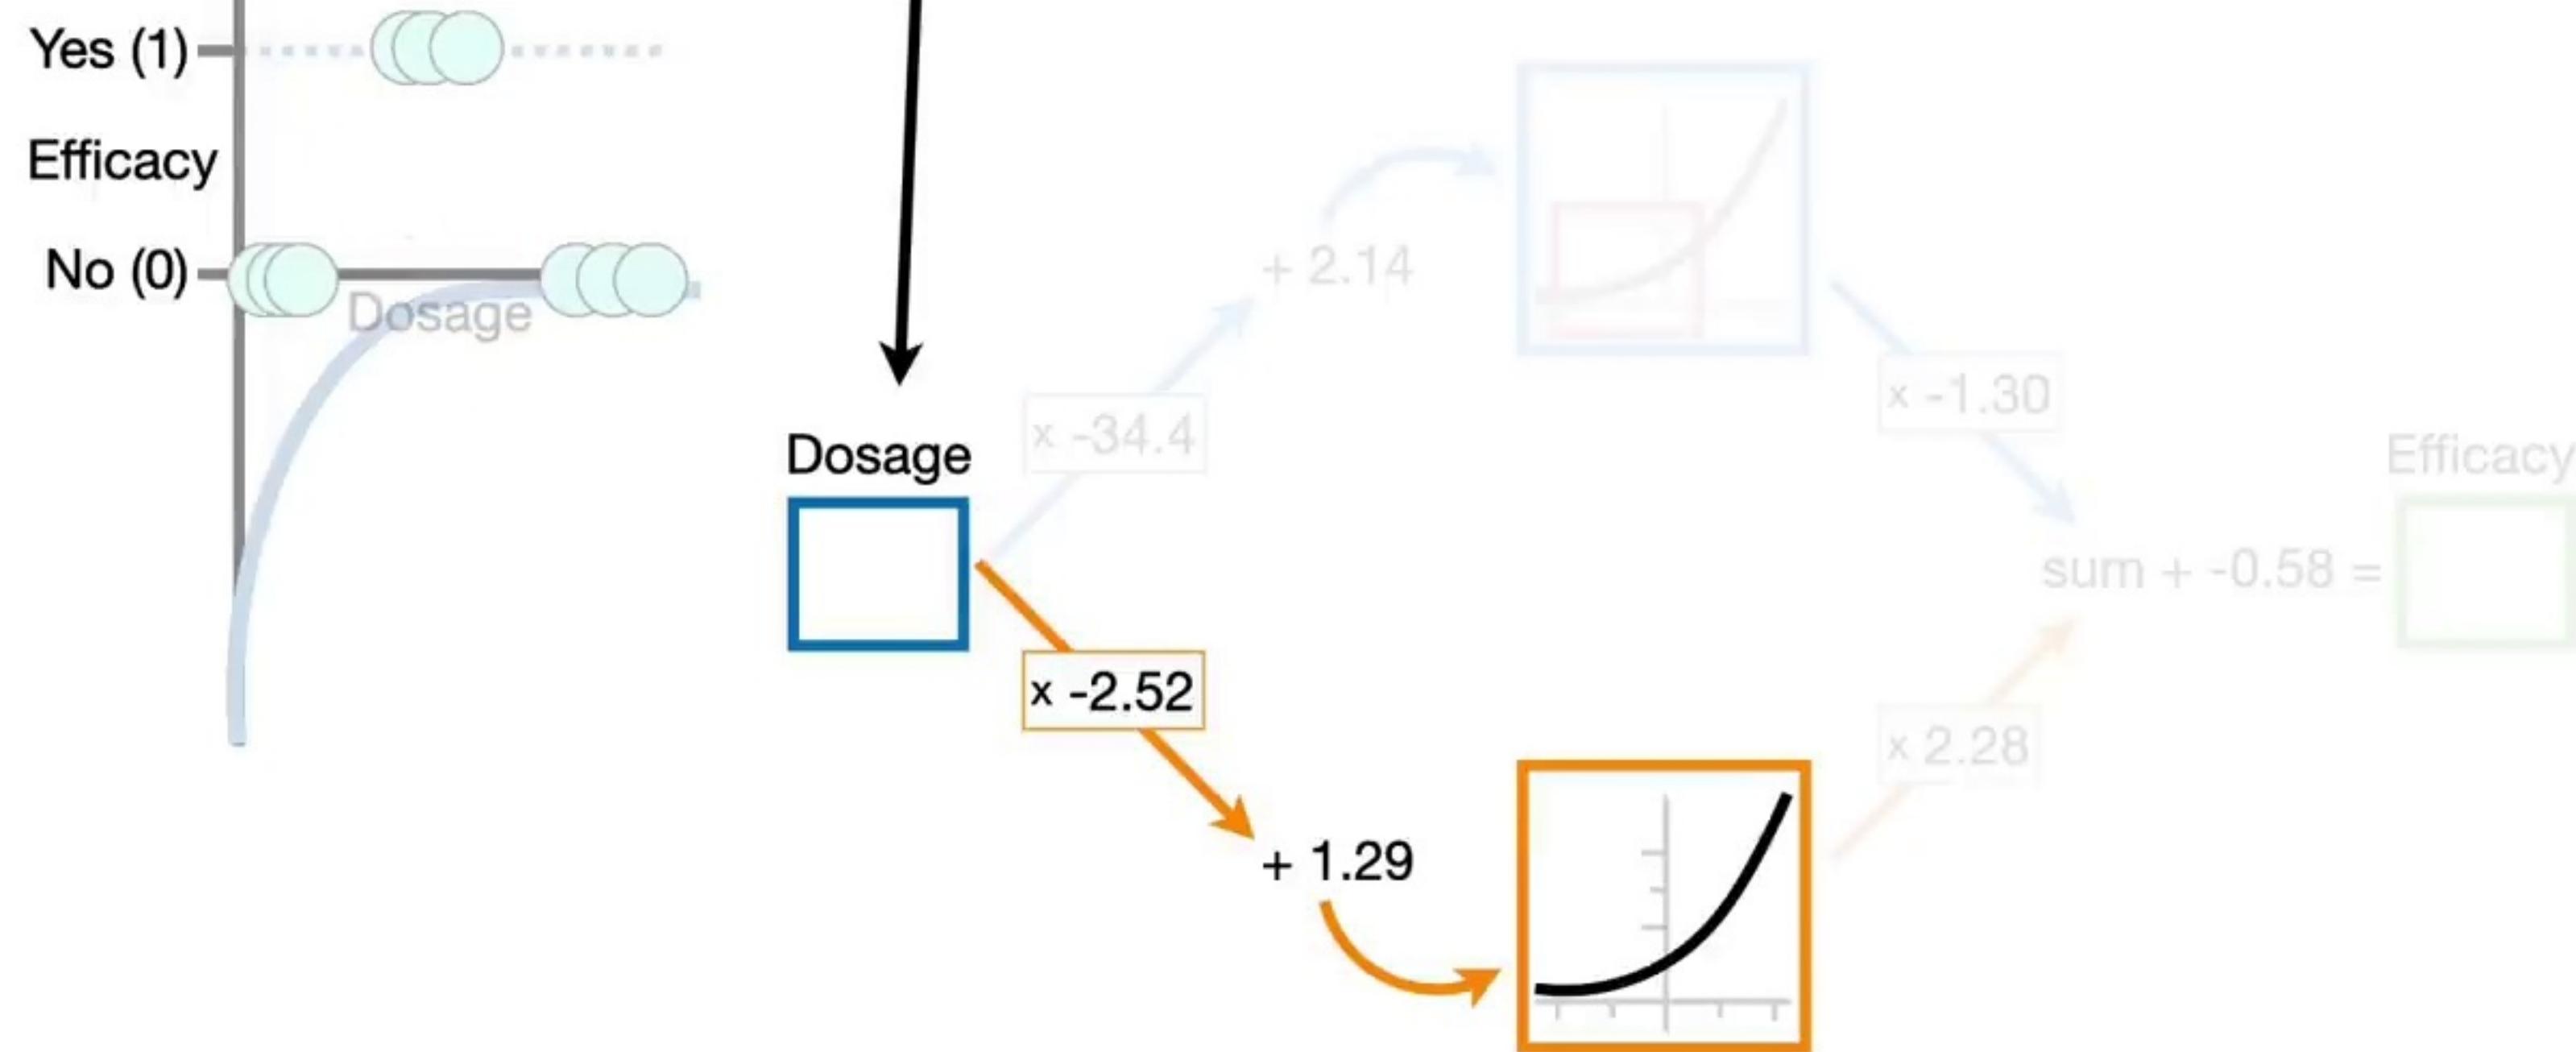


...and we end up with a new **blue curve**.





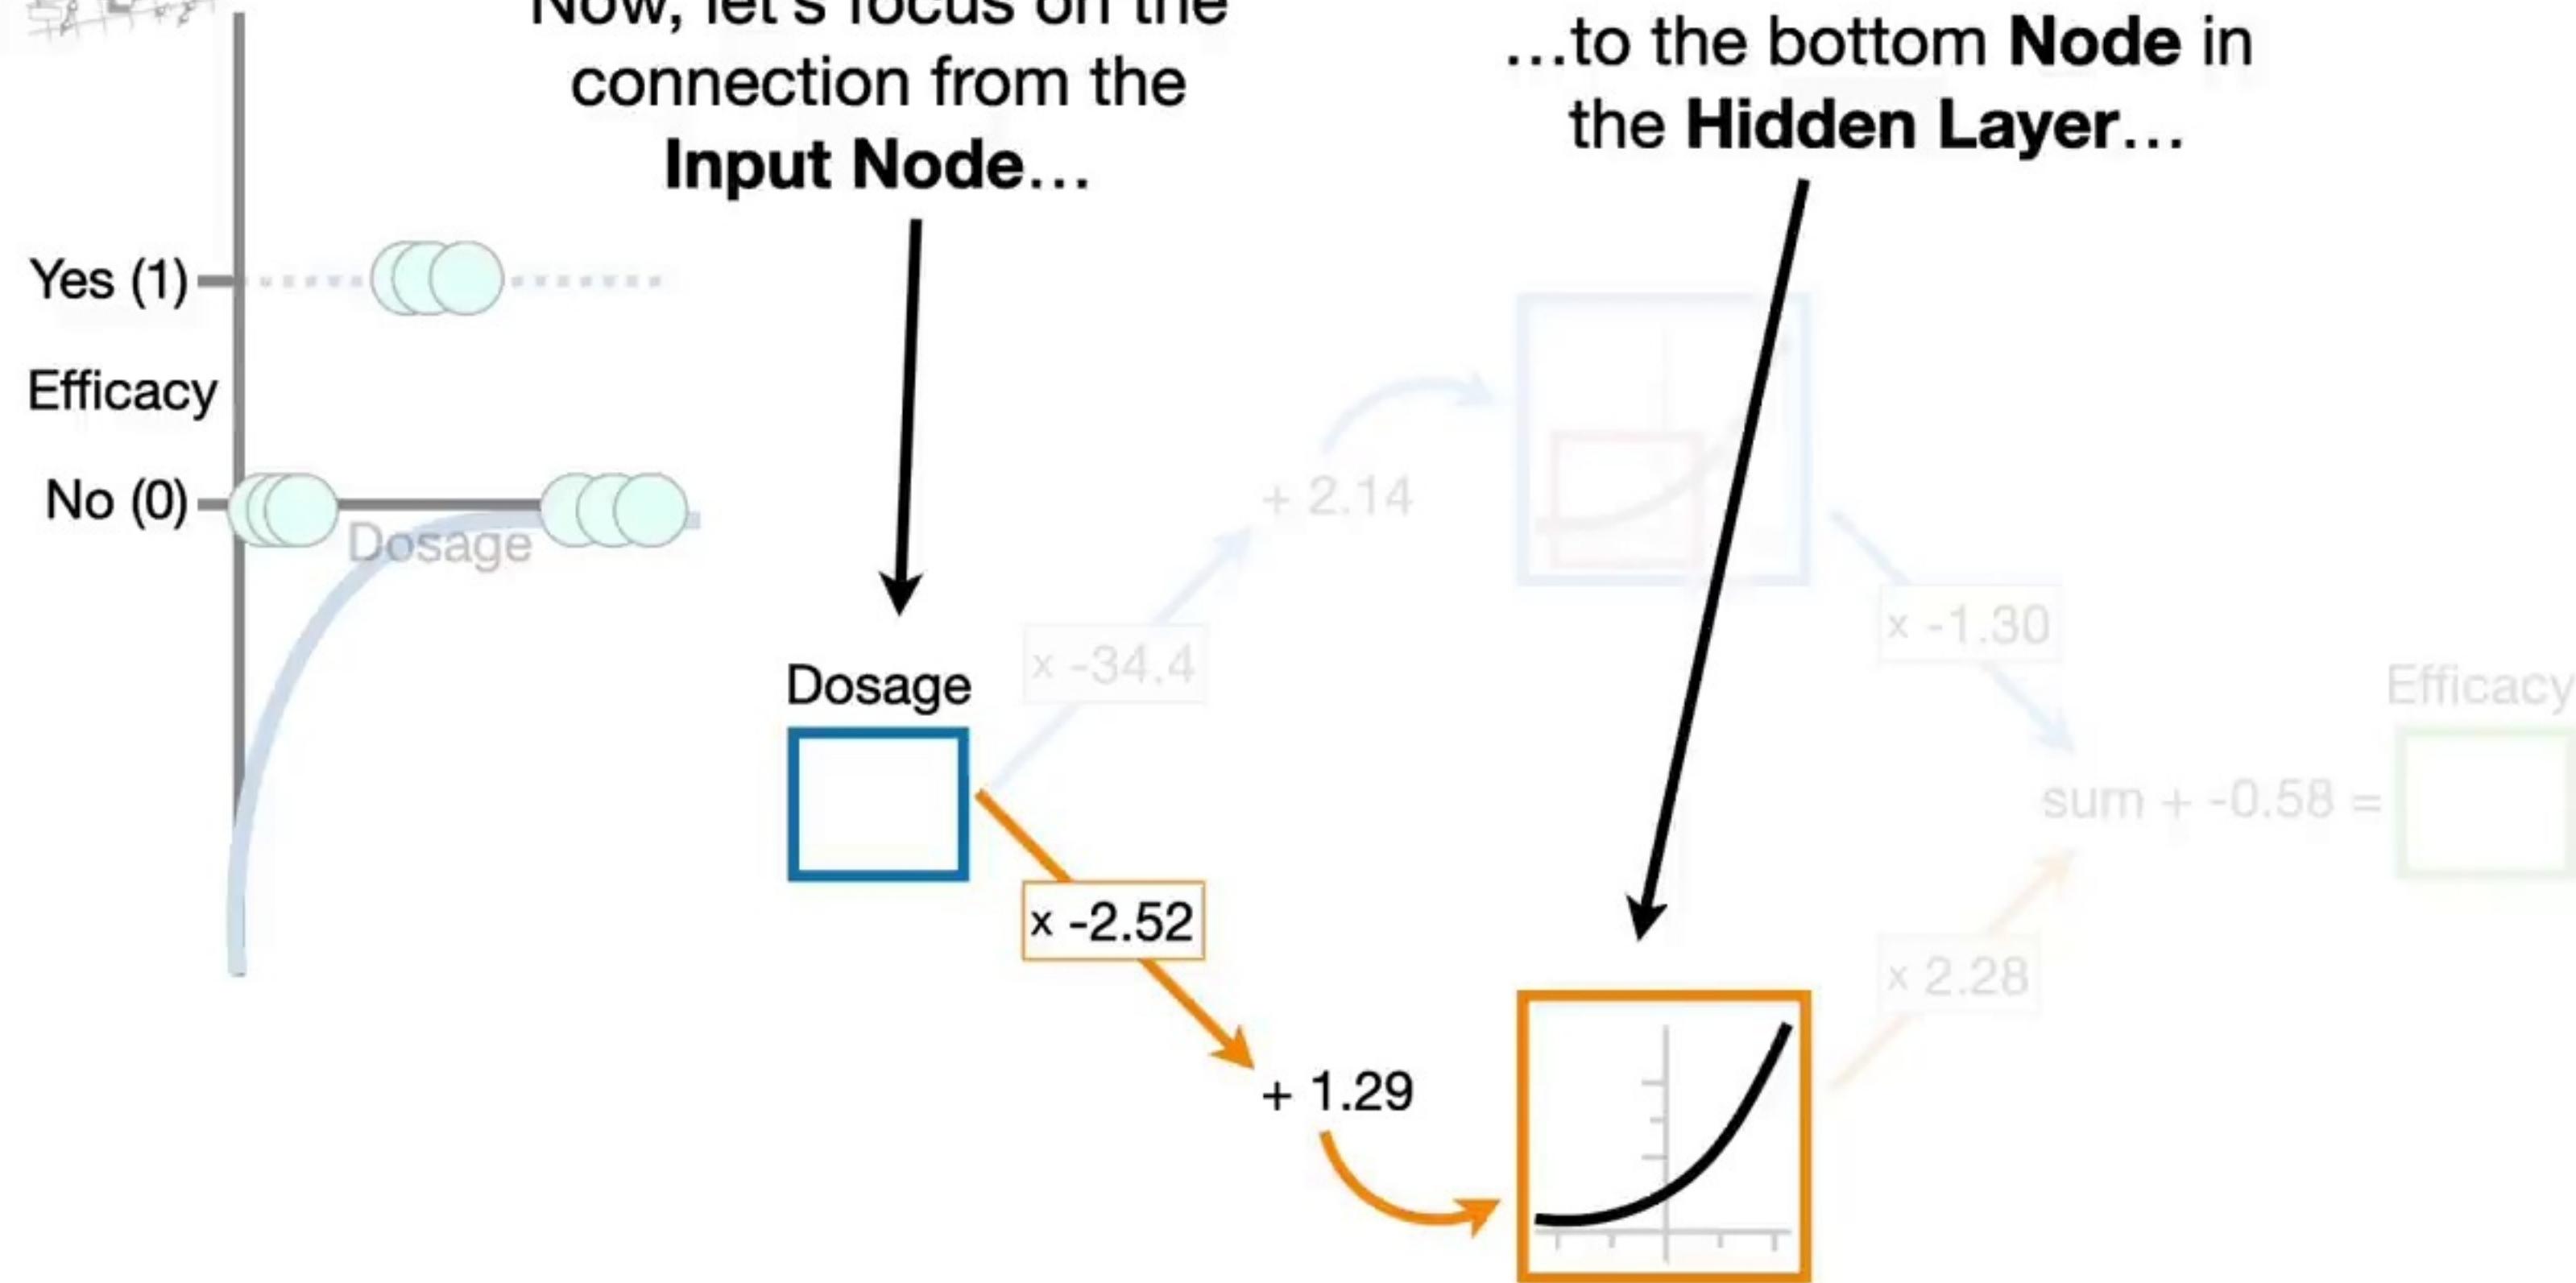
Now, let's focus on the connection from the  
**Input Node...**





Now, let's focus on the connection from the **Input Node...**

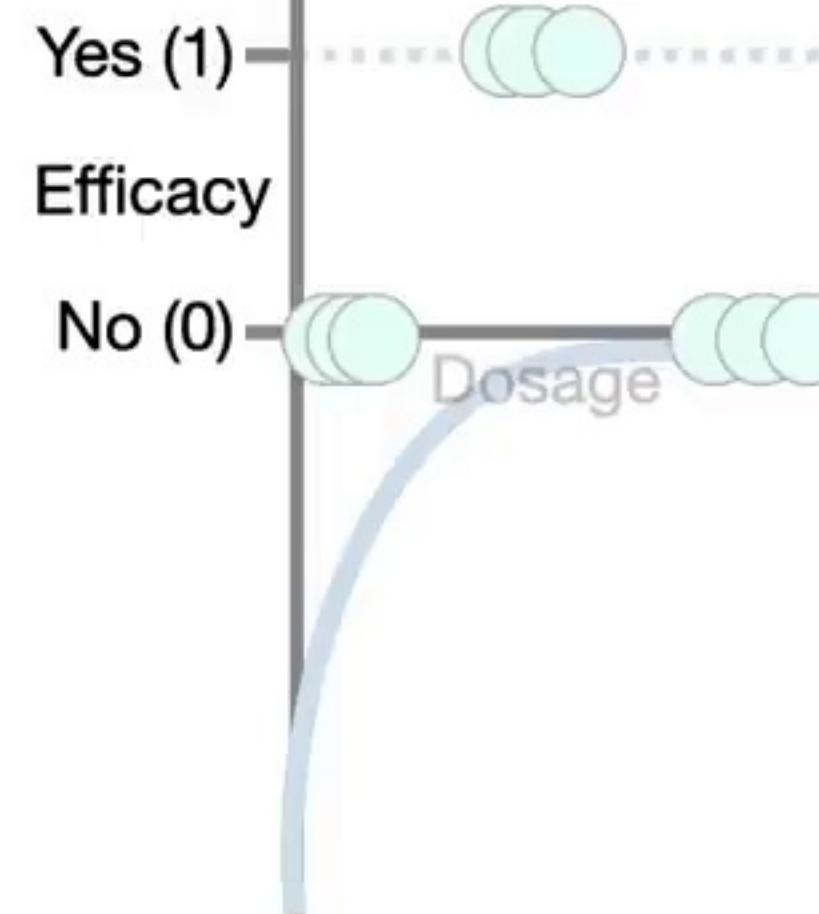
...to the bottom **Node** in the **Hidden Layer...**





(Dosage  $\times$  -2.52)

However, this time, we  
multiply the **Dosage**  
by **-2.52...**



Dosage

$\times -34.4$

$\times -2.52$

+ 1.29

+ 2.14

$\times -1.30$

$\times 2.28$

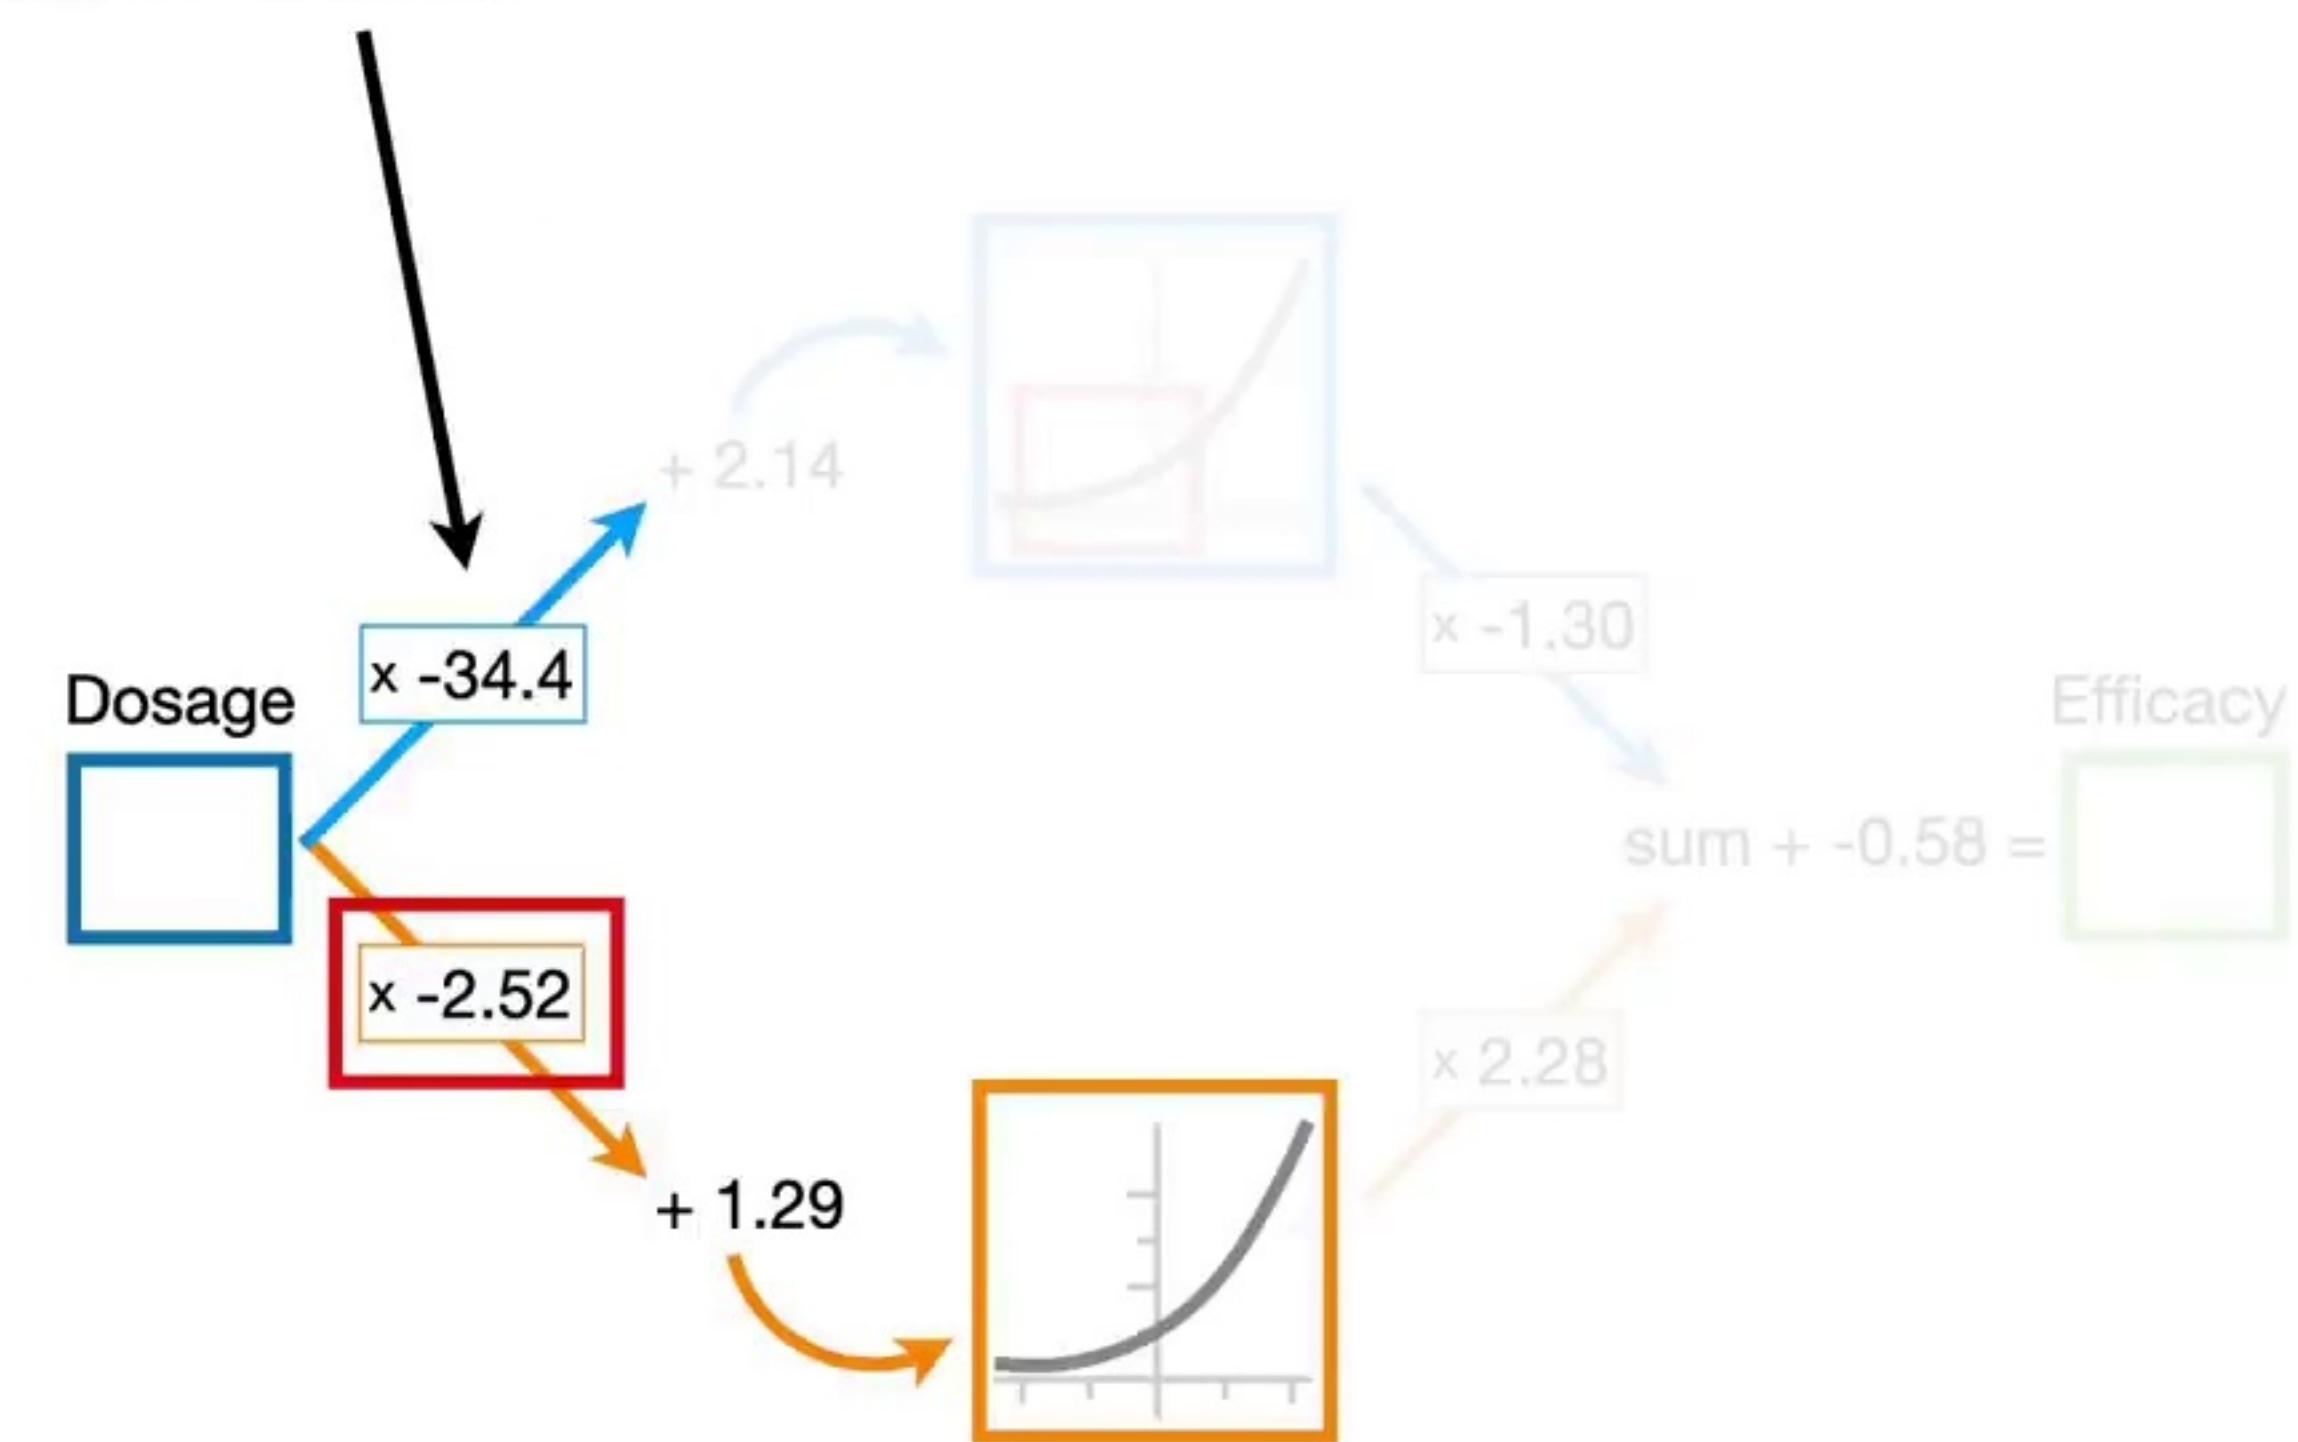
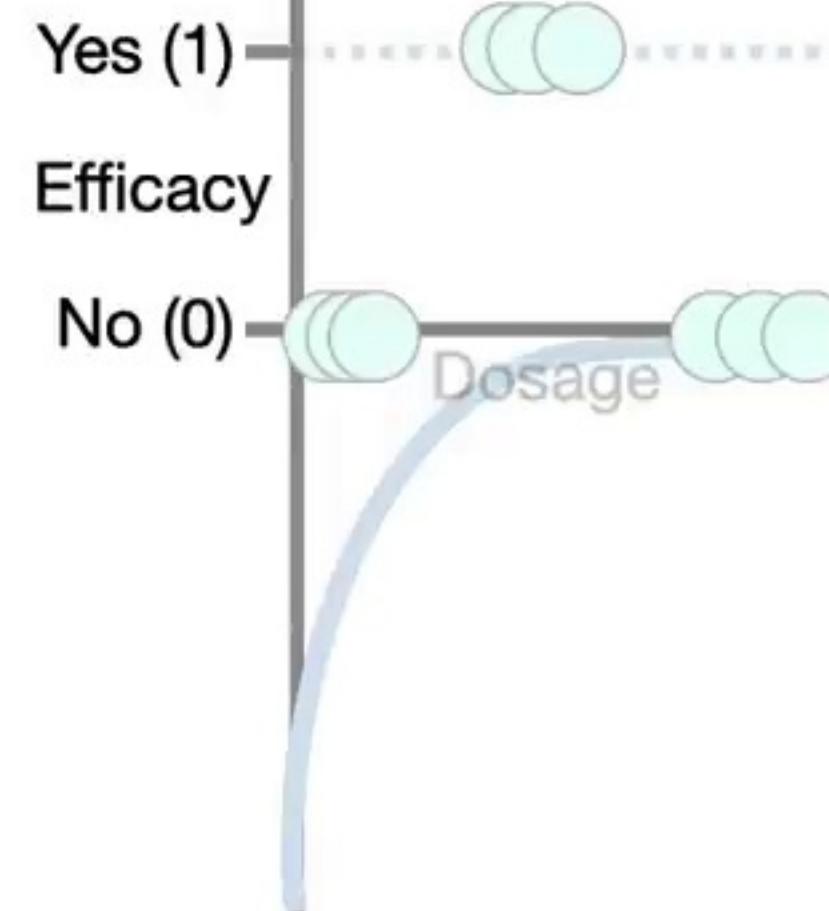
sum + -0.58 =





(Dosage  $\times$  -2.52)

... instead of -34.4...





$$(\text{Dosage} \times -2.52) + 1.29$$

...and we add 1.29...

Yes (1)  
Efficacy  
No (0)



Dosage

Dosage

$\times -2.52$

$+ 1.29$

$+ 2.14$

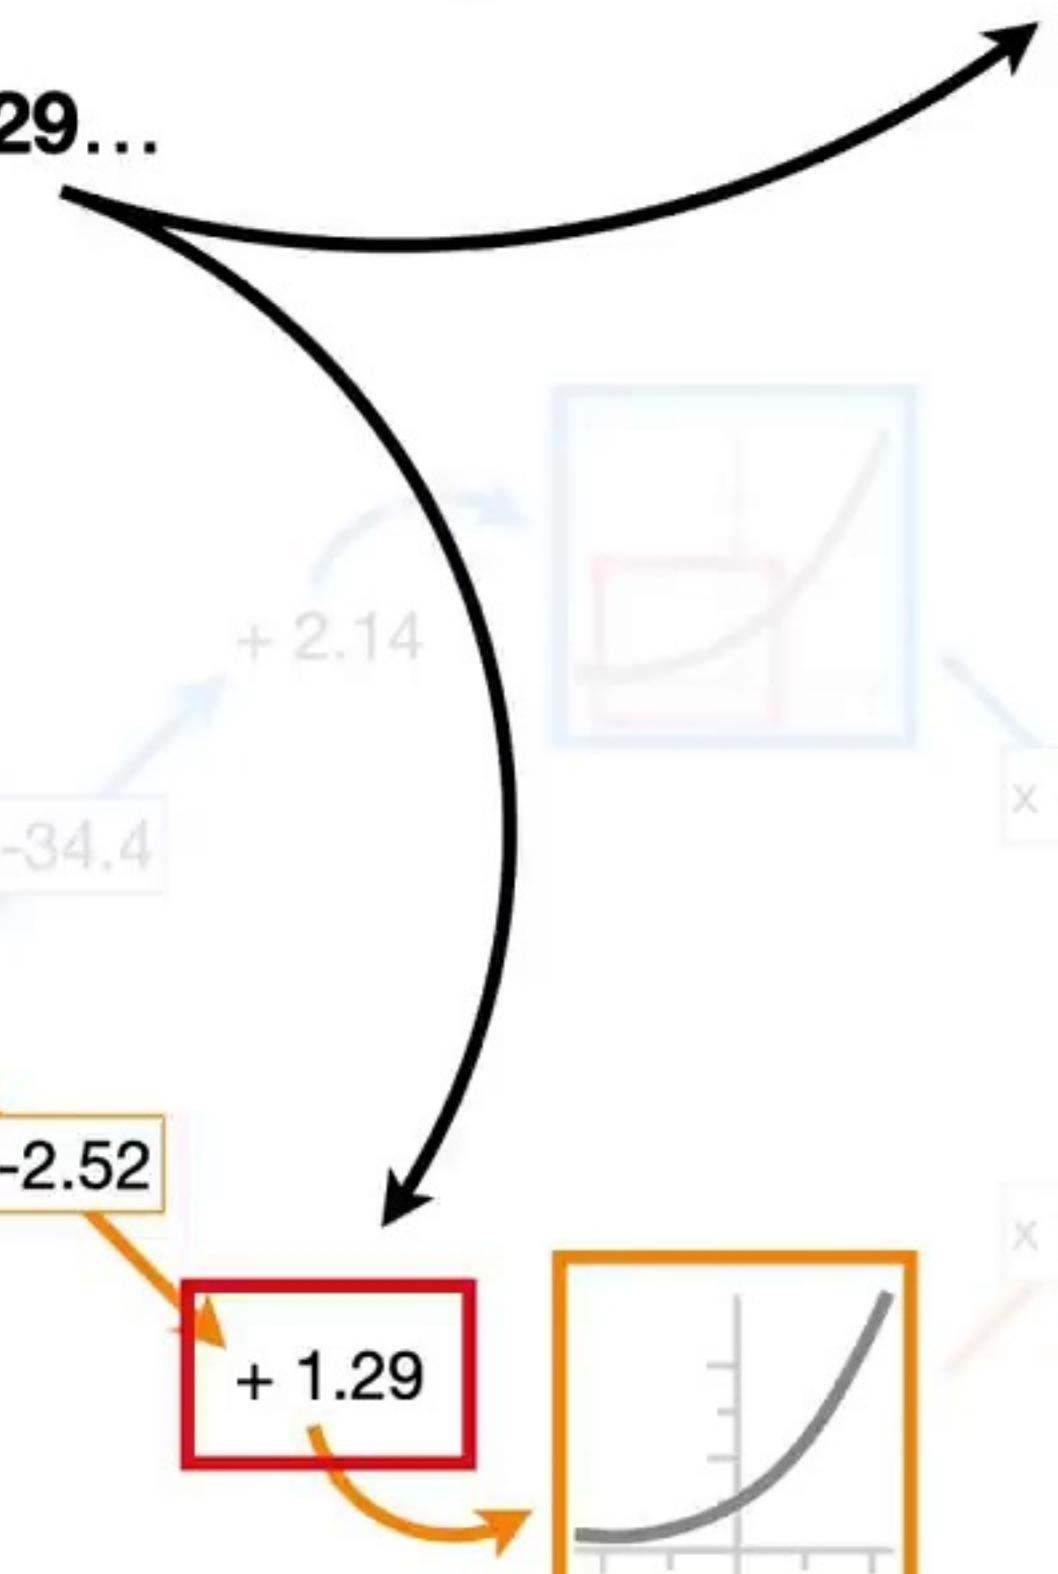
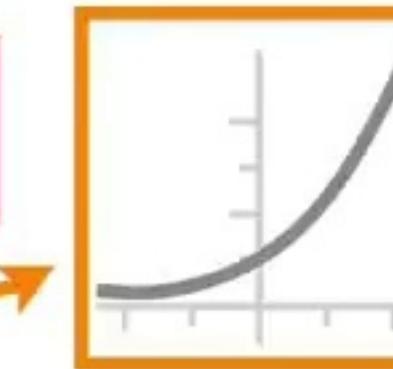
$\times -34.4$

$\times -1.30$

sum  $+ -0.58 =$

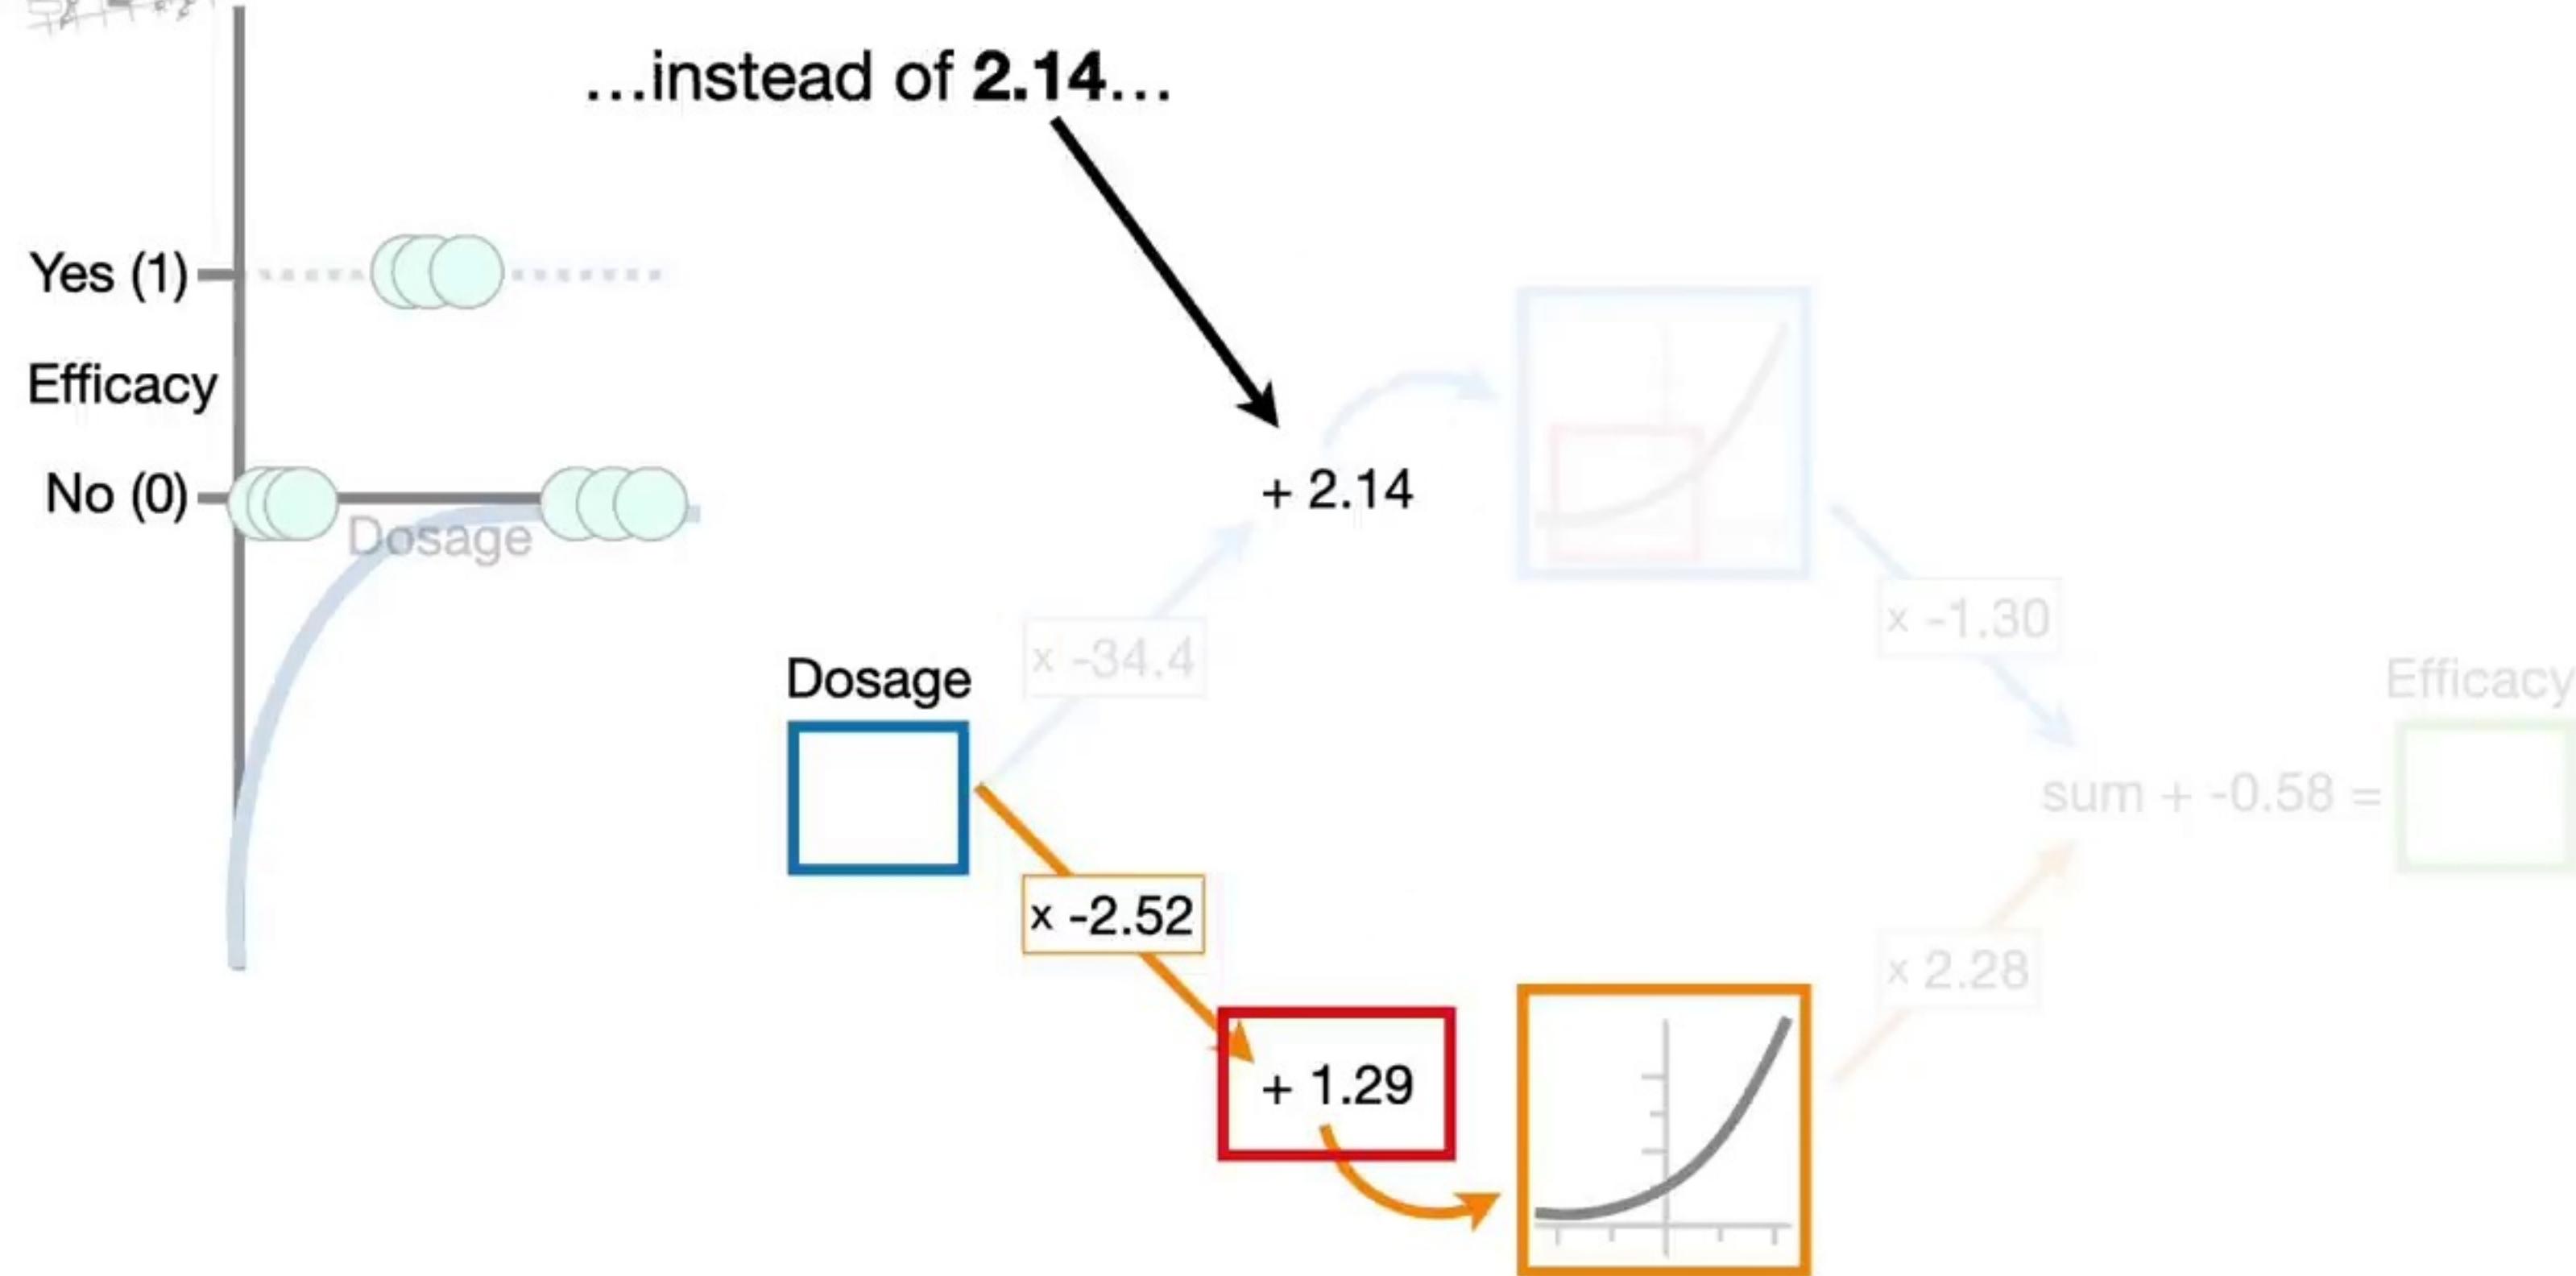
Efficacy

$\times 2.28$



Double  
BAM!!  
**SO!**

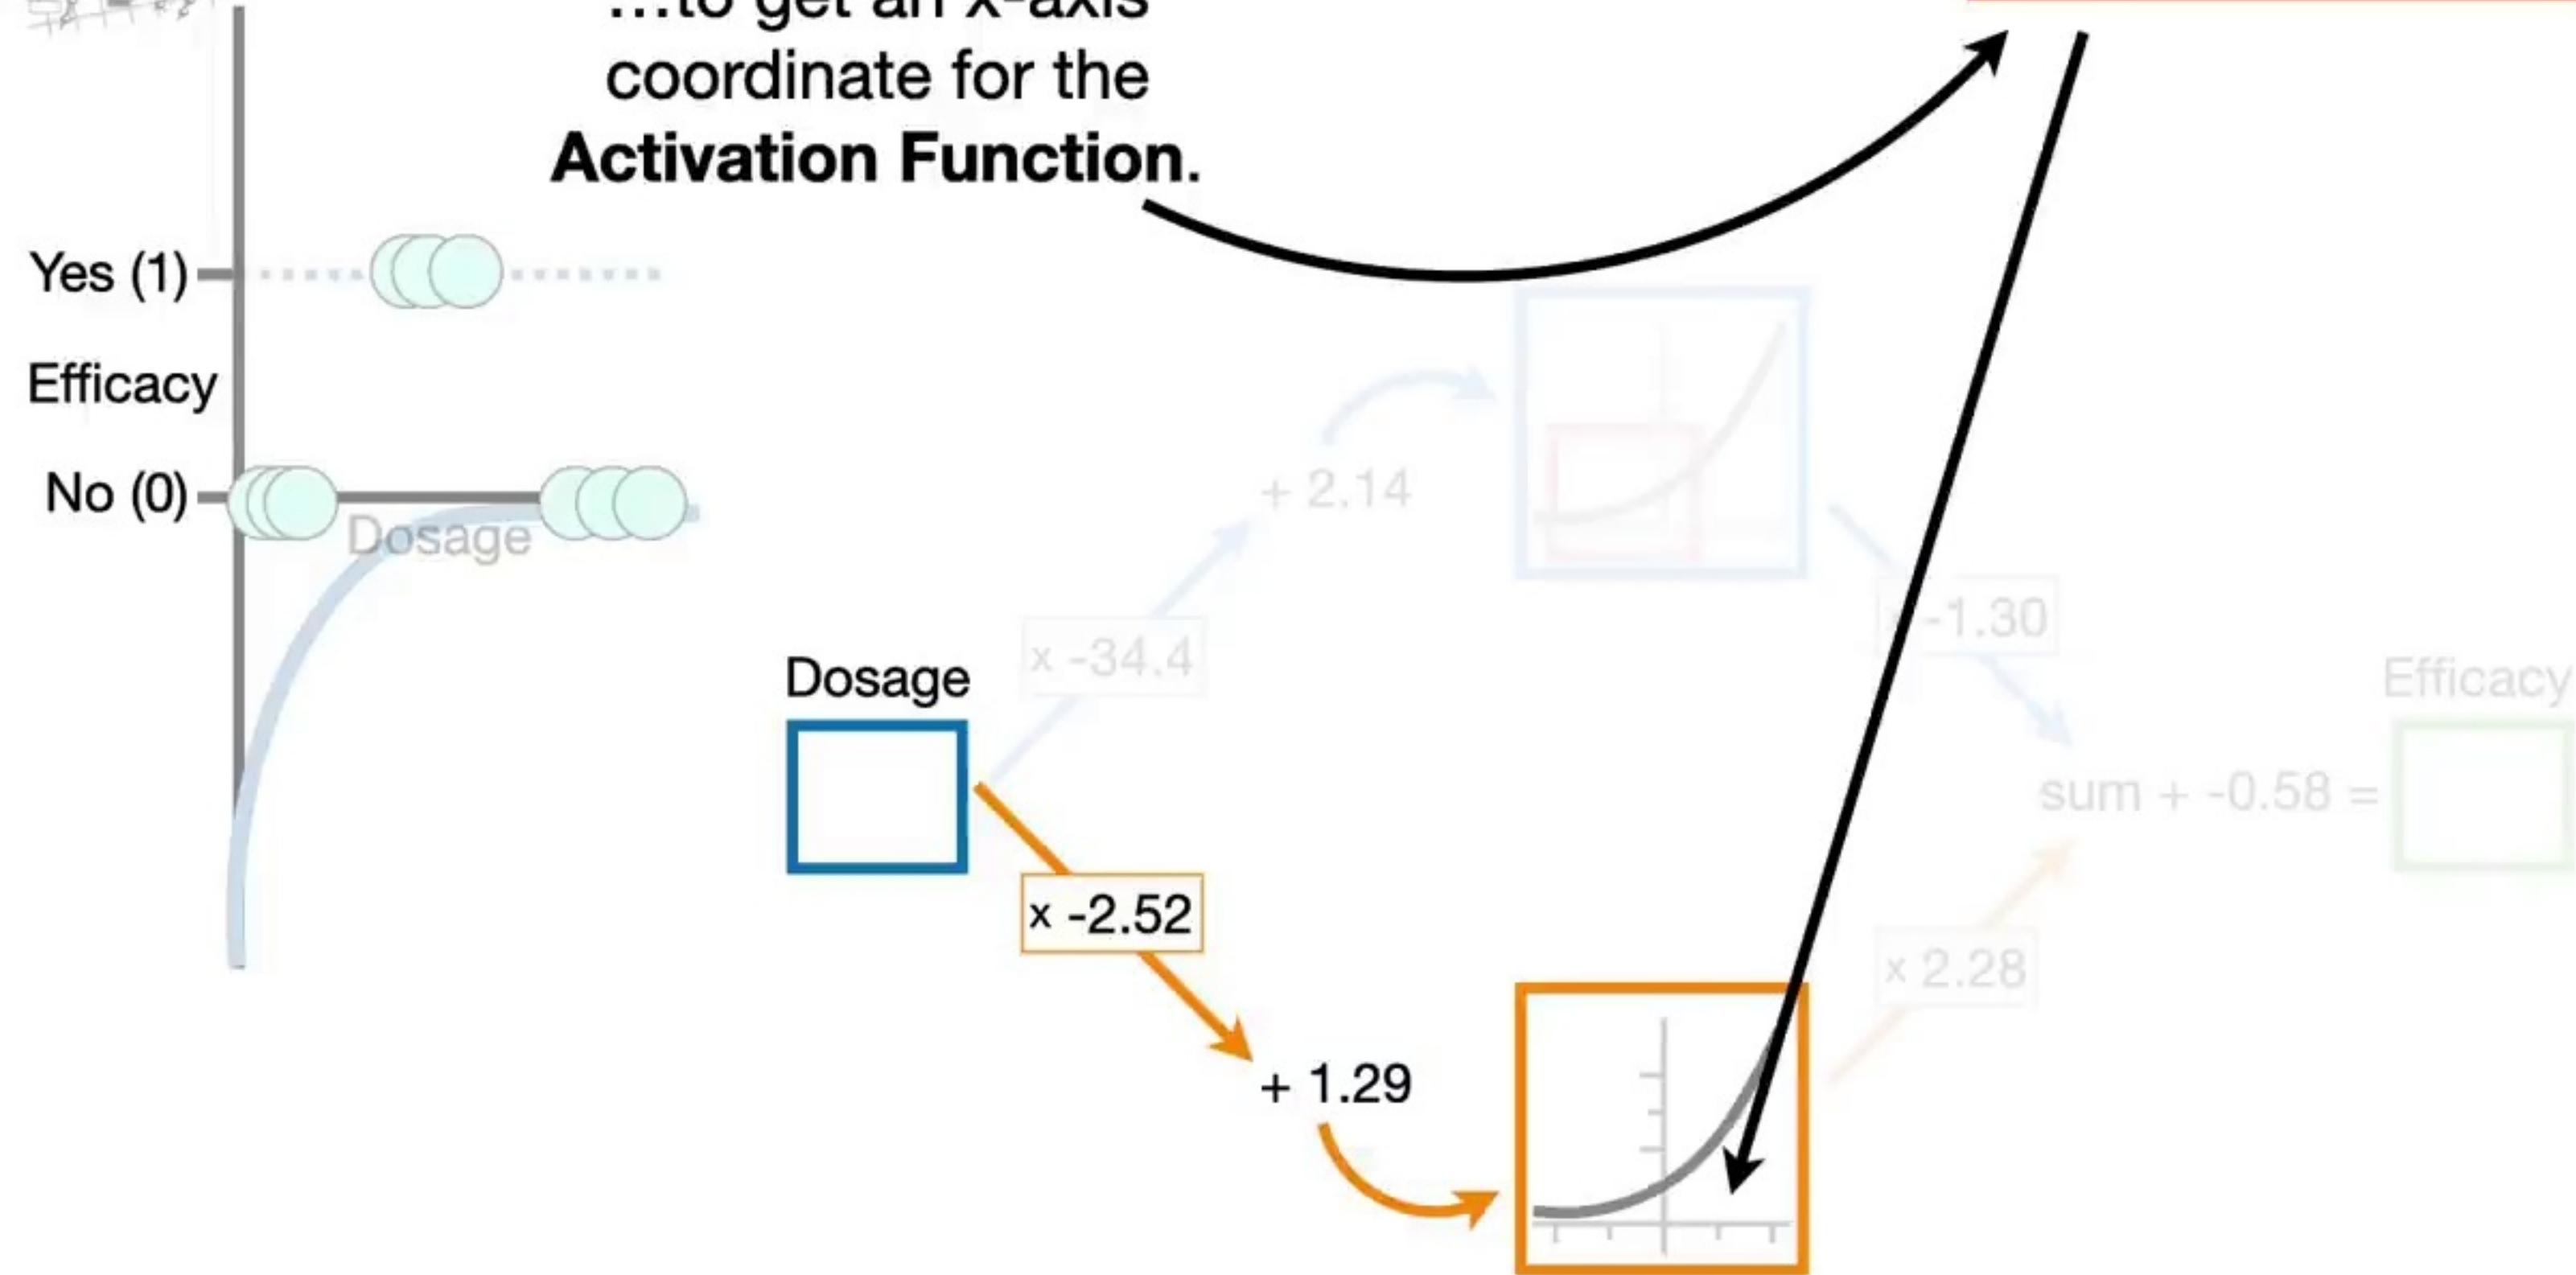
$$(\text{Dosage} \times -2.52) + 1.29$$



Double  
BAM!!  
**SQ!**

$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

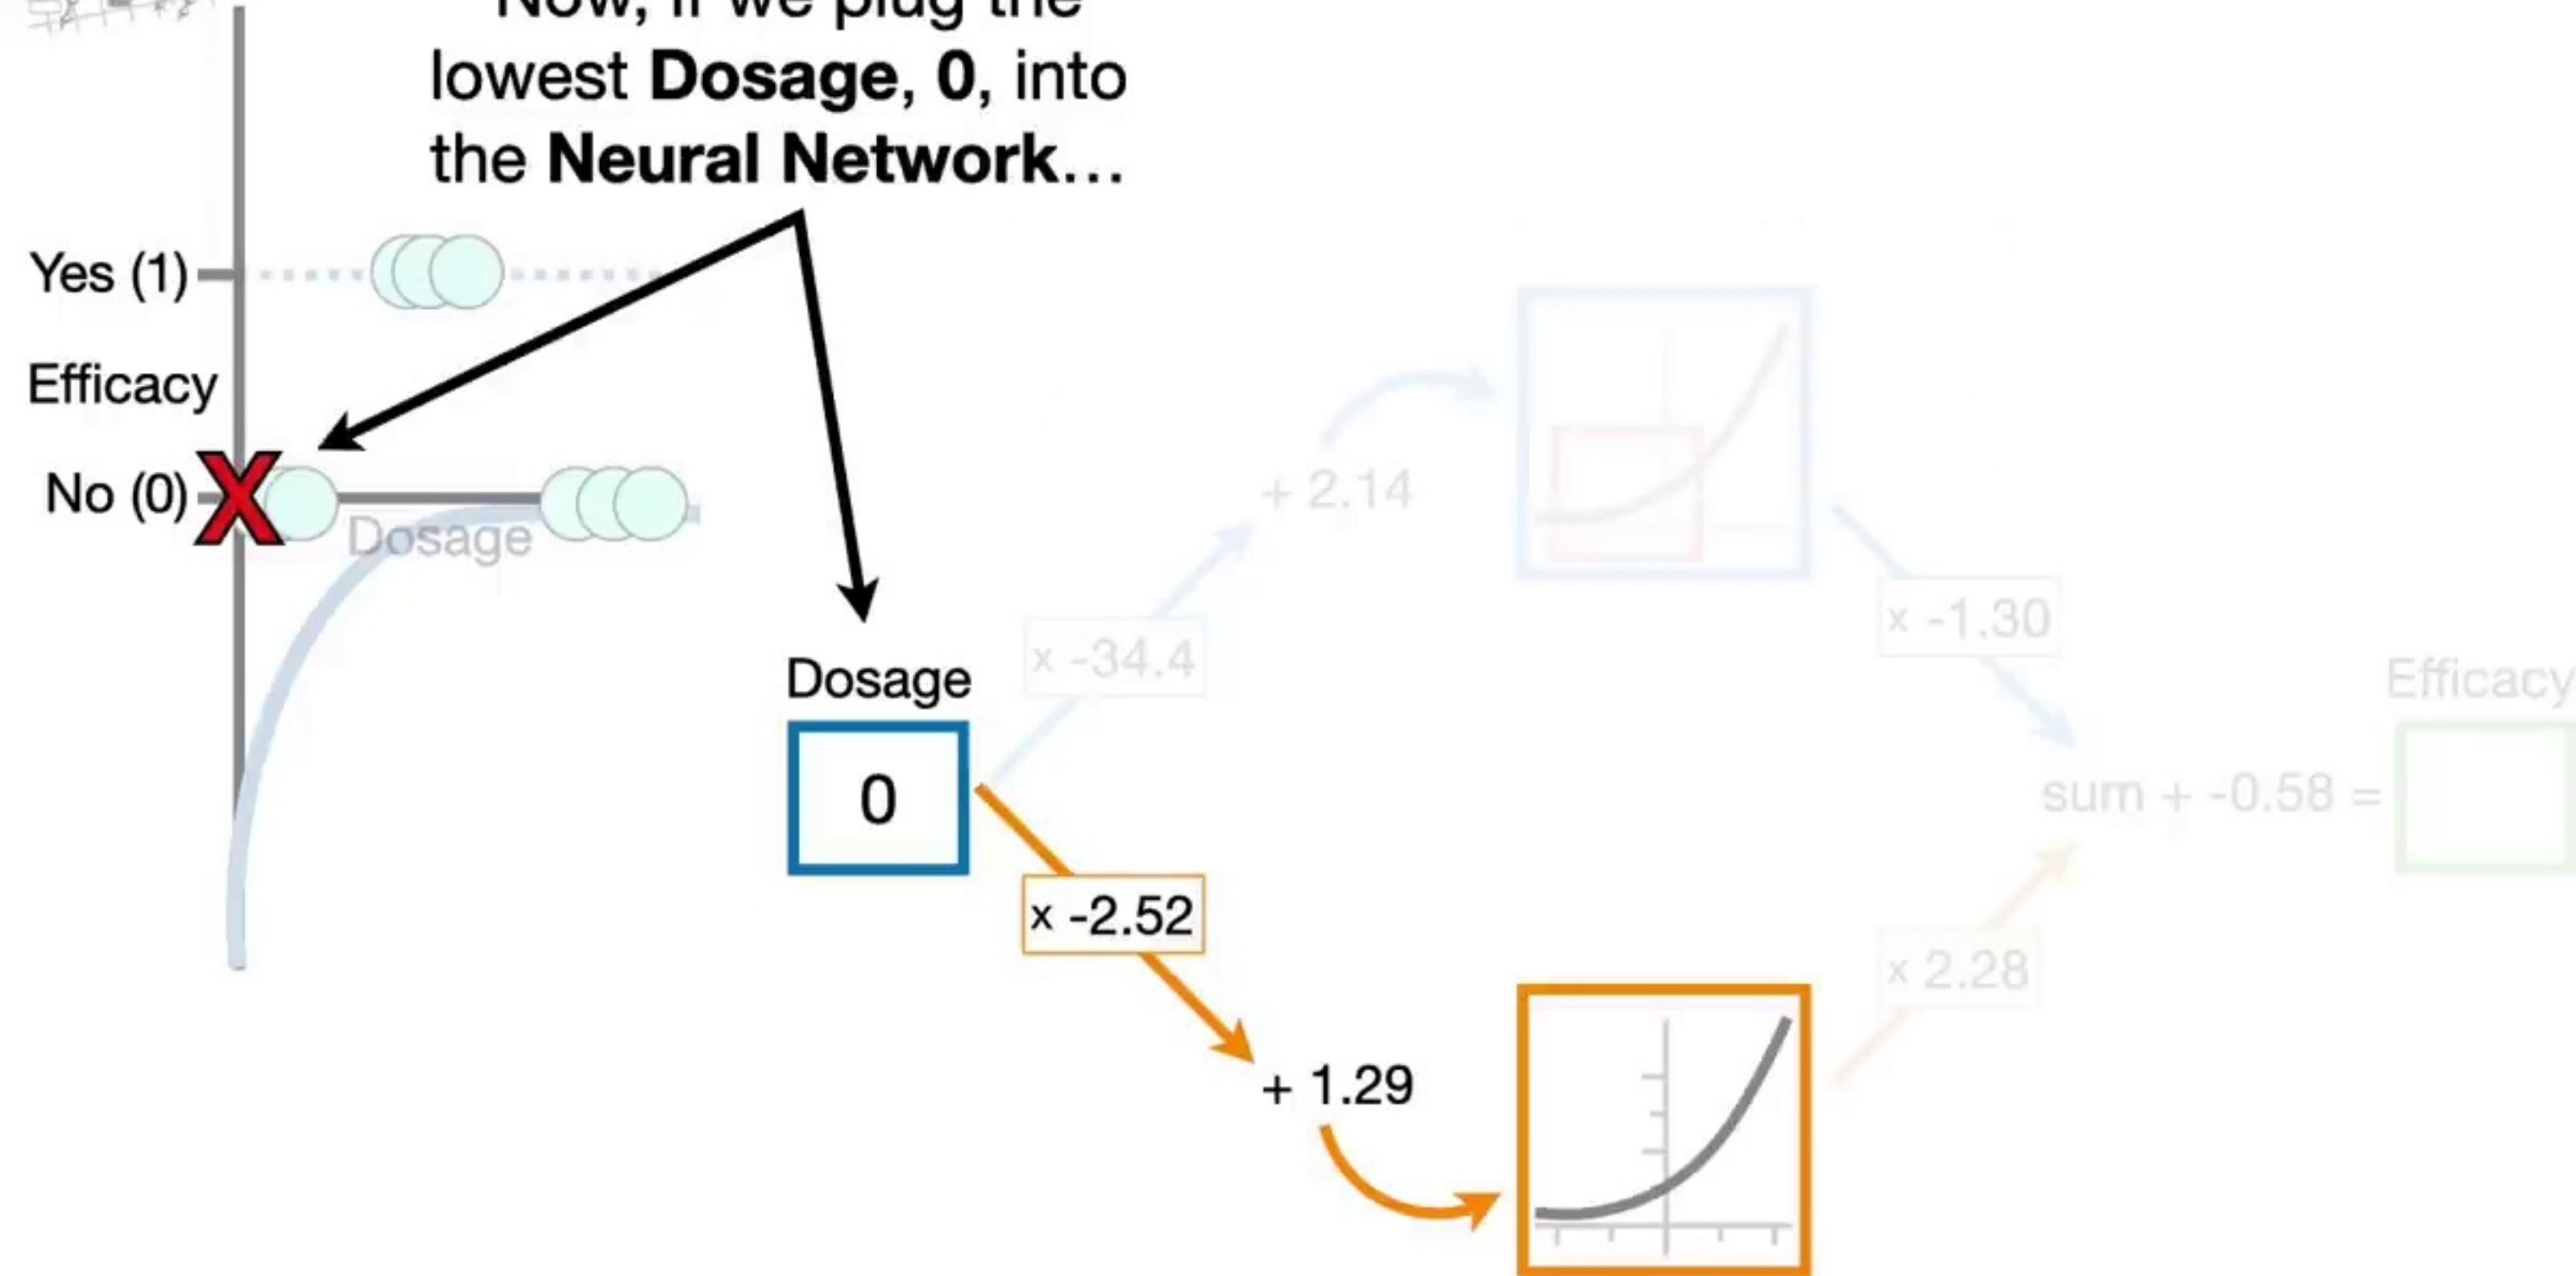
...to get an x-axis  
coordinate for the  
**Activation Function.**





$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

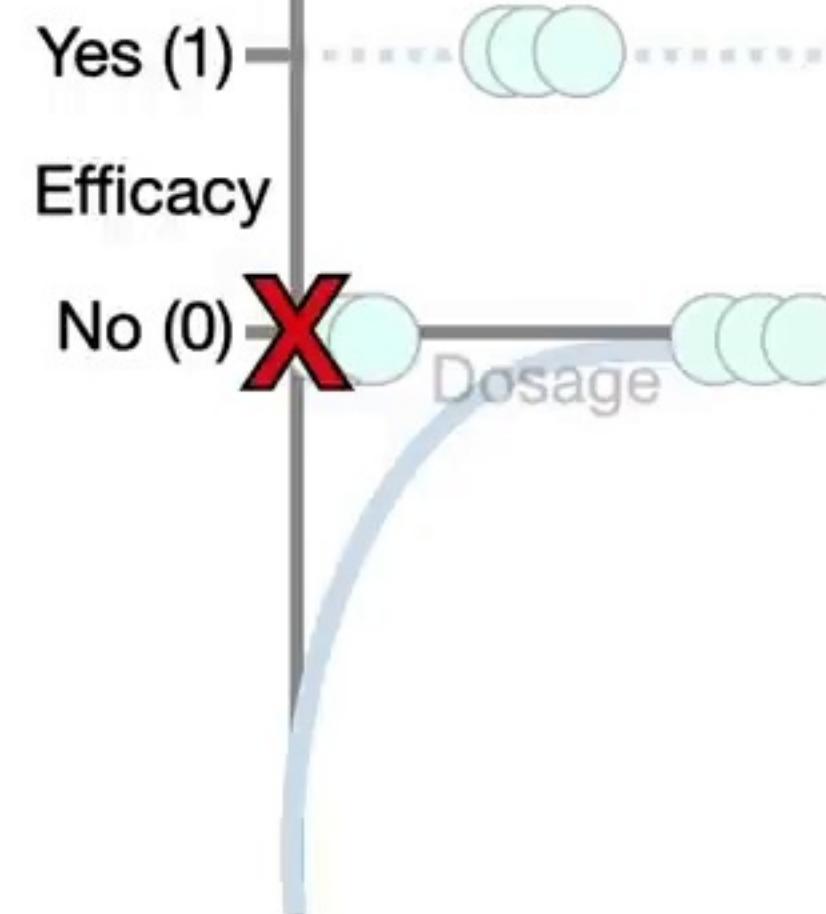
Now, if we plug the lowest **Dosage**, 0, into the **Neural Network**...





$$(0 \times -2.52) + 1.29 = 1.29$$

...then the x-axis coordinate  
for the **Activation Function**  
is **1.29**.



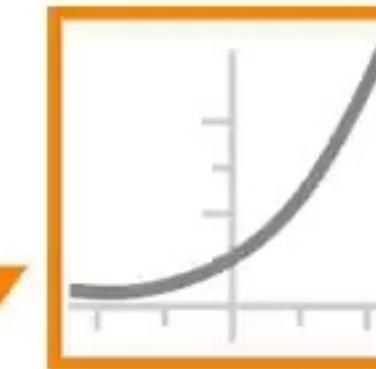
Dosage

**0**

**x -2.52**

**x -34.4**

**+ 1.29**



**x 2.28**

**x -1.30**

sum + -0.58 =

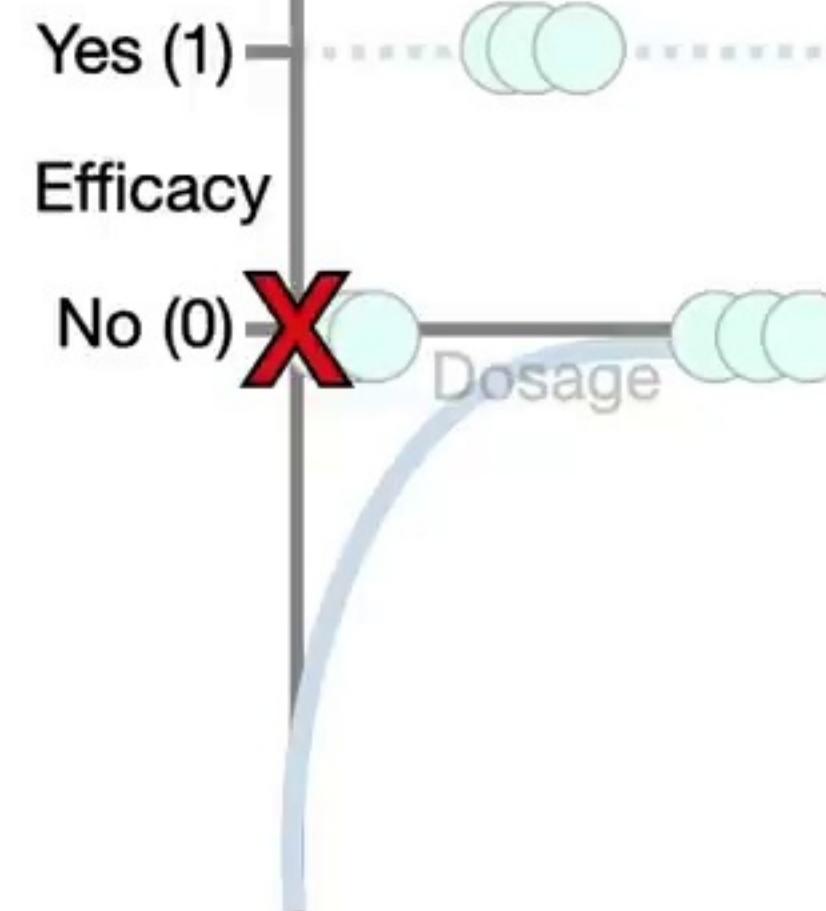
Efficacy

**[ ]**



$$(0 \times -2.52) + 1.29 = 1.29$$

...then the x-axis coordinate  
for the **Activation Function**  
is **1.29**.



Dosage

0

x -2.52

x -34.4

+ 1.29

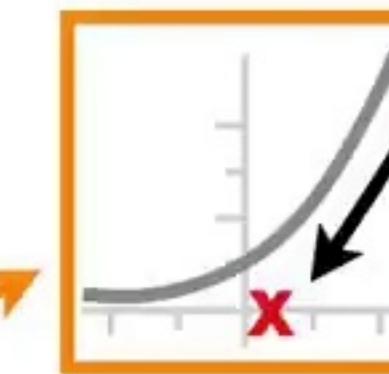
+ 2.14

x -1.00

sum + -0.58 =

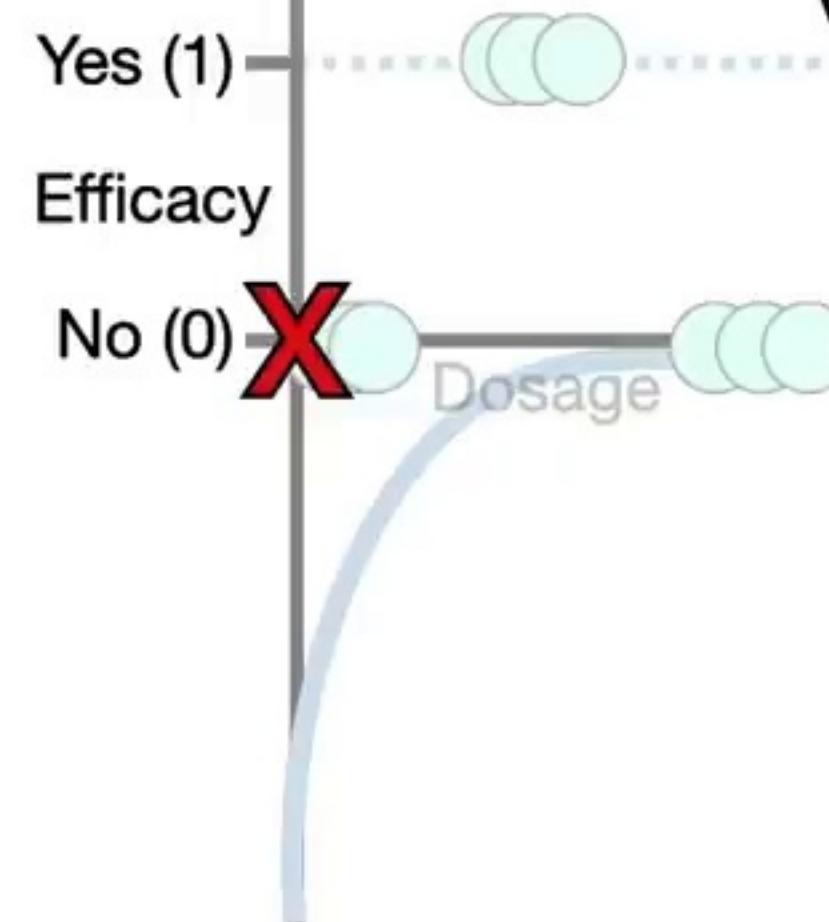
x 2.28

Efficacy





Now we plug **1.29** into the **Activation Function** to get the corresponding y-axis value...



$$(0 \times -2.52) + 1.29 = 1.29$$

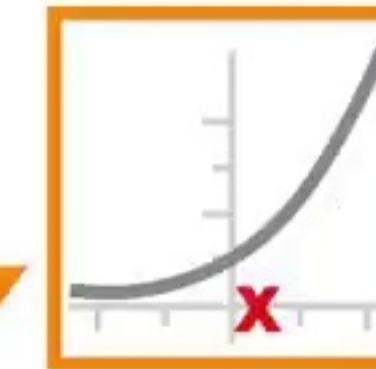
$$f(x) = \log(1 + e^x)$$

Dosage

0

x -2.52

+ 1.29



x -34.4

+ 2.14

x -1.30

sum + -0.58 =

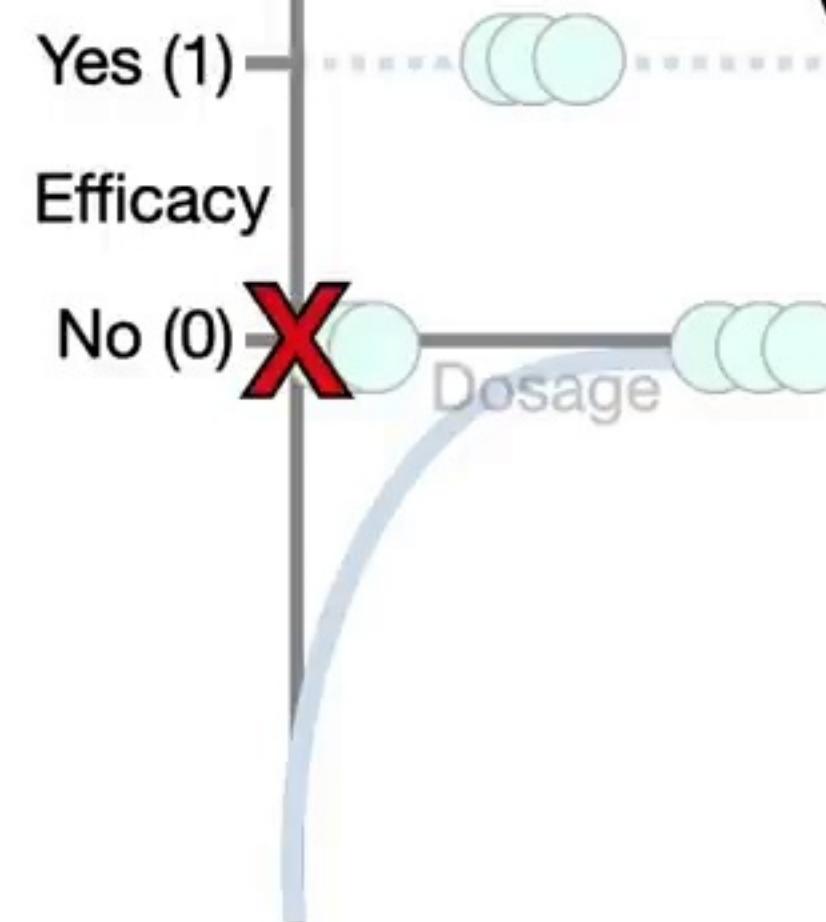


Efficacy

x 2.28



Now we plug **1.29** into the **Activation Function** to get the corresponding y-axis value...



Dosage

$x - 2.52$

$+ 1.29$



$$(0 \times -2.52) + 1.29 = 1.29$$

$$f(1.29) = \log(1 + e^{1.29})$$



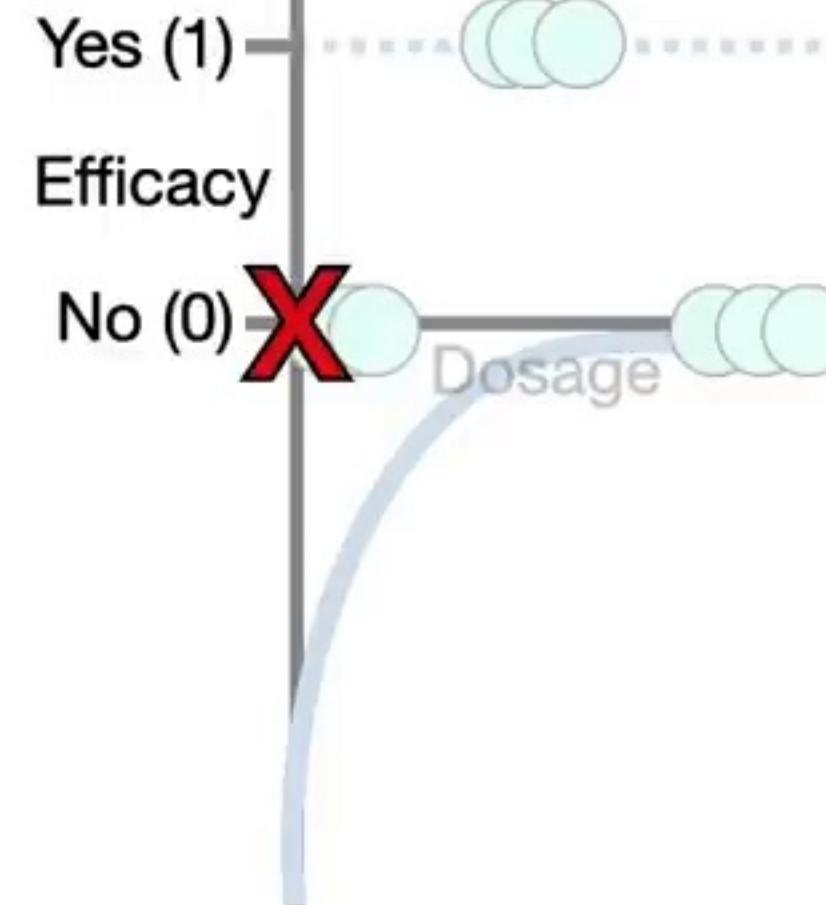
Efficacy

Double  
BAM!!  
**SQ!**

$$(0 \times -2.52) + 1.29 = 1.29$$

...and get 1.53...

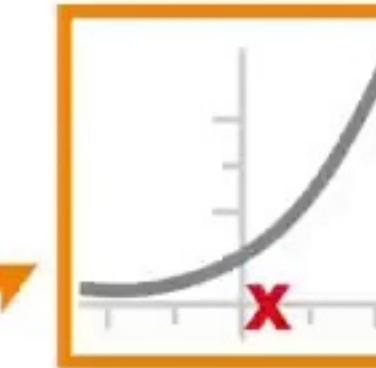
$$f(1.29) = \log(1 + e^{1.29}) = 1.53$$



Dosage

$x - 2.52$

$+ 1.29$



$x - 34.4$

$+ 2.14$

$x - 1.30$

sum + -0.58 =

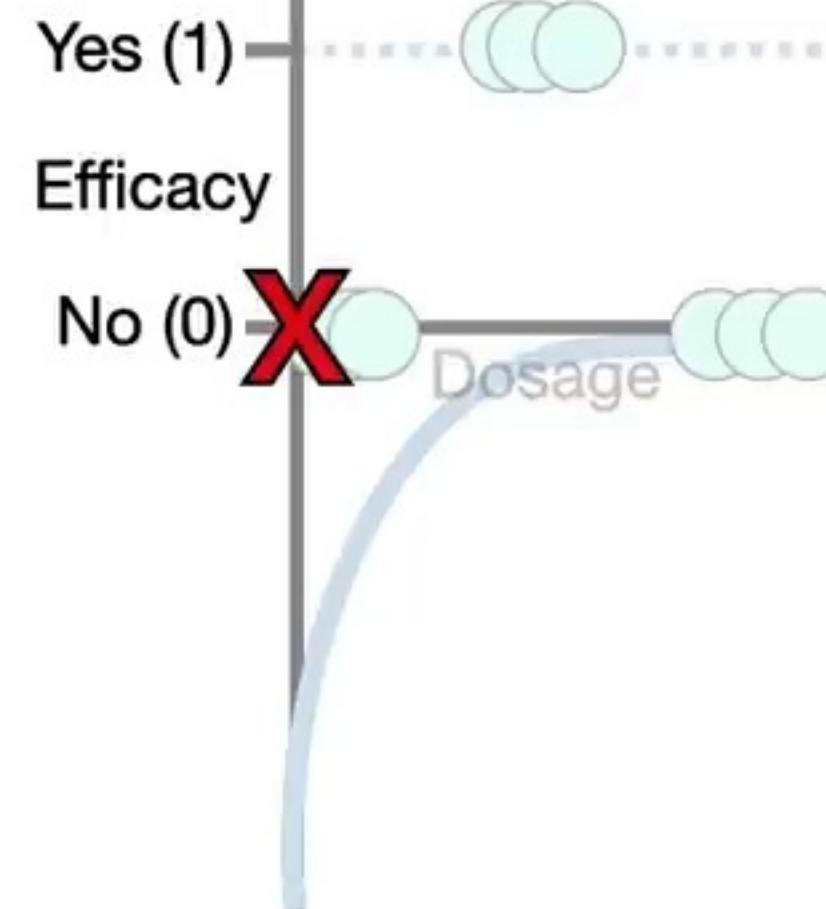
$x 2.28$

Double  
BAM!!  
**SQ!**

$$(0 \times -2.52) + 1.29 = 1.29$$

...and get 1.53...

$$f(1.29) = \log(1 + e^{1.29}) = 1.53$$

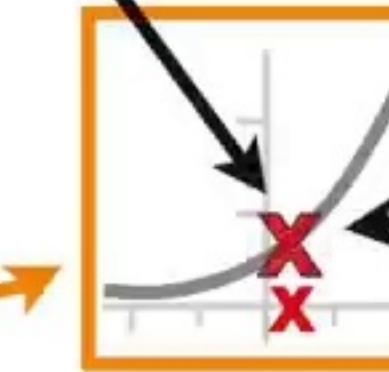


Dosage

**0**

**x -2.52**

+ 1.29



**x -1.30**

sum + -0.58 =

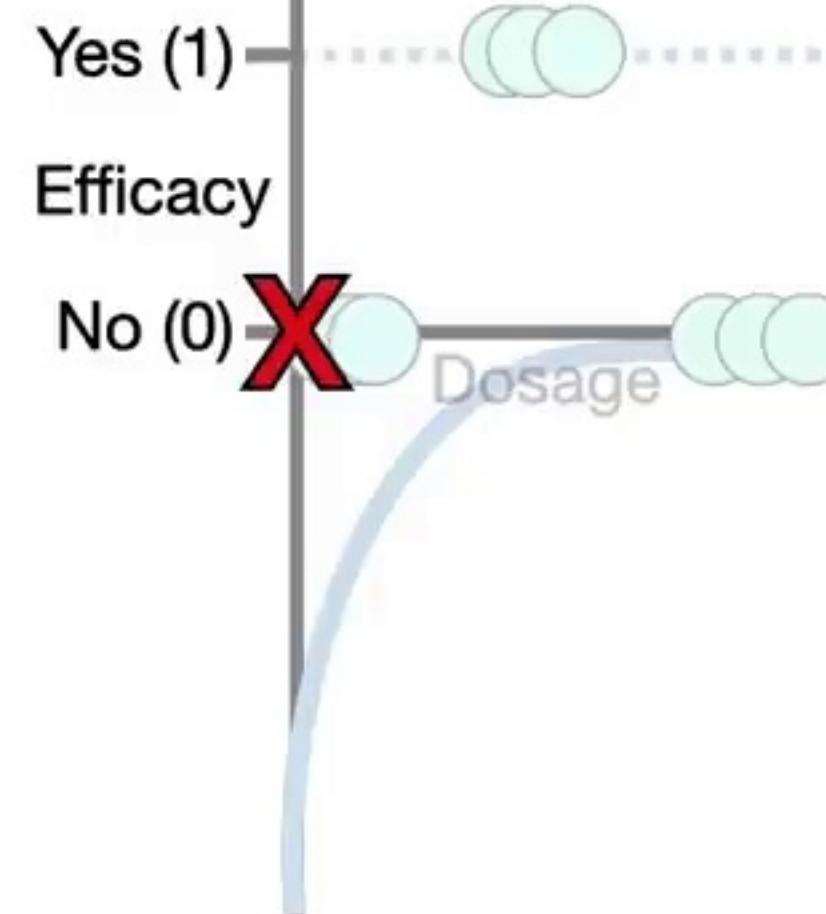
**x 2.28**

Efficacy



$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

Now we just plug in **Dosage** values from **0** to **1** to get the corresponding y-axis values...



Dosage

$x - 2.52$

+ 1.29

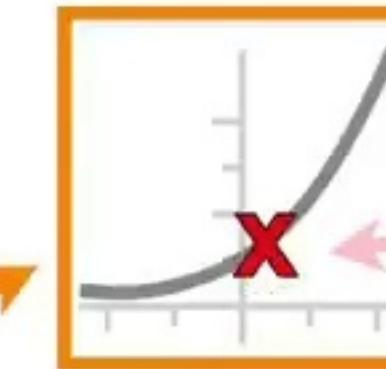
$x - 34.4$

+ 2.14

$x - 1.30$

sum + -0.58 =

Efficacy

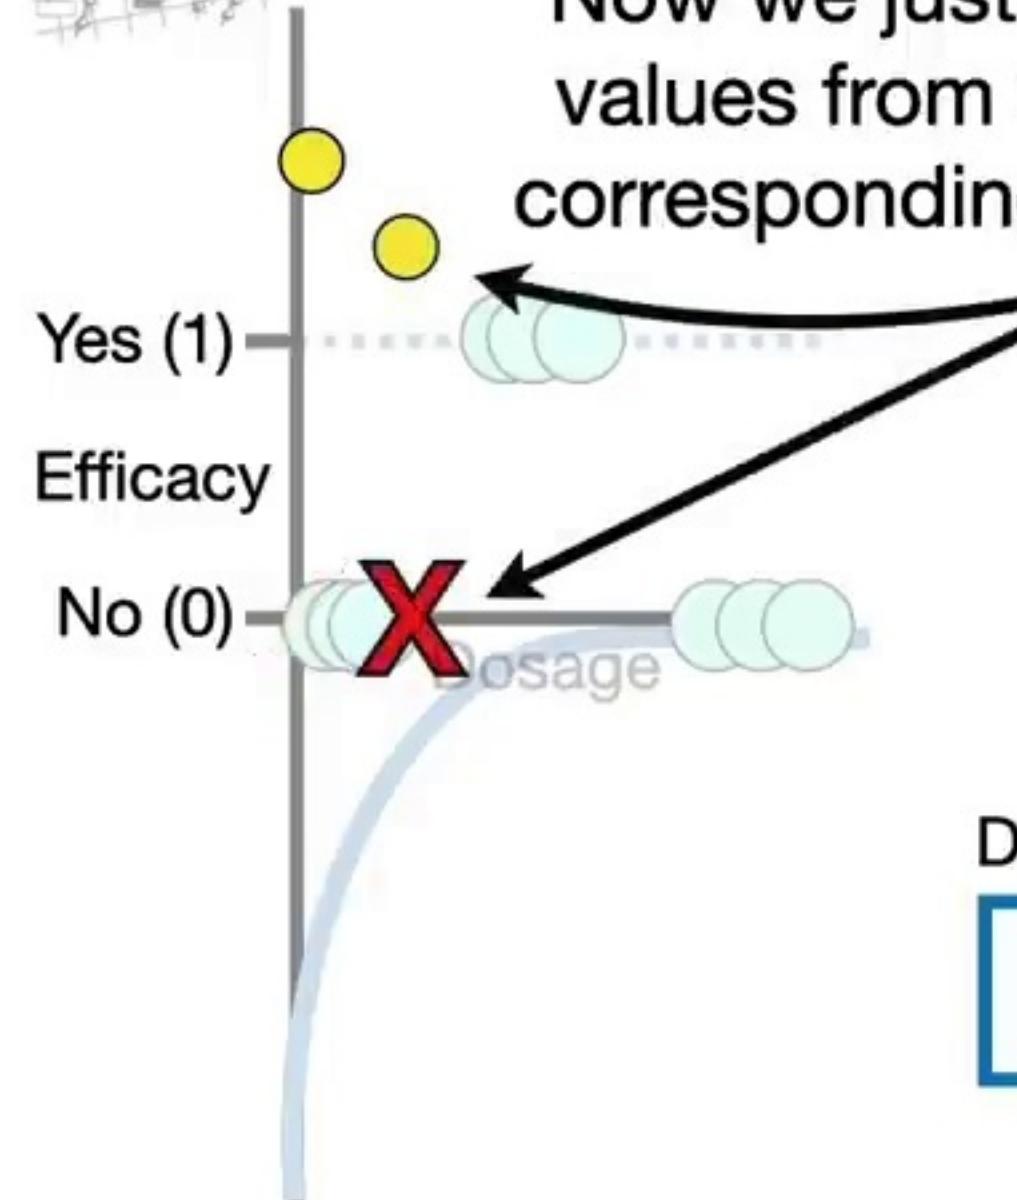


$$f(x) = \log(1 + e^x)$$

Double  
BAM!!  
**SO!**

$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

Now we just plug in **Dosage** values from **0** to **1** to get the corresponding y-axis values...



Dosage

0.2

x -2.52

+ 1.29

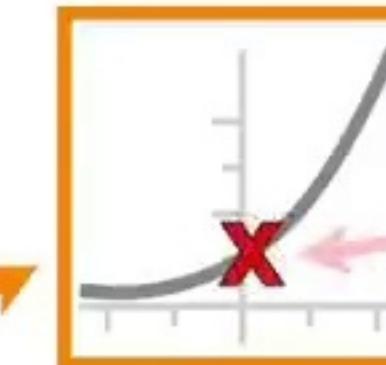
x -34.4

+ 2.14

x -1.30

sum + -0.58 =

Efficacy

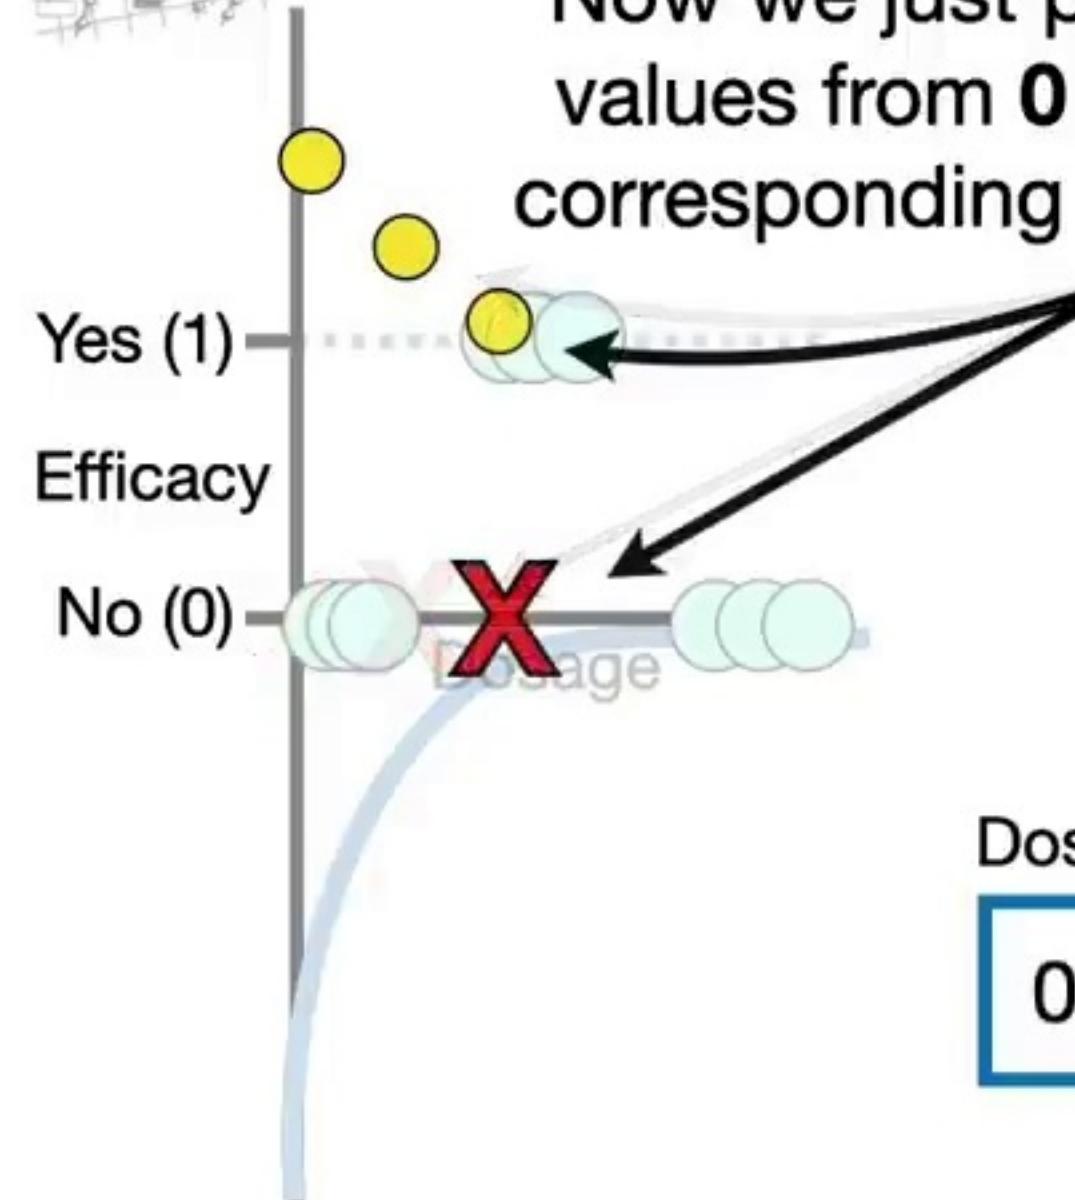


$$f(x) = \log(1 + e^x)$$



$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

Now we just plug in **Dosage** values from **0** to **1** to get the corresponding y-axis values...



Dosage

0.4

x -2.52

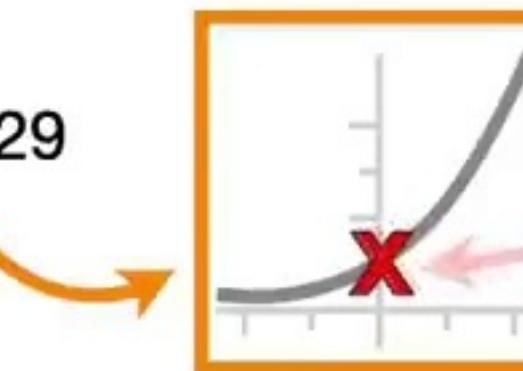
+ 1.29

$$x - 34.4 + 2.14$$

x -1.30

sum + -0.58 =

Efficacy

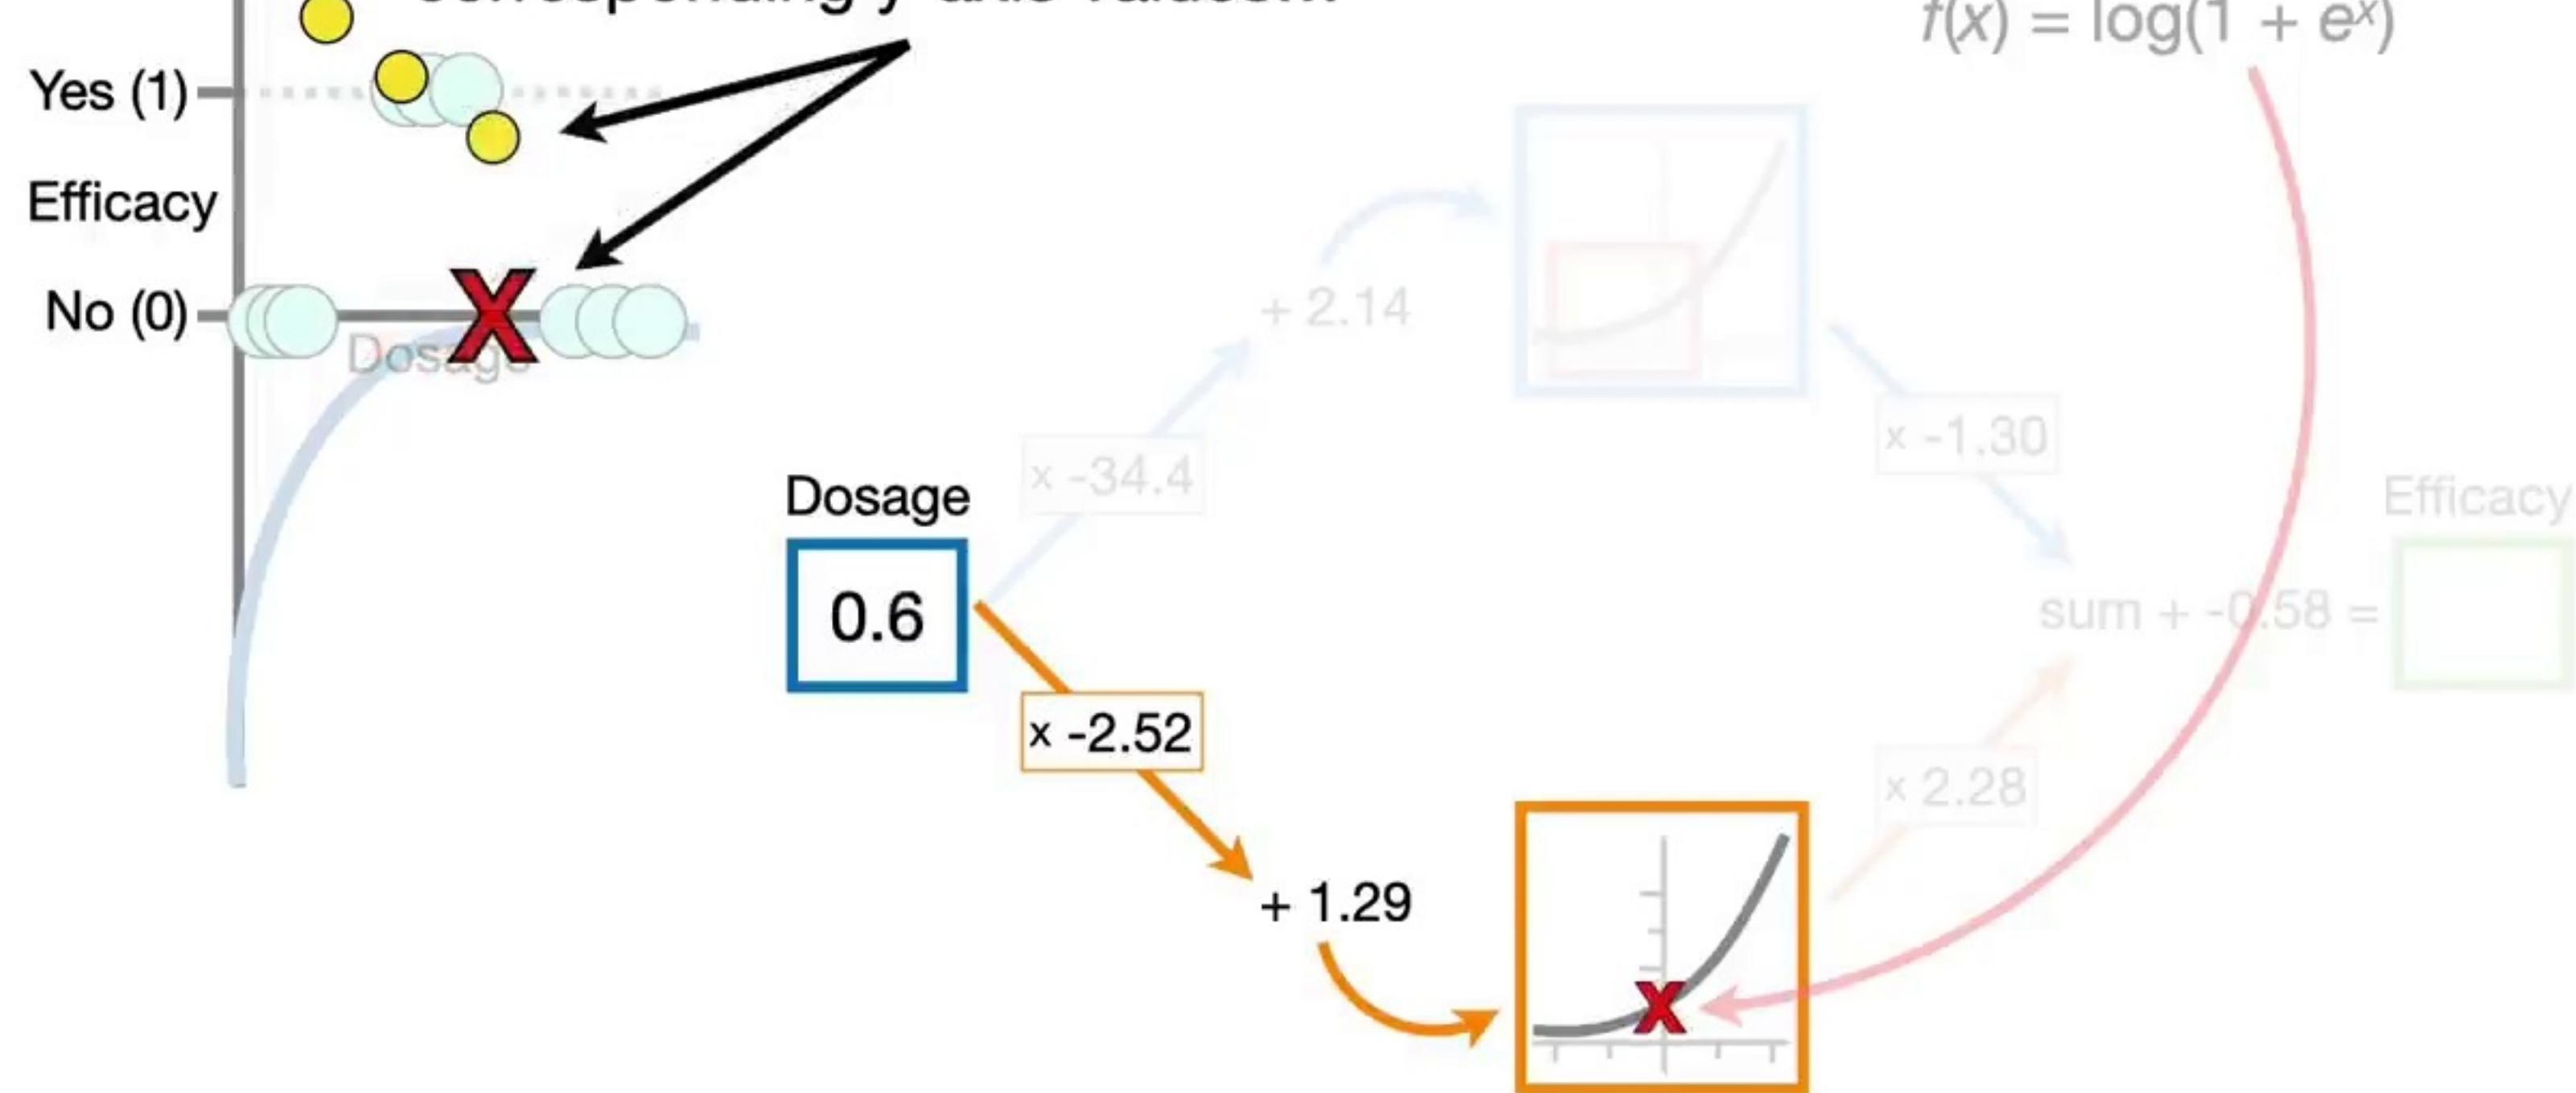


$$f(x) = \log(1 + e^x)$$



$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

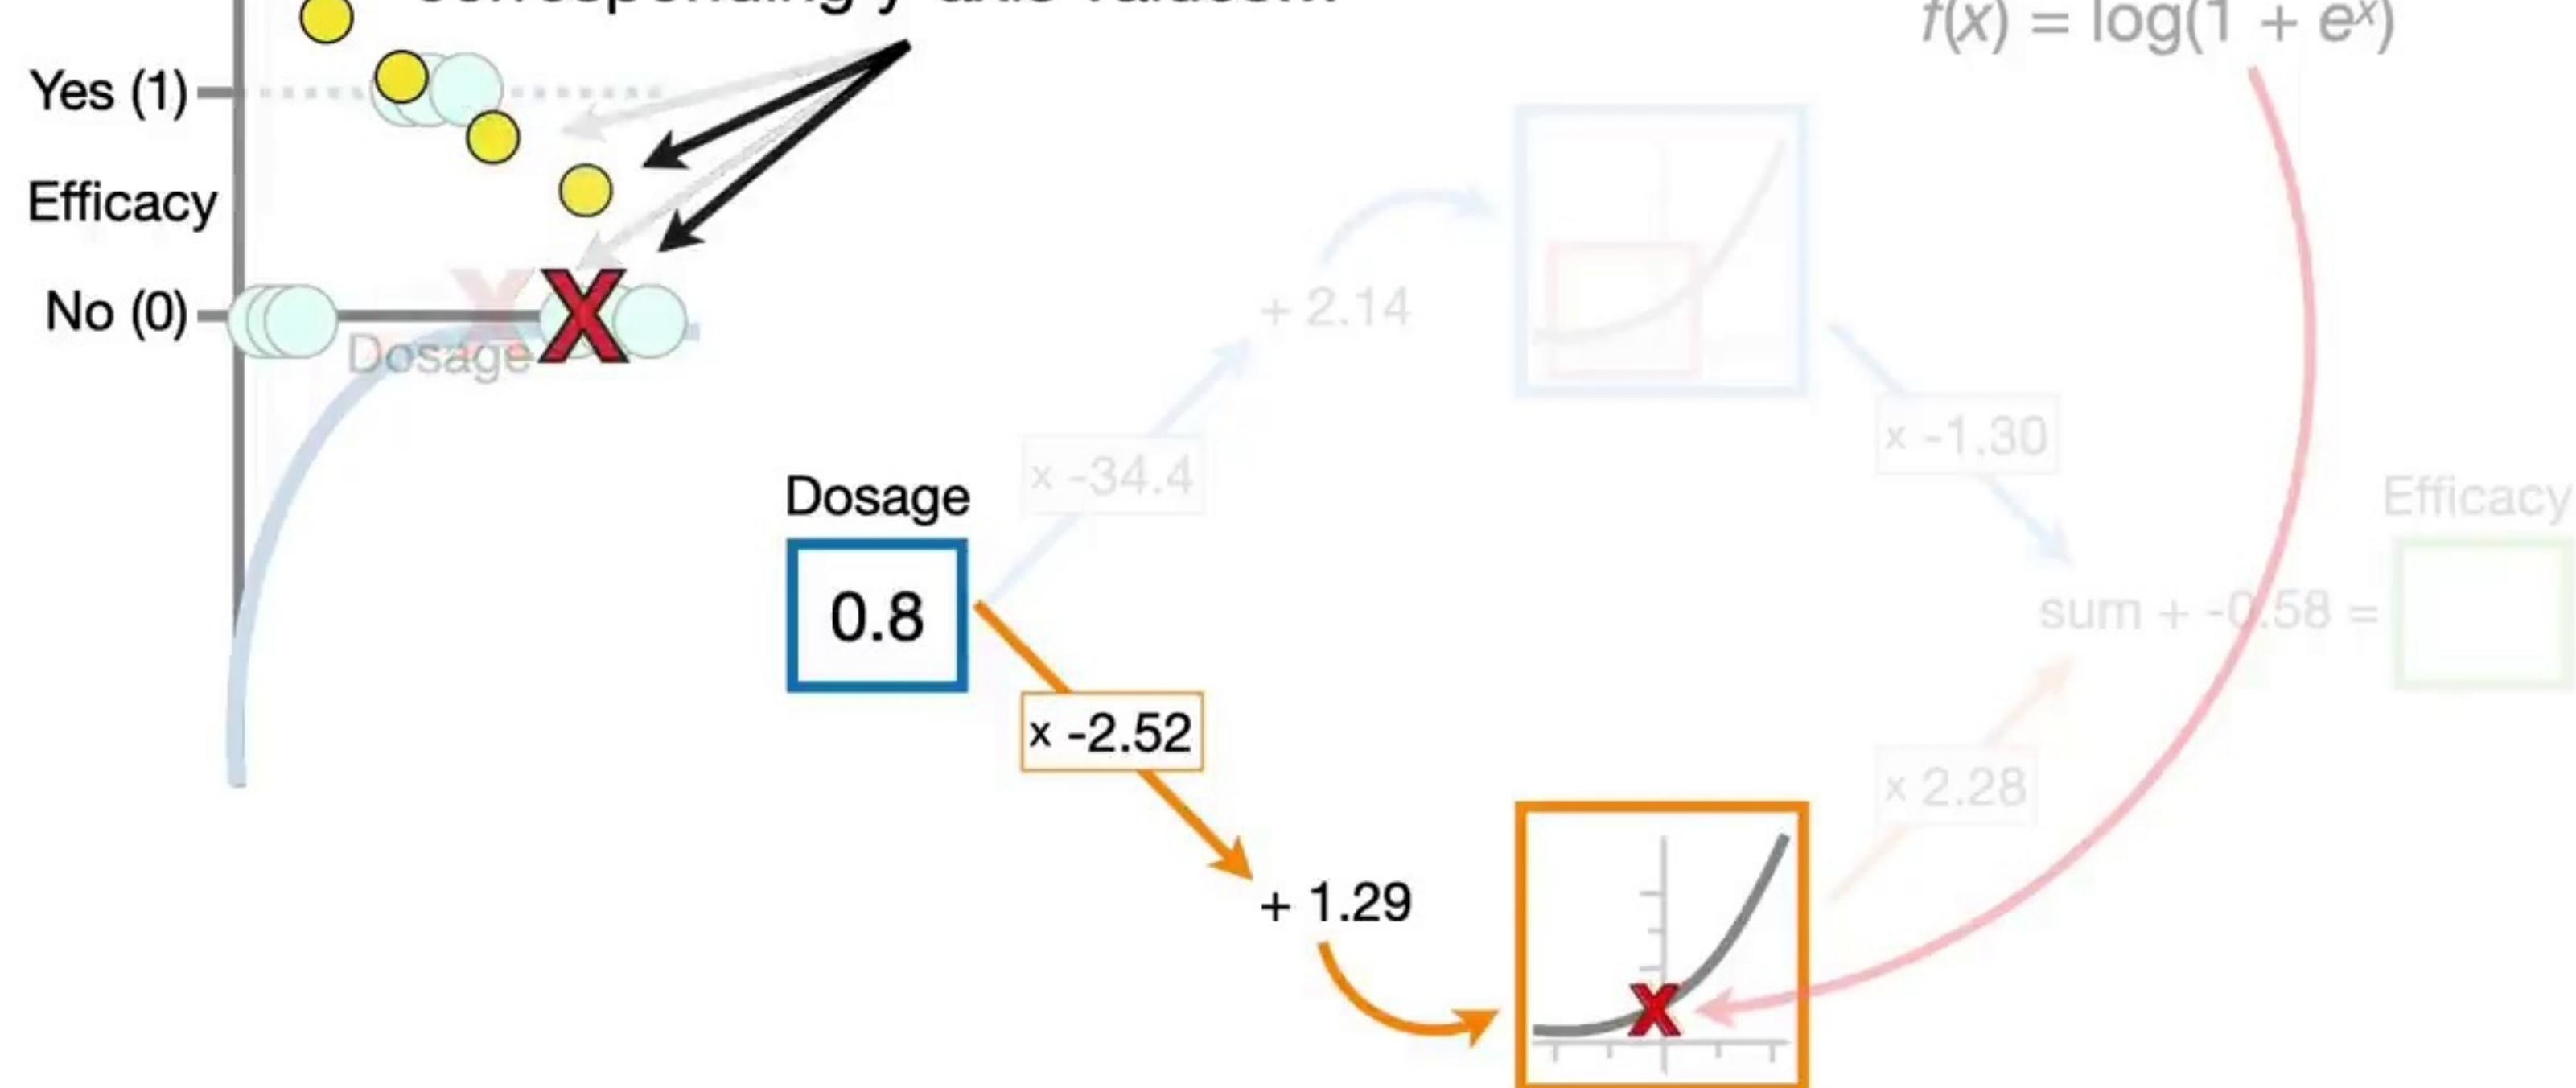
Now we just plug in **Dosage** values from **0** to **1** to get the corresponding y-axis values...





$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

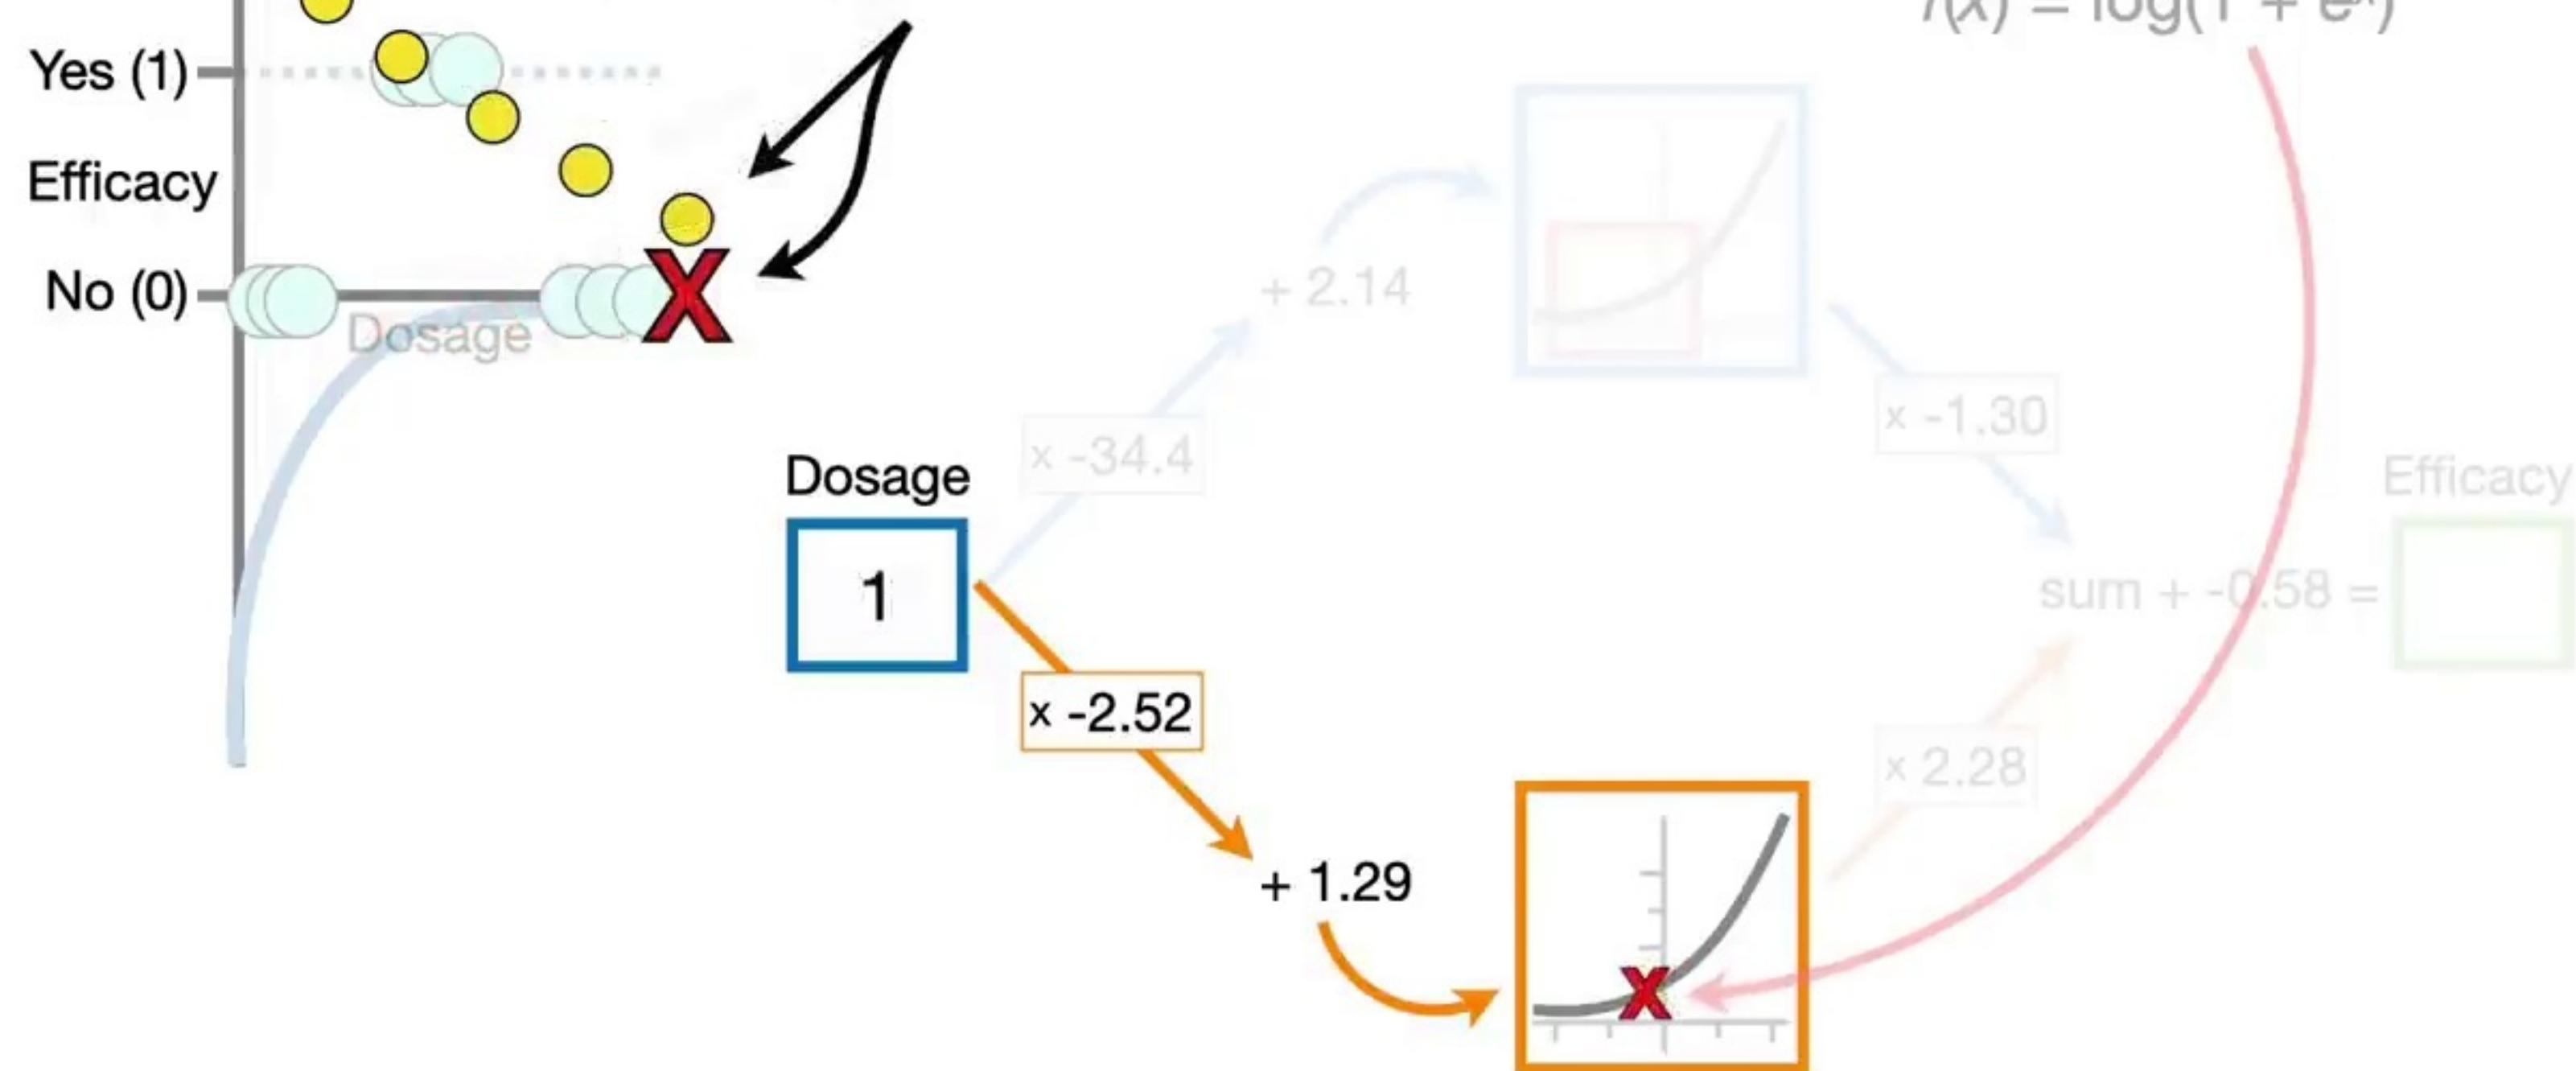
Now we just plug in **Dosage** values from **0** to **1** to get the corresponding y-axis values...





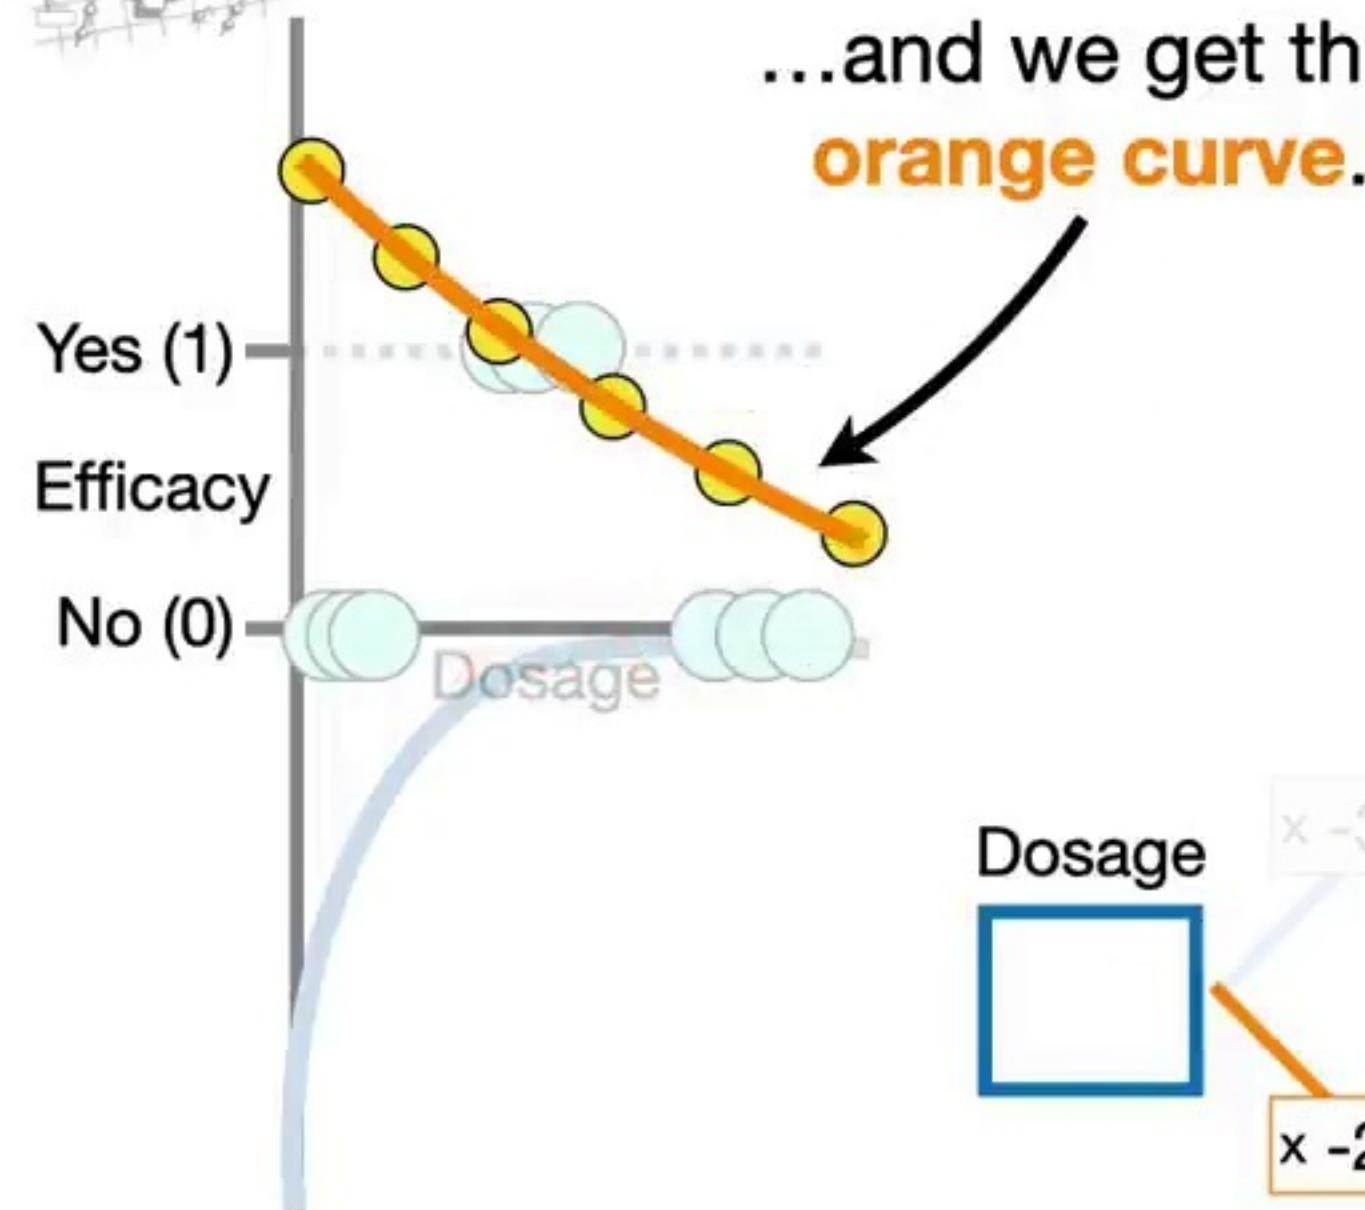
$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$

Now we just plug in **Dosage** values from **0** to **1** to get the corresponding y-axis values...



Double  
BAM!!  
**SQ!**

$$(\text{Dosage} \times -2.52) + 1.29 = \text{x-axis coordinate}$$



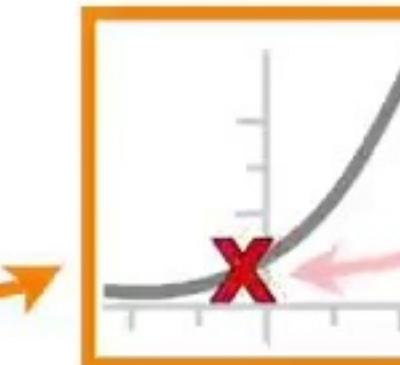
Dosage

$x - 2.52$

+ 1.29

+ 2.14

$x - 34.4$



$$f(x) = \log(1 + e^x)$$

$x - 1.30$

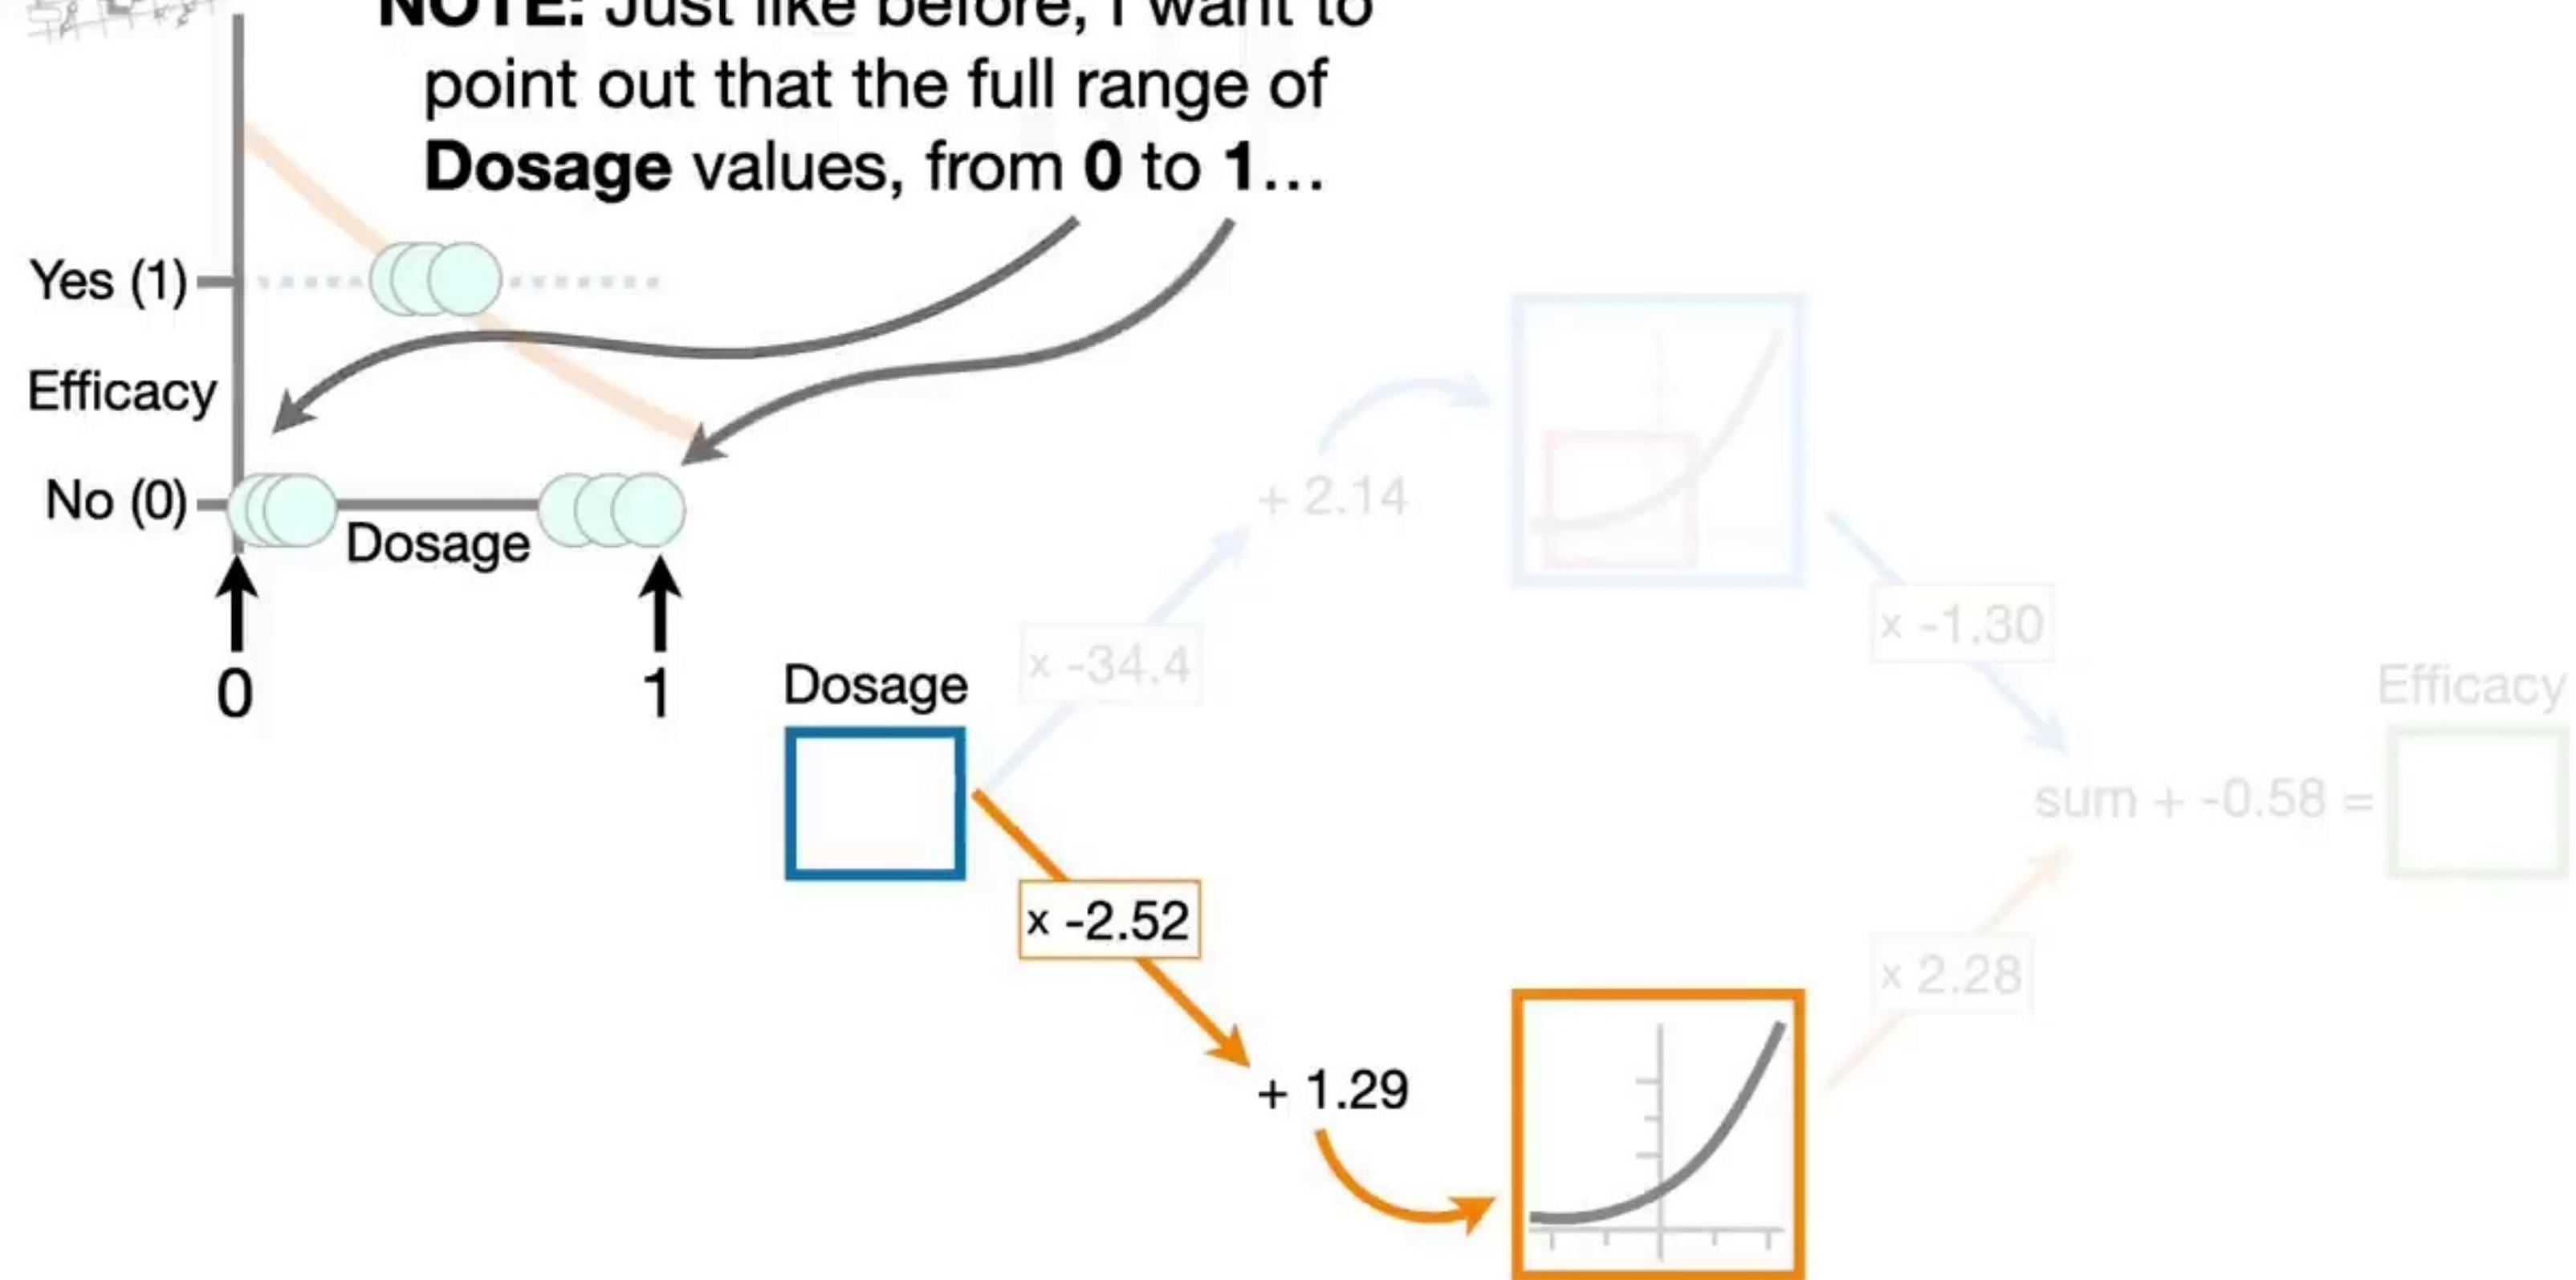
sum + -0.58 =

$x 2.28$

Efficacy

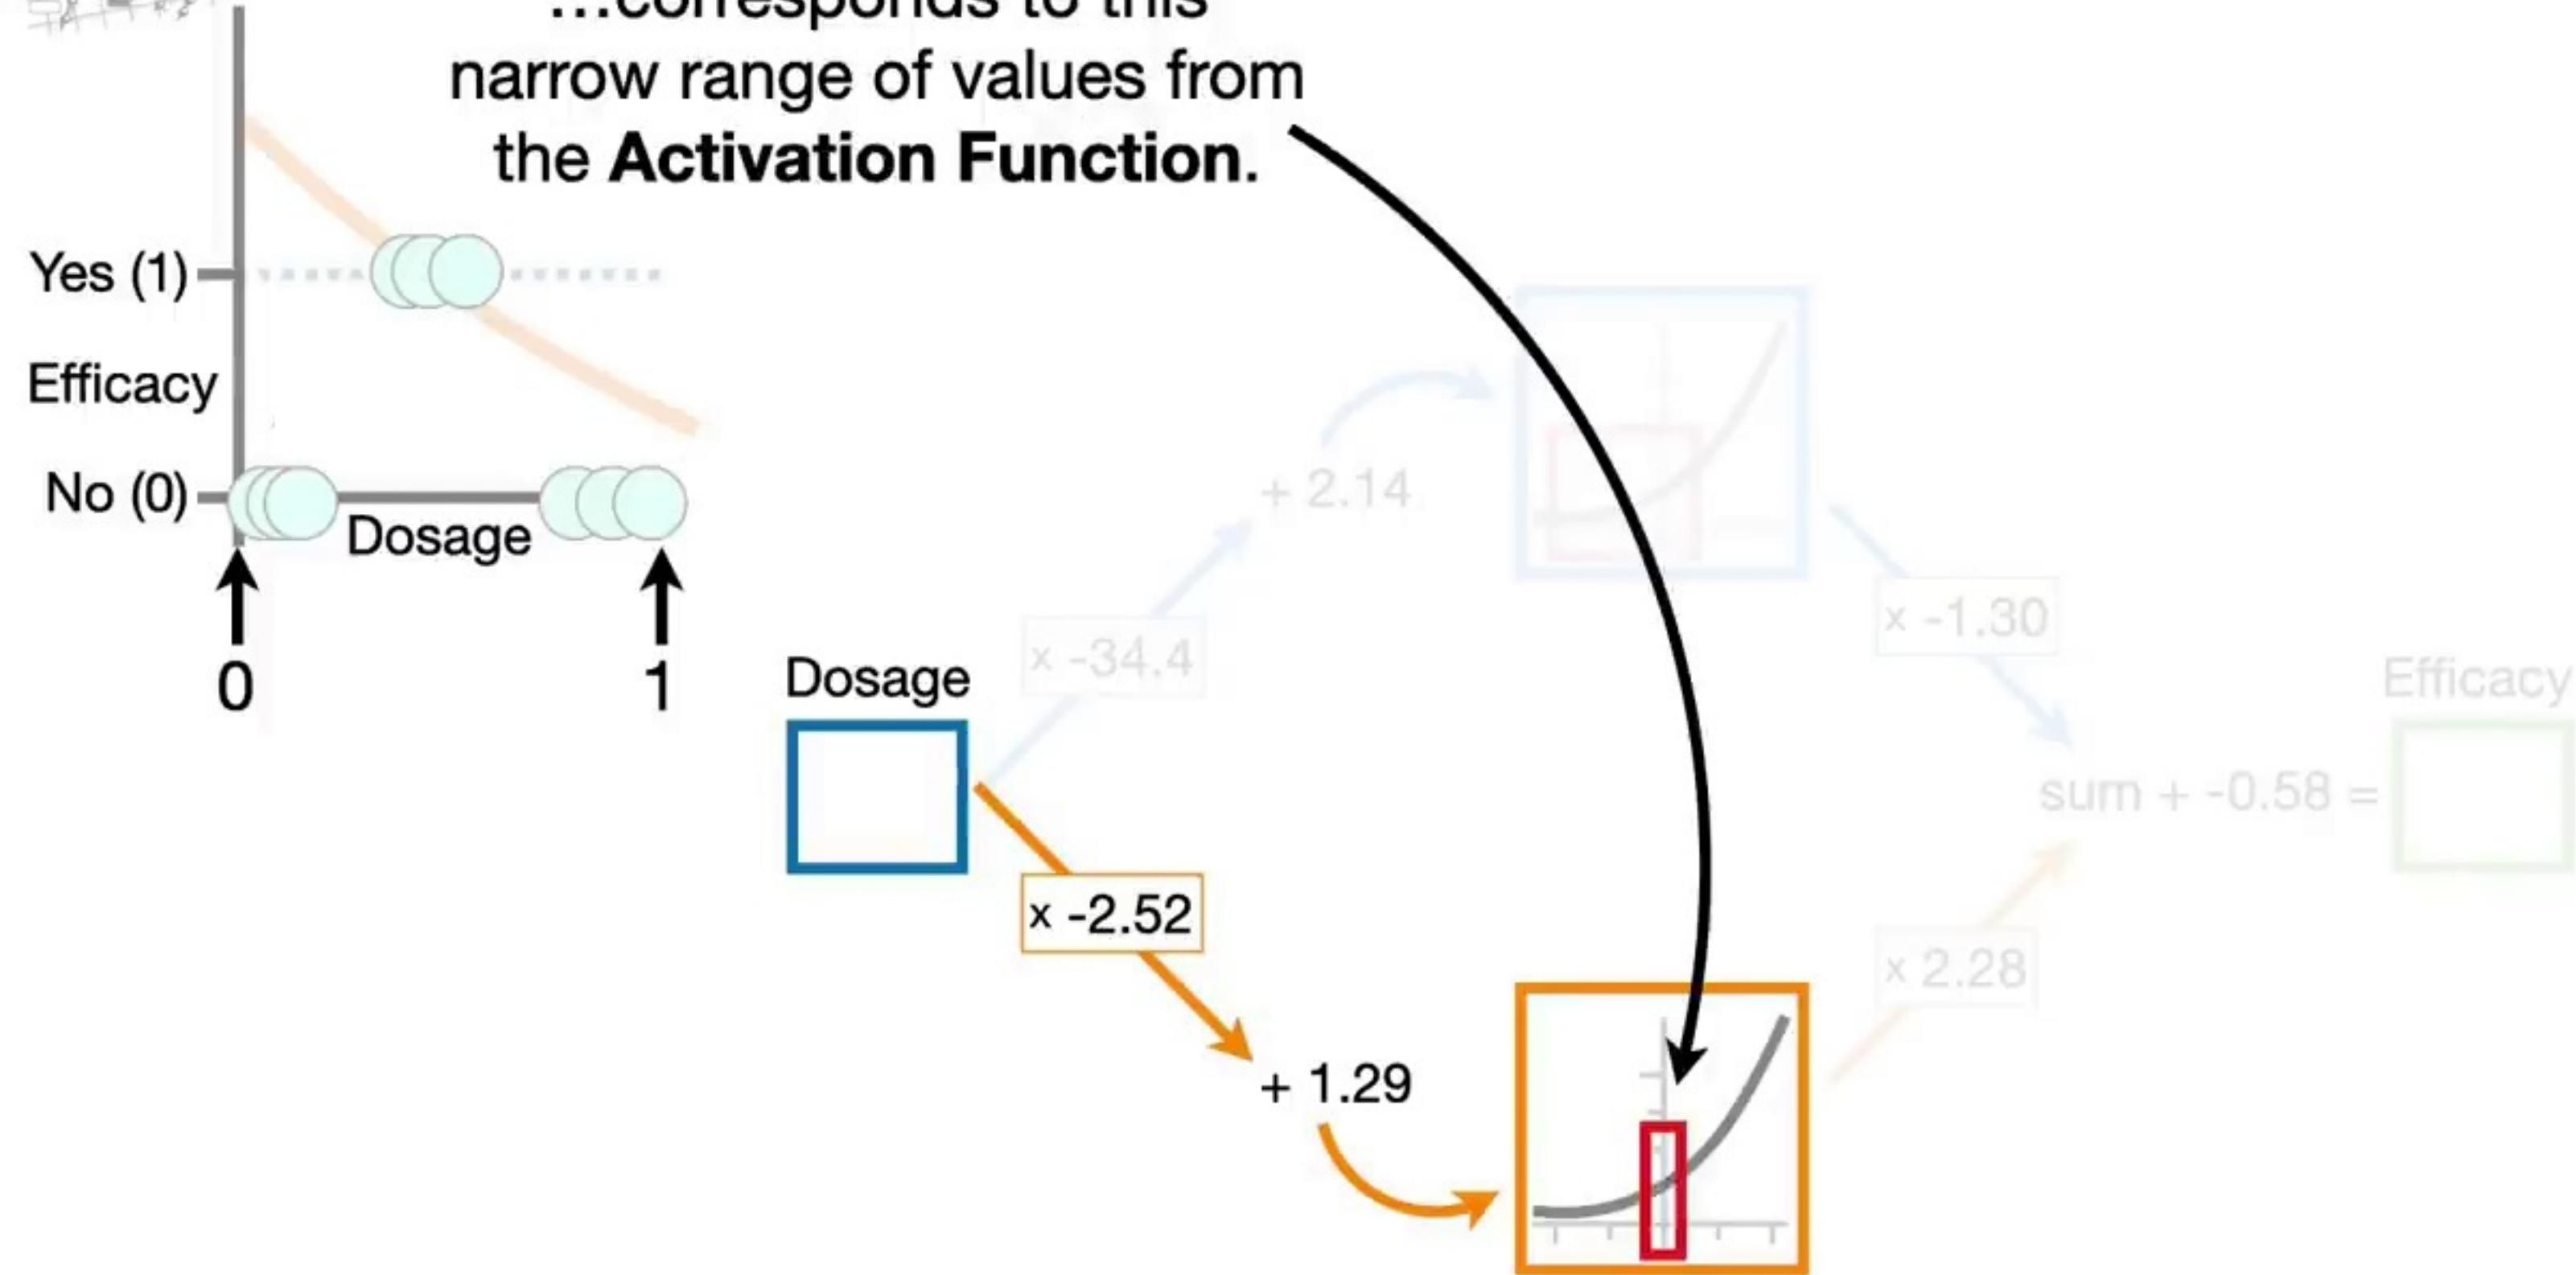


**NOTE:** Just like before, I want to point out that the full range of **Dosage** values, from **0** to **1**...



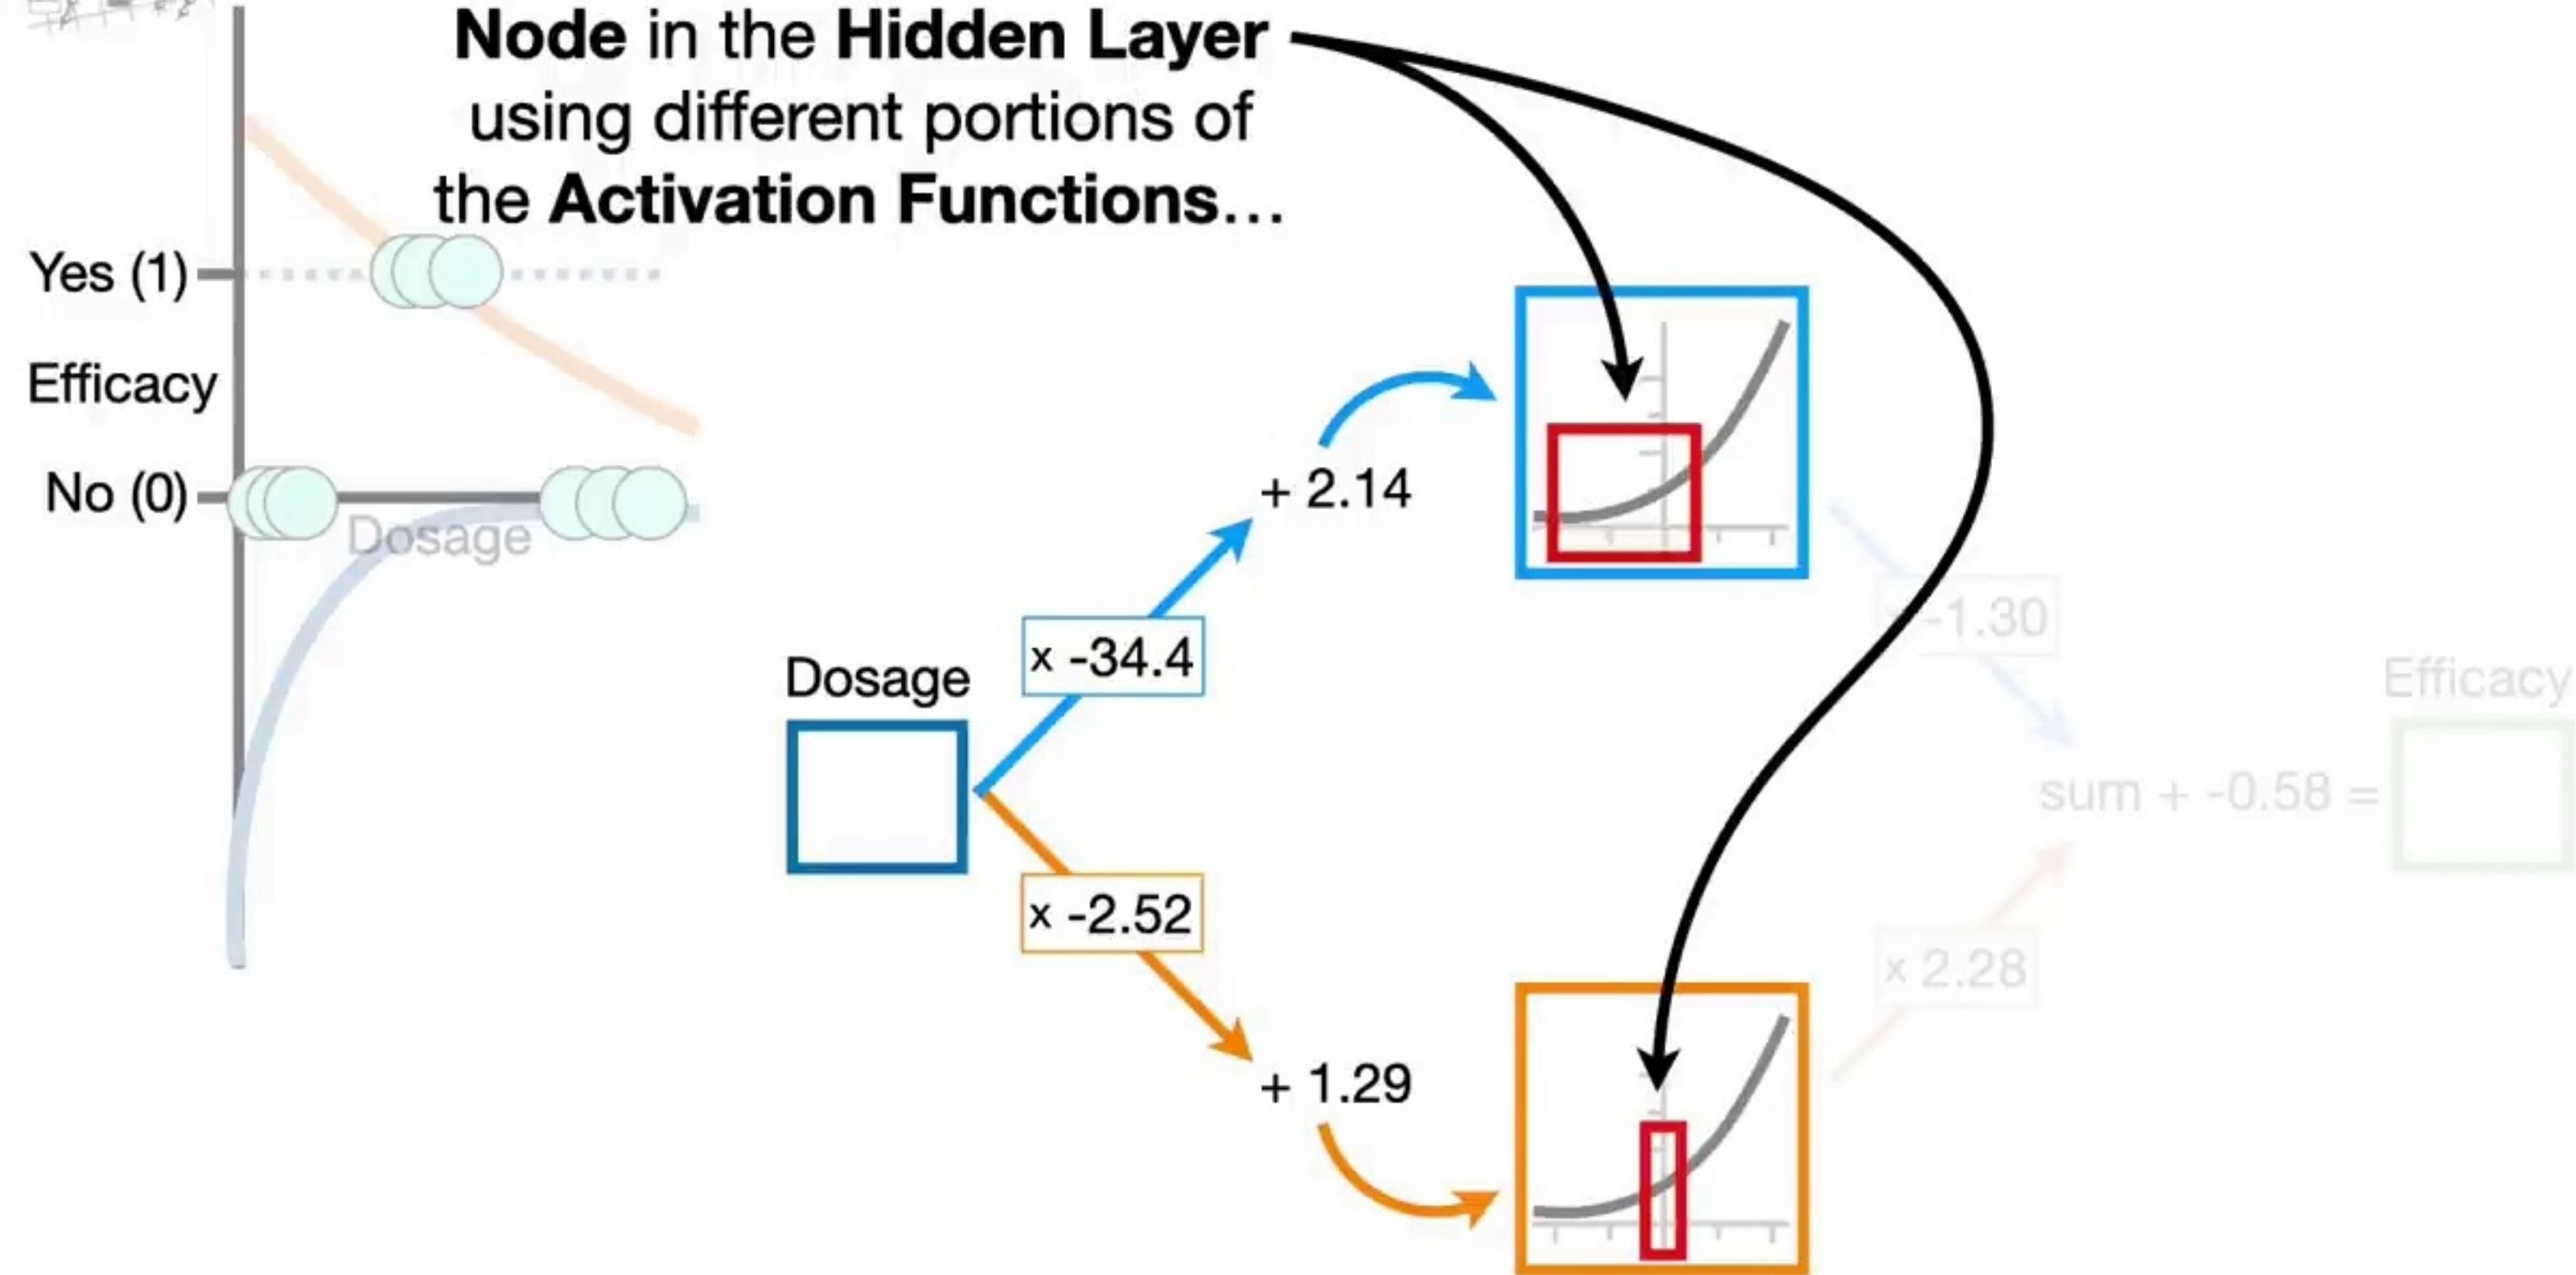


...corresponds to this narrow range of values from the **Activation Function**.



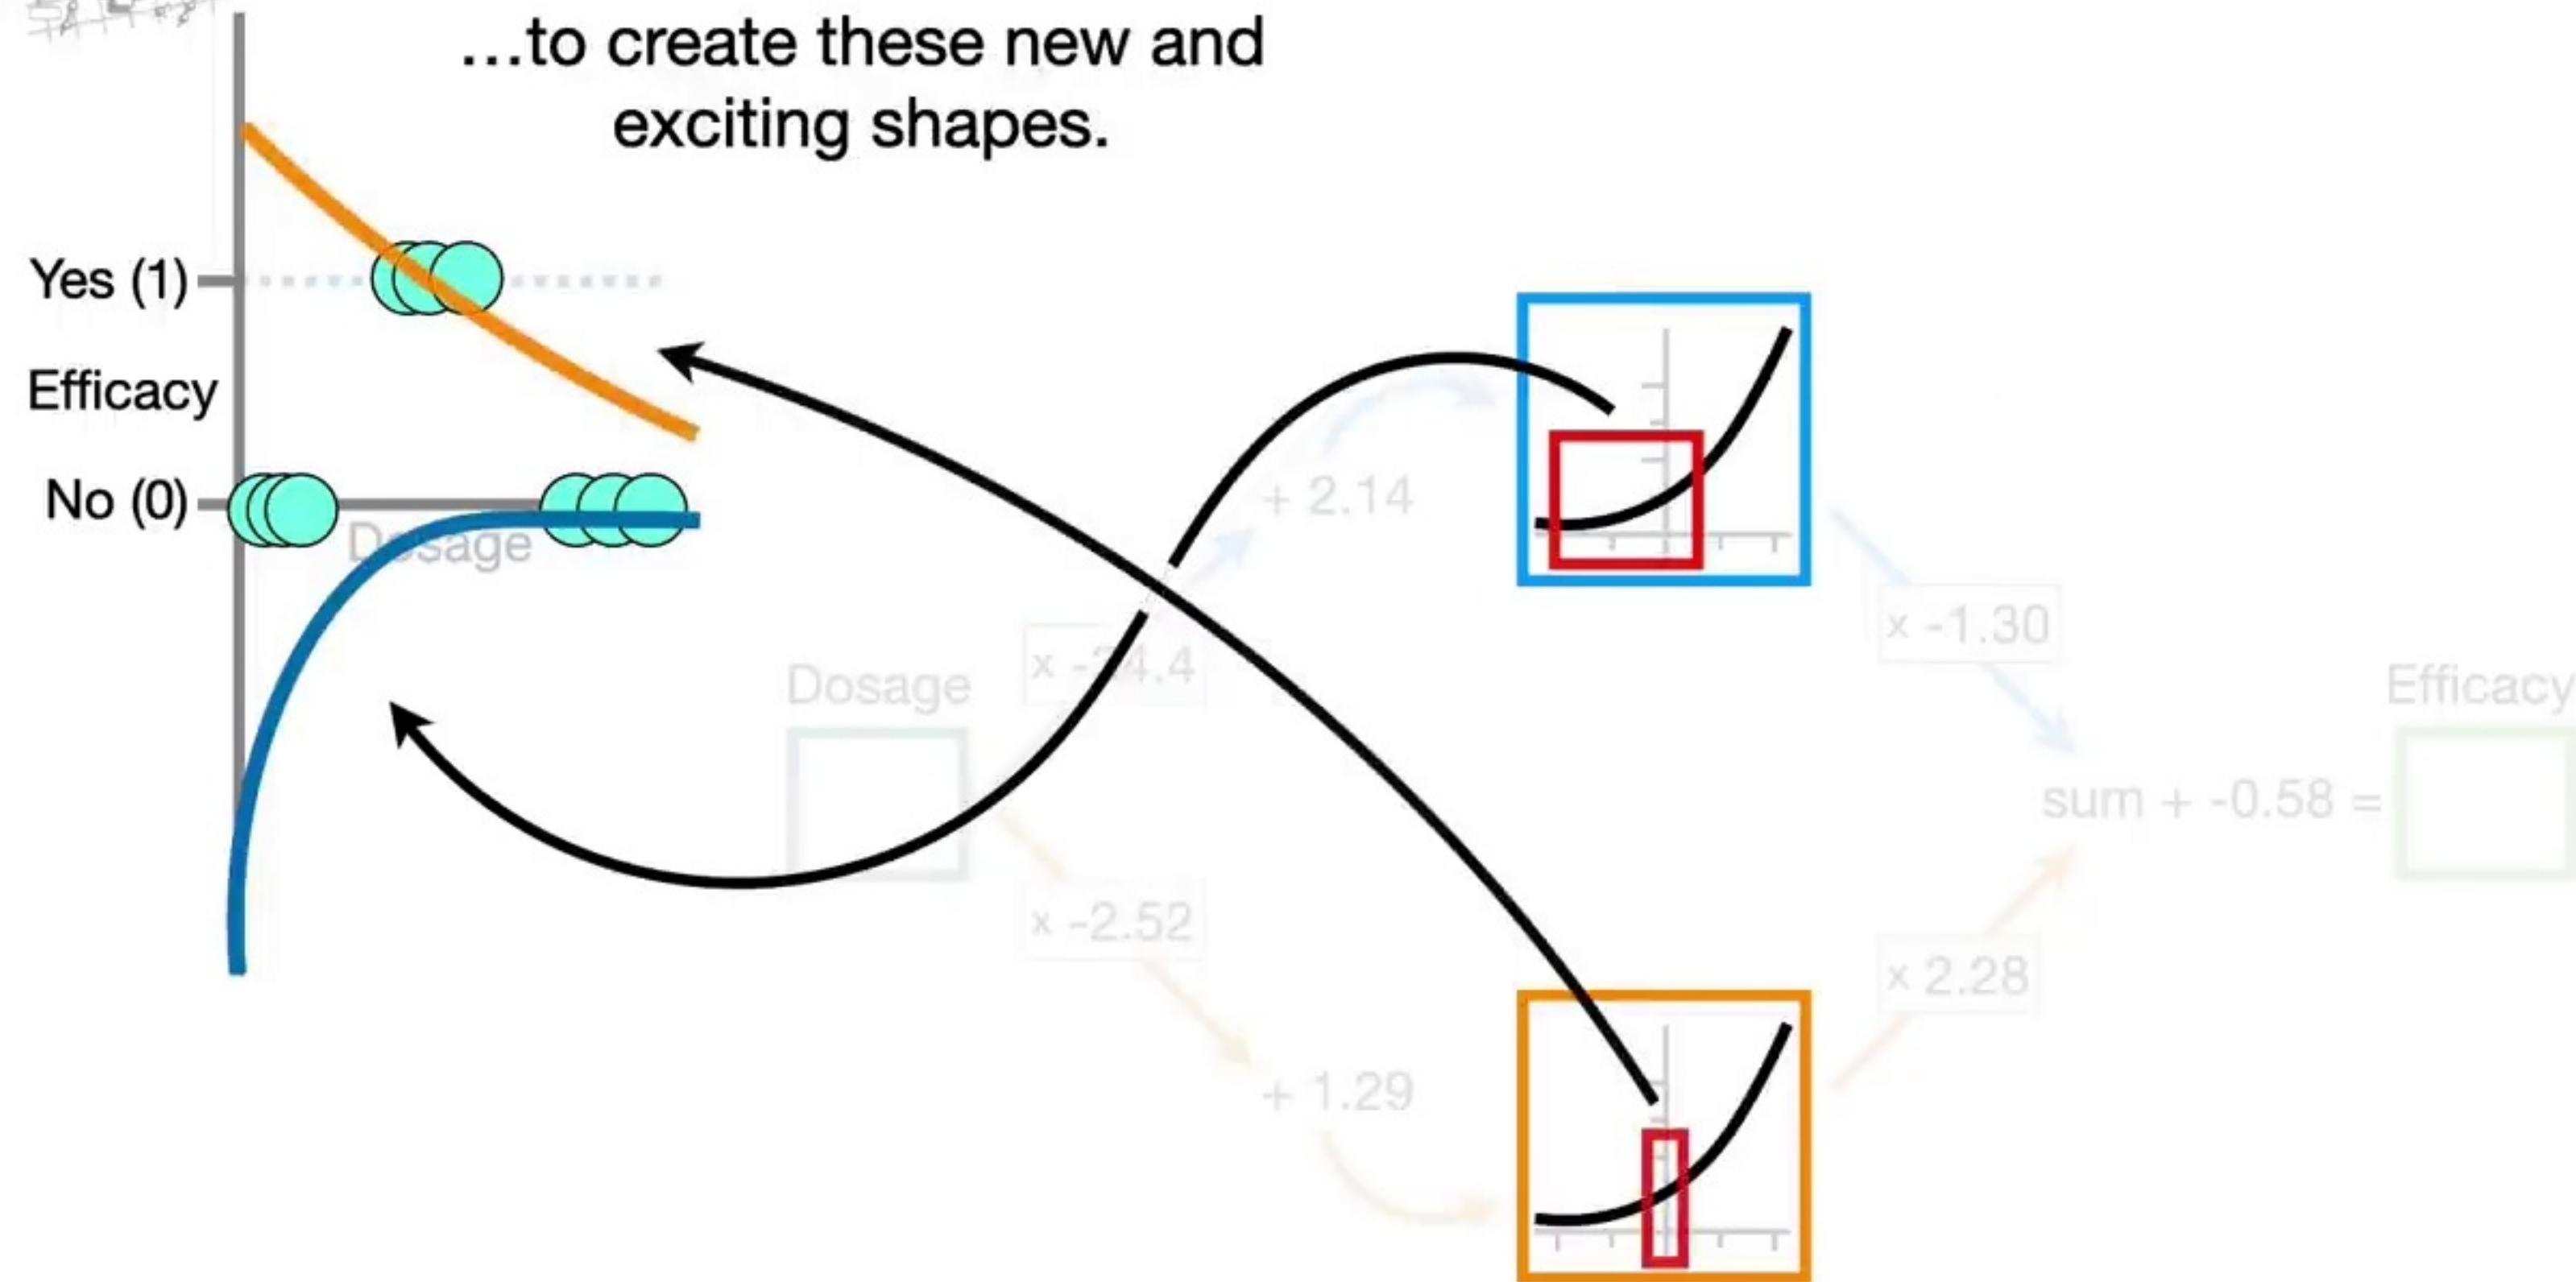


...and that results in each **Node** in the **Hidden Layer** using different portions of the **Activation Functions**...



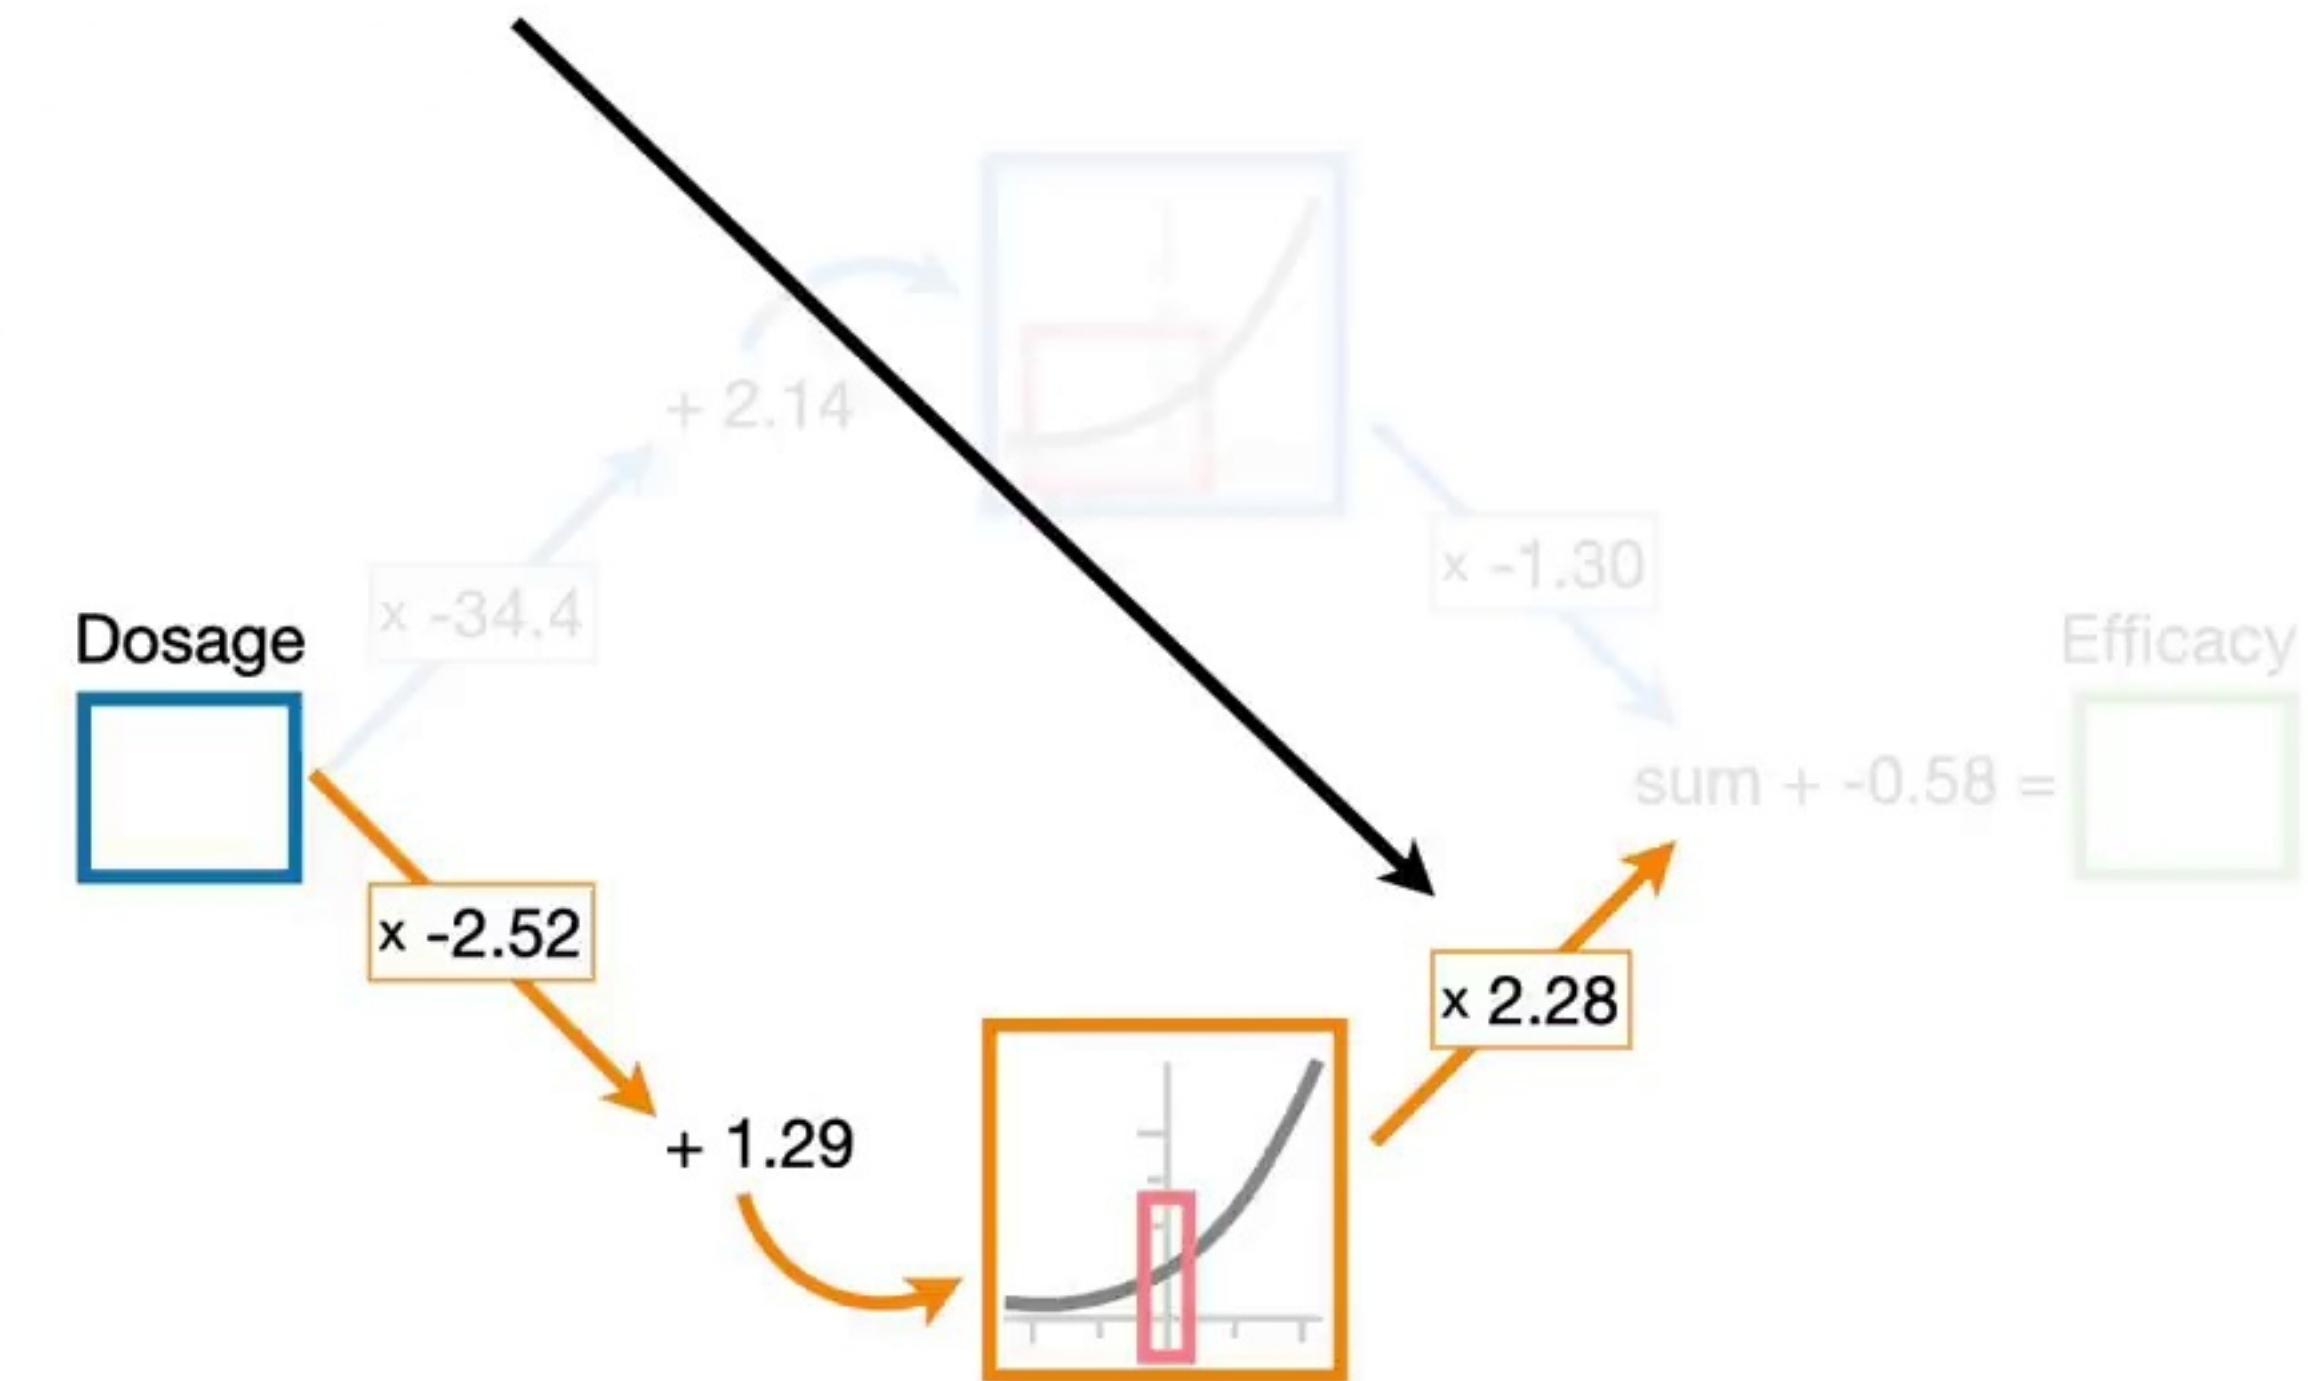
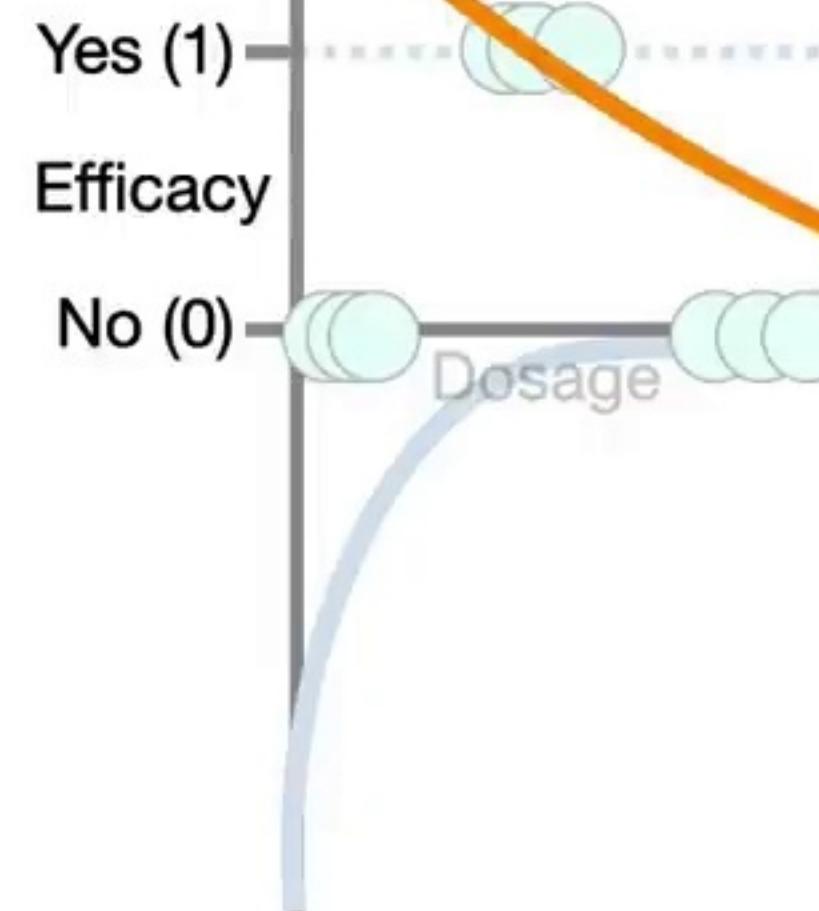


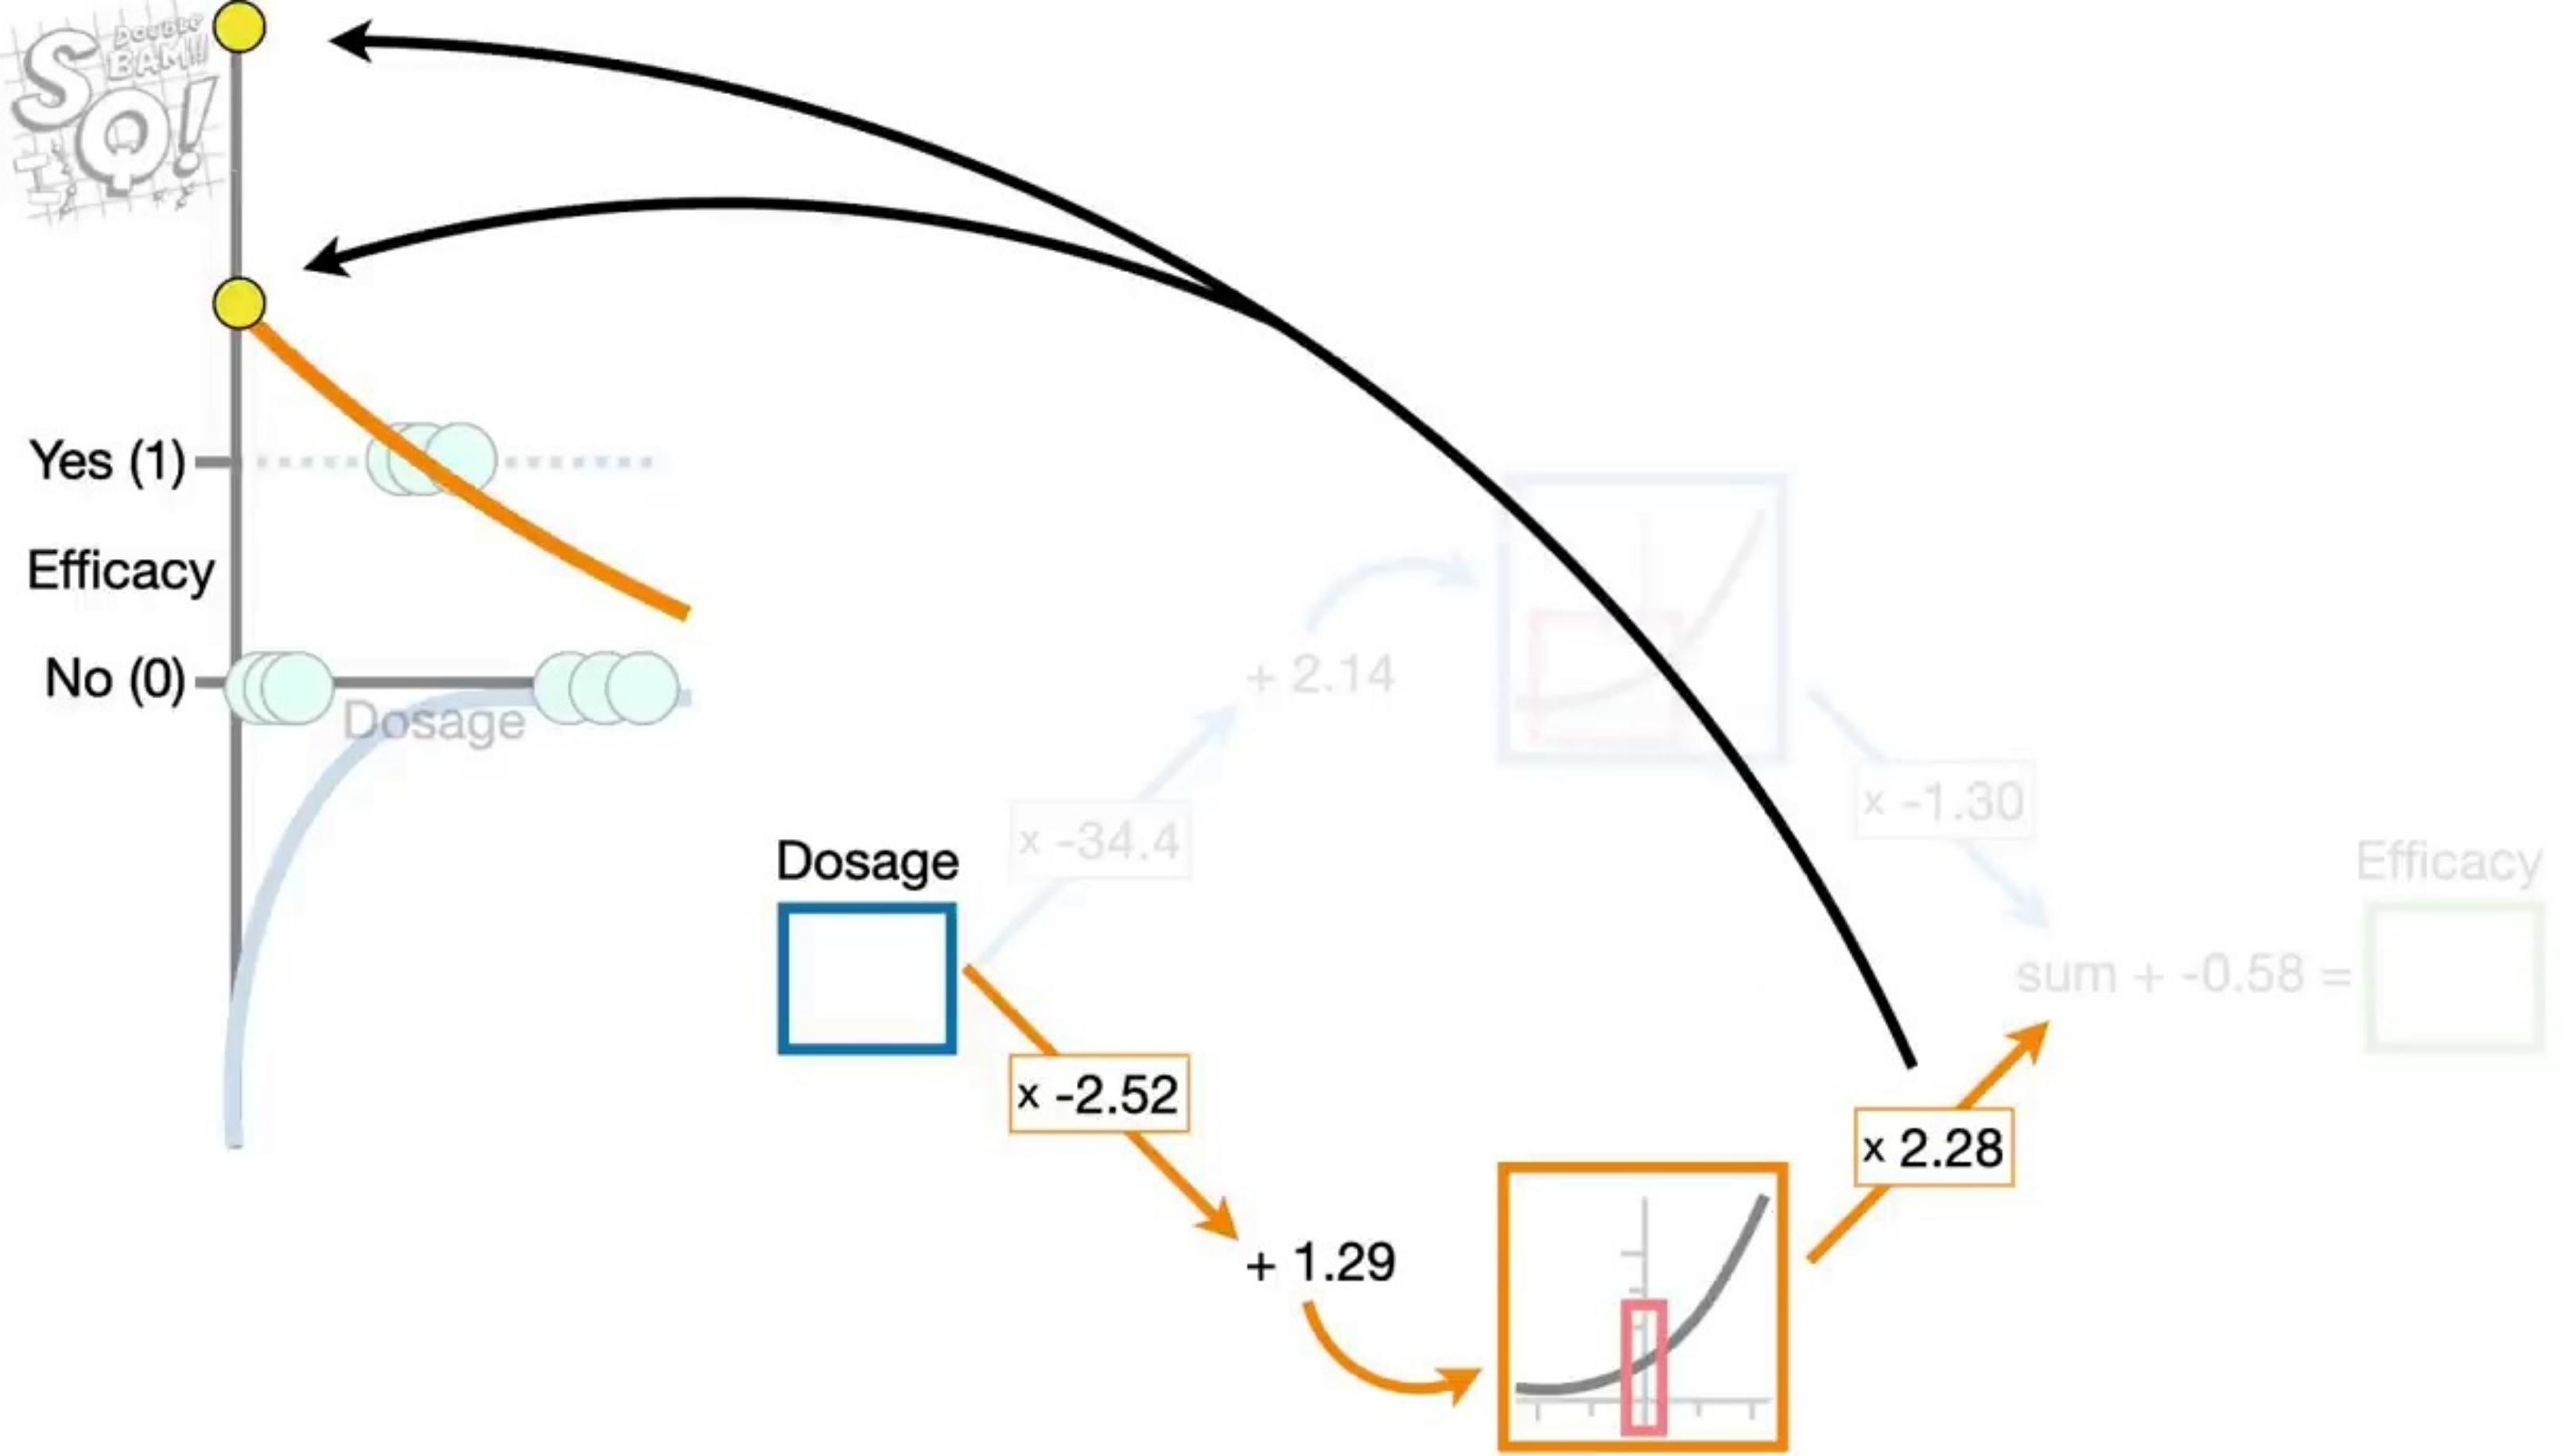
...to create these new and exciting shapes.

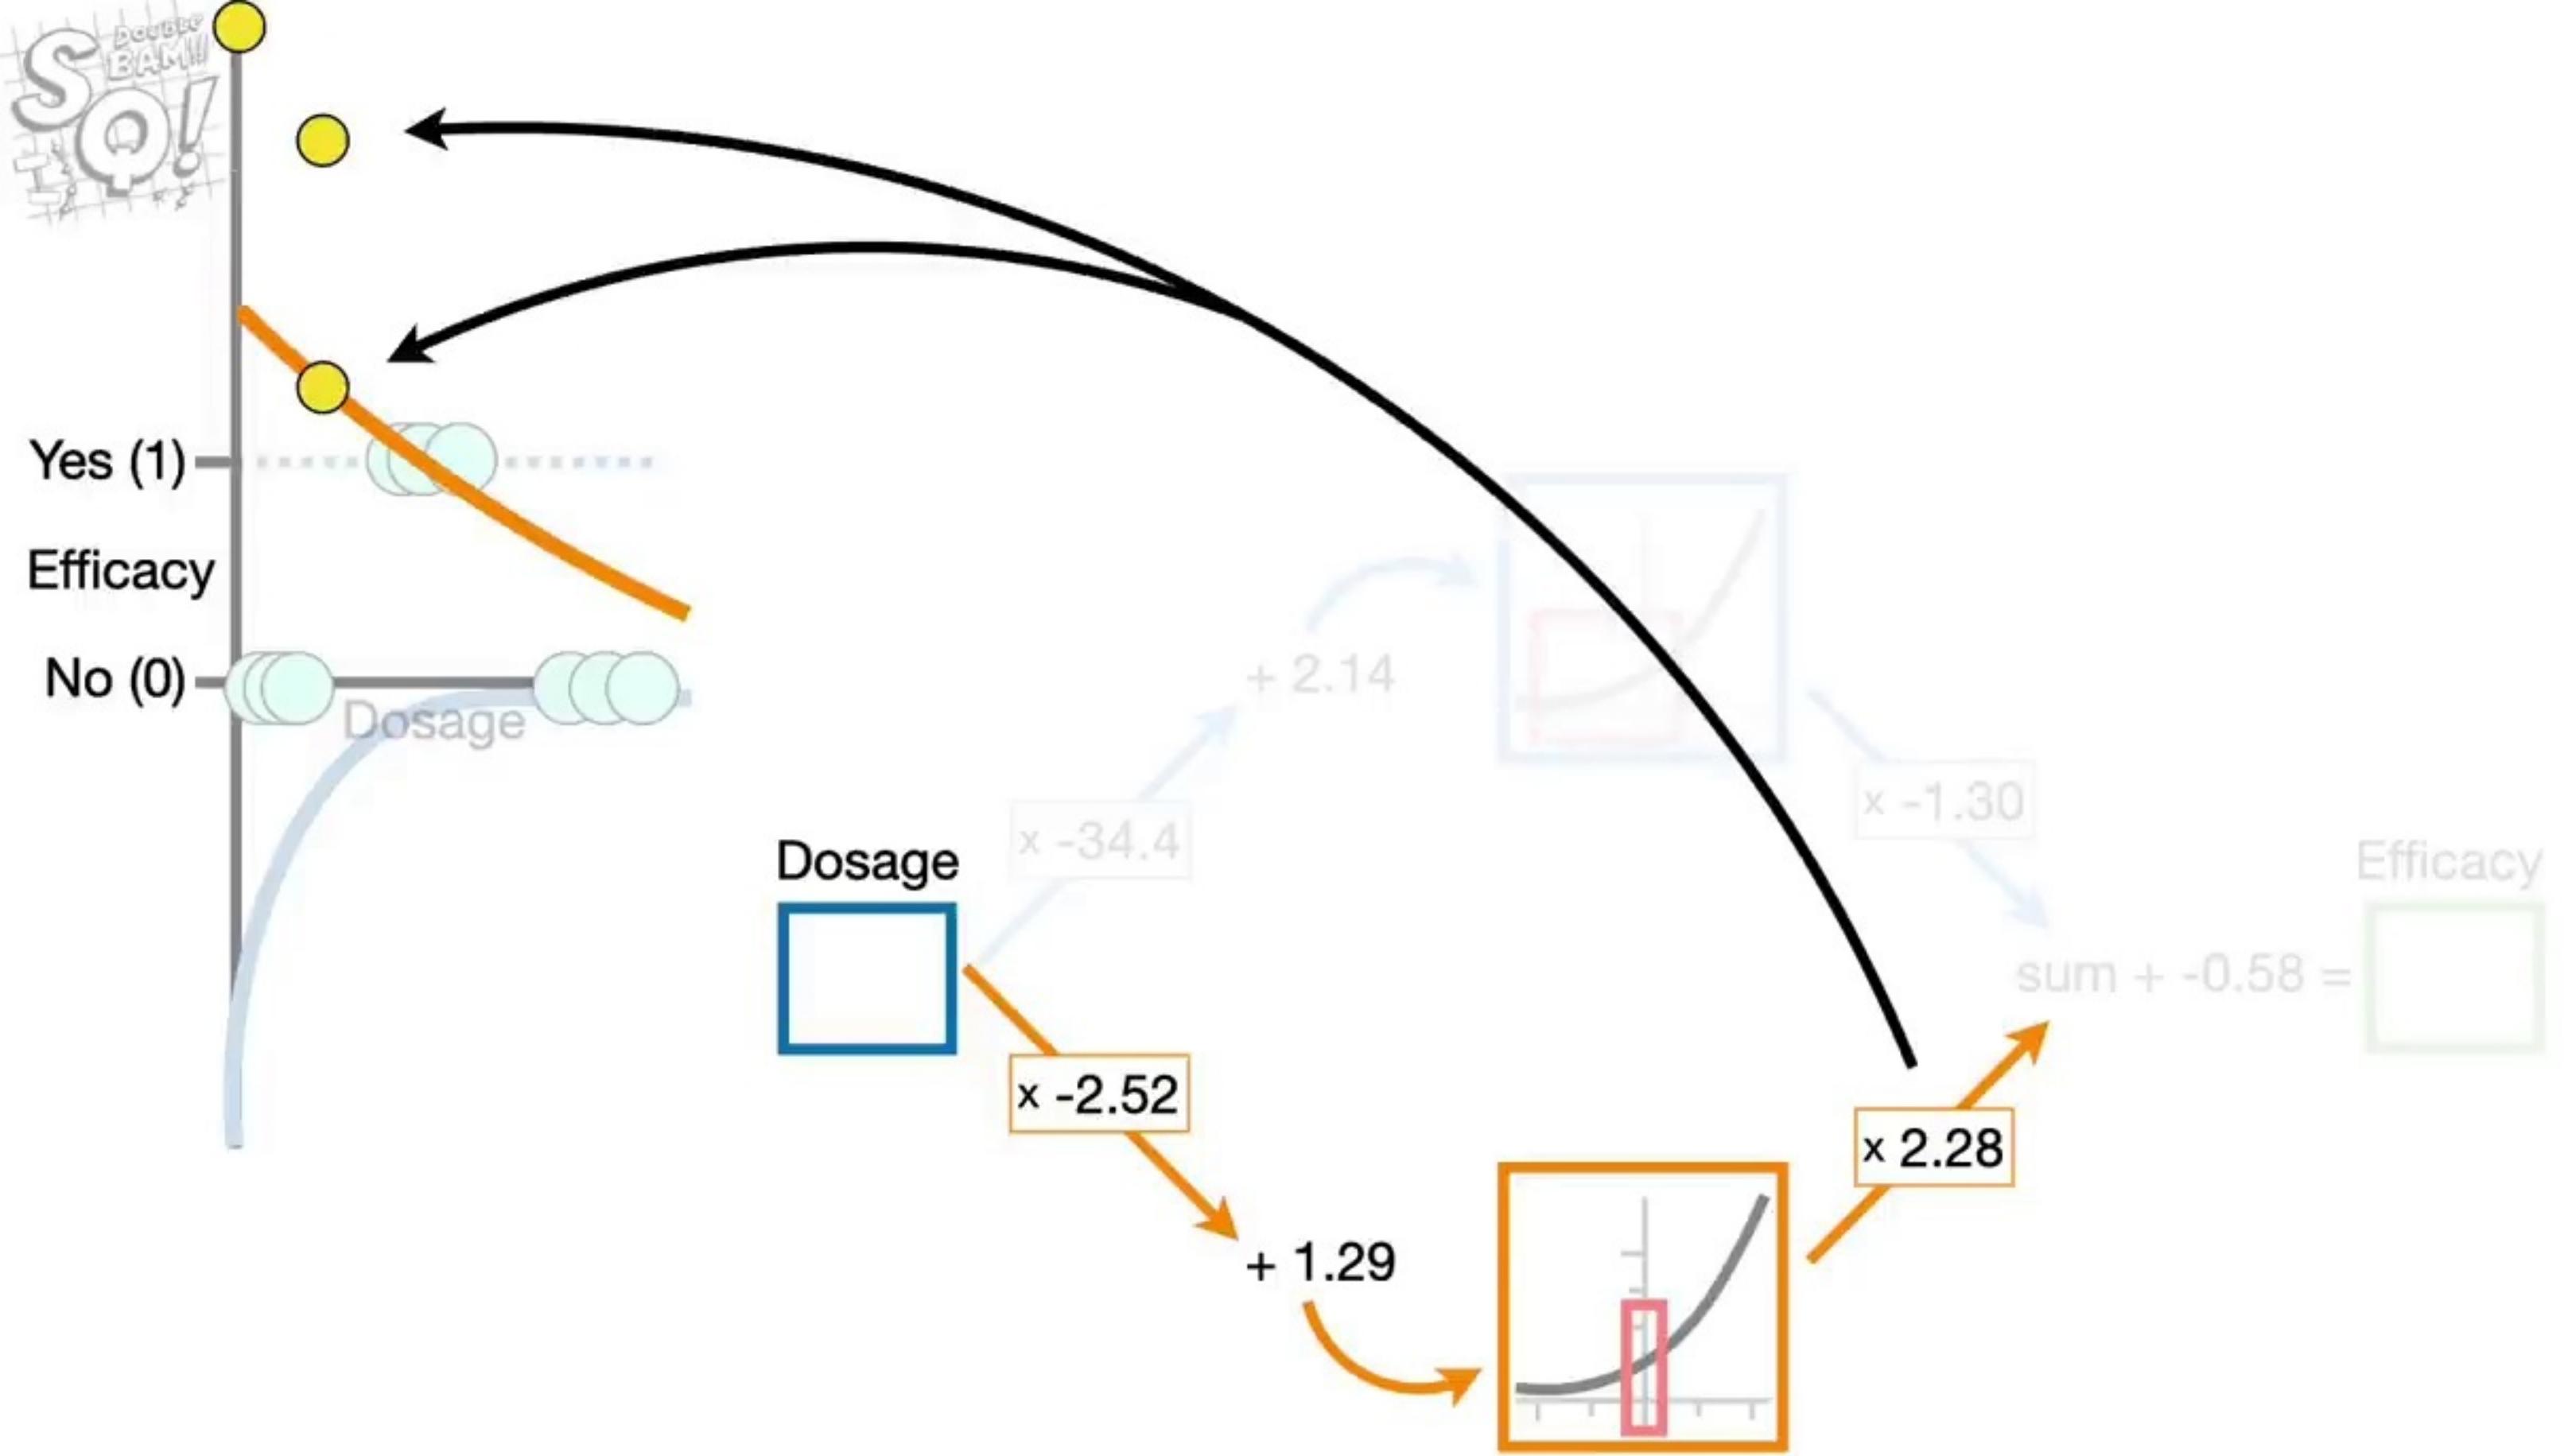


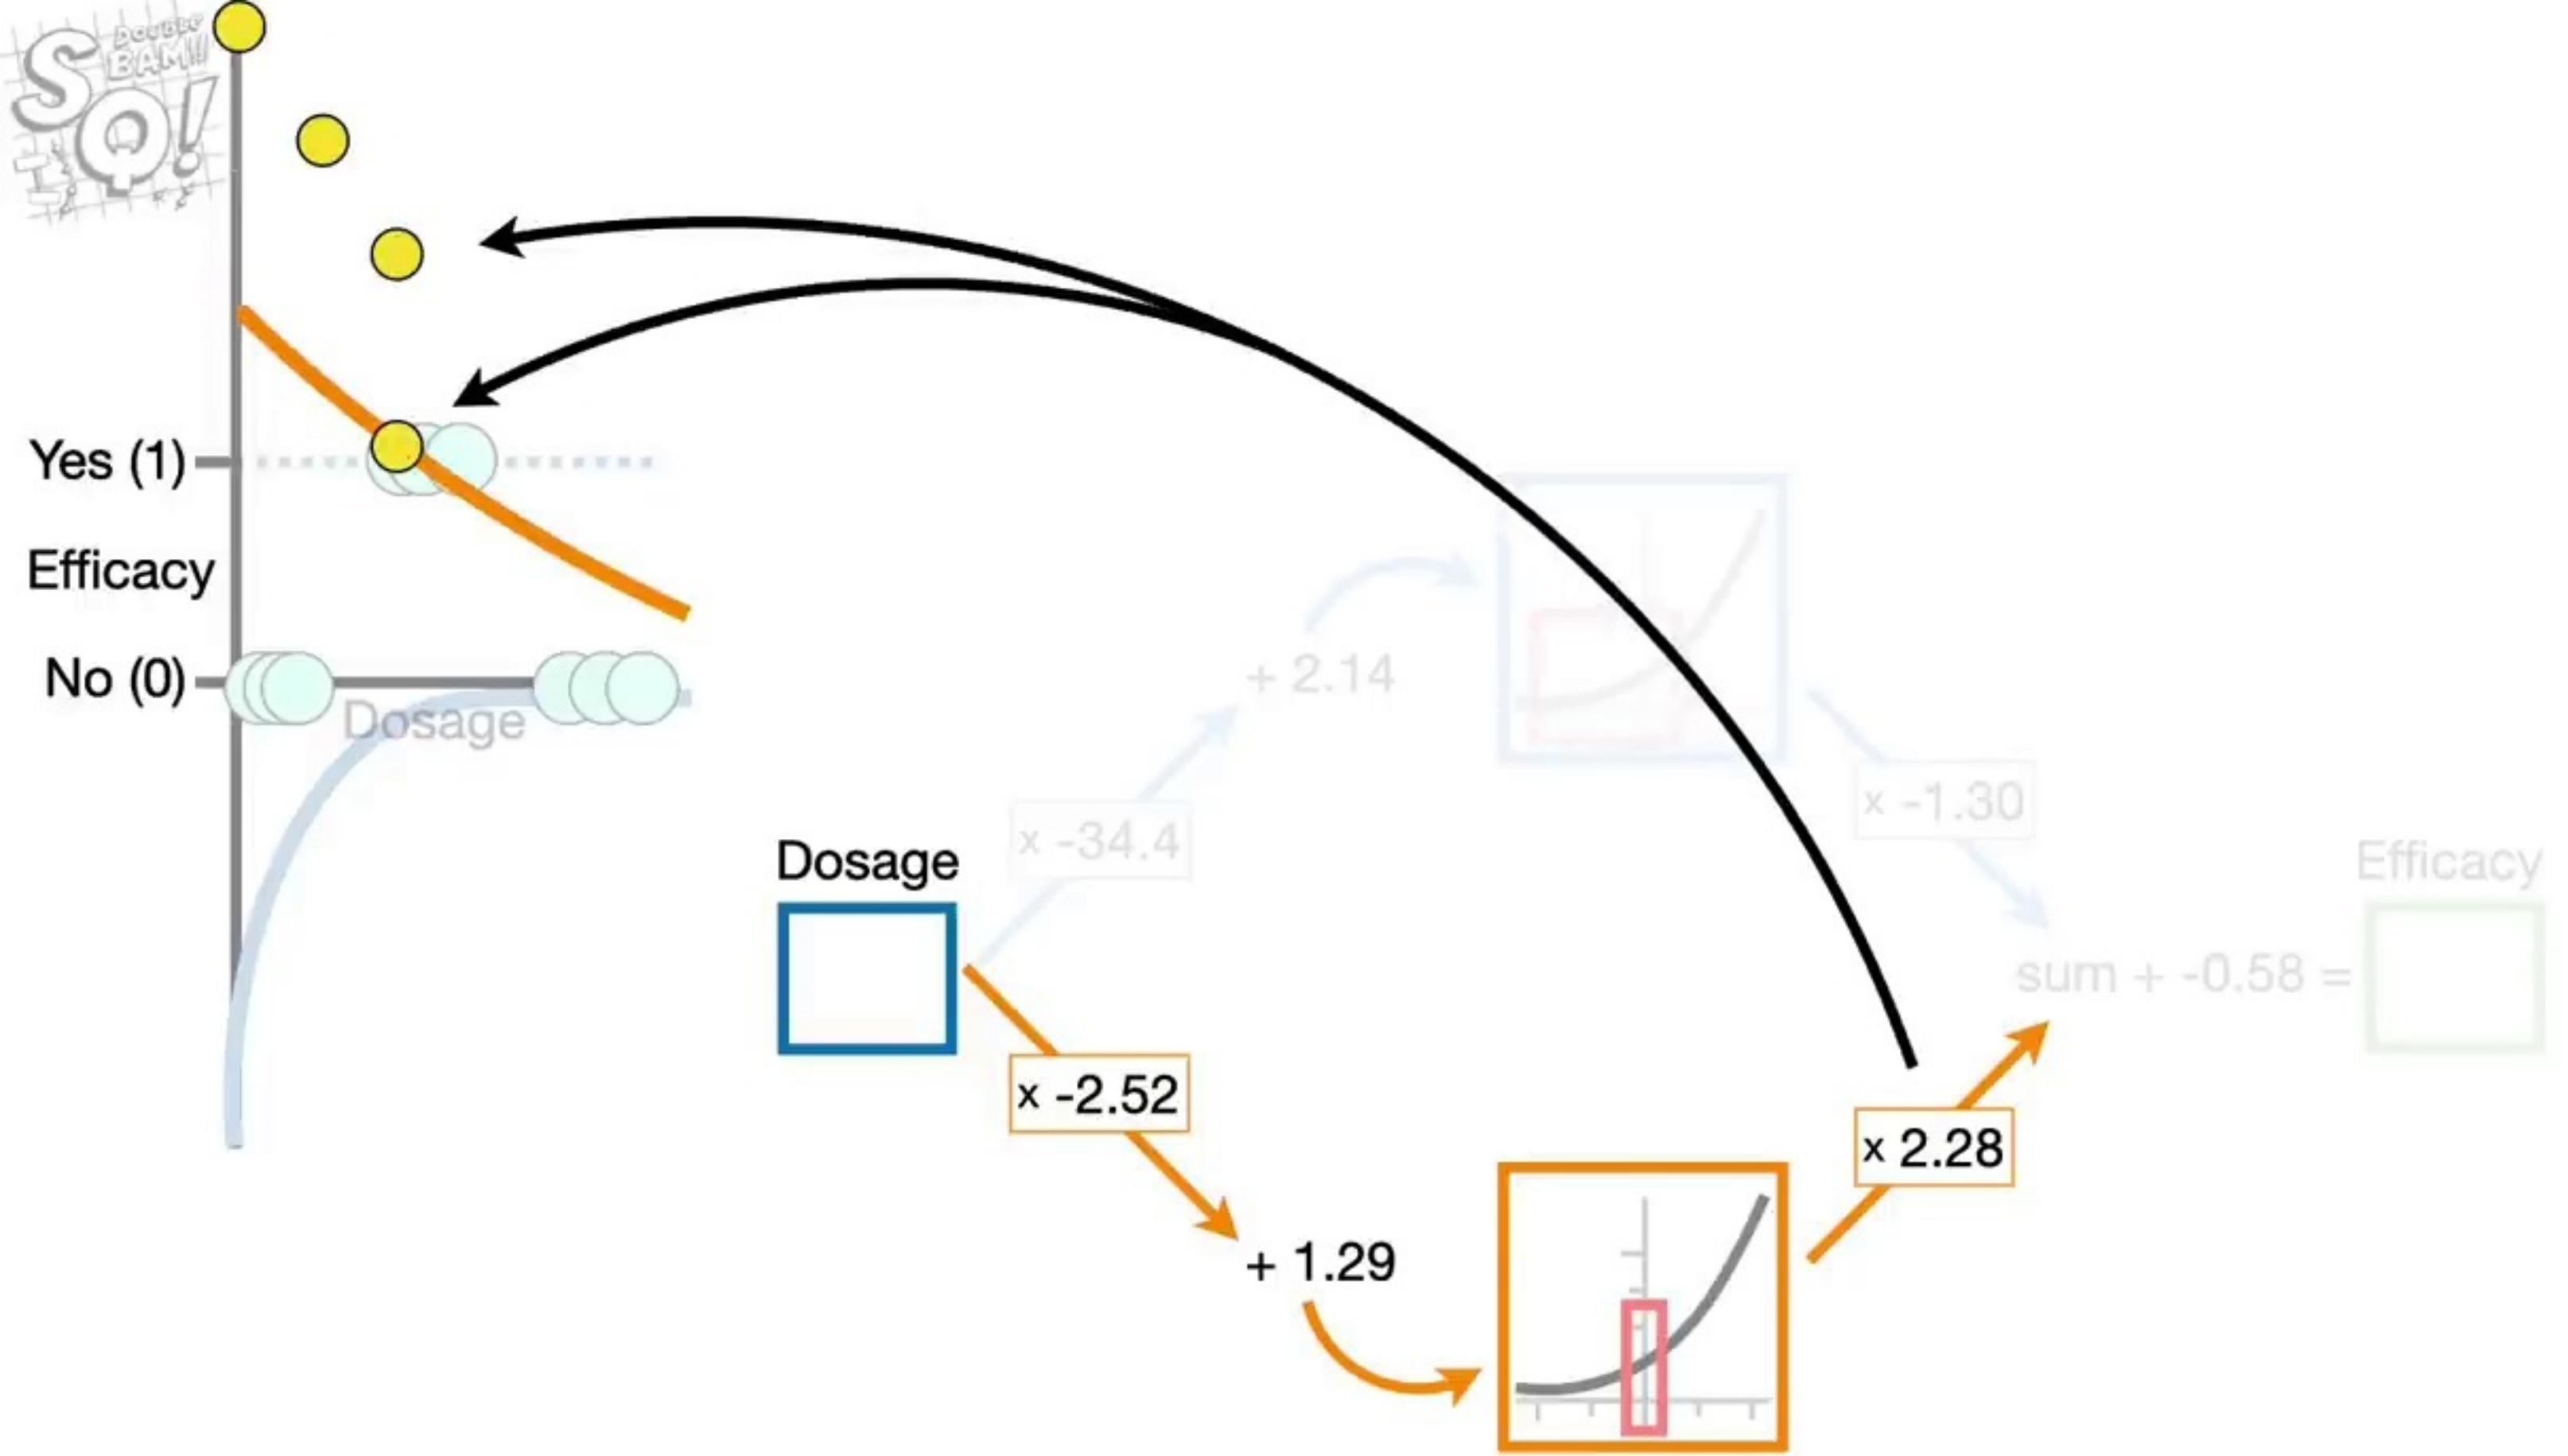


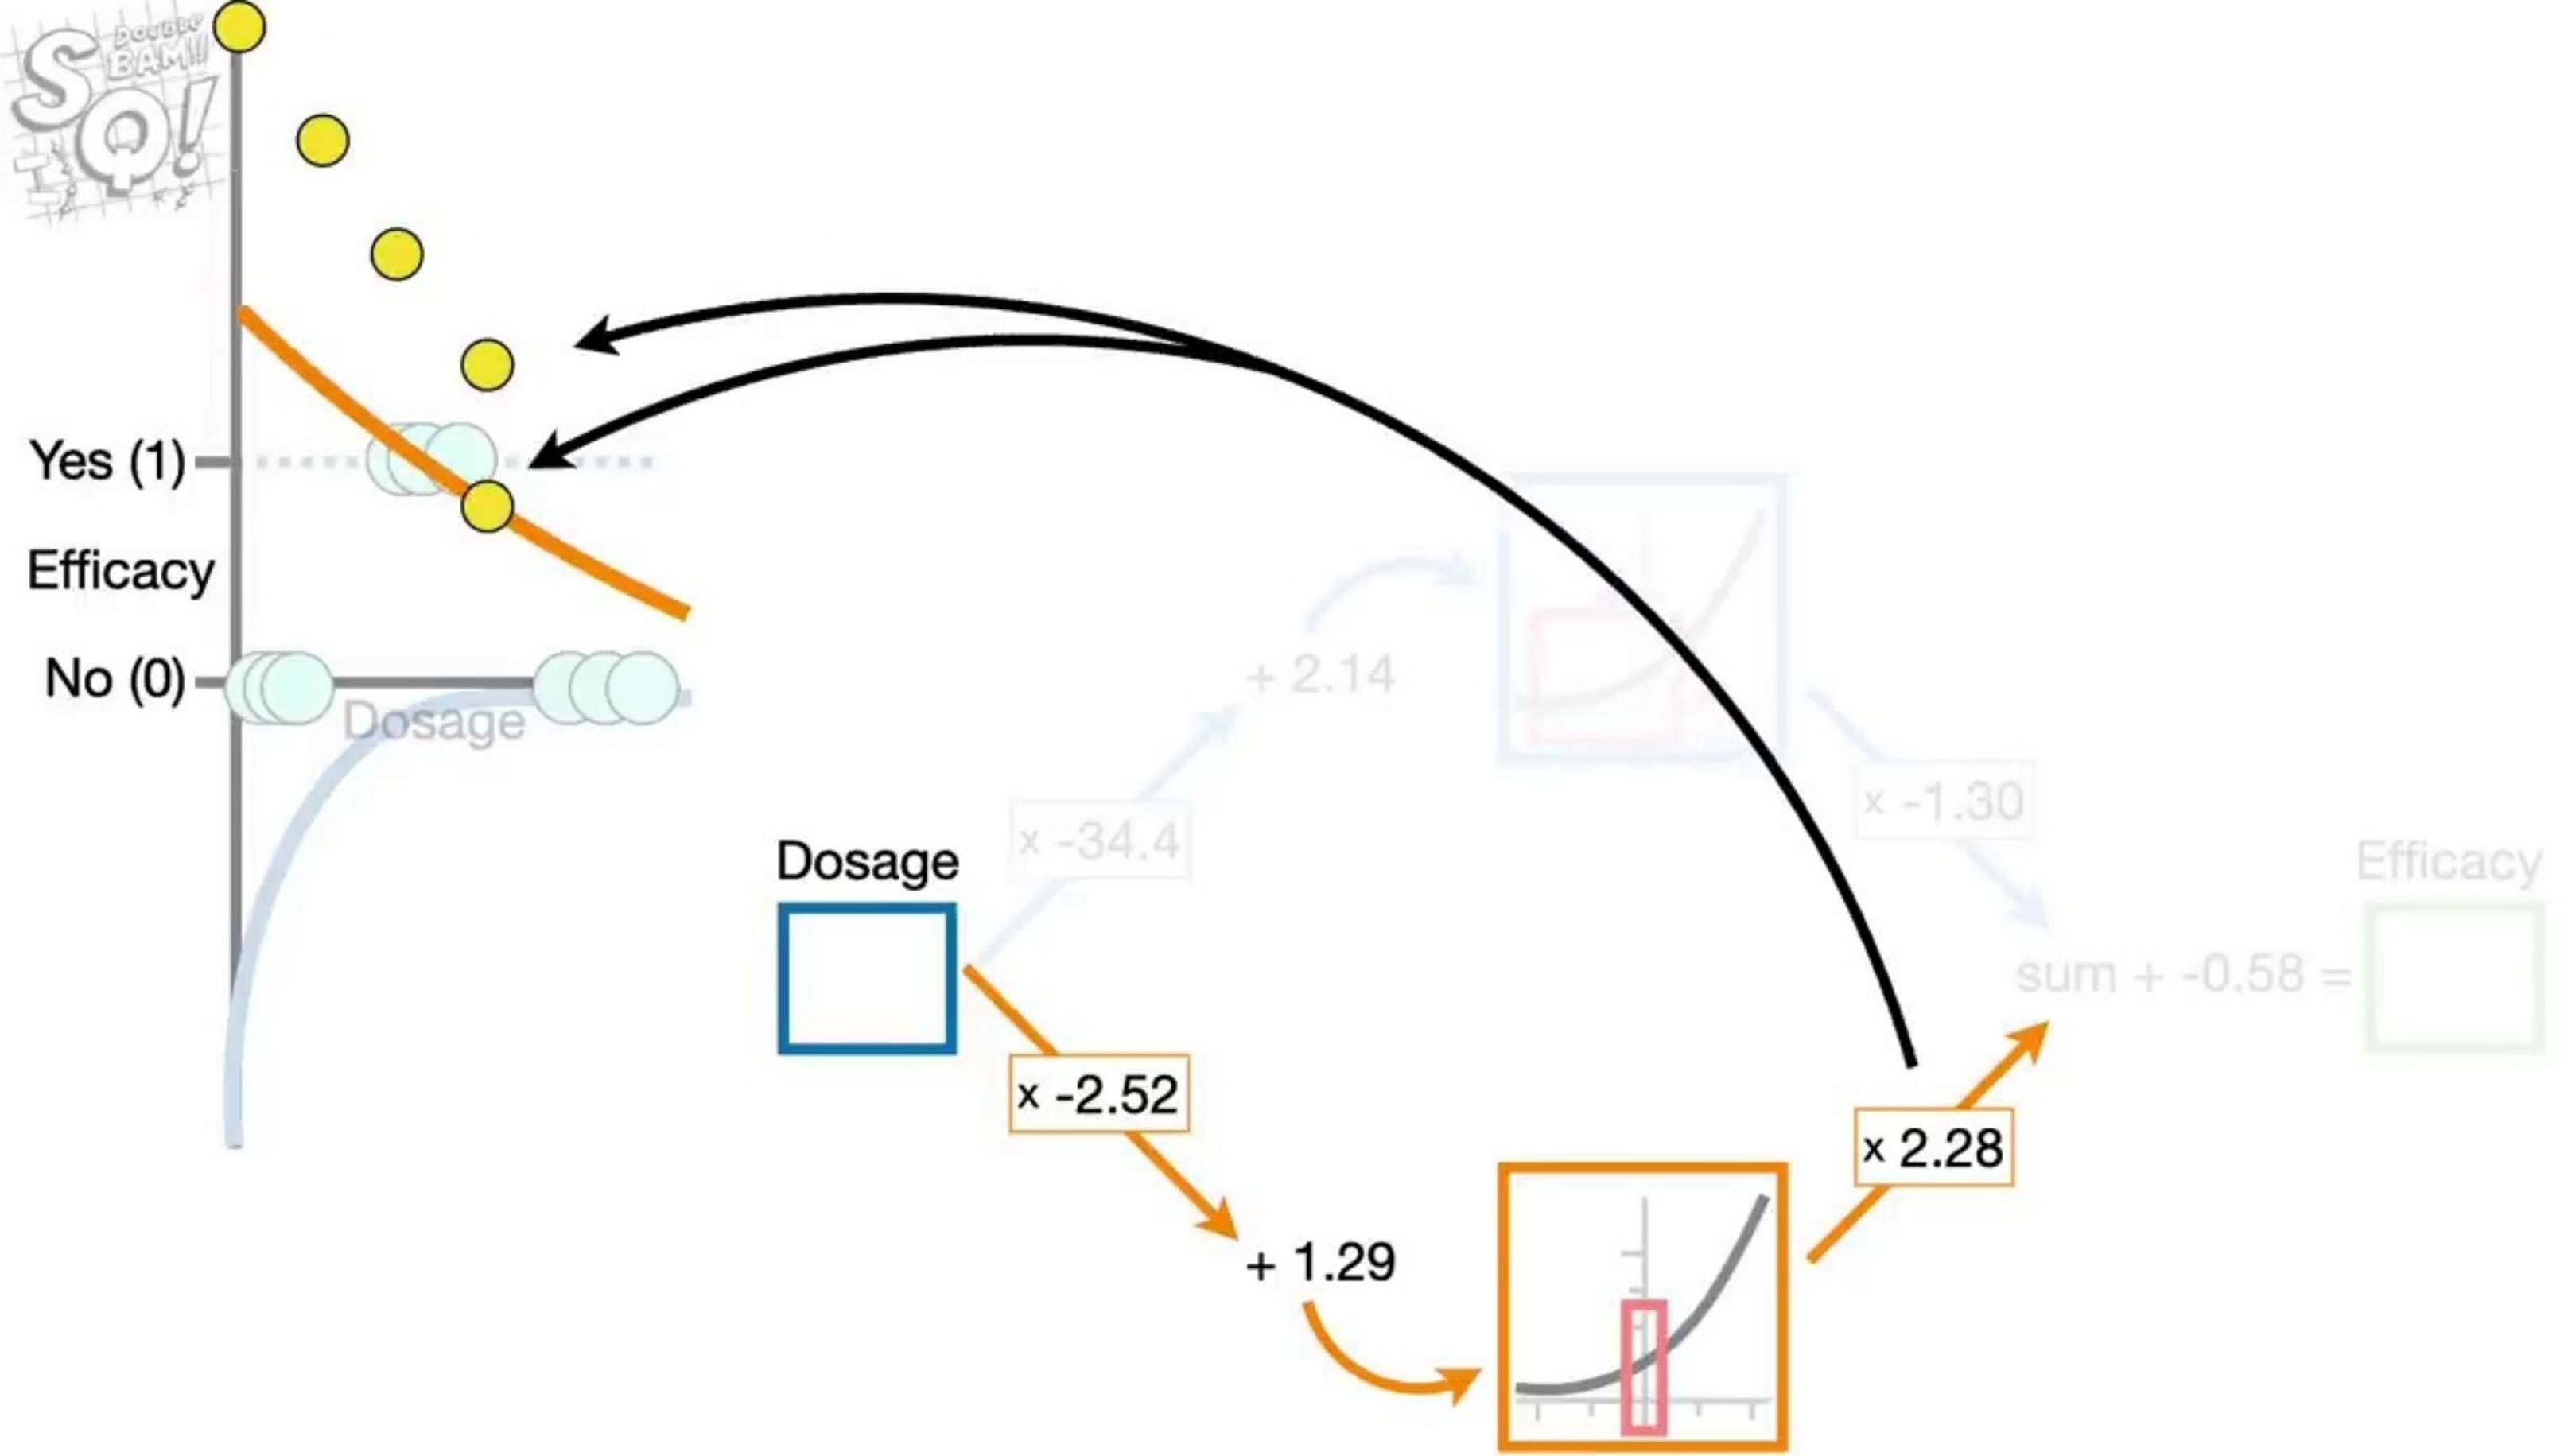
Only this time, we scale by  
a positive number, **2.28**.

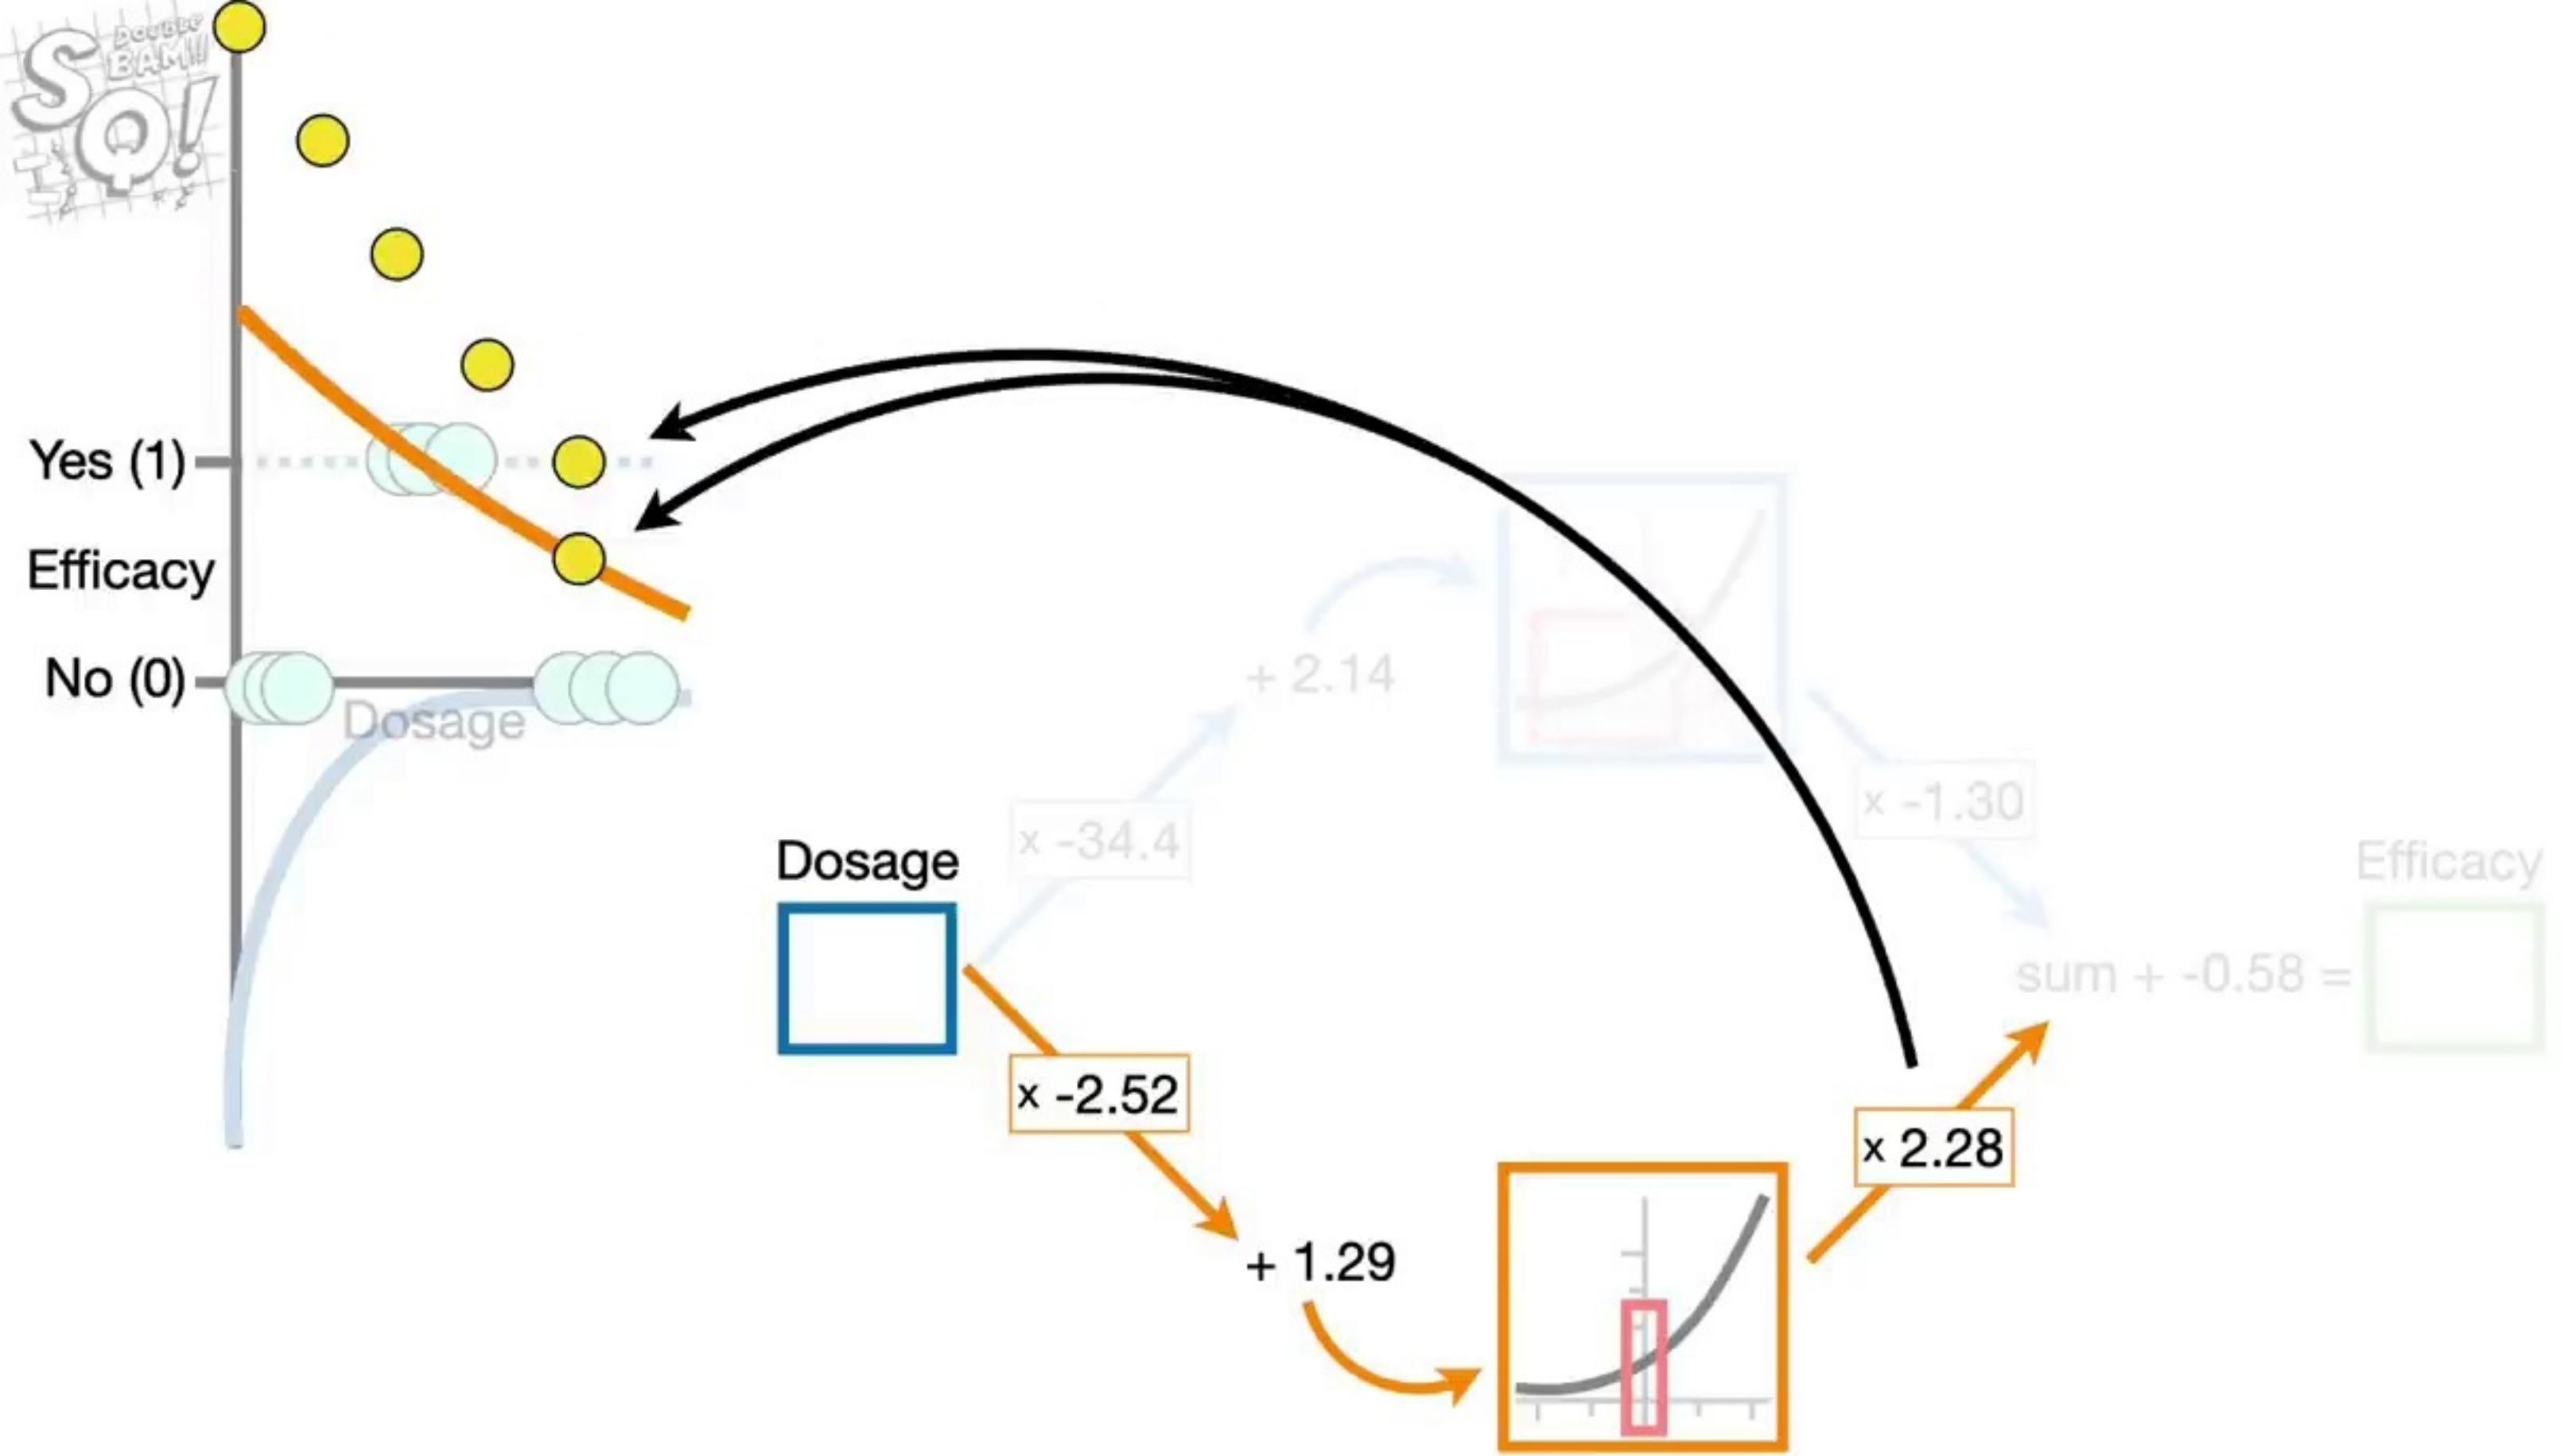


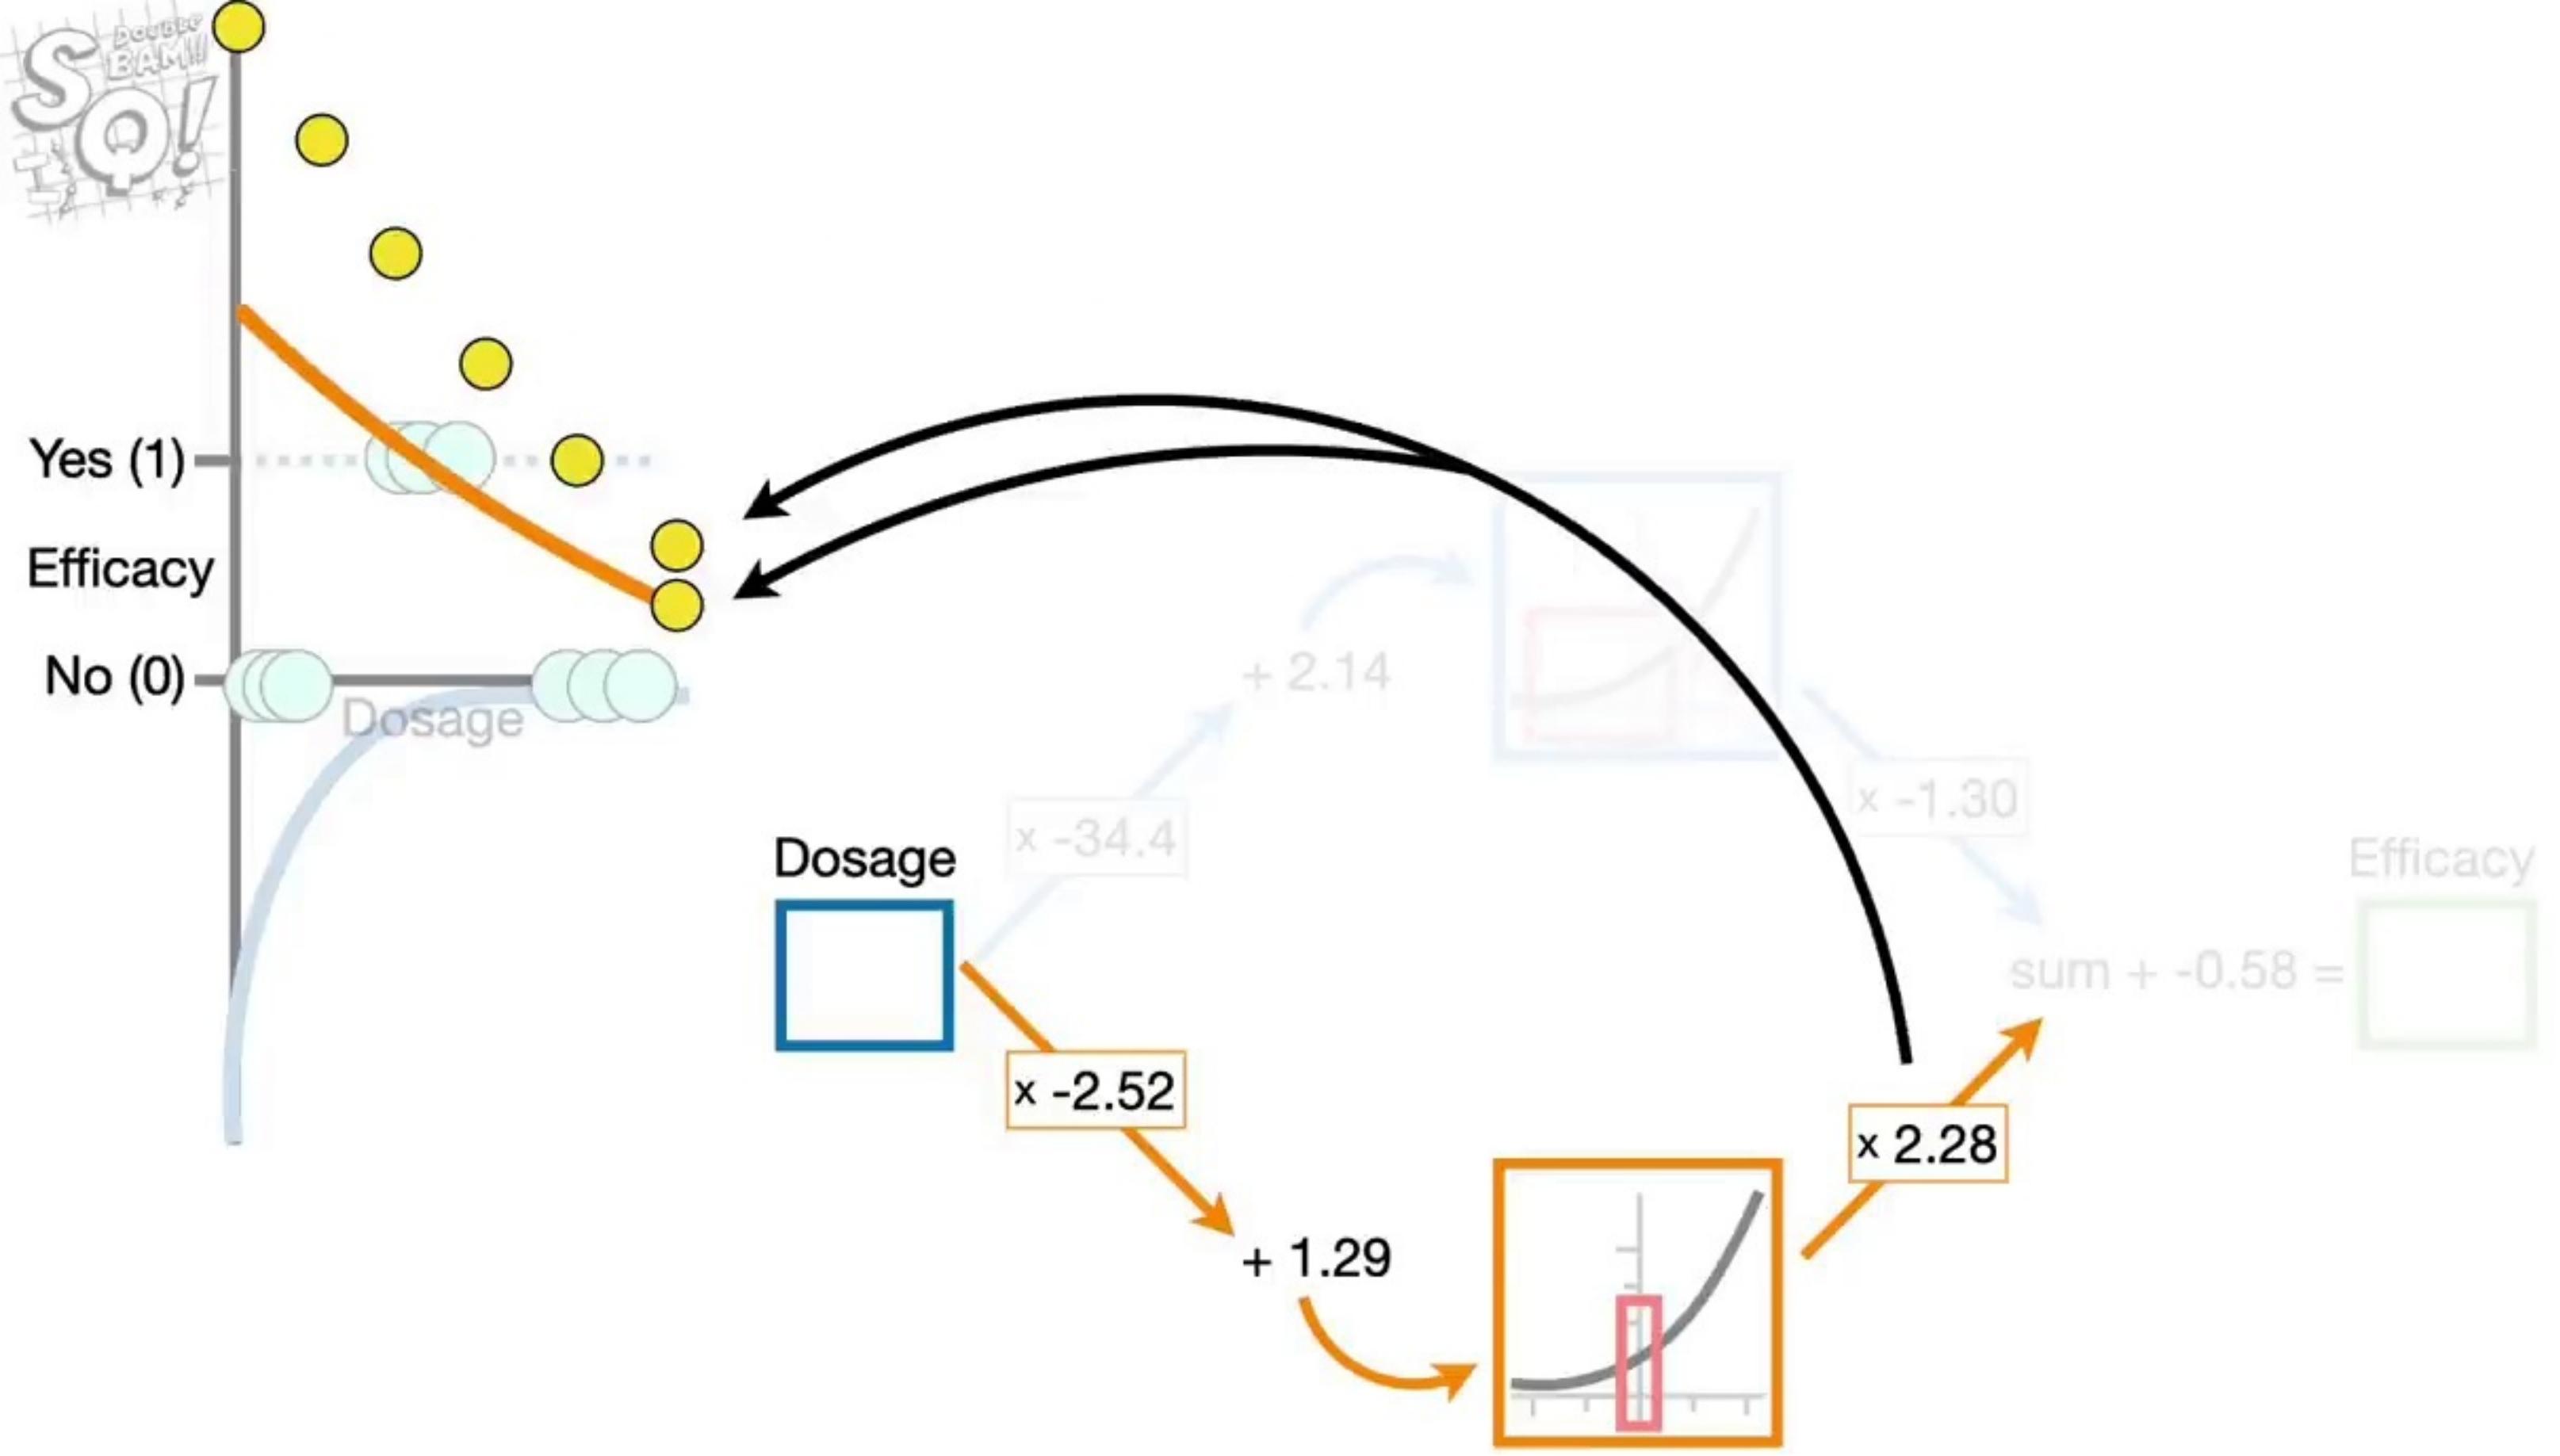


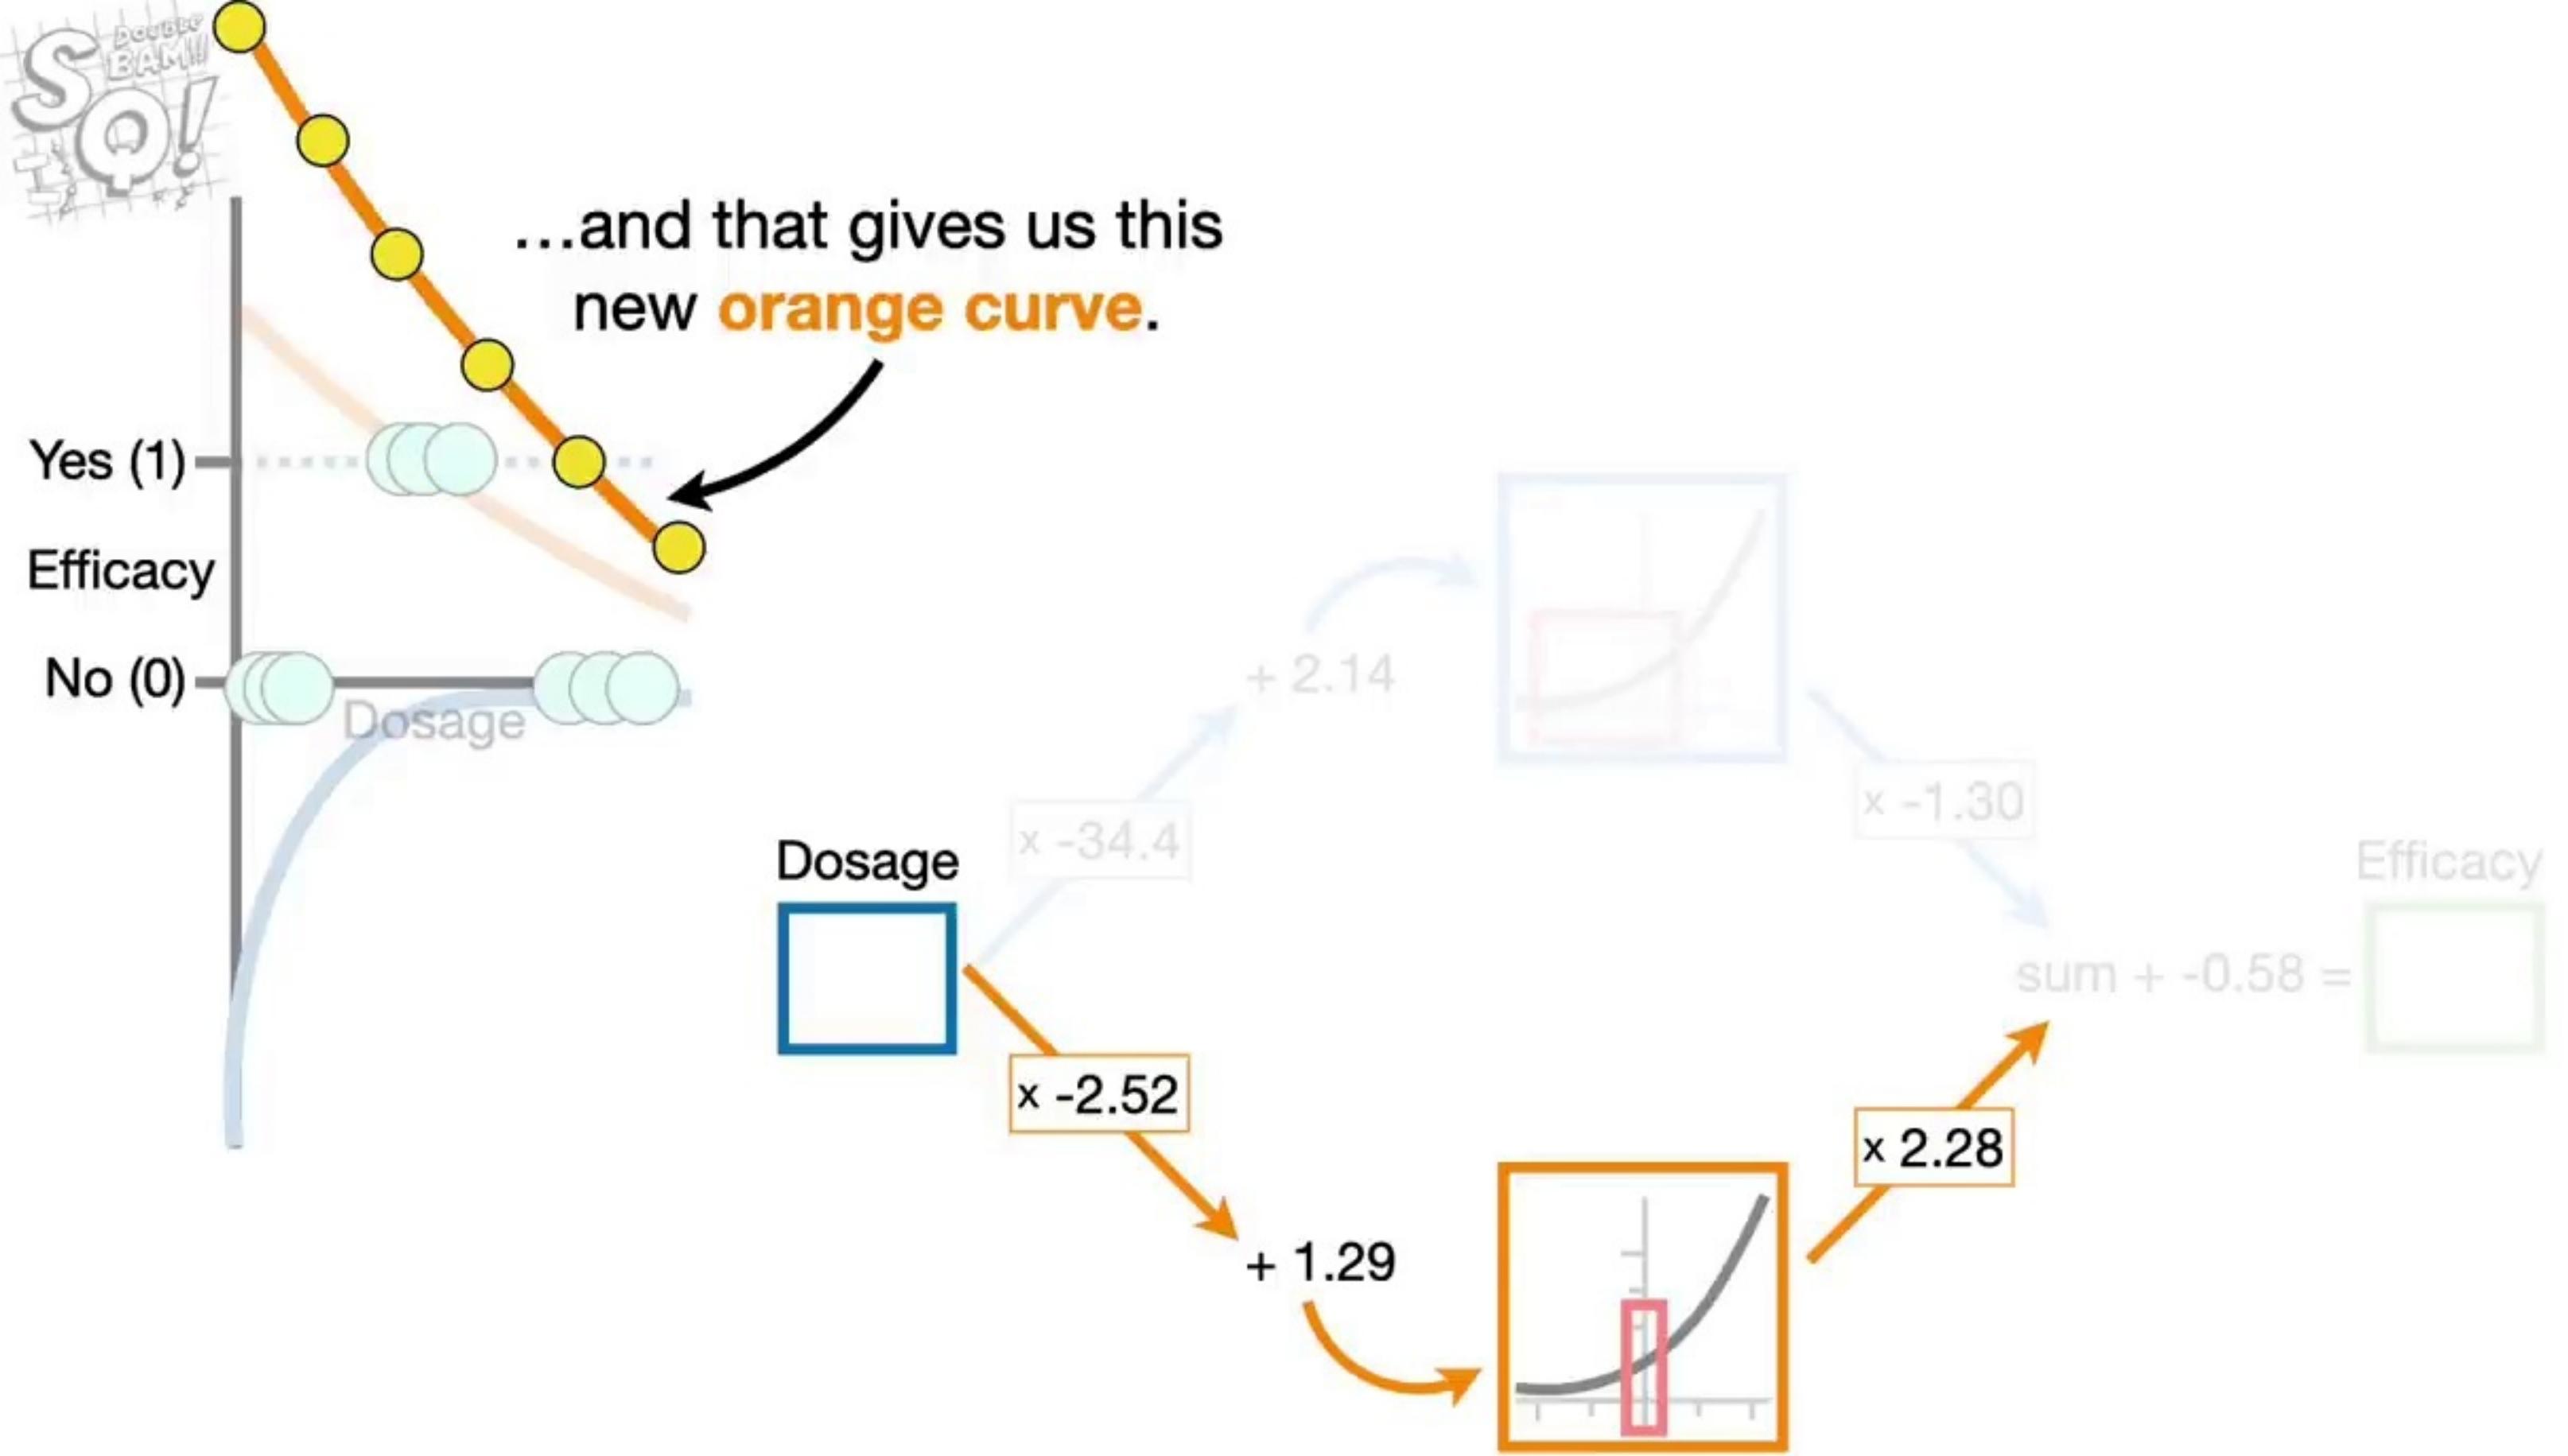






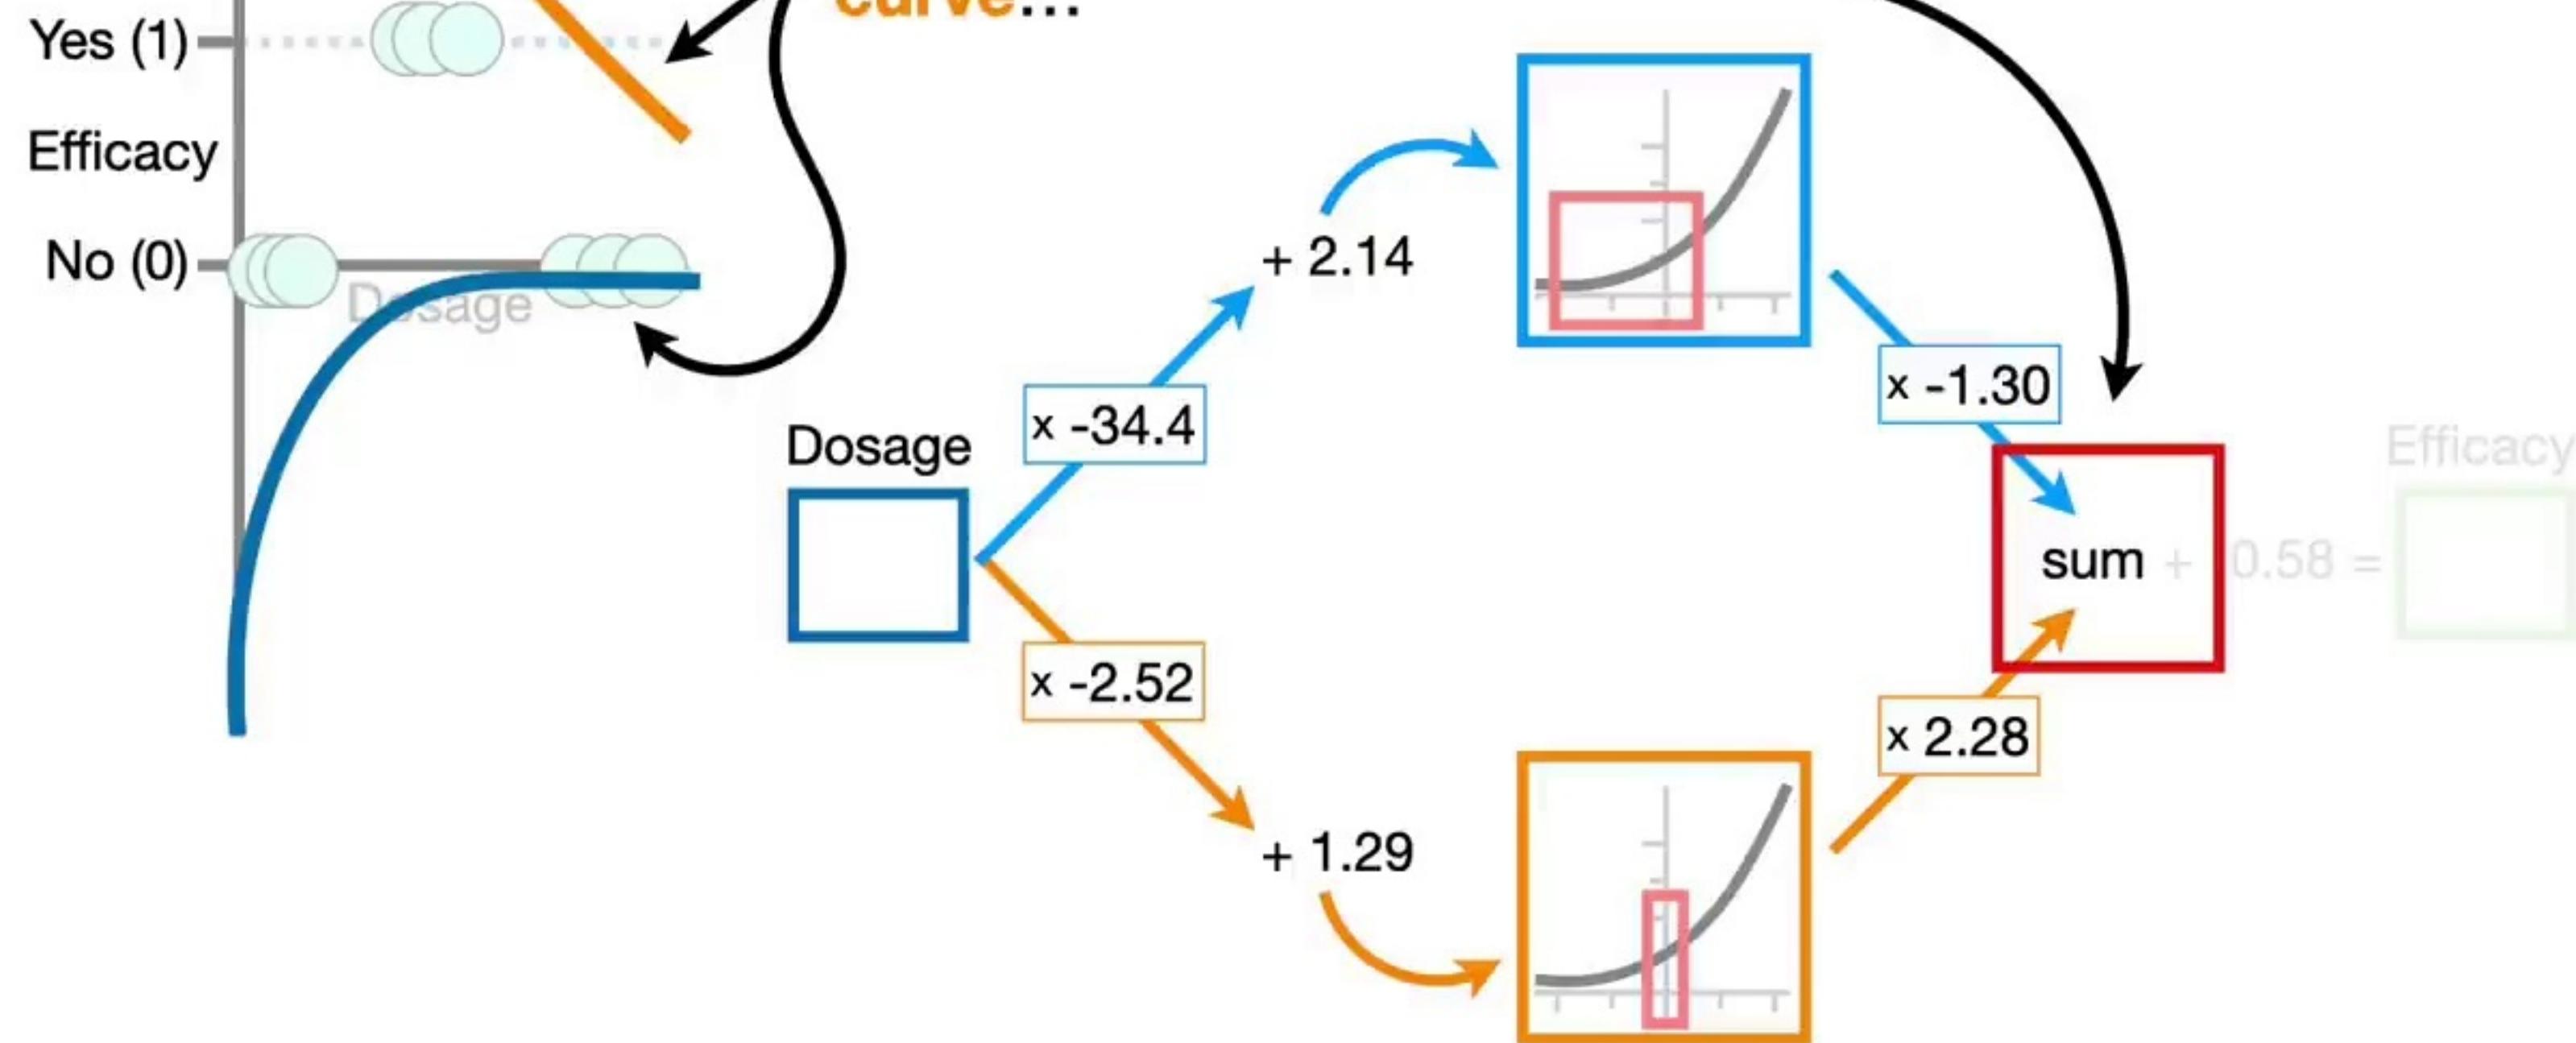






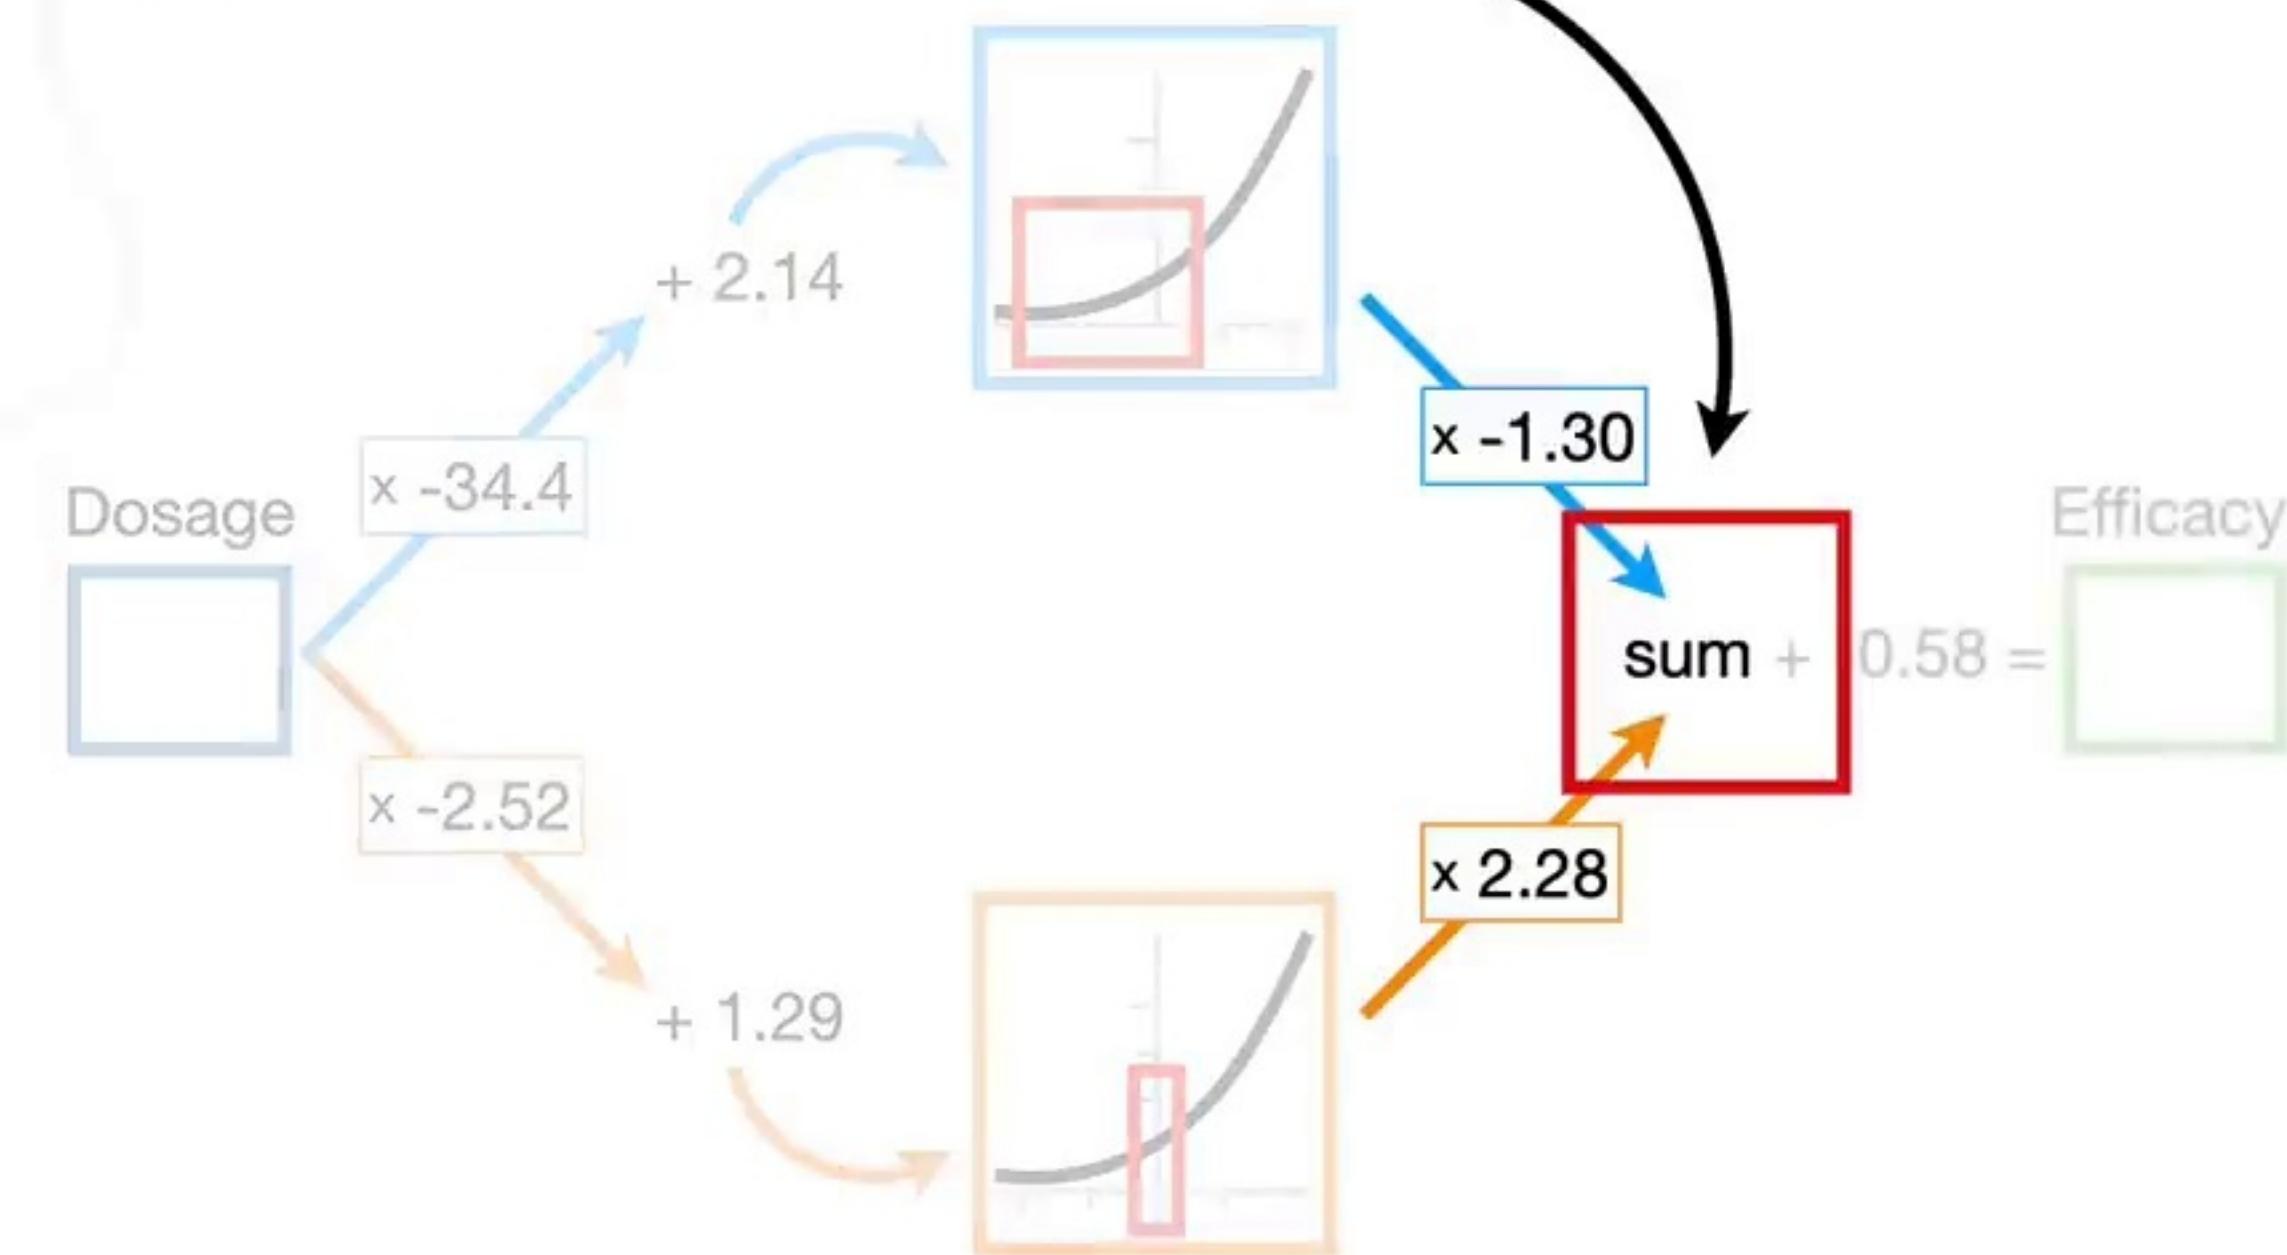
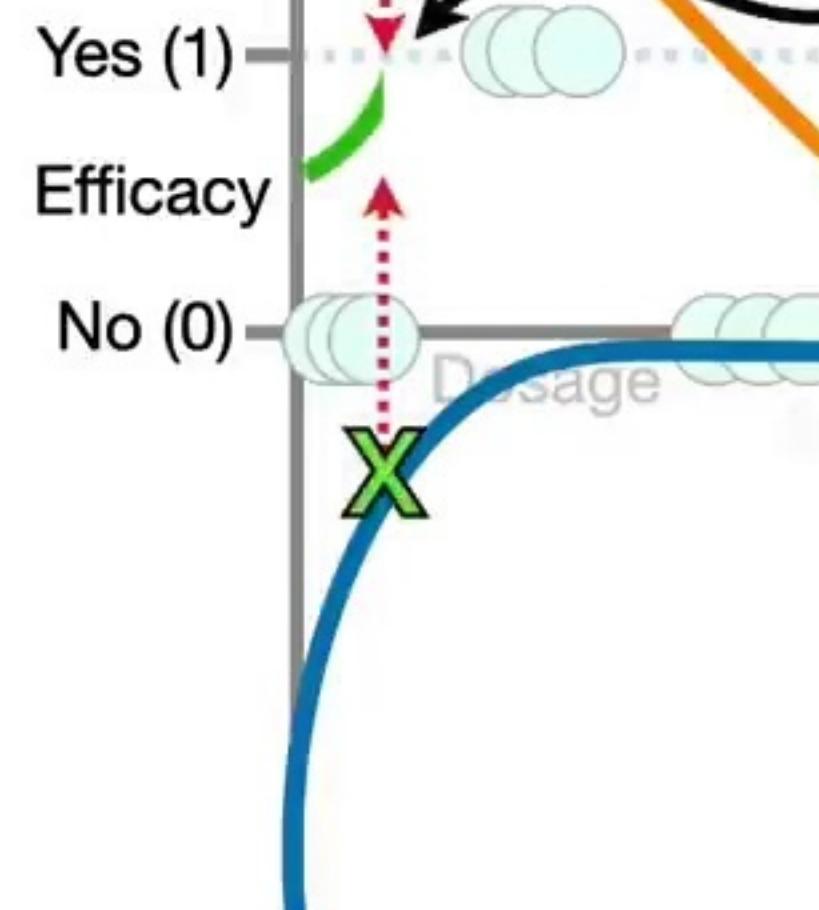


Now the **Neural Network**  
tells us to add the y-axis  
coordinates from **blue**  
**curve** to the **orange**  
**curve**...



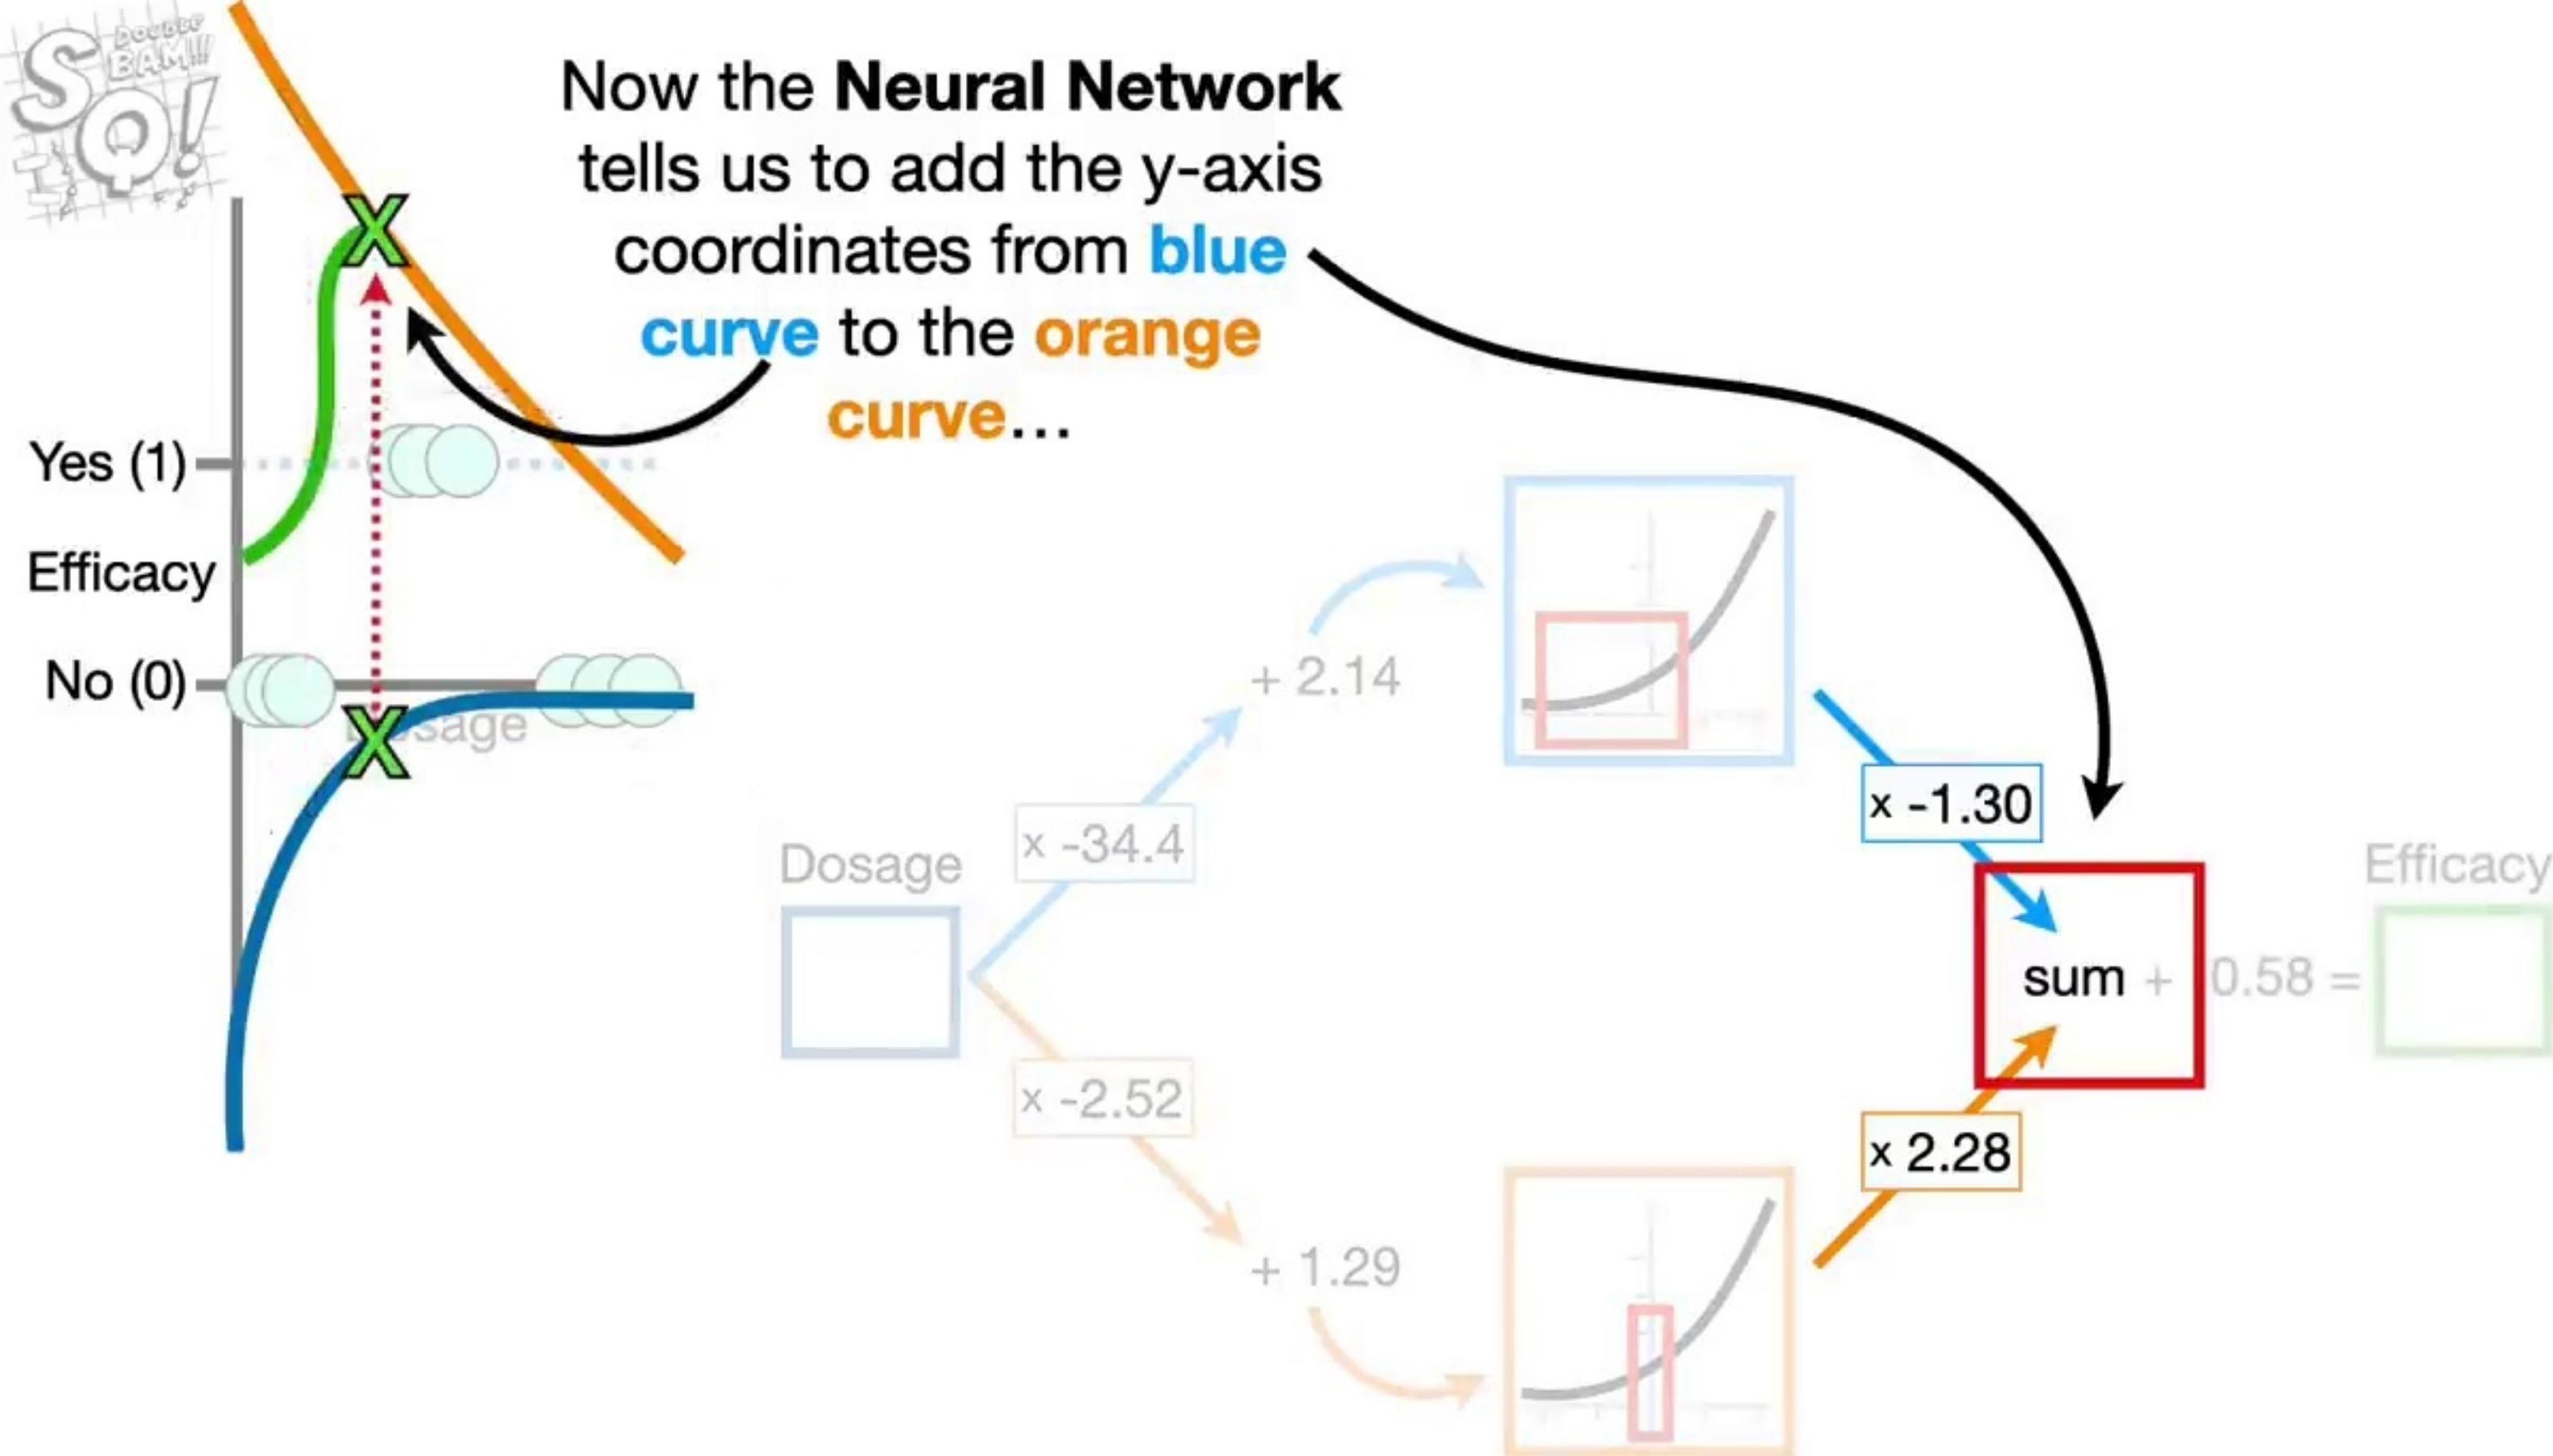
Double  
BAM!!  
**SQ!**

Now the **Neural Network**  
tells us to add the y-axis  
coordinates from **blue**  
**curve** to the **orange**  
**curve**...



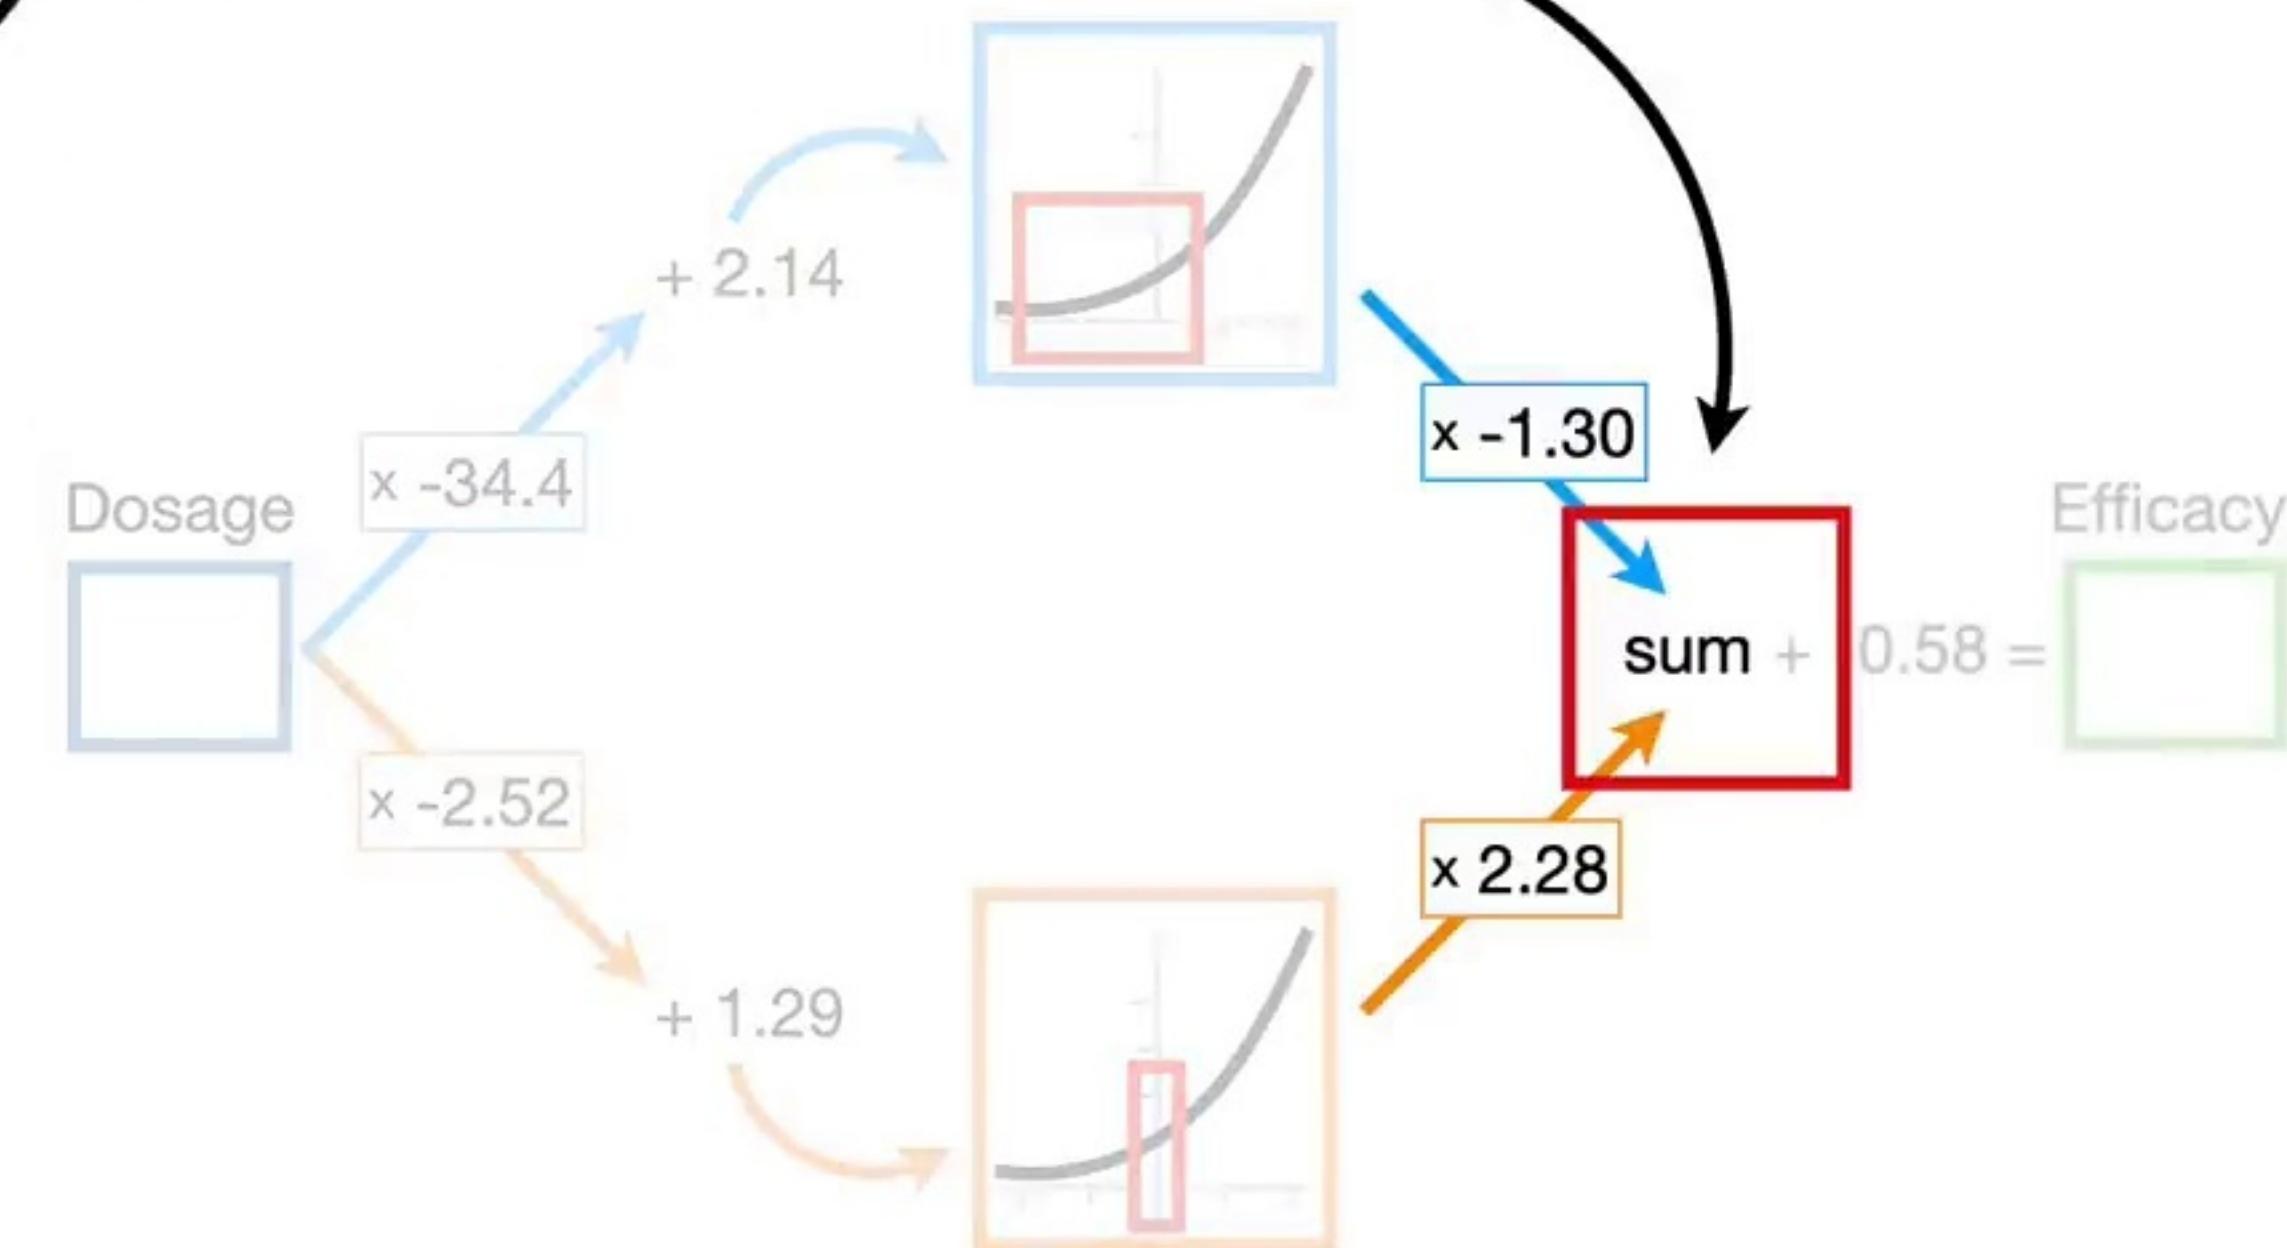
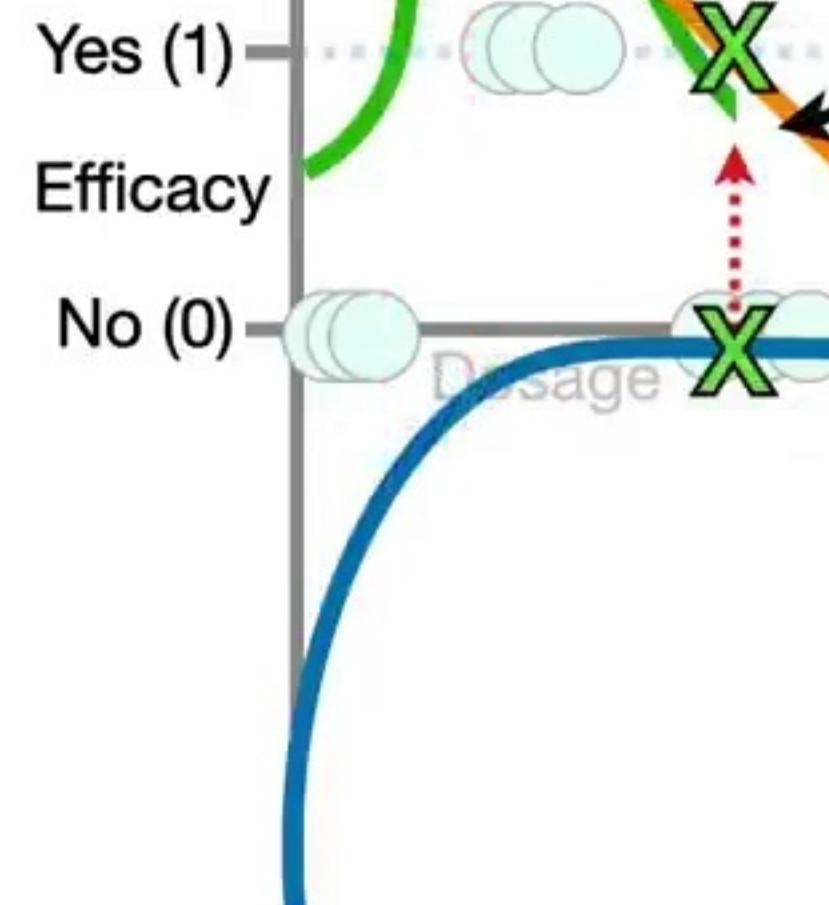
Double  
BAM!!  
**SQ!**

Now the **Neural Network**  
tells us to add the y-axis  
coordinates from **blue**  
**curve** to the **orange**  
**curve**...



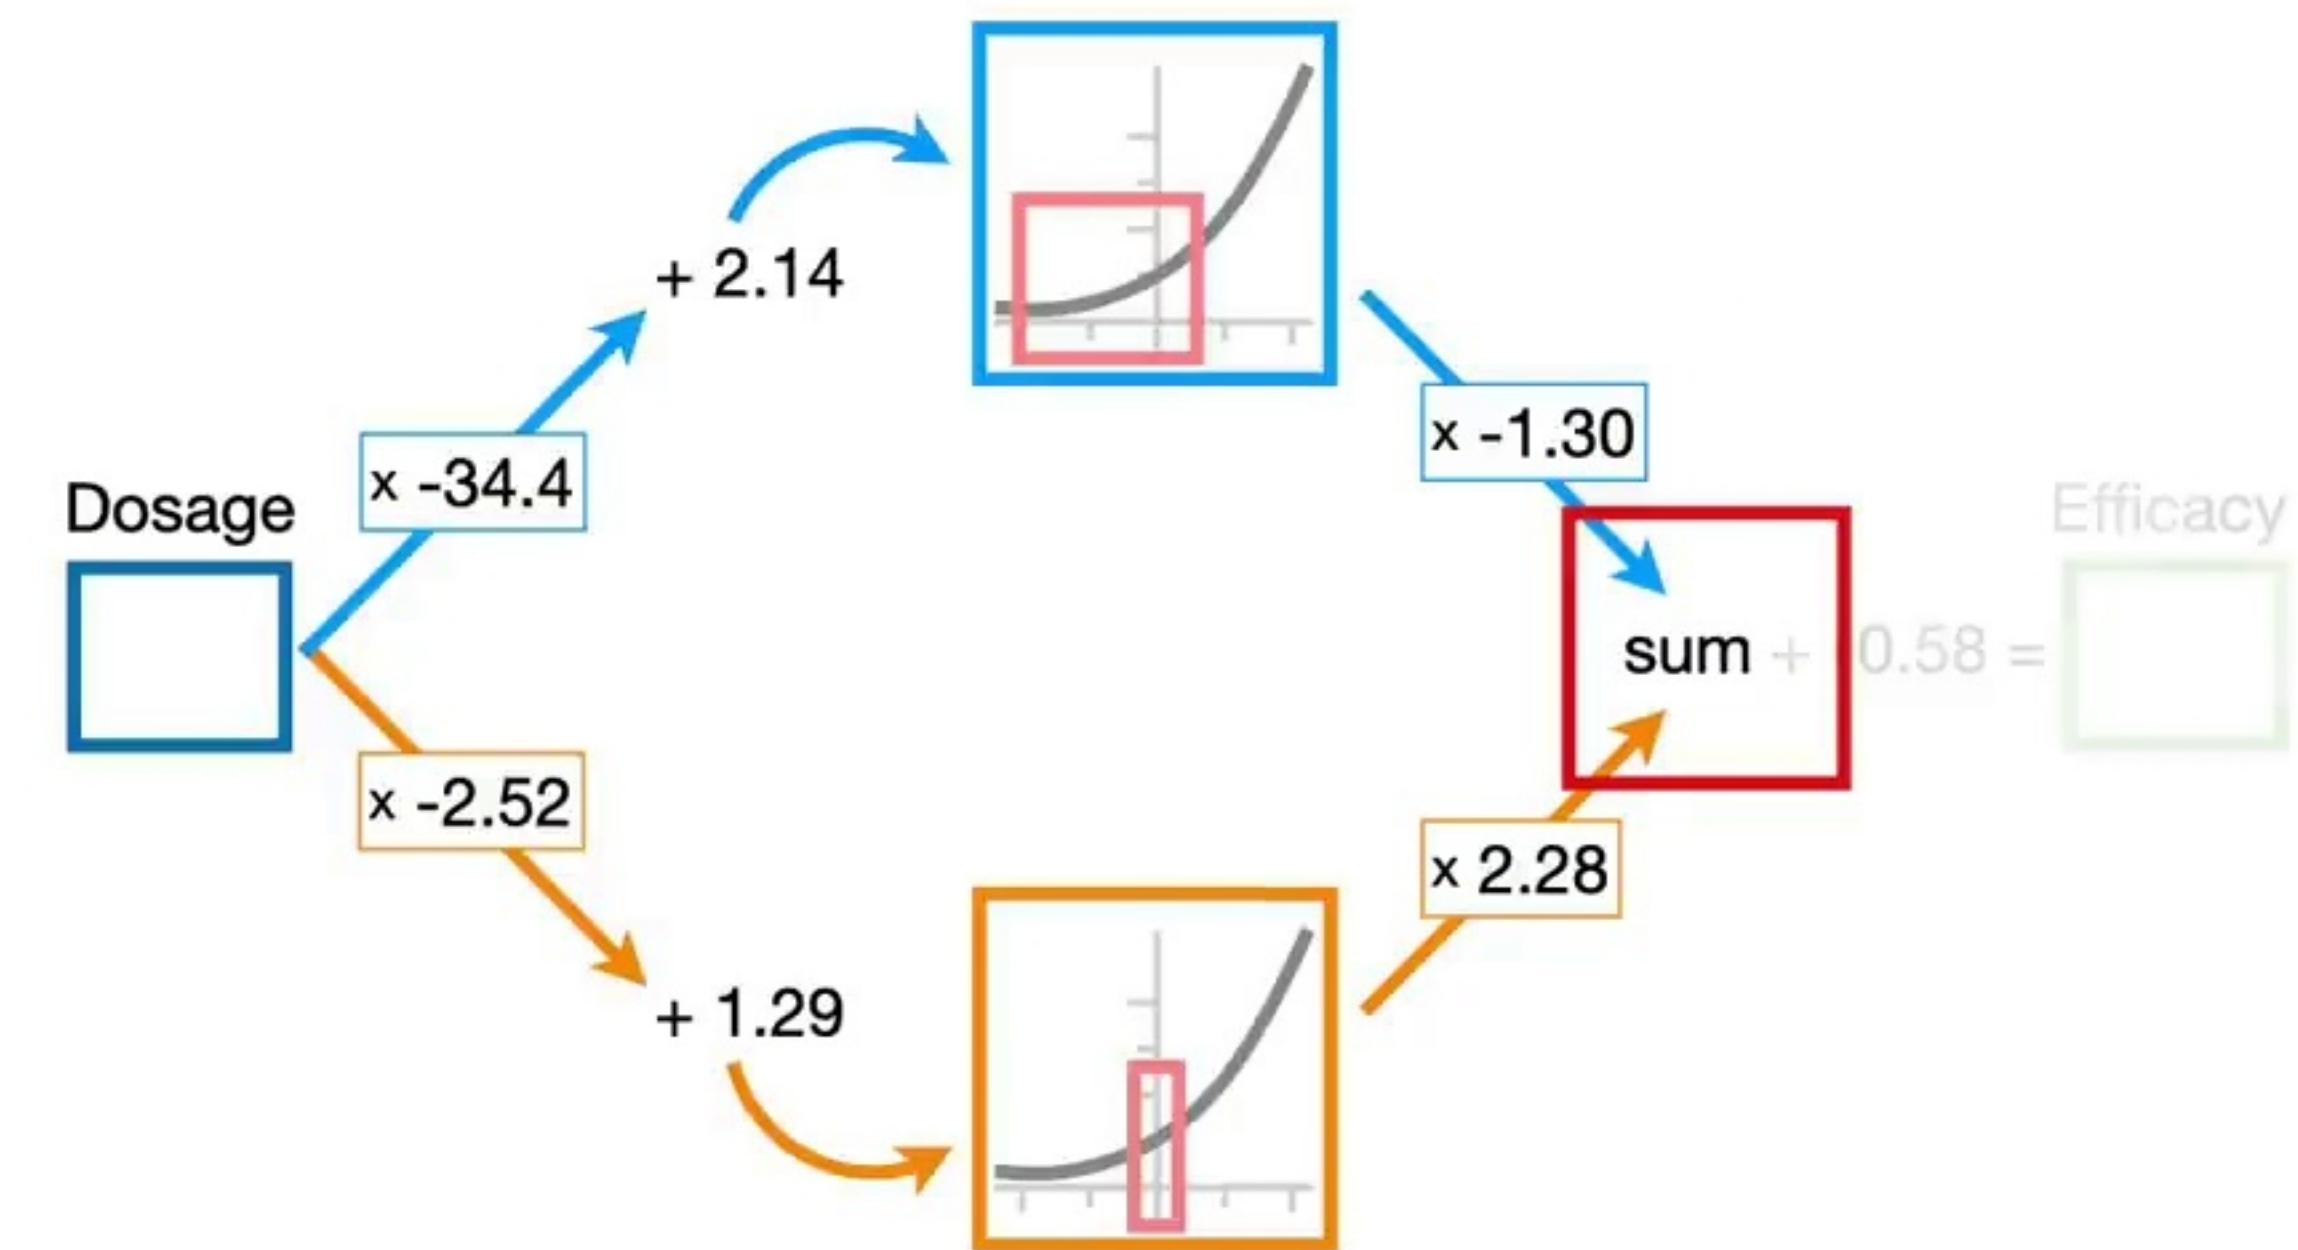
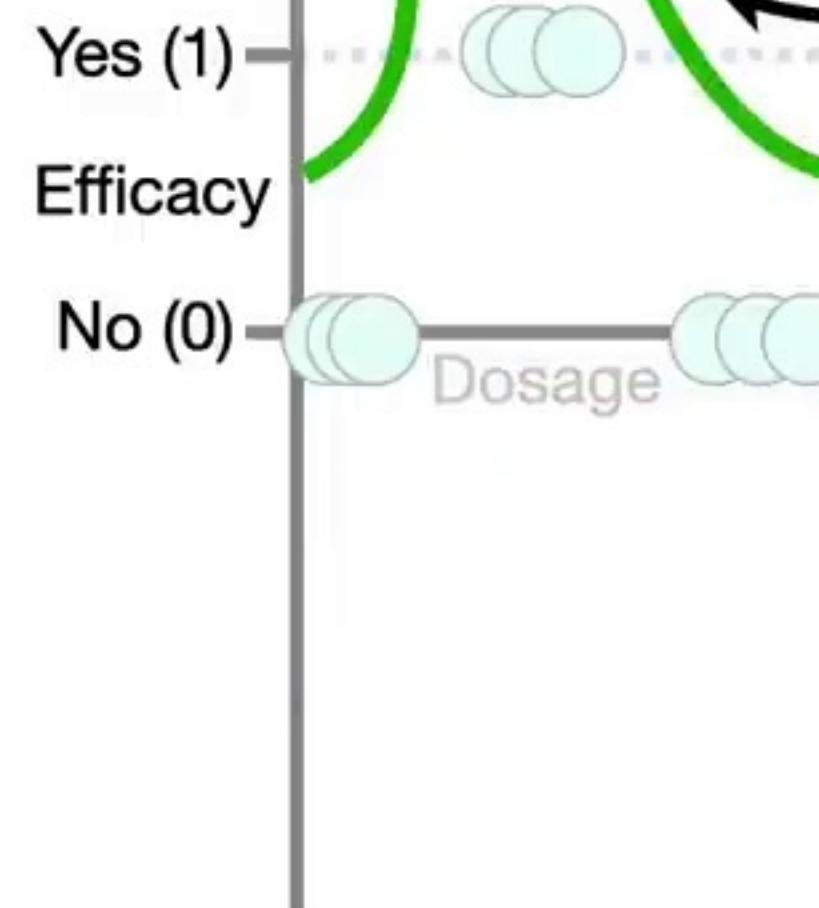
Double  
BAM!!  
**SQ!**

Now the **Neural Network**  
tells us to add the y-axis  
coordinates from **blue**  
**curve** to the **orange**  
**curve**...



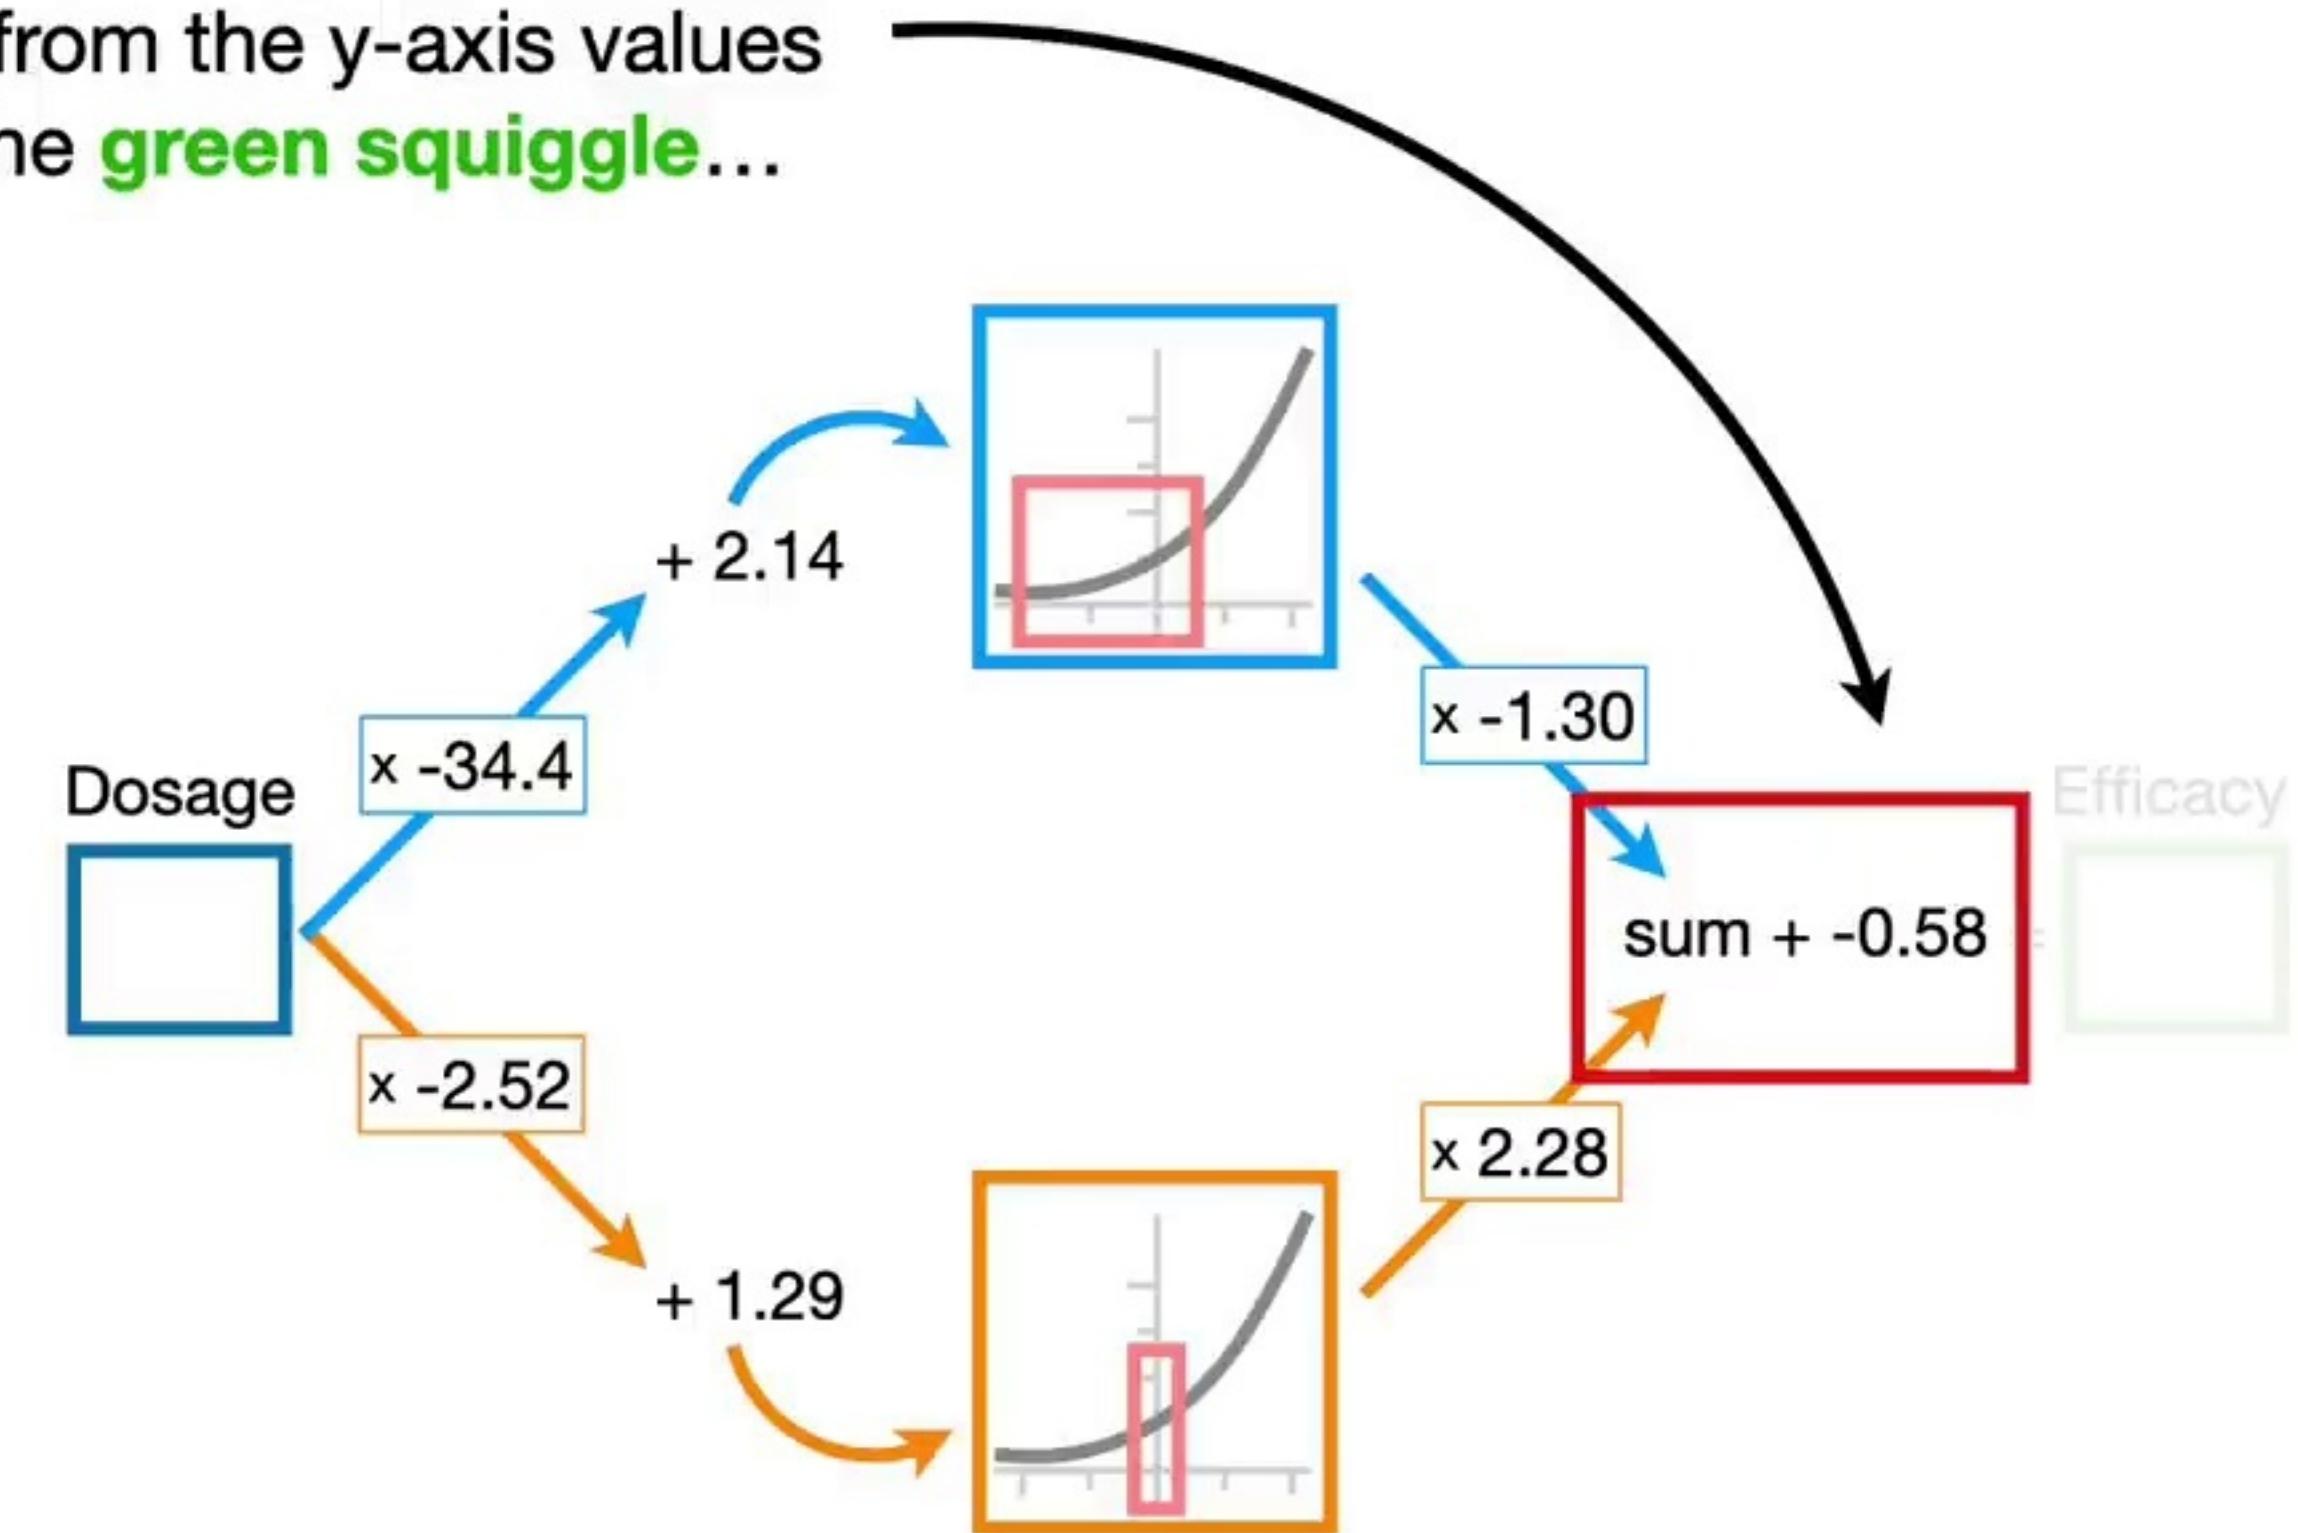
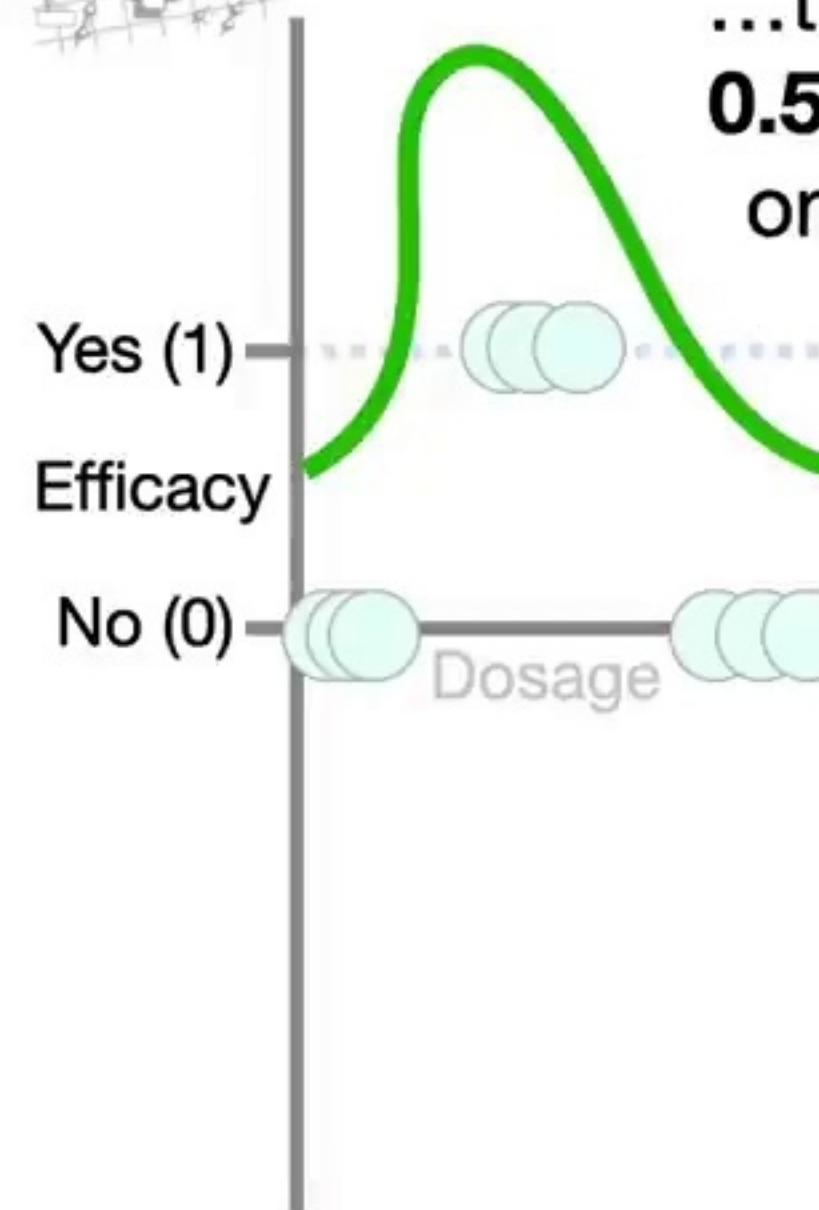
Double  
BAM!!  
**SO!**

...and that gives us this  
**green squiggle**...



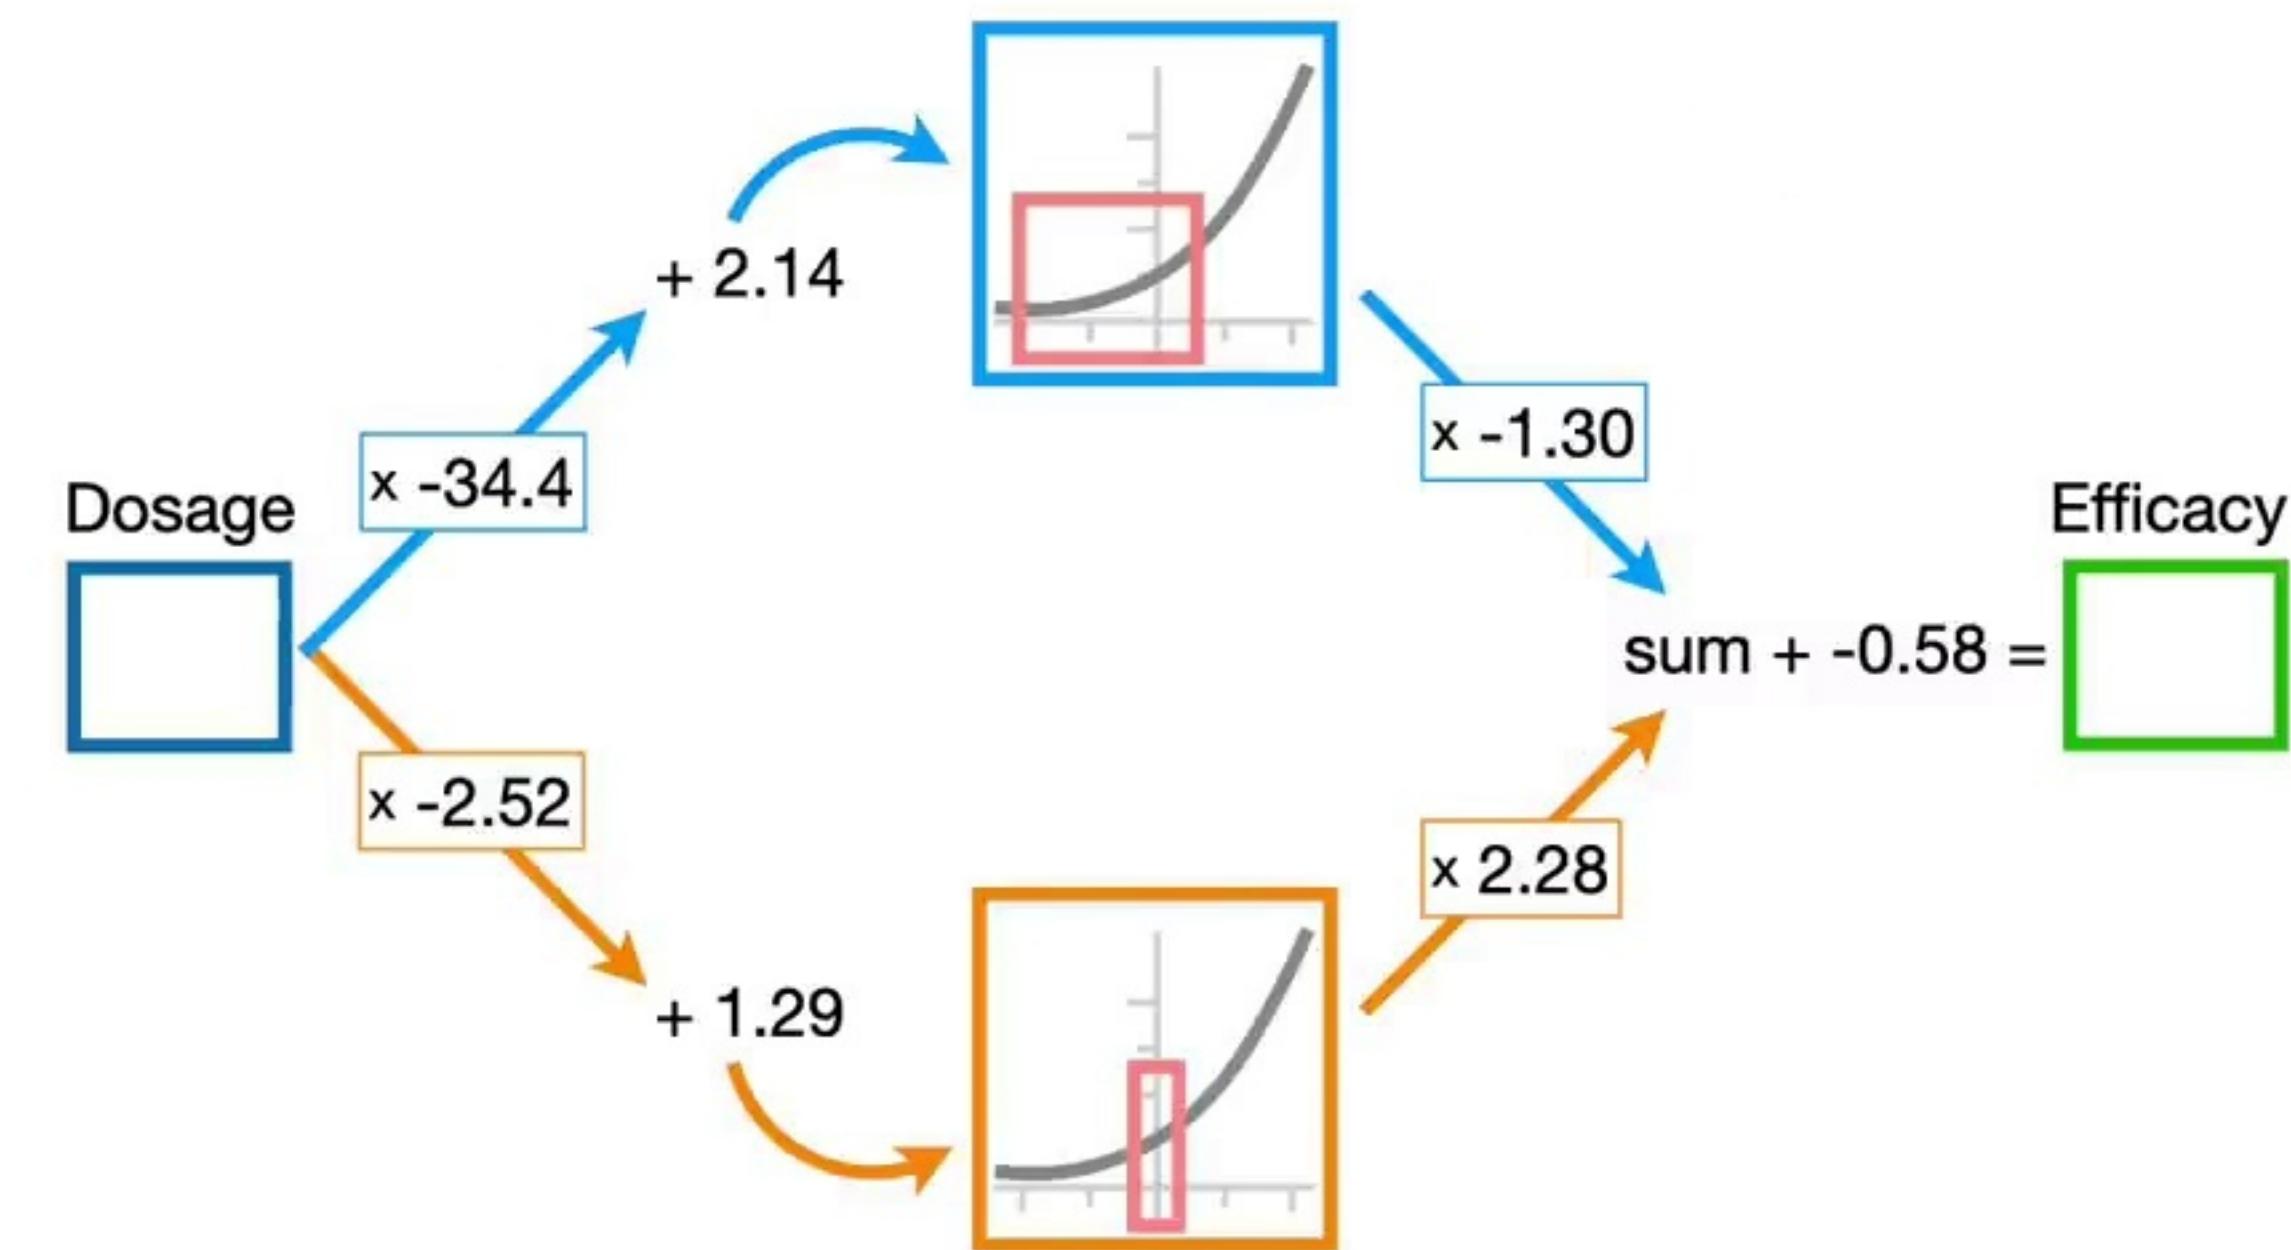
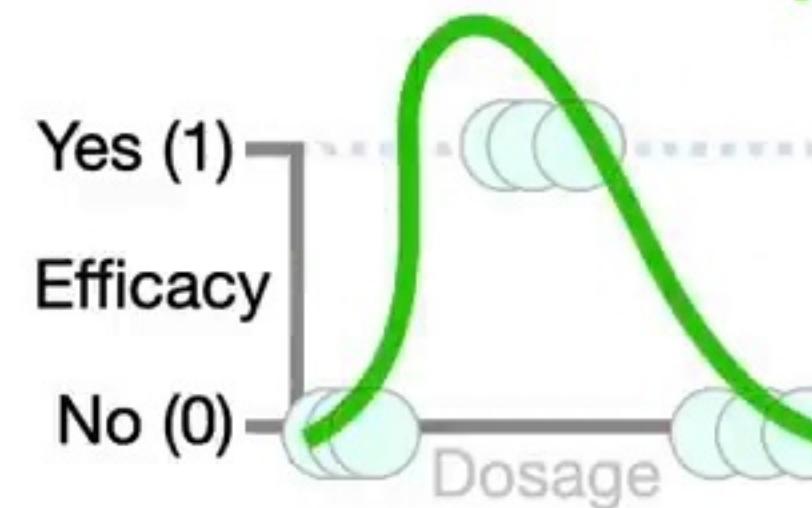


...then, finally, we subtract,  
**0.58** from the y-axis values  
on the **green squiggle**...



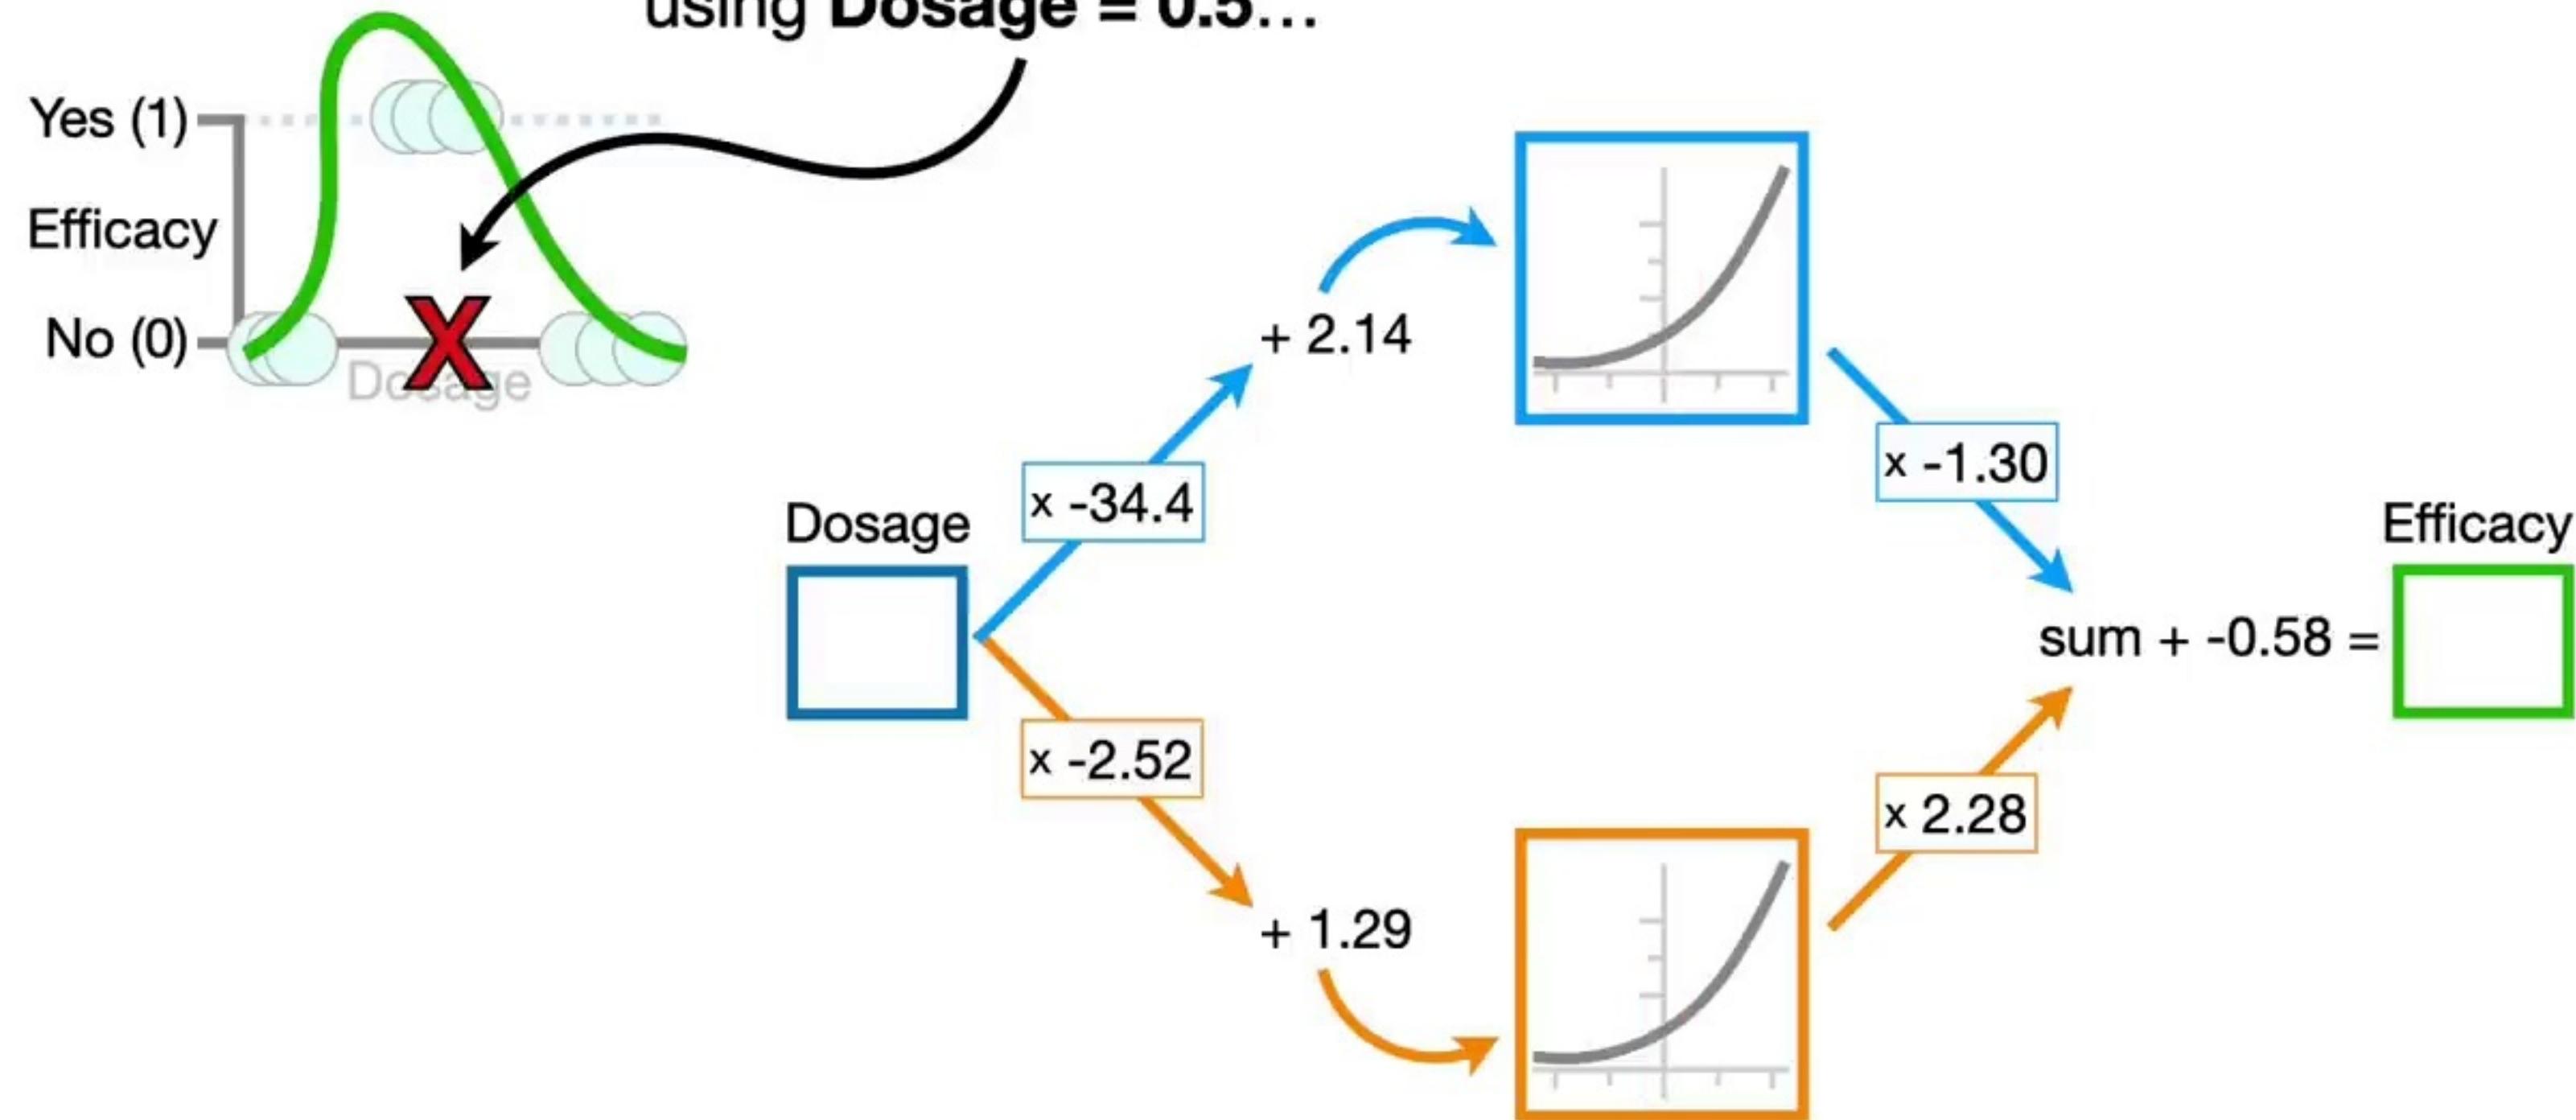


...and we have a **green squiggle** that fits the data.



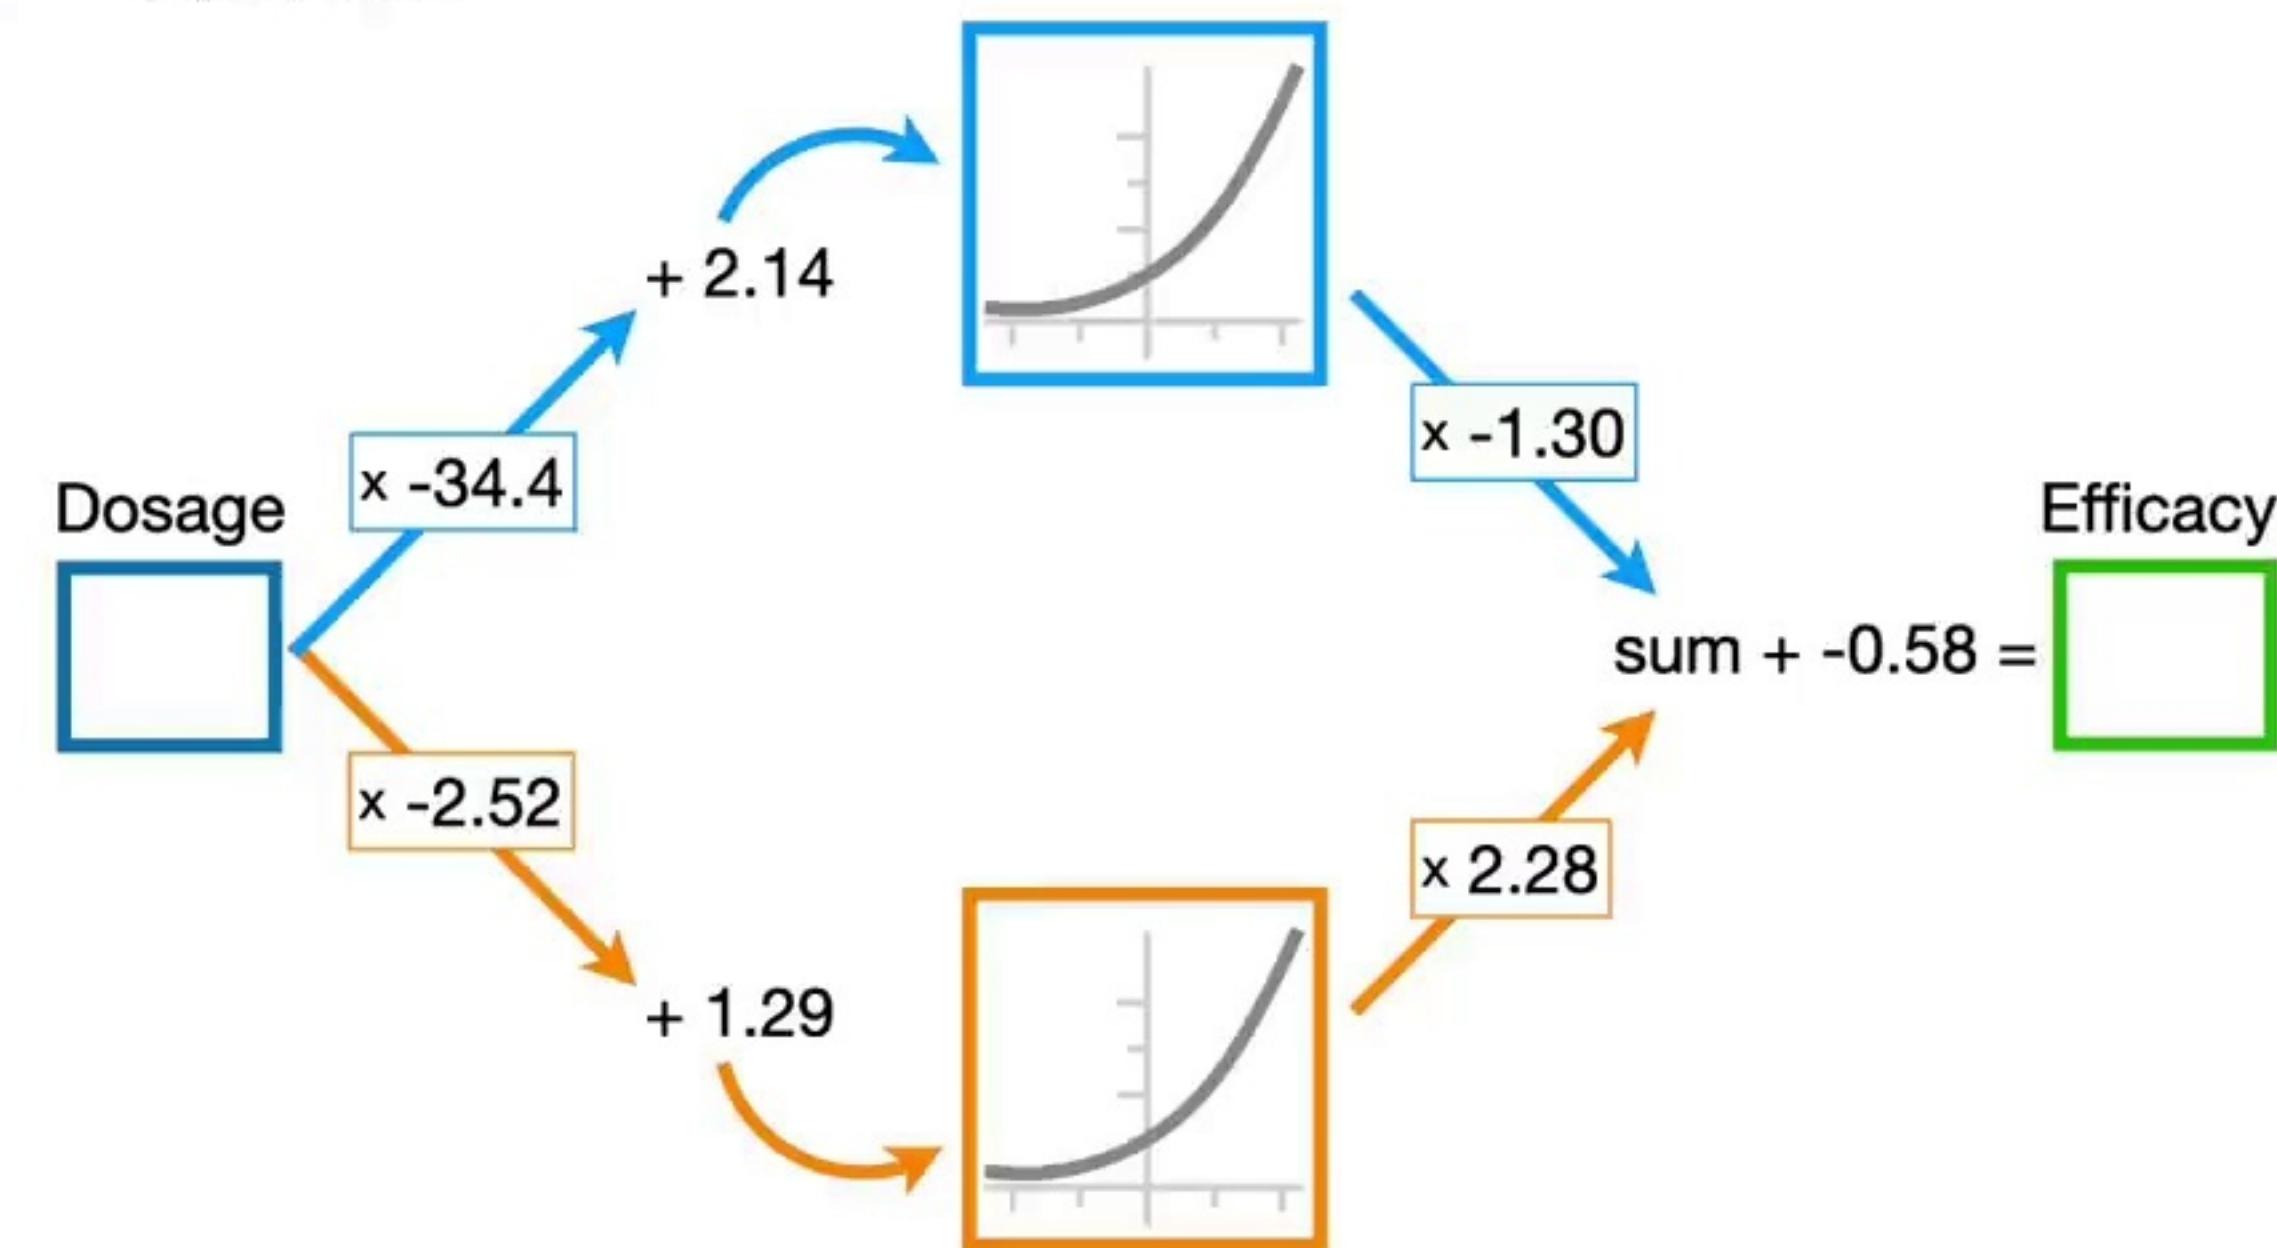
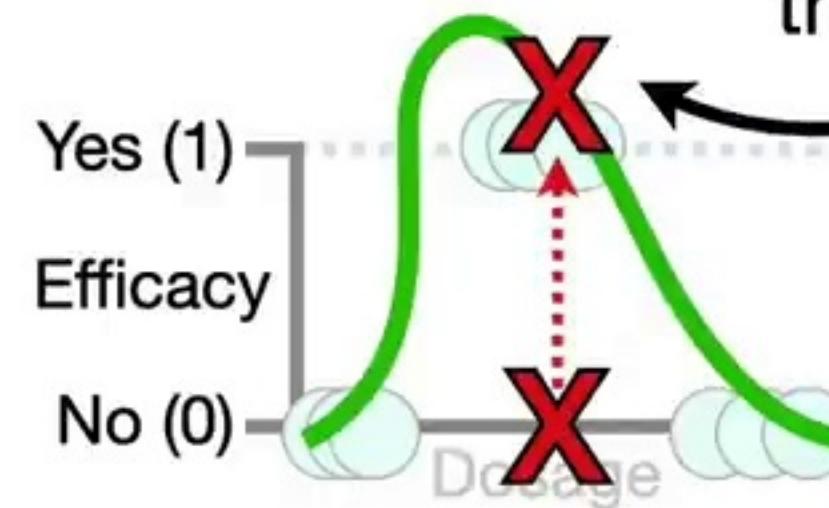


Now, if someone comes along and says that they are using **Dosage = 0.5...**



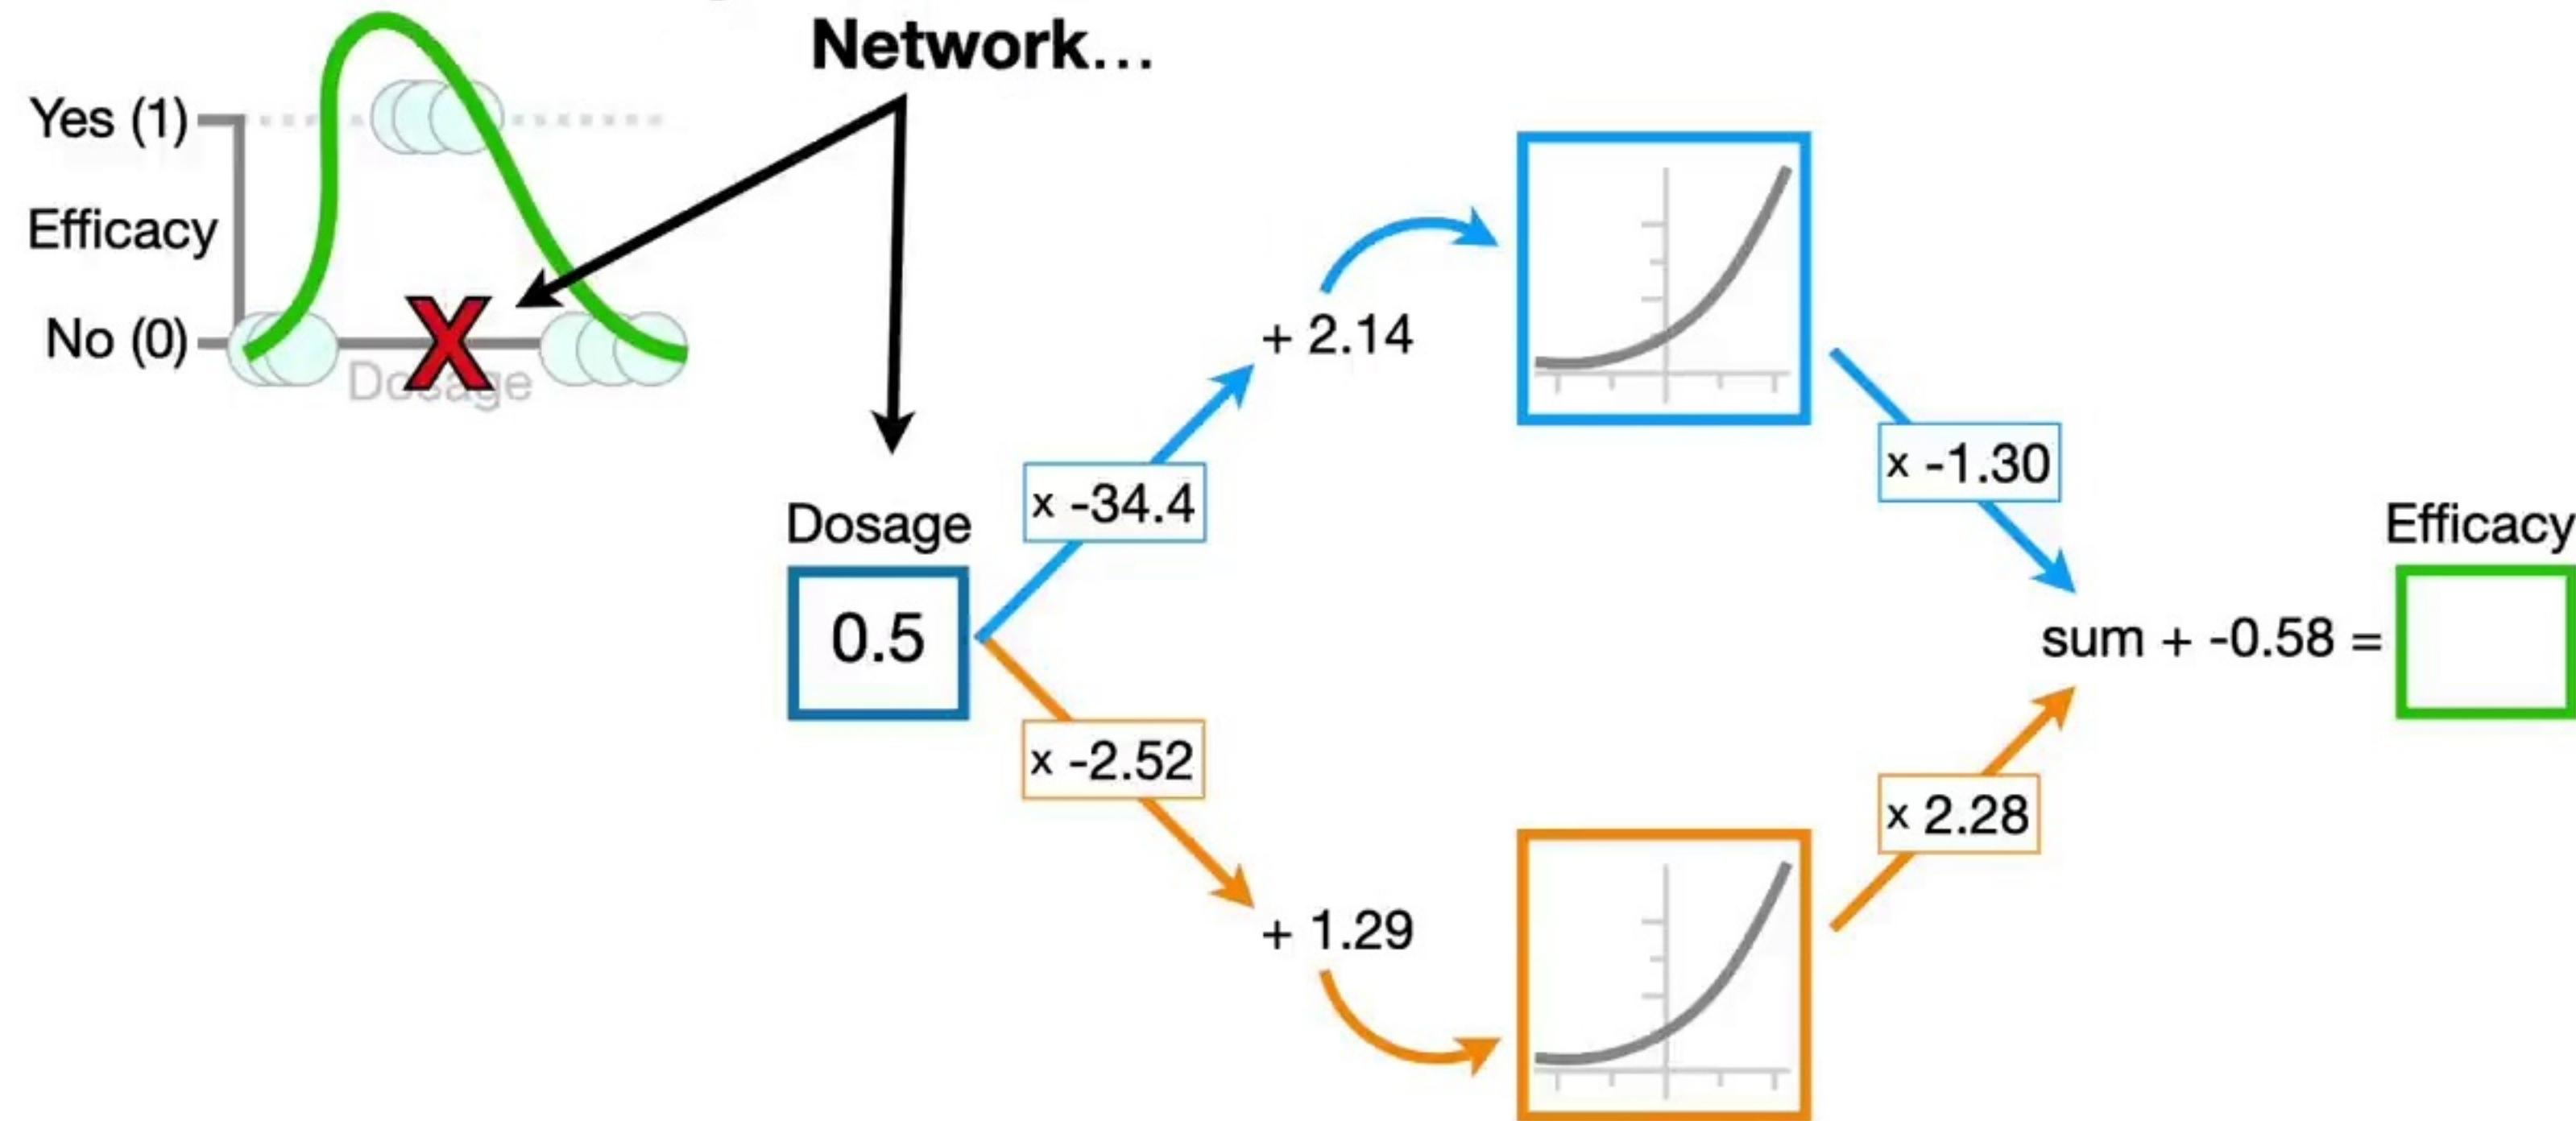


...we can look at the corresponding y-axis coordinate on the **green squiggle** and see that the **Dosage** will be effective.



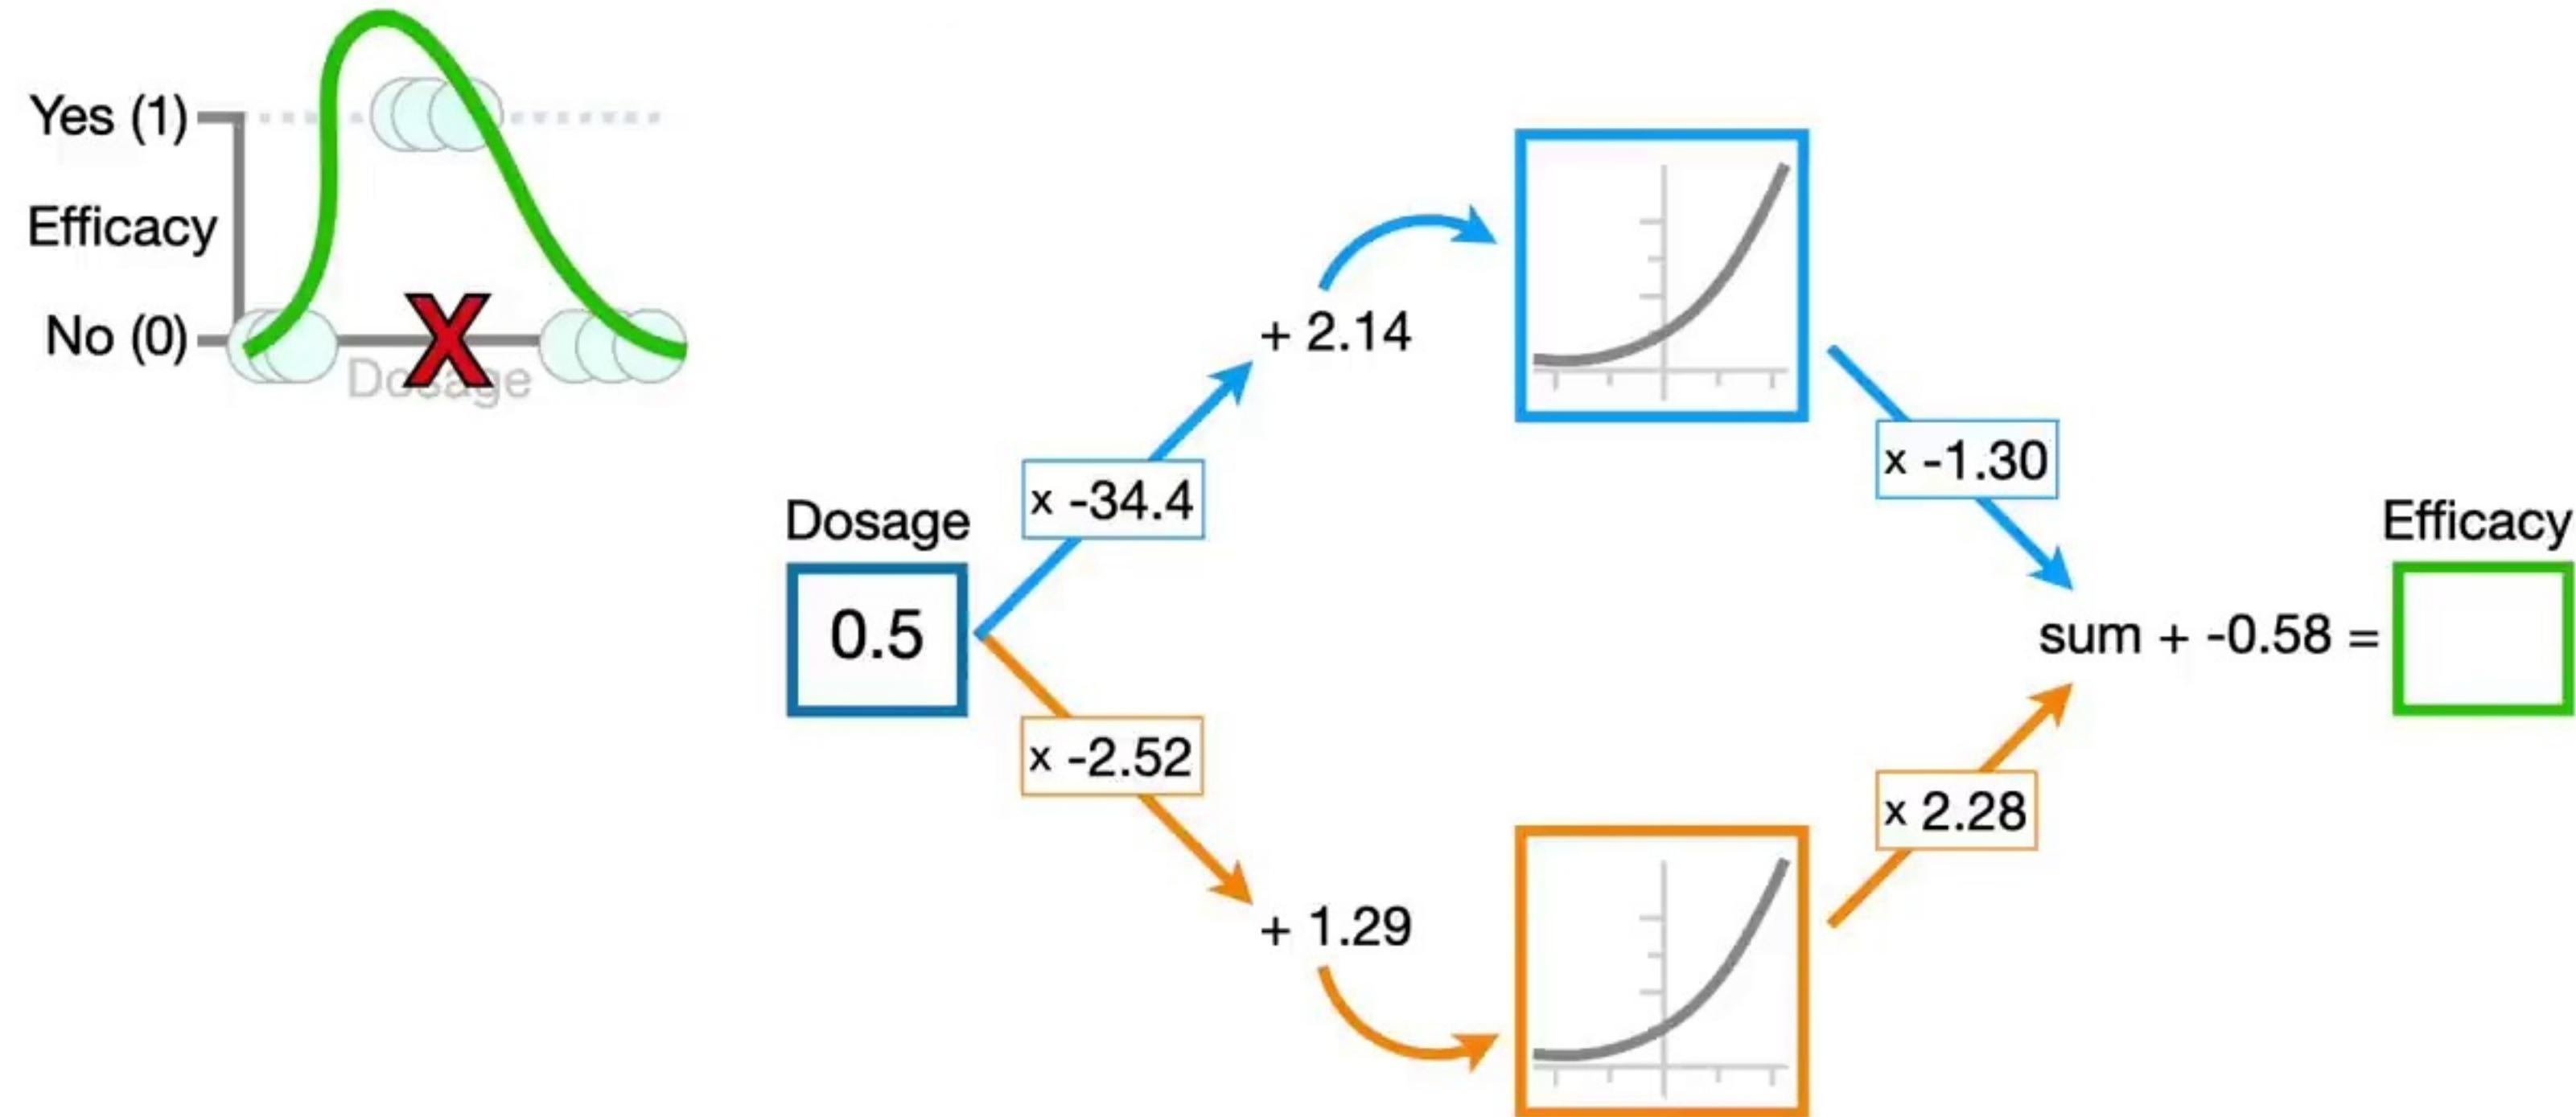


Or we can solve for the y-axis coordinate by plugging  
**Dosage = 0.5** into the **Neural Network...**



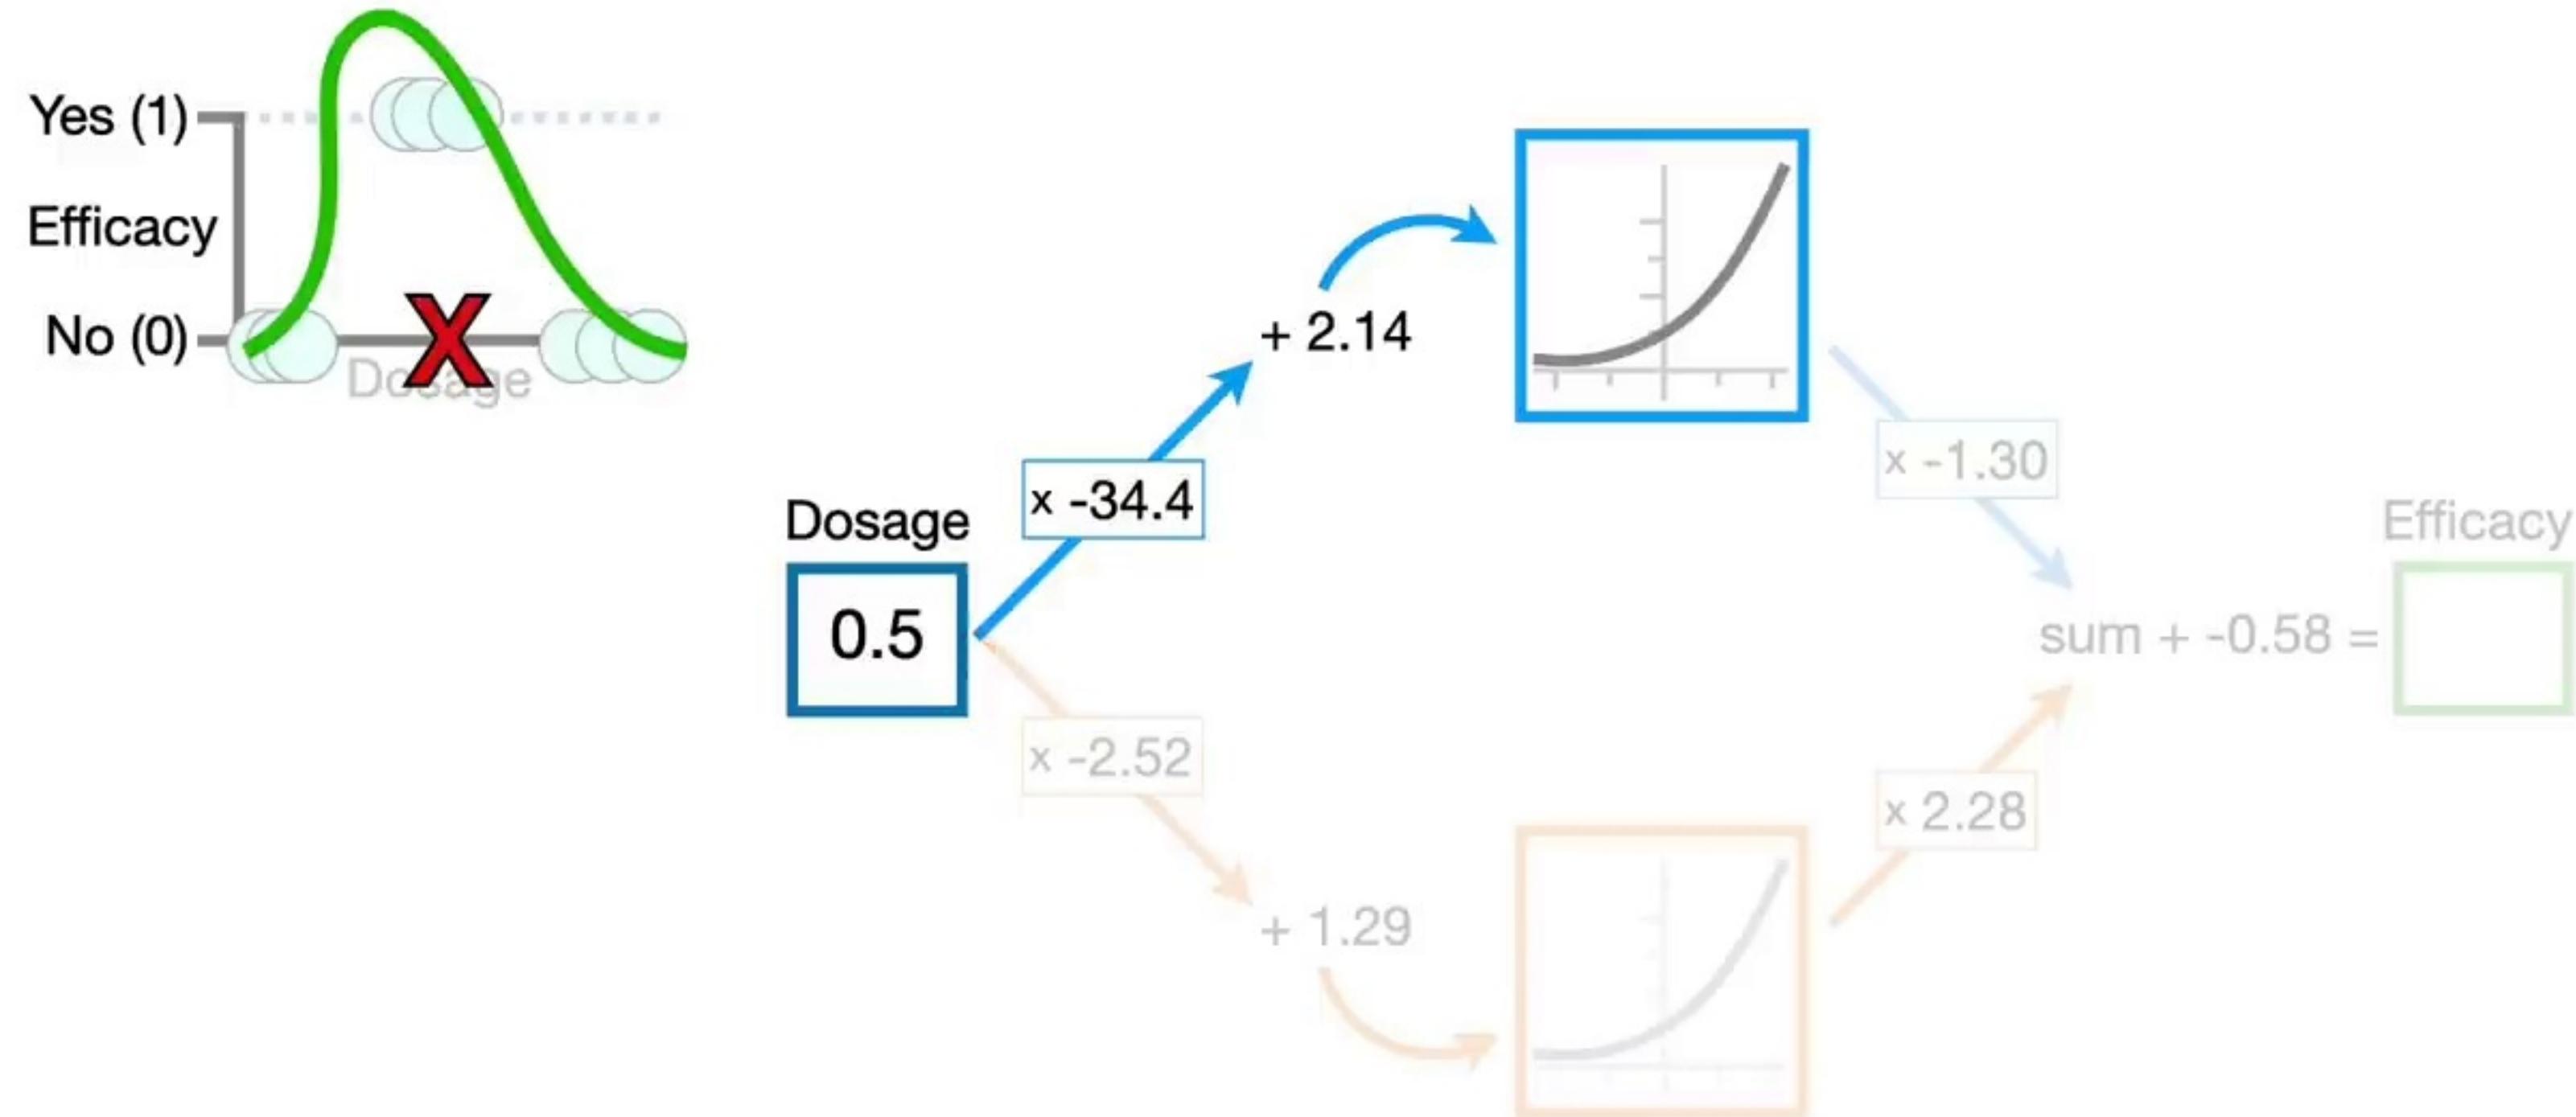


...and do the math...



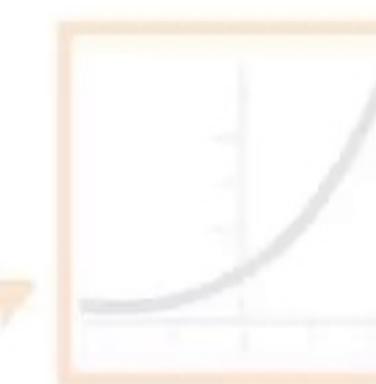
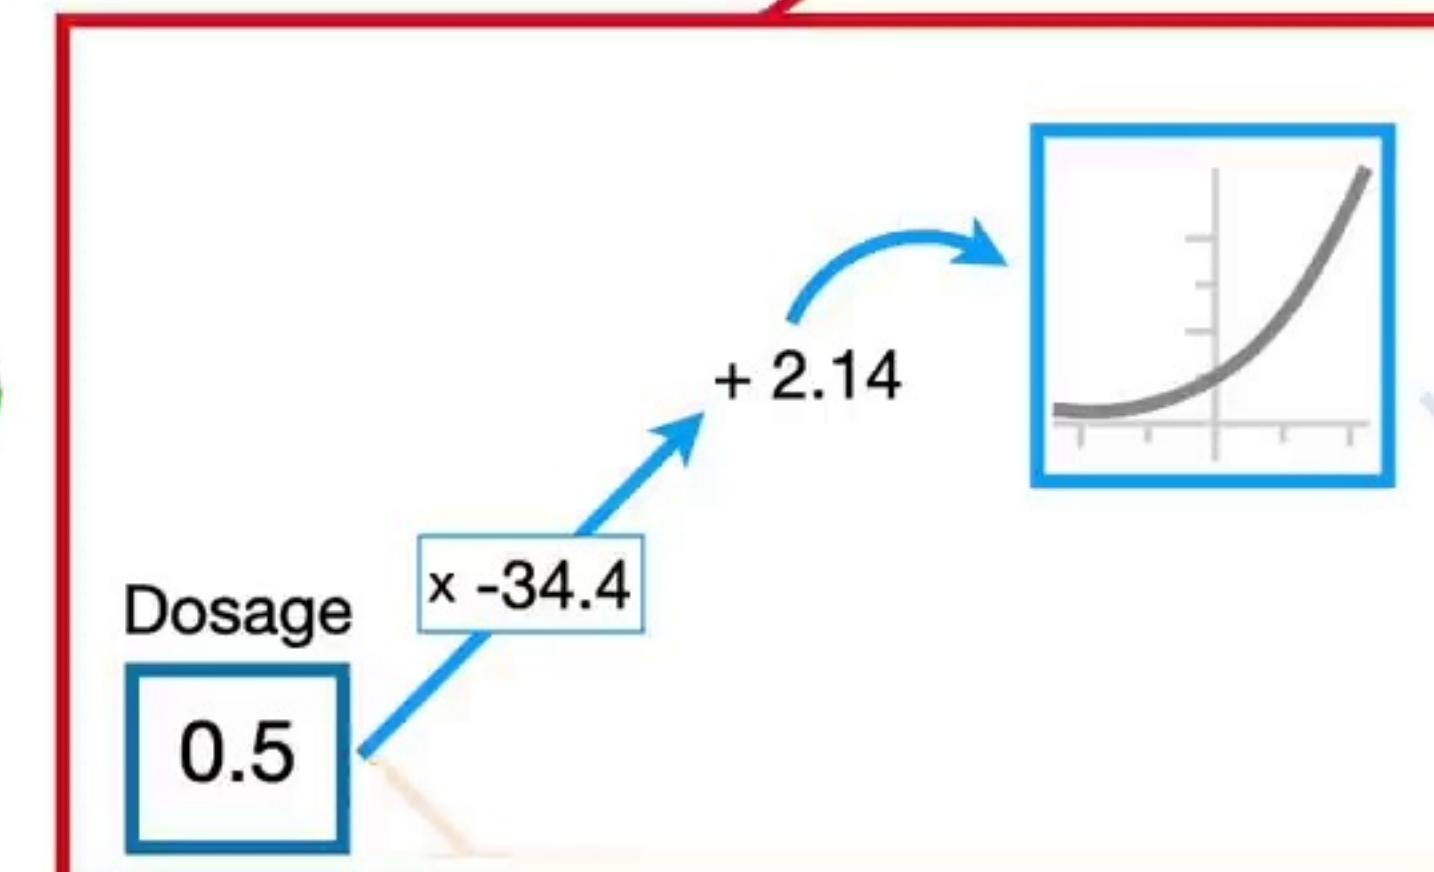
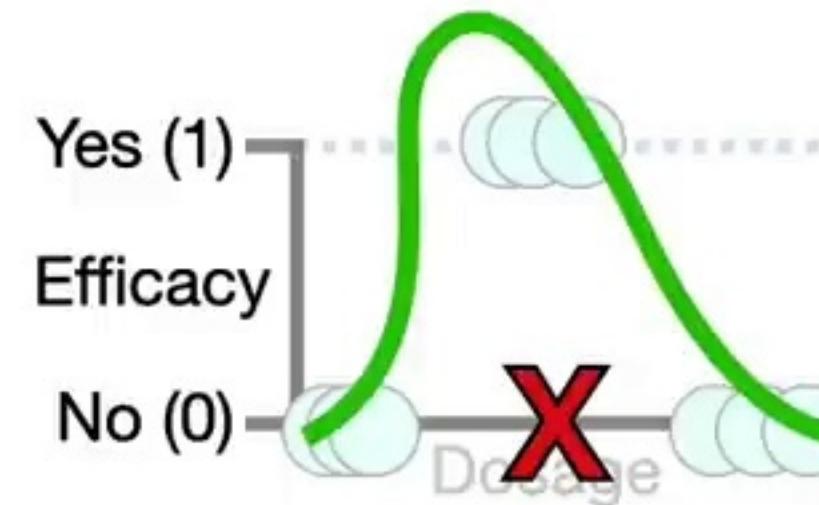


...and do the math...



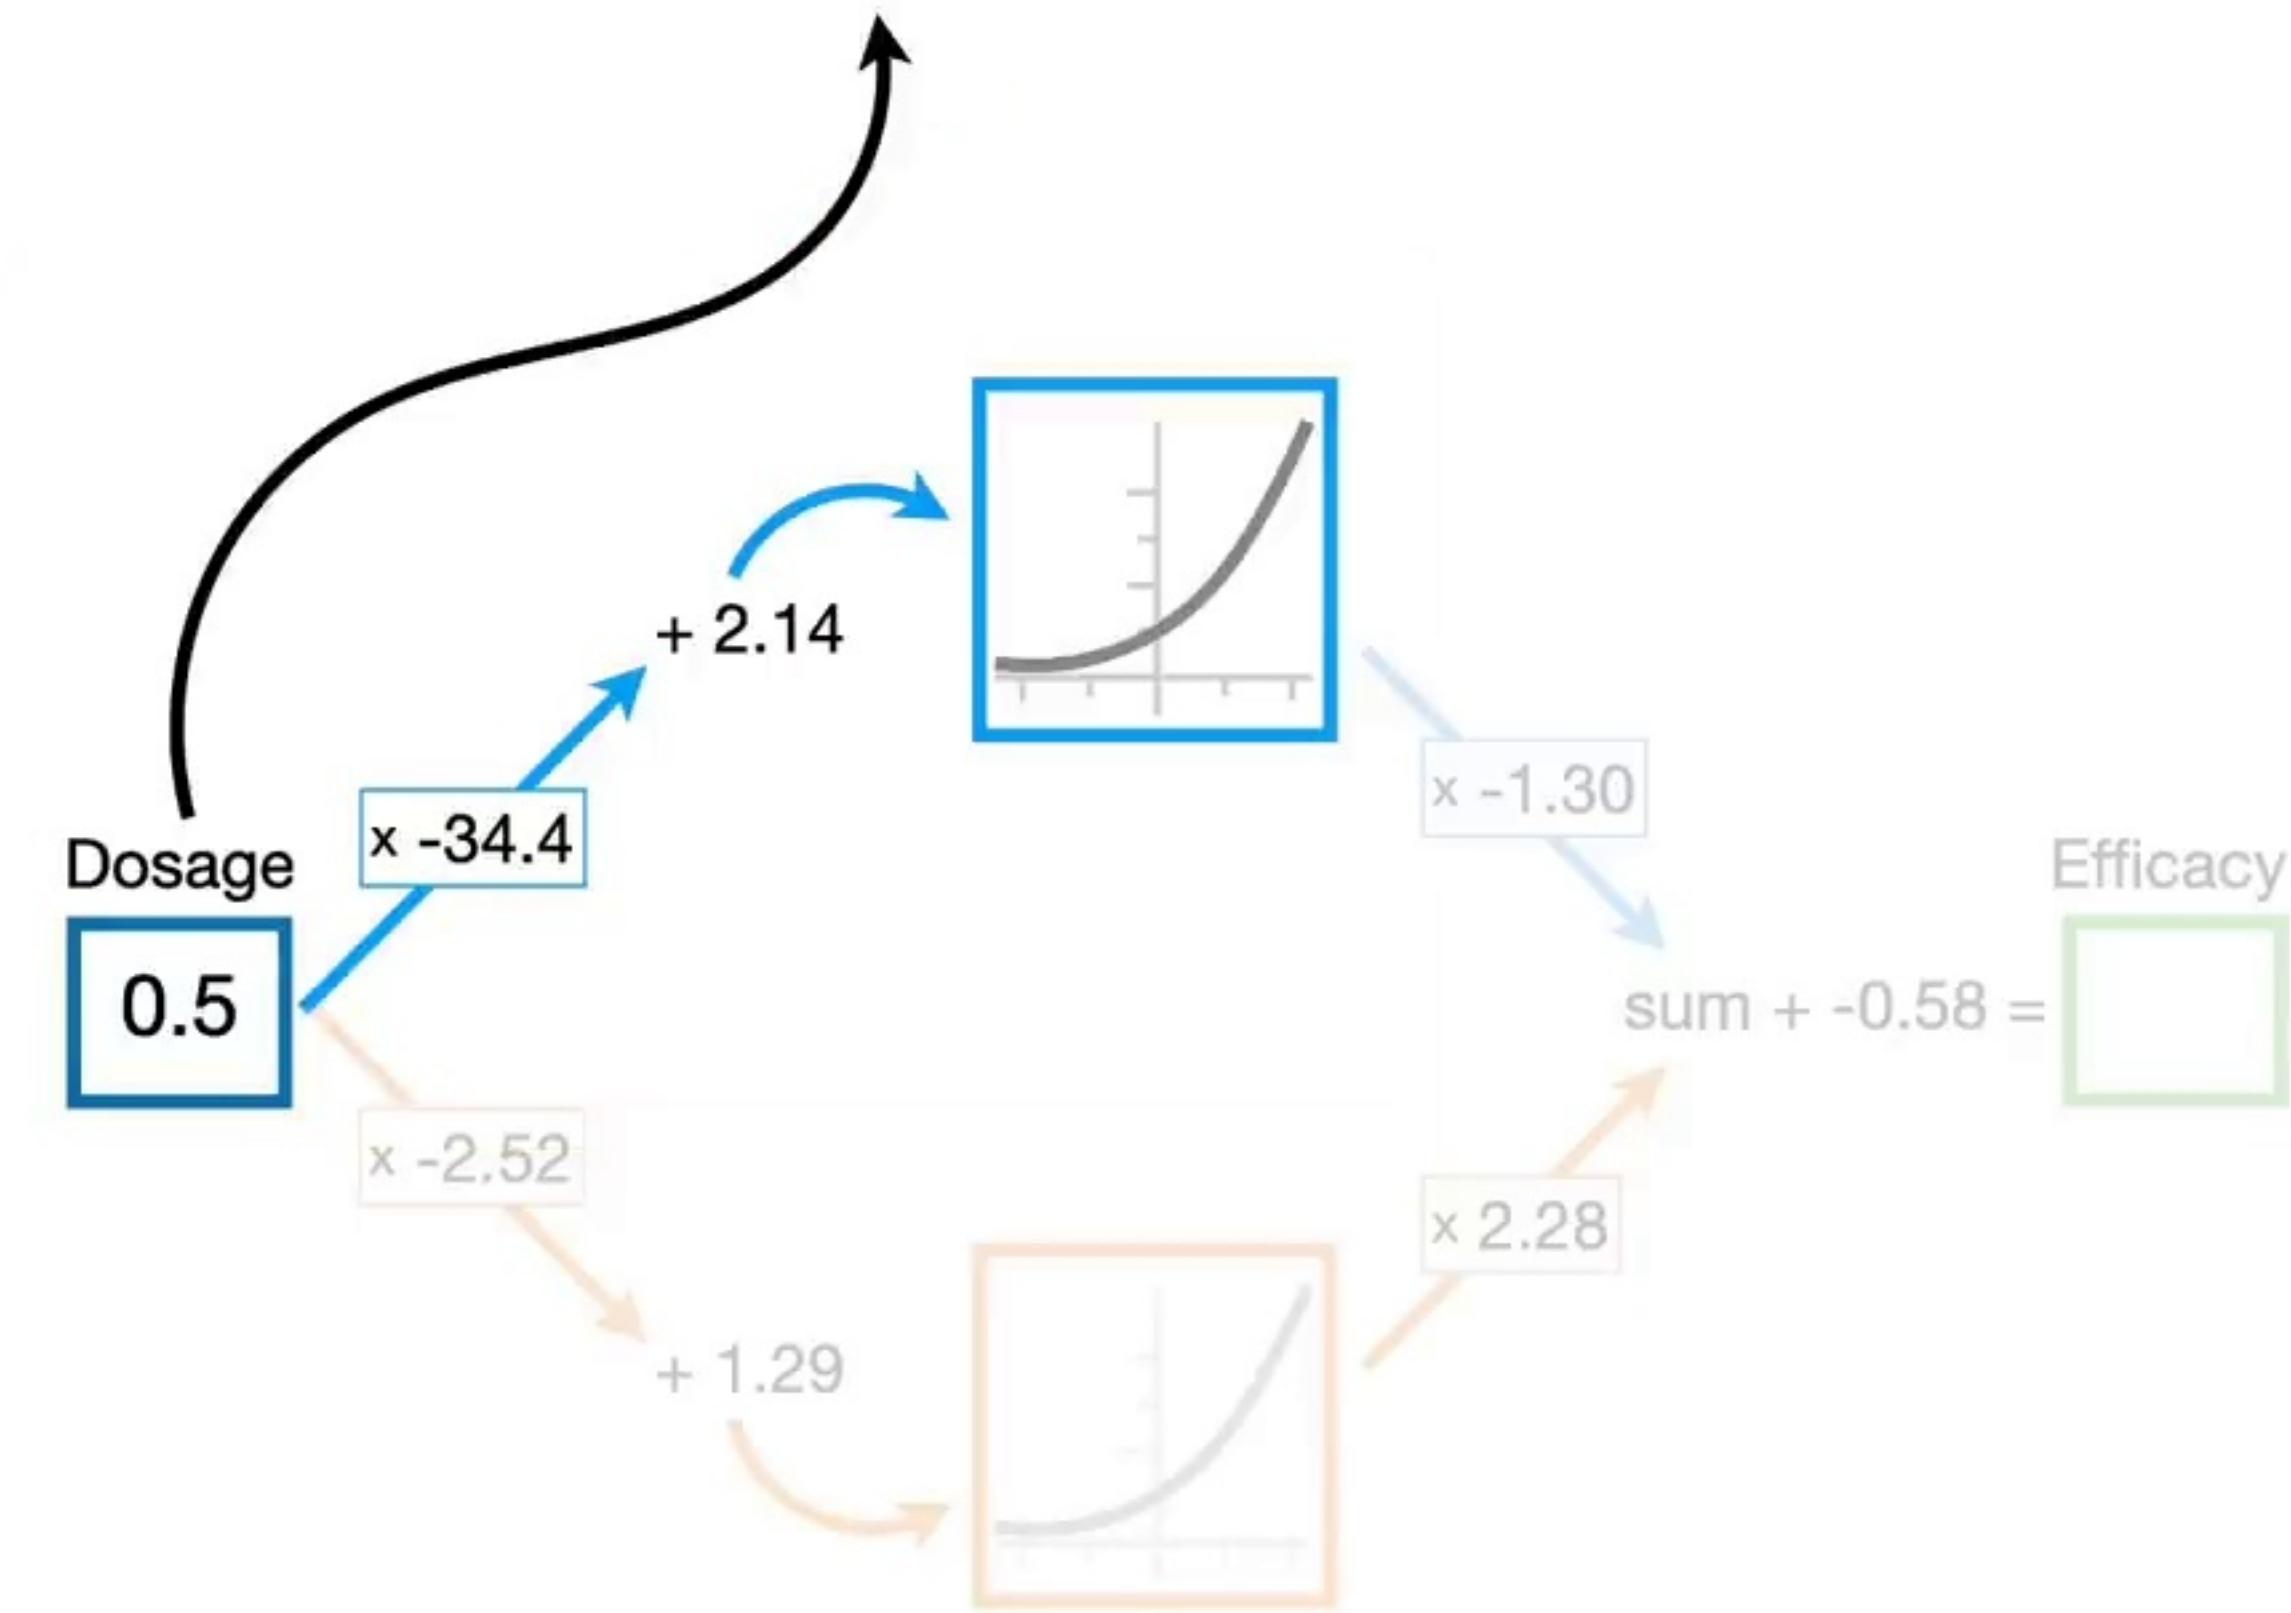
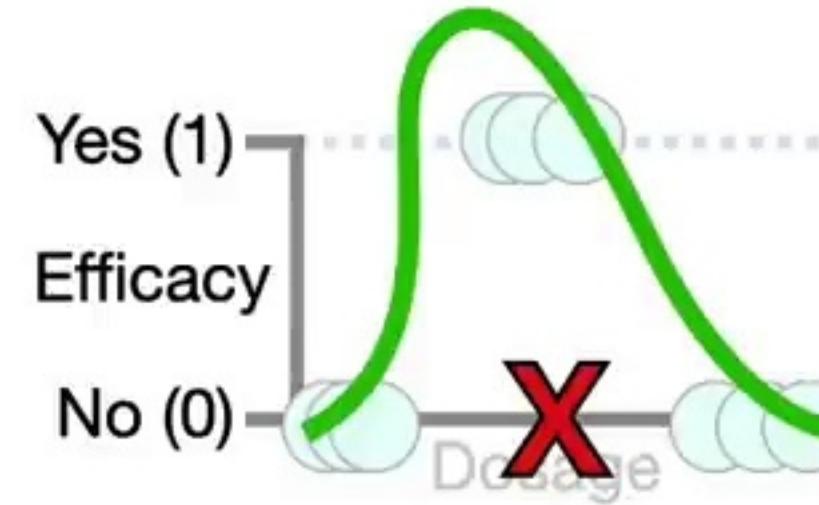


**(Dosage × -34.4) + 2.14 = x-axis coordinate**



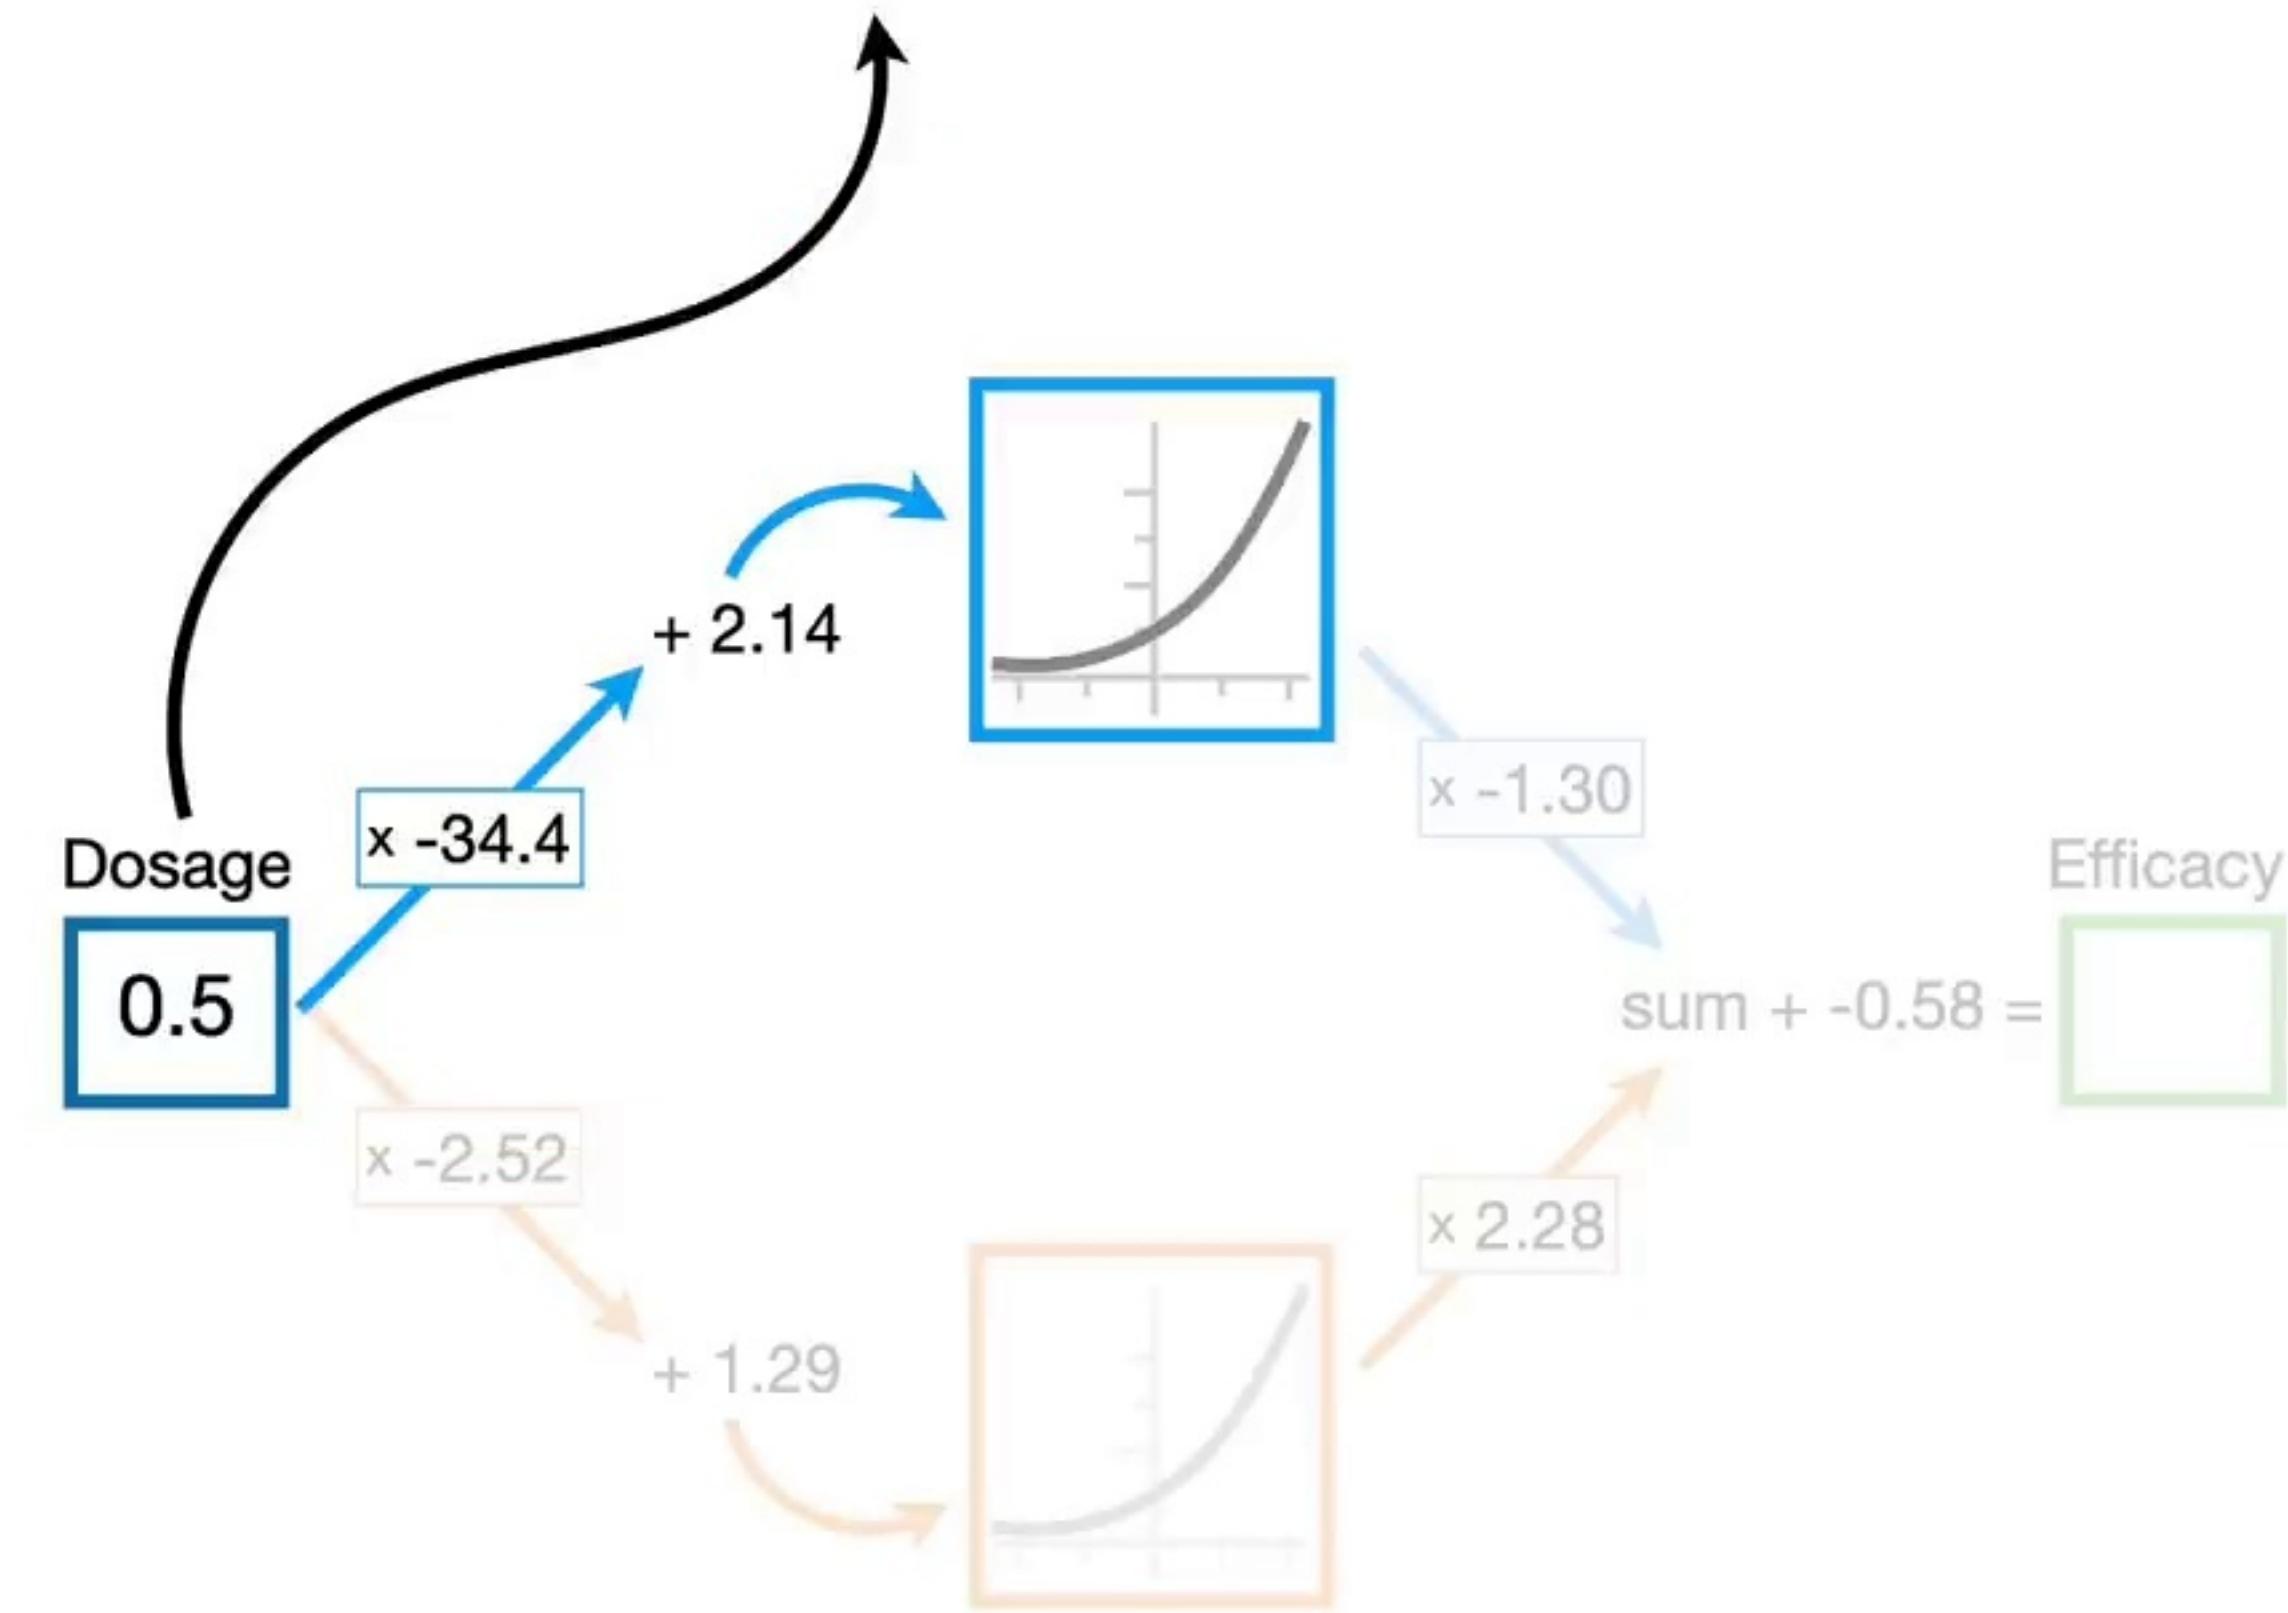
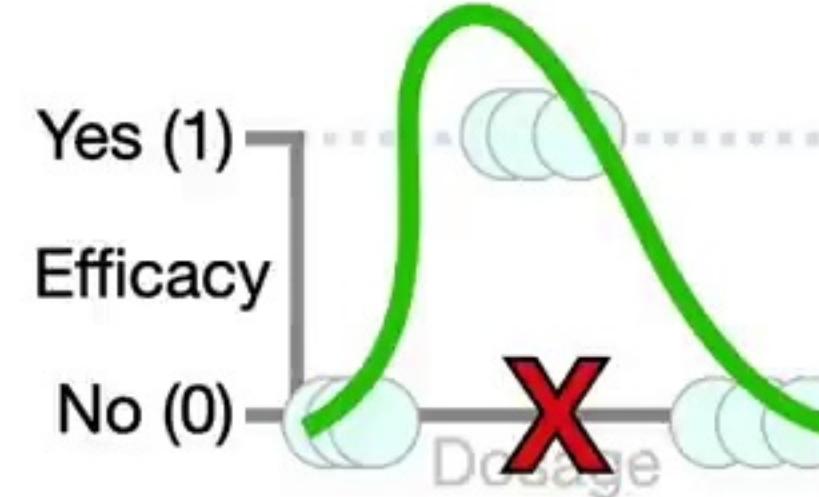


**(Dosage × -34.4) + 2.14 = x-axis coordinate**



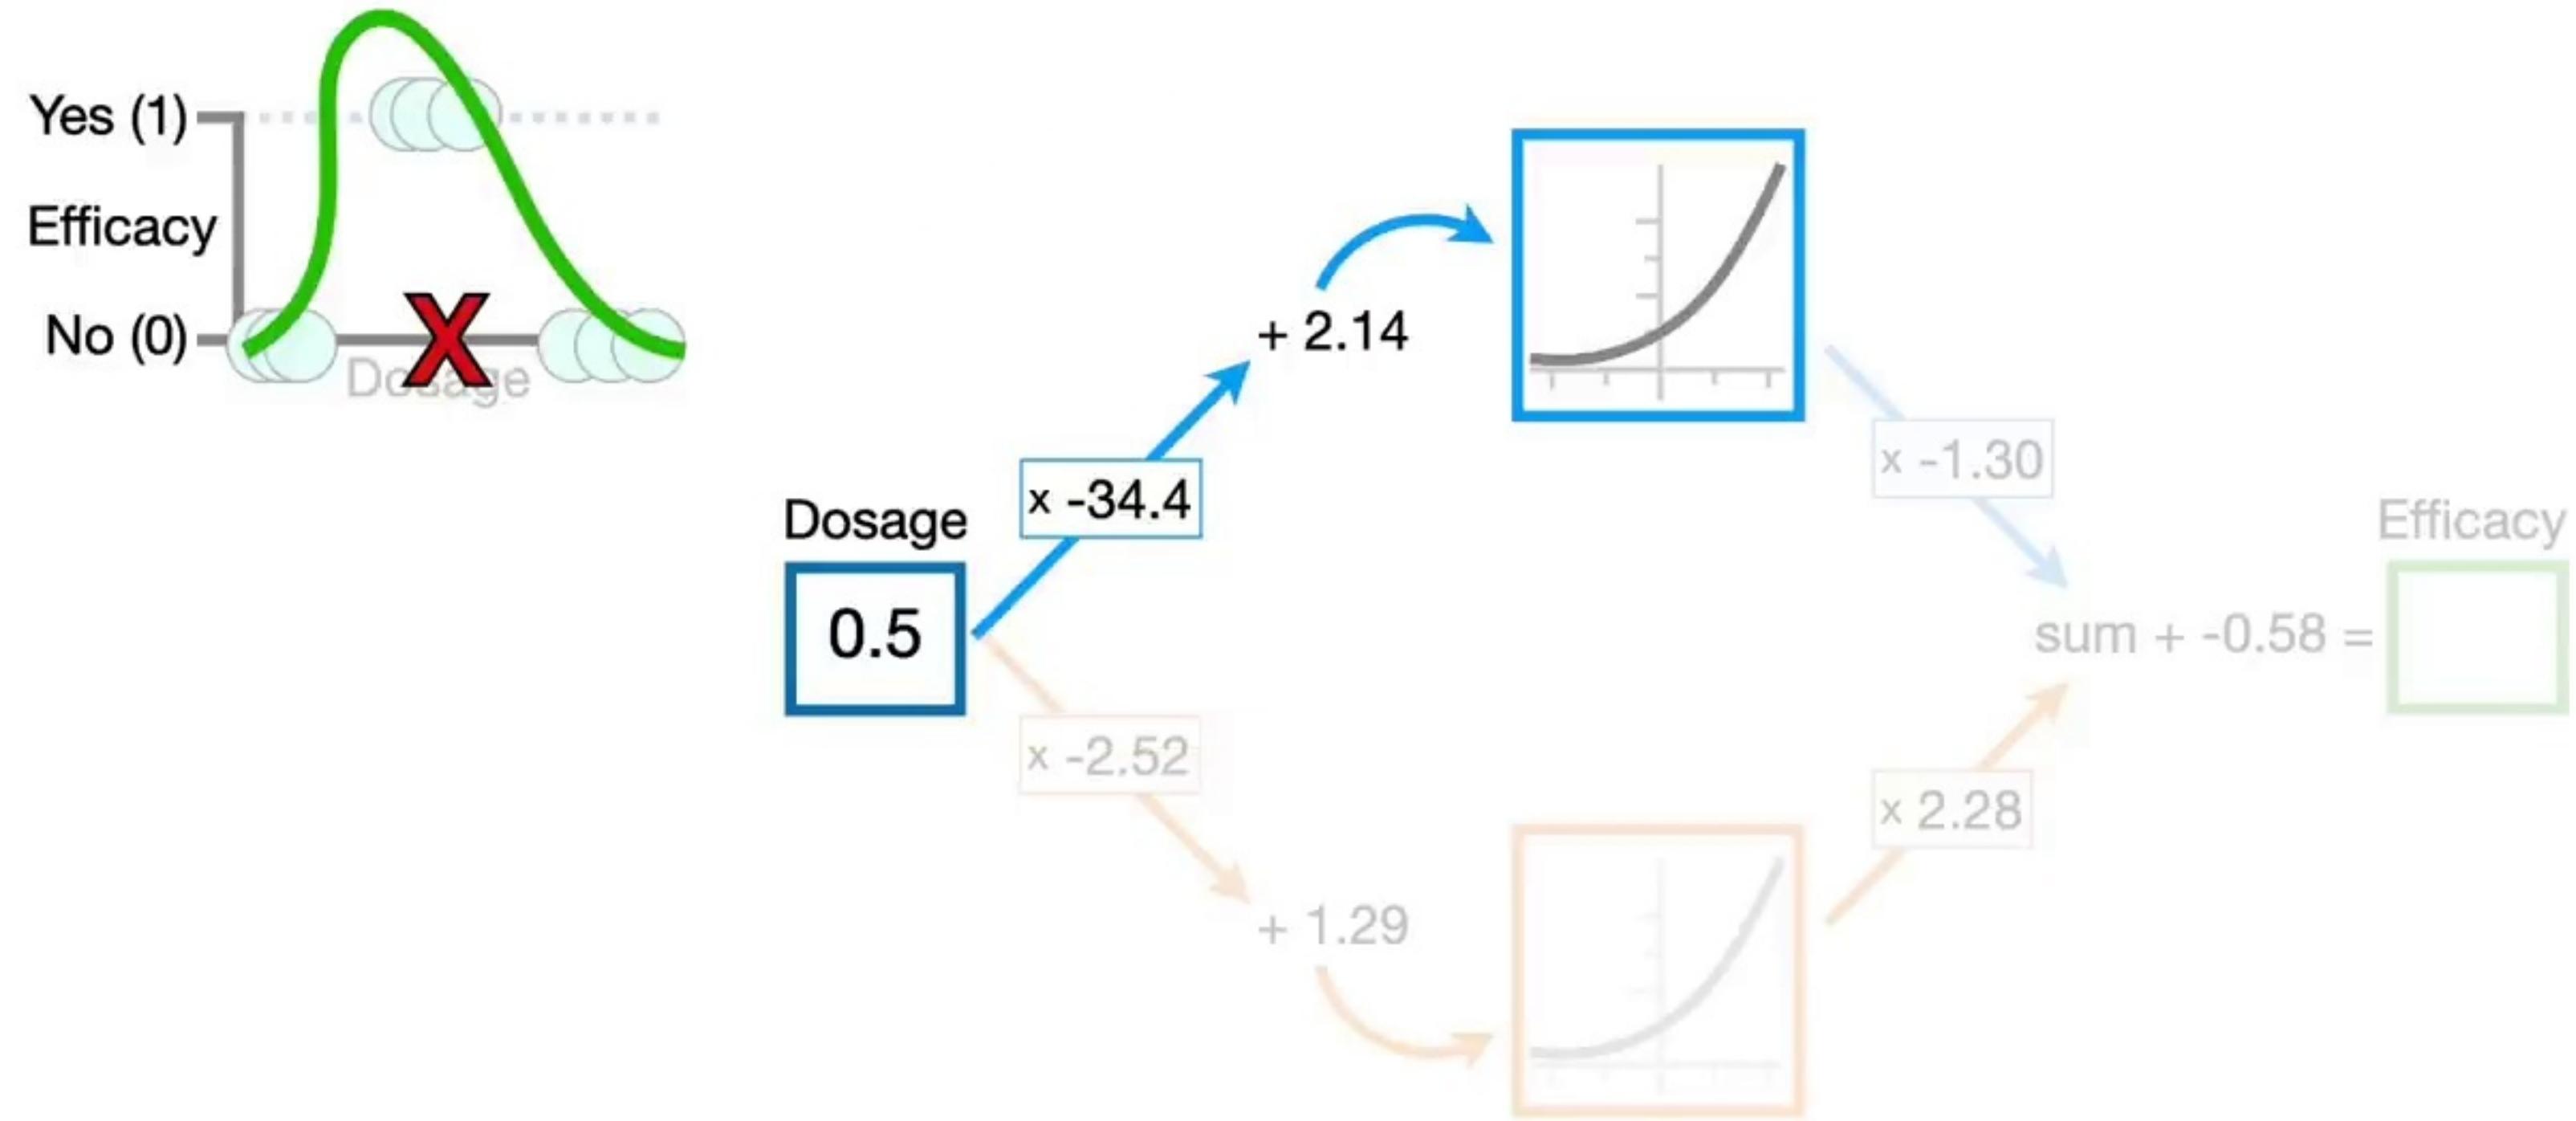


**(0.5 × -34.4) + 2.14 = x-axis coordinate**



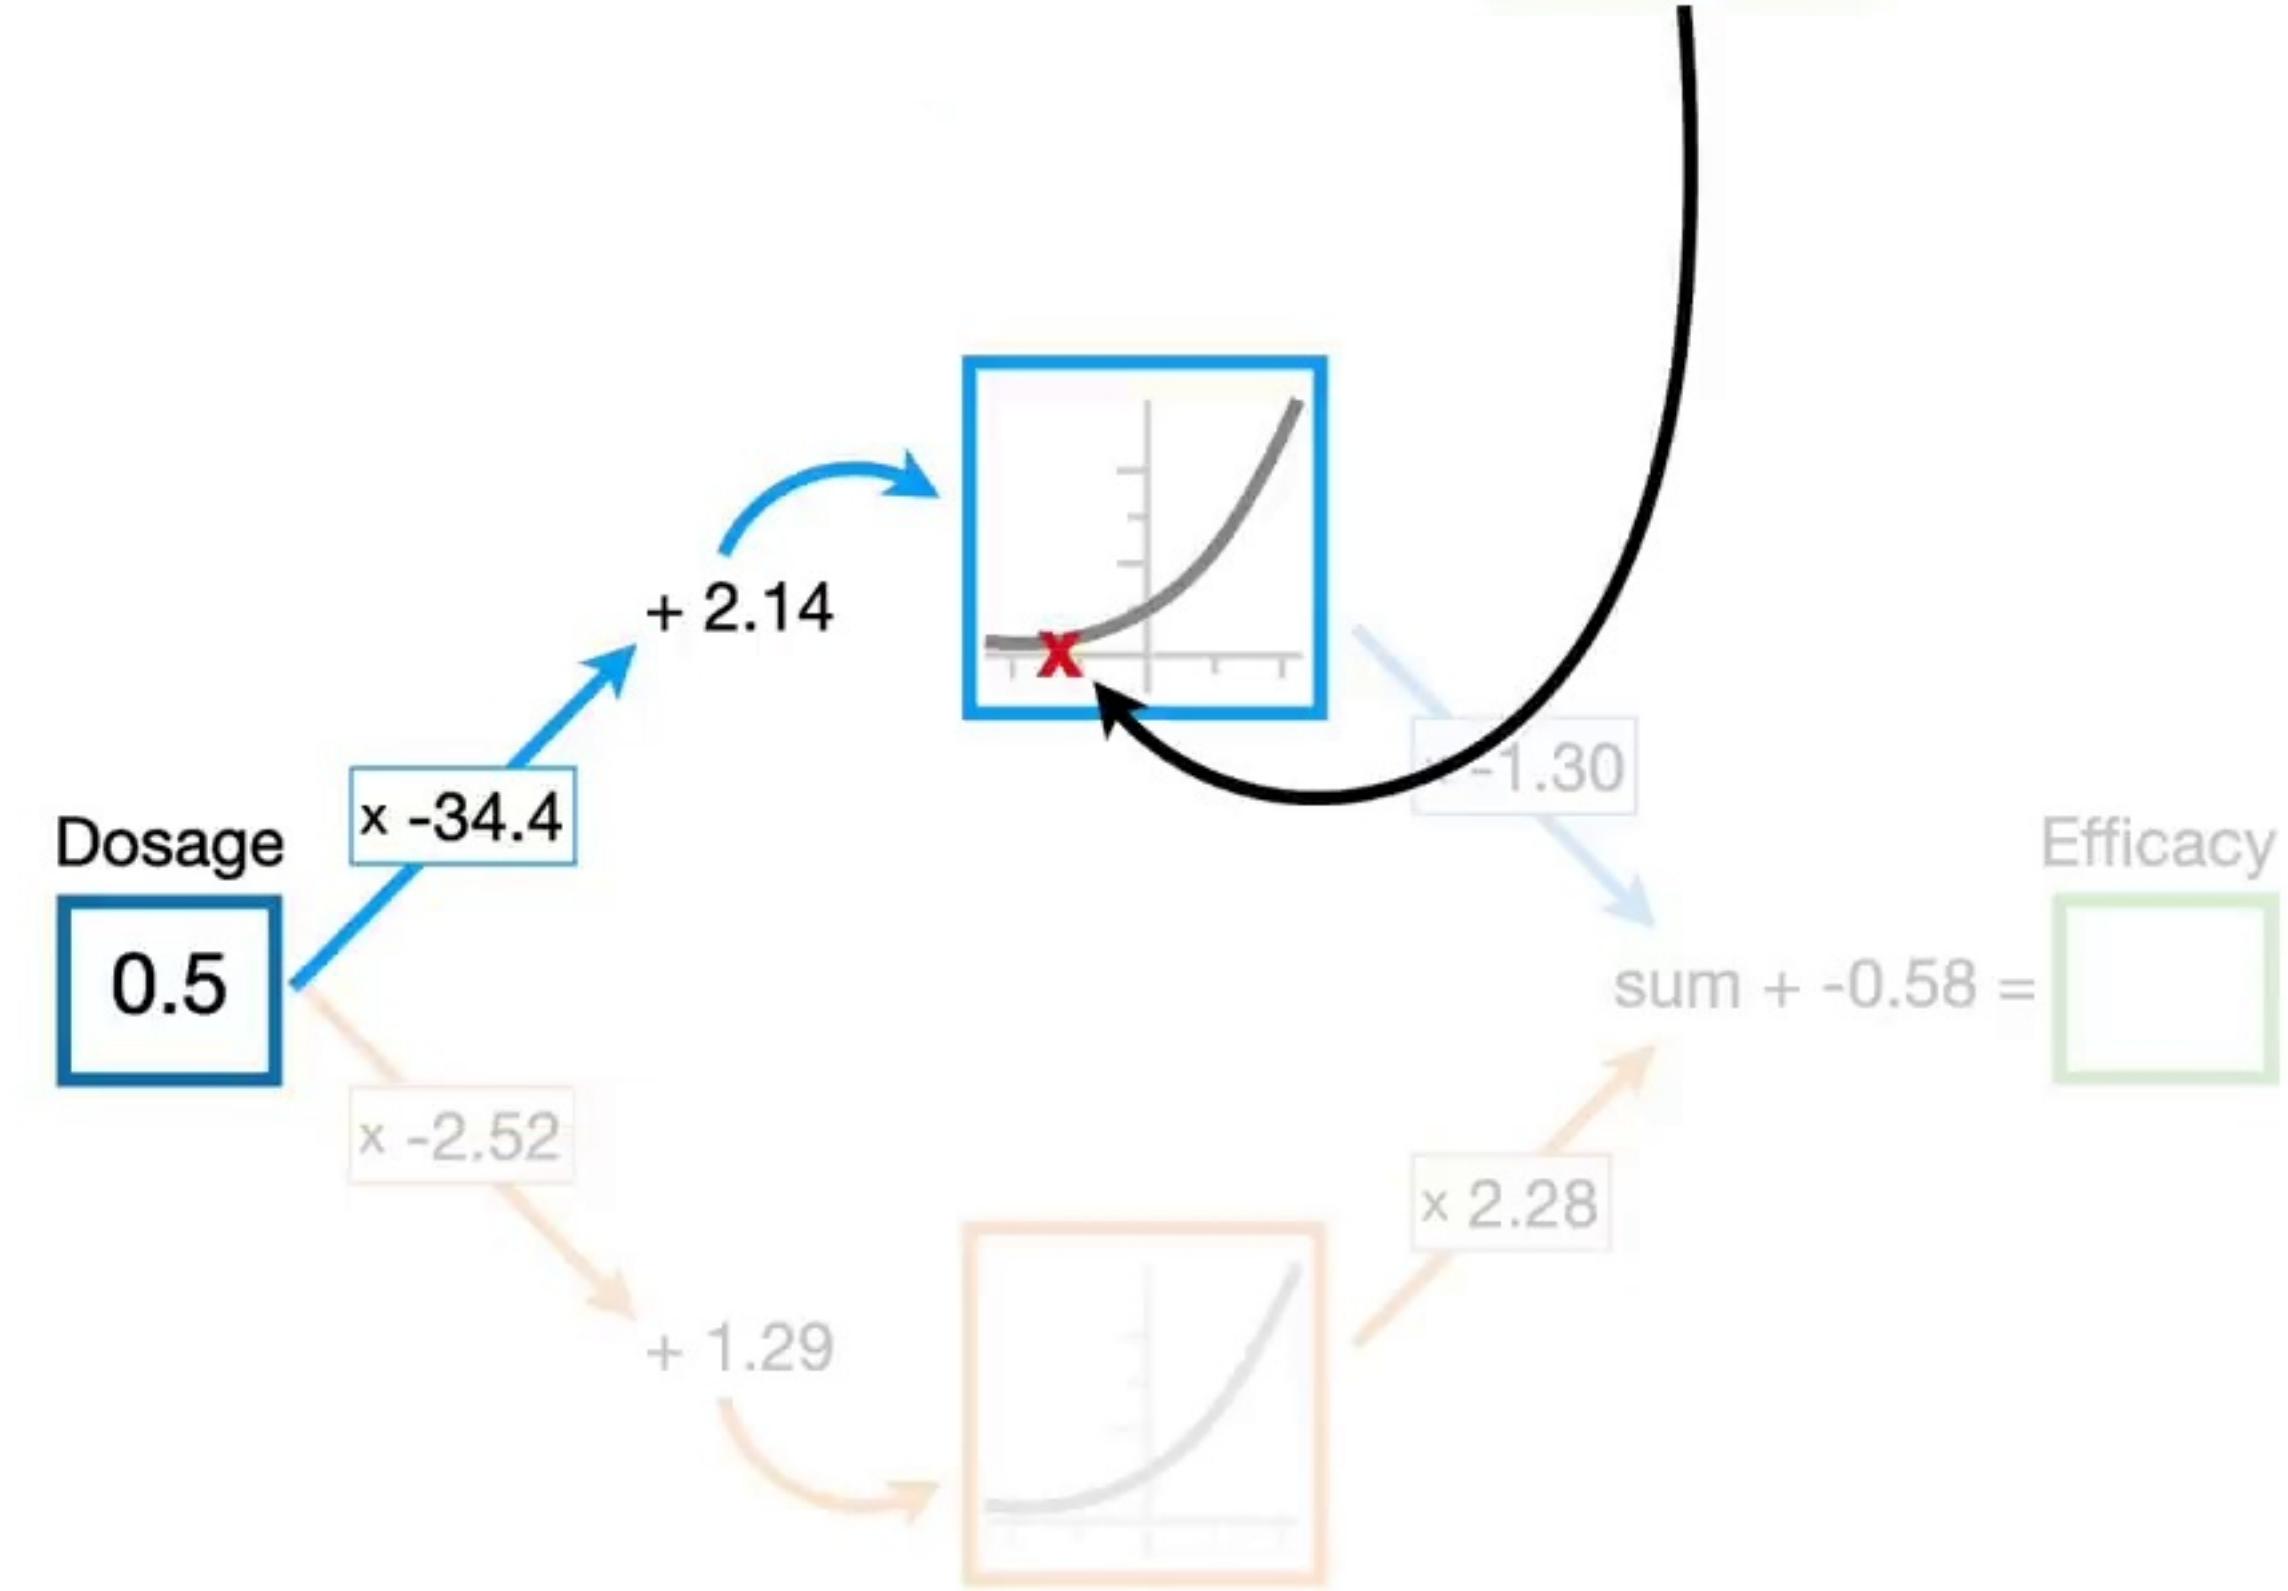
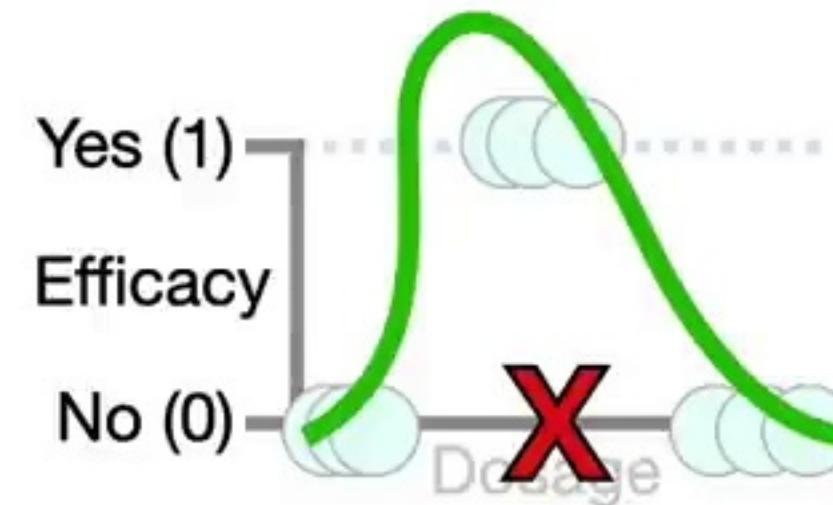


$$(0.5 \times -34.4) + 2.14 = -15.06$$





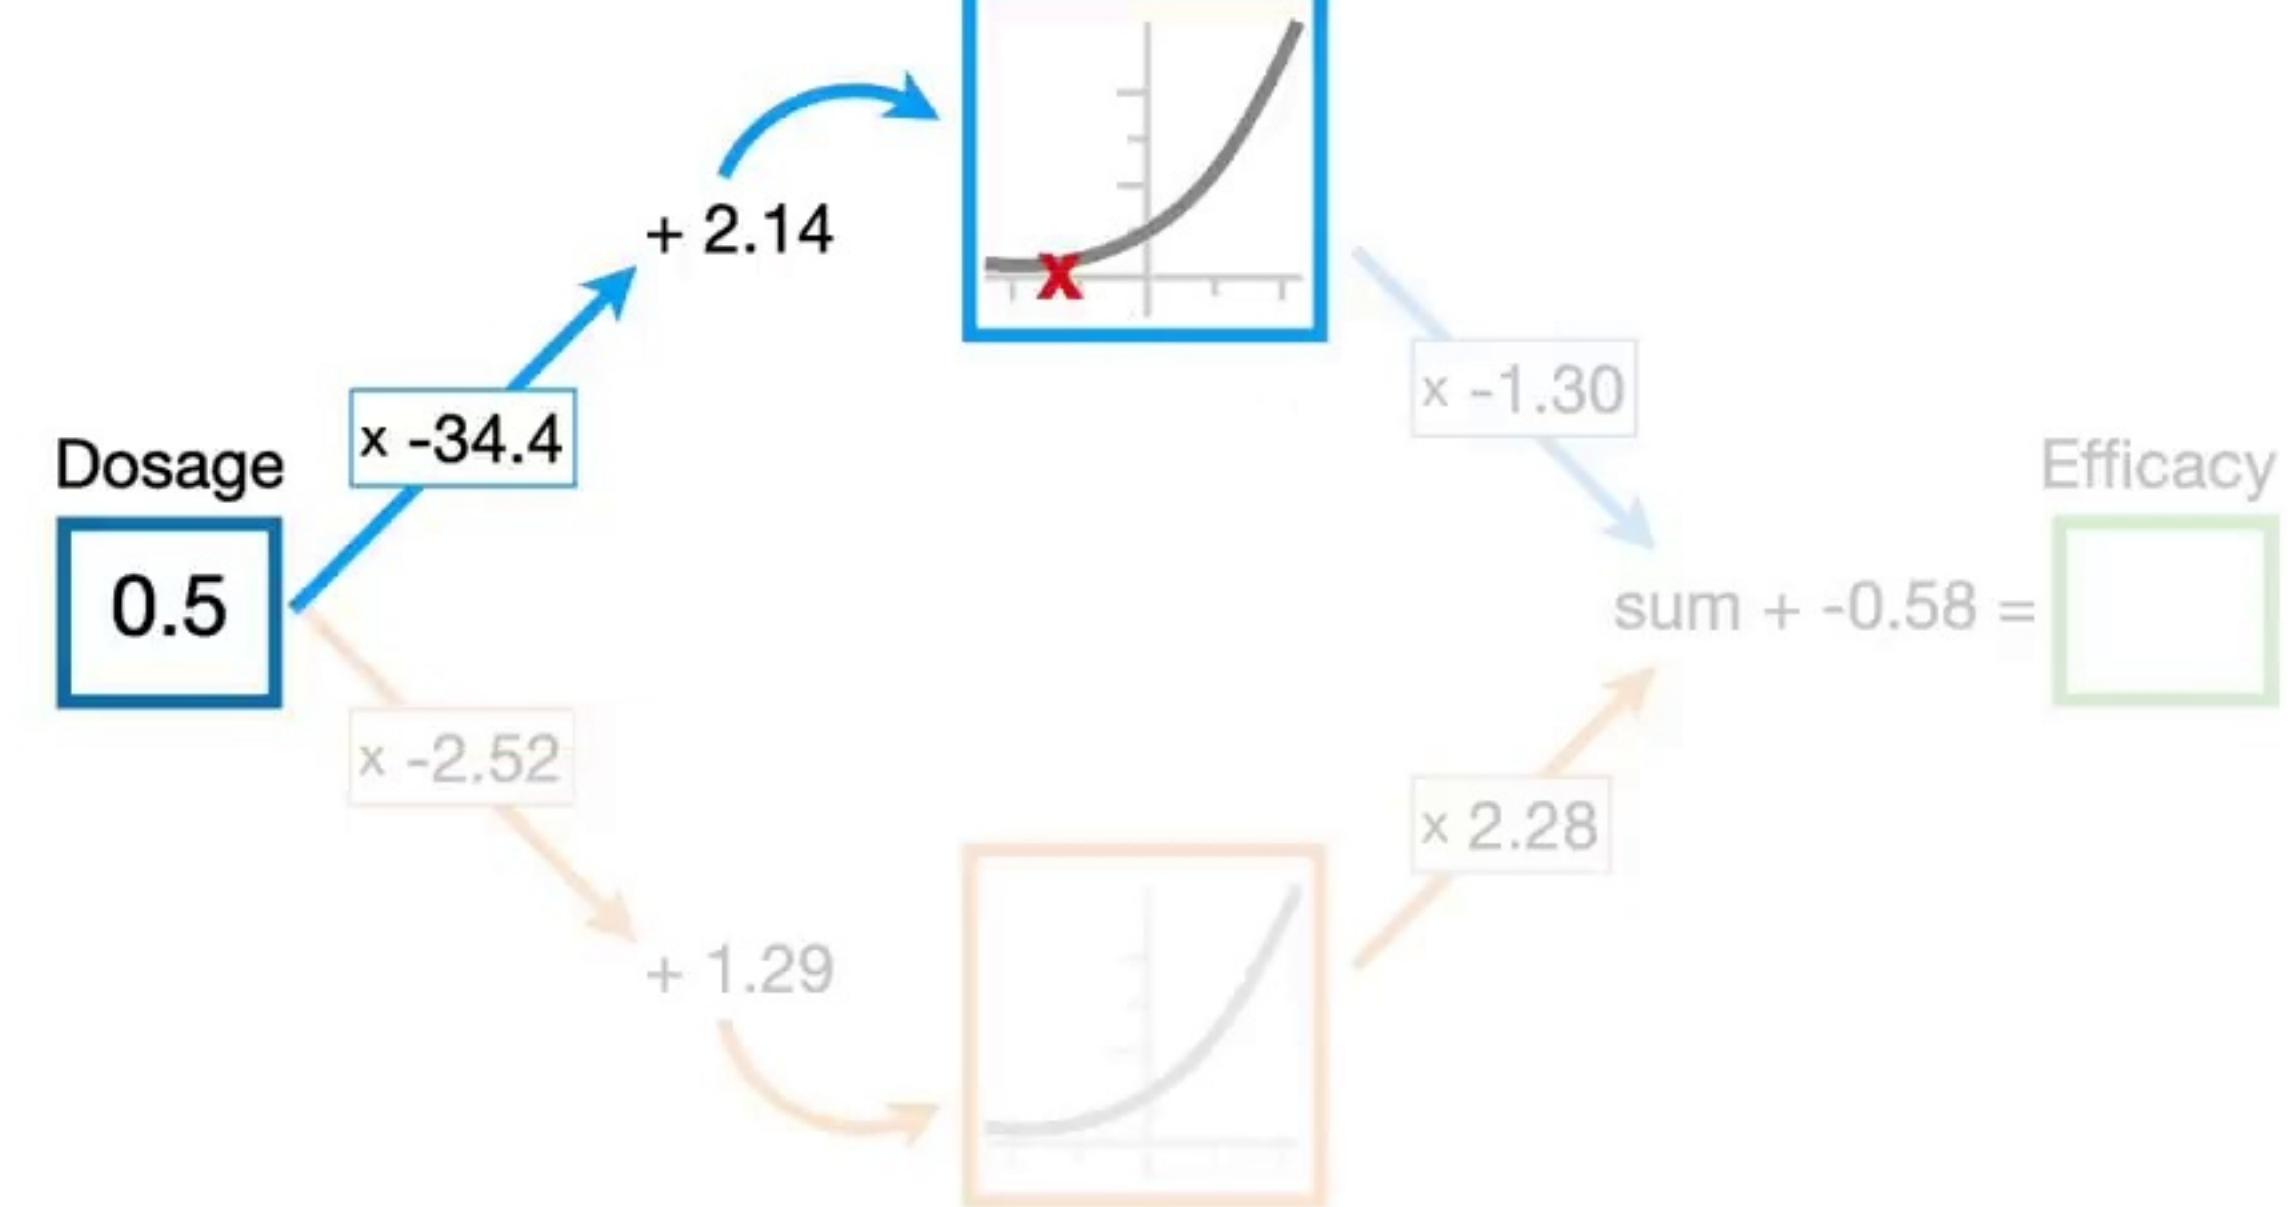
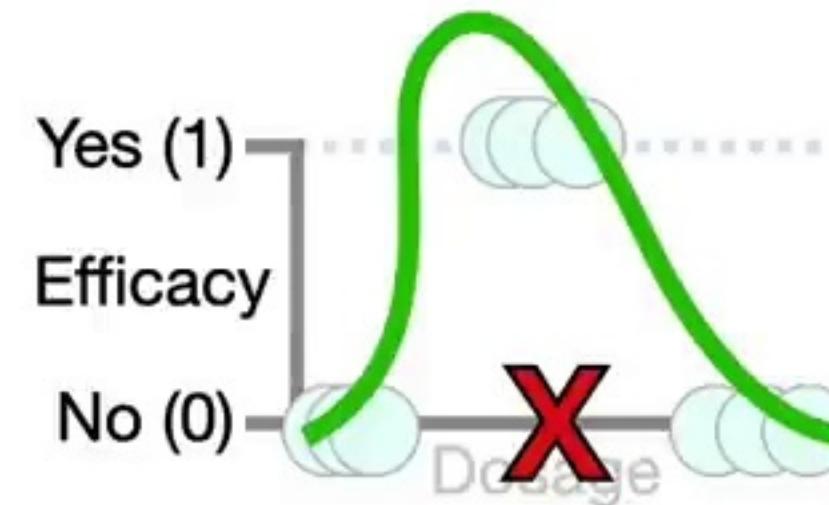
$$(0.5 \times -34.4) + 2.14 = -15.06$$





$$(0.5 \times -34.4) + 2.14 = -15.06$$

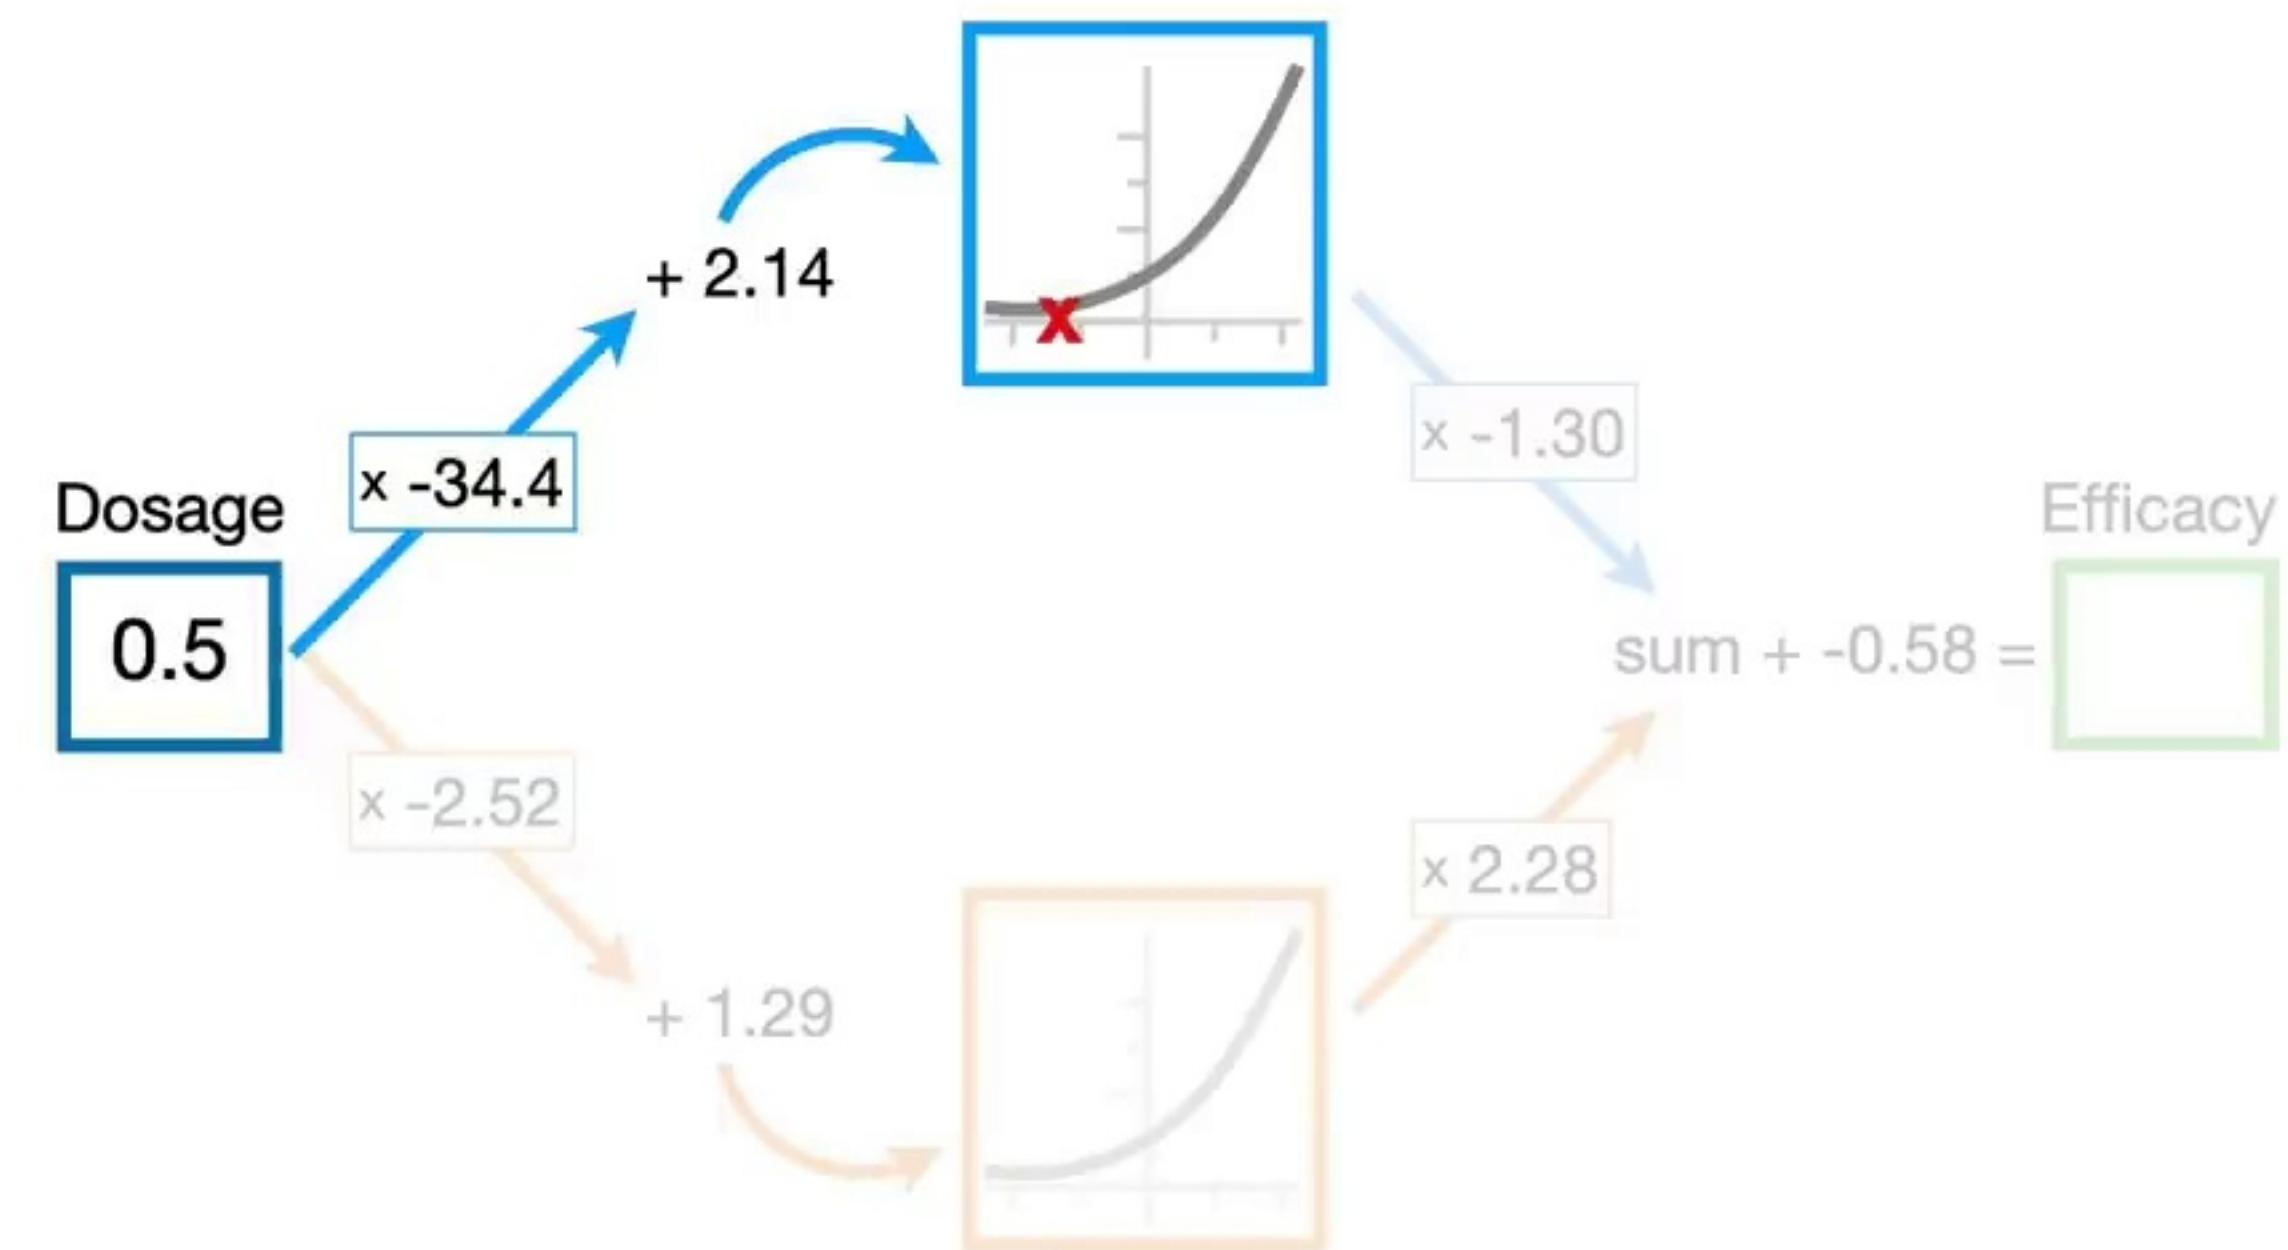
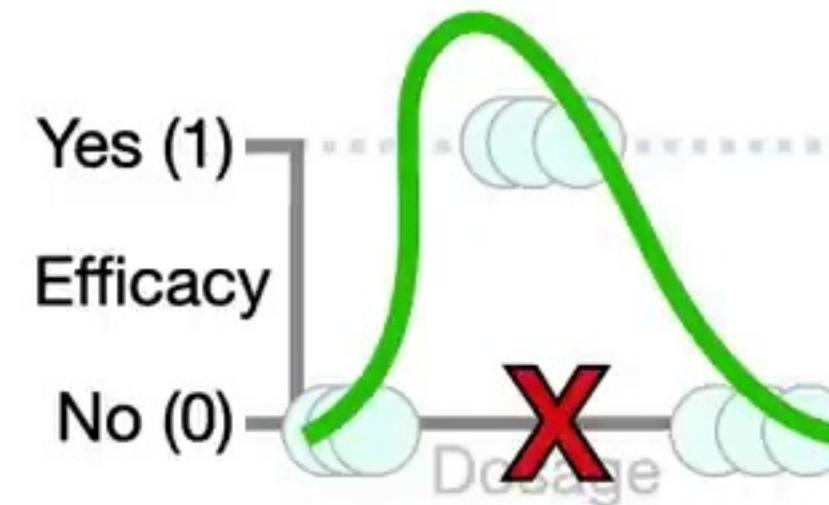
$$f(x) = \log(1 + e^x)$$



Double  
BAM!!  
**SQ!**

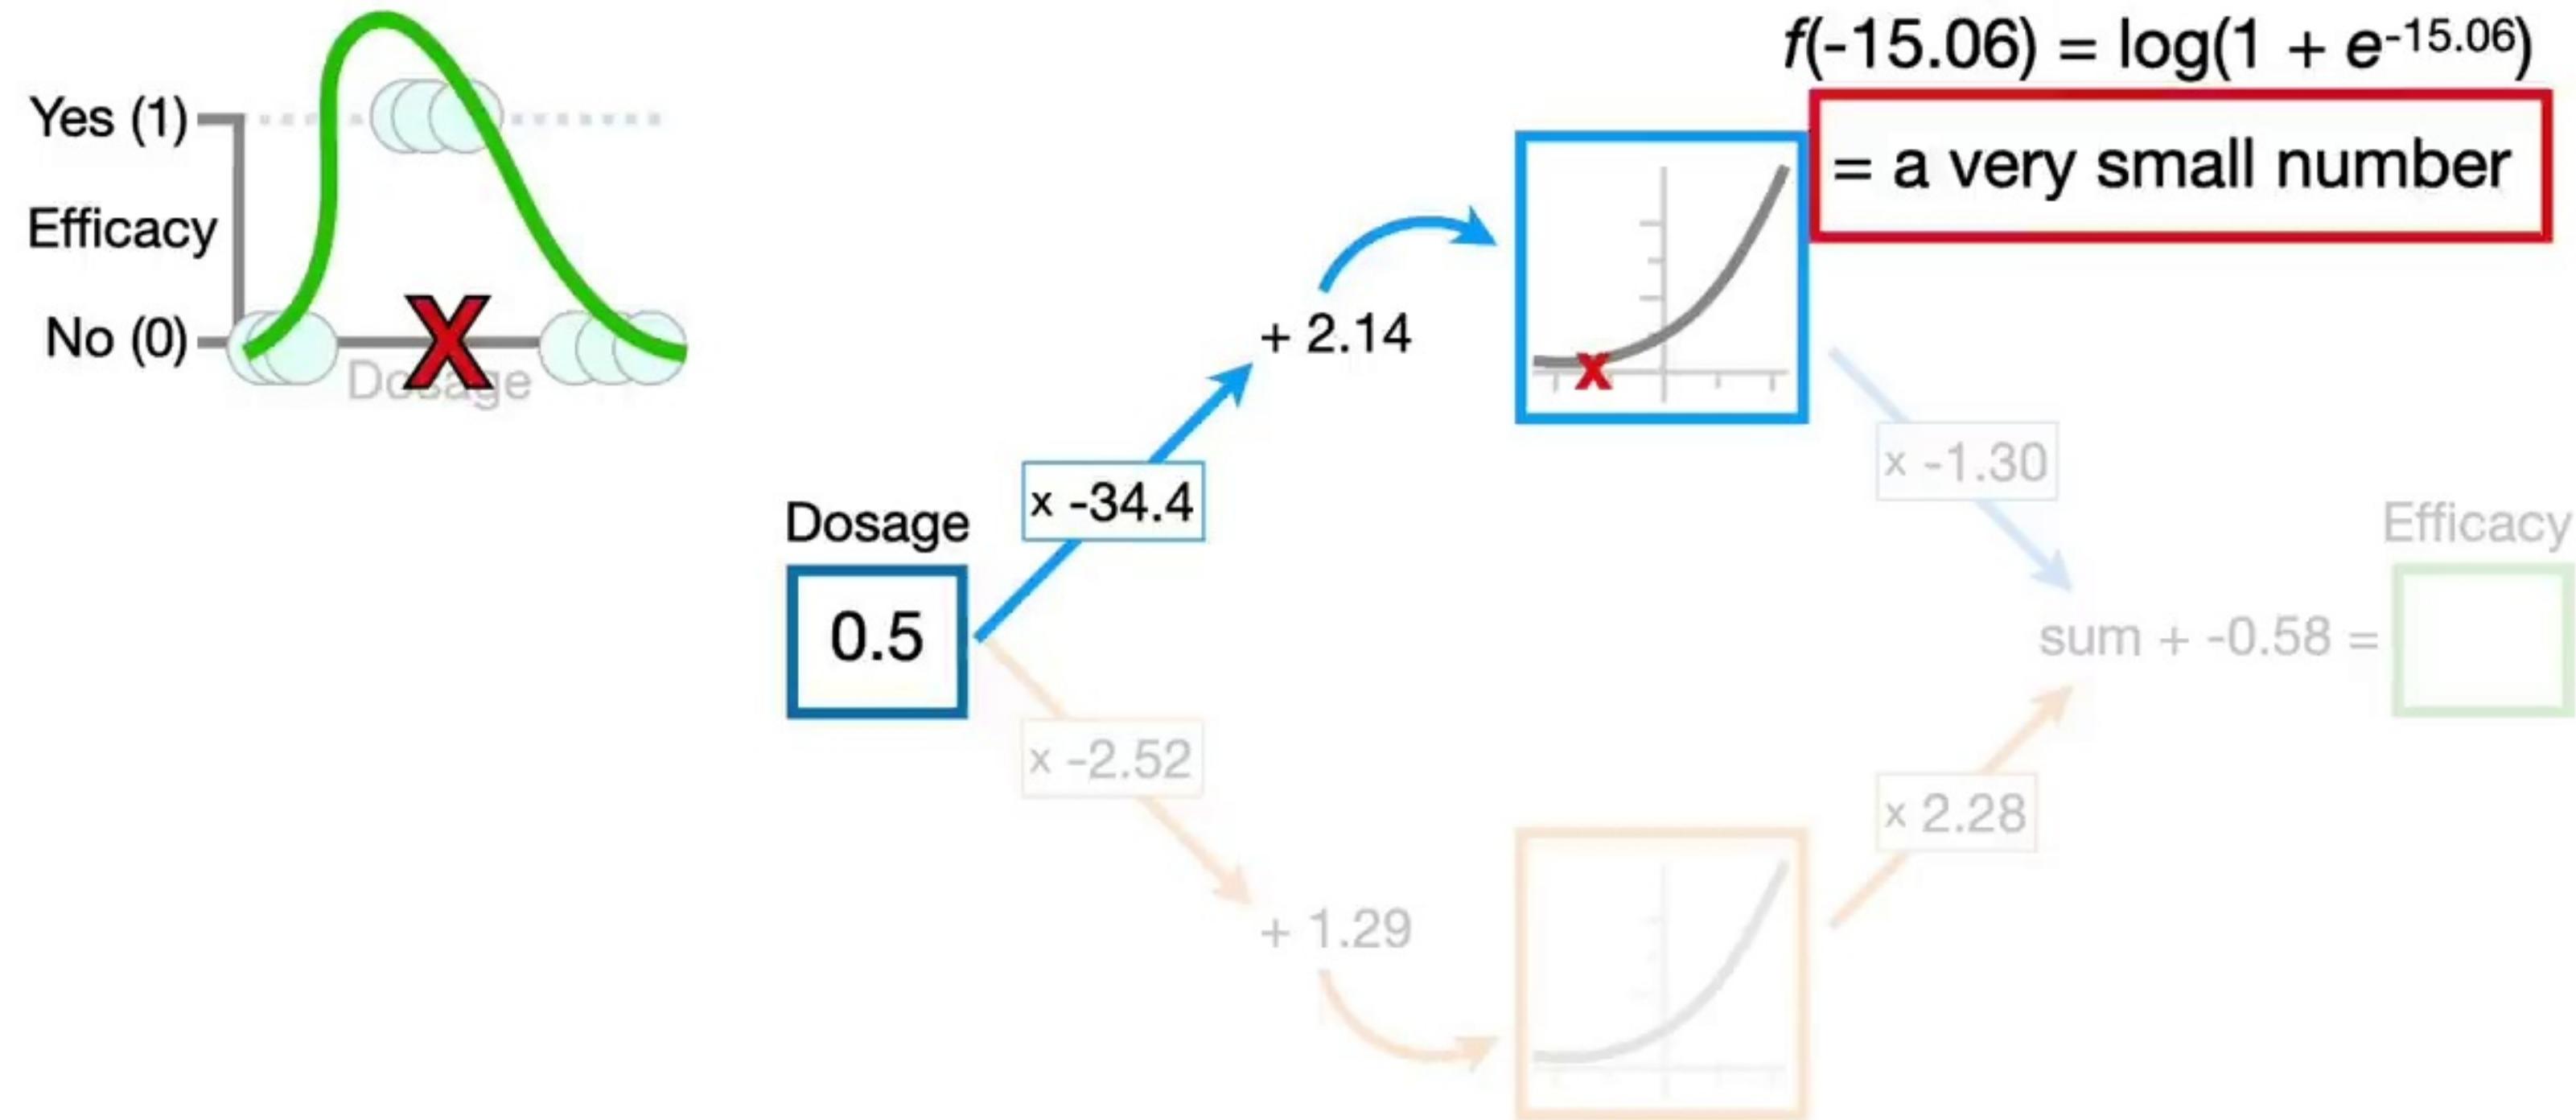
$$(0.5 \times -34.4) + 2.14 = -15.06$$

$$f(-15.06) = \log(1 + e^{-15.06})$$



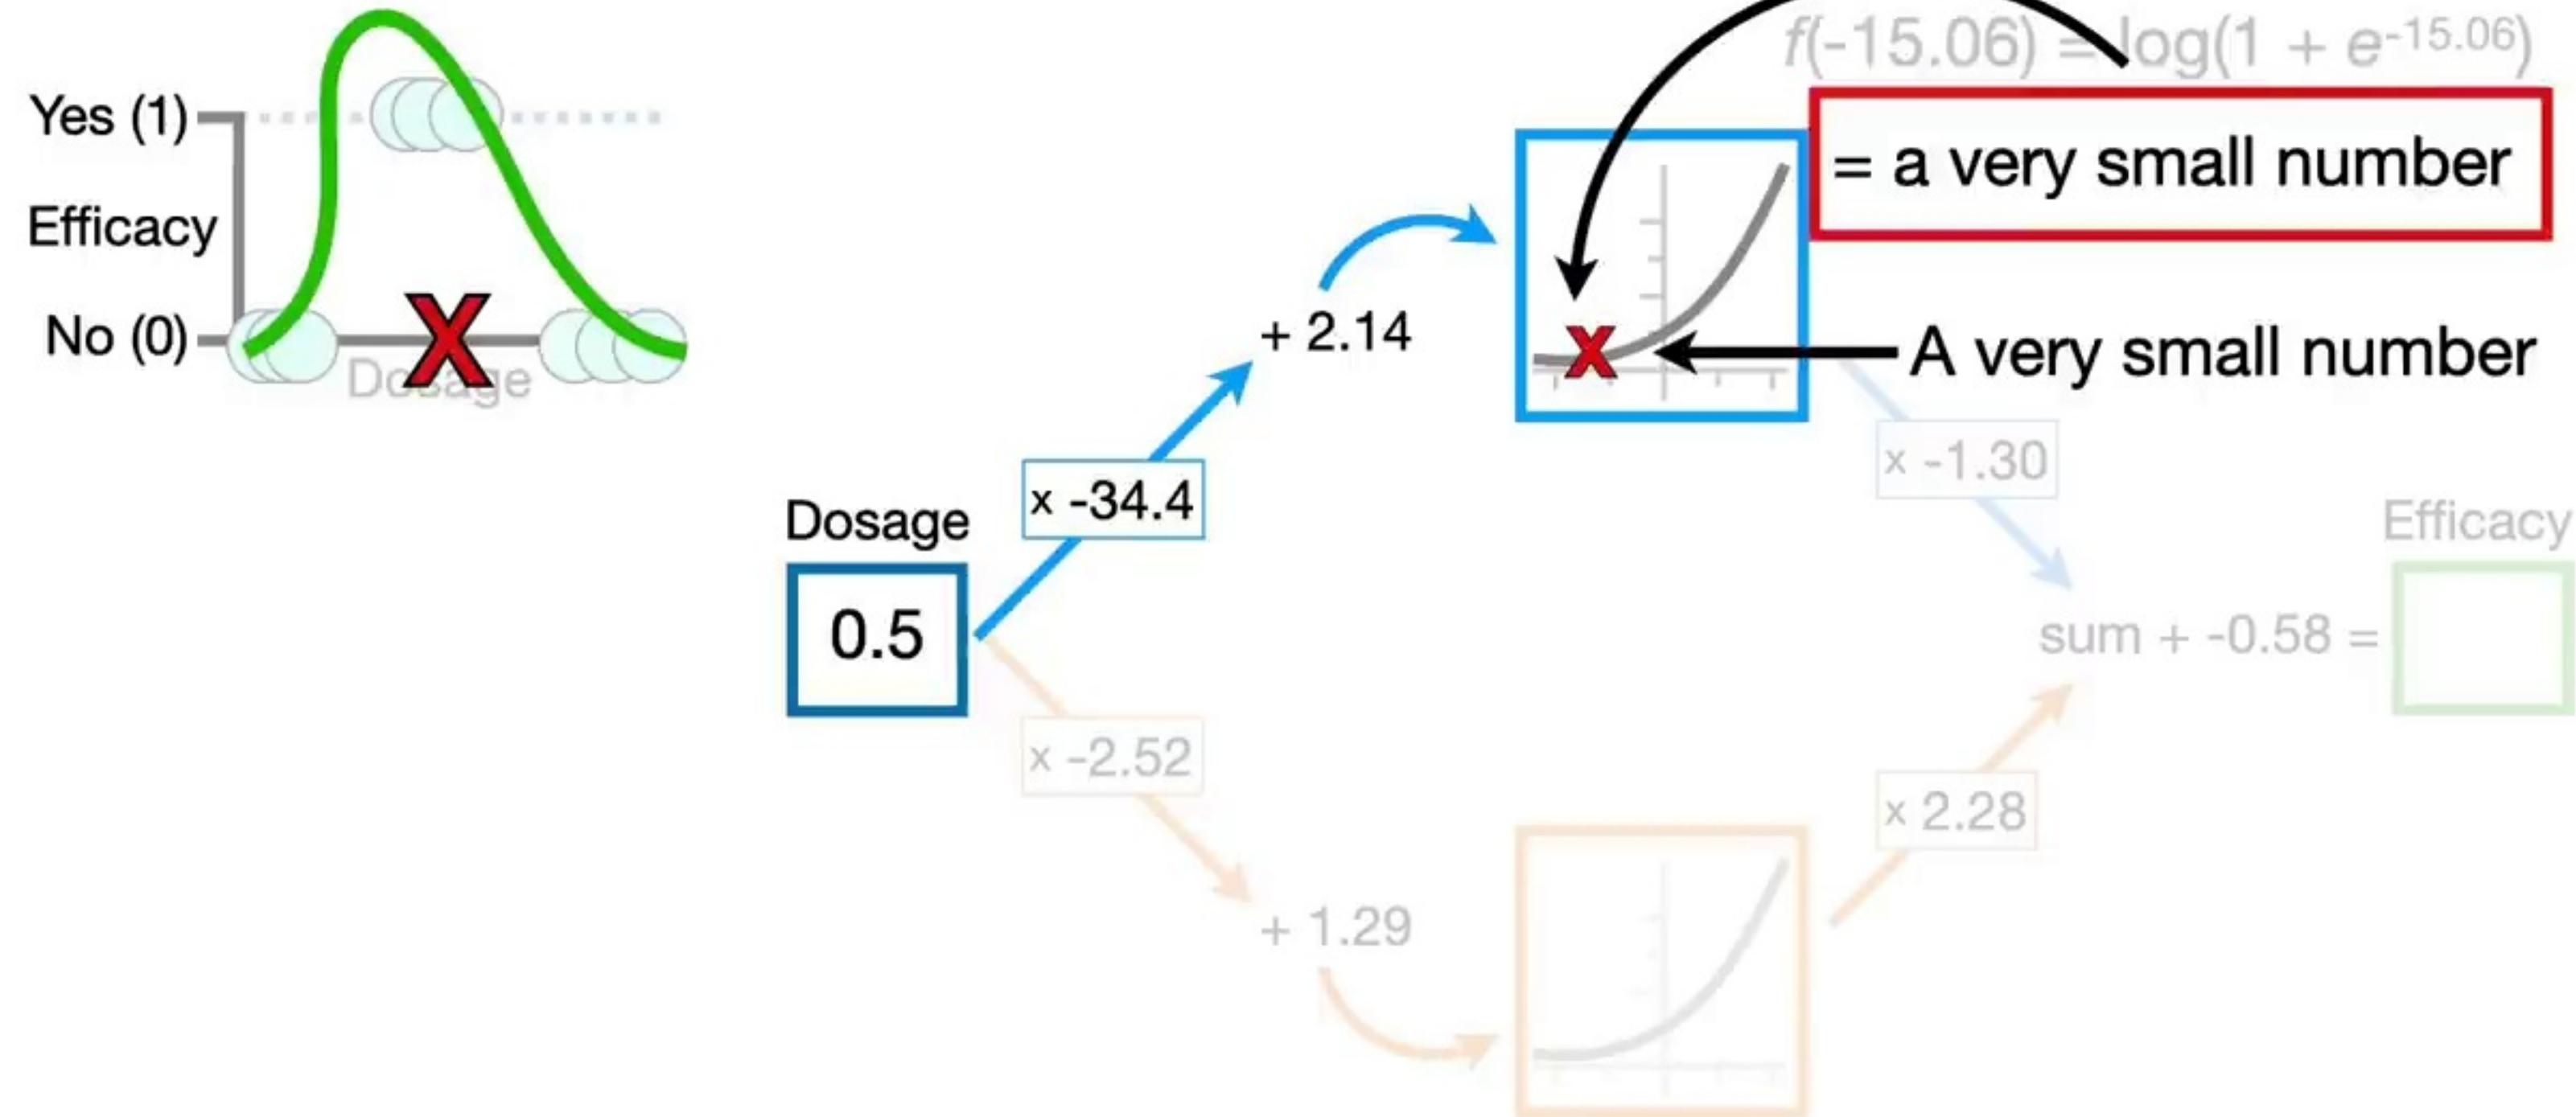
Double  
BAM!!  
**SQ!**

$$(0.5 \times -34.4) + 2.14 = -15.06$$



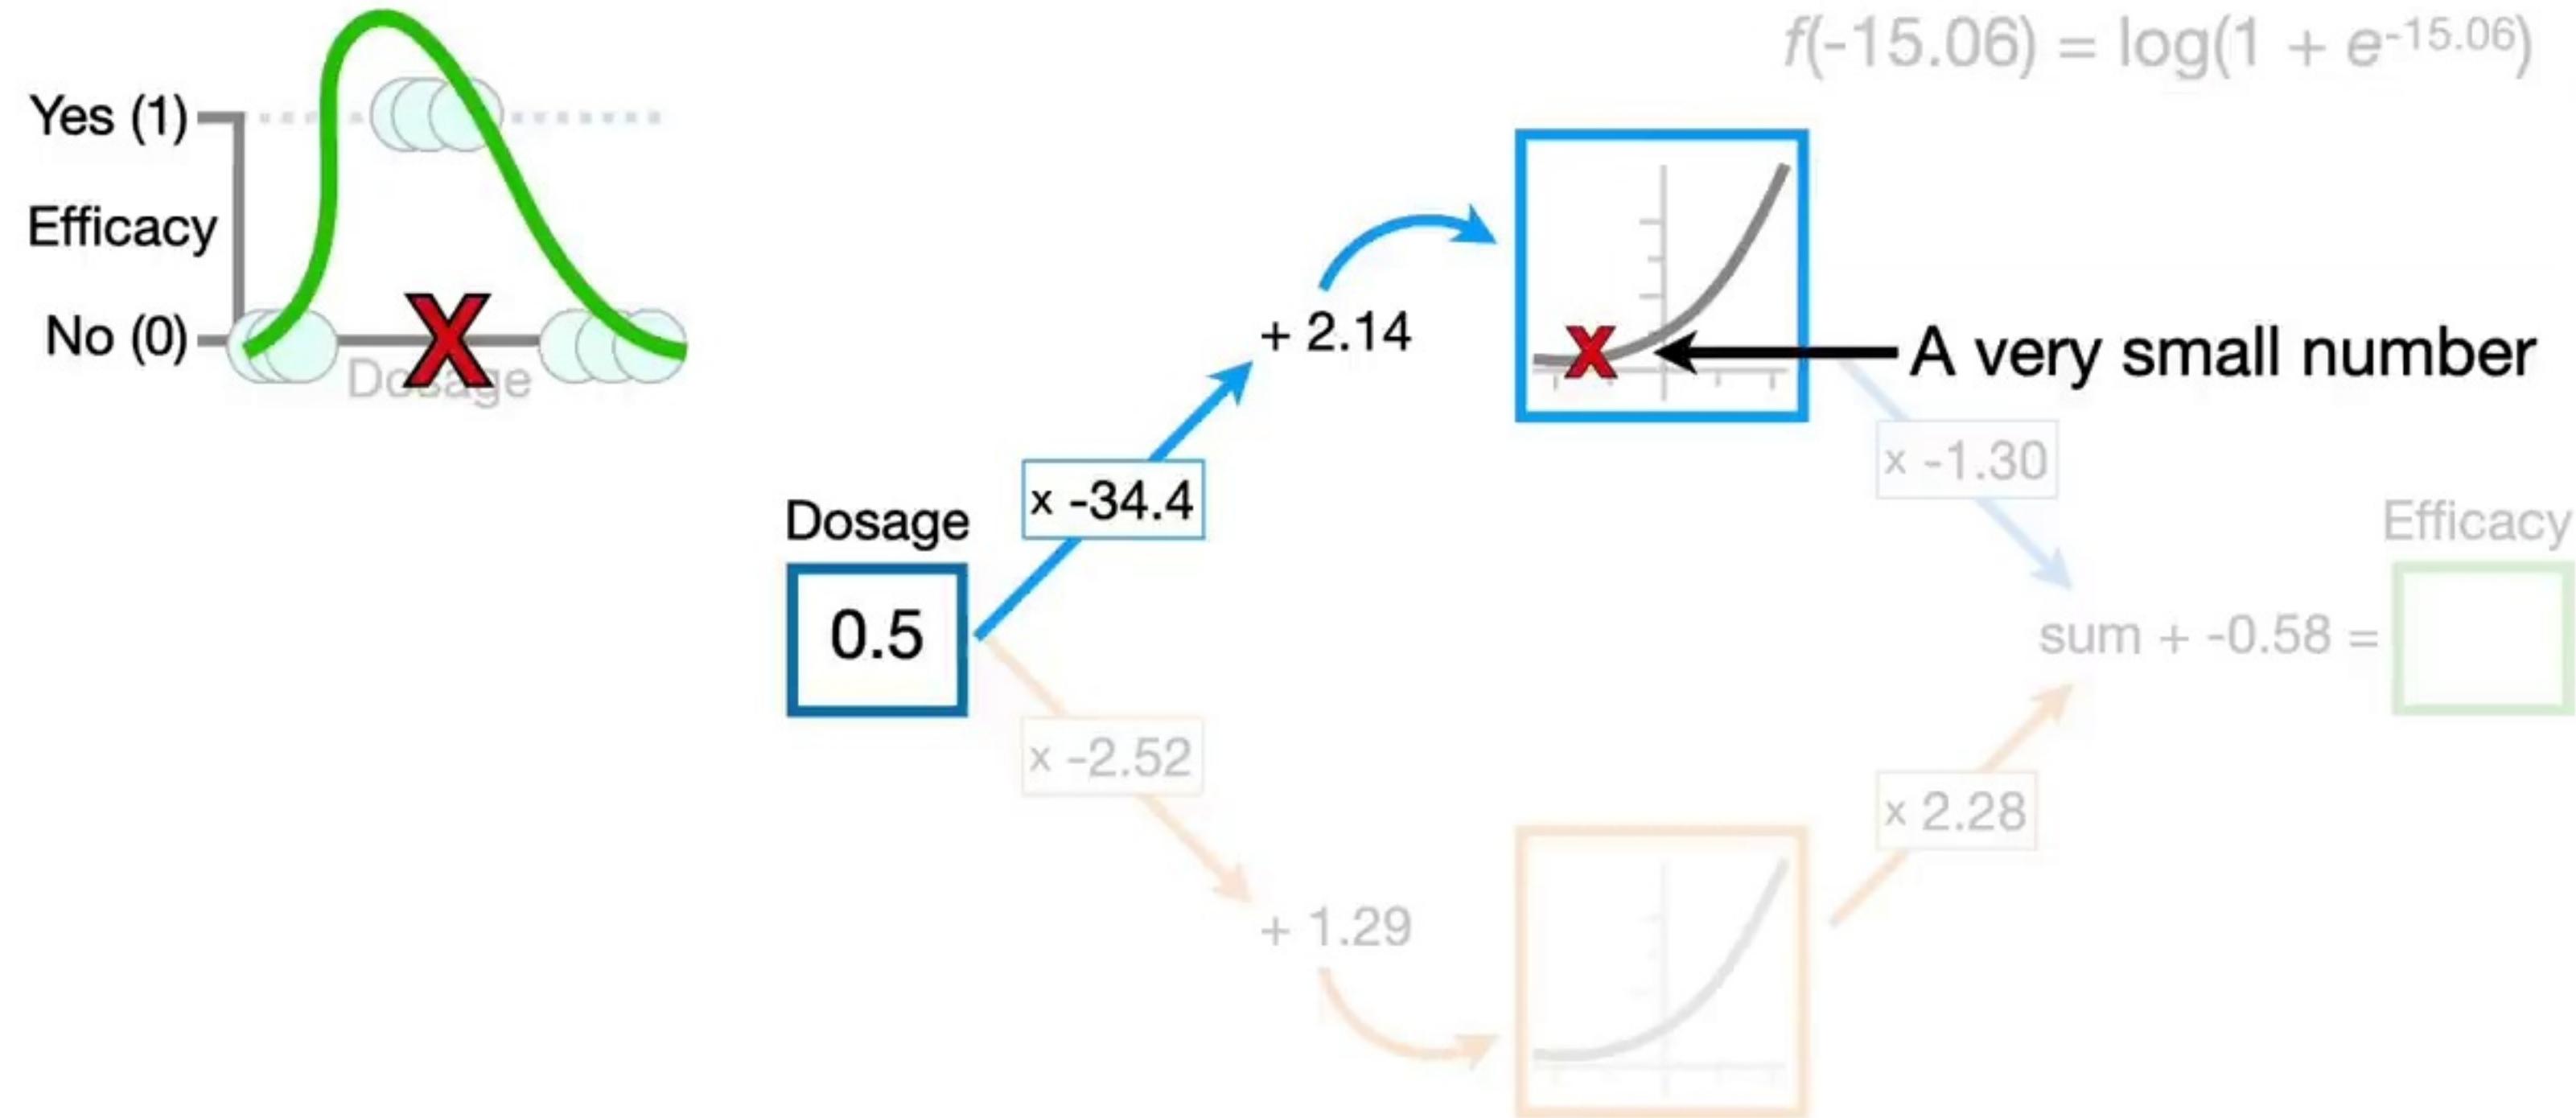
Double  
BAM!!  
**SO!**

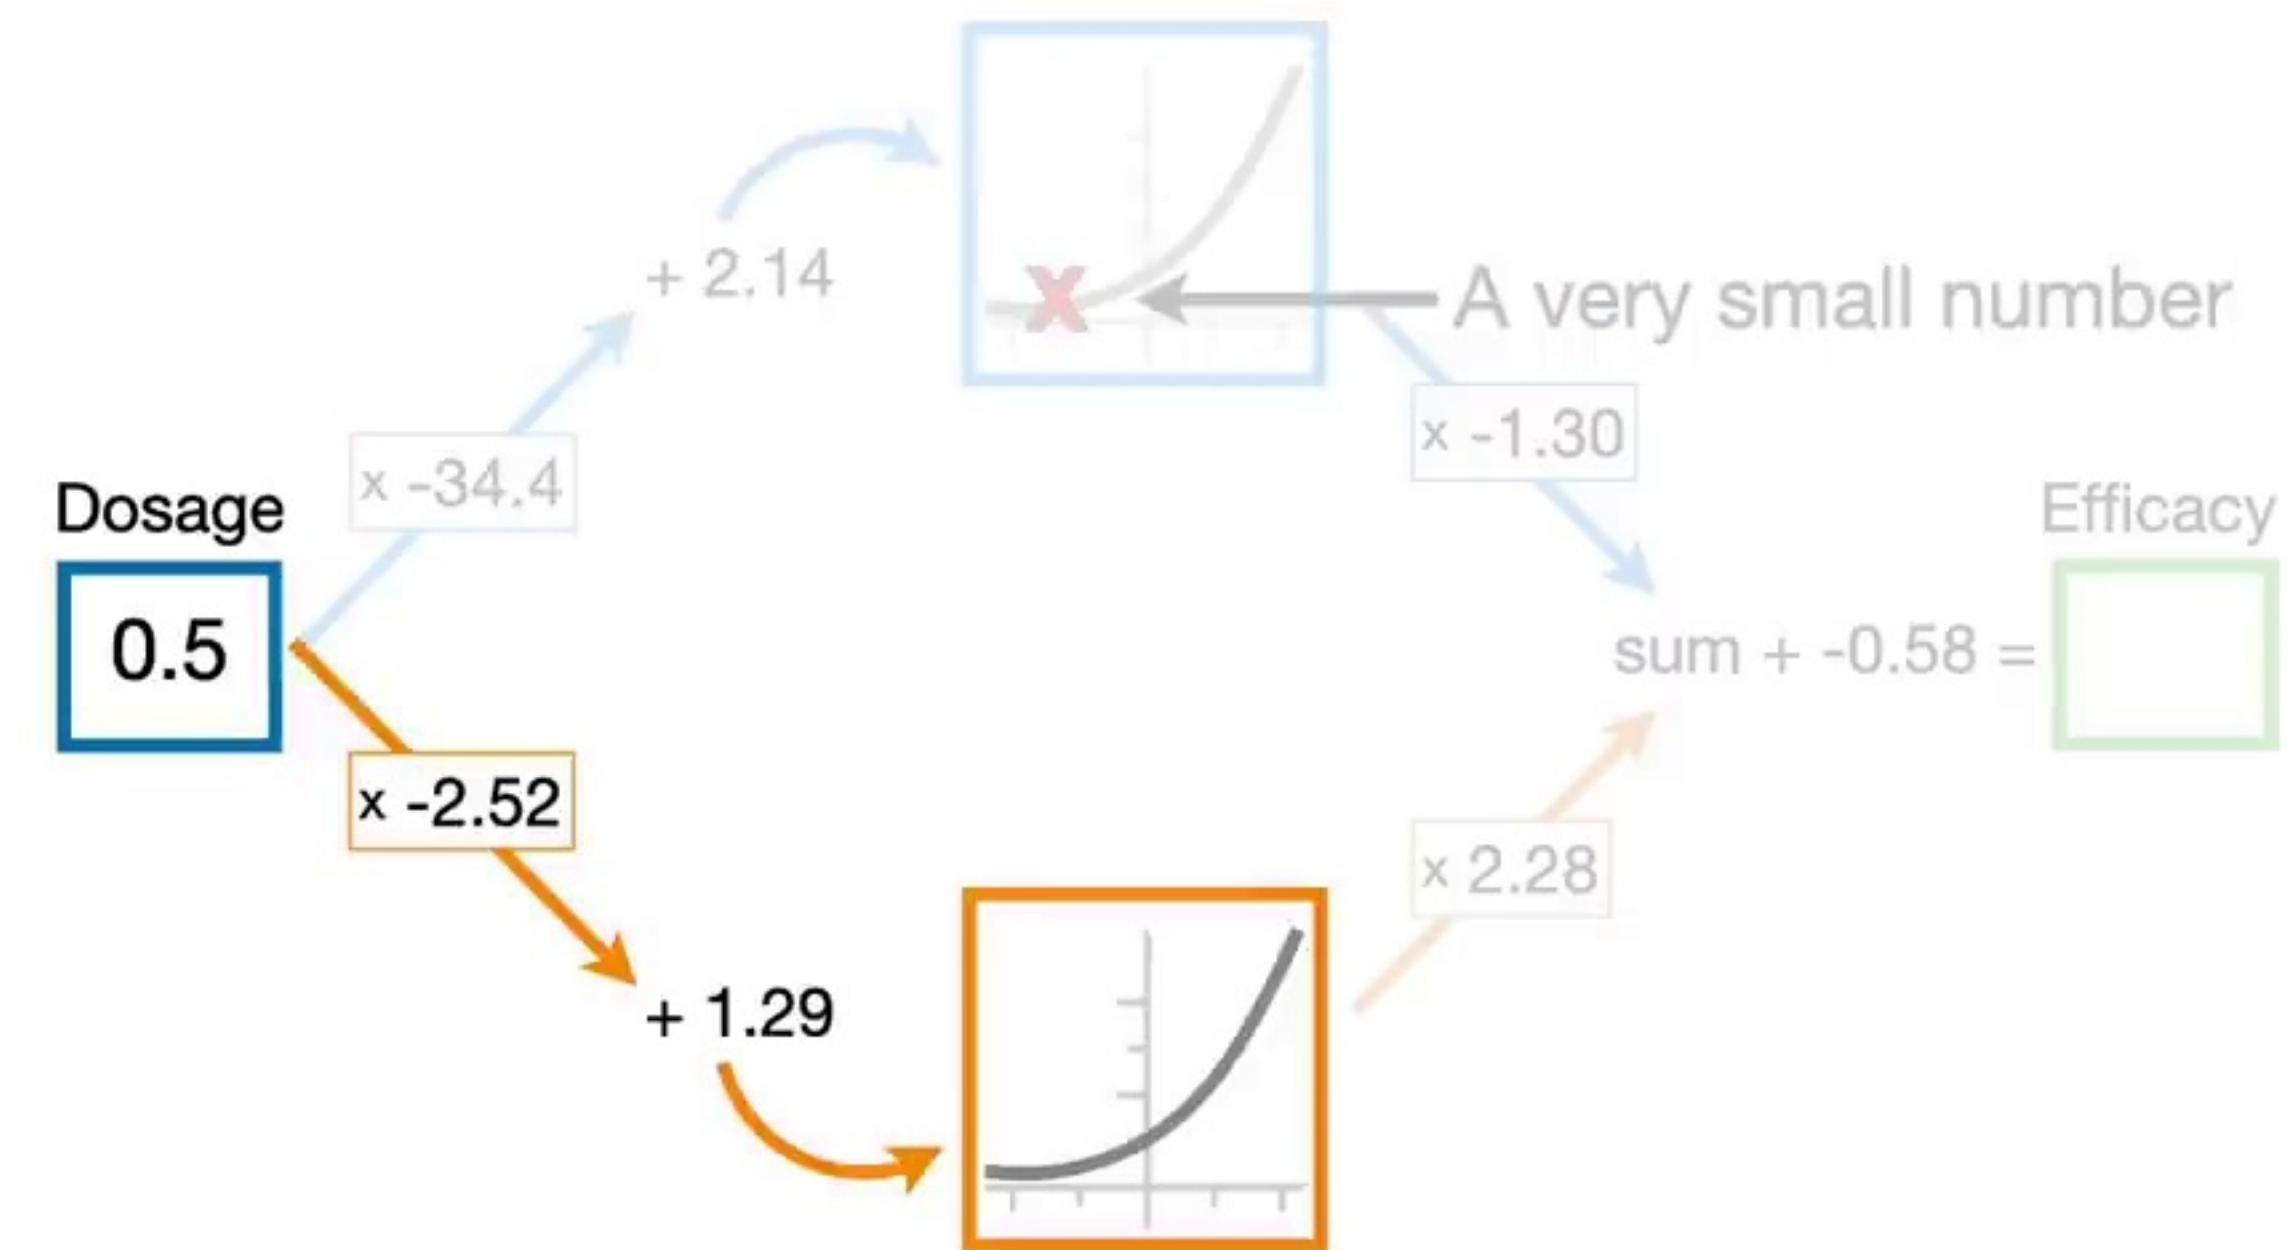
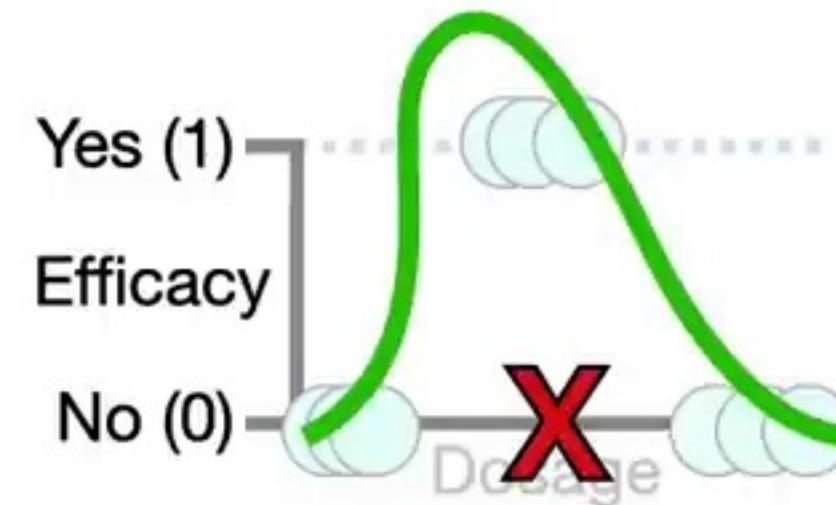
$$(0.5 \times -34.4) + 2.14 = -15.06$$



**SQ!**  
Double  
BAM!!

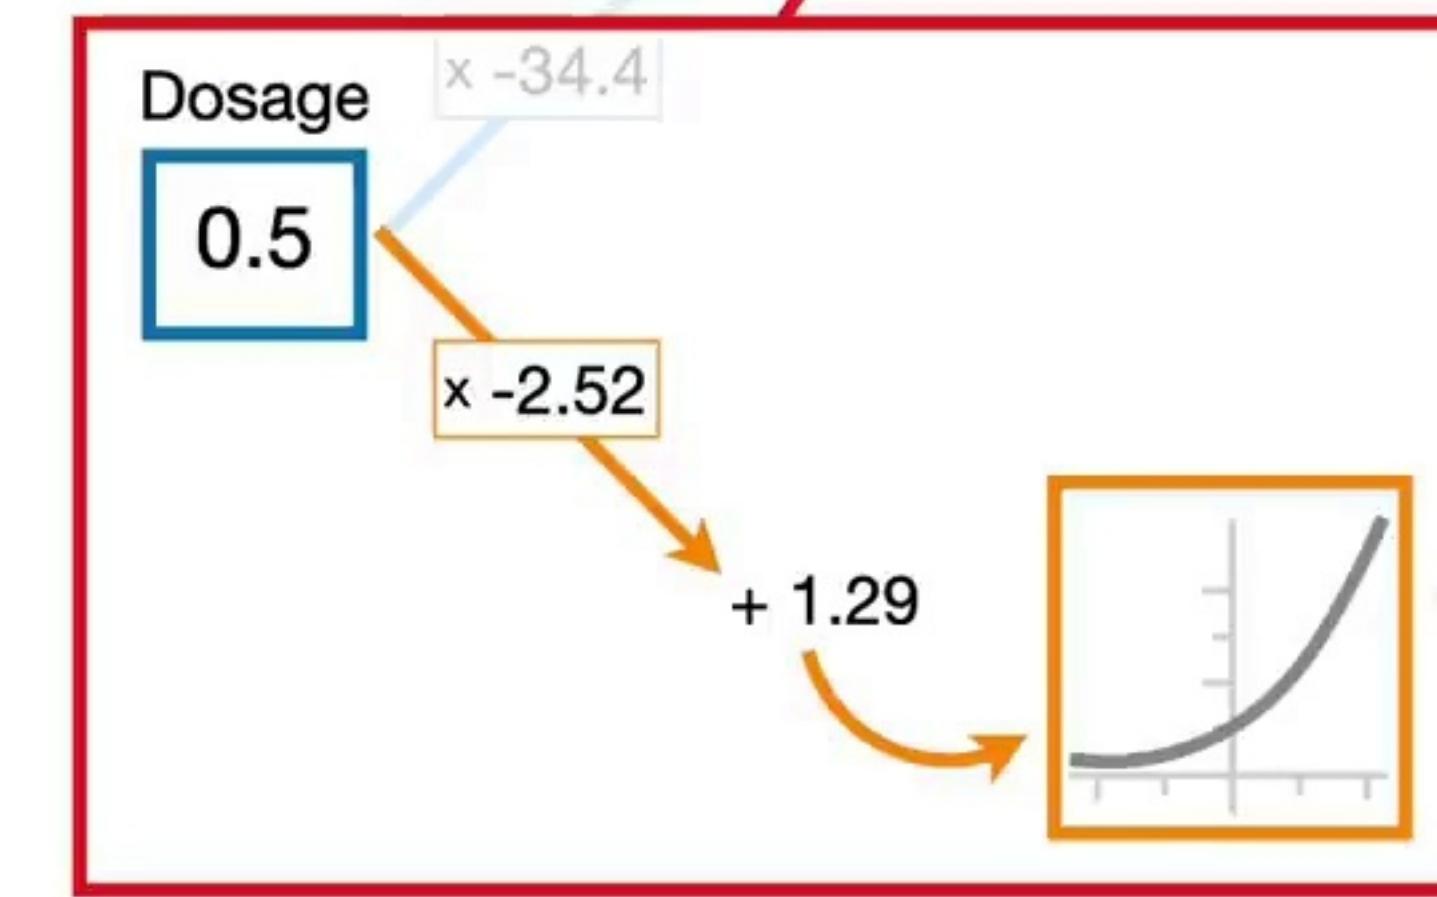
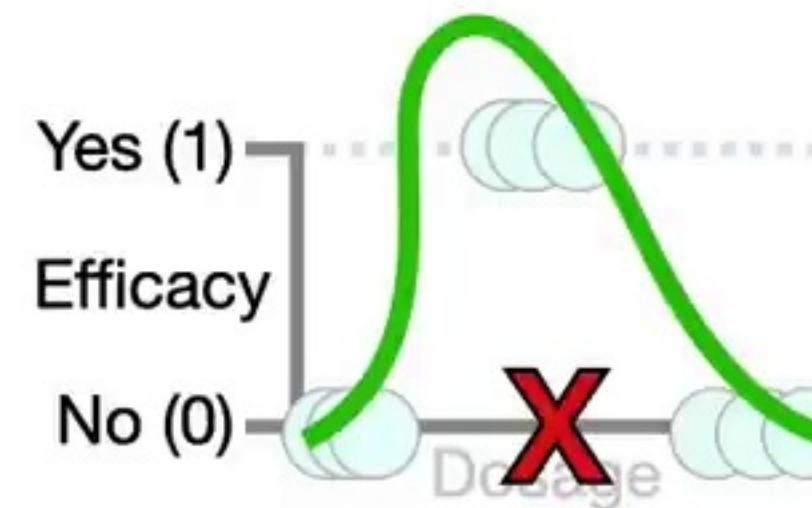
$$(0.5 \times -34.4) + 2.14 = -15.06$$





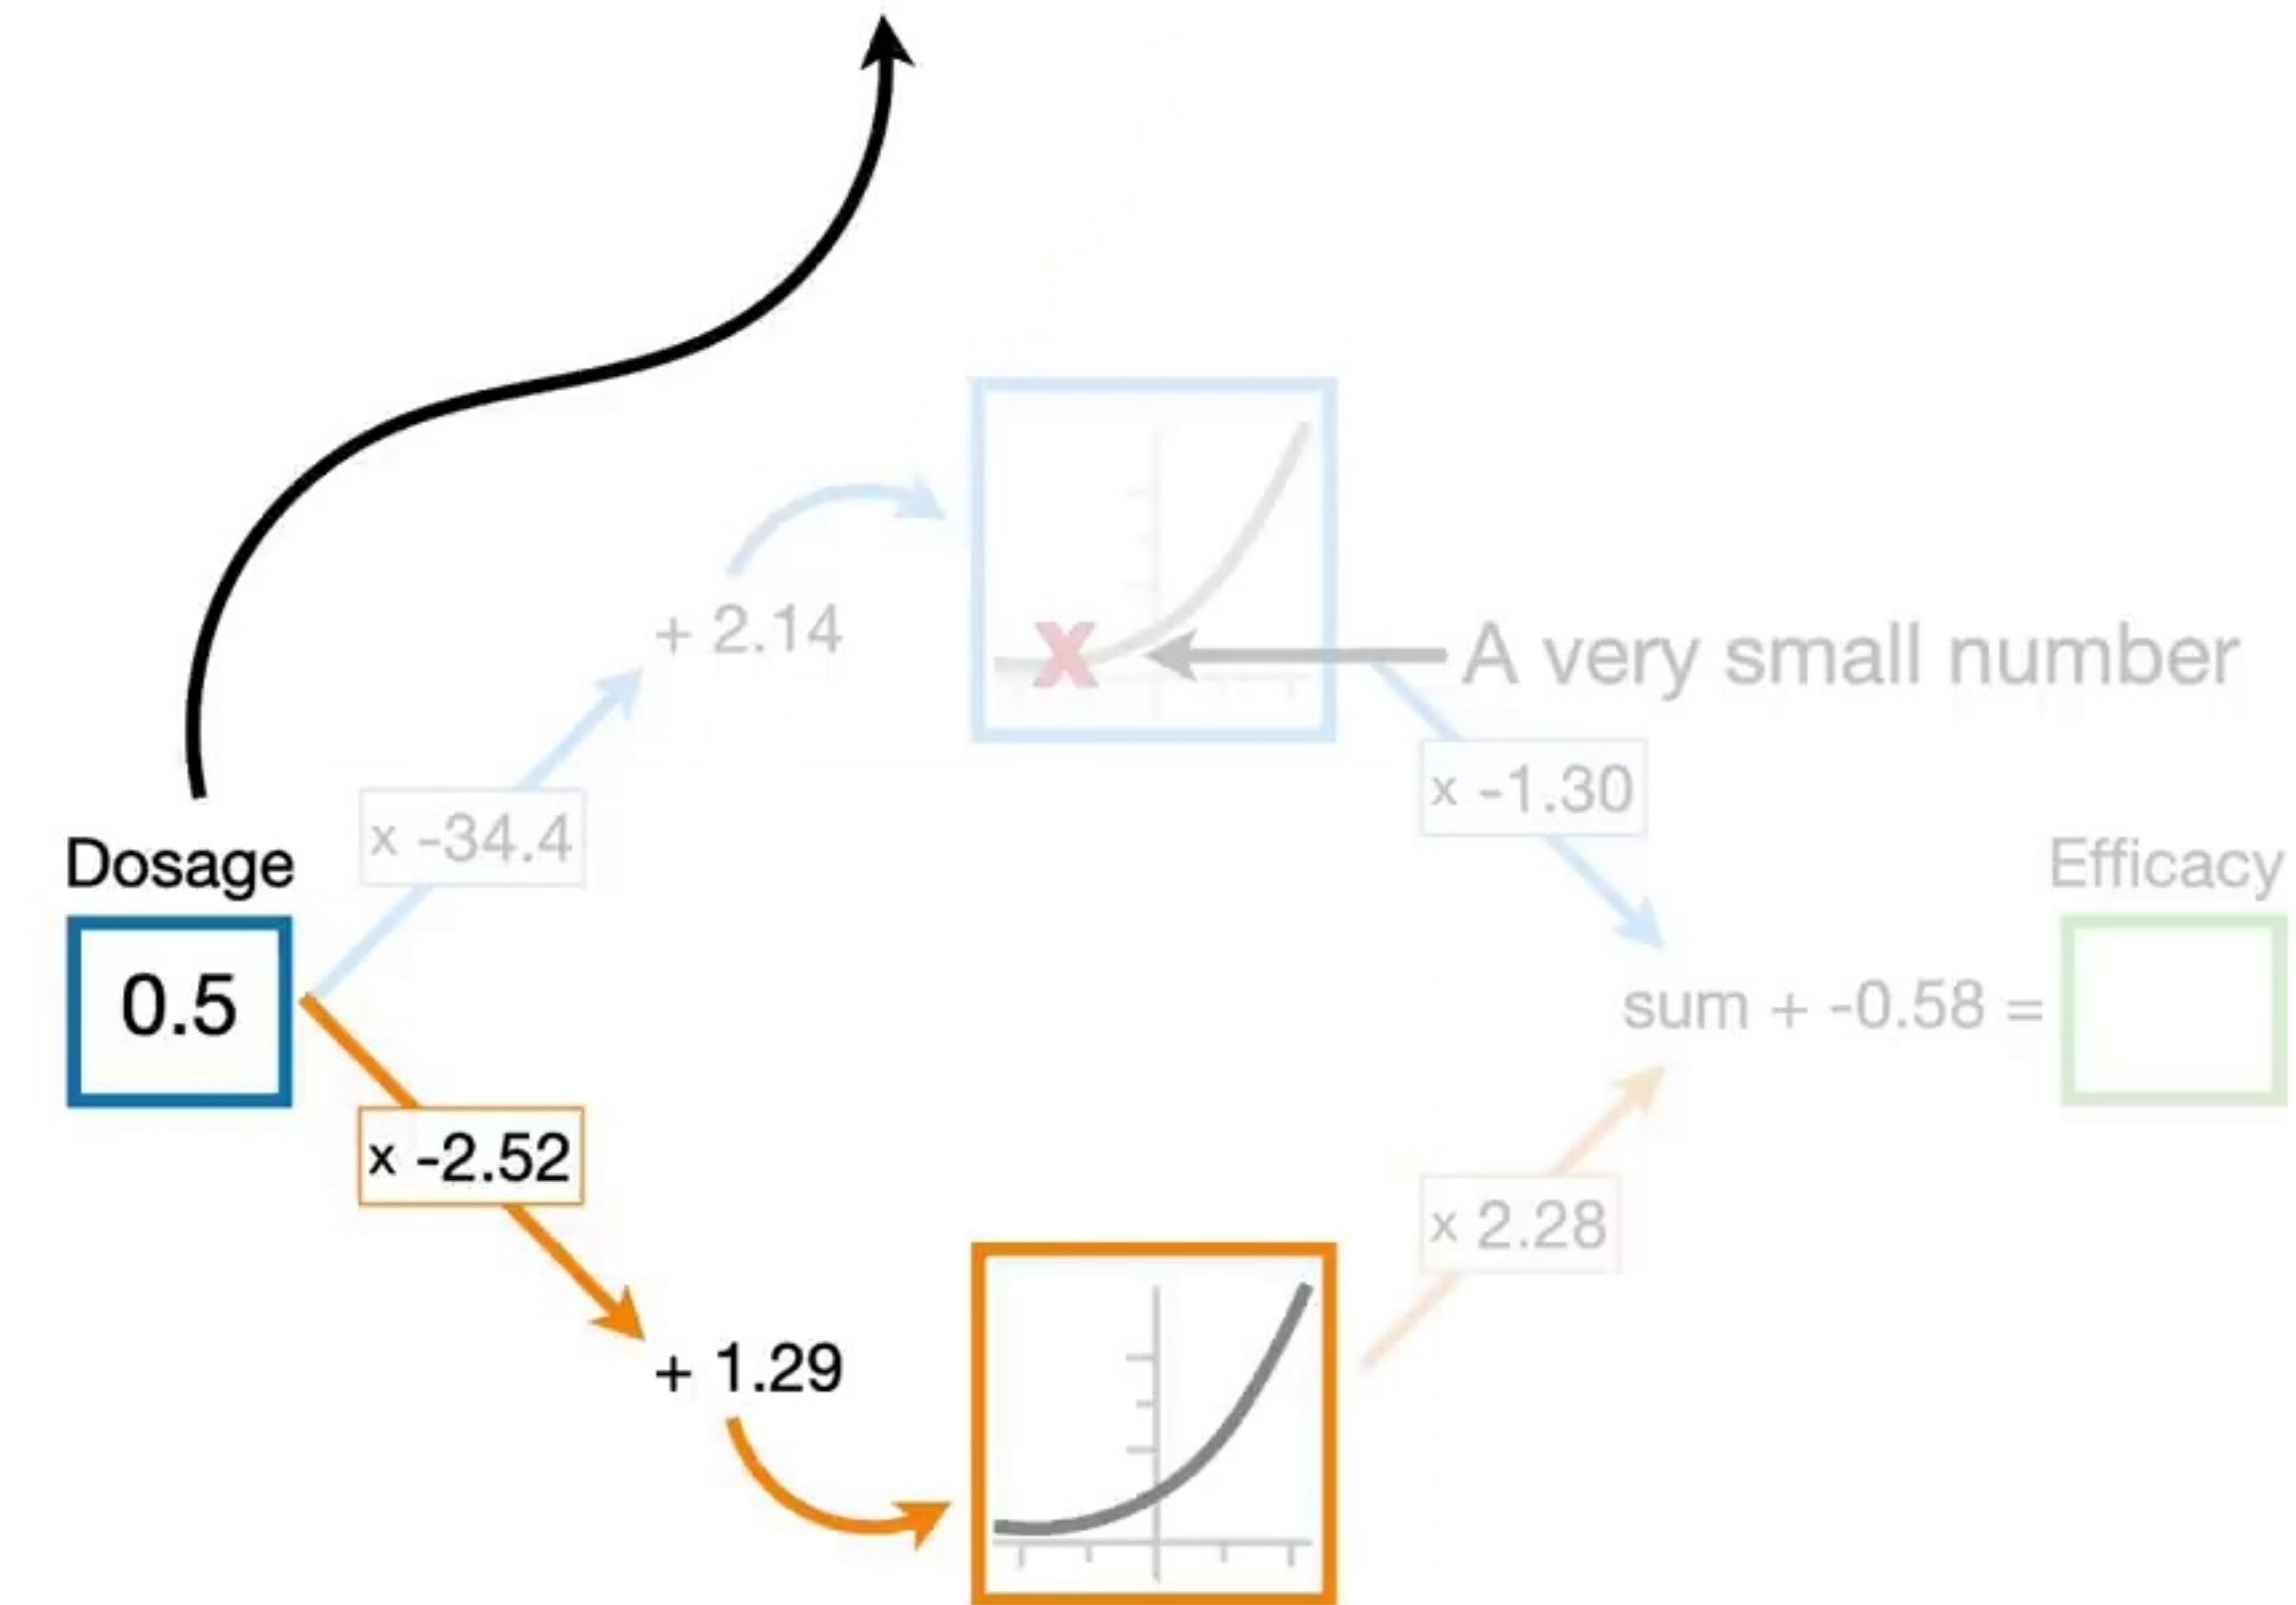
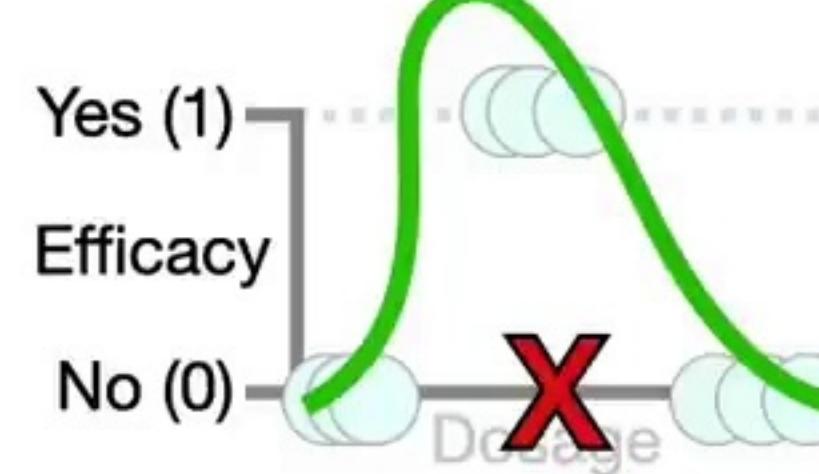
Double  
BAM!!  
**SQ!**

**(Dosage × -2.52) + 1.29 = x-axis coordinate**

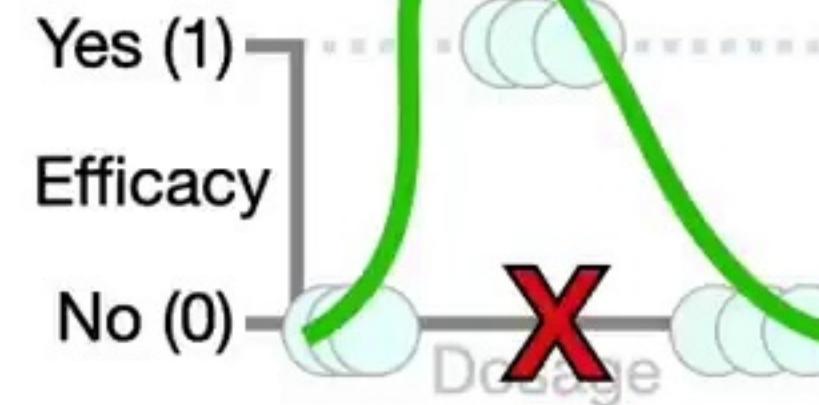


Double  
BAM!!  
**SO!**

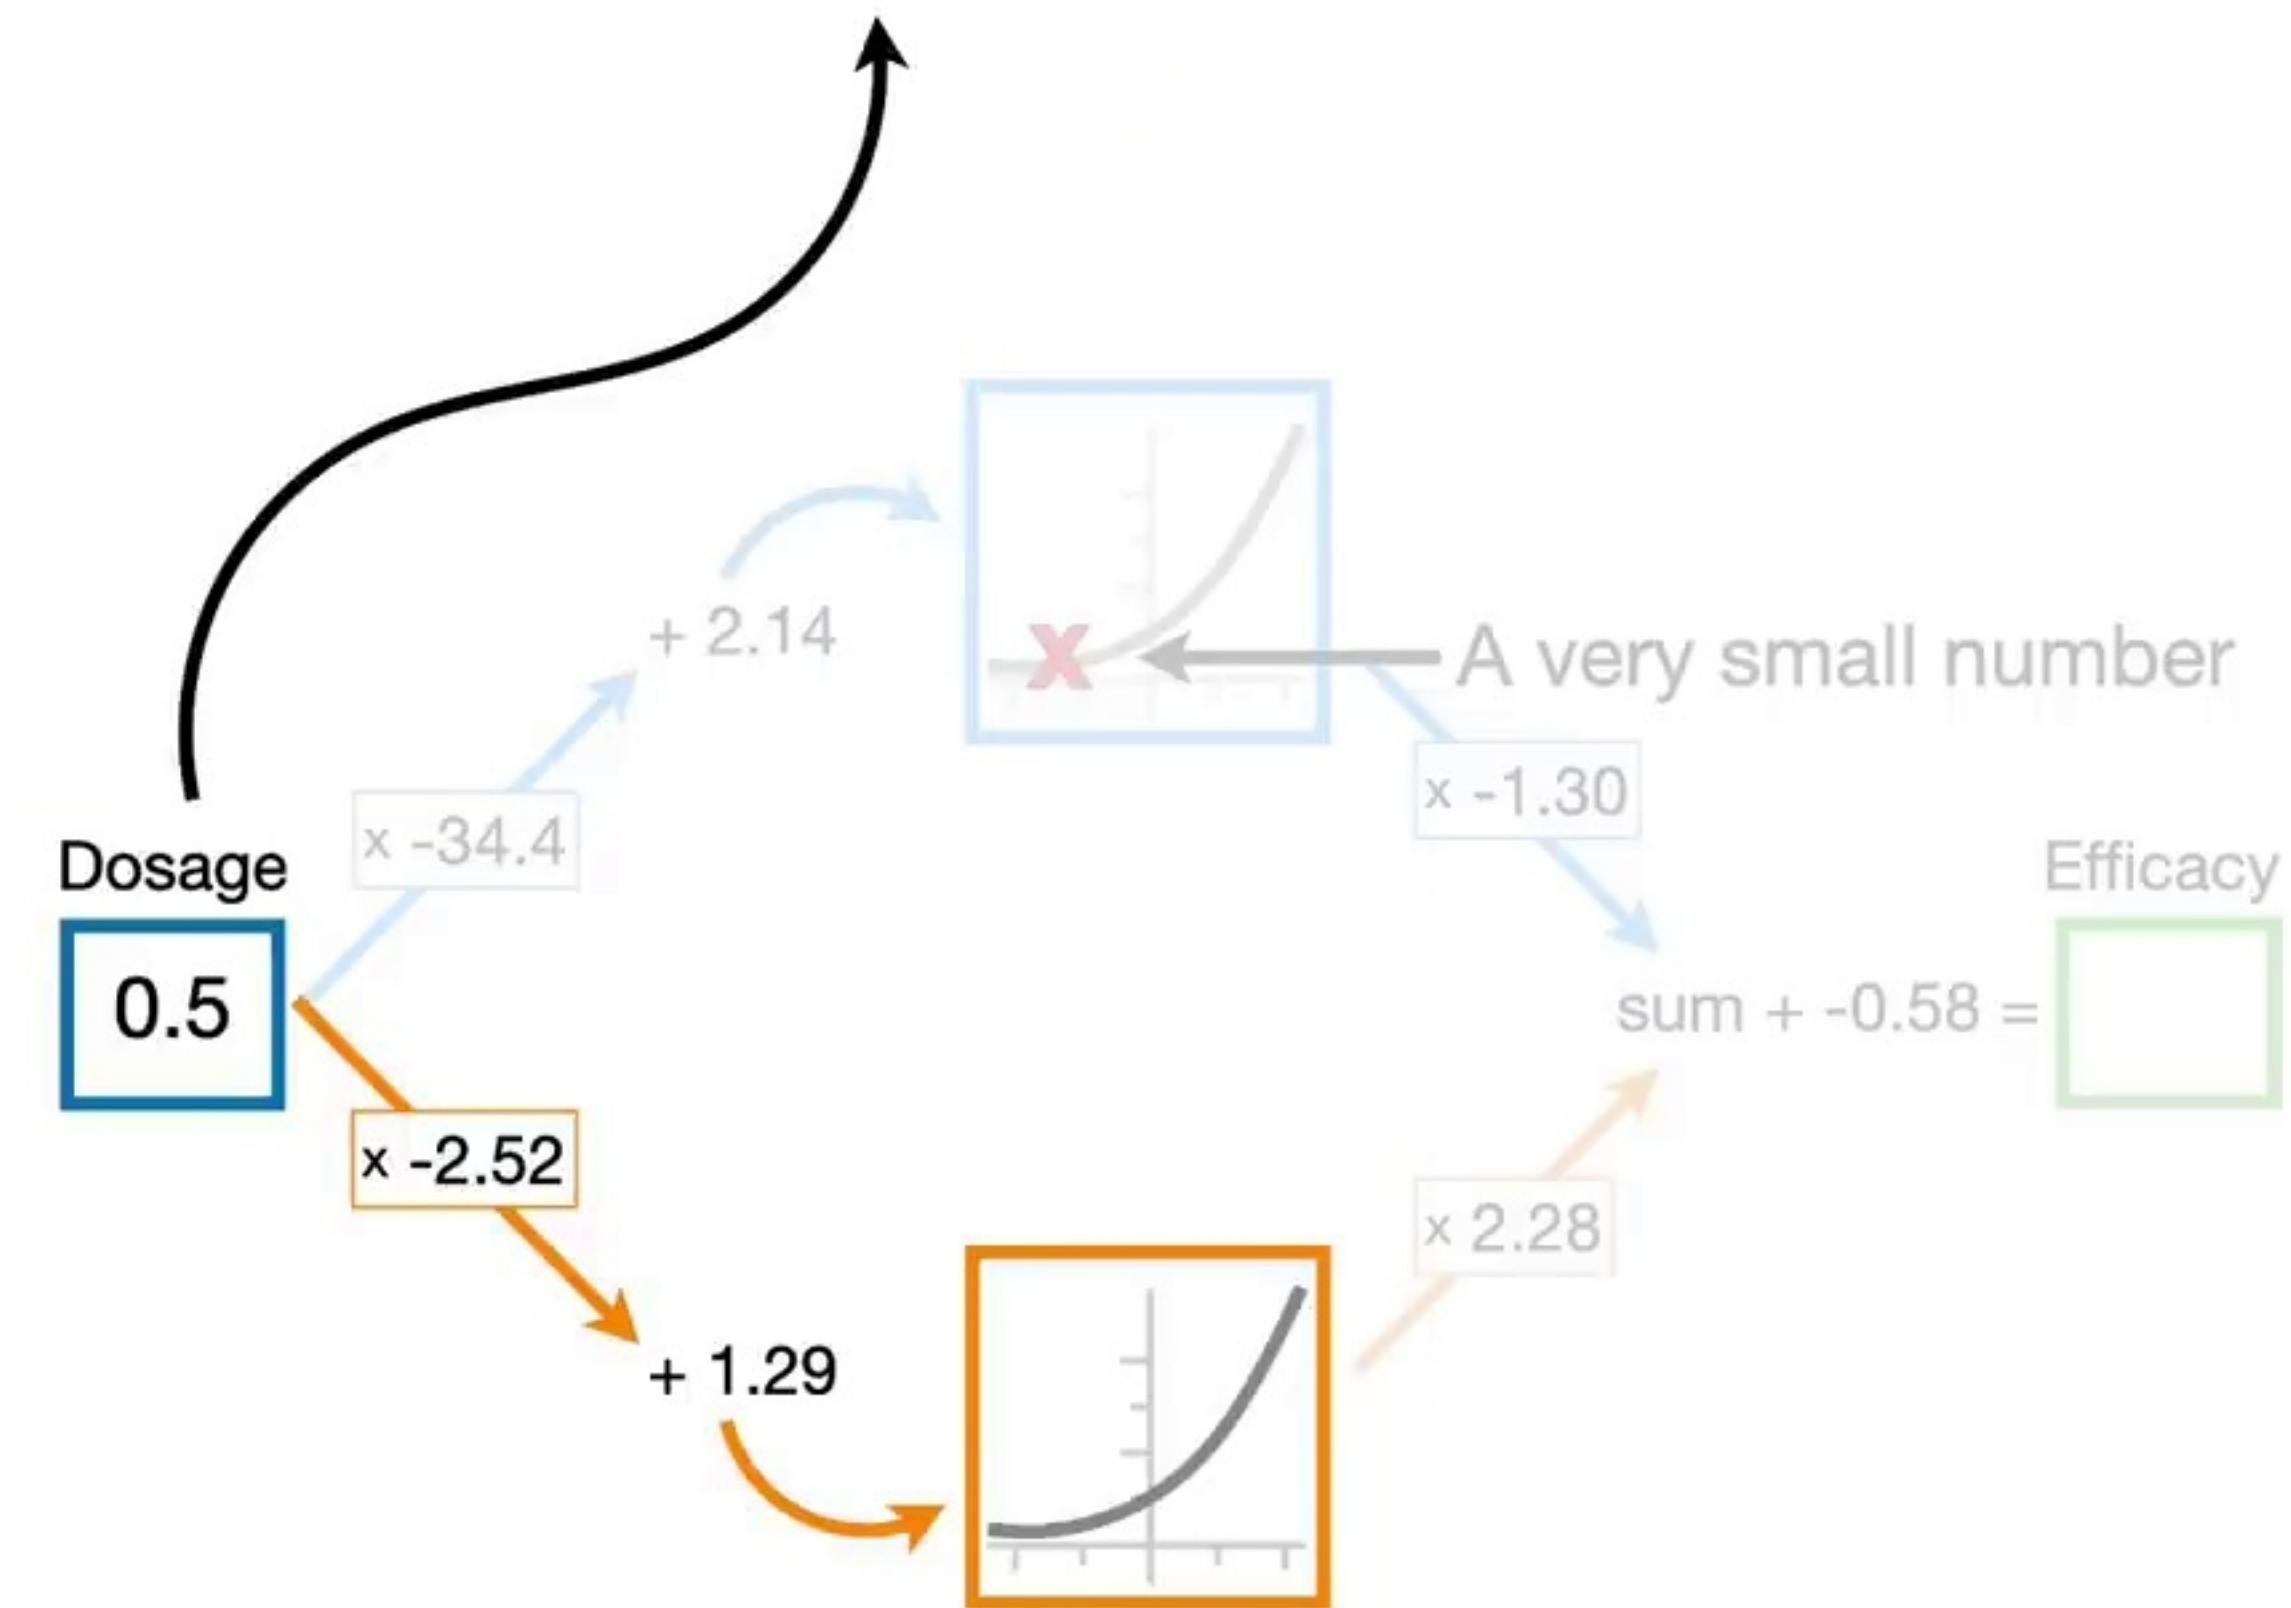
**(Dosage × -2.52) + 1.29 = x-axis coordinate**



Double  
BAM!!  
**SQ!**

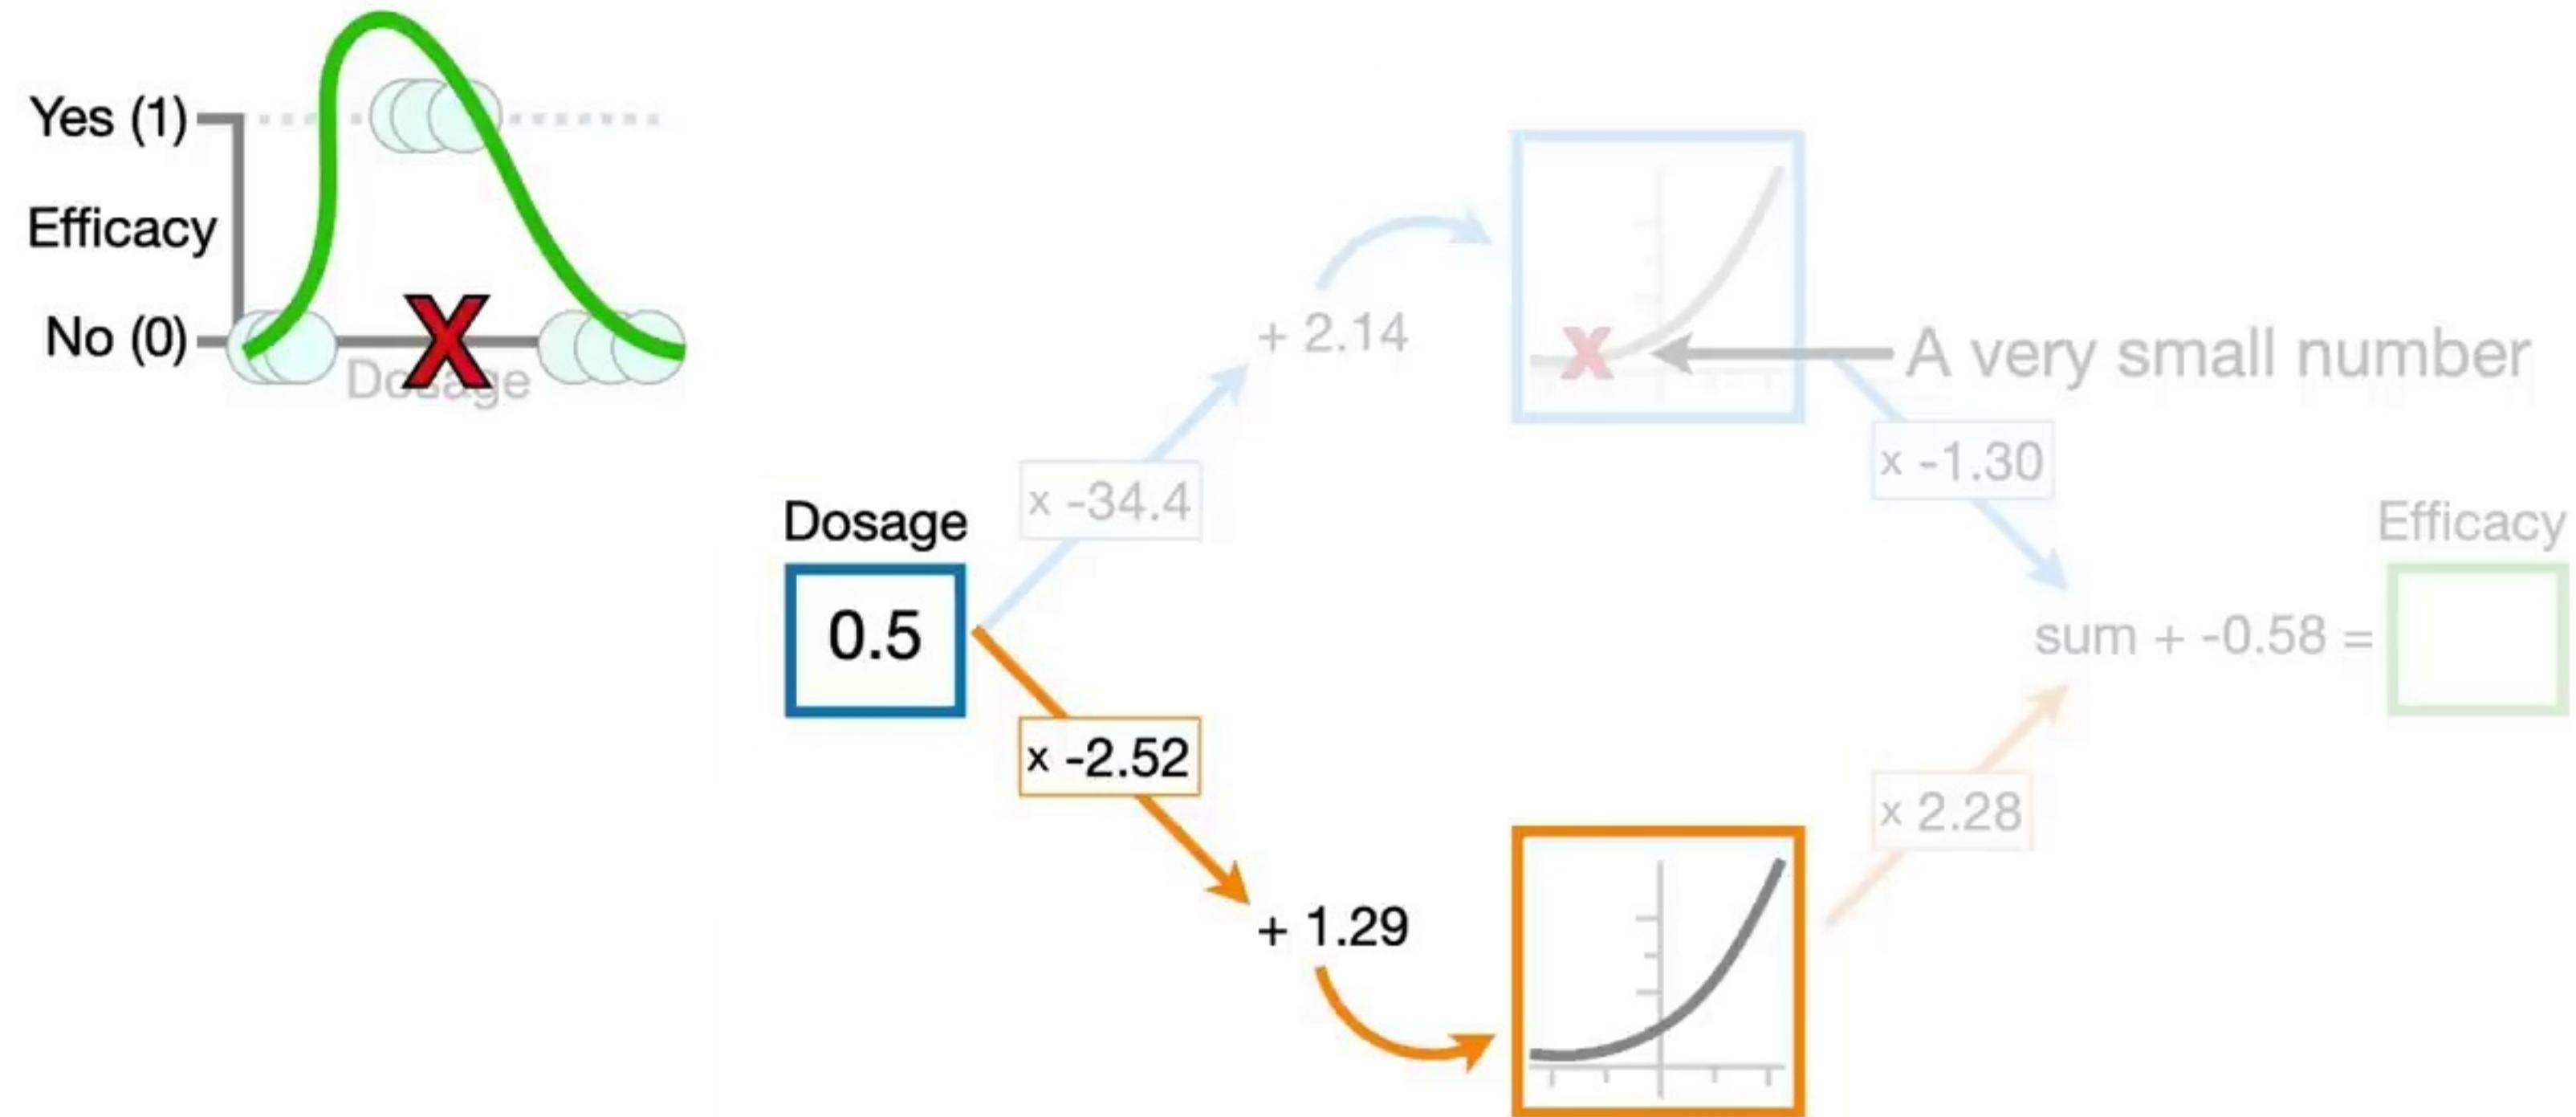


$$(0.5 \times -2.52) + 1.29 = \text{x-axis coordinate}$$



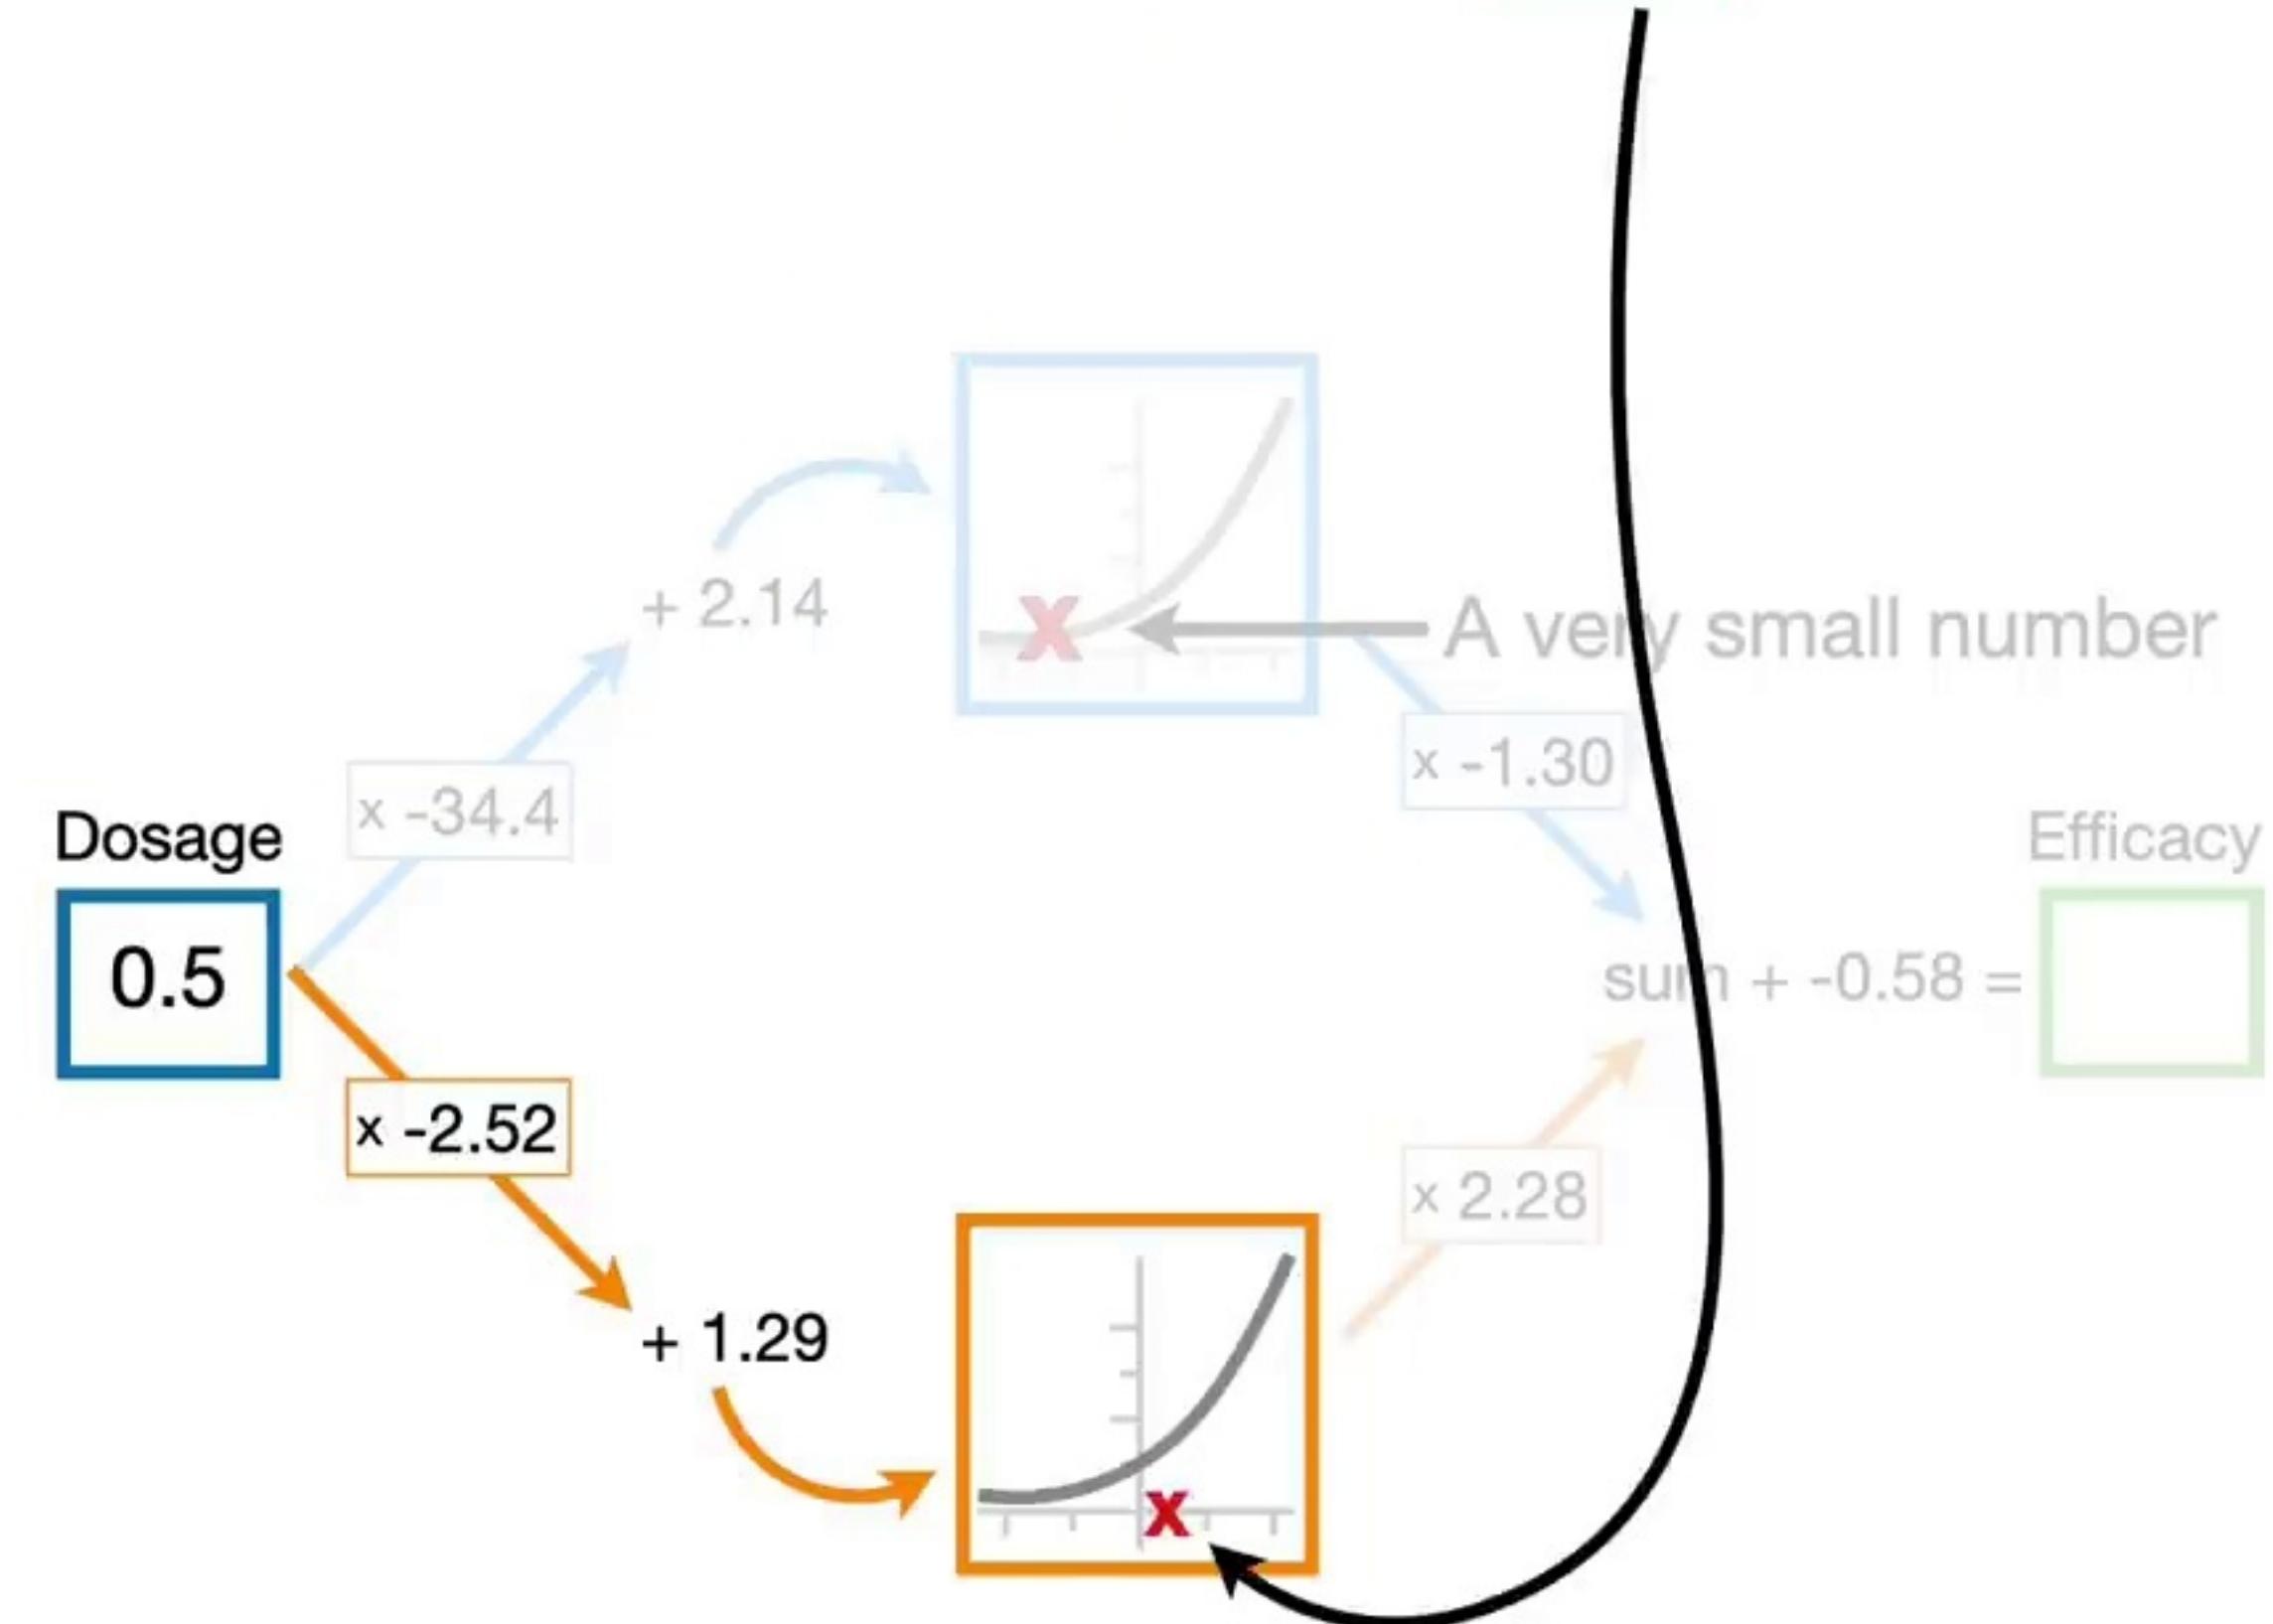
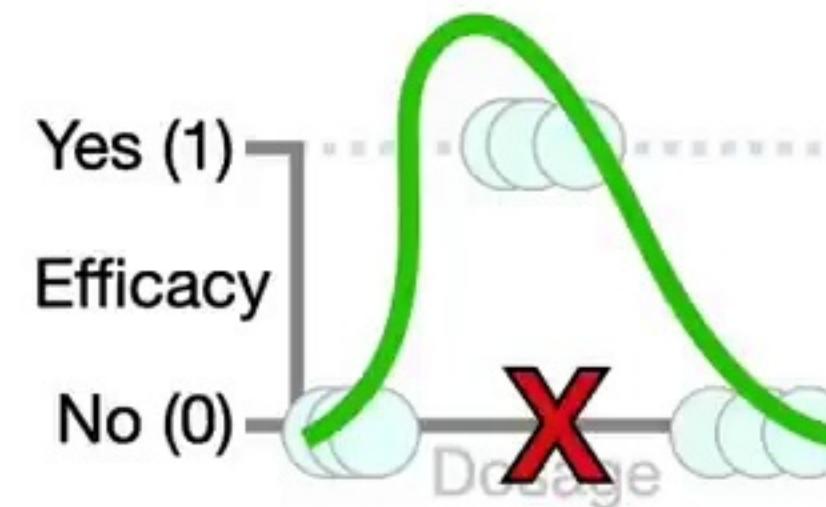


$$(0.5 \times -2.52) + 1.29 = 0.3$$



Double  
BAM!!  
**SO!**

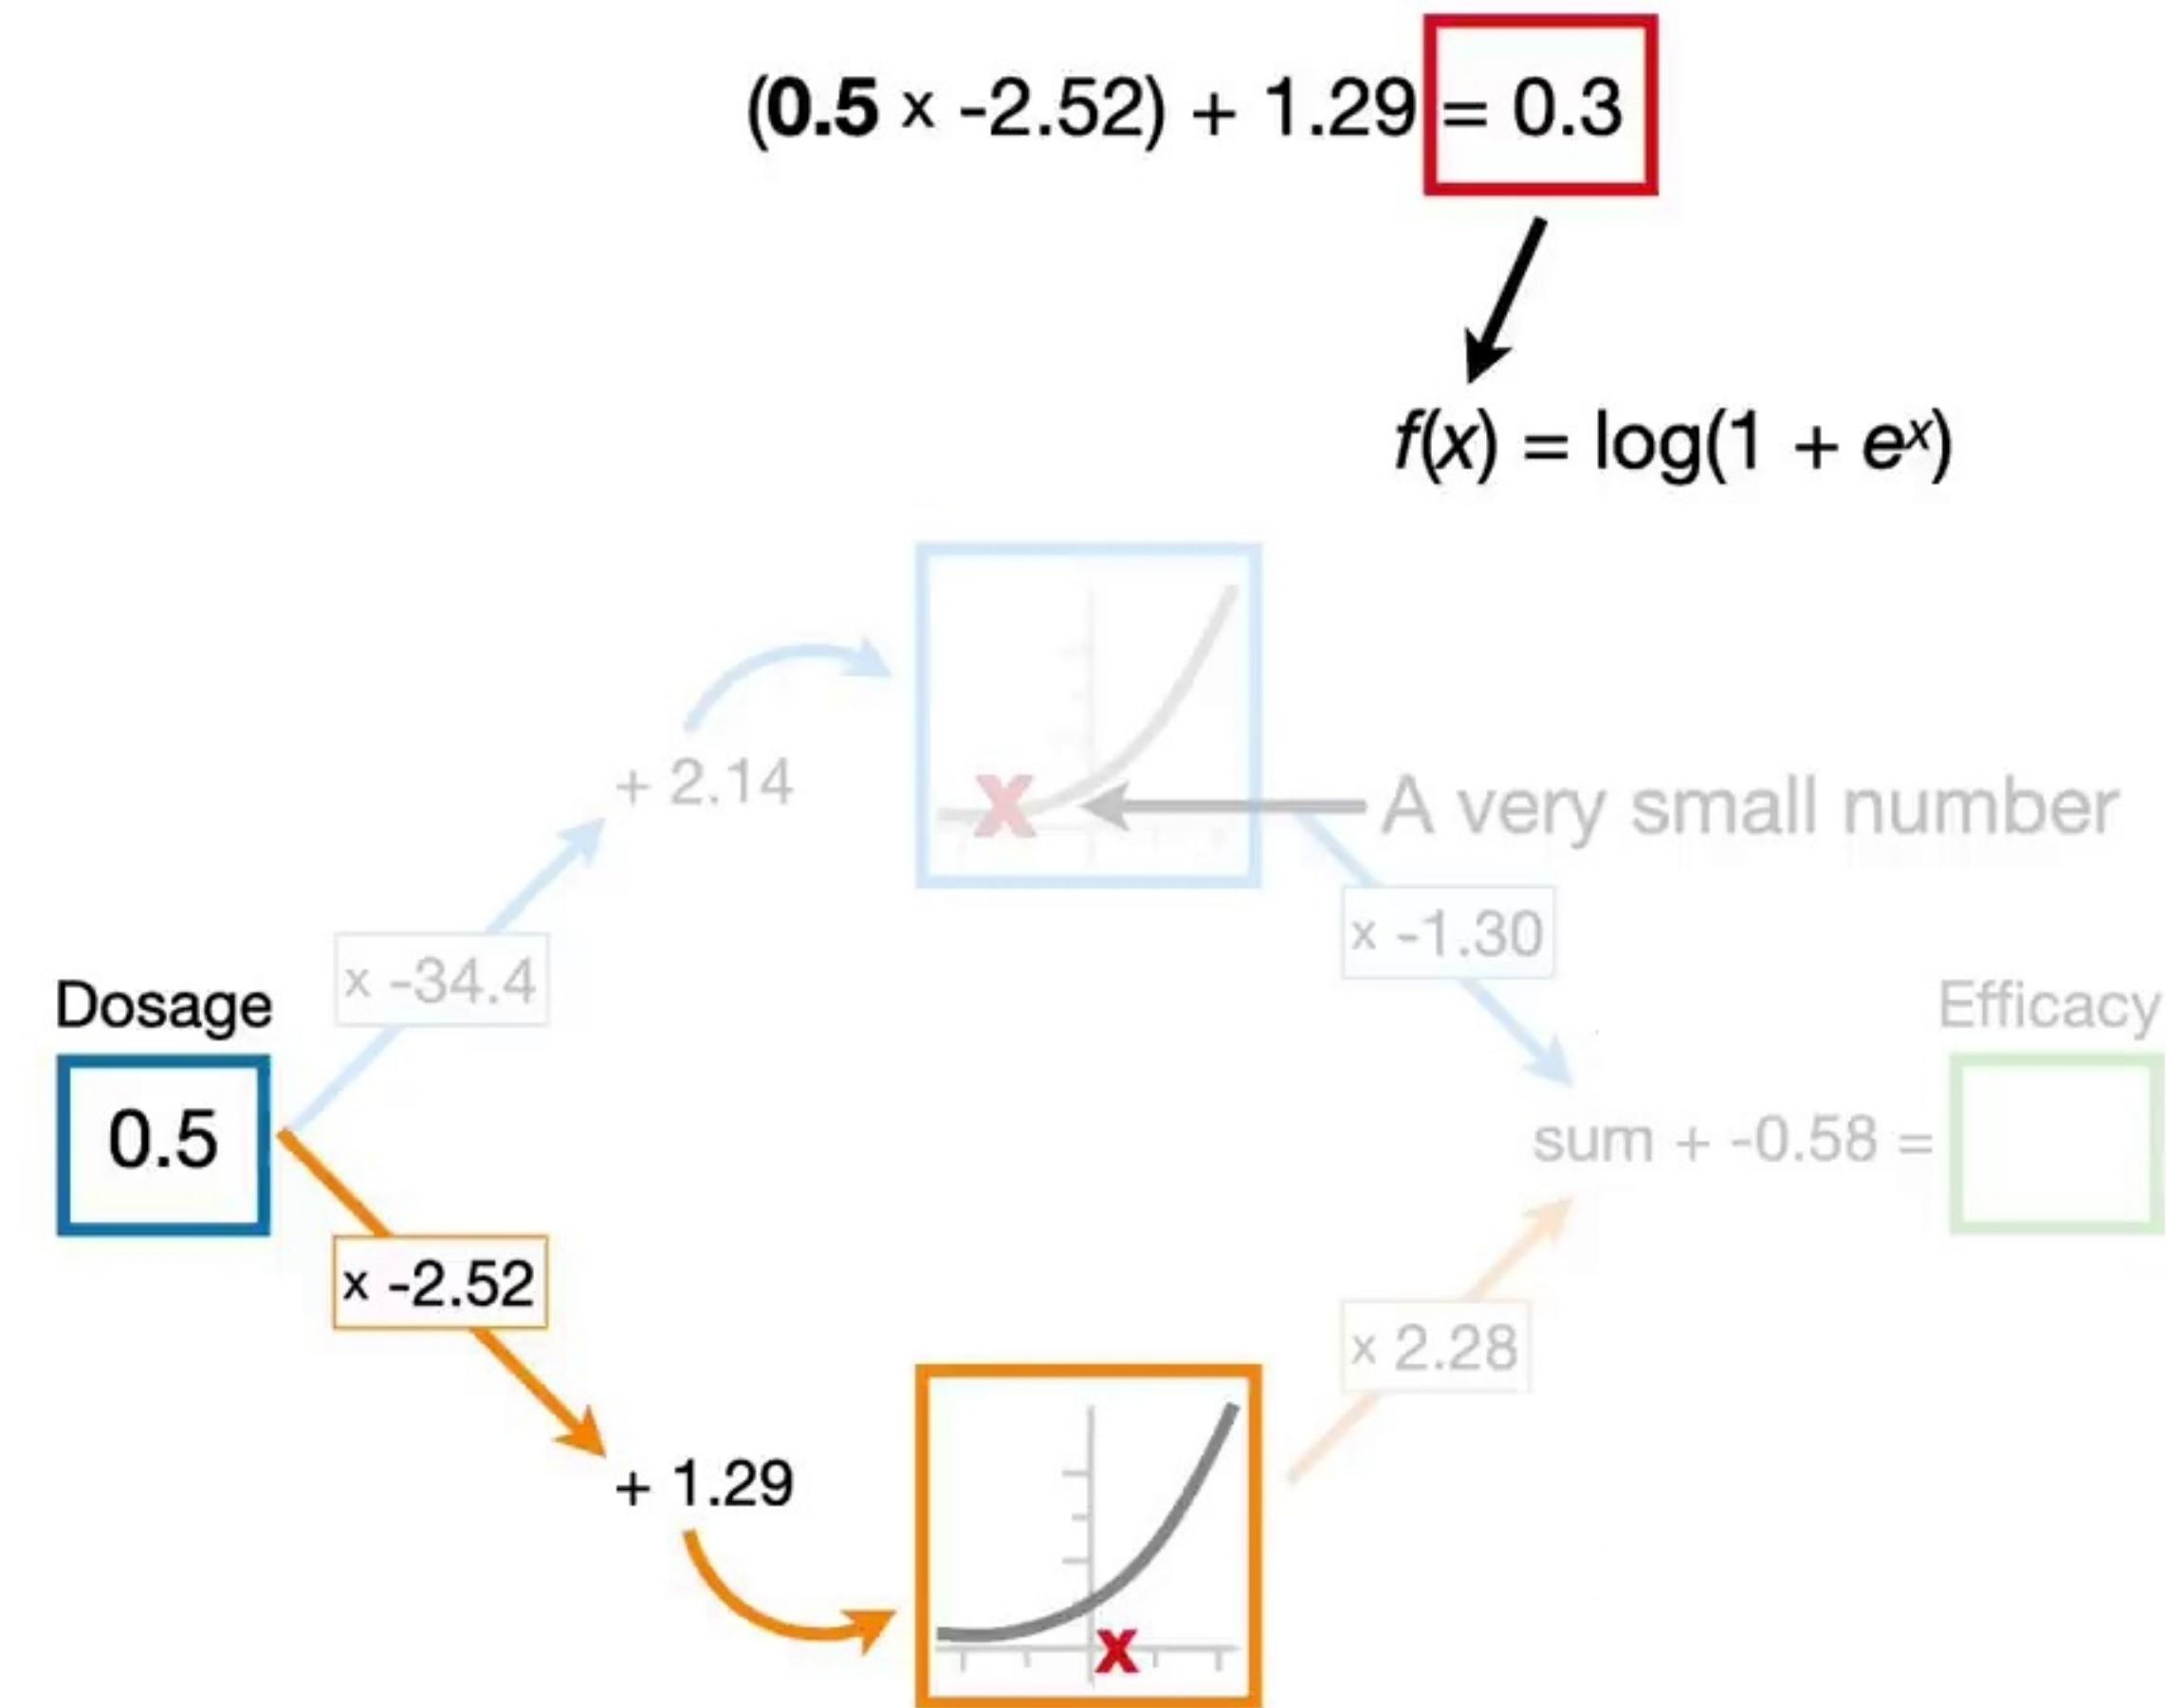
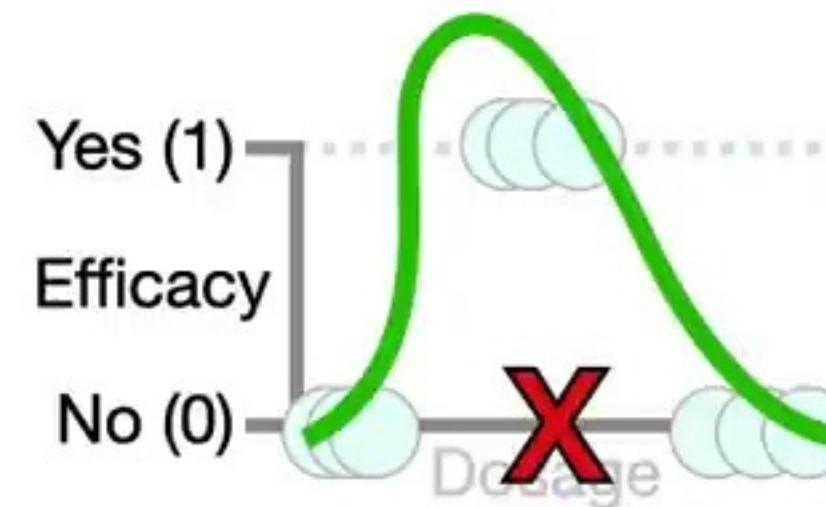
$$(0.5 \times -2.52) + 1.29 = 0.3$$



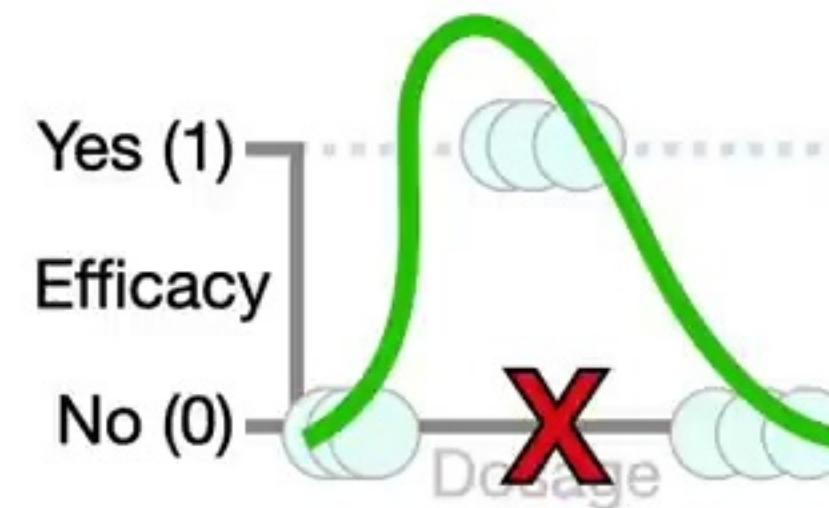
double  
BAM!!  
**SQ!**

$$(0.5 \times -2.52) + 1.29 = 0.3$$

$$f(x) = \log(1 + e^x)$$

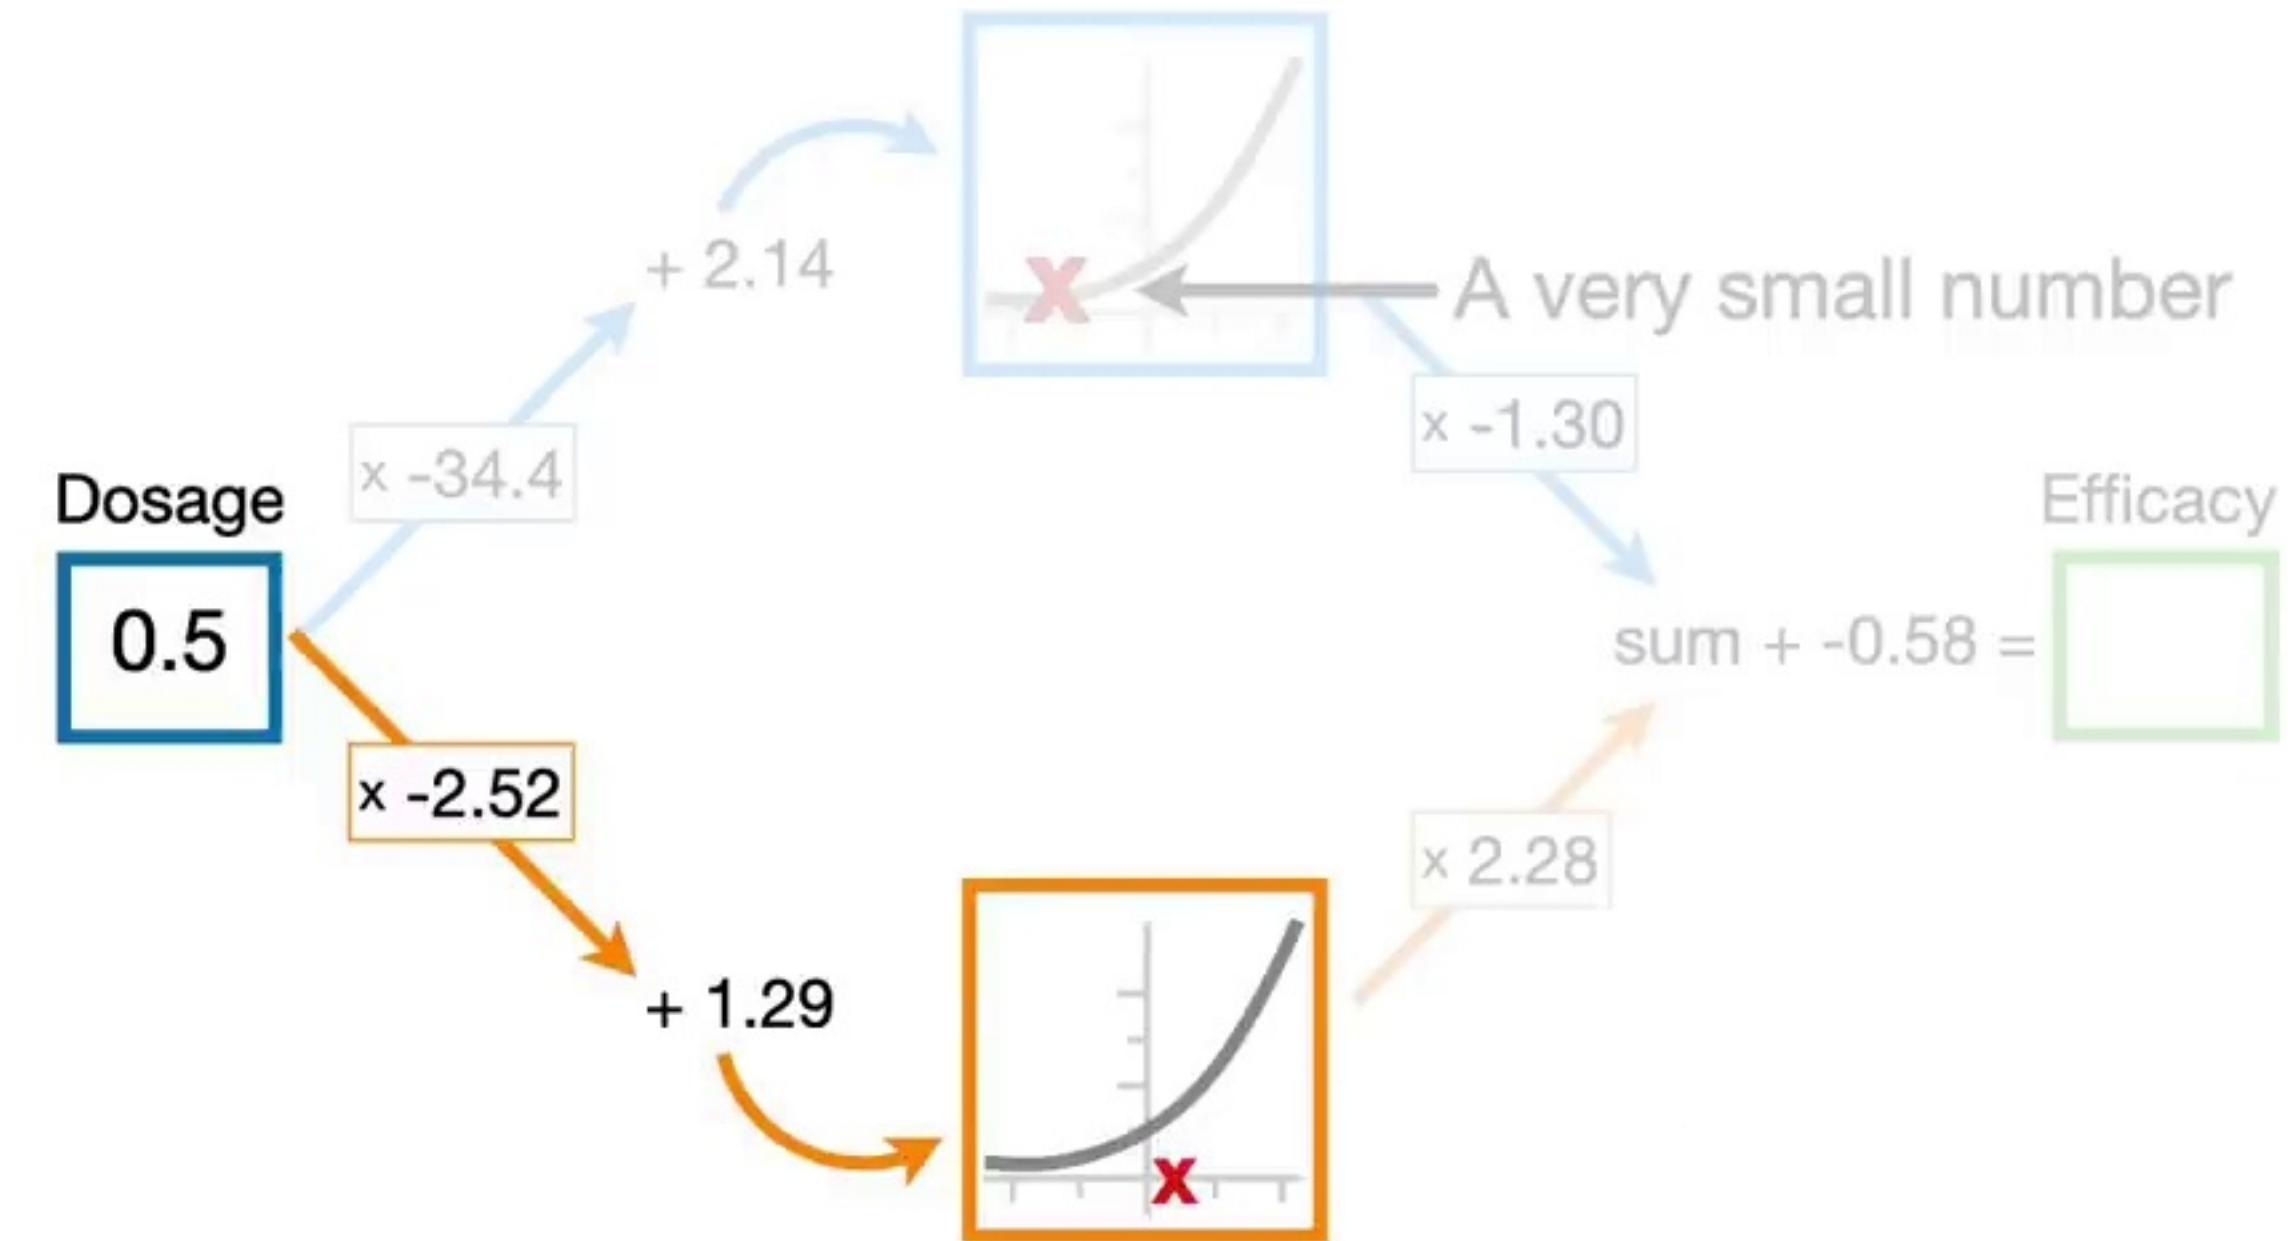


**Double  
BAM!!**



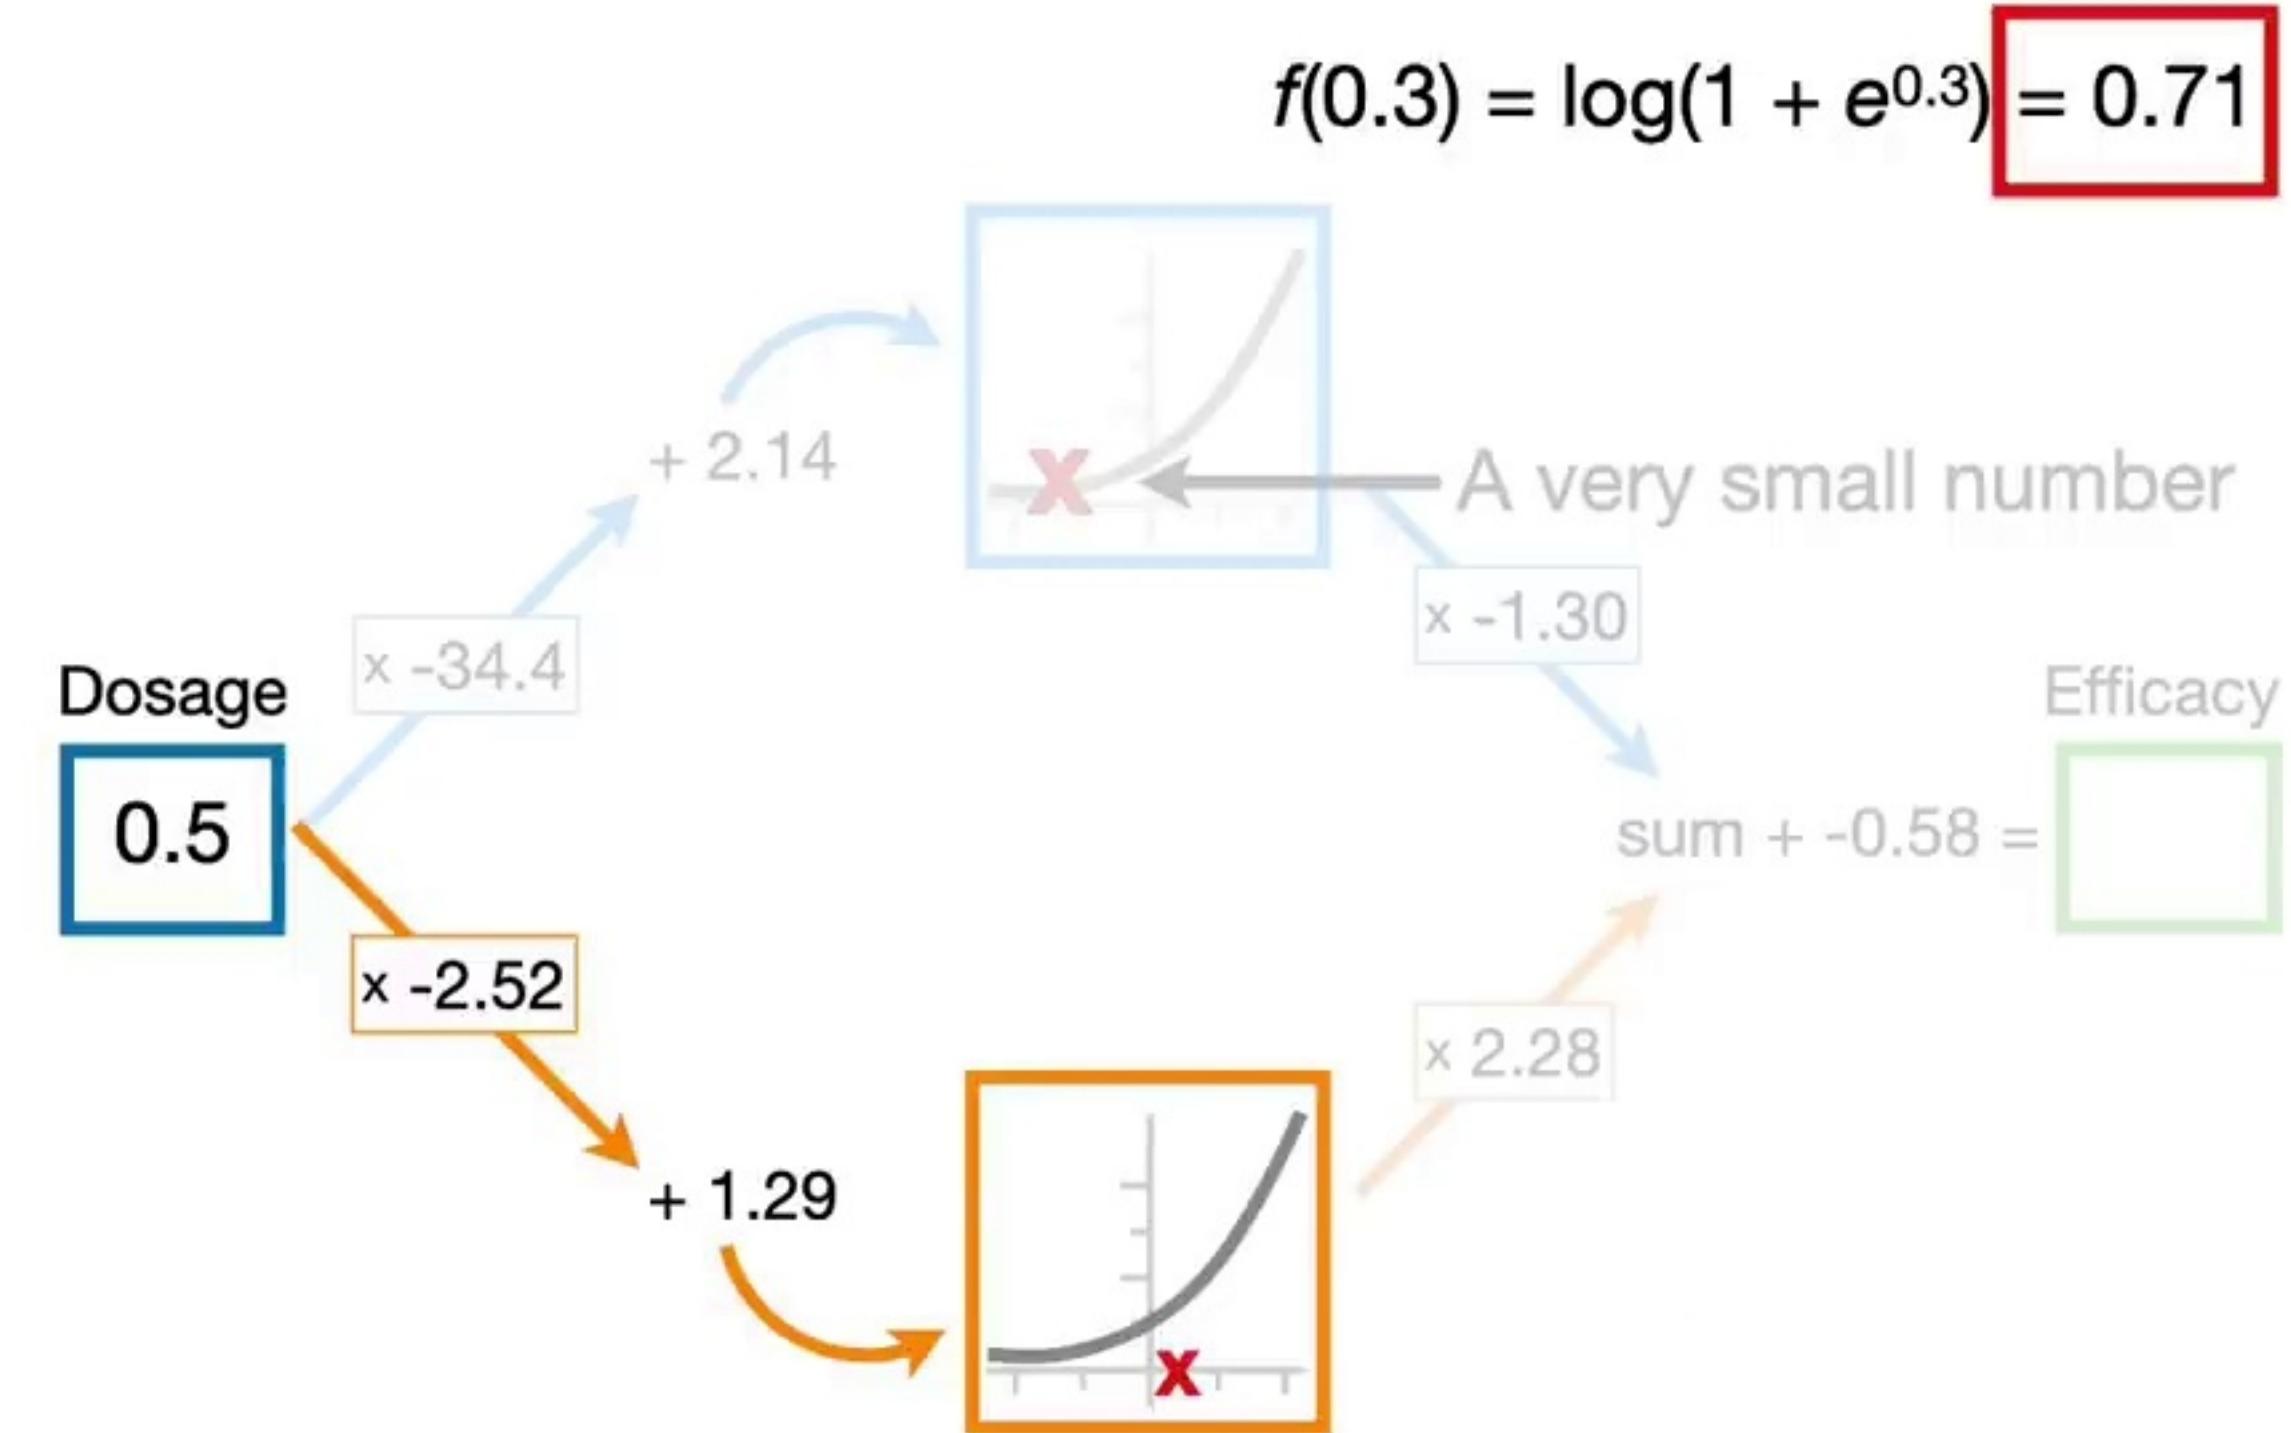
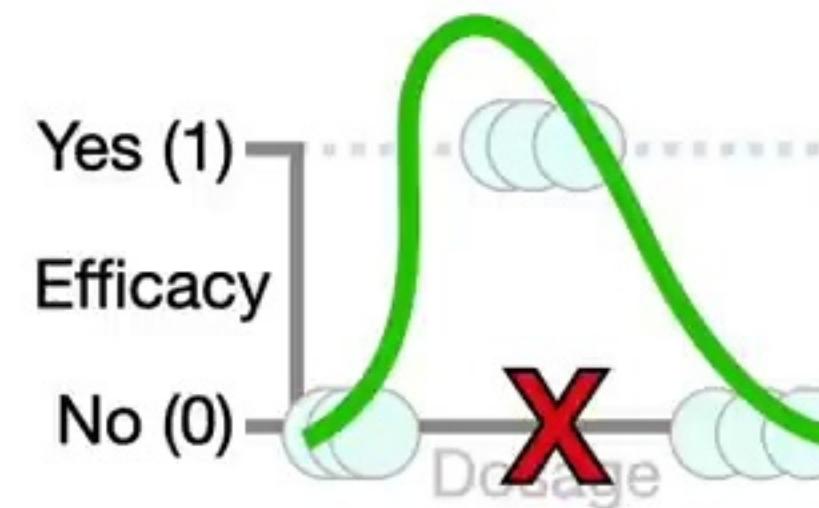
$$(0.5 \times -2.52) + 1.29 = 0.3$$

$$f(0.3) = \log(1 + e^{0.3})$$

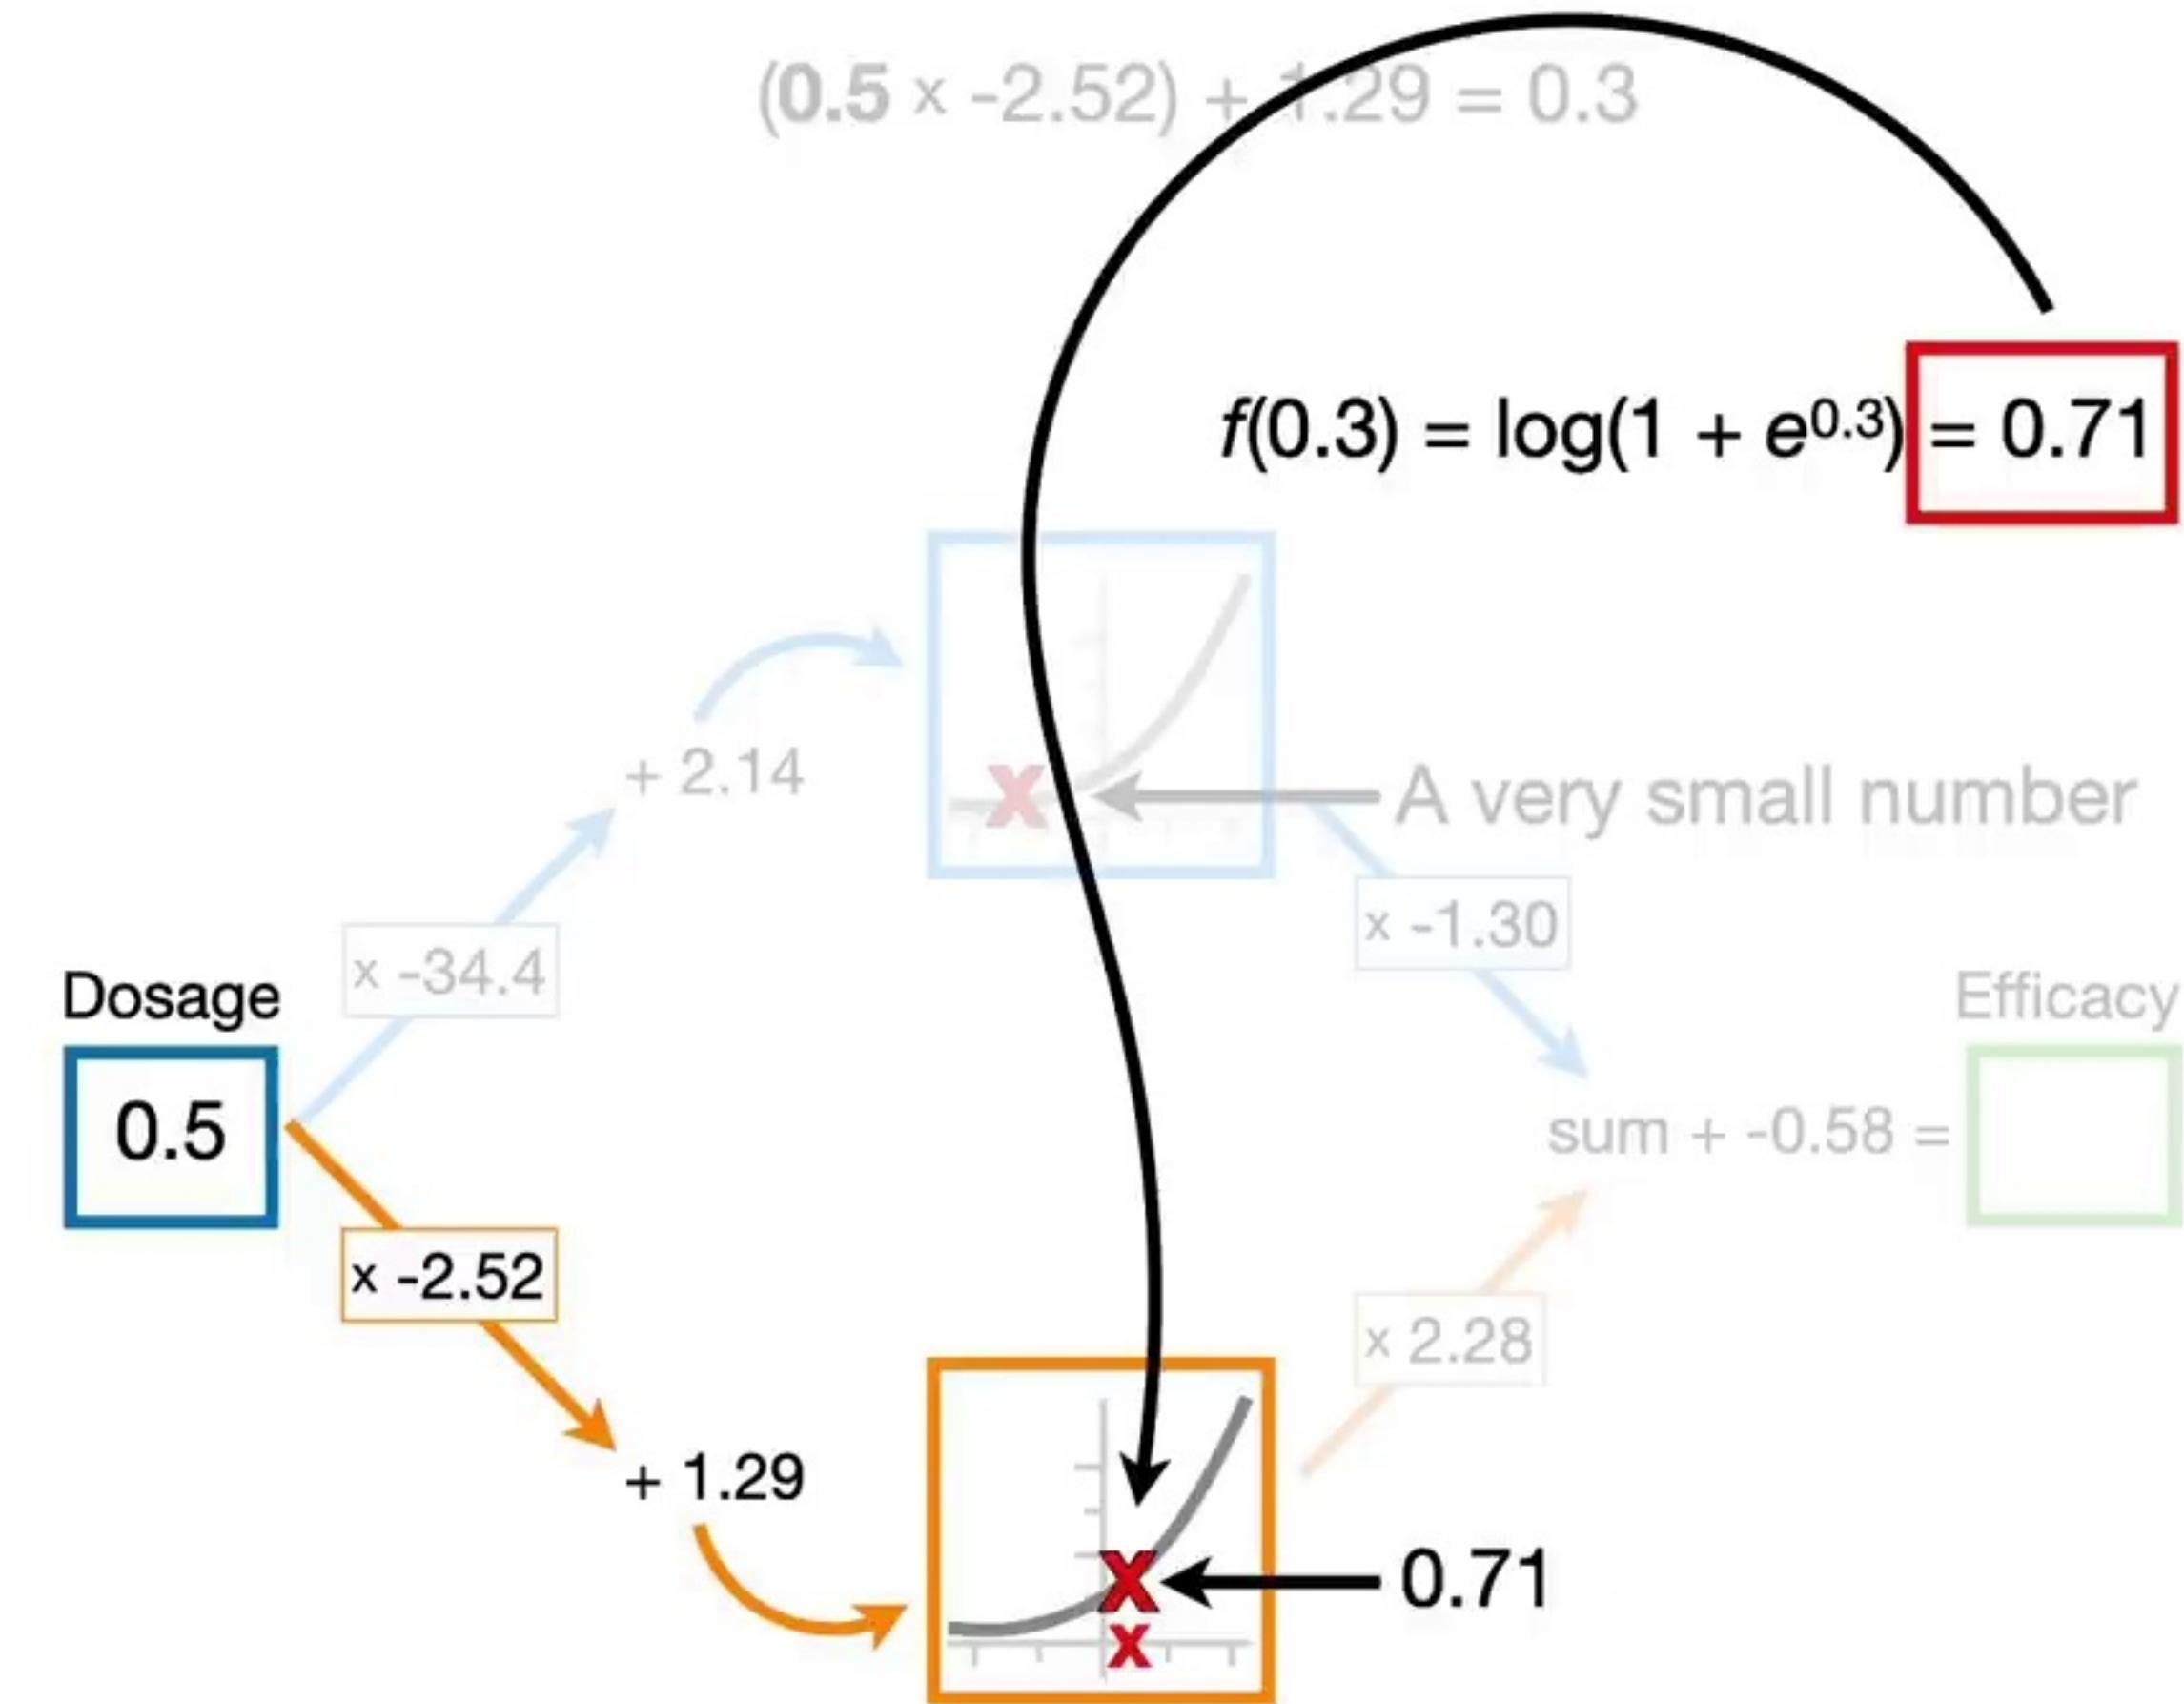
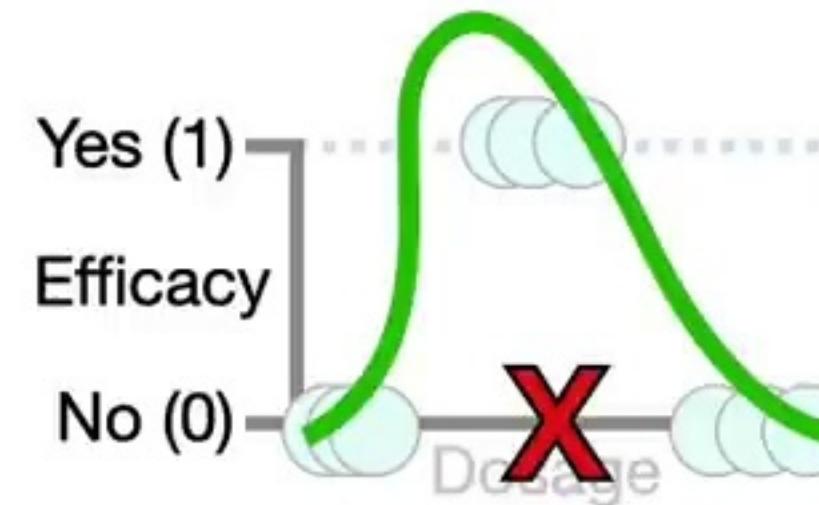


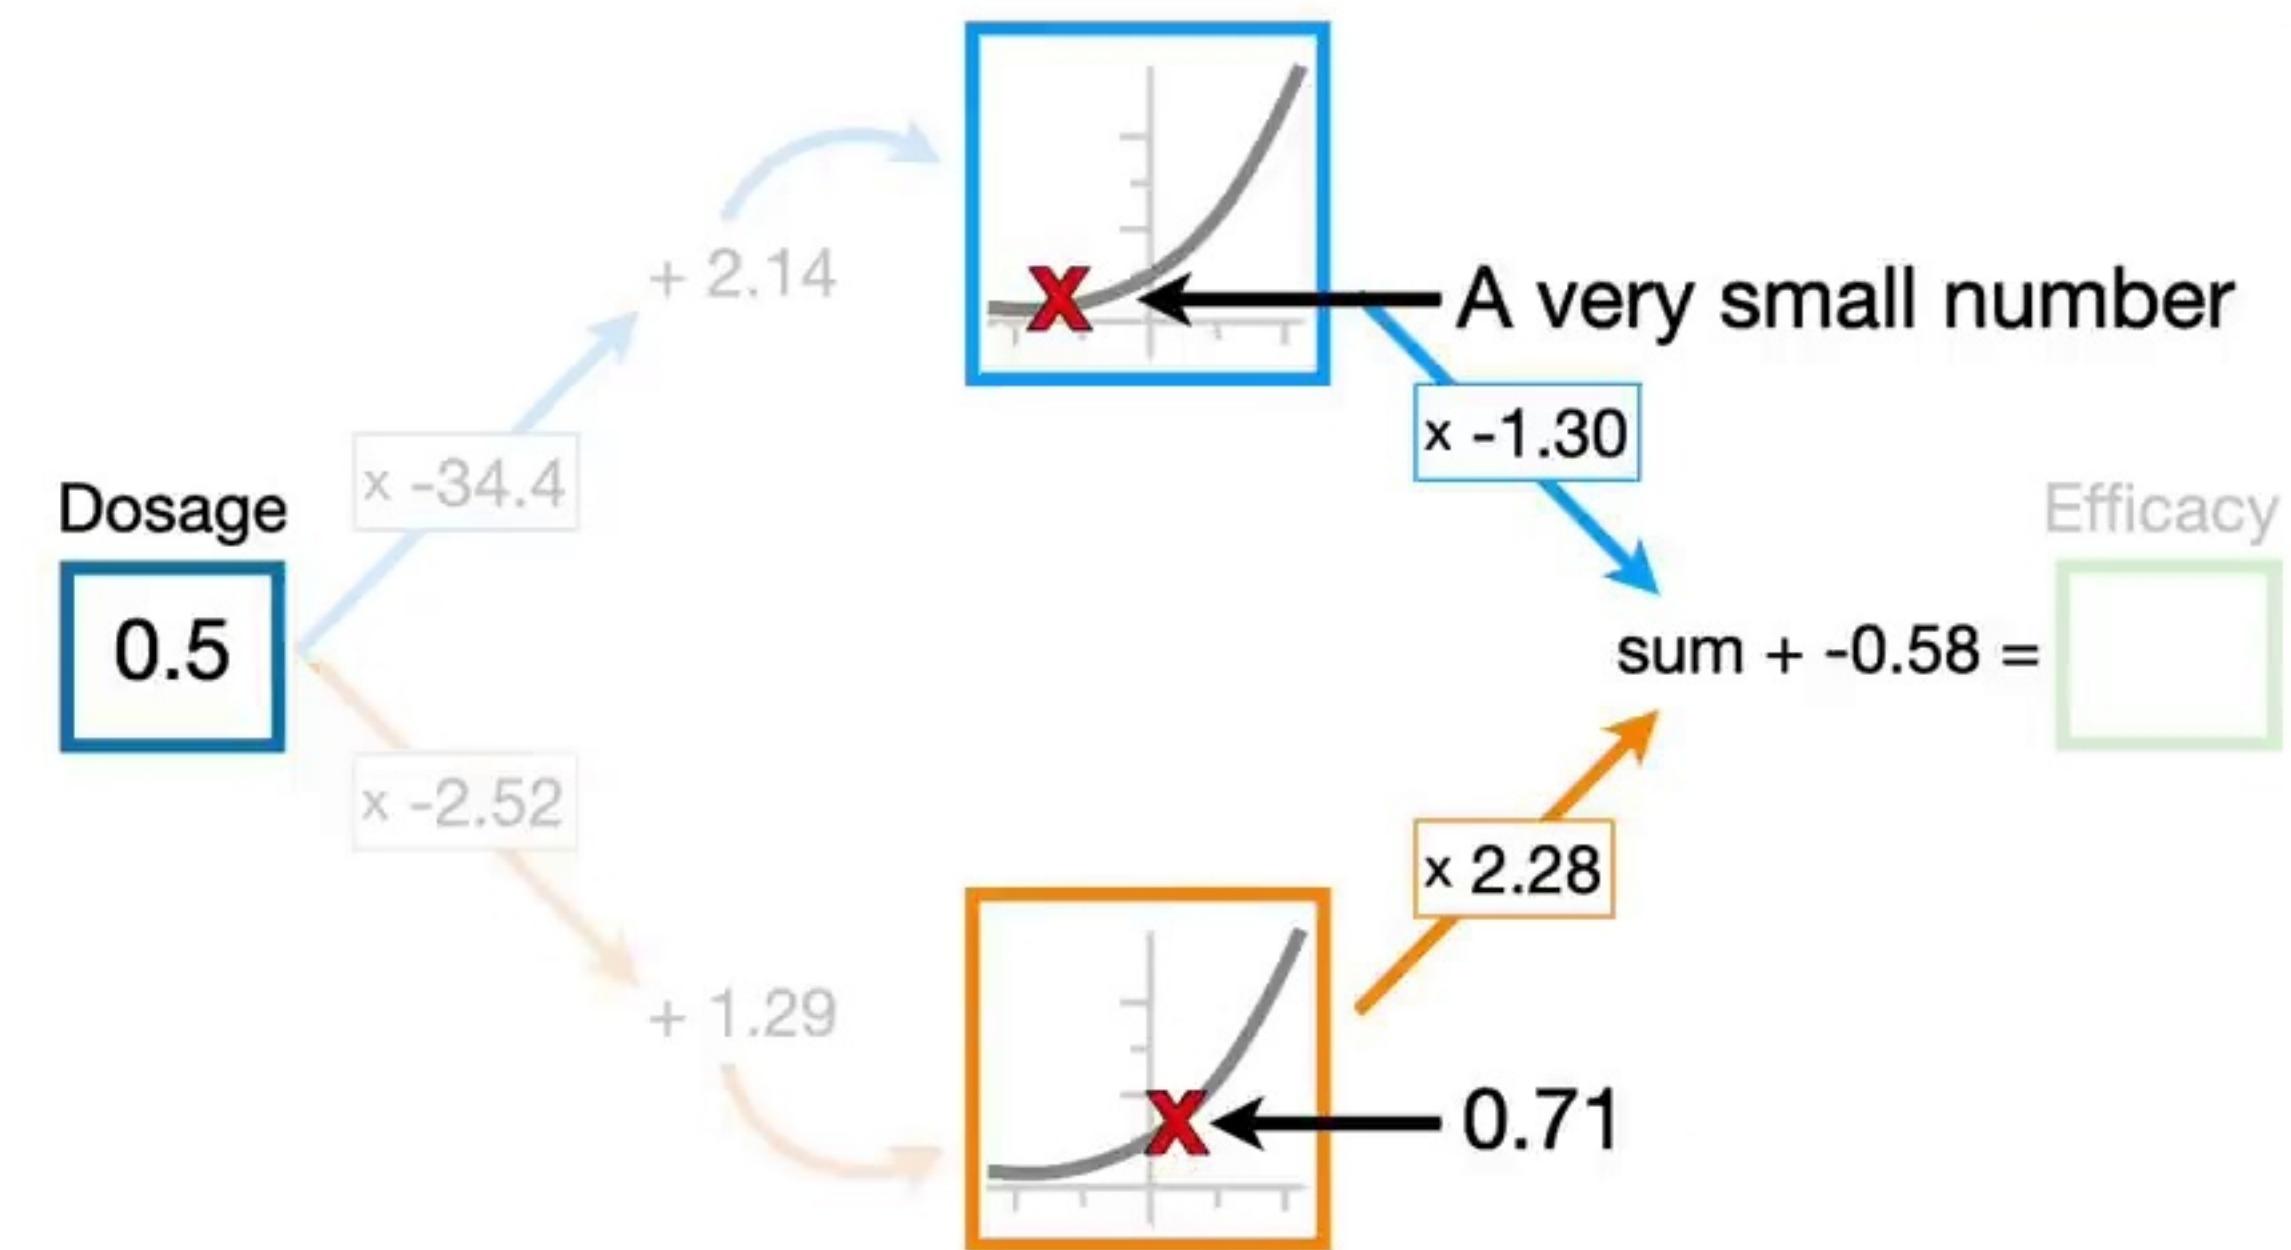
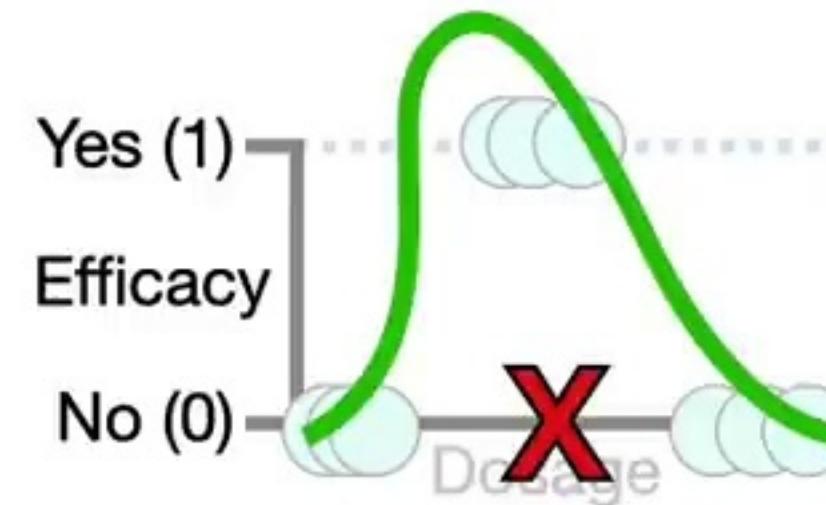
double  
BAM!!  
**SQ!**

$$(0.5 \times -2.52) + 1.29 = 0.3$$



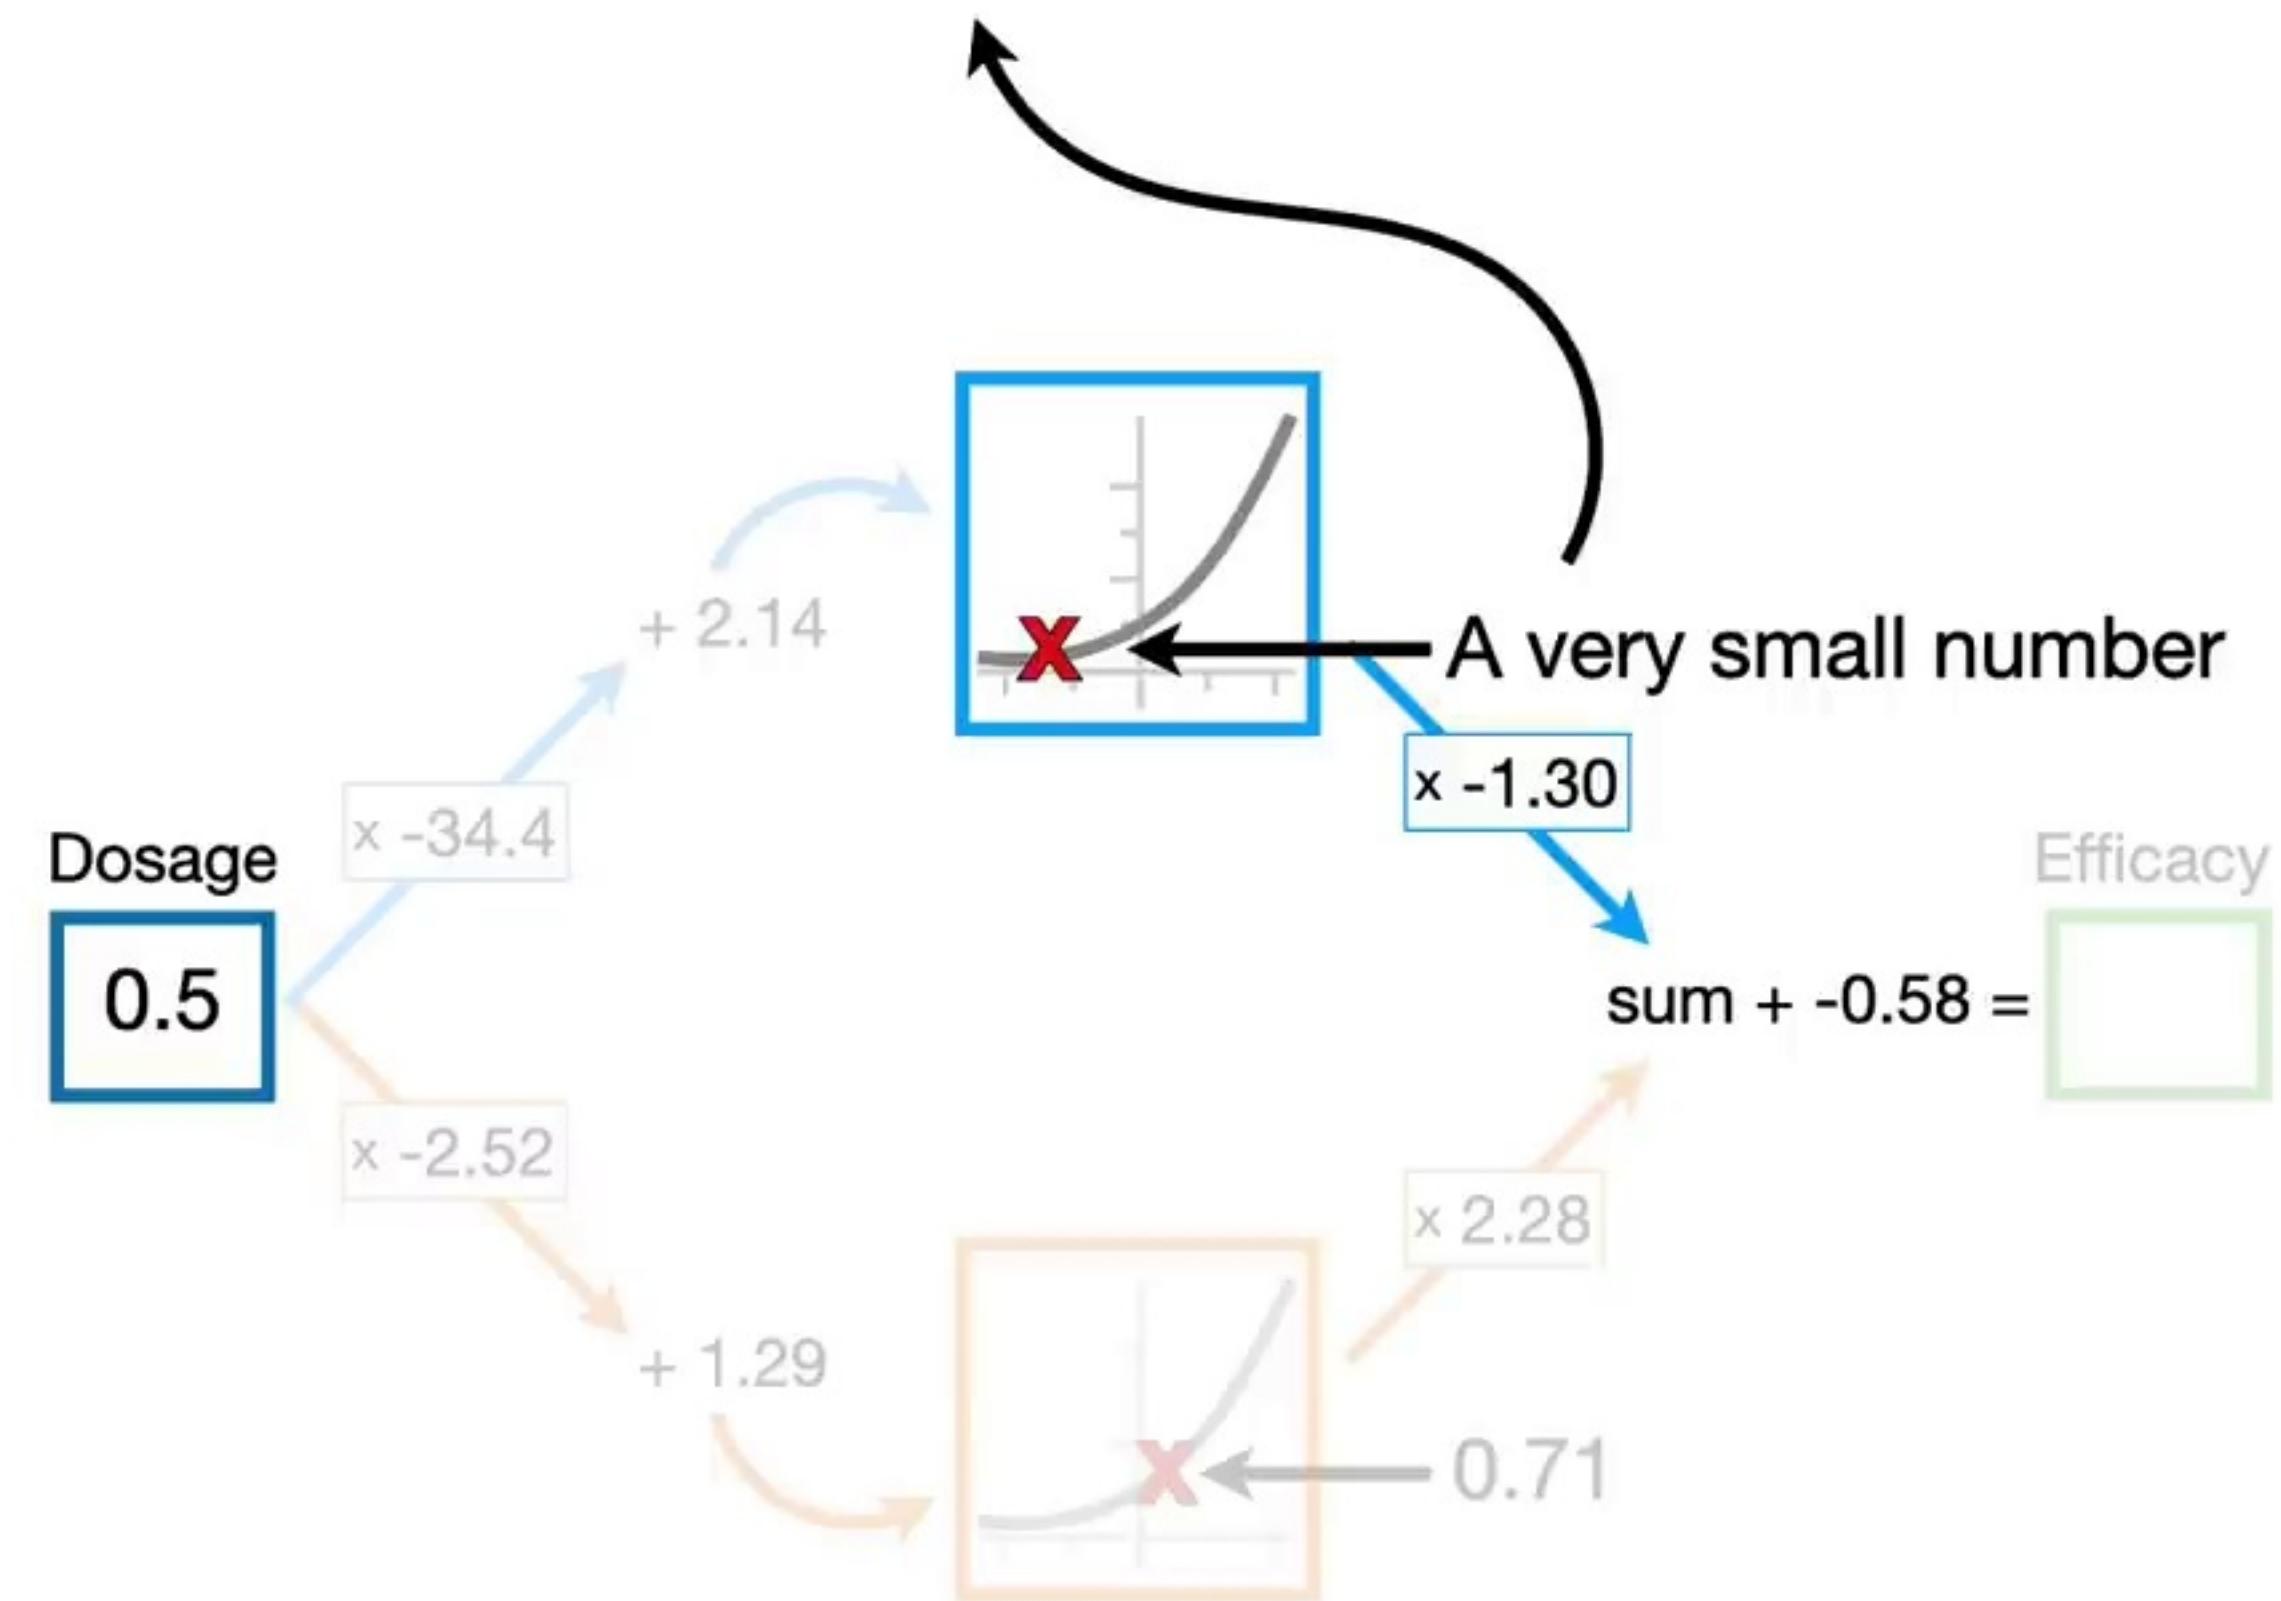
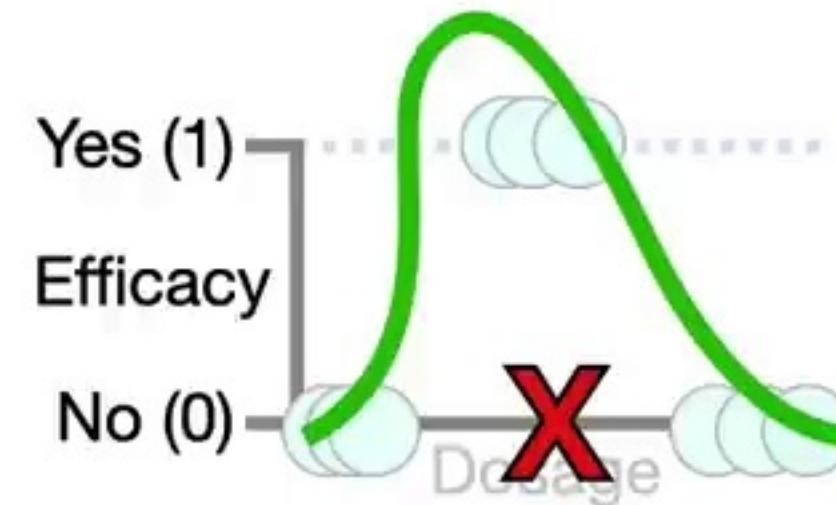
double  
BAM!!  
**SQ!**





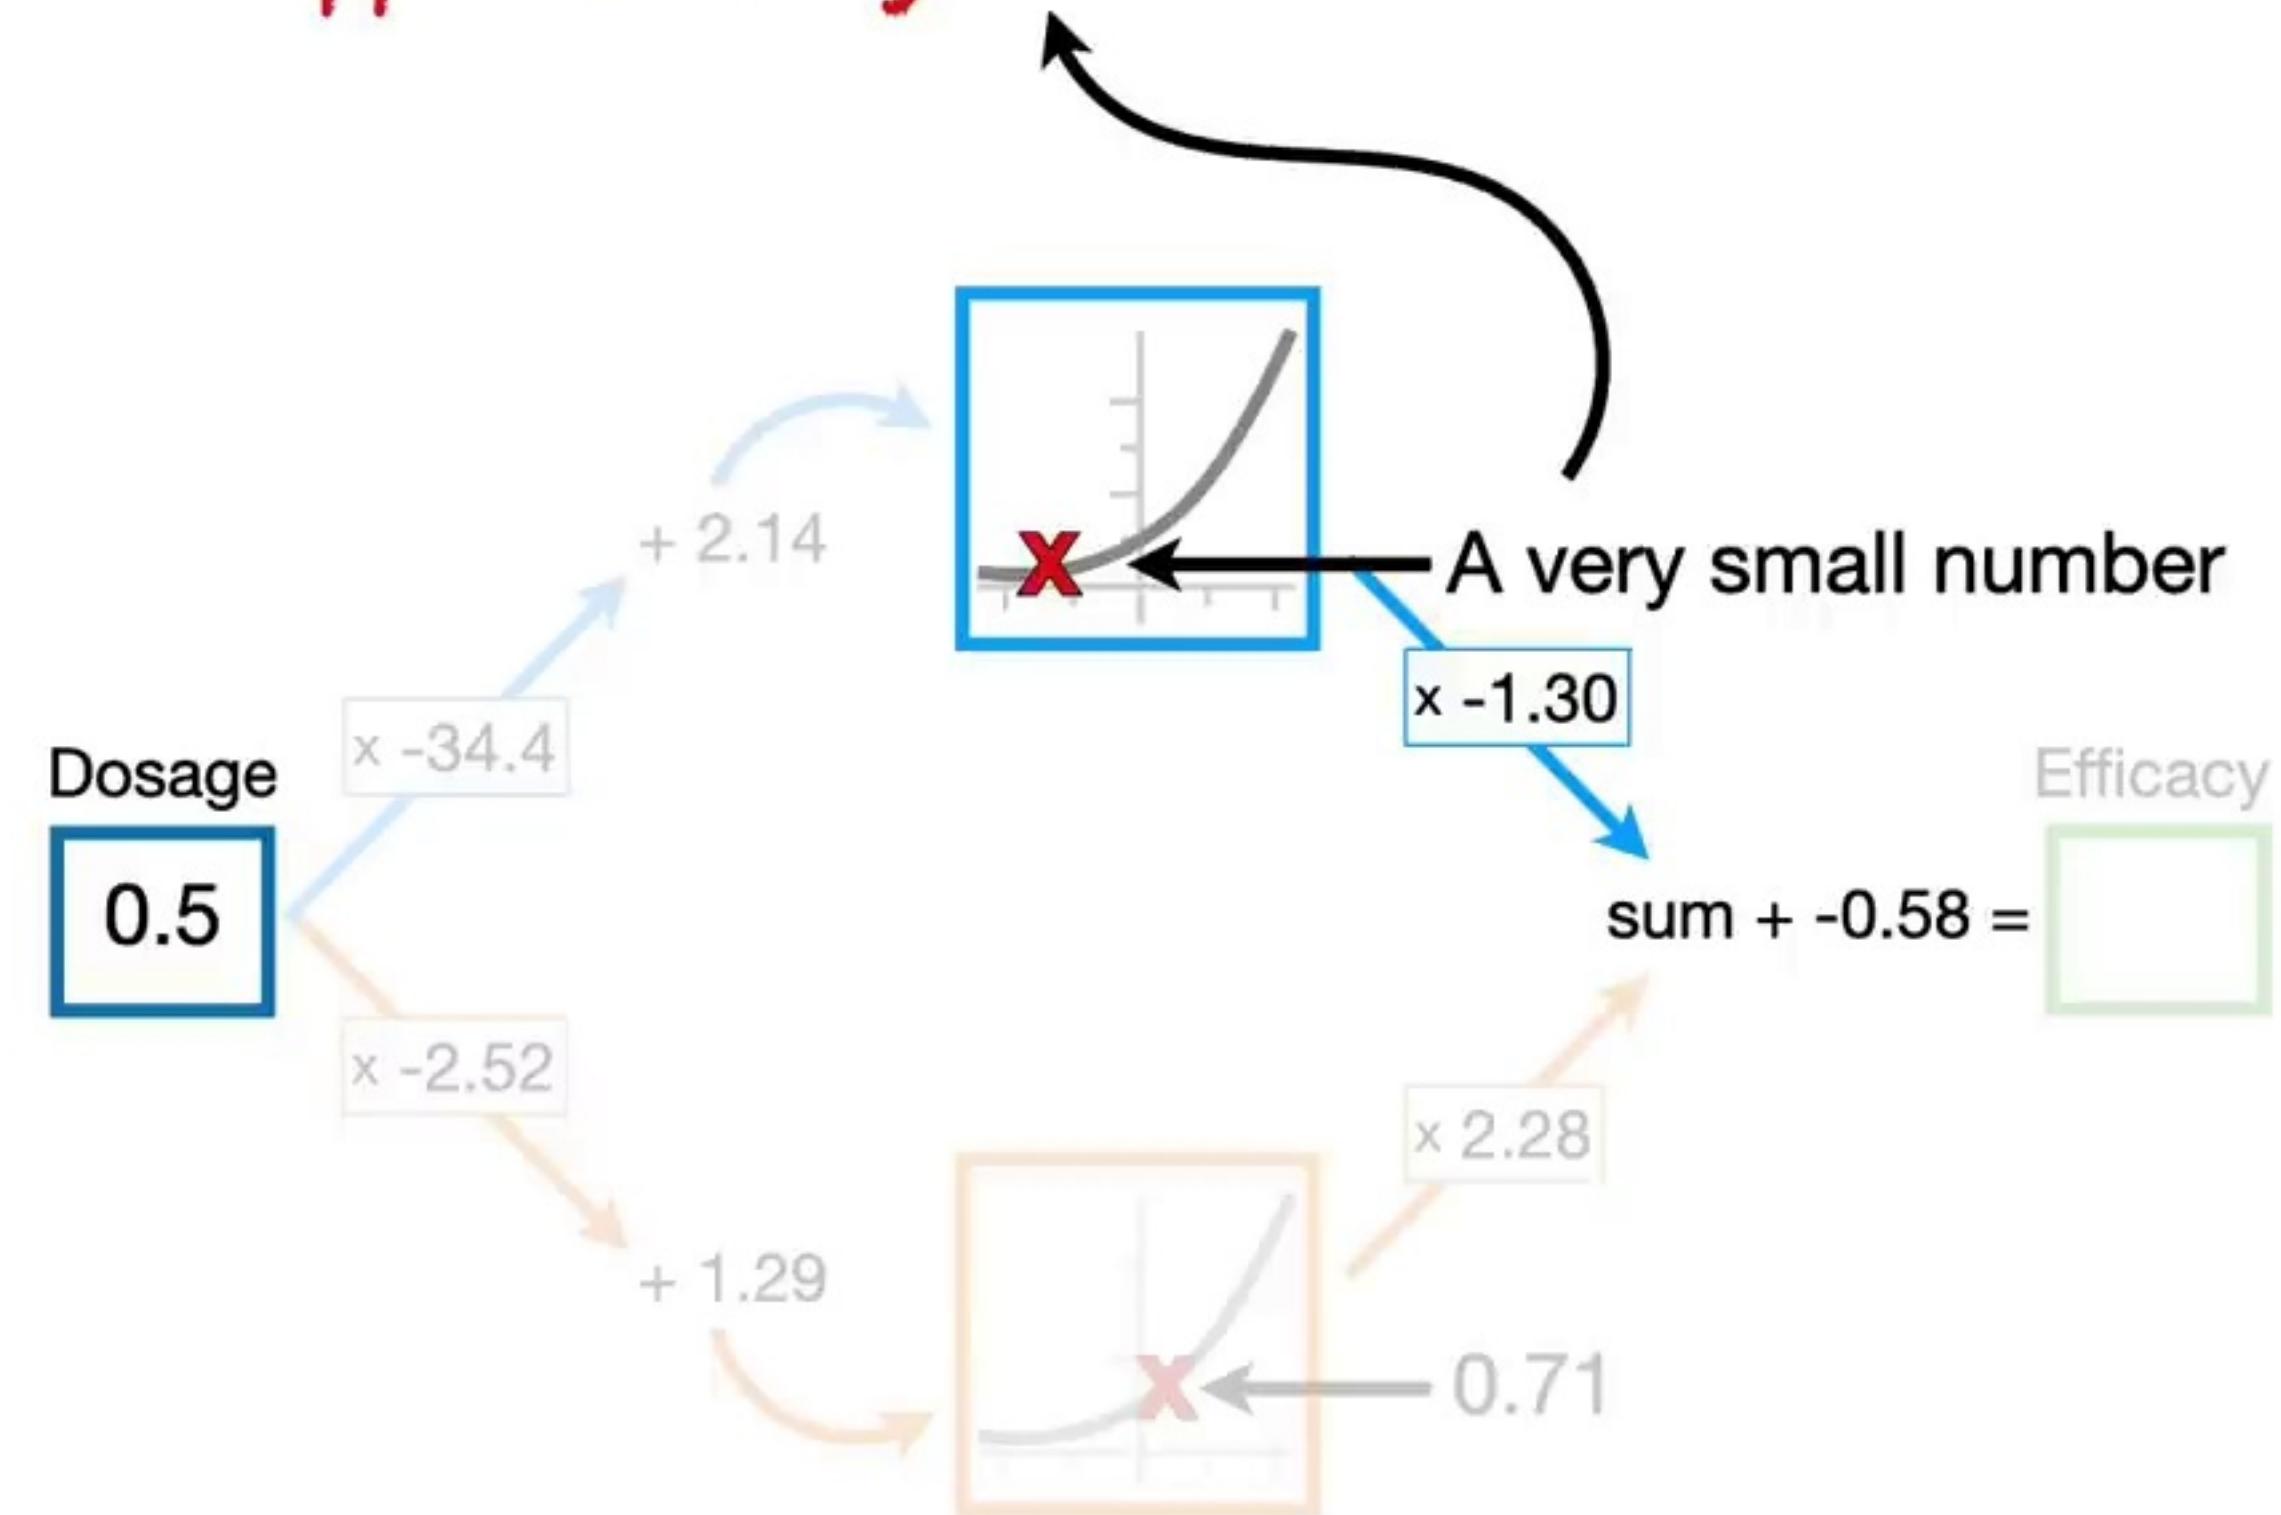
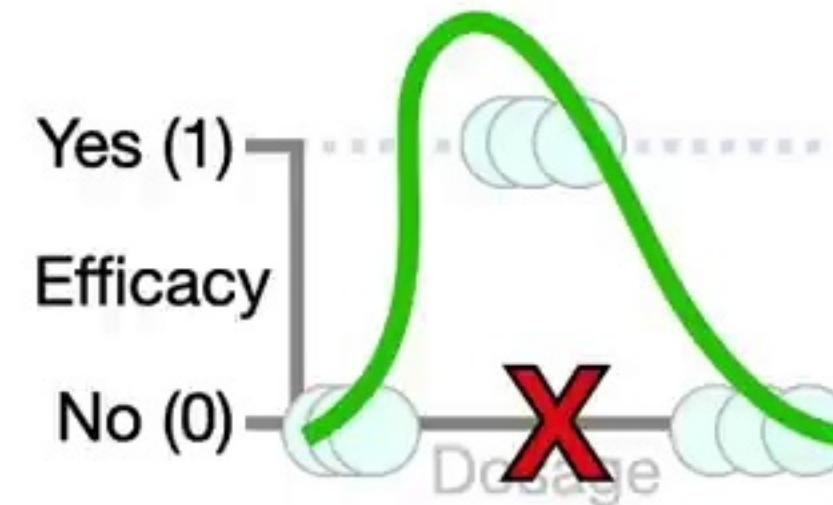


(A really small number  $\times$  -1.30)



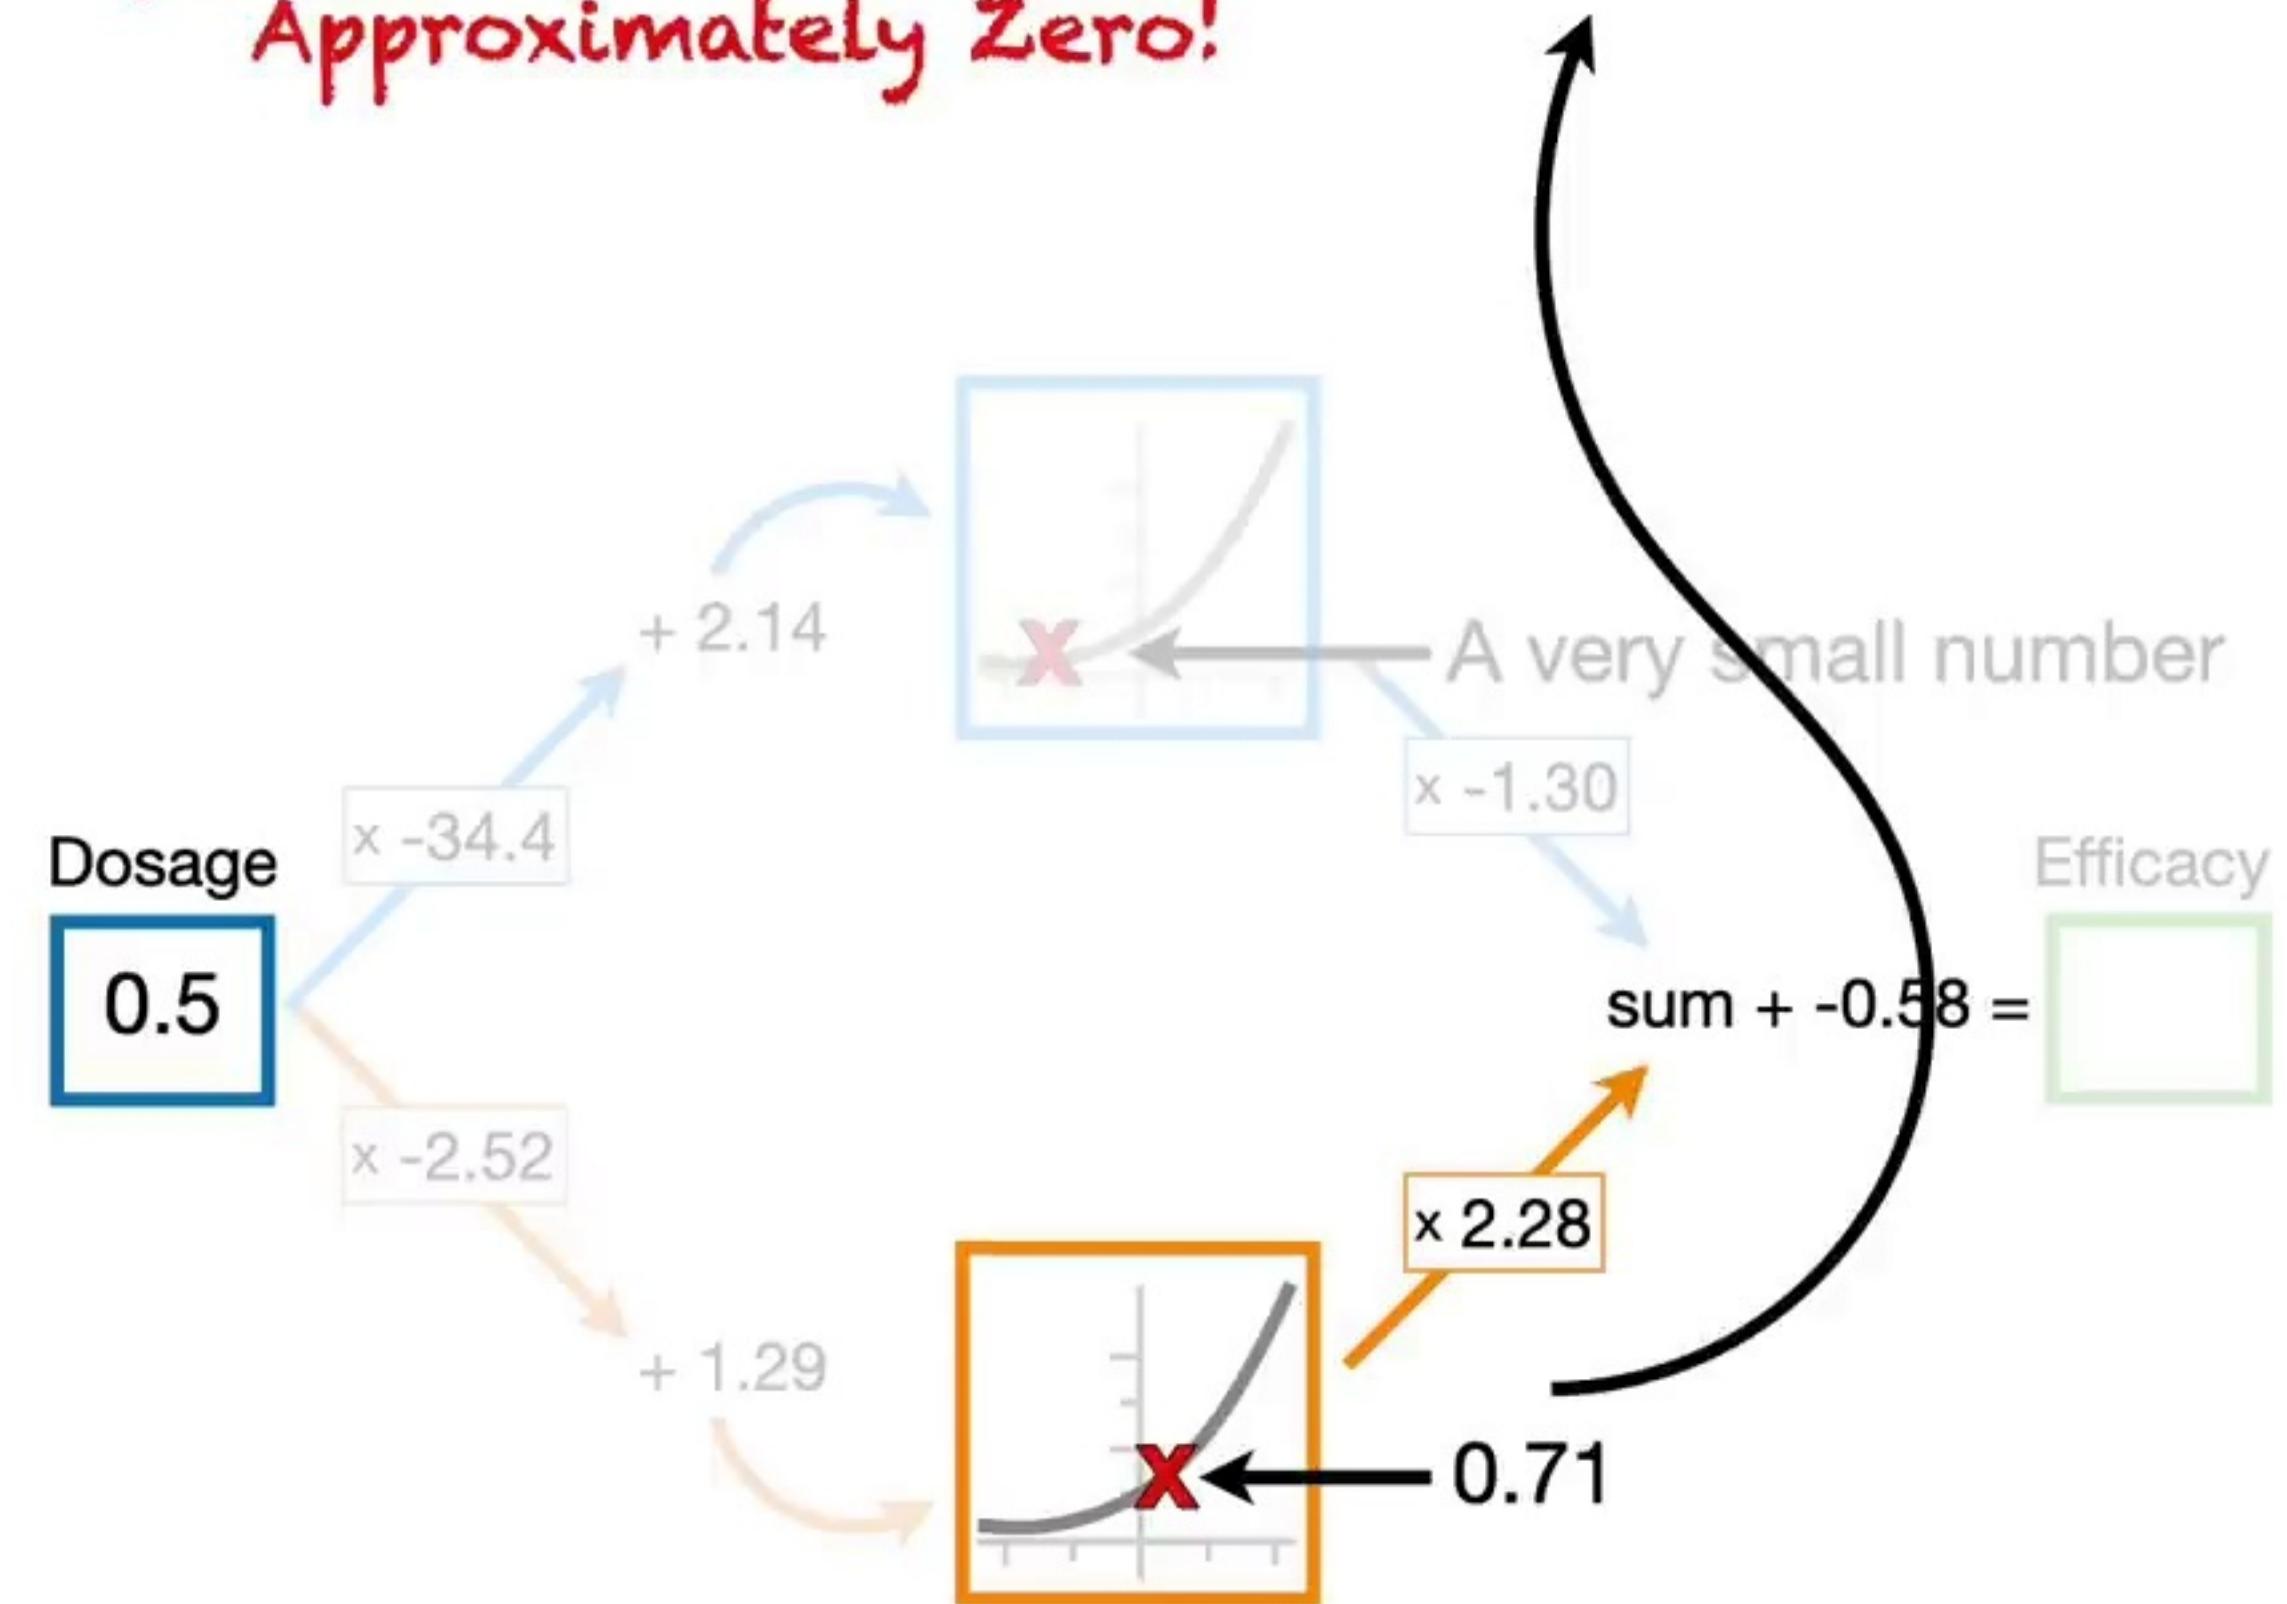
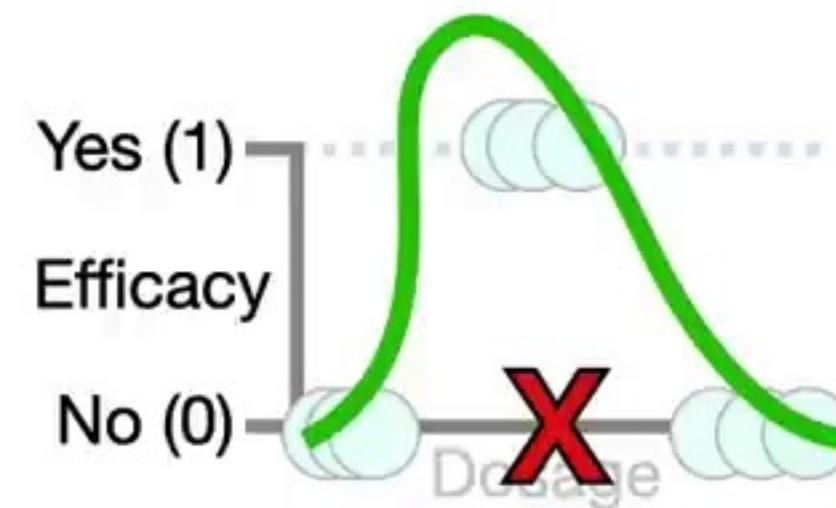


(A really small number  $x - 1.30$ )  
Approximately Zero!



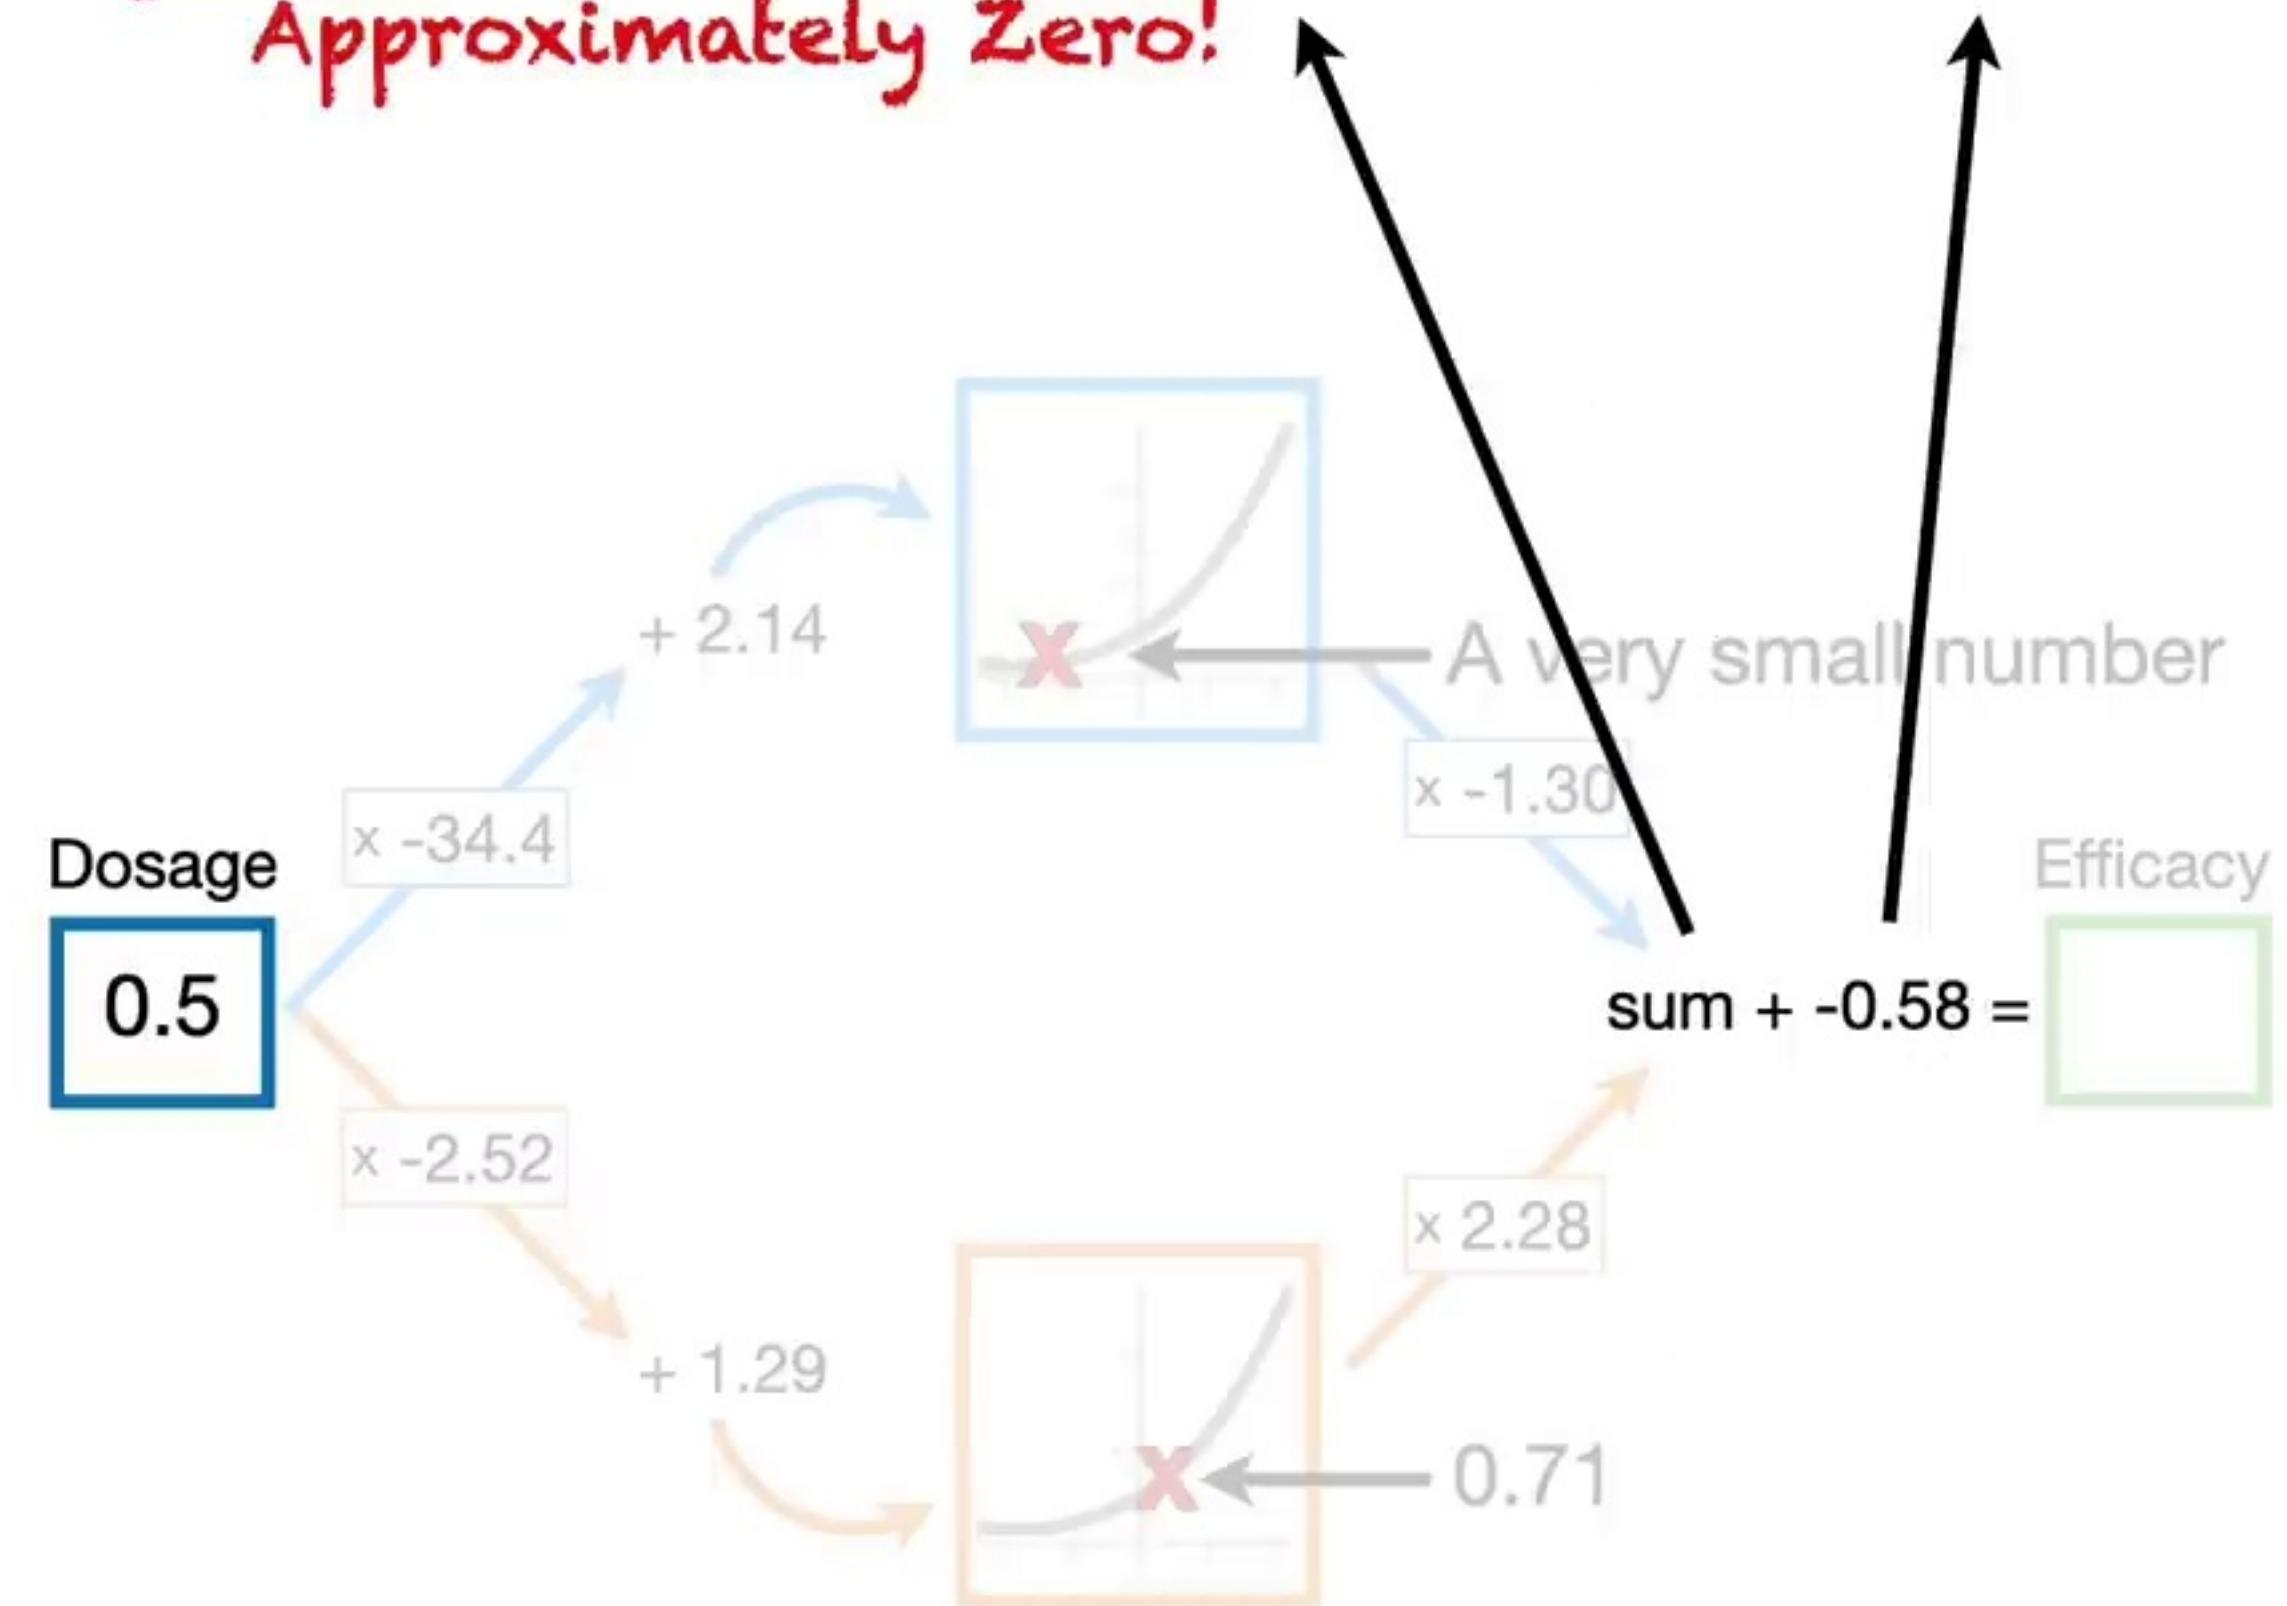
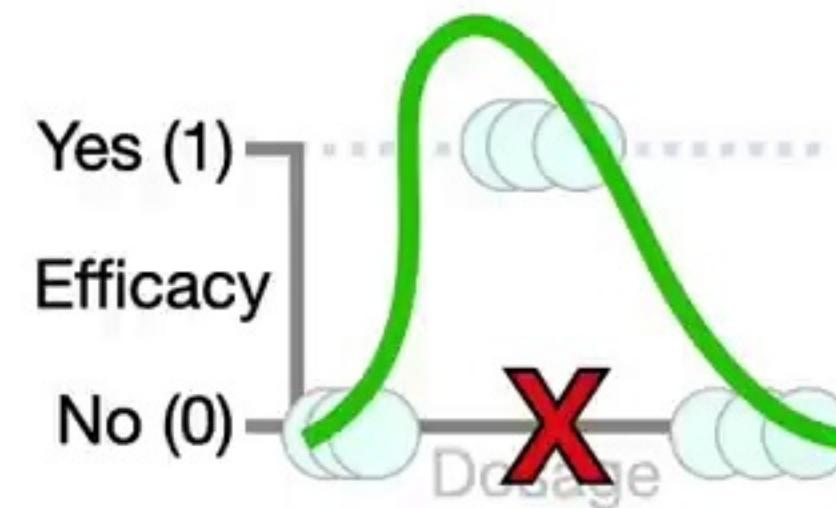


(A really small number  $\times -1.30$ )  $(0.71 \times 2.28)$   
Approximately Zero!



Double  
BAM!!  
**SQ!**

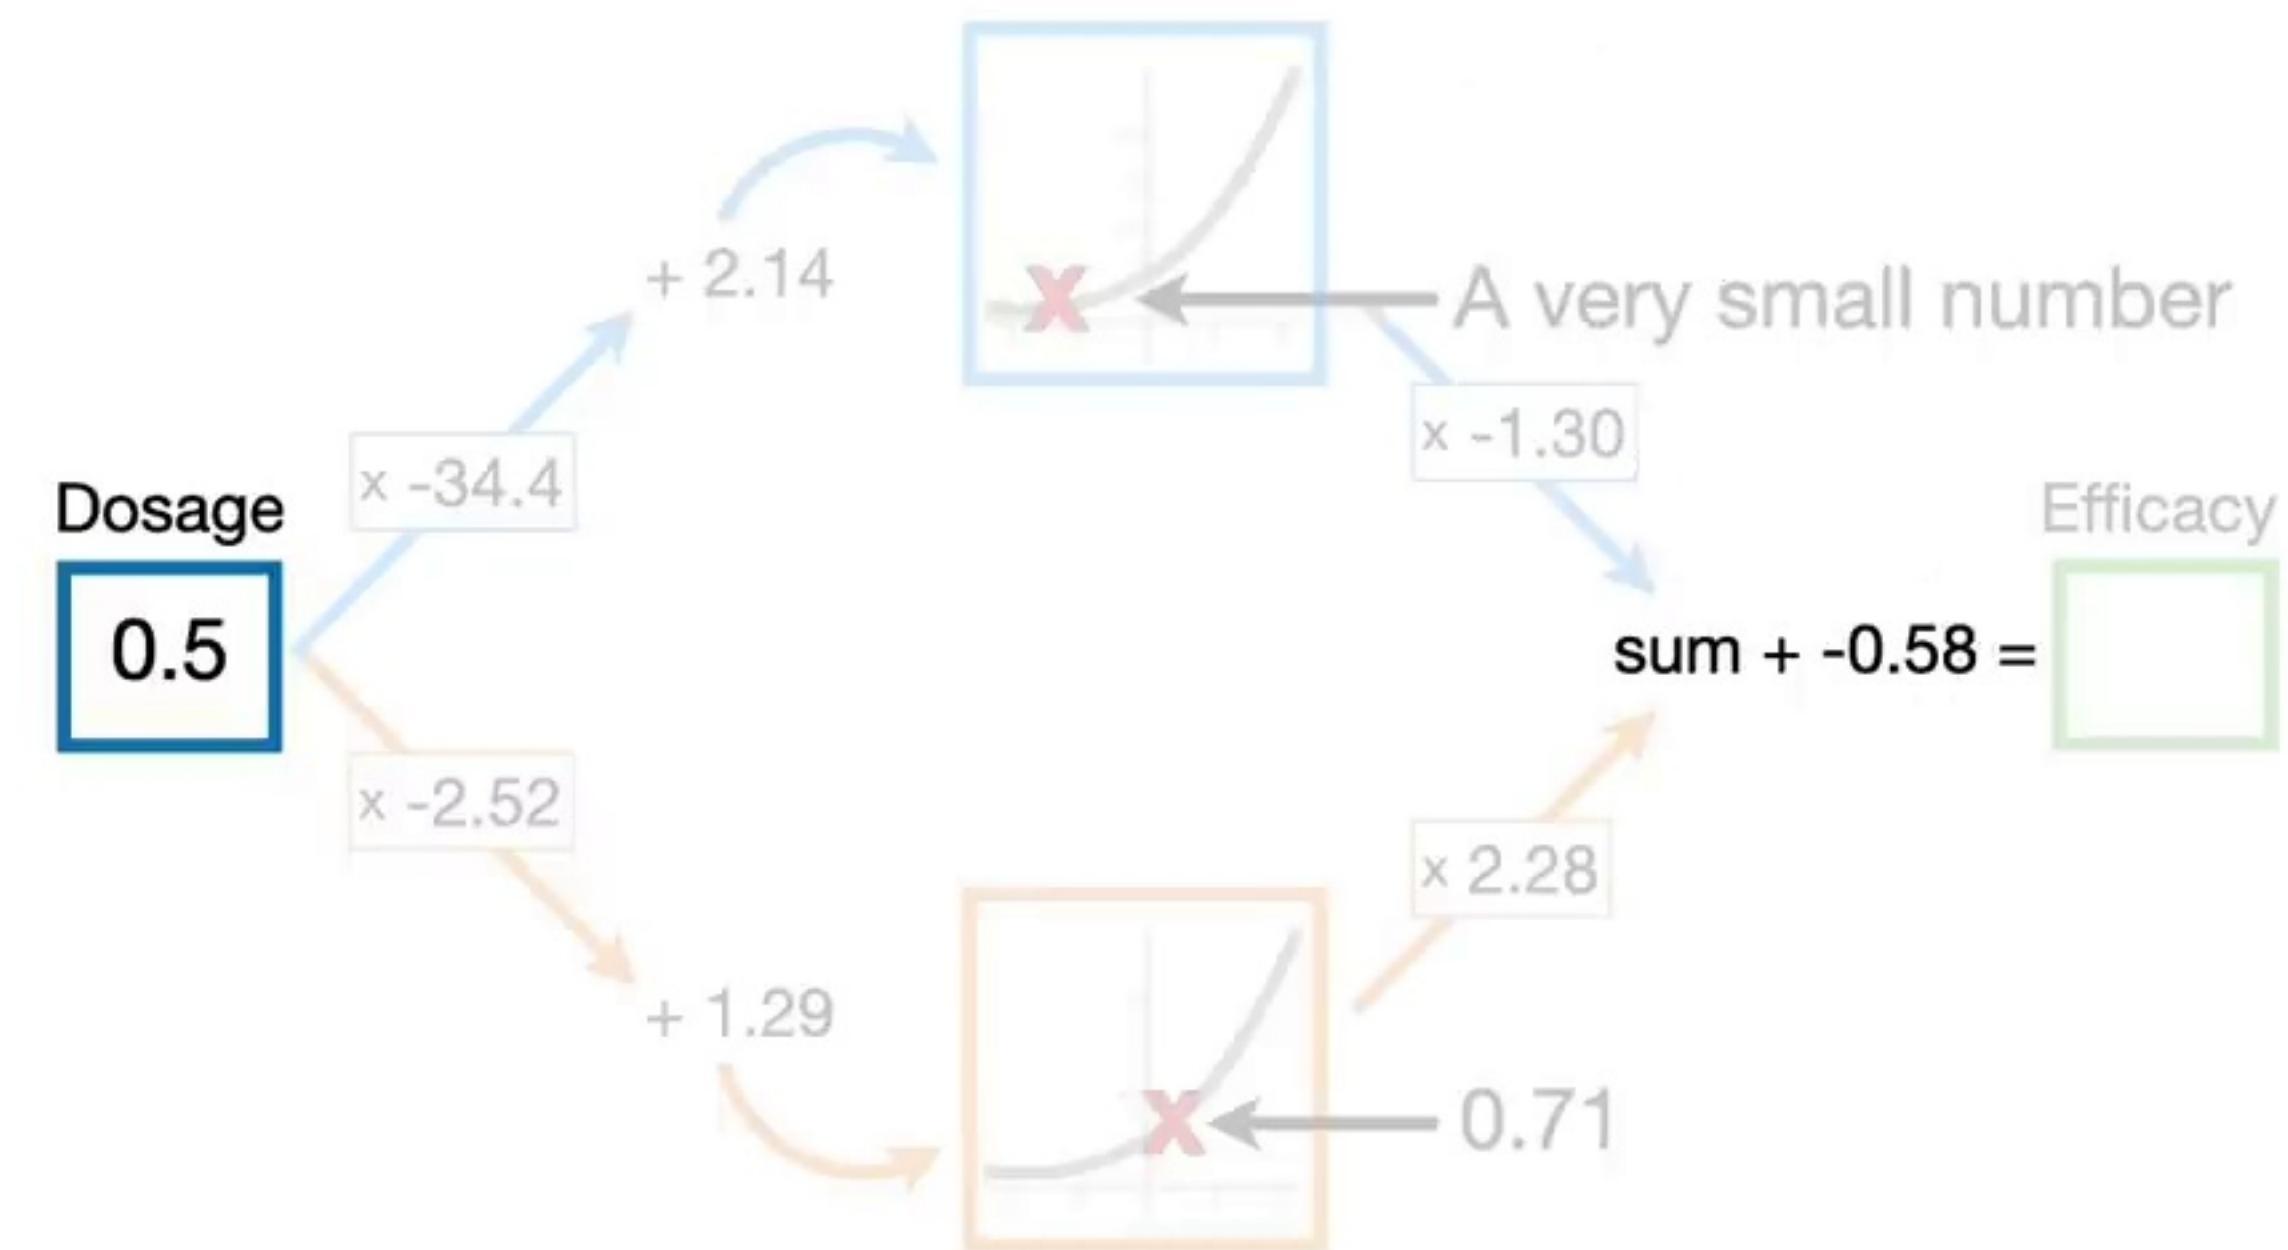
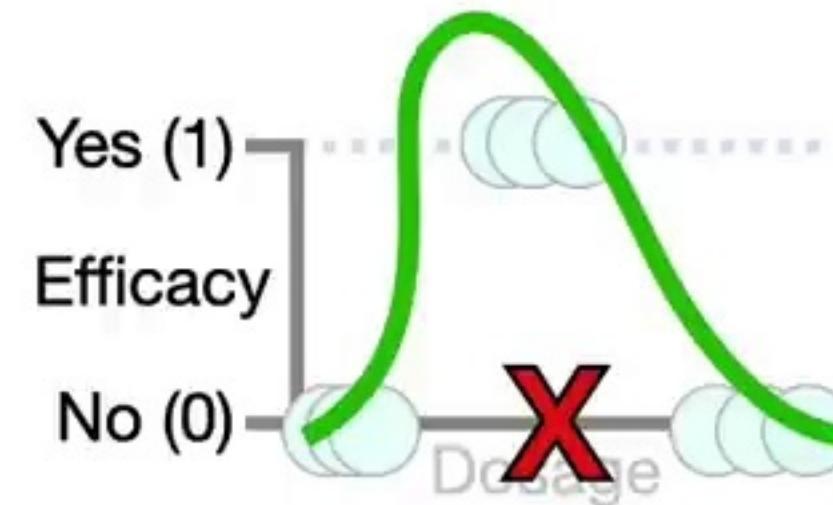
~~(A really small number  $x - 1.30$ ) + (0.71 × 2.28) - 0.58~~  
**Approximately Zero!**



Double  
BAM!!  
**SQ!**

~~(A really small number  $x - 1.30$ ) + (0.71 \times 2.28) - 0.58 = 1.03~~

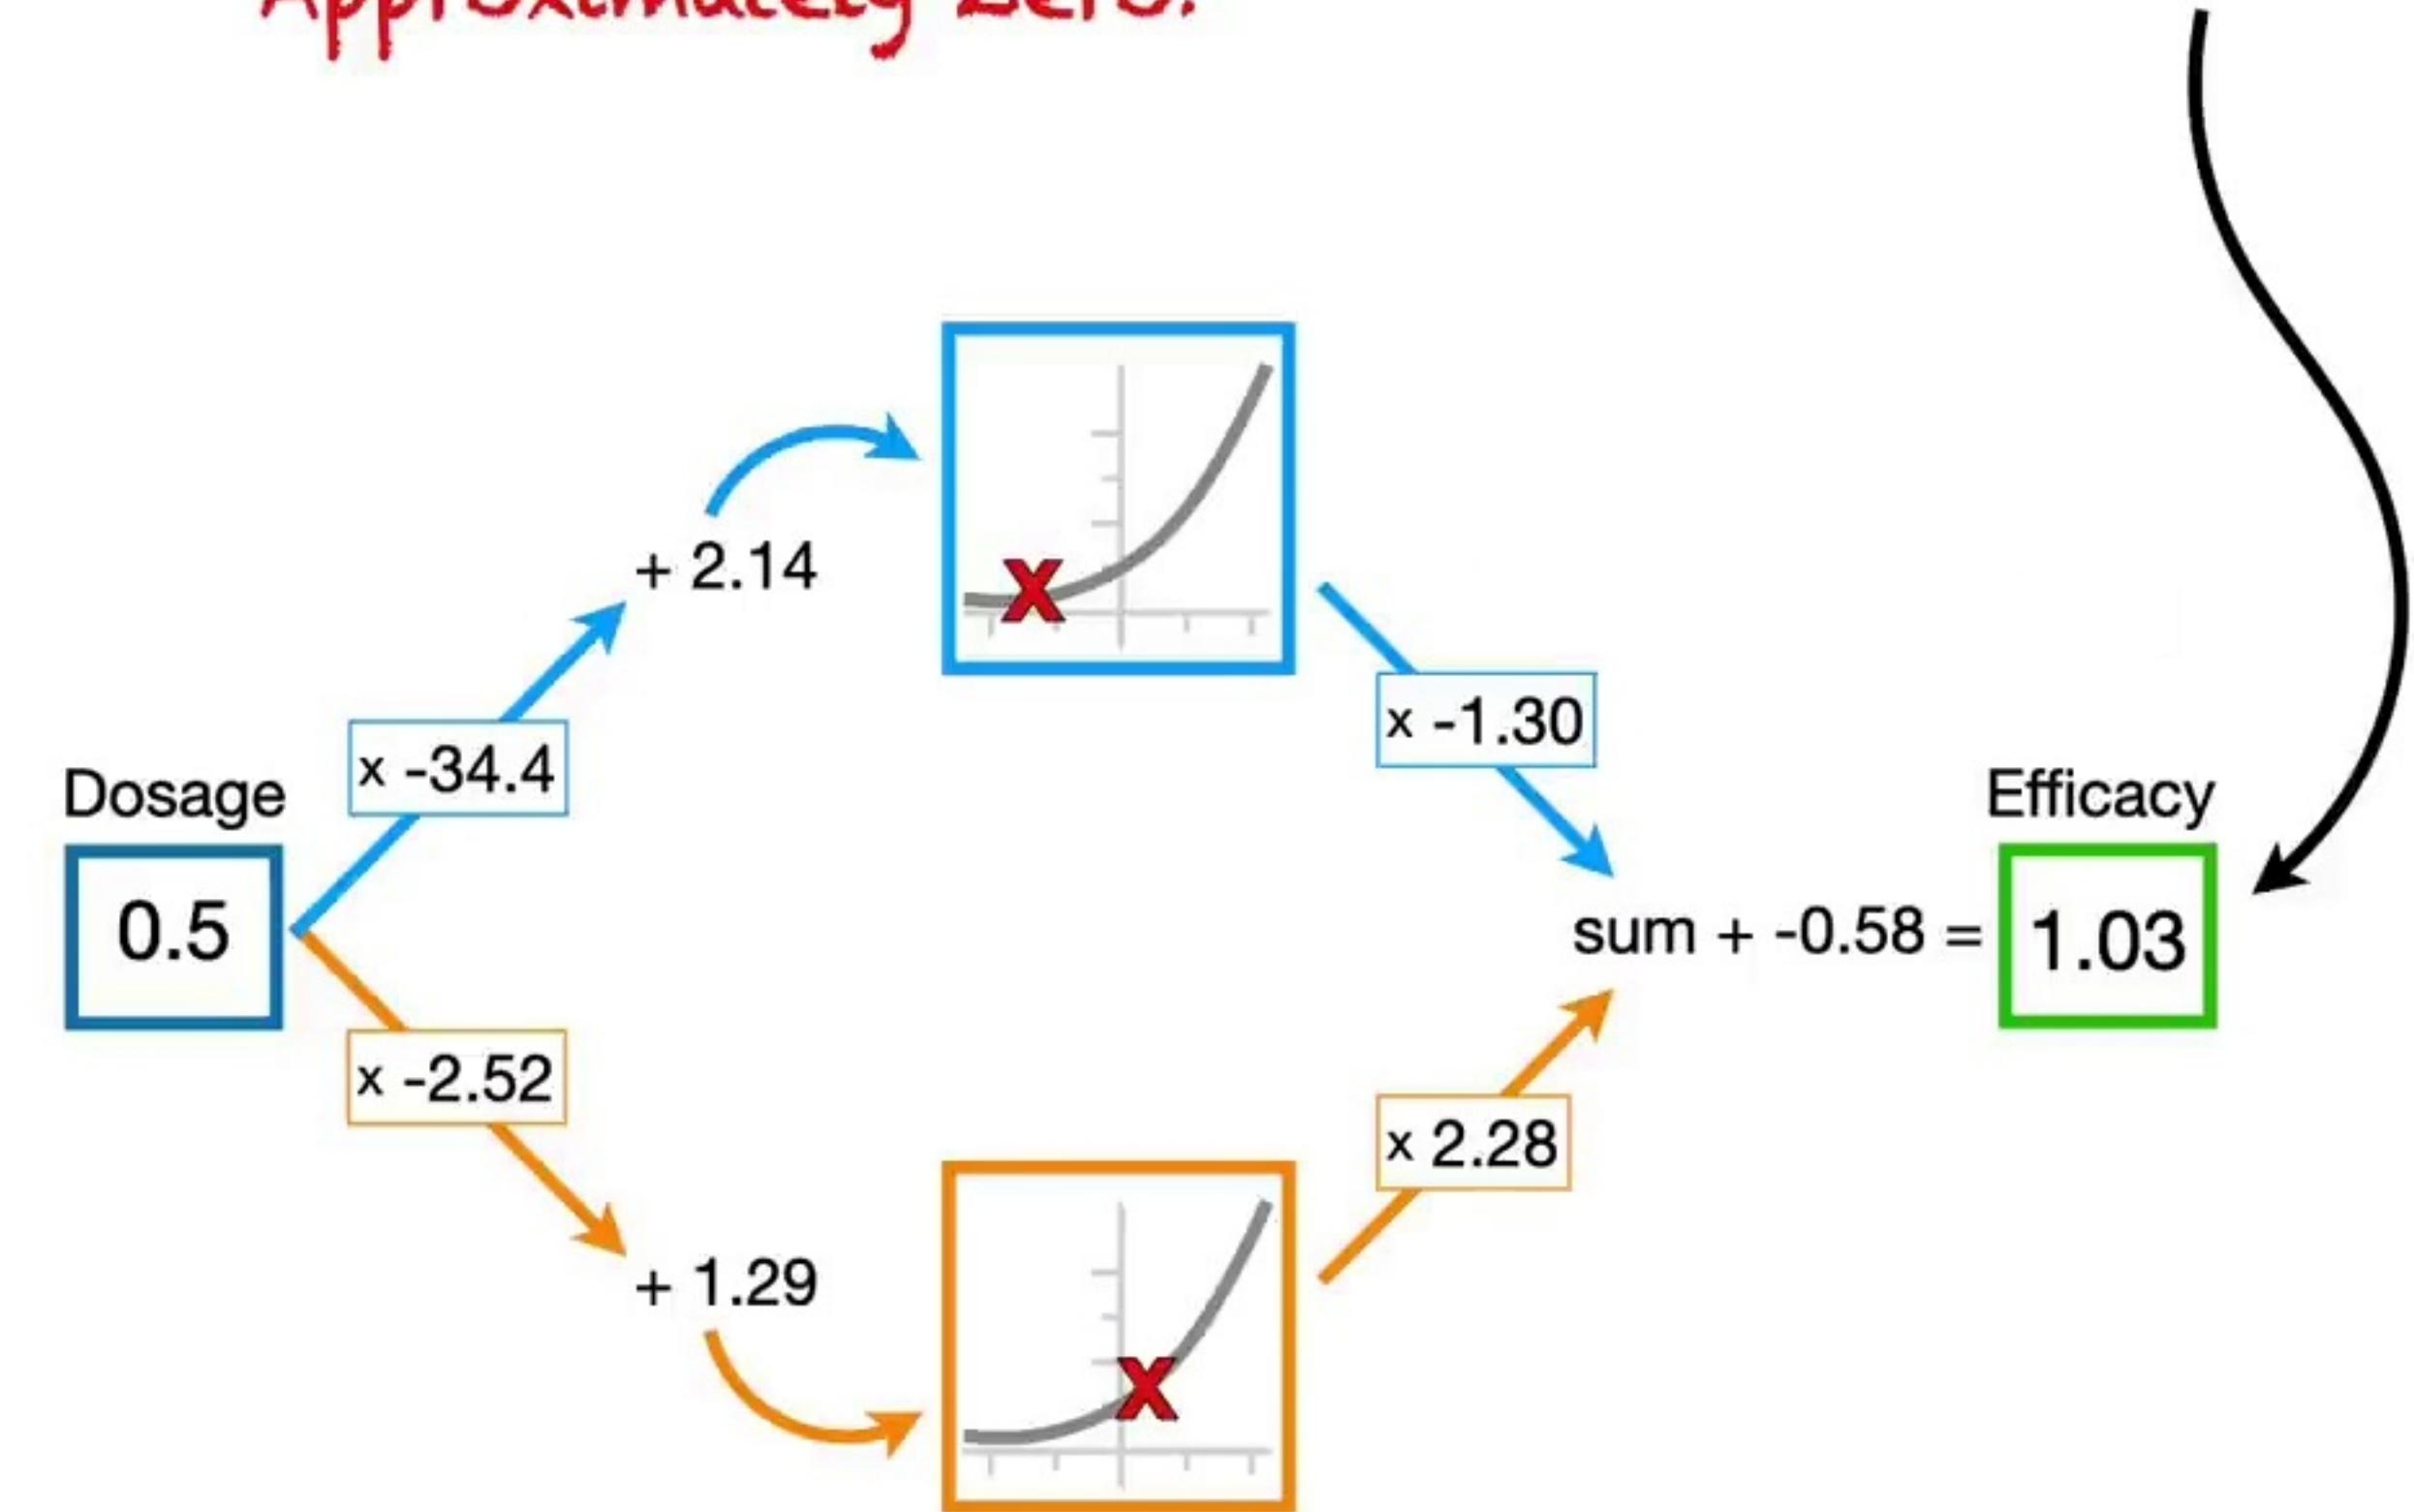
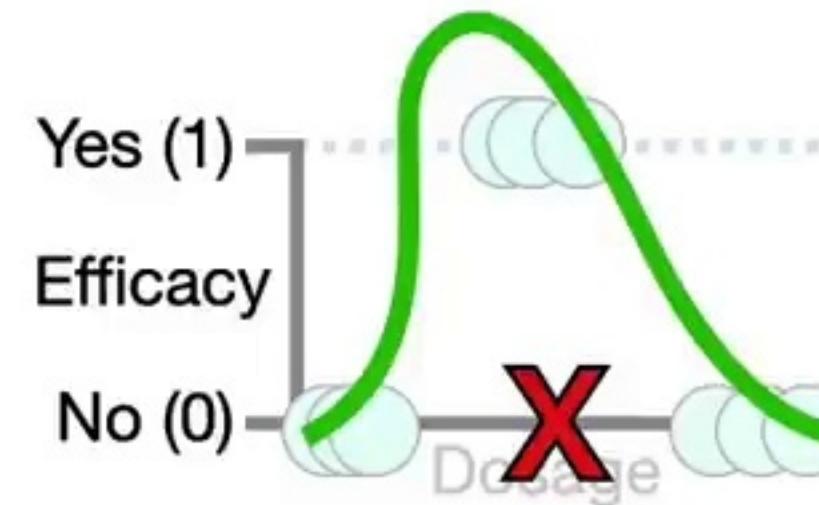
**Approximately Zero!**



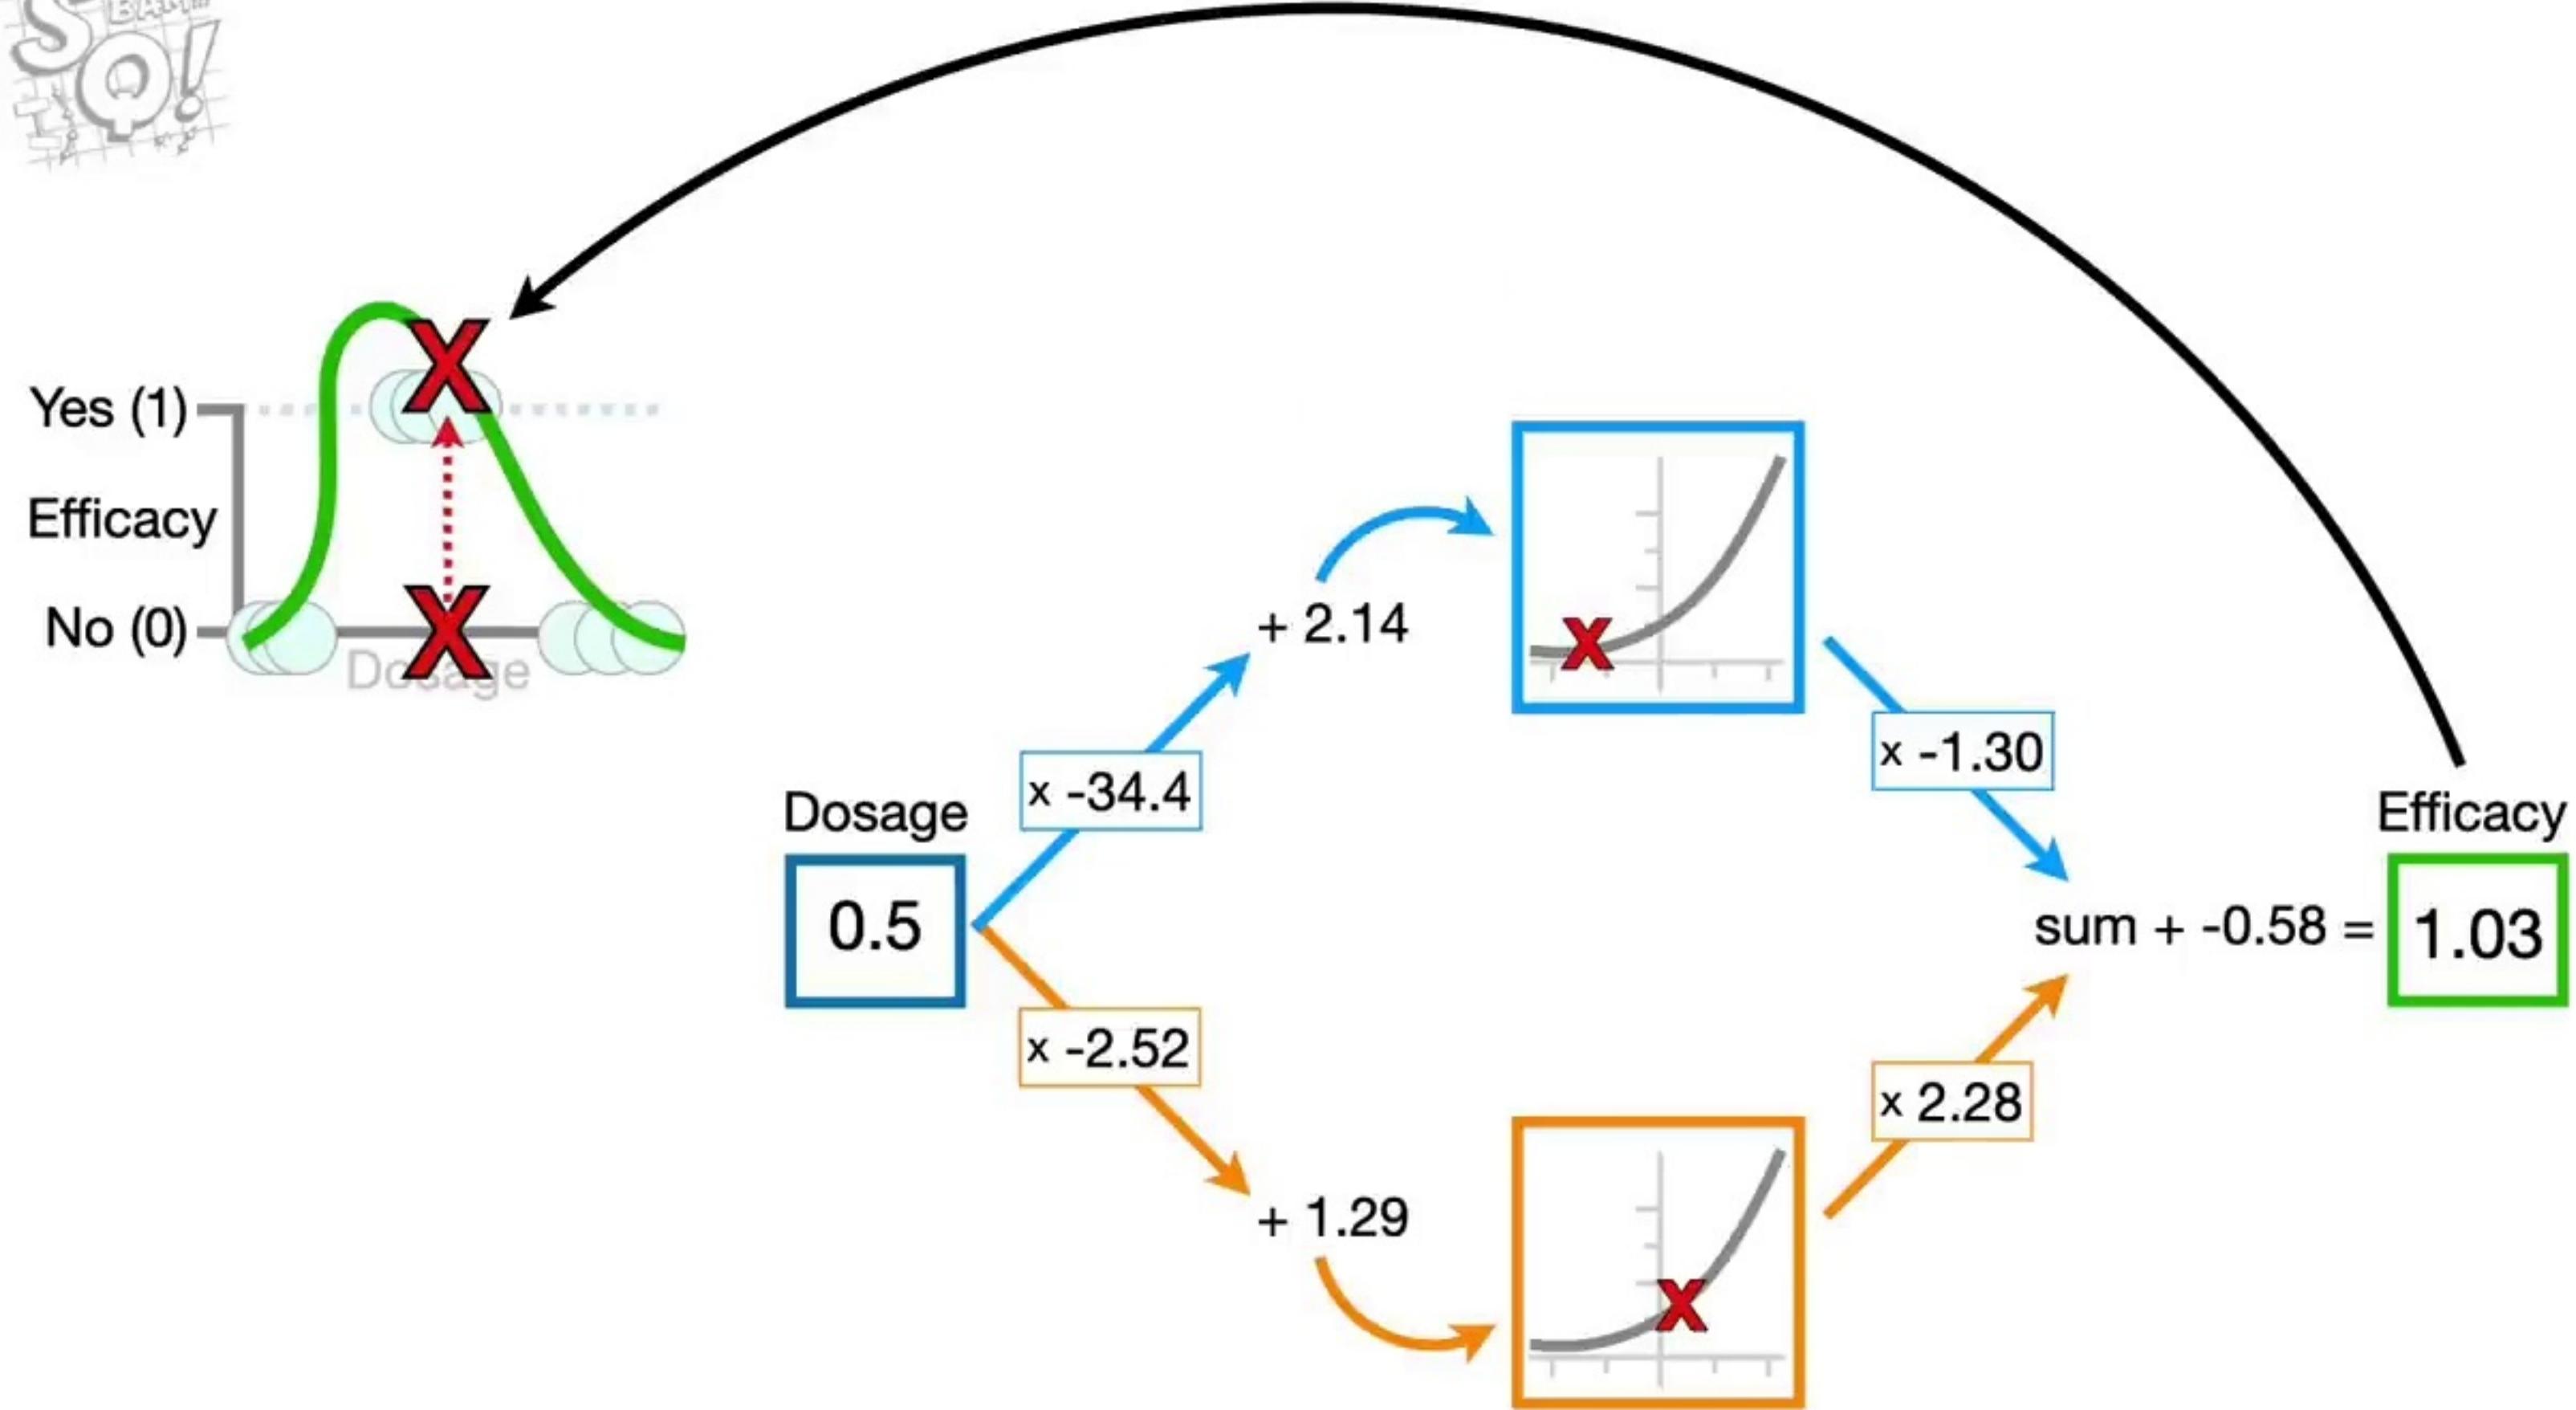
Double  
BAM!!  
**SQ!**

~~(A really small number  $x - 1.30$ ) + (0.71 \times 2.28) - 0.58 = 1.03~~

Approximately Zero!

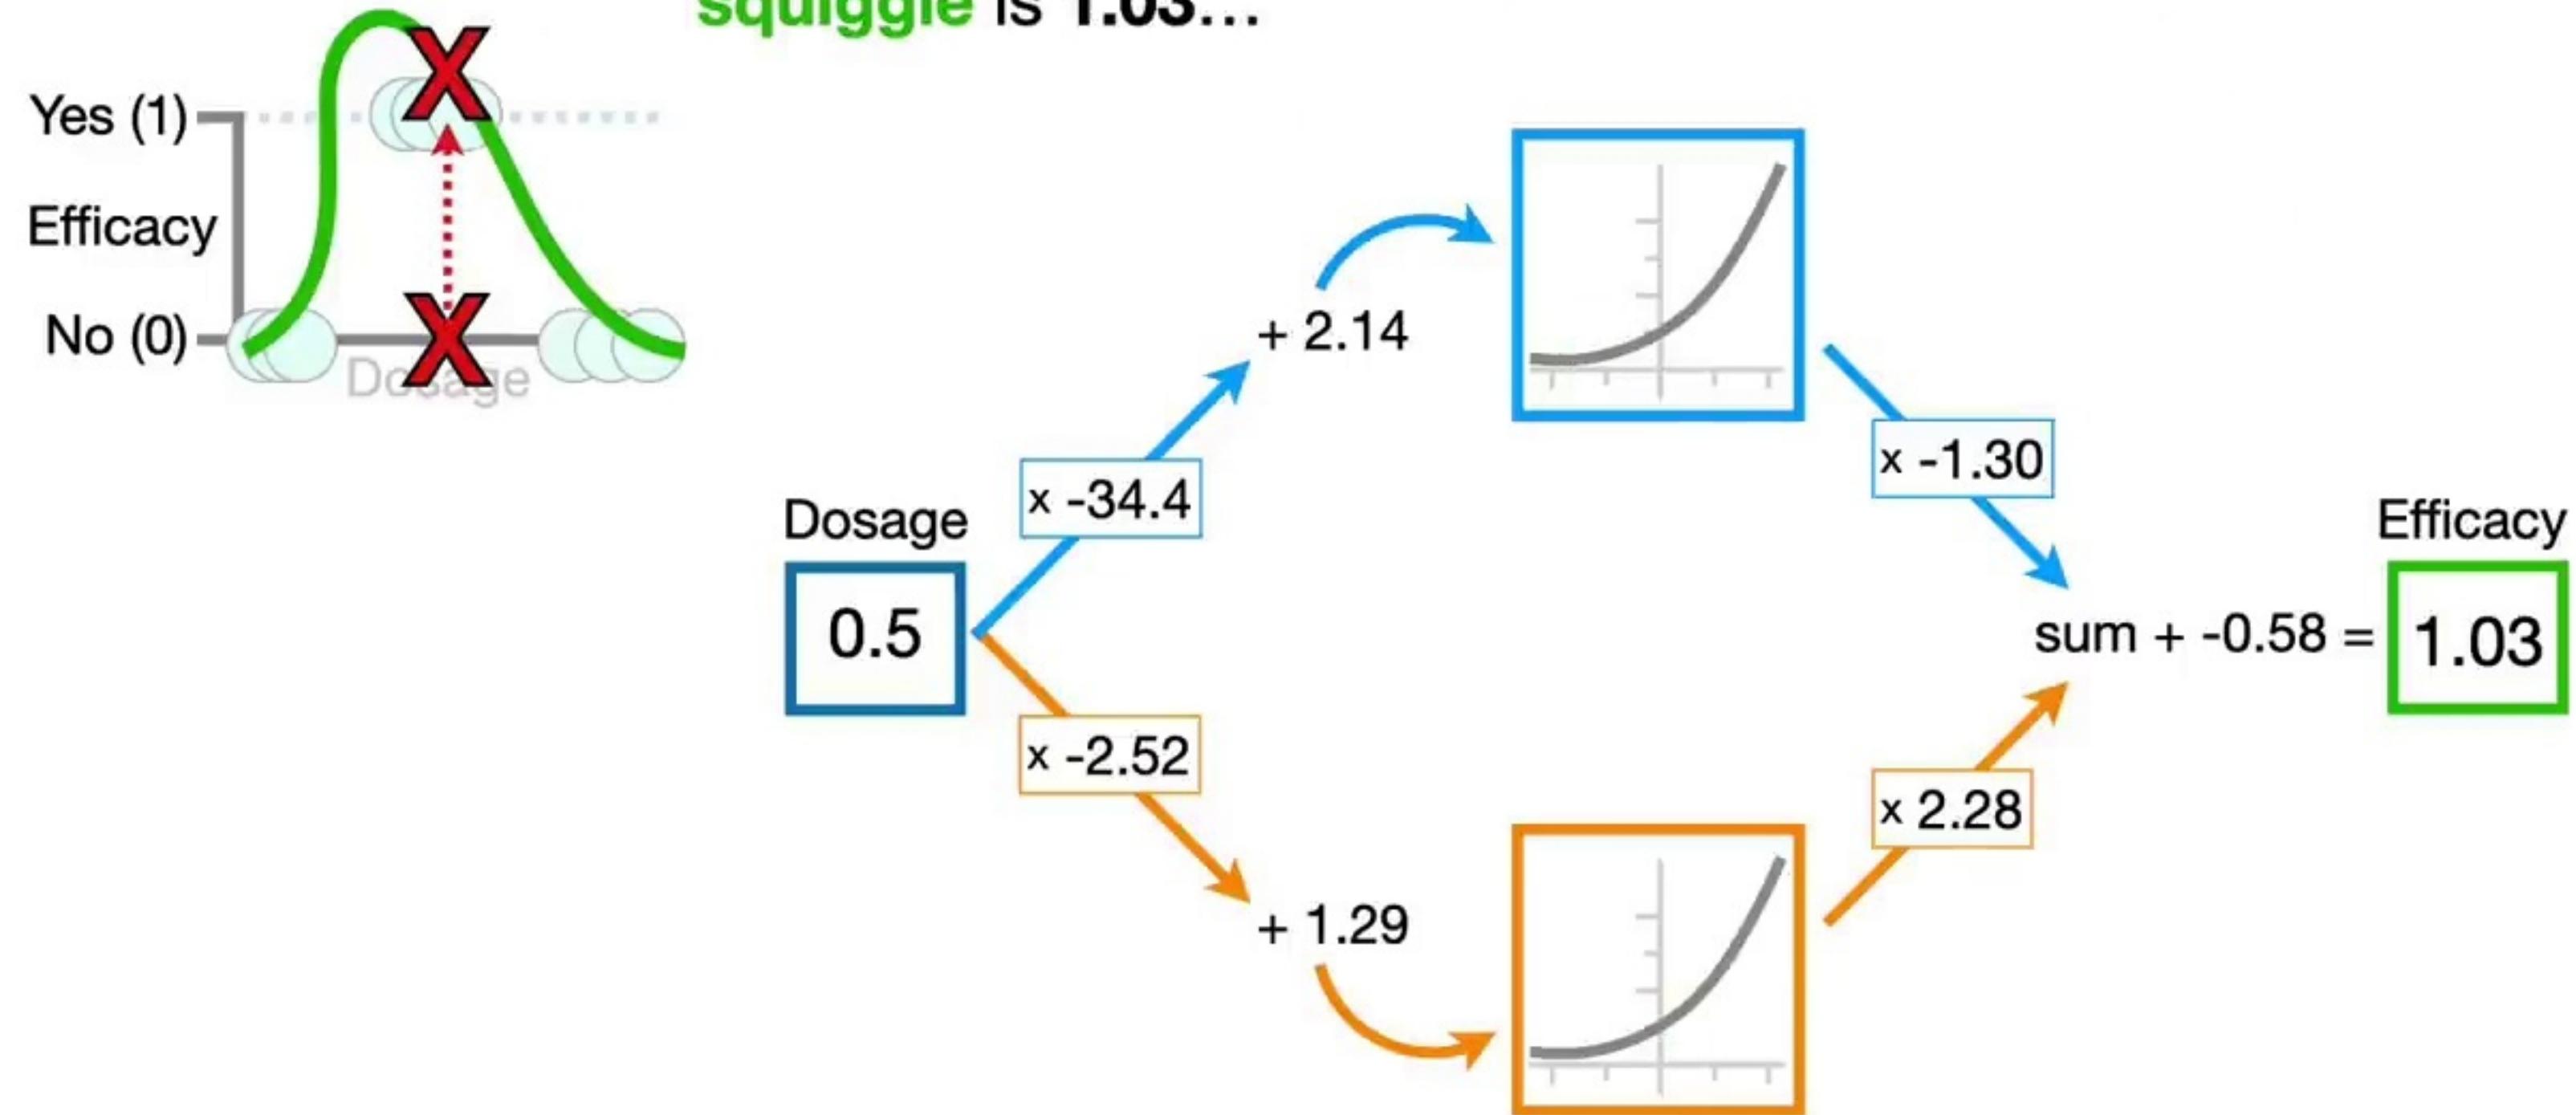


Double  
BAM!!  
**SQ!**



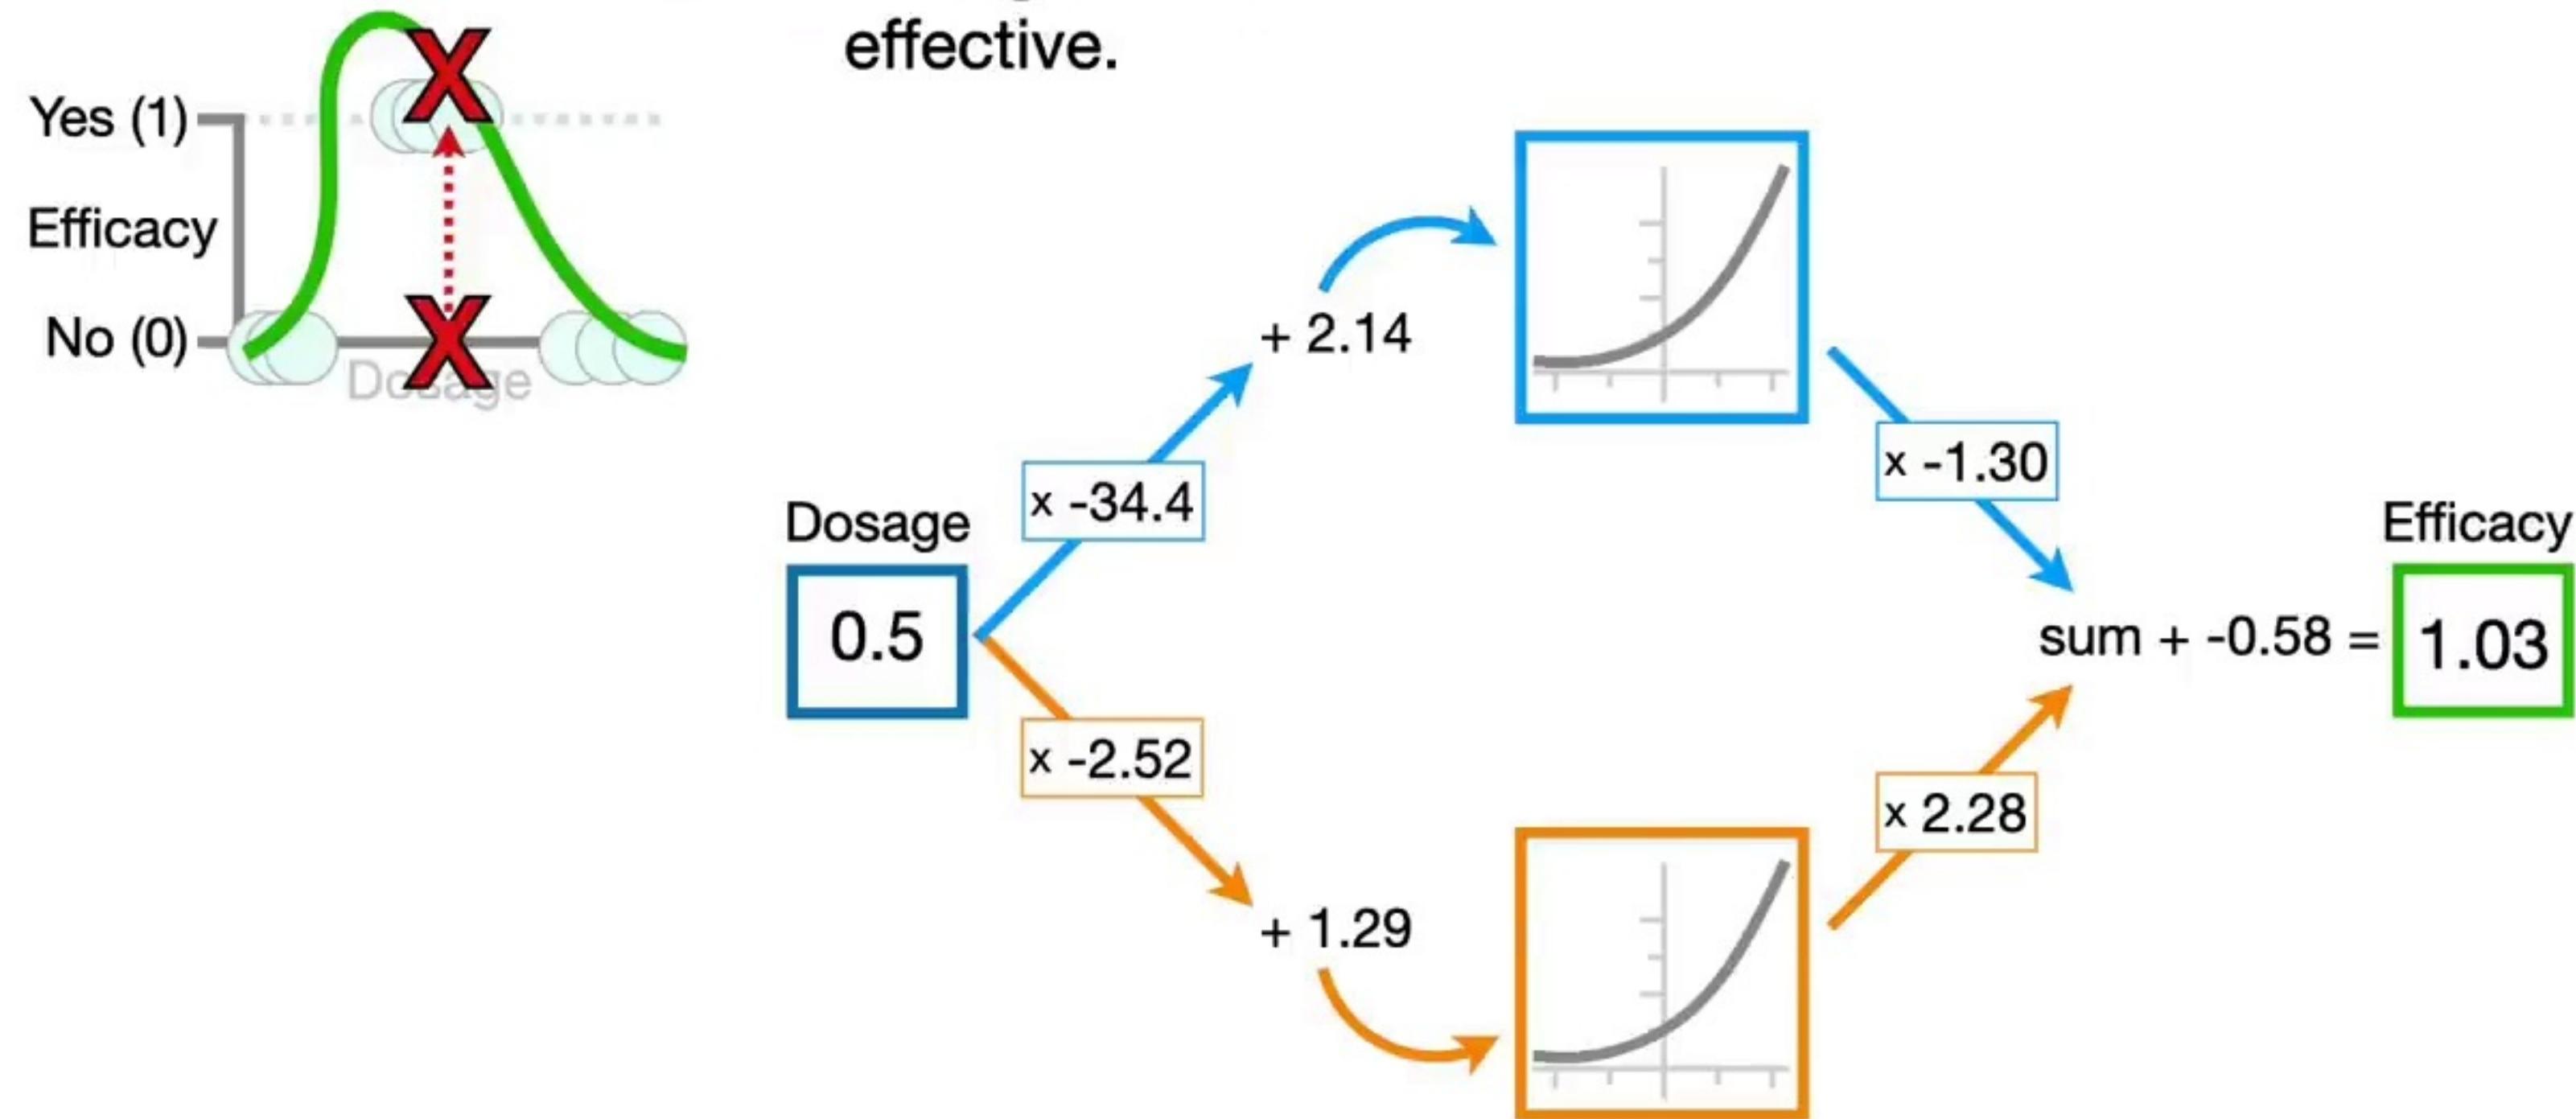


...and we see that the y-axis coordinate on the **green squiggle** is **1.03**...



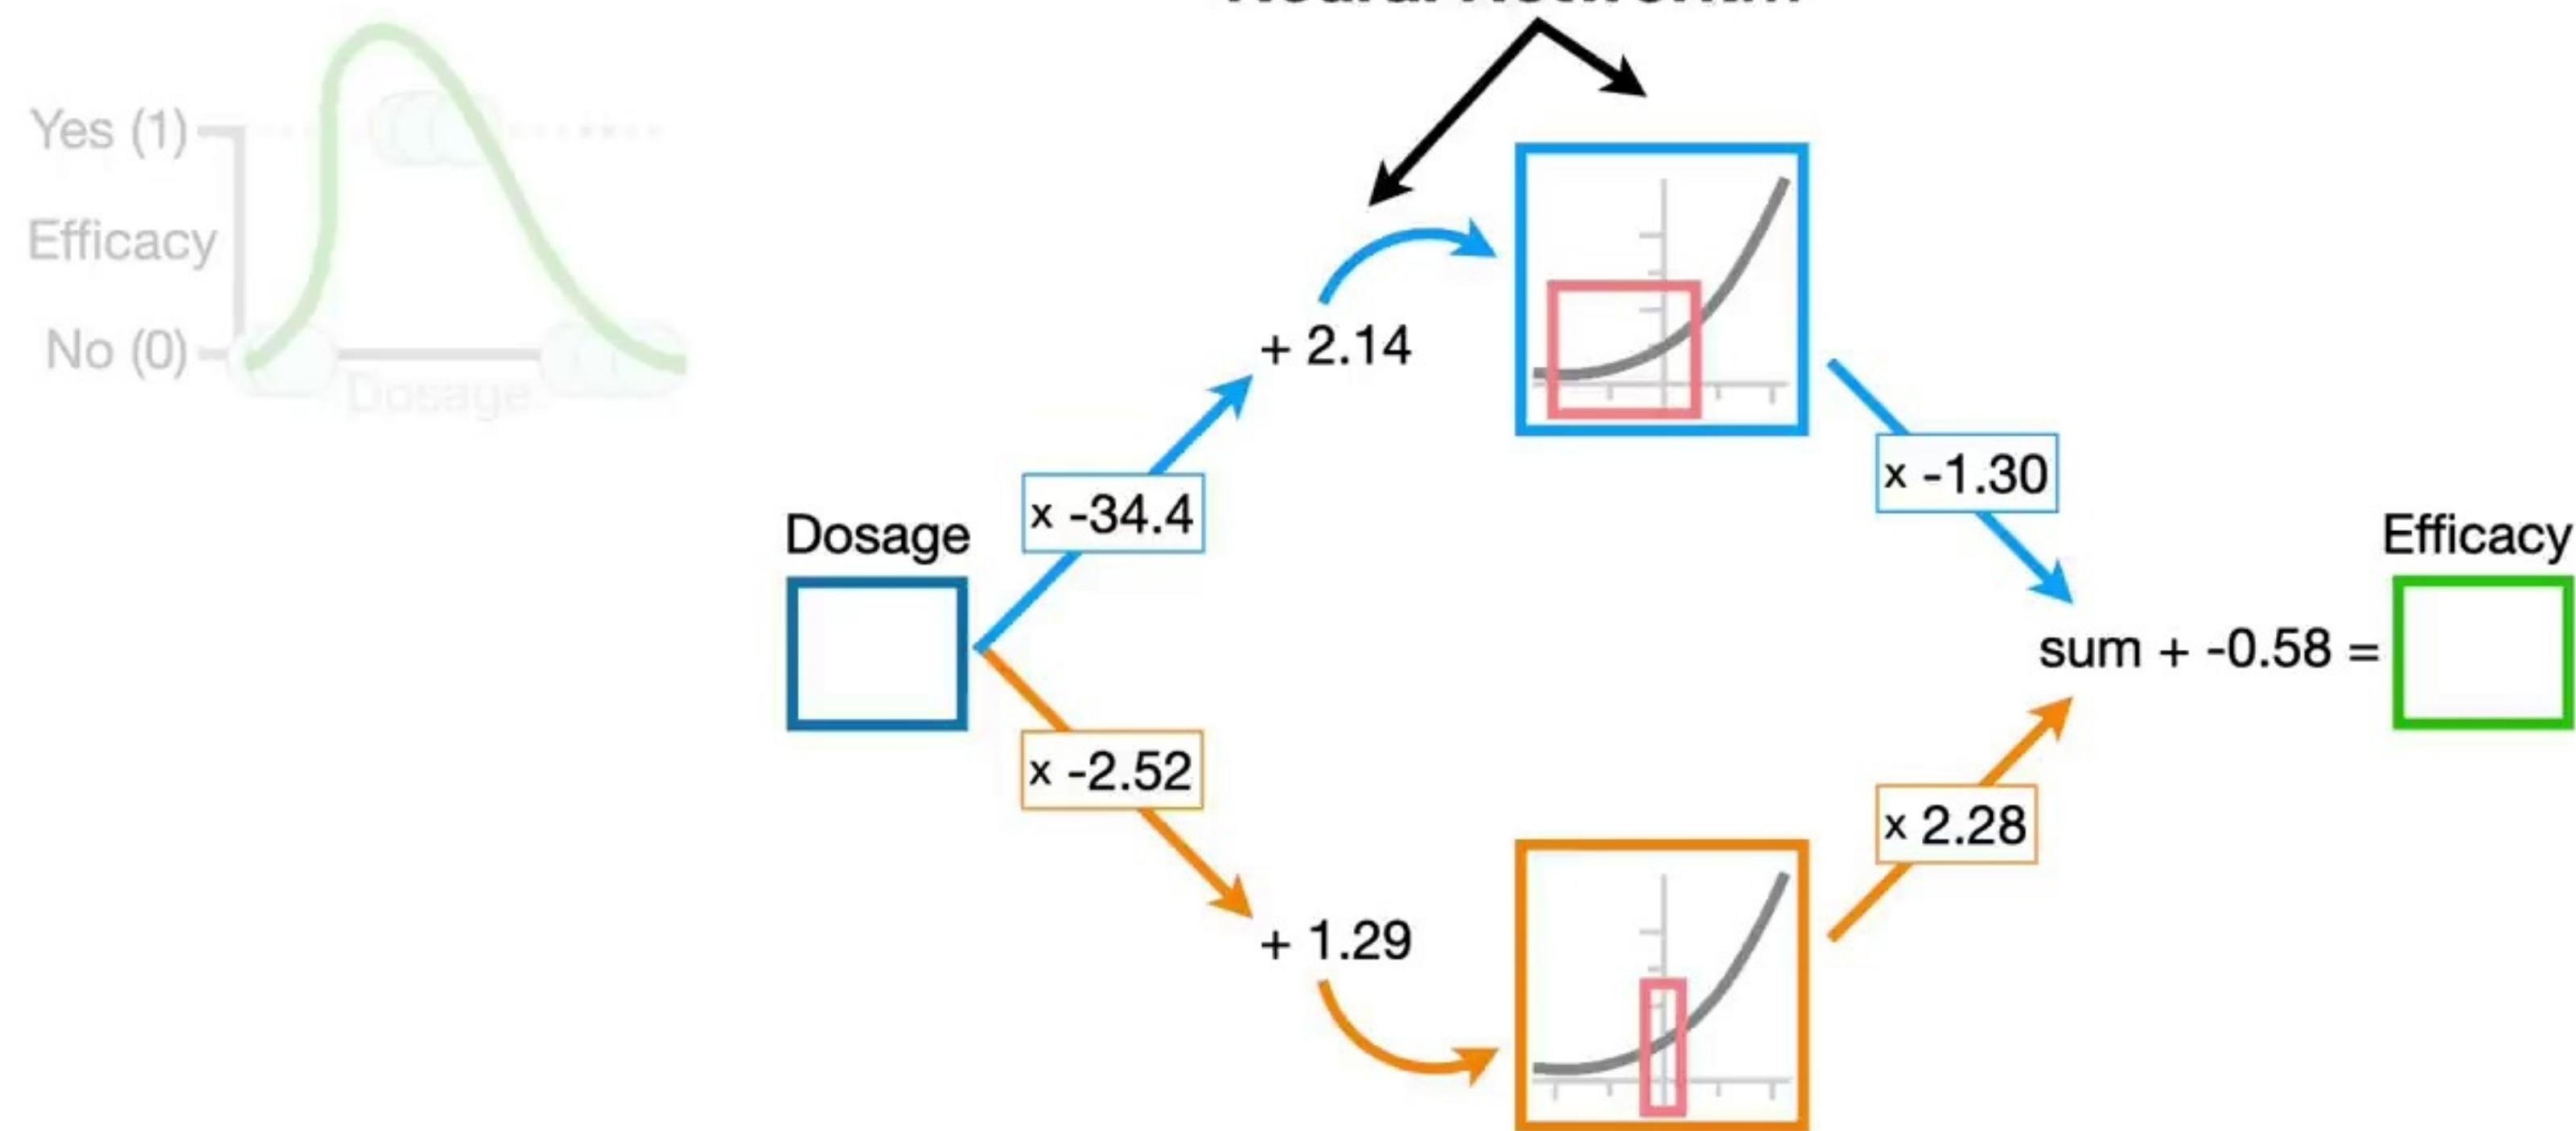


...and since **1.03** is closer to **1** than **0**, we will conclude that a **Dosage = 0.5** is effective.



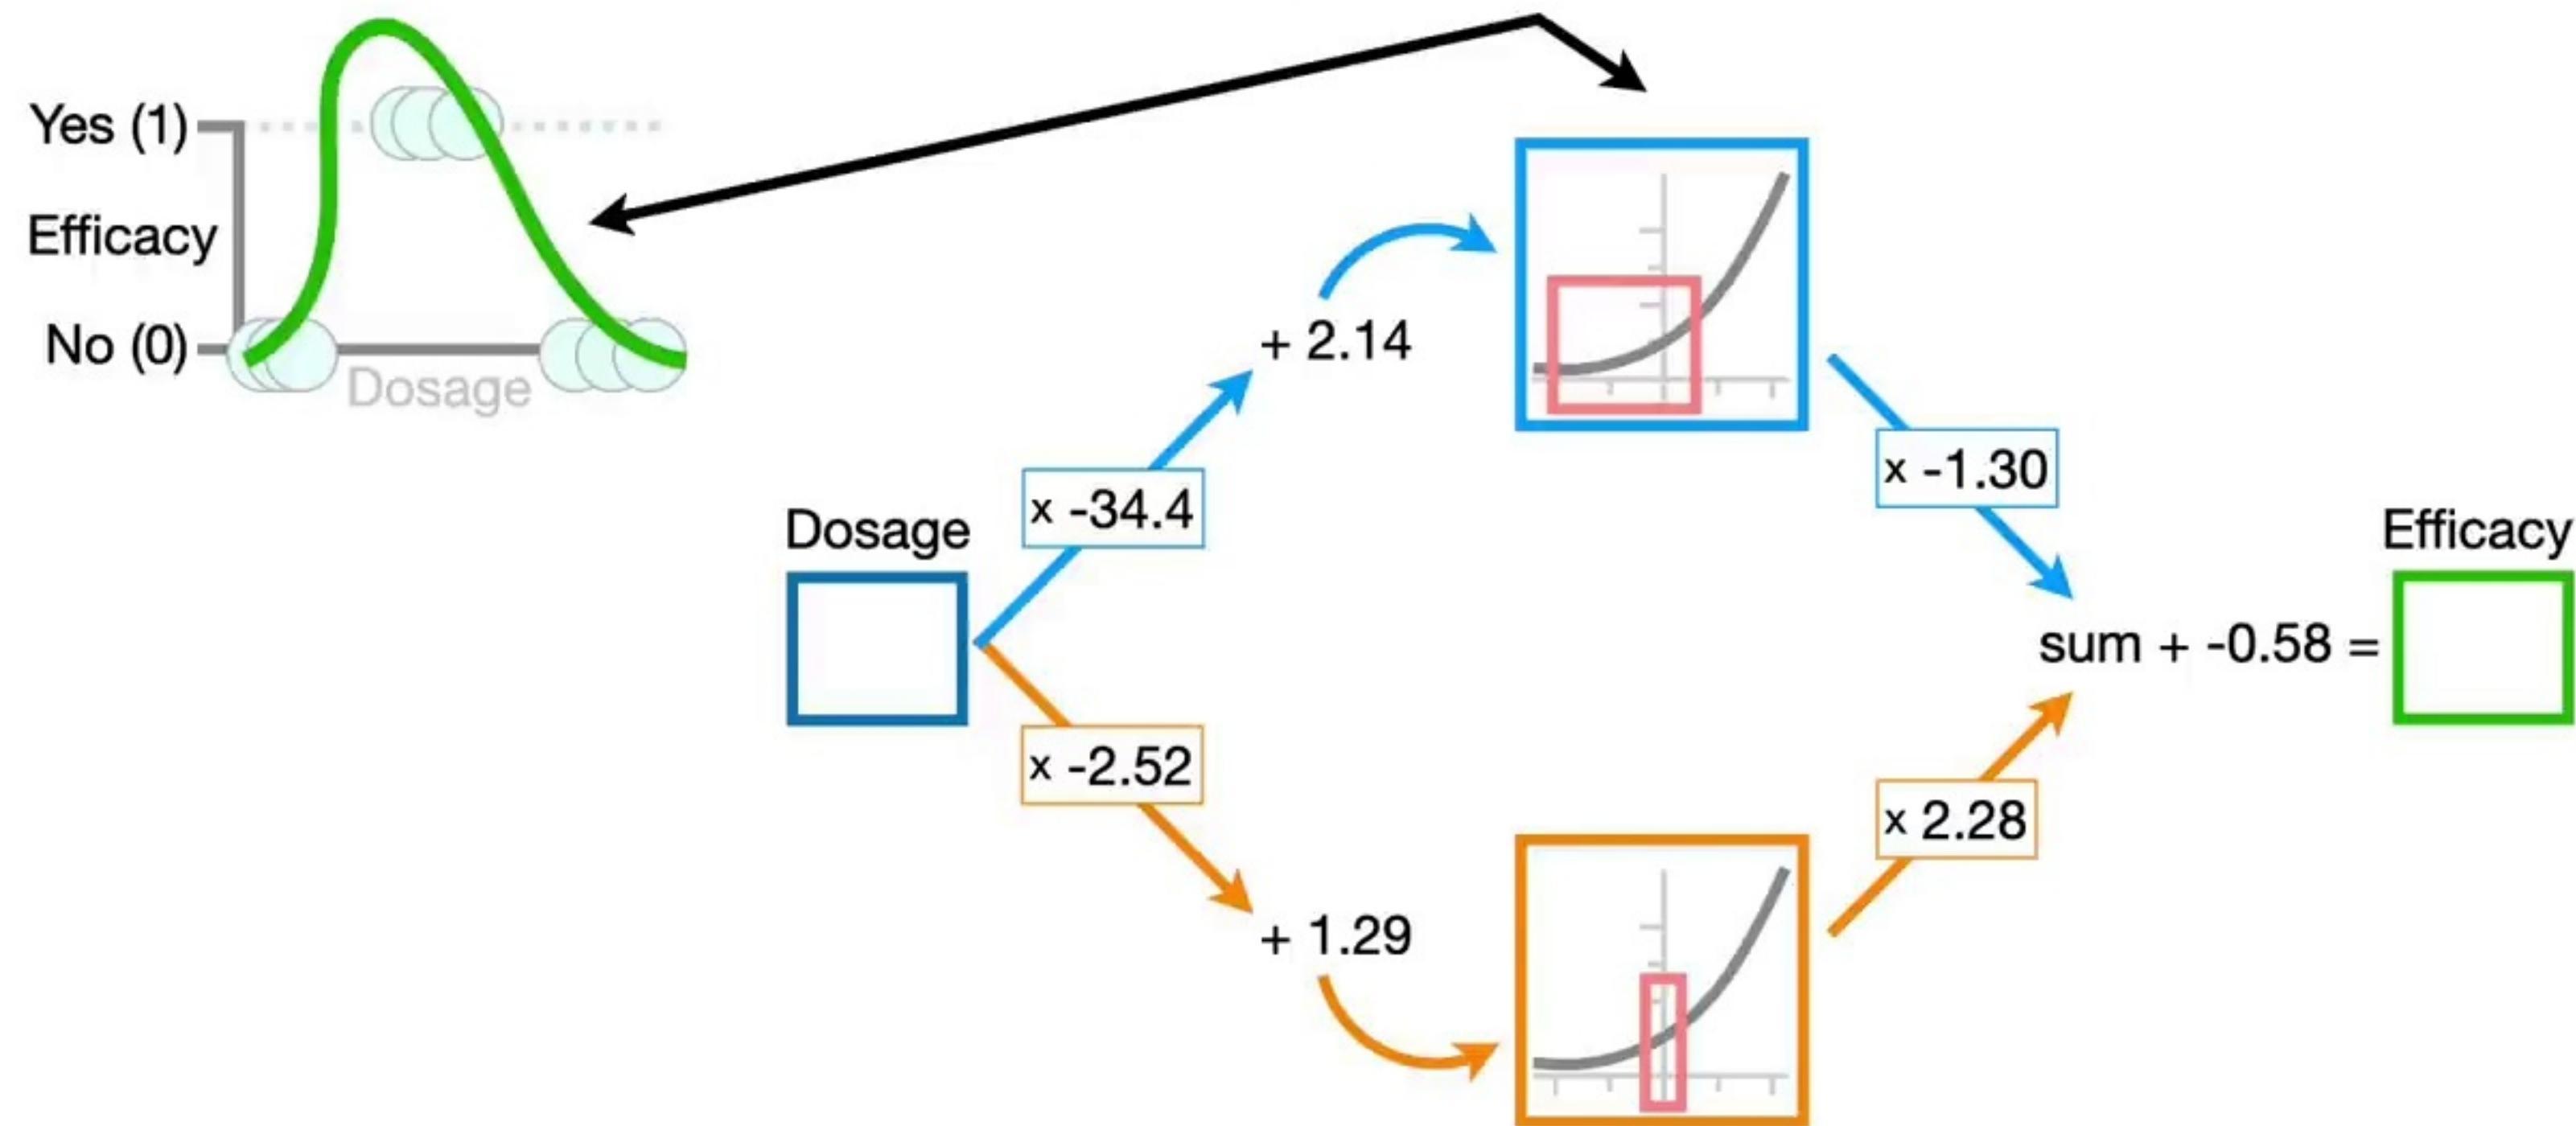


Now, if you've made it this far, you may  
be wondering why this is called a  
**Neural Network...**



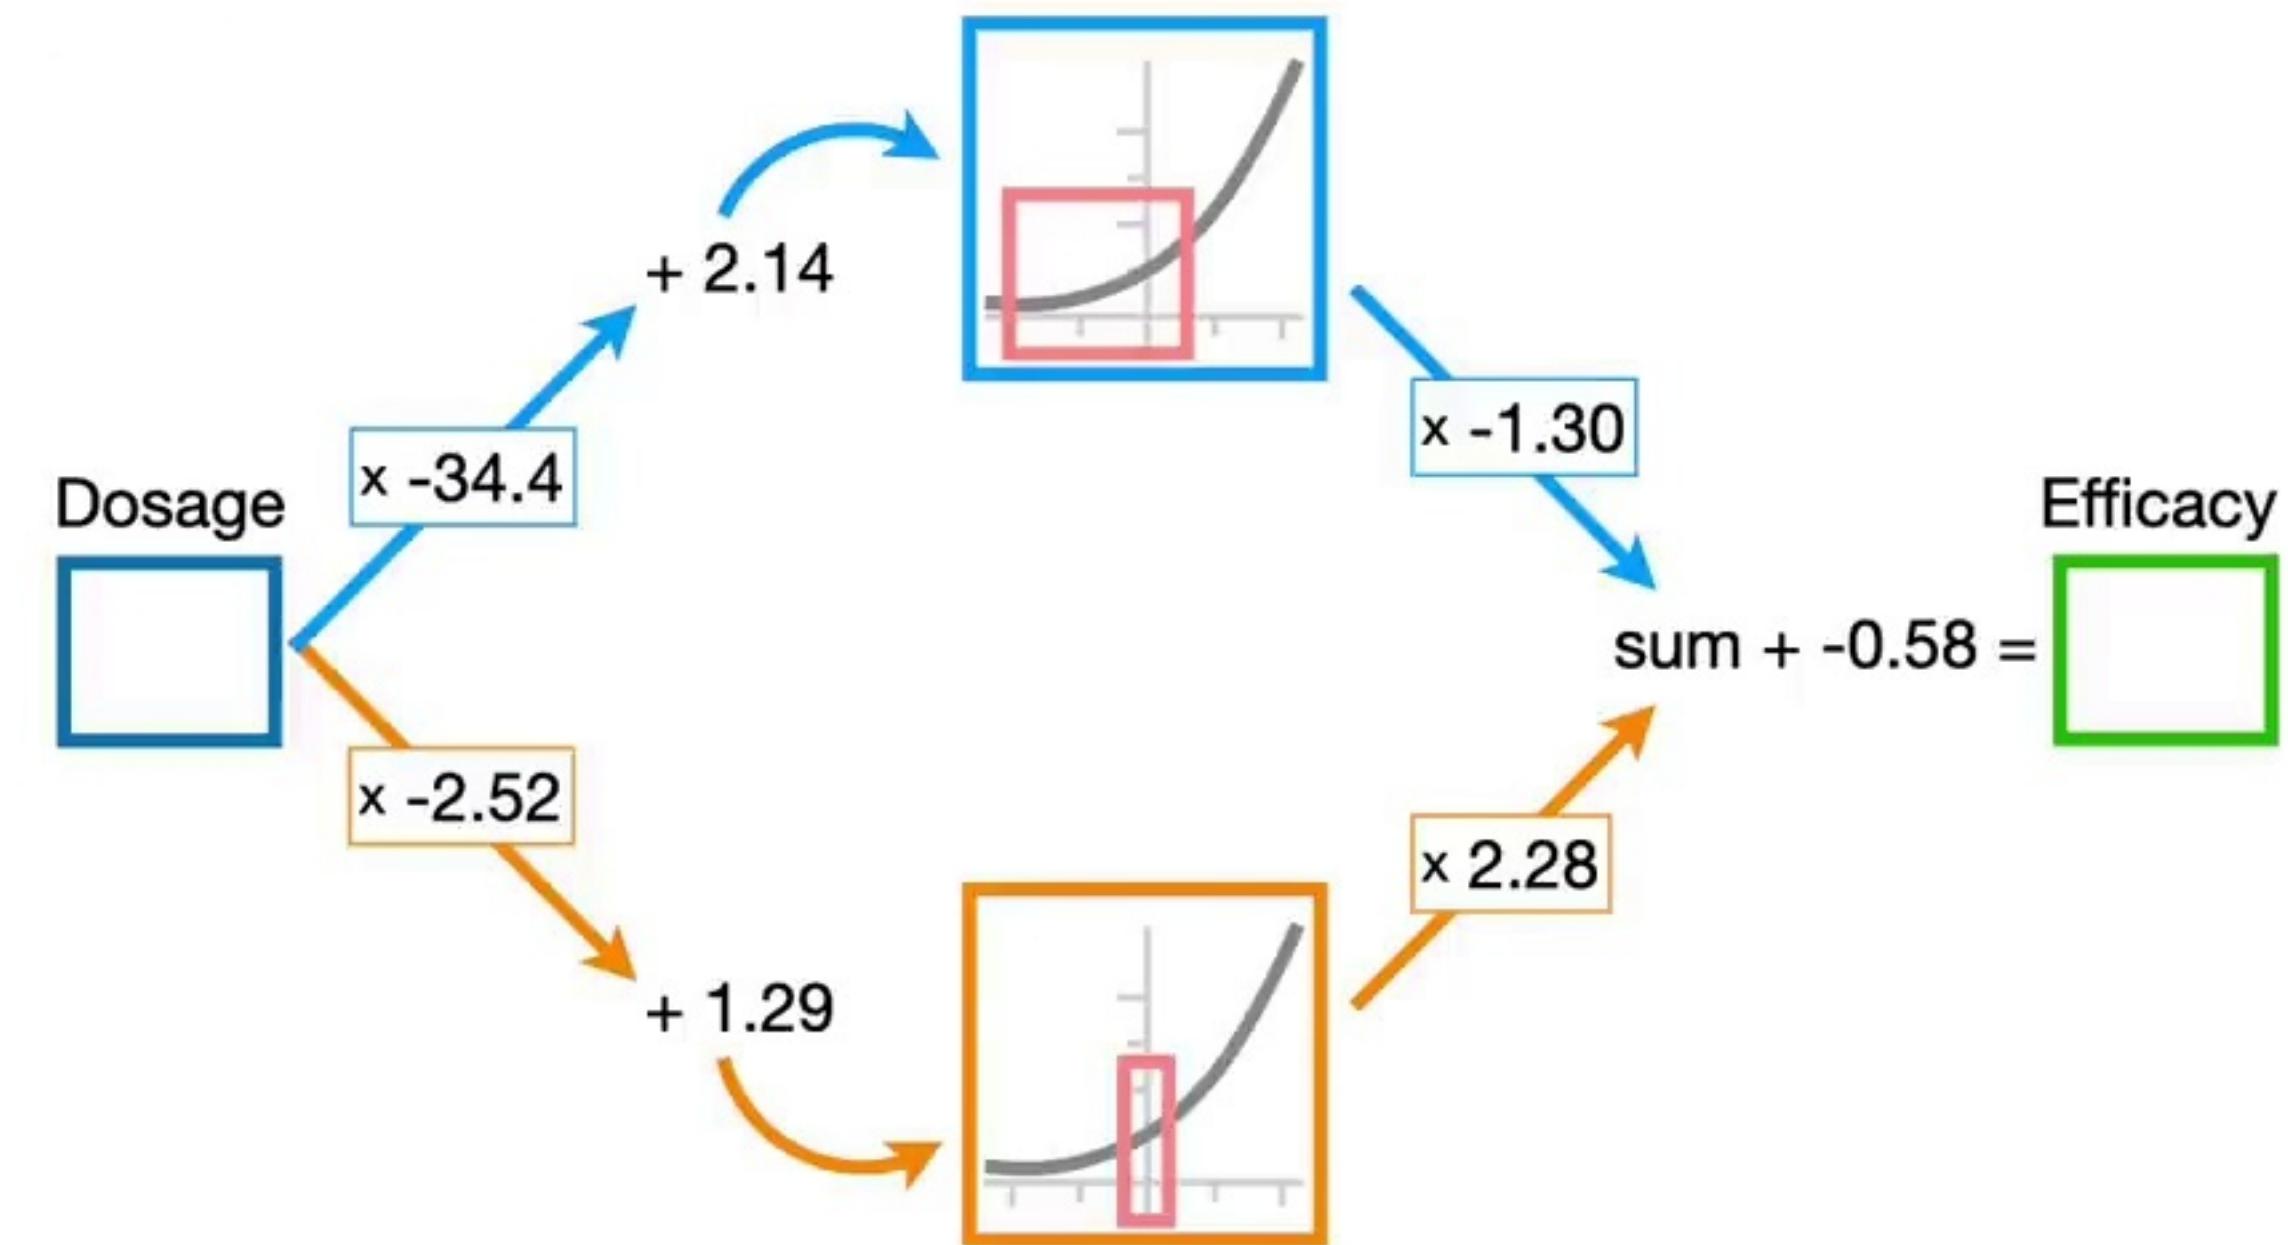


...instead of a **Big Fancy Squiggle Fitting Machine.**



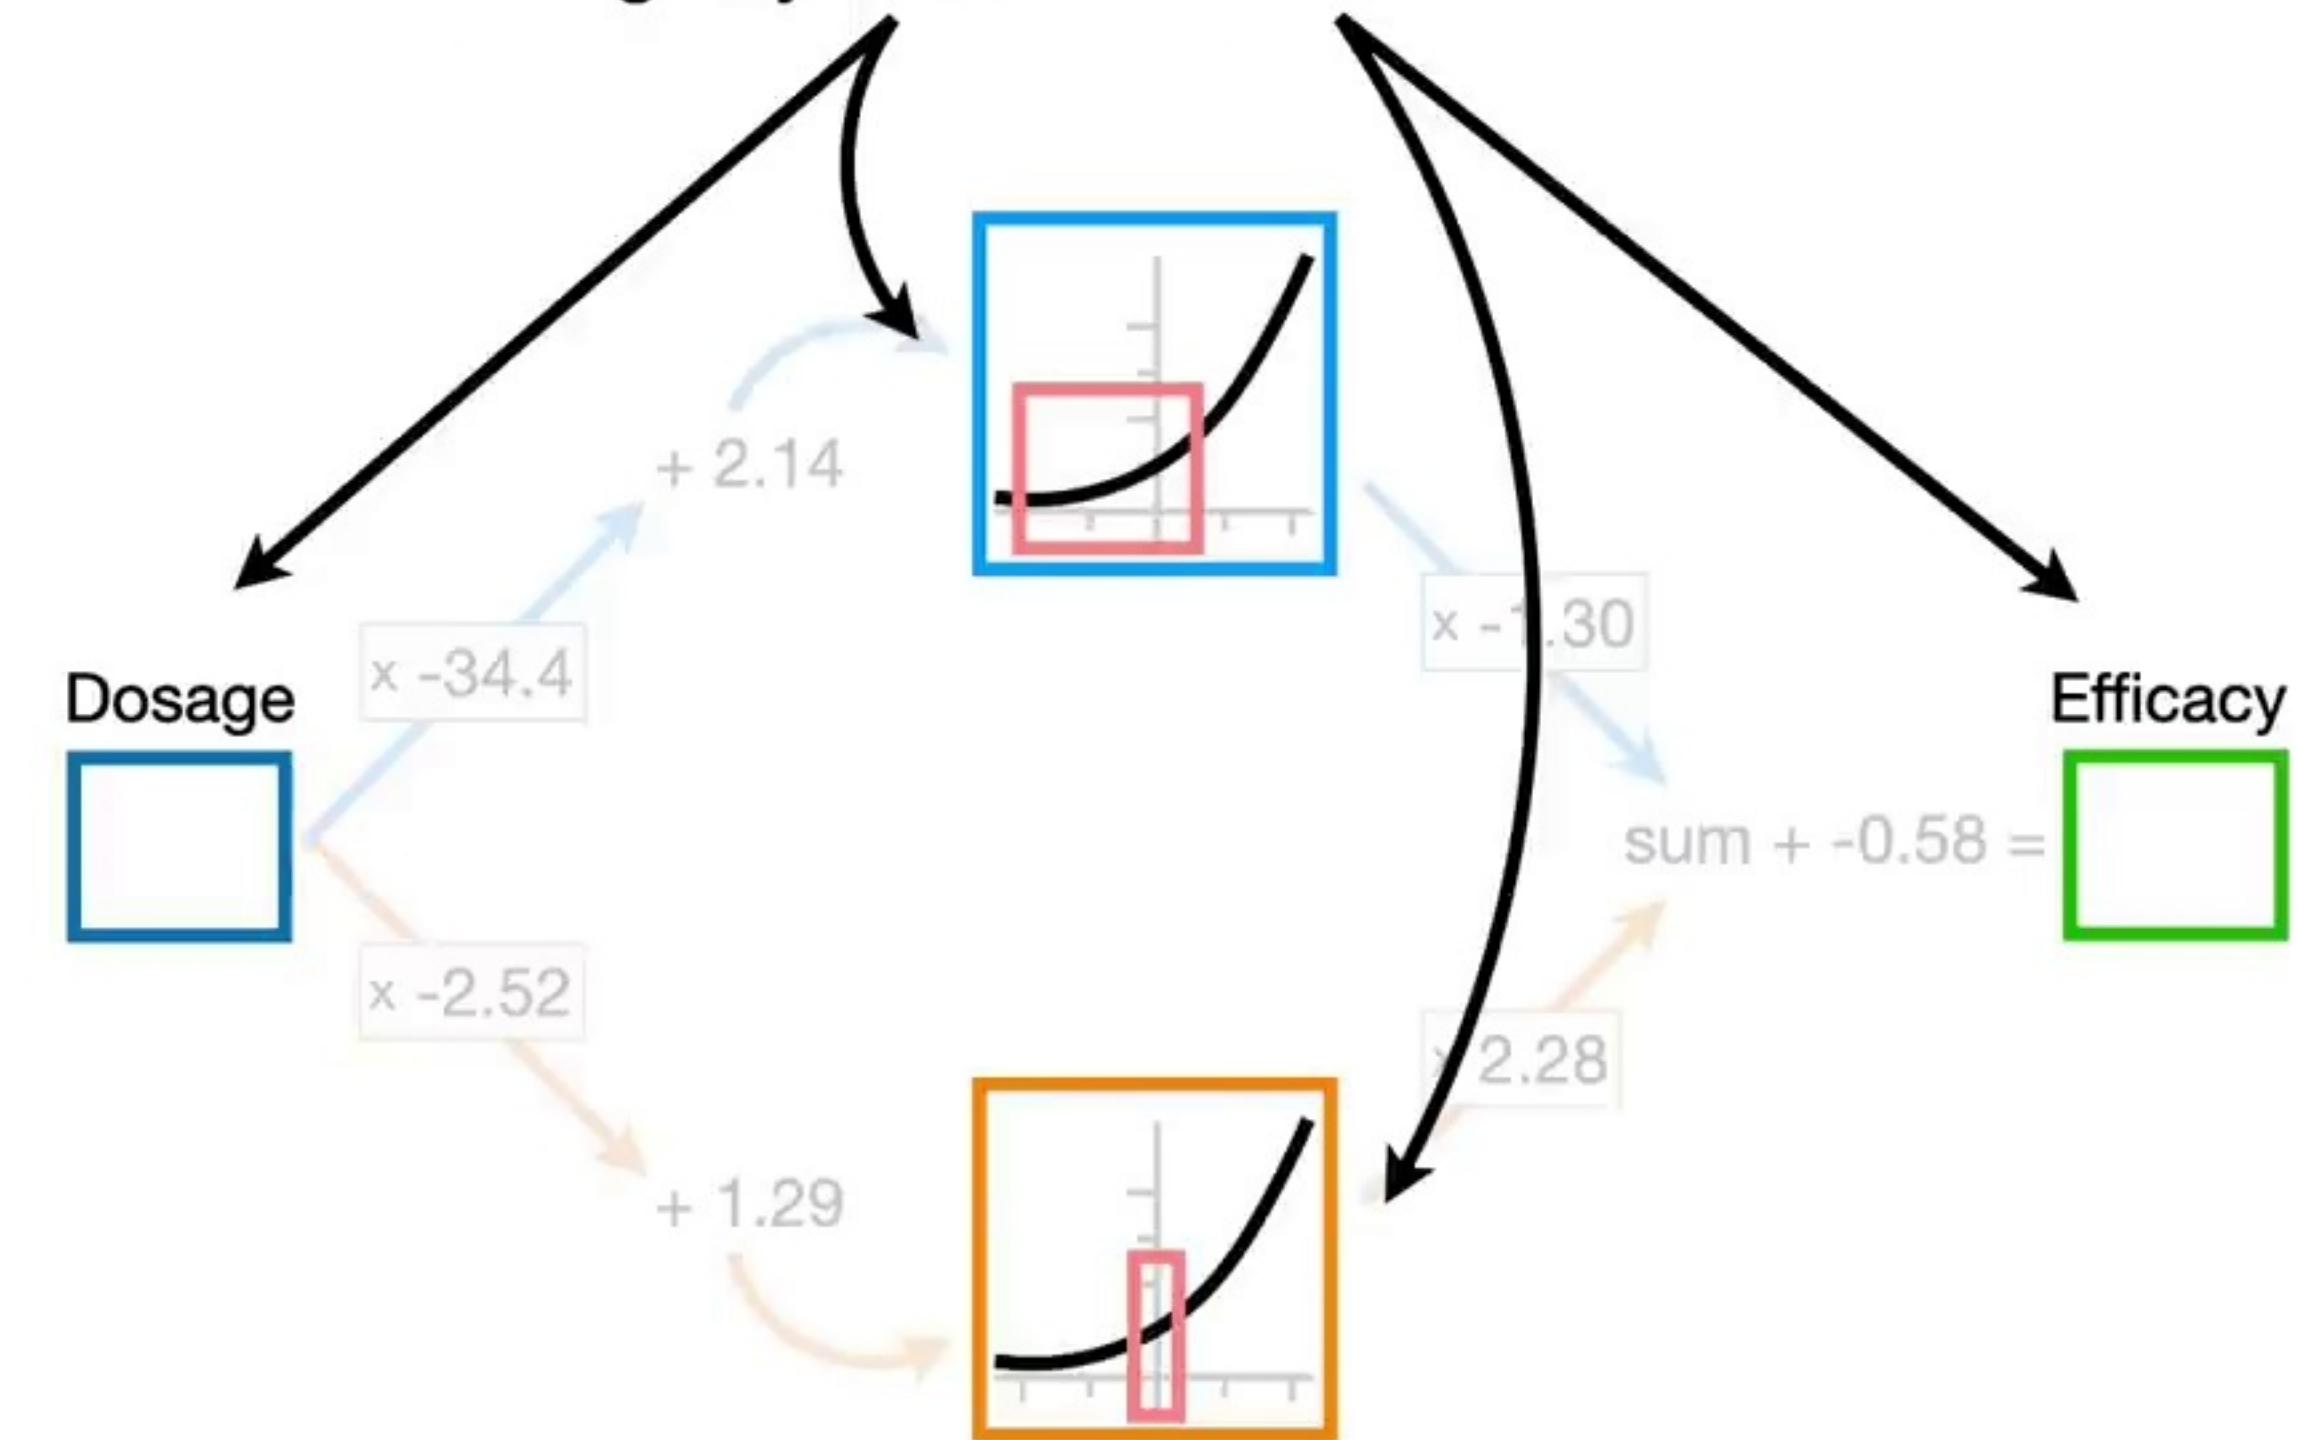


The reason is that way back in the  
**1940s and 50s** when **Neural  
Networks** were invented...



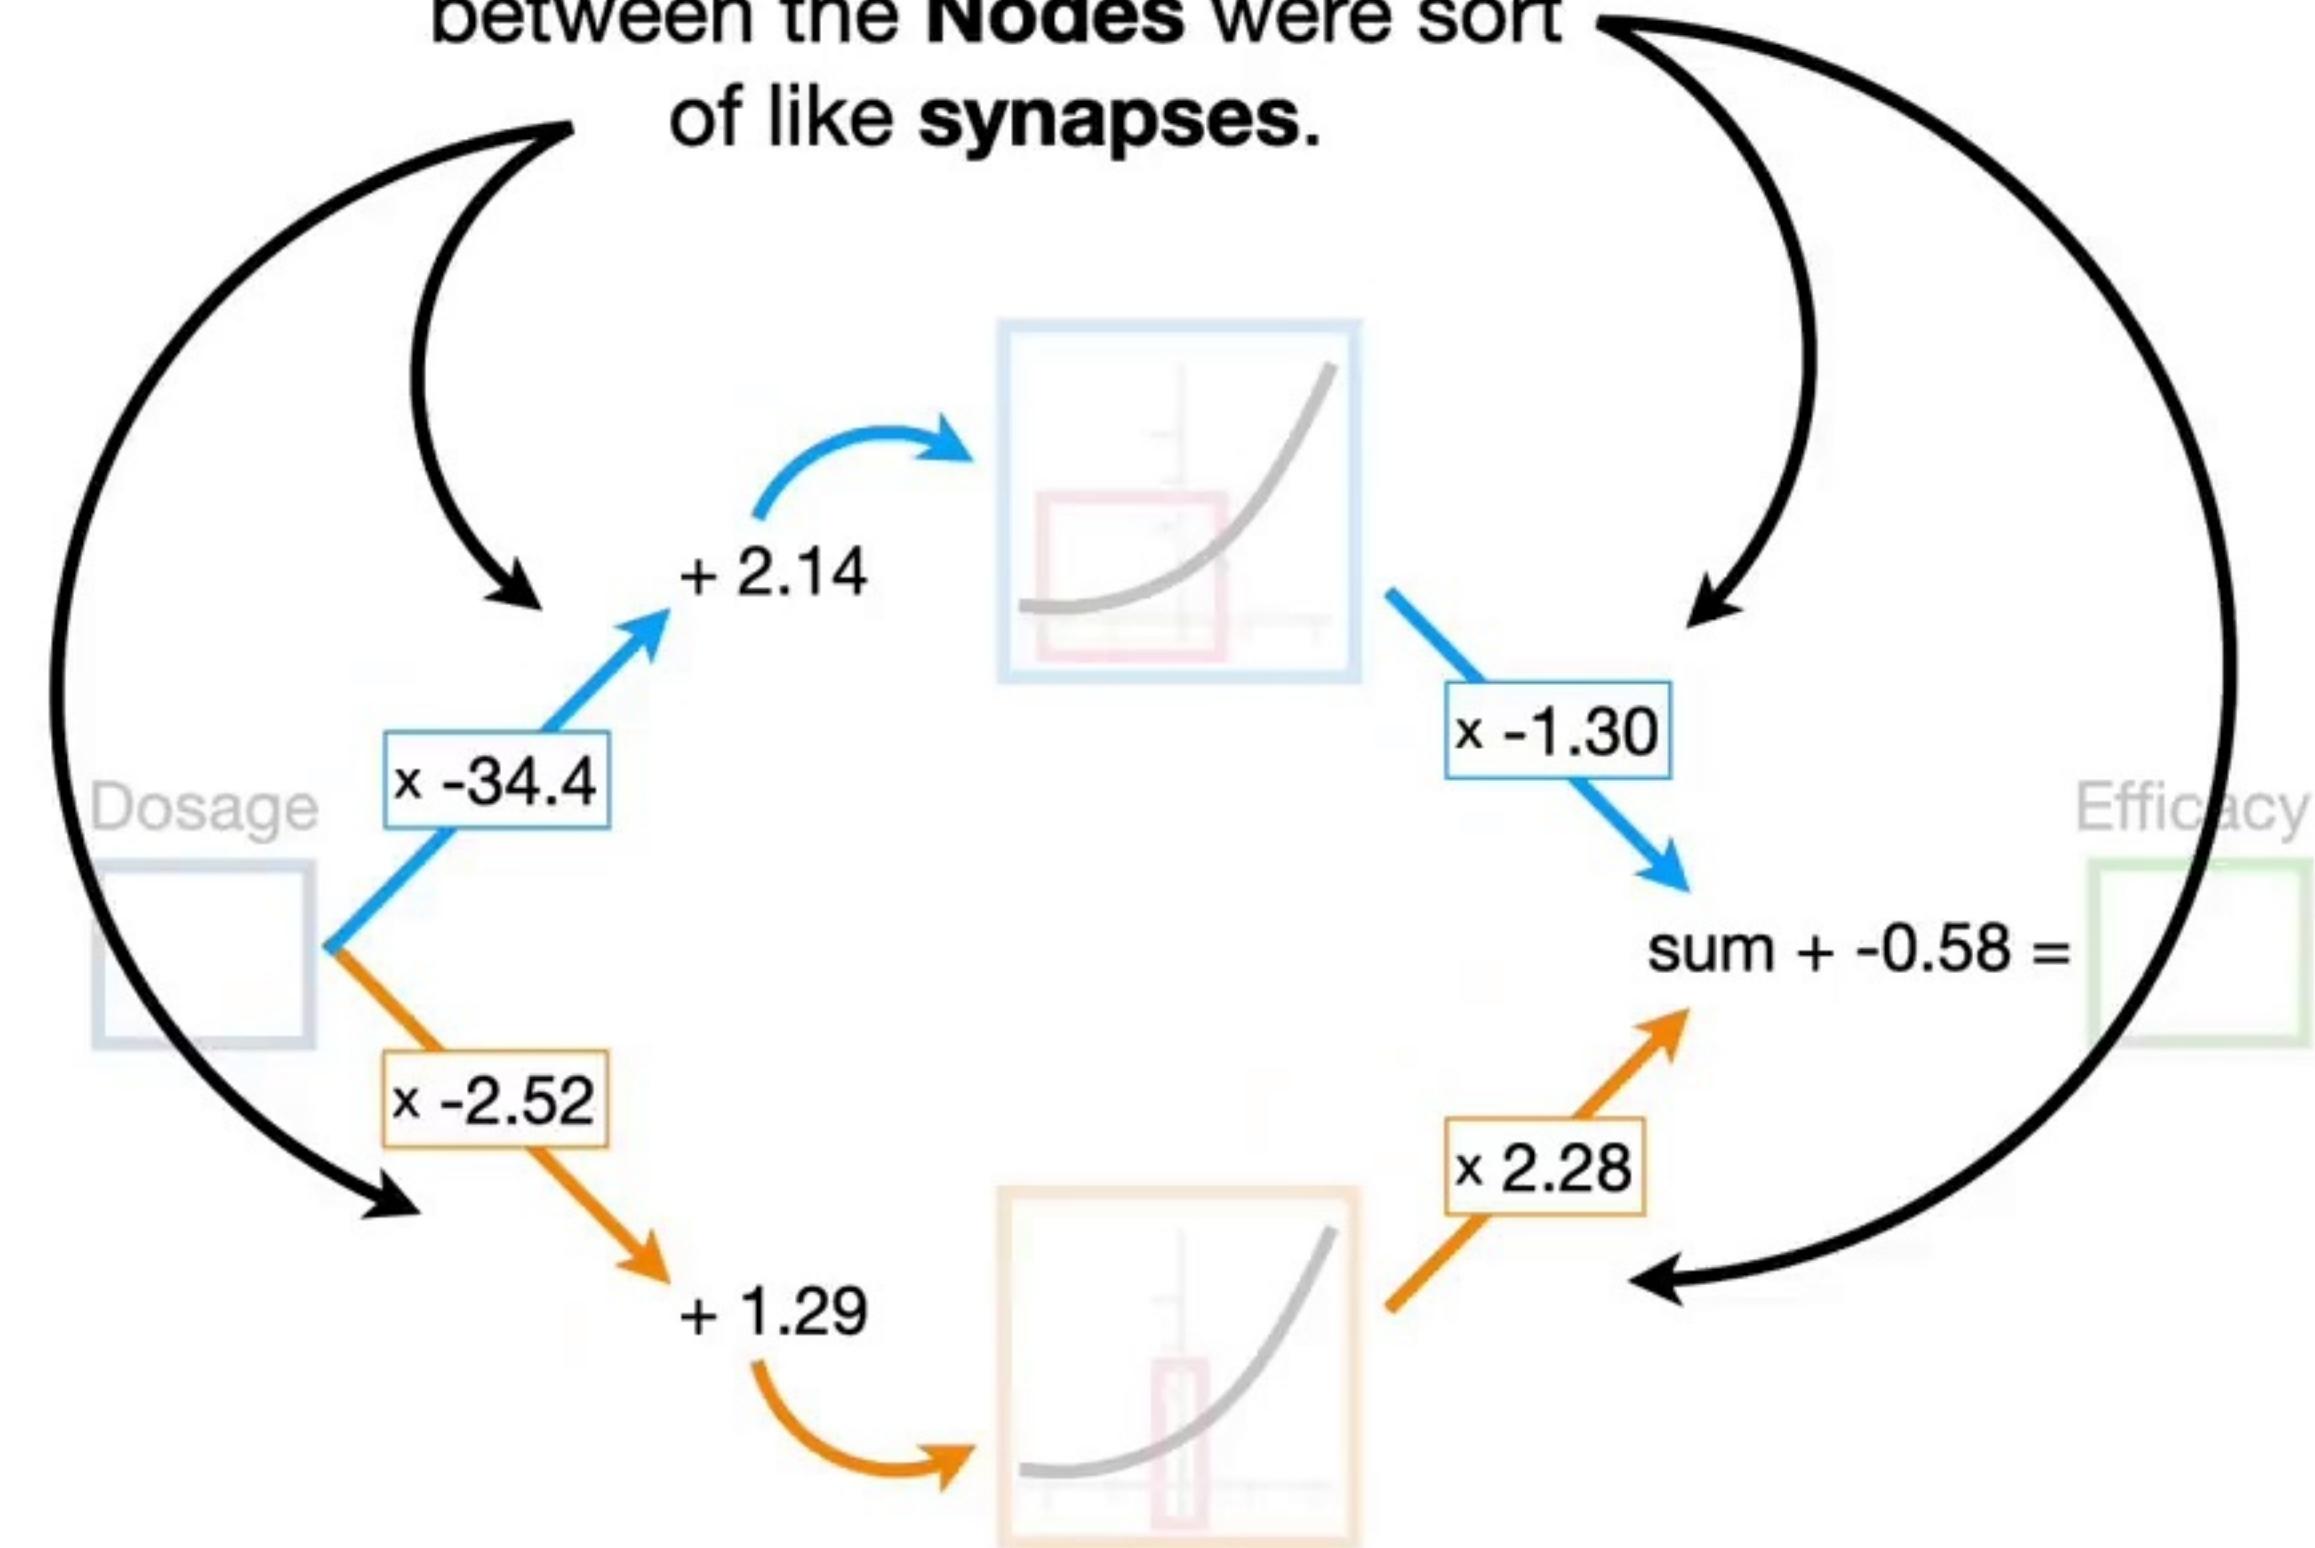


...they thought the **Nodes** were  
vaguely like **Neurons**...



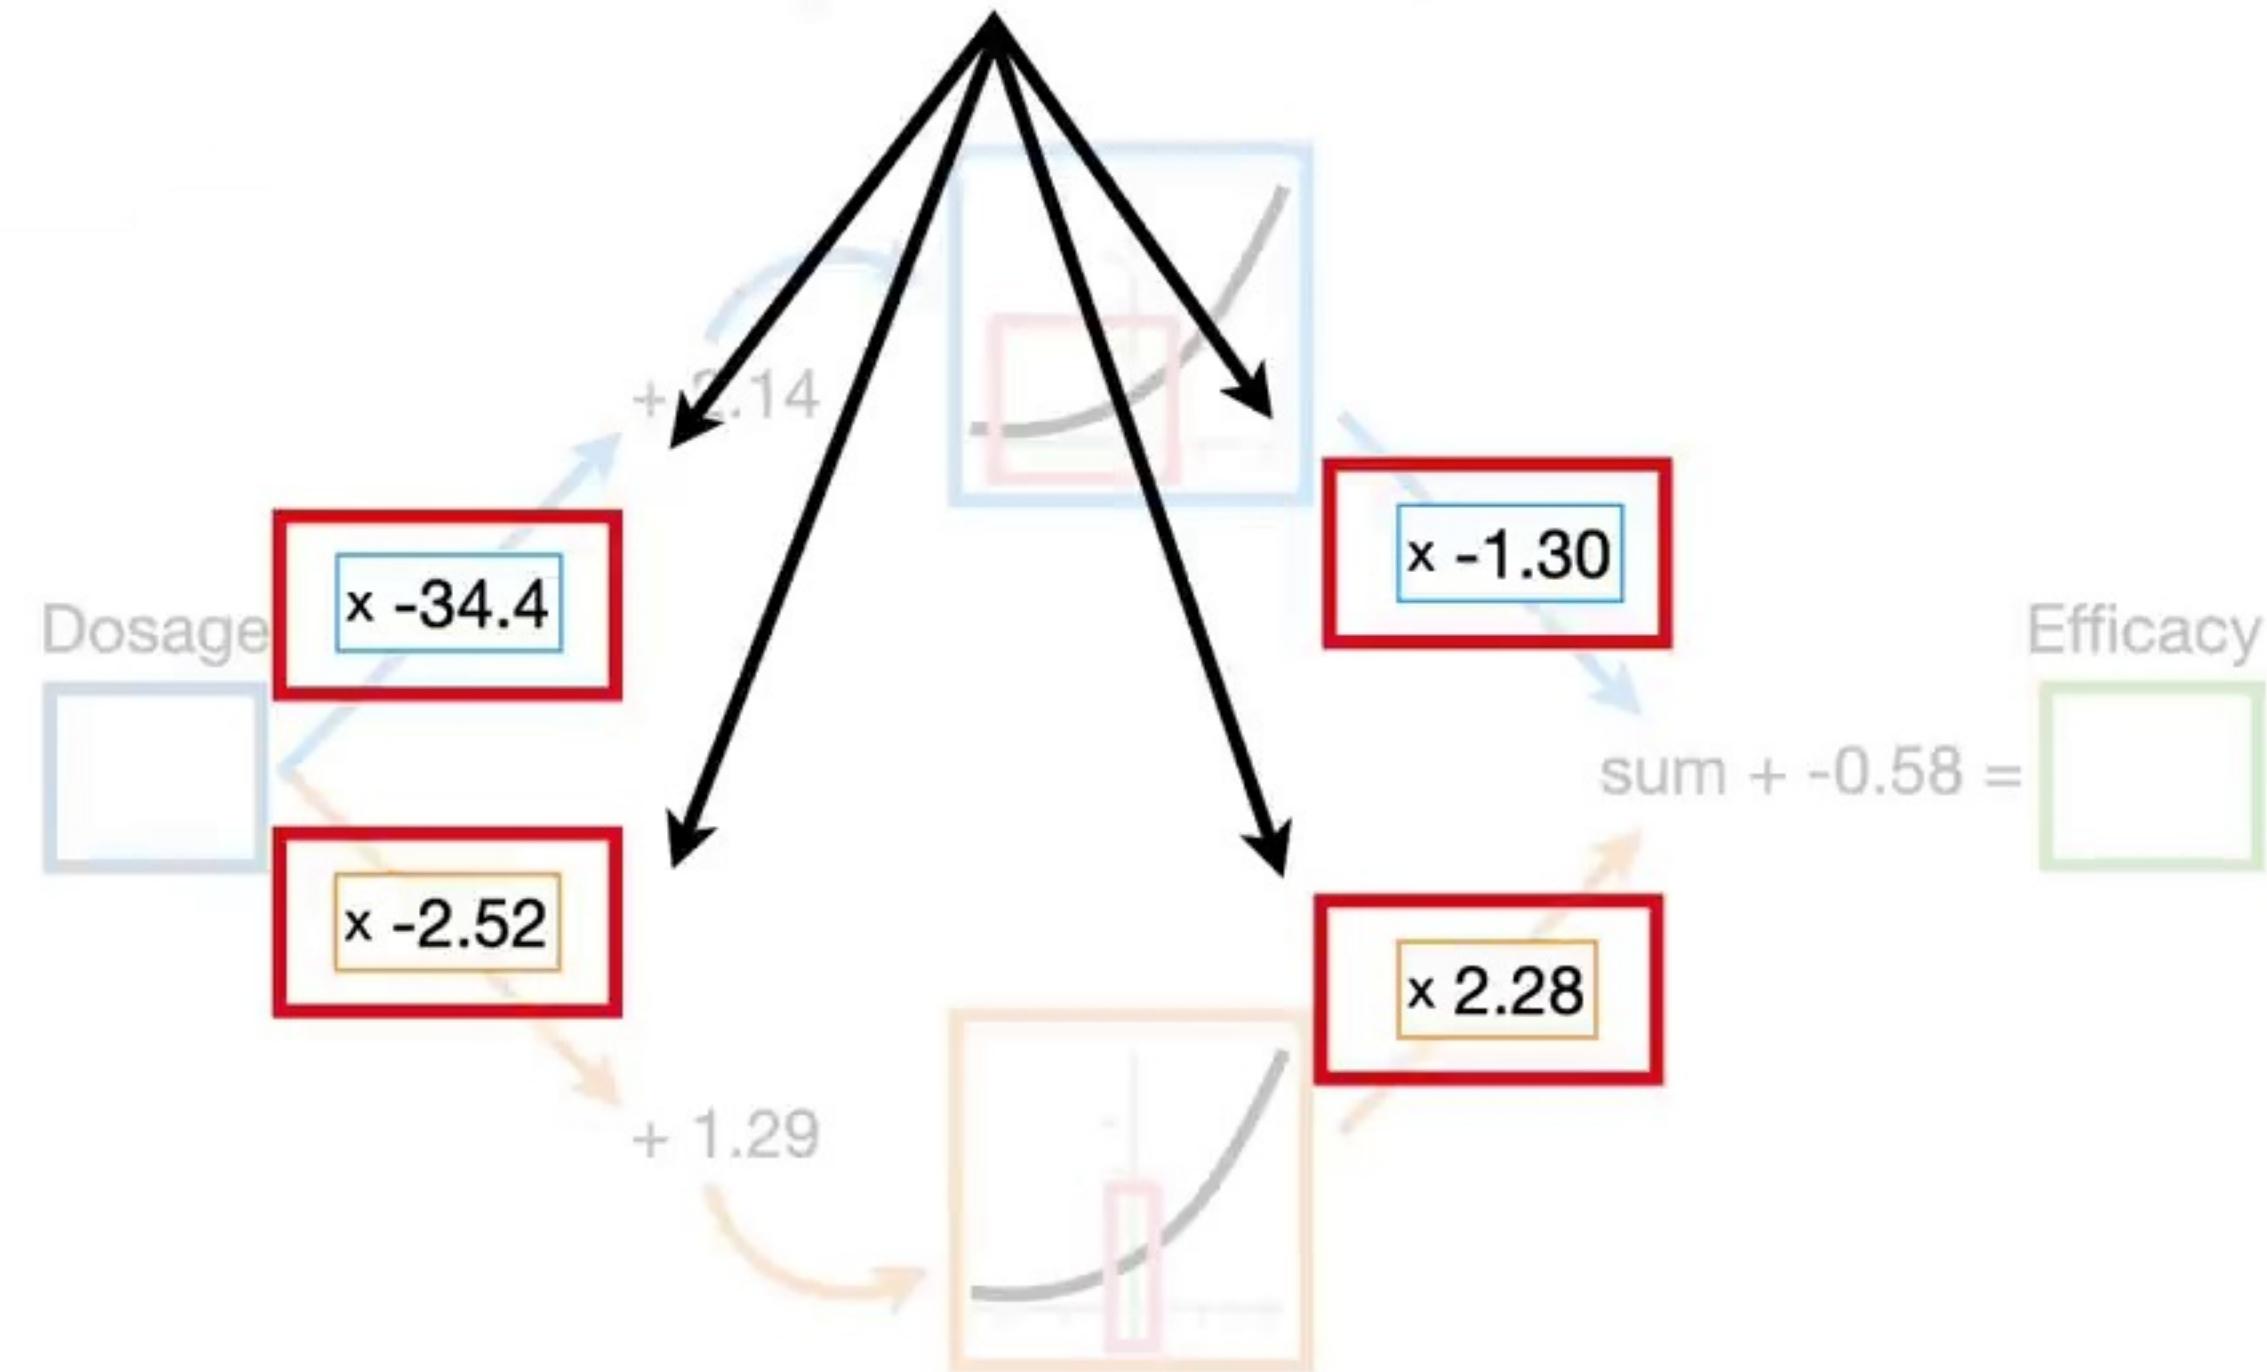


...and the **connections**  
between the **Nodes** were sort  
of like **synapses**.



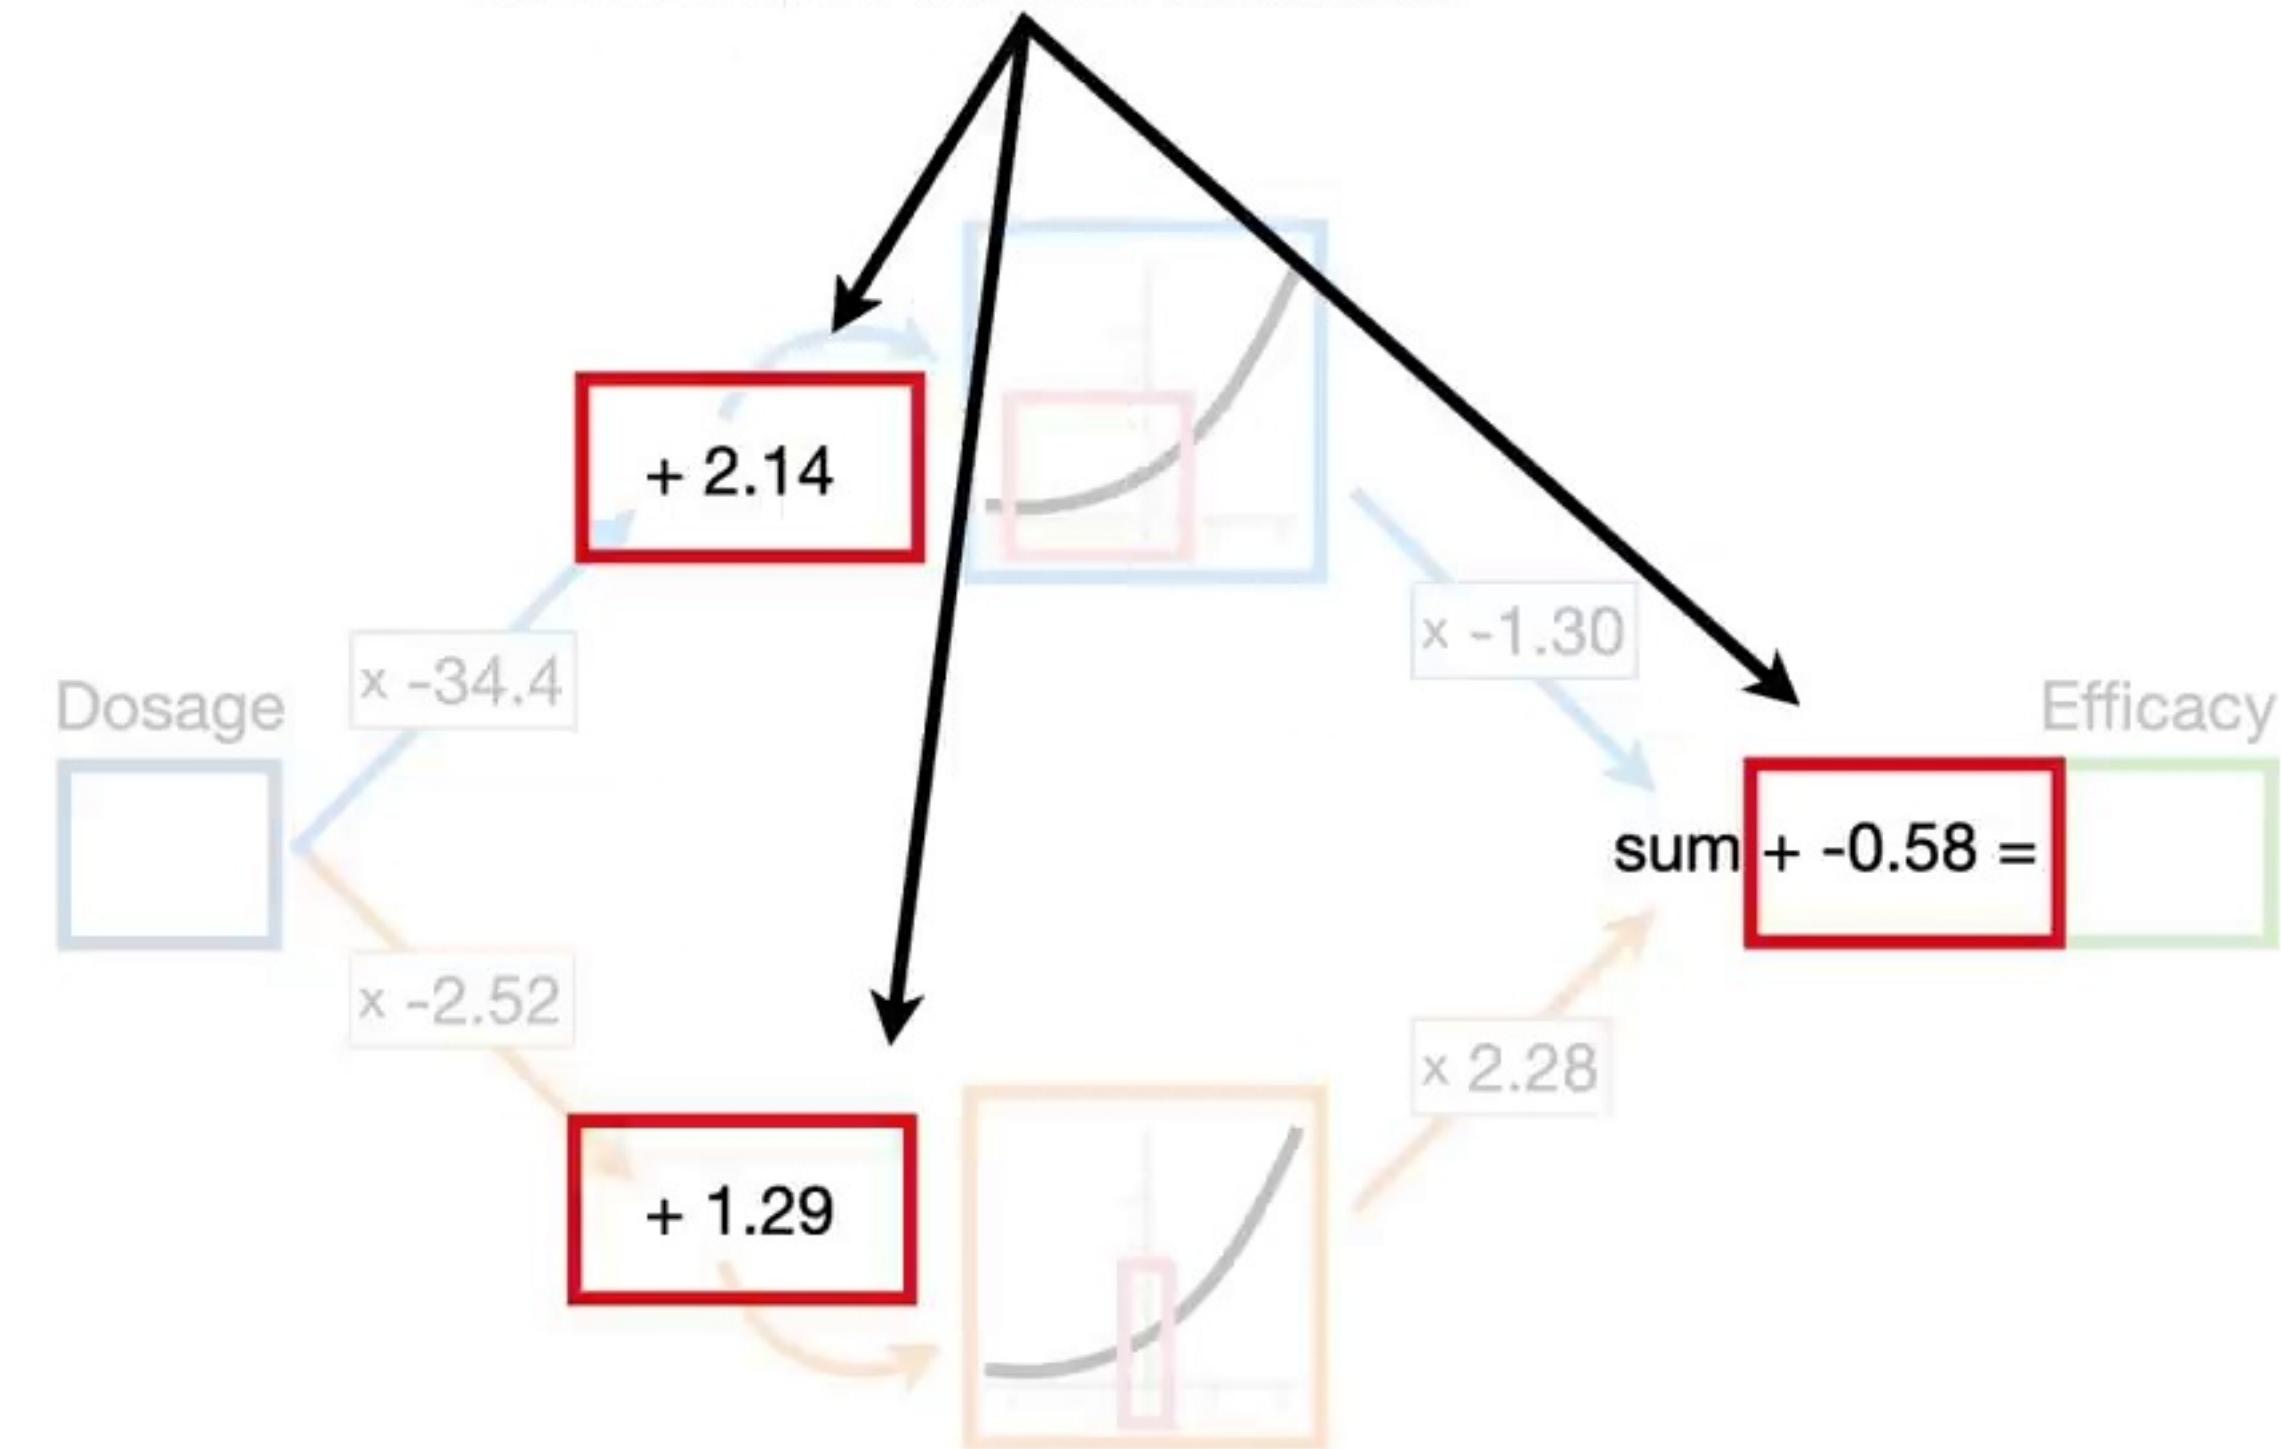


**NOTE:** Whether or not you call it a **Squiggle Fitting Machine**, the parameters that we multiply are called **weights**...



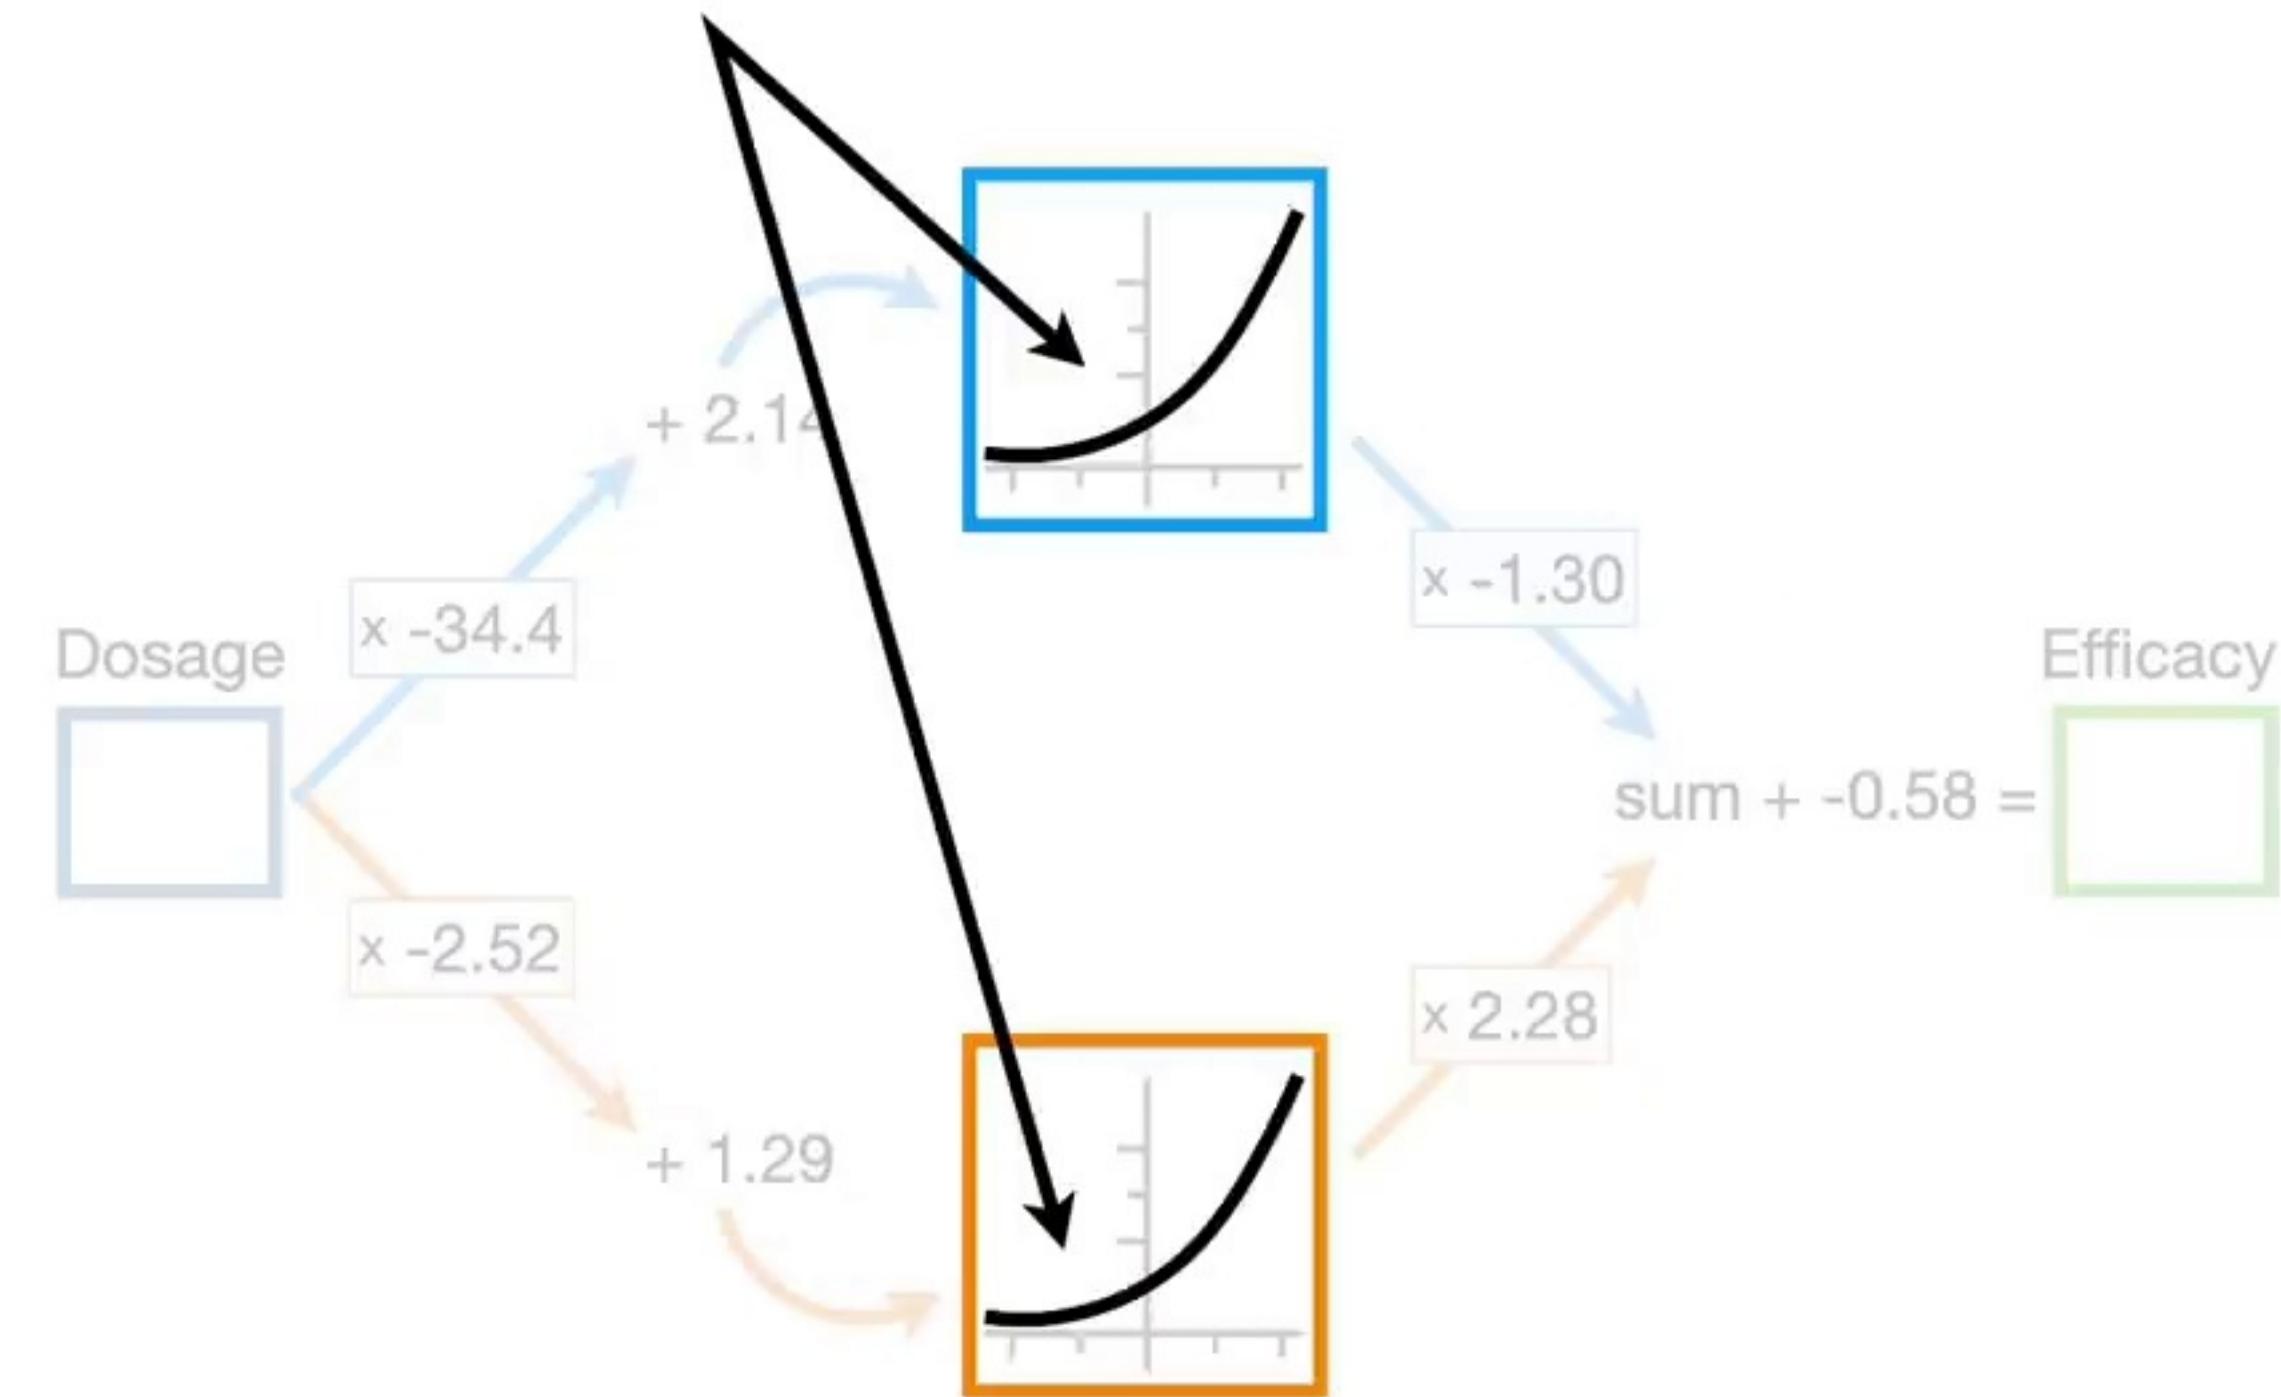


...and the parameters that we add are called **biases**.



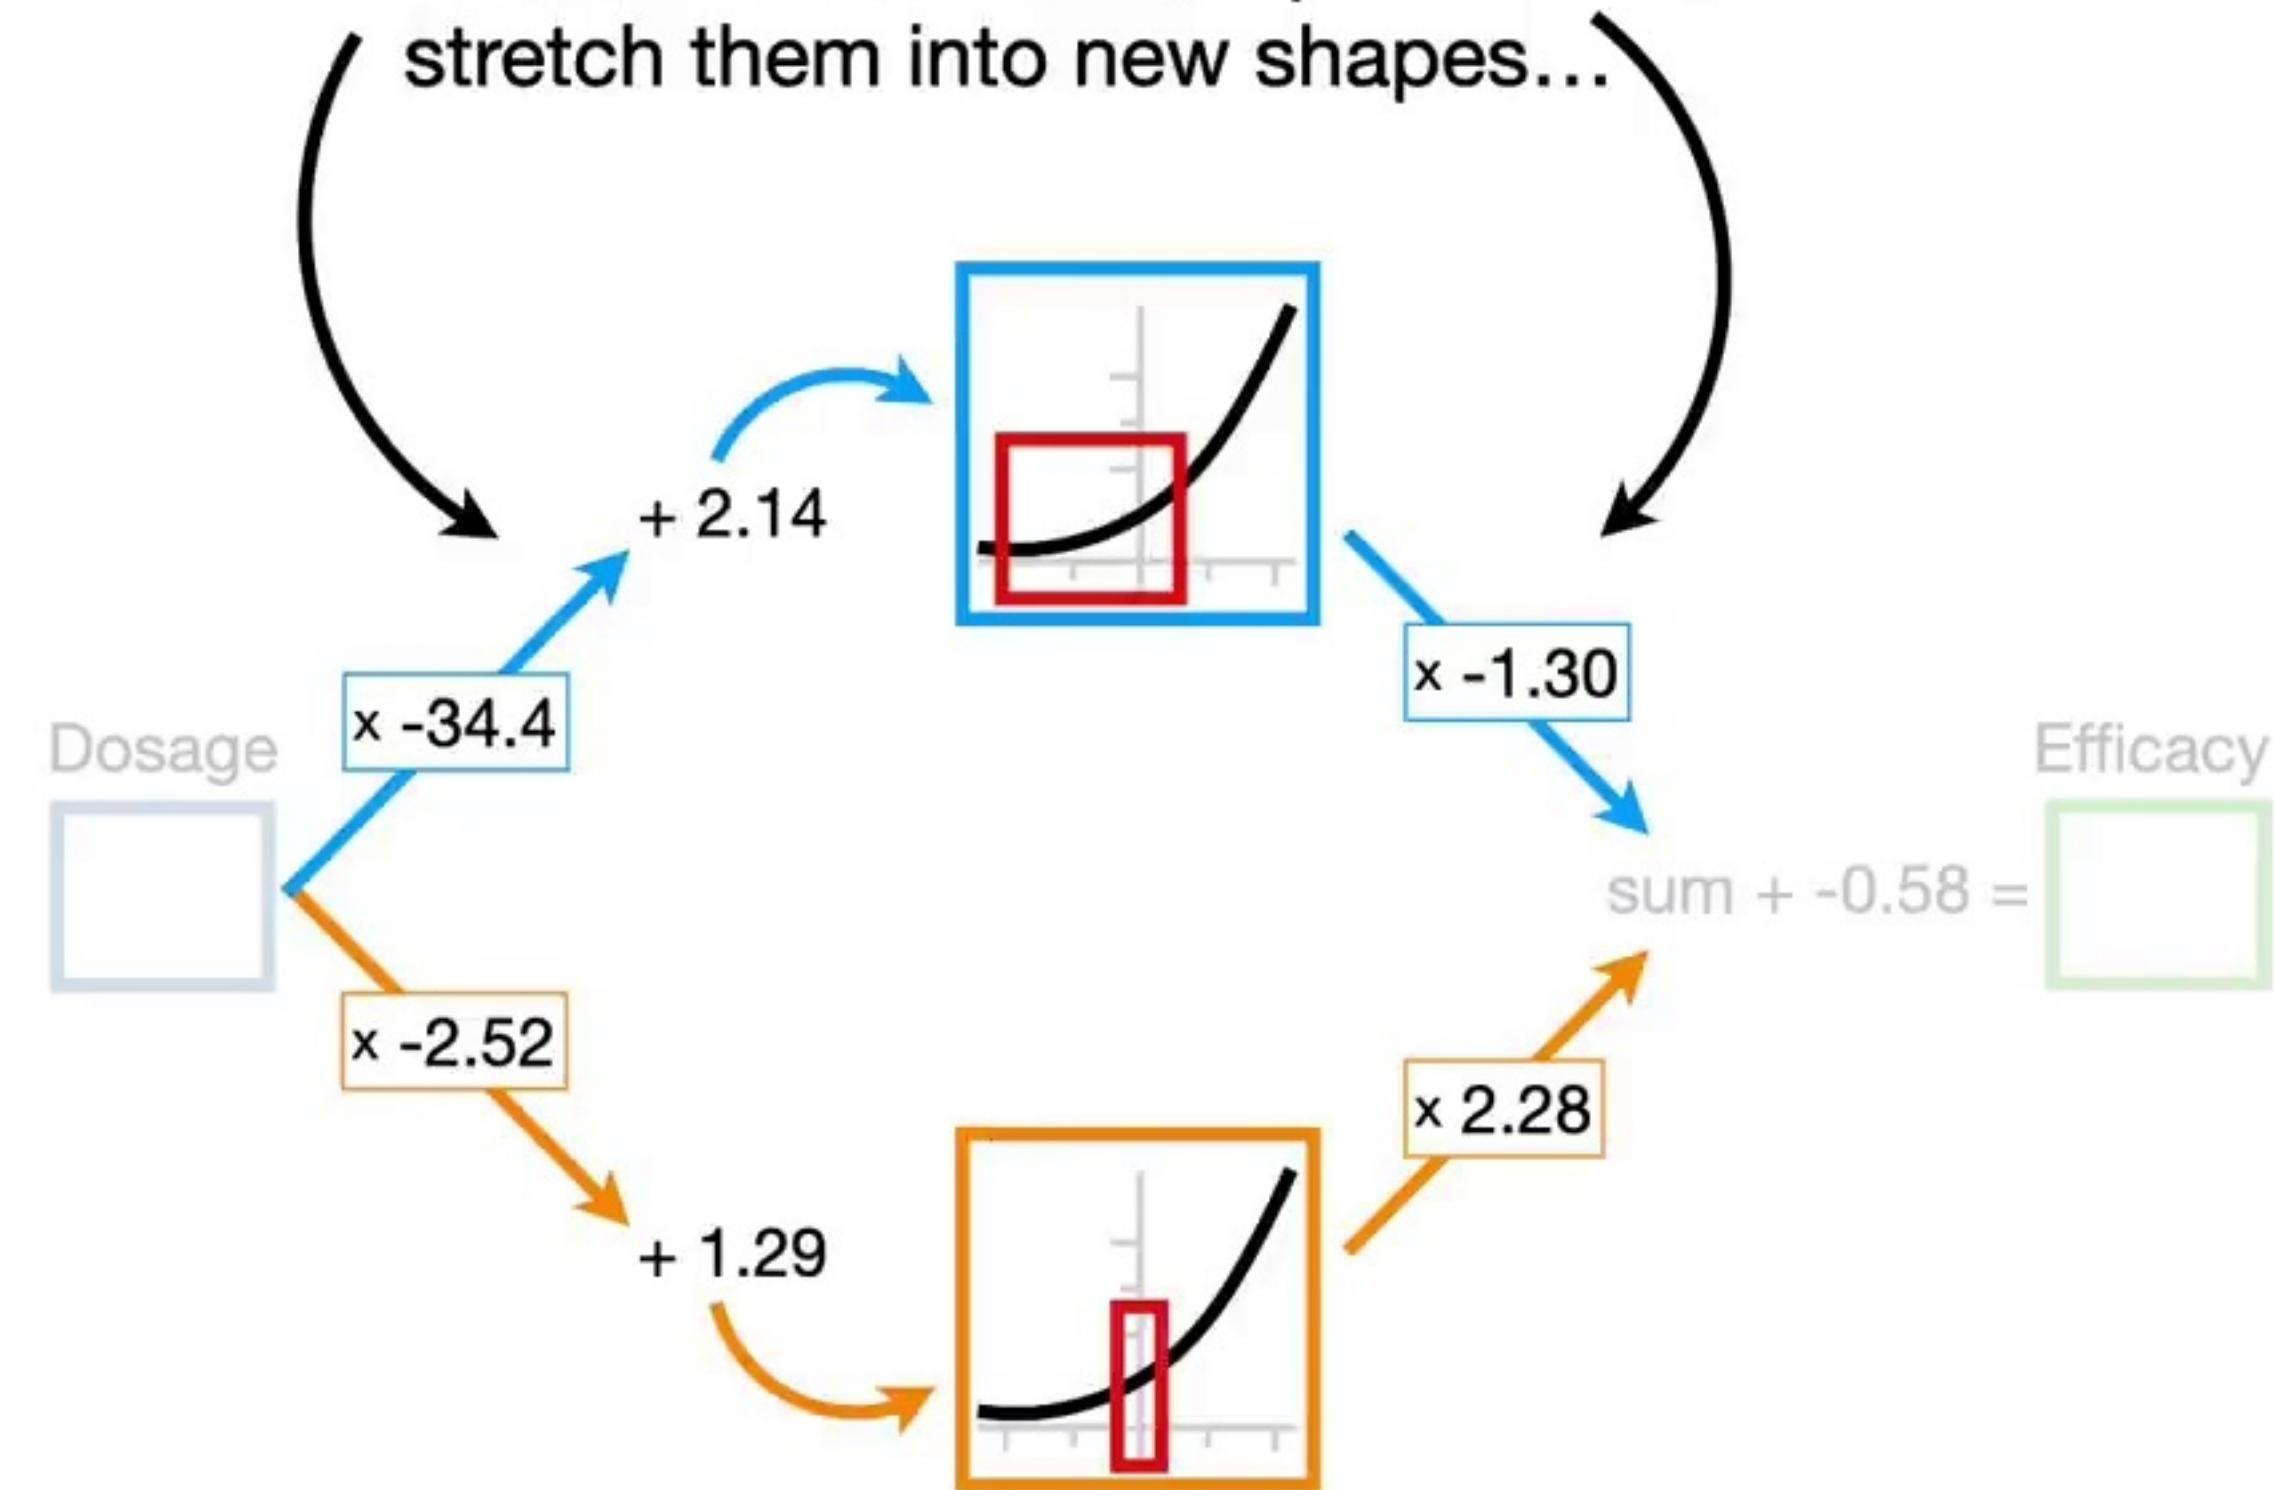


**NOTE: This Neural Network starts with two identical Activation Functions...**



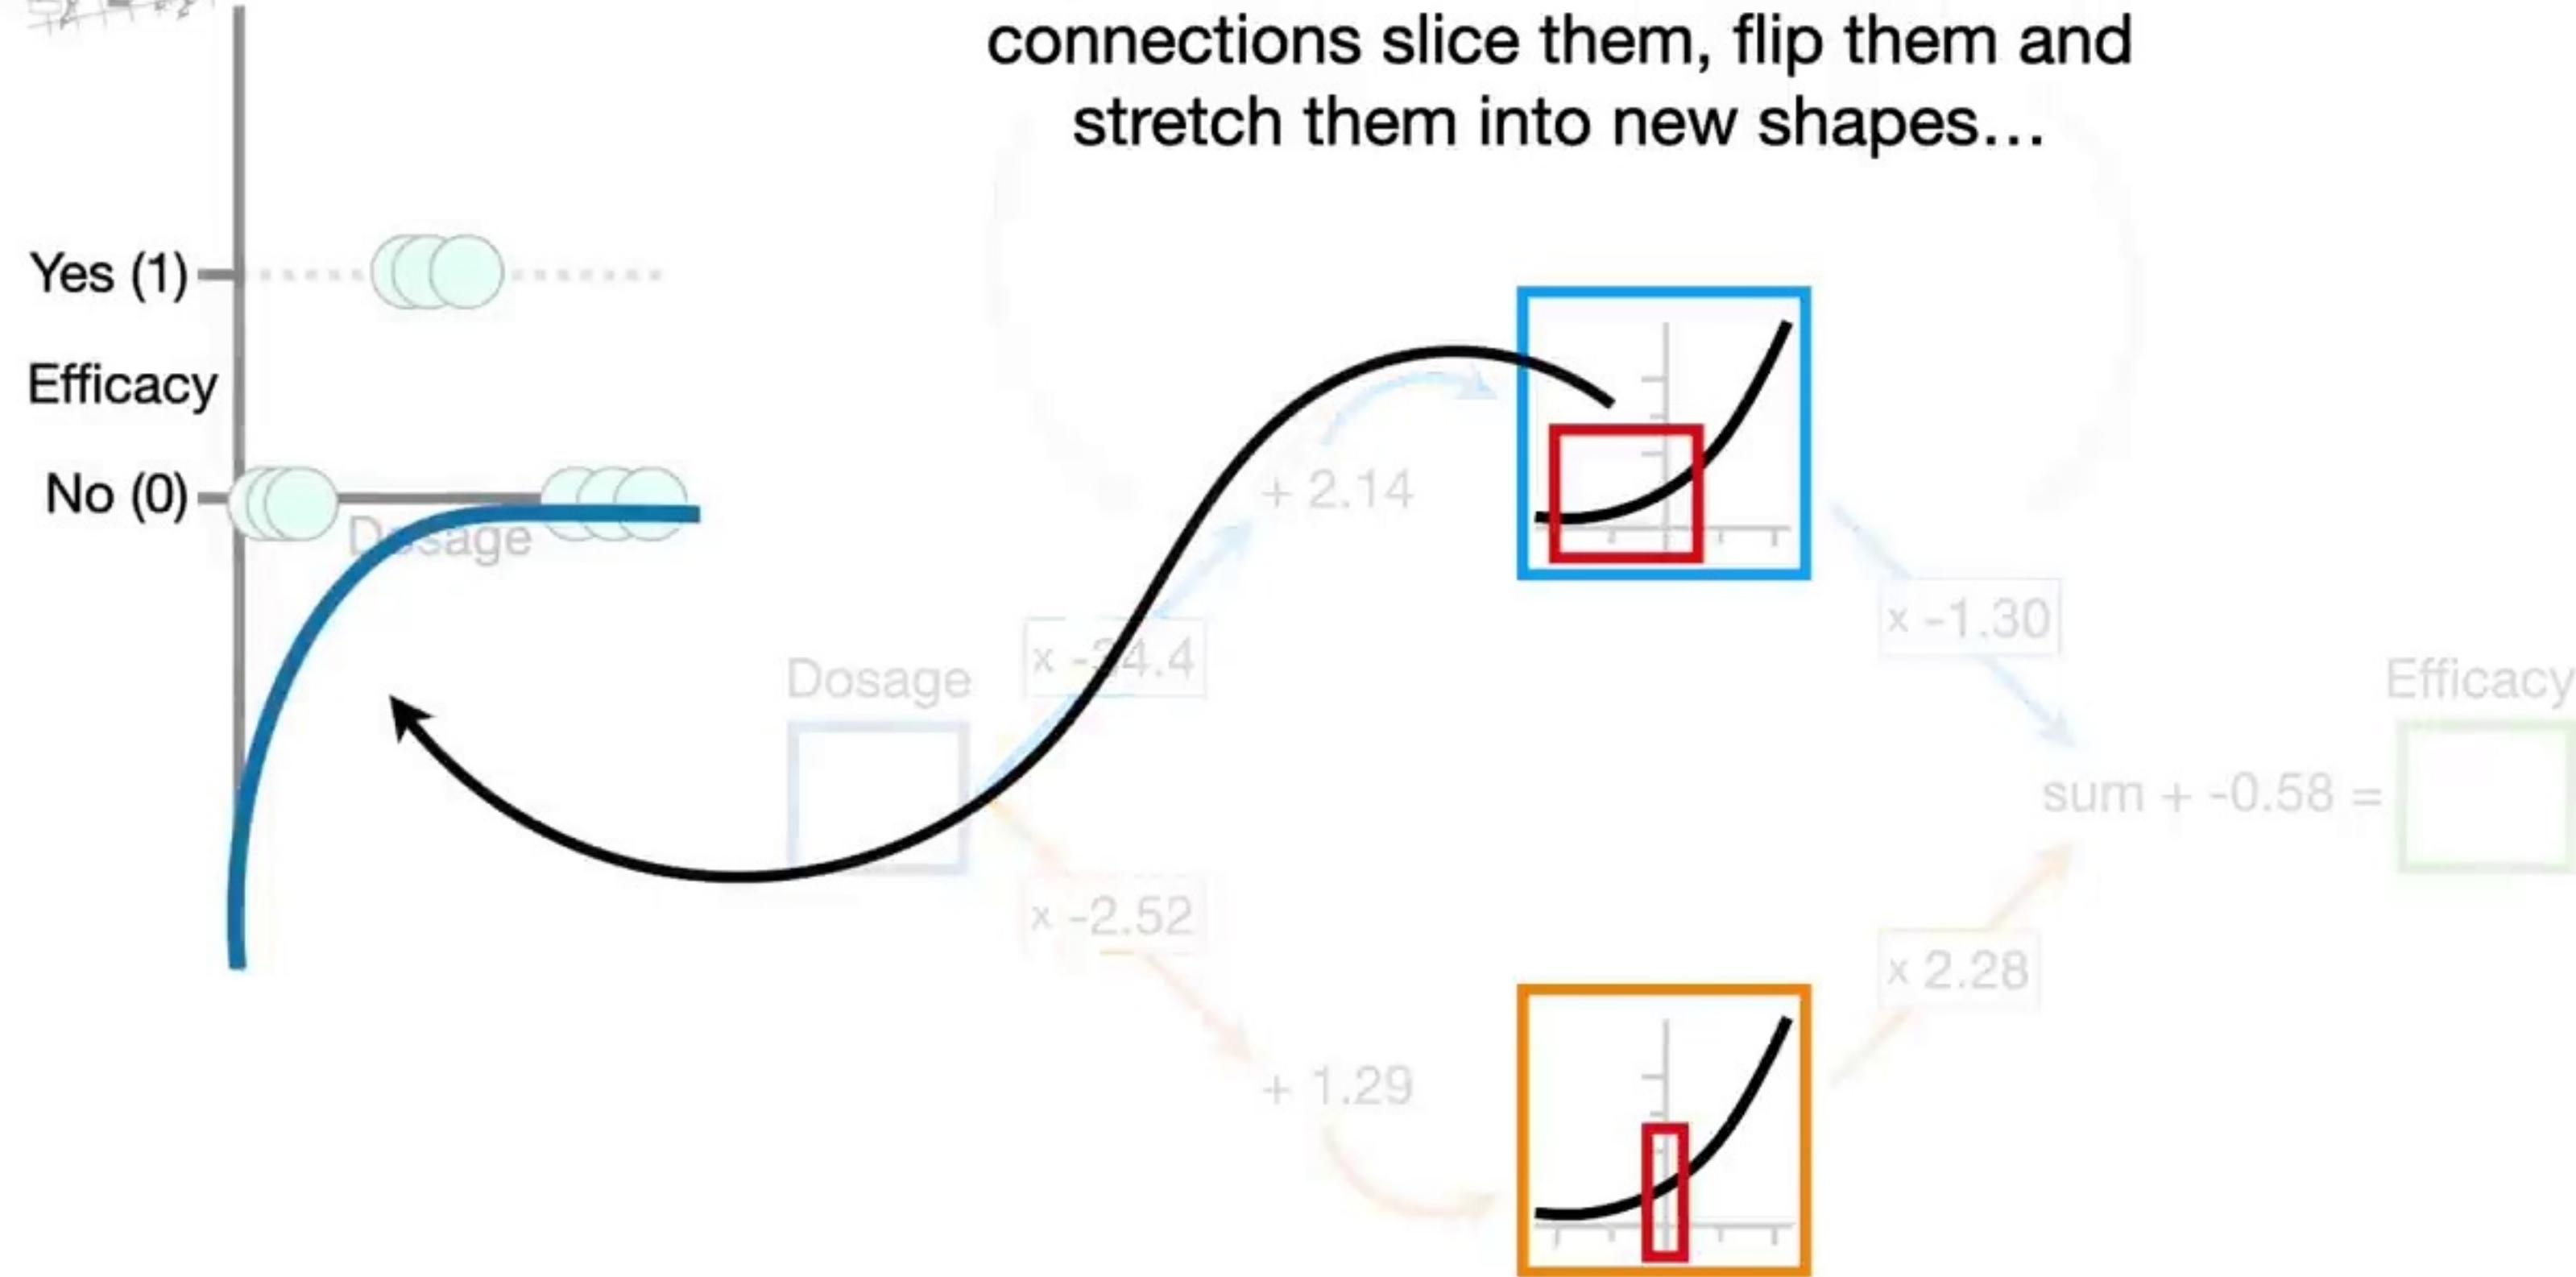


...but the **weights** and **biases** on the connections slice them, flip them and stretch them into new shapes...



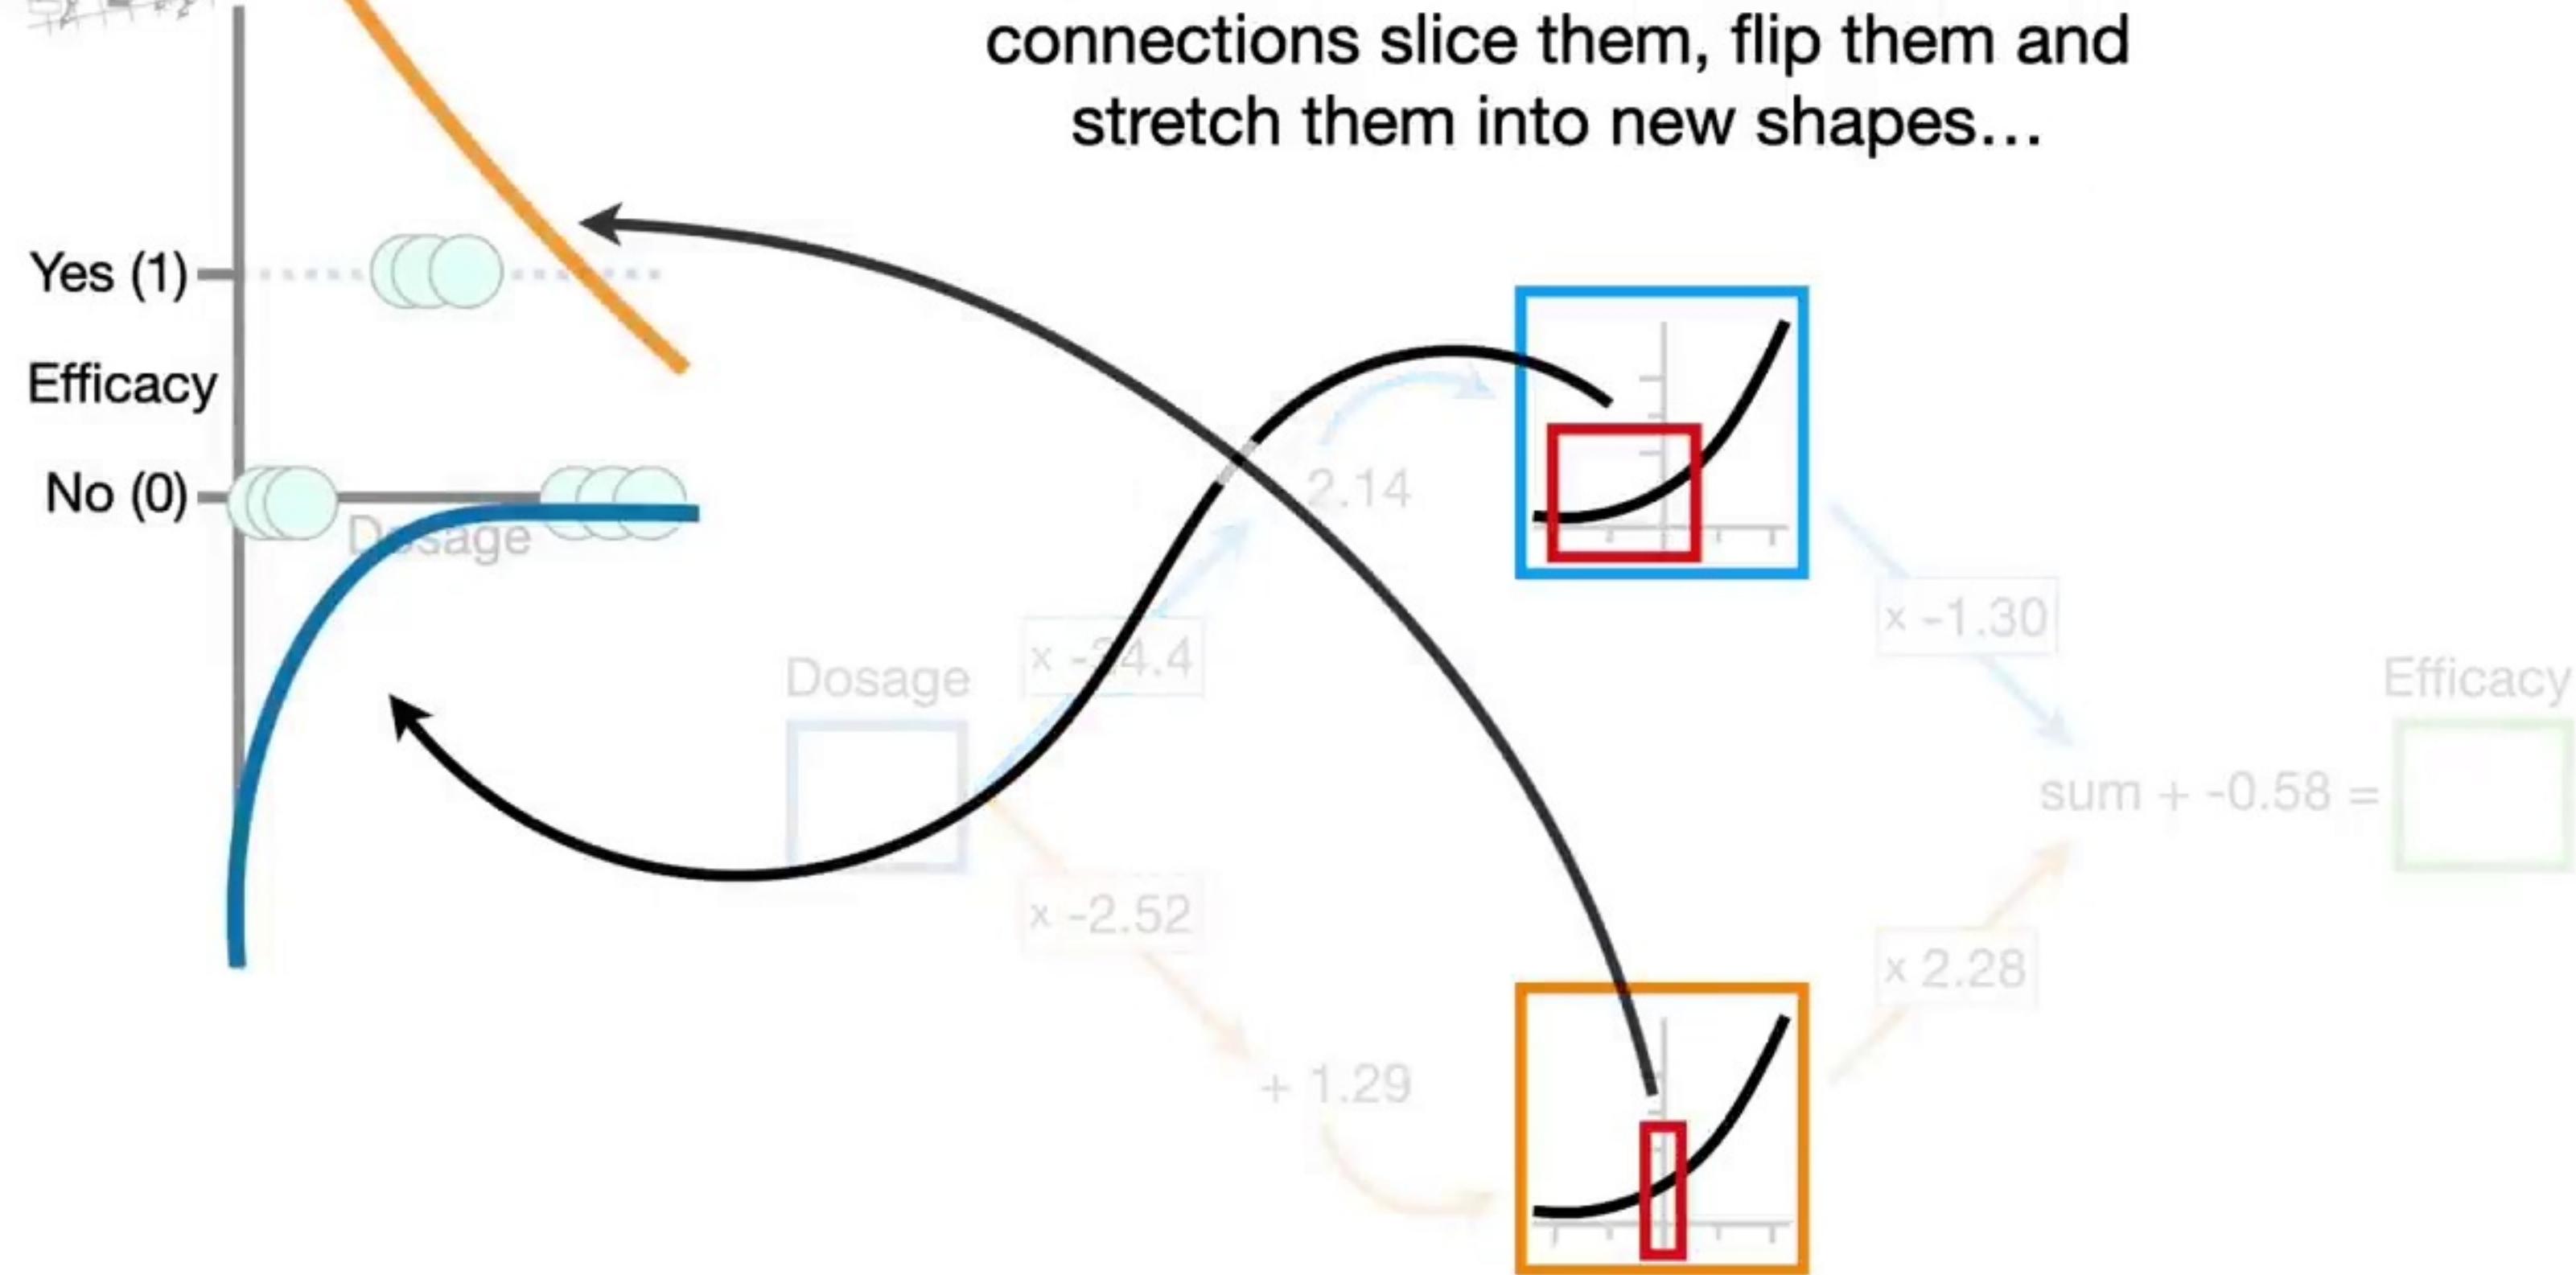


...but the **weights** and **biases** on the connections slice them, flip them and stretch them into new shapes...



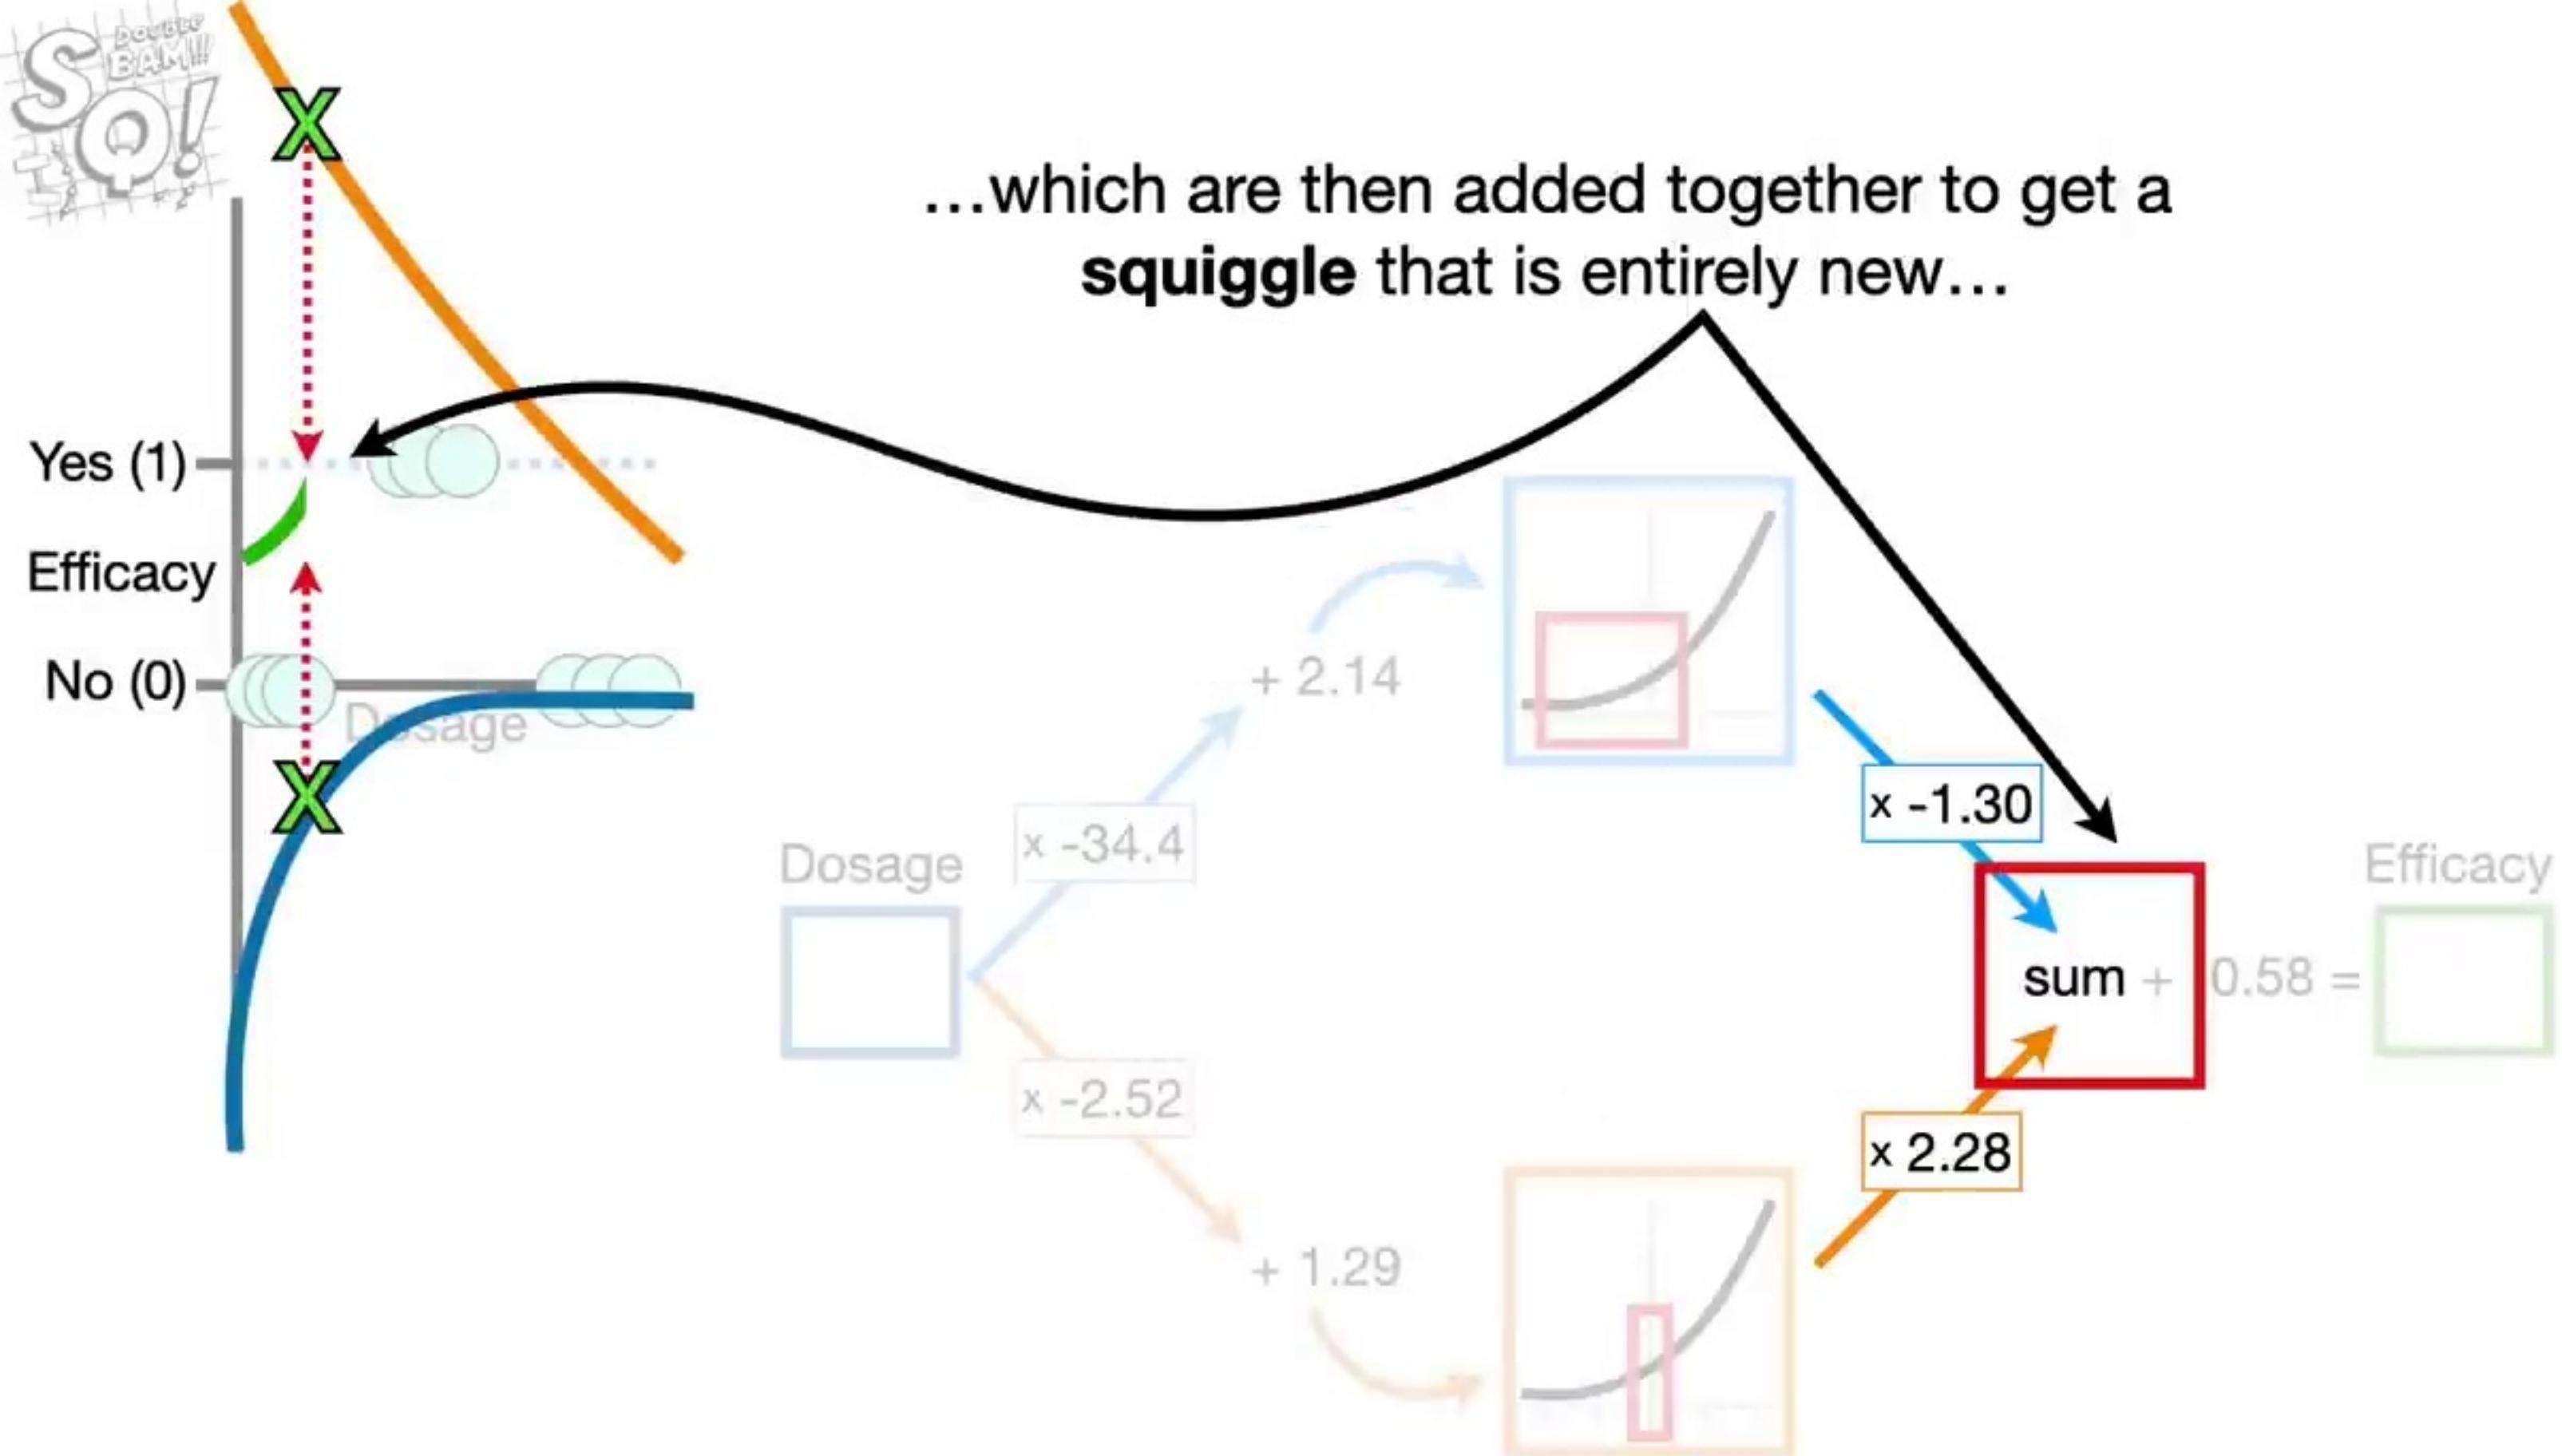
double  
BAM!!  
**SQ!**

...but the **weights** and **biases** on the connections slice them, flip them and stretch them into new shapes...



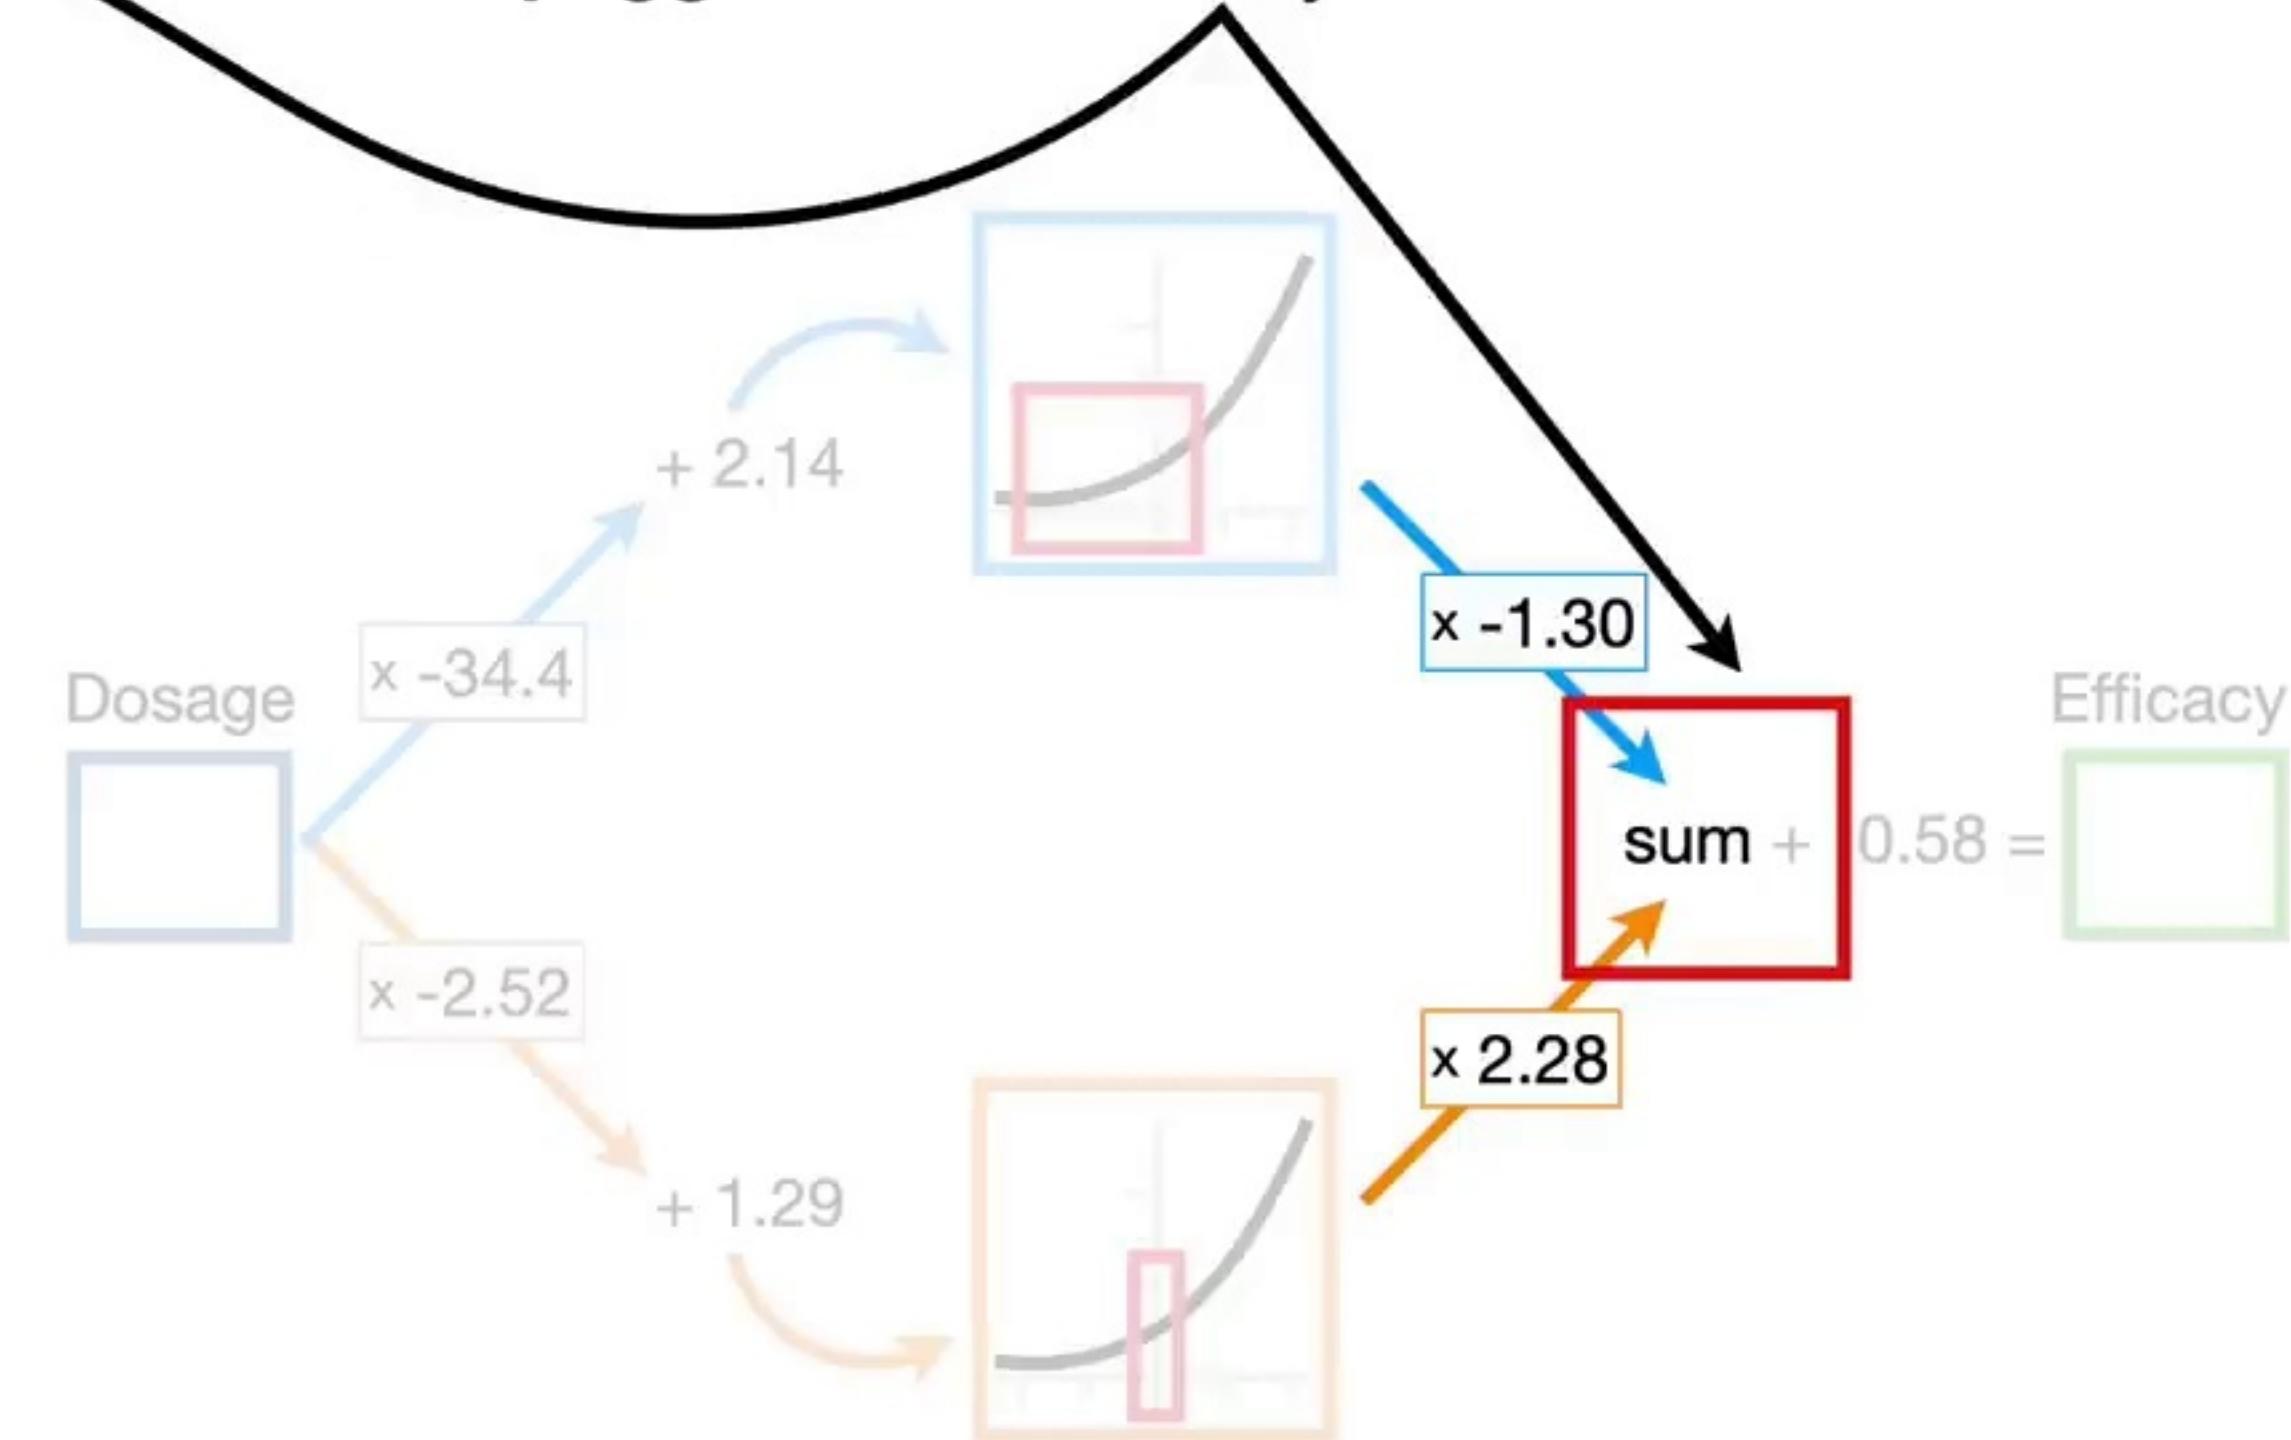
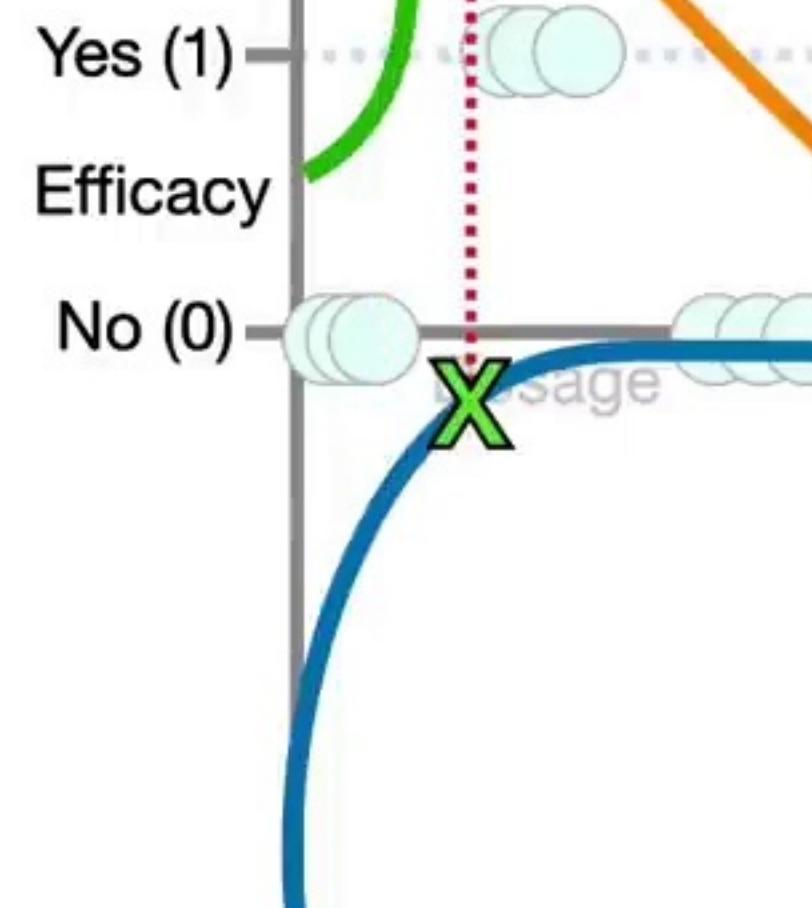
double  
BAM!!!  
**SQ!**

...which are then added together to get a **squiggle** that is entirely new...



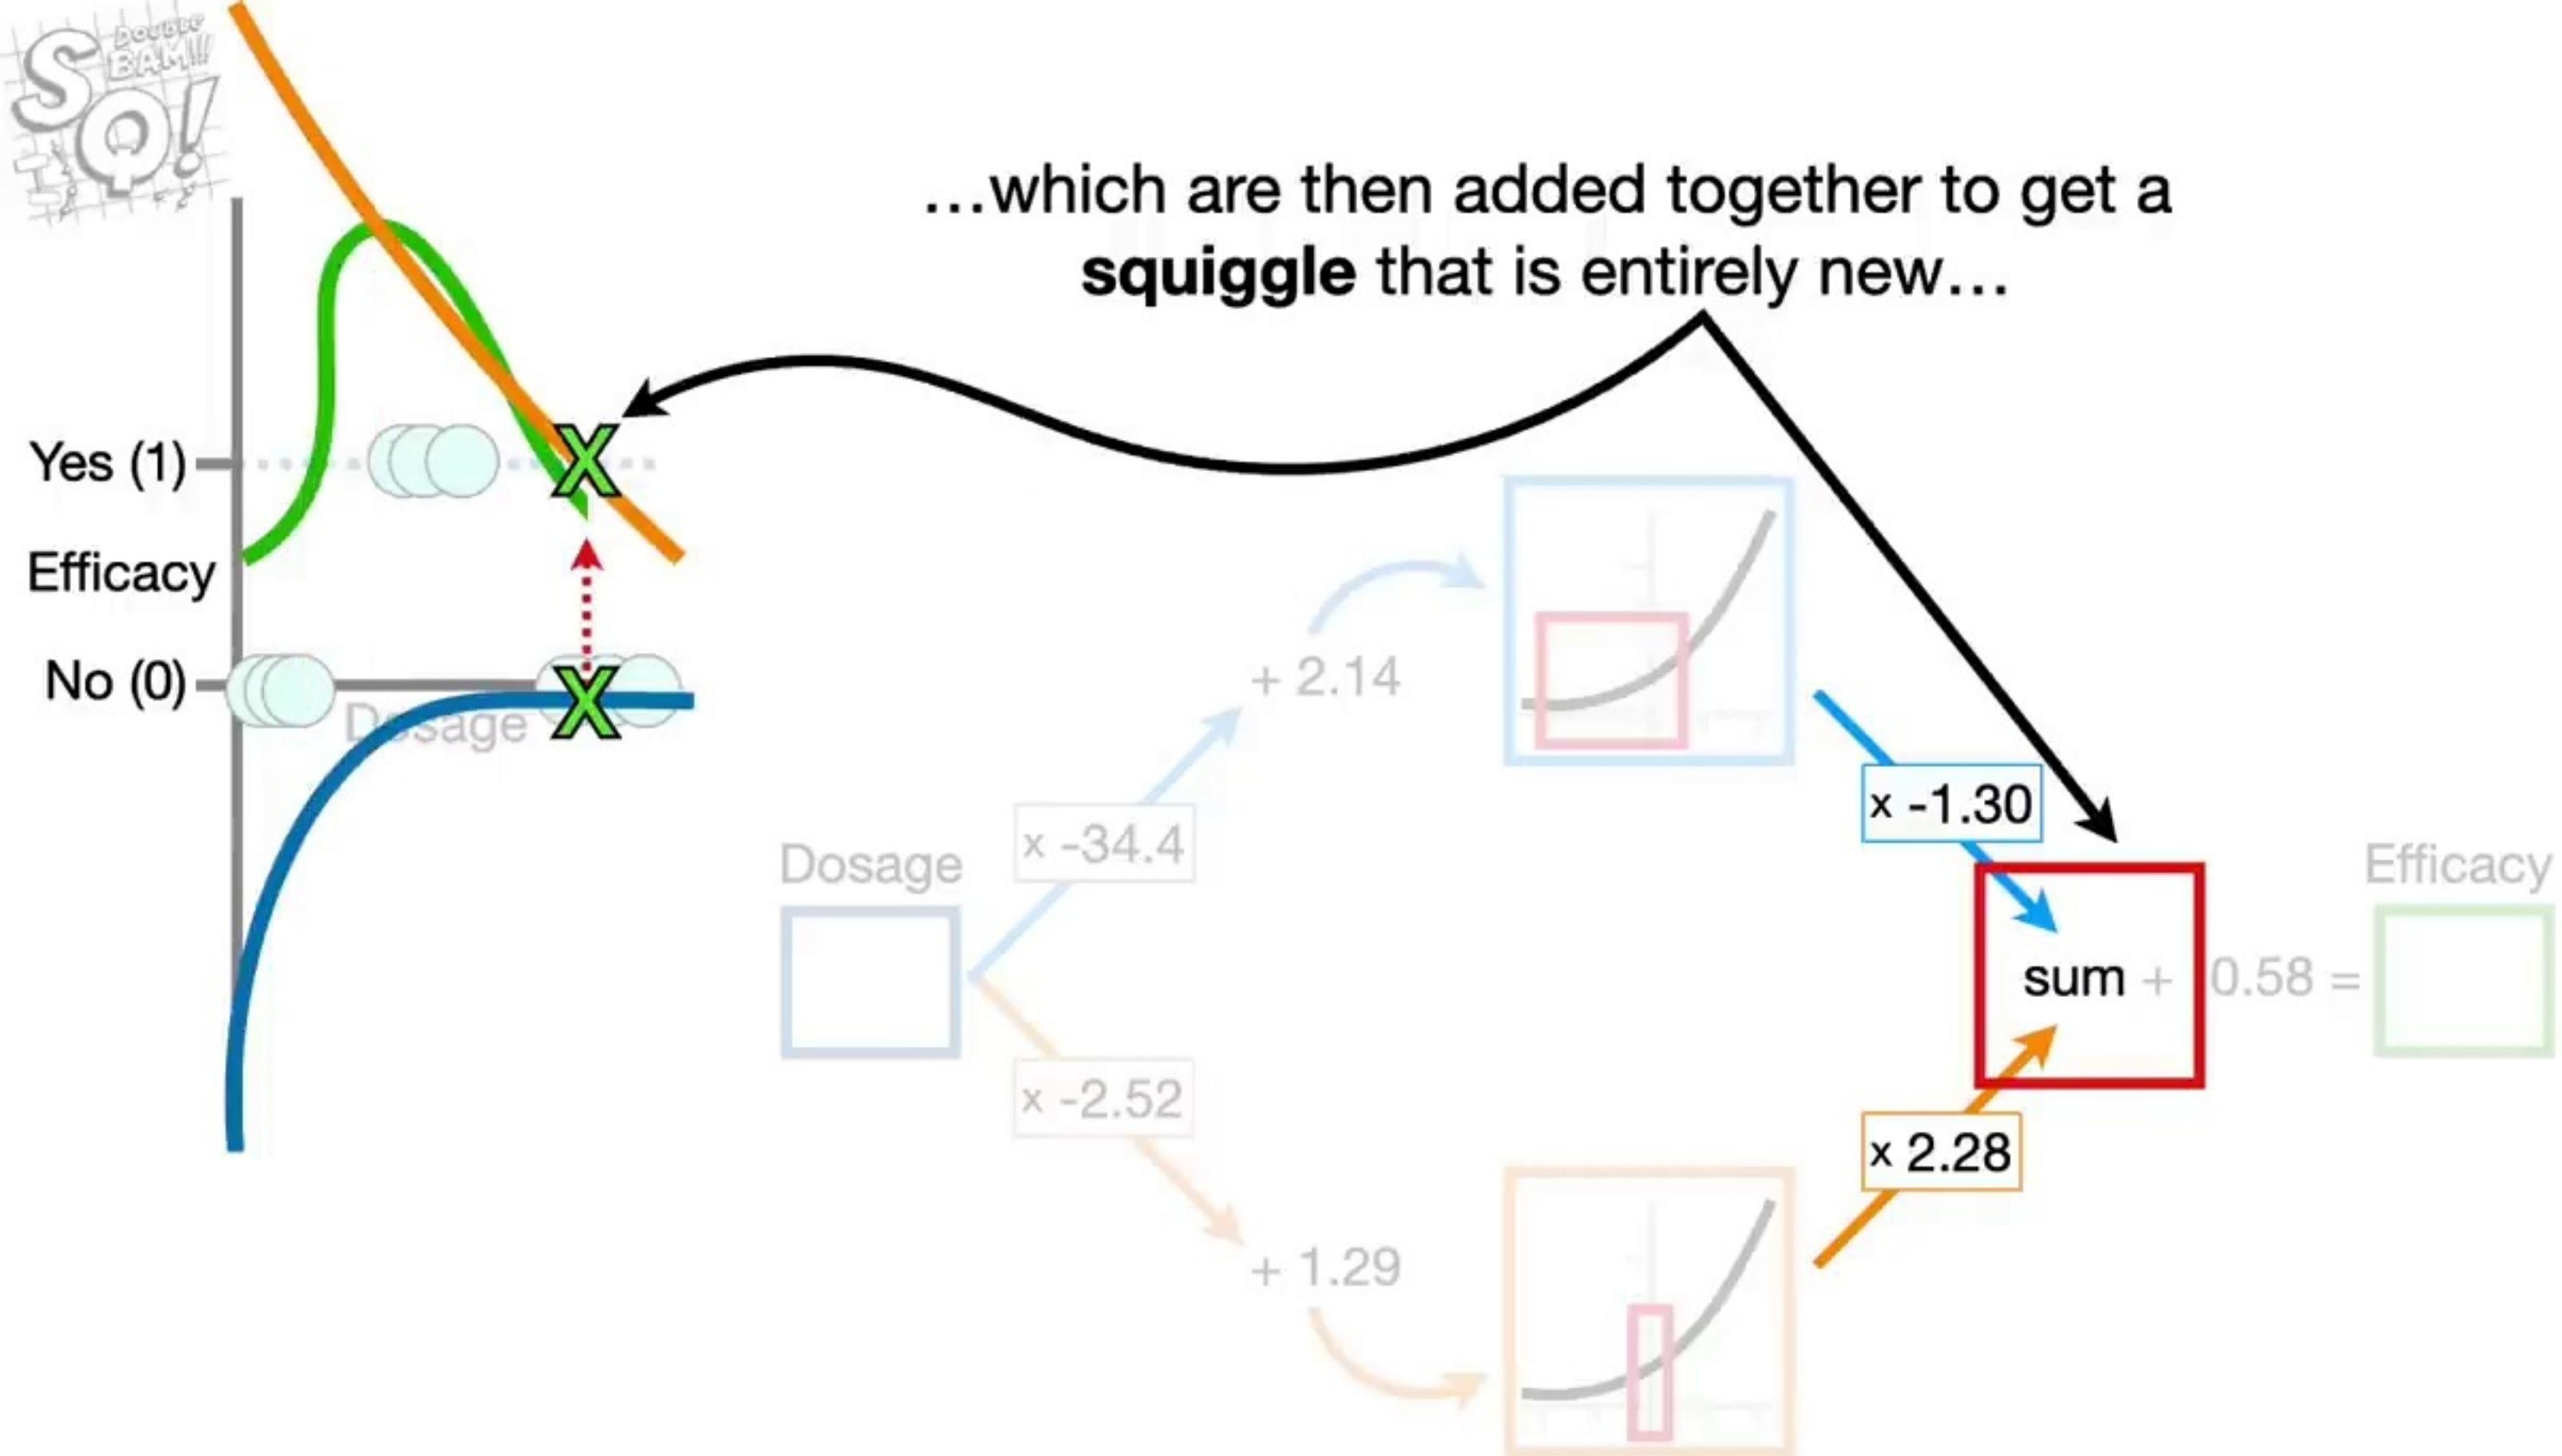
Double  
BAM!!  
**SQ!**

...which are then added together to get a **squiggle** that is entirely new...



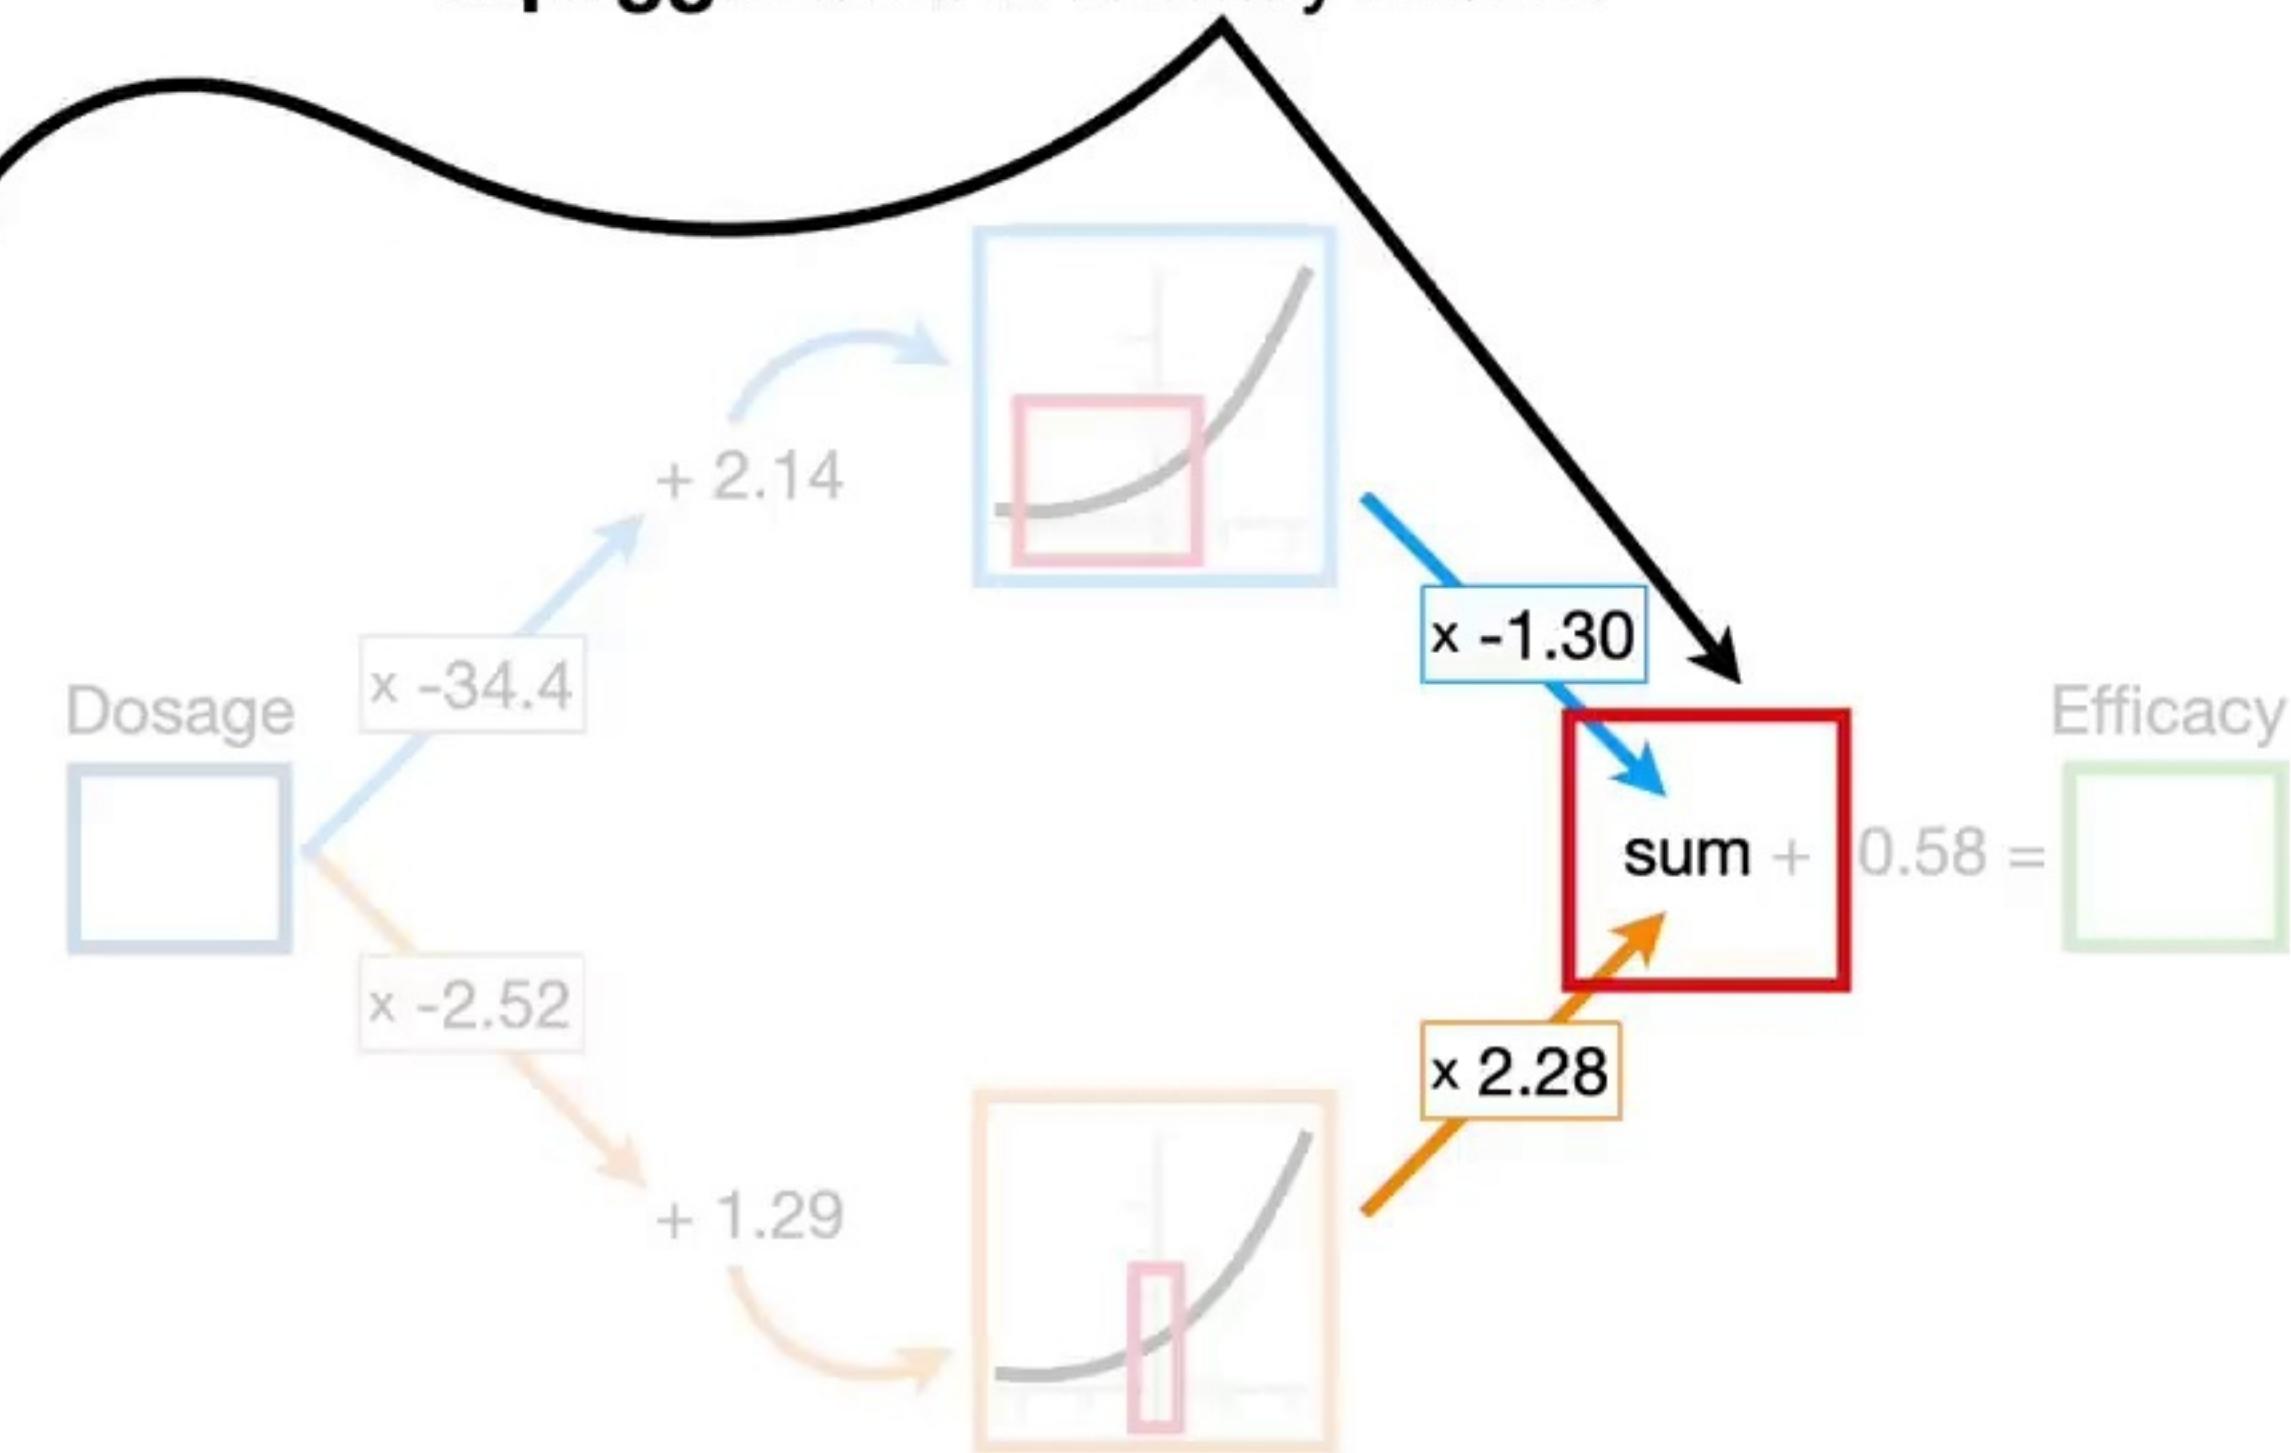
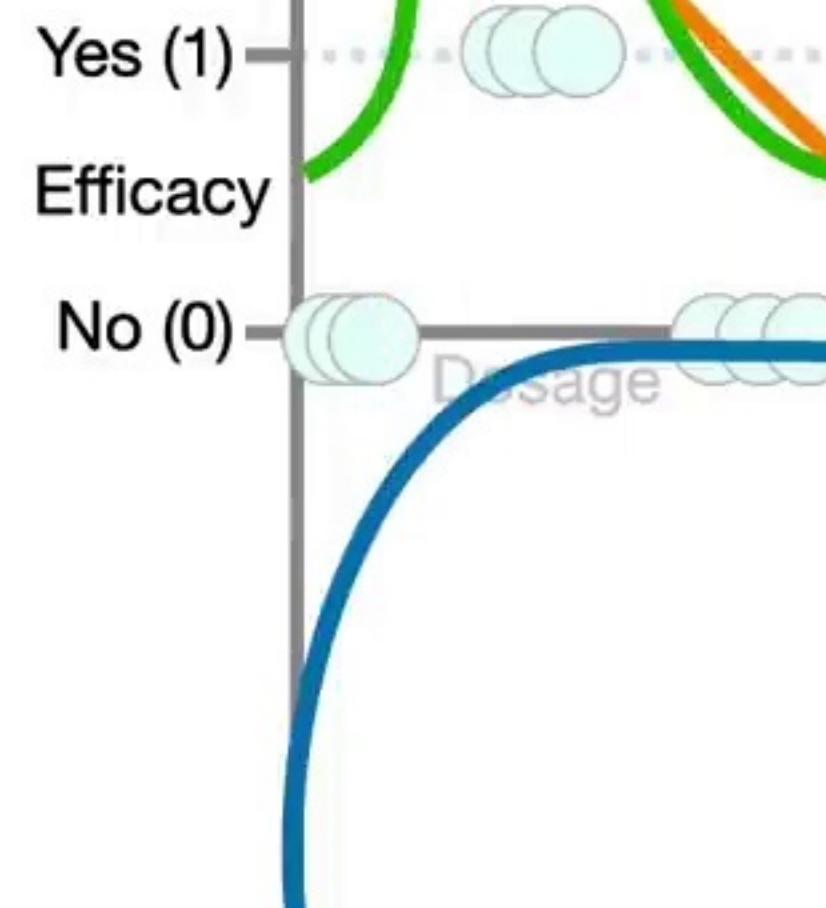
Double  
BAM!!  
**SQ!**

...which are then added together to get a **squiggle** that is entirely new...



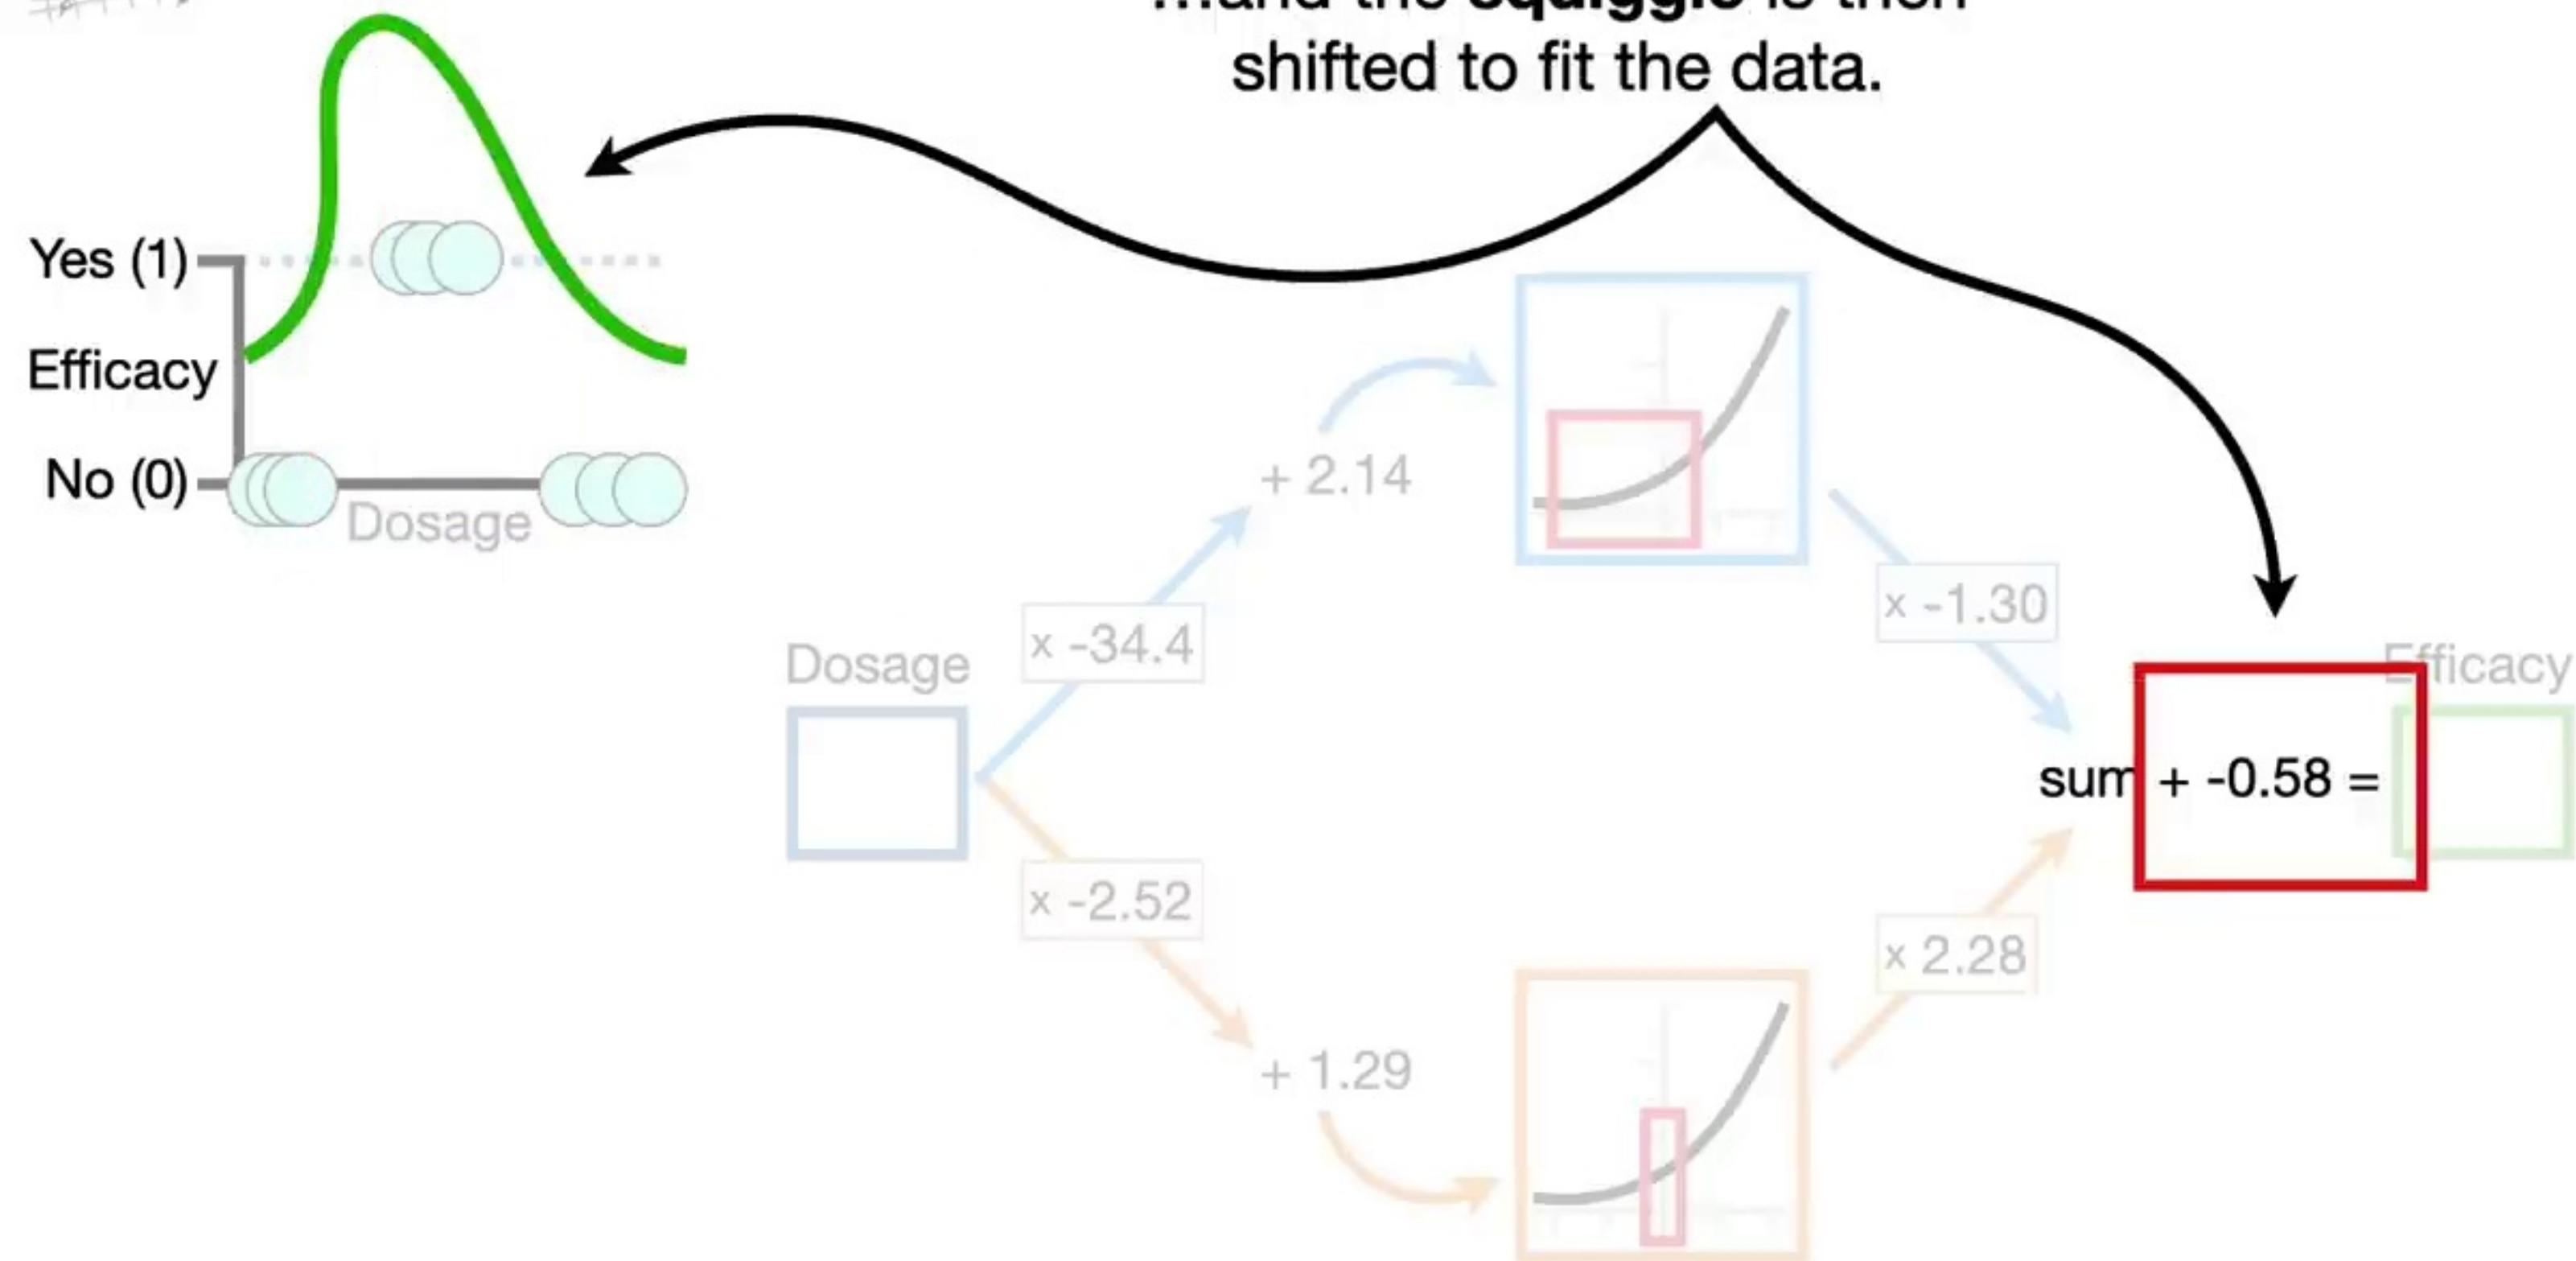
Double  
BAM!!  
**SQ!**

...which are then added together to get a **squiggle** that is entirely new...



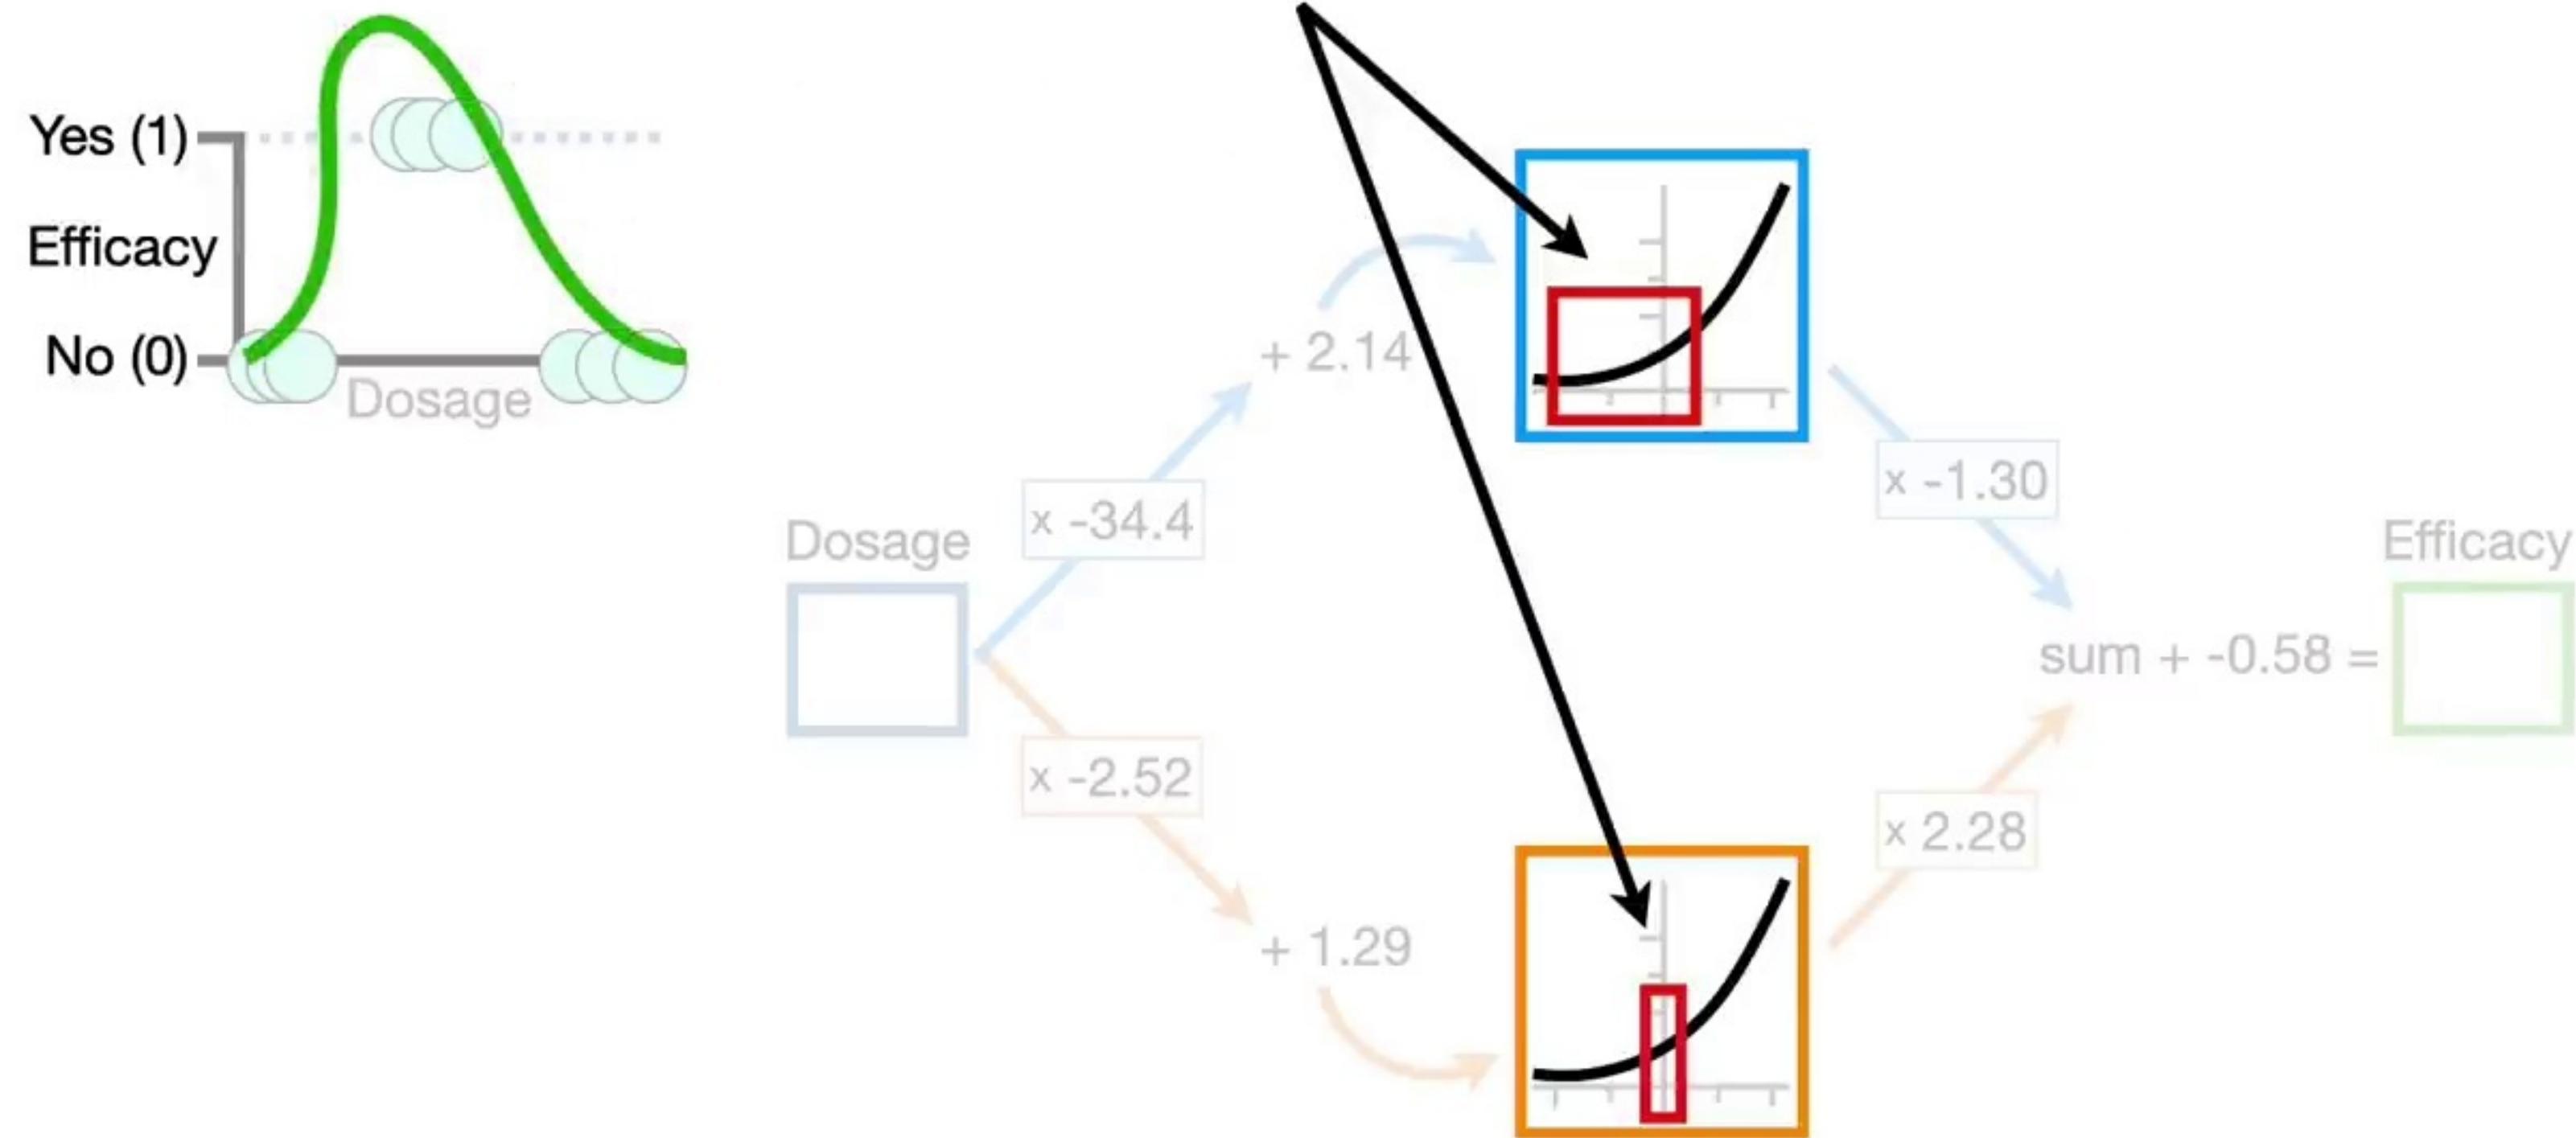


...and the **squiggle** is then shifted to fit the data.



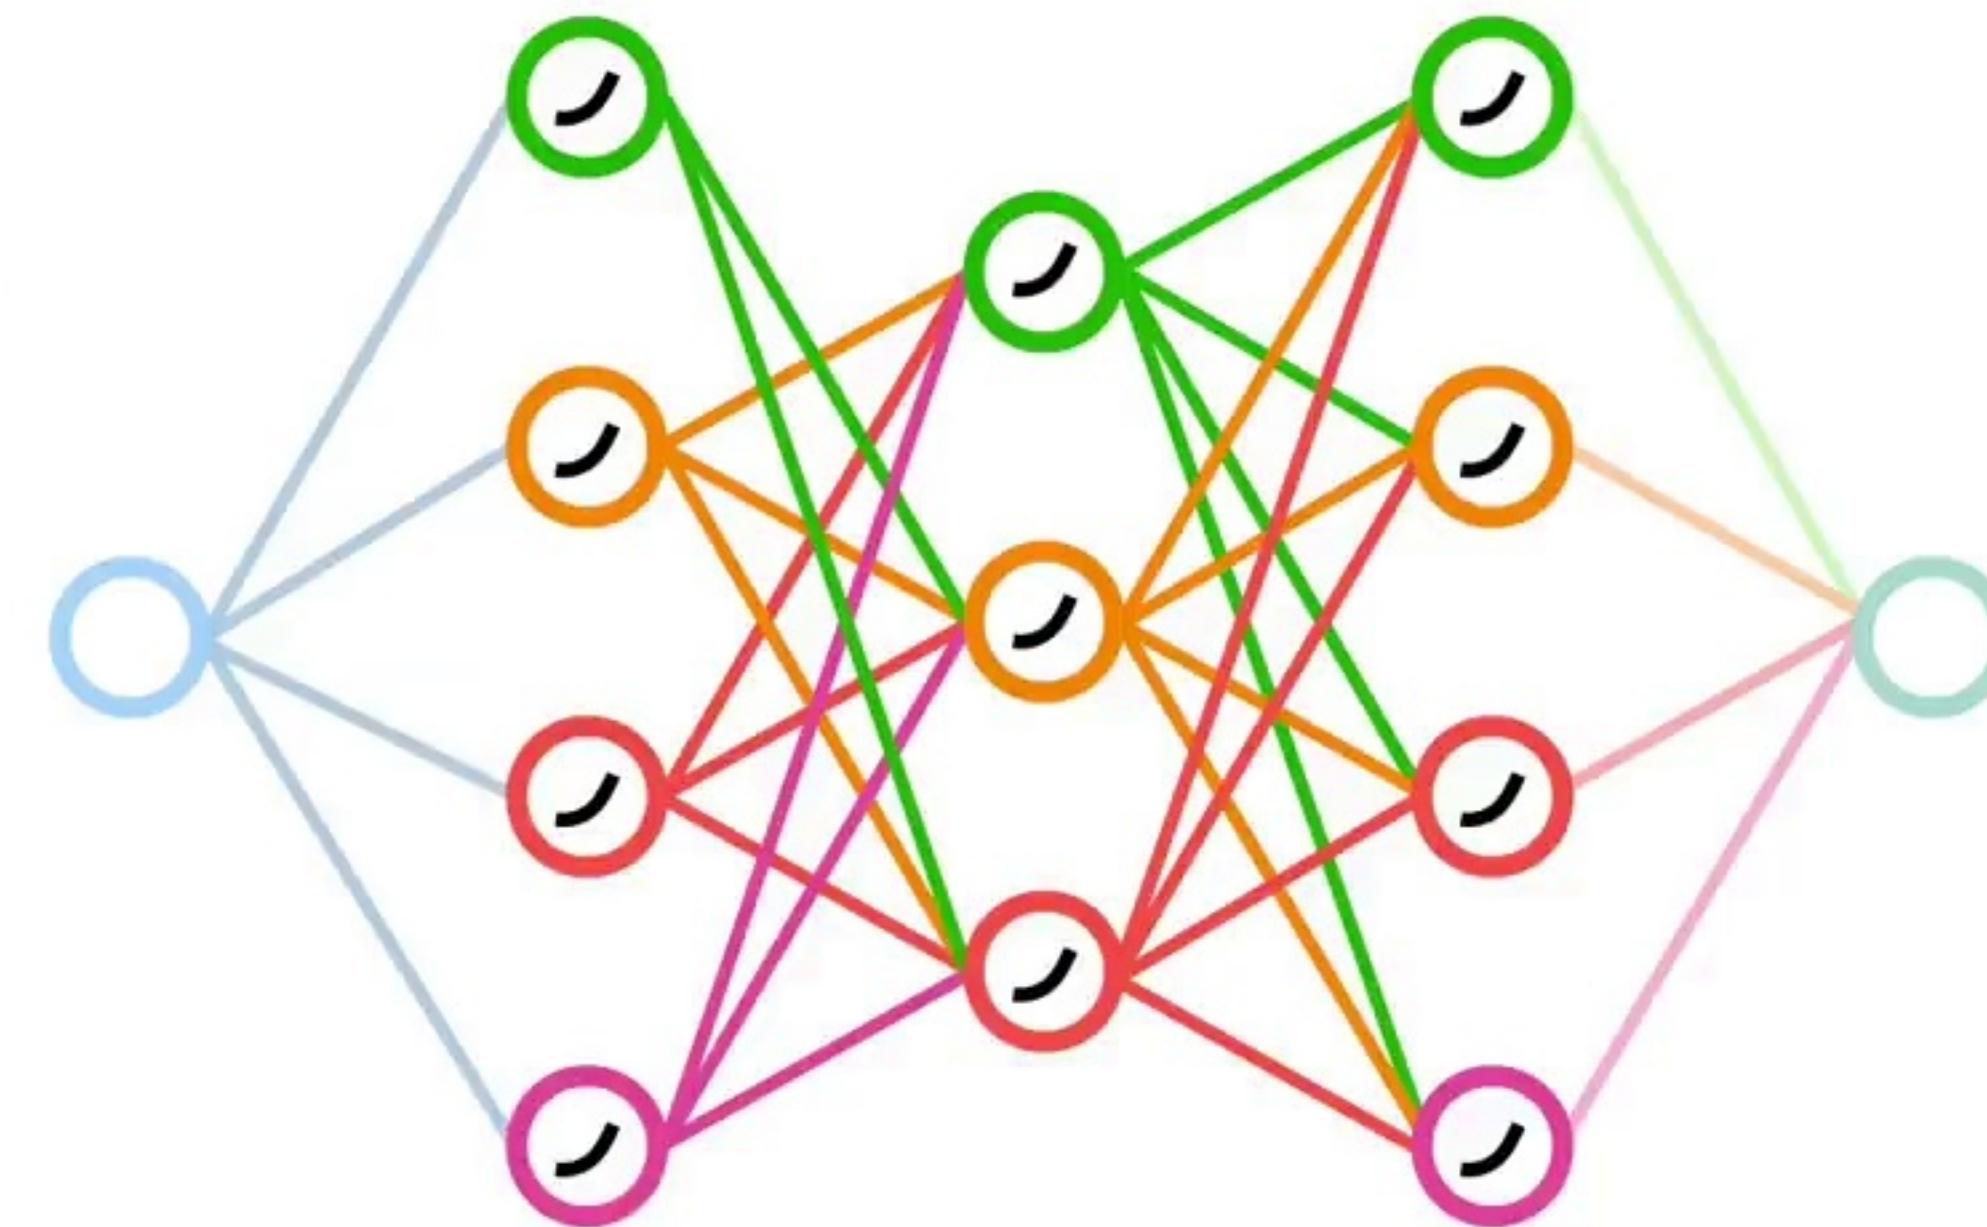
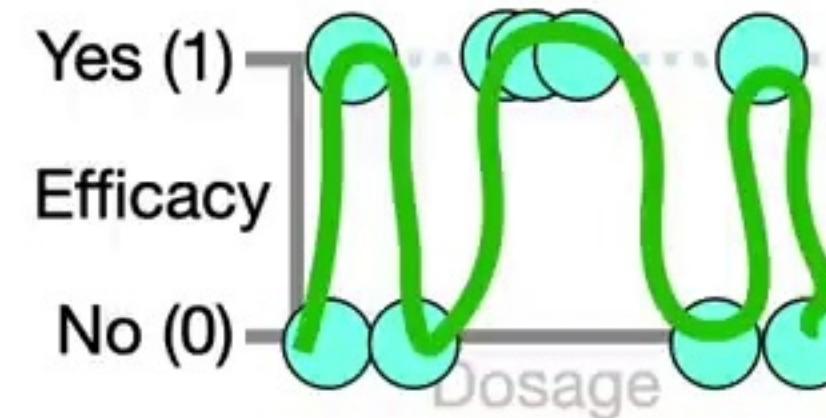


Now, if we can create this **green squiggle** with just two **Nodes** in a single **Hidden Layer**...



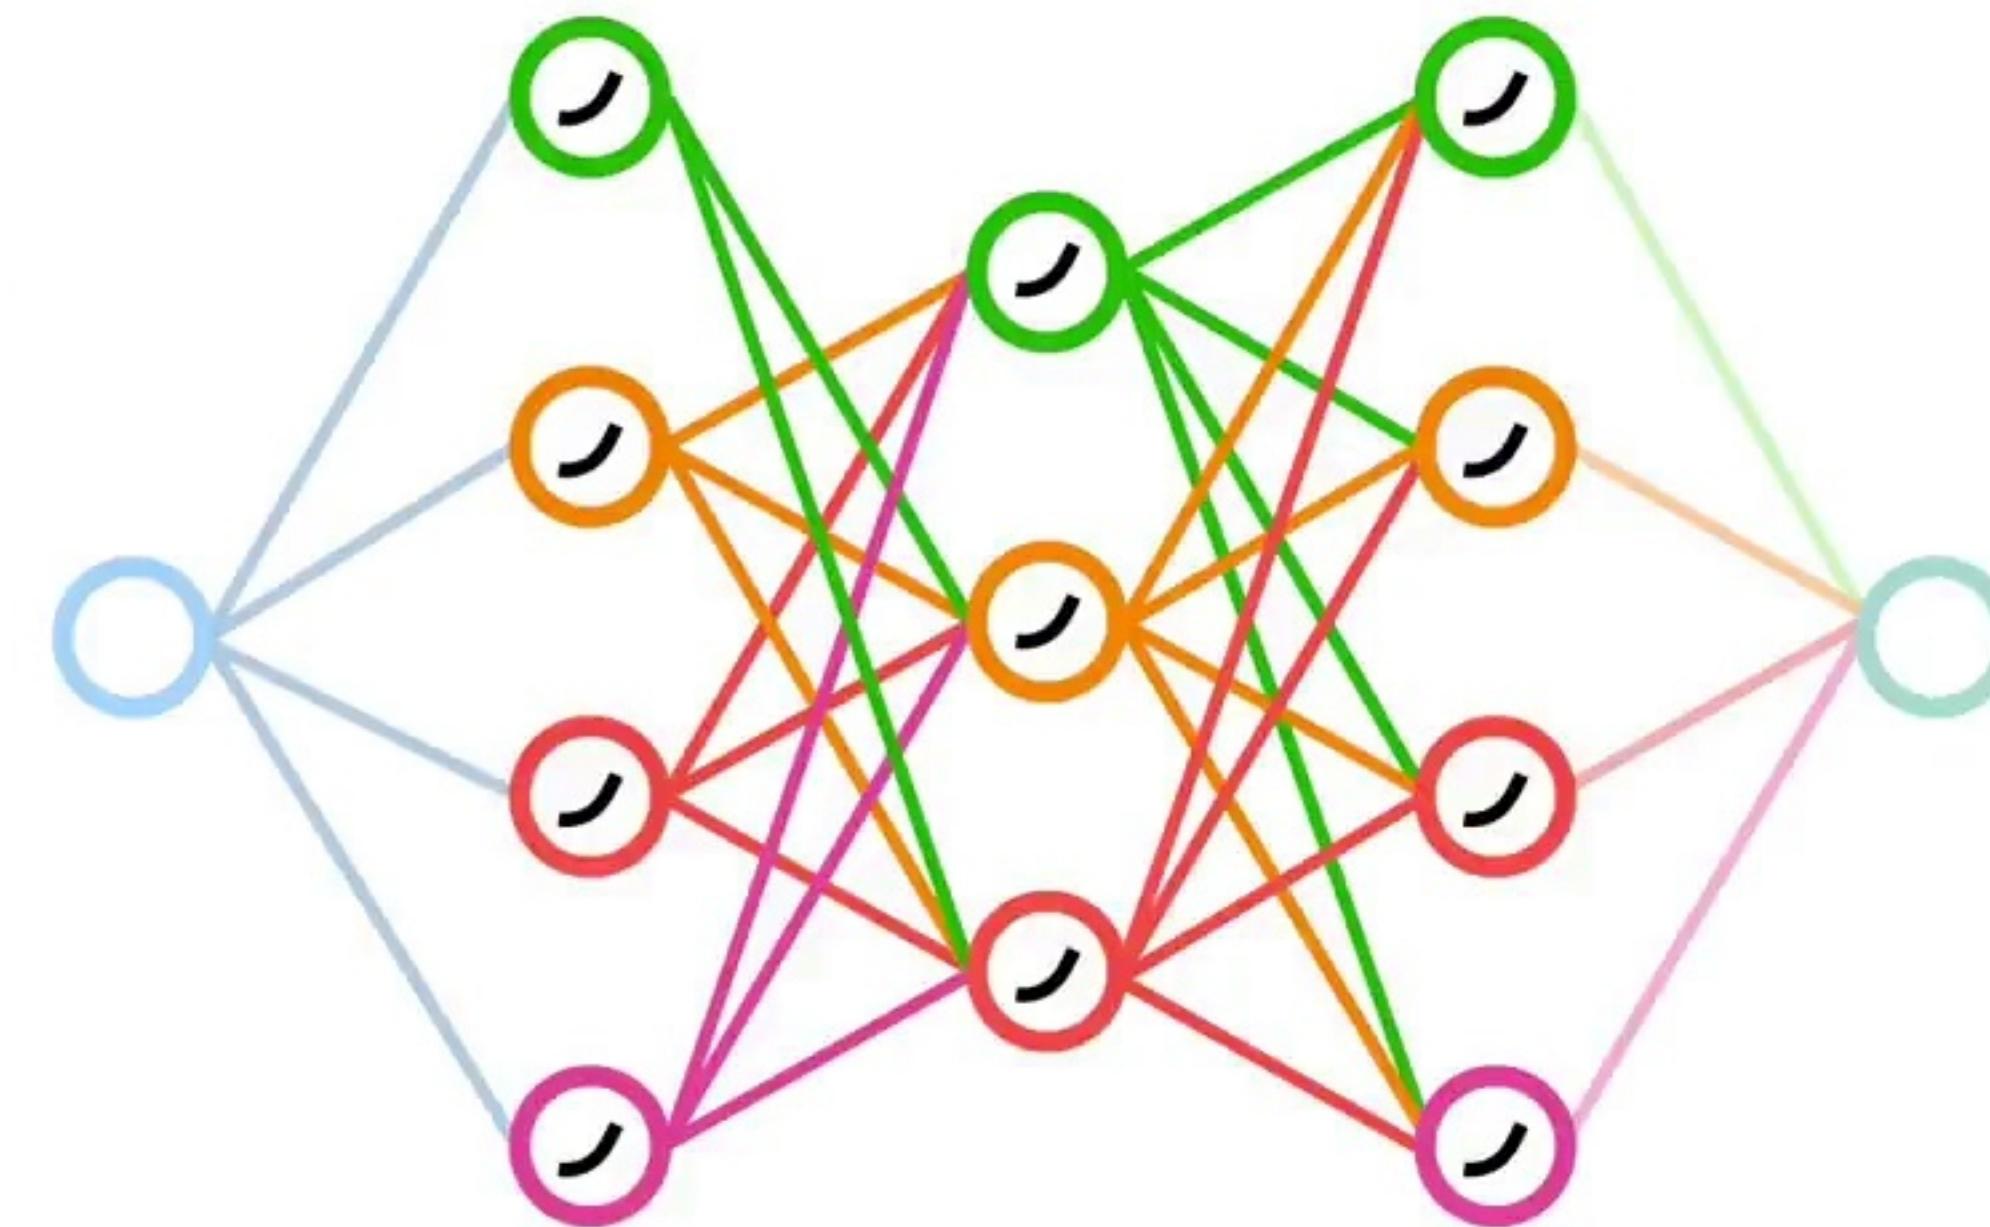
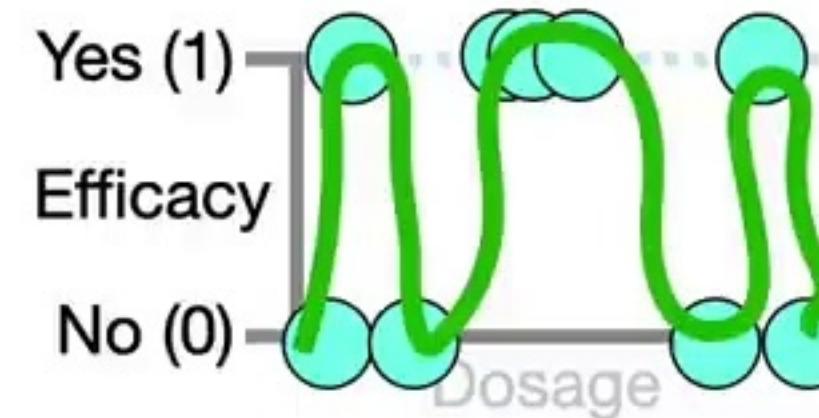


...just imagine what types of **green squiggles**  
we could fit with more **Hidden Layers** and  
**more Nodes** in each **Hidden Layer**.





In theory, **Neural Networks** can fit a **green squiggle** to just about any dataset, no matter how complicated!





And I think that's pretty cool!

